

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12386541
Kind Code	B1
Date of Patent	August 12, 2025
Inventor(s)	Slavchev; Slavcho Georgiev et al.

Data item and witness service partitioning in a distributed storage system

Abstract

Systems and methods are provided for using request rate data to determine when and how to split storage partitions in a distributed storage system. Data objects in a given partition are stored and managed using a set of computing resources, such as caches, a witness service, and persistent storage nodes. When the request rate for the partition is too great, then the partition may be split in order to distribute the request load over two sets of computing resources. Some requests that would be able to be handled by the cache are passed through to the persistent storage nodes, which can track the requests at a granular level. Based on this data, specific partition splitting decisions can be made.

Inventors: Slavchev; Slavcho Georgiev (Seattle, WA), Dun; Tianxing (Seattle, WA), Narendra; Vishwas (Evanston, IL), Okumusoglu; Ozan (Vancouver, CA), Koshelev; Alexander (Richmond, CA), Jung; Tyler Michael (Seattle, WA), Choudhary; Sameer (Seattle, WA), Shah; Nikhil (Redmond, WA), Wilkinson; Bryan T. (Seattle, WA)

Applicant: Amazon Technologies, Inc. (Seattle, WA)

Family ID: 1000005955369

Assignee: Amazon Technologies, Inc. (Seattle, WA)

Appl. No.: 17/449585

Filed: September 30, 2021

Publication Classification

Int. Cl.: G06F3/06 (20060101); G06F12/08 (20160101); G06F12/0813 (20160101)

U.S. Cl.:

CPC **G06F3/0644** (20130101); **G06F3/0604** (20130101); **G06F3/0631** (20130101);
G06F3/0659 (20130101); **G06F3/067** (20130101); **G06F12/0813** (20130101);

Field of Classification Search

CPC: G06F (3/0644); G06F (3/0604); G06F (3/0631); G06F (3/0659); G06F (3/067); G06F (12/0813)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
6405315	12/2001	Burns	713/193	G06F 16/10
7269648	12/2006	Krishnan et al.	N/A	N/A
8949535	12/2014	Hunter	N/A	N/A
11487442	12/2021	Bhat et al.	N/A	N/A
11537516	12/2021	Jung et al.	N/A	N/A
11640240	12/2022	Narendra et al.	N/A	N/A
2007/0177724	12/2006	Fuse et al.	N/A	N/A
2007/0263861	12/2006	Kiyomoto et al.	N/A	N/A
2012/0110113	12/2011	Lee et al.	N/A	N/A
2013/0054869	12/2012	Tolia et al.	N/A	N/A
2014/0019490	12/2013	Roy	707/798	G06F 16/20
2014/0067884	12/2013	Tata	N/A	N/A
2020/0387367	12/2019	Eklund et al.	N/A	N/A

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
106462545	12/2016	CN	N/A
WO 2017/152768	12/2016	WO	N/A

Primary Examiner: Khan; Masud K

Attorney, Agent or Firm: Knobbe, Martens, Olson & Bear, LLP

Background/Summary

CROSS-REFERENCE TO CONCURRENTLY-FILED APPLICATIONS

(1) This application is being filed concurrently with the following U.S. applications, each of which is incorporated by reference in its entirety and made part of this specification:

(2) TABLE-US-00001 U.S. application No. Title Filing Date 17/491004 MULTI-TIER CACHE FOR A Sep. 30, 2021 DISTRIBUTED STORAGE SYSTEM 17/449574 REQUEST ROUTING MANAGEMENT Sep. 30, 2021 FOR A DISTRIBUTED STORAGE SYSTEM 17/449597 WITNESS SERVICE FOR ENSURING Sep. 30, 2021 DATA CONSISTENCY IN A DISTRIBUTED STORAGE SYSTEM 17/449625 TRANSACTION MANAGEMENT FOR Sep. 30, 2021 MONOTONIC WRITE CONSISTENCY IN A DISTRIBUTED STORAGE SYSTEM 17/491029 DATA EVENT MANAGEMENT FOR Sep. 30, 2021 MONOTONIC READ CONSISTENCY IN A DISTRIBUTED STORAGE SYSTEM

BACKGROUND

(3) Computing devices can utilize communication networks to exchange data. Companies and organizations operate computer networks that interconnect a number of computing devices to support operations or to provide services to third parties. The computing devices can be located in a single geographic location or located in multiple, distinct geographic locations (e.g., interconnected via private or public communication networks). Specifically, data centers or data processing centers, herein generally referred to as a “data center,” may include a number of interconnected computing systems to provide computing resources to users of the data center. The data centers may be private data centers operated on behalf of an organization or public data centers operated on behalf of, or for the benefit of, the general public.

(4) Data centers may provide data storage services configured to store data submitted by client devices, and enable retrieval of that data over a network. A variety of types of data storage services can be provided, often varying according to their input/output (I/O) mechanisms. For example, database services may allow I/O based on a database query language, such as the Structured Query Language (SQL). Block storage services may allow I/O based on modification to one or more defined-length blocks, in a manner similar to how an operating system interacts with local storage, and may thus facilitate virtualized disk drives usable, for example, to store an operating system of a virtual machine. Object storage services may allow I/O at the level of individual objects or resources, such as individual files, which may vary in content and length. For example, an object storage service may provide an interface compliant with the Representational State Transfer (REST) architectural style, such as by allowing I/O based on calls designating input data and a hypertext transport protocol request method (e.g., GET, PUT, POST, DELETE, etc.) to be applied to that data. By transmitting a call designating input data and a request method, a client can thus retrieve the data from an object storage service, write the data to the object storage service as a new object, modify an existing object, etc. Some storage services provide caching services. A cache is a high-speed data store which stores a subset of data, typically transient in nature, so that future requests for that data are served up faster than is possible by accessing the data's primary storage location.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) Embodiments of various inventive features will now be described with reference to the following drawings. Throughout the drawings, reference numbers may be re-used to indicate correspondence between referenced elements. The drawings are provided to illustrate example embodiments described herein and are not intended to limit the scope of the disclosure.

(2) FIG. 1 is a block diagram depicting an illustrative environment in which a distributed storage system can operate using a caching index service and a witness service according to some embodiments.

(3) FIG. 2 is a block diagram of illustrative data flows and interactions of components of a distributed storage system to use the witness service to verify that stored data is update to date according to some embodiments.

(4) FIG. 3 depicts a flow diagram of an illustrative routine for using a witness service to verify whether data is up to date according to some embodiments.

(5) FIG. 4 is a diagram of illustrative cached data items and corresponding records of a witness service according to some embodiments.

(6) FIG. 5 is a block diagram showing illustrative data flows and interactions between a cache node and multiple nodes of the witness service according to some embodiments.

(7) FIG. 6 is a flow diagram of an illustrative routine for storing data such that sequencers

associated with data increase monotonically according to some embodiments.

(8) FIG. 7 is a block diagram of illustrative data flows and interactions between components of a caching index service to ensure monotonic increase of sequencers during storage of data according to some embodiments.

(9) FIG. 8 is a block diagram of illustrative data flows and interactions between components of a caching index service to ensure monotonic increase of sequencers during storage of data according to some embodiments.

(10) FIG. 9 is a flow diagram of an illustrative routine for persistent storage of data across a set of persistent storage nodes such that subsequent read requests return the most recent version of requested data regardless of the node to which the read requests are routed according to some embodiments.

(11) FIG. 10 is a block diagram of illustrative data flows and interactions between components of a distributed storage system to ensure that a future read request returns the most recent version of requested data regardless of which node in the distributed storage system is servicing the request according to some embodiments.

(12) FIG. 11 is a block diagram of illustrative data flows and interactions between components of a distributed storage system to ensure that a future read request returns the most recent version of requested data regardless of which node in the distributed storage system is servicing the request according to some embodiments.

(13) FIG. 12 is a block diagram of illustrative data flows and interactions between components of a distributed storage system including an event notification system according to some embodiments.

(14) FIG. 13A is a block diagram of illustrative data flows and interactions between components of multi-level cache architecture according to some embodiments.

(15) FIG. 13B is a block diagram of illustrative data flows and interactions between components of multi-level cache architecture according to some embodiments.

(16) FIG. 14 is a flow diagram of an illustrative routine for use of a multi-level cache architecture according to some embodiments.

(17) FIG. 15 is a block diagram of illustrative data flows and interactions between components of multi-level cache architecture according to some embodiments.

(18) FIG. 16 is block diagram of illustrative data flows and interactions between components of a distributed storage system to track request rates and make partition splitting decisions according to some embodiments.

(19) FIG. 17 is a flow diagram of an illustrative routine for splitting a storage partition based on request data according to some embodiments.

(20) FIG. 18 is a flow diagram of illustrative data flows and interactions between components of a distributed storage system to split a storage partition according to some embodiments.

(21) FIG. 19 is a block diagram of an illustrative multi-tier routing system in a distributed storage system according to some embodiments.

(22) FIG. 20 is a flow diagram of an illustrative routine for modifying multi-tier routing data and implementing routing changes in a distributed storage system according to some embodiments.

(23) FIG. 21 is a block diagram of illustrative data flows and interactions between systems and services associated with a partition according to some embodiments.

DETAILED DESCRIPTION

(24) Generally described, the present disclosure relates to managing storage nodes (e.g., cache nodes and/or persistent storage nodes) in a distributed storage system such that data consistency is ensured across the distributed storage system. Consistency is ensured in part through use of a consistency verification service, also referred to as a witness service. The witness service maintains data regarding the most recent change to—or otherwise most up-to-date version of—each data item of the distributed storage system that may be cached or otherwise stored. In high volume distributed storage systems, various optimizations and related features may be implemented to

reduce the load on the consistency witness service.

(25) Some distributed storage systems provide caches for high-speed retrieval of data that would otherwise require reading the data from comparatively slower persistent storage, performing additional network hops, etc. In a high-volume storage system that is distributed among a large number of separate physical devices (e.g., thousands or tens of thousands) and across sometimes great geographic distances (e.g., hundreds or thousands of miles), caches can significantly reduce the time for responding to a data request. Some distributed storage systems improve data availability by replicating data across any number of storage nodes that are necessarily implemented as caches. However, cached and other replicated data does not immediately reflect changes to the corresponding data stored in an authoritative source. When cached and other replicated data is out of sync with the corresponding data stored in the authoritative storage node, the data may be referred to as “stale.” Stale data in distributed systems can result in inconsistent data retrieval results whereby a data item is modified and, after the modification, different read requests may return different results depending upon whether the cache node and/or replica storage node involved in servicing the request has been updated with the modified data.

(26) Conventional distributed systems attempt to improve the “freshness” of caches and replicas through various techniques. For example, a distributed system may introduce latency to ensure that, after data is modified, the data replication process is able to complete before serving read requests for the data. To avoid introduction of latency, some conventional systems allow eventual consistency of the persistently-stored data (including data replicated across multiple storage nodes) and the corresponding cached and replicated data.

(27) Some aspects of the present disclosure provide for strong consistency of cached and persistently-stored data in a distributed storage system, while also minimizing or eliminating the latency that would otherwise be associated with such strong consistency. Generally described, strong consistency refers to a property of distributed storage systems whereby any request for a data item occurring after the data item has been changed in one location will result in access to the changed data item or a copy thereof, even if the subsequent request is made shortly after the data item has been changed on one storage node and/or the subsequent request is being serviced by another storage node. When a cache is used in addition to a distributed persistent storage system, the issue of cache staleness can interfere with providing strong consistency. For example, a subsequent request that is serviced from a cache node in which a “time to live” (“TTL”) for the requested item has not yet expired can result in providing the stale data item from the cache that does not reflect the change to the persistently-stored data item. The issue of cache staleness can be exacerbated in cases where caches are freshened from replicas that may themselves be stale.

(28) To address the issues noted above, among others, a consistency witness service may be implemented. The consistency witness service (also referred to as a “cache freshness witness service” or simply as a “witness service”) may be implemented as a streamlined service that provides data regarding the most recent instance of each data item that has been stored in a distributed storage system. The data regarding the most recent instance of storage of a given data item may be related to time or may otherwise be sequential and monotonically-increasing. In some embodiments, a sequence number (e.g., a transaction sequencer) may be used to indicate the relative time or sequence of a data change event or other storage operation for a given data item. Each cached version of the data item may be stored in connection with data regarding the transaction sequencer for the persistently-stored data item to which the cached data item corresponds. When a request for the data item is received by a cache node for servicing, the cache node may check with the witness service to see whether the cached data item stored at the cache node is the most up-to-date version of the requested data item based on the transaction sequencer. If the cached data item is stale (e.g., the transaction sequencer stored at the witness service is greater than the transaction sequencer associated with the cached data item at the cache node), then the cached data item may be invalidated and the cache node may obtain the latest version of the data

item from persistent storage. In this way, the benefits of a cache may be realized without the risk of serving stale data (e.g., based on an unexpired TTL).

(29) In distributed storage systems that store a large number of data items (hundreds of millions, billions, or more individual data items) each identified by a unique key, maintaining a record in the witness service for each individual data item may be difficult or impossible, particularly if it is desired that the records be maintained in volatile memory for fast access. To reduce the amount of memory and other computing resources required to store the records, the witness service may compress the relevant data using a hash table in which x data items are stored using y individual records, where $y < x$. In some embodiments, the degree of compression (e.g., the degree to which $y < x$) may be high, thus resulting in relatively frequent hash collisions. For example, two or more data items may both be mapped to the same record in the hash table once the keys of the respective data items are hashed. Thus, when one of the data items is updated and the new transaction sequencer is sent to the witness service, the new transaction sequencer will apply to each of the two or more data items mapped to the same record in the hash table. However, this effect of the hash table compression may be acceptable because as long as the sequence number mapped to any given data item via the hash table is equal to or greater than the latest sequence number that actually applies to the data item in persistent storage, then consistency will be maintained. One side effect will be that some cached data items will appear to be stale when they are not actually stale.

(30) In some embodiments, to address some drawbacks of the compression implemented by the witness service, the witness service may include multiple witness nodes to host multiple instances of the hash table. Each witness node may use different hash functions to maintain their respective instances of the hash table, and therefore the different instances of the hash table are less likely to experience the exact same hash collisions even though each instance may experience a similar number of hash collisions overall. For example, an odd number of hash table instances may be implemented (e.g., 3, 5, 7, or 9) and each may be maintained by a different witness node implemented on a different physical computing device. A majority of the hash table instances may be considered a quorum when performing consistency verification checks. When a request for the data item is received by a cache node for servicing, the cache node may check with each witness node of the witness service (or a subset thereof) to see whether the cached data item stored at the cache node is the most up-to-date version of the requested data item based on the transaction sequencer. If a majority of the witness nodes return a sequence number for the data item that is greater than the sequence number associated with the data item at the cache node, then the cached data item is likely to be stale, even if there are hash collisions for the data item at one or more of the witness nodes. The cached data item may be invalidated and the cache node may obtain the latest version of the data item from persistent storage. However, if a majority of the witness nodes return a transaction sequence number for the data item that matches the transaction sequence number associated with the data item at the cache node, then cached data item is guaranteed to be fresh.

(31) Additional aspects of the present disclosure relate to ensuring the witness service has the most up-to-date information regarding the most up-to-date versions of data items stored in the distributed storage system. When a request is made to store a data item in the storage system (either a new data item, or a modified version of a previously-stored data item), the data item may initially be stored in a transaction log from which persistent storage nodes are then updated. In some embodiments, once the data item is stored in the transaction log, an update is sent to the witness service indicating the transaction sequencer (e.g., transaction log sequence number) for the data item. Successful storage of the data item may then be acknowledged to the source of the storage request. Thus, the witness service can have an updated record of the latest transaction sequencer for the data item even before the data item is stored to persistent storage. When a request for the data item is received by a cache node, the cache node can query the witness service for the latest transaction sequencer and determine that its cached version of the data item is stale. Moreover, the cache node

is informed of the transaction sequencer that must be stored to a persistent storage node before the data item can be retrieved from the persistent storage node to service a request and freshen the cache. In this way, the cache node can be prevented from retrieving and serving an out-of-date data item from the persistent storage node when the persistent storage node has not yet been caught up on the corresponding transaction from the transaction log.

(32) In some embodiments, storage of a data item from the transaction log to a persistent data store may trigger an additional update to the witness service. For example, some updates of the witness service during initial storage of the data item to the transaction log may fail (e.g., due to network issues or other issues that may arise in distributed computing systems). In such cases, the distributed storage system may not acknowledge storage of the data item in the transaction log to the source of the data item storage request. This can prevent the source from expecting to access the new version of the data item on subsequent requests. When the transaction log transaction is later stored to the persistent storage nodes, the nodes may first update the witness service with the transaction sequencer, optionally acknowledge storage of the data item to the source of the request, and then store the data items in the respective storage nodes. Thus, now that the new version of the data item is available at the persistent storage nodes, it can be obtained by the cache nodes in response to future requests.

(33) In distributed storage systems, multiple persistent storage nodes may store a copy of the data item. Such persistent storage nodes (also referred to as replica nodes, or simply as storage nodes, to distinguish them from cache nodes) are typically used for load balancing, backup, and general availability and durability of the stored data. However, if each storage node communicates an updated transaction sequencer to the witness service for each data item that is stored to the distributed storage system, the witness service and/or network may be overwhelmed. For example, whereas a single update to the witness service may be made when the data item is saved to the transaction log, the addition of witness service updates when the data item is saved from the transaction log to the storage nodes will result in $1+n$ updates to the witness service, where n is the number of storage nodes to which the data item is stored. To reduce update traffic to the witness service, the first node that successfully updates the witness service may inform the other storage nodes of the successful update (e.g., using a gossip protocol). The other storage nodes may then omit updating the witness service when the data item is stored from the transaction log to the other storage nodes. Additionally, or alternatively, the storage nodes may batch witness service updates. For example, a storage node may wait until there are x data items ready to be stored from the transaction log to the storage node. The storage node may then send a single witness service update regarding all x data items prior to storing the data items to the storage node. Thus, traffic to the witness service can be reduced on the order of a factor of x , where x is the size of the batch. In some cases, the batch may additionally or alternatively be managed using a time-based criterion. For example, the storage node may wait y units of time between batch updates to the witness service. If only a single data item has been updated in the prior y units of time, the storage node may nevertheless send the update to the witness service and store the data item to persistent storage so as not to introduce a large degree of latency in storage of any given data item. In some cases, the batch may additionally or alternatively be managed using a request-based criterion. For example, the storage node may wait until a request for a data item in the batch has been received, and then send the update to the witness service and store the data item to persistent storage so as not to introduce a large degree of latency to the current requestor.

(34) Further aspects of the present disclosure relate to ensuring a monotonic increase in transaction sequencers for individual data items in a distributed storage system. As described above and in greater detail below, when a request is received to store a data item (e.g., a changed version of an existing data item), the data item is written to a transaction log from which various persistent storage nodes of the distributed storage system are updated. To distinguish and impart order on the transactions in the transaction log, each transaction may be associated with a transaction sequencer.

For example, the sequencer may be based on a system clock at the time a write request is received. However, write requests may be received by various nodes, each of which may write (or initiate writing of) the change to the transaction log. Small inconsistencies in the system clocks of the nodes writing the changes to the transaction log can result in some changes being written with a transaction sequencer that is less than or otherwise sequentially prior to the transaction sequencer of a change previously written to the transaction log by a different node. Thus, it may appear on the transaction log that the later-saved transaction occurred before the prior-saved transaction, which can have a negative effect on data consistency when the transactions each apply to the same data item. To prevent this scenario, the transaction log may be read to determine the sequencer of the most recent transaction associated with a data item prior to writing a new data record regarding the data item to the transaction log. If the transaction sequencer for the most-recent prior transaction is greater than the sequencer generated for the current transaction, then the current transaction may be rejected. The source of the request may choose to resubmit the request, during which time another transaction sequencer will be generated and evaluated against the transaction sequencer of the most-recent prior record associated with the data item.

(35) In some embodiments, the content of the data item in the current request may be evaluated against the content of the most recent transaction log record associated with the data item. If the value being requested to be written in the current request is the same value as written in the most recent transaction log record associated with the data item, then the source of the request would have no way of distinguishing between the different transactions because the value stored in persistent storage for the data item after both requests are processed would be the same regardless of the order in which the requests are processed. In this case, rather than reject the request, the request may be acknowledged and the greatest transaction sequencer may be used to store the data to persistent storage and update the witness service.

(36) In some embodiments, if the current request is a complex request that would cause a significant degree of latency and/or inefficient use of computing resources to reject and be resubmitted, the transaction sequencer assigned to the request may be altered. For example, the current request may be a multi-part request. If the transaction sequencer assigned to such a multi-part request happens to be prior to the transaction sequence of the most-recent prior transaction log record associated with the data item (e.g., due to differencing system clocks), then the transaction sequencer assigned to the current request may be adjusted to be greater than the transaction sequencer of the most-recent prior transaction log record. By adjusting the transaction sequencer in this manner, the monotonically-increasing nature of the transaction sequencers for the data item may be maintained, thus contributing to consistency of the data item. Moreover, the transaction is not rejected, thereby avoiding the latency and otherwise inefficient use of computing resources that would otherwise be the case in resubmitting the storage request.

(37) Some aspects of the present disclosure relate to a multi-tier caching architecture in which data items that are requested at a high frequency are cached in a first caching layer, also referred to as a “hot key cache,” and other data items that are requested relatively less frequently but which may still benefit from caching are cached at a second caching layer, also referred to as a “key cache.”

(38) In some embodiments, the key cache may be implemented using multiple cache nodes configured as a consistent hash ring. In the consistent hash ring, each node may be mapped to a set of keys (representing a corresponding set of data items) using a hash function such that a contiguous portion of the key space, corresponding to data items that are related and often requested close in time with each other, is distributed among multiple cache nodes rather than being concentrated on a single cache node. Thus, this mechanism serves to balance the caching load. In this configuration, requests for a given key will be mapped to the same cache node from request to request, in some cases without regard to the passage of time between requests for the given key.

(39) The hot key cache may be implemented using a different architecture and/or caching scheme

than the key cache described above. For example, the hot key cache may be implemented using a randomized load balancing mechanism such that requests for a given key that is associated with a high volume or rate of requests—a “hot key”—will be mapped to different hot key cache nodes from request-to-request in order to balance the load for that specific hot key.

(40) In some embodiments, if the data item cached at a hot key cache node for a requested hot key is invalid or there is otherwise a cache miss, the request can be passed to the key cache in the second caching layer. However, to avoid a significant rate of requests being routed to the consistent hash ring cache node that would be mapped by the hash function to the hot key, a cache miss request from the hot key cache layer can be routed to a pseudo-randomly selected node of the second caching layer. For example, an alternative routing mechanism may be used such that each hot key may be mapped to a subset of nodes in the second caching layer using a function (e.g., a hash function, a modulo function, etc.). When a request for a hot key is received by the second caching layer, the request may be identified as a hot key request. A routing subsystem of the second caching layer may implement the alternative routing mechanism to determine the subset of nodes mapped to the hot key, and randomly select a node of the subset (e.g., using a pseudo-random number generation algorithm). The selected node can then obtain the current version of the data item from persistent storage for caching and responding to client requests. In this way, a single node of the consistent hash ring is less likely to experience such a high demand that it is effectively unable to service any requests (similar to a denial-of-service attack) when cached data items for hot keys are invalidated.

(41) Additional aspects of the present disclosure relate to separating the set of data items stored in a distributed storage system into different subsets, or “partitions,” to further organize and manage the large number of data items and volume or requests. Each partition may be associated with a different set of computing resources, such as a transaction log onto which data is written for persistent storage, and a set of persistent storage nodes to service requests for data items in the partition and provide load balancing. Moreover, different instances of the witness service may be used for different subsets of partitions of the distributed storage system. For example, one instance of the witness service may be used to maintain sequence data and respond to verification requests regarding data items in a first subset of partitions, a second instance of the witness service may be implemented on separate hardware than the first instance so that the computing resources of the second instance can maintain sequence data and respond to verification requests regarding data items in a second subset of partitions, and so on.

(42) In some embodiments, the key space for keys of data items stored in the distributed storage system may be used to define and manage partitions such that a first subset of keys may be assigned to a first partition, a second subset of keys may be assigned to a second partition, and so on. The subsets may be defined with respect to the values of the keys themselves. For example, keys of data items may be defined using hexadecimal numbers of preset or varying length, and a partition may be associated with a particular key prefix. Thus, a partition may be defined as the set of all data items with key prefixes matching the particular key prefix of the partition, and not matching any other partition's key prefix that is longer (e.g., more hexadecimal symbols). Illustratively, one partition may be defined as the “f” partition, and another partition may be defined as the “foo” partition, without any partition with a key length longer than three hexadecimal symbols beginning with “foo.” In this illustrative example, all data items with keys beginning with “foo” would be part of the “foo” partition, while all other data items with keys beginning with “f” would be part of the “f” partition.

(43) After an initial set of partitions is defined, new data items may continue to be added to the distributed storage system, and different data items may be associated with increasing and/or decreasing rates of requests. In some embodiments, to address the potential growth in data items and changing degree of requests, a control system may monitor requests or otherwise obtain data regarding requests. Based on an analysis of request data, the control system may dynamically

generate new partitions by splitting existing partitions. For example, the control system may access request rate data stored by or otherwise associated with an instance of a witness service to determine which partition(s) are associated with a high rate of requests (e.g., as measured or estimated in transactions per second or “TPS”). If a particular partition has a request rate that satisfies a partition splitting criterion (e.g., TPS for the partition exceeds a maximum TPS threshold), then the control system may determine to split the partition into two or more partitions. However, because the witness service is optimized to respond to a high volume of verification requests associated with multiple partitions (e.g., 50,000 partitions or more) rather than providing detailed request metrics and analytics, the amount and nature of request metric data available from the witness service may not be adequate to determine how to optimally split a partition. To aid the control system in determining how to split a partition, the persistent storage nodes of the partition may maintain data regarding request rates for various subsets of the key space of the partition. To ensure that the persistent storage nodes have enough data about the entire key space, including keys that have been cached and for which few or no requests may have been serviced by the persistent storage nodes, the caching layers above the storage nodes may forward, to the persistent storage nodes, requests that are otherwise serviceable by the cache nodes. For example, a small percentage of requests (e.g., 1%, 5%, etc.) may be passed through to the persistent storage nodes even if there is no cache miss, and/or without first making a verification request to the witness service. Thus, the persistent storage nodes may have access to a statistically significant number of requests from which the total request volume can be derived, but without requiring the persistent storage nodes to actually service or even see the vast majority of all requests. The control system may use the request data, maintained by the persistent storage nodes, to identify a key prefix (or subset thereof) that can be split from the current partition in order to split the total request volume equally (or as equally as possible). Thus, the load on the witness service assigned to a particular partition can be effectively balanced. In addition, the risk of overloading and potentially browning out a persistent storage node assigned to keys of a particular partition can be managed.

(44) Further aspects of the present disclosure relate to management of partitions and, in particular, the routing data used by the distributed storage system to route requests to the proper caching layers, witness services, persistent storage nodes, and the like. In some embodiments, data items may be managed in a multi-tier configuration in which they are grouped into different partitions based on their keys (e.g., based on the longest key matching key prefix for which there is a partition), and partitions are grouped into different cells based on their key prefixes (e.g., based on ranges of key prefixes for the individual partitions). When a request is received, an inter-cellular routing layer may use routing data to route the request to a particular cell based on the key of the requested data item and the ranges assigned to the different cells. Inter-cellular routing data can be used to determine the identity of the proper cell for a given request, and may include additional information such as a host name and/or network address to which the request is to be routed. Once routed to a particular cell, an intra-cellular routing layer may use further routing data to route the request to a particular partition based on the prefix of the key and the prefixes assigned to the different partitions. Intra-cellular routing data can be used to determine the identity of the proper partition for a given request, and may include additional information such as a host name and/or network address to which the request is to be routed. Once routed to a particular partition, the request may be handled by a cache node and/or passed to a persistent storage node.

(45) In some embodiments, in addition to partition splitting for request load management as described above, cells may be split and/or partitions moved from cell-to-cell. The transaction log for affected partitions can be used to inform storage nodes, witness services, and the like regarding changes to routing data. For example, once a decision is made to move a partition or set of partitions to a different cell, and/or once a decision is made to split a partition, the routing data for the cells and/or partitions is to be updated. To ensure that requests are being processed using the correct routing data, special transactions may be written to the transaction logs of affected

partitions. These special barrier transactions, also referred to as “barrier” transactions, can serve to inform persistent storage nodes of the proper version of routing data that must be used to submit a request that will be handled by the persistent storage nodes. In addition, the barrier transactions can serve to invalidate the caches of data items in the affected partitions so that subsequent requests for data items, even if initially handled by the cache, are eventually routed to the prior persistent storage node to obtain the latest data from the current location. For example, when a barrier transaction is written to a transaction log of a particular partition, the witness service may be updated to reflect the transaction sequencer of the barrier transaction as the latest sequence number for all data items in the partition. Thus, subsequent verification requests for any data item in the partition will result in a cache invalidation and cause retrieval of the data item from its new location (e.g., the computing resources assigned to the new cell and/or partition of the data item).

(46) Various aspects of the disclosure will now be described with regard to certain examples and embodiments, which are intended to illustrate but not limit the disclosure. Although aspects of some embodiments described in the disclosure will focus, for the purpose of illustration, on particular examples of network configurations, data storage architectures, data access protocols, and processing operations, the examples are illustrative only and are not intended to be limiting. In some embodiments, the techniques described herein may be applied to additional or alternative network configurations, data storage architectures, data access protocols, processing operations, and the like

Overview of Example Computing Environment

(47) FIG. 1 is a block diagram of an illustrative computing environment **100** in which client devices **102** may store data objects to a distributed storage system **104** that includes frontend nodes **110**, an index service **112**, and an object storage service **114**.

(48) By way of illustration, various example client devices **102** (also referred to simply as “clients”) are shown in communication with the distributed storage system **104**, including a server, a desktop computer, and a laptop. In general, the client devices **102** can be any computing device such as a sever computer, desktop, laptop or tablet computer, personal computer, wearable computer, server, personal digital assistant (PDA), hybrid PDA/mobile phone, mobile phone, electronic book reader, set-top box, voice command device, camera, digital media player, and the like. In some embodiments, a client device **102** may be a host computing device executing one or more virtual computing systems within a data center, which may or may not be a same data center that includes one or more computing systems of the distributed storage system **104**.

(49) Generally described, the distributed storage system **104** can operate to enable clients to read, write, modify, and delete data objects or other data items, each of which represents a set of data associated with an identifier, such as an object key (also referred to simply as a key). Individual data objects may be interacted with as individual resources. For example, an object may represent a single file submitted by a client device **102** (though the distributed storage system **104** may or may not store such an object as a single file). This object-level interaction can be contrasted with other types of storage services, such as block-based storage services providing data manipulation at the level of individual blocks or database storage services providing data manipulation at the level of tables (or parts thereof) or the like.

(50) The distributed storage system **104** illustratively includes one or more frontend nodes **110**, which provide an interface (a command-line interface (CLI), application programming interface (API), or other programmatic interface) through which client devices **102** can interface with the system **104**. For example, a client device **102** may interact with a frontend node **110** to create a collection of data objects on the system **104** (e.g., a “bucket” of objects) and to configure permissions for that collection. Client devices **102** may thereafter create, read, update, or delete objects within the collection based on the interfaces of the frontend nodes. In one embodiment, the frontend node provides a REST-compliant HTTP interface supporting a variety of request methods, each of which corresponds to a requested I/O operation on the system **104**. By way of non-limiting

example, request methods may include: a GET operation requesting retrieval of an object stored on the system **104** by reference to an identifier of the object; a PUT operation requesting storage of an object to be stored on the system **104**, including an identifier of the object and input data to be stored as the object; a DELETE operation requesting deletion of an object stored on the system **104** by reference to an identifier of the object; and a LIST operation requesting listing of objects within an object collection stored on the system **104** by reference to an identifier of the collection.

(51) During general operation, frontend nodes **110** may be configured to obtain a call to a request method, and apply that request method to input data for the method. For example, a frontend node **110** can respond to a request to PUT input data into the system **104** as an object by storing that input data as the object on the system **104**. Objects may be stored, for example, in one or more of the object storage service's object storage nodes **140**, which correspond to any persistent or substantially persistent storage (including hard disk drives (HDDs), solid state drives (SSDs), network accessible storage (NAS), storage area networks (SANs), non-volatile random access memory (NVRAM), or any of a variety of storage devices known in the art). As a further example, the frontend node **110** can respond to a request to GET an object from the system **104** by retrieving the object from the object storage service **114** (the object representing input data to the GET resource request), and returning the object to a requesting client device **102**.

(52) The index service **112** can maintain information about the specific storage location (e.g., network address) at which each data object is stored and accessible in the object storage service **114**. For example, data objects may be exposed to clients and interacted with according to a unique identifier for the object, such as a key. In some embodiments, the key may be an alphanumeric string for ease of reference. Use of the key by clients outside of the distributed storage system **104** can shield the clients from the complexity of storing objects in, and accessing objects stored in, the distributed storage system **104**. This shielding and simplification for clients can be particularly beneficial when the object storage system is a scalable storage system.

(53) The index service **112** may include a set of persistent storage nodes **120** to store the key map data that represents mappings of data object keys to storage locations at which the data objects are stored in the object storage service **114**. The persistent storage nodes **120** may include any number of replica nodes to provide durability and availability of the key map data. When the frontend node **110** receives a request to store, retrieve, or otherwise interact with a data object stored in the object storage server **114**, the index service **112** may be employed to determine the specific storage location of the data object identified by key.

(54) In some embodiments, the object storage service **114** may provide an immutable data store. For example, when a data object is stored at a particular location in the object storage service, only that data object, with the data value(s) assigned to the data object at the time of storage, will be stored at that particular location. When a change is made to the data object, the changed data object is stored at a new storage location. Thus, the mapping of the object key to storage location may change each time the data object itself is changed. The index service **112** may therefore serve an important role in maintaining the mapping of data object keys to current storage locations without which it may be difficult or impossible to retrieve the latest version of a data object using the data object key.

(55) To provide faster access to frequently-used key map data, the index service **112** may include a key cache **122** to cache key map records stored in the persistent storage nodes **120**. Whenever a data object is changed, the key map data stored in the cache becomes stale because the changed version of the data object is stored at a different location in the object storage service **114**. Thus, verifying the freshness of the key cache **122** may become necessary to avoid accessing an old version of the data object after it has been changed.

(56) To provide freshness verification, a witness service **116** may be implemented. The witness service **116** may be notified of any change to the key map data stored in the persistent storage nodes **120**. When the key cache **122** is accessed to obtain key map data for a particular key, the key

cache **122** may request that the witness service **116** provide data regarding the latest version of the key map data stored in the persistent data store for the key. For example, a transaction sequencer assigned to the key map data when the key map data was written to a transaction log may be used to determine whether a particular cached item of key map data is stale. Cached key map data items maintained at the key cache **122** may be augmented or otherwise associated with additional information, such as the transaction sequencer of the key map data item when it was cached. If the transaction sequencer of the cached key map data item is older (e.g., has a smaller value or otherwise prior value in a sequence) than the transaction sequencer obtained from the witness service **116**, then the cached key map data item is stale and the cache node may obtain the latest version of the key map data item from the persistent storage nodes **120**.

(57) In some cases, certain data objects may be accessed at an extremely high rate. To provide even more responsive service to requests for such data objects, a hot key cache **150** may be implemented at the front end. Like the key cache **122**, the hot key cache **150** may request that the witness service **116** provide data regarding the latest version of the key map data stored in the persistent data store for a requested hot key. Any request associated with a hot key for which the hot key cache **150** has stale data, or for a key that is not maintained in the hot key cache **150**, may be passed through to the index service **112**.

(58) Although the index service **112** and witness service **116** are shown in FIG. **1** and described herein with reference to examples of caching and verifying the freshness of key map data that maps data object keys to data object storage locations in an object storage service **114**, the examples are illustrative only and are not intended to be limiting, required, or exhaustive. In some embodiments, the witness service **116**, data caches, and techniques for ensuring consistency of stored data may be applied to various other examples and use cases which may or may not involve location data for objects stored at a separate object storage service **114**.

(59) The client devices **102**, distributed storage system **104**, and/or individual devices thereof may communicate via a network **180**, which may include any wired network, wireless network, or combination thereof. For example, the network may be or include a personal area network, local area network, wide area network, over-the-air broadcast network (e.g., for radio or television), cable network, satellite network, cellular telephone network, or combination thereof. As a further example, the network may be a publicly accessible network of linked networks, possibly operated by various distinct parties, such as the Internet. In some embodiments, the network may be a private or semi-private network, such as a corporate or university intranet. The network may include one or more wireless networks, such as a Global System for Mobile Communications (GSM) network, a Code Division Multiple Access (CDMA) network, a Long Term Evolution (LTE) network, or any other type of wireless network. The network can use protocols and components for communicating via the Internet or any of the other aforementioned types of networks. For example, the protocols used by the network may include Hypertext Transfer Protocol (HTTP), HTTP Secure (HTTPS), Message Queue Telemetry Transport (MQTT), Constrained Application Protocol (CoAP), and the like. Protocols and components for communicating via the Internet or any of the other aforementioned types of communication networks are well known to those skilled in the art and, thus, are not described in more detail herein.

(60) The distributed storage system **104** is depicted in FIG. **1** as operating in a distributed computing environment including several computer systems that are interconnected using one or more computer networks. The distributed storage system **104** could also operate within a computing environment having a fewer or greater number of devices than are illustrated in FIG. **1**. Thus, the depiction of the distributed storage system **104** in FIG. **1** should be taken as illustrative and not limiting to the present disclosure. For example, the distributed storage system **104** or various components thereof could implement various Web services components, hosted or “cloud” computing environments, or peer to peer network configurations to implement at least a portion of the processes described herein. In some instances, the distributed storage system **104** and clients

102 may be part of a same cloud computing service. Further, the distributed storage system **104** and clients **102** may be implemented directly in hardware, or as a combination of software executed by hardware devices and may, for instance, include one or more physical or virtual servers implemented on physical computer hardware configured to execute computer executable instructions for performing various features that will be described herein. The one or more servers may be geographically dispersed or geographically co-located, for instance, in one or more data centers. In some instances, the one or more servers may operate as part of a system of rapidly provisioned and released computing resources, often referred to as a “cloud computing environment.”

(61) In some embodiments, any of the components within the distributed storage system **104** can communicate with other components of the distributed storage system **104** via a network. In other embodiments, not all components of the distributed storage system **104** are capable of communicating with all other components of the distributed storage system. In addition, while some functionalities are generally described herein with reference to an individual component of the distributed storage system **104**, other components or a combination of components may additionally or alternatively implement such functionalities. Thus, the specific configuration of elements within FIG. 1 is intended to be illustrative only, and is not intended to be limiting, required, or exhaustive.

(62) Consistency Verification

(63) FIG. 2 illustrates example data flows and interactions between components of a distributed storage system **104** when using a witness service **116** to ensure that the most recent version of a data item is provided in response to a request.

(64) As shown, a request may be initially received at [A] by a cache node **202** (e.g., a node of the hot key cache **150** or key cache **122**). Illustratively, the request may be a request for a data item associated with a unique key. Although the cache node **202** may have a cached data item associated with the unique key, the cache node **202** may obtain verification data from the witness service **116** at [B] prior to providing a response. The witness service **116** may provide, to the cache node **202**, verification data that includes a sequence number associated with the key. The cache node **202** may determine at [C] that the sequence number received from the witness service **116** is greater than the sequence number associated with the data item cached at the cache node **202**. In response to this determination, the cache node **202** can request an up-to-date version of the data item from a persistent storage node **120** at [D].

(65) The persistent storage node **120**, or another component of the index service **112** such as a coordinator subsystem **124**, may perform a process similar to that performed by the cache node **202** to ensure that an up-to-date version of the requested data item is provided. For example, the persistent storage node **120** may obtain verification data from the witness service **116** at [E] prior to providing a response. The witness service **116** may provide, to the persistent storage node **120**, verification data that includes a sequence number associated with the key. The persistent storage node **120** may determine at [F] that the sequence number received from the witness service **116** is greater than the sequence number associated with the data item currently stored at the persistent storage node **120**. In response to this determination, the persistent storage node **120** can wait [G] to provide a response until persistent storage of data from the transaction log **200** to the persistent storage node **120** has reached at least the sequence number received from the witness service **116** at [H]. When that happens, the persistent storage node **120** can provide a response to the cache node **202** at [I] including the up-to-date version of data associated with the key.

(66) The cache node **202** can receive and cache the up-to-date version of the data at [J], and then provide a response to the request at [K].

(67) FIG. 3 is a flow diagram of an illustrative routine **300** that may be executed by a cache node **202** to communicate with the witness service **116** and determine whether cached data (e.g., key map data) is stale or fresh. In some embodiments, the routine **300** may be implemented by a cache

node of the hot key cache **150** or key cache **122**. Aspects of the routine **300** will be described with further reference to the illustrative cache objects and corresponding records of the verification shown in FIG. **4**, and the illustrative data flows and interactions between the cache node and multiple nodes of the witness service shown in FIG. **5**.

(68) The routine **300** may begin at block **302** in response to an event, such as when the cache node **202** begins operation. When the routine **300** is initiated, a set of executable program instructions stored on one or more non-transitory computer-readable media (e.g., hard drive, flash memory, removable media, etc.) may be loaded into memory (e.g., random access memory or “RAM”) of a computing device and executed. In some embodiments, the routine **300** or portions thereof may be implemented on multiple processors, serially or in parallel.

(69) At block **304**, the cache node **202** may receive a request for data stored in the index service **112**. In some embodiments, the request may be a request for data associated with a key. For example, the request may be a GET request, and may include the key. The request may be initiated by a client **102** or some other device or system.

(70) In one specific non-limiting embodiment, the request may be received by a frontend node **110**. The frontend node **110** may determine whether the key is a “hot key” (e.g., a key that is the subject of a large volume of requests, a key that is assigned to a hot key cache, etc.). If the key is a hot key, the request may be sent to the hot key cache **150** where it may be assigned to a random cache node by a load balancer. In this case, the cache node executing the current instance of the routine **300** is a cache node of the hot key cache **150**. Otherwise, if the key that is the subject of the request is not a hot key, then the request may be sent to the key cache **122** where it may be assigned to a particular node based on a hash of the key (e.g., if the key cache **122** is implemented as a consistent hash ring). In this case, the cache node executing the current instance of the routine **300** is a cache node of the key cache **122**.

(71) At block **306**, the cache node **202** may obtain cached data associated with the key in the request. A typical cache node **202** may maintain cached data for many keys. In some embodiments, the cached data may be cached key map data. An item of cached key map data may include a key and value that corresponds to the key. A typical request that is serviced by the cache node (e.g., a GET request) may be a request for the value that corresponds to the key. The cache node may maintain the cached key map data in memory for fast access and response to such requests.

(72) Although some examples herein are described with respect to a key being mapped to a value representing the location of another data object, the examples are illustrative only and are not intended to be limiting, required, or exhaustive. In some embodiments the value may be any value or other data that may be associated with a key.

(73) To facilitate various cache freshness verification features, cached data maintained by the cache node **202** may be augmented with additional data (e.g., a cached data item may include or otherwise be associated with data in addition to the key and corresponding value). The cached data may include or otherwise be associated with sequence data, such as a sequence number. For example, a sequence number may represent a relative time of an update event associated with key, or an ordinal position relative other update events associated with the key. In some embodiments, the sequence number may be a transaction log sequence number (also referred to as a “transaction sequencer”) used when recording update events associated with the key to a transaction log **200** for persistent storage.

(74) As described in greater detail below, the sequence number may be used to determine whether an update event associated with key has occurred after the update event reflected in the cached data maintained by the cache node **202**, in which case the cached data maintained by the cache node **202** is stale. In some embodiments, the cached data may be augmented with routing information to make a remote procedure call in order to verify the cached information with the witness service **116** (e.g., pre-resolved network addresses of the specific witness node(s) **160** to be queried).

(75) FIG. **4** illustrates an example cached data set **400** of cached data items **402** that may be

maintained by a cache node **202**. Each cached data item **402** may include one or more fields. As shown, cached data items **402** may include a unique key **410**, a value **412**, a sequence number **414**, and verification routing information **416**.

(76) Although many examples described herein use a sequence number such as a transaction sequencer to perform cache freshness verification queries and to implement other features for ensuring data consistency, the examples are illustrative only and are not intended to be limiting, required, or exhaustive. In some embodiments other forms of sequence data may be used to maintain data consistency, determine sequential update events associated with persistently-stored data and cached data, etc.

(77) At block **308** of routine **300**, the cache node **202** may obtain data from the witness service **116** for use in determining the freshness of the cached data item **402** for the key **410** that is the subject of the current request. Such data may be referred to as “freshness data.” In some embodiments, the cache node **202** may access verification routing information **416** associated with the key **410** for the cached data item **402**. The verification routing information **416** may identify a particular witness node **160** (or multiple witness nodes **160**) to query. The query may be a query for the sequence number that is associated with the latest update event associated with the key **410**.

(78) FIG. **4** illustrates an example verification data set **450** of verification data items **452** that may be maintained by a witness node **160**. Each verification data item **452** may include one or more fields. As shown, verification data items **452** may include a unique index **460** and a sequence number **462**. The index **460** may correspond to a key (e.g., the key **410** of a cached data item **412**). The value of the sequence number **462** field for a given verification data item **452** may be the sequence number assigned to the most recent update event associated with the index **460**.

(79) In some embodiments, the verification data set **450** may be implemented as a hash table. The witness node **160** may a set one or more hash functions **470** to hash keys and generate hash table indices. Due to the nature of hashing, use of a hash table introduces a degree of compression through hash collisions in which two or more keys **410** are mapped to the same index **460**. The degree of compression may depend upon the particular hash function **470** being used and the number of individual verification data items **452** in the hash table.

(80) Verification data set **450** compression may be desirable in order to keep the size of the hash table manageable (e.g., able to be maintained in random access memory of the witness node **160** for fast retrieval and low latency responses to freshness verification queries). Moreover, compression that results in hash collisions may not have negative effect from a data consistency standpoint, as long as the sequence numbers **462** associated with each data event are monotonically increasing. For example, if two keys **410** are mapped to the same index **460** by the hash function **470**, then they will both be associated with the same verification data item **452** and queries for both keys will return the same sequence number **462**. When an update event occurs for one key **410**, the sequence number associated with that key will be used to update the data item **452** as described in greater detail below, and will become the sequence number that is associated with both keys **410**. However, even though no recent update event may be associated with the other key **410**, use of a later sequence number **462** will not impact data consistency because a cache invalidation will result in retrieval of the current value, which would be the same value prior to the update event for the other key.

(81) At block **310** of routine **300**, the cache node **202** may receive verification results from the witness service **116**. The results may be or include the log sequence number **462** associated with the key **410** (e.g., via the index **460** generated using hash function **470**).

(82) In some embodiments, the cache node **202** may request and receive verification results from multiple witness nodes **160**. The witness service **116** may include any number of witness nodes **160** for redundancy and/or quorum-based operations. To reduce the possibility of hash collisions involving the same keys occurring at each witness node **160**, the witness nodes **160** may use different hash functions **470** to map keys to hash table indices.

(83) In some embodiments, different subsets of witness nodes **160** may manage different subsets of the set of all keys available to be cached. For example, the set of all keys available to be cached may be too large to be maintained entirely in available memory of individual witness nodes **160**. Additionally, or alternatively, the volume of cache freshness verification requests may be too great for individual witness nodes **160** to respond to such requests for all keys available to be cached. To address these issues, the set of all keys available to be cached may be partitioned into two or more partitions as described in greater detail below.

(84) FIG. 5 illustrates a cache node **202** (e.g., the cache node **202** executing the current instance of routine **300**) obtaining verification results from multiple witness nodes **160A**, **160B**, and **160C**. The cache node **202** may use routing information **416** associated with the key **410** to make cache freshness verification requests to the appropriate witness node or subset of witness nodes. Each witness node **160A**, **160B**, **160C** may include its own copy or version of verification data set **450A**, **450B**, **460C**, respectively. When the verification data sets are implemented as hash tables, each witness node **160A**, **160B**, **160C** may use a different set of one or more hash functions **470A**, **470B**, **470C**, respectively.

(85) At block **312**, the cache node **202** may determine the sequence number for the latest update event associated with the key **310**. If a single sequence number is requested and obtained for the key, then that sequence number may be used for subsequent portions of the routine **300**. If multiple sequence numbers are requested and obtained for the key, the cache node **202** may select one for use in subsequent portions of the routine **300**. The specific method used to select the sequence number from the multiple sequence numbers that have been received may depend on the reason the cache node is configured to obtain multiple verification results.

(86) The cache node **202** may obtain verification results from multiple witness nodes for any of a variety of different reasons, such as for use in quorum operations.

(87) In some embodiments, the cache node **202** may request verification results from multiple nodes, and use the results that are obtained first to proceed with the rest of the routine **300** and minimize latency.

(88) In some embodiments, the cache node **202** may request verification results from multiple witness nodes, and use the greatest or latest sequence number received in verification results to proceed with the result of the routine **300**. For example, when an update event occurs for a particular key, storage of the updated data may only be completed if at least a majority of the witness nodes are also updated with the sequence number associated with the update event (e.g., if there are 9 witness nodes for the current key, then at least 5 witness nodes must be updated). By obtaining verification results from at least a majority of the witness nodes (e.g., a quorum), the cache node **202** can be assured that it will receive verification results from at least one of the nodes updated in connection with the update event.

(89) In some embodiments, an asymmetric quorum may be used in which the quantity or proportion of witness nodes queried for a given request may differ depending upon the type of request received. The different quantities/proportions may be selected so as to guarantee an overlap of one or more witness nodes. For example, the asymmetric quorum may be an asymmetric read/write quorum. When a request is received to write data, then sequence numbers may be obtained from a super majority of witness nodes (e.g., eight out of nine witness nodes assigned to the partition to which the requested key belongs) and/or the write may not be acknowledged until a super majority of witness nodes are updated with the new sequence number that corresponds to the write transaction. When a request is received to read data, then sequence numbers may be obtained from less than a majority of witness nodes, where the particular quantity of witness nodes in the sub-majority subset is selected to guarantee overlap with the write quorum (e.g., two out of nine witness nodes).

(90) In some embodiments, the cache node **202** may request verification results from multiple witness nodes, and use the smallest or oldest sequence number that is nevertheless equal to or

greater/later than a majority of the witness nodes. This process may be used to mitigate the effect of hash collisions at the witness nodes, such as in case where high degrees of compression are used. By using the sequence number that is equal to or greater/later than a majority of all sequence numbers maintained for the key, the cache node **202** can be assured that it is using a sequence number that equal to or greater/later than the actual sequence number associated with the last update event for the key, even if there are one or more other sequence numbers in the verification results that are even greater/later than the selected sequence number. Those sequence numbers would be the result of hash collisions with other keys, and can be ignored to reduce the number of needless cache invalidation events.

(91) At decision block **314**, the cache node **202** may determine whether the sequence number of the cached data is less than the sequence number received from the witness service **116** (or less than the sequence number selected from the numbers received from multiple witness nodes, as described above). If so, then the cached data item may not include the most up-to-date value for the requested key and the routine **300** may proceed to block **316**. Otherwise, if the sequence number of the cached data is not less than the sequence number received from the witness service **116** (or selected as described above), then the routine **300** may proceed to block **320**.

(92) At block **316**, the cache node **202** may obtain updated data associated with the key from a persistent storage node.

(93) At block **318**, the cache node **202** may generate updated cached data associated with the key. In some embodiments, the cache node **202** may update one or more fields of the cached data item **402** associated with the key **410**. For example, the cache node **202** may update the value **412** with the new value obtained from the persistent storage node **120**. The cache node **202** may also update the sequence number **414** with a new sequence number. The new sequence number may be the number determined above based on verification results from the witness service **116**, or it may be received from the persistent storage node **120** or some other service in connection with obtaining the updated value for the key.

(94) At block **320**, the cache node **202** may provide results back to the requestor. Illustratively, the results may be or include the value **412** currently assigned to the key **410** in the cached data item **402** associated with the key **410**. In some embodiments, when a new value is obtained from a persistent storage node **120**, block **320** may be performed prior to block **318**. Thus, the value obtained from the persistent storage node may be provided back to the request prior to the cache node generating updated cached data associated with the key.

(95) At block **322**, the routine **300** may end.

(96) Although routine **300** is described with respect to execution by a cache **202**, the description is illustrative only and is not intended to be limiting, required, or exhaustive. In some embodiment, the routine **300** or portions thereof may be performed by any storage node to verify whether a stored data item is up to date, as described above with respect to FIG. 2.

(97) Ensuring Monotonic Writes

(98) FIG. 6 is a flow diagram of an illustrative routine **600** that may be executed by a distributed storage service **104** to ensure that data update events are applied or otherwise persisted in the proper sequence, even when originating from different nodes of the distributed storage service **104**. Advantageously, ensuring monotonic writes using the routine **600** helps to ensure data consistency among caches of the persistent data, such that stale data is less likely to be accessed.

(99) Although the description that follows is with respect to an example implementation in which the routine **600** is performed by the distributed storage system **104** shown in FIG. 1, including an index service **112**, the example is for purposes of illustration only and is not intended to be limiting. In some embodiments, the routine **600** may be used in other environments and to manage other types of data. Aspects of the routine **600** will be described with further reference to the illustrative data flows and interactions shown in FIGS. 7 and 8.

(100) The routine **600** may begin at block **602** in response to an event, such as when the index

service **112** begins operation. When the routine **600** is initiated, a set of executable program instructions stored on one or more non-transitory computer-readable media (e.g., hard drive, flash memory, removable media, etc.) may be loaded into memory (e.g., random access memory or “RAM”) of a computing device and executed. In some embodiments, the routine **600** or portions thereof may be implemented on multiple processors, serially or in parallel.

(101) At block **604**, the distributed storage system **104** may receive a write request to store data. In some embodiments, the request may be a request to change or otherwise store data associated with a particular key. For example, the request may be a PUT request, and may include the key and the data to be written. The request may be initiated by a client **102** or some other device or system, and may be received by a frontend node **110**.

(102) At block **606**, the frontend node **110** or some other component of the distributed storage system **104** may assign a sequence number to the request. In some embodiments, the sequence number may be based on a system clock of the frontend node **110**. For example, the frontend node **110** may include a system clock that provides a monotonically-increasing value from a particular point in time, such as the number of nanoseconds that have passed since a particular point in time. The frontend node may use this output from the system clock as the sequence number, or may otherwise derive the sequence number from the output of the system clock (e.g., by subtracting a particular value to reduce the magnitude of the sequence number and consume less memory). Sequence numbers based on the system clock can thus be compared to find out which sequence number corresponds to a later point in time based on which sequence number is greater.

(103) At decision block **608**, the frontend node **110** can determine whether to pass the write request to another subsystem, or to perform additional operations first. The decision may be made based on type of write requests, where there are multiple different write request types. In some embodiments, a write request type is based on the complexity of the write request (e.g., how computationally expensive the write request is), which can have an effect on costs of rejecting the write request. For example, if the write request is a complex write request, such as a multi-phase request, then the cost of rejecting the request in terms of time and computing resources consumed by the client **102** to re-submit the request or otherwise address the rejection may be relatively high compared to rejecting non-complex write requests, such as single part requests. The frontend node **110** may therefore perform additional processing in order to avoid a rejection of such complex requests, and the routine **600** may proceed to block **610**. Other types of requests, such as single-phase requests, may be less complex and/or less computationally expensive in terms of time and computing resources consumed by the client **102** to re-submit the request or otherwise address a rejection. Moreover, performing additional processing up front for each of such requests may introduce additional latency and/or other computational expenses that are unexpected or undesirable. For these types of requests, or if the frontend node **110** otherwise determines to pass the write request to another subsystem for the next operations in the request path, the routine **600** may proceed to block **620**.

(104) At decision block **610**, the frontend node **110** can determine whether the current sequence number generated above is greater than the last sequence number for the key. The frontend node **110** can obtain the latest sequence number for the key from the persistent data store and/or the transaction log. If current sequence number is not greater than the last sequence number for the key, then the routine **600** may proceed to block **612**. Otherwise, if the current sequence is greater than last sequence number for the key, then the routine **600** may proceed to block **614**.

(105) In some cases, the current sequence number may not be greater than the last sequence number for the key due to a difference in system clocks among different frontends **110**. Due to the nature of the distributed system, different write requests for a given key may be handled by different frontends **110**. Although the system clocks among the frontends **110** may be substantially synchronized, there may nevertheless be a slight skew in the time kept by different system clocks of different frontend nodes **110**. Thus, if two write requests are submitted for the same key in a very

short period of time (e.g., write requests from two different clients **102**, or from two different threads of a single client **102**), then even a slight clock skew may result in a later-arriving request being assigned a sequence number that is prior to the sequence number assigned to a prior-arriving request.

(106) At block **612**, in response to determining that the current sequence number is not greater than the last sequence number, the frontend node **110** may modify the current sequence number such that it is greater than the last sequence number. Thus, the write request may proceed using the modified sequence number, and monotonic increase of the sequence number may be maintained.

(107) At block **614**, the frontend node **110** may pass the write request with the modified sequence data to the coordinator subsystem **124** for further processing. At block **616**, the write request may be made to the transaction log to persist the data associated with the key and the current sequence number. At block **616**, the coordinator subsystem **124** can send a message to the witness service **116** regarding the new sequence number associated with the key. At block **618**, the coordinator subsystem **124** can update the witness service **116** with the new sequence number. At block **620**, the current instance of the routine **600** may end with the coordinator system **124** acknowledging successful completion of the write request to the source of the request.

(108) FIG. 7 illustrates example data flows and interactions between components of the distributed storage system **104** of FIG. 1 when handling a complex write request for a key and determining that a current sequence number is not greater than the last sequence number associated with the key.

(109) As shown, a write request is received by a frontend node **110** at [A]. The frontend node **110** uses a sequencer generated **700** to assign a sequence number to the write request at [B]. At [C], the frontend node **110** determines that the write request is a complex request, and the frontend node **110** is to perform additional processing prior to sending the request to another subsystem for the next operation(s) in the request path.

(110) At [D], the frontend node **110** can obtain the last sequence number associated with the key. For example, the frontend node **110** may obtain the last sequence number from a subsystem that manages the transaction log **200**, such as a coordinator subsystem **124**. The coordinator subsystem **124** may inspect the transaction log **200**, identify the last sequence number associated with the key, and provide the last sequence number to the frontend node **110**.

(111) At [E], the frontend node **110** can determine that the current sequence number, assigned at [B], is not greater than the latest sequence number associated with the key and received from the coordinator subsystem **124**. At [F], the frontend node **110** can modify the current sequence number, and submit the write request with the modified sequence number to the coordinator subsystem **124** at [G]. The coordinator subsystem **124** can write the key, data, and modified sequence number to the transaction log **200** at [H], from where it will be persisted to one or more persistent storage nodes **120** at [I].

(112) Returning to FIG. 6, an alternate request path will be described. At block **622**, the frontend node **110** passes the write request to another subsystem for performing the next operation(s) in the request path (e.g., based on the write request being a non-complex write request).

(113) At decision block **624**, the coordinator subsystem **124** can determine whether the current sequence number is greater than the last sequence number associated with the key. If so, the routine **600** can proceed to block **616**, where the write request may be made to the transaction log to persist the data associated with the key and the current transaction number. Otherwise, if the current sequence number is not greater than the last sequence number associated with the key, then the routine **600** can proceed to decision block **626**.

(114) At decision block **626**, the coordinator subsystem **124** can determine whether the value associated with the last sequence number is the same as the value of the current write request. If so, then the routine **600** may in some embodiments proceed to block **630**, where data is written to the

transaction log **200**, and then to block **632**, where the current instance of the routine **600** terminates with the coordinator subsystem **124** acknowledging success of the write request without sending an update to the witness service **116**. In some embodiments, data may be written to the transaction log **200** in the same manner as other write requests described herein, even though the current sequence number is less than the last sequence number. This may be done so that other services that use the transaction log **200**, such as the event notification service **1200** described in greater detail below, may obtain information regarding the acknowledged write operation. In some embodiments, additional data may be included, such as a tag or metadata indicating that the data is only written to the transaction log without any corresponding write to a persistent storage node **120** and/or without any update to the witness service.

(115) Otherwise, if the value associated with the last sequence number is different than the value associated with the current sequence number, then the routine **600** may terminate at block **632** where the write request is rejected.

(116) FIG. **8** illustrates example data flows and interactions between components of the distributed storage system **104** of FIG. **1** when handling a non-complex write request for a key and determining that a current sequence number is not greater than the last sequence number associated with the key.

(117) As shown, a write request is received by a frontend node **110** at [A]. The frontend node **110** assigns a sequence number to the write request at [B]. At [C], the frontend node **110** determines that the write request is a non-complex request, and the frontend node **110** sends the request to another subsystem (e.g., the coordinator subsystem **124**) for the next operations in the request path at [D]

(118) At [E], the coordinator subsystem **124** can obtain the last sequence number and value associated with the key. For example, the frontend node **110** may obtain the last sequence number and value from the transaction log **200**. At [F], the coordinator subsystem **124** can determine that the current sequence number, assigned at [B], is not greater than the latest sequence number associated with the key and received from the transaction log **200**.

(119) At [G], the coordinator subsystem **124** can determine whether the value associated with the current request is different than the value associated with the last sequence number. If the values are the same, then a subsequent data accessor will not be able to determine whether the value is from the current write request, or from a prior write request associated with the key. Thus, successful completion of the write request can be acknowledged without actually completing the write request (e.g., without writing the current sequence number to the transaction log, thereby violating the goal of monotonically increasing sequence numbers associated with the key). By acknowledging successful completion of the write request, the source of the write request will not be triggered to retry the write request, which would not change the value assigned to the key.

(120) Alternatively, if the values are different, then the coordinator subsystem **124** can reject the write request. Because the write request was a non-complex write request, the cost for the source of the write request to retry the write request is not as high as a complex write request.

(121) Ensuring Monotonic Reads

(122) FIG. **9** is a flow diagram of an illustrative routine **900** that may be executed by a persistent storage node to update the witness service as part of the process of persisting data from the transaction log to persistent storage. Advantageously, updating the witness service as part of the process of persisting the data to persistent storage can help to further ensure that the witness service can provide the latest sequence numbers to cache nodes, even if an update to the witness service during the original write request path processing to the transaction log fails for some reason. Moreover, to reduce the amount of additional network traffic and witness service processing that would result from multiple persistent storage nodes updating the witness service for each data update event (e.g., due to replication of data across multiple nodes), a batching and/or messaging protocol may be implemented.

(123) Although the description that follows is with respect to an example implementation in which the routine **900** is performed in the context of the distributed storage system **104** shown in FIG. **1**, including an index service **112**, the example is for purposes of illustration only and is not intended to be limiting. In some embodiments, the routine **900** may be used in other environments and to manage other types of data. Aspects of the routine **900** will be described with further reference to the illustrative data flows and interactions shown in FIGS. **10** and **11**.

(124) The routine **900** may begin at block **902** in response to an event, such as when a persistent storage node **120** begins operation. When the routine **900** is initiated, a set of executable program instructions stored on one or more non-transitory computer-readable media (e.g., hard drive, flash memory, removable media, etc.) may be loaded into memory (e.g., random access memory or “RAM”) of a computing device and executed by one or more processors. In some embodiments, the routine **900** or portions thereof may be implemented on multiple processors, serially or in parallel.

(125) At block **904**, a persistent storage node **120** may receive data from the transaction log **200** to be persisted in persistent storage. In some embodiments, the persistent storage node **120** may include or be associated with a component that manages persistent storage of data from the transaction log **200** to the persistent storage node. For example, the coordinator subsystem **124** may provide records from the transaction log **200** to one or more persistent storage nodes **120** for storage. As another example, an applier component that is part of the persistent storage node **120** or otherwise separate from the coordinator subsystem **124** and associated with a persistent storage node **120** may manage obtaining records from the transaction log **200** for storage in the persistent storage node **120**.

(126) At block **906**, the persistent storage node **120** can in some embodiments add the current transaction log record to a batch of records. The persistent storage node **120** may maintain transaction log records in a batch until a batch update criterion is satisfied, at which time the persistent storage node **120** may write data from the batched transaction log records to persistent storage (e.g., a high density magnetic disk drive, a solid state drive, etc. of the persistent storage node **120**). For example, the persistent storage node **120** may maintain the transaction log records, or data derived therefrom, in memory until the persistent storage write operation is to be performed.

(127) At decision block **908**, the persistent storage node **120** can determine whether a batch update criterion has been satisfied. The batch update criterion may relate to a variety of different events and/or properties of the batch, as described in greater detail below. If the batch update criterion is satisfied, the routine **900** may proceed to block **910**. Otherwise, if the batch update criterion has not been satisfied, the routine **900** may return to block **904** for possible receipt of another transaction log record. In some embodiments, the routine **900** may also or alternatively wait until the batch update criterion has been satisfied. For example, a batch update criterion analysis thread may sleep and periodically wake up to determine whether the batch update criterion has been satisfied. As another example, an event-based model may be implemented in which certain events trigger evaluation of the batch update criterion even if other portions of the routine **900** are being executed.

(128) In some embodiments, the batch update criterion may relate to a maximum threshold for the quantity of log records (or data derived therefrom) permitted to be maintained in the batch. If the quantity satisfies the threshold (e.g., meets or exceeds the threshold), then the batch update criterion may be satisfied and the routine **900** may proceed to block **910**.

(129) In some embodiments, the batch update criterion may relate to a maximum period of time that a batch may be maintained before being written. For example, when a first transaction log record is received after a prior batch update operation, then a new batch may be created and the time of creation may be recorded. There may be a maximum threshold period of time specified in a quantity of time units (e.g., n milliseconds) that the batch or any individual transaction log record of the batch may be maintained in memory before being written to persistent storage. If the

quantity satisfies the threshold (e.g., meets or exceeds the threshold), then the batch update criterion may be satisfied and the routine **900** may proceed to block **910**.

(130) In some embodiments, the batch update criterion may relate to receipt of a request to access data in the batch. For example, the persistent storage node **120** may receive a read request (e.g., a GET request) for a particular key or otherwise for data that is associated with a sequence number. The persistent storage node **120** may check whether data associated with that key and/or sequence number is currently in a batch waiting to be written to persistent storage. If so, the batch update criterion may be satisfied and the routine **900** may proceed to block **910**. Thus, data accessors are not required to wait for some other batch update criterion to be satisfied before they access their data.

(131) In some embodiments, multiple batch update criteria may be evaluated, conjunctively or disjunctively. For example, a persistent storage node may evaluate multiple criteria during a single execution of decision block **908**: if the threshold quantity of records has been satisfied, or if the threshold period of time has been satisfied (or if both have been satisfied), then the routine **900** may proceed to block **910**. As another example, the two thresholds may be evaluated each time decision block **908** is reached from block **906**, and a request-related criterion may be evaluated on an event-driven basis as indicated by the dotted arrow from block **902** and the dotted circular arrow at decision block **908**. The examples of batch update criteria and evaluations described herein are illustrative only, and are not intended to be limiting, required, or exhaustive. In some embodiments, additional, fewer, and/or alternative batch update criteria and methods of evaluation may be used.

(132) At block **910**, the persistent storage node **120** can send a batch witness service update message to the witness service **116**. In some embodiments, the batch witness service update message may include a listing of all keys and corresponding sequence numbers included in the current batch. The witness service **116** may apply the updated sequence numbers to its verification data set **350**. For example, if the verification data set **350** is implemented as a hash table, the witness service **116** may generate hashes of each key, and use the hashes as indices **360** to access the appropriate verification data items **352** and update the corresponding sequence numbers **362**. If the witness service **116** includes multiple witness nodes **160**, then each node that is associated with the keys in the batch witness service update message (or a subset thereof) may update its verification data set **350** accordingly.

(133) FIG. **10** illustrates an example in which a persistent storage node **120** batches transaction log records (or data derived therefrom) for a batch witness service update message. As shown, the persistent storage node **120** receives records from the transaction log **200** at [1]. The persistent storage node **120** can maintain a batch **1000** of the incoming transaction log records at [2] until one or more batch update criteria are satisfied. When the one or more batch update criteria are satisfied, the persistent storage node **120** can send a batch witness service update message **1002** to the witness service **116** at [3]. Upon successful receipt by the witness service **116** (e.g., via an acknowledgment from the witness service **116**), the persistent storage node **120** can proceed with persistent storage of data from the batched transaction log records at [4], as described below.

(134) At block **912**, the persistent storage node **120** can send a message to one or more other persistent storage nodes regarding the witness service update message that was sent to the witness service **116**. The message can be a notification to the other persistent storage node(s) that the witness service **116** has already been updated with the sequence numbers included in the witness service update message, and the other persistent storage nodes do not need to also update the witness service regarding those keys and sequence numbers when performing their own instances and iterations of routine **900**. In this way, traffic from the persistent storage nodes **120** to the witness service **116**, and processing at the witness service **116**, can be reduced while still obtaining the benefit of an update at the time of persistent storage, even if that update is a secondary confirmatory update.

(135) In some embodiments, the persistent storage node **120** executing the current instance of

routine **900** may maintain a list of the all keys with data update events which the persistent storage node **120** has updated the witness service **116**, and the sequence number before which all data update events have been updated to the witness service **116**. The persistent storage node **120** may generate the message to one or more other persistent storage nodes using the maintained list. For example, the message may comprise a “best effort” signal that includes the sequence number before which all data update events have been updated to the witness service **116**. A recipient of the message may then determine to exclude any such key and/or sequence number from its own witness service update message. Accordingly, the frequency of communications and volume of duplicative information sent to the witness service **116** may be reduced, in some cases substantially. (136) In some embodiments, the message that the persistent storage node **120** sends to other persistent storage nodes may be a notification broadcast to each other persistent storage node of the distributed storage system **104**, or each other persistent storage node associated with a given key or set of keys (e.g., when the set of persistent storage nodes **120** are partitioned as described in greater detail below). In some embodiments, peer-to-peer messages may be sent using a gossip protocol in which a persistent storage node sends a gossip message to one or more predetermined or dynamically-determined persistent storage nodes, those nodes forward the gossip message to one or more other predetermined or dynamically-determined persistent storage nodes, and so on until all nodes have received the gossip message.

(137) At block **914**, the persistent storage node can store data from the batch to persistent storage. The routine **900** may terminate at block **916**.

(138) FIG. **11** illustrates an example in which a persistent storage node **120** sends notifications, such as gossip messages, to other persistent storage nodes regarding a witness service update message sent to the witness service **116**. As shown, the persistent storage node **120** receives a record from the transaction log **200** at [1]. The persistent storage node **120** can update the witness service **116** at [2] by sending a witness service update message **1102**. In addition, the persistent storage node **120** can maintain a list **1100** of keys that have been verified to the to the witness service **116**, and the sequence number before which all keys have been verified. Upon successful receipt by the witness service **116** of the witness service update message **1102** (e.g., via an acknowledgment from the witness service **116**), the persistent storage node **120** can proceed with using the list **1100** to generate and send a gossip message **1104** to one or more other persistent storage nodes at [3]. Prior to, in parallel with, or subsequent to sending the gossip messages **1104**, the persistent storage node **120** can proceed with persistent storage of data from the transaction log records at [4].

(139) In some embodiments, an additional service may consume data update event data from the transaction log **200** and expose the data to consumers. This service may provide an additional opportunity to update the witness service **116**. In some cases, such a service may be able to update the witness service **116** more quickly than the persistent storage nodes **120**. FIG. **12** illustrates an example of an event notification service **1200** that provides information to consumers about data update events that have been written to the transaction log **200**. As shown, the event notification service **1200** is separate from the persistent storage nodes **120**, and operates independently of any persistent storage operations, witness service update messaging, gossip messaging, etc. performed by the persistent storage nodes **120**. At [1], the event notification service **1200** obtains event data **1202** from the transaction log **200**. The event data **1202** may be received on an event-by-event basis (e.g., one message per event), or as a batch (e.g., a message comprising event data for n events). To ensure that consumers of event notifications can access the data that is the subject of the events, the event notification service **1200** can send a witness service update message **1202** to the witness service **116** at [2] prior to sending event notifications to consumers or otherwise exposing event information at [3]. This witness service update message **1202** from the event notification service **1200** can help to ensure data consistency for read requests in cases where the event notification service **1200** provides event notifications prior to any persistent storage node **120**

having updated the witness service **116** and persistently stored data regarding the event. As indicated by the label [0, 1'-3',4], the persistent storage nodes **120** may obtain transaction log data, generate witness service update messages and gossip messages, and persistently store data at any time before, during, or after the illustrated operations of the event notification service **1200**.

(140) Multi-Tier Caching Architecture

(141) FIG. **13A** is a block diagram of an illustrative multi-tier caching architecture. In some embodiments, shown, the multi-tier caching architecture includes a first tier **1302** for caching the most frequently requested objects, and a second tier **1304** for caching other objects. The first tier **1302** may be referred to as a hot key cache, and the second tier **1304** may be referred to as a key cache.

(142) In some embodiments, the key cache **1304** may be implemented using multiple key cache nodes **122** configured as a consistent hash ring. In the consistent hash ring, each key cache node **122** may be mapped to a set of keys using a hash function such that a contiguous portion of the key space, corresponding to data items that are related and often requested close in time with each other, is distributed among multiple nodes. For example, each data item of the distributed storage system **104** may be assigned a unique key, as described above. Illustratively, the keys may be hexadecimal numbers represented by strings of hexadecimal symbols (e.g., numerals 0-9 and letters a-f, where “a” represents a value of 10, “b” represents a value of 11, and so on up to “f” which represents a value of 15). Keys for objects used by the same client **102** or associated with the same workload may be similar, such as being assigned contiguous or substantially contiguous keys. In the consistent hash ring configuration, requests for a given key will be mapped to the same cache node from request to request, in some cases without regard to the passage of time between requests. Thus, distributing keys in contiguous portions of the key space among multiple nodes serves as a load balancing mechanism.

(143) The hot key cache **1302** may be implemented using a randomized load balancing mechanism such that requests for a given hot key will be mapped to different hot key cache nodes **150** from request-to-request in order to balance the load for that specific hot key. For example, the hot key cache **1302** may include the frontend nodes **110** that initially evaluate requests to determine whether they are for hot keys. When a request is received by the distributed storage system **104**, it may initially be assigned to frontend node **110** by a load balancer (e.g., on a randomized basis). If a request is for a hot key, the frontend node **110** may attempt to service the request from its hot key cache node **150**. Otherwise, the request may be passed to the key cache **1304**.

(144) As shown in FIG. **13A**, when a frontend node **110** determines at [1] that the key of a requested data object is not one of a set of hot keys, then the request may be passed to the key cache **1304** at [2]. At [3], a management subsystem such as a key cache manager **1310** may determine the proper key cache node **122** to which the request is to be routed. For example, the key cache manager **1310** may evaluate a hash function or use some other mapping technique to determine the particular key cache node **122** of the set of key cache nodes in the key cache **1304** to which the key of the currently-requested data object is assigned. The key cache manager **1310** may forward the request to the assigned key cache node **122** at [4]. The key cache node **122** may then service the request by determining whether a cached version of the requested data object is available and valid, obtaining an up-to-date version from a persistent storage node **120** at [5] if needed, etc. For example, the key cache node **122** may perform some or all of the operations shown in FIG. **2** and described in greater detail above.

(145) As shown in FIG. **13B**, if a frontend node **110** determines at [A] that the key of the requested data object is one of a set of hot keys, then the frontend node **110** may attempt to service the request from a hot key cache node **150**. The particular frontend node **110** to which the current data object request is initially routed may be selected at random, selected according to a round robin scheme, or selected using some other load balancing method such that requests for particular hot keys are not always routed to the same hot key cache node **150** but rather are distributed among all

hot key cache nodes **150**. At [B], the hot key cache node **150** may determine whether a cached version of the requested data object is available, and if so, whether the cached version is valid. For example, the hot key cache node **150** may submit a verification request to the witness service **116** to obtain sequence data regarding a latest change to the data object, an update-to-date version of the data object, or otherwise the minimum required version of the data object to be distributed. If the cached version of the requested data object is not valid or not available, the hot key cache node **150** can pass the request to the key cache **1302** at [C].

(146) At [D], the key cache manager **1310** may determine the proper key cache node **122** to which the request is to be routed. The key cache management **1310** may then route the request to the selected key cache node **122** at [E]. In some embodiments, the key cache manager **1310** may implement a load balancing method to distribute requests for a particular hot key among multiple key cache nodes **122**, rather than determining a specific mapping as may normally be done in a consistent hash ring implementation. An example of a routine that the key cache management **1310** may use to route requests to key cache nodes **122** is shown in FIG. **14**.

(147) FIG. **14** is a flow diagram of an illustrative routine **1400** that may be executed by the key cache manager **1310** to route requests to key cache nodes **122**. Advantageously, the key cache manager **1310** may route non-hot key requests to particular nodes in a consistent hash ring using a predetermined mapping of keys to nodes, while distributing hot key requests among different nodes in a dynamic load-balanced manner.

(148) Although the description that follows is with respect to an example implementation in which the routine **1400** is performed in the context of the multi-tier caching architecture shown in FIGS. **13A** and **13B**, the example is for purposes of illustration only and is not intended to be limiting. In some embodiments, the routine **1400** may be used in other environments and to manage other types of data.

(149) The routine **1400** may begin at block **1402** in response to an event, such as when a key cache manager **1310** begins operation. When the routine **1400** is initiated, a set of executable program instructions stored on one or more non-transitory computer-readable media (e.g., hard drive, flash memory, removable media, etc.) may be loaded into memory (e.g., random access memory or “RAM”) of a computing device and executed by one or more processors. In some embodiments, the routine **1400** or portions thereof may be implemented on multiple processors, serially or in parallel.

(150) At block **1404**, the key cache manager **1310** may receive a data object request. The data object request may include, reference, or otherwise be associated with a key of the requested data object.

(151) At decision block **1406**, the key cache manager **1310** may determine whether the request is a hot key request. In some embodiments, the request may be tagged as a hot key request or a non-hot key request. In some embodiments, hot key requests and non-hot key requests may be separate requests (e.g., separate request signatures, message types, or the like). In some embodiments, the key cache manager **1310** may determine whether the key of requested data object is a hot key based on a listing of hot keys. The key cache manager **1310** may use the tag, request type, key, another signal, or some combination thereof to determine whether the request is a hot key request. If the request is not a hot key request, the routine **1400** may proceed to block **1408**. Otherwise, if the request is a hot key request, the routine **1400** may proceed to block **1410**.

(152) At block **1408**, because the request is not a hot key request, the key cache manager **1310** may determine the key cache node **122** to which the key of the requested data object is assigned. In some embodiments, the key cache manager **1310** may evaluate a hash function or use some other mapping technique to determine the particular key cache node **122** of the set of key cache nodes in the key cache **1304**. For example, the output of the hash function may be an index or identifier of a key cache node **122**. The routine **1400** may then proceed to block **1414**, where the key cache manager **1310** may forward the request to the assigned key cache node **122**.

(153) In some embodiments, when the request is not a hot key request, the key cache manager **1310** may determine a key cache node **122** from a set of key cache nodes **122** to which the object is assigned or with which the object is otherwise associated. For example, the key cache manager **1310** may use one of the methods described below at blocks **1410** and **1412** for processing hot key requests. In such cases, the quantity of nodes in the subset of key cache nodes that may service a request that is not a hot key request may be smaller than the quantity of nodes in the subset of key cache nodes that may service a hot key request. The smaller subset may be used for non-hot key requests in order to provide a desired cache hit rate, while the larger subset may be used for hot key requests to balance the larger load of such requests, as described below.

(154) At block **1410**, because the request is a hot key request, the key cache manager **1310** may determine a subset of the key cache nodes that may service the request. In some embodiments, there may be x hot key cache nodes **150** in the hot key cache **1302**, and y key cache nodes **122** in the key cache **1304**, where $x \gg y$ (e.g., $x=50,000$ and $y=15,000$). Because of the comparatively large number of hot key cache nodes **150** in such implementations, when a cached data object associated with a hot key is invalidated, there may be a large number of requests passed from the hot key cache **1302** to the key cache **1304** in a short period of time (e.g., up to 50,000). If all requests were sent to a single key cache node **122** as described above for non-hot key requests, then the single key cache node **122** may not be able to service the requests and a situation similar to a denial-of-service attack may be experienced. This effect would be exacerbated if multiple hot keys were invalidated simultaneously or in a short period of time, and more than one of those hot keys is assigned to the same key cache node **122**. To reduce or eliminate these issues, hot key requests may be distributed among multiple key cache nodes **122** to balance the load for handling such requests. However, rather than distributing hot key requests for a single hot key among all of the key cache nodes, the hot key requests may be distributed among a subset of all available key cache nodes **122**. The subset may be used in order to improve the likelihood of cache hits, while still distributing the load for servicing hot key requests among multiple key cache nodes **122** rather than assigning each key to a particular key cache node, as described above for non-hot key requests.

(155) In some embodiments, the key cache manager **1310** may use a plurality of hash functions to identify the subset of the key cache nodes **122** that may service the request. For example, if each key is to be associated with a different respective subset of z nodes, where $z \ll y$, (e.g., $z=400$ and $y=15,000$), then the key cache manager **1310** may evaluate z different hash functions using the key to determine the z nodes that may service the request for the key.

(156) In some embodiments, the key cache manager **1310** may use a single hash function with varied input, rather than multiple separate hash functions, to identify the subset of the key cache nodes **122** that service the request. For example, if each key is to be associated with a different respective subset of z nodes, then the key cache manager **1310** may append z different values to the key and evaluate the hash function once for each of the z different appended values (e.g., if $z=400$ and the data object key="abc" then the hash function may be evaluated for "abc1," "abc2," and so on to "abc400").

(157) In some embodiments, rather than identifying the subset of key cache nodes **122** that are associated with a given hot key each time a hot key request is received, the key cache manager **1310** may generate a random number in a range that corresponds to the size of the subset, append the number to the hot key, and evaluate a hash function to determine the key cache node **122** that is to service this particular hot key request. For example, if the subset is z key cache nodes **122**, then the key cache manager may generate a random number in the range 1- z and append the number to the hot key prior to evaluating the hash function (e.g., a random number **67** is generated in the range 1-400, and appended to the key "abc" to obtain hash function input "abc67"). In this way, the hash function is only evaluated once per request, but the results are still randomized within a particular subset size and tied to the particular key by use the key and random number as input to the hash function.

(158) The example methods of determining the subset of nodes to handle a hot key request are illustrative only, and are not intended to be limiting, required, or exhaustive. In some embodiments, additional or alternative methods may be used.

(159) At block **1412**, the cache manager **1310** may select a particular key cache node **122** from the subset of key cache nodes determined above. The key cache node **122** selected from the subset may be a randomly-selected key cache node **122**. In some embodiments, the cache manager **1310** may generate randomized selection data to select the key cache node **122**. For example, the cache manager **1310** may use a pseudo random number generator to generate a random number (e.g., between 1 and z inclusive). The subset of key cache nodes **122** identified in block **1410** may be assigned indices, and the key cache node **122** with the index that corresponds to the random number may be selected.

(160) At block **1414**, the cache manager **1310** may route the request to the selected key cache node **122**, which can service the request. For example, the key cache node **122** may perform some or all of the operations shown in FIG. 2 and described in greater detail above. In some embodiments, the key cache node **122** may utilize a throttling technique to reduce the load on persistent storage nodes in retrieving updated versions of hot key data items. For example, the key cache node **122** may use a token scheme, such as one management by a token bucket algorithm or leaky bucket algorithm to control the flow of requests to the persistent storage nodes **120**.

(161) In some embodiments, there may be no hot key cache **1302**. In other embodiments, the hot key cache **1302** may not actually cache data objects but may instead merely identify a key of a requested object as a hot key before passing the request to the key cache **1302**. In these embodiments, each request for a hot key is serviced by the key cache **1304**.

(162) FIG. 15 illustrates an architecture **1500** in which the requests are passed through the frontend nodes **110**/hot key cache nodes **150** to the key cache **1304** without caching. Such requests may be referred to as pass-through requests.

(163) As shown, when a data object request is received by the distributed storage system **104**, it may first be evaluated to determine whether the key of the requested data object is a hot key. If a frontend node **110** determines at [1] that the key of the requested data object is one of a set of hot keys, then the request may be passed to the key cache **1304** at [2]. Illustratively, the request may be passed with a tag indicating that the request is for a hot key, or a separate hot key request may be generated and sent to the key cache **1304**. In contrast, if the frontend node **110** determines that the key is not a hot key, then the request may be passed to the key cache **1304** without a tag indicating that the request is for a hot key, or a non-hot key request may be generated and sent to the key cache **1304**, where it may be processed as described above.

(164) The key cache **1304** may process a hot key request [3]. For example, the hot key request may be processed as described above, such as by selecting a particular key cache node **122** from a subset of nodes of the key cache **1304** using a pseudo-random selection process, and forwarding the request on to the selected key cache node **122** at [4].

(165) Performing a verification request from a key cache node **122** to the witness service **116** each time a hot key request is processed may cause a degradation in performance of the witness service **116**. To address and mitigate these situations, verification requests for hot keys may be batched and performed as a single hot key request.

(166) In some embodiments, as shown in FIG. 15, a key cache **1304** may batch verification requests for a given hot key (or different hot keys) at [5]. The batch may be maintained until a batch request criterion is satisfied. For example, the batch request criterion may relate to a maximum size of the batch before a verification request is to be sent. As another example, the batch request criterion may relate to a maximum amount of time that a given request may be batched (e.g., the maximum amount of time since a batch was created) before a verification request is to be sent. As another example, multiple criteria may be used, such as by triggering the verification request when either (1) the batch size meets a threshold, or (2) the length of time since creation of

the batch meets a threshold. When the batch request criterion/criteria are satisfied, a batch verification request may be sent to the witness service **116** at [6]. The batch request may be a single request regarding a single hot key, and the key cache node **122** may apply the results of the request to all pending requests for the hot key that the key cache node **122** is currently servicing. In some embodiments, a batch may include requests for multiple different hot keys, and the batch request may include a separate request for each of the hot keys.

(167) In some embodiments, the keys that are considered to be hot keys are determined based on the rate at which the corresponding objects are requested. For example, keys for objects that are requested at the highest rates, such as the top *i* objects by request rate or the top *j* percent of objects by request rate, may be considered hot keys. As another example, keys for objects that are requested at a threshold rate, such as *k* requests per unit of time, may be considered hot keys. The key cache nodes **122** that service requests may track the request rates for object requests that they service. An individual key cache node **122** can provide, to upper layers in the request path (e.g., key cache manager **1310**, frontend nodes **110**, etc.), a signal regarding the keys that the key cache node **122** considers to be hot (e.g., based on satisfaction of a request rate criterion). For example, as shown in FIG. **15**, a key cache node **122** can provide data regarding hot keys along with a response to a data object request. In this way, prior layers of the request path can be notified of new hot keys, and can then determine whether subsequent requests are hot key requests.

(168) Although the hot key determination at the key cache nodes **122** and signaling to prior layers is shown and described with respect to the embodiment illustrated in FIG. **15**, the example is illustrative only and is not intended to be limiting or required. In some embodiments, other methods of determining and signaling hot keys may be used. In addition, a hot key determination and signaling method may be part of any other embodiment. Illustratively, the embodiments illustrated in FIGS. **13A-13B** may use the same or similar hot key determination and signaling methods as shown in FIG. **15**. For example, key cache nodes **122** may maintain rate data regarding request rates observed by the respective key cache nodes **122** for individual keys and the key cache nodes **122** may signal the request rates or other information regarding likely hot keys to prior layers. In addition, the nodes of the hot key cache **1302** may maintain objects in cache according to a least-recently-used (“LRU”) protocol in which the hot key cache nodes **150** are configured to store a maximum number of objects or dedicate a maximum amount of storage space to objects. When a new object or updated version of an object is obtained and is to be added to the cache, the least-recently-used object (e.g., the least-recently requested and therefore “coldest” object) is purged from the cache. Thus, the hottest keys are maintained by the hot key cache nodes **150**.

(169) Dynamic Partition Management

(170) The distributed storage system **104** may store any number of data objects. In some cases, the number may be quite high, such as millions, billions, or more individual data objects. A single set of computing resources may be unable to adequately maintain and respond to requests for information regarding every object, such as routing information regarding the persistent storage nodes on which a current version of each object may be obtained, verification information regarding the current version of each object, and the like. In order effectively maintain and respond to requests for data objects and related information, the data objects may be separated into discrete subsets, also referred to as partitions, that may be managed by different sets of computing resources.

(171) In some embodiments, the key space of data item keys may be used to define and manage partitions such that a first subset of keys may be assigned to a first partition, a second subset of keys may be assigned to a second partition, and so on. The subsets may be defined with respect to the keys themselves. For example, keys of data items may be defined using hexadecimal numbers of preset or varying length, while a partition may be associated with a particular key prefix. Thus, a partition may be defined as the set of all data items with key prefixes matching a particular key prefix of the partition, and not matching any other partition's key prefix that is longer (e.g., more

hexadecimal symbols). Illustratively, one partition may be defined as the “f” partition, and another partition may be defined as the “foo” partition. There may not be any partition with a key length longer than three hexadecimal symbols beginning with “foo.” In this illustrative example, all data items with keys beginning with “foo” would be part of the “foo” partition, while all other data items with keys beginning with “f” would be part of the “f” partition.

(172) Each partition may be associated with a set of computing resources, some of which may be uniquely assigned to the partition, and some of which may be shared by multiple partitions. In some embodiments, a set of computing resources used to manage a partition of data objects may include a key cache **1304**, a set of persistent storage nodes **120**, a transaction log **200**, and a witness service **116** instance. Illustratively, the transaction log **200** and persistent storage nodes **120** may be specific to a single partition, while the key cache **1304** and witness service **116** instance may be shared among several partitions.

(173) In a distributed storage system **104** in which the number of data objects and requests for the data objects increases over the course of time, and sometimes in short periods of time, the load on a set of computing resources assigned to a partition can increase and cause performance degradation. To balance the load and otherwise improve performance, new partitions can be split from partitions that are experiencing a heavy volume of requests.

(174) FIG. **16** is a block diagram of illustrative data flows and interactions between cache nodes **202**, a witness service **116**, a set of persistent storage nodes **120**, and a control system **1600** (also referred to as a “control plane system”) to manage data object partitions. The cache nodes **202** may be hot key cache nodes **150** and/or key cache nodes **122**.

(175) In some embodiments, as shown, the cache nodes **202** may receive data object requests. Prior to servicing the requests, the cache nodes **202** may make verification requests to the witness service **116** instance for the partition. The witness service **116** instance may track such requests. For example, the witness service **116** instance may maintain request rate information regarding verification requests associated with the partition. The request rate may be determined using an exponential decay algorithm, for example, in which requests are weighted and those weights are decayed over time (e.g., according to a half-life).

(176) In addition to handling data object requests as described in greater detail herein, the cache nodes **202** may in some embodiments send a subset of data object requests to the persistent storage nodes **120** even if the cache nodes **202** would otherwise be able to service the requests out of their caches. For example, the cache nodes **202** may randomly select individual requests to pass through to the persistent storage nodes **120** such that a statistically significant portion of data item requests are passed through in a given period of time. Illustratively, the cache nodes **202** may use a random number generation algorithm to select a percentage of requests (e.g., 0.1, 0.5, 1.0, 5.0 etc.) to pass through to the persistent storage nodes **120** without first checking to see whether cached versions of the requested data objects are available at the cache nodes **202** and/or without first generating verification requests to the witness service **116** instance regarding the requested data objects.

(177) The cache nodes **202** may pass the portion of requests to the persistent storage nodes **120** so that the persistent storage nodes may generate more granular data than maintained by the witness service **116** instance. Such requests may be referred to as pass-through requests. In some embodiments, the storage nodes **120** may store data regarding requests associated with particular key portions (e.g., key prefixes) that are longer or otherwise more specific than the prefix of the current partition with which the storage nodes **120** are associated. For example, if the current partition is the “f” partition, a storage node **120** may store request rate data regarding requests for data objects, such as using an exponential decay algorithm in which requests are weighted and those weights are decayed over time (e.g., according to a half-life). The storage node **120** may determine when the request rate for data objects associated with keys having a more specific prefix (e.g., a prefix beginning with “f” followed by one or more additional hexadecimal symbols) satisfies a criterion, such as a threshold. If the request rate satisfies the criterion, then the storage

node **120** may split out the more specific prefix and begin tracking request rates for prefixes that are even more specific. The request rate information may be used to determine how to split the current partition into multiple partitions such that the overall request rates of the partitions may be close to equal.

(178) FIG. **16** shows an example of a rate tracking structure **1610** that may be used to track the request rates (e.g., as measured or estimated in transactions per second or “TPS”) for various subsets of a partition. As described above, the partition may be associated with a particular prefix. The rate tracking structure **1610** may be used to track request requests for data objects associated with keys that are more specific than the prefix of the current partition.

(179) In some embodiments, as shown, the rate tracking structure **1610** may be a heap tree (e.g., a minimum heap tree or maximum heap tree) with a set of elements. The root element **1612** may be associated with the prefix of the current partition. Each child, grandchild, and other descendent element may be associated with a more specific prefix that begins with the prefix of the current partition. In addition, each element may be associated with a rate tracker. As requests are received by the persistent storage node **120**, the rate trackers for the most specific prefix matching the keys of the requested data objects are updated to capture the incoming request rate. For example, when a request is received for a key associated with a particular element **1614** (e.g., the element of the rate tracking data structure **1610** that has the most specific prefix corresponding to the requested key), then the rate tracker for that element **1614** is evaluated (e.g., using an exponential decay algorithm) to determine whether the rate exceeds a threshold or satisfies some other criterion. If so, a child element **1616** may be created to track a request rate for an even more specific key prefix, such the more specific key prefix that has the highest request rate of all prefixes more specific that the prefix of the existing element **1614**.

(180) In some embodiments, the rate tracking data structure **1610** may have a fixed number of elements. Thus, when the rate tracker for a particular element exceeds a threshold or satisfies some other criterion, then the persistent storage node **120** may replace an existing element with a new element to track the request rate for a more specific prefix.

(181) The control system **1600** may determine, based on data from the witness service **116** instance, that a particular partition is experiencing too high of a volume of requests (e.g., there is too much “heat” on the partition). The level of heat may cause a degradation in performance of the witness service **116** instance to respond to verification requests associated with the partition or other partitions that share the same witness service **116** instance. In response, the control system **1600** may obtain more granular request rate data from the persistent storage nodes **120** assigned to the partition. For example, the control system **1600** may obtain the rate tracking structure **1610** or individual values derived therefrom to determine how to split the current partition and reduce the heat on the partition.

(182) FIG. **17** is a flow diagram of an illustrative routine **1700** that may be executed by the control system **1600** to make scaling decisions. The scaling decisions may be made, and scaling actions may be implemented, to reduce heat on the witness service **116** instance to which a partition is assigned, and/or to reduce the likelihood of overloading persistent storage nodes when large volume of cache invalidations occur for keys assigned to the persistent storage nodes.

Advantageously, the specific manner in which the computing resources for a partition are scaled (e.g., the specific manner in which the partition is split) to reduce heat on the witness service **116** instance may be determined using data obtained from an entirely separate system (e.g., the persistent storage nodes **120**), and in some cases based on requests never verified with the witness service **116** instance (e.g., based on requests randomly passed through to the persistent storage nodes **120**).

(183) The routine **1700** may begin at block **1702** in response to an event, such as when a control system **1600** begins operation. When the routine **1700** is initiated, a set of executable program instructions stored on one or more non-transitory computer-readable media (e.g., hard drive, flash

memory, removable media, etc.) may be loaded into memory (e.g., random access memory or “RAM”) of a computing device and executed by one or more processors. In some embodiments, the routine **1700** or portions thereof may be implemented on multiple processors, serially or in parallel.

(184) At block **1704**, the control system **1600** may receive partition request rate data from a witness service **116** instance. In some embodiments, the request rate data may be obtained from a priority queue of the witness service **116** instance. For example, the control system **1600** may receive request rate data associated with the top n partitions (e.g., where n=10) in terms of experiencing the greatest current request rates.

(185) At decision block **1706**, the control system **1600** may determine whether any partition satisfies a scaling criterion. If the scaling criterion is satisfied for any partition, the routine **1700** may proceed to block **1708**. Otherwise, if the scaling criterion is not satisfied for any partition, then the routine **1700** may return to block **1704** to continue monitoring request rate data from the witness service **116** instance.

(186) In some embodiments, the determination may be made based on an analysis of the request rate data for individual partitions with respect to a threshold. For example, if a current request rate for a particular partition exceeds a maximum threshold, then the scaling criterion may be satisfied for that partition. In some embodiments, such as when a scaling action to be implemented is the splitting of the partition, the scaling criterion may also be referred to as a partition splitting criterion.

(187) At block **1708**, the control system **1600** can obtain key prefix request rate data from one or more persistent storage nodes **120** associated with the partition that is to be split. In some embodiments, the control system **1600** can obtain the rate tracking structure **1610**, or data derived therefrom. For example, the control system **1600** may request or otherwise receive specific key prefixes and corresponding request rate metrics.

(188) At block **1710**, the control system **1600** can determine how to scale the computing resources assigned to the partition. In some embodiments, the scaling action is to split the partition and assign at least one portion of the partition to a different set of computing resources than another portion of the partition.

(189) The control system **1600** may evaluate the key prefix request rate data to determine an even split of request rate across the split partitions. In some embodiments, the control system **1600** may be configured to split one specific key prefix off of the current partition, or to otherwise split the partition into two. For example, the control system **1600** may determine a split that is as close to 50/50 as possible based on the request rate data for the more specific prefixes available.

Illustratively, the control system **1600** may select the prefix associated with as close to 50% of the overall request rate for the partition as possible. In some cases, the selected prefix may have greater than 50% of the overall request rate, and in other cases the selected prefix may be associated with less than 50% of the overall request for the partition. In some cases, the control system **1600** may not be able to determine an even or substantially even split of request rate for the partition. For example, if a particular key is experiencing an extremely high rate of requests (e.g., the key is a hot key), then that key or the most specific prefix including the key for which request rate data is available may be associated with significantly more than 50% of the requests. In such cases, the control system **1600** may determine to split that specific prefix from the current partition so that the heat associated with that specific prefix does not continue to negatively affect the rest of the partition.

(190) At block **1712**, the control system **1600** can split the partition as determined above. Splitting the partition may include assigning a new set of computing resources to manage the more specific key prefix that is being split from the current partition, while keeping the current partition with the existing set of computing resources.

(191) FIG. **18** illustrates an example of splitting a partition off of an existing partition and moving

it to a new set of computing resources. As shown, a first partition may be associated with a prefix “f” and may be managed by a first set of computing resources, including a transaction log **200A**, a persistent storage node pool **1800A**, and a witness service instance **116A**. A control system **1600** may determine to split off a subset of data objects from the first partition, such as data objects associated with keys having the more specific prefix “foo,” into a second partition.

(192) The control system **1600** may assign the second partition to a new transaction log **200B**, persistent storage node pool **1800B**, and witness service instance **116B**. For example, the control system **1600** may copy data associated with the second partition from persistent storage node pool **1800A** to persistent storage pool **1800B**. Transaction log **200B** may be the log to which new writes associated with the partition are written, and from which the data is eventually persisted to the persistent storage pool **1800B**. The control system **1600** may also cause all of the cache records associated with the second partition to be invalidated so that cache requests are eventually sent to the persistent storage pool **1800B** to receive the most up-to-date versions of data objects from the partition. In addition to updating the data objects themselves, the various components of the distributed storage system **104** may obtain and use updated routing data to access the data in the partition that has been split out. Examples of managing routing data updates are described in greater detail below.

(193) Routing Data Management

(194) When a request for a data object is received by the distributed storage system **104**, the request may require routing to the computing resources that manage the partition to which the requested data object is assigned. When a data object is moved from one set of computing resources to another set of computing resources (e.g., when a partition is split into two partitions), routing data used to route request for that data object may be changed and any system or component using out-of-date routing data will not be able to access the data object, or may access an out-of-date version of the object.

(195) To address these issues, among others, a routing data management protocol may be implemented such that when routing data associated with a data object, partition of data objects, or a group of partitions is changed, access to the data object using prior versions of the routing data is prevented.

(196) FIG. **19** is a diagram of an illustrative multi-tier architecture for organizing partitions of data objects, managing changes to routing data, and managing access to data objects using the routing data. In the illustrated multi-tier architecture, data objects may be assigned to partitions, and partitions may be managed as groups or “cells.”

(197) In some embodiments, cells may be defined according to the keys of the partitions within the cells, similar to partitions being defined according to the keys of the data objects within the partitions. In contrast to the most-specific key prefix assignment of data objects to partitions (e.g., whereby a data object is assigned to the partition with the longest or otherwise most-specific key prefix matching the key prefix of the data object), cells may be defined using a range of keys such that any partition with a key that is within a cell's range of keys is assigned to the cell.

(198) In the illustrated example, two cells are shown: cell **1902** and cell **1904**. The cells are each associated with a different contiguous range of keys in the key space **1950** of the distributed storage system **104**. Cell **1902** is associated with the range of keys beginning with any of the symbols “0”-“9.” Cell **1904** is associated with the range of keys beginning with any of the symbols “a”-“f.”

(199) Cell **1904** is shown with two illustrative partitions: partition **1942** and **1944**. Partition **1942** is associated with key beginning with a key prefix of “f.” Partition **1944** is associated with key beginning with a key prefix of “foo.” In this example, any data object of the distributed storage system **104** with a key beginning with “foo” will be assigned to partition **1944**, assuming there are no other partitions with a matching key prefix that is more specific than “foo.” Any data object of the distributed storage system **104** will be assigned to partition **1942**, assuming there are no other

partitions with a matching key prefix that is more specific than “f.” All of these data objects are assigned to cell **1904**.

(200) When a request is received by the distributed storage system **104**, it may be handled by an inter-cellular routing subsystem **1900**. The inter-cellular routing subsystem **1900** may use inter-cellular routing data **1910** to route the request to the proper cell. In some embodiments, the inter-cellular routing data **1910** may specify computing resources to which the request is to be routed. For example, the inter-cellular routing data **1910** may specify the ranges of keys assigned to or otherwise associated with each cell of the distributed storage system **104**, or some subset thereof. Thus, the inter-cellular routing subsystem **1900** can resolve the proper cell for an incoming request based on the key(s) in the request and the ranges of keys in the inter-cellular routing data **1910**. In some embodiments, the inter-cellular routing data **1910** may also include the network address, host names, and/or other information that may be used to route a request in a network to the appropriate system, device, subnet, or other network location.

(201) Once routed to a particular cell, an intra-cellular routing subsystem may use further routing metadata—intra-cellular routing data—to route the request to a particular partition based on the prefix of the key and the prefixes assigned to the different partitions. In some embodiments, intra-cellular routing metadata can be used to determine the network address, host name, and/or other information regarding the destination to which the request is to be routed. Once routed to a particular partition, the request may be handled by a cache node and/or passed to a persistent storage node.

(202) With reference to an illustrative example, if a request is received for a data object with a key of “foobar,” then the inter-cellular routing subsystem **1900** may use inter-cellular routing data **1910** to route the request to cell **1904**, which is associated with keys in the range “a”-“f.” The intra-cellular routing subsystem **1930** of cell **1904** may use intra-cellular routing data **1932** to route the request to the computing resources assigned to partition **1944**. Illustratively, the intra-cellular routing data **1932** may specify the key prefix assigned to each partition of the cell **1904**. In some embodiments, the intra-cellular routing data **1932** may also include the network address, host names, and/or other information that may be used to route a request in a network to the appropriate system, device, subnet, or other network location.

(203) In contrast, if the key of the requested data object was “89f,” then the inter-cellular routing subsystem **1900** may use inter-cellular routing data **1910** to route the request to cell **1902**, which is associated with keys in the range “0”-“9.” The intra-cellular routing subsystem **1920** of cell **1902** may use intra-cellular routing data **1922** to route the request to the computing resources assigned to a particular partition.

(204) FIG. **20** is a flow diagram of an illustrative routine **2000** that may be executed by the control system **1600** to modify routing data, such as inter-cellular routing data **1910** and/or intra-cellular routing data **1932**, based on partition splitting operations, movement of partitions from one cell to another cell, etc. Advantageously, the control system **1600** may use special transaction log records, also referred to as “barrier” records, to cause storage nodes to reject requests associated with outdated routing data, to invalidate caches of data associated with outdated routing data, and the like.

(205) The routine **2000** will be described with regards to “assigning” and changes to “assignments,” which may be used for both cell assignments and partition assignments interchangeably. In addition, the routine **2000** will also be described with regards to changes to “locations,” which may be used to refer to the computing resources on which data objects of cells and partitions, interchangeably, are stored. In general, changes to either inter-cellular routing data or intra-cellular routing data—or both—may be managed using the routine **2000**. Aspects of the routine **2000** will be described with further reference to the illustrative data flows and interactions shown in FIG. **21**.

(206) The routine **2000** may begin at block **2002** in response to an event, such as when a partition is

split, when a partition is moved from one cell another cell, when a new cell is split from an existing cell, or the like. When the routine **2000** is initiated, a set of executable program instructions stored on one or more non-transitory computer-readable media (e.g., hard drive, flash memory, removable media, etc.) may be loaded into memory (e.g., random access memory or “RAM”) of a computing device and executed by one or more processors. In some embodiments, the routine **2000** or portions thereof may be implemented on multiple processors, serially or in parallel.

(207) At block **2004**, the control system **1600** may update inter-cellular routing data for an affected cell or partition. In some embodiments, an affected cell may include partitions that are part of a cell being split off of an existing cell to form a new cell, or part of an existing cell that is being moved to another existing cell. For example, if a prefix range such as “f0” to “ff” (e.g., all of the prefixes beginning with “f”) is being moved from an “a” to “f” prefix range cell to a different cell, the control system **1600** may update the inter-cellular routing metadata **1910** to include the prefix range. Initially, the prefix range may be assigned to an existing cell, such as cell **1904** shown in FIG. **19**. After subsequent processing operations are performed, the inter-cellular routing metadata **1910** may be updated to assign the prefix range to the new cell, as described in greater detail below.

(208) In some embodiments, there may be separate items of routing data (whether inter-cellular or intra-cellular) for write requests and for read requests. This can be done to permit a phased migration of cells and partitions whereby data may continue to be read from a prior location (e.g., a prior set of computing resources) while writes are performed with respect to both the prior location and the new location (e.g., a new set of computing resources) until the new location is caught up or otherwise ready for use. In the present example, the inter-cellular routing data may be modified such that read requests continue to be routed to the prior/existing cell **1904**, while write requests are to be routed to both the prior cell **1904** and to the new cell to which the partitions for the prefixes are being moved. This version of routing data may be referred to as an intermediate version of routing data, or a first intermediate version of routing data if there are multiple intermediate steps as described below.

(209) At block **2006**, the control system **1600** can notify computing resources associated with affected partitions that the applicable routing data has been changed. For example, the change made to the routing data above may cause a version of inter-cellular routing data to be incremented. The affected partitions may include any partitions with prefixes in the range of the affected cell. The control system **1600** can notify the computing resources associated with the affected partitions regarding the new version of routing data so that the computing resources can ensure that any requests associated with data in the affected partitions are made based on the new version of the routing data.

(210) In some embodiments, the control system **1600** may cause a particular type of record, such as a barrier record, to be written to the transaction log of an affected partition. The barrier record may indicate the new version of routing data that is to be used to make any subsequent requests regarding keys in the affected partition. When such a barrier record is written to the transaction log **200**, a coordinator subsystem **124** may not permit any requests to a persistent storage node **120** of the partition until the persistent storage node **120** has read the barrier record and applied its meaning.

(211) FIG. **21** illustrates an example in which a barrier record **2100** is written to a transaction log **200**, from where it may be read by a coordinator subsystem **124** at [1]. The coordinator subsystem **124** thereafter prevents a persistent storage node **120** at [2] from responding to any request until after the persistent storage node **120** has caught up on the transaction log **200** to the barrier record **2100**. At [3], the coordinator subsystem **124** may notify the witness service **116** for the current partition that the sequence numbers for keys in the current partition should be updated to the sequence number of the barrier record **2100**. Once the persistent storage node **120** reads the barrier record **2100** at [4] (e.g., during the course of reading the records of the transaction log **200** and making any needed changes to persistent storage to reflect the changes recorded in the transaction

log **200**), the persistent storage node **120** may determine at [5] that a new version of routing data to be used to make requests. The persistent storage node **120** may thereafter reject any request that is made based on an outdated version of the routing data. For example, at [6] a cache node **202** may obtain the updated sequence number for a requested key from the witness service **116** in response to a request for a data object associated with the key. The updated sequence number provided by the witness service **116** may be the sequence number of the barrier record **2100**. The cache node **202** may determine at [7] that the sequence number associated with its cached version of the requested data object is less than the sequence number returned from the witness service **116**, and may thereafter request an updated version of the data object from the persistent storage node **120** at [8]. The request may be tagged with or otherwise reference the version of routing data used to route the request to the partition. Based on this tag and the new version data **2102** read by the persistent storage node **120** from the barrier record **2100**, the persistent storage node **120** may determine at [9] whether to reject the request as being associated with out-of-date routing data.

(212) At block **2008**, the control system **1600** can perform a backfill operation in which data is copied from the prior computing resources associated with affected partitions to new computing resources that will manage the affected partitions going forward. In some embodiments, the backfill operation may be performed as a background process while data continues to be read from the prior location, and while new write requests are preformed to both the prior location and new location. For example, the control system **1600** may copy data from persistent storage nodes of the prior set of computing resources into persistent storage nodes of the newly-assigned set of computing resources. Any new write requests that are received during the period in which the intermediate version of routing data is being used are written to both the transaction log of the prior set of computing resources and the transaction log of the new set of computing resources.

(213) At block **2010**, the control system **1600** can update the routing data to a transitional state. The transitional state—also referred to as a second intermediate state in cases where it occurs after a first intermediate state—may serve as a transition between (1) the first intermediate state in which read requests are routed to the prior location while write requests are routed to both locations, and (2) the desired end state in which both read and write requests are routed exclusively to the new location.

(214) In some embodiments, control system **1600** can update the routing data to the transitional state by updating both write and read request routing assignments to a null value. In this state, any other system using the routing data will be unable to determine the location to which a request is to be routed, and the request may therefore be rejected. By using this transitional state, the control system **1600** can be assured that there are no pending requests—read or write requests—routed to the old location during/after the routing change to the desired end state in which both read and write requests are routed exclusively to the new location. Once the backfill operation has completed, and because all write requests made since the beginning of the routine **2000** have been routed to both the new location and the prior location, the new location has all data for the affected partitions.

(215) At block **2012**, the control system **1600** can notify computing resources associated with the affected partitions that the applicable routing data has been changed. For example, the control system **1600** may write another barrier record to the transaction log of the affected partitions to indicate the newly-updated routing data version number, as described above.

(216) At block **2014** the control system **1600** can update the routing data to read/write exclusively to/from the new location.

(217) At block **2016**, the control system **1600** can notify computing resources associated with the affected partitions that the applicable routing data has been changed. For example, the control system **1600** may write another barrier record to the transaction log of the affected partitions to indicate the updated routing data version number, as described above.

(218) In some embodiments, when a barrier record is written to the transaction log, the barrier may

trigger an invalidation of cached versions of data assigned to the transaction log's partition. For example, the coordinator subsystem **124** may generate a notification to the witness service **116** indicating that the sequence number to be applied to all data objects associated with the partition is the transaction log sequence number of the barrier record. Because cache nodes make verification requests to the witness service **116** before serving cached data objects, this update of the sequence number at the witness service can have the effect of invalidating all cached versions of data objects in the partition. Such an operation may be desirable to ensure that cache nodes eventually obtain updated versions from the correct location and communicate with the correct witness service instance based on updated routing information.

(219) Terminology and Additional Considerations

(220) All of the methods and tasks described herein may be performed and fully automated by a computer system. The computer system may, in some cases, include multiple distinct computers or computing devices (e.g., physical servers, workstations, storage arrays, cloud computing resources, etc.) that communicate and interoperate over a network to perform the described functions. Each such computing device typically includes a processor (or multiple processors) that executes program instructions or modules stored in a memory or other non-transitory computer-readable storage medium or device (e.g., solid state storage devices, disk drives, etc.). The various functions disclosed herein may be embodied in such program instructions, or may be implemented in application-specific circuitry (e.g., ASICs or FPGAs) of the computer system. Where the computer system includes multiple computing devices, these devices may, but need not, be co-located. The results of the disclosed methods and tasks may be persistently stored by transforming physical storage devices, such as solid-state memory chips or magnetic disks, into a different state. In some embodiments, the computer system may be a cloud-based computing system whose processing resources are shared by multiple distinct business entities or other users.

(221) Depending on the embodiment, certain acts, events, or functions of any of the processes or algorithms described herein can be performed in a different sequence, can be added, merged, or left out altogether (e.g., not all described operations or events are necessary for the practice of the algorithm). Moreover, in certain embodiments, operations or events can be performed concurrently, e.g., through multi-threaded processing, interrupt processing, or multiple processors or processor cores or on other parallel architectures, rather than sequentially.

(222) The various illustrative logical blocks, modules, routines, and algorithm steps described in connection with the embodiments disclosed herein can be implemented as electronic hardware, or combinations of electronic hardware and computer software. To clearly illustrate this interchangeability, various illustrative components, blocks, modules, and steps have been described above generally in terms of their functionality. Whether such functionality is implemented as hardware, or as software that runs on hardware, depends upon the particular application and design constraints imposed on the overall system. The described functionality can be implemented in varying ways for each particular application, but such implementation decisions should not be interpreted as causing a departure from the scope of the disclosure.

(223) Moreover, the various illustrative logical blocks and modules described in connection with the embodiments disclosed herein can be implemented or performed by a machine, such as a processor device, a digital signal processor (DSP), an application specific integrated circuit (ASIC), a field programmable gate array (FPGA) or other programmable logic device, discrete gate or transistor logic, discrete hardware components, or any combination thereof designed to perform the functions described herein. A processor device can be a microprocessor, but in the alternative, the processor device can be a controller, microcontroller, or state machine, combinations of the same, or the like. A processor device can include electrical circuitry configured to process computer-executable instructions. In another embodiment, a processor device includes an FPGA or other programmable device that performs logic operations without processing computer-executable instructions. A processor device can also be implemented as a combination of computing devices,

e.g., a combination of a DSP and a microprocessor, a plurality of microprocessors, one or more microprocessors in conjunction with a DSP core, or any other such configuration. Although described herein primarily with respect to digital technology, a processor device may also include primarily analog components. For example, some or all of the algorithms described herein may be implemented in analog circuitry or mixed analog and digital circuitry. A computing environment can include any type of computer system, including, but not limited to, a computer system based on a microprocessor, a mainframe computer, a digital signal processor, a portable computing device, a device controller, or a computational engine within an appliance, to name a few.

(224) The elements of a method, process, routine, or algorithm described in connection with the embodiments disclosed herein can be embodied directly in hardware, in a software module executed by a processor device, or in a combination of the two. A software module can reside in RAM memory, flash memory, ROM memory, EPROM memory, EEPROM memory, registers, hard disk, a removable disk, a CD-ROM, or any other form of a non-transitory computer-readable storage medium. An exemplary storage medium can be coupled to the processor device such that the processor device can read information from, and write information to, the storage medium. In the alternative, the storage medium can be integral to the processor device. The processor device and the storage medium can reside in an ASIC. The ASIC can reside in a user terminal. In the alternative, the processor device and the storage medium can reside as discrete components in a user terminal.

(225) Conditional language used herein, such as, among others, “can,” “could,” “might,” “may,” “e.g.,” and the like, unless specifically stated otherwise, or otherwise understood within the context as used, is generally intended to convey that certain embodiments include, while other embodiments do not include, certain features, elements and/or steps. Thus, such conditional language is not generally intended to imply that features, elements and/or steps are in any way required for one or more embodiments or that one or more embodiments necessarily include logic for deciding, with or without other input or prompting, whether these features, elements and/or steps are included or are to be performed in any particular embodiment. The terms “comprising,” “including,” “having,” and the like are synonymous and are used inclusively, in an open-ended fashion, and do not exclude additional elements, features, acts, operations, and so forth. Also, the term “or” is used in its inclusive sense (and not in its exclusive sense) so that when used, for example, to connect a list of elements, the term “or” means one, some, or all of the elements in the list.

(226) Disjunctive language such as the phrase “at least one of X, Y, Z,” unless specifically stated otherwise, is otherwise understood with the context as used in general to present that an item, term, etc., may be either X, Y, or Z, or any combination thereof (e.g., X, Y, and/or Z). Thus, such disjunctive language is not generally intended to, and should not, imply that certain embodiments require at least one of X, at least one of Y, or at least one of Z to each be present.

(227) Unless otherwise explicitly stated, articles such as “a” or “an” should generally be interpreted to include one or more described items. Accordingly, phrases such as “a device configured to” are intended to include one or more recited devices. Such one or more recited devices can also be collectively configured to carry out the stated recitations. For example, “a processor configured to carry out recitations A, B and C” can include a first processor configured to carry out recitation A working in conjunction with a second processor configured to carry out recitations B and C.

(228) While the above detailed description has shown, described, and pointed out novel features as applied to various embodiments, it can be understood that various omissions, substitutions, and changes in the form and details of the devices or algorithms illustrated can be made without departing from the spirit of the disclosure. As can be recognized, certain embodiments described herein can be embodied within a form that does not provide all of the features and benefits set forth herein, as some features can be used or practiced separately from others. The scope of certain embodiments disclosed herein is indicated by the appended claims rather than by the foregoing

description. All changes which come within the meaning and range of equivalency of the claims are to be embraced within their scope.

Claims

1. A distributed storage system comprising: a persistent storage layer comprising a plurality of persistent storage nodes, wherein the persistent storage layer stores a plurality of data objects; a caching layer comprising a plurality of cache nodes, wherein the caching layer stores cached versions of the plurality of data objects; a first witness service instance storing freshness data regarding a first partition of the plurality of data objects, wherein the first partition is associated with a first key prefix, wherein individual data objects in the first partition are associated with corresponding keys having the first key prefix, and wherein the first witness service instance maintains first request rate data regarding freshness requests associated with the first partition; and a management subsystem comprising one or more computing devices; wherein the plurality of cache nodes are configured to: send, to the first witness service instance, freshness requests for a first portion of data object requests received by the caching layer; and determine to pass a second portion of the data object requests to the persistent storage layer without sending corresponding freshness requests to the first witness service instance; wherein a first subset of the plurality of persistent storage nodes are configured to: provide, to the caching layer, data objects in the first partition; and maintain second request rate data regarding requests for data objects in the first partition and associated with a second key prefix longer than the first key prefix; and wherein the management subsystem is configured to: determine that the first request rate data satisfies a splitting criterion; determine, based at least partly on the second request rate data, to split a second partition from the first partition, wherein the second partition comprises data objects associated with the second key prefix; and generate the second partition associated with a second witness service instance storing freshness data regarding the second partition.
2. The distributed storage system of claim 1, wherein the plurality of cache nodes are further configured to randomly select data object requests to be in the second portion of data object requests passed to the persistent storage layer, and wherein the second portion of the data object requests comprises a statistically significant portion of data object requests received by the plurality of cache nodes.
3. The distributed storage system of claim 1, wherein subsequent to generation of the second partition, the first partition does not include the data objects associated with the second key prefix.
4. The distributed storage system of claim 1, wherein the first partition is associated with a first transaction log, the first subset of the plurality of persistent storage nodes, and the first witness service instance, and wherein the second partition is associated with a second transaction log, a second subset of the plurality of persistent storage nodes, and the second witness service instance.
5. A computer-implemented method comprising: under control of a distributed storage system comprising one or more computing devices configured to execute specific instructions: generating first request rate data regarding pass-through requests received by a set of persistent storage nodes, wherein the pass-through requests comprise requests randomly-selected by a set of cache nodes to be passed to the set of persistent storage nodes, and wherein the pass-through requests comprise requests for data objects in a first partition of a plurality of partitions of data objects; determining, based on second request rate data associated with a first witness service, that a scaling criterion has been satisfied for the first partition, wherein the first witness service is configured to provide, to the set of cache nodes, data regarding current versions of data objects cached by the set of cache nodes; determining, based on the first request rate data, a scaling action to scale computing resources associated with the first partition; and performing the scaling action to scale the computing resources associated with the first partition.
6. The computer-implemented method of claim 5, wherein determining the scaling action

comprises determining to generate a second partition of the plurality of partitions, the second partition comprising a subset of data objects in the first partition, wherein the subset of data objects satisfies a request rate criterion.

7. The computer-implemented method of claim 6, wherein performing the scaling action comprises generating the second partition, wherein the second partition is associated with a second witness service separate from the first witness service.

8. The computer-implemented method of claim 5, wherein generating the first request rate data comprises: generating a first request rate associated with a first subset of the data objects in the first partition; and generating a second request rate associated with a second subset of the data objects of the first partition.

9. The computer-implemented method of claim 8, wherein generating the first request rate comprises generating the first request rate based on requests for data objects with keys having a first prefix; and wherein generating the second request rate comprises generating the second request rate based on requests for data objects with keys having a second prefix different from the first prefix.

10. The computer-implemented method of claim 9, wherein performing the scaling action comprises generating a second partition of the plurality of partitions, the second partition comprising a subset of data objects in the first partition, wherein the subset of data objects satisfies a request rate criterion, and wherein generating the second partition comprises assigning data objects associated with the second prefix to the second partition, wherein data objects associated with the first prefix remain assigned to the first partition.

11. The computer-implemented method of claim 10, wherein generating the second partition comprises moving data objects in the second partition from the set of persistent storage nodes to a second set of persistent storage nodes.

12. The computer-implemented method of claim 5, further comprising sending, by the set of cache nodes to the first witness service, verification requests for sequence data associated with changes to data objects cached by the set of cache nodes.

13. The computer-implemented method of claim 12, further comprising: receiving, by the first witness service from a cache node of the set of cache nodes, a verification request comprising a key of a data object; generating a hash of the key; obtaining a first sequence number from a hash table based at least partly on the hash; and sending the first sequence number to the cache node.

14. The computer-implemented method of claim 13, further comprising: determining, by the cache node, that the first sequence number is greater than a second sequence number associated with a cached data regarding the data object; and requesting the data object from the set of persistent storage nodes.

15. The computer-implemented method of claim 5, further comprising passing, by the set of cache nodes to the set of persistent storage nodes, a randomly-selected statistically significant portion of requests for data objects cached by the set of cache nodes, wherein the pass-through requests comprise the randomly-selected statistically significant portion of requests.

16. The computer-implemented method of claim 5, further comprising maintaining, by the first witness service, the second request rate data in a priority queue, wherein a first portion of the second request rate data is associated with the first partition, and wherein a second portion of the second request rate data is associated with a second partition of the plurality of partitions.

17. A system comprising: computer-readable memory storing executable instructions; and one or more computing devices programmed by the executable instructions to at least: generate first request rate data regarding pass-through requests received by a set of persistent storage nodes, wherein the pass-through requests comprise requests randomly-selected by a set of cache nodes to be passed to the set of persistent storage nodes, and wherein the pass-through requests comprise requests for data objects in a first partition of a plurality of partitions of data objects; determine, based on second request rate data associated with a first witness service instance, that a partition

splitting criterion has been satisfied for the first partition, wherein the first witness service instance is configured to provide, to the set of cache nodes, data regarding current versions of data objects cached by the set of cache nodes; determine, based on the first request rate data, to generate a second partition of the plurality of partitions, the second partition comprising a subset of data objects in the first partition, wherein the subset of data objects satisfies a request rate criterion; and generate the second partition, wherein the second partition is associated with a second witness service instance separate from the first witness service instance.

18. The system of claim 17, wherein to generate the first request rate data, the one or more computing devices are programmed by further executable instructions to: generate a first request rate associated with the subset of data objects in the first partition, wherein the first request rate is generated based on requests for data objects with keys having a first prefix; and generate a second request rate associated with a second subset of data objects of the first partition, wherein the second request rate is generated based on requests for data objects with keys having a second prefix different from the first prefix.

19. The system of claim 17, further comprising the set of cache nodes, wherein the set of cache nodes is configured to pass, to the set of persistent storage nodes, a randomly-selected statistically significant portion of requests for data objects cached by the set of cache nodes without sending corresponding verification requests to the first witness service instance, wherein the pass-through requests comprise the randomly-selected statistically significant portion of requests.

20. The system of claim 17, further comprising the first witness service instance, wherein the first witness service instance is configured to maintain the second request rate data in a priority queue, wherein a first portion of the second request rate data is associated with the first partition, and wherein a second portion of the second request rate data is associated with a second partition of the plurality of partitions.
