

# US Patent & Trademark Office

## Patent Public Search | Text View

---

United States Patent	12386557
Kind Code	B2
Date of Patent	August 12, 2025
Inventor(s)	Saluja; Mohinder

---

### Non-volatile storage device offloading of host tasks

---

#### Abstract

Various implementations relate to receiving, by a first non-volatile memory device from a host, a host command including device context information of a plurality of non-volatile memory devices. The device context includes an address of a buffer of each of the plurality of non-volatile memory devices, in response to receiving the host command. The first non-volatile memory device divides portions of host data corresponding to the host command among the plurality of non-volatile memory devices. The first non-volatile memory device sends to the host a transfer request indicating transfer of each of the portions of the host data to a respective one of the plurality of non-volatile memory devices. The first non-volatile memory device sends to each of the plurality of non-volatile memory devices other than the first non-volatile memory device, a peer command based on the device context information.

---

<b>Inventors:</b>	<b>Saluja; Mohinder (Cupertino, CA)</b>
<b>Applicant:</b>	<b>Kioxia Corporation (Tokyo, JP)</b>
<b>Family ID:</b>	<b>1000008747579</b>
<b>Assignee:</b>	<b>KIOXIA CORPORATION (Tokyo, JP)</b>
<b>Appl. No.:</b>	<b>17/955014</b>
<b>Filed:</b>	<b>September 28, 2022</b>

#### Prior Publication Data

<b>Document Identifier</b>	<b>Publication Date</b>
US 20240103765 A1	Mar. 28, 2024

---

#### Publication Classification

**Int. Cl.:** G06F3/06 (20060101)

U.S. Cl.:

CPC        **G06F3/0659** (20130101); **G06F3/0638** (20130101); **G06F3/0679** (20130101);

Field of Classification Search

CPC:        G06F (3/0638); G06F (3/0659); G06F (3/0679)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
9395921	12/2015	Lu	N/A	G06F 3/0613
2010/0293440	12/2009	Thatcher et al.	N/A	N/A
2015/0142996	12/2014	Lu	N/A	N/A
2015/0143031	12/2014	Lu	N/A	N/A
2016/0306581	12/2015	Belgaied et al.	N/A	N/A
2017/0147265	12/2016	Kingdon et al.	N/A	N/A
2019/0205244	12/2018	Smith	N/A	N/A
2019/0310913	12/2018	Helmick et al.	N/A	N/A
2020/0042215	12/2019	Roberts	N/A	G06F 3/0634
2021/0397373	12/2020	Muthiah	N/A	G06F 3/0604
2023/0418490	12/2022	Muthiah et al.	N/A	N/A

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
201727653	12/2016	TW	N/A
201732612	12/2016	TW	N/A

OTHER PUBLICATIONS

Non-Final Office Action in U.S. Appl. No. 17/955,040 dated Feb. 15, 2024. cited by applicant  
Final Office Action issued in related U.S. Appl. No. 17/955,040, dated Aug. 30, 2024. cited by applicant  
Non-Final Office Action in U.S. Appl. No. 17/955,040 dated Jan. 3, 2025. cited by applicant

Primary Examiner: Ruiz; Aracelis  
Attorney, Agent or Firm: Foley & Lardner LLP

Background/Summary

TECHNICAL FIELD

(1) The present disclosure generally relates to systems, methods, and non-transitory processor-readable media for data processing using multiple non-volatile memory devices.

BACKGROUND

(2) A general system that provides data storage can include a host coupled to multiple non-volatile memory devices via a bus such as a Peripheral Component Interconnect Express (PCIe) bus. The host can include a processing unit such as a Central Processing Unit (CPU) coupled a memory unit

such as a Dynamic Random Access Memory (DRAM). The CPU is coupled to the bus via a PCIe root complex. Redundant Array of Inexpensive Drives (RAID) can be implemented on the non-volatile memory devices to achieve protection from drive failures.

## SUMMARY

(3) Various arrangements disclosed herein relate to systems, methods, apparatuses, and non-transitory processor-readable media for receiving, by a first non-volatile memory device from a host, a host command including device context information of a plurality of non-volatile memory devices. The device context includes an address of a buffer of each of the plurality of non-volatile memory devices, in response to receiving the host command. The first non-volatile memory device divides portions of host data corresponding to the host command among the plurality of non-volatile memory devices. The first non-volatile memory device sends to the host a transfer request indicating transfer of each of the portions of the host data to a respective one of the plurality of non-volatile memory devices. The first non-volatile memory device sends to each of the plurality of non-volatile memory devices other than the first non-volatile memory device, a peer command based on the device context information.

---

## Description

### BRIEF DESCRIPTION OF THE FIGURES

(1) FIG. 1 shows a block diagram of examples of a system including non-volatile memory devices and a host, according to some implementations.

(2) FIGS. 2A, 2B, 2C, and 2D are schematic diagrams illustrating an example replication method according to various arrangements.

(3) FIG. 2E is a flowchart diagram illustrating the replication method according to various arrangements.

(4) FIGS. 3A, 3B, 3C, 3D, 3E, and 3F are schematic diagrams illustrating an example propagation method according to various arrangements.

(5) FIG. 3G is a flowchart diagram illustrating the propagation method according to various arrangements.

(6) FIGS. 4A, 4B, and 4C are schematic diagrams illustrating an example division method according to various arrangements.

(7) FIG. 4D is a flowchart diagram illustrating the division method according to various arrangements.

(8) FIG. 5A is a flowchart diagram illustrating an example method for group leader-based computation according to various arrangements.

(9) FIG. 5B is a diagram illustrating grouping non-volatile memory devices according to various arrangements.

### DETAILED DESCRIPTION

(10) The host includes a CPU and memory device such as DRAM. The host can be coupled to multiple non-volatile memory devices via a bus (e.g., PCIe bus) and PCIe root complex. The PCIe root complex provides the CPU of the host with knowledge of addresses of all non-volatile memory devices attached to the bus. The CPU can explicitly instruct the non-volatile memory devices to perform various operations. Non-trivial data communications between the CPU of the host and a non-volatile memory device can be performed using the memory device of the host. Such non-trivial data communications have to traverse across the interface (e.g., the bus and the PCIe root complex). Conventionally, the non-volatile memory devices do not communicate with each other, and for a write operations, the CPU of the host issues a write command to each of the non-volatile memory devices. In response to the write command, each of the non-volatile memory devices performs a DMA transfer from the same DRAM of the host to write the data to the non-

volatile storage (e.g., flash media) of each of the non-volatile memory devices. Each time a non-volatile memory device performs a DMA transfer, data has to pass across the PCIe root complex and the bus.

(11) The controller of a non-volatile memory device has processing capabilities (e.g., a CPU), a Direct Memory Access (DMA) engine, and memory, among other elements. Conventional controllers of non-volatile memory devices are not exposed to the host, and the non-volatile memory devices are not directly addressable.

(12) The arrangements disclosed herein allow the non-volatile memory devices to be directly addressable by the host and other devices. In some examples, the Controller Memory Buffer (CMB), Subsystem Level Memory (SLM), or Persistent Memory Region (PMR) can be implemented to allow the non-volatile memory devices to be directly addressable by the host and other devices. Such implementations allow improved autonomy to the non-volatile memory devices and relieve the CPU of the host of some of the burdens of managing the non-volatile memory. The CPU costs on the host side to manage the non-volatile memory devices can be accordingly reduced, and bottlenecks in the data path can be reduce.

(13) In some arrangements, the processor (e.g., CPU) of the host, being configured by device driver code running thereon, issues a command to at least one of multiple non-volatile memory devices. The command includes device context information (e.g., a set of device contexts) that instructs one or more of the at least one of multiple non-volatile memory devices on how to operate other ones of the multiple non-volatile memory devices. In some examples, the device context information can specify one or more methods for contacting at least one of the multiple non-volatile memory devices. The multiple non-volatile memory devices may be parts of a RAID protection scheme.

(14) FIG. 1 shows a block diagram of a system including non-volatile memory devices **100a**, **100b**, **100n** (collectively, non-volatile memory devices **100**) coupled to a host **101** according to some examples. The host **101** can be a user device operated by a user or an autonomous central controller of the non-volatile memory devices **100**, where the host **101** and non-volatile memory devices **100** correspond to a storage subsystem or storage appliance. Examples of such a storage subsystem or appliance include an All Flash Array (AFA) or a Network Attached Storage (NAS) device. As shown, the host **101** includes a memory **102**, a processor **104**, and an bus **106**. The processor **104** is operatively coupled to the memory **102**. The processor **104** and the memory **102** are operatively coupled to the bus **106** through a root complex **105**. The processor **104** is sometimes referred to as a CPU of the host **101**, and configured to perform processes of the host **101** as described herein.

(15) The memory **102** is a local memory of the host **101**. In some examples, the memory **102** is or includes a buffer, sometimes referred to as a host buffer. In some examples, the memory **102** includes at least one of a volatile storage or non-volatile persistent storage. Examples of the memory **102** include but are not limited to, Random Access Memory (RAM), DRAM, Static RAM (SRAM), Magnetic RAM (MRAM), Phase Change Memory (PCM), and so on.

(16) The bus **106** includes one or more of software, firmware, and hardware that allows the processor **104**, network cards, storage devices, the memory **102**, graphic cards, the non-volatile memory devices **100**, and so on to communicate with each other. In some examples, the non-volatile memory devices **100** are directly attached or communicably coupled to the bus **106**. The bus **106** is one or more of a serial, a PCIe bus or network, an internal switch (e.g., PCIe switch), and so on.

(17) In some examples, the root complex **105** (e.g., PCIe root complex) includes one or more of software, firmware, and hardware that connect to the memory **102** and the processor **104** to the bus **106**. In that regard, the root complex **105** can include at least one of one or more controllers, one or more physical connectors, one or more data transfer protocols including namespaces, one or more ports, one or more switches, one or more bridges, one or more transport mechanisms, connectivity thereof, and so on. The root complex **105** can create transaction requests for operation tasks of the processor **104** and send the same to the non-volatile memory devices **100** via the bus **106** according

to the addresses of the non-volatile memory devices **100** on the bus **106**. In some examples, the root complex **105** can be implemented on the hardware (e.g., chip) of the processor **104**. In some examples, the root complex **105** and the bus **106** can be collectively referred to as the interface **140** between the host processor **104**/memory **102** and the non-volatile memory devices **100**.

(18) During start up, the root complex **105** scans the bus **106** for any attached devices (e.g., physically connected or connected via a network such as a network fabric) and obtain the device addresses of the non-volatile memory devices **100**, the processor **104**, and the memory **102**. In some examples, the root complex **105** scans the bus **106** also for the buffer **112** on the non-volatile memory devices **100**. The non-volatile memory devices **100**, the buffers **112**, and the memory **102** are each assigned an address space within the logical address space of the processor **104**. In some examples, SLM and PMR namespaces can be used for addressing the buffers **112**. Accordingly, the processor **104** can perform operations such as read and write using the logical address space. The addresses of the buffers **112** are therefore exposed to the processor **104** and the non-volatile memory devices **100**. Other methods of exposing the addresses of the buffers **112**, such as memory map (e.g., memory-mapped Input/Output (I/O) space) can be likewise implemented. The memory-mapped I/O space allows any memory coupled to the bus **106** to be mapped to a address recognizable by the processor **104**.

(19) The processor **104** can execute an Operating System (OS), which provides a filesystem and applications which use the filesystem. The processor **104** can communicate with the non-volatile memory devices **100** (e.g., a controller **110** of each of the non-volatile memory devices **100**) via a communication link or network. In that regard, the processor **104** can send data to and receive data from one or more of the non-volatile memory devices **100** using the interface **140** to the communication link or network. While the connection between the host **101** and the non-volatile memory devices **100** is shown as a direct link, in some implementations the link may comprise a network fabric which may include networking components such as bridges and switches.

(20) To send and receive data, the processor **104** (the software or filesystem run thereon) communicates with the non-volatile memory devices **100** using a storage data transfer protocol running on the interface **140**. Examples of the protocol include but is not limited to, the SAS, Serial ATA (SATA), and NVMe protocols. In some examples, the interface **140** includes hardware (e.g., controllers) implemented on or operatively coupled to the bus **106**, the non-volatile memory devices **100** (e.g., the controllers **110**), or another device operatively coupled to the bus **106** and/or the non-volatile memory device **100** via one or more suitable networks. The interface **140** and the storage protocol running thereon also includes software and/or firmware executed on such hardware.

(21) In some examples the processor **104** can communicate, via the bus **106** and a network interface, with a communication network. Other host systems or non-volatile memory devices attached or communicably coupled to the communication network can communicate with the host **101** using a suitable network storage protocol, examples of which include, but are not limited to, NVMe over Fabrics (NVMeoF), Internet Small Computer System Interface (iSCSI), Fibre Channel (FC), Network File System (NFS), Server Message Block (SMB), and so on. The network interface allows the software (e.g., the storage protocol or filesystem) running on the processor **104** to communicate with the external hosts attached to the communication network via the bus **106**. In this manner, network storage commands may be issued by the external hosts and processed by the processor **104**, which can issue storage commands to the non-volatile memory devices **100** as needed. Data can thus be exchanged between the external hosts and the non-volatile memory devices **100** via the communication network.

(22) In some examples, the non-volatile memory devices **100** are located in a datacenter (not shown for brevity). The datacenter may include one or more platforms or rack units, each of which supports one or more storage devices (such as but not limited to, the non-volatile memory devices **100**). In some implementations, the host **101** and non-volatile memory devices **100** together form a

storage node, with the host **101** acting as a node controller. One or more storage nodes within a platform are connected to a Top of Rack (TOR) switch, each storage node connected to the TOR via one or more network connections, such as Ethernet, Fiber Channel or InfiniBand, and can communicate with each other via the TOR switch or another suitable intra-platform communication mechanism. In some implementations, non-volatile memory devices **100** may be network attached storage devices (e.g. Ethernet SSDs) connected to the TOR switch, with host **101** also connected to the TOR switch and able to communicate with the non-volatile memory devices **100** via the TOR switch. In some implementations, at least one router may facilitate communications among the non-volatile memory devices **100** in storage nodes in different platforms, racks, or cabinets via a suitable networking fabric. Examples of the non-volatile memory devices **100** include non-volatile devices such as but are not limited to, Solid State Drive (SSDs), Ethernet attached SSDs, Non-Volatile Dual In-line Memory Modules (NVDIMMs), a Universal Flash Storage (UFS), Secure Digital (SD) devices, Compute Express Link (CXL) devices, and so on.

(23) Each of the non-volatile memory devices **100** includes at least a controller **110** and a memory array **120**. Other components of the non-volatile memory devices **100** are not shown for brevity. The memory array **120** includes NAND flash memory devices **130a-130n**. Each of the NAND flash memory devices **130a-130n** includes one or more individual NAND flash dies, which are NVM capable of retaining data without power. Thus, the NAND flash memory devices **130a-130n** refer to multiple NAND flash memory devices or dies within the flash memory device **100**. Each of the NAND flash memory devices **130a-130n** includes one or more dies, each of which has one or more planes. Each plane has multiple blocks, and each block has multiple pages.

(24) While the NAND flash memory devices **130a-130n** are shown to be examples of the memory array **120**, other examples of non-volatile memory technologies for implementing the memory array **120** include but are not limited to, non-volatile (battery-backed) DRAM, Magnetic Random Access Memory (MRAM), Phase Change Memory (PCM), Ferro-Electric RAM (FeRAM), and so on. The arrangements described herein can be likewise implemented on memory systems using such memory technologies and other suitable memory technologies.

(25) Examples of the controller **110** include but are not limited to, an SSD controller (e.g., a client SSD controller, a datacenter SSD controller, an enterprise SSD controller, and so on), a UFS controller, or an SD controller, and so on.

(26) The controller **110** can combine raw data storage in the plurality of NAND flash memory devices **130a-130n** such that those NAND flash memory devices **130a-130n** function logically as a single unit of storage. The controller **110** can include processors, microcontrollers, a buffer memory (e.g., buffer **112**, **114**, and/or **116**), error correction systems, data encryption systems, Flash Translation Layer (FTL) and flash interface modules. Such functions can be implemented in hardware, software, and firmware or any combination thereof. In some arrangements, the software/firmware of the controller **110** can be stored in the memory array **120** or in any other suitable computer readable storage medium.

(27) The controller **110** includes suitable processing and memory capabilities for executing functions described herein, among other functions. As described, the controller **110** manages various features for the NAND flash memory devices **130a-130n** including but not limited to, I/O handling, reading, writing/programming, erasing, monitoring, logging, error handling, garbage collection, wear leveling, logical to physical address mapping, data protection (encryption/decryption, Cyclic Redundancy Check (CRC)), Error Correction Coding (ECC), data scrambling, and the like. Thus, the controller **110** provides visibility to the NAND flash memory devices **130a-130n**.

(28) The buffer memory is a memory device local to, and operatively coupled to, the controller **110**. For instance, the buffer memory can be an on-chip SRAM memory located on the chip of the controller **110**. In some implementations, the buffer memory can be implemented using a memory device of the storage device **110** external to the controller **110**. For instance, the buffer memory can

be DRAM located on a chip other than the chip of the controller **110**. In some implementations, the buffer memory can be implemented using memory devices both internal and external to the controller **110** (e.g., both on and off the chip of the controller **110**). For example, the buffer memory can be implemented using both an internal SRAM and an external DRAM, which are transparent/exposed and accessible by other devices via the bus **106**, such as the host **101** (with the assistance of the root complex **105**) and other non-volatile memory devices **100**. In this example, the controller **110** includes an internal processor that uses memory addresses within a single address space and the memory controller, which controls both the internal SRAM and external DRAM, selects whether to place the data on the internal SRAM and an external DRAM based on efficiency. In other words, the internal SRAM and external DRAM are addressed like a single memory. The buffer memory includes at least one of the buffer **112**, the write buffer **114**, or the read buffer **116**.

(29) The controller **110** includes a buffer **112**, which is sometimes referred to as a drive buffer, an example of which can be a CMB, SLM, or PMR. Besides being accessible by the controller **110**, the buffer **112** can be accessible by other devices, such as the host **101** and other non-volatile memory devices **100a**, **100b**, . . . **100n**, via the bus **106**. In that manner, the buffer **112** (e.g., addresses of memory locations within the buffer **112**) is exposed across the bus **106**, and any device operatively coupled to the bus **106** can issue commands (e.g., read commands, write commands, and so on) using addresses that correspond to memory locations within the buffer **112** in order to read data from those memory locations within the buffer **112** and write data to those memory locations within the buffer **112**. In some examples, the buffer **112** is a volatile storage. In some examples, the buffer **112** is a non-volatile persistent storage, which may offer improvements in protection against unexpected power loss of one or more of the non-volatile memory devices **100**. Examples of the buffer **112** include but are not limited to, RAM, DRAM, SRAM, MRAM, PCM, and so on. The buffer **112** may refer to multiple buffers each configured to store data of a different type, as described herein.

(30) In some implementations, as shown in FIG. **1**, the buffer **112** is a local memory of the controller **110**. For instance, the buffer **112** can be an on-chip SRAM memory located on the chip of the controller **110**. In some implementations, the buffer **112** can be implemented using a memory device of the storage device **110** external to the controller **110**. For instance, the buffer **112** can be DRAM located on a chip other than the chip of the controller **110**. In some implementations, the buffer **112** can be implemented using memory devices both internal and external to the controller **110** (e.g., both on and off the chip of the controller **110**). For example, the buffer **112** can be implemented using both an internal SRAM and an external DRAM, which are transparent/exposed and accessible by other devices **100** via the bus **106**, such as the host **101** and other non-volatile memory devices **100**. In this example, the controller **110** includes an internal processor uses memory addresses within a single address space and the memory controller, which controls both the internal SRAM and external DRAM, selects whether to place the data on the internal SRAM and an external DRAM based on efficiency. In other words, the internal SRAM and external DRAM are addressed like a single memory.

(31) In one example concerning a write operation, in response to receiving data from the host **101** (via the host interface **140**), the controller **110** acknowledges the write commands to the host **101** after writing the data to a write buffer **114**. In some implementations the write buffer **114** may be implemented in a separate, different memory than the buffer **112**, or the write buffer **114** may be a defined area or part of the memory comprising buffer **112**, where only the CMB, SLM, or PMR part of the memory is accessible by other devices, but not the write buffer **114**. The controller **110** can write the data stored in the write buffer **114** to the memory array **120** (e.g., the NAND flash memory devices **130a-130n**). Once writing the data to physical addresses of the memory array **120** is complete, the FTL updates mapping between logical addresses (e.g., Logical Block Address (LBAs)) used by the host **101** to associate with the data and the physical addresses used by the

controller **110** to identify the physical locations of the data. In another example concerning a read operation, the controller **110** includes another buffer **116** (e.g., a read buffer) different from the buffer **112** and the buffer **114** to store data read from the memory array **120**. In some implementations the read buffer **116** may be implemented in a separate, different memory than the buffer **112**, or the read buffer **116** may be a defined area or part of the memory comprising buffer **112**, where only the CMB, SLM, or PMR part of the memory is accessible by other devices, but not the read buffer **116**.

(32) While non-volatile memory devices (e.g., the NAND flash memory devices **130a-130n**) are presented as examples herein, the disclosed schemes can be implemented on any storage system or device that is connected to the host **101** over an interface, where such system temporarily or permanently stores data for the host **101** for later retrieval.

(33) In some examples, the non-volatile memory devices **100** form a RAID group for parity protection. That is, one or more of the non-volatile memory devices **100** stores parity data (e.g., parity bits) for data stored on those devices and/or data stored on other ones of the non-volatile memory devices **100**.

(34) In some implementations, the processor **104** of the host **101** sends a host command to a first non-volatile memory device (e.g., the non-volatile memory device **100a**) of the non-volatile memory devices **100**, where the host command includes device context information for at least another one of the non-volatile memory devices **100** in the set. The first non-volatile memory device can be referred to as the source device. In addition to executing the host command, the first non-volatile memory device is also responsible for executing this command on at least another one of the non-volatile memory devices **100**, while containing all further communication within the PCIe subsystem without reaching the processor **104** or the memory **102**.

(35) FIGS. 2A, 2B, 2C, and 2D are schematic diagrams illustrating an example replication method **200** according to various arrangements. FIG. 2E is a flowchart diagram illustrating the replication method **200** according to various arrangements. FIGS. 2A-2E show elements of the non-volatile memory devices **100** and the host **101** in FIG. 1, with certain elements omitted in FIGS. 2A-2E for clarity. In some examples, the method **200** relates to a replication operation where data from the host is replicated across multiple non-volatile memory devices **100a-100n**. The method **200** shown in FIG. 2E can be performed by the non-volatile memory device **100a**.

(36) At **210**, the controller **110** of a first non-volatile memory device (e.g., the non-volatile memory device **100a**) receives a host command from the host **101**, via the interface **140** including the root complex **105** and the bus **106**. In some examples, the host command is a write command (e.g., a first write command) addressed to the non-volatile memory device **100a** (e.g., its buffer **112**). The root complex **105** can route the host command to the controller **110** of the non-volatile memory device **100a** on the bus **106** using the address of the non-volatile memory device **100a** (e.g., its buffer **112**). The host command includes an address of the memory **102**. The address of the memory **102** of the host **101** includes a buffer address, an address descriptor, an identifier, a pointer, or another suitable indicator that identifies the memory **102** of the host **101**.

(37) The host command includes the device context information, which can be a data structure including various information related to at least one other non-volatile memory device (e.g., the addresses of the non-volatile memory devices **100b-100n**), referred to as a second non-volatile memory device. The second non-volatile memory device can be referred to as a source non-volatile memory device. In some examples, the device context information can include an address of second non-volatile memory device.

(38) In some examples, the device context information can include an address of the buffer **112** of the second non-volatile memory device. The address of the buffer **112** can be a CMB address, a SLM address, a PMR address, an address descriptor, an identifier, a pointer, or another suitable indicator that identifies the buffer **112** of that non-volatile memory device. The addresses of the buffers **112** of the non-volatile memory devices **100** are stored within a shared address register



(e.g., a shared PCIe Base Address Register) known to the host **101** through the root complex **105**. The addresses can be used to send data, instructions, commands, and so on the bus **106**. For example, in NVMe, the CMB is defined by a NVMe controller register Controller Memory Buffer Location (CMBLOC) which defines the PCI address location of the start of the CMB and a controller register Controller Memory Buffer Size (CMBSZ) which defines the size of the CMB. The SLM address and the PMR address can be similarly implemented. These controller registers can reside in the root complex **105**, the bus **106**, or another suitable entity coupled to the root complex **105**.

(39) In some examples, the device context information can include permission information of at least one of the second non-volatile memory device or its buffer **112**. The permission information of at least one of the second non-volatile memory device or its buffer **112** includes information related to whether or how the first non-volatile memory device can access at least one of the second non-volatile memory device or its buffer **112**. For example, the permission information can include one or more types of operations or tasks (e.g., read, write, replication, propagation, division, reduction, etc.) for which the at least one of the second non-volatile memory device or its buffer **112** can be accessed. The permission can include authentication credentials used to access the at least one of the second non-volatile memory device or its buffer **112**, such that after the first non-volatile memory device provides the authentication credentials (along with any commands or instructions described herein) to the at least one of the second non-volatile memory device or its buffer **112**, the second non-volatile memory device can perform operations according to the commands or instructions. In some examples, the authentication credentials can be implemented using a security token. The security token can be separately issued for different operations and tasks. In some examples, the authentication credentials can be implemented using an access control list (e.g., a whitelist, a blacklist, or a combination thereof) that specifies whether the first non-volatile memory device can access the second non-volatile memory device, and whether the first non-volatile memory device can access the second non-volatile memory device for certain operations and tasks.

(40) In some examples, the device context information can include priority information for processing the data associated with the host command by the second non-volatile memory device. For example, the priority information can indicate a priority level associated with the data, so that the second non-volatile memory device can prioritize processing the data corresponding to the host command before another data if the data corresponding to the host command has higher priority level than that of the another data (e.g., processing the data associated with the host command faster). The second non-volatile memory device can prioritize processing the another data before processing the data corresponding to the host command if the data corresponding to the host command has a lower priority level than that of the another data (e.g., processing the data associated with the host command slower). In some examples, the priority information can indicate a priority level for the first non-volatile memory device to communicate with the second non-volatile memory device with respect to the data in connection with the host command. For instance, the first non-volatile memory device can prioritize processing the data corresponding to the host command (e.g., sending out the peer command as described herein) before processing another data (e.g., sending out the peer command for the another data) if the data corresponding to the host command has higher priority level than that of the another data (e.g., processing the data associated with the host command faster). The first non-volatile memory device can prioritize processing the another before processing the data corresponding to the host command if the data corresponding to the host command has a lower priority level than that of the another data (e.g., processing the data associated with the host command slower).

(41) At **220**, the controller **110** of the non-volatile memory device **100a** transfers data corresponding to the host command from the host to the buffer **112** of the non-volatile memory device **100a** via the interface **140**. For example, the controller **110** of the non-volatile memory

device **100a** can perform a DMA transfer by issuing a DMA request to the host **101** to transfer the data stored in the memory **102** (e.g., a host buffer) of the host **101** to the buffer **112**, where the DMA request includes the address of the memory **102**. The PCI root complex **105** can route the DMA request to the memory **102** on the bus **106** using the address of the memory **102**.

(42) At **230**, the controller **110** of the non-volatile memory device **100a** writes the data from the buffer **112** of the non-volatile memory device **100a** to the non-volatile memory (e.g., the memory array **120**) of the non-volatile memory device **100a**. For example, the controller **110** of the non-volatile memory device **100a** issues an internal command to write the buffer **112** to its own flash media (e.g., the memory array **120**). In some examples, the non-volatile memory device **100a** writes the data into the memory array **120** at the same time that other non-volatile memory devices **100b-100n** write the data into their respective memory arrays **120**, e.g., at **260**. In some examples, **230** can occur immediately in response to **220**, at the same time as **240**, after **240**, at the same time as **250**, or after **250**. In some examples, **230** can occur any time before the data is deleted or erased from the buffer **112**.

(43) At **240**, the controller **110** of the non-volatile memory device **100a** sends a peer command to a second non-volatile memory device (e.g., the non-volatile memory devices **100b-100n**) based on the device context information, via the bus **106**.

(44) For example, the controller **110** of the non-volatile memory device **100a** can send the peer command to the controller **110** of each of the non-volatile memory devices **100b-100n** according to (e.g., at) the address of each of the non-volatile memory devices **100b-100n** contained in the device context information received from the host **101** at **210**.

(45) For example, the controller **110** of the non-volatile memory device **100a** can determine, according to the permission information, whether there is permission for the non-volatile memory device **100a** to send the peer command to at least one of each of the non-volatile memory devices **100b-100n** or its buffer **112** for the type of operation or task (e.g., write or replication) specified in the host command. For example, the controller **110** of the non-volatile memory device **100a** can authenticate, using the authentication credentials, with at least one of each of the non-volatile memory devices **100b-100n** or its buffer.

(46) For example, the controller **110** of the non-volatile memory device **100a** can determine, according to the priority information, the priority level for sending the peer command to each of the non-volatile memory devices **100b-100n**. The controller **110** of the non-volatile memory device **100a** processes the data or the host command associated higher priority level before processing data or host command associated with a lower priority level. For example, the peer command can include a priority level for processing the data associated with the peer command and the host command. Each of the non-volatile memory devices **100b-100n** can process the data or the peer command associated higher priority level before processing data or peer command associated with a lower priority level. The priority level included in the peer command can correspond to the priority level included in the context information of the host command.

(47) The peer command includes another write command (e.g., a second write command). In some examples in which a set of non-volatile memory devices (e.g., RAID-1 devices, such as the non-volatile memory devices **100a-100n**) includes  $n$  devices, the controller **110** of the non-volatile memory device **100a** issues  $n-1$  peer commands to  $n-1$  non-volatile memory devices **100b-100n**. In some examples, the peer command includes an address of the data from which the second non-volatile memory device (e.g., the non-volatile memory devices **100b-100n**) can transfer the data. The address of the data includes the address of the buffer **112** of the non-volatile memory device **100a**, so the second non-volatile memory device (e.g., the non-volatile memory devices **100b-100n**) can transfer the data from the buffer **112** of the non-volatile memory device **100a** using the address thereof, e.g., at **260**.

(48) At **250**, the controller **110** of the non-volatile memory device **100a** participates transferring the data to the buffer **112** of the second non-volatile memory device from the buffer **112** of the non-

volatile memory device **100a**, via the bus **106**. For example, the controller **110** of the non-volatile memory device **100a** exposes its buffer **112** (e.g., the buffer address) to the bus **106** (e.g., to any device such as the non-volatile memory devices **100b-100n** coupled or connected to the bus **106**). The controller **110** of the second non-volatile memory device (e.g., each of the non-volatile memory devices **100b-100n**) perform a DMA transfer by issuing a DMA request to the controller **110** of the non-volatile memory device **100a** using the address of the buffer **112** of the non-volatile memory device **100a**. In response, the controller **110** of each of the non-volatile memory devices **100b-100n** transfers the data stored in the buffer **112** of the non-volatile memory device **100a** into the buffer **112** of each of the non-volatile memory devices **100b-100n**.

(49) At **260**, the controller **110** of the second non-volatile memory device (e.g., each of the non-volatile memory devices **100b-100n**) writes the data from the buffer **112** of the second non-volatile memory device to the non-volatile memory (e.g., the memory array **120**) of the second non-volatile memory device. For example, the controller **110** of the second non-volatile memory device issues an internal command to write the buffer **112** to its own flash media (e.g., the memory array **120**). In some examples, the non-volatile memory devices **100b-100n** write the data into their respective memory arrays **120** at the same time, in parallel. In some examples, in response to the controller **110** of the non-volatile memory device **100a** receiving write completion confirmation from each of the non-volatile memory devices **100b-100n**, the controller **110** of the non-volatile memory device **100a** sends a write complete message to the processor **104** across the interface **140**.

(50) Accordingly, following **210** and **220**, for the transfer of data in replication to the non-volatile memory devices **100b-100n**, no requests, commands, or data crosses the root complex **105** to the processor **104** and memory **102** of the host **101**. The routing of requests and data among the non-volatile memory devices **100** can be performed via the bus **106** as the root complex **105**, which provides an interface to the bus **106** for the processor **104** and memory **102**, is not involved. The method **200** can reduce workload of the processor **104** and the memory **102** by delegating repetitive operations across a set of devices **100**. The amount of data passing across the root complex **105** can be reduced. The load on system memory **102** can also be reduced.

(51) In some examples, the non-volatile memory devices **100** can be SSDs. In some examples, at least one or two or more of the non-volatile memory devices **100** can be NVMeoF devices, which are attached to the bus **106** via a communication network such as NVMeoF, iSCSI, FC, NFS, SMB, and so on. The NVMeoF devices may have peer-to-peer data transfer capability on NVMeoF target side, without routing communications through the root complex **105**.

(52) FIGS. **3A**, **3B**, **3C**, **3D**, **3E**, and **3F** are schematic diagrams illustrating an example propagation method **300** according to various arrangements. FIG. **3G** is a flowchart diagram illustrating the propagation method **300** according to various arrangements. FIGS. **3A-3G** show elements of the non-volatile memory devices **100** and the host **101** in FIG. **1**, with certain elements omitted for clarity. In some examples, the method **300** relates to a propagation operation where the multiple non-volatile memory devices **100a-100n** are part of a RAID-5 protection scheme and are a part of a stripe spanning across the non-volatile memory devices **100a-100n**. A write update updates a part of the strip that is entirely within the memory array **120** of the non-volatile memory device **100a**. The method **300** shown in FIG. **3G** can be performed by the non-volatile memory device **100a**.

(53) At **310**, the controller **110** of a first non-volatile memory device (e.g., the non-volatile memory device **100a**) receives a host command from the host **101**, via the interface **140** including the root complex **105** and the bus **106**. In some examples, the host command is a write command (e.g., a first write command) addressed to the non-volatile memory device **100a** (e.g., its buffer **112**). The root complex **105** can route the host command to the controller **110** of the non-volatile memory device **100a** on the bus **106** using the address of the non-volatile memory device **100a** (e.g., its buffer **112**). The host command includes an address of the memory **102**. The address of the memory **102** of the host **101** includes a buffer address, an address descriptor, an identifier, a pointer, or another suitable indicator that identifies the memory **102** of the host **101**.

(54) In some examples, the host command includes the device context information. In some examples, the device context information includes address of at least one other non-volatile memory device (e.g., the address of the parity device, the non-volatile memory device **100n**) in the RAID group. The address of a non-volatile memory device can be a CMB address, SLM address, a PMR address, an address descriptor, an identifier, a pointer, or another suitable indicator that identifies the buffer **112** of that non-volatile memory device, as described. In some examples, the host command includes the logical address of the data to be updated. In some examples, the device context information can include permission information of at least one of the second non-volatile memory device or its buffer **112**. In some examples, the device context information can include priority information including at least one of a priority level for processing the data associated with the host command by the second non-volatile memory device or a priority level for the first non-volatile memory device to communicate with the second non-volatile memory device with respect to the data in connection with the host command.

(55) At **320**, the controller **110** of the non-volatile memory device **100a** transfers new data corresponding to the host command from the host to the buffer **112** of the non-volatile memory device **100a** via the interface **140**. For example, the controller **110** of the non-volatile memory device **100a** can perform a DMA transfer by issuing a DMA request to the host **101** to transfer the new data stored in the memory **102** (e.g., a host buffer) of the host **101** to the buffer **112**, where the DMA request includes the address of the memory **102**. The PCI root complex **105** can route the DMA request to the memory **102** on the bus **106** using the address of the memory **102**.

(56) At **330**, the controller **110** reads old data corresponding to the logical address included in the host command from the non-volatile memory (e.g., the memory array **120**) of the non-volatile memory device **100a** into the buffer **112**. For example, the controller **110** can determine the physical address mapped to the logical address through suitable Logical-to-Physical (L2P) conversation and read the physical address at the non-volatile memory to obtain the old data and transfer the same to the buffer **112**. At **340**, the controller **110** computes result data based on at least one of the new data and the old data. For example, the result data can be a bitwise XOR of the old and the new data. The controller **110** stores the result data into the buffer **112** of the non-volatile memory device **100a**.

(57) This result data is sent to the member device holding the parity for the stripe (e.g., the non-volatile memory device **100n**). At **350**, the controller **110** of the non-volatile memory device **100a** sends a peer command to a second non-volatile memory device (e.g., the non-volatile memory device **100n**) based on the device context information, via the bus **106**. The second non-volatile memory device may be the parity device that stores the parity bit. For example, the controller **110** of the non-volatile memory device **100a** sends the peer command to the controller **110** of the non-volatile memory device **100n** according to (e.g., at) the address of the non-volatile memory device **100n** contained in the device context information received from the host **101**. The bus **106** can route the peer command to the controller **110** of the non-volatile memory device **100n** using the address of the non-volatile memory device **100n** (e.g., its buffer **112**).

(58) For example, the controller **110** of the non-volatile memory device **100a** can determine, according to the permission information, whether there is permission for the non-volatile memory device **100a** to send the peer command to the non-volatile memory device **100n** or its buffer **112** for the type of operation or task (e.g., write or propagation) specified in the host command. For example, the controller **110** of the non-volatile memory device **100a** can authenticate, using the authentication credentials, with at least one of the non-volatile memory device **100n** or its buffer.

(59) For example, the controller **110** of the non-volatile memory device **100a** can determine, according to the priority information, the priority level for sending the peer command to the non-volatile memory device **100n**. The controller **110** of the non-volatile memory device **100a** processes the data or the host command associated higher priority level before processing data or host command associated with a lower priority level. For example, the peer command can include a

priority level for processing the data associated with the peer command and the host command. The non-volatile memory device **100n** can process the data or the peer command associated higher priority level before processing data or peer command associated with a lower priority level. The priority level included in the peer command can correspond to the priority level included in the context information of the host command.

(60) The peer command includes another write command (e.g., a second write command). In some examples, the peer command includes an address of the result data from which the second non-volatile memory device (e.g., the non-volatile memory device **100n**) can transfer the result data. The address of the data includes the address of the buffer **112** of the non-volatile memory device **100a**, so the second non-volatile memory device can transfer the result data from the buffer **112** of the non-volatile memory device **100a** using the address thereof, e.g., at **360**.

(61) At **360**, the controller **110** of the non-volatile memory device **100a** participates transferring the data to the buffer **112** of the second non-volatile memory device from the buffer **112** of the non-volatile memory device **100a**, via the bus **106**. For example, the controller **110** of the non-volatile memory device **100a** exposes its buffer **112** (e.g., the buffer address) to the bus **106** (e.g., to any device such as the non-volatile memory device **100n** coupled or connected to the bus **106**). The controller **110** of the second non-volatile memory device (e.g., the non-volatile memory device **100n**) can perform a DMA transfer by issuing a DMA request to the controller **110** of the non-volatile memory device **100a** using the address of the buffer **112** of the non-volatile memory device **100a**. In response, the controller **110** of the non-volatile memory device **100n** transfers the data stored in the buffer **112** of the non-volatile memory device **100a** into the buffer **112** of the non-volatile memory device **100n**.

(62) The parity is updated to the bitwise XOR of the result data and the old parity data. For example, at **370**, the controller **110** of the non-volatile memory device **100n** reads old parity data from the non-volatile memory (e.g., the memory array **120**) of the non-volatile memory device **100n** into the buffer **112** of the non-volatile memory device **100n**. In some examples, the second write command includes the logical address associated with the new data, and the controller **110** of the non-volatile memory device **100n** determines the physical address mapped to the logical address through suitable L2P conversation and read the physical address at the non-volatile memory to obtain the old parity data and transfer the same to the buffer **112** of the non-volatile memory device **100n**.

(63) At **380**, the controller **110** computes new parity data based on at least one of the result data and the old parity data. For example, the new parity data can be a bitwise XOR of the old parity data and the result data. The controller **110** of the non-volatile memory device **100n** stores the result data into the buffer **112** of the non-volatile memory device **100n**.

(64) At **390**, the controller **110** of the second non-volatile memory device (e.g., the non-volatile memory device **100n**) writes the new parity data from the buffer **112** of the second non-volatile memory device to the non-volatile memory (e.g., the memory array **120**) of the second non-volatile memory device. For example, the controller **110** of the second non-volatile memory device issues an internal command to write the buffer **112** to its own flash media (e.g., the memory array **120**).

(65) At **395**, the controller **110** of the non-volatile memory device **100a** writes the new data from the buffer **112** of the non-volatile memory device **100a** to the non-volatile memory (e.g., the memory array **120**) of the non-volatile memory device **100a**. For example, the controller **110** of the non-volatile memory device **100a** issues an internal command to write the buffer **112** to its own flash media (e.g., the memory array **120**).

(66) In some examples, the non-volatile memory device **100a** writes the new data into the memory array **120** at **395** in parallel with one or more of **360**, **370**, **380**, or **390**. In some examples, **395** can occur immediately in response to **320**, at the same time as **330**, after **330**, at the same time as **340**, or after **340**. In some examples, **395** can occur any time before the new data is deleted or erased from the buffer **112** of the non-volatile memory device **100a**.

(67) In some examples, the propagation method **300** can be repeated for at least one additional device such as a third non-volatile memory device, which can be one of non-volatile memory devices **100b-100n-1**. In that regard, the second non-volatile memory device becomes the first non-volatile memory device, and the third non-volatile memory device becomes the second non-volatile memory device. The method **300** repeats at **330**, with the new data being the data computed at **380**. The method **300** can be repeated as the last device in the set (e.g., the RAID group) is reached.

(68) In some examples, in response to the controller **110** of the non-volatile memory device **100a** receiving write completion confirmation from the last non-volatile memory device in the set in the propagation chain, the controller **110** of the non-volatile memory device **100a** sends a write complete message to the processor **104** across the interface **140**. Accordingly, following **310** and **320**, for the transfer of data in propagation, no requests, commands, or data crosses the root complex **105** to the processor **104** and memory **102** of the host **101**. The routing of requests and data among the non-volatile memory devices **100** can be performed via the bus **106**, as the root complex **105**, which provides an interface to the bus **106** for the processor **104** and memory **102**, is not involved. The method **300** can reduce workload of the processor **104** and the memory **102** by offloading aggregation operations across a set of devices **100**. The amount of data passing across the root complex **105** can be reduced. The load on system memory **102** can also be reduced.

(69) FIGS. **4A**, **4B**, and **4C** are schematic diagrams illustrating an example division method **400** according to various arrangements. FIG. **4D** is a flowchart diagram illustrating the division method **400** according to various arrangements. FIGS. **4A-4D** show elements of the non-volatile memory devices **100** and the host **101** in FIG. **1**, with certain elements omitted for clarity. In some examples, the method **400** relates to a division mechanism where a non-volatile memory device (e.g., the non-volatile memory device **100a**) of the multiple non-volatile memory devices **100a-100n** divide host data among the multiple non-volatile memory devices **100a-100n**. The multiple non-volatile memory devices **100a-100n** are part of a RAID-5 protection scheme and are a part of a stripe spanning across the non-volatile memory devices **100a-100n**. The non-volatile memory device **100n** is the parity device that stores the parity bit for the data of the strip stored in the non-volatile memory devices **100a-100n-1**. The method **400** shown in FIG. **4D** can be performed by the non-volatile memory device **100a**. The work includes transferring host data from buffers **401a**, **401b**, . . . , **401n-1** of the memory **102** into the non-volatile memory devices **100a-100n-1**. The method **400** can reduce the load on processor **104** by allowing the division task to be performed by one of the non-volatile memory devices **100a-100n-1**.

(70) At **410**, the controller **110** of a first non-volatile memory device (e.g., the non-volatile memory device **100a**) receives a host command from the host **101**, via the interface **140** including the root complex **105** and the bus **106**. In some examples, the host command is a write command addressed to the non-volatile memory device **100a** (e.g., its buffer **112**). The root complex **105** can route the host command to the controller **110** of the non-volatile memory device **100a** on the bus **106** using the address of the non-volatile memory device **100a**. The host command includes addresses of the buffers **401a**, **401b**, . . . , **401n-1** of the memory **102**, which corresponds to the work that needs to be performed.

(71) In some examples, the host command includes the device context information of a plurality of non-volatile memory devices. In some examples, the device context information includes addresses of the non-volatile memory devices **100a-100n-1** in the RAID group. The address of a non-volatile memory device can be a CMB address, SLM address, a PMR address, an address descriptor, an identifier, a pointer, or another suitable indicator that identifies the buffer **112** of that non-volatile memory device, as described. In some examples, the host command includes the logical address of the data to be updated, including the logical address (e.g., a buffer address) corresponding to each of the buffers **401a-401n-1**. In some examples, the device context information can include permission information of at least one of the plurality of non-volatile memory devices or their

buffers **112**. In some examples, the device context information can include priority information including at least one of a priority level for processing the data associated with the host command by each of the plurality of non-volatile memory devices or a priority level for the first non-volatile memory device to communicate with each of the plurality of non-volatile memory devices with respect to the data in connection with the host command.

(72) At **420**, the controller **110** of the first non-volatile memory device (e.g., the non-volatile memory device **100a**) divides the host data among non-volatile memory devices and issues transfer requests to transfer the host data into the buffers **112** of the non-volatile memory devices. The non-volatile memory devices can include the non-parity non-volatile memory devices **100a-100n-1** in the RAID group.

(73) In some examples, the controller **110** of the first non-volatile memory device maps each of the addresses of the buffers **401a, 401b, . . . , 401n-1** of the memory **102** to one of the non-parity non-volatile memory devices **100a-100n-1**. In the example in which the number of the buffers **401a, 401b, . . . , 401n-1** is the same as the number of the non-parity non-volatile memory devices **100a-100n-1**, the controller **110** of the first non-volatile memory device maps the addresses of the buffers **401a, 401b, . . . , 401n-1** of the memory **102** to the non-parity non-volatile memory devices **100a-100n-1** according to an one-to-one relationship. For example, the addresses of the buffers **401a, 401b, . . . , 401n-1** of the memory **102** can be mapped to the non-parity non-volatile memory devices **100a-100n-1** according to an increasing order of the buffer addresses or ID and an order of the non-parity non-volatile memory devices **100a-100n-1** in the RAID group, e.g., the buffer **401a** is mapped to non-volatile memory device **100a**, the buffer **401b** is mapped to non-volatile memory device **100b**, . . . , the buffer **401n-1** is mapped to non-volatile memory device **100n-1**, and so on. In some examples, the addresses of the buffers **401a, 401b, . . . , 401n-1** of the memory **102** can be mapped to the non-parity non-volatile memory devices **100a-100n-1** according to a decreasing order of the buffer addresses or ID and the order of the non-parity non-volatile memory devices **100a-100n-1** in the RAID group, e.g., the buffer **401a** is mapped to non-volatile memory device **100n-1**, the buffer **401b** is mapped to non-volatile memory device **100n-2**, . . . , the buffer **401n-1** is mapped to non-volatile memory device **100a**, and so on.

(74) In the example in which the number of the buffers **401a, 401b, . . . , 401n-1** is the less than the number of the non-parity non-volatile memory devices **100a-100n-1**, the controller **110** of the first non-volatile memory device maps the addresses of the buffers of the memory **102** to the non-parity non-volatile memory devices based on the increasing or decreasing order of the buffer addresses or ID and the order of the non-parity non-volatile memory devices **100a-100n-1** in the RAID group, where not all of the non-parity non-volatile memory devices **100a-100n-1** in the RAID group is mapped to one of the buffers **401a, 401b, . . . , 401n-1**.

(75) In the example in which the number of the buffers **401a, 401b, . . . , 401n-1** is the greater than the number of the non-parity non-volatile memory devices **100a-100n-1**, the controller **110** of the first non-volatile memory device maps the addresses of the buffers of the memory **102** to the non-parity non-volatile memory devices based on the increasing or decreasing order of the buffer addresses or ID and the order of the non-parity non-volatile memory devices **100a-100n-1** in the RAID group, where one or more of the non-parity non-volatile memory devices **100a-100n-1** in the RAID group is mapped to two or more of the buffers **401a, 401b, . . . , 401n-1**. The mapping can be performed cyclically, e.g., after each of the non-volatile memory devices **100a-100n-1** in the RAID group is mapped to a buffer, the next unmapped buffer is mapped to the non-volatile memory device **100a**, and progressing to **100b, 100c**, and so on.

(76) The controller **110** of the non-volatile memory device **100a** issues transfer requests to transfer the data corresponding to the host data into the buffers **112** of the non-parity non-volatile memory devices based on the division. For example, the non-volatile memory device **100a** can issue a first DMA request to the host **101** (e.g., to the memory **102**) to transfer the data stored in the buffer **401a** to the buffer **112** of the mapped non-volatile memory device (e.g., the non-volatile memory device

**100a**), issue a second DMA request to the host **101** (e.g., to the memory **102**) to transfer the data stored in the buffer **401b** to the buffer **112** of the mapped non-volatile memory device (e.g., the non-volatile memory device **100b**), . . . , issue an (n-1)th DMA request to the host **101** (e.g., to the memory **102**) to transfer the data stored in the buffer **401n-1** to the buffer **112** of the mapped non-volatile memory device (e.g., the non-volatile memory device **100n-1**). Each transfer request includes the buffer address of one of the buffers **401a-401n-1** and an address of the buffer **112** of the corresponding non-volatile memory device. The data from the buffers **401a-401n-1** can be accordingly transferred to the buffers **112** according to the DMA requests.

(77) At **430**, the controller **110** of the non-volatile memory device **100a** sends a peer command to the plurality of non-volatile memory devices (e.g., each of the non-volatile memory devices **100b-100n-1** other than the non-volatile memory device **100a**) based on the device context information, via the bus **106**. For example, the controller **110** of the non-volatile memory device **100a** sends the peer command to the controller **110** of each of the non-volatile memory devices **100b-100n-1** according to (e.g., at) the address of each of the non-volatile memory devices **100b-100n-1** contained in the device context information received from the host **101**. The bus **106** can route the peer command to the controller **110** of each of the non-volatile memory devices **100b-100n-1** using the address of the each of the non-volatile memory devices **100b-100n-1** (e.g., its buffer **112**).

(78) For example, the controller **110** of the non-volatile memory device **100a** can determine, according to the permission information, whether there is permission for the non-volatile memory device **100a** to send the peer command to each of the non-volatile memory devices **100b-100n-1** or its buffer **112** for the type of operation or task (e.g., write) specified in the host command. For example, the controller **110** of the non-volatile memory device **100a** can authenticate, using the authentication credentials, with each of the non-volatile memory devices **100b-100n-1** or its buffer.

(79) For example, the controller **110** of the non-volatile memory device **100a** can determine, according to the priority information, the priority level for sending the peer command to each of the non-volatile memory devices **100b-100n-1**. The controller **110** of the non-volatile memory device **100a** processes the data or the host command associated higher priority level before processing data or host command associated with a lower priority level. For example, the peer command can include a priority level for processing the data associated with the peer command and the host command. Each of the non-volatile memory devices **100b-100n-1** can process the data or the peer command associated higher priority level before processing data or peer command associated with a lower priority level. The priority level included in the peer command can correspond to the priority level included in the context information of the host command.

(80) The peer command includes another write command (e.g., a second write command) that instructs each of the non-volatile memory devices **100b-100n-1** to write the data corresponding to a respective one of the buffers **401a-401n-1** stored in the buffer **112** to the memory array **120**. The peer command can include the address of the buffer **112** of each respective one of the non-volatile memory devices **100b-100n-1**.

(81) At **440**, the controller **110** of the plurality of non-volatile memory devices **100a-100n-1** writes the data corresponding to respective ones of the buffers **401a-401n-1** stored in the buffers **112** to the memory arrays **120**. For example, the controller **110** of each of the non-volatile memory devices **100a-100n-1** issues an internal command to write its buffer **112** to its own flash media (e.g., the memory array **120**). In some examples, the non-volatile memory devices **100a-100n-1** write the data into the memory arrays **120** at **440** in parallel.

(82) In some examples, a plurality of non-volatile memory devices can be grouped together, and a leader device can be assigned. In some examples, the grouped plurality of non-volatile memory devices corresponds to a group of non-volatile memory devices **100a-100n** that are part of the same RAID group. In some examples, the grouped plurality of non-volatile memory devices corresponds to a group of non-parity non-volatile memory devices **100a-100n-1** that are part of the same RAID



group. The leader device is the non-volatile memory device **100a** in the method **400**, shown in FIGS. **4A-4D**.

(83) In some examples, the grouped plurality of non-volatile memory devices can be determined according to any other methods as described herein. FIG. **5A** is a flowchart diagram illustrating an example method **500** for group leader-based computation, according to various arrangements. FIG. **5B** is a diagram illustrating grouping non-volatile memory devices, according to various arrangements. In FIG. **5B**, each the non-volatile memory devices **501a**, **501b**, **501c**, **501d**, **501e**, **501f**, **501g**, and **501h** can be a non-volatile memory device such as one of the non-volatile memory devices **100a**, **100b**, **100n**. For example, each of the non-volatile memory devices **501a**, **501b**, **501c**, **501d**, **501e**, **501f**, **501g**, and **501h** can include a controller **110**, a buffer **112**, and a memory array **120**. The method **500** can be performed by the controller **110** of one or more of the non-volatile memory devices **501a**, **501b**, **501c**, **501d**, **501e**, **501f**, **501g**, and **501h**, the processor **104**, or another suitable entity with processing capabilities.

(84) At **510**, the plurality of non-volatile memory devices are grouped into at least one group (e.g., at least one first group). For example, the  $n$  non-volatile memory devices, which may be part of a RAID group or are non-parity non-volatile memory devices of a RAID group, can be grouped according to an identifier (e.g.,  $1-n-1$ ) assigned to each non-volatile memory device. For example, the non-volatile memory device **501a** is assigned identifier 1, the non-volatile memory device **501b** is assigned identifier 2, the non-volatile memory device **501c** is assigned identifier 3, the non-volatile memory device **501d** is assigned identifier 4, the non-volatile memory device **501e** is assigned identifier 5, the non-volatile memory device **501f** is assigned identifier 6, the non-volatile memory device **501g** is assigned identifier 7, and the non-volatile memory device **501h** is assigned identifier 8. Starting from identifier 1, a predetermined number (e.g., 2) of adjacent identifiers are grouped in a same group. For example, the non-volatile memory devices **501a** and **501b** are grouped into a group **502a**. The non-volatile memory devices **501c** and **501d** are grouped into a group **502b**. The non-volatile memory devices **501e** and **501f** are grouped into a group **502c**. The non-volatile memory devices **501g** and **501h** are grouped into a group **502d**. The groups can include any other numbers (e.g., 3 or more) of non-volatile memory devices.

(85) At **520**, it is determined whether the number of total groups (e.g., the at least one first group) is a non-zero number that is greater than 1. As shown in FIG. **5B**, the number (4) of groups **502a**, **502b**, **502c**, and **502d** is greater than 1. In response to determining that the number of total groups is greater than 1 (**520: YES**), a leader device (e.g., a first leader device) is selected for each group (e.g., from first non-volatile memory devices in each of the at least one first group), at **530**. In some examples, the device with an odd identifier (e.g., 1, 3, 5, and 7) in each of the groups **502a**, **502b**, **502c**, and **502d** is selected as the leader. Accordingly, the non-volatile memory devices **501a**, **501c**, **501e**, and **501g** are selected for the groups **502a**, **502b**, **502c**, and **502d**, respectively. In some examples, the device with the even identifier in each group can be selected as the leader. In some examples, the device with the smallest identifier in each group can be selected as the leader. In some examples, the device with the greatest identifier in each group can be selected as the leader.

(86) At **540**, the leader device determines the result data by performing at least one operation for the data in each group. For example, the controller **110** of first leader device of each of the at least one first group determines the first result data by performing an operation based on first data from each of the non-volatile memory devices in each of the at least one first group.

(87) For example, a buffer (e.g., the buffer **112**) of the first leader device receives the first data from at least one of the first non-volatile memory devices in each of the at least one first group. The controller **110** of the first leader device **501a** determines the first result data by computing the first result data based on data stored in the buffer **112** of the first leader device and the first data from at least one of the non-volatile memory devices in each of the at least one first group.

(88) For example, the controller **110** of the first leader device **501a** of the group **502a** performs a DMA transfer by issuing a DMA request to the controller **110** of a first non-volatile memory device

**501b** in the group **502a** using the address of the buffer **112** of the first non-volatile memory device **501b**. The data at the address of the buffer **112** of the first non-volatile memory device **501b** can be read from the memory array **120** of the first non-volatile memory device **501b** into the buffer **112** of the first non-volatile memory device **501b** or can be received by the buffer **112** of the first non-volatile memory device **501b** from the memory **102** across the interface **140**. In response, the controller **110** of the first leader device **501a** transfers the data stored in the buffer **112** of the first non-volatile memory device **501b** into the buffer **112** of the first leader device **501a**. The controller **110** of the first leader device **501a** can read the data stored in the non-volatile memory (e.g., the memory array **120**) of the first leader device **501a** into the buffer **112** of the first leader device **501a**. The controller **110** of the first leader device **501a** can perform the operation using the data of the first leader device **501a** and the first data from the first non-volatile memory device **501b** to determine the first result data. The operation includes an XOR operation where the data of the first leader device **501a** is bitwise XORed with the first data from the first non-volatile memory device **501b**. The first result data is stored in the buffer **112** of the first leader device **501a**. The result data for the groups **502b**, **502c**, and **502d** can be similarly determined by the leader devices **501c**, **501e**, and **501g**, respectively.

(89) In the examples in which there are three or more non-volatile memory devices in the same group, the first result data can be XORed with data from another non-volatile memory devices in the same group to determine an updated result data, which can in turn be XORed with data from another non-volatile memory devices in the same group if any, and so on.

(90) At **550**, the leader devices from the at least one groups are further grouped. For example, the first leader devices **501a**, **501c**, **501e**, and **501g** for the first groups **502a**, **502b**, **502c**, and **502d** can be grouped into at least one second group **504a** and **504b**. As shown, the second group **504a** includes devices **501a** and **501c**, and the second group **504b** includes devices **501e** and **501g**. The grouping can be performed in the manner as described with respect to **510**. New identifiers can be assigned for the devices **501a**, **501c**, **501e**, and **501g**. For example, the non-volatile memory device **501a** is assigned identifier 1, the non-volatile memory device **501c** is assigned identifier 2, the non-volatile memory device **501e** is assigned identifier 3, the non-volatile memory device **501g** is assigned identifier 4. The grouping can be performed based on the updated identifiers or using another suitable method as described.

(91) In response to grouping the leader devices at **550** (e.g., in response to grouping the first leader device into the at least one second group), the method **500** returns to **520** to determine whether the number of the groups (e.g., the at least one second group) is greater than one. In response to determining that the number of groups is greater than 1 (**520**: YES), the method **500** proceeds to **530**.

(92) For example, a leader device (e.g., a second leader device) is selected for each group (e.g., from second non-volatile memory devices in each of the at least one second group), at **530**. In some examples, the device with an odd identifier (e.g., 1 and 3) in each of the second groups **504a** and **504b** is selected as the leader. Accordingly, the non-volatile memory devices **501a** and **501e** are selected for the groups **504a** and **504b**, respectively. In some examples, the device with the even identifier in each group can be selected as the leader. In some examples, the device with the smallest identifier in each group can be selected as the leader. In some examples, the device with the greatest identifier in each group can be selected as the leader.

(93) At **540**, the leader device determines the result data by performing at least one operation for the data in each group. For example, the controller **110** of second leader device of each of the at least one second group determines the second result data by performing an operation based on second data from at least one of the second non-volatile memory devices in each of the at least one second group.

(94) For example, a buffer (e.g., the buffer **112**) of the second leader device receives the second data from at least one of the second non-volatile memory devices in each of the at least one second

group. The controller **110** of the second leader device **501a** determines the second result data by computing the second result data based on data stored in the buffer **112** of the second leader device **501a** (which can be the first result data determined by the leader device **501a**) and the second data from each of the second non-volatile memory devices in each of the at least one second group. Accordingly, the first result may be an intermediate result.

(95) For example, the controller **110** of the second leader device **501a** of the group **504a** performs a DMA transfer by issuing a DMA request to the controller **110** of a second non-volatile memory device **501c** in the group **504a** using the address of the buffer **112** of the second non-volatile memory device **501c**. The data at the address of the buffer **112** of the second non-volatile memory device **501c** can be read from the memory array **120** of the second non-volatile memory device **501c** into the buffer **112** of the second non-volatile memory device **501b**, can be received by the buffer **112** of the second non-volatile memory device **501c** from the memory **102** across the interface **140**, or can be a result data (e.g., first result data) from a previous iteration of the method **500**. In response, the controller **110** of the second leader device **501a** transfers the data stored in the buffer **112** of the second non-volatile memory device **501c** into the buffer **112** of the second leader device **501a**. The controller **110** of the second leader device **501a** can read the data stored in the non-volatile memory (e.g., the memory array **120**) of the second leader device **501a** into the buffer **112** of the second leader device **501a**, or the data can be a result data (e.g., the first result data) from a previous iteration of the method **500**. The controller **110** of the second leader device **501a** can perform the operation using the data of the second leader device **501a** and the second data from the second non-volatile memory device **501c** to determine the second result data. The operation includes an XOR operation where the data of the second leader device **501a** is bitwise XORed with the second data from the second non-volatile memory device **501c**. The second result data is stored in the buffer **112** of the second leader device **501a**. The result data for the group **504b** can be similarly determined by the leader device **501e**.

(96) At **550**, the second leader devices **501a** and **510e** for the second groups **502a**, **502b**, **502c**, and **502d** can be grouped into at least one third group **506**. As shown, the second group **506** includes devices **501a** and **501e**. The grouping can be performed in the manner as described with respect to **510**. New identifiers can be assigned for the devices **501a** and **510e**. For example, the non-volatile memory device **501a** is assigned identifier 1, the non-volatile memory device **501e** is assigned identifier 2. The grouping can be performed based on the updated identifiers or using another suitable method as described.

(97) In response to grouping the leader devices at **550** (e.g., in response to grouping the second leader devices into the third group), the method **500** proceeds to **560**, in which the leader device for the group **506** is selected, according to any method disclosed herein relative to **530**. At **570**, the leader device determines the final result data by performing an operation based on data in the remaining group, as described relative to **540**. The result data determined by the leader device of the last remaining group is deemed as the final result data.

(98) For example, the controller **110** of the third leader device **501a** of the group **506** performs a DMA transfer by issuing a DMA request to the controller **110** of a third non-volatile memory device **501e** in the group **506** using the address of the buffer **112** of the c. The data at the address of the buffer **112** of the third non-volatile memory device **501e** can be read from the memory array **120** of the third non-volatile memory device **501e** into the buffer **112** of the third non-volatile memory device **501e**, can be received by the buffer **112** of the third non-volatile memory device **501e** from the memory **102** across the interface **140**, or can be a result data (e.g., second result data) from a previous iteration of the method **500**. In response, the controller **110** of the third leader device **501a** transfers the data stored in the buffer **112** of the third non-volatile memory device **501e** into the buffer **112** of the third leader device **501a**. The controller **110** of the third leader device **501a** can read the data stored in the non-volatile memory (e.g., the memory array **120**) of the third leader device **501a** into the buffer **112** of the third leader device **501a**, or the data can be a result

data (e.g., the second result data) from a previous iteration of the method **500**. The controller **110** of the third leader device **501a** can perform the operation using the data of the third leader device **501a** and the third data from the third non-volatile memory device **501e** to determine the third and final result data. The operation includes an XOR operation where the data of the third leader device **501a** is bitwise XORed with the third data from the third non-volatile memory device **501e**. The third and final result data is stored in the buffer **112** of the third leader device **501a**, and can be subsequently stored in the non-volatile memory (e.g., the memory array **120**) or transferred to the memory **102** over the interface **140**.

(99) In some arrangements, **510**, **520**, and **530** can be performed by one or more of the plurality of non-volatile memory devices. For example, at least one managing non-volatile memory device of the plurality of non-volatile memory devices or at least one managing non-volatile memory device other than the plurality of non-volatile memory devices.

(100) The managing non-volatile memory device can receive from the host device context information for each of the plurality of non-volatile memory devices. In some examples, the device context information includes addresses or identifiers of the plurality of non-volatile memory devices which may be in the RAID group. The address of a non-volatile memory device can be a CMB address, SLM address, a PMR address, an address descriptor, an identifier, a pointer, or another suitable indicator that identifies the buffer **112** of that non-volatile memory device of the non-volatile memory device itself, as described. In some examples, the host command includes the logical address of the data subject to the operation, including the logical address (e.g., a buffer address) of the host for the data. In some examples, the device context information can include permission information of at least one of the plurality of non-volatile memory devices or their buffers **112**. In some examples, the device context information can include priority information including at least one of a priority level for processing the data (e.g., the group-based calculations) by each of the plurality of non-volatile memory devices or a priority level for the managing non-volatile memory device or the leader device to communicate with each of the plurality of non-volatile memory devices with respect to the data in connection with the group-based calculations. The managing non-volatile memory device can communicate with one or more of the plurality of non-volatile memory device using the addresses and permission information in the device context information. The managing non-volatile memory device can communicate the group (e.g., the identifier assignment) to all of the plurality of non-volatile memory devices or the non-volatile memory devices in the same group as the managing non-volatile memory device. The managing non-volatile memory device can determine whether the number of group is greater than 1 at **520** and notify the remaining leader devices whether additional operations need to be performed. At **530**, one managing non-volatile memory device may select the leaders for all groups, or a different managing non-volatile memory device may select the leader for each group.

(101) The method **500** can reduce the workload of the processor **104** by dividing the work of the processor **104** among device controllers **110**. The method **500** can reduce the amount of data passing across the root complex **105** due to the transfer of the data being within the PCIe subsystem (e.g., in the bus **106**). The method **500** can reduce the load on system memory **102** by using the buffer **112** to restore intermediate results.

(102) The previous description is provided to enable any person skilled in the art to practice the various aspects described herein. Various modifications to these aspects will be readily apparent to those skilled in the art, and the generic principles defined herein may be applied to other aspects. Thus, the claims are not intended to be limited to the aspects shown herein, but is to be accorded the full scope consistent with the language claims, wherein reference to an element in the singular is not intended to mean “one and only one” unless specifically so stated, but rather “one or more.” Unless specifically stated otherwise, the term “some” refers to one or more. All structural and functional equivalents to the elements of the various aspects described throughout the previous description that are known or later come to be known to those of ordinary skill in the art are

expressly incorporated herein by reference and are intended to be encompassed by the claims. Moreover, nothing disclosed herein is intended to be dedicated to the public regardless of whether such disclosure is explicitly recited in the claims. No claim element is to be construed as a means plus function unless the element is expressly recited using the phrase “means for.”

(103) It is understood that the specific order or hierarchy of steps in the processes disclosed is an example of illustrative approaches. Based upon design preferences, it is understood that the specific order or hierarchy of steps in the processes may be rearranged while remaining within the scope of the previous description. The accompanying method claims present elements of the various steps in a sample order, and are not meant to be limited to the specific order or hierarchy presented.

(104) The previous description of the disclosed implementations is provided to enable any person skilled in the art to make or use the disclosed subject matter. Various modifications to these implementations will be readily apparent to those skilled in the art, and the generic principles defined herein may be applied to other implementations without departing from the spirit or scope of the previous description. Thus, the previous description is not intended to be limited to the implementations shown herein but is to be accorded the widest scope consistent with the principles and novel features disclosed herein.

(105) The various examples illustrated and described are provided merely as examples to illustrate various features of the claims. However, features shown and described with respect to any given example are not necessarily limited to the associated example and may be used or combined with other examples that are shown and described. Further, the claims are not intended to be limited by any one example.

(106) The foregoing method descriptions and the process flow diagrams are provided merely as illustrative examples and are not intended to require or imply that the steps of various examples must be performed in the order presented. As will be appreciated by one of skill in the art the order of steps in the foregoing examples may be performed in any order. Words such as “thereafter,” “then,” “next,” etc. are not intended to limit the order of the steps; these words are simply used to guide the reader through the description of the methods. Further, any reference to claim elements in the singular, for example, using the articles “a,” “an” or “the” is not to be construed as limiting the element to the singular.

(107) The various illustrative logical blocks, modules, circuits, and algorithm steps described in connection with the examples disclosed herein may be implemented as electronic hardware, computer software, or combinations of both. To clearly illustrate this interchangeability of hardware and software, various illustrative components, blocks, modules, circuits, and steps have been described above generally in terms of their functionality. Whether such functionality is implemented as hardware or software depends upon the particular application and design constraints imposed on the overall system. Skilled artisans may implement the described functionality in varying ways for each particular application, but such implementation decisions should not be interpreted as causing a departure from the scope of the present disclosure.

(108) The hardware used to implement the various illustrative logics, logical blocks, modules, and circuits described in connection with the examples disclosed herein may be implemented or performed with a general purpose processor, a DSP, an ASIC, an FPGA or other programmable logic device, discrete gate or transistor logic, discrete hardware components, or any combination thereof designed to perform the functions described herein. A general-purpose processor may be a microprocessor, but, in the alternative, the processor may be any conventional processor, controller, microcontroller, or state machine. A processor may also be implemented as a combination of computing devices, e.g., a combination of a DSP and a microprocessor, a plurality of microprocessors, one or more microprocessors in conjunction with a DSP core, or any other such configuration. Alternatively, some steps or methods may be performed by circuitry that is specific to a given function.

(109) In some exemplary examples, the functions described may be implemented in hardware,

software, firmware, or any combination thereof. If implemented in software, the functions may be stored as one or more instructions or code on a non-transitory computer-readable storage medium or non-transitory processor-readable storage medium. The steps of a method or algorithm disclosed herein may be embodied in a processor-executable software module which may reside on a non-transitory computer-readable or processor-readable storage medium. Non-transitory computer-readable or processor-readable storage media may be any storage media that may be accessed by a computer or a processor. By way of example but not limitation, such non-transitory computer-readable or processor-readable storage media may include RAM, ROM, EEPROM, FLASH memory, CD-ROM or other optical drive storage, magnetic drive storage or other magnetic storages, or any other medium that may be used to store desired program code in the form of instructions or data structures and that may be accessed by a computer. Drive and disc, as used herein, includes compact disc (CD), laser disc, optical disc, digital versatile disc (DVD), floppy drive, and blu-ray disc where drives usually reproduce data magnetically, while discs reproduce data optically with lasers. Combinations of the above are also included within the scope of non-transitory computer-readable and processor-readable media. Additionally, the operations of a method or algorithm may reside as one or any combination or set of codes and/or instructions on a non-transitory processor-readable storage medium and/or computer-readable storage medium, which may be incorporated into a computer program product.

(110) The preceding description of the disclosed examples is provided to enable any person skilled in the art to make or use the present disclosure. Various modifications to these examples will be readily apparent to those skilled in the art, and the generic principles defined herein may be applied to some examples without departing from the spirit or scope of the disclosure. Thus, the present disclosure is not intended to be limited to the examples shown herein but is to be accorded the widest scope consistent with the following claims and the principles and novel features disclosed herein.

## Claims

1. A method, comprising: receiving, by a first non-volatile memory device from a host, a host command comprising device context information of a plurality of non-volatile memory devices, wherein the device context information comprises an address of a buffer of each of the plurality of non-volatile memory devices; in response to receiving the host command, dividing, by the first non-volatile memory device, portions of host data corresponding to the host command among the plurality of non-volatile memory devices; sending, by the first non-volatile memory device to the host, a transfer request indicating transfer of each of the portions of the host data to a respective one of the plurality of non-volatile memory devices; and sending, by the first non-volatile memory device to each of the plurality of non-volatile memory devices other than the first non-volatile memory device, a peer command based on the device context information.
2. The method of claim 1, wherein each of the portions of the host data comprises a buffer in a memory of the host; and dividing the portions of the host data comprises mapping the buffer in the memory of the host to the respective one of the plurality of non-volatile memory devices.
3. The method of claim 2, wherein the transfer request comprises an address of the buffer of the memory of the host and the address of the buffer of the respective one of the plurality of non-volatile memory devices.
4. The method of claim 1, wherein the device context information comprises permission information of at least one of the plurality of non-volatile memory devices or the buffer of the at least one of the plurality of non-volatile memory devices.
5. The method of claim 4, wherein the permission information comprises one or more types of operations for which the second non-volatile memory device or a buffer of the second non-volatile memory device is permitted to be accessed by the first non-volatile memory device.

6. The method of claim 4, wherein the permission information comprises authentication credentials to the second non-volatile memory device or a buffer of the second non-volatile memory device.
7. The method of claim 1, wherein the device context information comprises priority information for processing the host data corresponding to the host command.
8. The method of claim 7, wherein the priority information comprises a priority level of the host data in processing the host data by the first non-volatile memory device.
9. The method of claim 7, wherein the priority information comprises a priority level of the host data in processing the host data by the plurality of non-volatile memory devices.
10. The method of claim 1, wherein the host command comprises a first write command; and the peer command comprises a second write command comprising the address of the buffer of one of the plurality of non-volatile memory devices other than the first non-volatile memory device, wherein the second write command instructs each of the plurality of non-volatile memory devices to transfer one of the portions of the host data from the address of the buffer of the one of the plurality of non-volatile memory devices to a non-volatile memory of the one of the plurality of non-volatile memory devices.
11. A system, comprising: a first non-volatile memory device comprising a buffer, a non-volatile memory, and a controller, wherein the controller is configured to: receive, from a host, a host command comprising device context information of a plurality of non-volatile memory devices, wherein the device context information comprises an address of a buffer of each of the plurality of non-volatile memory devices; in response to receiving the host command, divide portions of host data corresponding to the host command among the plurality of non-volatile memory devices; send to the host a transfer request indicating transfer of each of the portions of the host data to a respective one of the plurality of non-volatile memory devices; and send to each of the plurality of non-volatile memory devices other than the first non-volatile memory device, a peer command based on the device context information.
12. The system of claim 11, wherein each of the portions of the host data comprises a buffer in a memory of the host; and dividing the portions of the host data comprises mapping the buffer in the memory of the host to the respective one of the plurality of non-volatile memory devices.
13. The system of claim 12, wherein the transfer request comprises an address of the buffer of the memory of the host and the address of the buffer of the respective one of the plurality of non-volatile memory devices.
14. The system of claim 11, wherein the device context information comprises permission information of at least one of the plurality of non-volatile memory devices or the buffer of the at least one of the plurality of non-volatile memory devices.
15. The system of claim 11, wherein the device context information comprises priority information for processing the host data corresponding to the host command.
16. The system of claim 11, wherein the host command comprises a first write command; and the peer command comprises a second write command comprising the address of the buffer of one of the plurality of non-volatile memory devices other than the first non-volatile memory device, wherein the second write command instructs each of the plurality of non-volatile memory devices to transfer one of the portions of the host data from the address of the buffer of the one of the plurality of non-volatile memory devices to a non-volatile memory of the one of the plurality of non-volatile memory devices.
17. A non-transitory processor-readable medium comprising processor-readable instructions, such that, when executed, causes a controller of a first non-volatile memory device to: receive, from a host, a host command comprising device context information of a plurality of non-volatile memory devices, wherein the device context information comprises an address of a buffer of each of the plurality of non-volatile memory devices; in response to receiving the host command, divide portions of host data corresponding to the host command among the plurality of non-volatile memory devices; send to the host a transfer request indicating transfer of each of the portions of the

host data to a respective one of the plurality of non-volatile memory devices; and send to each of the plurality of non-volatile memory devices other than the first non-volatile memory device, a peer command based on the device context information.

18. The non-transitory processor-readable medium of claim 17, wherein each of the portions of the host data comprises a buffer in a memory of the host; and dividing the portions of the host data comprises mapping the buffer in the memory of the host to the respective one of the plurality of non-volatile memory devices.

19. The non-transitory processor-readable medium of claim 18, wherein the transfer request comprises an address of the buffer of the memory of the host and the address of the buffer of the respective one of the plurality of non-volatile memory devices.

20. The non-transitory processor-readable medium of claim 17, wherein the host command comprises a first write command; and the peer command comprises a second write command comprising the address of the buffer of one of the plurality of non-volatile memory devices other than the first non-volatile memory device, wherein the second write command instructs each of the plurality of non-volatile memory devices to transfer one of the portions of the host data from the address of the buffer of the one of the plurality of non-volatile memory devices to a non-volatile memory of the one of the plurality of non-volatile memory devices.

---