

# US Patent & Trademark Office

## Patent Public Search | Text View

---

United States Patent	12386629
Kind Code	B2
Date of Patent	August 12, 2025
Inventor(s)	Duval; Olivier

---

### Concurrent image measurement and execution

---

#### Abstract

Apparatuses and methods related to concurrently measuring and executing images. An apparatus for concurrently measuring and executing images can include a memory device, a first processing resource and a second processing resource. The first processing resource can execute instructions stored in the memory device to execute a first portion of an image responsive to measuring the first portion of the image and execute a second portion of the image responsive to measuring the second portion of the image. The second processing resource can execute instructions stored in the memory device to measure the first portion of the image and measure the second portion of the image concurrently with an execution of the first portion of the image by the first processing resource.

---

<b>Inventors:</b>	Duval; Olivier (Pacifica, CA)
<b>Applicant:</b>	Lodestar Licensing Group, LLC (Evanston, IL)
<b>Family ID:</b>	1000008749106
<b>Appl. No.:</b>	18/232714
<b>Filed:</b>	August 10, 2023

#### Prior Publication Data

<b>Document Identifier</b>	<b>Publication Date</b>
US 20240111545 A1	Apr. 04, 2024

#### Related U.S. Application Data

continuation parent-doc US 17366916 20210702 US 11726795 child-doc US 18232714  
division parent-doc US 16118492 20180831 US 11055105 20210706 child-doc US 17366916

---

#### Publication Classification

**Int. Cl.:** G06F9/4401 (20180101)

**U.S. Cl.:**

**CPC** G06F9/4406 (20130101);

## Field of Classification Search

**CPC:** G06F (9/4406)

---

## References Cited

### U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
5765037	12/1997	Morrison et al.	N/A	N/A
6035374	12/1999	Panwar et al.	N/A	N/A
7007161	12/2005	Bradley	713/1	G06F 9/4405
7210127	12/2006	Rangachari	N/A	N/A
7624260	12/2008	Ethier et al.	N/A	N/A
8358653	12/2012	Attig	N/A	N/A
9165143	12/2014	Sanders	N/A	G06F 21/575
9361244	12/2015	Hiremane	N/A	N/A
9658858	12/2016	Koszek	N/A	G06F 9/4405
10057243	12/2017	Kumar	N/A	N/A
10565382	12/2019	Diamant	N/A	N/A
11734018	12/2022	Benisty	713/2	G06F 3/0673
2006/0095726	12/2005	Zaabab	N/A	N/A
2008/0165952	12/2007	Smith et al.	N/A	N/A
2009/0110190	12/2008	Dolgunov	N/A	N/A
2009/0327678	12/2008	Dutton	N/A	N/A
2009/0327680	12/2008	Dale et al.	N/A	N/A
2011/0246759	12/2010	Lai	N/A	N/A
2011/0271126	12/2010	Hill	N/A	N/A
2012/0042376	12/2011	Dolgunov	N/A	N/A
2012/0084547	12/2011	Jung	713/2	G06F 15/177
2012/0284772	12/2011	Kwon	N/A	N/A
2014/0013089	12/2013	Henry et al.	N/A	N/A
2014/0089651	12/2013	Yao	N/A	N/A
2015/0019793	12/2014	Dover	N/A	N/A
2015/0039876	12/2014	Baratam	713/2	G06F 9/4401
2015/0106609	12/2014	Koszek	713/1	G06F 9/4405
2015/0113258	12/2014	Grieco	713/2	G06F 21/572
2015/0199520	12/2014	Woolley	N/A	N/A
2016/0274918	12/2015	Moon	N/A	N/A
2016/0283250	12/2015	Fujiwara	N/A	G06F 11/2284
2017/0372074	12/2016	Gunti	N/A	G06F 21/575
2019/0005244	12/2018	Brasher	N/A	G06F 12/145
2019/0058588	12/2018	Chen	N/A	H04L 9/0863
2019/0078746	12/2018	Dubosc	N/A	N/A

2019/0087578	12/2018	Govindarajan	N/A	N/A
2019/0102539	12/2018	Durham	N/A	N/A
2019/0266331	12/2018	Sanchez Diaz	N/A	N/A
2019/0324806	12/2018	Javre	N/A	G06F 9/4406
2020/0134186	12/2019	Hong	N/A	N/A

## FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
10-2007-0108060	12/2006	KR	N/A
2012138957	12/2011	WO	N/A

## OTHER PUBLICATIONS

Implementing a ARM-based secure Boot Scheme for the isolated Execution Environment, Hang Jiang, Rui Chang, Lu Ren, Weiyu Dong (Year: 2017). cited by examiner  
International Search Report & Written Opinion from related International Application No. PCT/US2019/048945, dated Dec. 12, 2019, 12 pages. cited by applicant

---

*Primary Examiner:* Prifti; Aurel

*Attorney, Agent or Firm:* Brooks, Cameron & Huebsch, PLLC

---

## Background/Summary

**PRIORITY INFORMATION** (1) This application is a Continuation of U.S. application Ser. No. 17/366,916, filed on Jul. 2, 2021, which is a Divisional of U.S. application Ser. No. 16/118,492, filed on Aug. 31, 2018, now issued as U.S. Pat. No. 11,055,105 on Jul. 6, 2021, the contents of which are incorporated herein by reference.

### TECHNICAL FIELD

(1) The present disclosure relates generally to memory, and more particularly, to apparatuses and methods associated with measuring and executing an image currently.

### BACKGROUND

(2) Memory devices are typically provided as internal, semiconductor, integrated circuits in computers or other electronic devices. There are many different types of memory including volatile and non-volatile memory. Volatile memory can require power to maintain its data and includes random-access memory (RAM), dynamic random access memory (DRAM), and synchronous dynamic random access memory (SDRAM), among others. Non-volatile memory can provide persistent data by retaining stored data when not powered and can include NAND flash memory, NOR flash memory, read only memory (ROM), Electrically Erasable Programmable ROM (EEPROM), Erasable Programmable ROM (EPROM), ferroelectric RAM (FeRAM), and resistance variable memory such as phase change random access memory (PCRAM), three dimensional (3D) cross-point (e.g., 3D XPoint), resistive random access memory (RRAM), and magnetoresistive random access memory (MRAM), among others.

(3) Memory is also utilized as volatile and non-volatile data storage for a wide range of electronic applications. Non-volatile memory may be used in, for example, personal computers, portable memory sticks, digital cameras, cellular telephones, portable music players such as MP3 players, movie players, and other electronic devices. Memory cells can be arranged into arrays, with the arrays being used in memory devices.

(4) Various computing systems include a number of processing resources that are coupled to memory (e.g., a memory system), which is accessed in association with executing a set of

instruction (e.g., a program, applications, etc.). A computing system can generate an image corresponding to an application. The computing system can also execute applications that are part of an image. The application can be, for example, an operating system (OS). The operating system execution can begin at boot time. The application can be provided as source code. The source code can be compiled to generate executable files. The image can be created in part from one or more executable files.

---

## Description

### BRIEF DESCRIPTION OF THE DRAWINGS

- (1) FIG. 1 is flow chart illustrating a method for loading and executing an image in accordance with the prior art.
- (2) FIG. 2 is a block diagram of an apparatus in the form of a computing system including a processing resource and a memory device capable of concurrently measuring and executing an image accordance with a number of embodiments of the present disclosure.
- (3) FIG. 3 illustrates a flow chart associated with a method for concurrently measuring and executing an image according with a number of embodiments of the present disclosure.
- (4) FIG. 4 illustrates a flow chart associated with a method for concurrently measuring and executing an image according with a number of embodiments of the present disclosure.
- (5) FIG. 5 illustrates a flow chart associated with generating an image according with a number of embodiments of the present disclosure.
- (6) FIG. 6 is a block diagram of an address space comprising timing constraints in accordance with a number of embodiments of the present disclosure.

### DETAILED DESCRIPTION

- (7) The present disclosure includes apparatuses and methods related to the concurrent measurement and execution of an image. An example apparatus can include a memory device and a first processing resource to execute a first portion of an image responsive to measuring the first portion of the image and execute a second portion of the image responsive to measuring the second portion of the image. The example apparatus can also include a second processing resource to measure the first portion of the image and measure the second portion of the image concurrently with an execution of the first portion of the image by the first processing resource.
- (8) Secure booting can include associating an image stored in memory with a signature that can be verified with a cryptographic key. The signature can be assessed to determine whether the image has been produced by a trusted agent. As used herein, a trusted agent is an entity authorized to generate and/or store images in the memory. The trusted agent can include an authorized user of the computing system and/or authorized computer readable instructions configured to generate an image and/or store the image in memory.
- (9) The process of utilizing the cryptographic key and the signature to determine whether the image has been produced by a trusted agent can be referred to as measuring. The image can be measured to determine whether the image has been produced by a trusted agent.
- (10) Measuring an image may delay boot time. The delay may exceed a tolerance threshold in time-critical applications. A time-critical application can include, for example, an application operating in an automobile. An automobile application can be time-critical because a duration of time used to measure an image can have an impact on the state of the automobile and/or the application operating in the automobile. For example, a delay in measuring an image can have negative impact on the automobile. An image may be authenticated by authenticating a portion of the image and running the authenticated portions sequentially. Portions of the image may be authenticated and executed sequentially to verifying that the portions of the image have been produced by a trusted agent. For example, the image may be copied from memory to cache

corresponding to the processing resource. The processing resource may retrieve the image from cache and execute different portions of the image sequentially.

(11) In a number of examples, the image can be measured and executed directly from memory. For example, the image can be measured and executed directly from non-volatile memory such as flash memory, among other types of non-volatile memory. Measuring and/or executing an image from memory can be referred to as execution-in-place. Said differently, a processing resource can execute an image from memory. Executing the image directly from memory may include refraining from copying the image before the image is executed. For example, executing the image directly can include refraining from copying the image to volatile memory before the image is executed from non-volatile memory.

(12) Images executed-in-place can also be authenticated to perform a secure execution of the image. The image can be divided into portions. The portions of the image can be measured and executed at different times without compromising the security.

(13) For example, the image can be measured while the image is being executed. Portions of the image can be measured while different portions of the image are being executed. Concurrently measuring and executing the image can reduce a delay to running an application. For example, concurrently measuring and executing the image can reduce the delay to a boot time.

(14) In the following detailed description of the present disclosure, reference is made to the accompanying drawings that form a part hereof, and in which is shown by way of illustration how a number of embodiments of the disclosure may be practiced. These embodiments are described in sufficient detail to enable those of ordinary skill in the art to practice the embodiments of this disclosure, and it is to be understood that other embodiments may be utilized and that process, electrical, and/or structural changes may be made without departing from the scope of the present disclosure.

(15) As used herein, “a number of” something can refer to one or more of such things. For example, a number of memory devices can refer to one or more memory devices. A “plurality” of something intends two or more. Additionally, designators such as “N,” as used herein, particularly with respect to reference numerals in the drawings, indicates that a number of the particular feature so designated can be included with a number of embodiments of the present disclosure.

(16) The figures herein follow a numbering convention in which the first digit or digits correspond to the drawing figure number and the remaining digits identify an element or component in the drawing. Similar elements or components between different figures may be identified by the use of similar digits. As will be appreciated, elements shown in the various embodiments herein can be added, exchanged, and/or eliminated so as to provide a number of additional embodiments of the present disclosure. In addition, the proportion and the relative scale of the elements provided in the figures are intended to illustrate various embodiments of the present disclosure and are not to be used in a limiting sense.

(17) FIG. 1 is flow chart illustrating prior art of a method **100** for loading and executing an image. At **102**, the method **100** includes powering (e.g., turning on) a computing device. At **104**, the method includes initiating a board. A board refers to a circuit board. The circuit board is a component of the computing device. The board can facilitate communication between devices of the computing system. The board can also provide power to the devices comprising the computing system. The board can facilitate communications provide powers through electrical connections by which the devices of the computing system are coupled. Initiating a circuit board can include providing power to non-volatile memory (e.g., flash memory) integrated in the circuit board and providing power to the electrical connections/components of the circuitry board. Initiating the circuitry board can also include initiating a basic input/output system stored in the non-volatile memory. Initiating the circuit board can further include conducting a number of tests to ensure system health. Initializing the circuit board can include loading an operating system using machine-readable instructions stored in the non-volatile memory.

(18) At **106**, the method **100** includes loading a full image. At **108**, the method includes executing the image.

(19) Powering a device may include providing power to a computing device. Upon providing power to the computing device, a board can be initiated. In the example of FIG. **1**, loading an image includes loading a full image. The image can be loaded from a non-volatile memory device (e.g., flash or hard disk, among others) to volatile memory (e.g., DRAM and/or SRAM, among others). Loading can include storing or copying the image. Loading a full image includes loading an entire image without making alterations and/or changes to the image. The image is loaded before the image is executed. For example, the full image is loaded before the full image is executed.

(20) Loading an image and executing an image can be part of executing an operating system (OS) which can be done upon executing a power on event. FIG. **1** shows a boot sequence with no delay due to the unsecure nature of the boot sequence. The boot sequence is unsecure because the image is not measured.

(21) To secure an image, a signature of the image can be created; the image and its signature are then stored in a memory device. The image can also be measured before the image is executed. The integrity of the image can be verified before the image is executed. Verifying that the content of an image has not been modified since its signature was generated can be referred to as measuring the image. The signature can be created using a cryptographic key. A signature generated with a cryptographic key can be referred to as a cryptographic measurement. As such, there are not delays introduced due to cryptographic measurements in FIG. **1** because no cryptographic measurements are performed in FIG. **1**.

(22) Cryptographic measurements can introduce delays to the execution of an image. For example, measuring the image before the image is run at **108** can introduce a delay corresponding to the duration of time used to measure the image.

(23) FIG. **2** is a block diagram of an apparatus in the form of a computing system **211** including a processing resource **213**, a memory device **215** and a separate processing resource **217** capable of concurrently measuring and executing an image in accordance with a number of embodiments of the present disclosure. The computing system **211** comprises a host **209**, a memory device **215** and a processing resource. The host includes the processing resource **213**. The memory device **215** includes a processing resource **217** and an array **219**. In some examples, the processing resource **217** can be a standalone processing resource. For example, the processing resource **217** can be external to the memory device **215**.

(24) As used herein, an “apparatus” can refer to, but is not limited to, a variety of structures or combinations of structures, such as a circuit or circuitry, a die or dice, a module or modules, a device or devices, or a system or systems. For example, the host **209** and/or the memory device **215** may separately be referred to as an “apparatus.”

(25) The memory device **215** can comprise multiple channels each having a number of devices corresponding thereto. The devices can collectively be referred to as memory devices **215**. Each of the devices **215** can be, for example, a “chip” comprising multiple banks. A controller can service multiple memory channels or the controller can be implemented as multiple separate controllers (e.g., one for each channel). The memory device **215** can be a non-volatile memory device or a volatile memory device. Volatile memory can require power to maintain its data and includes random-access memory (RAM), dynamic random access memory (DRAM), and synchronous dynamic random access memory (SDRAM), among others. Non-volatile memory can provide persistent data by retaining stored data when not powered and can include NAND flash memory, NOR flash memory, read only memory (ROM), Electrically Erasable Programmable ROM (EEPROM), Erasable Programmable ROM (EPROM), ferroelectric RAM (FeRAM), and resistance variable memory such as phase change random access memory (PCRAM), three dimensional (3D) cross-point (e.g., 3D XPoint), resistive random access memory (RRAM), and

magnetoresistive random access memory (MRAM), among others. In examples that load the image, the image can be loaded from non-volatile memory to volatile memory. In examples that do not load the image, the image can be executed directly from non-volatile memory.

(26) In this example, the computing system includes the host **211** coupled to a controller (e.g., via an interface), which is coupled to the memory device **215** (e.g., via an interface **205**). The computing system **211** can be a laptop computer, personal computers, digital camera, digital recording and playback device, mobile telephone, PDA, memory card reader, interface hub, sensor, Internet-of-Things (IoT) enabled device, among other systems, and the host **209** can include a processing resource **213** (e.g., one or more processors) capable of accessing the memory device **215** (e.g., via a controller). The host **209** may be responsible for execution of an OS and/or various applications that can be loaded thereto (e.g., from memory device **215** via a controller).

(27) The memory device **215** through the controller may receive memory access requests (e.g., in the form of read and write commands, which may be referred to as load and store commands, respectively) from the host **209**. The controller can transfer commands and/or data between the host **209** and the memory device **215** over the interface **205**, which can comprise physical interfaces such as buses, for example, employing a suitable protocol. Such protocol may be custom or proprietary, or the interface **205** may employ a standardized protocol, such as Peripheral Component Interconnect Express (PCIe), Gen-Z, CCIX, or the like.

(28) As an example, the interface **205** may comprise combined address, command, and data buses or separate buses for respective address, command, and data signals. The controller can comprise control circuitry, in the form of hardware, firmware, or software, or any combination of the three. As an example, the controller can comprise a state machine, a sequencer, and/or some other type of control circuitry, which may be implemented in the form of an application specific integrated circuit (ASIC) coupled to a printed circuit board. In a number of embodiments, the controller may be co-located with the host **209** (e.g., in a system-on-chip (SOC) configuration). Also, the controller may be co-located with the memory device **215**.

(29) The memory device **215** may be referred to as a memory system. The memory device **215** can include a number of memory devices which may be referred to collectively as a memory device **215**. The memory device **215** can include a memory array **219** of memory cells. For example, the memory device **215** can comprise a number of physical memory “chips,” or dice which can each include a number of arrays (e.g., banks) of memory cells and corresponding support circuitry (e.g., address circuitry, I/O circuitry, control circuitry, read/write circuitry, etc.) associated with accessing the array(s) **219** (e.g., to read data from the arrays and write data to the arrays). As an example, the memory device **215** can include a number of DRAM devices, SRAM devices, PCRAM devices, RRAM devices, FeRAM, phase-change memory, 3DXPoint, and/or Flash memory devices. In a number of embodiments, the memory device **215** can serve as main memory for the computing system.

(30) In some examples, the processing resource **213** can comprise a central processing unit (CPU). The CPU is an electronic circuitry within the host **209** that executes instructions of an application by performing arithmetic, logical, control and/or input/output (I/O) operations specific to the application. An application can include executable computer readable instructions in the form of an image.

(31) The processing resource **217** can be, for example, a cryptoprocessor. In some examples, the cryptoprocessor can reside external to the memory device **215**. A cryptoprocessor can be a processing resource that performs cryptographic operations. In some examples, the cryptoprocessor can be a dedicated processing resource that performs cryptographic operations. In a number of examples, the processing resource **217** can be a processing resource other than a cryptoprocessor such as a graphical processing unit (GPU) among other types of processing resources.

(32) Cryptographic operations can be performed to measure data in order to secure the data. The cryptographic operations can include signature and signature verification operations. Signature

operations can sign data utilizing an algorithm and a cryptographic key. Signature verification operations can verify data utilizing the a similar and a cryptographic key. The cryptographic keys used to sign data and verify data can be different keys. For example, data can be signed using a private key while data can be verified using a public key. Data can be signed using an algorithm such as elliptic curve digital signature, involving a cryptographic hash function of the data and an elliptic function of the hash and the private key. The data and its signature are combined and stored. The same data can then be measured by verifying its signature using similar steps, but now with the public key. If the signature verification succeeds, then the image can be trusted. A trusted image is an image produced by a trusted agent, holding a cryptographic key.

(33) In some instances, the cryptographic key can be a symmetric key. The symmetric key can be used for encryption and decryption. The symmetric encryption key and the symmetric decryption key can be a same key or can utilize a transformation to generate one key from the other key.

(34) In some instances, the processing resource(s) can measure an image and execute (e.g., run) an image concurrently. For example, the processing resource **213** can execute an image while the processing resource **217** measures the image. The processing resource **213** can execute a first portion of the image while the processing resource **217** measures a second portion of the image.

(35) In some examples, the processing resource **213** can execute an image that is stored in the memory device **215** (e.g., array **219**). For example, the processing resource **213** can execute an image without storing (e.g., copying) the image to volatile memory. The image can be copied to the cache **221** which includes a hardware and/or software device that stores data local to the processing resource **213**. In some examples, the image and/or portions of the image can also be copied to program registers of the processing resource **213**. The data stored in the cache **221** can be a duplicate (e.g., copy) of data stored in the array **219**.

(36) In some examples, the image can be loaded (e.g., copied) from the memory device **215** to volatile memory before the image is copied from volatile memory to the cache **221**/program registers of the processing resource **213**. The processing resource **213** can execute an image stored in the array **219** without storing the data corresponding to the image in non-volatile memory other than the cache **221** and/or the program registers.

(37) FIG. 3 illustrates a flow chart **310** associated with a method for concurrently measuring and executing an image according with a number of embodiments of the present disclosure. In some examples, a first processing resource that can be part of a memory device can receive instructions from a second processing resource which is part of a host. The instructions can request the authentication and the delivery of data comprising an image. For example, the request can include a request for a portion of the data comprising the image. In some examples, the request can include a memory address that corresponds to the portion of the data comprising the image. The request can also include a range of addresses comprising a base memory address and a range of memory space which defines a range of addresses corresponding to the portion of the data comprising the image.

(38) At **312-1**, a first processing resource (e.g., processing resource **217** in FIG. 2) can measure the first portion of an image. The first processing resource can retrieve the image from a memory array local to the first processing resource. For example, the first processing resource, which is included in the memory device **215** can retrieve the first portion of data comprising a first portion of an image from a memory array which is also included in the memory device.

(39) The first processing resource can measure the first portion of the data comprising the image. The measurement can verify that the image is authentic based on an authentication scheme implemented by the first processing resource. The first processing resource can respond to the second processing resource with the result of the measurement.

(40) The second processing resource can determine whether to load the first portion of the image. If the image is not authenticated, as measured by the first processing resource, then the second processing resource can refrain from proceeding. At **314-1**, the second processing resource can load the first portion of the image based on a determination that the first portion of the image was



authenticated by the first processing resource. The first portion of the image can be copied from non-volatile memory to volatile memory. The first portion of the image can further be copied from volatile memory to cache. For example, the first portion of the image can be loaded to level 1 cache, level 2 cache, and/or level 3 cache, among other types of cache.

(41) In examples where the image is not loaded, the first portion of the image can be retrieved from the non-volatile memory device and can be provided directly to the second processing resource without storing the first portion of the image in volatile memory. Accordingly, in examples where the image is not copied to volatile memory, the loading (**314-1**, **314-2**, . . . , and **314-N**) can be substituted with retrieving the portion of the image from the memory device.

(42) At **316-1** and **312-2**, the first portion of the image can be executed concurrently with the measuring of a second portion of the image. The concurrent measurement and execution of the first portion of the image can constitute a bifurcation of an execution stream. For example, a first branch of the execution stream can be controlled by the first processing resource while a second branch of the execution stream is controlled by the second processing resource. As used herein, an execution stream is an execution of an ordered set of instructions as referenced by time. As such, the bifurcation of the execution stream can constitute the execution of different sets of instructions or the refraining from executing instructions at a given time. For example, the second processing resource can execute a first portion of the image while the first processing resource measures a second portion of the image during a same time period.

(43) The first portion of the image can be executed by the second processing resource at **316-1**. The second portion of the image can be measured at **312-2**. The first portion of the image can be executed concurrently with the measuring of the second portion of the image because the second processing resource executes the first portion of the image while the first processing resource measures the second portion of the image.

(44) In some examples, the second processing resource can direct the measurement of the second portion of the image and the concurrent execution of the first portion of the image. For example, the second processing resource can execute the first portion of the image and wait for the measurement of the second portion of the image to conclude. Responsive to concluding the measurement of the second portion of the image, the first processing resource can notify the second processing resource of the conclusion of the measurement of the second portion of the image.

(45) The first processing resource can provide a notice to the second processing resource utilizing a register. In some examples, the register utilized to provide the notice can be hosted by the memory device hosting the first processing resource and/or the host hosting the second processing resource. The register can be referred to as a status register. The status register can indicate that the second portion of the image is ready to be accessed and/or that the second portion of the image has been successfully measured. Successfully measuring a portion of an image can include a successful termination of the measurement regardless of whether the measurement concluded that the second portion was authentic or not.

(46) In some examples, the status register can include a plurality of status registers. The plurality of status register can describe whether the measurement terminated, whether the measurement determined an authentic portion of the image, and/or whether the measurement determined that the portion of the image was not authentic.

(47) The second processing resource, upon terminating the execution of the first portion of the image, can access the status register. Responsive to determining whether the measurement of the second portion of the image has concluded, the second processing resource can wait for the conclusion of the second portion of the image or can load the second portion of the image at **314-2**.

(48) To wait for the conclusion of the measurement of the second portion of the image, the second processing resource can execute a loop. A loop is a control flow statement that repeatedly provides access to instructions. The repeated access to instructions is described as an iteration. As such, each iteration of a loop provides access to a same set of instructions. In some examples, the loop may

provide access to no instructions and can be utilized to retain control of the instruction flow of the second processing resource to wait for the conclusion of the measurement of the second portion of the image.

(49) Before each iteration of the loop, the second processing resource can determine whether the measurement of the second portion of the image has terminated. The determination of the termination of the measurement of the second portion of the image can be done utilizing the status register(s). The loop can be any type of loop including a while loop and a for loop, among other types of loops. The determination can be performed before the initiation of an iteration of a loop, during the iteration of a loop, or after each iteration of a loop.

(50) In other examples, the second processing resource can direct the measurement of the second portion of the image and the concurrent execution of the first portion of the image using timing constraints. The first portion of the image can include time restrictions and/or time guidelines that can be used by the second processing resource to determine the execution of the first portion of the image, the measurement of the second portion of the image, the loading of the second portion of the image, and/or the execution of the second portion of the image.

(51) For example, the first portion of the image can include a first duration of time that corresponds to the execution of the first portion of the image. The first duration of time describes a duration of time utilized to execute the first portion of the image by the second processing resource. The first portion of the image can also include a second duration of time that corresponds to the measurement of the second portion of the image. The second portion of the image describes a duration of time utilized to measure the second portion of the image by the first processing resource. The second processing resource can compare the first duration of time and the second duration of time to determine a longest duration of time. The second processing resource can refrain from loading and/or executing the second portion of the image until the longest duration of time has expired.

(52) The longest duration of time can be tracked from a time during the execution of the first portion of the image or from the beginning of the execution of the first portion of the image. For example, the longest duration of time can be measured from a time after the first duration of time and the second duration of time are compared. The longest duration of time can also be measured from a time marking the beginning of the execution of the first portion of the image. If the longest duration of time is the second duration of time, then the longest duration of time can be tracked from a time marking the beginning of the measurement of the first portion of the image.

(53) In some examples, the longest duration of time is not calculated by the second processing resource but can be included in the first portion of the image. In such examples, the second processing resource can extract the longest duration of time during the execution of the first portion of the image and can begin tracking the expiration of the longest duration of time during the execution of the first portion of the image.

(54) The duration of time included in the first portion of the image can be included, for example, in a header of the first portion of the image and/or in the body of the first portion of the image. In some examples, each portion of the image can include a header comprising at least a duration of time. In other examples, not every portion of the image includes a duration of time. For example, a last portion of the image may not include a duration of time due the lack of a need for synchronicity. The last (P.sub.N) portion of the image may be executed without concurrently measuring a different portion of the image.

(55) At 312-2, the second processing resource can instruct the first processing resource to measure a second portion of the image. The second processing resource can provide an address, a range of addresses, and/or an identifier of the second portion of the image. The first processing resource can identify and measure the second portion of the image using the address, the range of addresses, and/or the identifier provided by the first processing resource. The first processing resource can provide an indication of the result of the measurement to the second processing resource.

(56) Upon determining that the measurement of the second portion of the image has concluded and that the second portion of the image is authenticated, the second processing resource can load the second portion of the image similarly to the loading of the first portion of the image at **314-1**. At **316-2** and **312-3**, the second processing resource can execute the second portion of the image and the first processing resource can concurrently measure a third portion of image as described at **316-1** and **312-2**. The loading and/or execution of the second portion of the image can constitute a merger of the execution stream of the first processing resource with the second processing resource.

(57) The second processing resource can execute and the first processing resource can concurrently measure a plurality of portion of the image. The second processing resource can execute a P.sub.N-1 portion of the image and direct the concurrent measurement of the P.sub.N of the image by the first processing resource. Upon the successful measurement of the P.sub.N portion of the image, the second processing resource can load the P.sub.N portion of the image at **314-N** and execute the P.sub.N of the image at **316-N**.

(58) FIG. 4 illustrates a flow chart **410** associated with a method for concurrently measuring and executing an image according with a number of embodiments of the present disclosure. In the example provided by the flow chart **410**, the image is an image of an OS. The image can include a second stage boot loader (SSBL). In some examples, the image can include a first stage boot loader (FSBL) and/or a SSBL. FIG. 4 can include a flow chart **410** used during a boot sequence.

(59) A first stage boot loader (FSBL) can be included in the master boot record. The FSBL can be retrieved from secure memory such that in some instances the FSBL is not measured. The FSBL can comprise information used by the second processing resource to measure, load, and execute a SSBL. For example, the FSBL can include an address of the image. As another example, the FSBL can include an address of the SSBL. The SSBL can load an OS and transfer execution to the OS.

(60) At **412-1** the SSBL can be measured. The second processing resource can access an address or identifier of the SSBL from the FSBL. The second processing resource can provide the address of the SSBL to the first processing resource. The first processing resource can retrieve the SSBL from the memory device and can measure the SSBL. The first processing resource can provide the determination of the measurement to the second processing resource.

(61) Based on the determinant of the measurement of the SSBL, the first processing resource can load the SSBL at **414-1**. For example, if the SSBL is determined to be secure, then the second processing resource can load or move the SSBL. If the SSBL is loaded, then the SSBL is copied from volatile memory from non-volatile memory. If the SSBL is not moved (e.g., loaded), then a copy of the SSBL is provided to the second processing resource without storing the SSBL in volatile memory.

(62) At **416-2** and **412-2**, the SSBL is executed and the rest of the image is measured concurrently. In the example provided in FIG. 4, the image can comprise two portions, the first portion can be a SSBL and the second portion can be the rest of the image. In other examples, the rest of the image can be further divided into additional portions.

(63) At **414-2**, the rest of the image can be loaded to volatile memory from non-volatile memory. At **416-2**, the rest of the image can be executed. The rest of the image can be an OS. As such, the OS can be measured at **414-2** and loaded at **416-2**.

(64) FIG. 5 illustrates a flow chart **520** associated with generating an image according with a number of embodiments of the present disclosure. The flow chart **520** includes a compiler **524**, a profiler **566**, a reorder unit **570**, a driver **574**, and a packer **580**. The packer **580** can be a secure packing tool. The compiler **524** can include a compiler and a linker. Each of the compiler **524**, the profiler **566**, the reorder unit **570**, the driver **574**, and the packer **580** can be implemented in hardware and/or software.

(65) The compiler **524** can receive source code **522** corresponding to the image. The source code can include a plurality of tags. The tags can be added to a source code during development. The

tags can identify different section of the source code **522**. The different sections of the source code **522** can contain code that is able to be measure and executed without the need for code outside of a particular section.

(66) The compiler **524** can include software or hardware to transform machine readable instructions written in a first programming language into machine readable instructions written in a second programming language. For example, the compiler **524** can receive the source code **522** that comprises machine readable instructions written in a first programming language and can generate the executable instructions which can comprise machine readable instructions in a second programming language.

(67) The compiler **524** can generate the executable instructions **528**. The executable instructions **528** are unordered and comprise tags. For example, the executable instructions **528** can include the tags provided with the source code **522** and the different section generated with the tags may be unordered. The order of the tags can correspond to the execution stream of the different sections of the executable instructions **528**. For example, a second section may proceed a third section in the executable instructions **528** but may be executed after the execution of the third section. As such, the occurrence of the second section before the third section, given that the second section may be executed after the third section, can render the executable instructions **528** as unordered.

(68) The profiler **566** can receive the executable instructions **529**, a root file system **562**, and boot sequence files **564**. The file system **562** can describe the file system in which the image **532** is executed. The boot sequence files **564** can describe the boot sequence and can provide additional information used to execute the final image **532** in view of the boot sequence.

(69) The profiler **566** can be an emulated profiler. The emulated profile can be a profiler implemented in machine executable instructions. The profiler **566** can also be hardware.

(70) The profiler **566** can label code section based on the tags implemented in the executable instructions **529**. For example, the beginning of the executable instructions **529** to the first tag can be labeled a first section, from the first tag to the second tag can be labeled a second section, etc. There can be one more section than there are tags. The profiler **566** can generate the executable instructions **568** comprising tags and sections, where the executable instructions **568** are unordered.

(71) The profiler **566** can also identify the non-volatile memory load commands called between each tag from the executable instructions **529**. In some examples, the compiler **524** can generate a single set of executable instructions **528** and the profiler **566** can receive multiple sets of executable instructions **529**. The profiler **566** can label the areas of the non-volatile memory targeted by each load command with a section number. The profiler **566** can store the labels of the non-volatile memory and can provide the labels in a data vector **578**.

(72) The reorder unit **570** can reorder the executable instructions **568** using the tags and the code sections. The reorder unit **570** can generate the ordered executable instructions **572** comprising tags. The order of the executable instructions **572** can correspond to an execution stream of the executable instructions **568**. For example, if a third section is executed before a second section, then the reorder unit **570** can reorder the third section and the second section such the third section occurs and/or is executed before the second section.

(73) The order of the third section and the second section can determine the order in which the third section and the second section are measured, loaded, and/or executed. For example, reordering the third section before the second section provides for the measurement, loading, and/or execution of the third section before the second section. The third section can be measured and/or loaded before the second section is measured. The reordering provides for the concurrent execution and measurement of the third section and the second section, respectively.

(74) The reorder unit **570** can reorder the executable instructions **568** and can modify the executable instructions **568** to correspond with the reordering. For example, the reorder unit **570** can update jump commands and/or linker information, among other portion of the executable instructions, that can be associated with each code section. Updating the jump command and/or

linker information can provide for the jump commands and/or the linker information reflecting the updated ordering of the code sections. For example, the reorder unit **570** can update pointer information based on the reordering of the code sections.

(75) The driver **574** can receive the executable instructions **572** and can generate the executable instructions **576**. The executable instructions **576** can be ordered and can include memory calls (e.g., memory device calls). The driver **574** can be specific to the platform that is targeted by the executable. The platform describes the host, the processing resources, the memory device and/or other components of the computing device executing the image **532**. The driver **574** can include low level commands used to communicate with the memory device. The driver **574** can inject the commands directly in the executable instructions **572** to generate the executable instructions **576**. The commands can be placed where tags exist in the executable instructions **572**. The driver **574** can replace the tag with a set of commands (e.g., calls) to the memory device. In some examples, the driver **574** can insert measure commands and load command in the executable instructions **572** to replace the tags in the executable instructions **572**, where the tags identify measure commands and loads commands of sections (e.g., portions) of an image.

(76) The packer **580** can receive the root file system **562**, the boot sequence files **564**, and the executable instructions **576**. In some examples, the packer **580** can also receive the data vector **578** with tagged memory device areas. The packer **580** can generate an image **532** comprising separate code sections (e.g., portions). The packer **580** can also generate a configuration file **582**. The configuration file **582** can be used by a memory device to measure synchronously with a processing resource external to the memory device. The measure commands implemented by a memory device utilize a configuration implemented by the memory device. For example, the memory device can be configured to perform measurements utilizing the configuration file **582**.

(77) FIG. **6** is a block diagram of an address space **640** comprising timing constraints in accordance with a number of embodiments of the present disclosure. The address space **640** can store a plurality of portions **651-1**, **651-2**, **651-3**, and **651-4** of an image, referred to as portions **651**.

(78) The portion **651-1** can be stored in a range of addresses beginning with address 0x0000000. The portion **651-2** can be stored in a range of addresses beginning with address 0x0080000. The portion **651-3** can be stored in a range of addresses beginning with address 0x0010000. The portion **651-4** can be stored in a range of addresses beginning with address 0x0100000.

(79) The reorder unit **570** of FIG. **5** can order an execution of the portions **651** such that the portion **651-1** is executed before the portion **651-2** is executed, the portion **651-2** is executed before the portion **651-3** is executed, and the portion **651-3** is executed before the portion **651-4** is executed. In some examples, the reorder unit **570** can order the concurrent execution and measurement of the portions **651** such that the portion **651-1** is executed concurrently with the measurement of the portion **651-2**, the portion **651-2** is executed concurrently with the measurement of the portion **651-3**, and the portion **651-3** is executed concurrently with the measurement of the portion **651-4**.

(80) The reorder unit **570** can order the concurrent execution and measurement of the portions **651** using the addresses corresponding to each of the portions **651**. For example, the reorder unit **570** can associate the portion **651-1** with the address 0x0080000 to indicate that the portion **651-1** is to be executed concurrently with the measurement of the portion **651-2**. The processing resource executing the portion **651-1** can provide the address 0x0080000 to the cryptoprocessor and continue with the execution of the portion **651-1**. The cryptoprocessor can retrieve the portion **651-2** using the address 0x0080000 and can measure the portion **651-2** concurrently with the execution of the portion **651-1** as previously defined.

(81) If the cryptoprocessor, responsive to measuring the portion **651-2**, determines that the portion **651-2** is not secure, then the cryptoprocessor can indicate such to the processing resource. The processing resource can prevent unsecure memory access by refraining from loading or executing the portion **651-2** based on the results of the measurement.

(82) In some embodiments, the processing resource can poll the cryptoprocessor regarding the

measurement of a portion of the image. The processing resource can poll the cryptoprocessor instead of waiting for the cryptoprocessor to respond with a result of the measurement. The processing resource can poll the memory device and/or the cryptoprocessor by accessing a register of the memory device to determine whether the portions 651 of the image have been measured and/or to retrieve a result of the measurement.

(83) Although specific embodiments have been illustrated and described herein, those of ordinary skill in the art will appreciate that an arrangement calculated to achieve the same results can be substituted for the specific embodiments shown. This disclosure is intended to cover adaptations or variations of various embodiments of the present disclosure. It is to be understood that the above description has been made in an illustrative fashion, and not a restrictive one. Combinations of the above embodiments, and other embodiments not specifically described herein will be apparent to those of skill in the art upon reviewing the above description. The scope of the various embodiments of the present disclosure includes other applications in which the above structures and methods are used. Therefore, the scope of various embodiments of the present disclosure should be determined with reference to the appended claims, along with the full range of equivalents to which such claims are entitled.

(84) In the foregoing Detailed Description, various features are grouped together in a single embodiment for the purpose of streamlining the disclosure. This method of disclosure is not to be interpreted as reflecting an intention that the disclosed embodiments of the present disclosure have to use more features than are expressly recited in each claim. Rather, as the following claims reflect, inventive subject matter lies in less than all features of a single disclosed embodiment. Thus, the following claims are hereby incorporated into the Detailed Description, with each claim standing on its own as a separate embodiment.

## Claims

1. An apparatus, comprising: a memory device; a first processing resource executing instructions stored in the memory device to: measure a second stage boot loader (SSBL) of an image responsive to execution of a first stage boot loader (FSBL) by a second processing resource; execute the SSBL responsive to measuring the SSBL; and the second processing resource executing instructions stored in the memory device to: measure a remainder of the image, concurrently with an execution of the SSBL by the first processing resource, to determine whether the remainder of the image has been produced by a trusted agent, wherein the second processing resource executes the FSBL and measures the remainder of the image; and execute the remainder of the image.
2. The apparatus of claim 1, wherein the FSBL is included in a master boot record.
3. The apparatus of claim 1, wherein the second processing resource is further configured to measure the FSBL prior to the first processing resource measuring the SSBL.
4. The apparatus of claim 1, wherein the second processing resource is further configured to execute the FSBL prior to the first processing resource executing the SSBL.
5. The apparatus of claim 1, wherein the second processing resource is configured to execute the FSBL to obtain information used by the first processing resource to measure and execute the SSBL.
6. The apparatus of claim 5, wherein the second processing resource is configured to execute the FSBL to obtain an address of the image.
7. The apparatus of claim 5, wherein the second processing resource is configured to execute the FSBL to obtain an address of the SSBL.
8. The apparatus of claim 1, wherein the second processing resource is further configured to load an operating system (OS) responsive to executing the SSBL.
9. The apparatus of claim 8, wherein the second processing resource is further configured to transfer execution to the OS.
10. A method, comprising: measuring, using a first processing resource coupled to a memory

device, a second stage boot loader (SSBL) of an image responsive to execution of a first stage boot loader (FSBL) by a second processing resource, wherein the SSBL is measured to determine whether the SSBL is secure; executing, using the first processing resource, the SSBL responsive to a determination that the SSBL is secure; measuring, using a second processing resource coupled to the memory device, a remainder of the image, concurrently with an execution of the SSBL by the first processing resource, to determine whether the remainder of the image has been produced by a trusted agent, wherein the second processing resource executes the FSBL and measures the remainder of the image; and executing, using the second processing resource, the remainder of the image.

11. The method of claim 10, further comprising retrieving the FSBL from secure memory.

12. The method of claim 10, further comprising retrieving the FSBL from secure memory without measuring the FSBL.

13. The method of claim 10, further comprising accessing, using the first processing resource, an address of the SSBL from the FSBL.

14. The method of claim 10, further comprising providing, by the second processing resource, an address of the SSBL retrieved from the FSBL to the first processing resource.

15. The method of claim 10, further comprising providing the determination that the SSBL is secure from the first processing resource to the second processing resource.

16. The method of claim 15, further comprising measuring, using the second processing resource, the remainder of the image responsive to the determination that the SSBL is secure.

17. An apparatus, comprising: a memory device; a first processing resource executing instructions stored in the memory device to: measure a second stage boot loader (SSBL) of an image of an operating system (OS) responsive to an execution of a first stage boot loader (FSBL) by a second processing resource; execute the SSBL responsive to measuring the SSBL; and the second processing resource executing instructions stored in the memory device to: measure a remainder of the image, concurrently with an execution of the SSBL by the first processing resource, to determine whether the remainder of the image has been produced by a trusted agent, wherein the second processing resource executes the FSBL and measures the remainder of the image; and execute the remainder of the image responsive to a determination that the remainder of the image has been produced by the trusted agent.

18. The apparatus of claim 17, wherein second processing resource is further configured to measure the FSBL to determine whether the FSBL is secure.

19. The apparatus of claim 17, wherein the second processing resource is further configured to execute the FSBL prior to the first processing resource measuring the SSBL.

20. The apparatus of claim 17, wherein the second processing resource is further configured to execute the FSBL prior to the second processing resource measuring the remainder the image.

---