

# US Patent & Trademark Office

## Patent Public Search | Text View

---

United States Patent	12386742
Kind Code	B2
Date of Patent	August 12, 2025
Inventor(s)	Ray; Anirban et al.

---

### Memory virtualization for accessing heterogeneous memory components

---

#### Abstract

A computing system having at least one bus, a plurality of different memory components, and a processing device operatively coupled with the plurality of memory components through the at least one bus. The different memory components include first memory and second memory having different memory access speeds. The computing system further includes a memory virtualizer operatively to: store an address map between first addresses used by the processing device to access memory and second addresses used to access the first memory and the second memory; monitor usages of the first memory and the second memory; adjust the address map based on the usages to improve speed of the processing device in memory access involving the first memory and the second memory; and swap data content in the first memory and the second memory according to adjustments to the address map.

---

**Inventors:** Ray; Anirban (Santa Clara, CA), Maharana; Parag R. (Dublin, CA), Anand; Gurpreet (Pleasanton, CA)

**Applicant:** Micron Technology, Inc. (Boise, ID)

**Family ID:** 1000008748372

**Assignee:** Micron Technology, Inc. (Boise, ID)

**Appl. No.:** 17/888398

**Filed:** August 15, 2022

#### Prior Publication Data

<b>Document Identifier</b>	<b>Publication Date</b>
US 20220398194 A1	Dec. 15, 2022

#### Related U.S. Application Data

## Publication Classification

**Int. Cl.:** **G06F12/0802** (20160101); **G06F9/50** (20060101); **G06F12/02** (20060101); **G06F16/958** (20190101)

**U.S. Cl.:**

**CPC** **G06F12/0802** (20130101); **G06F9/5077** (20130101); **G06F12/0238** (20130101);  
G06F16/972 (20190101)

## Field of Classification Search

**CPC:** G06F (12/0802); G06F (12/0238); G06F (9/5077); G06F (16/972)

---

## References Cited

### U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
5390334	12/1994	Harrison	N/A	N/A
5727150	12/1997	Laudon et al.	N/A	N/A
5909540	12/1998	Carter et al.	N/A	N/A
5918229	12/1998	Davis et al.	N/A	N/A
6026475	12/1999	Woodman	N/A	N/A
6230260	12/2000	Luick	N/A	N/A
6247097	12/2000	Sinharoy	N/A	N/A
6279138	12/2000	Jadav et al.	N/A	N/A
6473845	12/2001	Hornung et al.	N/A	N/A
7376681	12/2007	Todd et al.	N/A	N/A
8082400	12/2010	Chang et al.	N/A	N/A
8117373	12/2011	Berlin	N/A	N/A
8131814	12/2011	Schlansker et al.	N/A	N/A
8135933	12/2011	Fisher et al.	N/A	N/A
8316187	12/2011	Pothireddy	N/A	N/A
8352709	12/2012	Glasco et al.	N/A	N/A
8560761	12/2012	Tzeng	N/A	N/A
8700724	12/2013	McDaniel et al.	N/A	N/A
8799554	12/2013	Vincent et al.	N/A	N/A
8825937	12/2013	Atkisson et al.	N/A	N/A
8838887	12/2013	Burke et al.	N/A	N/A
8868842	12/2013	Yano et al.	N/A	N/A
8930647	12/2014	Smith	711/E12.008	G11C 7/1072
8965819	12/2014	Tirunagari	N/A	N/A
8996834	12/2014	Brenner et al.	N/A	N/A
9043530	12/2014	Sundaram et al.	N/A	N/A

9047017	12/2014	Dolan et al.	N/A	N/A
9104555	12/2014	Liebowitz et al.	N/A	N/A
9122503	12/2014	Hoff	N/A	N/A
9342453	12/2015	Nale et al.	N/A	N/A
9535740	12/2016	Graham et al.	N/A	N/A
9619408	12/2016	Nale et al.	N/A	N/A
9817739	12/2016	Pise et al.	N/A	N/A
9910618	12/2017	Curley et al.	N/A	N/A
9996370	12/2017	Khafizov et al.	N/A	N/A
10120797	12/2017	Foley et al.	N/A	N/A
10223371	12/2018	Thirumal	N/A	N/A
10241943	12/2018	Nale et al.	N/A	N/A
10282322	12/2018	Nale et al.	N/A	N/A
10282323	12/2018	Nale et al.	N/A	N/A
10289566	12/2018	Dalmatov et al.	N/A	N/A
10298496	12/2018	Nakibly et al.	N/A	N/A
10394789	12/2018	Animesh et al.	N/A	N/A
10430723	12/2018	Tzur et al.	N/A	N/A
10514855	12/2018	Voigt	N/A	N/A
10725663	12/2019	Awasthi et al.	N/A	N/A
10782908	12/2019	Mittal et al.	N/A	N/A
10852949	12/2019	Frolikov et al.	N/A	N/A
10877892	12/2019	Ray et al.	N/A	N/A
10880401	12/2019	Maharana et al.	N/A	N/A
11099789	12/2020	Maharana et al.	N/A	N/A
11354056	12/2021	Mittal et al.	N/A	N/A
11416395	12/2021	Ray et al.	N/A	N/A
2002/0145919	12/2001	Lamb et al.	N/A	N/A
2002/0196659	12/2001	Hurst	365/175	G11C 17/16
2003/0126232	12/2002	Mogul et al.	N/A	N/A
2004/0186960	12/2003	Poggio	N/A	N/A
2006/0087893	12/2005	Nishihara et al.	N/A	N/A
2006/0095679	12/2005	Edirisooriya	N/A	N/A
2006/0206658	12/2005	Hendel et al.	N/A	N/A
2007/0033367	12/2006	Sakarda et al.	N/A	N/A
2007/0074093	12/2006	Lasser	N/A	N/A
2007/0112864	12/2006	Ben-natan	N/A	N/A
2008/0016297	12/2007	Bartley et al.	N/A	N/A
2008/0276038	12/2007	Tanaka et al.	N/A	N/A
2009/0113422	12/2008	Kani	N/A	N/A
2009/0150639	12/2008	Ohata	N/A	N/A
2009/0248957	12/2008	Tzeng	N/A	N/A
2010/0011169	12/2009	Pothireddy	N/A	N/A
2010/0017650	12/2009	Chin et al.	N/A	N/A
2010/0082899	12/2009	Nakajima et al.	N/A	N/A
2010/0281230	12/2009	Rabii et al.	N/A	N/A
2010/0293412	12/2009	Sakaguchi et al.	N/A	N/A
2011/0161554	12/2010	Selinger et al.	N/A	N/A
2011/0231857	12/2010	Zaroo et al.	N/A	N/A

2011/0238887	12/2010	Bazzani	N/A	N/A
2011/0238899	12/2010	Yano et al.	N/A	N/A
2012/0047312	12/2011	Nathuji	711/170	G06F 9/3832
2012/0054419	12/2011	Chen et al.	N/A	N/A
2012/0084497	12/2011	Subramaniam et al.	N/A	N/A
2012/0117304	12/2011	Worthington et al.	N/A	N/A
2012/0124572	12/2011	Cunningham et al.	N/A	N/A
2012/0226850	12/2011	Nakanishi	N/A	N/A
2012/0246403	12/2011	Mchale et al.	N/A	N/A
2012/0297121	12/2011	Gorobets et al.	N/A	N/A
2012/0297122	12/2011	Gorobets et al.	N/A	N/A
2013/0024609	12/2012	Gorobets et al.	N/A	N/A
2013/0111113	12/2012	Harari et al.	N/A	N/A
2013/0124811	12/2012	Hamilton et al.	N/A	N/A
2013/0145095	12/2012	McKean et al.	N/A	N/A
2013/0151761	12/2012	Kim et al.	N/A	N/A
2013/0152086	12/2012	Yoo	718/1	G06F 9/455
2013/0179632	12/2012	Ben-shemesh et al.	N/A	N/A
2013/0226837	12/2012	Lymberopoulos et al.	N/A	N/A
2014/0006740	12/2013	Tokusho et al.	N/A	N/A
2014/0032818	12/2013	Chang et al.	N/A	N/A
2014/0040550	12/2013	Bill et al.	N/A	N/A
2014/0089631	12/2013	King	N/A	N/A
2014/0164676	12/2013	Borchers et al.	N/A	N/A
2014/0207954	12/2013	Calo	709/226	G06F 9/00
2015/0016046	12/2014	Shaeffer	N/A	N/A
2015/0026509	12/2014	Zhang et al.	N/A	N/A
2015/0032921	12/2014	Malkin	N/A	N/A
2015/0067087	12/2014	Guerin et al.	N/A	N/A
2015/0082062	12/2014	Saraswat et al.	N/A	N/A
2015/0097851	12/2014	Anderson et al.	N/A	N/A
2015/0106656	12/2014	Bian et al.	N/A	N/A
2015/0134927	12/2014	Varanasi	N/A	N/A
2015/0160858	12/2014	Okada et al.	N/A	N/A
2015/0169238	12/2014	Lee et al.	N/A	N/A
2015/0199276	12/2014	Radhakrishnan et al.	N/A	N/A
2015/0212741	12/2014	Lee et al.	N/A	N/A
2015/0227465	12/2014	Sundaram et al.	N/A	N/A
2015/0268875	12/2014	Jeddeloh	N/A	N/A
2015/0278091	12/2014	Wilkerson et al.	N/A	N/A
2015/0356125	12/2014	Golander et al.	N/A	N/A
2016/0042005	12/2015	Liu et al.	N/A	N/A
2016/0054922	12/2015	Awasthi et al.	N/A	N/A
2016/0125048	12/2015	Hamada	N/A	N/A
2016/0152086	12/2015	Kawakita	N/A	N/A
2016/0188218	12/2015	Gray et al.	N/A	N/A

2016/0188700	12/2015	Kleinschnitz, Jr et al.	N/A	N/A
2016/0210167	12/2015	Bolic et al.	N/A	N/A
2016/0210251	12/2015	Nale et al.	N/A	N/A
2016/0210465	12/2015	Craske et al.	N/A	N/A
2016/0212214	12/2015	Rahman et al.	N/A	N/A
2016/0253263	12/2015	Takada	N/A	N/A
2016/0294710	12/2015	Sreeramoju	N/A	N/A
2016/0306557	12/2015	Koseki et al.	N/A	N/A
2016/0350236	12/2015	Tsirkin et al.	N/A	N/A
2016/0371014	12/2015	Roberts	N/A	N/A
2016/0371496	12/2015	Sell	N/A	G06F 21/604
2017/0039164	12/2016	Nikolas et al.	N/A	N/A
2017/0060754	12/2016	Nakra et al.	N/A	N/A
2017/0060769	12/2016	Wires et al.	N/A	N/A
2017/0123796	12/2016	Kumar et al.	N/A	N/A
2017/0131902	12/2016	Goss et al.	N/A	N/A
2017/0147427	12/2016	Nero	N/A	N/A
2017/0177486	12/2016	Horn	N/A	N/A
2017/0187621	12/2016	Shalev et al.	N/A	N/A
2017/0199666	12/2016	Sundaram et al.	N/A	N/A
2017/0249266	12/2016	Nale et al.	N/A	N/A
2017/0255383	12/2016	Chang et al.	N/A	N/A
2017/0262215	12/2016	Banerjee et al.	N/A	N/A
2017/0285967	12/2016	Pandurangan et al.	N/A	N/A
2017/0285992	12/2016	Vogt	N/A	N/A
2017/0302734	12/2016	Liang et al.	N/A	N/A
2017/0316321	12/2016	Whitney et al.	N/A	N/A
2017/0364422	12/2016	Antony et al.	N/A	N/A
2017/0364450	12/2016	Struttmann	N/A	N/A
2018/0004441	12/2017	Takamura et al.	N/A	N/A
2018/0011790	12/2017	Gaur et al.	N/A	N/A
2018/0018379	12/2017	Eda	N/A	G06F 16/185
2018/0024853	12/2017	Warfield et al.	N/A	N/A
2018/0046581	12/2017	Banerjee et al.	N/A	N/A
2018/0089087	12/2017	Chang et al.	N/A	N/A
2018/0121366	12/2017	Tian	N/A	N/A
2018/0150219	12/2017	Chen	N/A	G06F 3/0685
2018/0189207	12/2017	Nale et al.	N/A	N/A
2018/0260135	12/2017	Hayashida et al.	N/A	N/A
2018/0293163	12/2017	Bergeron	N/A	N/A
2018/0316569	12/2017	Cilfone et al.	N/A	N/A
2018/0332366	12/2017	Paduroiu	N/A	N/A
2019/0004841	12/2018	Starks	N/A	N/A
2019/0018809	12/2018	Nale et al.	N/A	N/A
2019/0079689	12/2018	Cherubini et al.	N/A	N/A
2019/0129834	12/2018	Purkayastha et al.	N/A	N/A

2019/0129847	12/2018	Roh	N/A	N/A
2019/0179760	12/2018	Bhargava et al.	N/A	N/A
2019/0196996	12/2018	Balakrishnan et al.	N/A	N/A
2019/0243552	12/2018	Maharana et al.	N/A	N/A
2019/0243570	12/2018	Mittal et al.	N/A	N/A
2019/0243756	12/2018	Ray et al.	N/A	N/A
2019/0243771	12/2018	Mittal et al.	N/A	N/A
2019/0243787	12/2018	Mittal et al.	N/A	N/A
2019/0253520	12/2018	Maharana et al.	N/A	N/A
2019/0303283	12/2018	McGlaughlin	N/A	G06F 3/0616
2019/0310780	12/2018	Gholamipour et al.	N/A	N/A
2019/0332556	12/2018	Nale et al.	N/A	N/A
2019/0370043	12/2018	Olderdissen	N/A	G06F 3/0604
2020/0019506	12/2019	Ray et al.	N/A	N/A
2020/0073827	12/2019	Wallach	N/A	N/A
2020/0074094	12/2019	Wallach	N/A	N/A
2020/0319813	12/2019	Mittal et al.	N/A	N/A
2020/0326851	12/2019	Frolikov et al.	N/A	N/A
2021/0034524	12/2020	Leidel et al.	N/A	N/A
2021/0048947	12/2020	Frolikov et al.	N/A	N/A
2021/0117326	12/2020	Ray et al.	N/A	N/A
2021/0120099	12/2020	Maharana et al.	N/A	N/A
2021/0349638	12/2020	Maharana et al.	N/A	N/A

## FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
20140041408	12/2013	KR	N/A
20160116533	12/2015	KR	N/A

## OTHER PUBLICATIONS

“Solid-State Drive”, Wikipedia, printed on Mar. 14, 2018. cited by applicant

“Write combining,” Wikipedia encyclopedia entry located at

[https://en.wikipedia.org/wiki/Write\\_combining](https://en.wikipedia.org/wiki/Write_combining), Feb. 12, 2018. cited by applicant

A. Romanow, J. Mogul, T. Talpey, S. Bailey, “Remote Direct Memory Access (RDMA) over IP Problem Statement”, RFC 4297, Dec. 2005 ([https://www.rfc-editor.org/search/rfc\\_search\\_detail.php](https://www.rfc-editor.org/search/rfc_search_detail.php)). cited by applicant

Abhijith, et al. “The Efficient Use of Storage Resources in SAN for Storage Tiering and Caching,” 2015 International Conference on Computational Intelligence and Networks, IEEE, Jan. 11, 2016. cited by applicant

Cai, “Error Characterization, Mitigation and Recovery”, 2017. cited by applicant

Demand paging, Wikipedia, printed on Apr. 18, 2018. cited by applicant

Device driver, Wikipedia, printed on Apr. 19, 2018. cited by applicant

Extended European Search Report, EP19746665.9, mailed on May 3, 2021. cited by applicant

Extended European Search Report, EP19747107.1, mailed on Apr. 30, 2021. cited by applicant

Extended European Search Report, EP19746836.6, mailed on May 19, 2021. cited by applicant

Extended European Search Report, EP19748086.6, mailed on Jun. 8, 2021. cited by applicant

Extended European Search Report, EP19746972.9, mailed on May 19, 2021. cited by applicant

Extended European Search Report, EP19833924.4, mailed on Mar. 18, 2022. cited by applicant

Garbage collection (computer science), Wikipedia, printed on Sep. 26, 2018. cited by applicant

Graphics processing unit, Wikipedia, printed on Jul. 26, 2018. cited by applicant

Handy, Jim. "Where does NVRAM Fit?" Objective Analysis, Aug. 15, 2014. cited by applicant

Hypervisor, Wikipedia, printed on Apr. 19, 2018. cited by applicant

InfiniBand, Wikipedia, printed on Jul. 26, 2018. cited by applicant

Integrated circuit, Wikipedia, printed on Sep. 24, 2018. cited by applicant

Interlaken (networking), Wikipedia, printed on Sep. 21, 2018. cited by applicant

International Search Report and Written Opinion, PCT/US2019/013841, Apr. 23, 2019. cited by applicant

International Search Report and Written Opinion, PCT/US2019/013836, Apr. 22, 2019. cited by applicant

International Search Report and Written Opinion, PCT/US2019/015602, mailed May 3, 2019. cited by applicant

International Search Report and Written Opinion, PCT/US2019/014205, mailed May 1, 2019. cited by applicant

International Search Report and Written Opinion, PCT/US2019/014285, mailed May 8, 2019. cited by applicant

International Search Report, PCT/US2019/040413, mailed Oct. 25, 2019. cited by applicant

International Search Report and Written Opinion, PCT/US2019/014275, mailed May 8, 2019. cited by applicant

International Search Report and Written Opinion, PCT/US2020/021825, mailed Jul. 6, 2020. cited by applicant

Jacob, et al. "Optimized utilization of disks in storage area network by storage tiering," 2017 International Conference on Computing, Communication, and Automation 9ICCCA), IEEE, May 5, 2017. cited by applicant

Jeff Hawkins et al., "Sequence memory for prediction, inference and behaviour", In: Philosophical Transactions of the Royal Society B: Biological Sciences, pp. 1203-1209, May 1, 2009. cited by applicant

Kalnoskas, Aimee. "SSD controllers integrate DDR4 DRAM on chip." Microsemi, Aug. 5, 2016. cited by applicant

Message Passing Interface, Wikipedia, printed on Jul. 26, 2018. cited by applicant

Microsequencer, Wikipedia, printed on Sep. 21, 2018. cited by applicant

Oracle, "NFS Over ROMA", Retrieved on Feb. 18, 2020 from [https://web.archive.org/web/20150920225910/https://docs.oracle.com/cd/E23824\\_01/html/821-1454/rfsrefer-154.html#](https://web.archive.org/web/20150920225910/https://docs.oracle.com/cd/E23824_01/html/821-1454/rfsrefer-154.html#) (Year: 2015). cited by applicant

Operating system, Wikipedia, printed on Apr. 18, 2018. cited by applicant

PCI Express, Wikipedia, printed on Sep. 21, 2018. cited by applicant

Page cache, Wikipedia, printed on Apr. 18, 2018. cited by applicant

Page replacement algorithm, Wikipedia, printed on Jul. 31, 2018. cited by applicant

Page table, Wikipedia, printed on Jul. 31, 2018. cited by applicant

Paging, Wikipedia, printed on Apr. 18, 2018. cited by applicant

Palmer, Mark, et al. "Fido: A Cache That Learns to Fetch." Proceedings of the 17th International Conference on Very arge Data Bases, Sep. 1991. cited by applicant

Ppeled, Leeor, et al. "A neural network memory prefetcher using semantic locality." arXiv:1804.00478v2, Jul. 26, 2018. cited by applicant

RAM drive, Wikipedia, printed on Apr. 18, 2018. cited by applicant

Salkhordeh, et al. "Operating system level data tiering using online workload characterization," Journal of Supercomputing, Kluwer Academic Publishers, Jan. 31, 2015. cited by applicant

SerDes, Wikipedia, printed on Sep. 21, 2018. cited by applicant

Switched fabric, Wikipedia, printed on Jul. 26, 2018. cited by applicant  
Translation lookaside buffer, Wikipedia, printed on Apr. 18, 2018. cited by applicant  
Vatto, "Analyzing Intel-Micron 3D XPoint: The NextGeneration Non-Volatile Memory", 2015.  
cited by applicant  
Virtual memory, Wikipedia, printed on Apr. 18, 2018. cited by applicant  
Wilson, Tracy. "How PCI Express Works." Retrieved from the Internet  
<<https://computer.howstuffworks.com/pci-express.htm>>, Mar. 30, 2021. cited by applicant

---

*Primary Examiner:* Thammavong; Prasith

*Attorney, Agent or Firm:* Greenberg Traurig

---

## **Background/Summary**

RELATED APPLICATIONS (1) The present application is a continuation application of U.S. patent application Ser. No. 16/054,719, filed Aug. 3, 2018, which claims the benefit of the filing date of Prov. U.S. Pat. App. Ser. No. 62/626,491, filed on Feb. 5, 2018 and entitled "Memory Virtualization for Two Tier Memory Systems," the entire disclosure of which is hereby incorporated herein by reference. (2) The present application relates to U.S. patent application Ser. No. 16/032,331, filed Jul. 11, 2018 and entitled "Predictive Paging to Accelerate Memory Access," and U.S. patent application Ser. No. 16/035,469, filed Jul. 13, 2018 and entitled "Isolated Performance Domains in a Memory System," the entire disclosures of which applications are hereby incorporated herein by reference.

## **FIELD OF THE TECHNOLOGY**

(1) At least some embodiments disclosed herein relate to memory systems in general, and more particularly, but not limited to memory virtualization for accessing heterogeneous memory components.

## **BACKGROUND**

(2) A memory sub-system can be a memory module, such as a dual in-line memory module (DIMM), a small outline DIMM (SO-DIMM), or a non-volatile dual in-line memory module (NVDIMM). A memory sub-system can be a storage system, such as a solid-state drive (SSD), or a hard disk drive (HDD). A memory sub-system can include one or more memory components that store data. The memory components can be, for example, non-volatile memory components and volatile memory components. Examples of memory components include memory integrated circuits. Some memory integrated circuits are volatile and require power to maintain stored data. Some memory integrated circuits are non-volatile and can retain stored data even when not powered. Examples of non-volatile memory include flash memory, Read-Only Memory (ROM), Programmable Read-Only Memory (PROM), Erasable Programmable Read-Only Memory (EPROM) and Electronically Erasable Programmable Read-Only Memory (EEPROM) memory, etc. Examples of volatile memory include Dynamic Random-Access Memory (DRAM) and Static Random-Access Memory (SRAM). In general, a host system can utilize a memory sub-system to store data at the memory components and to retrieve data from the memory components.

(3) For example, a computer can include a host system and one or more memory sub-systems attached to the host system. The host system can have a central processing unit (CPU) in communication with the one or more memory sub-systems to store and/or retrieve data and instructions. Instructions for a computer can include operating systems, device drivers, and application programs. An operating system manages resources in the computer and provides common services for application programs, such as memory allocation and time sharing of the



resources. A device driver operates or controls a particular type of devices in the computer; and the operating system uses the device driver to offer resources and/or services provided by the type of devices. A central processing unit (CPU) of a computer system can run an operating system and device drivers to provide the services and/or resources to application programs. The central processing unit (CPU) can run an application program that uses the services and/or resources. For example, an application program implementing a type of applications of computer systems can instruct the central processing unit (CPU) to store data in the memory components of a memory sub-system and retrieve data from the memory components.

(4) An operating system of a computer system can allow an application program to use virtual addresses of memory to store data in, or retrieve data from, memory components of one or more memory sub-systems of the computer system. The operating system maps the virtual addresses to physical addresses of one or more memory sub-systems connected to the central processing unit (CPU) of the computer system. The operating system implements the memory accesses specified at virtual addresses using the physical addresses of the memory sub-systems.

(5) A virtual address space can be divided into pages. A page of virtual memory can be mapped to a page of physical memory in the memory sub-systems. The operating system can use a paging technique to access a page of memory in a storage device via a page of memory in a memory module. At different time instances, the same page of memory in a memory module can be used as proxy to access different pages of memory in the storage device or another storage device in the computer system.

(6) A computer system can include a hypervisor (or virtual machine monitor) to create or provision virtual machines. A virtual machine is a computing device that is virtually implemented using the resources and services available in the computer system. The hypervisor presents the virtual machine to an operating system as if the components of virtual machine were dedicated physical components. A guest operating system runs in the virtual machine to manage resources and services available in the virtual machine, in a way similar to the host operating system running in the computer system. The hypervisor allows multiple virtual machines to share the resources of the computer system and allows the virtual machines to operate on the computer substantially independently from each other.

---

## Description

### BRIEF DESCRIPTION OF THE DRAWINGS

(1) The embodiments are illustrated by way of example and not limitation in the figures of the accompanying drawings in which like references indicate similar elements.

(2) FIG. 1 illustrates an example computing system having a memory sub-system in accordance with some embodiments of the present disclosure.

(3) FIG. 2 shows a computing system having different types of memory and running instructions to virtualize memory access in accordance with at least some embodiments disclosed herein.

(4) FIG. 3 shows a system having a device driver configured to virtualize memory access made via a hypervisor.

(5) FIG. 4 shows a prediction engine coupled to a device driver to improve memory access through virtualization.

(6) FIG. 5 shows method to virtualize memory access.

(7) FIG. 6 is a block diagram of an example computer system in which embodiments of the present disclosure can operate.

### DETAILED DESCRIPTION

(8) At least some aspects of the present disclosure are directed to a memory virtualization technique that can improve memory access performance in a computer system having heterogeneous memory

components. A memory sub-system is also hereinafter referred to as a “memory device”. An example of a memory sub-system is a memory module that is connected to a central processing unit (CPU) via a memory bus. Examples of memory modules include a dual in-line memory module (DIMM), a small outline DIMM (SO-DIMM), a non-volatile dual in-line memory module (NVDIMM), etc. Another example of a memory sub-system is a storage device that is connected to the central processing unit (CPU) via a peripheral interconnect (e.g., an input/output bus, a storage area network). Examples of storage devices include a solid-state drive (SSD), a flash drive, a universal serial bus (USB) flash drive, and a hard disk drive (HDD). In some embodiments, the memory sub-system is a hybrid memory/storage sub-system that provides both memory functions and storage functions. In general, a host system can utilize a memory sub-system that includes one or more memory components. The host system can provide data to be stored at the memory sub-system and can request data to be retrieved from the memory sub-system.

(9) A memory sub-system can have different types of memory components that offer different advantages in different aspects. For example, a memory module can have volatile DRAM for fast access and non-volatile memory for persistent storage. However, it is a challenge for applications to program the usages of the different types of memory components for full utilization of the different advantages offered by different types of memory components.

(10) At least some aspects of the present disclosure address the above and other deficiencies by virtualizing memory access to the different types of memory components such that the differences in the memory components are shielded from the applications. Memory accesses to the different types of memory components are made through a memory device driver. Based on the current and/or past memory usage patterns, the device drive adjusts what data is stored in what memory components to improve the performance of the memory sub-system as a whole for a current time period. The device drive automatically manages the usage of the different types of memory components based on past or current usage patterns and/or predicted usage patterns, such that the different advantages of the different types of memory components can be fully utilizes.

(11) FIG. 1 illustrates an example computing system **100** having a memory sub-system **110** in accordance with some embodiments of the present disclosure. The memory sub-system **110** can include media, such as memory components **109A** to **109N**. The memory components **109A** to **109N** can be volatile memory components, non-volatile memory components, or a combination of such. In some embodiments, the memory sub-system **110** is a memory module. Examples of a memory module includes a DIMM, NVDIMM, and NVDIMM-P. In some embodiments, the memory sub-system is a storage system. An example of a storage system is an SSD. In some embodiments, the memory sub-system **110** is a hybrid memory/storage sub-system. In general, the computing environment can include a host system **120** that uses the memory sub-system **110**. For example, the host system **120** can write data to the memory sub-system **110** and read data from the memory sub-system **110**.

(12) The host system **120** can be a computing device such as a desktop computer, laptop computer, network server, mobile device, or such computing device that includes a memory and a processing device. The host system **120** can include or be coupled to the memory sub-system **110** so that the host system **120** can read data from or write data to the memory sub-system **110**. The host system **120** can be coupled to the memory sub-system **110** via a physical host interface. As used herein, “coupled to” generally refers to a connection between components, which can be an indirect communicative connection or direct communicative connection (e.g., without intervening components), whether wired or wireless, including connections such as electrical, optical, magnetic, etc. Examples of a physical host interface include, but are not limited to, a serial advanced technology attachment (SATA) interface, a peripheral component interconnect express (PCIe) interface, universal serial bus (USB) interface, Fibre Channel, Serial Attached SCSI (SAS), a double data rate (DDR) memory bus, etc. The physical host interface can be used to transmit data between the host system **120** and the memory sub-system **110**. The host system **120** can further

utilize an NVM Express (NVMe) interface to access the memory components **109A** to **109N** when the memory sub-system **110** is coupled with the host system **120** by the PCIe interface. The physical host interface can provide an interface for passing control, address, data, and other signals between the memory sub-system **110** and the host system **120**. FIG. **1** illustrates a memory sub-system **110** as an example. In general, the host system **120** can access multiple memory sub-systems via a same communication connection, multiple separate communication connections, and/or a combination of communication connections.

(13) The host system **120** includes a processing device **118** and a controller **116**. The processing device **118** of the host system **120** can be, for example, a microprocessor, a central processing unit (CPU), a processing core of a processor, an execution unit, etc. In some instances, the controller **116** can be referred to as a memory controller, a memory management unit, and/or an initiator. In one example, the controller **116** controls the communications over a bus coupled between the host system **120** and the memory sub-system **110**.

(14) In general, the controller **116** can send commands or requests to the memory sub-system **110** for desired access to memory components **109A** to **109N**. The controller **116** can further include interface circuitry to communicate with the memory sub-system **110**. The interface circuitry can convert responses received from memory sub-system **110** into information for the host system **120**.

(15) The controller **116** of the host system **120** can communicate with controller **115** of the memory sub-system **110** to perform operations such as reading data, writing data, or erasing data at the memory components **109A** to **109N** and other such operations. In some instances, the controller **116** is integrated within the same package of the processing device **118**. In other instances, the controller **116** is separate from the package of the processing device **118**. The controller **116** and/or the processing device **118** can include hardware such as one or more integrated circuits and/or discrete components, a buffer memory, a cache memory, or a combination thereof. The controller **116** and/or the processing device **118** can be a microcontroller, special purpose logic circuitry (e.g., a field programmable gate array (FPGA), an application specific integrated circuit (ASIC), etc.), or another suitable processor.

(16) The memory components **109A** to **109N** can include any combination of the different types of non-volatile memory components and/or volatile memory components. An example of non-volatile memory components includes a negative-and (NAND) type flash memory. Each of the memory components **109A** to **109N** can include one or more arrays of memory cells such as single level cells (SLCs) or multi-level cells (MLCs) (e.g., triple level cells (TLCs) or quad-level cells (QLCs)). In some embodiments, a particular memory component can include both an SLC portion and a MLC portion of memory cells. Each of the memory cells can store one or more bits of data (e.g., data blocks) used by the host system **120**. Although non-volatile memory components such as NAND type flash memory are described, the memory components **109A** to **109N** can be based on any other type of memory such as a volatile memory. In some embodiments, the memory components **109A** to **109N** can be, but are not limited to, random access memory (RAM), read-only memory (ROM), dynamic random access memory (DRAM), synchronous dynamic random access memory (SDRAM), phase change memory (PCM), magnetoresistive random access memory (MRAM), Spin Transfer Torque (STT)-MRAM, ferroelectric random-access memory (FeTRAM), ferroelectric RAM (FeRAM), conductive bridging RAM (CBRAM), resistive random access memory (RRAM), oxide based RRAM (OxRAM), negative-or (NOR) flash memory, electrically erasable programmable read-only memory (EEPROM), nanowire-based non-volatile memory, memory that incorporates memristor technology, and a cross-point array of non-volatile memory cells. A cross-point array of non-volatile memory can perform bit storage based on a change of bulk resistance, in conjunction with a stackable cross-gridded data access array. Additionally, in contrast to many flash-based memories, cross-point non-volatile memory can perform a write in-place operation, where a non-volatile memory cell can be programmed without the non-volatile memory cell being previously erased. Furthermore, the memory cells of the memory components **109A** to

**109N** can be grouped as memory pages or data blocks that can refer to a unit of the memory component used to store data.

(17) The controller **115** of the memory sub-system **110** can communicate with the memory components **109A** to **109N** to perform operations such as reading data, writing data, or erasing data at the memory components **109A** to **109N** and other such operations (e.g., in response to commands scheduled on a command bus by controller **116**). The controller **115** can include hardware such as one or more integrated circuits and/or discrete components, a buffer memory, or a combination thereof. The controller **115** can be a microcontroller, special purpose logic circuitry (e.g., a field programmable gate array (FPGA), an application specific integrated circuit (ASIC), etc.), or another suitable processor. The controller **115** can include a processing device **117** (processor) configured to execute instructions stored in local memory **119**. In the illustrated example, the local memory **119** of the controller **115** includes an embedded memory configured to store instructions for performing various processes, operations, logic flows, and routines that control operation of the memory sub-system **110**, including handling communications between the memory sub-system **110** and the host system **120**. In some embodiments, the local memory **119** can include memory registers storing memory pointers, fetched data, etc. The local memory **119** can also include read-only memory (ROM) for storing micro-code. While the example memory sub-system **110** in FIG. 1 has been illustrated as including the controller **115**, in another embodiment of the present disclosure, a memory sub-system **110** may not include a controller **115**, and can instead rely upon external control (e.g., provided by an external host, or by a processor or controller separate from the memory sub-system).

(18) In general, the controller **115** can receive commands or operations from the host system **120** and can convert the commands or operations into instructions or appropriate commands to achieve the desired access to the memory components **109A** to **109N**. The controller **115** can be responsible for other operations such as wear leveling operations, garbage collection operations, error detection and error-correcting code (ECC) operations, encryption operations, caching operations, and address translations between a logical block address and a physical block address that are associated with the memory components **109A** to **109N**. The controller **115** can further include host interface circuitry to communicate with the host system **120** via the physical host interface. The host interface circuitry can convert the commands received from the host system into command instructions to access the memory components **109A** to **109N** as well as convert responses associated with the memory components **109A** to **109N** into information for the host system **120**.

(19) The memory sub-system **110** can also include additional circuitry or components that are not illustrated. In some embodiments, the memory sub-system **110** can include a cache or buffer (e.g., DRAM) and address circuitry (e.g., a row decoder and a column decoder) that can receive an address from the controller **115** and decode the address to access the memory components **109A** to **109N**.

(20) The computing system **100** includes a memory virtualizer **113** in the host system **120** that can virtualize memory access to the memory components **109A** to **109N** by shielding the differences in the memory components **109A** to **109N** from applications running in the processing device **118**. In some embodiments, the controller **116** and/or the processing device **118** in the host system **120** includes at least a portion of the memory virtualizer **113**. For example, the controller **116** and/or the processing device **118** can include logic circuitry implementing the memory virtualizer **113**. For example, the processing device **118** (processor) of the host system **120** can be configured to execute instructions stored in memory for performing the operations of the memory virtualizer **113** described herein. In some embodiments, the memory virtualizer **113** is part of an operating system of the host system **120**, a device driver, or an application. Optionally, the memory virtualizer **113** can be implemented partially, or entirely, in the controller **115** of the memory sub-system **110**.

(21) The memory virtualizer **113** offers to access the memory components **109A** to **109N** using a set of virtualized memory addresses. Some of the virtualized memory addresses are used to access

one type of memory components (e.g., faster but volatile memory **109A**); and some of the virtualized memory addresses are used to access another type of memory components (e.g., non-volatile but slower memory **109N**). However, there is no predetermined association between groups of virtualized memory addresses and the different types of memory components **109A** to **109N**. The memory virtualizer **113** can adjust the mapping between the virtualized memory addresses and the different types of memory components **109A** to **109N** dynamically to fully utilizes the different advantages offered by the different types of memory components. For example, a predictive model can be trained in predicting hot pages of virtual memory that will be used soon and cold pages of virtual memory that will not be used for a period of time. Based on the prediction, the memory virtualizer **113** automatically adjusts the mapping and instructs the controller **115** of the memory sub-system **110** to swap the hot pages to the faster memory (e.g., **109A**) and swap the cold pages to the slower memory (e.g., **109N**). The page swapping can be performed by the controller **115** within the memory sub-system **110** without utilizing the communication bandwidth on the connection between the memory sub-system **110** and the host system **120**. After the mapping adjustment and page swapping, the memory access performance of the memory sub-system **110** for the predicted memory usage is improved. The memory virtualizer **113** shields the memory differences from applications. Thus, the applications running in the processing device **118** can make memory access requests as if the memory components **109A** to **109N** were a pool of uniform virtual memory having the same characteristics. The memory virtualizer **113** dynamically adjusts the data distribution in the memory components **109A** to **109N** to utilize the advantages offered by the different types of memory components **109A** to **109N**. Further details with regards to the operations of the memory virtualizer **113** are described below.

(22) In one example, the central processing unit (CPU) can access two sets of memory provided in one or more memory sub-systems connected to the CPU. For example, one set of memory can be slower than the other set of memory; and the central processing unit (CPU) uses the memory virtualizer **113** to access the slower set of memory via the faster set of memory using a paging technique. In some instances, one set of memory cannot be directly addressable by the CPU and is coupled to the other set of memory that is directly addressable by the CPU; and the central processing unit (CPU) uses the memory virtualizer **113** to access a set of memory that is not directly addressable via the set of memory that is directly addressable in a way similar to the use of the paging technique.

(23) For example, when a page of virtual memory being accessed is currently in the slower set of memory but not in the faster set of memory, a page can be allocated from the faster set of memory to service the page in the slower set of memory; and the memory virtualizer **113** is to fetch the data of the page from the slower set of memory, store the fetched page of data in the allocated page in the faster set of memory, such that the data access of the page of the virtual memory can be made via accessing the allocated page in the faster set of memory in subsequent operations. In some instances, swapping a page takes a time period longer than simply accessing a requested data element from the slower memory. Thus, the requested data element is first serviced to the requester, while the page swapping is performed to speed up subsequent access to the data elements in the hot page. Using such an approach, the overall performance can be better than holding the request for the data element until the page swap is completed.

(24) Further, the memory virtualizer **113** can receive information related to the use of the pages in the slower set of memory, and train a self-learning prediction engine to predict the use of the pages based on such information. For example, a supervised machine learning technique can be used to train, using the information, an artificial neural network to predict the use of the pages in the slower set of memory by reducing the errors between predictions and the actual use of the pages. After the training of the artificial neural network, the prediction engine can use the current information to predict the next pages to be used. Further, the training, prediction, and feedback from the actual usage following the prediction for further training can be performed in a continuous fashion to

adapt the prediction model of the artificial neural network to the most recent usage patterns of memory pages.

(25) In response to the memory usage prediction, made based on the trained artificial neural network and the current information, that a page in the slower set of memory is to be used soon, the memory virtualizer **113** can instruct the controller **115** to proactively swap the page of data from the slower set of memory to the faster set of memory, such that when needed for processing, the page of data is already in the faster set of memory, which arrangement improves the data access speed of the page of data.

(26) The accuracy of the prediction can be measured against the subsequent actual page use; and the prediction and the subsequent actual page use can be used to further train or adjust the artificial neural network to track the most recent usage patterns of memory pages.

(27) Alternatively, or in combination, the machine learning-based prediction can be replaced or augmented with policy based prediction rules. For example, pages storing resident codes (e.g., in lower addresses) can be maintained in the faster set of memory when possible to reduce swapping of frequently used pages. For example, a huge page can be loaded into the faster set of memory when a page that is a portion of the huge page is being accessed. For example, predictions can be made at least in part using heuristic rules, based on indications such as whether the pages are accessed sequentially or randomly, whether the data access is in a steady state mode or in a bursty mode, and/or the logical relations between pages (and pages of different sizes).

(28) Some details and examples regarding the prediction techniques can be found in U.S. patent application Ser. No. 16/032,331, filed Jul. 11, 2018 and entitled “Predictive Paging to Accelerate Memory Access”.

(29) FIG. 2 shows a computing system having different types of memory and running instructions to virtualize memory access in accordance with at least some embodiments disclosed herein.

(30) The computing system of FIG. 2 includes a host system **120**, a memory module **205** connected to the host system **120** via a memory bus **203**, and a storage device **209** connected to the host system **120** via a peripheral interconnect **207**. The storage device **209** and the memory module **205** are examples of the memory sub-system **110** illustrated in FIG. 1.

(31) The host system **120** has a processing device **118**, which can be a central processing unit or a microprocessor with one or more processing cores. The host system **120** can have a memory management unit **213** and cache memory **211**. The memory management unit **213** and/or at least a portion of the cache memory **211** can be optionally integrated within the same integrated circuit package of the processing device **118**.

(32) The memory module **205** illustrated in FIG. 2 has multiple types of memory (e.g., **221** and **223**). For example, memory of type A **221** is faster than memory of type B **223**.

(33) The memory bus **203** typically operates at a speed faster than the peripheral interconnect **207**. For example, the memory bus **203** can be a double data rate bus; and the peripheral interconnect **207** can be a peripheral component interconnect express (PCIe) bus, a serial advanced technology attachment (SATA) bus, a universal serial bus (USB) bus, and/or a storage area network. Memory of type B **223** in the memory module **205** can be accessed at a speed faster than accessing memory of type B **223** in the storage device **209**.

(34) The storage device **209** illustrated in FIG. 2 has multiple types of memory (e.g., **223** and **225**). For example, memory of type B **223** is faster than memory of type C **225**.

(35) In general, a plurality of memory modules (e.g., **205**) can be coupled to the memory bus **203**; and a plurality of storage devices (e.g., **209**) can be coupled to the peripheral interconnect **207**. In some instances, the peripheral interconnect **207** and the storage devices (e.g., **209**) are optional and can be absent from the computing system. In other instances, the memory bus **203** and the memory modules (e.g., **205**) can be optional and can be absent from the computing system.

(36) In a possible configuration when a plurality of memory modules (e.g., **205**) are coupled to the memory bus **203**, one of the memory modules (e.g., **205**) has memory of type A **221**; and another

of the memory modules has memory of type B **223** that is accessible at a speed lower than the memory of type A **221** in a separate memory module (e.g., **205**).

(37) Similarly, in a possible configuration when a plurality of storage devices (e.g., **209**) are coupled to the peripheral interconnect **207**, one of the storage device (e.g., **209**) has memory of type B **223**, and another of the storage devices has memory of type C **225** that is accessible at a speed lower than the memory of type B **223** in a separate storage device (e.g., **209**).

(38) The processing device **118** and/or the MMU **213** are configured via instructions (e.g., an operating system and/or one or more device drivers) to access a portion of memory in the computer system via another portion of memory in the computer system using a paging technique and/or a memory map interface.

(39) For example, memory of type B **223** of the memory module **205** can be accessed via memory of type A **221** of the memory module **205** (or another memory module).

(40) For example, memory of type B **223** of the storage device **207** can be accessed via memory of type A **221** of the memory module **205** and/or via memory of type B **223** of the memory module **205**.

(41) For example, memory of type C **225** of the storage device **207** can be accessed via memory of type A **221** of the memory module **205**, via memory of type B **223** of the memory module **205**, and/or via memory of type B **223** of the storage device **207** (or another storage device).

(42) For example, in some instances, memory of type A **221** and memory of type B **223** in the same memory module **205** (or different memory modules) are addressable directly and separately over the memory bus **203** by the memory management unit **213** of the processing device **118**. However, since the memory of type B **223** is slower than memory of type A **221**, it is desirable to access the memory of type B **223** via the memory of type A **221**.

(43) In other instances, memory of type B **223** of the memory module **205** is accessible only through addressing the memory of type A **221** of the memory module **205** (e.g., due to the size restriction in the address portion of the memory bus **203**).

(44) A controller X **227** can be provided in the memory module **205** to manage data transfer between the memory of type A **221** and the memory of type B **223** within the memory module **205**, especially when the memory of type B **223** of the memory module **205** is not directly addressable using the memory bus **203**.

(45) In one variation, the memory (e.g., **221** and **223**) of the memory module **205** can have the same performance individually within the memory module **205**; however, the memory management unit **213** and/or the processing device **118** are restricted to access via the memory **223** via the memory **221** (e.g., due to the size restriction in the address portion of the memory bus **203**). Thus, the memory **223** appears to be slower than the memory **221** to the processing device **118**.

(46) In general, the memory sub-systems (e.g., **205** and **209**) can include media, such as memory (e.g., **221**, . . . , **223**, . . . , **225**). The memory (e.g., **221**, . . . , **223**, . . . , **225**) can includes volatile memory, non-volatile memory (NVM), and/or a combination of such. In some embodiments, the computer system includes at least one memory sub-system that is a storage device **209**. An example of a storage device **209** is a solid-state drive (SSD). In some embodiments, the computer system includes at least one memory sub-system that is a hybrid memory/storage system configured as a memory module **205**. The processing device **118** can write data to each of the memory sub-systems (e.g., **205** and **209**) and read data from the memory sub-systems (e.g., **205** and **209**) directly or indirectly.

(47) The computing system of FIG. 2 can be used to implement a desktop computer, laptop computer, network server, mobile device, or such computing device that includes a memory and a processing device. The processing device **118** can read data from or write data to the memory sub-systems (e.g., **205** and **209**).

(48) The processing device **118** can be coupled to a memory sub-system (e.g., **209**) via one or more physical interface (e.g., **203**, **207**).

(49) As used herein, “coupled to” generally refers to a connection between components, which can be an indirect communicative connection or direct communicative connection (e.g., without intervening components), whether wired or wireless, including connections such as, electrical, optical, magnetic, etc.

(50) Examples of a physical host interface include, but are not limited to, a serial advanced technology attachment (SATA) interface, a peripheral component interconnect express (PCIe) interface, universal serial bus (USB) interface, Fibre Channel, Small Computer System Interface (SCSI), Serial Attached SCSI (SAS), etc.

(51) The physical host interface can be used to transmit data between the processing device **118** and the memory sub-system (e.g., **209**). The computer system can further utilize an NVM Express (NVMe) interface to access the memory (e.g., **223**, . . . , **225**) when the memory sub-system **209** is coupled with the peripheral interconnect **207** via the PCIe interface. The physical host interface can provide an interface for passing control, address, data, and other signals between the memory sub-system (e.g., **209**) and the processing device **118**.

(52) In general, a memory sub-system (e.g., **205** and **209**) includes a printed circuit board that connects a set of memory devices, such as memory integrated circuits, that provides the memory (e.g., **221**, . . . , **223**, . . . , **225**). The memory (e.g., **221**, . . . , **223**, . . . , **225**) on the memory sub-system (e.g., **205** and **209**) can include any combination of the different types of non-volatile memory devices and/or volatile memory devices.

(53) An example of non-volatile memory devices includes a negative-and (NAND) type flash memory or a negative-or (NOR) type flash memory. A memory integrated circuit can include one or more arrays of memory cells, such as single level cells (SLCs), multi-level cells (MLCs), triple level cells (TLCs), quad-level cells (QLCs), etc. In some implementations, a particular memory device can include both an SLC portion and a MLC (or TLC or QLC) portion of memory cells. Each of the memory cells can store one or more bits of data used by the host system **120**. Although non-volatile memory devices such as NAND type flash memory are described, the memory integrated circuits can be based on any other type of memory such as a volatile memory. In some implementations, the memory (e.g., **221**, . . . , **223**, . . . , **225**) can include, but are not limited to, random access memory (RAM), read-only memory (ROM), dynamic random access memory (DRAM), static random access memory (SRAM), synchronous dynamic random access memory (SDRAM), phase change memory (PCM), magneto random access memory (MRAM), negative-or (NOR) flash memory, electrically erasable programmable read-only memory (EEPROM), and/or a cross-point array of non-volatile memory cells. A cross-point array of non-volatile memory can perform bit storage based on a change of bulk resistance, in conjunction with a stackable cross-gridded data access array. Additionally, in contrast to many Flash-based memory, cross point non-volatile memory can perform a write in-place operation, where a non-volatile memory cell can be programmed without the non-volatile memory cell being previously erased. Furthermore, the memory cells of the memory devices can be grouped as memory pages or data blocks that can refer to a unit of the memory device used to store data.

(54) A memory sub-system (e.g., **205** or **209**) can have a controller (e.g., **227** or **229**) that communicate with the memory (e.g., **221**, . . . , **223**, . . . , **225**) to perform operations such as reading data, writing data, or erasing data in the memory (e.g., **221**, . . . , **223**, . . . , **225**) and other such operations, in response to requests, commands or instructions from the processing device **118** and/or the memory management unit (MMU) **213**. The controller (e.g., **227** or **229**) can include hardware such as one or more integrated circuits and/or discrete components, a buffer memory, or a combination thereof. The controller (e.g., **227** or **229**) can be a microcontroller, special purpose logic circuitry (e.g., a field programmable gate array (FPGA), an application specific integrated circuit (ASIC), etc.), or another suitable processor. The controller (e.g., **227** or **229**) can include one or more processors (processing devices) configured to execute instructions stored in local memory.

(55) The local memory of the controller (e.g., **227** or **229**) can include an embedded memory



configured to store instructions for performing various processes, operations, logic flows, and routines that control operation of the memory sub-system (e.g., **205** or **209**), including handling communications between the memory sub-system (e.g., **205** or **209**) and the processing device **118**/MMU **213**, and other functions described in greater detail below. The local memory **119** of the controller (e.g., **227** or **229**) can include read-only memory (ROM) for storing micro-code and/or memory registers storing, e.g., memory pointers, fetched data, etc.

(56) While the example memory sub-systems (e.g., **205** and **209**) in FIG. 2 have been illustrated as including controllers (e.g., **227** and **229**), in another embodiment of the present disclosure, a memory sub-system (e.g., **205** or **209**) may not include a controller (e.g., **227** or **229**), and can instead rely upon external control (e.g., provided by the MMU **213**, or by a processor or controller separate from the memory sub-system (e.g., **205** or **209**)).

(57) In general, the controller (e.g., **227** or **229**) can receive commands, requests or instructions from the processing device **118** or MMU **213** in accordance with a standard communication protocol for the communication channel (e.g., **203** or **207**) and can convert the commands, requests or instructions in compliance with the standard protocol into detailed instructions or appropriate commands within the memory sub-system (e.g., **205** or **209**) to achieve the desired access to the memory (e.g., **221**, . . . , **223**, . . . , **225**). For example, the controller (e.g., **227** or **229**) can be responsible for operations such as wear leveling operations, garbage collection operations, error detection and error-correcting code (ECC) operations, encryption operations, caching operations, and address translations between a logical block address and a physical block address that are associated with the memory (e.g., **221**, . . . , **223**, . . . , **225**). The controller (e.g., **227** or **229**) can further include host interface circuitry to communicate with the processing device **118** via the physical host interface. The host interface circuitry can convert the commands received from the processing device **118** into command instructions to access the memory devices (e.g., **221**, . . . , **223**, . . . , **225**) as well as convert responses associated with the memory devices (e.g., **221**, . . . , **223**, . . . , **225**) into information for the processing device **118**.

(58) The memory sub-system (e.g., **205** or **209**) can also include additional circuitry or components that are not illustrated. In some implementations, the memory sub-system (e.g., **205** or **209**) can include a cache or buffer (e.g., DRAM) and address circuitry (e.g., a row decoder and a column decoder) that can receive an address from the controller (e.g., **227** or **229**) or the MMU **213** and decode the address to access the memory (e.g., **221**, . . . , **223**, . . . , **225**).

(59) In one example, the peripheral interconnect **207**, or the memory bus **203**, has one or more connectors to provide the memory sub-system (e.g., **209** or **205**) with power and/or communicate with the memory sub-system (e.g., **209** or **205**) via a predetermined protocol; and the memory sub-system (e.g., **209** or **205**) has one or more connectors to receive the power, data and commands from the processing device **118**. For example, the connection between the connector on the peripheral interconnect **207** and the connector on a memory sub-system (e.g., **209**) can utilize a PCIe bus or a SATA bus.

(60) In general, the processing device **118** can execute one or more operating systems to provide services, including acceleration of memory access in which a portion of memory in the computer system is accessed via another portion of memory in the computer system using a paging technique and/or a memory map interface, as further discussed below.

(61) The memory virtualizer **113** can be implemented in a device driver of the memory module **205** to virtualize the memory **221**, . . . , **223** in the memory module **205**. Alternatively, the memory virtualizer **113** can be implemented in a device driver of the storage device **209** to virtualize the memory **223**, . . . , **225** in the storage device **209**.

(62) In some instances, a device driver of the memory module **205** virtualizes not only the memory **221**, . . . , **223** in the memory module **205**, but also similar memory modules coupled to the memory bus **203**. Thus, the page swapping controlled by the memory virtualizer **113** can across the memory bus **203** between two memory modules (e.g., **205** and another memory module).

(63) In some instances, a device driver of the storage device **209** virtualizes not only the memory **223**, . . . , **225** in the storage device **209**, but also similar storage devices to the peripheral interconnect **207**. Thus, the page swapping controlled by the memory virtualizer **113** can across the peripheral interconnect **209** between two storage devices (e.g., **209** and another storage device).

(64) In a further alternative, the memory virtualizer **113** can be implemented in a device driver of a combination of the memory module **205** and the storage device **209** to virtualize the memory **221**, . . . , **223**, . . . , **225** in the memory module **205** and the storage device **209**. Thus, the page swapping controlled by the memory virtualizer **113** can across the peripheral interconnect **209** and the memory bus (**203**) between the memory module **205** and the storage device **209**.

(65) In some instances, the memory virtualizer **113** can be implemented in the operating system of the host system **120** and/or a hypervisor of the host system **120**.

(66) FIG. **3** shows a system having a device driver **247** configured to virtualize memory access made via a hypervisor **245**. For example, the system of FIG. **3** can be implemented in a computer system of FIG. **1** or **2**.

(67) The system of FIG. **3** includes a host operating system **241** that can run in the processing device **118** of the computer system of FIG. **1** or **2**. The host operating system **241** includes one or more device drives (e.g., **247**) that provides memory services using the memory (e.g., **221**, . . . , **223**, . . . , **225**) of memory sub-systems, such as the memory module **205** and/or the storage device **209**.

(68) The host operating system **241** includes a hypervisor **245** that provisions a virtual machine **249**. The virtual machine **249** has virtual hardware implemented via the resources and services provided by the host operating system **241** using the hardware of the computing system of FIG. **1** or **2**. For example, the hypervisor **245** can provision virtual memory as part of the virtual machine **249** using a portion of the memory (e.g., **221**, . . . , **223**, . . . , **225**) of memory sub-systems, such as the memory module **205** and/or the storage device **209**.

(69) The virtual machine **249** allows a guest operating system **243** to provide resources and/or services to applications (e.g., **251**, . . . , **253**) running in the guest operating system **243**, in a way as the operating system **243** running on a physical computing machine that has the same or similar set of hardware as provisioning in the virtual machine. The hypervisor **245** manages the mapping between the virtual hardware provisioned in the virtual machine and the services of hardware in the computing system managed by the host operating system **241**.

(70) FIG. **3** illustrates an instance in which a virtual machine **249** is provisioned by the hypervisor **245**. In general, the hypervisor **245** can provision a plurality of virtual machines (e.g., **249**) that can run the same guest operating system **243**, or different guest operating systems (e.g., **243**). Different sets of users and/or application programs can be assigned to use different virtual machines.

(71) In some instances, the host operating system **241** is specialized to provide services for the provisioning of virtual machines and does not run other application programs. Alternatively, the host operating system **241** can provide additional services to support other application programs, such as applications (e.g., **251**, . . . , **253**).

(72) The device driver **247** can be configured to predict the use of a page of slower memory and uses the memory virtualizer **113** to swap the page into faster memory before the page is actually requested for use by the virtual machine **249**. The prediction reduces the time between a request to use the page and the availability of the page in the faster memory by loading and/or transferring the page of data from the slower memory to the faster before the request to use the page, which accelerates the data access of the page.

(73) For example, the slower memory can be the memory **223** in the memory module **205** and the faster memory be the memory **221** in the same memory module **205** (or another memory module connected to the same memory bus **203** as the memory module **205**).

(74) For example, the slower memory can be the memory **223** in the storage device **209**; and the faster memory can be the memory **223** of the same type in the memory module **205**, or the memory

**221** in the memory module **205**.

(75) For example, the slower memory can be the memory **225** in the storage device **209**; and the faster memory can be the memory **223** in the same storage device **209** or another storage device connected to the peripheral interconnect **207**, or memory (e.g., **223** or **221**) in the memory module **205**.

(76) Preferably, the predictive data swapping is performed within a same memory sub-system, such as within the same memory module **205** or the same storage device **209**, to avoid or reduce congestion in communication channels connected to the processing device **118**, such as the memory bus **203** and/or the peripheral interconnect **207**. For example, the data swapping-fetching can be performed to copy data from the slower memory **223** in the memory module **205** to the faster memory **221** in the memory module **205**, under the control of a controller **227** in the memory module **205** in response to one or more command, request, or instruction from the device driver **247** and/or memory management unit (MMU) **213**. For example, the pre-fetching can be performed to copy data from the slower memory **225** in the storage device **209** to the faster memory **223** in the storage device **209**, under the control of a controller **229** in the storage device **209** in response to one or more command, request, or instruction from the device driver **247** and/or memory management unit (MMU) **213**.

(77) In one embodiment, the hypervisor **245** not only requests the device driver **247** to access a memory (e.g., **221**, . . . , **223**, . . . , or **225**) in a memory sub-system (e.g., memory module **205** or storage device **209**) but also provides the device driver **247** with information that can be used in making predictions of which pages in the memory (e.g., **221**, . . . , **223**, . . . , or **225**) are likely to be used in a subsequent time period and which pages in the memory (e.g., **221**, . . . , **223**, . . . , or **225**) are unlikely to be used in the subsequent time period. Examples of information provided by the hypervisor **245** and used by the device driver **247** to make the predictions include: sequences of pages being used in a prior time period, instances of requests to load pages from the slower memory to the faster memory, content attributes of the pages, ownership attributes of the pages, identifications of users or applications of the pages, an indication of whether pages are accessed in a sequential mode in a virtual machine and/or in a user account, an indication of whether page accesses are in a steady state, an indication whether a page used is associated with a huge page, etc.

(78) A page that is likely to be used can be referred to as a hot page; and a page that is unlikely to be used can be referred to as a cold page. The likelihood of a page being used in the subsequent time period can be referred to as the temperature of the page. The device driver **247** uses the information provided by the hypervisor **245** to predict the temperatures of the pages, adjust the locations of the pages in one or more memory sub-systems to optimize the distribution of the pages in the memory (e.g., **221**, . . . , **223**, . . . , or **225**) and accelerate data access.

(79) In some instances, a hardware prediction engine with self-learning capability is used by the device driver **247** in making the predictions of the temperatures of the pages.

(80) FIG. 4 shows a prediction engine **265** coupled to a device driver **247** to improve memory access through virtualization. For example, the prediction engine **265** of FIG. 4 can be used with the host operating system **241** of FIG. 3 and the memory sub-systems (e.g., **205** and **209**) of FIG. 2.

(81) The hypervisor **245** in FIG. 4 is coupled to a translation lookaside buffer (TLB) **267** that defines the mapping between pages of virtual memory provisioned by the hypervisor **245** to the virtual machines (e.g., **249**) and the pages in the memory (e.g., **271** or **273**).

(82) The processing device **118** can access a page of memory **271** at an address **261** faster than a page of memory **273** at an address **263**. It is desirable, and in some instances required, to service the data **260** of a page of the slower memory **273** via a page of the faster memory **271**.

(83) When a virtual machine **249** is accessing the data **260** that is in the slower memory **273** but not in the faster memory **271**, the hypervisor **245** requests the device driver **247** to load the data **260** from the slower memory **273** to the faster memory **271**. The device driver **247** communicates with a controller **228** (e.g., the controller **227** of the memory module **205**, the controller **229** of the

storage device **229**, and/or the memory management unit (MMU) **213**) to load the data **260** to the faster memory **271**. Once the data **260** is in the faster memory **271**, the hypervisor **245** can update the translation lookaside buffer (TLB) **267** to allow the translation of the virtual address of the page to the address **261** of the faster memory **271** such that the data **260** can be read and/or modified in the faster memory **271**.

(84) When the data **260** is not being used, the memory virtualizer **113** can unload the data **260** back to the slower memory **273**, to save the changes made to the data **260** while the data **260** is being used in the faster memory **271**.

(85) The hypervisor **245** provides information to the device driver **247** and/or the self-learning prediction engine **265** to train a prediction model **269** in the prediction engine **265**. Such information can include page access history, the identifications of users, programs and/or virtual machines that use the pages, and attributes of pages, such as locality, time, ownership, usage mode, logical relations among the pages (e.g., identifications of pages that are used together to represent an object), the type of data contained in the pages (e.g., operating system, application program, software tool, system data, user data), and/or other meta data of pages.

(86) For example, the prediction model **269** can be constructed using an artificial neural network, or another machine learning model. A supervised machine learning technique can be used to predict the temperatures of the pages in the slower memory **273** and minimize the predicted temperatures and measured temperatures of the pages as indicated in the actual usages of the pages as observed in the hypervisor **245**. In some instances, the temperature of a page at a time instance is inversely proportional to the closest subsequent time instance the page is to be used. A time gap to the use smaller than a threshold corresponds to a very hot temperature and thus represents an elevated need or priority for fetching the page to the faster memory **271**. A time gap to the use larger than another threshold corresponds to a very cold temperature and thus represents an opportunity to unload the page from the faster memory **271** and make room in the faster memory **271** for other hot pages.

(87) Further, the prediction model **269** can be continuously trained based on the ongoing information about the page usages and the identifications of the actually used pages. The actual usages following the predictions provide feedback regarding the accuracy of the predictions; and the prediction model **269** can be continuously trained to minimize the prediction errors based on the most recent time period of usages.

(88) When the prediction engine **265** identifies a hot page that is not currently in the faster memory **271** and a cold page that is currently in the faster memory **271**, the device driver **247** loads the data **260** of the hot page into the faster memory **271** in anticipation of the use of the data **260** and/or unloads the data **260** of the cold page into the slower memory **273** to make room in the faster memory **271** for hot pages.

(89) Optionally or in combination, the prediction engine **265** uses heuristic rules to aid predictive data positioning. The heuristic rules can be based on whether an application or virtual machine is accessing pages of data sequentially or randomly, whether the application or virtual machine is accessing pages of data in a steady state mode or in a bursty mode, whether the application or virtual machine is accessing pages of data through small pages or huge pages, etc.

(90) In some instances, the prediction engine **265** identifies the most frequently used pages of a subsequent time period; and in response, the memory virtualizer **113** instructs the controller **228** to swap pages that are predicted to be the most frequently used to the faster memory **271**, while allowing the remaining pages, predicted to be less frequently used, to be accessed from the slower memory **273**.

(91) FIG. 4 illustrates the training of the self-learning prediction engine **265** through information provided by a hypervisor **245**. In general, the software tool in the host operating system **241** can be used to provide information about the page usage history and/or page information to train the self-learning prediction engine **265** with or without the hypervisor **245**. For example, the predictive data fetching technique can be used in an operating system that runs a plurality of application programs

but does not have the hypervisor **245** and does not offer virtual machines (e.g., **249**).

(92) Preferably, the prediction engine **265** can be implemented using a hardware circuitry separate from the processing device **118** and/or the controllers (e.g., **213**, **227**, **229**, **228**). For example, an Application-Specific Integrated Circuit (ASIC) or Field-Programmable Gate Array (FPGA) can be used to implement the prediction engine **265**.

(93) Alternatively, the prediction engine **265** can be implemented as a part of the processing device **118**, or a processor coupled to the processing device **118**.

(94) In some instances, the prediction engine **265** is implemented in the controller **228** of a memory sub-system, such as the controller **227** of a memory module **205** or the controller **229** of a storage device **209**.

(95) FIG. 5 shows method to virtualize memory access. The method of FIG. 5 can be performed by processing logic that can include hardware (e.g., processing device, circuitry, dedicated logic, programmable logic, microcode, hardware of a device, integrated circuit, etc.), software (e.g., instructions run or executed on a processing device), or a combination thereof. In some embodiments, the method of FIG. 5 is performed at least in part by the memory virtualizer **113** of FIG. 1 or 2. Although shown in a particular sequence or order, unless otherwise specified, the order of the processes can be modified. Thus, the illustrated embodiments should be understood only as examples, and the illustrated processes can be performed in a different order, and some processes can be performed in parallel. Additionally, one or more processes can be omitted in various embodiments. Thus, not all processes are required in every embodiment. Other process flows are possible.

(96) For example, the method of FIG. 5 can be implemented in a computing system of FIG. 1 or 2 with a host operating system **241** of FIG. 3 and a self-learning prediction engine **265** of FIG. 4. For example, the memory virtualizer **113** can be implemented at least in part via the self-learning prediction engine **265** of FIG. 4 and the host operating system **241** of FIG. 3.

(97) At block **301**, a computing system (e.g., as illustrated in FIG. 1 or 2) provides first memory (e.g., **109A**, **221**, or **271**) and second memory (e.g., **109N**, **223**, or **273**) having different memory access speeds.

(98) For example, the first memory is faster but volatile (e.g., implemented using DRAM); and the second memory is non-volatile but slower (e.g., implemented using flash memory or cross-point non-volatile memory).

(99) In some instances, the first memory and the second memory are separately accessible in the computing system and/or on a bus (e.g., memory bus **203**). However, managing the differences in the first memory and the second memory to full utilizes the different advantages offered by the first memory and the second memory can be a challenge for individual applications (e.g., **251 . . . , 253** in FIG. 3) and/or for individual virtual machines (e.g., **249** in FIG. 3). Thus, a memory virtualizer **113** can be provided in the computing system to virtualize the first memory and the second memory and shield the different characteristics of the first memory and the second memory from the applications (e.g., **251 . . . , 253** in FIG. 3) and/or virtual machines (e.g., **249** in FIG. 3). For example, the memory virtualizer **113** can be implemented as a device driver of a memory module **205** having the first memory (e.g., **221**) and the second memory (e.g., **223**) and thus automatically optimizes the usages of the first memory (e.g., **221**) and the second memory (e.g., **223**) for the current usage patterns of the applications (e.g., **251 . . . , 253** in FIG. 3) and/or virtual machines (e.g., **249** in FIG. 3).

(100) At block **303**, the memory virtualizer **113** stores an address map between first addresses used by the processing device to access memory and second addresses used to access the first memory and the second memory.

(101) For example, the address map can be stored in the cache memory **211** of the host system **120** or in a dedicated/reserved portion of the first memory. The address map can be stored in the translation lookaside buffer (TLB) (**267**). For example, the address map defines the relation

between the virtual memory addresses used by applications/virtual machines to address memory locations in the memory module **205** and the logical memory addresses transmitted on the memory bus **203** to the memory module **205**. Alternatively, the address map defines the relation between the logical memory addresses of the memory module **205** and the physical memory addresses in the memory module **205**.

(102) At block **305**, the memory virtualizer **113** monitors usages of the first memory and the second memory.

(103) For example, the memory virtualizer **113** can be implemented in a device driver **247** of a memory module (e.g., **205**) having the first memory (e.g., **221**) and the second memory (e.g., **223**). Thus, the usages of the first memory and the second memory are made via the memory virtualizer **113**. Past usages and/or current usages can be used to predict future usages in a subsequent period of time. The memory virtualizer **113** can automatically adjust the distribution of data within the first memory and second memory to fully utilize the different advantages offered by the first memory and the second memory.

(104) At block **307**, the memory virtualizer **113** adjusts the address map based on the usages to improve speed of memory access involving the first memory and the second memory.

(105) At block **309**, the memory virtualizer **113** swaps data content in the first memory and the second memory according to adjustments to the address map.

(106) For examples, the first memory (e.g., **109A**, **221**, or **271**) is faster than the second memory (e.g., **109N**, **223** or **273**). Temperatures of pages of data are predicted based on how frequently the pages of data will be used in a subsequent time period. Hotter pages have data that will be used more frequently in the subsequent time period than colder pages. The memory virtualizer **113** instructs the controller **227** of the memory module **205** to swap, adjust, or change storage locations such that hotter pages of data are stored in the faster memory (e.g., **109A**, **221**, or **271**), and colder pages of data are stored in the slower memory (e.g., **109N**, **223** or **273**). The memory virtualizer **113** adjusts the address map such that the changes are shielded from the applications and/or virtual machines that use the first memory and second memory.

(107) In some implementations, the processing device **118** can access the second memory (e.g., **109N**, **223**, or **273**) over a bus (e.g., **203**) without accessing the first memory (e.g., **109A**, **221**, or **271**) at least for a portion of memory access requests processed in the computing system. For example, the frequent access to hotter pages made in the first memory (e.g., **109A**, **221**, or **271**); and the infrequent access to colder pages are made in the second memory (e.g., **109N**, **223**, or **273**). When the performance cost of data swapping is taken into account, swapping colder pages from the slower memory to the faster memory only serve a few infrequent access requests may not improve the overall performance. Thus, the memory virtualizer **113** can avoid such swapping of a colder page to a faster memory and service data access to the colder page directly from the slower memory. In other implementations, the first memory (e.g., **109A**, **221**, or **271**) may be used as cache or gateway for accessing the second memory (e.g., **109N**, **223**, or **273**).

(108) Preferably, the first memory and the second memory are disposed in a same memory module (e.g., **205**) and share a common physical host interface to the bus (e.g., **203**). Thus, the adjustment of data storage locations within the memory module (e.g., **205**) can be performed within the memory module (e.g., **205**) without using the communication bandwidth of the memory bus (e.g., **203**). Preferably, the memory module (e.g., **205**) has a broader communication bandwidth than the memory bus (e.g., **203**) for internal data swapping.

(109) Optionally, the memory module (e.g., **205**) can be implemented as a dual in-line memory module (DIMM) that includes a controller **227**. The controller **227** can implement at least in part the memory virtualizer **113**. For example, in response to an instruction from the host system **120**, the controller performs internal data swapping to reposition data in different types of memories (**221**, . . . , **223**). The memory virtualizer **113** can include a device driver **247** of the memory module (e.g., **205**), where the device driver **247** runs in the processing device **118** of the host

system **120**. Alternatively, the memory virtualizer **113** can be entirely implemented in the controller **227** with hardware and/or firmware.

(110) For example, the memory virtualizer **113** can determinate that data content in the portion of the first memory is accessed less frequently than the portion of the second memory in a subsequent time period. In response to a request to access a memory location currently in the second memory (e.g., **109N**, **223**, or **273**), the memory virtualizer **113** instructs the controller (e.g., **227** or **228**) to swap data content between a portion of the first memory (e.g., **109A**, **221**, or **271**) and a portion of the second memory (e.g., **109N**, **223**, or **273**). Thus, the swapping is triggered by the request to access the memory location in the second memory. The controller (e.g., **227** or **228**) can service the request before initiating the internal operation of swapping the data content between the portion of the first memory and the portion of the second memory such that the host system **120** can obtain a response to the access request prior to the completion of data swapping.

(111) The memory virtualizer **113** is configured to perform address translation to virtualize memory access. For example, after receiving an address of a memory access request to be made by the processing device **118** of the host system **120**, the memory virtualizer **113** converts the address to a converted address based on the address map (e.g., using the translation lookaside buffer **267**); The memory virtualizer **113** makes the memory access request on the memory bus **203** according to the converted address.

(112) In some implementations, the first memory (e.g., **109A**, **221**, or **271**) may be used as cache or gateway for accessing the second memory (e.g., **109N**, **223**, or **273**) in a memory module (**205**) when the second memory is not directly addressable on the memory bus (**203**). Alternatively, the first memory and the second memory can be in separate memory modules that are coupled to a central processing unit via a same memory bus **203**. Alternatively, the first memory and the second memory can be in separate memory systems that are coupled to a central processing unit via separate communication channels.

(113) In some implementations, a communication channel between the processing device **118** and a memory sub-system includes a computer network, such as a local area network, a wireless local area network, a wireless personal area network, a cellular communications network, a broadband high-speed always-connected wireless communication connection (e.g., a current or future generation of mobile network link); and the processing device **118** and the memory sub-system can be configured to communicate with each other using data storage management and usage commands similar to those in NVMe protocol.

(114) A memory sub-system in general can have non-volatile storage media. Examples of non-volatile storage media include memory cells formed in an integrated circuit and magnetic material coated on rigid disks. Non-volatile storage media can maintain the data/information stored therein without consuming power. Memory cells can be implemented using various memory/storage technologies, such as NAND logic gate, NOR logic gate, phase-change memory (PCM), magnetic memory (MRAM), resistive random-access memory, cross point storage and memory devices (e.g., 3D XPoint memory). A cross point memory device uses transistor-less memory elements, each of which has a memory cell and a selector that are stacked together as a column. Memory element columns are connected via two perpendicular lays of wires, where one lay is above the memory element columns and the other lay below the memory element columns. Each memory element can be individually selected at a cross point of one wire on each of the two layers. Cross point memory devices are fast and non-volatile and can be used as a unified memory pool for processing and storage.

(115) The controller (e.g., **228**, **227**, or **229**) of a memory sub-system (e.g., **205** or **209**) can run firmware to perform operations responsive to the communications from the processing device **118**. Firmware in general is a type of computer program that provides control, monitoring and data manipulation of engineered computing devices.

(116) Some embodiments involving the operation of the controller **228** can be implemented using

computer instructions executed by the controller **228**, such as the firmware of the controller **228**. In some instances, hardware circuits can be used to implement at least some of the functions. The firmware can be initially stored in the non-volatile storage media, or another non-volatile device, and loaded into the volatile DRAM and/or the in-processor cache memory for execution by the controller **228**.

(117) A non-transitory computer storage medium can be used to store instructions of the firmware of a memory sub-system (e.g., **209** or **205**) and/or the instructions of the operating system (e.g., **241**, **243**) in general and the device driver **247** and the hypervisor **245** in particular. When the instructions are executed by the controller **228** and/or the processing device **118**, the instructions cause the controller **228** and/or the processing device **118** to perform a method discussed above.

(118) FIG. **6** illustrates an example machine of a computer system **600** within which a set of instructions, for causing the machine to perform any one or more of the methodologies discussed herein, can be executed. In some embodiments, the computer system **600** can correspond to a host system (e.g., the host system **120** of FIG. **1**) that includes, is coupled to, or utilizes a memory sub-system (e.g., the memory sub-system **110** of FIG. **1**) or can be used to perform the operations of a memory virtualizer **113** (e.g., to execute instructions to perform operations corresponding to the memory virtualizer **113** described with reference to FIGS. **1-5**). In alternative embodiments, the machine can be connected (e.g., networked) to other machines in a LAN, an intranet, an extranet, and/or the Internet. The machine can operate in the capacity of a server or a client machine in client-server network environment, as a peer machine in a peer-to-peer (or distributed) network environment, or as a server or a client machine in a cloud computing infrastructure or environment. (119) The machine can be a personal computer (PC), a tablet PC, a set-top box (STB), a Personal Digital Assistant (PDA), a cellular telephone, a web appliance, a server, a network router, a switch or bridge, or any machine capable of executing a set of instructions (sequential or otherwise) that specify actions to be taken by that machine. Further, while a single machine is illustrated, the term “machine” shall also be taken to include any collection of machines that individually or jointly execute a set (or multiple sets) of instructions to perform any one or more of the methodologies discussed herein.

(120) The example computer system **600** includes a processing device **602**, a main memory **604** (e.g., read-only memory (ROM), flash memory, dynamic random access memory (DRAM) such as synchronous DRAM (SDRAM) or Rambus DRAM (RDRAM), static random access memory (SRAM), etc.), and a data storage system **618**, which communicate with each other via a bus **630** (which can include multiple buses).

(121) Processing device **602** represents one or more general-purpose processing devices such as a microprocessor, a central processing unit, or the like. More particularly, the processing device can be a complex instruction set computing (CISC) microprocessor, reduced instruction set computing (RISC) microprocessor, very long instruction word (VLIW) microprocessor, or a processor implementing other instruction sets, or processors implementing a combination of instruction sets. Processing device **602** can also be one or more special-purpose processing devices such as an application specific integrated circuit (ASIC), a field programmable gate array (FPGA), a digital signal processor (DSP), network processor, or the like. The processing device **602** is configured to execute instructions **626** for performing the operations and steps discussed herein. The computer system **600** can further include a network interface device **608** to communicate over the network **620**.

(122) The data storage system **618** can include a machine-readable storage medium **624** (also known as a computer-readable medium) on which is stored one or more sets of instructions **626** or software embodying any one or more of the methodologies or functions described herein. The instructions **626** can also reside, completely or at least partially, within the main memory **604** and/or within the processing device **602** during execution thereof by the computer system **600**, the main memory **604** and the processing device **602** also constituting machine-readable storage media.



The machine-readable storage medium **624**, data storage system **618**, and/or main memory **604** can correspond to the memory sub-system **110** of FIG. **1**.

(123) In one embodiment, the instructions **626** include instructions to implement functionality corresponding to a memory virtualizer **113** (e.g., the memory virtualizer **113** described with reference to FIGS. **1-5**). While the machine-readable storage medium **624** is shown in an example embodiment to be a single medium, the term “machine-readable storage medium” should be taken to include a single medium or multiple media that store the one or more sets of instructions. The term “machine-readable storage medium” shall also be taken to include any medium that is capable of storing or encoding a set of instructions for execution by the machine and that cause the machine to perform any one or more of the methodologies of the present disclosure. The term “machine-readable storage medium” shall accordingly be taken to include, but not be limited to, solid-state memories, optical media, and magnetic media.

(124) Some portions of the preceding detailed descriptions have been presented in terms of algorithms and symbolic representations of operations on data bits within a computer memory. These algorithmic descriptions and representations are the ways used by those skilled in the data processing arts to most effectively convey the substance of their work to others skilled in the art. An algorithm is here, and generally, conceived to be a self-consistent sequence of operations leading to a desired result. The operations are those requiring physical manipulations of physical quantities. Usually, though not necessarily, these quantities take the form of electrical or magnetic signals capable of being stored, combined, compared, and otherwise manipulated. It has proven convenient at times, principally for reasons of common usage, to refer to these signals as bits, values, elements, symbols, characters, terms, numbers, or the like.

(125) It should be borne in mind, however, that all of these and similar terms are to be associated with the appropriate physical quantities and are merely convenient labels applied to these quantities. The present disclosure can refer to the action and processes of a computer system, or similar electronic computing device, that manipulates and transforms data represented as physical (electronic) quantities within the computer system's registers and memories into other data similarly represented as physical quantities within the computer system memories or registers or other such information storage systems.

(126) The present disclosure also relates to an apparatus for performing the operations herein. This apparatus can be specially constructed for the intended purposes, or it can include a general purpose computer selectively activated or reconfigured by a computer program stored in the computer. Such a computer program can be stored in a computer readable storage medium, such as, but not limited to, any type of disk including floppy disks, optical disks, CD-ROMs, and magnetic-optical disks, read-only memories (ROMs), random access memories (RAMs), EPROMs, EEPROMs, magnetic or optical cards, or any type of media suitable for storing electronic instructions, each coupled to a computer system bus.

(127) The algorithms and displays presented herein are not inherently related to any particular computer or other apparatus. Various general purpose systems can be used with programs in accordance with the teachings herein, or it can prove convenient to construct a more specialized apparatus to perform the method. The structure for a variety of these systems will appear as set forth in the description below. In addition, the present disclosure is not described with reference to any particular programming language. It will be appreciated that a variety of programming languages can be used to implement the teachings of the disclosure as described herein.

(128) The present disclosure can be provided as a computer program product, or software, that can include a machine-readable medium having stored thereon instructions, which can be used to program a computer system (or other electronic devices) to perform a process according to the present disclosure. A machine-readable medium includes any mechanism for storing information in a form readable by a machine (e.g., a computer). In some embodiments, a machine-readable (e.g., computer-readable) medium includes a machine (e.g., a computer) readable storage medium such

as a read only memory (“ROM”), random access memory (“RAM”), magnetic disk storage media, optical storage media, flash memory components, etc.

(129) In this description, various functions and operations are described as being performed by or caused by computer instructions to simplify description. However, those skilled in the art will recognize what is meant by such expressions is that the functions result from execution of the computer instructions by one or more controllers or processors, such as a microprocessor.

Alternatively, or in combination, the functions and operations can be implemented using special purpose circuitry, with or without software instructions, such as using Application-Specific Integrated Circuit (ASIC) or Field-Programmable Gate Array (FPGA). Embodiments can be implemented using hardwired circuitry without software instructions, or in combination with software instructions. Thus, the techniques are limited neither to any specific combination of hardware circuitry and software, nor to any particular source for the instructions executed by the data processing system.

(130) In the foregoing specification, embodiments of the disclosure have been described with reference to specific example embodiments thereof. It will be evident that various modifications can be made thereto without departing from the broader spirit and scope of embodiments of the disclosure as set forth in the following claims. The specification and drawings are, accordingly, to be regarded in an illustrative sense rather than a restrictive sense.

## Claims

1. A computing system, comprising: at least one bus; a plurality of memory components, comprising first memory and second memory, wherein the first memory and the second memory have different speeds when accessed via the at least one bus; a processing device, operatively coupled with the plurality of memory components through the at least one bus; wherein the computing system further comprises a memory virtualizer operatively to: store an address map between first addresses used by the processing device to access memory and second addresses used to access the first memory and the second memory; monitor usages of the first memory and the second memory, and train a self-learning prediction engine to generate a prediction associated with a use of the first memory and the second memory; and adjust the address map based on the prediction to improve speed of the processing device in memory access involving the first memory and the second memory; and swap data content in the first memory and the second memory according to adjustments to the address map; wherein the adjusting of the address map based on the prediction are shielded from an application, virtual machine, or a combination thereof, associated with a request to access a memory location currently in the second memory.
2. The computing system of claim 1, wherein the first memory is faster than the second memory; and the processing device accesses the second memory over the at least one bus without accessing the first memory for a portion of memory access requests in the computing system.
3. The computing system of claim 1, wherein the first memory is faster than the second memory; and the first memory and the second memory share a common physical host interface to the bus.
4. The computing system of claim 3, wherein the first memory and the second memory are disposed in a same memory module.
5. The computing system of claim 4, wherein the bus is a memory bus; and the memory module is a dual in-line memory module (DIMM).
6. The computing system of claim 5, wherein the memory module includes a controller that implements at least in part the memory virtualizer.
7. The computing system of claim 6, wherein in response to a request to access a memory location currently in the second memory, the memory virtualizer instructs the controller to swap data content between a portion of the first memory and a portion of the second memory, in response to a determination that data content in the portion of the first memory is accessed less frequently than

the portion of the second memory.

8. The computing system of claim 7, wherein swapping of the data content between the portion of the first memory and the portion of the second memory is performed within the memory module without going through the bus.

9. The computing system of claim 7, wherein the swapping of the data content between the portion of the first memory and the portion of the second memory is performed after accessing the memory location in the second memory.

10. The computing system of claim 5, wherein the memory virtualizer comprises a device driver of the memory module running in the processing device.

11. The computing system of claim 10, wherein the device driver runs in the processing device to: receive an address of a memory access request to be made by the processing device; convert the address to a converted address based on the address map; and make the memory access request on the bus according to the converted address.

12. The computing system of claim 10, wherein the memory virtualizer adjusts the address map to place first data in the first memory and second data in the second memory based on a determination that first data is accessed more frequently in a time period than the second data.

13. The computing system of claim 1, wherein the first memory and the second memory are in a same memory module; and at least a portion of the second memory is not directly addressable on the at least one bus.

14. The computing system of claim 13, wherein the first memory is volatile dynamic random-access memory and the second memory is non-volatile cross-point memory.

15. The computing system of claim 1, wherein the first memory and the second memory are in separate memory modules that are coupled to a central processing unit via a same memory bus.

16. The computing system of claim 1, wherein the first memory and the second memory are in separate memory systems that are coupled to a central processing unit via separate communication channels.

17. A method, comprising: providing a plurality of memory components on at least one bus coupled to a processing device, the plurality of memory components comprising first memory and second memory, wherein the first memory and the second memory have different speeds when accessed via the at least one bus; storing an address map between first addresses used by the processing device to access memory and second addresses used to access the first memory and the second memory; monitoring usages of the first memory and the second memory and training a self-learning prediction engine to predict a use of the first memory and the second memory; and adjusting the address map based on the prediction to improve speed of the processing device in memory access involving the first memory and the second memory; and swapping data content in the first memory and the second memory according to adjustments to the address map; wherein the adjusting of the address map based on the prediction, are shielded from an application, virtual machine, or a combination thereof, associated with a request to access a memory location currently in the second memory.

18. The method of claim 17, further comprising generating, by a hypervisor, information that is used to predict the usages of the first memory and the second memory.

19. The method of claim 18, wherein the information comprises at least one of: a sequence of pages being used in a period of time; instances of requests to load pages from the second memory to the first memory; content attributes of the pages loaded from the second memory to the first memory; ownership attributes of the pages loaded from the second memory to the first memory; identifications of users of the pages loaded from the second memory to the first memory; identifications of applications loaded from the second memory to the first memory; an identification of pages that are accessed in a sequential mode in a virtual machine; an identification of pages that are accessed in a sequential mode in a user account; and an identification of page accesses that are in a steady state.

20. A non-transitory computer storage medium storing instructions which, when executed by a computing system having first memory, and second memory provided on at least one bus coupled to a processing device, cause the computing system to perform a method, the method comprising: storing an address map between first addresses used by the processing device to access memory and second addresses used to access the first memory and the second memory; monitoring usages of the first memory and the second memory and training a self-learning prediction engine to generate a prediction associated with a use of the first memory and the second memory; adjusting the address map based on the prediction to improve speed of the processing device in memory access involving the first memory and the second memory; and swapping data content in the first memory and the second memory according to adjustments to the address map; wherein the adjusting of the address map based on the prediction, are shielded from an application, virtual machine, or a combination thereof, associated with a request to access a memory location currently in the second memory.

---