

# US Patent & Trademark Office

## Patent Public Search | Text View

---

United States Patent	12386744
Kind Code	B2
Date of Patent	August 12, 2025
Inventor(s)	Desai; Keyur B. et al.

---

### Ghost list cache eviction

---

#### Abstract

One example method includes a cache eviction operation. Entries in a cache are maintained in an entry list that includes a recent list, a recent ghost list, a frequent list, and a frequent ghost list. When an eviction operation is initiated or triggered, timestamps of last access for the entries in the entry list are adjusted by corresponding adjustment values. Candidate entries for eviction are identified based on the adjusted timestamps of last access. At least some of the candidates are evicted from the cache.

---

<b>Inventors:</b>	<b>Desai; Keyur B. (Cary, NC), Zhang; Xiaobing (Saratoga, CA)</b>
<b>Applicant:</b>	<b>Dell Products L.P. (Round Rock, TX)</b>
<b>Family ID:</b>	<b>1000008751765</b>
<b>Assignee:</b>	<b>Dell Products L.P. (Round Rock, TX)</b>
<b>Appl. No.:</b>	<b>18/327493</b>
<b>Filed:</b>	<b>June 01, 2023</b>

#### Prior Publication Data

<b>Document Identifier</b>	<b>Publication Date</b>
US 20230305958 A1	Sep. 28, 2023

#### Related U.S. Application Data

continuation parent-doc US 17457000 20211130 US 11669449 child-doc US 18327493

---

#### Publication Classification

**Int. Cl.:** G06F12/08 (20160101); G06F12/0802 (20160101)

**U.S. Cl.:**CPC            **G06F12/0802** (20130101); G06F2212/60 (20130101)**Field of Classification Search****CPC:**            G06F (12/0802); G06F (2212/60)**USPC:**        711/118

---

**References Cited****U.S. PATENT DOCUMENTS**

<b>Patent No.</b>	<b>Issued Date</b>	<b>Patentee Name</b>	<b>U.S. Cl.</b>	<b>CPC</b>
8713261	12/2013	Aharoni et al.	N/A	N/A
8924647	12/2013	Evans	N/A	N/A
9569367	12/2016	Wigmore et al.	N/A	N/A
9582421	12/2016	Agarwala et al.	N/A	N/A
9921963	12/2017	Li et al.	N/A	N/A
9990301	12/2017	Kanaujia et al.	N/A	N/A
11599462	12/2022	Hawkins	N/A	N/A
2006/0143393	12/2005	Petev	711/E12.071	G06F 12/122
2008/0059707	12/2007	Makineni et al.	N/A	N/A
2012/0137081	12/2011	Shea	N/A	N/A
2012/0221774	12/2011	Atkisson et al.	N/A	N/A
2012/0254543	12/2011	Sintorn et al.	N/A	N/A
2013/0227218	12/2012	Chang et al.	N/A	N/A
2014/0115244	12/2013	Maybee et al.	N/A	N/A
2016/0062916	12/2015	Das et al.	N/A	N/A
2017/0031823	12/2016	Ross et al.	N/A	N/A
2017/0235681	12/2016	Kaburaki et al.	N/A	N/A
2017/0329720	12/2016	Bedi et al.	N/A	N/A
2018/0275899	12/2017	Munsil et al.	N/A	N/A
2020/0213627	12/2019	Ra	N/A	H04N 21/23113
2020/0225868	12/2019	Dalmatov et al.	N/A	N/A
2021/0263862	12/2020	Gupta et al.	N/A	N/A

**OTHER PUBLICATIONS**

John L. Hennessy and David Patterson, Computer Architecture: A Quantitative Approach, 2019, Morgan Kaufmann, Sixth Edition, Appendix pp. B13 and B39 (Year: 2019). cited by applicant

---

*Primary Examiner:* Yu; Jae U

*Attorney, Agent or Firm:* Workman Nydegger

---

**Background/Summary**

## FIELD OF THE INVENTION

(1) Embodiments of the present invention generally relate to caching in computing systems. More particularly, at least some embodiments of the invention relate to systems, hardware, software, computer-readable media, and methods for evicting data from caches in computing systems.

## BACKGROUND

(2) In computing systems, data is often stored on a storage device such as a hard drive. When needed, the hard drive is accessed to retrieve and use the data. In order to improve this process, caches are used. A cache is typically implemented in a memory that is faster than the hard drive. Data from the hard drive can be cached in the cache and accessed more quickly. As a result, caches tend to improve the performance of the computing system at least by avoiding the need to access the hard drive when the data is cached.

(3) The size of the cache, however, is limited or fixed in size and is typically smaller than the hard drive. As a result, it is necessary to remove or evict data from the cache when the cache is full or for other reasons. Various cache eviction techniques exist, including least recently used (LRU) and least frequently used (LFU). Another method for evicting data from a cache is a time threshold-based eviction. Entries in the cache whose latest access was earlier than a threshold time are candidates for eviction.

(4) The time threshold-based eviction method has some problems. For example, some entries in the cache, even though frequently used and frequently accessed, become candidates for eviction when they satisfy the eviction criteria. Further, in a cache that limits the number of evictions, some entries may stay in the cache at the expense of other more frequently used entries because the eviction thread does not reach these entries due to the limit. Improvements in caching methodologies are needed.

---

## Description

### BRIEF DESCRIPTION OF THE DRAWINGS

(1) In order to describe the manner in which at least some of the advantages and features of the invention may be obtained, a more particular description of embodiments of the invention will be rendered by reference to specific embodiments thereof which are illustrated in the appended drawings. Understanding that these drawings depict only typical embodiments of the invention and are not therefore to be considered to be limiting of its scope, embodiments of the invention will be described and explained with additional specificity and detail through the use of the accompanying drawings, in which:

(2) FIG. 1A discloses aspects of a cache where cache evictions are determined by adjusting characteristics of the cache entries;

(3) FIG. 1B discloses additional aspects of a cache configured to evict entries based on adjusted entry characteristics;

(4) FIG. 2 discloses aspects of an eviction engine configured to identify and evict entries from a cache;

(5) FIG. 3A discloses aspects of managing a cache directory or a cache entry list;

(6) FIG. 3B discloses aspects of performing an eviction operation;

(7) FIG. 3C discloses aspects of performing an eviction operation based on adjusted cache entries;

(8) FIG. 4 discloses aspects of performing operations such as an eviction operation;

(9) FIG. 5 discloses aspects of using adjustment factors in ghost lists;

(10) FIG. 6 discloses aspects of an eviction operation using adjustment factors in ghost lists; and

(11) FIG. 7 discloses aspects of a computing device or a computing system.

### DETAILED DESCRIPTION OF SOME EXAMPLE EMBODIMENTS

(12) Embodiments of the present invention generally relate to caching systems and caching

operations. More particularly, at least some embodiments of the invention relate to systems, hardware, software, computer-readable media, and methods for caching operations including cache entry operations and cache eviction operations.

(13) In general, example embodiments of the invention relate to cache eviction operations.

Effective eviction strategies can improve the performance of the cache, applications that use the cache, and the operation of the computing system in general.

(14) Embodiments of the invention relate to caching operations such that items that are most recently used and/or most frequently used are able to obtain or are given a longer time to live (TTL) in the cache. By giving at least some of the entries in the cache a longer TTL, some of the entries do not become candidates for eviction using time threshold-based eviction operations. This helps improve cache operations by evicting entries that are less likely to be accessed compared to other entries in the cache.

(15) This is generally achieved by adjusting the timestamps of at least some of the cached entries. By adjusting the timestamps and creating adjusted timestamps, at least some of these entries are no longer candidates for certain cache eviction operations.

(16) Embodiments of the invention are described in the context of various caching methodologies including methodologies based on frequency such as least frequently used (LFU) and methodologies based on recency such as least recently used (LRU). By way of example and without limitation, recency refers to when an entry was accessed relative to other entries, and frequency refers to how many times an entry has been accessed.

(17) Embodiments of the invention are also disclosed in the context of adaptive replacement cache (ARC) and time threshold-based methodologies. ARC is a method of cache management that uses aspects of both recency and frequency. A time threshold-based methodology identifies candidates for eviction based on a time threshold.

(18) Embodiments of the invention improve caching methodologies by introducing an adjustment mechanism to improve caching performance. The following discussion described aspects of these caching methodologies.

(19) FIG. 1A illustrates an example of a cache that is managed using a least recently used (LRU) methodology. The cache **102** includes an entry list **104**. The cache **102** operates by maintaining an ordered list of cache entries (e.g., a cache directory) in the entry list **104**. The entries in the list may include {key,value} pairs. The entry list **104** may identify the data in different manners.

(20) The entry list **104** is sorted based on time of the most recent access. Thus, new entries to the cache **102** are added at the top of the list **104**, after the bottom entry has been evicted from the list **104**. If an existing entry in the cache **102** is hit (or accessed), the corresponding entry in the list **104** is moved to the top of the list **104**. This ensures that the least recently used entries move towards the bottom of the list and are eventually evicted. Periodically or for other reasons, an eviction operation may be performed. The eviction operation may use an adjustment factor **106** to identify candidates.

(21) FIG. 1B illustrates an example of a cache that is managed using adaptive replacement cache (ARC) in accordance with embodiments of the invention. The cache **110** similarly includes an entry list **112**. In this example, the entry list **112** is divided into two lists: a recent list **114** and a frequent list **116**. The sizes of the lists **116** and **118** may or may not be fixed in size. The respective sizes of the lists **116** and **118** can change over time in some embodiments.

(22) When data is added to the entries in the cache **110** for the first time, an entry is made for that data in the recent list **114**. This entry then moves down the recent list **114** as newer entries are entered into the recent list **116**. If that entry is accessed a second time, that entry is moved to the frequent list **116**. In effect, the recent list **114** is an LRU list. The frequent list **116** allows the cache **110** to account for entries that are frequently accessed. The frequent list **116** also has aspects of an LRU list. The entry list **112** can account for both recency and frequency. During an eviction operation, an adjustment factor **118** may be used to identify candidates for eviction as discussed

herein.

(23) FIG. 2 illustrates a more detailed example of a cache directory or entry list in accordance with embodiments of the invention. FIG. 2 illustrates a cache **200** that is configured to store data or cache data that may be used, for example, by an application. The cache **200** typically has a fixed size, represented by the box **210**. The contents or data stored in the cache **200** is referenced in a cache directory or entry list **220**. The entry list **220** includes a list **216** and a list **218**. The list **216** is typically configured to store or list entries with respect to recency. The list **218** is configured to store or list entries in the cache **200** with respect to frequency and/or recency.

(24) The list **216** includes a recent list **204** and an associated ghost list **202** (also referred to herein as a recent ghost list). The recent list **204** includes entries whose data is present in the cache **200**. The ghost list **202** includes metadata associated with data that was recently in the cache **200**. The list **218** similarly includes a frequent list **206** and an associated ghost list **208** (also referred to as a frequent ghost list).

(25) The top of the list **216** and the top of the list **218** are associated with the marker **214** (illustrated as an “!”). The marker **212** (illustrated as an “{circumflex over ( )}”), represents a target size for the list **216** or, more specifically, for the recent list **204**. In one example, the target size of the list **216** (or more specifically the recent list **216**) is larger than the target size of the list **218** (or more specifically the frequent list **206**). In effect, the marker **214** conveys the actual sizes of the lists **216** and **218** in the cache **200**. As illustrated, the actual sizes of the lists **216** and **218** (or the lists **204** and **206**) may differ from the target sizes.

(26) As entries are added to and/or removed from the lists **216** and **218**, their relative sizes may change. In other words, the marker **214** may move or change during operation of the cache. The marker **214** can reference a location or position relative to the entry list **220**. Entries in the list **216** are illustrated to the left of the marker **214**. When an entry is added to the list **216**, the entry is added at the top (e.g., next to the marker **214**) and the remaining entries in the list **216** shift to the left (or away from the marker) as needed. Similarly, entries added to the list **218** are added to the right of the marker **214** and the remaining entries in the list **218** shift to the right as needed. Also, it is understood that the implementation of the entry list **220** may have no concept of direction but may be implemented using an appropriate data structure. Generally, the structure or list is configured such that the entries are arranged in a time-related manner. However, this structure or arrangement is not limiting and other structures or list arrangements are within the scope of embodiments of the invention. The description of FIG. 2 aids in understanding the relationship between the cache size, entries in the cache, and the manner in which the entries are identified in the directory or entry list **220**.

(27) As previously stated, the list **216** includes two portions in this example: the recent list **204** and the associated ghost list **202**. The ghost list **202** is an extension of the recent list **204**. The recent list **204** operates much like an LRU list. The recent list **204** is for recent entries into the cache **200** and the ghost list **202** is for entries that have been recently evicted from the recent list **204** but are still tracked. The ghost list **202** keeps a history of recently evicted entries from the recent list **204**. The ghost list **202** only contains metadata (e.g., keys for cache entries) and not the resource or data itself. When an entry moves into the ghost list **202** from the recent list **204**, the data or actual resource is discarded. The ghost list **202** and the recent list **204** represent a history of resources that have been accessed a single time.

(28) The list **218** includes two portions in this example: the frequent list **206** and the associated ghost list **208**. The ghost list **208** is an extension of the frequent list **206**. Entries evicted from the frequent list **206** are inserted or added to the ghost list **208**. The ghost list **208** stores recently evicted entries from the frequent list **206**. The ghost list **208** only contains metadata (e.g., keys for cache entries) and not the resource or itself. When an entry moves from the frequent list **206** to the ghost list **208**, the metadata is retained, and the data or resource is discarded. The list **218** is configured to store entries that have been accessed multiple times.

(29) In sum, new entries are added to the recent list **204** at the top and are gradually pushed to the left (in this visual example). Eventually, these entries are pushed out of the recent list **204** into the ghost list **202** and then dropped out of the ghost list **202**. Similarly, entries to the list **218** are added at the top of the frequent list **206**. These entries move to the right and are evicted to the ghost list **208** and then evicted from the from the ghost list **208**.

(30) In some embodiments, entries in the list **216** have only been accessed once. When entries in the list **216** are accessed a second time, these entries are moved to the list **218**. Entries reentering the cache **200** cause the marker **214** to move towards the marker **212**. If no free space exists in the cache **200**, the relative position of the marker **214** to the marker **212** determines whether an entry is evicted from the list **218** or the list **216** (or from the recent list **204** or the frequent list **206**).

(31) For example, a hit in the ghost list **202** will increase the size of the recent list **204** and push the marker **212** to the right. In this example, the last entry in the ghost list **208** is evicted.

(32) A hit in the ghost list **208** shrink the recent list **204** and push the marker **212** back to the left. The last entry in the list **216** is evicted. More specifically, the last entry of the recent list **204** is evicted to the ghost list **202** and the last entry of the ghost list **202** is evicted and discarded in this example.

(33) A cache miss does not affect the marker **212**, but the marker **214** moves closer to the marker **212**.

(34) In one example each of the lists **216** and **218** (or the portions thereof) is substantially operated as a least recently used list. As previously stated, however, the list **218** also accounts for frequency of access. The arrow **222** illustrates that the cache **200**, which may be fixed in size, can move freely across the history represented by the lists **216** and **218**.

(35) At times, the eviction engine **224** may perform an eviction operation. The eviction engine **224** applies an adjustment factor **226** to the entry list **220** and generates an output **228**, which includes candidates for eviction. At least some of the candidates may be evicted.

(36) FIG. 3A illustrates operation of a cache in accordance with embodiments of the invention. FIG. 3A illustrates a cache **302** having a fixed cache size (c). The recent list is represented by T1, and the associated ghost list is represented as B1. The frequent list and associated ghost list are represented, respectively, by T2 and B2. The commands **304** are an example of how entries are handled in the cache **302** and in the entry list. The commands **306** illustrate another example of how entries are handled in the cache **302** or in the entry list. The ghost lists (B1 and B2) are ignored in the example shown in FIG. 3A.

(37) In this example, the cache **302** can handle 4 entries. The first command of get **02** in the commands **204** results in **02** being placed in the T1 list at the top. The next get **02** command causes the **02** entry to move to the top of the list T2. This occurs because the entry **02** is accessed recently and more than once. A second access moves the entry **02** from the T1 list to the T2 list. The entry **02** is not accessed again until the last command in the commands **304**. This command illustrates that the entry **02** moves from the bottom of the T2 list to the top of the T2 list. The T1 list can account for recency. The T2 list can account for recency and frequency. A review of the entries in the commands **304** illustrates that the sizes of the T1 and T2 lists change as entries are added or moved from one list to another list. If a new command get **05** were added to the commands **304**, the entry **05** would be inserted at the top of the T1 list and the **03** entry would be evicted from the T2 list in one example.

(38) The commands **306** illustrate another example of commands that impact the contents or entries in the cache. Generally, the first time data is accessed, an entry is added to the T1 list. Entries in the T1 list gradually move to the bottom and are eventually evicted unless hit a second time. When hit the second time, the entries are moved to the top of the T2 list. Entries in the T2 list gradually move to the bottom and are evicted eventually unless hit again. Each hit to an entry in the T2 list moves the entry back to the top of the T2 list.

(39) FIG. 3B discloses aspects of a time threshold-based cache management. In this example, a

table **310** illustrates a cache that stores four entries: **00**, **01**, **02** and **03**. The table **310** illustrates, for each entry, a timestamp of last access, a frequency of use, and a recency of use. Entries whose frequency of use is greater than 1 are in the T2 or frequent list while entries whose frequent of use is equal to 1 are in the T1 or recent list.

(40) When evicting entries based on time (e.g., earlier than a certain time), it is possible for entries that are more recently/frequently used to be evicted due to their access timestamp compared to the eviction criteria. If the eviction process is limited to evicting N entries, items that are less recently/frequently used may stay in the cache because the eviction process never processed those entries (i.e., N entries have already been evicted prior to reaching those entries).

(41) The table **310** in FIG. **3B** illustrates the status of cache entries at a point in time when an eviction engine **312** executes an eviction process to evict stale entries from the cache. The eviction process is configured to evict two entries using a threshold time of 10:03:01 in this example. Using this threshold, there are three entries that are candidates for eviction (**00**, **01**, **03**) because each has a timestamp of last access that is less than (or earlier than) the threshold time. However, the eviction process is configured to evict a maximum of two entries. If the cache entries are processed linearly and the number of evictions is limited to 2, the entries **00** and **01** will be evicted from the cache.

(42) As illustrated in the table **314**, the result is not optimal. The table **314** illustrates the cache after running the eviction process. As illustrated, only the entries **02** and **03** remain in the cache. This may not be an optimal result.

(43) For example, an evaluation of the recency and frequency of the entries remaining in the cache demonstrates that the timestamp of the entry **03** was earlier than or occurred before the timestamps of the entries **00** and **01**. The remaining entries **02** and **03** were less frequently used than the evicted entries **00** and **01**. The remaining entry **03** has a less recency of use compared to evicted entries **00** and **01**. This demonstrates that the eviction thread evicted entries from the cache when there were better candidates for eviction. More specifically, this suggests that the entry **03** should have been evicted before the entries **00** and **01**.

(44) FIG. **3C** discloses aspects of an eviction operation using an adjustment factor. The adjustment factor is configured to improve the efficiency of the cache eviction operations and to ensure that more optimal candidates are identified and evicted by the eviction engine. The adjustment factor is used to adjust at least one characteristic or metadata of the cache entries.

(45) FIG. **3C** illustrates a cache **320** that includes four entries: **00**, **01**, **02**, and **03**. Entries **02** and **03** are in the T1 list while the entries **00** and **01** are in the T2 list. When the eviction engine executes on the table **322**, which is representative of the cache entries, the entries are provided with an adjustment factor. The adjustment factor is applied to the timestamp of last access to generate a timestamp of last access with adjustment factor. Thus, the adjustment factor is used to generate adjusted timestamps.

(46) The adjustment factor can be implemented in many different ways and configurations. By way of example only and not limitation, the adjustment factors for entries in the T1 list may be determined as follows. The top entry (which is **02** in T1 list of the cache **320**) is given an adjustment factor of 30 seconds. The next entry in the T1 list (**03**) in this example, is given an adjustment factor of 25 seconds. In this example, the adjustment factor decreases for entries further from the top of the list or further from the marker **326**. As a general rule by way of example only, entries in the T1 or recent list nearer to the top are adjusted using an adjustment factor that is larger than entries further from the top of the list. The manner in which the adjustment factor decreases can vary. In this example, the adjustment factor decreases by 5 seconds for each entry.

(47) The adjustment factor for entries in the T2 or frequent list may be the same or different. In this example, the adjustment factor for the entries in the T2 list are double the adjustment factor for the entries in the T1 list. Thus, the entry **01** (at the top of the T2 list) has an adjustment factor of 60 seconds and the entry **02** has an adjustment factor of 50 seconds. The adjustment factors applied to entries in the T2 list may decrease by a factor, a constant, or the like as the entry index increases

(the index of the entry at the top of the list is 0 and the index for the nth entry is n).

(48) When the eviction engine **328** processes the entries in the cache **320** represented in the table **322** whose characteristics or metadata have been adjusted with an adjustment factor, the eviction engine **328** identifies a single candidate for eviction—entry **03**. This is true, in this example, even if the eviction engine **328** is configured to identify a maximum of two entries for eviction. Thus, as illustrated in the table **324**, the entry **03** is evicted from the cache **320**.

(49) More specifically in this example, assuming that the threshold time was 10:03:01, all of the entries except the entry **03** had an adjusted timestamp that was greater than or later than the threshold time. Further, the entry **03** has the lowest frequency of use. Further, the entry **03** has a lower recency than the other entries in the cache **320** and the frequency of use for the entry **03** was equal to that of the entry **02**. Thus, the eviction engine **328**, compared to FIG. **3B**, evicted an entry that was a better candidate for eviction than the entries **00** and **01**, which were evicted by the eviction engine **312** that did not use an adjustment factor. The frequency of use is not affected by the adjustment factor in one embodiment.

(50) Embodiments of the invention, thus provide novel cache eviction strategies and are capable of evicting multiple entries at the same time. Entries that are frequently or recently used are not evicted because the adjustment factor prevents these entries from being identified as candidates for eviction based on the time threshold.

(51) FIG. **4** discloses aspects of a method for performing an eviction operation. In one example, the cache may implement an eviction strategy based on one or more aspects or characteristics of the entries in the cache including, but not limited to, recency, frequency, and time (i.e., a timestamp). In one example, the eviction operation may adjust at least some of the characteristics or metadata of the entries prior to identifying candidates for eviction. Some of the elements may be performed on an ongoing basis, periodically, or less frequently than other elements of the method **400**.

(52) In this example method **400**, the cache operates **402** normally and includes an entry list or cache directory. The cache entry list may include a recent list and a frequent list. Normal operation may include adding new entries to the cache or accessing existing entries in the cache. When these actions occur, the entry list is adjusted as previously described.

(53) Next, a cache eviction operation may be initiated **404** or triggered. In one example, the cache eviction operation includes aspects of a time threshold-based eviction operation. Thus, a threshold time may be established or determined. The time threshold may be determined by default, by user input, or the like.

(54) Once the eviction operation is initiated, adjustments for the cache entries may be generated **406**. The adjustments to the cache entries can all be determined by the eviction process.

Alternatively, adjustments may be determined or recalculated as entries are newly added or rearranged. Thus, the adjustment values may already be generated when an eviction operation is generated. The adjustments may need to be refreshed in some examples. Generating **406** the adjustments can include generating an adjusted timestamp (i.e., adjusted timestamp since last access) by adjusting or altering the timestamp since last access by the adjustment factor.

(55) Once the adjustments or adjusted timestamps are generated, candidates for eviction are identified **408** based on the adjusted entries. Because the eviction operation identifies candidates based on a time threshold, the candidates are identified by comparing the time threshold to the adjusted timestamps of the cache entries.

(56) In one example, the eviction operation may identify all candidates for eviction. In this case, the candidates are identified **408** and the earliest entries are evicted **410**. This number may be limited by a maximum. In another example, the eviction operation identifies **408** candidates until the maximum number of candidates have been found. In either case or in other embodiments, the identified entries (up to the maximum) are then evicted **410**. In one example, it is not necessary to evict the maximum number of entries from the cache.

(57) In these examples, embodiments of the invention allow the time threshold-based cache



eviction operation to account for recency and/or frequency of the cache entries rather than simply relying on identified candidates based on the time threshold alone.

(58) The adjustment factor can further be incorporated into the ghost lists. This allows for a situation where an entry in the recent list or the frequent list is a better candidate for eviction than an entry in one of the ghost lists.

(59) FIG. 5 discloses aspects of evicting cache entries by applying adjustment factors to entries in the cache entry list, including the recent list (T1), the recent ghost list (B1), the frequent list (T2), and the frequent ghost list (B2). FIG. 5 illustrates a cache 502 that includes four entries: 00, 01, 02 and 03. The cache 502 includes entries 03 and 02 in the T1 or recent list and entries 01 and 00 in the T2 or frequent list. The recent ghost list B1 and the recent ghost list B2 are currently empty. This is the same cache as the cache 320 illustrated in FIG. 3C.

(60) At time 10:04:01, entry 02 is hit or accessed 504. Because this is the second hit or access to the entry 02, the entry 02 is moved to the top of the frequent list and the entry 00 is moved to the frequent ghost list B2 as illustrated in the cache 506. In this example, the data associated with the entry 00 may not be removed from the cache. Although data is often discarded when an entry moves to the ghost list, there are situations in embodiments of the invention where the entry is on the ghost list and the data is not discarded. If, for example, a new entry (e.g., entry 04) were added to the cache, it may be necessary to discard data associated with the entry 00 in the ghost list if the cache is full.

(61) In one example, when the entry 00 moves to the ghost list B2, the adjustment factor may be further reduced or may reduce at a faster rate than the adjustment factor of the frequent list T2. The adjustment factors (or values) for the recent and frequent lists were previously described. The adjustment factors for the ghost lists B1 and B2 are now described. The adjustment factors are described by way of example and not limitation.

(62) For example, as an entry in the recent or T1 list moves to the bottom of the recent list, the adjustment factor may change by a constant. The recent list adjustment factor calculation from  $i=0$  to  $i=n$ , where the top entry in the T1 list is an index  $i=0$  is as follows:  
$$\text{adjustment factor.sub.i+1} = \text{adjustment factor.sub.i} - \text{constant.}$$

(63) Similarly, the adjustment factor for entries in the frequent list may be expressed (in terms of the constant for the recent list by way of example only) as follows:  
$$\text{adjustment factor.sub.i+1} = \text{adjustment factor.sub.i} - 2 * \text{constant.}$$

(64) These examples illustrate that the adjustment factors for the entries in the T1 and T2 lists can be related, but this is not required. Further, the constant can be replaced with a variable constant such that each step reduces the adjustment factor or value by a greater amount than the previous reduction.

(65) The adjustment factors for the ghost lists may be different from the adjustment factors of the recent and frequent lists. For example, the constant for the T1 list may be 5 seconds. If the adjustment factor for the first (index 0) entry is 25 seconds, the adjustment factor for the next entry in the T1 list is 20 seconds. If an entry falls to the ghost list, the adjustment factor may reduce at a faster pace (e.g., double). For example, if the adjustment factor of the entry 03 in the cache 506 is 25 seconds and the entry 03 is moved to the ghost list or the B1 list, the adjustment factor would decrease by 10 seconds (instead of 5 seconds) to 15 seconds.

(66) Thus, as an entry moves to the left in the T1 list, the adjustment factor decreases by 5 seconds for each move. The adjustment factor increases to 10 seconds once the entry is on the ghost list. The adjustment factor for the ghost list B2 may also be 10 seconds, but other factors could be used.

(67) The tables 508 and 510 illustrate an example of using an adjustment factor for entries in the ghost lists.

(68) When the adjustment factor is applied, as illustrated in the table 508, the adjustment factor for the entry 00 (which is in the ghost list of the cache 506) is 40 seconds. The table 510 considers an example where no adjustment factor is applied to entries in the ghost lists.

(69) If the cache **506** is full and an eviction operation is performed, providing an adjustment factor for ghost list entries improves the performance of the cache **506**. For example, if a time threshold-based operation is performed using a time threshold of 10:02:59 on the cache **506**, the adjustment factor is applied and results in the adjusted timestamps (timestamp of last access with adjustment factor) shown in the table **508**. This results in the eviction of the entry **03** because only the adjusted timestamp of the element **503** is earlier than the time threshold (**10:01:25** is earlier than 10:02:59).

(70) This is a good result because the entry **03** has the worst recency value (4) and was only accessed once (frequency of access). In contrast, the entry **00**, which is in the ghost list, has a frequency use of 2 and a recency value of (3). The adjustment factor thus improves the performance by identifying a better candidate for eviction than the entry in the ghost list.

(71) The table **510** illustrates that, without the adjustment factor and using the same time threshold, entries **00** and **03** are evicted from the cache **506**. This may not be as optimal as the result illustrated in the table **508**.

(72) Using the adjustment factor for the recent and frequent ghost lists in addition to the recent and frequent lists provides a gradual reduction of the adjustment factor that is applied to the entries in the entry list, which includes the recent list, the recent ghost list, the frequent list, and the frequent ghost list, and allows entries to have a longer time to live in the cache.

(73) FIG. **6** discloses aspects of an eviction operation. The method is similar to the method **400** and further accounts for entries in the ghost lists. In the method **600**, the cache is operating **602**. Next, a cache eviction operation is initiated **604**. Once initiated, the relevant adjustment factors are applied to the entries in the entry list to generate **606** adjustments such as adjusted timestamps. In this example, the entry list includes a recent list and its associated ghost list and a frequent list and its associated ghost list.

(74) Candidates for eviction are identified **608** based on the adjusted timestamps compared to a time threshold value. Identified candidates are then evicted **610**. As previously stated, the actual number of entries evicted may be limited. However, there is no requirement to evict more than one entry, although embodiments of the invention advantageously allow multiple entries to be identified and evicted at the same time.

(75) The following is a discussion of aspects of example operating environments for various embodiments of the invention. This discussion is not intended to limit the scope of the invention, or the applicability of the embodiments, in any way.

(76) In general, embodiments of the invention may be implemented in connection with systems, software, and components, that individually and/or collectively implement, and/or cause the implementation of, cache operations including cache eviction operations, least frequently used operations, frequency and/or recency-based operations, time threshold-based operations, or combination thereof. More generally, the scope of the invention embraces any operating environment in which the disclosed concepts may be useful.

(77) At least some embodiments of the invention provide for the implementation of the disclosed functionality in existing backup platforms, examples of which include the Dell-EMC NetWorker and Avamar platforms and associated backup software, and storage environments such as the Dell-EMC DataDomain storage environment. In general, however, the scope of the invention is not limited to any data backup platforms or data storage environments.

(78) Example cloud computing environments, which may or may not be public, include storage environments that may provide data protection functionality for one or more clients. Another example of a cloud computing environment is one in which processing, data protection, and other, services may be performed on behalf of one or more clients. Some example cloud computing environments in connection with which embodiments of the invention may be employed include, but are not limited to, Microsoft Azure, Amazon AWS, Dell EMC Cloud Storage Services, and Google Cloud. More generally however, the scope of the invention is not limited to employment of any particular type or implementation of cloud computing environment. Cache operations may be

implemented therein.

(79) In addition to the cloud environment, the operating environment may also include one or more clients that are capable of collecting, modifying, and creating, data. These clients may use caches during these operations. As such, a particular client may employ, or otherwise be associated with, one or more instances of each of one or more applications that perform such operations with respect to data including cache related operations. Such clients may comprise physical machines, or virtual machines (VM) or containers.

(80) Particularly, devices in the operating environment may take the form of software, physical machines, or VMs, or containers, or any combination of these, though no particular device implementation or configuration is required for any embodiment. Similarly, data protection system components such as databases, storage servers,

(81) As used herein, the term 'data' is intended to be broad in scope. Thus, that term embraces, by way of example and not limitation, data segments such as may be produced by data stream segmentation processes, data chunks, data blocks, atomic data, emails, objects of any type, files of any type including media files, word processing files, spreadsheet files, and database files, as well as contacts, directories, sub-directories, volumes, and any group of one or more of the foregoing.

(82) Example embodiments of the invention are applicable to any system capable of storing and handling various types of objects, in analog, digital, or other form. Although terms such as document, file, segment, block, or object may be used by way of example, the principles of the disclosure are not limited to any particular form of representing and storing data or other information. Rather, such principles are equally applicable to any object capable of representing information.

(83) It is noted that any of the disclosed processes, operations, methods, and/or any portion of any of these, may be performed in response to, as a result of, and/or, based upon, the performance of any preceding process(es), methods, and/or, operations. Correspondingly, performance of one or more processes, for example, may be a predicate or trigger to subsequent performance of one or more additional processes, operations, and/or methods. Thus, for example, the various processes that may make up a method may be linked together or otherwise associated with each other by way of relations such as the examples just noted. Finally, and while it is not required, the individual processes that make up the various example methods disclosed herein are, in some embodiments, performed in the specific sequence recited in those examples. In other embodiments, the individual processes that make up a disclosed method may be performed in a sequence other than the specific sequence recited.

(84) Following are some further example embodiments of the invention. These are presented only by way of example and are not intended to limit the scope of the invention in any way.

(85) Embodiment 1. A method, comprising: initiating a cache eviction operation for a cache, wherein entries associated with the cache are identified in an entry list that includes a recent list, a recent ghost list, a frequent list, and a frequent ghost list, generating an adjusted characteristic for each entry in the entry list, identifying candidate entries for eviction from the cache based at least on the adjusted characteristics of the entries compared to a time threshold, and evicting at least one of the identified candidate entries from the cache based on the comparisons.

(86) Embodiment 2. The method of embodiment 1, wherein the characteristic is a timestamp of last access, further comprising adjusting the timestamp of last access of each entry by an adjustment value to generate an adjusted timestamp for each entry, wherein the adjusted characteristics comprise the adjusted timestamps.

(87) Embodiment 3. The method of embodiment 1 and/or 2, wherein the adjustment value for different entries is different.

(88) Embodiment 4. The method of embodiment 1, 2, and/or 3, wherein the adjustment values for entries in the recent list are different from adjustment values in the recent ghost list and wherein adjustment values for entries in the frequent list are different from adjustment values in the

frequent ghost list.

(89) Embodiment 5. The method of embodiment 1, 2, 3, and/or 4, wherein the adjustment values for entries in the frequent list are larger than the adjustment values for entries in the recent list.

(90) Embodiment 6. The method of embodiment 1, 2, 3, 4, and/or 5, wherein the adjustment values for entries near a top of the recent list are greater than the adjustment values for entries further from the top of the recent list and wherein the adjustment values for entries near a top of the frequent list are greater than the adjustment values for entries further from the top of the frequent list.

(91) Embodiment 7. The method of embodiment 1, 2, 3, 4, 5, and/or 6, wherein the adjustments values for entries in the recent list decrease by a constant value moving from a top of the recent list to a bottom of the recent list, wherein the adjustment values for entries in the recent ghost list decrease by a second constant value, wherein the adjustments values for entries in the frequent list decrease by a third constant value moving from a top of the frequent list to a bottom of the frequent list, and wherein adjustment values for entries in the frequent ghost list decrease by a fourth constant value.

(92) Embodiment 8. The method of embodiment 1, 2, 3, 4, 5, 6, and/or 7, wherein the first, second, third, and fourth constant values can be the all the same, partially the same, or all different.

(93) Embodiment 9. The method of embodiment 1, 2, 3, 4, 5, 6, 7, and/or 8, further comprising operating the cache, wherein the cache includes an entry list having a recent list and a frequent list, by: adding new entries to a top of the recent list, moving entries in the recent list that are accessed a second time to a top of the frequent list, and moving entries in the frequent list that are accessed another time to the top of the frequent list.

(94) Embodiment 10. The method of embodiment 1, 2, 3, 4, 5, 6, 7, 8, and/or 9, further comprising storing, for each entry, a recency value and a frequency value.

(95) Embodiment 11. A method for performing any of the operations, methods, or processes, or any portion of any of these, or any combination thereof disclosed herein.

(96) Embodiment 12. A non-transitory storage medium having stored therein instructions that are executable by one or more hardware processors to perform operations comprising the operations of any one or more of embodiments 1 through 11.

(97) The embodiments disclosed herein may include the use of a special purpose or general-purpose computer including various computer hardware or software modules, as discussed in greater detail below. A computer may include a processor and computer storage media carrying instructions that, when executed by the processor and/or caused to be executed by the processor, perform any one or more of the methods disclosed herein, or any part(s) of any method disclosed.

(98) As indicated above, embodiments within the scope of the present invention also include computer storage media, which are physical media for carrying or having computer-executable instructions or data structures stored thereon. Such computer storage media may be any available physical media that may be accessed by a general purpose or special purpose computer.

(99) By way of example, and not limitation, such computer storage media may comprise hardware storage such as solid state disk/device (SSD), RAM, ROM, EEPROM, CD-ROM, flash memory, phase-change memory ("PCM"), or other optical disk storage, magnetic disk storage or other magnetic storage devices, or any other hardware storage devices which may be used to store program code in the form of computer-executable instructions or data structures, which may be accessed and executed by a general-purpose or special-purpose computer system to implement the disclosed functionality of the invention. Combinations of the above should also be included within the scope of computer storage media. Such media are also examples of non-transitory storage media, and non-transitory storage media also embraces cloud-based storage systems and structures, although the scope of the invention is not limited to these examples of non-transitory storage media.

(100) Computer-executable instructions comprise, for example, instructions and data which, when executed, cause a general-purpose computer, special purpose computer, or special purpose

processing device to perform a certain function or group of functions. As such, some embodiments of the invention may be downloadable to one or more systems or devices, for example, from a website, mesh topology, or other source. As well, the scope of the invention embraces any hardware system or device that comprises an instance of an application that comprises the disclosed executable instructions.

(101) Although the subject matter has been described in language specific to structural features and/or methodological acts, it is to be understood that the subject matter defined in the appended claims is not necessarily limited to the specific features or acts described above. Rather, the specific features and acts disclosed herein are disclosed as example forms of implementing the claims.

(102) As used herein, the term ‘module’ or ‘component’ or ‘engine’ may refer to software objects or routines that execute on the computing system. The different components, modules, engines, and services described herein may be implemented as objects or processes that execute on the computing system, for example, as separate threads. While the system and methods described herein may be implemented in software, implementations in hardware or a combination of software and hardware are also possible and contemplated. In the present disclosure, a ‘computing entity’ may be any computing system as previously defined herein, or any module or combination of modules running on a computing system.

(103) In at least some instances, a hardware processor is provided that is operable to carry out executable instructions for performing a method or process, such as the methods and processes disclosed herein. The hardware processor may or may not comprise an element of other hardware, such as the computing devices and systems disclosed herein.

(104) In terms of computing environments, embodiments of the invention may be performed in client-server environments, whether network or local environments, or in any other suitable environment. Suitable operating environments for at least some embodiments of the invention include cloud computing environments where one or more of a client, server, or other machine may reside and operate in a cloud environment.

(105) With reference briefly now to FIG. 7, any one or more of the entities disclosed, or implied, by the Figures, and/or elsewhere herein, may take the form of, or include, or be implemented on, or hosted by, a physical computing device, one example of which is denoted at **700**. As well, where any of the aforementioned elements comprise or consist of a virtual machine (VM), that VM may constitute a virtualization of any combination of the physical components disclosed in FIG. 7.

(106) In the example of FIG. 7, the physical computing device **700** includes a memory **702** which may include one, some, or all, of random-access memory (RAM), non-volatile memory (NVM) **704** such as NVRAM for example, read-only memory (ROM), and persistent memory, one or more hardware processors **706**, non-transitory storage media **708**, UI device **710**, and data storage **712**. One or more of the memory components **702** of the physical computing device **700** may take the form of solid-state device (SSD) storage. As well, one or more applications **714** may be provided that comprise instructions executable by one or more hardware processors **706** to perform any of the operations, or portions thereof, disclosed herein. One or more of the memory components **702** may include one or more caches. The caches may include multiple levels of caches.

(107) Such executable instructions may take various forms including, for example, instructions executable to perform any method or portion thereof disclosed herein, and/or executable by/at any of a storage site, whether on-premises at an enterprise, or a cloud computing site, client, datacenter, data protection site including a cloud storage site, or backup server, to perform any of the functions disclosed herein. As well, such instructions may be executable to perform any of the other operations and methods, and any portions thereof, disclosed herein.

(108) The present invention may be embodied in other specific forms without departing from its spirit or essential characteristics. The described embodiments are to be considered in all respects only as illustrative and not restrictive. The scope of the invention is, therefore, indicated by the

appended claims rather than by the foregoing description. All changes which come within the meaning and range of equivalency of the claims are to be embraced within their scope.

## Claims

1. A method, comprising: initiating a cache eviction operation for a cache, wherein entries associated with the cache are identified in an entry list, each of the entries having a timestamp of last access; generating an adjusted characteristic for each entry in the entry list, the adjusted characteristic is an adjusted timestamp of last access, the adjusted timestamp of last access of each entry being determined by adding an adjustment value to the timestamp of last access, wherein the adjustment value for different entries is different; identifying candidate entries for eviction from the cache based at least on the adjusted characteristics of the entries compared to a time threshold; and evicting at least one of the identified candidate entries from the cache based on the comparisons.
2. The method of claim 1, wherein generating an adjusted characteristic includes giving entries that are most frequently used and/or most recently used a longer time to live in the cache.
3. The method of claim 1, wherein the entry list comprises a recent list, a recent ghost list, a frequent list, and a frequent ghost list, wherein the adjustment values for entries in the recent list are different from adjustment values in the recent ghost list and wherein adjustment values for entries in the frequent list are different from adjustment values in the frequent ghost list.
4. The method of claim 3, further comprising evicting an entry from the recent ghost list or the frequent ghost list.
5. The method of claim 4, further comprising reordering entries in the frequent list and the frequent list based on number of times the entries are accessed.
6. The method of claim 4, wherein the adjustments values for entries in the recent list decrease by a constant value moving from a top of the recent list to a bottom of the recent list, wherein the adjustment values for entries in the recent ghost list decrease by a second constant value, wherein the adjustments values for entries in the frequent list decrease by a third constant value moving from a top of the frequent list to a bottom of the frequent list, and wherein adjustment values for entries in the frequent ghost list decrease by a fourth constant value.
7. The method of claim 3, further comprising operating the cache, wherein the entry list includes a recent list and a frequent list, by: adding new entries to a top of the recent list; moving entries in the recent list that are accessed a second time to a top of the frequent list; and moving entries in the frequent list that are accessed another time to the top of the frequent list.
8. The method of claim 1, further comprising storing, for each entry, a recency value and a frequency value.
9. The method of claim 8, wherein the frequency value is not affected by the adjustment characteristic.
10. The method of claim 1, wherein the adjustment characteristic prevents at least one of the entries from being a candidate for eviction based when evicting using time-based eviction operations.
11. A non-transitory storage medium having stored therein instructions that are executable by one or more hardware processors to perform operations comprising: initiating a cache eviction operation for a cache, wherein entries associated with the cache are identified in an entry list, each of the entries having a timestamp of last access; generating an adjusted characteristic for each entry in the entry list, the adjusted characteristic is an adjusted timestamp of last access, the adjusted timestamp of last access of each entry being determined by adding an adjustment value to the timestamp of last access, wherein the adjustment value for different entries is different; identifying candidate entries for eviction from the cache based at least on the adjusted characteristics of the entries compared to a time threshold; and evicting at least one of the identified candidate entries from the cache based on the comparisons.
12. The non-transitory storage medium of claim 11, wherein generating an adjusted characteristic

includes giving entries that are most frequently used and/or most recently used a longer time to live in the cache.

13. The non-transitory storage medium of claim 11, wherein the entry list comprises a recent list, a recent ghost list, a frequent list, and a frequent ghost list, wherein the adjustment values for entries in the recent list are different from adjustment values in the recent ghost list and wherein adjustment values for entries in the frequent list are different from adjustment values in the frequent ghost list.

14. The non-transitory storage medium of claim 3, further comprising evicting an entry from the recent ghost list or the frequent ghost list.

15. The non-transitory storage medium of claim 13, further comprising reordering entries in the frequent list and the frequent list based on number of times the entries are accessed.

16. The non-transitory storage medium of claim 13, wherein the adjustments values for entries in the recent list decrease by a constant value moving from a top of the recent list to a bottom of the recent list, wherein the adjustment values for entries in the recent ghost list decrease by a second constant value, wherein the adjustments values for entries in the frequent list decrease by a third constant value moving from a top of the frequent list to a bottom of the frequent list, and wherein adjustment values for entries in the frequent ghost list decrease by a fourth constant value.

17. The non-transitory storage medium of claim 14, further comprising operating the cache, wherein the entry list includes a recent list and a frequent list, by: adding new entries to a top of the recent list; moving entries in the recent list that are accessed a second time to a top of the frequent list; and moving entries in the frequent list that are accessed another time to the top of the frequent list.

18. The non-transitory storage medium of claim 11, further comprising storing, for each entry, a recency value and a frequency value.

19. The non-transitory storage medium of claim 18, wherein the frequency value is not affected by the adjustment characteristic.

20. The non-transitory storage medium of claim 11, wherein the adjustment characteristic prevents at least one of the entries from being a candidate for eviction based when evicting using time-based eviction operations.

---