

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12386863
Kind Code	B2
Date of Patent	August 12, 2025
Inventor(s)	Ganesh; Varun et al.

Micro-partition clustering based on expression property metadata

Abstract

A method for selecting micro-partitions for a clustering operation includes: storing table data in a plurality of micro-partitions of a storage device, wherein each of the plurality of micro-partitions comprises a portion of the table data, wherein subsets of the plurality of micro-partitions are associated with a respective one of a plurality of expression property (EP) files, and wherein each of the plurality of EP files comprises an EP data region that represents the portions of the table data of the subset of the plurality of micro-partitions associated with the EP file; determining sub-ranges of the table data based on the EP data regions of the plurality of EP files; selecting a subset of the plurality of EP files for a clustering operation based on the sub-ranges of the table data; and performing the clustering operation on the micro-partitions associated with the subset of the EP files.

Inventors: Ganesh; Varun (San Carlos, CA), Jou; Alvin E. (Woodinville, WA), Kang; Donghe (Columbus, OH), Shelly; Ryan Michael Thomas (San Francisco, CA), Yan; Jiaqi (Menlo Park, CA), Zhu; Yizhi (Bellevue, WA)

Applicant: SNOWFLAKE INC. (Bozeman, MT)

Family ID: 1000008749036

Assignee: Snowflake Inc. (Bozeman, MT)

Appl. No.: 18/302234

Filed: April 18, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20240354315 A1	Oct. 24, 2024

Publication Classification

Int. Cl.: G06F16/28 (20190101); G06F16/2455 (20190101)

U.S. Cl.:

CPC G06F16/285 (20190101); G06F16/24556 (20190101);

Field of Classification Search

USPC: None

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
2016/0162599	12/2015	Dickie et al.	N/A	N/A
2017/0264640	12/2016	Narayanaswamy et al.	N/A	N/A
2018/0068008	12/2017	Cruanes	N/A	G06F 16/211
2020/0167361	12/2019	Princehouse et al.	N/A	N/A
2020/0272628	12/2019	Dageville	N/A	G06F 16/221
2022/0253441	12/2021	Baptist	N/A	N/A
2022/0382751	12/2021	Dhuse et al.	N/A	N/A

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
110100242	12/2018	CN	N/A
110825794	12/2019	CN	N/A
20210135548	12/2020	KR	N/A

Primary Examiner: Singh; Amresh

Attorney, Agent or Firm: WOMBLE BOND DICKINSON (US) LLP

Background/Summary

TECHNICAL FIELD

(1) The present disclosure relates to databases and database management, and, more particularly, relates to clustering micro-partitions containing database data.

BACKGROUND

(2) Databases are widely used for data storage and access in computing applications. Databases may include one or more tables that include data that can be joined, read, modified, or deleted using queries. Databases can store small or extremely large sets of data within one or more tables. This data can be accessed by various users in an organization or even be used to service public users, such as via a website or an application programming interface (API). The large amount of data that can be contained within a database can often be useful for various types of data analytics, which involves the attempt to determine conclusions and/or predictions based on analysis of the information contained in the data. When working with large volumes of data, improvements to the

efficiency of database queries can provide significant improvements to the cost and time associated with the data analysis.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) The described embodiments and the advantages thereof may best be understood by reference to the following description taken in conjunction with the accompanying drawings. These drawings in no way limit any changes in form and detail that may be made to the described embodiments by one skilled in the art without departing from the scope of the described embodiments.
- (2) FIG. 1 is a block diagram depicting an example computing environment in which the methods disclosed herein may be implemented.
- (3) FIG. 2 is a schematic diagram of a data structure for storage of database metadata, according to some embodiments of the present disclosure.
- (4) FIG. 3 illustrates an example grouping of micro-partitions, micro-partition metadata, and EP files of FIG. 2, according to some embodiments of the present disclosure.
- (5) FIG. 4 illustrates an example modification of a micro-partition, according to some embodiments of the present disclosure.
- (6) FIG. 5 is a block diagram illustrating techniques for the selection of micro-partitions for clustering, according to some embodiments of the present disclosure.
- (7) FIG. 6 is a flow diagram of one embodiment of a method for selecting micro-partitions for clustering, according to some embodiments of the present disclosure.
- (8) FIG. 7 is a diagram of a data structure illustrating an example scenario of EP files for the selection of micro-partitions for clustering, according to some embodiments of the present disclosure.
- (9) FIG. 8A is a flow diagram of one embodiment of a method for selecting a seed EP data region, according to some embodiments of the present disclosure.
- (10) FIG. 8B is a block diagram of the data structure illustrating an example scenario of selecting the seed EP data region, according to some embodiments of the present disclosure.
- (11) FIG. 9A is a flow diagram of one embodiment of a method for selecting clustering EP data regions, according to some embodiments of the present disclosure.
- (12) FIG. 9B is a block diagram of the data structure illustrating an example scenario of selecting clustering EP data regions, according to some embodiments of the present disclosure.
- (13) FIG. 10A is a flow diagram of one embodiment of a method for selecting clustering EP data regions, according to some embodiments of the present disclosure.
- (14) FIG. 10B is a diagram of the data structure illustrating an example scenario of selecting the clustering EP data regions, according to some embodiments of the present disclosure.
- (15) FIG. 11 is a flow diagram of one embodiment of a method for selecting micro-partitions for a clustering operation, according to some embodiments of the present disclosure.
- (16) FIG. 12 is a block diagram of an example computing device that may perform one or more of the operations described herein, in accordance with some embodiments of the present disclosure.

DETAILED DESCRIPTION

- (17) Databases are widely used for data storage and access in computing applications. A goal of database storage is to provide enormous sums of information in an organized manner so that it can be accessed, managed, and updated. In some embodiments of a database, data may be organized into rows, columns, and tables. Different database storage systems may be used for storing different types of content, such as bibliographic, full text, numeric, and/or image content. Further, in computing, different database systems may be classified according to the organization approach of the database. There are many different types of databases, including relational databases,

distributed databases, cloud databases, object-oriented databases, and others.

(18) Queries can be executed against database data to find certain data within the database and respond to a question about the database data. A database query extracts data from the database and formats it into a readable form. For example, when a user wants data from a database, the user may write a query in the language required by the database. The query may request specific information from the database. For example, if the database includes information about sales transactions made by a retail store, a query may request all transactions for a certain product during a certain time frame. The query may request any pertinent information that is stored within the database. If the appropriate data can be found to respond to the query, the database has the potential to reveal complex trends and activities.

(19) However, when a database becomes very large and includes vast sums of data, it can be very difficult to respond to a database query. Further to the above example, if the database includes a record of all sales transactions for the retail store over an extensive time period, the database may include multiple tables that each include billions of rows of information divided into hundreds or thousands of columns. If a user requests all transactions for a certain product across the entire history of the database, it can require extensive computing resources and time to scan the entire database to find each of the requested transactions.

(20) Databases may further include metadata, or information about the database data, to aid in organizing the database and responding to queries on the database. Metadata is data information that provides information about other data. For example, metadata about an image file may include information such as the date and time the image was captured, the camera that captured the image, the camera settings when the image was captured, a file size of the image, a name of the image, and so forth. Further for example, metadata about a table in a database may include information such as the minimum value in the table, the maximum value in the table, the number of rows in the table, the number of columns in the table, the type of data stored in the table, the subject of the data stored in the table, and so forth.

(21) Metadata can be useful for responding to a database query. For example, metadata may be helpful in performing various types of performance enhancements for queries. One such enhancement is pruning. Pruning is described, for example, in U.S. Pat. No. 10,437,780, entitled "Data pruning based on metadata," the entire contents of which are incorporated herein by reference. In some embodiments of data pruning, the predicates of a database query are examined to determine the ranges and/or types of data being queried. Metadata for the data stored in the table may then be used to rule out certain portions of the data storage from being scanned. By reducing the amount of data being scanned, the overall performance of the database will improve.

(22) Pruning can be a very important part of query compilation in complex databases, such as those described herein. However, particularly in the context of a large database, the metadata itself can become very large and require extensive computing resources and time to just scan the metadata without scanning any of the database data. In certain implementations it can be useful to employ an organized and efficient metadata structure. With multi-level metadata, a database can effectively have metadata about groups of metadata. In such a hierarchical structure, higher level metadata may be queried before lower level metadata, which may allow for less of the metadata to be loaded.

(23) However, even with such multi-level metadata, additional improvements may still be made. By aligning the groups of the metadata with underlying organization strategies of the database tables, such as data clustering, additional improvements to the metadata grouping may be achieved. Described herein are systems, methods, and computer program products that intelligently organize groups of metadata so as to reduce the amount of metadata that is loaded when query pruning and other types of database management are performed. Methods described herein reduce the amount of memory utilized for query compilation and allow for fewer resources to be used to perform other database tasks such as data compaction.

(24) FIG. 1 is a block diagram of an example computing environment **100** in which the systems

and methods disclosed herein may be implemented. In particular, a cloud computing platform **110** may be implemented, such as AMAZON WEB SERVICES™ (AWS), MICROSOFT AZURE™, GOOGLE CLOUD™ or GOOGLE CLOUD PLATFORM™, or the like. As known in the art, a cloud computing platform **110** provides computing resources and storage resources that may be acquired (purchased) or leased and configured to execute applications and store data. The cloud computing platform **110** may be accessed by a client **101** (e.g., a client device). Non-limiting examples of client devices include a smart phone **101A**, personal computer **101B**, laptop **101C**, tablet computer **101D**, server computer **101E**, and/or another type of device that can process data.

(25) FIG. 1 and the other figures may use like reference numerals to identify like elements. A letter after a reference numeral, such as “**101A**,” indicates that the text refers specifically to the element having that particular reference numeral. A reference numeral in the text without a following letter, such as “**101**,” refers to any or all of the elements in the figures bearing that reference numeral.

(26) In some embodiments, client devices **101** may access the cloud computing platform **110** over a network **105**. Network **105** may be a public network (e.g., the internet), a private network (e.g., a local area network (LAN) or wide area network (WAN)), or a combination thereof. In one embodiment, network **105** may include a wired or a wireless infrastructure, which may be provided by one or more wireless communications systems, such as a WIFI® hotspot connected with the network **105** and/or a wireless carrier system that can be implemented using various data processing equipment, communication towers (e.g., cell towers), etc. The network **105** may carry communications (e.g., data, message, packets, frames, etc.) between the various components of the cloud computing platform **110** and one more of the client devices **101**.

(27) The cloud computing platform **110** may host a cloud computing service **112** that facilitates storage of data on the cloud computing platform **110** (e.g., data management and access) and analysis functions (e.g., SQL queries, analysis), as well as other computation capabilities (e.g., secure data sharing between users of the cloud computing platform **110**). The cloud computing platform **110** may include a three-tier architecture: data storage **140**, query processing **130**, and cloud services **120**.

(28) Data storage **140** may facilitate the storing of data on the cloud computing platform **110** in one or more cloud databases **141**. Data storage **140** may use a storage service such as AMAZON S3 to store data and query results on the cloud computing platform **110**. In particular embodiments, to load data into the cloud computing platform **110**, data tables may be horizontally partitioned into large, immutable files which may be analogous to blocks or pages in a traditional database system. Within each file, the values of each attribute or column are grouped together and compressed using a scheme sometimes referred to as hybrid columnar. Each table has a header which, among other metadata, contains the offsets of each column within the file.

(29) In addition to storing table data, data storage **140** facilitates the storage of metadata, temp data generated by query operations (e.g., joins), as well as the data contained in large query results. This may allow the system to compute large queries without out-of-memory or out-of-disk errors. Storing query results this way may simplify query processing as it removes the need for server-side cursors found in traditional database systems.

(30) Query processing **130** may handle query execution by compute nodes within elastic clusters of virtual machines, referred to herein as virtual warehouses or data warehouses. Thus, query processing **130** may include one or more virtual warehouses **131** having one or more compute nodes **132**, which may also be referred to herein as data warehouses. The virtual warehouses **131** may be one or more virtual machines operating on the cloud computing platform **110**. The virtual warehouses **131** may be compute resources that may be created, destroyed, or resized at any point, on demand. This functionality may create an “elastic” virtual warehouse **131** that expands, contracts, or shuts down according to the user's needs. Expanding a virtual warehouse **131** involves generating one or more compute nodes **132** to the virtual warehouse **131**. Contracting a virtual warehouse **131** involves removing one or more compute nodes **132** from the virtual warehouse **131**.

More compute nodes **132** may lead to faster compute times. For example, a data load which takes fifteen hours on a system with four nodes might take only two hours with thirty-two nodes.

(31) Cloud services **120** may be a collection of services (e.g., computer instruction executing on a processing device) that coordinate activities across the cloud computing service **112**. These services tie together all of the different components of the cloud computing service **112** in order to process user requests, from login to query dispatch. Cloud services **120** may operate on compute instances provisioned by the cloud computing service **112** from the cloud computing platform **110**. Cloud services **120** may include a collection of services that manage virtual warehouses, queries, transactions, data exchanges, and the metadata associated with such services, such as database schemas, access control information, encryption keys, and usage statistics. Cloud services **120** may include, but not be limited to, authentication engine **121**, infrastructure manager **122**, optimizer **123**, exchange manager **124**, security engine **125**, and metadata storage **126**. Though metadata storage **126** is illustrated as being separate from the data storage **140** in FIG. 1, this is merely for schematic purposes and is not intended to limit the present disclosure. In some embodiments, the metadata storage **126** is co-located with the data storage **140**.

(32) In one embodiment, the cloud computing service **112** can perform, or cause to be performed, data management operations on the data storage **140**. In some embodiments, the data management operations may include the generation of metadata for data of the data storage **140**. In some embodiments, multiple levels of metadata may be formed, one level of the metadata containing metadata values for a group of metadata in another level of the metadata. In some embodiments, the groups of the metadata may be organized at least partially based on data organization metrics of the data storage **140**. For example, in some embodiments, the groups of the metadata may be organized at least partially based on a clustering key of the data storage **140**.

(33) FIG. 2 is a schematic diagram of a data structure **200** for storage of database metadata, according to some embodiments of the present disclosure. The data structure **200** includes table metadata **202** pertaining to database data stored across a table of the database. The table may be composed of multiple micro-partitions (MP) **210**. Each of the multiple micro-partitions **210** may include numerous rows and columns making up cells of database data.

(34) In some embodiments, the plurality of micro-partitions **210** of the table may be provided as immutable storage devices. When a transaction is executed on such a table, all impacted micro-partitions **210** are recreated to generate new micro-partitions **210** that reflect the modifications of the transaction. After a transaction is fully executed, any original micro-partitions **210** that were recreated may then be removed from the database. In some embodiments, a new version of the table is generated after each transaction that is executed on the table. The table may undergo many versions over a time period if the data in the table undergoes many changes, such as inserts, deletes, updates, and/or merges.

(35) The table metadata **202** may include global information about the table of a specific version. In some embodiments, the table metadata **202** may include a table identification and versioning information indicating, for example, how many versions of the table have been generated over a time period, which version of the table includes the most up-to-date information, how the table was changed over time, and so forth. A new table version may be generated each time a transaction is executed on the table, where the transaction may include a data manipulation language (DML) statement such as an insert, delete, merge, and/or update command. Each time a DML statement is executed on the table, and a new table version is generated, one or more new micro-partitions **210** may be generated that reflect the DML statement.

(36) The data structure **200** further includes micro-partition metadata (MP MD) **214** (also referred to herein as a micro-partition metadata file **214**) that includes metadata about a micro-partition **210** of the table. In some embodiments, each of the micro-partition metadata **214** may include information about at least one respective micro-partition **210** of the table. The micro-partitions **210** illustrated in FIG. 2 may include database data that is separate from the metadata storage. In some

embodiments, the micro-partition metadata **214** may be stored for each column of each micro-partition **210** of the table. The micro-partition metadata **214** pertaining to a particular micro-partition **210** may include any suitable information about data of a portion of the table (e.g., a column) stored in the micro-partition **210**, including for example, a minimum and maximum for the data stored in a portion of the database (e.g., a column), a type of data stored in the portion of the database, a subject of the data stored in the portion of the database, versioning information for the data stored in the portion of the database, file statistics for all micro-partitions **210** in the table, global cumulative expressions for columns of the table, and so forth. Each column of each micro-partition **210** of the table may be associated with one or more micro-partition metadata elements **214**. As illustrated in the example embodiment shown in FIG. 2, the table metadata **202** includes micro-partition metadata **214** for each micro-partition **210**. However, the embodiments of the present disclosure are not limited thereto. It should be appreciated that the table may include any number of micro-partitions **210** and any number of micro-partition metadata **214**, and each micro-partition **210** may include any number of columns. The micro-partitions **210** may have the same or different columns and may have different types of columns storing different information.

(37) The data structure **200** further includes a plurality of expression property files (EP or EP files) **220A**, **220B**, and **220C** (which may be collectively referenced herein as EP files **220**). The EP files **220** may include aggregated micro-partition statistics, cumulative column properties, and so forth about a micro-partition **210** or a collection of micro-partitions **210** of the table. A micro-partition metadata file **214** may be stored for each column of each micro-partition **210** of the table, and an EP file **220** may be stored for a collection of micro-partition metadata files **214** associated with a plurality of micro-partitions **210** of the table as illustrated in FIG. 2.

(38) The data structure **200** may include micro-partition metadata **214** for each micro-partition **210** of the table. The micro-partition metadata **214** may include a minimum/maximum data point for the corresponding micro-partition **210**, a type of data stored in the corresponding micro-partition **210**, a micro-partition structure of the corresponding micro-partition **210**, and so forth. The micro-partition metadata **214** may be stored as part of a micro-partition that is allocated for metadata (e.g., as opposed to data of the database). A micro-partition allocated for metadata may be persisted in immutable storage and the EP files **220** may also be stored within the metadata micro-partition in immutable storage. A metadata manager may maintain all metadata micro-partitions, including metadata micro-partitions comprising the EP files **220** and/or the micro-partition metadata **214**.

(39) The cumulative table metadata **202** may include global information about all micro-partitions **210** within the applicable table. For example, the cumulative table metadata **202** may include a global minimum and global maximum for the entire table, which may include millions or even hundreds of millions of micro-partitions **210**. The cumulative table metadata **202** may include any suitable information about the data stored in the table, including, for example, minimum/maximum values, null count, a summary of the database data collectively stored across the table, a type of data stored across the table, a distinct version for the data stored in the table, and so forth.

(40) The EP files **220A-220C** may include information about database data stored in an associated grouping **240** of micro-partitions **210**. In the example of FIG. 2, an example EP file **220A** is associated with a grouping **240** of three micro-partition metadata **214** that are respectively associated with three micro-partitions **210**. The example EP file **220A** may include information about those three different micro-partitions **210**. An EP file **220** may include any suitable information about the grouping **240** of the micro-partitions **210** and/or micro-partition metadata **214** with which it is associated. For example, an EP file **220** may include a cumulative minimum/maximum for the collective grouping **240** of micro-partitions **210**, a minimum/maximum for each of the micro-partitions **210** within the grouping **240**, a cumulative null count, a null count for each of the micro-partitions **210** within the grouping **240**, a cumulative summary of data collectively stored across the grouping **240** of micro-partitions **210**, a summary of data stored in each of the micro-partitions **210** in the grouping **240**, and so forth.

(41) The EP files **220** illustrated in FIG. 2 provide valuable cumulative metadata pertaining to a collection of micro-partition metadata **214** and micro-partitions **210** of the database. Further, each of the micro-partition metadata **214** provide valuable information about respective ones of the micro-partitions **210** of the table.

(42) FIG. 3 illustrates an example grouping **240** of micro-partitions **210**, micro-partition metadata **214**, and EP files **220** of FIG. 2, according to some embodiments of the present disclosure. The grouping **240** is an example to illustrate the EP files **220**, and is not intended to limit the embodiments of the present disclosure.

(43) Referring to FIG. 3, the data structure **200** may include grouping **240** of micro-partitions **210**. More specifically, the example grouping **240** may include micro-partition **210A**, micro-partition **210B**, and micro-partition **210C** (referred to collectively as micro-partitions **210**). Each of the micro-partitions may be respectively associated with a micro-partition metadata file **214**. For example, micro-partition metadata file **214A** may be associated with micro-partition **210A**, micro-partition metadata file **214B** may be associated with micro-partition **210B**, and micro-partition metadata file **214C** may be associated with micro-partition **210C**. An EP file **220A** may be associated with the micro-partition **210A**, micro-partition **210B**, and micro-partition **210C**, as well as their associated micro-partition metadata files **214** (e.g., micro-partition metadata files **214A**, **214B**, **214C**).

(44) As described herein, the micro-partitions **210** may contain one or more portions of data from a database. For example, the micro-partitions **210** may contain one or more columns and/or rows from a table of the database. Example data is provided in FIG. 3 to aid in understanding. For example, micro-partition **210A** contains a first column including data elements **180**, **200**, and **150** along with a second column including data elements “Alpha,” “Beta,” and “Gamma.” Micro-partition **210B** contains the first column including data element values **450**, **500**, and **400** along with the second column including data element values “Delta,” “Epsilon,” and “Zeta.” Micro-partition **210C** contains the first column including data element values **175**, **220**, and **180** along with the second column including data elements “Eta,” “Theta,” and “Iota.” As previously noted, these data values are merely examples and not intended to limit the embodiments of the present disclosure.

(45) Each of the micro-partitions **210** may be associated with at least one micro-partition metadata file **214**. The micro-partition metadata file **214** may maintain, in part, metadata describing the data values of the micro-partitions **210**. In FIG. 3, an example is shown for metadata of the micro-partition metadata file **214** regarding the maximum (max) and minimum (min) of data of the micro-partition **210**. However, as described herein, a number of different types of metadata may be maintained in the micro-partition metadata file **214**. For example, the metadata **214** of the micro-partition **210** may include values tracking null values of the data of the micro-partition **210**, length of the data, and so on.

(46) In the example of FIG. 3, micro-partition metadata file **214A** is associated with micro-partition **210A**. The metadata of the micro-partition metadata file **214A** tracks that the maximum (max) of the data of the first column of the micro-partition **210A** is 200 and the minimum (min) of the data of the first column of the micro-partition **210A** is 150. Micro-partition metadata file **214B** is associated with micro-partition **210B**. The metadata of the micro-partition metadata file **214B** tracks that the maximum (max) of the data of the first column of the micro-partition **210B** is 500 and the minimum (min) of the data of the first column of the micro-partition **210B** is 400. Micro-partition metadata file **214C** is associated with micro-partition **210C**. The metadata of the micro-partition metadata file **214C** tracks that the maximum (max) of the data of the first column of the micro-partition **210C** is 220 and the minimum (min) of the data of the first column of the micro-partition **210C** is 175. The minimum and maximum data values maintained by the respective micro-partition metadata **214** represents a smallest value and a largest value, respectively, of the data for a particular column (e.g., the first column) for a respective micro-partition **210**.

(47) The micro-partition metadata files **214A**, **214B**, and **214C** may be associated with an EP file **220**. In the example of FIG. 3, micro-partition metadata files **214A**, **214B**, and **214C** are each associated with EP file **220A**. The metadata of the EP file **220A** tracks that the maximum (max) of the metadata of the first column tracked by the micro-partition metadata files **214A**, **214B**, and **214C** is 500 and the minimum (min) of the metadata of the first column tracked by the micro-partition metadata files **214A**, **214B**, and **214C** is 150. The minimum and maximum data values maintained by the EP file **220** represents a smallest value and a largest value, respectively of the metadata for a particular column (e.g., the first column) for the micro-partition metadata files **214** within the grouping **240**.

(48) The use of the micro-partition metadata files **214** and the EP files **220** may increase the efficiency of query processing. For example, as part of data pruning for a query, it may be determined that a particular query utilizes a predicate that restricts data values returned to those in which the first column has a value less than 100. By reviewing the min and max of the EP file **220**, it can be determined that none of the micro-partition metadata files **214** associated with the EP file **220** have a data value in the first column that would match this query. Subsequently, the micro-partitions **210** of the grouping **240** may be skipped in a scan of data to find values that should be returned for the query.

(49) As another example, as part of data pruning for a query, it may be determined that a particular query utilizes a predicate that restricts data values returned to those in which the first column has a value less than 200. By reviewing the min and max of the EP file **220**, it can be determined that at least one of the micro-partition metadata files **214** associated with the EP file **220** has a data value in the first column that would match this query. Responsive to this determination, the metadata for the micro-partition metadata files **214** associated with the EP file **220** may be examined. From that examination, it can be determined that the micro-partition metadata files **214A** and **214C** have min and max values indicating that their respective micro-partitions **210** contain data that would match this query. Subsequently, the micro-partition **210B** of the grouping **240** may be skipped and the micro-partitions **210A** and **210C** may be examined in a scan of data to find values that should be returned for the query.

(50) The data structure **200** shown in FIGS. 2 and 3 increases efficiency when responding to database queries. A database query may request any collection of data from the database and may be used to create advanced analyses and metrics about the database data. Some queries, particularly for a very large database, can be extremely costly to run both in time and computing resources. When it is necessary to scan metadata and/or database data for each file or micro-partition of each table of a database, it can take many minutes or even hours to respond to a query. In certain implementations this may not be an acceptable use of computing resources. The data structure **200** disclosed herein provides increased metadata granularity and enables multi-level pruning of database data.

(51) During compilation and optimization of a query on the database, a processing device may scan the cumulative table metadata **202** to determine if the table includes information pertaining to the query. In response to determining, based on the cumulative table metadata **202**, that the table includes information pertaining to the query, the processing device may scan each of the EP files **220** to determine which grouping of micro-partitions **210** of the table include information pertaining to the query. In response to determining, based on a first EP file **220**, that a first grouping **240** of micro-partitions **210** does not include information pertaining to the query, the processing device may discontinue database scanning of that first grouping **240** of micro-partitions **210**.

(52) In response to determining, based on a second EP file **220**, that a second grouping **240** of micro-partitions **210** includes information pertaining to the query, the processing device may proceed to scan micro-partition metadata files **214** for that second grouping **240** of micro-partitions **210**. The processing device may efficiently determine which micro-partitions **210** include pertinent data and which columns of which micro-partitions **210** do not include pertinent data. The

processing device may proceed to scan only the relevant column(s) and micro-partition(s) **210** that include information relevant to a database query. This provides a cost efficient means for responding to a database query by way of multi-level pruning based on multi-level table metadata. (53) Further to increase the cost efficiency of database queries, a resource manager (e.g., cloud services **120** of FIG. 1) may store the cumulative table metadata **202** in a cache for faster retrieval. Metadata for the database may be stored in a metadata store separate and independent of a plurality of shared storage devices collectively storing database data. In some embodiments, metadata for the database may be stored within the plurality of shared storage devices collectively storing database data. In some embodiments, metadata may be stored in metadata-specific micro-partitions that do not include database data, and/or may be stored within micro-partitions that also include database data. The metadata may be stored across disk storage, such as the plurality of shared storage devices, and it may also be stored in cache within the resource manager.

(54) Though the use of multi-level metadata may provide benefits, it may still introduce challenges. For example, the generation of the groupings **240** may be made in a naive manner. For example, as new micro-partitions **210** are created (along with new micro-partition metadata files **214**), they may be grouped based on time of creation and/or registration. This may result in configurations that are suboptimal. FIG. 4 illustrates an example modification of a micro-partition **210**, according to some embodiments of the present disclosure.

(55) The example of FIG. 4 illustrates a scenario in which the micro-partition metadata **214** and the micro-partitions **210** are stored in immutable storage. In such a scenario, a modification to data within a micro-partition **210** may result in the formation of a new micro-partition **210** rather than an in-situ modification of the current micro-partition **210** containing the data. Though the example of FIG. 4 concerns immutable storage, the embodiments of the present disclosure are not limited to such a configuration.

(56) In the example of FIG. 4, it is assumed that a data modification is performed to the data of micro-partition **210B** (e.g., as a result of a DML transaction). For example, the data of the first column that previously had a value of 450 (illustrated by reference number **410**) may be changed to have a value of 50. In such a scenario, new micro-partitions **210D** and **210E** may be created.

(57) For example, micro-partition **210D** contains a first column including data element values **500** and **400** along with a second column including data elements “Epsilon” and “Zeta.” The data of micro-partition **210D** may be the data from the original micro-partition **210B** that has been moved, e.g., due to the immutability of micro-partition **210B**. Micro-partition **210E** contains the first column including data element **50** along with the second column including the data element “Delta.” The data of micro-partition **210E** may be the data from the original micro-partition **210B** that has been updated, e.g., due to the immutability of micro-partition **210B**. In some embodiments, two new micro-partitions **210D**, **210E** may be created rather than forming a single new micro-partition **210** for a number of potential reasons. For example, as will be discussed further herein, the data of the micro-partitions **210** may be organized based on clustering keys associated with the data.

(58) Each of the new micro-partitions **210D**, **210E** will each be associated with a micro-partition metadata file **214**. For example, micro-partition metadata **214D** is associated with micro-partition **210D**. The metadata of the micro-partition metadata **214D** tracks that the maximum (max) of the data of the first column of the micro-partition **210D** is 500 and the minimum (min) of the data of the first column of the micro-partition **210D** is 400. Micro-partition metadata **214E** is associated with micro-partition **210E**. The metadata of the micro-partition metadata **214E** tracks that the maximum (max) of the data of the first column as well as the minimum (min) of the data of the first column of the micro-partition **210E** is 50.

(59) The micro-partition metadata files **214D**, **214E** may also be associated with EP file **220A'**. Due to the update of the data that resulted in the formation of the new micro-partition **210E**, the minimum (min) of the metadata tracked by the micro-partition metadata file **214E** may now be 50.

The EP file **220A'** may either be an updated form of the previous EP file **220A**, or may be a new EP file **220A'** formed due to the EP file **220A** being formed in immutable storage. The updated EP file **220A'** may include micro-partition metadata files **214A**, **214C**, **214D**, and **214E**.

(60) The update of the minimum value of the EP file **220A'** to 50 now means that the minimum and maximum values of the EP file **220A'** cover a much larger range. As a result of covering a much larger range, the efficiency of the pruning may be impacted because the pruning may not be able to deselect as many micro-partition metadata files **214** (and their associated micro-partitions **210**) for scanning. For example, in the prior-discussed scenario in which a predicate of a query selected data based on the data of the first column being less than 100, pruning would allow all of the micro-partitions **210** to be skipped in the example of FIG. 3. With this update to the data, the same query will now require all of the micro-partition metadata **214** to be scanned, despite the fact that only one of the micro-partitions **210**, micro-partition **210E**, meets this predicate.

(61) In some embodiments, since the data represented by the micro-partitions **210** may include different types of data having different values, clustering may be performed based on metadata associated with a particular data value and/or values of the micro-partitions **210**. For example, in some embodiments, the metadata that is used for determining membership in an EP file **220** may be based on a clustering key used to organize the micro-partitions **210**. A clustering key is a subset of columns in a table (or expressions on a table, such as SQL expressions) that are designated to co-locate the data in the table in the same micro-partitions **210**. When specified, the clustering key may be utilized to identify portions of the table which are to be used for making decisions about which data to co-locate within a micro-partition **210**. For example, if a particular column of a table is identified as a clustering key, values of the data of that particular column may be utilized to identify rows of the table that are to be co-located in a same micro-partition **210**. For example, rows of the table having values of the first particular column that are relatively close to one another (e.g., within a threshold numeric distance if the first column is a numeric value) may be located in a same micro-partition **210**. In some embodiments, the clustering key may be specified by the administrator, such as when a database table is created. In some embodiments, the clustering keys that are utilized to co-locate data in micro-partitions **210** may also be utilized for determining the membership in EP files **220**. Clustering keys, as well as clustering in general, are discussed in U.S. Pat. No. 10,817,540, entitled "Incremental clustering maintenance of a table," the entire contents of which are incorporated herein by reference.

(62) In some embodiments, to avoid and/or reduce issues in which the EP files **220** and/or the micro-partitions **210** cover large ranges, rendering pruning inefficient, the EP files **220** and/or the micro-partitions **210** may be re-clustered (e.g., periodically). Clustering may include reorganizing the data of the EP files **220** and/or the micro-partitions **210** to form EP files **220** and/or the micro-partitions **210** that cover smaller ranges. In a system in which the EP files **220** and/or the micro-partitions **210** are stored in immutable storage, clustering may include forming new EP files **220** and/or new micro-partitions **210** having different data and/or micro-partition metadata **214**, as described, for example, with respect to FIG. 4.

(63) In some embodiments, the selection of micro-partitions **210** for clustering members of an EP file **220** may be performed based on how well-clustered the EP files **220** of the database are determined to be. A table may be defined as clustered based on a certain order-preserving function which takes data in each row as input if rows that are close in evaluation of this function are also close together in their physical ordering. The degree of clustering (clustering ratio) of a table may be determined by the proportion of rows in the physical layout of the table that satisfy such ordering criteria. Perfect clustering is achieved when for any two rows in the table that are adjacent in their physical layout, no third row can be found that yields a closer distance to both rows according to the ordering function. For tables stored as micro-partitions **210**, clustering improves the probability that rows closer according to the ordering function should reside in the same micro-partition **210**.

(64) In some embodiments, the micro-partitions **210** may be scanned to determine if the clustering ratio may be improved and, if so, clustering may be performed on the micro-partitions **210**. If the micro-partitions **210** are stored in immutable storage, the clustering of the micro-partitions **210** may involve the deletion and recreation of micro-partitions **210**. When the micro-partitions **210** are recreated, they may be associated with micro-partition metadata files **214**, which may subsequently be organized into EP files **220**, as described herein.

(65) In a similar manner, the EP files **220** may be scanned to determine if the clustering ratio of the metadata range encompassed by the EP file **220** may be improved. If so, the members of the EP file **220** may be re-clustered. The clustering of the EP file **220** may involve the movement of assignments of micro-partition metadata files **214** between EP files **220**. For example, a micro-partition metadata file **214** may be removed from one EP file **220** to another EP file **220**. In some embodiments, selection of micro-partition metadata files **214** for membership in an EP file **220** may be performed in a similar matter as described herein (e.g., based on a metadata range for a data value, such as a clustering key, and/or a time of creation/registration). If the EP files **220** are stored in immutable storage, the clustering of the EP files **220** may involve the deletion and recreation of EP files **220**. In some embodiments, an improvement to the clustering quality of the micro-partitions **210** may also involve the deletion and recreation of the micro-partitions **210**, which may cause the EP files **220** to be modified based on modifications to the underlying micro-partitions **210**.

(66) Clustering the EP files **220** and/or the micro-partitions **210** may be performed by looking at how storage of a clustering key is represented within the EP files **220** and/or the micro-partitions **210** and reforming the EP files **220** and/or the micro-partitions **210** to more closely locate values of the clustering keys stored in the micro-partitions **210** within a same EP file **220** and/or the micro-partition **210**. In a theoretical scenario, all of the micro-partitions **210** of a given database could be analyzed during clustering, and new micro-partitions **210** formed in which values of the clustering key that were close to one another within a data range were also clustered within a same micro-partition **210**. New micro-partition metadata **214** could then be formed for the data ranges of the new micro-partitions **210**, and new EP files **220** could be formed based on the new micro-partition metadata **214**. However, in a practical sense, analysis of all of the micro-partitions **210** and/or micro-partition metadata **214** may take an unreasonable amount of time and/or computing resources. For example, to understand how the different ranges of the clustering key are stored across large numbers of micro-partitions **210**, large amounts of data may be cached, which may make such an analysis impractical. Moreover, if all of the micro-partitions **210** are potentially re-clustered, it may set up a scenario in which a large number of the EP files **220** associated with the micro-partitions **210** are also re-clustered due to the changes in the micro-partitions **210**. This may cause extensive churn within the database and negatively affect performance.

(67) Embodiments of the present disclosure arise, at least in part, from a recognition that clustering may be improved by selectively analyzing a subset of the micro-partitions **210** for clustering. By focusing clustering on micro-partitions **210** that are more likely to have overlapping values (e.g., with respect to the data values of a clustering key) micro-partitions **210** may be selected that have a high potential of providing improvements to the clustering of the data storage. Embodiments according to the present disclosure may be far more memory efficient and may reduce churn across the entire table, which can affect query performance when EP files **220** are to be reloaded into the cache. By selecting micro-partitions **210** from a subset of EP files **220**, a number of the EP files **220** that are to be reloaded may be limited, and this may also improve the efficiency of clustering of the micro-partitions **210** because it may reduce the amount of work associated with the clustering and allow the clustering of the micro-partitions **210** to converge.

(68) Embodiments of the present disclosure may improve an operation of the computer by reducing an amount of memory and/or processing resources utilized to provide a database operation. Moreover, embodiments of the present disclosure may improve the technology associated with

databases and database processing by providing a more effective mechanism to cluster data within the database storage and, thereby, improve an efficiency and/or speed of a data query (e.g., via pruning).

(69) It may be understood that the clustering of the EP files **220** may be independent of a clustering of the micro-partitions **210**. For example, in some embodiments, a clustering ratio of the micro-partitions **210** may be at least partially independent of the clustering ratio of the EP file **220**. For example, in some embodiments, clustering of the EP files **220** may be performed regardless of a clustering state of the micro-partitions **210**. In some embodiments, clustering of the micro-partitions **210** may be performed regardless of a clustering state of the EP files **220**.

(70) In some embodiments, clustering of the micro-partitions **210** may be performed based on one or more clustering statistics of the EP files **220** and/or the micro-partition metadata files **214** of the micro-partitions **210**. Statistics that may be used to determine a clustering level of an EP file **220** may include: a number of micro-partitions **210** per micro-partition metadata file **214**; an average depth and overlap per micro-partition metadata file **214**; an average depth and overlap for the micro-partition metadata files **214** of the EP file **220**; a maximum depth and/or overlap per micro-partition metadata file **214**; a maximum depth and/or overlap for the micro-partition metadata files **214** of the EP file **220**; a median depth and/or overlap per micro-partition metadata file **214**; a median depth and/or overlap for the micro-partition metadata files **214** of the EP file **220**; a distribution (e.g., a percentile) of the depth and/or overlap for the micro-partition metadata files **214**; a distribution (e.g., a percentile) of the depth and/or overlap for the micro-partition metadata files **214** of the EP file **220**; and/or an average range of the EP file **220**. The depth for a micro-partition **210** may be the maximum number of intersected micro-partitions **210** at any given range. The overlap for a particular micro-partition **210** is the total number of other micro-partitions **210** intersected with the particular micro-partition **210**. A first micro-partition **210** may overlap a second micro-partition **210** if the data values between the minimum and maximum data values (e.g., of the clustering key) of the first micro-partition **210** intersect the data values between the minimum and maximum data values (e.g., of the clustering key) of the second micro-partition intersection **210**. The above characteristics are merely examples, and not intended to limit the embodiments of the present disclosure.

(71) FIG. 5 is a block diagram illustrating techniques for the selection of micro-partitions **210** for clustering, according to some embodiments of the present disclosure. A description of elements of FIG. 5 that have been previously described will be omitted for brevity. The data structures of FIG. 5 are for the purposes of an example, and are not intended to limit the embodiments of the present disclosure.

(72) In some embodiments, the selection of the micro-partition metadata files **214** and/or EP files **220** may be performed based on an average depth and an overlap of the micro-partition metadata files **214** and/or EP files **220**. Once selected for clustering, a given micro-partition metadata file **214** and/or EP file **220** may be recreated and/or reorganized according to embodiments described herein. For example, a selection of data for membership in a micro-partition **210** and/or a selection of micro-partition metadata files **214** and their associated micro-partitions **210** for membership in an EP file **220** during clustering may be performed in a similar matter as described herein (e.g., based on a metadata range for a data value, such as a clustering key).

(73) FIG. 5 illustrates an example of a data structure **500** having a plurality of EP files **220**, each having micro-partition metadata **214**, as described herein. In some embodiments, each of the micro-partition metadata files **214** may respectively represent a micro-partition **210**. However, for purposes of description, only the micro-partition metadata files **214** and the EP files **220** are illustrated in FIG. 5. For example, the data structure **500** may include a first EP file **220A**, a second EP file **220B**, a third EP file **220C**, and a fourth EP file **220D**. The first EP file **220A** may include a first micro-partition metadata file **214A**, a second micro-partition metadata file **214B**, a third micro-partition metadata file **214C**, and a fourth micro-partition metadata file **214D**. The second EP

file **220B** may include a fifth micro-partition metadata file **214E** and a sixth micro-partition metadata file **214F**. The third EP file **220C** may include a seventh micro-partition metadata file **214G**, an eighth micro-partition metadata file **214H**, a ninth micro-partition metadata file **214I**, and a tenth micro-partition metadata file **214J**. The fourth EP file **220D** may include an eleventh micro-partition metadata file **214K**, a twelfth micro-partition metadata file **214L**, a thirteenth micro-partition metadata file **214M**, and a fourteenth micro-partition metadata file **214N**.

(74) In FIG. 5, each of the micro-partition metadata files **214** is illustrated along a representative data range **510**. The length of a line associated with each of the micro-partition metadata files **214** is intended to illustrate a given range (e.g., a minimum data value to a maximum data value) for the data values of the clustering key represented by a particular micro-partition metadata file **214** (e.g., the data values of the clustering key stored in the micro-partitions **210** associated with the particular micro-partition metadata file **214**). For example, when two of the micro-partition metadata files **214** are illustrated as being overlapping, that is intended to convey that the data ranges (e.g., of the clustering key) covered by the micro-partitions **210** represented by the micro-partition metadata files **214** may overlap. In FIG. 5, the illustrations are shown with respect to the range of the micro-partition metadata files **214**, but it will be understood that the micro-partition metadata files **214** also represent micro-partitions **210**.

(75) In some embodiments, whether to perform clustering a given EP file **220** and/or micro-partition metadata file **214** may be made based on a calculation of an average depth of the micro-partition metadata files **214** within an EP file **220** and/or an overlap of the EP files **220**. However, the embodiments of the present disclosure are not limited to such a configuration.

(76) A depth of a first micro-partition metadata file **214** and/or micro-partition **210** may be defined to be, at any given data point, the maximum number of other micro-partition metadata files **214** within the EP file **220** covering that same data. For a given EP file **220**, a depth may be calculated for each micro-partition metadata file **214** (and its associated micro-partition **210**) of the EP file **220**, and an average depth for the EP file **220** may be calculated based on the calculated depth for each of the member micro-partitions **210**. An overlap of a first EP file **220** may be defined to be the number of other EP file **220** whose data ranges overlap the first EP file **220**.

(77) Referring to FIG. 5, within the first EP file **220A**, the first micro-partition metadata file **214A** may have a depth of 1 (e.g., itself). The second micro-partition metadata file **214B** may have a depth of 2. The third micro-partition metadata file **214C** may have a depth of 2. The fourth micro-partition metadata file **214D** may have a depth of 2. The first EP file **220A** may have an average depth of 1.75.

(78) Within the second EP file **220B**, both the fifth micro-partition metadata file **214E** and the sixth micro-partition metadata file **214F** may have a depth of 1. The second EP file **220B** may have an average depth of 1.

(79) Within the third EP file **220C**, each of the seventh micro-partition metadata file **214G**, the eighth micro-partition metadata file **214H**, the ninth micro-partition metadata file **214I**, and the tenth micro-partition metadata file **214J** may have a depth of 3. The third EP file **220C** may have an average depth of 3.

(80) Within the fourth EP file **220D**, the eleventh micro-partition metadata file **214K** may have a depth of 2. The twelfth micro-partition metadata file **214L** may have a depth of 3. The thirteenth micro-partition metadata file **214M** may have a depth of 3. The fourteenth micro-partition metadata file **214N** may have a depth of 3. The fourth EP file **220D** may have an average depth of 2.75.

(81) In addition to being calculated at a level of the micro-partition metadata file **214**, the depth and the overlap may also be calculated at the level of the EP file **220** as well. In some embodiments, the range of the EP files **220** may extend from the minimum value of all of the data ranges of the member micro-partition metadata files **214** of the EP file **220** to the maximum value of all of the data ranges of the member micro-partition metadata files **214** of the EP file **220**. Thus, an EP-file-level overlap and an EP-file-level depth may be calculated as well. For example, the EP-file-level

overlap for a particular EP file **220** may be calculated as the number of other EP files **220** containing member micro-partition metadata files **214** overlapping the particular EP file **22**. In addition, the EP-file-level depth for a particular EP file **220** may be calculated as the maximum number of other EP files **220** covering a same data point within the particular EP file **220**.

(82) Referring to FIG. 5, an EP-file-level overlap value of the first EP file **220A** may be calculated as 3 and the EP-file-level depth value of the first EP file **220A** may be calculated as 3. An EP-file-level overlap of the second EP file **220B** may be calculated as 2 and the EP-file-level depth of the second EP file **220B** may be calculated as 3. An EP-file-level overlap of the third EP file **220C** may be calculated as 2 and the EP-file-level depth of the third EP file **220C** may be calculated as 3. An EP-file-level overlap of the fourth EP file **220D** may be calculated as 3 and the EP-file-level depth of the fourth EP file **220D** may be calculated as 3.

(83) In some embodiments, the average depth of the micro-partition metadata files **214** and the depth and overlap of the EP files **220** may be utilized to select which of the EP files **220** and or micro-partition metadata files **214** (and their associated micro-partitions **210**) are to be selected for clustering (e.g., based on a clustering key). As will be described further herein, an average depth of the various micro-partition metadata files **214** of an EP file **220** may be utilized to select a particular target sub-ranges of the data range **510**, and EP files **220**, and their associated micro-partitions **210**, that overlap these ranges may be selected for clustering.

(84) FIG. 6 is a flow diagram of one embodiment of a method **600** for selecting micro-partitions **210** for clustering, according to some embodiments of the present disclosure. In general, the method **600** may be performed by processing logic that may include hardware (e.g., processing device, circuitry, dedicated logic, programmable logic, microcode, hardware of a device, integrated circuit, etc.), software (e.g., instructions run or executed on a processing device), or a combination thereof.

(85) With reference to FIG. 6, method **600** illustrates example functions used by various embodiments. Although specific function blocks (“blocks”) are disclosed in method **600**, such blocks are examples. That is, embodiments are well suited to performing various other blocks or variations of the blocks recited in method **600**. It is appreciated that the blocks in method **600** may be performed in an order different than presented, and that not all of the blocks in method **600** may be performed.

(86) The operations of FIG. 6 are described with respect to the data structure of FIG. 7. FIG. 7 is a diagram of a data structure **700** illustrating an example scenario of EP files **220** for the selection of micro-partitions **210** for clustering, according to some embodiments of the present disclosure. The data structure **700** of FIG. 7 illustrates a configuration of EP files **220** whose configuration will be used to provide examples for the operations provided herein. The configuration of FIG. 7 is merely an example, and is not intended to limit the embodiments of the present disclosure. A description of elements of FIG. 7 that have been previously described will be omitted for brevity.

(87) In FIG. 7, a plurality of EP files **220** are illustrated, each having an EP data region **720**. The EP files **220** are examples of the EP files **220** similar to those described herein. Thus, each of the plurality of EP files **220** may contain a plurality of micro-partition metadata **214**, each corresponding to a micro-partition **210**, as described herein. For ease of description, the micro-partition metadata **214** and the micro-partitions **210** are not illustrated in FIG. 7. The EP data region **720** may encompass a range of data values, within the data range **710**, between a minimum data value (e.g., of a clustering key) of the micro-partitions **210** of the EP file **220** and a maximum data value (e.g., of the clustering key) of the micro-partitions **210** of the EP file **220**.

(88) For example, the data structure **700** may include a first EP file **220A** having a first EP data region **720A**, a second EP file **220B** having a second EP data region **720B**, a third EP file **220C** having a third EP data region **720C**, a fourth EP file **220D** having a fourth EP data region **720D**, a fifth EP file **220E** having a fifth EP data region **720E**, a sixth EP file **220F** having a sixth EP data region **720F**, a seventh EP file **220G** having a seventh EP data region **720G**, and an eighth EP file

220H having an eighth EP data region 720H. In FIG. 7, each of the EP files 220 is illustrated along a representative data range 710. The length of a line associated with each of the EP files 220 is intended to illustrate a given range (e.g., a minimum data value to a maximum data value) of the EP data region 720 for the data values of the clustering key represented by the micro-partitions 210 of the EP file 220 (e.g., the data values of the clustering key stored in the micro-partitions 210 associated with the particular EP file 220). For example, when two of the EP files 220 are illustrated as having EP data regions 720 that overlap, that is intended to convey that the data ranges (e.g., of the clustering key) covered by the EP files 220 may overlap.

(89) The EP files 220 of FIG. 7 are illustrated in a different configuration than those of FIG. 5 and, as such, will have different metadata. For purposes of the following description, the data structure 700 is presumed to have the following configuration:

(90) TABLE-US-00001 TABLE 1 220A/ 220B/ 220C/ 220D/ 220E/ 220F/ 220G/ 220H/ 720A
720B 720C 720D 720E 720F 720G 720H Average Depth 3 7 2 1 3 3 2 2 Number of 100 400 150
50 200 100 500 100 Micro-partitions

(91) In the table above, values for the average depth and number of micro-partitions 210 are shown for each EP file 220. The average depth for a micro-partition 210 may be calculated as described herein with respect to FIG. 5. The number of micro-partitions 210 may represent the number of micro-partitions 210 (and associated micro-partition metadata files 214) that are grouped within the particular EP file 220.

(92) Referring to FIGS. 6 and 7, method 600 may begin at operation 610, where a plurality of sub-ranges 730 may be determined based on the respective EP data regions 720 of a plurality of EP files 220. To determine the sub-ranges 730, a minimum data value 740 and a maximum data value 745 may be selected for each of the EP data regions 720 of each of the plurality of EP files 220. The minimum data value 740 may represent a minimum value of the data of the micro-partitions 210 of the EP file 220 (e.g., the left side of the EP data region 720). The maximum data value 745 may represent a maximum value of the data of the micro-partitions 210 of the EP file 220 (e.g., the right side of the EP data region 720).

(93) The collections of minimums 740 and maximums 745 may be utilized to delineate sub-ranges 730 within the overall data range 710. For example, a sub-range 730 may be defined between a first minimum value 740 or maximum value 745 and the nearest adjacent minimum value 740 or maximum value 745. Thus, the sub-ranges 730 are the intervals of the data range 710 between adjacent minimum and maximum values 740, 745 of the EP data regions 720. In the example of FIG. 7, fifteen sub-ranges 730 are illustrated. For example, a first sub-range R1, a second sub-range R2, a third sub-range R3, a fourth sub-range R4, a fifth sub-range R5, a sixth sub-range R6, a seventh sub-range R7, an eighth sub-range R8, a ninth sub-range R9, a tenth sub-range R10, an eleventh sub-range R11, a twelfth sub-range R12, a thirteenth sub-range R13, a fourteenth sub-range R14, and a fifteenth sub-range R15 are illustrated in FIG. 7.

(94) Referring back to FIG. 6, in operation 620, a seed EP data region 820 is selected based on the sub-ranges 730 determined in operation 610. FIG. 8A is a flow diagram of one embodiment of a method 800 for selecting a seed EP data region 820, according to some embodiments of the present disclosure. In general, the method 800 may be performed by processing logic that may include hardware (e.g., processing device, circuitry, dedicated logic, programmable logic, microcode, hardware of a device, integrated circuit, etc.), software (e.g., instructions run or executed on a processing device), or a combination thereof. FIG. 8B is a block diagram of the data structure 700 illustrating an example scenario of selecting the seed EP data region 820, according to some embodiments of the present disclosure.

(95) With reference to FIG. 8A, method 800 illustrates example functions used by various embodiments. Although specific function blocks ("blocks") are disclosed in method 800, such blocks are examples. That is, embodiments are well suited to performing various other blocks or variations of the blocks recited in method 800. It is appreciated that the blocks in method 800 may

be performed in an order different than presented, and that not all of the blocks in method **800** may be performed.

(96) With reference to FIGS. **8A** and **8B**, the method **800** may begin at operation **810** in which a deepest sub-range **830** of the sub-ranges **730** is selected. The deepest sub-range **830** may be selected based on a total average depth of the sub-range **730**. The total average depth for a given sub-range **730** may be a sum of the average depths of the EP data regions **720** that overlap the sub-range **730**. For example, referring to FIG. **8B**, the sixth sub-range **R6** may be selected as the deepest sub-range **830**. This selection may be based on the arrangement where the sixth sub-range **R6** is overlapped by the first EP data region **720A** which has an average depth of 3, the second EP data region **720B** which has an average depth of 7, the sixth EP data region **720F** which has an average depth of 3, and the eighth EP data region **720H** which has an average depth of 2, for a total average depth for the overlapping EP data regions of 15. This total average depth is higher than the total average depth for any of the other sub-ranges **730**. As a result, the sixth sub-range **R6** will be selected as the deepest sub-range **830**.

(97) Referring to FIG. **8A**, in operation **812**, a candidate EP data region **720** may be selected from the EP data regions **720** that overlap the deepest sub-range **830**. The EP data regions **720** that overlap the deepest sub-range **830** may also be referred to herein as deep EP data regions. Referring to the example of FIG. **8B**, four EP data regions **720** overlap the deepest sub-range **830**: the first EP data region **720A**, the second EP data region **720B**, the sixth EP data region **720F**, and the eighth EP data region **720H**. As an example, the first EP data region **720A** may be selected as the candidate EP data region **720**.

(98) Referring to FIG. **8A**, in operation **814**, the deep EP data regions **720** that overlap the candidate EP data region **720** and the deepest sub-range **830** may be determined, and a total number of micro-partitions **210** represented by those overlapping EP data regions **720** may be calculated. For example, referring to the example of FIG. **8B**, if the first EP data region **720A** is selected as the candidate EP data region **720**, it is overlapped by seven other EP data regions **720** (**720B**, **720C**, **720D**, **720E**, **720F**, **720G**, and **720H**), for a total number of represented micro-partitions **210** of **1600** (per Table 1, above).

(99) In operation **816**, it may be determined if all of the EP data regions **720** that overlap the deepest sub-range **830** determined in operation **810** have been processed. If not, operation **812** and **814** may be repeated for each of the EP data regions **720** that overlap the deepest sub-range **830**. In the example illustrated in FIG. **8B**, these operations would further result in analysis of the second EP data region **720B** (overlapped by **720A**, **720E**, **720F**, and **720H**, for 500 total micro-partitions **210**), the sixth EP data region **720F** (overlapped by **720A**, **720B**, **720C**, and **720H**, for 750 total micro-partitions **210**), and the eighth EP data region **720H** (overlapped by **720A**, **720B**, **720C**, and **720F**, for 750 total micro-partitions **210**).

(100) After the total number of micro-partitions **210** of all of the EP data regions **720** that overlap the deepest sub-range **830** have been processed, the seed EP data region **820** may be selected as the EP data region **720** that overlaps the deepest sub-range **830** (e.g., a deep EP data region) and also overlaps the highest number of micro-partitions **210** of the EP data regions **720**. In the example illustrated in FIG. **8B**, the seed EP data region **820** is selected as the first EP data region **720A** based on its overlap of the deepest sub-range **R6** and the number of micro-partitions **210** represented by the EP data regions **720** (e.g., **1600**) that overlap the first EP data region **720A**.

(101) Referring back to FIG. **6**, in operation **630**, a plurality of clustering EP data regions may be selected based on the seed EP data region **820** determined in operation **620**. FIG. **9A** is a flow diagram of one embodiment of a method **900** for selecting clustering EP data regions **920**, according to some embodiments of the present disclosure. In general, the method **900** may be performed by processing logic that may include hardware (e.g., processing device, circuitry, dedicated logic, programmable logic, microcode, hardware of a device, integrated circuit, etc.), software (e.g., instructions run or executed on a processing device), or a combination thereof. FIG.

9B is a block diagram of the data structure 700 illustrating an example scenario of selecting the clustering EP data regions 920, according to some embodiments of the present disclosure.

(102) With reference to FIG. 9A, method 900 illustrates example functions used by various embodiments. Although specific function blocks (“blocks”) are disclosed in method 900, such blocks are examples. That is, embodiments are well suited to performing various other blocks or variations of the blocks recited in method 900. It is appreciated that the blocks in method 900 may be performed in an order different than presented, and that not all of the blocks in method 900 may be performed.

(103) With reference to FIGS. 9A and 9B, the method 900 may begin at operation 950 in which deep EP data regions 920 are selected as EP data regions 720 that overlap the deepest sub-range 830. These deep EP data regions 920 may be the same as the overlapping EP data regions 720 described with respect to operation 812 of FIG. 8A. In the example illustrated in FIG. 9B, the deep EP data regions 920 may be identified as the first EP data region 720A, 920A, which is also the seed EP data region 820, the second EP data region 720B, 920B, the sixth EP data region 720F, 920F, and the eighth EP data region 720H, 920H.

(104) Referring back to FIG. 9A, the number of micro-partitions 210 represented by the deep EP data regions 920 may be counted. For example, in the configuration of FIG. 9B, the number of micro-partitions 210 represented by the deep EP data regions 920 is 700 (per Table 1). At operation 952, it may be determined if the number of micro-partitions 210 represented by the deep EP data regions 920 is greater than a defined threshold (e.g., one thousand micro-partitions 210).

(105) If the number of micro-partitions 210 represented by the deep EP data regions 920 is not greater than the threshold (operation 952:N), the method 600 continues to expand the number of EP data regions 720 that are used for the clustering operation. Those operations will be described with respect to FIGS. 10A and 10B herein.

(106) If the number of micro-partitions 210 represented by the deep EP data regions 920 is greater than the threshold (operation 952:Y), enough micro-partitions 210 may be available for clustering. In such an event, the method 900 may continue to operation 954 in which the deep EP data regions 920 may be sorted by their average local depth. In the example of FIG. 9B, the deep EP data regions 920 include the first deep EP data region 920A with an average depth of 3, the second deep EP data region 920B with an average depth of 7, the sixth deep EP data region 920F with an average depth of 3, and the eighth deep EP data region 920H with an average depth of 2. Sorting these deep EP data regions 920 may result in an ordering of 920H, 920A/920B, 920F.

(107) At operation 956, the deep EP data region 920 having the smallest average local depth (the eighth deep EP data region 920H in the example of FIG. 9B) may be removed from the deep EP data regions 920, and the number of micro-partitions 210 remaining in the deep EP data regions 920 may be examined to determine if the number of micro-partitions 210 represented by the remaining deep EP data regions 920 is less than the threshold. If not (operation 958:N) the method 900 may revert to operation 956 where another of the deep EP data regions 920 is removed. The process may continue until the number of micro-partitions 210 represented by the remaining deep EP data regions 920 is less than the threshold (operation 958:N).

(108) For example, referring to the example of FIG. 9B, if it were to be assumed that the threshold was 650, the number of micro-partitions 210 represented by the deep EP data regions 920 calculated for operation 952 is 700. Following the operations of operation 954, the eighth deep EP data region 920H would have the smallest average local depth, and would be removed from the deep EP data regions 920 in operation 956. This would leave the first deep EP data region 920A, the second deep EP data region 920B, and the sixth deep EP data region 920F with a total number of 600 represented micro-partitions 210, which would be less than the threshold at operation 958.

(109) In operation 960, once the deep EP data regions 920 are adjusted to be below the threshold, the deep EP data regions 920 may be returned as the clustering EP data regions 720 for clustering (e.g., as a result of operation 630 of FIG. 6).

(110) As previously noted, if the number of micro-partitions **210** of the deep EP data regions **920** determined in operation **952** are not greater than the threshold, additional EP data regions **920** may be identified for inclusion. FIG. **10A** is a flow diagram of one embodiment of a method **1000** for selecting clustering EP data regions **1020**, according to some embodiments of the present disclosure. In general, the method **1000** may be performed by processing logic that may include hardware (e.g., processing device, circuitry, dedicated logic, programmable logic, microcode, hardware of a device, integrated circuit, etc.), software (e.g., instructions run or executed on a processing device), or a combination thereof. FIG. **10B** is a diagram of the data structure **700** illustrating an example scenario of selecting the clustering EP data regions **1020**, according to some embodiments of the present disclosure.

(111) With reference to FIG. **10A**, method **1000** illustrates example functions used by various embodiments. Although specific function blocks (“blocks”) are disclosed in method **1000**, such blocks are examples. That is, embodiments are well suited to performing various other blocks or variations of the blocks recited in method **1000**. It is appreciated that the blocks in method **1000** may be performed in an order different than presented, and that not all of the blocks in method **1000** may be performed.

(112) With reference to FIGS. **10A** and **10B**, method **1000** may begin operation (after the operations described herein with respect to FIG. **9A**) at operation **1050** in which a plurality of target EP data regions **1020** are selected. The target EP data regions **1020** may be selected as those EP data regions **720** that overlap the seed EP data region **820** (described herein with respect to FIGS. **8A** and **8B**) but not the deepest sub-range **830**. In the example of FIG. **10B**, the target EP data regions **1020** include the third EP data region **720C**, **1020C**, the fourth EP data region **720D**, **1020D**, the fifth EP data region **720E**, **1020E**, and the seventh EP data region **720G**, **1020G**.

(113) In operation **1052**, the number of micro-partitions **210** represented by the deep EP data regions **920** (see FIGS. **9A** and **9B**) and the target EP data regions **1020** may be compared to the threshold value. In the example of FIG. **10B**, the number of micro-partitions **210** represented by the deep EP data regions **920** (see FIGS. **9A** and **9B**) and the target EP data regions **1020** is 1600 (per Table 1). If the number of micro-partitions **210** represented by the deep EP data regions **920** (see FIGS. **9A** and **9B**) and the target EP data regions **1020** is not greater than the threshold (operation **1052:N**), the method **1000** may continue to operation **1080** in which the deep EP data regions **920** and the target EP data regions **1020** may be returned as the clustering EP data regions **720** for clustering (e.g., as a result of operation **630** of FIG. **6**).

(114) If the number of micro-partitions **210** represented by the deep EP data regions **920** (see FIGS. **9A** and **9B**) and the target EP data regions **1020** is greater than the threshold (operation **1052:Y**), the method **1000** may continue to operation **1054** to adjust the members of the target EP data regions **1020**. In operation **1054**, a lower target EP data region **1020_L** may be determined. The lower target EP data region **1020_L** may be determined by identifying a lowest maximum data value **745_L** (e.g., of the data range **710** of the clustering key) of the maximum data values **745** of the target EP data regions **1020** that overlap with the seed EP data region **820**. The lower target EP data region **1020_L** may be determined as the target EP data region **1020** that does not overlap with the lowest maximum data value **745_L**. In the example of FIG. **10B**, the lower target EP data region **1020_L** is the fifth EP data region **720E**. Once the lower target EP data region **1020_L** is determined, the number of micro-partitions **210** represented by the members of the deep EP data regions **920** (see FIGS. **9A** and **9B**) that overlap the lower target EP data region **1020_L** may be calculated. In the example of FIG. **10B**, the first deep EP data region **920A** and the second deep EP data region **920B** overlap the fifth EP data region **720E**, so the number of micro-partitions **210** represented by the members of the deep EP data regions **920** that overlap the fifth EP data region **720E** is 500 (per Table 1). This number will be designated as L.sub.D.

(115) In operation **1056**, an upper target EP data region **1020_U** may be determined. The upper target EP data region **1020_U** may be determined by identifying a highest minimum data value

740_U (e.g., of the data range **710** of the clustering key) of the minimum data values **740** of the target EP data regions **1020** that overlap with the seed EP data region **820**. The upper target EP data region **1020_U** may be determined as the target EP data region **1020** that does not overlap with the highest minimum data value **740_U**. In the example of FIG. **10B**, the upper target EP data region **1020_U** is the seventh EP data region **720G**. Once the upper target EP data region **1020_U** is determined, the number of micro-partitions **210** represented by the members of the deep EP data regions **920** that overlap the upper target EP data region **1020_U** may be calculated. In the example of FIG. **10B**, only the first deep EP data region **920A** overlaps the seventh EP data region **720G**, so the number of micro-partitions **210** represented by the members of the deep EP data regions **920** that overlap the seventh EP data region **720G** is 100 (per Table 1). This number will be designated as U.sub.D.

(116) In operation **1058**, U.sub.D is compared to L.sub.D. If U.sub.D is not equal to L.sub.D (operation **1058:N**), the EP data region **720** that is associated with the lower of U.sub.D or L.sub.D is removed from the target EP data regions **1020** in operation **1070**. In the example of FIG. **10B**, the seventh EP data region **720G** would be removed from the target EP data regions **1020**. After the target EP data regions **1020** are adjusted in operation **1070**, the method **1000** may revert back to operation **1052** to compare the number of micro-partitions **210** represented by the deep EP data regions **920** and the target EP data regions **1020** to the threshold value.

(117) If U.sub.D is equal to L.sub.D (operation **1058:Y**), the method **1000** may continue to operation **1060**. In operation **1060**, the number of micro-partitions **210** represented by the EP data regions **720** that overlap the lower target EP data region **1020_L** and the seed EP data region **820** may be calculated. In the example of FIG. **10B**, the second EP data region **720B** overlaps the lower target EP data region **1020_L** (the fifth EP data region **720E**), so the number of micro-partitions **210** represented by the members of the EP data regions **720** that overlap the lower target EP data region **1020_L** and the seed EP data region **820** is 400 (per Table 1). This number will be designated as L.sub.C.

(118) In operation **1062**, the number of micro-partitions **210** represented by the EP data regions **720** that overlap the upper target EP data region **1020_U** and the seed EP data region **820** may be calculated. In the example of FIG. **10B**, the fourth EP data region **720D** overlaps the upper target EP data region **1020_U** (the seventh EP data region **720G**), so the number of micro-partitions **210** represented by the members of the EP data regions **720** that overlap the upper target EP data region **1020_U** and the seed EP data region **820** is 50 (per Table 1). This number will be designated as U.sub.C.

(119) In operation **1064**, the EP data region **720** that is associated with the lower of L.sub.C or U.sub.C is removed from the target EP data regions **1020** in operation **1064**. In the example of FIG. **10B**, had L.sub.D been equal to U.sub.D, the seventh EP data region **720G** would have been removed from the target EP data regions **1020**. After the target EP data regions **1020** are adjusted in operation **1064**, the method **1000** may revert back to operation **1052** to compare the number of micro-partitions **210** represented by the deep EP data regions **920** and the target EP data regions **1020** to the threshold value. The operations of the method **1000** may continue until the number of micro-partitions **210** represented by the deep EP data regions **920** and the target EP data regions is less than the threshold value and operation **1080** is executed, in which the remaining target EP data regions **1020** and the deep EP data regions **920** are returned as the clustering EP data regions **720** for clustering in the method **600** of FIG. **6**.

(120) Referring back to FIG. **6**, as a result of the methods **900**, **1000** described with respect to FIGS. **7** to **10B**, a plurality of clustering EP data regions **720** may be identified. In some embodiments, the method **600** may continue with operation **640**, in which additional EP data regions **720** may be added to the clustering EP data regions **720**. In some embodiments, additional EP data regions **720** may be identified by performing the operations of FIGS. **7** to **10B** again utilizing a higher threshold and/or a different range. For example, the embodiments of FIGS. **7** to

10B may operate on the range of the seed EP data region **820** identified in method **800** of FIG. **8A**. In some embodiments, the range may be increased beyond that of the seed EP data region **820**, and the operations of FIGS. **7** to **10B** may be repeated.

(121) In some embodiments, rather than adjusting the range and/or threshold, additional clustering EP data regions **720** by repeating the operations of FIGS. **7** to **10B** after removing the EP data regions **720** identified in operation **630** from the group of EP data regions **720** being considered. The clustering EP data region selection algorithm may work on the adjusted batch of EP data regions **720** without touching the EP data regions **720** selected in the previous iteration.

(122) This solution may have two advantages. First, each iteration may pick the EP data regions **720** overlapping with the deepest data sub-range **830**. As a result, this solution can reduce the depth of multiple ranges in one micro-partition clustering. Second, this solution is more efficient because one EP data region **720** is only processed once in the EP data regions **720** selection phase. In each iteration, the EP data regions **720** selection phase processes a limited number of EP data regions **720** such that it will not utilize large amounts of memory.

(123) In operation **650**, a clustering operation may be performed on the micro-partitions **210** that are represented by the clustering EP data regions **720** selected in operations **630** and/or **640**. The clustering operation may examine the data ranges of the micro-partitions **210** represented by the clustering EP data regions **720** that were selected to determine if the ranges of some of the micro-partitions **210** may be reduced and/or micro-partitions can be combined. For example, data of the micro-partitions may be moved and/or recombined between micro-partitions **210** so there is less overlap between micro-partitions **210** and/or EP data regions **720**. In some embodiments, the clustering may be performed based on the data ranges of the clustering key. In some embodiments, the micro-partitions **210** are stored in immutable storage, so moving and/or recombining the data between micro-partitions **210** may include generating new micro-partitions **210** that contain different combinations of the data.

(124) Clustering the micro-partitions **210** so that they have less overlap may allow for more efficient data query pruning, as described herein. Moreover, by utilizing the mechanisms described herein to select the micro-partitions **210**, only a subset of all of the EP data regions **720** may be analyzed during the clustering operation. By reducing the number of EP data regions **720** that are analyzed for the clustering operation, less memory may be used and the clustering operation may be performed more quickly. In addition, embodiments of the present disclosure attempt to select the micro-partitions **210** that will have the highest impact on clustering by focusing the selection of the EP data regions **720** for clustering based on the deepest data sub-range **830**. The deepest data sub-range **830** may include the micro-partitions **210** having a large amount of overlap.

(125) The operations of the method **600** may be repeated during the operation of the database. For example, once a clustering operation has been performed, the ranges of the EP files **220** and their associated micro-partitions **210** will have changed. As a result, a repeat of the operations of the method **600** may select different EP files **220** with different respective EP data regions **720** for clustering on a subsequent iteration. Because the method **600** selects only a subset of the EP files **220** for clustering in each iteration, it may utilize less memory and/or processing time while continually improving the organization of the data in the database.

(126) FIG. **11** is a flow diagram of one embodiment of a method **1100** for selecting micro-partitions **210** for a clustering operation, according to some embodiments of the present disclosure. In general, the method **1100** may be performed by processing logic that may include hardware (e.g., processing device, circuitry, dedicated logic, programmable logic, microcode, hardware of a device, integrated circuit, etc.), software (e.g., instructions run or executed on a processing device), or a combination thereof.

(127) With reference to FIG. **11**, method **1100** illustrates example functions used by various embodiments. Although specific function blocks (“blocks”) are disclosed in method **1100**, such blocks are examples. That is, embodiments are well suited to performing various other blocks or

variations of the blocks recited in method **1100**. It is appreciated that the blocks in method **1100** may be performed in an order different than presented, and that not all of the blocks in method **1100** may be performed.

(128) Method **1100** begins at operation **1110**, which includes storing table data in a plurality of micro-partitions of a storage device. Each of the plurality of micro-partitions may include a portion of the table data. Subsets of the plurality of micro-partitions may be associated with a respective one of a plurality of EP files. Each of the plurality of EP files may include an EP data region that represents the portions of the table data of the subset of the plurality of micro-partitions associated with the EP file. The micro-partitions may be similar to the micro-partitions **210** described herein with respect to FIGS. **1** to **10B**. The EP files may be similar to the EP files **220** described herein with respect to FIGS. **2** to **10B**. The EP data regions may be similar to the EP data regions **720** described herein with respect to FIGS. **7** to **10B**. The EP data regions may be similar to the EP data regions **720** described herein with respect to FIGS. **7** to **10B**.

(129) In operation **1120**, the method **1100** may include determining sub-ranges of the table data based on the plurality of EP data regions of the plurality of EP files. The sub-ranges may be similar to the sub-ranges **730** described herein with respect to FIGS. **7** to **10B**.

(130) In operation **1130**, the method **1100** may include selecting (e.g., by a processing device) a subset of the plurality of EP files for a clustering operation based on the sub-ranges of the table data. The subset of the plurality of EP files may be similar to the EP files **220** associated with the clustering EP data regions **720** selected for clustering described herein with respect to FIGS. **6** to **10B**.

(131) In some embodiments, selecting the subset of the plurality of EP files for the clustering operation based on the sub-ranges of the table data includes selecting EP files of the plurality of EP files that comprise table data having values that overlap a sub-range that is associated with highest average micro-partition depth, as described herein with respect to FIGS. **1** to **10B**.

(132) In some embodiments, selecting the subset of the plurality of EP files for the clustering operation based on the sub-ranges of the table data may include, for each sub-range of the sub-ranges, calculating a total average micro-partition depth of ones of the plurality of EP data regions comprising table data that overlaps the sub-range, selecting the sub-range having a highest total average micro-partition depth as a deepest sub-range of the sub-ranges, and selecting the subset of the plurality of EP files for the clustering operation based on the deepest sub-range. The deepest sub-range may be similar to the deepest sub-range **830** described herein with respect to FIGS. **8** to **10B**.

(133) In some embodiments, selecting the subset of the plurality of EP files for the clustering operation based on the deepest sub-range includes determining a group of overlapping EP data regions that comprise table data that overlap the deepest sub-range, selectively removing at least one of the overlapping EP data regions from the group of the overlapping EP data regions until a sum of the micro-partitions associated with the group of the overlapping EP data regions is less than a threshold value, and returning the EP files associated with the group of overlapping EP data regions as the subset of the EP files. The selective removal of the overlapping EP data regions is described herein, for example, with respect to FIGS. **9A** and **9B**.

(134) In some embodiments, selecting the subset of the plurality of EP files for the clustering operation based on the deepest sub-range includes determining deep EP data regions of the plurality of EP data regions that overlap the deepest sub-range, selecting a seed EP data region from the deep EP data regions, and selecting the EP files associated with the EP data regions of the plurality of EP data regions that overlap the seed EP data region as the subset of the plurality of EP files for the clustering operation. The determination of the deep EP data regions and the selection of the seed EP data region is described herein, for example, with respect to FIGS. **8A** to **10B**.

(135) In some embodiments, selecting the seed EP data region includes, for each deep EP data region of the deep EP data regions, determining a set of the plurality of EP data regions comprising

table data that overlaps with table data of the deep EP data region, calculating a sum of micro-partitions associated with each of the sets of the EP data regions, and selecting the seed EP data region as the deep EP data region associated with the set of EP data regions having a highest sum of micro-partitions. The determination of the seed EP data region is described herein, for example, with respect to FIGS. **8A** and **8B**.

(136) In some embodiments, selecting the EP files associated with the EP data regions of the plurality of EP data regions that overlap the seed EP data region as the subset of the plurality of EP files for the clustering operation includes determining a group of target EP data regions that comprise table data that overlaps the table data of the seed EP data region but not the deepest sub-range, selectively removing at least one of the target EP data regions from the group of the target EP data regions based on a number of micro-partitions associated with the at least one target EP data region to generate adjusted target EP data regions, and selecting the EP files associated with the adjusted target EP data regions and the deep EP data regions as the subset of the plurality of EP files for the clustering operation. The selective removal of the target EP data regions is described herein, for example, with respect to FIGS. **10A** and **10B**.

(137) In operation **1140**, the method **1100** may include performing the clustering operation on the micro-partitions associated with the subset of the plurality of EP files. The clustering operation may be similar to the clustering operation described herein, such as with respect to FIGS. **2** to **5**.

(138) FIG. **12** is a block diagram of an example computing device **1200** that may perform one or more of the operations described herein, in accordance with some embodiments. Computing device **1200** may be connected to other computing devices in a LAN, an intranet, an extranet, and/or the internet. The computing device may operate in the capacity of a server machine in a client-server network environment or in the capacity of a client in a peer-to-peer network environment. The computing device may be provided by a personal computer (PC), a set-top box (STB), a server, a network router, switch or bridge, or any machine capable of executing a set of instructions (sequential or otherwise) that specify actions to be taken by that machine. Further, while only a single computing device is illustrated, the term “computing device” shall also be taken to include any collection of computing devices that individually or jointly execute a set (or multiple sets) of instructions to perform the methods discussed herein.

(139) The example computing device **1200** may include a processing device (e.g., a general purpose processor, a PLD, etc.) **1202**, a main memory **1204** (e.g., synchronous dynamic random access memory (DRAM), read-only memory (ROM)), a static memory **1206** (e.g., flash memory) and a data storage device **1218**, which may communicate with each other via a bus **1230**.

(140) Processing device **1202** may be provided by one or more general-purpose processing devices such as a microprocessor, central processing unit, or the like. In an illustrative example, processing device **1202** may comprise a complex instruction set computing (CISC) microprocessor, reduced instruction set computing (RISC) microprocessor, very long instruction word (VLIW) microprocessor, or a processor implementing other instruction sets or processors implementing a combination of instruction sets. Processing device **1202** may also comprise one or more special-purpose processing devices such as an application specific integrated circuit (ASIC), a field programmable gate array (FPGA), a digital signal processor (DSP), network processor, or the like. The processing device **1202** may be configured to execute the operations described herein, in accordance with one or more aspects of the present disclosure, for performing the operations and steps discussed herein. In one embodiment, processing device **1202** represents a processing device of cloud computing platform **110** of FIG. **1**.

(141) Computing device **1200** may further include a network interface device **1208** which may communicate with a network **1220**. The computing device **1200** also may include a video display unit **1210** (e.g., a liquid crystal display (LCD) or a cathode ray tube (CRT)), an alphanumeric input device **1212** (e.g., a keyboard), a cursor control device **1214** (e.g., a mouse) and an acoustic signal generation device **1216** (e.g., a speaker). In one embodiment, video display unit **1210**,

alphanumeric input device **1212**, and cursor control device **1214** may be combined into a single component or device (e.g., an LCD touch screen).

(142) Data storage device **1218** may include a computer-readable (also referred to herein as machine-readable) storage medium **1228** on which may be stored one or more sets of instructions **1225**, such as instructions for executing the cloud services component **120**, e.g., instructions for carrying out the operations described herein, in accordance with one or more aspects of the present disclosure. Cloud services instructions **120** may also reside, completely or at least partially, within main memory **1204** and/or within processing device **1202** during execution thereof by computing device **1200**, the main memory **1204** and processing device **1202** also constituting computer-readable media. The instructions **1225** may further be transmitted or received over a network **1220** via network interface device **1208**.

(143) While computer-readable storage medium **1228** is shown in an illustrative example to be a single medium, the term “computer-readable storage medium” should be taken to include a single medium or multiple media (e.g., a centralized or distributed database and/or associated caches and servers) that store the one or more sets of instructions. The term “computer-readable storage medium” shall also be taken to include any medium that is capable of storing, encoding, or carrying a set of instructions for execution by the machine and that cause the machine to perform the methods described herein. The term “computer-readable storage medium” shall accordingly be taken to include, but not be limited to, solid-state memories, optical media, and magnetic media.

(144) Unless specifically stated otherwise, terms such as “storing,” “determining,” “selecting,” “performing,” “calculating,” “removing,” or the like, refer to actions and processes performed or implemented by computing devices that manipulates and transforms data represented as physical (electronic) quantities within the computing device's registers and memories into other data similarly represented as physical quantities within the computing device memories or registers or other such information storage, transmission or display devices. Also, the terms “first,” “second,” “third,” “fourth,” etc., as used herein are meant as labels to distinguish among different elements and may not necessarily have an ordinal meaning according to their numerical designation.

(145) Examples described herein also relate to an apparatus for performing the operations described herein. This apparatus may be specially constructed for the required purposes, or it may comprise a general purpose computing device selectively programmed by a computer program stored in the computing device. Such a computer program may be stored in a computer-readable non-transitory storage medium.

(146) The methods and illustrative examples described herein are not inherently related to any particular computer or other apparatus. Various general purpose systems may be used in accordance with the teachings described herein, or it may prove convenient to construct more specialized apparatus to perform the required method steps. The required structure for a variety of these systems will appear as set forth in the description above.

(147) The above description is intended to be illustrative, and not restrictive. Although the present disclosure has been described with references to specific illustrative examples, it will be recognized that the present disclosure is not limited to the examples described. The scope of the disclosure should be determined with reference to the following claims, along with the full scope of equivalents to which the claims are entitled.

(148) As used herein, the singular forms “a,” “an,” and “the” are intended to include the plural forms as well, unless the context clearly indicates otherwise. It will be further understood that the terms “comprises,” “comprising,” “includes,” and/or “including,” when used herein, specify the presence of stated features, integers, steps, operations, elements, and/or components, but do not preclude the presence or addition of one or more other features, integers, steps, operations, elements, components, and/or groups thereof. Therefore, the terminology used herein is for the purpose of describing particular embodiments only and is not intended to be limiting.

(149) It should also be noted that in some alternative implementations, the functions/acts noted

may occur out of the order noted in the figures. For example, two figures shown in succession may in fact be executed substantially concurrently or may sometimes be executed in the reverse order, depending upon the functionality/acts involved.

(150) Although the method operations were described in a specific order, it should be understood that other operations may be performed in between described operations, described operations may be adjusted so that they occur at slightly different times or the described operations may be distributed in a system which allows the occurrence of the processing operations at various intervals associated with the processing.

(151) Various units, circuits, or other components may be described or claimed as “configured to” or “configurable to” perform a task or tasks. In such contexts, the phrase “configured to” or “configurable to” is used to connote structure by indicating that the units/circuits/components include structure (e.g., circuitry) that performs the task or tasks during operation. As such, the unit/circuit/component can be said to be configured to perform the task, or configurable to perform the task, even when the specified unit/circuit/component is not currently operational (e.g., is not on). The units/circuits/components used with the “configured to” or “configurable to” language include hardware—for example, circuits, memory storing program instructions executable to implement the operation, etc. Reciting that a unit/circuit/component is “configured to” perform one or more tasks, or is “configurable to” perform one or more tasks, is expressly intended not to invoke 35 U.S.C. 112, sixth paragraph, for that unit/circuit/component. Additionally, “configured to” or “configurable to” can include generic structure (e.g., generic circuitry) that is manipulated by software and/or firmware (e.g., an FPGA or a general-purpose processor executing software) to operate in manner that is capable of performing the task(s) at issue. “Configured to” may also include adapting a manufacturing process (e.g., a semiconductor fabrication facility) to fabricate devices (e.g., integrated circuits) that are adapted to implement or perform one or more tasks. “Configurable to” is expressly intended not to apply to blank media, an unprogrammed processor or unprogrammed generic computer, or an unprogrammed programmable logic device, programmable gate array, or other unprogrammed device, unless accompanied by programmed media that confers the ability to the unprogrammed device to be configured to perform the disclosed function(s).

(152) Any combination of one or more computer-usable or computer-readable media may be utilized. For example, a computer-readable medium may include one or more of a portable computer diskette, a hard disk, a random access memory (RAM) device, a read-only memory (ROM) device, an erasable programmable read-only memory (EPROM or Flash memory) device, a portable compact disc read-only memory (CDROM), an optical storage device, and a magnetic storage device. Computer program code for carrying out operations of the present disclosure may be written in any combination of one or more programming languages. Such code may be compiled from source code to computer-readable assembly language or machine code suitable for the device or computer on which the code will be executed.

(153) Embodiments may also be implemented in cloud computing environments. In this description and the following claims, “cloud computing” may be defined as a model for enabling ubiquitous, convenient, on-demand network access to a shared pool of configurable computing resources (e.g., networks, servers, storage, applications, and services) that can be rapidly provisioned (including via virtualization) and released with minimal management effort or service provider interaction and then scaled accordingly. A cloud model can be composed of various characteristics (e.g., on-demand self-service, broad network access, resource pooling, rapid elasticity, and measured service), service models (e.g., Software as a Service (“SaaS”), Platform as a Service (“PaaS”), and Infrastructure as a Service (“IaaS”)), and deployment models (e.g., private cloud, community cloud, public cloud, and hybrid cloud). The flow diagrams and block diagrams in the attached figures illustrate the architecture, functionality, and operation of possible implementations of systems, methods, and computer program products according to various embodiments of the

present disclosure. In this regard, each block in the flow diagrams or block diagrams may represent a module, segment, or portion of code, which comprises one or more executable instructions for implementing the specified logical function(s). It will also be noted that each block of the block diagrams or flow diagrams, and combinations of blocks in the block diagrams or flow diagrams, may be implemented by special purpose hardware-based systems that perform the specified functions or acts, or combinations of special purpose hardware and computer instructions. These computer program instructions may also be stored in a computer-readable medium that can direct a computer or other programmable data processing apparatus to function in a particular manner, such that the instructions stored in the computer-readable medium produce an article of manufacture including instruction means which implement the function/act specified in the flow diagram and/or block diagram block or blocks.

(154) The foregoing description, for the purpose of explanation, has been described with reference to specific embodiments. However, the illustrative discussions above are not intended to be exhaustive or to limit the invention to the precise forms disclosed. Many modifications and variations are possible in view of the above teachings. The embodiments were chosen and described in order to best explain the principles of the embodiments and its practical applications, to thereby enable others skilled in the art to best utilize the embodiments and various modifications as may be suited to the particular use contemplated. Accordingly, the present embodiments are to be considered as illustrative and not restrictive, and the invention is not to be limited to the details given herein, but may be modified within the scope and equivalents of the appended claims.

Claims

1. A method comprising: storing table data in a plurality of micro-partitions of a storage device, wherein each of the plurality of micro-partitions comprises a portion of the table data, wherein subsets of the plurality of micro-partitions are associated with a respective one of a plurality of expression property (EP) files, and wherein each of the plurality of EP files comprises an EP data region that represents the portions of the table data of the subset of the plurality of micro-partitions associated with the EP file; determining sub-ranges of the table data based on the plurality of EP data regions of the plurality of EP files, wherein each sub-range corresponds to an interval between adjacent minimum values or adjacent maximum values among the plurality of EP data regions; for each sub-range of the sub-ranges, calculating a total average micro-partition depth of ones of the plurality of EP data regions comprising table data that overlaps the sub-range; determining from the plurality of EP data regions, a set of EP data regions that each comprise table data that overlap a deepest sub-range of the sub-ranges, the deepest sub-range having a highest total average micro-partition depth; selecting a subset of the plurality of EP files for a clustering operation based at least in part on a number of micro-partitions represented by the set of EP data regions; and performing the clustering operation on the micro-partitions associated with the subset of the plurality of EP files to reduce an amount of memory required for the clustering operation.
2. The method of claim 1, wherein selecting the subset of the plurality of EP files for the clustering operation comprises: selectively removing at least one of the set of EP data regions until a sum of the micro-partitions associated with the set of EP data regions is less than a threshold value; and selecting the EP files associated with the set of EP data regions as the subset of the plurality of EP files for the clustering operation.
3. The method of claim 1, wherein selecting the subset of the plurality of EP files for the clustering operation comprises: selecting a seed EP data region from the set of EP data regions; and selecting as the subset of the plurality of EP files for the clustering operation, the EP files associated with the EP data regions of the plurality of EP data regions that overlap the seed EP data region.
4. The method of claim 3, wherein selecting the seed EP data region comprises: for each EP data region of the set of EP data regions, determining a subset of the plurality of EP data regions

comprising table data that overlaps with table data of the EP data region of the set of EP data regions; calculating a sum of micro-partitions associated with each of the subsets of the EP data regions; and selecting the seed EP data region as the EP data region of the set of EP data regions associated with the subset of EP data regions having a highest sum of micro-partitions.

5. The method of claim 3, wherein selecting the EP files associated with the EP data regions of the plurality of EP data regions that overlap the seed EP data region as the subset of the plurality of EP files for the clustering operation comprises: determining a group of target EP data regions that comprise table data that overlaps the table data of the seed EP data region but not the deepest sub-range; selectively removing at least one of the target EP data regions from the group of the target EP data regions based on a number of micro-partitions associated with the at least one target EP data region to generate adjusted target EP data regions; and selecting the EP files associated with the adjusted target EP data regions and the set of EP data regions as the subset of the plurality of EP files for the clustering operation.

6. The method of claim 1, wherein the micro-partitions associated with the subset of the plurality of EP files for the clustering operation have a large degree of overlap.

7. A system comprising: a memory; and a processing device, operatively coupled to the memory, to: store table data in a plurality of micro-partitions of a storage device, wherein each of the plurality of micro-partitions comprises a portion of the table data, wherein subsets of the plurality of micro-partitions are associated with a respective one of a plurality of expression property (EP) files, and wherein each of the plurality of EP files comprises an EP data region that represents the portions of the table data of the subset of the plurality of micro-partitions associated with the EP file; determine sub-ranges of the table data based on the plurality of EP data regions of the plurality of EP files, wherein each sub-range corresponds to an interval between adjacent minimum values or adjacent maximum values among the plurality of EP data regions; for each sub-range of the sub-ranges, calculate a total average micro-partition depth of ones of the plurality of EP data regions comprising table data that overlaps the sub-range; determine from the plurality of EP data regions, a set of EP data regions that each comprise table data that overlap a deepest sub-range of the sub-ranges, the deepest sub range having a highest total average micro-partition depth; select a subset of the plurality of EP files for a clustering operation based at least in part on a number of micro-partitions represented by the set of EP data regions; and perform the clustering operation on the micro-partitions associated with the subset of the plurality of EP files to reduce an amount of memory required for the clustering operation.

8. The system of claim 7, wherein, to select the subset of the plurality of EP files for the clustering operation, the processing device is to: selectively remove at least one of the set of EP data regions until a sum of the micro-partitions associated with the set of EP data regions is less than a threshold value; and select the EP files associated with the set of EP data regions as the subset of the plurality of EP files for the clustering operation.

9. The system of claim 7, wherein, to select the subset of the plurality of EP files for the clustering operation, the processing device is to: select a seed EP data region from the set of EP data regions; and select as the subset of the plurality of EP files for the clustering operation, the EP files associated with the EP data regions of the plurality of EP data regions that overlap the seed EP data region.

10. The system of claim 9, wherein, to select the seed EP data region, the processing device is to: for each EP data region of the set of EP data regions, determine a subset of the plurality of EP data regions comprising table data that overlaps with table data of the EP data region of the set of EP data regions; calculate a sum of micro-partitions associated with each of the subsets of the EP data regions; and select the seed EP data region as the EP data region of the set of EP data regions associated with the subset of EP data regions having a highest sum of micro-partitions.

11. The system of claim 9, wherein, to select the EP files associated with the EP data regions of the plurality of EP data regions that overlap the seed EP data region as the subset of the plurality of EP

files for the clustering operation, the processing device is to: determine a group of target EP data regions that comprise table data that overlaps the table data of the seed EP data region but not the deepest sub-range; selectively remove at least one of the target EP data regions from the group of the target EP data regions based on a number of micro-partitions associated with the at least one target EP data region to generate adjusted target EP data regions; and select the EP files associated with the adjusted target EP data regions and the set of EP data regions as the subset of the plurality of EP files for the clustering operation.

12. The system of claim 7, wherein the micro-partitions associated with the subset of the plurality of EP files for the clustering operation have a large degree of overlap.

13. A non-transitory computer-readable storage medium including instructions that, when executed by a processing device, cause the processing device to: store table data in a plurality of micro-partitions of a storage device, wherein each of the plurality of micro-partitions comprises a portion of the table data, wherein subsets of the plurality of micro-partitions are associated with a respective one of a plurality of expression property (EP) files, and wherein each of the plurality of EP files comprises an EP data region that represents the portions of the table data of the subset of the plurality of micro-partitions associated with the EP file; determine sub-ranges of the table data based on the plurality of EP data regions of the plurality of EP files, wherein each sub-range corresponds to an interval between adjacent minimum values or adjacent maximum values among the plurality of EP data regions; for each sub-range of the sub-ranges, calculate a total average micro-partition depth of ones of the plurality of EP data regions comprising table data that overlaps the sub-range; determine from the plurality of EP data regions, a set of EP data regions that each comprise table data that overlap a deepest sub-range of the sub-ranges, the deepest sub-range having a highest total average micro-partition depth; select a subset of the plurality of EP files for a clustering operation based at least in part on a number of micro-partitions represented by the set of EP data regions; and perform the clustering operation on the micro-partitions associated with the subset of the plurality of EP files to reduce an amount of memory required for the clustering operation.

14. The non-transitory computer-readable storage medium of claim 13, wherein, to select the subset of the plurality of EP files for the clustering operation, the processing device is to: selectively remove at least one of the set of EP data regions until a sum of the micro-partitions associated with the set of EP data regions is less than a threshold value; and select the EP files associated with the set of EP data regions as the subset of the plurality of EP files for the clustering operation.

15. The non-transitory computer-readable storage medium of claim 13, wherein, to select the subset of the plurality of EP files for the clustering operation, the processing device is to: select a seed EP data region from the set of EP data regions; and select as the subset of the plurality of EP files for the clustering operation, the EP files associated with the EP data regions of the plurality of EP data regions that overlap the seed EP data region.

16. The non-transitory computer-readable storage medium of claim 15, wherein, to select the seed EP data region, the processing device is to: for each EP data region of the set of EP data regions, determine a subset of the plurality of EP data regions comprising table data that overlaps with table data of the EP data region of the set of EP data regions; calculate a sum of micro-partitions associated with each of the subsets of the EP data regions; and select the seed EP data region as the EP data region of the set of EP data regions associated with the subset of EP data regions having a highest sum of micro-partitions.

17. The non-transitory computer-readable storage medium of claim 13, wherein the micro-partitions associated with the subset of the plurality of EP files for the clustering operation have a large degree of overlap.
