

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent
Kind Code
Date of Patent
Inventor(s)

12387028
B2
August 12, 2025
Roy; Rajarshi et al.

Data path circuit design using reinforcement learning

Abstract

Apparatuses, systems, and techniques for designing a data path circuit such as a parallel prefix circuit with reinforcement learning are described. A method can include receiving a first design state of a data path circuit, inputting the first design state of the data path circuit into a machine learning model, and performing reinforcement learning using the machine learning model to output a final design state of the data path circuit, wherein the final design state of the data path circuit has decreased area, power consumption and/or delay as compared to conventionally designed data path circuits.

Inventors: Roy; Rajarshi (Milpitas, CA), Godil; Saad (Milpitas, CA), Raiman; Jonathan (Seattle, WA), Kant; Neel (Santa Clara, CA), Elkin; Ilyas (Sunnyvale, CA), Siu; Ming Y. (Santa Clara, CA), Kirby; Robert (San Francisco, CA), Oberman; Stuart (Sunnyvale, CA), Catanzaro; Bryan (Los Altos Hills, CA)

Applicant: NVIDIA Corporation (Santa Clara, CA)

Family ID: 1000008749102

Assignee: NVIDIA Corporation (Santa Clara, CA)

Appl. No.: 17/517612

Filed: November 02, 2021

Prior Publication Data

Document Identifier	Publication Date
US 20230139623 A1	May. 04, 2023

Publication Classification

Int. Cl.: G06F30/394 (20200101); G06F30/327 (20200101); G06N20/00 (20190101)

U.S. Cl.:

CPC G06F30/394 (20200101); G06F30/327 (20200101); G06N20/00 (20190101);

Field of Classification Search

CPC: G06F (30/394); G06F (30/327); G06F (30/32); G06F (30/337); G06F (2115/10); G06F (30/27); G06N (20/00); G06N (3/0495); G06N (3/084); G06N (3/006); G06N (3/044); G06N (3/0464); G06N (3/092)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
8527920	12/2012	Choudhury	716/104	G06F 30/327
9792397	12/2016	Nagaraja	N/A	G06F 30/32
11599699	12/2022	Roberto	N/A	G06F 30/398
11984910	12/2023	Fan	N/A	G06N 3/044
2021/0073456	12/2020	Nath	N/A	G06F 30/3953
2021/0342516	12/2020	Ren	N/A	G06N 3/045
2022/0229960	12/2021	Nath	N/A	G06F 30/3315
2023/0139623	12/2022	Roy	716/104	G06N 3/044

OTHER PUBLICATIONS

“Approximate Logic Synthesis: A Reinforcement Learning-Based Technology Mapping Approach”, Ghasem Pasandi, Shahin Nazarian, and Massoud Pedram, Department of Electrical and Computer Engineering, Feb. 1, 2019. cited by examiner

Main reference (Year: 2019). cited by examiner

Primary Examiner: Nguyen; Nha T

Attorney, Agent or Firm: Lowenstein Sandler LLP

Background/Summary

TECHNICAL FIELD

(1) At least one embodiment pertains to use of machine learning to perform and facilitate circuit design. For example, at least one embodiment pertains to technology for data path circuit design utilizing reinforcement learning, according to various novel techniques described herein.

BACKGROUND

(2) Many types of circuits can include data paths or data path circuits—e.g., arithmetic logic units or multipliers that can perform data processing operations. For example, data path circuits can include parallel prefix circuits (e.g., gray to binary converter, adder, decrementer, incrementer, priority encoder, etc.) that process or synthesize data over data paths. Prefix computation is a very useful and basic operation that is used in various applications, such as image processing, cryptography, processor allocation, biological sequence comparison, binary addition, design of silicon compilers, job scheduling, loop parallelization, polynomial evaluation, and sorting.

(3) Design of data path circuits look to decrease a delay (e.g., an amount of time the data path circuit takes to output a value given an input) and an area (e.g., an amount of space the data path circuit takes) while avoiding increases in power consumption of the circuit. As delay of a data path circuit is decreased, the area and power consumption of the data path circuit can be affected—e.g., as delay of a data path circuit decreases, the area of the data path circuit can increase. Accordingly, a design of data path circuits looks to optimize the delay, area, and power consumption of the data path circuit.

(4) The optimization of prefix circuits is challenging as their large design space grows exponentially with input length and is intractable to enumerate. As a result, exhaustive search approaches do not scale beyond small input lengths. Several regular prefix circuit structures have been proposed that trade off logic level, maximum fanout and wiring tracks. Another set of algorithms optimize prefix circuit size and level properties. However, prefix circuit level and maximum fanout properties do not map to circuit area, power and delay due to physical design complexities such as capacitive loading and congestion. Conventional methods for data path circuit design do not fully optimize the delay, area, and power consumption of the data path circuit.

Description

BRIEF DESCRIPTION OF DRAWINGS

- (1) Various embodiments in accordance with the present disclosure will be described with reference to the drawings, in which:
- (2) FIG. 1 is an example system architecture, in accordance with at least some embodiments;
- (3) FIG. 2 illustrates an example system architecture for reinforcement learning, in accordance with at least some embodiments;
- (4) FIG. 3 illustrates an example prefix graph modification, in accordance with at least some embodiments;
- (5) FIG. 4 illustrates example grid representations of parallel prefix graphs, in accordance with at least some embodiments;
- (6) FIG. 5 illustrates an example data path circuit design calculation, in accordance with at least some embodiments;
- (7) FIG. 6 illustrates an example constraint for data path circuit design, in accordance with at least some embodiments;
- (8) FIG. 7 illustrates a diagram of an example method for data path circuit design with reinforcement learning, in accordance with at least some embodiments;
- (9) FIG. 8 illustrates an example system architecture, in accordance with at least some embodiments;
- (10) FIGS. 9A and 9B illustrate flow diagrams of example methods for data path circuit design with reinforcement learning, in accordance with at least some embodiments;
- (11) FIG. 10A illustrates an inference and/or training logic, in accordance with at least some embodiments.
- (12) FIG. 10B illustrates an inference and/or training logic, in accordance with at least some embodiments.
- (13) FIG. 11 illustrates an example data center system, in accordance with at least some embodiments.
- (14) FIG. 12 illustrates a computer system, in accordance with at least some embodiments.
- (15) FIG. 13 illustrates a computer system, in accordance with at least some embodiments.
- (16) FIG. 14 illustrates at least portions of a graphics processor, in accordance with at least some embodiments.
- (17) FIG. 15 illustrates at least portions of a graphic processor, in accordance with at least some embodiments.
- (18) FIG. 16 illustrates an example data flow diagram for an advanced computing pipeline, in accordance with at least some embodiments.
- (19) FIG. 17 illustrates a system diagram for an example system for training, adapting, instantiating and deploying machine learning models in an advanced computing pipeline, in accordance with at least some embodiment.
- (20) FIGS. 18A and 18B illustrate a data flow diagram for a process to train a machine learning model, as well as client-server architecture to enhance annotation tools with pre-trained annotation models, in accordance with at least some embodiments.

DETAILED DESCRIPTION

- (21) Several fundamental digital design building blocks such as adders, priority encoders, incrementers, decrementers and gray-to-binary code converters can be reduced to prefix-sum computations and implemented as data path circuits such as prefix circuits (e.g., parallel prefix circuits). The optimization of data path circuits such as prefix circuits for area, delay and power can be important in digital hardware design. Embodiments described herein provide a system and method to design data path circuits (e.g., parallel prefix circuits) that are optimized for area, power and/or delay using reinforcement learning.
- (22) Some memory systems store or communicate data—e.g., for a host system. The memory system can include data paths that communicate data from one component of the memory system to another component of the memory system. In some embodiments, the memory system can process or synthesize data at data path circuits. For example, the memory system can include prefix circuits—e.g., adders, decrementers, incrementers, gray to binary converters, priority encoders, etc.—to process and synthesize data. For example, an adder can be utilized to calculate addresses or table indices for a process of the memory system. In one example, the prefix circuits can be utilized in a parallel configuration to reduce a time to perform prefix computations—e.g., as parallel prefix circuits.
- (23) Each data path circuit can have an associated area (e.g., size of the data path circuit in the memory system), power consumption (e.g., an amount of power consumed by the data path circuit while in operation), and delay (e.g., an amount of time to generate an output from a given number of inputs). To increase the performance of a circuit (e.g., for a memory system), data path circuits can be designed to reduce the area, power consumption and/or delay. In some examples, though, reducing one property of the data path circuit can affect another property of the data path circuit—e.g., reducing the delay of the data path circuit can cause the data path circuit to be larger in area or consume more power. Accordingly, in embodiments data path circuits can be designed to balance the tradeoffs between reduced area, power consumption and/or delay—e.g., designed or optimized for

the smallest amount of area for a respective delay or power consumption.

(24) In some examples, conventional methods for design of data path circuits are ineffective at designing optimized data path circuits. For example, some conventional methods propose prefix circuit structures that optimize logic level, maximum fanout, and wiring tracks. However, optimizing logic level, fanout, and wiring tracks can fail to optimize for area, delay, and/or power due to physical design complexities of prefix structures—e.g., due to capacitive loading or congestion. Some conventional methods for data path circuit design can include utilizing heuristic rules for circuit design that either attempt to predict physical metrics for circuit designs or that are perform random modifications to circuit designs. The circuit designs generated using such heuristics may be evaluated using inaccurate analytical models. These analytical models and heuristic rules can be ineffective at producing optimal circuits as they rely on hand-crafted heuristics or are limited by analytical evaluation metrics. For example, prefix circuits designed using analytical evaluation metrics degrade in quality (e.g., see an increase in delay, area, or power consumption) when they undergo physical synthesis—e.g., when the analytical model is put through a simulation and converted to a predicted physical model. Because physical synthesis is more intensive than analytical evaluations, conventional methods generally cannot be scaled for physical synthesis. Accordingly, conventional circuit design techniques for data path circuits are not optimized and reduce the performance of systems that use these data path circuits (e.g., a memory system).

(25) Aspects of the present disclosure address the above and other deficiencies by designing data path circuits (e.g., parallel prefix circuits) using reinforcement learning with a machine learning model. For example, a first processing device (e.g., executing an agent) can give a machine learning model an initial design of a data path circuit. The machine learning model can modify the design to generate a second design for the data path circuit. In some examples, modifying the design can include modifying a prefix graph associated with the data path circuit—e.g., modifying a node of the prefix graph representing the data path circuit where each node represents one or more components of the data path circuit. For example, the machine learning model can modify the prefix graph represented by the initial design by adding or removing a node to generate a second prefix graph representing the second design. After the modification, the first processing device or a second processing device can process the first design and the second design to determine one or more metrics associated with the respective designs. For example, the first or second processing device can determine an area, power, and/or delay associated with the initial design and the second design generated by the machine learning model. In some embodiments, a second machine learning model is used to predict the area, power and/or delay associated with the initial and second circuit designs. In some embodiments, a circuit synthesis tool is used to predict the area, power, and/or delay of the initial and second circuit designs.

(26) The first or second processing device can generate a reward (e.g., parameter) that indicates a net change in area, delay, and/or power based transitioning from the initial design to the second design. The reward can be used to update a training of the machine learning model e.g., the machine learning model can undergo reinforcement learning based on the received reward. For example, the machine learning model can receive the reward, determine that the previous change caused the area, delay, and/or power of the data path circuit to decrease, and update weights of one or more nodes of the machine learning model. The updated machine learning model may then modify the second design of the data path circuit to generate a third data path circuit. This process may be repeated, and the reinforcement learning of the data path circuit design using the machine learning model can continue until a determination is made that no further improvements are being output (e.g., there are no additional modifications to the design that reduce the delay, area, and/or power of the data path circuit) or that the delay, area, and/or power of the data path circuit satisfies a target delay, area, and/or power for the data path circuit—e.g., until the data path circuit design is optimized for a respective delay, area, and/or power constraint.

(27) Embodiments avoid the use of hand-crafted heuristics (e.g., such as the heuristics for pruning) that are applied in conventional circuit design techniques. In embodiments, a machine learning model can be trained to perform circuit design for data path circuits via reinforcement learning techniques, as described herein. That is, the machine learning model learns to modify the data path circuit in a way that optimizes the current design of the data path circuit—e.g., to modify the data path circuit to reduce the area, power consumption and/or delay of the data path circuit. For example, the machine learning model may be trained based on receiving a reward indicating one or more improved or decreased circuit design optimization metric values and determining whether a previous modification of the data path circuit design resulted in an improved or decreased circuit design optimization metric value, and then adjusting nodes of the machine learning model based on the reward. In some examples, the machine learning model can be trained using model (based/free), value/policy based, or on/off policy reinforcement learning techniques. In other embodiments, the machine learning model can utilize deep neural networks (e.g., a convolutional neural network, transformer, graph neural network etc.) and/or

decision trees. By utilizing reinforcement learning, the design of the data path circuit can be more optimized more thoroughly than compared with other solutions—e.g., the delay, area, and/or power of the data path circuit can be reduced as compared to data path circuits designed using traditional approaches. Embodiments produce circuit designs that have smaller area and power for given delays and that have lower delays for given area and power as compared to prior approaches at circuit design of data path circuits. Accordingly, the overall performance of systems such as a memory system can increase with use of embodiments of the present disclosure, as designed data path circuits can consume less power, have less delay, and/or consume less area as compared to traditionally designed data path circuits.

(28) FIG. 1 illustrates a computer system **100**, in accordance with at least one embodiment. In at least one embodiment, computer system **100** may be a system with interconnected devices and components, an SOC, or some combination thereof. In at least one embodiment, computer system **100** is formed with a processor **102** that may include execution units to execute an instruction. In at least one embodiment, computer system **100** may include, without limitation, a component, such as processor **102** to employ execution units including logic to perform algorithms for processing data. In at least one embodiment, computer system **100** may include processors, such as PENTIUM® Processor family, Xeon™, Itanium®, XScale™ and/or StrongARM™, Intel® Core™, or Intel® Nervana™ microprocessors available from Intel Corporation of Santa Clara, Calif., although other systems (including PCs having other microprocessors, engineering workstations, set-top boxes and like) may also be used.

(29) In at least one embodiment, computer system **800** may be used in other devices such as handheld devices and embedded applications. Some examples of handheld devices include cellular phones, Internet Protocol devices, digital cameras, personal digital assistants (“PDAs”), and handheld PCs. In at least one embodiment, embedded applications may include a microcontroller, a digital signal processor (DSP), an SoC, network computers (“NetPCs”), set-top boxes, network hubs, wide area network (“WAN”) switches, or any other system that may perform one or more instructions. In an embodiment, computer system **100** may be used in devices such as graphics processing units (GPUs), network adapters, central processing units and network devices such as switch (e.g., a high-speed direct GPU-to-GPU interconnect such as the NVIDIA GH100 NVLINK or the NVIDIA Quantum 2 64 Ports InfiniBand NDR Switch).

(30) In at least one embodiment, computer system **100** may include, without limitation, processor **102** that may include, without limitation, one or more execution units **107** that may be configured to execute a Compute Unified Device Architecture (“CUDA”) (CUDA® is developed by NVIDIA Corporation of Santa Clara, Calif.) program. In at least one embodiment, a CUDA program is at least a portion of a software application written in a CUDA programming language. In at least one embodiment, computer system **100** is a single processor desktop or server system. In at least one embodiment, computer system **100** may be a multiprocessor system. In at least one embodiment, processor **102** may include, without limitation, a CISC microprocessor, a RISC microprocessor, a VLIW microprocessor, a processor implementing a combination of instruction sets, or any other processor device, such as a digital signal processor, for example. In at least one embodiment, processor **102** may be coupled to a processor bus **110** that may transmit data signals between processor **102** and other components in computer system **100**.

(31) In at least one embodiment, processor **102** may include, without limitation, a Level 1 (“L1”) internal cache memory (“cache”) **104**. In at least one embodiment, processor **102** may have a single internal cache or multiple levels of internal cache. In at least one embodiment, cache memory may reside external to processor **102**. In at least one embodiment, processor **102** may also include a combination of both internal and external caches. In at least one embodiment, a register file **106** may store different types of data in various registers including, without limitation, integer registers, floating point registers, status registers, and instruction pointer register.

(32) In at least one embodiment, processor **102** can include data path circuits **108**. In some examples, data path circuits **108** can be examples of parallel prefix circuits. For example, data path circuits **108** can be adders, incrementers, decrementers, priority encoders, and/or gray to binary converters, possibly including connected logic. In some examples, data path circuits **108** can also be located in other components of computer system **100**. In some embodiments, data path circuits **108** can consume power, take up a respective area, and have a respective delay. In some embodiments, the data path circuit **108** delay can be inversely related to a clock frequency of a component of computer system **108**—e.g., the delay of data path circuit **108** can be utilized to set a clock frequency for a component of computer system **108**. In some embodiments, the design of data path circuits **108** can be performed via reinforcement learning using a machine learning model that is trained over time to reduce or optimize the area, power consumption and/or delay of the data path circuit **108**. For example, during design of the data patch circuit **108**, the machine learning model can modify a design of the data path circuit **108** and determine if the modification resulted in a reduction of area, power, or delay of the given data

path circuit **108**. The computer system **100** can benefit from optimized data path circuits that were designed using reinforcement learning. That is, machine learning model can be updated based on whether a resulting modification to the design of the data path circuit resulted in a reduction of the area, delay, and/or power of the data path circuit **108**. Over several iterations, the machine learning model can be trained to choose modifications that result in the best reductions of area, power consumption and/or delay of the data path circuit. In some embodiments, the machine learning model can be used until the area, power consumption and/or delay of the data path circuit **108** satisfies a target metric—e.g., the area, power consumption and/or delay of the data path circuit **108** satisfies a target area, power consumption and/or delay. By utilizing reinforcement learning to design the data path circuit **108**, the design of data path circuit **108** can be optimized and the performance of the computer system **100** can be improved.

(33) In at least one embodiment, execution unit **107**, including, without limitation, logic to perform integer and floating point operations, also resides in processor **102**. Processor **102** may also include a microcode (“ucode”) read only memory (“ROM”) that stores microcode for certain macro instructions. In at least one embodiment, execution unit **102** may include logic to handle a packed instruction set **109**. In at least one embodiment, by including packed instruction set **109** in an instruction set of a general-purpose processor **102**, along with associated circuitry to execute instructions, operations used by many multimedia applications may be performed using packed data in a general-purpose processor **102**. In at least one embodiment, many multimedia applications may be accelerated and executed more efficiently by using full width of a processor's data bus for performing operations on packed data, which may eliminate a need to transfer smaller units of data across a processor's data bus to perform one or more operations one data element at a time.


(34) In at least one embodiment, an execution unit may also be used in microcontrollers, embedded processors, graphics devices, DSPs, and other types of logic circuits. In at least one embodiment, computer system **100** may include, without limitation, a memory **120**. In at least one embodiment, memory **120** may be implemented as a DRAM device, an SRAM device, flash memory device, or other memory device. Memory **120** may store instruction(s) **119** and/or data **121** represented by data signals that may be executed by processor **102**.

(35) In at least one embodiment, a system logic chip may be coupled to processor bus **110** and memory **120**. In at least one embodiment, the system logic chip may include, without limitation, a memory controller hub (“MCH”) **116**, and processor **102** may communicate with MCH **116** via processor bus **110**. In at least one embodiment, MCH **116** may provide a high bandwidth memory path **118** to memory **120** for instruction and data storage and for storage of graphics commands, data and textures. In at least one embodiment, MCH **116** may direct data signals between processor **102**, memory **120**, and other components in computer system **100** and to bridge data signals between processor bus **110**, memory **120**, and a system I/O **122**. In at least one embodiment, system logic chip may provide a graphics port for coupling to a graphics controller. In at least one embodiment, MCH **116** may be coupled to memory **120** through high bandwidth memory path **118** and graphics/video card **112** may be coupled to MCH **116** through an Accelerated Graphics Port (“AGP”) interconnect **114**.

(36) In at least one embodiment, computer system **100** may use system PO **122** that is a proprietary hub interface bus to couple MCH **116** to I/O controller hub (“ICH”) **130**. In at least one embodiment, ICH **130** may provide direct connections to some I/O devices via a local I/O bus. In at least one embodiment, local I/O bus may include, without limitation, a high-speed I/O bus for connecting peripherals to memory **120**, a chipset, and processor **102**. Examples may include, without limitation, an audio controller **129**, a firmware hub (“flash BIOS”) **128**, a wireless transceiver **126**, a data storage **124**, a legacy I/O controller **123** containing a user input interface **125** and a keyboard interface, a serial expansion port **127**, such as a USB, and a network controller **134**. Data storage **124** may comprise a hard disk drive, a floppy disk drive, a CD-ROM device, a flash memory device, or other mass storage device.

(37) FIG. 2 illustrates an example system **200** for performing reinforcement learning to generate an improved design of a data path circuit, according to at least one embodiment. In some embodiments, the system is or includes a Q-network. In some embodiments, the system **200** is or includes a deep Q-network.

(38) Reinforcement learning (RL) is a class of algorithms applicable to sequential decision making tasks. RL makes use of the Markov Decision Process (MDP) formalism wherein an agent **202** attempts to optimize a function in its environment **204**. An MDP can be completely described by a state space S (with states $s \in S$), an action space A ($a \in A$), a transition function $T: S \times A \rightarrow S$ and a reward function $R: S \times A \rightarrow \mathbb{R}$.

 In an MDP, an episode evolves over discrete time steps $t=0, 1, 2, \dots, n$, where the agent **202** observes a state $s_{sub.t}$ (**206**) and responds with an action $a_{sub.t}$ (**210**) using a policy $\pi(a_{sub.t}|s_{sub.t})$. The environment **204** provides to the agent **202** the next state $s_{sub.t+1} \sim T(s_{sub.t}, a_{sub.t})$ (**212**) and the reward $r_{sub.t} = R(s_{sub.t}, a_{sub.t})$ (**214**). The agent **202** is tasked with maximizing the return (cumulative future rewards) by learning an optimal policy π^* .

(39) A Q-network may be trained via a process referred to as Q-learning. Q learning is a reinforcement learning process that causes a model to perform a sequence of actions that will eventually generate a maximum total reward. This total reward is also called the Q-value. A function for computing the Q value may be as follows:

$$Q(s_{\text{sub.t}}, \alpha_{\text{sub.t}}) \leftarrow r_{\text{sub.t}} + \gamma \max_{\alpha} Q(s_{\text{sub.t}+1}, \alpha)$$

(40) The above equation states that the Q-value yielded from being at state $s_{\text{sub.t}}$ (206) and performing action $\alpha_{\text{sub.t}}$ (210) is the immediate reward $r(s_{\text{sub.t}}, a_{\text{sub.t}})$ (214) plus the highest Q-value possible from the next state $s_{\text{sub.t}+1}$ (212), state where γ is the discount factor which controls the contribution of rewards further in the future. The recursive definition of Q-functions allows the expression of $Q(s_{\text{sub.t}+1}, a_{\text{sub.t}+1})$ to be unrolled into future states, as follows:

$$Q(s_{\text{sub.t}}, \alpha_{\text{sub.t}}) = r_{\text{sub.t}} + \gamma r_{\text{sub.t}+1} + \dots + \gamma^{n-1} r_{\text{sub.t}+n-1} + \gamma^n Q(s_{\text{sub.t}+n}, \alpha_{\text{sub.t}+n})$$

The machine learning model 240 of the agent 202 learns to predict $Q(s_{\text{sub.t}}, a_{\text{sub.t}})$ by performing the following update step:

$$(41) \quad Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha [r_{t+1} + \gamma \max_a Q(s_{t+1}, a) - Q(s_t, a_t)]$$

where α represents learning rate or step size, which controls to what extent newly acquired information overrides old information.

(42) In embodiments, the Q value of a state-action pair $(s_{\text{sub.t}}, a_{\text{sub.t}})$ under a policy π is defined to be the expected return if the action $a_{\text{sub.t}}$ 210 is taken at state $s_{\text{sub.t}}$ 206 and future actions are taken using the policy π , as set forth below:

$$Q_{\pi}(s_{\text{sub.t}}, \alpha_{\text{sub.t}}) = \mathbb{E}[r_{\text{sub.t}} + \gamma r_{\text{sub.t}+1} + \gamma^2 r_{\text{sub.t}+2} + \dots], \gamma \in [0, 1]$$

(43) In embodiments, the discount factor $\gamma \in [0, 1]$ balances short-term versus long-term rewards. The Q-learning algorithm may start the agent 202 with a random policy and uses the experience gathered during its interaction with the environment $(s_{\text{sub.t}}, a_{\text{sub.t}}, r_{\text{sub.t}}, s_{\text{sub.t}+1})$ 204 to iterate towards an optimal policy by updating Q with a learning rate $\alpha \in [0, 1]$:

$$(44) \quad Q(s_t, a_t) \leftarrow (1 - \alpha) * Q(s_t, a_t) + \alpha * (r_t + \gamma \max_a Q(s_{t+1}, a))$$

(45) The policy for a Q-learning agent 202 may be represented as $\pi(\cdot | s_{\text{sub.t}}) = \arg \max_{\alpha} Q(s_{\text{sub.t}}, \alpha)$. In one embodiment, a ϵ -greedy policy, where random actions a are chosen with a probability ϵ to increase exploration in the state space is used. In one embodiment, ϵ is annealed to zero during the course of training and is zero when performing evaluation. In one embodiment, multiple explorations can be done in parallel with a range of ϵ values.

(46) Deep Q-learning is an extension of Q-learning that implements one or more machine learning models (e.g., machine learning model 240) such as neural networks to essentially approximate the aforementioned Q values. In deep Q-learning, one or more artificial neural networks (e.g., machine learning model 240) may be used to approximate the aforementioned Q-value function. Artificial neural networks generally include a feature representation component with a classifier or regression layers that map features to a desired output space. A convolutional neural network (CNN), for example, hosts multiple layers of convolutional filters. Pooling is performed, and non-linearities may be addressed, at lower layers, on top of which a multi-layer perceptron is commonly appended, mapping top layer features extracted by the convolutional layers to decisions (e.g. modifications to design state of prefix circuits). Deep learning is a class of machine learning algorithms that use a cascade of multiple layers of nonlinear processing units for feature extraction and transformation. Each successive layer uses the output from the previous layer as input. Deep neural networks may learn in a supervised (e.g., classification) and/or unsupervised (e.g., pattern analysis) manner. Deep neural networks include a hierarchy of layers, where the different layers learn different levels of representations that correspond to different levels of abstraction. In deep learning, each level learns to transform its input data into a slightly more abstract and composite representation. Notably, a deep learning process can learn which features to optimally place in which level on its own. The “deep” in “deep learning” refers to the number of layers through which the data is transformed. More precisely, deep learning systems have a substantial credit assignment path (CAP) depth. The CAP is the chain of transformations from input to output. CAPs describe potentially causal connections between input and output. For a feedforward neural network, the depth of the CAPs may be that of the network and may be the number of hidden layers plus one. For recurrent neural networks, in which a signal may propagate through a layer more than once, the CAP depth is potentially unlimited.

(47) In embodiments, a deep Q-network (DQN) may stabilize training using a second target network to estimate the Q values of $(s_{\text{sub.t}+1}, a')$. The second target network may be updated less frequently than a first network. In embodiments, the DQN may sample an experience replay buffer. In one embodiment, a first machine learning model is used to determine a prediction and a second neural network is used to determine a target. The second neural network may have a same architecture as the first neural network in embodiments. However, in an

embodiment of the second neural network may have frozen parameters while the first neural network may have variable parameters. In an embodiment, the second neural network is updated less frequently than the first neural network. In one embodiment, a double-DQN algorithm is used, which may further improve training by reducing overestimations in the DQN.

(48) In some embodiments, system **200** can include agent (e.g., an actor, a circuit modifier, etc.) **202** and environment (e.g., a simulation environment, a circuit synthesizer, etc.) **204**. In some embodiments, the agent **202** can include one or more machine learning model **240**. The machine learning model **240** may be, for example a deep neural network (e.g., a convolutional neural network, transformer, graph neural network etc.) or decision trees. For example, the machine learning model **240** may be a neural network of a deep Q-network.

(49) In some examples, system **200** can be utilized to design a data path circuit **108** as described with reference to FIG. 1. In embodiments, the optimization of data path circuits (e.g., prefix circuits) is framed as an RL task by creating an MDP for their construction. For example, the RL system **200** may be trained to select a design for parallel-prefix adders (e.g., to design an area-delay minimized pareto frontier of adders with possible connected logic included). In some embodiments, agent **202** is configured to modify a design of a data path circuit and train the machine learning model **240** based on the modifications. In some embodiments, agent **202** can execute on a processing device such as a graphical processing unit (GPU) or a central processing unit (CPU). In some embodiments, the system **200** can include multiple agents **202** that may operate in parallel and share learning. Each agent **202** may execute on the same or a different processing device and/or the same or a different core of a processing device. Each agent may perform a modification of the data path circuit in parallel—e.g., multiple agents **202** may concurrently modify the data path circuit. This may decrease an amount of time that it takes to find a target or optimal data path circuit design. In some embodiments, the agent **202** (or agents) can receive a design state **206** of the data path circuit. Upon receiving the design state **206**, the agent **202** (or agents) can modify the design state **206** via an action **210**. Each agent **202** may output a different action **210** in embodiments. In some embodiments, the action **210** can be determined for the agent **202** by the machine learning model **240** (e.g., such as a deep neural network). The agent **202** can also be configured to output a modified state of the data path circuit to the environment **204**.

(50) In embodiments, the design of a data path circuit **108** (e.g., a prefix circuit such as a parallel prefix adder) can be represented using a prefix graph. A prefix graph is a directed acyclic graph (DAG) in which edges may represent signals or signal pairs and nodes may represent logic operators. For example, parallel prefix computations can be represented as a directed acyclic parallel prefix graph where every computation unit $z_{sub.i:j}$ is a graph node that performs a single operation on two inputs: $z_{sub.i:j} = z_{sub.i:k} \circ z_{sub.k-1:j}$, where \circ represents an operation such as a sum operation, a carry operation, a difference operation, and so on.

(51) Accordingly, the agent **202** can receive a prefix graph representing an initial design state **206** of a data path circuit. The agent **202** may then modify the prefix graph (e.g., using a machine learning model **240**) and output a new prefix graph that represents a modified state of the data path circuit via action **210**. In embodiments, modifying the prefix graph may include adding a node, removing a node, or moving a node in the prefix graph.

(52) FIG. 3 illustrates an example prefix graph modification, in accordance with at least some embodiments. A data path circuit can be represented by a prefix graph **305**, **310**, where prefix graph **305** represents an initial state of a circuit design for a data path circuit (e.g., such as an adder), and prefix graph **310** represents a modified or updated state of the data path circuit. In a prefix problem, n inputs $x_{sub.n-1}, x_{sub.n-2}, \dots, x_{sub.0}$ and an arbitrary associative operator are used to compute n outputs $y_{sub.i} = x_{sub.i} \circ x_{sub.i-1} \circ \dots \circ x_{sub.0}$, $i=0, \dots, n-1$. Thus, each output $y_{sub.i}$ is dependent on all inputs $x_{sub.j}$ of the same or lower index ($j \leq i$).

(53) In an example, an N -input prefix-sum computation can be performed in several ways due to the associativity of the operator. For example, two of the ways the 4-input prefix sum can be computed are:

$$y_{sub.0} = x_{sub.0}, y_{sub.1} = x_{sub.1} \circ y_{sub.0}, y_{sub.2} = x_{sub.2} \circ y_{sub.1}, y_{sub.3} = x_{sub.3} \circ y_{sub.2}$$

$$y_{sub.0} = x_{sub.0}, y_{sub.1} = x_{sub.1} \circ y_{sub.0}, y_{sub.2} = x_{sub.2} \circ y_{sub.1}, z_{sub.3:2} = x_{sub.3} \circ x_{sub.2}, y_{sub.3} = z_{sub.3:2} \circ y_{sub.1}$$

(54) In this example, introducing the additional term $z_{sub.3:2}$ breaks the dependency of $y_{sub.3}$ on $y_{sub.2}$ and allows it to be computed in parallel with $y_{sub.2}$, thus the term parallel prefix. Embodiments denote $z_{sub.i:j}$ to represent $x_{sub.i} \circ x_{sub.i-1} \circ \dots \circ x_{sub.j}$. Then the outputs $y_{sub.i}$ can be rewritten as $z_{sub.i:0}$ and inputs $x_{sub.i}$ can be rewritten as $z_{sub.i:i}$. Note that $y_{sub.0}$ and $x_{sub.0}$ may both correspond to $z_{sub.0:0}$.

(55) Parallel prefix computations can be represented as a directed acyclic parallel prefix graph where every computation unit $z_{sub.i:j}$ is a graph node that performs a single operation on two inputs: $z_{sub.i:j} = z_{sub.i:k} \circ z_{sub.k-1:j}$. In embodiments, most and least significant bits (MSB, LSB) of computation node $z_{sub.i:j}$ may be (i, j) . Using this notation, the node (i, k) may be the upper parent of (i, j) and the node $(k-1, j)$ may be the lower parent of (i, j) . The prefix graphs corresponding to the 4-input prefix sum computations of the above example are shown in FIG. 4 as prefix graph **305** and prefix graph **310**. In both graphs, the upper and lower parents of

node (2, 0) are (2, 2) and (1, 0).

(56) In at least one embodiment, every valid N-input prefix graph has input nodes (i, i), output nodes (i, 0) for $1 \leq i \leq N-1$, and the input/output node (0, 0). Furthermore, in at least one embodiment every non-input node has exactly one upper parent (up) and one lower parent (lp) such that:

$LSB(node) = LSB(lp(node))$

$LSB(lp(node)) \leq MSB(lp(node))$

$MSB(lp(node)) = LSB(up(node)) - 1$

$LSB(up(node)) \leq MSB(up(node))$

$MSB(up(node)) = MSB(node)$

(57) In an example, a data path circuit can receive inputs **315-a** through **315-d** and produce outputs **325-a** through **325-d**. In some examples, each input **315** and output **325** can represent signals received or produced by the data path circuit or wires coming to the data path circuit. In some examples, each input **315** can represent an element (or some data) the data path circuit is configured to synthesize. For example, in embodiments where the data path circuit is a binary adder, each input can represent a bits of the inputs with which the adder circuit performs a computational sum—e.g., the input **315-a** can represent the zeroth bits or one's place and the input **315-d** can represent the third bit or eight's place of numbers the adder circuit is to find the sum of. Accordingly, each output **325** can represent a value or bit generated by the data path circuit during its operation. In some embodiments, each output **325** can be generated from each previous input—e.g., output **325-c** can be generated from input **315-a** through **315-c** while output **325-d** can be generated from input **315-a** through **315-d**.

(58) In some embodiments, each prefix graph can also include one or more nodes. For example, prefix graph **305** can include nodes **320-a** through **320-c**. In some embodiments, each node represents or is associated with one or more components of data path circuit that perform one or more operations. For example, node **320-a** can represent one or more logic gates (e.g., AND gate, NOR gate, XOR gate, etc.) of the data path circuit **108**. In some embodiments, the nodes **320** can also represent buffers or other types of gates.

(59) In some embodiments, the area of a data path circuit can be related to the number of nodes **320** in the prefix graph representing the data path circuit—e.g., the greater the number of nodes **320**, the greater the area associated with the data path circuit. For example, the area of a data path circuit represented by prefix graph **305** can be less than the area of a data path circuit represented by prefix graph **310**—e.g., there are less nodes **320** in prefix graph **305**. In some embodiments, the prefix graphs can also illustrate the delay associated with the data path circuit. For example, the delay of the data path circuit can be related to a longest path an input **315** takes to be generated as an output **325**. For example, input **315-a** goes through three nodes **320** before being utilized in the generating of output **325-d** in prefix graph **305** while each input **315** of prefix graph **310** goes through no more than two nodes **320**. In such embodiments, the delay of the data path circuit associated with prefix graph **310** can be less than the delay of the data path circuit **108** associated with prefix graph **305**.

(60) Referring to FIG. 2 and FIG. 3, in some examples, each design state **206** of the data path circuit can be represented by a unique prefix graph. In such embodiments, the agent **202** can receive a prefix graph that represents the current design state **206** of the data path circuit **108**. In some embodiments, a grid representation of the prefix graph is received, as set forth in detail below and with reference to FIG. 4. In an example, if the prefix graph **305** represents an initial design state **206** of the data path circuit, the agent **202** can receive the prefix graph **305** (or a grid representation of the prefix graph) from the environment **204** before performing a modification thereto.

(61) In some embodiments, the agent **202** can be configured to take an action **210** that modifies a node **320** of a prefix graph representing the current design state **206** of the data path circuit. In at least one embodiment, the agent **202** can add or remove a node **320** from the prefix graph representing the current design state **206** of the data path circuit. For example, the agent **202** can add a node **320-d** to the prefix graph **305**. In some embodiments, the machine learning model **240** can determine which node **320** to remove or determine where to add a node **320** to the prefix graph. In some embodiments, the machine learning model **240** receives an input grid representation of a prefix graph and outputs a grid representation of a modified prefix graph in which a node has been added or removed.

(62) In some embodiments, agent **202** includes a graph validity determiner **244**. Alternatively, graph validity determiner **244** may be included in environment **204**. Graph validity determiner **244** may assess the validity of an updated prefix graph (e.g., that is output by machine learning model **240**). If an updated prefix graph is invalid, then graph validity determiner **244** performs one or more further modifications to the prefix graph to cause it to be valid. Such modifications may be the addition or removal of one or more nodes to/from the prefix graph. In some embodiments, if the graph validity determiner **244** determines the action **210** results in an invalid state (e.g., invalid prefix graph), the graph validity determiner **244** can modify the design state **206** again. That

is, the graph validity determiner **244** can validate the updated state following the action. In some embodiments, the graph validity determiner **244** can validate the state by adding or removing invalid nodes and/or making sure each node follows the rules as specified below.

(63) In embodiments, a valid N-input prefix graph has input nodes (i, i), output nodes (1,0) for $1 \leq i \leq N-1$, and the input/output node (0,0) as described above. For example, in prefix graph **310**, input nodes (1,1) to (3,3) correspondent to inputs **315-b** to **315-d** at index one (1) to three (3); output nodes (1,0) to (3,0) correspond to nodes **320-a** to **320-c** that feed outputs **325-b** to **325-d** at index one (1) to three (3); input/output node (0,0) correspond to input **315-a** at index zero (0) that also feeds output **325-a** at index zero (0).

(64) In some embodiments, a valid prefix graph can be one where each non-input node (e.g., each node that is not (0,0), (1,1), (2,2), etc.) has exactly two parents, an upper and a lower parent, that it directly receives values from. These parents may be another non-input node **320** or an input **315**. That is, a prefix graph where a non-input node **320** has one parent or more than two parents is invalid. In some embodiments, the valid prefix graph also has each node **320** having a sequential contiguous range of input indices it directly or indirectly receives values from. For example, node **320-b** is a valid node with a range (2,0) as it receives values from inputs at index zero (0), one (1) and two (2). In the (MSB, LSB) notation for the node, The most-significant bit (MSB) represents the upper end or first element of the node's range and the least-significant bit (LSB) represents the lower end or second element of the node's range. A node **320** with range (0,2) would be invalid as the range increases from the MSB to the LSB. In some embodiments, a prefix graph with a node **320** having a range (6,3) but not receiving a value from input at index five (5) is invalid—e.g., the range is not contiguous as the node **320** receives inputs at index three (3), four (4), and six (6) but not five (5).

(65) Additionally, each non-input node can directly receive values from exactly one upper parent and one lower parent that are contiguous. For example, if a node **320** has a range (3,1), and has an upper parent having a range (3,3), then the node **320** must also have a lower parent having a range (2,1). That is, node **320**'s upper parent's range must have the MSB as node **320**'s range (e.g., 3) and node **320**'s lower parent's range must have the same LSB as node **320**'s range (e.g., 1) while also being contiguous—e.g., including the input at index 2. Accordingly, the prefix graph can follow the rules above regarding the upper and lower parents.

(66) In embodiments, the action space A for an N-input prefix graph consists of two actions (add or delete) for any non-input/output node—e.g., where $LSB \in [1, N-2]$ and $MSB \in [LSB+1, N-1]$. Hence, $|A| = (N-1) \times (N-2) / 2$. The environment evolution through T may maintain valid prefix graphs by: 1) applying a validation procedure (e.g., performed by graph validity determiner **244**) after an action that may add or delete additional nodes to maintain validity; and 2) forbidding redundant actions that would become undone by the validation procedure.

(67) In one embodiment, during validation, the upper parent of a node, $up(node)$, is the existing node with same MSB and the next highest LSB. In one embodiment, the lower parent of a node, $lp(node)$ is computed using the node and its upper parent according to:

$(MSB.sub.lp(node), LSB.sub.lp(node)) = (LSB.sub.up(node) - 1, LSB.sub.node)$

(68) In one embodiment, an invalid condition happens when the lower parent of a node does not exist. In such a condition, graph validity determiner **244** performs the validation procedure to add any missing lower parent nodes.

(69) In one embodiment, system **200** (e.g., environment **204** or agent **202**) maintains a list of all nodes **nodelist** in a valid prefix graph. In one embodiment, the action of adding a node that already exists in **nodelist** is redundant and is forbidden. In one embodiment, system **200** (e.g., environment **204** or agent **202**) maintains a minimal list **minlist** of nodes from **nodelist** that are not lower parents of other nodes. In one embodiment, the action of deleting a node is limited to those in **minlist**, otherwise a node may be deleted that either does not exist, or that is the lower parent of another node. For such modifications, the deleted node may be added back during validation.

(70) Environment **204** receives action **210** (e.g., a prefix graph update) and determines a next design state **212** and a reward **214** associated with the action **210** on current design state **206**. The environment **204** may start with an initial state $S.sub.0$, which may be randomly chosen in some embodiments. In some examples, the environment **204** can execute on one or more processing device, such as a CPU or GPU. In some embodiments, the system **200** can include multiple environments **204**, each of which may receive a different action (e.g., a different modified state of the data path circuit). In an example, multiple environments **204** may concurrently receive modified states of the data path circuit as output by different agents **202** to decrease an amount of time to find a target data path circuit design. In at least one embodiment, the environment **204** can update or modify the design state **206** based on the action **210** chosen by the machine learning model **240**. In one embodiment, the environment **204** can generate a prefix graph for the state **212** after the action **210**—e.g., the environment can generate a prefix graph **310** when the action **210** specifies to add a node **320** with a value (3,2) to the prefix

graph **305**. In some embodiments, the agent **202** can receive the prefix graph associated with the modified state —e.g., the agent **202** can receive the prefix graph **310** or a grid representation of the prefix graph.

(71) For every action **210** (e.g., for every design state of a data path circuit), a prefix graph assessor **242** of the environment **204** may estimate one or more parameters of the data path circuit. Such estimated parameters may include an area of the data path circuit, power consumption of the data path circuit, and/or delay of the data path circuit, for example. The prefix graph assessor **242** may then compare the determined parameters to one or more goals and/or constraints for the data path circuit. Additionally, the prefix graph assessor **242** may estimate (or may have previously estimated) similar parameters for a previous state of the data path circuit (e.g., design state **206**). Prefix graph assessor **242** may additionally compare the parameters of for the previous state of the data path circuit to the constraints and/or goals. Prefix graph assessor **242** may compare the parameters and/or a distance between the parameters and the goals of the initial design state **206** to the parameters and/or distance between the parameters and the goals of the updated design state **212**. Based on such comparison, prefix graph assessor **242** may output a reward **214**. For example, if the parameters associated with the updated design state are closer to the goal than the parameters associated with the previous design state, then a positive reward **214** may be output. On the other hand, if the parameters associated with the updated design state are further from the goal than the parameters associated with the previous design state, then a negative reward **214** may be output. Environment **204** may also output a next design state **212** to be input into the agent **202**.

(72) If the environment **204** determines the next design state **212** following the action **210** to modify the current design state **206** is valid or has validated the modified design state, the environment can calculate a next reward **214**. In some embodiments, the reward **214** predicts the net change in area, power consumption and/or delay of the data path circuit **108** as a result of the action **210**. That is, the environment **204** can calculate the delay, area, and power for the initial state (or current design state **206**) and the delay, area, and power for the next design state **212** and determine a difference between the two to calculate the next reward **214**. In one embodiment, the environment **204** can determine the reward as described with reference to FIG. 8. In some embodiments, the environment **204** can determine the reward via a second machine learning model as described with reference to FIG. 7. That is, the environment **204** can include a second machine learning model that predicts the change in area, delay, and power of the data path circuit following the action **210**. In some examples, the environment **204** can determine the reward **208** (or next reward **214**) as described with reference to FIG. 5.

(73) For example, the environment **204** can determine an area, delay, and/or power consumed by a data path circuit **108** having a design represented by prefix graph **305**. In some examples, the prefix graph assessor **242** of the environment **204** can utilize a synthesis tool **505** that computes the area, delay, and/or power of a physical data path circuit given a prefix graph (e.g., prefix graph **305**). In some examples, the synthesis tool **505** can determine types of logic gates to use, determine sizes of a logic gates to use, determine the connectivity of logic gates, determine if buffers or other components will optimize the prefix graph, and so on. In some examples, generating a potential physical data path circuit from the prefix graph via the synthesis tool **505** can increase or decrease the area, delay, or power consumption of the data path circuit. That is, generating the potential physical data path circuit can cause changes and modifications to the specific circuit implementation from the prefix graph due to physical and manufacturing constraints. For example, if a node of a prefix graph was an input to four other nodes (e.g., four other nodes receive value from the node), the synthesis tool **505** could insert a buffer when generating the potential physical data path circuit, causing an increase in area.

(74) In some embodiments, the prefix graph assessor **242** of environment **204** can also determine the area, delay, and/or power consumed by a data path circuit having a design represented by prefix graph (e.g., prefix graph **310**) that is output by agent **202** (e.g., that is the result of the modification and validation of an input prefix graph). To determine the area, delay, and/or power of the data path circuit, the environment can use the synthesis tool **505** to generate a predicted physical data path circuit from the prefix graph. For each calculation (e.g., for calculation of the initial prefix circuit such as prefix circuit **305** and calculation of the modified or updated prefix circuit such as prefix circuit **310**), the prefix graph assessor **242** can determine an area/delay curve for a graph **510**. That is, the environment **204** can determine a curve **515** that represents the delay of a data path circuit associated with the initial prefix graph (e.g., prefix graph **305**) for respective areas. Similarly, the environment **204** can determine a curve **520** that represents the delay of a data path circuit associated with an updated prefix graph (e.g., prefix graph **310**) for respective areas.

(75) In some embodiments, the environment can calculate the reward of the action **210** by determining a difference in the delay and area based on the weight constraint described with reference to FIG. 6. For example, if the weight is one (1), the reward can be calculated by taking a difference between points **530** and **535**—e.g., based on the weight curve described with reference to FIG. 6. In this embodiment, the modification from prefix graph **305** to prefix graph **310** via the action **210** reduced the area and the delay. Accordingly, the environment

204 can calculate a next reward **214** that indicates the amount of area reduced and the amount of delay reduced via the action **210**. In other embodiments, the environment **204** can calculate a next reward **214** that indicates the area increased, the area decreased, the delay increased, the delay decreased, or a combination thereof. Although area vs delay is shown in FIG. 5, the environment **204** can determine the difference in area, power, delay, or any combination thereof when calculating the next reward **214**. In some embodiments, the environment **204** can send the next reward to the agent **202**—e.g., to the machine learning model **240**.

(76) The agent **202** can be configured to receive the next reward **214**. In some embodiments, the agent **202** can determine whether the action **210** taken optimized the data path circuit—e.g., whether the action **210** caused a reduction in area, power, and/or delay of the data path circuit. In some embodiments, the agent **202** can use reinforcement learning to train the machine learning model **240**. That is, the agent **202** can train the machine learning model **240** based on the next rewards **214** received. In one embodiment, agent **202** trains the machine learning model(s) **240** based on the reward **214** and the previously output action **210**. Training can be performed by defining an error (e.g., based on the action and reward), and using techniques such as stochastic gradient descent and backpropagation to tune the weights of the network across all its layers and nodes such that the error is minimized.

(77) In an example, if the next reward **214** indicates a reduction in area, power, or delay, the machine learning model **240** can continue taking actions **210** similar to those that resulted in the reduction. If the next reward **214** indicates an increase in area, power or delay, the machine learning model **240** can be trained to take other actions **210** by adjusting the weights of nodes in the machine learning model. In some embodiments, the machine learning model **240** can be trained using model (based/free), value/policy based, or on/off policy reinforcement learning techniques.

(78) As discussed above, in embodiments where the machine learning model **240** is used as a component of a deep Q-network, the machine learning model **240** may receive an input of the design state **206** predict an area and delay reduction for an action **210**. For example, the machine learning model **240** may receive state $s_{sub.t}$ **206** as input and predict:

$$\forall \alpha \in A: [Q_{sub.area}(s_{sub.t}, \alpha), Q_{sub.delay}(s_{sub.t}, \alpha)].$$

(79) The input to the machine learning model can be a $N \times N \times 4$ tensor where the four channels encode node features as: 1) 1 if node (MSB, LSB) in nodelist, 0 otherwise; 2) 1 if node (MSB, LSB) in minlist, 0 otherwise; 3) Level of node (MSB, LSB) in nodelist, 0 otherwise; 4) Fanout of node (MSB, LSB) in nodelist, 0 otherwise where the nodelist is all nodes **320** in a valid prefix graph, minlist is all nodes **320** in the nodelist that are not lower parents of other nodes, level of node is topological depth from input nodes in the prefix graph (e.g., the number of nodes **320** between it and a respective input **315**) and fanout of a node refers to the number of nodes **320** that depend from it. Over time, the machine learning model **240** via reinforcement learning techniques is trained to take actions **210** that most optimize the design of the data path circuit **108**—e.g., reduce the area, power and/or delay of the data path circuit **108**.

(80) In some embodiments, the agent **202** can be configured to separately train different instances of the machine learning model **240** for each unique circuit that the agent **202** is to optimize. For example, the agent **202** can train the machine learning model **240** for a 32 bit adder circuit and separately train another instance of the machine learning model **240** for a 64 bit adder circuit. In some embodiments, different instances of the machine learning model **240** can be trained for each respective property—e.g., trained for each target delay time of the data path circuit to reduce the area as much as possible. By utilizing reinforcement learning with the machine learning model, the design of the data path circuit can be more optimized than compared with other solutions—e.g., the delay, area, and/or power of the data path circuit can be reduced. Accordingly, the overall performance of any system (e.g., a memory system) that incorporates the data path circuit can increase as the data path circuits consume less power, have less delay, and take up less area in the memory system.

(81) In some examples, the agent **202** and/or environment **204** can use reinforcement learning to train the machine learning model **240** to take actions **210** that optimize the area, delay, and/or power of the data path circuit **108** as described herein. In some embodiments, the machine learning model **240** can determine whether to add or remove a node **320** based on programmed constraints or values. For example, the machine learning model **240** can be constrained by a maximum circuit width, a maximum arrival time of data path circuit, a weight between area and delay (or area and power, power and delay, power and area, or any combination thereof), technology library, synthesis tool, circuit generator options, and/or other target parameters for data path circuit. Such constraints may be provided by environment **204** and used to determine reward **214** in embodiments.

(82) In an example, the machine learning model **240** can be trained not to add or remove nodes that cause the delay of the data path circuit to not exceed a maximum circuit delay. In some embodiments, the machine

learning model **240** can be constrained by an arrival time of the inputs **315** of the data path circuit—e.g., some delay associated with receiving the inputs **315** rather than a delay associated with the data path circuit itself. In some embodiments the machine learning model **240** can be constrained by a target parameter. For example, the machine learning model **240** can be trained to modify (e.g., design) the data path circuit until it reaches a given area, power consumption and/or delay. In another example, the machine learning model **240** can be trained to modify (e.g., design) the data path circuit to optimize for weighted balance between improvements in area, power consumption, and delay objectives using weight parameters. In other embodiments, the machine learning model **240** can be trained to modify the data path circuit for a given area, a given power consumption, and/or a given delay.

(83) In some embodiments, the machine learning model **240** can be constrained by circuit generator options. That is, there can be different ways to generate a data path circuit for same parallel prefix graphs—e.g., the predicted physical implementation of the data path circuit can change based on which options are selected to predict the physical implementation. For example, for an adder data path circuit, the prefix adder circuit generally performs the function:

$$O(n+1)=A(n)+B(n)+C(1)$$

Where “O” is the output having length $n+1$ bits, “A” and “B” are binary inputs having length “n” bits, and “C” is an input having a length one (1) bit. In some embodiments, the circuit generator can different options that cause different functions and results. For example, the following functions are possible: 1) Not having a carry-in (e.g., no “C” value): $O(n+1)=A(n)+B(n)$; 2) Not having a carry-out: $O(n)=A(n)+B(n)+C(1)$; 3) Not having a carry-in or carry out: $O(n)=A(n)+B(n)$; 4) When one or more bits of the inputs “A/B/C” are fixed at values one (1) and zero(0);

In some embodiments, the system **200** can select between any of the options listed above. In such examples, the agent **202** or environment **204** can generate different types of circuits from a same parallel prefix graph based on which option is selected—e.g., the same prefix graph can result in a different area, delay, power consumption, or any combination thereof based on the selected settings for the circuit generator. Accordingly, the agent **202** is trained based on the specific settings and options selected for the data path circuit. In some embodiments, the prefix adder generation can also be affected by a choice of recurrence (e.g., Weinberger or Ling) or a choice of bitwise carry propagation logic (XOR or OR). In some embodiments, the various options for the circuit generator listed above can apply to adder circuits—e.g., not to other data path circuits such as priority encoder circuits.

(84) In some embodiments, the system **200** can choose to include connected logic in the data path circuit **108**. For example, the environment **204** can use a circuit generator that generates a prefix circuit corresponding to a prefix graph along with additional circuit logic connected to the inputs and/or outputs of the prefix circuit. In such an example, the agent **202** is trained based on the specific additional circuit logic that is connected to the prefix circuit. In such an example, the prefix graph assessor **242** may assess the area, power consumption and/or delay of the prefix circuit and the additional circuit logic combined.

(85) In other embodiments, the machine learning model **240** can be constrained by a weighted balance between optimizing area, power consumption and/or delay. For example, prefix graph assessor may use a weighted balance between estimated area, power consumption and/or delay when computing a reward **214**. That is, as described above, modifying a data path circuit to reduce an area, or power, or delay associated with the data path circuit can cause a different property of the data path circuit to increase—e.g., reducing delay can cause the area of the data path circuit to increase. For example, modifying the prefix graph **305** to generate prefix graph **310** can cause the delay of the data path circuit to decrease but the area of the data path circuit to increase. Accordingly, the machine learning model **240** can be trained to optimize the data path circuit according to respective weights assigned to delay and area.

(86) In an example, FIG. 6 illustrates possible weighted values on a curve between delay and area for a data path circuit. In some embodiments, a weighted value **610** represents a weight value of one (1). In such embodiments, rewards can be determined that cause the machine learning model **240** to seek to optimize (minimize) just the delay—e.g., the reinforcement learning trains the machine learning model **240** to take actions **210** that reduce the delay without regard to the resultant decrease or increase in the area of the data path circuit. In some embodiments, a weighted value **602** represents a weight value of zero (0). In such embodiments, the machine learning model **240** be trained to optimize (minimize) just the area of the data path circuit—e.g., the reinforcement learning trains the machine learning model **240** to take actions that **210** that reduce the area without regard to the resultant increase or decrease in the delay of the data path circuit. Weighted value **608** can represent optimizing more for delay than for area while weighted value **604** can represent optimizing more for area than for delay. In some embodiments, weighted value **606** can represent optimizing for delay and area

equally.

(87) In some embodiments, the machine learning model **240** and/or environment **204** can use the following formula to determine how much to optimize between delay and area:

$$(88) \begin{bmatrix} Q_{\text{area}}(s_{t+1}, a) \\ Q_{\text{delay}}(s_{t+1}, a) \end{bmatrix} \times \begin{bmatrix} w \\ 1 - w \end{bmatrix}$$

wherein Q_{area} is associated with optimizing the area of the data path circuit, Q_{delay} is associated with optimizing the delay of the data path circuit, s_{t+1} indicates the next design state **212**, “a” indicates the action taken (e.g., which node **320** was added or removed from the prefix graphs) and “w” represents a predetermined weight. In such embodiments, the weight “w” can have a value between zero (0) and one (1). By utilizing the weighted balance, the machine learning model **240** can be trained to optimize for delay, or area, or power, or any combination thereof. That is, the environment can be updated **204** to train the machine learning model **240** to modify the design for a data path circuit such that the delay for the data path circuit is the same while the area is reduced or vice versa. Although FIG. 6 illustrates the balance between optimizing delay and area, a weighted balance between area and power or delay and power is also possible. In some embodiments, the weighted balance between area, delay, and power is also possible—e.g., the graph **600** can be a three-dimensional graph that represents how much the machine learning model **240** and/or environment **204** should weigh optimizing between power, delay, and area.

(89) In some embodiments, a grid representation is used to represent a prefix graph, as shown in FIG. 4. Use of the grid representation of the prefix graph enables the prefix graph to be processed by machine learning models. The grid representation is a concise representation of prefix graphs in a grid (e.g., in a two-dimensional grid). Each row and each column of the grid representation of a prefix graph may be associated with a different input of the prefix graph, and each intersection of a row and column may be associated with a node of the prefix graph. In one embodiment, the state space S of the system **200** consists of all valid N -input prefix graphs. N -input graphs can be represented in a $N \times N$ grid with rows representing MSB and columns representing LSB. Note that in embodiments the input nodes (MSB=LSB) lie on the diagonal, output nodes will lie on the first column (LSB=0) and locations above the diagonal (LSB>MSB) cannot contain a node. In embodiments, the remaining $(N-1)(N-2)/2$ locations where non-input/output nodes may or may not exist define the $O(2^{\text{sup.}}(N-1)(N-2)/2) = O(2^{\text{sup.}}N \cdot 2^{\text{sup.}}2)$ state space of N -input prefix graphs. For example, 32-input graphs may have a state space of $|S| = O(2^{\text{sup.}}465)$, where the exact value is lower due to some of the possible combinations of nodes not being valid.

(90) With reference to FIGS. 3-4, each value of the grid can represent an input **315** (e.g., an input node) or a potential node **320** on the prefix graph **305**. In such embodiments, each node **320** can have a range (e.g., position) on the grid **405** that corresponds to its location on the prefix graph **305**. Each range can include a row index as its first element and a column index as its second element. For example, node **320-a** can be represented by the range **420-a** (1,0)—e.g., the node **320-a** receives a first input (e.g., 1 or **315-b**) and a second input (e.g., 0 or **315-a**).

(91) As described with reference to FIG. 3, each output **325** of the prefix graph receives values from directly or indirectly every previous input **315**. A valid prefix graph will have nodes **320** that are associated with contiguous ranges—e.g., (2,0) or (3,1) where the node **320** receives inputs at index 0-2 or 1-3 respectively. In such embodiments, ranges where the first element is less than the second element are not possible—e.g., (0,3) is not a possible range for a node.

(92) In some embodiments, the machine learning model **240** can receive a grid representation of an initial state of a data path circuit, and may select a node **320** for addition or removal from a prefix graph. The machine learning model **240** may then output a grid representation of nodes for addition or removal. For example, a machine learning model **240** for a Q network may populate the output grid representation with Q values at every node position. The Q value at any position will correspond to the Q value for the addition action (if the node does not exist) or removal action (if the node exists) action for a node **320** corresponding to that position. In some embodiments, the machine learning model **240** may output multiple grids representations of the same dimensions. For example, if the machine learning model **240** is a Q network that is optimizing for area and delay of a data path circuit **108**, it may output a grid representation for $Q_{\text{sub.area}}$ and a grid representation for $Q_{\text{sub.delay}}$. In such an example, the action with the highest weighted Q value will be chosen to add or remove a node **320** from the prefix graph.

(93) In some embodiments, the machine learning model **240** can receive the grid representation of the current design state **206**. For example, the machine learning model **240** can receive the generated grid **405** when the current design state **206** is represented by prefix graph **305**. In some embodiments, the machine learning model **240** can modify the prefix graph **305** by selecting a node to add or remove from the grid **405**. In one

embodiment, the machine learning model **240** can select to add a node **420-d** (e.g., (3,2)) to grid **405** to generate a grid **410**. In at least some embodiments, the grid **410** can represent a modification to the prefix graph **305**. For example, by adding the range **420-d** to the grid for the prefix graph, the prefix graph **305** can be modified to generate a grid for prefix graph **310** with an additional node **320-d**.

(94) In embodiments, the agent **202** (e.g., the machine learning model **240** of the agent **202**) can output the action. The output action may be an action to update the prefix graph (e.g., a new node to add or an existing node to delete). In an example, a grid representation in which a node to be added at (3,2) is output. The environment **204** receives the output action **210** (e.g., the addition action along with the node location (3,2) in the grid representation). The environment may be configured to operate on prefix graphs, on grid representations of prefix graphs, or on other representations of prefix graphs. In some embodiments, the environment **204** receives a node location on the grid representation of a prefix graph (e.g., (3,2)) and generates the prefix graph **310** after updating to the grid (e.g., grid **410**) from the grid of the previous state (e.g., grid **405**).

(95) As described elsewhere in the present disclosure, the machine learning model **240** can be trained using reinforcement learning so that the action **210** taken optimizes the data path circuit for area, power consumption and/or delay. That is, the machine learning model **240** can be trained so that it, for example, determines to add node **320-d** having a range (3,2) to optimize one or more property associated with the data path circuit—e.g., so it reduces the delay, area, power, and/or any combination thereof according to a weighted constraint. As described elsewhere, in some embodiments the certain parameters in the reinforcement learning algorithm such as the discount factor can be configured to train the machine learning model **240** be trained to choose a sequence of actions over time optimize one or more property associated with the data path circuit **108**.

(96) FIG. 7 illustrates an example diagram **700** of data path circuit design using reinforcement and machine learning, in accordance with at least one embodiment. The reinforcement learning shown in diagram **700** can be performed by processing logic comprising hardware, software, firmware, or any combination thereof. In at least one embodiment, operations shown in diagram **700** include operations performed by agent **202**, machine learning model **240**, and/or environment **204** as described with reference to FIG. 2. Although shown in a particular sequence or order, unless otherwise specified, the order of the processes can be modified. Thus, the illustrated embodiments should be understood only as examples, and the illustrated processes can be performed in a different order, and some processes can be performed in parallel. Additionally, one or more processes can be omitted in various embodiments. Thus, not all processes are required in every embodiment.

(97) At operation **705**, an agent **202** can receive a current design state **206** of a data path circuit. In some examples, the agent **202** can receive a prefix graph corresponding to the current design state **206** of the data path circuit. In at least one embodiment, the agent **202** can receive or generate a grid representation of the prefix graph as described with reference to FIG. 4. In some embodiments, the agent **202** can send the grid representation of the current design state **206** to the machine learning model **240**. In some embodiments, the agent **202** can also receive a reward (e.g., reward **208** as described with reference to FIG. 2). In such embodiments, the agent **202** can train the machine learning model **240** with reinforcement techniques based on the received reward. For example, if the reward indicates the previous action taken by the machine learning model **240** resulted in a reduction of an area, delay, or power associated with the data path circuit, the agent can train the machine learning model **240** to continue taking similar actions.

(98) At operation **710**, the machine learning model **240** can receive the grid representation of the parallel prefix circuit. In some embodiments, the machine learning model **240** can select an action to modify the grid representation of the prefix graph. In some embodiments, the machine learning model **240** can select an action that adds or removes a node from the prefix graph—e.g., the machine learning model can select a range on the grid representation that corresponds to adding or removing a node on the prefix graph as described with reference to FIG. 2 and FIG. 4.

(99) At operation **715**, the agent **202** can transmit the action (e.g., adding a node at (3,2)) to the environment **204**.

(100) At operation **720**, the environment **204** can calculate a next reward **214** based on the modification to the grid representation of the prefix graph—e.g., the environment **204** can modify the grid representation based on the action received from the agent **202** and then modify the prefix graph accordingly. In some embodiments, the environment **204** can first determine if the received action causes the design state **206** to be valid or invalid as described with reference to FIG. 2. In some embodiments, if the environment **204** determines the action causes the state to be invalid, the environment **204** can validate the state by adding or removing invalid nodes to/from the prefix graph. In some embodiments, the environment **204** can generate a prefix graph representing the next design state **212** after determining the modified state after applying action received from the agent **202** is valid or after validating the modified state after applying the action received from the agent **202**. In some

embodiment, the environment **204** can calculate the area, delay, and/or power for the current design state **206** and the modified next design state **212** after generating the prefix graph representing the next design state **212**. In one embodiment, the environment **204** can calculate the next reward **214** by performing circuit synthesis on each of the design state **206** and the next design state **212**. In such embodiments, the environment **204** can determine the area, delay, and/or power for the design state **206** and for the next design state **212** and determine a difference between the two to determine the next reward **214** as described with reference to FIG. 5.

(101) In some embodiments, the environment **204** can include a metrics predictor model (e.g., a second machine learning model or second model) to calculate the next reward **214**. In such embodiments, the metrics predictor model can be trained to predict the delay, area, and/or power for each state of the data path circuit—e.g., for the design state **206** and for the next design state **212**. For example, initially, the environment **204** can determine the area, power, and/or delay of the modified data path circuit by using the circuit synthesis. A database can store each state generated by the environment, and the area, power, and/or delay metrics associated with a respective state. For example, the database can store the design state **206**, the next design state **212**, and the area, power, and/or delay for the design state **206** and the next design state **212**. The metrics predictor model can process the data stored at the database and be trained to predict an area, power, and delay for a respective state. That is, the metrics predictor model can be trained to receive a state as an input and predict the area, delay, and/or power associated with the state based on processing the data stored at the database—e.g., based on processing previous states and their respective delay, area, and power. The metrics predictor model may be trained to receive a grid representation of a prefix graph as an input in embodiments. Accordingly, when the metrics predictor model is trained, the environment can send the prefix graph (or the grid representation) associated with a given state to the metrics predictor model. In such embodiments, the metrics predictor model can predict the delay, area, and/or power for the respective state. To calculate the reward, the environment **204** can find the difference between the predicted delay, area, and/or powers output by the metrics predictor model. For example, the environment **204** can provide the design state **206** and the next design state **212** to the metrics predictor model. In such embodiments, the metrics predictor model can predict the delay, area, and/or power for the design state **206** and for the next design state **212**. The environment **204** can then find the difference between the delay, area, and/or power for the next design state **212** and the delay, area, and/or power of the design state **206** to determine the next reward **214**. In some embodiments, the environment **204** can then send the reward **214** back to the agent **202**. In such embodiments, the agent **202** can utilize the next reward **214** to train the machine learning model **240** via the reinforcement learning techniques as described herein. In some embodiments, using machine learning for predicting the area, delay, and power can take less resources and consume less time than performing a circuit synthesis.

(102) FIG. 8 illustrates an example system **800** that performs reinforcement learning using a machine learning model **240**, according to at least one embodiment. In some embodiments, system **800** can include an agent (e.g., an actor, a circuit modifier, etc.) **202** and an environment (e.g., a simulation environment, a circuit synthesizer, etc.) **204** as described with reference to FIG. 2. In some embodiments, the system **800** can include a parallel circuit synthesis **810**, a database **810**, and an optimizer **825**. In some embodiments, the agent **202** can include the machine learning model **240** as described with reference to FIG. 2. In some examples, the optimizer **825** can be included in the agent **202**. In at least one embodiment, system **200** can be utilized to design a data path circuit **108** as described with reference to FIG. 1.

(103) As described with reference to FIG. 2, an agent **202** can be configured to select an action to modify a state of a data path circuit. In some embodiments, the environment **204** can be configured to apply the modification, validate the modified state if needed and return a next state (e.g., a new state **805**) to the agent **202**. In some examples, the environment **204** can also convert the state generated after applying action **210** taken by the agent **202** into a prefix graph—e.g., convert the grid generated by the applying the action from agent **202** into a prefix graph as described with reference to FIG. 4. In some embodiments, the system **800** can continue this process until a design of the data path circuit is optimized as described with reference to FIG. 2. In some embodiments, the system **800** can include multiple agents **202** and multiple environments **204** to perform the process.

(104) In at least one embodiment, system **800** can calculate the reward for a transition from an initial state **805** to a new state **805** separately from applying the action **805** from the agent **202**. That is, system **800** illustrates an alternative method to utilize reinforcement learning to design a data path circuit **108** as compared with the method illustrated in FIG. 2 and FIG. 7.

(105) For example, in one embodiment, the agent **202** can send the action taken (e.g., adding or removing a node) to both the environment **204** and to a database **820**. Similarly, environment **204** can send a current state to the agent **202** as well to the database **820** and to a parallel circuit synthesis **810**. In such embodiments, the environment **204** and agent **202** can cycle through actions and states without waiting on a circuit synthesis **810**.

to convert the prefix graph to a predicted physical data path circuit to determine the reward—e.g., the agent **202** and environment **204** can cycle through states and actions in a shorter duration.

(106) In some embodiments, parallel circuit synthesis **810** can include one or more CPUs that synthesize the prefix graph into a predicted physical data path circuit. In such embodiments, the parallel circuit synthesis **810** can also calculate a reward for each state **805** received from the environment **204**—e.g., for each prefix graph associated with a respective state **805**. That is, the parallel circuit synthesis **810** can determine an area, power, and delay for each prefix graph and state **805** received from the environment **204**. In some embodiments, the parallel circuit synthesis **810** can calculate multiple rewards **815** and synthesize multiple prefix graphs concurrently—e.g., each CPU included in the parallel circuit synthesis **810** can calculate a reward **815** for a different state **805**. In some examples, utilizing parallel circuit synthesis **810** can reduce a time to calculate a reward **815** for each state **805**. In some embodiments, the parallel circuit synthesis **810** can send the calculated rewards **815** for each state **805** to the database **820**.

(107) In some embodiments, database **820** can be configured to store states **805**, actions **210**, and rewards **815**—e.g., a reward **815** for each action **210** taken on a transition from a first state **805** to a second state **805**. In some embodiments, the database can send the states **805**, actions **210**, and rewards **815** to the database—e.g., send the first state **805**, the second state **805**, the respective action **210** that was used to go from the first state **805** to the second state **805**, and the reward calculated for the modification. In some embodiments, the optimizer **825** can access the database for the states **805**, actions **210**, and rewards **815**.

(108) In some embodiments, optimizer **825** can receive states **805**, actions **210**, and rewards **815** from the database. For example, the optimizer **825** can receive the first state **805**, the second state **805**, the respective action **210** that was used to go from the first state **805** to the second state **805**, and the reward **815** calculated for the modification. The optimizer **825** can determine if the action **210** taken resulted in a decrease in the area, delay, or power of the data path circuit. In some examples, the optimizer **825** can train the machine learning model **240** in response to determining if the action **210** taken reduced the area, power, and/or delay of the data path circuit **108**. For example, if the optimizer **825** determines the action **210** failed reduce the area, delay, and/or power of the data path circuit, the optimizer **825** can update or train the machine learning model **240** to avoid retaking such actions **210**. Accordingly, the machine learning model **240** can be trained using reinforcement learning. In such embodiments, the updated machine learning model **240** can take different actions **210**—e.g., cause the agent **202** to take different actions **210** in response to being trained or updated. In that, initially the agent **202** can take multiple actions **802** using an initial machine learning algorithm. While the agent **202** cycles through multiple actions, the parallel circuit synthesis **810** can concurrently calculate the rewards **815** for each action **210** taken. Because the synthesis takes longer than performing the action **210**, the agent **202** can continue to cycle several actions **210** before any reward **815** is calculated. In some embodiments, when the reward **815** is calculated, the database can send the information to the optimizer **825**. Accordingly, the machine learning model **240** can be updated and cause the agent to take actions **210** using the updated machine learning algorithm. The agent **202** can then cycle through actions **210** using the updated machine learning algorithm until additional rewards **815** are calculated and used by the optimizer **825** to update the machine learning algorithm a second time. The system **800** can continue using the method—e.g., calculating rewards independently of the agent **202** and environment **204** action **210** cycles while periodically updating the machine learning model **240**—until an optimal target parameter for the data path circuit **108** is satisfied.

(109) FIG. 9A illustrates a flow diagram of a method **900** for designing a data path circuit with reinforcement learning. The method **900** can be performed by processing logic comprising hardware, software, firmware, or any combination thereof. In at least one embodiment, the method **900** is performed by system **200** as described with reference to FIG. 2. Although shown in a particular sequence or order, unless otherwise specified, the order of the processes can be modified. Thus, the illustrated embodiments should be understood only as examples, and the illustrated processes can be performed in a different order, and some processes can be performed in parallel. Additionally, one or more processes can be omitted in various embodiments. Thus, not all processes are required in every embodiment. Other diagrams illustrating a method design a data path circuit are possible.

(110) At operation **905**, the processing logic can receive a first design state of a data path circuit. In one embodiment, the data path circuit can comprise a parallel prefix circuit. In at least one embodiment, the processing logic can generate a first parallel prefix graph that represents the first design state in response to receiving the first design state. That is, the processing logic can represent the first design state of the parallel prefix circuit utilizing a first parallel prefix graph. In at least one embodiment, the processing logic can generate a grid representation of the first parallel prefix graph.

(111) At operation **910**, the processing logic can input the first design state of the data path circuit into a machine learning model. In at least one embodiment, the processing logic inputs the grid representation of the first

parallel prefix graph into the machine learning model after generating the grid representation.

(112) At operation **915**, the processing logic performs reinforcement learning using the machine learning model to cause the machine learning model to output a final design state of the data path circuit. The final design state may be achieved after multiple iterations of reinforcement learning, where for each iteration a different design state is generated and assessed, and where for each iteration the machine learning model is trained to produce new design states that are improved over previous design states. In some embodiments, the final design state of the data path circuit is associated with a final parameter value that is closer to a target parameter value associated with the first design state. That is, the processing logic can use reinforcement learning to optimize the design of the data path circuit as described with reference to FIG. 2. In some embodiments, the first parameter value and the final parameter value can represent a prediction of an area associated with the data path circuit, a delay associated with the data path circuit, a power consumption associated with the data path circuit, or any combination thereof. Accordingly, the final design state of the data path circuit is associated with the final parameter value having an area, delay, or power consumption that is less than the first parameter value. In some embodiments, the first parameter value and the second parameter value can represent a prediction of a weighted value—e.g., a weighted value as described with reference to FIG. 6. In some embodiments, the machine learning model outputs an action that causes the processing logic to construct a grid representation of a final prefix graph that represents the final design state. That is, the machine learning model can output the action and the processing logic can utilize the action to generate a grid representation and convert the grid representation to a prefix graph that represents the final design state. In some embodiments, the machine learning model can iteratively modify the design state of the data path circuit from the first design state to the final design state, where during each iteration the machine learning model removes or adds a node of a graph of the data path circuit, each node of the graph associated with one or more components of the data path circuit—e.g., with one or more logical gates of the data path circuit.

(113) FIG. 9B illustrates a flow diagram of a method **902** for designing a data path circuit with reinforcement learning. The method **902** can be performed by processing logic comprising hardware, software, firmware, or any combination thereof. In at least one embodiment, the method **902** is performed by system **200** as described with reference to FIG. 2. Although shown in a particular sequence or order, unless otherwise specified, the order of the processes can be modified. Thus, the illustrated embodiments should be understood only as examples, and the illustrated processes can be performed in a different order, and some processes can be performed in parallel. Additionally, one or more processes can be omitted in various embodiments. Thus, not all processes are required in every embodiment. Other diagrams illustrating a method design a data path circuit are possible.

(114) At operation **920** of method **902**, the processing logic can receive a design state of a data path circuit. In some embodiments, the processing logic can receive a first design state of the data path circuit. In some embodiments, the processing logic can generate a parallel prefix graph that represents the design state—e.g., a first parallel prefix graph that represents the first design state. In some embodiments, the processing logic can generate a grid representation of the parallel prefix graph—e.g., a first grid representation of the first parallel prefix graph. In some embodiments, the processing logic can input the grid representation to the machine learning model.

(115) At operation **925**, the processing logic can process the design state of the data path circuit using the machine learning model to generate a next design state—e.g., a second design state. In some embodiments, the machine learning model outputs the action that the environment applies to the design state of the data path circuit to generate the second design state—e.g., the modified design state. In some embodiments, the machine learning model can output the node in the grid representation to add to or delete from the first design state to generate the second design state. That is, the machine learning model can output an action the processing logic can utilize to construct a grid representation of a second graph that is associated with the second design state—e.g., a second grid representation of a second parallel prefix graph associated with the next design state. In some embodiments, the processing logic can convert the second grid representation of the graph to the second parallel prefix graph.

(116) At operation **928**, the processing logic can determine whether the second design state is valid. If the second design state is invalid, the method proceeds to operation **929**. If the second design state is determined to be valid, the method continues to operation **930**. At operation **929**, processing logic modifies the second design state of the data path circuit to produce an updated second design state that is valid—e.g., validates the parallel prefix graph by adding and/or removing additional nodes. In some embodiments, the processing logic can convert the first and second design states to data path circuit implementations. The method then continues to operations **930**.

(117) At operation **930**, the processing logic can determine a first parameter value for the first design state and

determine a second parameter value for the second design state—e.g., determine the first parameter value for the first design state and determine the second parameter value for the second design state. In some embodiments, the processing logic can determine the area, delay, and power associated with the design state and the next design state—e.g., the first parameter value and the second parameter value can represent the area, delay, and power. Processing logic can then determine the first parameter value based on a first area, delay and/or power (e.g., based on a first weighted combination of these values) and determine the second parameter value based on a second area, delay and/or power (e.g., based on a first weighted combination of these values).

(118) In some embodiments, the processing logic can process the first design state using a second model (e.g., a second machine learning model or metrics predictor model), where the second model outputs the first parameter value associated with the first design state. Similarly, the processing logic can process the second design state using the second model, where the second model outputs the second parameter value associated with the second design state—e.g., the metrics predictor model can process the first design state and second design state to output the first parameter value and the second parameter value. In such embodiments, the processing device can send the grid representation of the first graph of the first design state as a first input to the second model to receive the first parameter value and send the grid representation of the second graph of the second design state as a second input to the second model to receive the second parameter value—e.g., the metric predictor model can output the first parameter value and second parameter value based on receiving the grid representation of the first and second graphs. In some embodiments, the second model can receive a circuit implementation of the first design state as a first input to determine the first parameter value and receive a circuit implementation of the second design state as a second input to determine the second parameter value—e.g., the metrics predictor model can receive the first parallel prefix graph and the second parallel prefix graph. In some embodiments, the processing logic can process the first design state and the second design state using a circuit synthesis tool. In such embodiments, the circuit synthesis tool can output the first parameter value and the second parameter value. In some embodiments, the circuit synthesis tool can process the first design state concurrent with the machine learning model outputting the second design state of the data path circuit.

(119) At operation **935**, the processing logic can update the machine learning model—e.g., based on whether the modification selected by the machine learning model reduced or increased the area, delay, power consumption, or any combination thereof associated with the data path circuit. After updating the machine learning model, at operation **938** the processing logic can determine whether one or more stopping criteria have been met. A stopping criterion may be met, for example, after a threshold number of iterations of the design state, after one or more target goals (e.g., for area, power and/or delay) are met, after a threshold number of design state iterations have been performed without further improvement to the parameter value(s), and so on. If the stopping criterion has been met, the method may proceed to operation **945**, at which a final design state may be selected for a data path circuit. The final design state may be, for example, the most recent design state or a design state encountered during iterations that obtained the best data path circuit metrics (e.g., best area, delay, power consumption, or any combination thereof) or a parameter closest to the target. If a stopping criterion has not been met, then at operation **940** processing logic may select a next design state to be input into the machine learning model. The next design state may be, for example, a most recent design state or another design state encountered during iterations. The method may return to operation **925**, in which the next design state is processed using the machine learning model. The processing logic can utilize the updated machine learning model. For example, the processing logic can process the design state or next design state using the updated machine learning model—e.g., the processing logic can process the first design state or the second design state using the machine learning model. In some embodiments, the updated machine learning model outputs a third design state of the data path circuit that is a modification of the first design state or the second design state. In at least one embodiment, the processing logic can determine a third parameter value associated with the third design state, the third parameter value closer to the target than the first parameter value or the second parameter value. The process can repeat for fourth, fifth and more design states until the stopping criterion is met.

(120) FIG. **10A** illustrates inference and/or training logic **1015** used to perform inferencing and/or training operations associated with one or more embodiments. Details regarding inference and/or training logic **1015** are provided below in conjunction with FIGS. **10A** and/or **10B**.

(121) In at least one embodiment, inference and/or training logic **1015** may include, without limitation, code and/or data storage **1001** to store forward and/or output weight and/or input/output data, and/or other parameters to configure neurons or layers of a neural network trained and/or used for inferencing in aspects of one or more embodiments. In at least one embodiment, training logic **1015** may include, or be coupled to code and/or data storage **1001** to store graph code or other software to control timing and/or order, in which weight and/or other parameter information is to be loaded to configure, logic, including integer and/or floating point units

collectively, arithmetic logic units (ALUs). In at least one embodiment, code, such as graph code, loads weight or other parameter information into processor ALUs based on an architecture of a neural network to which the code corresponds. In at least one embodiment, code and/or data storage **1001** stores weight parameters and/or input/output data of each layer of a neural network trained or used in conjunction with one or more embodiments during forward propagation of input/output data and/or weight parameters during training and/or inferencing using aspects of one or more embodiments. In at least one embodiment, any portion of code and/or data storage **1001** may be included with other on-chip or off-chip data storage, including a processor's L1, L2, or L3 cache or system memory.

(122) In at least one embodiment, any portion of code and/or data storage **1001** may be internal or external to one or more processors or other hardware logic devices or circuits. In at least one embodiment, code and/or code and/or data storage **1001** may be cache memory, dynamic randomly addressable memory ("DRAM"), static randomly addressable memory ("SRAM"), non-volatile memory (e.g., Flash memory), or other storage. In at least one embodiment, choice of whether code and/or code and/or data storage **1001** is internal or external to a processor, for example, or comprised of DRAM, SRAM, Flash or some other storage type may depend on available storage on-chip versus off-chip, latency requirements of training and/or inferencing functions being performed, batch size of data used in inferencing and/or training of a neural network, or some combination of these factors.

(123) In at least one embodiment, inference and/or training logic **1015** may include, without limitation, a code and/or data storage **1005** to store backward and/or output weight and/or input/output data corresponding to neurons or layers of a neural network trained and/or used for inferencing in aspects of one or more embodiments. In at least one embodiment, code and/or data storage **1005** stores weight parameters and/or input/output data of each layer of a neural network trained or used in conjunction with one or more embodiments during backward propagation of input/output data and/or weight parameters during training and/or inferencing using aspects of one or more embodiments. In at least one embodiment, training logic **1015** may include, or be coupled to code and/or data storage **1005** to store graph code or other software to control timing and/or order, in which weight and/or other parameter information is to be loaded to configure, logic, including integer and/or floating point units (collectively, arithmetic logic units (ALUs)). In at least one embodiment, code, such as graph code, loads weight or other parameter information into processor ALUs based on an architecture of a neural network to which the code corresponds. In at least one embodiment, any portion of code and/or data storage **1005** may be included with other on-chip or off-chip data storage, including a processor's L1, L2, or L3 cache or system memory. In at least one embodiment, any portion of code and/or data storage **1005** may be internal or external to on one or more processors or other hardware logic devices or circuits. In at least one embodiment, code and/or data storage **1005** may be cache memory, DRAM, SRAM, non-volatile memory (e.g., Flash memory), or other storage. In at least one embodiment, choice of whether code and/or data storage **1005** is internal or external to a processor, for example, or comprised of DRAM, SRAM, Flash or some other storage type may depend on available storage on-chip versus off-chip, latency requirements of training and/or inferencing functions being performed, batch size of data used in inferencing and/or training of a neural network, or some combination of these factors.

(124) In at least one embodiment, code and/or data storage **1001** and code and/or data storage **1005** may be separate storage structures. In at least one embodiment, code and/or data storage **1001** and code and/or data storage **1005** may be same storage structure. In at least one embodiment, code and/or data storage **1001** and code and/or data storage **1005** may be partially same storage structure and partially separate storage structures. In at least one embodiment, any portion of code and/or data storage **1001** code and/or data storage **1005** may be included with other on-chip or off-chip data storage, including a processor's L1, L2, or L3 cache or system memory.

(125) In at least one embodiment, inference and/or training logic **1015** may include, without limitation, one or more arithmetic logic unit(s) ("ALU(s)") **1010**, including integer and/or floating point units, to perform logical and/or mathematical operations based, at least in part on, or indicated by, training and/or inference code (e.g., graph code), a result of which may produce activations (e.g., output values from layers or neurons within a neural network) stored in an activation storage **1020** that are functions of input/output and/or weight parameter data stored in code and/or data storage **1001** and/or code and/or data storage **1005**. In at least one embodiment, activations stored in activation storage **1020** are generated according to linear algebraic and or matrix-based mathematics performed by ALU(s) **1010** in response to performing instructions or other code, wherein weight values stored in code and/or data storage **1005** and/or code and/or data storage **1001** are used as operands along with other values, such as bias values, gradient information, momentum values, or other parameters or hyperparameters, any or all of which may be stored in code and/or data storage **1005** or code and/or data storage

1001 or another storage on or off-chip.

(126) In at least one embodiment, ALU(s) **1010** are included within one or more processors or other hardware logic devices or circuits, whereas in another embodiment, ALU(s) **1010** may be external to a processor or other hardware logic device or circuit that uses them (e.g., a co-processor). In at least one embodiment, ALUs **1010** may be included within a processor's execution units or otherwise within a bank of ALUs accessible by a processor's execution units either within same processor or distributed between different processors of different types (e.g., central processing units, graphics processing units, fixed function units, etc.). In at least one embodiment, code and/or data storage **1001**, code and/or data storage **1005**, and activation storage **1020** may be on same processor or other hardware logic device or circuit, whereas in another embodiment, they may be in different processors or other hardware logic devices or circuits, or some combination of same and different processors or other hardware logic devices or circuits. In at least one embodiment, any portion of activation storage **1020** may be included with other on-chip or off-chip data storage, including a processor's L1, L2, or L3 cache or system memory. Furthermore, inferencing and/or training code may be stored with other code accessible to a processor or other hardware logic or circuit and fetched and/or processed using a processor's fetch, decode, scheduling, execution, retirement and/or other logical circuits.

(127) In at least one embodiment, activation storage **1020** may be cache memory, DRAM, SRAM, non-volatile memory (e.g., Flash memory), or other storage. In at least one embodiment, activation storage **1020** may be completely or partially within or external to one or more processors or other logical circuits. In at least one embodiment, choice of whether activation storage **1020** is internal or external to a processor, for example, or comprised of DRAM, SRAM, Flash or some other storage type may depend on available storage on-chip versus off-chip, latency requirements of training and/or inferencing functions being performed, batch size of data used in inferencing and/or training of a neural network, or some combination of these factors. In at least one embodiment, inference and/or training logic **1015** illustrated in FIG. **10A** may be used in conjunction with an application-specific integrated circuit ("ASIC"), such as Tensorflow® Processing Unit from Google, an inference processing unit (IPU) from Graphcore™, or a Nervana® (e.g., "Lake Crest") processor from Intel Corp. In at least one embodiment, inference and/or training logic **1015** illustrated in FIG. **10A** may be used in conjunction with central processing unit ("CPU") hardware, graphics processing unit ("GPU") hardware or other hardware, such as data processing unit ("DPU") hardware, or field programmable gate arrays ("FPGAs").

(128) FIG. **10B** illustrates inference and/or training logic **1015**, according to at least one or more embodiments. In at least one embodiment, inference and/or training logic **1015** may include, without limitation, hardware logic in which computational resources are dedicated or otherwise exclusively used in conjunction with weight values or other information corresponding to one or more layers of neurons within a neural network. In at least one embodiment, inference and/or training logic **1015** illustrated in FIG. **10B** may be used in conjunction with an application-specific integrated circuit (ASIC), such as Tensorflow® Processing Unit from Google, an inference processing unit (IPU) from Graphcore™, or a Nervana® (e.g., "Lake Crest") processor from Intel Corp. In at least one embodiment, inference and/or training logic **1015** illustrated in FIG. **10B** may be used in conjunction with central processing unit (CPU) hardware, graphics processing unit (GPU) hardware or other hardware, such as data processing unit ("DPU") hardware, or field programmable gate arrays (FPGAs). In at least one embodiment, inference and/or training logic **1015** includes, without limitation, code and/or data storage **1001** and code and/or data storage **1005**, which may be used to store code (e.g., graph code), weight values and/or other information, including bias values, gradient information, momentum values, and/or other parameter or hyperparameter information. In at least one embodiment illustrated in FIG. **10B**, each of code and/or data storage **1001** and code and/or data storage **1005** is associated with a dedicated computational resource, such as computational hardware **1002** and computational hardware **1006**, respectively. In at least one embodiment, each of computational hardware **1002** and computational hardware **1006** comprises one or more ALUs that perform mathematical functions, such as linear algebraic functions, only on information stored in code and/or data storage **1001** and code and/or data storage **1005**, respectively, result of which is stored in activation storage **1020**.

(129) In at least one embodiment, each of code and/or data storage **1001** and **1005** and corresponding computational hardware **1002** and **1006**, respectively, correspond to different layers of a neural network, such that resulting activation from one "storage/computational pair **1001/1002**" of code and/or data storage **1001** and computational hardware **1002** is provided as an input to "storage/computational pair **1005/1006**" of code and/or data storage **1005** and computational hardware **1006**, in order to mirror conceptual organization of a neural network. In at least one embodiment, each of storage/computational pairs **1001/1002** and **1005/1006** may correspond to more than one neural network layer. In at least one embodiment, additional storage/computation pairs (not shown) subsequent to or in parallel with storage computation pairs **1001/1002** and **1005/1006** may be

included in inference and/or training logic **1015**.

(130) FIG. **11** illustrates an example data center **1100**, in which at least one embodiment may be used. In at least one embodiment, data center **1100** includes a data center infrastructure layer **1110**, a framework layer **1120**, a software layer **1130**, and an application layer **1240**.

(131) In at least one embodiment, as shown in FIG. **11**, data center infrastructure layer **1110** may include a resource orchestrator **1112**, grouped computing resources **1114**, and node computing resources (“node C.R.s”) **1116(1)-1116(N)**, where “N” represents any whole, positive integer. In at least one embodiment, node C.R.s **1116(1)-1116(N)** may include, but are not limited to, any number of central processing units (“CPUs”) or other processors (including accelerators, field programmable gate arrays (FPGAs), data processing units, graphics processors, etc.), memory devices (e.g., dynamic read-only memory), storage devices (e.g., solid state or disk drives), network input/output (“NW I/O”) devices, network switches, virtual machines (“VMs”), power modules, and cooling modules, etc. In at least one embodiment, one or more node C.R.s from among node C.R.s **1116(1)-1116(N)** may be a server having one or more of above-mentioned computing resources.

(132) In at least one embodiment, grouped computing resources **1114** may include separate groupings of node C.R.s housed within one or more racks (not shown), or many racks housed in data centers at various geographical locations (also not shown). Separate groupings of node C.R.s within grouped computing resources **1114** may include grouped compute, network, memory or storage resources that may be configured or allocated to support one or more workloads. In at least one embodiment, several node C.R.s including CPUs or processors may grouped within one or more racks to provide compute resources to support one or more workloads. In at least one embodiment, one or more racks may also include any number of power modules, cooling modules, and network switches, in any combination.

(133) In at least one embodiment, resource orchestrator **1112** may configure or otherwise control one or more node C.R.s **1116(1)-1116(N)** and/or grouped computing resources **1114**. In at least one embodiment, resource orchestrator **1112** may include a software design infrastructure (“SDI”) management entity for data center **1100**. In at least one embodiment, resource orchestrator may include hardware, software or some combination thereof.

(134) In at least one embodiment, as shown in FIG. **11**, framework layer **1120** includes a job scheduler **1122**, a configuration manager **1124**, a resource manager **1126** and a distributed file system **1128**. In at least one embodiment, framework layer **1120** may include a framework to support software **1132** of software layer **1130** and/or one or more application(s) **1142** of application layer **1140**. In at least one embodiment, software **1132** or application(s) **1142** may respectively include web-based service software or applications, such as those provided by Amazon Web Services, Google Cloud and Microsoft Azure. In at least one embodiment, framework layer **1120** may be, but is not limited to, a type of free and open-source software web application framework such as Apache Spark™ (hereinafter “Spark”) that may utilize distributed file system **1128** for large-scale data processing (e.g., “big data”). In at least one embodiment, job scheduler **1122** may include a Spark driver to facilitate scheduling of workloads supported by various layers of data center **1100**. In at least one embodiment, configuration manager **1124** may be capable of configuring different layers such as software layer **1130** and framework layer **1120** including Spark and distributed file system **1128** for supporting large-scale data processing. In at least one embodiment, resource manager **1126** may be capable of managing clustered or grouped computing resources mapped to or allocated for support of distributed file system **1128** and job scheduler **1122**. In at least one embodiment, clustered or grouped computing resources may include grouped computing resource **1114** at data center infrastructure layer **1110**. In at least one embodiment, resource manager **1126** may coordinate with resource orchestrator **1112** to manage these mapped or allocated computing resources.

(135) In at least one embodiment, software **1132** included in software layer **1130** may include software used by at least portions of node C.R.s **1116(1)-1116(N)**, grouped computing resources **1114**, and/or distributed file system **1128** of framework layer **1120**. The one or more types of software may include, but are not limited to, Internet web page search software, e-mail virus scan software, database software, and streaming video content software.

(136) In at least one embodiment, application(s) **1142** included in application layer **1140** may include one or more types of applications used by at least portions of node C.R.s **1116(1)-1116(N)**, grouped computing resources **1114**, and/or distributed file system **1128** of framework layer **1120**. One or more types of applications may include, but are not limited to, any number of a genomics application, a cognitive compute, and a machine learning application, including training or inferencing software, machine learning framework software (e.g., PyTorch, TensorFlow, Caffe, etc.) or other machine learning applications used in conjunction with one or more embodiments.

(137) In at least one embodiment, any of configuration manager **1124**, resource manager **1126**, and resource orchestrator **1112** may implement any number and type of self-modifying actions based on any amount and type

of data acquired in any technically feasible fashion. In at least one embodiment, self-modifying actions may relieve a data center operator of data center **1100** from making possibly bad configuration decisions and possibly avoiding underutilized and/or poor performing portions of a data center.

(138) In at least one embodiment, data center **1100** may include tools, services, software, or other resources to train one or more machine learning models or predict or infer information using one or more machine learning models according to one or more embodiments described herein. For example, in at least one embodiment, a machine learning model may be trained by calculating weight parameters according to a neural network architecture using software and computing resources described above with respect to data center **1100**. In at least one embodiment, trained machine learning models corresponding to one or more neural networks may be used to infer or predict information using resources described above with respect to data center **1100** by using weight parameters calculated through one or more training techniques described herein.

(139) In at least one embodiment, data center may use CPUs, application-specific integrated circuits (ASICs), GPUs, DPUs, FPGAs, or other hardware to perform training and/or inferencing using above-described resources. Moreover, one or more software and/or hardware resources described above may be configured as a service to allow users to train or performing inferencing of information, such as image recognition, speech recognition, or other artificial intelligence services.

(140) Inference and/or training logic **1015** are used to perform inferencing and/or training operations associated with one or more embodiments. Details regarding inference and/or training logic **1015** are provided below in conjunction with FIGS. **10A** and/or **10B**. In at least one embodiment, inference and/or training logic **1015** may be used in system FIG. **11** for inferencing or predicting operations based, at least in part, on weight parameters calculated using neural network training operations, neural network functions and/or architectures, or neural network use cases described herein.

(141) Such components may be used to generate synthetic data imitating failure cases in a network training process, which may help to improve performance of the network while limiting the amount of synthetic data to avoid overfitting.

(142) FIG. **12** is a block diagram illustrating an exemplary computer system, which may be a system with interconnected devices and components, a system-on-a-chip (SOC) or some combination thereof **1200** formed with a processor that may include execution units to execute an instruction, according to at least one embodiment. In at least one embodiment, computer system **1200** may include, without limitation, a component, such as a processor **1202** to employ execution units including logic to perform algorithms for process data, in accordance with present disclosure, such as in embodiment described herein. In at least one embodiment, computer system **1200** may include processors, such as PENTIUM® Processor family, Xeon™, Itanium®, XScale™ and/or StrongARM™, Intel® Core™, or Intel® Nervana™ microprocessors available from Intel Corporation of Santa Clara, Calif., although other systems (including PCs having other microprocessors, engineering workstations, set-top boxes and like) may also be used. In at least one embodiment, computer system **1200** may execute a version of WINDOWS' operating system available from Microsoft Corporation of Redmond, Wash., although other operating systems (UNIX and Linux for example), embedded software, and/or graphical user interfaces, may also be used.

(143) Embodiments may be used in other devices such as handheld devices and embedded applications. Some examples of handheld devices include cellular phones, Internet Protocol devices, digital cameras, personal digital assistants ("PDAs"), and handheld PCs. In at least one embodiment, embedded applications may include a microcontroller, a digital signal processor ("DSP"), system on a chip, network computers ("NetPCs"), set-top boxes, network hubs, wide area network ("WAN") switches, edge devices, Internet-of-Things ("IoT") devices, or any other system that may perform one or more instructions in accordance with at least one embodiment.

(144) In at least one embodiment, computer system **1200** may include, without limitation, processor **1202** that may include, without limitation, one or more execution units **1208** to perform machine learning model training and/or inferencing according to techniques described herein. In at least one embodiment, computer system **1200** is a single processor desktop or server system, but in another embodiment computer system **1200** may be a multiprocessor system. In at least one embodiment, processor **1202** may include, without limitation, a complex instruction set computer ("CISC") microprocessor, a reduced instruction set computing ("RISC") microprocessor, a very long instruction word ("VLIW") microprocessor, a processor implementing a combination of instruction sets, or any other processor device, such as a digital signal processor, for example. In at least one embodiment, processor **1202** may be coupled to a processor bus **1210** that may transmit data signals between processor **1202** and other components in computer system **1200**.

(145) In at least one embodiment, processor **1202** may include, without limitation, a Level 1 ("L1") internal cache memory ("cache") **1204**. In at least one embodiment, processor **1202** may have a single internal cache or

multiple levels of internal cache. In at least one embodiment, cache memory may reside external to processor **1202**. Other embodiments may also include a combination of both internal and external caches depending on particular implementation and needs. In at least one embodiment, register file **1206** may store different types of data in various registers including, without limitation, integer registers, floating point registers, status registers, and instruction pointer register.

(146) In at least one embodiment, execution unit **1208**, including, without limitation, logic to perform integer and floating point operations, also resides in processor **1202**. In at least one embodiment, processor **1202** may also include a microcode (“ucode”) read only memory (“ROM”) that stores microcode for certain macro instructions. In at least one embodiment, execution unit **1208** may include logic to handle a packed instruction set **1209**. In at least one embodiment, by including packed instruction set **1209** in an instruction set of a general-purpose processor **1202**, along with associated circuitry to execute instructions, operations used by many multimedia applications may be performed using packed data in a general-purpose processor **1202**. In one or more embodiments, many multimedia applications may be accelerated and executed more efficiently by using full width of a processor's data bus for performing operations on packed data, which may eliminate need to transfer smaller units of data across processor's data bus to perform one or more operations one data element at a time.

(147) In at least one embodiment, execution unit **1208** may also be used in microcontrollers, embedded processors, graphics devices, DSPs, and other types of logic circuits. In at least one embodiment, computer system **1200** may include, without limitation, a memory **1220**. In at least one embodiment, memory **1220** may be implemented as a Dynamic Random Access Memory (“DRAM”) device, a Static Random Access Memory (“SRAM”) device, flash memory device, or other memory device. In at least one embodiment, memory **1220** may store instruction(s) **1219** and/or data **1221** represented by data signals that may be executed by processor **1202**.

(148) In at least one embodiment, system logic chip may be coupled to processor bus **1210** and memory **1220**. In at least one embodiment, system logic chip may include, without limitation, a memory controller hub (“MCH”) **1216**, and processor **1202** may communicate with MCH **1216** via processor bus **1210**. In at least one embodiment, MCH **1216** may provide a high bandwidth memory path **1218** to memory **1220** for instruction and data storage and for storage of graphics commands, data and textures. In at least one embodiment, MCH **1216** may direct data signals between processor **1202**, memory **1220**, and other components in computer system **1200** and to bridge data signals between processor bus **1210**, memory **1220**, and a system I/O **1222**. In at least one embodiment, system logic chip may provide a graphics port for coupling to a graphics controller. In at least one embodiment, MCH **1216** may be coupled to memory **1220** through a high bandwidth memory path **1218** and graphics/video card **1212** may be coupled to MCH **1216** through an Accelerated Graphics Port (“AGP”) interconnect **1214**.

(149) In at least one embodiment, computer system **1200** may use system I/O **1222** that is a proprietary hub interface bus to couple MCH **1216** to I/O controller hub (“ICH”) **1230**. In at least one embodiment, ICH **1230** may provide direct connections to some I/O devices via a local I/O bus. In at least one embodiment, local I/O bus may include, without limitation, a high-speed I/O bus for connecting peripherals to memory **1220**, chipset, and processor **1202**. Examples may include, without limitation, an audio controller **1229**, a firmware hub (“flash BIOS”) **1228**, a wireless transceiver **1226**, a data storage **1224**, a legacy I/O controller **1223** containing user input and keyboard interfaces **1225**, a serial expansion port **1227**, such as Universal Serial Bus (“USB”), and a network controller **1234**, which may include in some embodiments, a data processing unit. Data storage **1224** may comprise a hard disk drive, a floppy disk drive, a CD-ROM device, a flash memory device, or other mass storage device.

(150) In at least one embodiment, FIG. **12** illustrates a system, which includes interconnected hardware devices or “chips”, whereas in other embodiments, FIG. **12** may illustrate an exemplary System on a Chip (“SoC”). In at least one embodiment, devices may be interconnected with proprietary interconnects, standardized interconnects (e.g., PCIe) or some combination thereof. In at least one embodiment, one or more components of computer system **1200** are interconnected using compute express link (CXL) interconnects.

(151) Inference and/or training logic **1015** are used to perform inferencing and/or training operations associated with one or more embodiments. Details regarding inference and/or training logic **1015** are provided below in conjunction with FIGS. **10A** and/or **10B**. In at least one embodiment, inference and/or training logic **1015** may be used in system FIG. **12** for inferencing or predicting operations based, at least in part, on weight parameters calculated using neural network training operations, neural network functions and/or architectures, or neural network use cases described herein.

(152) Such components may be used to generate synthetic data imitating failure cases in a network training

process, which may help to improve performance of the network while limiting the amount of synthetic data to avoid overfitting.

(153) FIG. 13 is a block diagram illustrating an electronic device **1300** for utilizing a processor **1310**, according to at least one embodiment. In at least one embodiment, electronic device **1300** may be, for example and without limitation, a notebook, a tower server, a rack server, a blade server, a laptop, a desktop, a tablet, a mobile device, a phone, an embedded computer, an edge device, an IoT device, or any other suitable electronic device.

(154) In at least one embodiment, system **1300** may include, without limitation, processor **1310** communicatively coupled to any suitable number or kind of components, peripherals, modules, or devices. In at least one embodiment, processor **1310** coupled using a bus or interface, such as a 1° C. bus, a System Management Bus (“SMBus”), a Low Pin Count (LPC) bus, a Serial Peripheral Interface (“SPI”), a High Definition Audio (“HDA”) bus, a Serial Advance Technology Attachment (“SATA”) bus, a Universal Serial Bus (“USB”) (versions 1, 2, 3), or a Universal Asynchronous Receiver/Transmitter (“UART”) bus. In at least one embodiment, FIG. 13 illustrates a system, which includes interconnected hardware devices or “chips”, whereas in other embodiments, FIG. 13 may illustrate an exemplary System on a Chip (“SoC”). In at least one embodiment, devices illustrated in FIG. 13 may be interconnected with proprietary interconnects, standardized interconnects (e.g., PCIe) or some combination thereof. In at least one embodiment, one or more components of FIG. 13 are interconnected using compute express link (CXL) interconnects.

(155) In at least one embodiment, FIG. 13 may include a display **1324**, a touch screen **1325**, a touch pad **1330**, a Near Field Communications unit (“NEC”) **1345**, a sensor hub **1340**, a thermal sensor **1346**, an Express Chipset (“EC”) **1335**, a Trusted Platform Module (“TPM”) **1338**, BIOS/firmware/flash memory (“BIOS, FW Flash”) **1322**, a DSP **1360**, a drive **1320** such as a Solid State Disk (“SSD”) or a Hard Disk Drive (“HDD”), a wireless local area network unit (“WLAN”) **1350**, a Bluetooth unit **1352**, a Wireless Wide Area Network unit (“WWAN”) **1356**, a Global Positioning System (GPS) **1355**, a camera (“USB 3.0 camera”) **1354** such as a USB 3.0 camera, and/or a Low Power Double Data Rate (“LPDDR”) memory unit (“LPDDR3”) **1315** implemented in, for example, LPDDR3 standard. These components may each be implemented in any suitable manner.

(156) In at least one embodiment, other components may be communicatively coupled to processor **1310** through components discussed above. In at least one embodiment, an accelerometer **1341**, Ambient Light Sensor (“ALS”) **1342**, compass **1343**, and a gyroscope **1344** may be communicatively coupled to sensor hub **1340**. In at least one embodiment, thermal sensor **1339**, a fan **1337**, a keyboard **1336**, and a touch pad **1330** may be communicatively coupled to EC **1335**. In at least one embodiment, speaker **1363**, headphones **1364**, and microphone (“mic”) **1365** may be communicatively coupled to an audio unit (“audio codec and class d amp”) **1362**, which may in turn be communicatively coupled to DSP **1360**. In at least one embodiment, audio unit **1364** may include, for example and without limitation, an audio coder/decoder (“codec”) and a class D amplifier. In at least one embodiment, SIM card (“SIM”) **1357** may be communicatively coupled to WWAN unit **1356**. In at least one embodiment, components such as WLAN unit **1350** and Bluetooth unit **1352**, as well as WWAN unit **1356** may be implemented in a Next Generation Form Factor (“NGFF”).

(157) Inference and/or training logic **1015** are used to perform inferencing and/or training operations associated with one or more embodiments. Details regarding inference and/or training logic **1015** are provided below in conjunction with FIGS. 10A and/or 10B. In at least one embodiment, inference and/or training logic **1015** may be used in system FIG. 13 for inferencing or predicting operations based, at least in part, on weight parameters calculated using neural network training operations, neural network functions and/or architectures, or neural network use cases described herein.

(158) Such components may be used to generate synthetic data imitating failure cases in a network training process, which may help to improve performance of the network while limiting the amount of synthetic data to avoid overfitting.

(159) FIG. 14 is a block diagram of a processing system, according to at least one embodiment. In at least one embodiment, system **1400** includes one or more processors **1402** and one or more graphics processors **1408**, and may be a single processor desktop system, a multiprocessor workstation system, or a server system having a large number of processors **1402** or processor cores **1407**. In at least one embodiment, system **1400** is a processing platform incorporated within a system-on-a-chip (SoC) integrated circuit for use in mobile, handheld, edge, or embedded devices.

(160) In at least one embodiment, system **1400** may include, or be incorporated within a server-based gaming platform, a game console, including a game and media console, a mobile gaming console, a handheld game console, or an online game console. In at least one embodiment, system **1400** is a mobile phone, smart phone, tablet computing device or mobile Internet device. In at least one embodiment, processing system **1400** may also include, couple with, or be integrated within a wearable device, such as a smart watch wearable device, smart

eyewear device, augmented reality device, or virtual reality device. In at least one embodiment, processing system **1400** is a television or set top box device having one or more processors **1402** and a graphical interface generated by one or more graphics processors **1408**.

(161) In at least one embodiment, one or more processors **1402** each include one or more processor cores **1407** to process instructions which, when executed, perform operations for system and user software. In at least one embodiment, each of one or more processor cores **1407** is configured to process a specific instruction set **1409**. In at least one embodiment, instruction set **1409** may facilitate Complex Instruction Set Computing (CISC), Reduced Instruction Set Computing (RISC), or computing via a Very Long Instruction Word (VLIW). In at least one embodiment, processor cores **1407** may each process a different instruction set **1409**, which may include instructions to facilitate emulation of other instruction sets. In at least one embodiment, processor core **1407** may also include other processing devices, such as a Digital Signal Processor (DSP).

(162) In at least one embodiment, processor **1402** includes cache memory **1404**. In at least one embodiment, processor **1402** may have a single internal cache or multiple levels of internal cache. In at least one embodiment, cache memory is shared among various components of processor **1402**. In at least one embodiment, processor **1402** also uses an external cache (e.g., a Level-3 (L3) cache or Last Level Cache (LLC)) (not shown), which may be shared among processor cores **1407** using known cache coherency techniques. In at least one embodiment, register file **1406** is additionally included in processor **1402** which may include different types of registers for storing different types of data (e.g., integer registers, floating point registers, status registers, and an instruction pointer register). In at least one embodiment, register file **1406** may include general-purpose registers or other registers.

(163) In at least one embodiment, one or more processor(s) **1402** are coupled with one or more interface bus(es) **1410** to transmit communication signals such as address, data, or control signals between processor **1402** and other components in system **1400**. In at least one embodiment, interface bus **1410**, in one embodiment, may be a processor bus, such as a version of a Direct Media Interface (DMI) bus. In at least one embodiment, interface **1410** is not limited to a DMI bus, and may include one or more Peripheral Component Interconnect buses (e.g., PCI, PCI Express), memory busses, or other types of interface busses. In at least one embodiment processor(s) **1402** include an integrated memory controller **1416** and a platform controller hub **1430**. In at least one embodiment, memory controller **1416** facilitates communication between a memory device and other components of system **1400**, while platform controller hub (PCH) **1430** provides connections to I/O devices via a local I/O bus.

(164) In at least one embodiment, memory device **1420** may be a dynamic random access memory (DRAM) device, a static random access memory (SRAM) device, flash memory device, phase-change memory device, or some other memory device having suitable performance to serve as process memory. In at least one embodiment memory device **1420** may operate as system memory for system **1400**, to store data **1422** and instructions **1421** for use when one or more processors **1402** executes an application or process. In at least one embodiment, memory controller **1416** also couples with an optional external graphics processor **1412**, which may communicate with one or more graphics processors **1408** in processors **1402** to perform graphics and media operations. In at least one embodiment, a display device **1411** may connect to processor(s) **1402**. In at least one embodiment display device **1411** may include one or more of an internal display device, as in a mobile electronic device or a laptop device or an external display device attached via a display interface (e.g., DisplayPort, etc.). In at least one embodiment, display device **1411** may include a head mounted display (HMD) such as a stereoscopic display device for use in virtual reality (VR) applications or augmented reality (AR) applications.

(165) In at least one embodiment, platform controller hub **1430** enables peripherals to connect to memory device **1420** and processor **1402** via a high-speed I/O bus. In at least one embodiment, I/O peripherals include, but are not limited to, an audio controller **1446**, a network controller **1434**, a firmware interface **1428**, a wireless transceiver **1426**, touch sensors **1425**, a data storage device **1424** (e.g., hard disk drive, flash memory, etc.). In at least one embodiment, data storage device **1424** may connect via a storage interface (e.g., SATA) or via a peripheral bus, such as a Peripheral Component Interconnect bus (e.g., PCI, PCI Express). In at least one embodiment, touch sensors **1425** may include touch screen sensors, pressure sensors, or fingerprint sensors. In at least one embodiment, wireless transceiver **1426** may be a Wi-Fi transceiver, a Bluetooth transceiver, or a mobile network transceiver such as a 3G, 4G, or Long Term Evolution (LTE) transceiver. In at least one embodiment, firmware interface **1428** enables communication with system firmware, and may be, for example, a unified extensible firmware interface (UEFI). In at least one embodiment, network controller **1434** may enable a network connection to a wired network. In at least one embodiment, a high-performance network controller (not shown) couples with interface bus **1410**. In at least one embodiment, audio controller **1446** is a multi-channel

high definition audio controller. In at least one embodiment, system **1400** includes an optional legacy I/O controller **1440** for coupling legacy (e.g., Personal System 2 (PS/2)) devices to system. In at least one embodiment, platform controller hub **1430** may also connect to one or more Universal Serial Bus (USB) controllers **1442** connect input devices, such as keyboard and mouse **1443** combinations, a camera **1444**, or other USB input devices.

(166) In at least one embodiment, an instance of memory controller **1416** and platform controller hub **1430** may be integrated into a discreet external graphics processor, such as external graphics processor **1411**. In at least one embodiment, platform controller hub **1430** and/or memory controller **1416** may be external to one or more processor(s) **1402**. For example, in at least one embodiment, system **1400** may include an external memory controller **1416** and platform controller hub **1430**, which may be configured as a memory controller hub and peripheral controller hub within a system chipset that is in communication with processor(s) **1402**.

(167) Inference and/or training logic **1015** are used to perform inferencing and/or training operations associated with one or more embodiments. Details regarding inference and/or training logic **1015** are provided below in conjunction with FIGS. **10A** and/or **10B**. In at least one embodiment portions or all of inference and/or training logic **1015** may be incorporated into graphics processor **1408**. For example, in at least one embodiment, training and/or inferencing techniques described herein may use one or more of ALUs embodied in a graphics processor. Moreover, in at least one embodiment, inferencing and/or training operations described herein may be done using logic other than logic illustrated in FIG. **10A** or **10B**. In at least one embodiment, weight parameters may be stored in on-chip or off-chip memory and/or registers (shown or not shown) that configure ALUs of a graphics processor to perform one or more machine learning algorithms, neural network architectures, use cases, or training techniques described herein.

(168) Such components may be used to generate synthetic data imitating failure cases in a network training process, which may help to improve performance of the network while limiting the amount of synthetic data to avoid overfitting.

(169) FIG. **15** is a block diagram of a processor **1500** having one or more processor cores **1502A-1502N**, an integrated memory controller **1513**, and an integrated graphics processor **1508**, according to at least one embodiment. In at least one embodiment, processor **1500** may include additional cores up to and including additional core **1502N** represented by dashed lined boxes. In at least one embodiment, each of processor cores **1502A-1502N** includes one or more internal cache units **1504A-1504N**. In at least one embodiment, each processor core also has access to one or more shared cached units **1506**.

(170) In at least one embodiment, internal cache units **1504A-1504N** and shared cache units **1506** represent a cache memory hierarchy within processor **1500**. In at least one embodiment, cache memory units **1504A-1504N** may include at least one level of instruction and data cache within each processor core and one or more levels of shared mid-level cache, such as a Level 2 (L2), Level 3 (L3), Level 4 (L4), or other levels of cache, where a highest level of cache before external memory is classified as an LLC. In at least one embodiment, cache coherency logic maintains coherency between various cache units **1506** and **1504A-1504N**.

(171) In at least one embodiment, processor **1500** may also include a set of one or more bus controller units **1516** and a system agent core **1510**. In at least one embodiment, one or more bus controller units **1516** manage a set of peripheral buses, such as one or more PCI or PCI express busses. In at least one embodiment, system agent core **1510** provides management functionality for various processor components. In at least one embodiment, system agent core **1510** includes one or more integrated memory controllers **1513** to manage access to various external memory devices (not shown).

(172) In at least one embodiment, one or more of processor cores **1502A-1502N** include support for simultaneous multi-threading. In at least one embodiment, system agent core **1510** includes components for coordinating and operating cores **1502A-1502N** during multi-threaded processing. In at least one embodiment, system agent core **1510** may additionally include a power control unit (PCU), which includes logic and components to regulate one or more power states of processor cores **1502A-1502N** and graphics processor **1508**.

(173) In at least one embodiment, processor **1500** additionally includes graphics processor **1508** to execute graphics processing operations. In at least one embodiment, graphics processor **1508** couples with shared cache units **1506**, and system agent core **1510**, including one or more integrated memory controllers **1513**. In at least one embodiment, system agent core **1510** also includes a display controller **1511** to drive graphics processor output to one or more coupled displays. In at least one embodiment, display controller **1511** may also be a separate module coupled with graphics processor **1508** via at least one interconnect, or may be integrated within graphics processor **1508**.

(174) In at least one embodiment, a ring based interconnect unit **1512** is used to couple internal components of processor **1500**. In at least one embodiment, an alternative interconnect unit may be used, such as a point-to-

point interconnect, a switched interconnect, or other techniques. In at least one embodiment, graphics processor **1508** couples with ring interconnect **1512** via an I/O link **1513**.

(175) In at least one embodiment, I/O link **1513** represents at least one of multiple varieties of I/O interconnects, including an on package I/O interconnect which facilitates communication between various processor components and a high-performance embedded memory module **1518**, such as an eDRAM module. In at least one embodiment, each of processor cores **1502A-1502N** and graphics processor **1508** use embedded memory modules **1518** as a shared Last Level Cache.

(176) In at least one embodiment, processor cores **1502A-1502N** are homogenous cores executing a common instruction set architecture. In at least one embodiment, processor cores **1502A-1502N** are heterogeneous in terms of instruction set architecture (ISA), where one or more of processor cores **1502A-1502N** execute a common instruction set, while one or more other cores of processor cores **1502A-1502N** executes a subset of a common instruction set or a different instruction set. In at least one embodiment, processor cores **1502A-1502N** are heterogeneous in terms of microarchitecture, where one or more cores having a relatively higher power consumption couple with one or more power cores having a lower power consumption. In at least one embodiment, processor **1500** may be implemented on one or more chips or as an SoC integrated circuit.

(177) Inference and/or training logic **1015** are used to perform inferencing and/or training operations associated with one or more embodiments. Details regarding inference and/or training logic **1015** are provided below in conjunction with FIGS. **10A** and/or **10B**. In at least one embodiment portions or all of inference and/or training logic **1015** may be incorporated into processor **1500**. For example, in at least one embodiment, training and/or inferencing techniques described herein may use one or more of ALUs embodied in graphics processor **1508**, graphics core(s) **1502A-1502N**, or other components in FIG. **15**. Moreover, in at least one embodiment, inferencing and/or training operations described herein may be done using logic other than logic illustrated in FIG. **10A** or **10B**. In at least one embodiment, weight parameters may be stored in on-chip or off-chip memory and/or registers (shown or not shown) that configure ALUs of graphics processor **1500** to perform one or more machine learning algorithms, neural network architectures, use cases, or training techniques described herein.

(178) Such components may be used to generate synthetic data imitating failure cases in a network training process, which may help to improve performance of the network while limiting the amount of synthetic data to avoid overfitting.

(179) FIG. **16** is an example data flow diagram for a process **1600** of generating and deploying an image processing and inferencing pipeline, in accordance with at least one embodiment. In at least one embodiment, process **1600** may be deployed for use with imaging devices, processing devices, and/or other device types at one or more facilities **1602**. Process **1600** may be executed within a training system **1604** and/or a deployment system **1606**. In at least one embodiment, training system **1604** may be used to perform training, deployment, and implementation of machine learning models (e.g., neural networks, object detection algorithms, computer vision algorithms, etc.) for use in deployment system **1606**. In at least one embodiment, deployment system **1606** may be configured to offload processing and compute resources among a distributed computing environment to reduce infrastructure requirements at facility **1602**. In at least one embodiment, one or more applications in a pipeline may use or call upon services (e.g., inference, visualization, compute, AI, etc.) of deployment system **1606** during execution of applications.

(180) In at least one embodiment, some of applications used in advanced processing and inferencing pipelines may use machine learning models or other AI to perform one or more processing steps. In at least one embodiment, machine learning models may be trained at facility **1602** using data **1608** (such as imaging data) generated at facility **1602** (and stored on one or more picture archiving and communication system (PACS) servers at facility **1602**), may be trained using imaging or sequencing data **1608** from another facility(ies), or a combination thereof. In at least one embodiment, training system **1604** may be used to provide applications, services, and/or other resources for generating working, deployable machine learning models for deployment system **1606**.

(181) In at least one embodiment, model registry **1624** may be backed by object storage that may support versioning and object metadata. In at least one embodiment, object storage may be accessible through, for example, a cloud storage (e.g., cloud **1726** of FIG. **17**) compatible application programming interface (API) from within a cloud platform. In at least one embodiment, machine learning models within model registry **1624** may be uploaded, listed, modified, or deleted by developers or partners of a system interacting with an API. In at least one embodiment, an API may provide access to methods that allow users with appropriate credentials to associate models with applications, such that models may be executed as part of execution of containerized instantiations of applications.

(182) In at least one embodiment, training pipeline **1704** (FIG. **17**) may include a scenario where facility **1602** is

training their own machine learning model, or has an existing machine learning model that needs to be optimized or updated. In at least one embodiment, imaging data **1608** generated by imaging device(s), sequencing devices, and/or other device types may be received. In at least one embodiment, once imaging data **1608** is received, AI-assisted annotation **1610** may be used to aid in generating annotations corresponding to imaging data **1608** to be used as ground truth data for a machine learning model. In at least one embodiment, AI-assisted annotation **1610** may include one or more machine learning models (e.g., convolutional neural networks (CNNs)) that may be trained to generate annotations corresponding to certain types of imaging data **1608** (e.g., from certain devices). In at least one embodiment, AI-assisted annotations **1610** may then be used directly, or may be adjusted or fine-tuned using an annotation tool to generate ground truth data. In at least one embodiment, AI-assisted annotations **1610**, labeled clinic data **1612**, or a combination thereof may be used as ground truth data for training a machine learning model. In at least one embodiment, a trained machine learning model may be referred to as output model **1616**, and may be used by deployment system **1606**, as described herein.

(183) In at least one embodiment, training pipeline **1704** (FIG. 17) may include a scenario where facility **1602** needs a machine learning model for use in performing one or more processing tasks for one or more applications in deployment system **1606**, but facility **1602** may not currently have such a machine learning model (or may not have a model that is optimized, efficient, or effective for such purposes). In at least one embodiment, an existing machine learning model may be selected from a model registry **1624**. In at least one embodiment, model registry **1624** may include machine learning models trained to perform a variety of different inference tasks on imaging data. In at least one embodiment, machine learning models in model registry **1624** may have been trained on imaging data from different facilities than facility **1602** (e.g., facilities remotely located). In at least one embodiment, machine learning models may have been trained on imaging data from one location, two locations, or any number of locations. In at least one embodiment, when being trained on imaging data from a specific location, training may take place at that location, or at least in a manner that protects confidentiality of imaging data or restricts imaging data from being transferred off-premises. In at least one embodiment, once a model is trained—or partially trained—at one location, a machine learning model may be added to model registry **1624**. In at least one embodiment, a machine learning model may then be retrained, or updated, at any number of other facilities, and a retrained or updated model may be made available in model registry **1624**. In at least one embodiment, a machine learning model may then be selected from model registry **1624**—and referred to as output model **1616**—and may be used in deployment system **1606** to perform one or more processing tasks for one or more applications of a deployment system.

(184) In at least one embodiment, training pipeline **1704** (FIG. 17), a scenario may include facility **1602** requiring a machine learning model for use in performing one or more processing tasks for one or more applications in deployment system **1606**, but facility **1602** may not currently have such a machine learning model (or may not have a model that is optimized, efficient, or effective for such purposes). In at least one embodiment, a machine learning model selected from model registry **1624** may not be fine-tuned or optimized for imaging data **1608** generated at facility **1602** because of differences in populations, robustness of training data used to train a machine learning model, diversity in anomalies of training data, and/or other issues with training data. In at least one embodiment, AI-assisted annotation **1610** may be used to aid in generating annotations corresponding to imaging data **1608** to be used as ground truth data for retraining or updating a machine learning model. In at least one embodiment, labeled data **1612** may be used as ground truth data for training a machine learning model. In at least one embodiment, retraining or updating a machine learning model may be referred to as model training **1614**. In at least one embodiment, model training **1614**—e.g., AI-assisted annotations **1610**, labeled clinic data **1612**, or a combination thereof—may be used as ground truth data for retraining or updating a machine learning model. In at least one embodiment, a trained machine learning model may be referred to as output model **1616**, and may be used by deployment system **1606**, as described herein.

(185) In at least one embodiment, deployment system **1606** may include software **1618**, services **1620**, hardware **1622**, and/or other components, features, and functionality. In at least one embodiment, deployment system **1606** may include a software “stack,” such that software **1618** may be built on top of services **1620** and may use services **1620** to perform some or all of processing tasks, and services **1620** and software **1618** may be built on top of hardware **1622** and use hardware **1622** to execute processing, storage, and/or other compute tasks of deployment system **1606**. In at least one embodiment, software **1618** may include any number of different containers, where each container may execute an instantiation of an application. In at least one embodiment, each application may perform one or more processing tasks in an advanced processing and inferencing pipeline (e.g., inferencing, object detection, feature detection, segmentation, image enhancement, calibration, etc.). In at least one embodiment, an advanced processing and inferencing pipeline may be defined based on selections of different containers that are desired or required for processing imaging data **1608**, in addition to containers that

receive and configure imaging data for use by each container and/or for use by facility **1602** after processing through a pipeline (e.g., to convert outputs back to a usable data type). In at least one embodiment, a combination of containers within software **1618** (e.g., that make up a pipeline) may be referred to as a virtual instrument (as described in more detail herein), and a virtual instrument may leverage services **1620** and hardware **1622** to execute some or all processing tasks of applications instantiated in containers.

(186) In at least one embodiment, a data processing pipeline may receive input data (e.g., imaging data **1608**) in a specific format in response to an inference request (e.g., a request from a user of deployment system **1606**). In at least one embodiment, input data may be representative of one or more images, video, and/or other data representations generated by one or more imaging devices. In at least one embodiment, data may undergo pre-processing as part of data processing pipeline to prepare data for processing by one or more applications. In at least one embodiment, post-processing may be performed on an output of one or more inferencing tasks or other processing tasks of a pipeline to prepare an output data for a next application and/or to prepare output data for transmission and/or use by a user (e.g., as a response to an inference request). In at least one embodiment, inferencing tasks may be performed by one or more machine learning models, such as trained or deployed neural networks, which may include output models **1616** of training system **1604**.

(187) In at least one embodiment, tasks of data processing pipeline may be encapsulated in a container(s) that each represents a discrete, fully functional instantiation of an application and virtualized computing environment that is able to reference machine learning models. In at least one embodiment, containers or applications may be published into a private (e.g., limited access) area of a container registry (described in more detail herein), and trained or deployed models may be stored in model registry **1624** and associated with one or more applications. In at least one embodiment, images of applications (e.g., container images) may be available in a container registry, and once selected by a user from a container registry for deployment in a pipeline, an image may be used to generate a container for an instantiation of an application for use by a user's system.

(188) In at least one embodiment, developers (e.g., software developers, clinicians, doctors, etc.) may develop, publish, and store applications (e.g., as containers) for performing image processing and/or inferencing on supplied data. In at least one embodiment, development, publishing, and/or storing may be performed using a software development kit (SDK) associated with a system (e.g., to ensure that an application and/or container developed is compliant with or compatible with a system). In at least one embodiment, an application that is developed may be tested locally (e.g., at a first facility, on data from a first facility) with an SDK which may support at least some of services **1620** as a system (e.g., system **1700** of FIG. 17). In at least one embodiment, because DICOM objects may contain anywhere from one to hundreds of images or other data types, and due to a variation in data, a developer may be responsible for managing (e.g., setting constructs for, building pre-processing into an application, etc.) extraction and preparation of incoming data. In at least one embodiment, once validated by system **1700** (e.g., for accuracy), an application may be available in a container registry for selection and/or implementation by a user to perform one or more processing tasks with respect to data at a facility (e.g., a second facility) of a user.

(189) In at least one embodiment, developers may then share applications or containers through a network for access and use by users of a system (e.g., system **1700** of FIG. 17). In at least one embodiment, completed and validated applications or containers may be stored in a container registry and associated machine learning models may be stored in model registry **1624**. In at least one embodiment, a requesting entity—who provides an inference or image processing request—may browse a container registry and/or model registry **1624** for an application, container, dataset, machine learning model, etc., select a desired combination of elements for inclusion in data processing pipeline, and submit an imaging processing request. In at least one embodiment, a request may include input data (and associated patient data, in some examples) that is necessary to perform a request, and/or may include a selection of application(s) and/or machine learning models to be executed in processing a request. In at least one embodiment, a request may then be passed to one or more components of deployment system **1606** (e.g., a cloud) to perform processing of data processing pipeline. In at least one embodiment, processing by deployment system **1606** may include referencing selected elements (e.g., applications, containers, models, etc.) from a container registry and/or model registry **1624**. In at least one embodiment, once results are generated by a pipeline, results may be returned to a user for reference (e.g., for viewing in a viewing application suite executing on a local, on-premises workstation or terminal).

(190) In at least one embodiment, to aid in processing or execution of applications or containers in pipelines, services **1620** may be leveraged. In at least one embodiment, services **1620** may include compute services, artificial intelligence (AI) services, visualization services, and/or other service types. In at least one embodiment, services **1620** may provide functionality that is common to one or more applications in software **1618**, so functionality may be abstracted to a service that may be called upon or leveraged by applications. In at least one

embodiment, functionality provided by services **1620** may run dynamically and more efficiently, while also scaling well by allowing applications to process data in parallel (e.g., using a parallel computing platform **1730** (FIG. 17)). In at least one embodiment, rather than each application that shares a same functionality offered by a service **1620** being required to have a respective instance of service **1620**, service **1620** may be shared between and among various applications. In at least one embodiment, services may include an inference server or engine that may be used for executing detection or segmentation tasks, as non-limiting examples. In at least one embodiment, a model training service may be included that may provide machine learning model training and/or retraining capabilities. In at least one embodiment, a data augmentation service may further be included that may provide GPU accelerated data (e.g., DICOM, RIS, CIS, REST compliant, RPC, raw, etc.) extraction, resizing, scaling, and/or other augmentation. In at least one embodiment, a visualization service may be used that may add image rendering effects—such as ray-tracing, rasterization, denoising, sharpening, etc.—to add realism to two-dimensional (2D) and/or three-dimensional (3D) models. In at least one embodiment, virtual instrument services may be included that provide for beam-forming, segmentation, inferencing, imaging, and/or support for other applications within pipelines of virtual instruments.

(191) In at least one embodiment, where a service **1620** includes an AI service (e.g., an inference service), one or more machine learning models may be executed by calling upon (e.g., as an API call) an inference service (e.g., an inference server) to execute machine learning model(s), or processing thereof, as part of application execution. In at least one embodiment, where another application includes one or more machine learning models for segmentation tasks, an application may call upon an inference service to execute machine learning models for performing one or more of processing operations associated with segmentation tasks. In at least one embodiment, software **1618** implementing advanced processing and inferencing pipeline that includes segmentation application and anomaly detection application may be streamlined because each application may call upon a same inference service to perform one or more inferencing tasks.

(192) In at least one embodiment, hardware **1622** may include GPUs, CPUs, DPUs, graphics cards, an AI/deep learning system (e.g., an AI supercomputer, such as NVIDIA's DGX), a cloud platform, or a combination thereof. In at least one embodiment, different types of hardware **1622** may be used to provide efficient, purpose-built support for software **1618** and services **1620** in deployment system **1606**. In at least one embodiment, use of GPU processing may be implemented for processing locally (e.g., at facility **1602**), within an AI/deep learning system, in a cloud system, and/or in other processing components of deployment system **1606** to improve efficiency, accuracy, and efficacy of image processing and generation. In at least one embodiment, software **1618** and/or services **1620** may be optimized for GPU processing with respect to deep learning, machine learning, and/or high-performance computing, as non-limiting examples. In at least one embodiment, at least some of computing environment of deployment system **1606** and/or training system **1604** may be executed in a datacenter one or more supercomputers or high performance computing systems, with GPU optimized software (e.g., hardware and software combination of NVIDIA's DGX System). In at least one embodiment, hardware **1622** may include any number of GPUs that may be called upon to perform processing of data in parallel, as described herein. In at least one embodiment, cloud platform may further include GPU processing for GPU-optimized execution of deep learning tasks, machine learning tasks, or other computing tasks. In at least one embodiment, cloud platform may further include DPU processing to transmit data received over a network and/or through a network controller or other network interface directly to (e.g., a memory of) one or more GPU(s). In at least one embodiment, cloud platform (e.g., NVIDIA's NGC) may be executed using an AI/deep learning supercomputer(s) and/or GPU-optimized software (e.g., as provided on NVIDIA's DGX Systems) as a hardware abstraction and scaling platform. In at least one embodiment, cloud platform may integrate an application container clustering system or orchestration system (e.g., KUBERNETES) on multiple GPUs to enable seamless scaling and load balancing.

(193) FIG. 17 is a system diagram for an example system **1700** for generating and deploying an imaging deployment pipeline, in accordance with at least one embodiment. In at least one embodiment, system **1700** may be used to implement process **1600** of FIG. 16 and/or other processes including advanced processing and inferencing pipelines. In at least one embodiment, system **1700** may include training system **1604** and deployment system **1606**. In at least one embodiment, training system **1604** and deployment system **1606** may be implemented using software **1618**, services **1620**, and/or hardware **1622**, as described herein.

(194) In at least one embodiment, system **1700** (e.g., training system **1604** and/or deployment system **1606**) may be implemented in a cloud computing environment (e.g., using cloud **1726**). In at least one embodiment, system **1700** may be implemented locally with respect to a healthcare services facility, or as a combination of both cloud and local computing resources. In at least one embodiment, access to APIs in cloud **1726** may be restricted to authorized users through enacted security measures or protocols. In at least one embodiment, a security

protocol may include web tokens that may be signed by an authentication (e.g., AuthN, AuthZ, Gluecon, etc.) service and may carry appropriate authorization. In at least one embodiment, APIs of virtual instruments (described herein), or other instantiations of system **1700**, may be restricted to a set of public IPs that have been vetted or authorized for interaction.

(195) In at least one embodiment, various components of system **1700** may communicate between and among one another using any of a variety of different network types, including but not limited to local area networks (LANs) and/or wide area networks (WANs) via wired and/or wireless communication protocols. In at least one embodiment, communication between facilities and components of system **1700** (e.g., for transmitting inference requests, for receiving results of inference requests, etc.) may be communicated over data bus(es), wireless data protocols (Wi-Fi), wired data protocols (e.g., Ethernet), etc.

(196) In at least one embodiment, training system **1604** may execute training pipelines **1704**, similar to those described herein with respect to FIG. **16**. In at least one embodiment, where one or more machine learning models are to be used in deployment pipelines **1710** by deployment system **1606**, training pipelines **1704** may be used to train or retrain one or more (e.g. pre-trained) models, and/or implement one or more of pre-trained models **1706** (e.g., without a need for retraining or updating). In at least one embodiment, as a result of training pipelines **1704**, output model(s) **1616** may be generated. In at least one embodiment, training pipelines **1704** may include any number of processing steps, such as but not limited to imaging data (or other input data) conversion or adaption. In at least one embodiment, for different machine learning models used by deployment system **1606**, different training pipelines **1704** may be used. In at least one embodiment, training pipeline **1704** similar to a first example described with respect to FIG. **16** may be used for a first machine learning model, training pipeline **1704** similar to a second example described with respect to FIG. **16** may be used for a second machine learning model, and training pipeline **1704** similar to a third example described with respect to FIG. **16** may be used for a third machine learning model. In at least one embodiment, any combination of tasks within training system **1604** may be used depending on what is required for each respective machine learning model. In at least one embodiment, one or more of machine learning models may already be trained and ready for deployment so machine learning models may not undergo any processing by training system **1604**, and may be implemented by deployment system **1606**.

(197) In at least one embodiment, output model(s) **1616** and/or pre-trained model(s) **1706** may include any types of machine learning models depending on implementation or embodiment. In at least one embodiment, and without limitation, machine learning models used by system **1700** may include machine learning model(s) using linear regression, logistic regression, decision trees, support vector machines (SVM), Naïve Bayes, k-nearest neighbor (Knn), K means clustering, random forest, dimensionality reduction algorithms, gradient boosting algorithms, neural networks (e.g., auto-encoders, convolutional, recurrent, perceptrons, Long/Short Term Memory (LSTM), Hopfield, Boltzmann, deep belief, deconvolutional, generative adversarial, liquid state machine, etc.), and/or other types of machine learning models.

(198) In at least one embodiment, training pipelines **1704** may include AI-assisted annotation, as described in more detail herein with respect to at least FIG. **16B**. In at least one embodiment, labeled data **1612** (e.g., traditional annotation) may be generated by any number of techniques. In at least one embodiment, labels or other annotations may be generated within a drawing program (e.g., an annotation program), a computer aided design (CAD) program, a labeling program, another type of program suitable for generating annotations or labels for ground truth, and/or may be hand drawn, in some examples. In at least one embodiment, ground truth data may be synthetically produced (e.g., generated from computer models or renderings), real produced (e.g., designed and produced from real-world data), machine-automated (e.g., using feature analysis and learning to extract features from data and then generate labels), human annotated (e.g., labeler, or annotation expert, defines location of labels), and/or a combination thereof. In at least one embodiment, for each instance of imaging data **1608** (or other data type used by machine learning models), there may be corresponding ground truth data generated by training system **1604**. In at least one embodiment, AI-assisted annotation may be performed as part of deployment pipelines **1710**; either in addition to, or in lieu of AI-assisted annotation included in training pipelines **1704**. In at least one embodiment, system **1700** may include a multi-layer platform that may include a software layer (e.g., software **1618**) of diagnostic applications (or other application types) that may perform one or more medical imaging and diagnostic functions. In at least one embodiment, system **1700** may be communicatively coupled to (e.g., via encrypted links) PACS server networks of one or more facilities. In at least one embodiment, system **1700** may be configured to access and referenced data from PACS servers to perform operations, such as training machine learning models, deploying machine learning models, image processing, inferencing, and/or other operations.

(199) In at least one embodiment, a software layer may be implemented as a secure, encrypted, and/or

authenticated API through which applications or containers may be invoked (e.g., called) from an external environment(s) (e.g., facility **1602**). In at least one embodiment, applications may then call or execute one or more services **1620** for performing compute, AI, or visualization tasks associated with respective applications, and software **1618** and/or services **1620** may leverage hardware **1622** to perform processing tasks in an effective and efficient manner.

(200) In at least one embodiment, deployment system **1606** may execute deployment pipelines **1710**. In at least one embodiment, deployment pipelines **1710** may include any number of applications that may be sequentially, non-sequentially, or otherwise applied to imaging data (and/or other data types) generated by imaging devices, sequencing devices, genomics devices, etc.—including AI-assisted annotation, as described above. In at least one embodiment, as described herein, a deployment pipeline **1710** for an individual device may be referred to as a virtual instrument for a device (e.g., a virtual ultrasound instrument, a virtual CT scan instrument, a virtual sequencing instrument, etc.). In at least one embodiment, for a single device, there may be more than one deployment pipeline **1710** depending on information desired from data generated by a device. In at least one embodiment, where detections of anomalies are desired from an MRI machine, there may be a first deployment pipeline **1710**, and where image enhancement is desired from output of an MRI machine, there may be a second deployment pipeline **1710**.

(201) In at least one embodiment, an image generation application may include a processing task that includes use of a machine learning model. In at least one embodiment, a user may desire to use their own machine learning model, or to select a machine learning model from model registry **1624**. In at least one embodiment, a user may implement their own machine learning model or select a machine learning model for inclusion in an application for performing a processing task. In at least one embodiment, applications may be selectable and customizable, and by defining constructs of applications, deployment, and implementation of applications for a particular user are presented as a more seamless user experience. In at least one embodiment, by leveraging other features of system **1700**—such as services **1620** and hardware **1622**—deployment pipelines **1710** may be even more user friendly, provide for easier integration, and produce more accurate, efficient, and timely results.

(202) In at least one embodiment, deployment system **1606** may include a user interface **1714** (e.g., a graphical user interface, a web interface, etc.) that may be used to select applications for inclusion in deployment pipeline(s) **1710**, arrange applications, modify, or change applications or parameters or constructs thereof, use and interact with deployment pipeline(s) **1710** during set-up and/or deployment, and/or to otherwise interact with deployment system **1606**. In at least one embodiment, although not illustrated with respect to training system **1604**, user interface **1714** (or a different user interface) may be used for selecting models for use in deployment system **1606**, for selecting models for training, or retraining, in training system **1604**, and/or for otherwise interacting with training system **1604**.

(203) In at least one embodiment, pipeline manager **1712** may be used, in addition to an application orchestration system **1728**, to manage interaction between applications or containers of deployment pipeline(s) **1710** and services **1620** and/or hardware **1622**. In at least one embodiment, pipeline manager **1712** may be configured to facilitate interactions from application to application, from application to service **1620**, and/or from application or service to hardware **1622**. In at least one embodiment, although illustrated as included in software **1618**, this is not intended to be limiting, and in some examples (e.g., as illustrated in FIG. **15**) pipeline manager **1712** may be included in services **1620**. In at least one embodiment, application orchestration system **1728** (e.g., Kubernetes, DOCKER, etc.) may include a container orchestration system that may group applications into containers as logical units for coordination, management, scaling, and deployment. In at least one embodiment, by associating applications from deployment pipeline(s) **1710** (e.g., a reconstruction application, a segmentation application, etc.) with individual containers, each application may execute in a self-contained environment (e.g., at a kernel level) to increase speed and efficiency.

(204) In at least one embodiment, each application and/or container (or image thereof) may be individually developed, modified, and deployed (e.g., a first user or developer may develop, modify, and deploy a first application and a second user or developer may develop, modify, and deploy a second application separate from a first user or developer), which may allow for focus on, and attention to, a task of a single application and/or container(s) without being hindered by tasks of another application(s) or container(s). In at least one embodiment, communication, and cooperation between different containers or applications may be aided by pipeline manager **1712** and application orchestration system **1728**. In at least one embodiment, so long as an expected input and/or output of each container or application is known by a system (e.g., based on constructs of applications or containers), application orchestration system **1728** and/or pipeline manager **1712** may facilitate communication among and between, and sharing of resources among and between, each of applications or containers. In at least one embodiment, because one or more of applications or containers in deployment

pipeline(s) **1710** may share same services and resources, application orchestration system **1728** may orchestrate, load balance, and determine sharing of services or resources between and among various applications or containers. In at least one embodiment, a scheduler may be used to track resource requirements of applications or containers, current usage or planned usage of these resources, and resource availability. In at least one embodiment, a scheduler may thus allocate resources to different applications and distribute resources between and among applications in view of requirements and availability of a system. In some examples, a scheduler (and/or other component of application orchestration system **1728**) may determine resource availability and distribution based on constraints imposed on a system (e.g., user constraints), such as quality of service (QoS), urgency of need for data outputs (e.g., to determine whether to execute real-time processing or delayed processing), etc.

(205) In at least one embodiment, services **1620** leveraged by and shared by applications or containers in deployment system **1606** may include compute services **1716**, AI services **1718**, visualization services **1720**, and/or other service types. In at least one embodiment, applications may call (e.g., execute) one or more of services **1620** to perform processing operations for an application. In at least one embodiment, compute services **1716** may be leveraged by applications to perform super-computing or other high-performance computing (HPC) tasks. In at least one embodiment, compute service(s) **1716** may be leveraged to perform parallel processing (e.g., using a parallel computing platform **1730**) for processing data through one or more of applications and/or one or more tasks of a single application, substantially simultaneously. In at least one embodiment, parallel computing platform **1730** (e.g., NVIDIA's CUDA) may enable general purpose computing on GPUs (GPGPU) (e.g., GPUs **1722**). In at least one embodiment, a software layer of parallel computing platform **1730** may provide access to virtual instruction sets and parallel computational elements of GPUs, for execution of compute kernels. In at least one embodiment, parallel computing platform **1730** may include memory and, in some embodiments, a memory may be shared between and among multiple containers, and/or between and among different processing tasks within a single container. In at least one embodiment, inter-process communication (IPC) calls may be generated for multiple containers and/or for multiple processes within a container to use same data from a shared segment of memory of parallel computing platform **1730** (e.g., where multiple different stages of an application or multiple applications are processing same information). In at least one embodiment, rather than making a copy of data and moving data to different locations in memory (e.g., a read/write operation), same data in same location of a memory may be used for any number of processing tasks (e.g., at a same time, at different times, etc.). In at least one embodiment, as data is used to generate new data as a result of processing, this information of a new location of data may be stored and shared between various applications. In at least one embodiment, location of data and a location of updated or modified data may be part of a definition of how a payload is understood within containers.

(206) In at least one embodiment, AI services **1718** may be leveraged to perform inferencing services for executing machine learning model(s) associated with applications (e.g., tasked with performing one or more processing tasks of an application). In at least one embodiment, AI services **1718** may leverage AI system **1724** to execute machine learning model(s) (e.g., neural networks, such as CNNs) for segmentation, reconstruction, object detection, feature detection, classification, and/or other inferencing tasks. In at least one embodiment, applications of deployment pipeline(s) **1710** may use one or more of output models **1616** from training system **1604** and/or other models of applications to perform inference on imaging data. In at least one embodiment, two or more examples of inferencing using application orchestration system **1728** (e.g., a scheduler) may be available. In at least one embodiment, a first category may include a high priority/low latency path that may achieve higher service level agreements, such as for performing inference on urgent requests during an emergency, or for a radiologist during diagnosis. In at least one embodiment, a second category may include a standard priority path that may be used for requests that may be non-urgent or where analysis may be performed at a later time. In at least one embodiment, application orchestration system **1728** may distribute resources (e.g., services **1620** and/or hardware **1622**) based on priority paths for different inferencing tasks of AI services **1718**.

(207) In at least one embodiment, shared storage may be mounted to AI services **1718** within system **1700**. In at least one embodiment, shared storage may operate as a cache (or other storage device type) and may be used to process inference requests from applications. In at least one embodiment, when an inference request is submitted, a request may be received by a set of API instances of deployment system **1606**, and one or more instances may be selected (e.g., for best fit, for load balancing, etc.) to process a request. In at least one embodiment, to process a request, a request may be entered into a database, a machine learning model may be located from model registry **1624** if not already in a cache, a validation step may ensure appropriate machine learning model is loaded into a cache (e.g., shared storage), and/or a copy of a model may be saved to a cache. In at least one embodiment, a scheduler (e.g., of pipeline manager **1712**) may be used to launch an application

that is referenced in a request is not already running or if there are not enough instances of an application. In at least one embodiment, if an inference server is not already launched to execute a model, an inference server may be launched. Any number of inference servers may be launched per model. In at least one embodiment, in a pull model, in which inference servers are clustered, models may be cached whenever load balancing is advantageous. In at least one embodiment, inference servers may be statically loaded in corresponding, distributed servers.

(208) In at least one embodiment, inferencing may be performed using an inference server that runs in a container. In at least one embodiment, an instance of an inference server may be associated with a model (and optionally a plurality of versions of a model). In at least one embodiment, if an instance of an inference server does not exist when a request to perform inference on a model is received, a new instance may be loaded. In at least one embodiment, when starting an inference server, a model may be passed to an inference server such that a same container may be used to serve different models so long as inference server is running as a different instance.

(209) In at least one embodiment, during application execution, an inference request for a given application may be received, and a container (e.g., hosting an instance of an inference server) may be loaded (if not already), and a start procedure may be called. In at least one embodiment, pre-processing logic in a container may load, decode, and/or perform any additional pre-processing on incoming data (e.g., using a CPU(s) and/or GPU(s) and/or DPU(s)). In at least one embodiment, once data is prepared for inference, a container may perform inference as necessary on data. In at least one embodiment, this may include a single inference call on one image (e.g., a hand X-ray), or may require inference on hundreds of images (e.g., a chest CT). In at least one embodiment, an application may summarize results before completing, which may include, without limitation, a single confidence score, pixel level-segmentation, voxel-level segmentation, generating a visualization, or generating text to summarize findings. In at least one embodiment, different models or applications may be assigned different priorities. For example, some models may have a real-time (TAT<1 min) priority while others may have lower priority (e.g., TAT<12 min). In at least one embodiment, model execution times may be measured from requesting institution or entity and may include partner network traversal time, as well as execution on an inference service.

(210) In at least one embodiment, transfer of requests between services **1620** and inference applications may be hidden behind a software development kit (SDK), and robust transport may be provided through a queue. In at least one embodiment, a request will be placed in a queue via an API for an individual application/tenant ID combination and an SDK will pull a request from a queue and give a request to an application. In at least one embodiment, a name of a queue may be provided in an environment from where an SDK will pick it up. In at least one embodiment, asynchronous communication through a queue may be useful as it may allow any instance of an application to pick up work as it becomes available. Results may be transferred back through a queue, to ensure no data is lost. In at least one embodiment, queues may also provide an ability to segment work, as highest priority work may go to a queue with most instances of an application connected to it, while lowest priority work may go to a queue with a single instance connected to it that processes tasks in an order received. In at least one embodiment, an application may run on a GPU-accelerated instance generated in cloud **1726**, and an inference service may perform inferencing on a GPU.

(211) In at least one embodiment, visualization services **1720** may be leveraged to generate visualizations for viewing outputs of applications and/or deployment pipeline(s) **1710**. In at least one embodiment, GPUs **1722** may be leveraged by visualization services **1720** to generate visualizations. In at least one embodiment, rendering effects, such as ray-tracing, may be implemented by visualization services **1720** to generate higher quality visualizations. In at least one embodiment, visualizations may include, without limitation, 2D image renderings, 3D volume renderings, 3D volume reconstruction, 2D tomographic slices, virtual reality displays, augmented reality displays, etc. In at least one embodiment, virtualized environments may be used to generate a virtual interactive display or environment (e.g., a virtual environment) for interaction by users of a system (e.g., doctors, nurses, radiologists, etc.). In at least one embodiment, visualization services **1720** may include an internal visualizer, cinematics, and/or other rendering or image processing capabilities or functionality (e.g., ray tracing, rasterization, internal optics, etc.).

(212) In at least one embodiment, hardware **1622** may include GPUs **1722**, AI system **1724**, cloud **1726**, and/or any other hardware used for executing training system **1604** and/or deployment system **1606**. In at least one embodiment, GPUs **1722** (e.g., NVIDIA's TESLA and/or QUADRO GPUs) may include any number of GPUs that may be used for executing processing tasks of compute services **1716**, AI services **1718**, visualization services **1720**, other services, and/or any of features or functionality of software **1618**. For example, with respect to AI services **1718**, GPUs **1722** may be used to perform pre-processing on imaging data (or other data types

used by machine learning models), post-processing on outputs of machine learning models, and/or to perform inferencing (e.g., to execute machine learning models). In at least one embodiment, cloud **1726**, AI system **1724**, and/or other components of system **1700** may use GPUs **1722**. In at least one embodiment, cloud **1726** may include a GPU-optimized platform for deep learning tasks. In at least one embodiment, AI system **1724** may use GPUs, and cloud **1726**—or at least a portion tasked with deep learning or inferencing—may be executed using one or more AI systems **1724**. As such, although hardware **1622** is illustrated as discrete components, this is not intended to be limiting, and any components of hardware **1622** may be combined with, or leveraged by, any other components of hardware **1622**.

(213) In at least one embodiment, AI system **1724** may include a purpose-built computing system (e.g., a super-computer or an HPC) configured for inferencing, deep learning, machine learning, and/or other artificial intelligence tasks. In at least one embodiment, AI system **1724** (e.g., NVIDIA's DGX) may include GPU-optimized software (e.g., a software stack) that may be executed using a plurality of GPUs **1722**, in addition to DPUs, CPUs, RAM, storage, and/or other components, features, or functionality. In at least one embodiment, one or more AI systems **1724** may be implemented in cloud **1726** (e.g., in a data center) for performing some or all of AI-based processing tasks of system **1700**.

(214) In at least one embodiment, cloud **1726** may include a GPU-accelerated infrastructure (e.g., NVIDIA's NGC) that may provide a GPU-optimized platform for executing processing tasks of system **1700**. In at least one embodiment, cloud **1726** may include an AI system(s) **1724** for performing one or more of AI-based tasks of system **1700** (e.g., as a hardware abstraction and scaling platform). In at least one embodiment, cloud **1726** may integrate with application orchestration system **1728** leveraging multiple GPUs to enable seamless scaling and load balancing between and among applications and services **1620**. In at least one embodiment, cloud **1726** may be tasked with executing at least some of services **1620** of system **1700**, including compute services **1716**, AI services **1718**, and/or visualization services **1720**, as described herein. In at least one embodiment, cloud **1726** may perform small and large batch inference (e.g., executing NVIDIA's TENSOR RT), provide an accelerated parallel computing API and platform **1730** (e.g., NVIDIA's CUDA), execute application orchestration system **1728** (e.g., KUBERNETES), provide a graphics rendering API and platform (e.g., for ray-tracing, 2D graphics, 3D graphics, and/or other rendering techniques to produce higher quality cinematics), and/or may provide other functionality for system **1700**.

(215) FIG. **18A** illustrates a data flow diagram for a process **1800** to train, retrain, or update a machine learning model, in accordance with at least one embodiment. In at least one embodiment, process **1800** may be executed using, as a non-limiting example, system **1700** of FIG. **17**. In at least one embodiment, process **1800** may leverage services **1620** and/or hardware **1622** of system **1700**, as described herein. In at least one embodiment, refined models **1812** generated by process **1800** may be executed by deployment system **1606** for one or more containerized applications in deployment pipelines **1710**.

(216) In at least one embodiment, model training **1614** may include retraining or updating an initial model **1804** (e.g., a pre-trained model) using new training data (e.g., new input data, such as customer dataset **1806**, and/or new ground truth data associated with input data). In at least one embodiment, to retrain, or update, initial model **1804**, output or loss layer(s) of initial model **1804** may be reset, or deleted, and/or replaced with an updated or new output or loss layer(s). In at least one embodiment, initial model **1804** may have previously fine-tuned parameters (e.g., weights and/or biases) that remain from prior training, so training or retraining **1614** may not take as long or require as much processing as training a model from scratch. In at least one embodiment, during model training **1614**, by having reset or replaced output or loss layer(s) of initial model **1804**, parameters may be updated and re-tuned for a new data set based on loss calculations associated with accuracy of output or loss layer(s) at generating predictions on new, customer dataset **1806** (e.g., image data **1608** of FIG. **16**).

(217) In at least one embodiment, pre-trained models **1706** may be stored in a data store, or registry (e.g., model registry **1624** of FIG. **16**). In at least one embodiment, pre-trained models **1706** may have been trained, at least in part, at one or more facilities other than a facility executing process **1800**. In at least one embodiment, to protect privacy and rights of patients, subjects, or clients of different facilities, pre-trained models **1706** may have been trained, on-premise, using customer or patient data generated on-premise. In at least one embodiment, pre-trained models **1706** may be trained using cloud **1726** and/or other hardware **1622**, but confidential, privacy protected patient data may not be transferred to, used by, or accessible to any components of cloud **1726** (or other off premise hardware). In at least one embodiment, where a pre-trained model **1706** is trained at using patient data from more than one facility, pre-trained model **1706** may have been individually trained for each facility prior to being trained on patient or customer data from another facility. In at least one embodiment, such as where a customer or patient data has been released of privacy concerns (e.g., by waiver, for experimental use, etc.), or where a customer or patient data is included in a public data set, a customer or patient data from any

number of facilities may be used to train pre-trained model **1706** on-premise and/or off premise, such as in a datacenter or other cloud computing infrastructure.

(218) In at least one embodiment, when selecting applications for use in deployment pipelines **1710**, a user may also select machine learning models to be used for specific applications. In at least one embodiment, a user may not have a model for use, so a user may select a pre-trained model **1706** to use with an application. In at least one embodiment, pre-trained model **1706** may not be optimized for generating accurate results on customer dataset **1806** of a facility of a user (e.g., based on patient diversity, demographics, types of medical imaging devices used, etc.). In at least one embodiment, prior to deploying pre-trained model **1706** into deployment pipeline **1710** for use with an application(s), pre-trained model **1706** may be updated, retrained, and/or fine-tuned for use at a respective facility.

(219) In at least one embodiment, a user may select pre-trained model **1706** that is to be updated, retrained, and/or fine-tuned, and pre-trained model **1706** may be referred to as initial model **1804** for training system **1604** within process **1800**. In at least one embodiment, customer dataset **1806** (e.g., imaging data, genomics data, sequencing data, or other data types generated by devices at a facility) may be used to perform model training **1614** (which may include, without limitation, transfer learning) on initial model **1804** to generate refined model **1812**. In at least one embodiment, ground truth data corresponding to customer dataset **1806** may be generated by training system **1604**. In at least one embodiment, ground truth data may be generated, at least in part, by clinicians, scientists, doctors, practitioners, at a facility (e.g., as labeled clinic data **1612** of FIG. **16**).

(220) In at least one embodiment, AI-assisted annotation **1610** may be used in some examples to generate ground truth data. In at least one embodiment, AI-assisted annotation **1610** (e.g., implemented using an AI-assisted annotation SDK) may leverage machine learning models (e.g., neural networks) to generate suggested or predicted ground truth data for a customer dataset. In at least one embodiment, user **1810** may use annotation tools within a user interface (a graphical user interface (GUI)) on computing device **1808**.

(221) In at least one embodiment, user **1810** may interact with a GUI via computing device **1808** to edit or fine-tune (auto)annotations. In at least one embodiment, a polygon editing feature may be used to move vertices of a polygon to more accurate or fine-tuned locations.

(222) In at least one embodiment, once customer dataset **1806** has associated ground truth data, ground truth data (e.g., from AI-assisted annotation, manual labeling, etc.) may be used by during model training **1614** to generate refined model **1812**. In at least one embodiment, customer dataset **1806** may be applied to initial model **1804** any number of times, and ground truth data may be used to update parameters of initial model **1804** until an acceptable level of accuracy is attained for refined model **1812**. In at least one embodiment, once refined model **1812** is generated, refined model **1812** may be deployed within one or more deployment pipelines **1710** at a facility for performing one or more processing tasks with respect to medical imaging data.

(223) In at least one embodiment, refined model **1812** may be uploaded to pre-trained models **1706** in model registry **1624** to be selected by another facility. In at least one embodiment, his process may be completed at any number of facilities such that refined model **1812** may be further refined on new datasets any number of times to generate a more universal model.

(224) FIG. **18B** is an example illustration of a client-server architecture **1832** to enhance annotation tools with pre-trained annotation models, in accordance with at least one embodiment. In at least one embodiment, AI-assisted annotation tools **1836** may be instantiated based on a client-server architecture **1832**. In at least one embodiment, annotation tools **1836** in imaging applications may aid radiologists, for example, identify organs and abnormalities. In at least one embodiment, imaging applications may include software tools that help user **1810** to identify, as a non-limiting example, a few extreme points on a particular organ of interest in raw images **1834** (e.g., in a 3D MRI or CT scan) and receive auto-annotated results for all 2D slices of a particular organ. In at least one embodiment, results may be stored in a data store as training data **1838** and used as (for example and without limitation) ground truth data for training. In at least one embodiment, when computing device **1808** sends extreme points for AI-assisted annotation **1610**, a deep learning model, for example, may receive this data as input and return inference results of a segmented organ or abnormality. In at least one embodiment, pre-instantiated annotation tools, such as AI-Assisted Annotation Tool **1836B** in FIG. **18B**, may be enhanced by making API calls (e.g., API Call **1844**) to a server, such as an Annotation Assistant Server **1840** that may include a set of pre-trained models **1842** stored in an annotation model registry, for example. In at least one embodiment, an annotation model registry may store pre-trained models **1842** (e.g., machine learning models, such as deep learning models) that are pre-trained to perform AI-assisted annotation on a particular organ or abnormality. These models may be further updated by using training pipelines **1704**. In at least one embodiment, pre-installed annotation tools may be improved over time as new labeled clinic data **1612** is added.

(225) Such components may be used to generate synthetic data imitating failure cases in a network training

process, which may help to improve performance of the network while limiting the amount of synthetic data to avoid overfitting.

(226) Other variations are within spirit of present disclosure. Thus, while disclosed techniques are susceptible to various modifications and alternative constructions, certain illustrated embodiments thereof are shown in drawings and have been described above in detail. It should be understood, however, that there is no intention to limit the disclosure to a specific form or forms disclosed, but on the contrary, the intention is to cover all modifications, alternative constructions, and equivalents falling within the spirit and scope of the disclosure, as defined in appended claims.

(227) Use of terms “a” and “an” and “the” and similar referents in the context of describing disclosed embodiments (especially in the context of following claims) are to be construed to cover both singular and plural, unless otherwise indicated herein or clearly contradicted by context, and not as a definition of a term. Terms “comprising,” “having,” “including,” and “containing” are to be construed as open-ended terms (meaning “including, but not limited to,”) unless otherwise noted. “Connected,” when unmodified and referring to physical connections, is to be construed as partly or wholly contained within, attached to, or joined together, even if there is something intervening. Recitations of ranges of values herein are merely intended to serve as a shorthand method of referring individually to each separate value falling within the range, unless otherwise indicated herein, and each separate value is incorporated into the specification as if it were individually recited herein. In at least one embodiment, the use of the term “set” (e.g., “a set of items”) or “subset” unless otherwise noted or contradicted by context, is to be construed as a nonempty collection comprising one or more members. Further, unless otherwise noted or contradicted by context, the term “subset” of a corresponding set does not necessarily denote a proper subset of the corresponding set, but subset and corresponding set may be equal.

(228) Conjunctive language, such as phrases of the form “at least one of A, B, and C,” or “at least one of A, B and C,” unless specifically stated otherwise or otherwise clearly contradicted by context, is otherwise understood with the context as used in general to present that an item, term, etc., may be either A or B or C, or any nonempty subset of the set of A and B and C. For instance, in an illustrative example of a set having three members, conjunctive phrases “at least one of A, B, and C” and “at least one of A, B and C” refer to any of the following sets: {A}, {B}, {C}, {A, B}, {A, C}, {B, C}, {A, B, C}. Thus, such conjunctive language is not generally intended to imply that certain embodiments require at least one of A, at least one of B and at least one of C each to be present. In addition, unless otherwise noted or contradicted by context, the term “plurality” indicates a state of being plural (e.g., “a plurality of items” indicates multiple items). In at least one embodiment, the number of items in a plurality is at least two, but can be more when so indicated either explicitly or by context. Further, unless stated otherwise or otherwise clear from context, the phrase “based on” means “based at least in part on” and not “based solely on.”

(229) Operations of processes described herein can be performed in any suitable order unless otherwise indicated herein or otherwise clearly contradicted by context. In at least one embodiment, a process such as those processes described herein (or variations and/or combinations thereof) is performed under control of one or more computer systems configured with executable instructions and is implemented as code (e.g., executable instructions, one or more computer programs or one or more applications) executing collectively on one or more processors, by hardware or combinations thereof. In at least one embodiment, code is stored on a computer-readable storage medium, for example, in the form of a computer program comprising a plurality of instructions executable by one or more processors. In at least one embodiment, a computer-readable storage medium is a non-transitory computer-readable storage medium that excludes transitory signals (e.g., a propagating transient electric or electromagnetic transmission) but includes non-transitory data storage circuitry (e.g., buffers, cache, and queues) within transceivers of transitory signals. In at least one embodiment, code (e.g., executable code or source code) is stored on a set of one or more non-transitory computer-readable storage media having stored thereon executable instructions (or other memory to store executable instructions) that, when executed (i.e., as a result of being executed) by one or more processors of a computer system, cause a computer system to perform operations described herein. In at least one embodiment, a set of non-transitory computer-readable storage media comprises multiple non-transitory computer-readable storage media and one or more of individual non-transitory storage media of multiple non-transitory computer-readable storage media lack all of the code while multiple non-transitory computer-readable storage media collectively store all of the code. In at least one embodiment, executable instructions are executed such that different instructions are executed by different processors.

(230) Accordingly, in at least one embodiment, computer systems are configured to implement one or more services that singly or collectively perform operations of processes described herein and such computer systems are configured with applicable hardware and/or software that enable the performance of operations. Further, a computer system that implements at least one embodiment of present disclosure is a single device and, in

another embodiment, is a distributed computer system comprising multiple devices that operate differently such that distributed computer system performs operations described herein and such that a single device does not perform all operations.

(231) Use of any and all examples, or exemplary language (e.g., “such as”) provided herein, is intended merely to better illuminate embodiments of the disclosure and does not pose a limitation on the scope of the disclosure unless otherwise claimed. No language in the specification should be construed as indicating any non-claimed element as essential to the practice of the disclosure.

(232) All references, including publications, patent applications, and patents, cited herein are hereby incorporated by reference to the same extent as if each reference were individually and specifically indicated to be incorporated by reference and were set forth in its entirety herein.

(233) In description and claims, terms “coupled” and “connected,” along with their derivatives, may be used. It should be understood that these terms may not be intended as synonyms for each other. Rather, in particular examples, “connected” or “coupled” may be used to indicate that two or more elements are in direct or indirect physical or electrical contact with each other. “Coupled” may also mean that two or more elements are not in direct contact with each other, but yet still co-operate or interact with each other.

(234) Unless specifically stated otherwise, it may be appreciated that throughout specification terms such as “processing,” “computing,” “calculating,” “determining,” or like, refer to action and/or processes of a computer or computing system, or similar electronic computing device, that manipulate and/or transform data represented as physical, such as electronic, quantities within computing system's registers and/or memories into other data similarly represented as physical quantities within computing system's memories, registers or other such information storage, transmission or display devices.

(235) In a similar manner, the term “processor” may refer to any device or portion of a device that processes electronic data from registers and/or memory and transform that electronic data into other electronic data that may be stored in registers and/or memory. A “computing platform” may comprise one or more processors. As used herein, “software” processes may include, for example, software and/or hardware entities that perform work over time, such as tasks, threads, and intelligent agents. Also, each process may refer to multiple processes, for carrying out instructions in sequence or in parallel, continuously or intermittently. In at least one embodiment, terms “system” and “method” are used herein interchangeably insofar as the system may embody one or more methods and methods may be considered a system.

(236) In the present document, references may be made to obtaining, acquiring, receiving, or inputting analog or digital data into a subsystem, computer system, or computer-implemented machine. In at least one embodiment, the process of obtaining, acquiring, receiving, or inputting analog and digital data can be accomplished in a variety of ways such as by receiving data as a parameter of a function call or a call to an application programming interface. In at least one embodiment, processes of obtaining, acquiring, receiving, or inputting analog or digital data can be accomplished by transferring data via a serial or parallel interface. In at least one embodiment, processes of obtaining, acquiring, receiving, or inputting analog or digital data can be accomplished by transferring data via a computer network from providing entity to acquiring entity. In at least one embodiment, references may also be made to providing, outputting, transmitting, sending, or presenting analog or digital data. In various examples, processes of providing, outputting, transmitting, sending, or presenting analog or digital data can be accomplished by transferring data as an input or output parameter of a function call, a parameter of an application programming interface or inter-process communication mechanism.

(237) Although descriptions herein set forth example embodiments of described techniques, other architectures may be used to implement described functionality, and are intended to be within the scope of this disclosure. Furthermore, although specific distributions of responsibilities may be defined above for purposes of description, various functions and responsibilities might be distributed and divided in different ways, depending on circumstances.

(238) Furthermore, although the subject matter has been described in language specific to structural features and/or methodological acts, it is to be understood that subject matter claimed in appended claims is not necessarily limited to specific features or acts described. Rather, specific features and acts are disclosed as exemplary forms of implementing the claims.

Claims

1. A method comprising: receiving, using a processing device, a first design state of a data path circuit; providing the first design state of the data path circuit as an input to a machine learning model; and performing one or more reinforcement learning techniques using the machine learning model to cause the machine learning

model to output a final design state of the data path circuit, the one or more reinforcement learning techniques comprising: determining a first parameter value on a first parameter curve of a first physical data path circuit synthesized from the first design state, wherein the first parameter curve indicates a plurality of parameter values at one or more areas of the first physical data path circuit; and determining the final design state of the data path circuit, the final design state associated with a final parameter value, wherein the final design state is closer to a target than the first parameter value associated with the first design state.

2. The method of claim 1, wherein performing the one or more reinforcement learning techniques comprises: processing the first design state of the data path circuit using the machine learning model, wherein the machine learning model outputs a second design state of the data path circuit that is a modification of the first design state; determining a second parameter value for the second design state; updating the machine learning model based on a comparison of the first parameter value and the second parameter value; and processing the first design state or the second design state of the data path circuit using the updated machine learning model, wherein the updated machine learning model outputs a third design state of the data path circuit that is a modification of the first design state or the second design state.

3. The method of claim 2, further comprising: processing the first design state using a second model, wherein the second model outputs the first parameter value associated with the first design state; processing the second design state using the second model, wherein the second model outputs the second parameter value associated with the second design state; and making the comparison between the second parameter value and the first parameter value to determine which of the first parameter value and the second parameter value is closer to the target.

4. The method of claim 3, wherein: the machine learning model receives as an input a grid representation of a first graph of the first design state and outputs a grid representation of a second graph of the second design state; and the second model is a metrics predictor model that receives as a first input the grid representation of the first graph of the first design state and receives as a second input the grid representation of the second graph of the second design state.

5. The method of claim 3, the machine learning model receives as an input a grid representation of a first graph of the first design state and outputs an action to generate the grid representation of a second graph of the second design state, the method further comprising: converting the grid representation of the second graph of the second design state to a circuit implementation of the second design state; wherein the second model receives as a first input a circuit implementation of the first design state and receives as a second input the circuit implementation of the second design state.

6. The method of claim 2, further comprising: determining a difference between the second parameter value and the first parameter value, wherein the machine learning model is updated based on the difference.

7. The method of claim 2, further comprising: processing the first design state using a circuit synthesis tool, wherein the circuit synthesis tool outputs the first parameter value associated with the first design state; processing the second design state using the circuit synthesis tool, wherein the circuit synthesis tool outputs the second parameter value associated with the second design state; and making the comparison between the second parameter value and the first parameter value to determine which of the first parameter value and the second parameter value is closer to the target.

8. The method of claim 7, wherein the circuit synthesis tool processes the first design state at least partially concurrent with the machine learning model outputting an action to generate the second design state of the data path circuit.

9. The method of claim 2, further comprising: determining that the second design state is an invalid state before determining the second parameter value; and modifying the second design state of the data path circuit to produce an updated second design state that is valid; wherein the second parameter value is determined from the updated second design state.

10. The method of claim 2, further comprising: generating a first parallel prefix graph that represents the first design state; and generating a grid representation of the first parallel prefix graph, wherein the grid representation of the first parallel prefix graph is provided as an input to the machine learning model, and wherein the machine learning model outputs an action or actions to construct a grid representation of a final parallel prefix graph that represents the final design state.

11. The method of claim 1, wherein the data path circuit comprises a parallel prefix circuit or a parallel prefix circuit with additional circuit logic connected to its inputs or outputs.

12. The method of claim 1, wherein the first parameter value and the final parameter value each represents at least one of: a prediction of an area associated with the data path circuit, a delay associated with the data path circuit, a power consumption associated with the data path circuit, or any combination thereof.

13. The method of claim 12, wherein the first parameter value and the final parameter value each represents at least one of: a prediction of a weighted value of the area associated with the data path circuit, the delay associated with the data path circuit, or the power consumption associated with the data path circuit.
14. The method of claim 1, wherein the machine learning model iteratively modifies a design state of the data path circuit from the first design state to the final design state, wherein each iteration comprises at least one of removing or adding a node of a graph of the data path circuit, wherein each node of the graph is associated with one or more components of the data path circuit.
15. A system comprising: a first processing device to: receive a first design state of a data path circuit; and process the first design state of the data path circuit using a machine learning model, wherein the machine learning model is to generate an output comprising a second design state of the data path circuit that is a modification of the first design state; and a second processing device configured to: determine a first parameter value on a first parameter curve of a first physical data path circuit synthesized from the first design state and a second parameter value on a second parameter curve of a second physical data path circuit synthesized from the second design state, wherein the first parameter curve and the second parameter curve indicate a plurality of parameter values at one or more areas of the first physical data path circuit and the second physical data path circuit respectively; determine a reward associated with the output based on a comparison of the first parameter value and the second parameter value; and wherein the first processing device is further to update the machine learning model based on the reward; and wherein after one or more updates to the machine learning model, the machine learning model is to output a final design state of the data path circuit, the final design state of the data path circuit being associated with a final parameter value that is closer to a target than a first parameter value associated with the first design state.
16. The system of claim 15, wherein the first processing device is further to: process the first design state or the second design state of the data path circuit using the updated machine learning model; and output a third design state of the data path circuit that is a modification of the first design state or the second design state.
17. The system of claim 15, wherein the second processing device is further to: process the first design state using a second model, wherein the second model outputs the first parameter value associated with the first design state; process the second design state using the second model, wherein the second model outputs the second parameter value associated with the first design state; and compare the second parameter value and the first parameter value to determine which of the first parameter value and the second parameter value is closer to the target.
18. The system of claim 15, wherein at least one of the second processing device or a third processing device is further configured to: generate the first physical data path circuit of the first design state using a circuit synthesis tool; and generate the second physical data path circuit of the second design state using the circuit synthesis tool.
19. The system of claim 15, wherein the first processing device comprises a graphical processing unit (GPU) and the second processing device comprises a central processing unit (CPU).
20. A non-transitory computer-readable medium storing instructions thereon, wherein the instructions, when executed by a processing device, cause the processing device to: receive a first design state of a data path circuit; provide the first design state of the data path circuit as an input to a machine learning model; and apply one or more reinforcement learning techniques using the machine learning model to cause the machine learning model to output a final design state of the data path circuit, the one or more reinforcement learning techniques comprising: determining a first parameter value on a first parameter curve of a first physical data path circuit synthesized from the first design state, wherein the first parameter curve indicates a plurality of parameter values at one or more areas of the first physical data path circuit; and determining the final design state of the data path circuit, the final design state associated with a final parameter value, wherein the final design state is closer to a target than the first parameter value associated with the first design state.
-