

# US Patent & Trademark Office

## Patent Public Search | Text View

---

United States Patent	12387619
Kind Code	B2
Date of Patent	August 12, 2025
Inventor(s)	Dillon; Tom

---

### Systems and methods for grapheme-phoneme correspondence learning

---

#### Abstract

Systems and methods are described for grapheme-phoneme correspondence learning. In an example, a display of a device is caused to output a grapheme graphical user interface (GUI) that includes a grapheme. Audio data representative of a sound made by the human user is received based on the grapheme shown on the display. A grapheme-phoneme model can determine whether the sound made by the human corresponds to a phoneme for the displayed grapheme based on the audio data. The grapheme-phoneme model is trained based on augmented spectrogram data. A speaker is caused to output a sound representative of the phoneme for the grapheme to provide the human with a correct pronunciation of the grapheme in response to the grapheme-phoneme model determining that the sound made by the human does not correspond to the phoneme for the grapheme.

---

<b>Inventors:</b>	<b>Dillon; Tom (Washington, DC)</b>
<b>Applicant:</b>	<b>617 Education Inc. (Washington, DC)</b>
<b>Family ID:</b>	<b>1000008749197</b>
<b>Assignee:</b>	<b>617 EDUCATION INC. (Washington, DC)</b>
<b>Appl. No.:</b>	<b>18/653629</b>
<b>Filed:</b>	<b>May 02, 2024</b>

#### Prior Publication Data

<b>Document Identifier</b>	<b>Publication Date</b>
US 20240282213 A1	Aug. 22, 2024

#### Related U.S. Application Data

continuation parent-doc US 18152625 20230110 US 12008921 child-doc US 18653629  
us-provisional-application US 63363406 20220422

---

## Publication Classification

**Int. Cl.:** **G09B7/04** (20060101); **G06F3/16** (20060101); **G09B19/04** (20060101); **G10L15/02** (20060101); **G10L15/06** (20130101); **G10L15/22** (20060101); **G10L25/18** (20130101); **G10L25/30** (20130101)

**U.S. Cl.:**

**CPC** **G09B7/04** (20130101); **G06F3/167** (20130101); **G09B19/04** (20130101); **G10L15/02** (20130101); **G10L15/063** (20130101); **G10L15/22** (20130101); **G10L25/18** (20130101); **G10L25/30** (20130101); G10L2015/025 (20130101); G10L2015/225 (20130101)

## Field of Classification Search

**CPC:** G10L (2015/025); G10L (2015/225); G06F (3/167); G09B (19/04); G09B (7/04)

---

## References Cited

### U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
5920838	12/1998	Mostow et al.	N/A	N/A
6963841	12/2004	Handal	704/251	G09B 19/04
7280964	12/2006	Wilson	704/251	G10L 15/187
9697822	12/2016	Naik	N/A	G10L 15/22
10199051	12/2018	Binder et al.	N/A	N/A
10210861	12/2018	Arel	N/A	G06N 20/00
10311980	12/2018	Kim et al.	N/A	N/A
10559299	12/2019	Arel	N/A	G06N 7/01
11004443	12/2020	Yu	N/A	G06N 3/044
11158320	12/2020	Scanlon	N/A	N/A
11282512	12/2021	Yang et al.	N/A	N/A
11308939	12/2021	Gao et al.	N/A	N/A
11373633	12/2021	Chicote	N/A	G06F 40/30
11410639	12/2021	Trueba	N/A	G10L 25/18
11410684	12/2021	Klimkov	N/A	G10L 25/78
11488489	12/2021	Burdis et al.	N/A	N/A
11580955	12/2022	Meng et al.	N/A	N/A
11676572	12/2022	Peddinti	704/260	G10L 13/08
11682318	12/2022	Jochim et al.	N/A	N/A
11741941	12/2022	Richards	N/A	N/A
12008921	12/2023	Dillon	N/A	G10L 15/22
2020/0286490	12/2019	Scanlon	N/A	N/A

### OTHER PUBLICATIONS

Mostow, "Why and How Our Automated Reading Tutor Listens," Project LISTEN ([www.cs.cmu.edu/~listen](http://www.cs.cmu.edu/~listen)), School of Computer Science, Carnegie Mellon University, Jun. 2012, 10 pages. cited by applicant

"Amira Is Science Of Reading," Amira Learning—White Paper, <https://amiralearning.com/documents/White-Paper-Amira-Is-The-Science-Of-Reading.pdf>, Jun. 15,

2023, 31 pages. cited by applicant

Mostow, "Lessons from Project LISTEN: What Have We Learned from a Reading Tutor That Listens?" AIED 2013: Artificial Intelligence in Education, pp. 557-558. cited by applicant

Mostow et al., "Generating Instruction Automatically for the Reading Strategy of Self-Questioning," Proceedings of the 2009 conference on Artificial Intelligence in Education: Building Learning Systems that Care: From Knowledge Representation to Affective Modelling, Jul. 2009, pp. 465-472. cited by applicant

"Research Foundations: Evidence Base," Amira Learning, Houghton Mifflin Harcourt, <https://www.hmhco.com/research/amira-learning-research-evidence-base>, 54 pages. cited by applicant

---

*Primary Examiner:* Thai; Xuan M

*Assistant Examiner:* Zaman; Sadaruz

*Attorney, Agent or Firm:* Vorys, Sater, Seymour and Pease LLP

---

## **Background/Summary**

RELATED APPLICATION (1) This application is a continuation of U.S. patent application Ser. No. 18/152,625, filed Jan. 10, 2023, which claims the benefit and priority of U.S. Provisional Application No. 63/363,406, filed Apr. 22, 2022, each of which is incorporated herein by reference in its entirety.

### **TECHNICAL FIELD**

(1) The present disclosure relates to systems and methods for speech repetition and more particularly to grapheme-phoneme correspondence learning.

### **BACKGROUND**

(2) A grapheme is a written symbol that represents a sound (e.g., phoneme). This can be a single letter, or could be a sequence of letters. When a human says a sound corresponding to a letter, for example, a spoken letter "t", this is a phoneme, and the written letter "t" is a grapheme. A digraph is a pair of characters or letters used together to represent a single sound, such as "ch" in English. A grapheme that consists of two letters is called a digraph, while one with three is called a trigraph. A collection of graphemes and/or digraphs can be used to represent a word and syllable. Phonemes can be combined to form syllables and words. For example, the word "kitty" is composed of four distinct sounds, or phonemes. Phonemes of graphemes and/or digraphs can be combined to represent morphemes, which is a small unit having a meaning (e.g., a base word, prefix, or suffix). Graphemes can also be arranged in no particular order to form a non-sensical word (e.g., "vang").

### **SUMMARY**

(3) In an example, a system can include memory to store machine-readable instructions, and one or more processors to access the memory and execute the machine-readable instructions. The machine-readable instructions can include a spectrogram generator that can be programmed to provide spectrogram data based on audio data representative of one or more sounds corresponding to one or more phonemes. The machine-readable instructions can further include a data augmentor that can be programmed to augment the spectrogram data to provide augmented spectrogram data, and a trainer that can be programmed to train a grapheme-phoneme model during a first training phase based on a first portion of the augmented spectrogram data, and re-train the grapheme-phoneme model during a second training phase based on a second portion of the augmented spectrogram data to provide a trained grapheme-phoneme model for determining whether a sound

made by a human is representative of a phoneme for a grapheme.

(4) In yet another example, a device can include a display, a speaker, memory to store machine-readable instructions, and one or more processors to access the memory and execute the machine-readable instructions. The machine-readable instructions can include a trained machine learning (ML) model that can be programmed to determine whether a sound made by a human corresponds to a phoneme for a grapheme displayed on the display, and a grapheme-phoneme module programmed to cause the speaker to output a sound representative of the phoneme for the grapheme in response to the trained ML model determining that the sound made by the human does not match the phoneme for the grapheme on the display.

(5) In a further example, a method can include causing a display of a device to output a grapheme graphical user interface (GUI) that includes a grapheme, receiving audio data representative of a sound made by the human in response to the grapheme being displayed on the display, providing the audio data to a trained neural network to determine whether the sound made by the human corresponds to a phoneme for the grapheme, and causing a speaker of the device to output a sound representative of the phoneme for the grapheme in response to determining that the sound made by the human does not correspond to the phoneme for the grapheme.

(6) In additional example, a computer-implemented system can include a tool that can be configured to output a user-interface display view that shows a user a series of graphemes, prompt the user to say the sound each grapheme makes, and capture one or more spoken responses from the user in an audio file. The system can further include a trained neural network model that can be configured to recognize individual grapheme sounds spoken out loud in isolation. The tool can output the audio file to the trained neural network model to evaluate whether a response was correct or mistaken. The tool can include a feedback mechanism that can be configured to provide modeling and repetition to the user when a mistaken response is detected.

(7) Further embodiments, features, and advantages of the invention, as well as the structure and operation of the various embodiments of the invention are described in detail below with reference to accompanying drawings.

---

## Description

### BRIEF DESCRIPTION OF THE DRAWINGS

(1) FIG. 1 is an example of a computing system for training a grapheme-phoneme model for grapheme-phoneme correspondence learning.

(2) FIG. 2 is an example of a user device for grapheme-phoneme correspondence learning.

(3) FIG. 3 is an example of a grapheme graphical user interface (GUI).

(4) FIG. 4 is an example of a grapheme-phoneme model.

(5) FIG. 5 is an example of a method for training a grapheme-phoneme model.

(6) FIG. 6 is an example of a method for grapheme-phoneme correspondence learning.

(7) FIG. 7 depicts an example computing environment.

### DETAILED DESCRIPTION

(8) Letter-sound correspondence, or a relationship of graphemes (e.g., in an alphabet) to phonemes (e.g., sounds) produced, is a component of alphabetic principle and learning to read. Letter-sound correspondence refers to an identification of sounds associated with individual letters and letter combinations. For example, teaching students' letter-sound correspondence is part of the curriculum and educational objectives of individuals, schools, teachers, learning centers, and other educational entities. Letter-sound correspondence (the linking in a brain of an abstract symbol ("A") with the sound: "/ah/") is learned through repetition. To learn proper letter-sound correspondence requires immediate correction (e.g., when a human makes a mistake) and modeling (e.g., demonstrating a correct sound for a letter for the human to repeat). Generally, an instructor

(e.g., teacher) is assigned a number of students (e.g., twenty or more students) and demonstrates the letter-sound correspondence to all the students, however, is unable to confirm that individual students are producing the correct letter sound.

(9) The present disclosure describes automated grapheme-phoneme practice with real-time feedback and modeling. The term “grapheme” as used herein can refer to a symbol (e.g., a letter), a combination of symbols (e.g., letters defining a digraph, a trigraph, or a blend (e.g., a consonant blend), etc.), and/or a word (e.g., sensical (e.g., having a known meaning) or non-sensical, a blend, a syllable, a morpheme, etc.). While the term “phoneme” as used herein can refer to a single unit of sound, a combination of units of sounds (e.g., combined to represent a word and/or syllable). Thus, in some examples, the system and methods described herein can be used for symbol-sound, word-sound practice, or any type of grapheme-sound association. In embodiments, computer-implemented methods, systems and devices for automated letter sound practice with real-time feedback and modeling are provided. In one embodiment, a computer-implemented tool outputs a user-interface display view that shows students a series of letters. Students are prompted to say the sound each letter makes, one after another. Beneath the user interface is a deep neural network that is trained to recognize individual letter sounds spoken out loud in isolation. This neural network model is linked to an error correction protocol—in the event of a mistake, the software immediately offers modeling (saying the letter sound correctly) followed by an opportunity for the student to repeat the sound correctly. A feedback mechanism is used to feed students a mix of letters with which they are struggling and those that they have mastered.

(10) In one embodiment, the tool can be provided by elementary school teachers, tutors, and parents to students to help students learn their letter sounds independently. Learning letter sounds is the foundation for all other reading skills. The tool can be used as part of an educational program or as part of an intervention to improve the phonics skills of students in need. The tool can also be used for diagnostic, formative, and summative assessment.

(11) One advantage is that a tool with a feedback mechanism as described herein can recognize and/or classify a spoken letter sound and provide immediate feedback in the event of an error, modeling of the correct sound, and an opportunity for the student to repeat the letter.

(12) Examples are described herein for speech repetition for grapheme-phoneme correspondence learning. For example, a grapheme-phoneme model can be trained to determine whether a sound made by a human corresponds to a phoneme for a grapheme based on augmented spectrogram data. A grapheme-phoneme module can output a grapheme GUI that includes a respective grapheme, which can be rendered on a display of a user device. Audio data representative of a sound made by the human is received based on the respective grapheme shown on the display. The grapheme-phoneme model can determine whether the sound made by the human corresponds to a given phoneme for the respective grapheme based on the audio data. The grapheme-phoneme module can cause a speaker to output a sound representative of the given phoneme for the respective grapheme to provide the human with a correct pronunciation of the respective grapheme in response to the grapheme-phoneme model determining that the sound made by the human does not correspond to the phoneme for the grapheme.

#### Computing Platform

(13) FIG. 1 is an example of a computing platform **100**. The computing platform **100** can be any type of computing device having one or more processors **102** and memory **104**. For example, the computing device can be a workstation, mobile device (e.g., a mobile phone, personal digital assistant, tablet or laptop), computer, server, computer cluster, server farm, game console, set-top box, kiosk, embedded device or system, or other device having at least one processor and computer-readable memory. In addition to at least one processor and memory, such a computing device may include software, firmware, hardware, or a combination thereof. Software may include one or more applications and an operating system. Hardware can include, but is not limited to, a processor, memory and user interface display or other input/output device.

## Grapheme-Phoneme Model Training

(14) As described herein, the computing platform **100** can be used for training a grapheme-phoneme model **106**, which as described herein is used for grapheme-phoneme correspondence learning (e.g., letter to sound, word to sound, and other language correspondences). The grapheme-phoneme model **106** can be trained by a trainer **108**, as described herein. By way of example, the grapheme-phoneme model **106** can be implemented as a neural network model, however, in other examples, a different ML model may be used and the trainer **108** can be configured to support training of this model. In some examples, the grapheme-phoneme model **106** may be implemented as a deep neural network, such as a residual network (e.g., a residual convolution network, such as MobileNetV2). Examples of neural networks can include a perceptron model, a feed forward neural network, a multilayer perceptron model, a convolutional neural network, a radial Basis function neural network, a recurrent neural network, a long short term memory neural network, a sequence to sequence model, and a modular neural network.

(15) By way of example, the memory **104** can be implemented, for example, as a non-transitory computer storage medium, such as volatile memory (e.g., random access memory), non-volatile memory (e.g., a hard disk drive, a solid-state drive, a flash memory, or the like) or a combination thereof. The processor **102** could be implemented, for example, as a processor core. The memory **104** can store machine-readable instructions that can be retrieved and executed by the processor **102** to implement training of the grapheme-phoneme model **106**. For example, the computing platform **100** could be implemented in a computing cloud. In such a situation, features of the computing platform **100**, such as the processor **102**, the memory **104**, and a network interface **110** could be representative of a single instance of hardware or multiple instances of hardware with applications executing across multiple of instances (e.g., distributed) of hardware (e.g., computers, routers, memory, processors, or a combination thereof). Alternatively, the computing platform **100** could be implemented on a single dedicated server or workstation.

(16) The network interface **110** (e.g., a network interface card) can be configured to communicate with a number of devices **112**, as shown in FIG. 1. In some examples, the devices **112** are user devices, such as described herein (e.g., a user device **200**, as shown in FIG. 2). The computing platform **100** can communicate with the devices over a network **114**. The network **114** can include a wired and/or wireless network. For example, the network **114** can include a public network (e.g., the Internet), a private network (e.g., a local area network (LAN)), or a combination thereof (e.g., a virtual private network). The computing platform **100** can provide the grapheme-phoneme model **106** using the network interface **110** over the network **114** to each device **112**. Each device **112** can employ the grapheme-phoneme model **106** for grapheme-phoneme correspondence learning. For example, during grapheme-phoneme correspondence learning, the grapheme-phoneme model **106** can be used to determine whether a sound made by a human corresponds to a phoneme for a respective grapheme. As described herein, the respective grapheme can be rendered on a display, for example, of the user device.

(17) The computing platform **100** can include an input device **112**, such as a keyboard, a mouse, and/or the like. The input device **112** can be used to provide relevant training parameters (e.g., weight values) for the trainer **108** during training of the grapheme-phoneme model. In some examples, the input device **112** can be used to initiate a tester **114** following training of the grapheme-phoneme model **106** to verify (e.g., test) a performance of the grapheme-phoneme model **106** to determine whether the grapheme-phoneme model is making accurate predictions (e.g., within a defined or specified range, for example, based on user input via the input device **112**).

(18) In some examples, the memory **102** includes a spectrogram generator **116** and a data augmentor **118**. The spectrogram generator **116** can be programmed to receive or retrieve audio data **120**, which can be stored in the memory **102**, or remotely (e.g., on another device). The audio data **120** can represent one or more sounds corresponding to one or more phonemes for corresponding graphemes. The recordings may be stored as m4a files, or in another file format. The

spectrogram generator **116** can be programmed to provide the spectrogram data **122** based on the audio data **120**, which can be provided or received by the data augmentor **118**. By way of example, the spectrogram generator **116** can be programmed to transform the audio data **120** into a Mel-Scaled spectrogram to provide the spectrogram data **122**. In some examples, the spectrogram data **122** includes spectrograms that are labeled (e.g., identifying a phoneme for a corresponding grapheme). For example, a user can employ the input device **112** to label corresponding spectrograms of spectrogram data **122**. Because in some instances the spectrogram data **122** includes labeled spectrograms, the grapheme-phoneme model **106** can be trained using supervised learning by the trainer **108**.

(19) The data augmentor **118** can be programmed to augment the spectrogram data **122** to provide augmented spectrogram data **124** for use in training and testing of the grapheme-phoneme model **106**. For example, the data augmentor **118** can be programmed to randomly augment the spectrogram data **122** to provide the augmented spectrogram data **124**. The augmentation can include scaling, shifts, noise and blanking using a spec-augmentation method, for example, as described in “SpecAugment A Simple Data Augmentation Method for Automatic Speech Recognition,” Daniel S. Park et al. Thus, in some examples, the data augmentor **118** can include a number of augmentation components (e.g., modules) for augmentation of the spectrogram data **122** to provide the augmented spectrogram data **124**.

(20) The augmented spectrogram data **124** can include augmented spectrogram training data **126** (referred to herein as “training data”) and augmented spectrogram testing data **128** (“referred to herein as “testing data”). In some examples, the data augmentor **118** can be programmed to tag (e.g., flag) a portion of the augmented spectrogram data **124** as the training data **126**, and another portion of the augmented spectrogram data **124** as the testing data **128**. In some examples, the flagging of the augmented spectrogram data **124** to provide the testing and training data **126** and **128** can be based on user input at the input device **112**.

(21) For example, the trainer **108** can be programmed to train the grapheme-phoneme model **106** over a number of training phases, such as two-training phases. During a first training phase, trainer **108** can be programmed to train the grapheme-phoneme model **106** based on a first portion of the training data **126** and re-train the grapheme-phoneme model **106** during a second training phase based on a second portion of the training data **126** to provide a trained grapheme-phoneme model for determining whether a sound made by a human is representative of a phoneme for a grapheme. In some examples, learning algorithms may be used by the trainer **108** during training of the grapheme-phoneme model **106**. For example, Stochastic Gradient Descent and Adam algorithms, which are gradient descent optimizers, can be used by the trainer **108** during the training of the grapheme-phoneme model **106**. The tester **114** can be programmed to execute the grapheme-phoneme model **106** to predict a corresponding grapheme-phoneme relationship based on the testing data **128**, as shown in FIG. 1, to verify a performance of the grapheme-phoneme model **106**. The grapheme-phoneme model **106** after training can be provided to a corresponding device **112** for use in grapheme-phoneme correspondence learning, such as described herein with respect to FIG. 2.

(22) Grapheme-Phoneme Correspondence Learning

(23) FIG. 2 is an example of a user device **200** for grapheme-phoneme correspondence learning. The user device **200** may be any type of computing device, such as a portable computing device (e.g., mobile phone, tablet, and/or the like), or stationary device (e.g., a desktop computer) that the user can access or use for learning grapheme-phoneme correspondences. In some examples, the user device **200** may be implemented on a device similar to a computing device as described herein with respect to FIG. 1. The user device **200** can include a processor **202** and a memory **204**. By way of example, the memory **204** can be implemented, for example, as a non-transitory computer storage medium, such as volatile memory (e.g., random access memory), non-volatile memory (e.g., a hard disk drive, a solid-state drive, a flash memory, or the like) or a combination thereof.

The processor **202** could be implemented, for example, as a processor core. The memory **204** can store machine-readable instructions that can be retrieved and executed by the processor **202** to implement grapheme-phoneme correspondence learning.

(24) For example, the user device **200** could be implemented in a computing cloud. In such a situation, features of the user device **200**, such as the processor **202**, the memory **204**, and a network interface **224** could be representative of a single instance of hardware or multiple instances of hardware with applications executing across multiple of instances (e.g., distributed) of hardware (e.g., computers, routers, memory, processors, or a combination thereof). Alternatively, the user device **200** could be implemented on a single dedicated server or workstation.

(25) The memory **204** can include a grapheme-phoneme module **206** that can be programmed for grapheme-phoneme correspondence learning. The grapheme-phoneme module **206** can communicate with a grapheme-phoneme database **208**. The grapheme-phoneme database **206** can store a number of graphemes. For example, the grapheme-phoneme database **206** can include an alphabet, such an English alphabet, Arabic alphabet, Chinese alphabet, or a different alphabet. The grapheme-phoneme database **206** can include a number of letters, sequence of letters representing a sound, words, syllables, and/or morphemes.

(26) For example, to teach a human a grapheme-phoneme correspondence (e.g., one or more letter and/or word sound correspondences), the grapheme-phoneme module **206** can be programmed to generate a grapheme GUI **210** that includes one or more graphemes (e.g., letters, words, etc.) for pronunciation by the user. The grapheme-phoneme module **206** can identify the one or more graphemes for generating the grapheme GUI **210** based on the grapheme-phoneme database **208**. For example, the grapheme-phoneme module **206** can identify one or more sequential graphemes (e.g., neighboring letters in an alphabet), or one or more random graphemes (e.g., non-neighboring letters in an alphabet), or a combination thereof.

(27) The user device **200** can include or communicate with a display **212** for rendering the grapheme GUI **210**. The display **212** can correspond to an output device, such as a screen, a touch-screen display, a monitor, a printer, a projector, wearable reality glasses, or another type of display. In some examples, the user device **200** can include or communicate with an input device **214** for interacting with elements of the GUI **210**. For example, the input device **214** can include a touchscreen, a keyboard, a mouse, a stylus pen, and/or the like. In some instances, the display **212** and the input device **214** may be implemented as a single device.

(28) In some examples, the grapheme GUI **210** may prompt the user to select one of the one or more graphemes rendered on the display **212** for pronunciation. A selected grapheme can be emphasized on the grapheme GUI **210** to distinguish the selected grapheme from other graphemes rendered on the display **212** so that the user can be visually alerted to a proper grapheme for grapheme-phoneme correspondences learning. The grapheme-phoneme module **206** can be programmed to receive grapheme selection data identifying the selected grapheme for grapheme-phoneme correspondences learning, which can be generated in response to the user (e.g., via the input device **210**). In some examples, the grapheme-phoneme module **206** can identify the selected grapheme of the one or more graphemes for pronunciation.

(29) The human, in response to being prompted to pronounce the selected grapheme, can speak the selected grapheme, which is represented as a user sound **216** in the example of FIG. 2. A microphone **218** can capture the user sound **216** and generate audio data **220** that is representative of the user sound **216**. The audio data **220** can be provided to a trained model **222** corresponding to the grapheme-phoneme model **106**, as shown in FIG. 1. Thus, reference can be made to the example of FIG. 1 in the example of FIG. 2.

(30) In some instances, the user device **200** includes a network interface **224**. The network interface **224** (e.g., a network interface card) can be configured to communicate with other computing platforms via a network (e.g., the network **114**, as shown in FIG. 1). In some examples, the network interface **224** is used to communicate with the computing platform **100**, as shown in FIG. 1, to



receive the grapheme-phoneme model **106**. While the example of FIG. 2 illustrates the trained model **222** as separate from the grapheme-phoneme module **206**, in some instances, the grapheme-phoneme module **206** includes the trained model **222**.

(31) The trained model **222** can process the audio data **220** to determine whether the sound made by the human sound **216** corresponds to a phoneme for the selected grapheme. The trained model **222** can communicate with the grapheme-phoneme module **206** to receive the selected grapheme and use this information to determine whether the sound made by human user sound **216** corresponds to the phoneme for the selected grapheme.

(32) In some instances, the trained model **222** can determine how closely the sound made by the user corresponds to the phoneme for the selected grapheme. That is, the trained model **222** can determine an accuracy or confidence in the accuracy of the pronunciation of the grapheme by the user relative to an actual or baseline pronunciation of the grapheme. The trained model **222** can output sound accuracy data indicating a sound similarity level for a sound made by a human matching or being similar to a phoneme for a grapheme. The sound similarity level can correspond to the accuracy of the pronunciation. Thus, the sound accuracy data can characterize how closely the pronunciation of the grapheme made by the user corresponding to the user sound **216** matches or is similar to the actual or baseline pronunciation of the grapheme corresponding to an actual or baseline phoneme. The accuracy may be represented as a percentage value, a whole number value, or a decimal number value.

(33) In some examples, the sound accuracy data can be provided to the grapheme-phoneme module **206** to update the grapheme GUI **210** to notify the user of the sound similarity level for the sound made by the human for the grapheme. Thus, the grapheme GUI **210** can be updated to visually indicate to the user how well the user is pronouncing the selected grapheme.

(34) In some instances, if the sound similarity level is below a sound similarity threshold (e.g., an accuracy threshold), the grapheme-phoneme module **206** can update the grapheme GUI **210** to alert the user to repeat the selected grapheme. The grapheme-phoneme module **206** can continuously compare the sound similarity level for the selected grapheme and one or more subsequent phonemes and alert the user to repeat the selected grapheme until the quality level is greater than or equal to the sound similarity threshold. While examples are described herein in which the trained model **222** determines the sound similarity level for the selected grapheme, in other examples, the grapheme-phoneme module **206** can be programmed to determine the sound similarity level in a same or similar manner as described herein.

(35) In additional or alternative examples, the grapheme-phoneme module **206** can be programmed to output artificial audio data **226**. The grapheme-phoneme module **206** can be programmed to output the artificial audio data **226** in response to determining that the user sound **216** does not correspond to the phoneme for the selected grapheme rendered on the display **212**. In some examples, the grapheme-phoneme module **206** can be programmed to output the artificial audio data **226** in response to determining that the sound similarity level is not within a given value (e.g., degree, percentage, etc.) of the sound similarity threshold, or is less than the sound similarity threshold.

(36) The artificial audio data **226** can represent sound referred to as an artificial sound **228** that can represent the phoneme for the selected grapheme and thus can be used to provide a correct or proper pronunciation for the selected grapheme. The term “artificial” as used herein relating to sound is used to indicate that the sound is generated by a speaker rather than a human. Thus, in some examples, the artificial sound **228** can represent a machine generated sound, or a previously captured sound for the selected grapheme made by a human.

(37) The artificial audio data **226** can be provided to a speaker **230** of the device **200**. The speaker **230** can convert the artificial audio data **226** into sound energy corresponding to the artificial sound **228**. In some examples, if the user provides the proper pronunciation (e.g., the phoneme) for the selected grapheme, the grapheme-phoneme module **206** can update the grapheme GUI **210**, such

that a different grapheme of the one more graphemes is selected for grapheme-phoneme correspondence learning. In some examples, the grapheme-phoneme module **206** can be programmed to query the grapheme-phoneme database **208** to identify a correct phoneme for the select grapheme and cause the speaker **230** to output the artificial sound **228** based on the identified phoneme. Thus, in some examples, the grapheme-phoneme database **208** can store audio data representative of different sounds corresponding to phonemes for respective graphemes. As such, in some examples, the graphemes in the grapheme-phoneme database **208** can be associated (e.g., logically linked) to audio data representative of a corresponding phoneme.

(38) Grapheme-Phoneme Correspondence Learning Tool

(39) Accordingly, the user device **200** can be implemented as a grapheme-phoneme correspondence learning tool enabling a user to learn through repetition grapheme-phoneme correspondences. The tool employs a feedback mechanism that models a proper sounding of a grapheme (e.g., as the artificial audio data **226**) so that the user can practice pronouncing the grapheme over a number of repetitions to learn a corresponding phoneme for the grapheme. For example, the tool can recognize and/or classify a spoken letter sound and provide feedback in the event the human mispronounces the letter, model the correct sound, and allow the user (e.g., a student) to repeat the letter, and then provide further feedback. In some examples, the tool can be used by elementary school teachers, tutors, and/or parents to assist students or children in learning letter sound independently (e.g., without an immediate instructor). The grapheme-phoneme correspondence learning tool as described herein can be used as part of an educational program or as part of an intervention to improve students in need. In some instances, the tool can be used as an assistance tool for diagnostics.

(40) User-Interface

(41) FIG. 3 is an example of a grapheme GUI **300**, which can correspond to the grapheme GUI **210**, as shown in FIG. 2. Thus, reference can be example of FIGS. 1-2 in example of FIG. 3. The grapheme GUI **300** includes a number of grapheme elements **302-308** representative of a corresponding grapheme. In the example of FIG. 3, the graphemes are English letters, such as “o,” “i,” “p,” “f,” and “s.” For example, to practice a grapheme-phoneme correspondence, the user or the grapheme-phoneme module **206** can select a given grapheme element of the grapheme elements **302-308**. In the example of FIG. 3, the grapheme element **304** is emphasized with a border to indicate the selection of the grapheme element **304**.

(42) In some examples, the grapheme GUI **300** includes a start element **312** and a stop element **314**. A user (e.g., human) can interact with the start element **312** of the grapheme GUI **300** to initiate learning one or more grapheme-phoneme correspondences, and the stop element **314** to terminate or stop learning the one or more grapheme-phoneme correspondences. The grapheme GUI **300** can further include a given number of star elements **316**, which can be used to provide a measure of how many correct grapheme-phoneme correspondences. If a user (e.g., student) gets all five (5) grapheme-phoneme correspondences in a given stage correct, the user can move onto a next stage.

(43) By way of further example, the grapheme-phoneme module **206** can be programmed to output the grapheme GUI **300** with a series of letters. A student can be prompted to speak each letter, one after another in response to the grapheme-phoneme module **206**. For example, the grapheme-phoneme module **206** can output instruction audio data requesting that the student speak a respective letter as identified on the grapheme GUI **300**, such as the grapheme element **304**. The trained model **222**, in some instances, implemented as a deep neural network, beneath the graphical GUI **300** is trained to recognize individual letter sounds spoken out loud (e.g., in isolation) by the student. The trained model **222** is linked to the grapheme-phoneme module **206**. If the student mispronounces the respective letter, the trained model **222** can provide data indicating an incorrect pronunciation, which the grapheme-phoneme module **206** can process to provide a correct pronunciation for the respective letter so that the student can repeat the respective letter. The

grapheme-phoneme module **206** can provide the student a mix of letters with which the student is struggling and letters that the student has mastered.

(44) Example Grapheme-Phoneme Model

(45) FIG. **4** is an example of a grapheme-phoneme model **400**. In some examples, the grapheme-phoneme model **400** is representative of the grapheme-phoneme model **106**, as shown in FIG. **1**. In the example of FIG. **4**, the grapheme-phoneme model **400** is implemented as a neural network model. The grapheme-phoneme model **400** can be trained to determine whether captured audio data (e.g., the captured audio data **218**, as shown in FIG. **2**) corresponds to a phoneme for a grapheme displayed on a grapheme GUI (e.g., the grapheme GUI **210**, as shown in FIG. **2**, or the grapheme GUI **300**, as shown in FIG. **1**). Thus, reference can be example of FIGS. **1-3** in example of FIG. **4**. For example, the grapheme-phoneme model **400** can be implemented as a residual convolution network (RCN) model, such as a MobileNetV2 RCN model, and in other examples, a combination of a Wav2Vec model and a feed-forward neural network model. The grapheme-phoneme model **400** can be implemented in TensorFlow 2.

(46) For example, the grapheme-phoneme model **400** can include an input layer **402** (identified as “L1” in the example of FIG. **4**). The input layer **402** can have a dimensionality similar to a dimensionality of a spectrogram image of spectrogram data **404**. The spectrogram data **404** can correspond to the spectrogram data **122**, as shown in FIG. **1**. The spectrogram image can include a plurality of pixels that can have a defined bit-width, and each node of the input layer **402** can be programmed to process a set of pixels of the plurality of pixels. In some examples, the input layer **402** corresponds to spectrograms of the spectrogram data **122** and can contain neurons that can accept 32-bit float values.

(47) The grapheme-phoneme model **400** can include an N number of intermediate layers **406** (identified as “L2,” “L3,” “LN” in the example of FIG. **4**), wherein “N” is an integer value. Each intermediate layer **406** can have a respective dimensionality (e.g., size) and include nodes for further processing of outputs provided by an upstream layer. In some examples, the intermediate layers **406** can be referred to as hidden layers. Each of the layers **406** may include a number of activation nodes. Nodes can be connected with another node of a similar layer or a different layer, and each connection can have a particular weight. The weights can be determined by the trainer **108**, or based on user input at the input device **112**, as shown in FIG. **1**.

(48) The grapheme-phoneme model **400** can further include an output layer (a feature map output layer) **408** and a classification layer **410**. The output layer **408** can be a feature vector output layer that can provide a feature vector representative of sound differences between two or more phonemes, shown as feature map data **410** in the example of FIG. **4**. The output layer **408** can be used for embedding with vectors for calculating differences between individual pronunciations. In an example, the output layer **408** can include 128 elements and the classification layer can include 26 elements (e.g., one per each letter of an alphabet.) The classification layer **410** provides a phoneme class mapping, shown as classifier data **412** in the example of FIG. **4**, for example, based on the feature map data **410**. The phoneme class mapping includes phoneme classes for phonemes. The feature vector map and the phoneme class mapping can be used by the trainer **108** during training of the grapheme-phoneme model **400**, such as during a first training phase of the grapheme-phoneme model **400**.

(49) During a second training phase of the grapheme-phoneme model **400**, the trainer **108** can be programmed to freeze non-output classification layers of the grapheme-phoneme model **400** to train only classification layers. During the second training phase, the grapheme-phoneme model **400** can be fine-tuned to improve the prediction accuracy of the grapheme-phoneme model **400**. The trainer **108** can be programmed to train the grapheme-phoneme model **400** during each of the first and second training phases by minimizing a cost function. In some examples, the grapheme-phoneme model **400** after being trained can be stored on a user device, such as the user device **200**, and used for grapheme-phoneme correspondence learning, such as described herein.

(50) In view of the foregoing structural and functional features described above, example methods will be better appreciated with reference to FIGS. 5-6. While, for purposes of simplicity of explanation, the example methods of FIGS. 5-6 are shown and described as executing serially, it is to be understood and appreciated that the present examples are not limited by the illustrated order, as some actions could in other examples occur in different orders, multiple times and/or concurrently from that shown and described herein. Moreover, it is not necessary that all described actions be performed to implement the methods.

#### (51) Training

(52) FIG. 5 is an example of a method **500** for training a grapheme-phoneme model (e.g., the grapheme-phoneme model **106**, as shown in FIG. 1, the grapheme-phoneme model **400**, as shown in FIG. 4) for determining whether a sound (e.g., the user sound **216**, as shown in FIG. 2) made by a human corresponds to a phoneme for a grapheme rendered on a display of a user device (e.g., the user device **200**, as shown in FIG. 2). The method **500** can be implemented by a computing platform, such as the computing platform **100**, as shown in FIG. 1. Thus, reference can be made to the example of FIGS. 1-4 in the example of FIG. 5.

(53) The method **500** can begin at **502** by providing (e.g., via the spectrogram generator **116**, as shown in FIG. 1) spectrogram data (e.g., the spectrogram data **122**, as shown in FIG. 1) based on audio data (e.g., the audio data **120**, as shown in FIG. 1) representative of one or more sounds corresponding to one or more phonemes. At **504**, augmenting (e.g., via the data augmentor **118**, as shown in FIG. 1) the spectrogram data to provide augmented spectrogram data (e.g., the augmented spectrogram data **126**, as shown in FIG. 1). At **506**, training (e.g., via the trainer **108**, as shown in FIG. 1) the grapheme-phoneme model during a first training phase based on a first portion of the augmented spectrogram data. At **508**, re-training (e.g., via the trainer **108**, as shown in FIG. 1) the grapheme-phoneme model during a second training phase based on a second portion of the augmented spectrogram data to provide a trained grapheme-phoneme model (e.g., the trained model **222**, as shown in FIG. 2).

#### (54) Operation

(55) FIG. 6 is an example of a method **600** for grapheme-phoneme correspondence learning (e.g., one or more letter and/or word sound correspondence learning). The method **600** can be implemented by a user device, such as the user device **112**, as shown in FIG. 1, or the user device **200**, as shown in FIG. 2. Thus, reference can be made to the example of FIGS. 1-4 in the example of FIG. 6. The method **600** can begin at **602** by causing a display (e.g., the display **212**, as shown in FIG. 2) of the user device to output a grapheme GUI (e.g., the grapheme GUI **210**, as shown in FIG. 2, or the grapheme GUI **300**, as shown in FIG. 3) with a grapheme.

(56) At **604**, receiving audio data (e.g., the captured audio data **220**, as shown in FIG. 2) representative of a sound (e.g., the user sound **216**, as shown in FIG. 2) made by the human. At **606**, providing the audio data to a grapheme-phoneme model (e.g., the trained model **222**, as shown in FIG. 2) to determine whether the sound made by the human corresponds to a phoneme for the grapheme. At **608**, causing a speaker (e.g., the speaker **230**, as shown in FIG. 2) of the user device to output an artificial sound (e.g., the artificial sound **228**, as shown in FIG. 2) representative of the phoneme for the grapheme in response to the grapheme-phoneme model determining that the sound made by the human does not correspond to the phoneme for the grapheme.

(57) The terminology used herein is for the purpose of describing particular embodiments only and is not intended to limit this disclosure. As used herein, for example, the singular forms “a,” “an,” and “the” are intended to include the plural forms as well, unless the context clearly indicates otherwise. It will be further understood that the terms “contains,” “containing,” “includes,” “including,” “comprises,” and/or “comprising,” and variations thereof, when used in this specification, specify the presence of stated features, integers, steps, operations, elements, and/or components, but do not preclude the presence or addition of one or more other features, integers, steps, operations, elements, components, and/or groups thereof. In addition, the use of ordinal

numbers (e.g., first, second, third, etc.) is for distinction and not counting. For example, the use of “third” does not imply there must be a corresponding “first” or “second.” Also, as used herein, the terms “coupled” or “coupled to” or “connected” or “connected to” or “attached” or “attached to” may indicate establishing either a direct or indirect connection, and is not limited to either unless expressly referenced as such.

(58) While the disclosure has described several exemplary embodiments, it will be understood by those skilled in the art that various changes can be made, and equivalents can be substituted for elements thereof, without departing from the spirit and scope of this disclosure. In addition, many modifications will be appreciated by those skilled in the art to adapt a particular instrument, situation, or material to embodiments of the disclosure without departing from the essential scope thereof. Therefore, it is intended that the invention herein not be limited to the particular embodiments disclosed, or to the best mode contemplated for carrying out this invention, but that the invention will include all embodiments falling within the scope of the appended claims. Moreover, reference in the appended claims to an apparatus or system or a component of an apparatus or system being adapted to, arranged to, capable of, configured to, enabled to, operable to, or operative to perform a particular function encompasses that apparatus, system, or component, whether or not it or that particular function is activated, turned on, or unlocked, as long as that apparatus, system, or component is so adapted, arranged, capable, configured, enabled, operable, or operative.

(59) In view of the foregoing structural and functional description, those skilled in the art will appreciate that portions of the embodiments may be embodied as a method, data processing system, or computer program product. Accordingly, these portions of the present embodiments may take the form of an entirely hardware embodiment, an entirely software embodiment, or an embodiment combining software and hardware, such as shown and described with respect to the computer system of FIG. 7.

(60) Furthermore, portions of the embodiments may be a computer program product on a computer-usable storage medium having computer readable program code on the medium. Any non-transitory, tangible storage media possessing structure may be utilized including, but not limited to, static and dynamic storage devices, hard disks, optical storage devices, and magnetic storage devices.

(61) As an example and not by way of limitation, a computer-readable storage media may include a semiconductor-based circuit or device or other IC (such, as for example, a field-programmable gate array (FPGA) or an ASIC), a hard disk, an HDD, a hybrid hard drive (HHD), an optical disc, an optical disc drive (ODD), a magneto-optical disc, a magneto-optical drive, a floppy disk, a floppy disk drive (FDD), magnetic tape, a holographic storage medium, a solid-state drive (SSD), a RAM-drive, a SECURE DIGITAL card, a SECURE DIGITAL drive, MEMS, nano-technological storage devices, or another suitable computer-readable storage medium or a combination of two or more of these, where appropriate. A computer-readable non-transitory storage medium may be volatile, nonvolatile, or a combination of volatile and non-volatile, where appropriate.

(62) Certain embodiments have also been described herein with reference to block illustrations of methods, systems, and computer program products. It will be understood that blocks of the illustrations, and combinations of blocks in the illustrations, can be implemented by computer-executable instructions. These computer-executable instructions may be provided to one or more processor of a general-purpose computer, special purpose computer, or other programmable data processing apparatus (or a combination of devices and circuits) to produce a machine, such that the instructions, which execute via the processor, implement the functions specified in the block or blocks. Embodiments also have been described herein with the aid of functional building blocks illustrating the implementation of specified functions and relationships thereof. The boundaries of these functional building blocks have been arbitrarily defined herein for the convenience of the description. Alternate boundaries can be defined so long as the specified functions and relationships

thereof are appropriately performed. The breadth and scope of the present invention should not be limited by any of the above-described exemplary embodiments.

(63) These computer-executable instructions may also be stored in computer-readable memory that can direct a computer or other programmable data processing apparatus to function in a particular manner, such that the instructions stored in the computer-readable memory result in an article of manufacture including instructions which implement the function specified in the flowchart block or blocks. The computer program instructions may also be loaded onto a computer or other programmable data processing apparatus to cause a series of operational steps to be performed on the computer or other programmable apparatus to produce a computer implemented process such that the instructions which execute on the computer or other programmable apparatus provide steps for implementing the functions specified in the flowchart block or blocks.

(64) In this regard, FIG. 7 illustrates one example of a computer system **700** that can be employed to execute one or more embodiments of the present disclosure. In some examples, the computer system **700** corresponds to the computing platform **100**, as shown in FIG. 1, and in other examples to the user device **200**, as shown in FIG. 2. Thus, reference can be made to the examples of FIGS. 1-2 in the example of FIG. 7. The computer system **700** can be implemented on one or more general purpose networked computer systems, embedded computer systems, routers, switches, server devices, client devices, various intermediate devices/nodes or standalone computer systems. Additionally, computer system **700** can be implemented on various mobile clients such as, for example, a personal digital assistant (PDA), laptop computer, pager, and the like, provided it includes sufficient processing capabilities.

(65) Computer system **700** includes processing unit **702**, system memory **704**, and system bus **706** that couples various system components, including the system memory **704**, to processing unit **702**. Dual microprocessors and other multi-processor architectures also can be used as processing unit **702**. System bus **706** may be any of several types of bus structure including a memory bus or memory controller, a peripheral bus, and a local bus using any of a variety of bus architectures. System memory **704** includes read only memory (ROM) **710** and random access memory (RAM) **712**. A basic input/output system (BIOS) **714** can reside in ROM **710** containing the basic routines that help to transfer information among elements within computer system **700**.

(66) Computer system **700** can include a hard disk drive **716**, magnetic disk drive **718**, e.g., to read from or write to removable disk **720**, and an optical disk drive **722**, e.g., for reading CD-ROM disk **724** or to read from or write to other optical media. Hard disk drive **716**, magnetic disk drive **718**, and optical disk drive **722** are connected to system bus **706** by a hard disk drive interface **726**, a magnetic disk drive interface **728**, and an optical drive interface **730**, respectively. The drives and associated computer-readable media provide nonvolatile storage of data, data structures, and computer-executable instructions for computer system **700**. Although the description of computer-readable media above refers to a hard disk, a removable magnetic disk and a CD, other types of media that are readable by a computer, such as magnetic cassettes, flash memory cards, digital video disks and the like, in a variety of forms, may also be used in the operating environment; further, any such media may contain computer-executable instructions for implementing one or more parts of embodiments shown and described herein.

(67) A number of program modules may be stored in drives and RAM **710**, including operating system **732**, one or more application programs **734**, other program modules **736**, and program data **738**. The application programs **734** and program data **738** can include functions and methods programmed for training a grapheme-phoneme model and/or learning grapheme-phoneme correspondences, such as shown and described herein. A user may enter commands and information into computer system **700** through one or more input devices **740**, such as a pointing device (e.g., a mouse, touch screen), keyboard, microphone, joystick, game pad, scanner, and the like. These and other input devices **740** are often connected to processing unit **702** through a corresponding port interface **742** that is coupled to the system bus, but may be connected by other

interfaces, such as a parallel port, serial port, or universal serial bus (USB). One or more output devices **744** (e.g., display, a monitor, printer, projector, or other type of displaying device) is also connected to system bus **706** via interface **746**, such as a video adapter.

(68) Computer system **700** may operate in a networked environment using logical connections to one or more remote computers, such as remote computer **748**. Remote computer **748** may be a workstation, computer system, router, peer device, or other common network node, and typically includes many or all the elements described relative to computer system **700**. The logical connections, schematically indicated at **750**, can include a local area network (LAN) and a wide area network (WAN). When used in a LAN networking environment, computer system **700** can be connected to the local network through a network interface or adapter **752**. When used in a WAN networking environment, computer system **700** can include a modem, or can be connected to a communications server on the LAN. The modem, which may be internal or external, can be connected to system bus **706** via an appropriate port interface. In a networked environment, application programs **734** or program data **738** depicted relative to computer system **300**, or portions thereof, may be stored in a remote memory storage device **754**.

(69) What have been described above are examples. It is, of course, not possible to describe every conceivable combination of components or methods, but one of ordinary skill in the art will recognize that many further combinations and permutations are possible. Accordingly, the present disclosure is intended to embrace all such alterations, modifications, and variations that fall within the scope of this application, including the appended claims. Where the disclosure or claims recite “a,” “an,” “a first,” or “another” element, or the equivalent thereof, it should be interpreted to include one or more than one such element, neither requiring nor excluding two or more such elements. As used herein, the term “includes” means includes but not limited to, the term “including” means including but not limited to. The term “based on” means based at least in part on.”

## Claims

1. A system for teaching grapheme-phoneme correspondence comprising: a data augmentor implemented on at least one processor and configured to augment the spectrogram data for one or more sounds corresponding to one or more phonemes to provide augmented spectrogram data; and a trainer implemented on at least one processor and configured to train a grapheme-phoneme model for grapheme-phoneme correspondence learning based on the augmented spectrogram data so that the grapheme-phoneme model determines whether a sound made by a human is representative of a phoneme for a grapheme.
2. The system of claim 1, wherein the grapheme-phoneme model is a neural network model.
3. The system of claim 2, wherein the trainer is configured to train the neural network model during a first training phase based on a first portion of the augmented spectrogram data, and re-train the neural network model during a second training phase based on a second portion of the augmented spectrogram data to provide a trained neural network model.
4. The system of claim 3, wherein the neural network model comprises a plurality of layers, the plurality of layers including a feature vector output layer to provide a feature vector representative of sound differences between two or more phonemes, and the trainer is configured to train the neural network model based on the feature vector.
5. The system of claim 4, wherein the plurality of layers further include at least one output classification layer that provides a phoneme class mapping, the phoneme class mapping comprising phoneme classes for phonemes, and the trainer is configured to train the neural network model further based on the phoneme class mapping.
6. The system of claim 5, wherein the trainer is configured to train the neural network model during each of the first and second training phases by minimizing a cost function.

7. The system of claim 2, further comprising a tester, and wherein the augmented spectrogram data includes augmented spectrogram testing data, and the tester is configured to run the neural network model to predict a corresponding grapheme-phoneme relationship based on the augmented spectrogram testing data.
8. The system of claim 1, wherein the grapheme-phoneme model is stored in a memory of a user device, or a cloud computing environment, the user device or the cloud computing environment comprising one or more processors to access the memory and execute machine readable instructions to: receive audio data representative of the sound made by the human during the grapheme-phoneme correspondence learning; and determine using the grapheme-phoneme model whether the sound made by the human is representative of a phoneme for the grapheme corresponding to determining whether the human correctly pronounced the grapheme.
9. The system of claim 8, wherein the machine readable instructions of the user device or the cloud computing environment further comprise a grapheme-phoneme module, and the grapheme-phoneme model is to provide an indication to the grapheme-phoneme module that the sound made by the human does not correspond to the phoneme for the grapheme.
10. The system of claim 9, wherein the user device comprises a speaker, and the grapheme-phoneme module is to query a grapheme-phoneme database to identify second audio data representative of the phoneme for the grapheme and cause the speaker to output a sound representative of the phoneme based on the second audio data to provide the human with a correct pronunciation.
11. The system of claim 10, wherein the grapheme-phoneme module is to output a grapheme graphical user interface (GUI) that includes the grapheme for rendering on a display of the user device.
12. A device for teaching grapheme-phoneme correspondence comprising memory to store machine-readable instructions; and one or more processors to access the memory and execute the machine-readable instructions, the machine-readable instructions comprising: a machine learning (ML) model configured to determine whether the sound made by the human corresponds to a phoneme for a grapheme; and a grapheme-phoneme module configured to cause a speaker to output a sound representative of the phoneme for the grapheme shown on the display in response to the ML model determining that the sound made by the human does not match the phoneme for the grapheme.
13. The device of claim 12, wherein the grapheme-phoneme module is to query a grapheme-phoneme database to identify the phoneme for the selected grapheme.
14. The device of claim 13, wherein the grapheme-phoneme module is to output a graphical user interface (GUI), the GUI being shown on a display and configured to prompt the human to select the grapheme for pronunciation.
15. The device of claim 14, wherein the ML model is a neural network model and is trained during a first training phase based on a first portion of augmented spectrogram data, and re-trained during a second training phase based on a second portion of the augmented spectrogram data.
16. The device of claim 14, wherein the device is a tablet, a mobile phone, a stationary computer, or a portable computer.
17. A method for teaching grapheme-phoneme correspondence comprising: causing a display of a device to output a graphical user interface (GUI) that includes a grapheme for pronunciation for grapheme-phoneme correspondence learning; providing audio data representative of a sound made by the human based on the grapheme shown on the display to a machine learning (ML) model to determine whether the sound made by the human corresponds to a phoneme for the grapheme; and outputting a sound representative of the phoneme for the grapheme in response to determining that the sound made by the human does not correspond to the phoneme for the grapheme.
18. The method of claim 17, further comprising querying a grapheme-phoneme database to identify the phoneme for the grapheme in response to an indication from the ML model that the sound made



by the human does not correspond to the phoneme for the grapheme.

19. The method of claim 18, wherein the ML model is a neural network model trained during a first training phase of the two-step training phase based on a first portion of augmented spectrogram data, and re-trained during a second training phase of the two-step training phase based on a second portion of the augmented spectrogram data.

20. The method of claim 17, wherein the ML model is configured to determine how closely the sound made by the human corresponds to a phoneme for a grapheme shown on a display of the user device to determine whether the human is correctly pronouncing the grapheme, and wherein sound accuracy data is generated indicating a sound similarity level for the sound made by the human matching or being similar to the phoneme for the grapheme, and the display is to visually indicate if the human is pronouncing the grapheme correctly based on the sound accuracy data.

---