

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12387715
Kind Code	B2
Date of Patent	August 12, 2025
Inventor(s)	Aggarwal; Vikram et al.

Learning offline voice commands based on usage of online voice commands

Abstract

Methods, systems, apparatus, including computer programs encoded on a computer storage medium, for a user device to learn offline voice actions. In one aspect, the method includes actions of detecting, by the user device, an utterance at a first time when the user device is connected to a server by a network, providing, by the user device, the utterance to the server using the network, receiving, by the user device and from the server, an update to the grammar of the user device, detecting, by the user device, a subsequent utterance of the utterance at a second time when the user device is not connected to the server by a network, and in response to detecting, by the user device, the subsequent utterance of the utterance at the second time, identifying, by the user device, an operation to perform based on (i) the subsequent utterance, and (ii) the updated grammar.

Inventors: Aggarwal; Vikram (Palo Alto, CA), Gali; Moises Morgenstern (San Francisco, CA)

Applicant: GOOGLE LLC (Mountain View, CA)

Family ID: 1000008749664

Assignee: GOOGLE LLC (Mountain View, CA)

Appl. No.: 18/380538

Filed: October 16, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20240038219 A1	Feb. 01, 2024

Related U.S. Application Data

Publication Classification

Int. Cl.: **G10L21/00** (20130101); **G10L15/07** (20130101); **G10L15/19** (20130101); **G10L15/22** (20060101); **G10L15/30** (20130101); **G10L15/32** (20130101); **G10L25/00** (20130101)

U.S. Cl.:

CPC **G10L15/07** (20130101); **G10L15/19** (20130101); **G10L15/22** (20130101); **G10L15/30** (20130101); **G10L15/32** (20130101); G10L2015/228 (20130101)

Field of Classification Search

CPC: G10L (15/22); G10L (15/30); G10L (15/063); G10L (2015/223); G10L (15/26); G10L (15/065); G10L (15/183); G10L (2015/0635); G10L (15/00); G10L (17/00); G10L (15/20); G10L (15/18); G10L (15/07); G10L (15/28); G10L (17/04); G10L (2015/0631)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
6195641	12/2000	Loring et al.	N/A	N/A
6873951	12/2004	Lin et al.	N/A	N/A
8185392	12/2011	Strope et al.	N/A	N/A
8370146	12/2012	Schalkwyk et al.	N/A	N/A
8805684	12/2013	Aleksic	704/270.1	G10L 15/07
9070367	12/2014	Hoffmeister et al.	N/A	N/A
9836527	12/2016	Gao et al.	N/A	N/A
10109273	12/2017	Rajasekaram	N/A	G10L 15/183
11170762	12/2020	Aggarwal et al.	N/A	N/A
2002/0077823	12/2001	Fox et al.	N/A	N/A
2003/0046074	12/2002	Ruback et al.	N/A	N/A
2005/0096753	12/2004	Arling et al.	N/A	N/A
2005/0105712	12/2004	Williams et al.	N/A	N/A
2007/0005360	12/2006	Huning et al.	N/A	N/A
2007/0073540	12/2006	Hirakawa et al.	N/A	N/A
2007/0276651	12/2006	Bliss et al.	N/A	N/A
2008/0235021	12/2007	Cross et al.	N/A	N/A
2009/0112605	12/2008	Gupta	N/A	N/A
2010/0179812	12/2009	Jang et al.	N/A	N/A
2012/0179471	12/2011	Newman et al.	N/A	N/A
2013/0073293	12/2012	Jang	704/E21.001	G10L 15/22
2013/0088419	12/2012	Kim et al.	N/A	N/A
2013/0132089	12/2012	Fanty et al.	N/A	N/A
2013/0204967	12/2012	Seo	709/224	H04L 67/63

2013/0211824	12/2012	Tseng	704/10	G06F 40/274
2014/0136013	12/2013	Wolverton et al.	N/A	N/A
2014/0200737	12/2013	Lortz et al.	N/A	N/A
2014/0244259	12/2013	Rosario et al.	N/A	N/A
2014/0274203	12/2013	Ganong, III et al.	N/A	N/A
2014/0337007	12/2013	Waibel et al.	N/A	N/A
2015/0045988	12/2014	Gusikhin et al.	N/A	N/A
2015/0255068	12/2014	Kim	704/246	G10L 17/04
2015/0281401	12/2014	Le et al.	N/A	N/A
2015/0371628	12/2014	Kreifeldt	704/254	G10L 15/02
2016/0094509	12/2015	Ye	N/A	N/A
2016/0109322	12/2015	Eguchi	73/40	G01N 29/036
2016/0155443	12/2015	Khan et al.	N/A	N/A
2016/0179462	12/2015	Bjorkengren	N/A	N/A
2016/0252972	12/2015	Kim	704/9	G10L 15/26
2016/0269524	12/2015	Stottlemeyer	N/A	N/A
2016/0322044	12/2015	Jung	N/A	G10L 15/065
2016/0350320	12/2015	Sung et al.	N/A	N/A
2017/0069312	12/2016	Sundararajan et al.	N/A	N/A
2017/0162204	12/2016	Wilson et al.	N/A	N/A
2017/0242914	12/2016	Gao	N/A	G10L 15/26
2018/0122365	12/2017	Kunitake et al.	N/A	N/A
2018/0174580	12/2017	Kim et al.	N/A	N/A
2018/0217810	12/2017	Agrawal	N/A	N/A
2019/0019516	12/2018	Lawrence et al.	N/A	N/A
2019/0034405	12/2018	Ivan	N/A	G06F 40/279
2019/0108832	12/2018	Tomar et al.	N/A	N/A
2019/0206388	12/2018	Aggarwal et al.	N/A	N/A
2020/0394465	12/2019	Gil Ramos	N/A	G06F 18/2155
2022/0157318	12/2021	Sharifi	N/A	G10L 15/22

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
101075434	12/2006	CN	N/A
101454775	12/2008	CN	N/A
101588417	12/2008	CN	N/A
107093425	12/2016	CN	N/A
2016191319	12/2015	WO	N/A

OTHER PUBLICATIONS

China National Intellectual Property Administration; Notification of Grant issued in Application No. 201980011141.4; 6pages; dated Nov. 26, 2023. cited by applicant

European Patent Office, Intention to Grant issued in Application No. 19701422.8, 44 pages, dated

May 28, 2024. cited by applicant

European Patent Office; Communication pursuant Article 94(3) EPC issued in Application No. 19701422.8; 4 pages; dated Jul. 6, 2023. cited by applicant

China National Intellectual Property Administration; Notification of Second Office Action issued in Application No. 201980011141.4; 7 pages; dated Jul. 20, 2023. cited by applicant

China National Intellectual Property Administration; Notification of First Office Action issued in Application No. 201980011141.4; 30 pages; dated Dec. 5, 2022. cited by applicant

European Patent Office; Communication issued in Application No. 19701422.8, 4 pages, dated Oct. 20, 2022. cited by applicant

European Patent Office; International Search Report issued in Application No. PCT/US2019/012399; 14 pages; dated mailed Apr. 11, 2019. cited by applicant

Primary Examiner: Adesanya; Olujimi A

Attorney, Agent or Firm: Gray Ice Higdon

Background/Summary

BACKGROUND

(1) People are relying on their user devices to perform a variety of different tasks. User devices may include, for example, a smartphone, a smartwatch, a tablet computer, a smart speaker, a headset, a car, a television, a thermostat, or the like. Such devices may include hardware and software that are configured to enable digital assistant functionality. Such applications are generally referred to as “digital assistants.”

(2) Digital assistant functionality may include using a speech recognizer to detect an utterance that is uttered by a user and using the grammar of the speech recognizer to identify an operation that should be performed based on the detected utterance. Such assistant functionality can be generally broken down into online and offline functionality. Online functionality may include commands such as “what is the weather like in Paris today?,” “Order a large pizza for me,” or the like. Such commands require an active network connection and information from one or more servers access via the network.

(3) In many systems, offline functionality is limited to a subset of commands that do not require network connectivity. For example, offline digital assistant functionality may be limited to a subset of commands such as “play 101.5 FM,” “increase volume,” “turn on air conditioner,” or the like.

SUMMARY

(4) The present disclosure is directed towards a system and method for updating the grammar of a speech recognizer that is stored on a user device and used to execute offline voice commands. In general, one innovative aspect of the subject matter described in this specification can be embodied in methods that include the actions of detecting, by a user device, an utterance at a first time when the user device is connected to a server by a network, wherein the utterance includes one or more terms, providing, by the user device, the utterance to the server using the network, receiving, by the user device and from the server, an update to the grammar of the user device based on the one or more terms, detecting, by the user device, a subsequent utterance at a second time when the user device is not connected to the server by a network, wherein the subsequent utterance includes the one or more terms, and in response to detecting, by the user device, the subsequent utterance at the second time: determining, by the user device, that the subsequent utterance is associated with an operation that can be performed without requesting information from the server, identifying, by the user device, an operation to perform based on (i) the one or more terms of the subsequent utterance,

and (ii) the updated grammar, and initiating, by the user device, performance of the operation without requesting information from the server. Other embodiments of this aspect include corresponding systems, apparatus, and computer programs, configured to perform the actions of the methods, encoded on computer storage devices.

(5) These and other versions may optionally include one or more of the following features. For example, in some implementations, initiating by the user device, performance of the operation without requesting information from the server may include generating, by the user device, an action comprising a data object that, when executed, causes the user device to initiate performance of an operation responsive to the subsequent utterance, and executing, by the user device, the action to trigger performance of the operation without requesting information from the server.

(6) In some implementations, the method may also include actions of receiving, by the user device, an action. The action may include a data object that, when executed, causes the user device to initiate performance of an operation responsive to the subsequent utterance, wherein the action was generated by the server, and storing, by the user device, the action on the user device.

(7) In some implementations, initiating by the user device, performance of the operation without requesting information from the server may include accessing, by the user device, the action that is stored on the user device, and executing, by the user device, the action to trigger performance of the operation without requesting information from the server.

(8) In some implementations, receiving, by the user device and from the server, an update to the grammar of the user device based on the one or more terms may include receiving, by the user device, the one or more terms, and updating the grammar of the user device to include the one or more terms.

(9) In some implementations, the one or more terms are associated with a threshold amount of time that indicates a maximum period of time the one or more terms will be included in the grammar absent the user device detecting the subsequent utterance of the one or more terms by a user.

(10) In some implementations, the method may also include actions of determining, by the user device, that the threshold period of time has been exceeded without the user device detecting the subsequent utterance of the one or more terms by the user, and removing, by the user device, the one or more terms from the grammar of the user device.

(11) Another innovative aspect of the subject matter described in this specification can be embodied in methods that include the actions of processing, by a server, a transcription of an utterance that was detected by a first user device when the server and the first user device are connected by a network, determining, by the server based on the processing of the transcription of the utterance, an operation that can be performed by the first user device without requesting information from the server, wherein the transcription includes one or more terms, and providing, by the server, an update to the grammar of the first user device based on the one or more terms. Other embodiments of this aspect include corresponding systems, apparatus, and computer programs, configured to perform the actions of the methods, encoded on computer storage devices.

(12) These and other versions may optionally include one or more of the following features. For example, in some implementations, providing, by the server, an update to the grammar of the first user device based on the one or more terms may include providing the one or more terms to the first user device.

(13) In some implementations, the method may also include actions of associating the one or more terms with a predetermined time period, wherein the predetermined time period indicates a maximum period of time the one or more terms will be included in the grammar absent the subsequent utterance of the one or more terms by a user to perform the operation.

(14) In some implementations, determining, by the server based on a transcription of the utterance, an operation that can be performed by the first user device without requesting information from the server may include generating, by the server, a data object that, when executed, causes the first user device to initiate performance of an operation responsive to the received utterance.

(15) In some implementations, providing, by the server, the update to the grammar of the first user device, based on the one or more terms may include providing the generated data object to the first user device.

(16) In some implementations, the update to the grammar enables the first user device to execute a data object that initiates performance of the operation in response to a subsequent receipt of the utterance without requesting information from the server.

(17) In some implementations, the method may also include actions of identifying, by the server, a user profile that is associated with the first user device, identifying, by the server based on the user profile, a second user device that is registered to a user of the first user device, and providing, by the server, an update to the grammar of the second user device based on the one or more terms.

(18) Particular embodiments of the subject matter described in this specification can be implemented so as to realize one or more of the following advantages. Updating the grammar of a speech recognizer using the system and method of the present disclosure enables the speech recognizer to learn personalized commands that can be used to execute offline voice commands. This provides the advantage of enabling a user to command a digital assistant using the user's own personalized language. If a user can utter commands that can be detected by a speech recognizer in the user's own personalized language and get positive results from applications that receive input commands detected by the speech recognizer, this will increase the likelihood that a user will use applications that can encourage voice input such as digital assistant applications relative to the use of such applications when a user is required memorize a set of standard commands made available to all users of the digital assistant.

(19) The present disclosure provides another advantage in that the learning of personalized commands is achieved as a result of automated processes employed the systems and methods of the present disclosure. That is, a user does not need to initiate the personalization processes. Instead, the personalization of the speech recognizer's grammar can be achieved based on the normal operation of the system based in part on user commands and system responses.

(20) In addition to encouraging a user to use applications that receive input commands detected by the speech recognizer, the present disclosure also improves the functionality of a computer. The updating of the grammar of the speech recognizer that is stored on the user device allows for an increase in the number commands that can be detected by the speech recognizer and executing using one or more control systems employed by the user device. As such, this requires less consultation with a server accessible via network to disambiguate a previously unknown voice command, prompt a user for feedback that is used to help disambiguate the unknown voice command, and then receive data object from the server that can be processed by one or more control systems on the user device to execute a particular action (e.g., turn on air conditioner, turn off air conditioner, tune radio to a particular radio station, increase volume, turn off lights, or the like). Less consultation with a networked server results in the usage of less network bandwidth and a reduction in latency that is involved to execute a command from the time the command is initially uttered.

(21) Yet other advantages may be achieved by the present disclosure. For example, a user device such as a phone, entertainment and climate system of an automobile, or the like, will have to contact a server less to disambiguate uttered commands once the local grammar has been updated. As a result, the user device has to communicate with one or more networks less, resulting in battery savings. By way of another example, a user device utilizing the present disclosure has to use less CPU resources and less storage space storing additional offline Language Packs. In some implementations, this may reduce the amount of downloads and storage of language packs to as little as 40 Mb downloads of Language pack. By way of another example, the server itself has to do less processing to create these large Language packs, which can have a cumulative effect as more languages are being added to grammar of a speech recognizer.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) FIG. 1 is a contextual diagram of an example of a system for executing an offline voice command.
- (2) FIG. 2A is a contextual diagram of an example of a system for learning offline voice commands
- (3) FIG. 2B is a contextual diagram of the system of FIG. 2A that has learned an offline voice command.
- (4) FIG. 3 is a flowchart of an example of a process for generating a personalized grammar update for a user device.
- (5) FIG. 4 is a flowchart of an example of a process for personalizing a grammar of a user device.

DETAILED DESCRIPTION

- (6) FIG. 1 is a contextual diagram of an example of a system **100** for executing an offline voice command **112**. The system **100** includes a user device **105**, a speech recognizer **110**, an instruction generation unit **129**, and one or more control systems **131**.
- (7) The user device **105** can include any device that includes one or more processors to and one or more memory units storing instructions that, when executed by the one or more processors, realize the functionality of a speech recognizer **110**, an instruction generation unit **129**, and one or more control system **131**, as described herein. In the example of FIG. 1, the user device **105** is a computing system of a vehicle. However, the present disclosure need not be so limited. For example, other types of user devices **105** may include a smartphone, a smartwatch, a tablet computer, a smart speaker, a headset, a television, a thermostat, or the like.
- (8) The speech recognition unit **110** includes an utterance detection unit **111** and a grammar **128**. The utterance detection unit **111** can analyze audio data detected by one or more microphones for the occurrence of an utterance **112** by a user, (in this case, the user's name is Bill). The utterance **112** may include a command such as "Turn on AC." Other types of commands may include, for example, "Tune to 99.7," "Increase the volume," "Roll down the window," or the like. In some implementations, the utterance detection unit **111** may be configured to perform additional functionality. For example, the utterance detection unit **111** may be configured to generate a transcription of audio data that the utterance detection unit **111** captures using the one or more microphones. The transcription may include a text version of the audio data that was captured by the one or more microphones in response to Bill's utterance.
- (9) The speech recognition unit **110** can use the grammar **128** to identify a particular operation that should be performed in response to the uttered command. The grammar **128** may include a set of one or more rules that can be applied to one or more words of the transcription of an utterance to determine the meaning of the utterance. In some implementations, the grammar **128** may use a set of one or more rules to associate one or more particular words with a particular operation. For example, the speech recognition unit **110** can identify a particular operation **128a** based on a comparison of one or more words of the transcription to one or more terminals and non-terminal of a grammar rule.
- (10) The speech recognizer **110** can provide data representing the particular operation **128a** to an instruction generation unit **129**. The instruction generation unit **129** can generate a data object **140** (e.g., instructions or commands), based on the received data representing the particular operation **128a**, that, when executed by control system **141**, results in the causing the user device to perform the particular operations.
- (11) The control system **141** may include one or more processing units, one or more memory units, and software code that is configured to receive a data object **140**, process the data objection **140**, and instruct one or more components of a user device **105** to perform the particular operation. In some implementations, the control system **141** may include one or more primary processing units,

the one or more primary memory units, and software code of the user device **105**. Alternatively, the one or more control units may be a separate subsystem of a user device such as climate control circuitry of a vehicle, a home, or the like. In such implementations, the control system **141** may include one or more processing units, one or more memory units, software code, or a combination thereof that is (or are) not the primary processing unit or the processing memory units. Alternatively, the control system **141** may be only software that is used to start, stop, or both, particular operations that can be performed, or are being performed, respectively, by the user device **105**.

(12) By way of example, the grammar **128** may associate the set of words “turn on AC” with an operation of turning on the air conditioner in a vehicle. Accordingly, in response to the speech recognizer **110** generating a transcription of an uttered voice command that includes the set of words “turn on AC,” the speech recognizer can apply the grammar **128** rules to determine that the vehicle should perform the operation of turning on the air conditioner in the vehicle. The speech recognizer can provide data representing the operation of “turn on the air conditioner” to the instruction generation unit **129**. The instruction generation unit **129** can generate a data object that, when executed by a control system **141**, results in the air conditioner being turned on and used to cool the vehicle. Digital assistant technology associated with the speech recognizer **110** may output a message to Bill through a speaker of the vehicle which recites “Thanks Bill. Turning on AC Now.”

(13) The system **100** of FIG. **1** can perform well offline for a limited set of offline commands that are programmed into the grammar for all users of the device so long as a user such as Bill knows the particular command to initiate performance of the particular operation that Bill wants performed by the user device **105**. However, the system **100** of FIG. **1** fails when a user such as Bill wants to use his own personalized language to perform an operation.

(14) For example, when Bill wants to turn on the air conditioner in his home, Bill may use personalized language to ask another family member in his home to turn on the air conditioner by telling the family member to “freeze me!” As a result, when Bill climbs into his vehicle on a hot day Bill may reflexively say “freeze me!” in an effort to get the air conditioner to turn on and start cooling down the inside of the vehicle. If Bill uttered such an offline command in a vehicle equipped with the system of FIG. **1**, then the vehicle's climate and entertainment system would not respond. Alternatively, vehicle's climate and entertainment system (or related digital assistant technology) may respond with a message to the effect of “Sorry, I do not understand that command. Please submit another command,” or the like. This is because the grammar of the speech recognizer used by the user device **105** does not include a set of rules for resolving the uttered command of “Freeze Me.” The system of FIG. **2A** provides a system that enables the speech recognizer's **110** grammar **128** to learn Bill's personalized command of “Freeze Me” so that Bill may utter an offline command “Freeze Me” and have the vehicle turn on the air conditioner in response.

(15) FIG. **2A** is a contextual diagram of an example of a system **200A** for learning offline voice commands. The system **200A** includes a user device **105**, a network **150**, and a server **160**. The user device **105** may include an utterance detection unit **111**, one or more logic determination units **113**, **114**, a network interface **115**, a user feedback input unit **119**, a local grammar **128**, an instruction generation unit **129**, and a deletion unit **131**.

(16) The server **160** may include an utterance receiving unit **116**, a command analysis unit **117**, a feedback data receiving unit **120**, one or more logic determination units **123**, **125**, a grammar update unit **124**, and an update delay unit **127**. In this specification, the term “unit” will be used broadly to refer to a software system or subsystem that, when executed by one or more computers, cause the performance of one or more specific functions. Generally, a unit will be implemented as one or more software modules or components installed on the same computer or computers; in other cases, one or more software modules or components may be installed on one or more computers in one or more locations.

(17) With reference to the contextual example of FIG. 2A, a user such as Bill can decide that the temperature is too hot and can utter a command to turn on the air conditioner. For example, a user Bill, speaking in his own vernacular, may utter the command “Freeze Me” **112** in an attempt to turn on the air conditioner. The command “Freeze Me” **112**, at the time of Bill's utterance of the command “Freeze Me” **112**, was not previously associated with one or more rules of the local grammar **128** of the user device **105**.

(18) One or more microphones of the user device **105** may detect the uttered command “Freeze Me” **112** and generate audio data **112a** based on the detected utterance. The generated audio data **112a** may include an electrical analog signal that is proportional to the acoustic sound wave created by Bill's voice when Bill utters the command “Freeze Me” **112**. The utterance detection unit **111** can analyze the audio data **112a** generated by the one or more microphones of the user device **105** and detect the occurrence of the command “Freeze Me” **112**. In some implementations, the utterance detection unit **111** can generate a transcription **112b** of the audio data **112a** and perform one or more of the operations below based on the transcription **112b** of the audio data **112a**. Alternatively, in other implementations, the user device **105** may perform one or more of the operations described below on the audio data **112a**, and then provide the audio data **112a** to the server **160** for transcription and further processing by one or more units of the server **160** described herein.

(19) The user device **105** may use a first logic unit **113** to determine whether the server **160** (or another server such as a content server) is required to disambiguate and perform one or more operations in response to the command “Freeze Me” **112**. For example, the first logic unit **113** may be configured to receive the audio data **112a**, the transcription **112b**, or both, and analyze the received audio data **112a**, the received transcription **112b**, or both, to determine whether the server **160** is required to disambiguate and perform one or more operations in response to the command “Freeze Me” **112**.

(20) The user device **105** may determine that the server **160** (or another server such as a content server) is required to execute the command “Freeze Me” **112** for a variety of different reasons. For example, the user device **105** may determine that the server **160** (or another server such as a content server) is required to execute a command that inherently requires data from a remote server such as server **160** (or another server such as a content server) to execute the command. Commands that inherently require data from a remote server such as server **160** (or another server such as a content server) in order to execute include commands such as “What is the weather like in Paris today,” “Order me a large pizza from Tony's Pizza,” “Tell me the score of the Steeler game,” or the like. Commands such as these inherently require data from a remote server such as server **160** (or another server such as a content server) because each of the respective commands require communication with one or more remote servers to fully execute the received command in a manner that a user such as Bill is likely to find satisfactory (e.g., by retrieving a real-time weather forecast for Paris from a remote server, by placing an order for Pizza with a server of a pizza shop, by retrieving the current score of the Steeler game from a server tracking scores of one or more sporting events).

(21) Another reason the server **160** (or another server such as a content server) may be required to execute the command “Freeze Me” **112** is because the user device **105** may determine that the local grammar **128** of the user device **105** does not include one or more rules that define the operation that should be performed based on the command “Freeze Me.” In such instances, the user device **105** may provide the data representing the uttered command such as a transcription **112b** of the uttered command **112**, audio data **112a** representing the uttered command **112** that can be transcribed by the server **160**, or both, to the server for disambiguation. Alternatively, if the user device **105** determines that the local grammar **128** does define one or more rules for disambiguating the command “Freeze Me” **112**, then the user device **105** may use the local grammar **128** to disambiguate the command “Freeze Me” **112**.

(22) If the uttered command is determined to not need a server **160** (or another server such as a content server) and the local grammar **128** defines one or more rules for disambiguating the command, then the user device can use the local grammar **128** to disambiguate and execute the received command **112**. For example, the user device **105** may perform a search of entries in the grammar **128** based on one or more portions of the transcription **112b** of the command **112** to identify one or more rules that define the operation(s) that should be performed based on the received command “Freeze Me” **112**. In the example of FIG. 2A, the local grammar **128** does not define one or more rules for disambiguating the command “Freeze Me” **112** at the time the command “Freeze Me” **112** is uttered by Bill. As a result, the local grammar **128** needs to be updated using the techniques described herein.

(23) Alternatively, if the user device **105** needs a server **160** (or another server such as a content server) to disambiguate and execute the command **112**, then the user device **105** may use a second logic unit **114** to determine whether the user device **105** has access to a network **150** to communicate with the server **160** (or another server such as a content server). For example, the user device **105** may determine whether the network interface unit **115** of the user device **105** is within range of, and has permission to access, one or more wireless networks of the network **150** such as a WiFi network, a cellular network, or the like.

(24) If the user device **105** determines that it does not have access to the network **150**, the user device can attempt to execute the command “Freeze Me” **112** using the local grammar **128** of the user device **105**. However, if the local grammar **128** does not include any rules for resolving the uttered command of “Freeze Me,” then the user device may provide an output to Bill indicating that the command cannot be executed. Alternatively, if the user device **105** determines that the user device **105** has access to the network **150**, the user device **105** can communicate with the server **160** for assistance related to disambiguating and executing the command “Freeze Me” **112**.

(25) In the example of FIG. 2A, the user device **105** determines that the user device **105** is connected to the network **150** and that the server **160** is required to disambiguate the command “Freeze Me” **112**. In the example of 2A, the server **160** is required to help disambiguate and execute the command “Freeze Me” **112** since the local grammar **128** does not include one or more rules for disambiguating the command “Freeze Me” **112**. The user device **105** may use the network interface unit **115** to transmit the transcription **112b** of the uttered command **112**, the audio data **112a** representing the uttered command **112**, or both, to the server **160** for disambiguation.

(26) With reference to FIG. 2A, the user device **105** is disclosed as including two logic units **113**, **114** for illustrative purposes. However, the present disclosure need not be so limited. For example, a single logic unit may be used to provide the functionality of logic unit **113**, **114** respectively.

(27) The server **160** can receive the transcription **112b** of the uttered command **112**, the audio data **112a** representing the uttered command **112**, or both, using an utterance data receiving unit **116**. The utterance data receiving unit **116** is an interface between the server **160** and the user device **105**. In some implementations, such as when the utterance receiving unit **116** receives audio data **112a** representing the uttered command **112**, the utterance data receiving unit **116** can be configured to generate a transcription of the uttered command **112** based on the audio data **112a**. The utterance data receiving unit **116** may provide the transcription **112b** of the uttered command **112** to the command analysis unit **117**. In other implementations, such as the example of FIG. 2A, the utterance data receiving unit **116** may receive a transcription **112b** of the uttered command **112** and then forward the transcription **112b** to the command analysis unit **117**.

(28) The command analysis unit **117** can be used to disambiguate a command that was not previously defined by the local grammar **130** of the user device. The command analysis unit **117** may disambiguate a command in a number of ways. For example, the command analysis **117** may associate the command “Freeze Me” with turning on the air conditioner based on an analysis of the keywords of the command, based on a user's response to a request for feedback from Bill regarding the command, based on historical aggregation of Bill's feedback related to the command, based on

a comparison of the keywords of the command to a grammar maintained by the server **160** that is more extensive than the grammar **130** stored on the user device **105** and has already learned that the command, or a combination thereof.

(29) With respect to the example of FIG. 2A, the command analysis unit **117** may determine, based on an analysis of the keywords in the command “Freeze Me” **112**, that the command **112** is associated with a particular operation such as turning on an air conditioner. Based on this determination, the command analysis unit **117** may generate and transmit prompt data **118** for user feedback from Bill that identifies a particular operation. For example, the command analysis unit may generate and transmit prompt data, that when processed for output by the user device **105**, asks Bill whether Bill wants to “Turn on the AC?”. The prompt data **118** is received by the network interface unit **115** and output by one or more interface units of the device **105**. The interface units may include a speaker, a graphical user interface, or both.

(30) For example, prompt data **118** may be received by the user device **105** and processed to generate the output **118a** phrase “Turn on AC?”. In some implementations, the output **118a** may be an audio prompt that is output by the one or more speakers of the user device **105**. In other implementations, the output may be text rendered on the display of a user interface of the user device **105**. In yet other implementations, the output **118a** may be provided via a speaker or display of a different user device that is associated with Bill and the user device **105** via a cloud service. The different user device may include, for example, a smartphone, a tablet, a smartwatch, or the like.

(31) Bill may provide feedback **119a** in response to the output **118a** generated based on the prompt data **118** that indicates whether or not the particular operation identified by the command analysis unit **117** should be associated with the command **112** uttered by Bill. For example, Bill can provide feedback indicating that the particular operation (e.g., turn on air conditioner) identified by the command analysis unit **117** should be associated with the command **112** (e.g., “Freeze Me”) uttered by Bill. Alternatively, Bill can provide feedback indicating that the particular operation (e.g., turn on air conditioner) identified by the command analysis unit **118** should not be associated with the command **112** (e.g., “Freeze Me”). In the example of FIG. 2A, Bill's feedback **119a** includes an utterance that is detected by one or more microphones of the user device **105**. Alternatively, or in addition, Bill's feedback may also include a selection of one or more options of a graphical user interface of the user device **105**, an utterance detected by a different user device such as a smartphone, a tablet, a smartwatch, or the like that is associated with Bill and the user device **105** via a cloud service.

(32) In the example of FIG. 2A, Bill may provide feedback to the output **118a** prompt “Turn on the AC?” by uttering the word “Yes” **119a**. One or more microphones of the user device **105** may detect the Bill's utterance and generate audio data based on the detected utterance. The generated audio data may include an electrical analog signal that is proportional to the acoustic sound wave created by Bill's voice when Bill utters the feedback **119** “Yes.” A user feedback input unit can receive the audio data generated based on the uttered feedback **119a**, analyze the audio data generated based on the uttered feedback **119a**, and detect the occurrence of the feedback **119** phrase “Yes.” In some implementations, the user feedback input unit **119** can generate a transcription **119bn** of the audio data generated based on the uttered feedback **119a** and perform one or more of the operations based on the transcription **119b** of the audio data generated based on the uttered feedback **119a**. The user device **119b** may use the network interface unit **115** to transmit the transcription **119b** to the server **160** using one or more networks **150**. Alternatively, in other implementations, the user device **105** provide the audio data generated based on the uttered feedback **119a** to the server **160** via the network **150**. In such instances, the server **160** may generate the transcription of the received audio data generated based on the uttered feedback **119a**.

(33) A feedback data receiving unit **120** may receive the transcription **119b** of the uttered feedback **119a** and provide the transcription **119b** to the command analysis unit **117**. Alternatively, the

feedback data receiving unit **120** may receive audio data generated based on the uttered feedback **119a**, generate a transcription **119b** of the audio data generated based on the uttered feedback **119a**, and then provide the transcription **119b** of the uttered feedback **119a** to the command analysis unit **117**.

(34) In response to receiving the transcription **119b** of the uttered feedback **119a**, the command analysis unit **117** may transmit response data **121** to the user device **105** via the network **150**. The user device **105** can receive the response data **121** using the network interface unit **115** and process response data **121** to generate an output **133** confirming the association of the command “Freeze Me” **112** with the particular operation of “turning on the air conditioner” of the environment associated with the user device **105**. The output **133** may include output provided via a speaker of the user device **105**, output provided via a display of the user device **105**, output provided by a speaker of a different user device (e.g., smartphone, tablet, smartwatch, or the like), output provided by a display of a different user device (e.g., a smartphone, tablet, smartwatch, or the like), or a combination thereof. In some implementations, the output **133** may include output that identifies the particular operation that the command analysis unit **117** associated with the command “Freeze Me.” For example, the output **133** may include output from a speaker of the user device **105** that vocalizes output of “Thanks Bill. Turning on the AC Now.”

(35) The command analysis unit **117** can initiate generation of an update to the local grammar **128** of the user device **105**. In some implementations, initiating generation of an update to the local grammar **128** of the user device **105** may include, for example, instructing a third logic unit **123** to evaluate data **122** identifying the uttered command **112** to determine whether the uttered command requires the server **160**. In some implementations, the determination of whether the uttered command **112** requires the server **160** may be based on the particular operation the command analysis unit **117** associated with the uttered command **112**. If the third logic unit **123** determines that the uttered command **112** (or the particular operation the command analysis unit **117** associated therewith) requires the server **160**, then the server **160** may cease generation of the update to the local grammar **128**. Alternatively, if the third logic unit **123** determines that the uttered command **112** and the particular operation the command analysis unit **117** associated therewith does not require the server **160**, then the result of the third logic unit's **123** processing may provide an indication to the grammar update unit **124** to generate an update for the local grammar **128** of the user device **105**.

(36) Generating the update can include generating a plurality of instructions that are configured to, upon execution of the instructions by the user device **105** during processing of the update, create a set of one or more rules in the local grammar **128** that associate the user command (e.g., “Freeze Me”) with the particular operation (e.g., “turn on the air conditioner”) in the local grammar **128** of the user device **105**. In some implementations, generating the local grammar update may also include a generating a data object that, when executed by the user device **105**, causes the user device **105** to initiate performance of the particular operation (e.g., “turn on air conditioner”) responsive to a subsequent utterance of the command **112** (e.g., “Freeze Me.”). The data object may include an XML data object, a binary data object, or the like. In some implementations, generating the local grammar update may also include generating expiration data that includes a maximum period of time that the set of one or more rules associated with the local grammar update will be stored in the local grammar **128** without being used.

(37) Each updating methodology provides particular advantages. For example, generating an update that includes a plurality of instructions to create a set of one or more rules in the local grammar **128** to associate the user command (e.g., “Freeze Me”) with the particular operation (e.g., “turning on the air conditioner”) without the generated data object results in an update package of a smaller size (relative to an update package that includes the generated data object). Accordingly, this update methodology may be used to conserve bandwidth less data would need to be provided to the user device **105** to facilitate updating the local grammar **128**. In addition, update download

times may be faster by downloading the plurality of instructions without the generated data object. In addition, this updating methodology also leads to better recognition as there are fewer possible matches. For example, an utterance such as “Freeze We,” “Sneeze Me,” or the like may also be corrected to “Freeze Me” if the words of the utterance do not match words in the local grammar **128**. However, this update methodology will require an instruction generation unit **129** to generate the data object.

(38) Alternatively, different benefits can be achieved by generating the data object and providing the data object with the grammar update for storage on the user device **105** may increase performance of the user device **105**. This is because it saves the user device **105** from having to generate the data object in response the user device **105** identifying the particular operation as being associated subsequent utterances of the command **112**. Because the user device **105** can merely access the data object (as opposed to generating the data object), such embodiments can reduce the latency involved in the user device executing the particular operation responsive to subsequent utterances of the command “Freeze Me.”

(39) Yet different benefits can be achieved by generating and providing expiration data with the local grammar update. For example, the system **200A** can use the expiration data as a mechanism for optimizing use of storage space by the user device **105**. That is, the user device **105** can free-up storage space allocated to store a set of rules associating an uttered command with a particular operation if Bill does not utter the command within the maximum period of time set forth in the expiration data.

(40) The grammar update unit **124** produces the local grammar update **126** and uses a fourth logic unit **125** to determine whether or not the server **160** has access to a network **150** that facilitates communication with the user device **105**. If the fourth logic unit **125** determines that the server **160** does not have access to a network **150** that facilitates communication with the user device **105**, then the fourth logic unit **125** instruct a local grammar update delay unit **127** that the server **160** cannot communicate with the user device **105** at this time. The local grammar update delay unit **127** can delay subsequent attempts to provide the local grammar update **126** to the user device **105** for a predetermined amount of time. In response to the expiration of the predetermined amount of time, the update delay unit **127** can use the fourth logic unit **125** to determine whether or not the server **160** has access to a network **150** that facilitates communication with the user device **105**.

(41) Alternatively, if the fourth logic unit **125** determines that the server **160** does have access to a network **150** that facilitates communication with the user device **105**, then the server **160** can provide the local grammar update **126** to the user device **105**. The user device **105** can receive the local grammar update via the network interface unit **115**.

(42) The user device **105** may process the instructions included in the local grammar update **126** and update the local grammar **128** by generating a set of one or more rules for storage in the local grammar **128**. The generation and storage of the set of one or more rules in the local grammar **128** enable the local grammar **128** to disambiguate subsequently received commands that correspond one of the commands that were part of the local grammar update **126**. For example, the one or more rules can be applied to one or more words of the transcription of a subsequent utterance to determine the meaning of the subsequent utterance. Determining a meaning of a subsequent utterance may include, for example, a identifying an operation that should be performed in response to the detection of the subsequent utterance. Accordingly, after receiving and processing the local grammar update **126**, the user device **150** can use the local grammar **128** to disambiguate subsequent utterances of the command “Freeze Me” by using one or more rules in the local grammar **128** to associate the “Freeze Me” command with the particular operation of “turning on the air conditioner.”

(43) The instruction generation unit **129** can also be updated based on the local grammar update **126**. In some implementations, for example, the local grammar update **126** may include an action that is to be initiated, performed, or both, by the user device **105** in response to the an uttered

command. The action may include a data object that, when executed by the user device **105**, causes the user device **105** to initiate performance of (or to perform) the particular operation (e.g., “turning on air conditioner”) that is associated with an uttered command (e.g., “Freeze Me.”) by the local grammar **128**. The data object may include an XML data object, a binary data object, or the like. If the local grammar update **126** includes the action that that is to be initiated, performed, or both, by the user device **105** in response to an uttered command, the action can be stored in the instruction generation unit **129**. Once stored in the instruction generation unit **129**, the action (e.g., data object that instructs the user device **105** to turn on the air conditioner) may be retrieved upon after receipt of the utterance of a corresponding command (e.g., “Freeze Me”) based on data provided by the local grammar **128** that identifies a particular operation associated with the utterance. Alternatively, the action (e.g., data object that instructs the user device **105** to turn on the air conditioner) may be generated by the instruction generation unit **129** after receipt of the utterance of a corresponding comment (e.g., “Freeze Me”) based on data provided by the local grammar **128** that identifies a particular operation associated with the utterance.

(44) Other units may also be updated based on the local grammar update **126**. For example, the deletion unit **131** may be updated to monitoring the usage of one or more sets of rules that are included in the local grammar **128**. In some implementations, a set of grammar rules may be associated with a maximum period of time for which the set of grammar rules may be idle (e.g., not used). The deletion unit **131** may determine whether or not a command that is associated with a set of grammar rules has been uttered within the maximum period of time. In response to determining that the command has not been uttered, or otherwise input, within the maximum period of time, the deletion unit **131** may instruct **132** the local grammar **128** to delete the set of rules from the local grammar **128**. This feature provides the benefit of optimizes the use of local user device storage, which may be finite.

(45) Though deletion unit **131** has been described as removing a set of one or more rules from the local grammar **128** based a determination that a command associated with the set of one or more rules has not been uttered, or otherwise input, within a maximum period of time, the present disclosure need not be so limited. Instead, the deletion unit **128** may detect the occurrence of one or more other triggers whose detection initiates removal (e.g., deletion) of a set of one or more rules associated with a particular command in the local grammar **128**. Such triggers may include changes in the state of the user device **105** such as removal of a user account, selection of a different language for the user device **105**, a deletion request received from a user (e.g., Bill), a request to reset the device received from a user (e.g., Bill), or the like.

(46) In the example of FIG. 2B, the user device that detected Bill's utterance of a previously unknown command (e.g., “Freeze Me”) and the user device that received and processed the local grammar update **126** were the same user device—i.e., user device **105**. However, the present disclosure need not be so limited. Instead, for example, multiple user devices may be associated with Bill via one or more cloud services. In such instances, the server **160** can transmit the local grammar update such as local grammar update **126** to each of the other devices associated with Bill via the one or more cloud services. For example, the local grammar update can also be transmitted to Bills smartphone, smartwatch, laptop computer, tablet computer, desktop computer, smart television, set top box, smart speaker, digital home assistant, or the like. In this way, the local grammar of each of Bill's user devices can be updated to respond to commands in Bill's own vernacular. Once updated, these user devices can be used to respond to commands in Bill's own vernacular even when offline.

(47) FIG. 2B is a contextual diagram of the system of FIG. 2A that has learned an offline voice command. In general, the system **200B** includes a user device **111**, a first logic unit **113**, a second logic unit **114**, a grammar **128**, an instruction generation unit **111**, a deletion unit **131**, a feedback input unit **119**, and a network interface unit **115**.

(48) In more detail, the system **200B** is the same as the system **200A** except the system **200B** is

shown after the local grammar of **128** has received and processed the local grammar update **126**. Processing of the local grammar update **126** enabled the local grammar **128** to learn a command (e.g., “Freeze Me”) that was not associated within one or more rules in the local grammar **128** prior to receiving and processing the local grammar update **126**. In addition, the user device **105** of FIG. 2B does not have access to the network **150**. Because the user device **105** of FIG. 2B does not have access to the network **150**, the user device **105** of FIG. 2B cannot communicate with the server **160** to avail itself of the command disambiguation services provided by the server **160**.

(49) With reference to FIG. 2B, the user device **105** may detect utterance of a command **112** that is uttered in Bill's own vernacular such as the command **112** “Freeze Me” that was unknown to the local grammar **128** prior to the user's device's **105** receipt and processing of the local grammar update **126**.

(50) The user device **105** may use one or more microphones to capture audio data **112a** that is provided to the utterance detection unit **111**. The utterance detection unit **111** may generate a transcription **112b** of the audio data **112a**. The user device **105** may use the first logic unit **113** to determine that the server is not needed to execute the command represented by the transcription **112b** because the command does not require data from the server **160** and the local grammar **128** that was updated using the local grammar update **126** now includes the command **112** represented by the transcription **112b**. The second logic unit **114** may determine that there is no access to the network **150**.

(51) The local grammar **128** may determine, based on the received a transcription **112b** and a set of one or more rules stored in the local grammar that were received as part of the local grammar update **126**, that the command **112** (e.g., “Freeze Me”) is associated with a particular operation (e.g., turning on the air conditioner). The local grammar may transmit data **228a** instruction generation unit **129** that identifies the particular operation (e.g., turning on the air conditioner) that is associated with the command **112**.

(52) The instruction generation unit **129** may obtain an action **130**, based on the particular operation identified in the data **228a** received from the local grammar **128**, that when, executed, initiates performance of a particular operation (e.g., turning on the air conditioner) that should be performed in response to the command **112**. In some implementations, the action **130** may have been previously received and stored by the instructions generation unit **129** as party of the local grammar update **126**. In such implementations, the latency for executing the particular operation in response to the command **112** may be reduced by storing the action **130** on the user device **150**. Alternatively, obtaining the action **130** may include, for example, generating the action **130**.

(53) The action **130** may include a data object that, when executed by the user device **105**, causes the user device **105** to initiate performance of the particular operation (e.g., “turning on air conditioner”) responsive to a subsequent utterance of the command **112** (e.g., “Freeze Me.”). The data object may include an XML data object, a binary data object, or the like. The action **130** may be processed by the user device **105** and trigger the performance of the particular operation (e.g., turning on the air conditioner).

(54) FIG. 3 is a flowchart of an example of a process **300** for generating a personalized grammar update for a user device. For convenience, the process **300** is described below as being performed by a system such as the system **200A** depicted in FIG. 2A.

(55) The system may use a server to receive **310** data representing at least a portion of an utterance from a first user device. The data representing the portion of the utterance may include a transcription of one or more portions of the utterance. Alternatively, the data representing the portion of the utterance may include audio data representing one or more portions of the utterance. The utterance may include a command in the unique vernacular of a user for a known operation. For example, the utterance may include a command to “Hurt My Ears.” In this example, the user may utter the command “Hurt My Ears” as a request to perform the operation of “increasing the volume of audio (or video) playback by the first user device.”

(56) The system may determine **320**, based on the data representing the utterance, an operation that can be performed responsive to the utterance. In some implementations, the operation may be an operation that can be performed without requesting server information. For example, the system may determine that the operation of “increasing the volume of audio (or video) playback by a device” is an operation that can be performed without requesting server information.

(57) The system may generate **330** an update to the grammar based on the (I) data representing at least a portion of the utterance and (II) the operation. Generating an update to the grammar may include generating a plurality of instructions that are configured to, upon execution of the instructions by the first user device during processing of the update, create a set of one or more rules in a grammar of the first user device that associates the user command (e.g., “Hurt My Ears”) with the particular operation (e.g., “increasing the volume of audio (or video) playback by the first”).

(58) In some implementations, generating the update to the grammar may also include generating a data object that, when executed by the first user device, causes the first user device to initiate performance of the particular operation (e.g., “increasing the volume of audio (or video) playback by the first user device”) responsive to a subsequent utterance of the command. The data object may include an XML data object, a binary data object, or the like. In some implementations, generating the grammar update may also include generating expiration data that includes a maximum period of time that the set of one or more rules associated with the grammar update will be stored in the grammar without being used.

(59) The system may determine **340** whether the first device and the first device and the server are connected by one or more networks. In response to determining that the first user device and the server are connected by one or more networks, then the system may provide **350** the grammar update to the first user device.

(60) Alternatively, in response to determining that the first user device and the server are not connected by one or more networks, the system may proceed to stage **342** and delay **342** the grammar update until the first user device and the server are connected by one or more networks. In such instances, the system may continue to stage **350** and provide the grammar update to the first user device after determining that a connection has been established between the first user device and the server.

(61) FIG. **4** is a flowchart of an example of a process **400** for personalizing a grammar of a user device. For convenience, the process **400** is described below as being performed by a system such as the system **200A** depicted in FIG. **2A**.

(62) The system can detect **410**, using a first user device, an utterance of one or more terms when a user is connected to a server by a network. For example, the first user device may use one or more microphones to capture audio data that represents the utterance. By way of example, the utterance may include a command in the user's own vernacular to “Hurt My Ears.”

(63) The system can provide **420**, using the first device, data representing the utterance to the server using the network. For example, the system may provide the captured audio data describing all, or a portion of, the utterance to the server. Alternatively, or in addition, the system may provide a transcription of all, or a portion of, the utterance to the server. By way of example, the first device may provide a transcription of audio data that represents the command “Hurt My Ears” to the server.

(64) The system can receive **430**, by the user device and from the server, an update to the grammar of the user device based on the data representing the utterance. Receiving the update to the grammar of the user device may include obtaining a copy of update to the grammar via a network. Alternatively, or in addition, receiving the update to the grammar may also include installing the update to the grammar. The update to the grammar of the user device may include a plurality of instructions that are configured to, upon execution of the instructions by the first user device during processing of the update, create a set of one or more rules in a grammar of the first user device that

associates the user command (e.g., “Hurt My Ears”) with the particular operation (e.g., “increasing the volume of audio (or video) playback by the first”).

(65) In some implementations, the update to the grammar may also include a data object that, when executed by the first user device, causes the first user device to initiate performance of the particular operation (e.g., “increasing the volume of audio (or video) playback by the first user device”) responsive to a subsequent utterance of the command. The data object may include an XML data object, a binary data object, or the like. In some implementations, the grammar update may also include expiration data that includes a maximum period of time that the set of one or more rules associated with the grammar update will be stored in the grammar without being used.

(66) The system can detect **440** a subsequent utterance that includes the one or more terms when the first user device is not connected to the server by a network. For example, the first user device may use one or more microphones to capture audio data that represents the subsequent utterance. The subsequent utterance may include an utterance that was received after the receipt and processing of the update to the grammar in stage **430**. For example, the subsequent utterance may include the command “Hurt My Ears” that was uttered by the user at some point in time after the update to the grammar was received in stage **430**.

(67) The system can determine **450**, using the first user device, that the subsequent utterance is associated with an operation that can be performed without requesting information from the server. For example, the first user device can evaluate data representing the utterance to determine whether information from a server is required to disambiguate the utterance and perform the operation that is responsive to the utterance.

(68) The system can obtain **460** a data object that can be used to trigger performance of the operation based on (I) the one or more terms of the subsequent utterance and (II) the updated grammar. The data object may include data that, when executed by the first user device, causes the first user device to initiate performance of an operation (e.g., “increasing the volume of audio (or video) playback by the first user device”) that was associated with the utterance in the updated grammar. The data object may include an XML data object, a binary data object, or the like. The system can initiate **470**, using the user device, performance of the operation without requesting information from the server by processing the data object.

(69) Embodiments of the subject matter, the functional operations and the processes described in this specification can be implemented in digital electronic circuitry, in tangibly-embodied computer software or firmware, in computer hardware, including the structures disclosed in this specification and their structural equivalents, or in combinations of one or more of them. Embodiments of the subject matter described in this specification can be implemented as one or more computer programs, i.e., one or more modules of computer program instructions encoded on a tangible nonvolatile program carrier for execution by, or to control the operation of, data processing apparatus. Alternatively, or in addition, the program instructions can be encoded on an artificially generated propagated signal, e.g., a machine-generated electrical, optical, or electromagnetic signal that is generated to encode information for transmission to suitable receiver apparatus for execution by a data processing apparatus. The computer storage medium can be a machine-readable storage device, a machine-readable storage substrate, a random or serial access memory device, or a combination of one or more of them.

(70) The term “data processing apparatus” encompasses all kinds of apparatus, devices, and machines for processing data, including by way of example a programmable processor, a computer, or multiple processors or computers. The apparatus can include special purpose logic circuitry, e.g., an FPGA (field programmable gate array) or an ASIC (application specific integrated circuit). The apparatus can also include, in addition to hardware, code that creates an execution environment for the computer program in question, e.g., code that constitutes processor firmware, a protocol stack, a database management system, an operating system, or a combination of one or more of them.

(71) A computer program (which may also be referred to or described as a program, software, a

software application, a module, a software module, a script, or code) can be written in any form of programming language, including compiled or interpreted languages, or declarative or procedural languages, and it can be deployed in any form, including as a standalone program or as a module, component, subroutine, or other unit suitable for use in a computing environment. A computer program may, but need not, correspond to a file in a file system. A program can be stored in a portion of a file that holds other programs or data (e.g., one or more scripts stored in a markup language document), in a single file dedicated to the program in question, or in multiple coordinated files (e.g., files that store one or more modules, sub programs, or portions of code). A computer program can be deployed to be executed on one computer or on multiple computers that are located at one site or distributed across multiple sites and interconnected by a communication network.

(72) The processes and logic flows described in this specification can be performed by one or more programmable computers executing one or more computer programs to perform functions by operating on input data and generating output. The processes and logic flows can also be performed by, and apparatus can also be implemented as, special purpose logic circuitry, e.g., an FPGA (field programmable gate array) or an ASIC (application specific integrated circuit).

(73) Computers suitable for the execution of a computer program include, by way of example, can be based on general or special purpose microprocessors or both, or any other kind of central processing unit. Generally, a central processing unit will receive instructions and data from a read-only memory or a random access memory or both. The essential elements of a computer are a central processing unit for performing or executing instructions and one or more memory devices for storing instructions and data. Generally, a computer will also include, or be operatively coupled to receive data from or transfer data to, or both, one or more mass storage devices for storing data, e.g., magnetic, magneto optical disks, or optical disks. However, a computer need not have such devices. Moreover, a computer can be embedded in another device, e.g., a mobile telephone, a personal digital assistant (PDA), a mobile audio or video player, a game console, a Global Positioning System (GPS) receiver, or a portable storage device (e.g., a universal serial bus (USB) flash drive), to name just a few.

(74) Computer readable media suitable for storing computer program instructions and data include all forms of nonvolatile memory, media and memory devices, including by way of example semiconductor memory devices, e.g., EPROM, EEPROM, and flash memory devices; magnetic disks, e.g., internal hard disks or removable disks; magneto optical disks; and CD-ROM and DVD-ROM disks. The processor and the memory can be supplemented by, or incorporated in, special purpose logic circuitry.

(75) To provide for interaction with a user, embodiments of the subject matter described in this specification can be implemented on a computer having a display device, e.g., a CRT (cathode ray tube), LCD (liquid crystal display), LED (light-emitting diode), or OLED (organic light-emitting diode) monitor, for displaying information to the user and a keyboard and a pointing device, e.g., a mouse, a trackball, or a touchscreen, by which the user can provide input to the computer. Other kinds of devices can be used to provide for interaction with a user as well; for example, feedback provided to the user can be any form of sensory feedback, e.g., visual feedback, auditory feedback, or tactile feedback; and input from the user can be received in any form, including acoustic, speech, or tactile input. In addition, a computer can interact with a user by sending documents to and receiving documents from a device that is used by the user; for example, by sending web pages to a web browser on a user's user device in response to requests received from the web browser.

(76) Embodiments of the subject matter described in this specification can be implemented in a computing system that includes a back end component, e.g., as a data server, or that includes a middleware component, e.g., an application server, or that includes a front end component, e.g., a client computer having a graphical user interface or a Web browser through which a user can interact with an implementation of the subject matter described in this specification, or any

combination of one or more such back end, middleware, or front end components. The components of the system can be interconnected by any form or medium of digital data communication, e.g., a communication network. Examples of communication networks include a local area network (“LAN”) and a wide area network (“WAN”), e.g., the Internet.

(77) The computing system can include clients and servers. A client and server are generally remote from each other and typically interact through a communication network. The relationship of client and server arises by virtue of computer programs running on the respective computers and having a client-server relationship to each other.

(78) While this specification contains many specific implementation details, these should not be construed as limitations on the scope of what may be claimed, but rather as descriptions of features that may be specific to particular embodiments. Certain features that are described in this specification in the context of separate embodiments can also be implemented in combination in a single embodiment. Conversely, various features that are described in the context of a single embodiment can also be implemented in multiple embodiments separately or in any suitable subcombination. Moreover, although features may be described above as acting in certain combinations and even initially claimed as such, one or more features from a claimed combination can in some cases be excised from the combination, and the claimed combination may be directed to a subcombination or variation of a subcombination.

(79) Similarly, while operations are depicted in the drawings in a particular order, this should not be understood as requiring that such operations be performed in the particular order shown or in sequential order, or that all illustrated operations be performed, to achieve desirable results. In certain circumstances, multitasking and parallel processing may be advantageous. Moreover, the separation of various system components in the embodiments described above should not be understood as requiring such separation in all embodiments, and it should be understood that the described program components and systems can generally be integrated together in a single software product or packaged into multiple software products.

(80) Particular embodiments of the subject matter have been described. Other embodiments are within the scope of the following claims. For example, the actions recited in the claims can be performed in a different order and still achieve desirable results. As one example, the processes depicted in the accompanying figures do not necessarily require the particular order shown, or sequential order, to achieve desirable results. In certain implementations, multitasking and parallel processing may be advantageous. Other steps or stages may be provided, or steps or stages may be eliminated, from the described processes. Accordingly, other implementations are within the scope of the following claims.

Claims

1. A method comprising: receiving a transcription of an utterance that was detected via a first user interface input, wherein the transcription of the utterance includes one or more terms corresponding to the utterance; identifying, based on the terms of the transcription of the utterance, an operation to be performed responsive to the utterance, wherein the operation to be performed responsive to the utterance comprises changing a state of a user device comprising the first user interface; determining, whether the operation can be performed at the user device responsive to the utterance and without requesting information from a server, wherein determining that the operation can be performed by the user device responsive to the utterance and without requesting information from the server comprises: determining that the user device does not require data from the server to change the state responsive to the utterance, and determining that the user device includes one or more rules associated with changing the state responsive to the utterance; and based on determining, that the operation can be performed by the user device responsive to the utterance and without requesting information from the server: causing, based on the one or more terms of the

transcription corresponding to the utterance and one or more of the rules associated with the operation to be performed responsive to the utterance, an update to a grammar stored locally on the user device to be provided, wherein, when it is determined, that the operation cannot be performed by the user device responsive to the utterance and without requesting information from the server, no update is provided to the grammar stored locally on the user device; subsequent to, and in response to, causing the update to the grammar stored locally on the user device to be provided: identifying a user profile that is associated with the user device, wherein identifying the user profile that is associated with the user device comprises identifying the user profile based on the user profile being associated with the user device and the user device being the user device to which the update to the grammar is provided; identifying, based on the user profile, one or more other user devices that are registered to a user of the user device, wherein identifying the one or more other user devices based on the user profile is in response to the user profile being identified as associated with the user device; determining, whether the operation can be performed by one or more of the other user devices responsive to the utterance and without requesting information from the server; and causing, based on the identification of an other user device of the other user devices that are registered to the user of the user device, based on determining the operation can be performed by the other user device responsive to the utterance and without requesting information from the server, and based on the one or more terms, an update to an additional grammar stored locally on the other user device, wherein, when it is determined, that the operation cannot be performed by the other user device responsive to the utterance and without requesting information from the server, no update is provided to the additional grammar stored locally on the other user device.

2. The method of claim 1, wherein causing the update to the grammar stored locally on the user device to be provided includes causing provision of the one or more terms to the user device.
3. The method of claim 2, further comprising: associating the one or more terms with a predetermined time period, wherein the predetermined time period indicates a maximum period of time the one or more terms will be included in the grammar absent receiving a subsequent utterance of the one or more terms by the user to perform the operation.
4. The method of claim 1, further comprising: causing generation, by the server, of a data object that, when executed, causes the user device to initiate performance of the operation responsive to the utterance.
5. The method of claim 4, wherein causing an update to the grammar stored locally on the user device to be provided further comprises causing provision of the generated data object to the user device.
6. The method of claim 1, wherein the update to the grammar enables the user device to execute a data object that initiates performance of the operation in response to subsequently receiving the utterance, and without requesting information from the server.
7. The method of claim 1, wherein the other user devices comprise one of: a smartphone, a smartwatch, a laptop computer, a tablet computer, a desktop computer, a smart television, a set top box, a smart speaker, and/or a digital home assistant.
8. A system comprising: one or more computers and one or more storage devices storing instructions that are operable, when executed by the one or more computers, to cause the one or more computers to perform operations comprising: receiving a transcription of an utterance that was detected via a first user interface input, wherein the transcription of the utterance includes one or more terms corresponding to the utterance; identifying, based on the terms of the transcription of the utterance, an operation to be performed responsive to the utterance, wherein the operation to be performed responsive to the utterance comprises changing a state of a user device comprising the first user interface; determining, whether the operation can be performed at the user device responsive to the utterance and without requesting information from a server, wherein determining that the operation can be performed by the user device responsive to the utterance and without

requesting information from the server comprises: determining that the user device does not require data from the server to change the state responsive to the utterance, and determining that the user device includes one or more rules associated with changing the state responsive to the utterance; and based on determining, that the operation can be performed by the user device responsive to the utterance and without requesting information from the server: causing, based on the one or more terms of the transcription corresponding to the utterance and one or more of the rules associated with the operation to be performed responsive to the utterance, an update to a grammar stored locally on the user device to be provided, wherein, when it is determined, that the operation cannot be performed by the user device responsive to the utterance and without requesting information from the server, no update is provided to the grammar stored locally on the user device; subsequent to, and in response to, causing the update to the grammar stored locally on the user device to be provided: identifying a user profile that is associated with the user device, wherein identifying the user profile that is associated with the user device comprises identifying the user profile based on the user profile being associated with the user device and the user device being the user device to which the update to the grammar is provided; identifying, based on the user profile, one or more other user devices that are registered to a user of the user device, wherein identifying the one or more other user devices based on the user profile is in response to the user profile being identified as associated with the user device; determining, whether the operation can be performed by one or more of the other user devices responsive to the utterance and without requesting information from the server; and causing, based on the identification of an other user device of the other user devices that are registered to the user of the user device, based on determining the operation can be performed by the other user device responsive to the utterance and without requesting information from the server, and based on the one or more terms, an update to an additional grammar stored locally on the other user device, wherein, when it is determined, that the operation cannot be performed by the other user device responsive to the utterance and without requesting information from the server, no update is provided to the additional grammar stored locally on the other user device.

9. The system of claim 8, wherein causing the update to the grammar stored locally on the user device to be provided includes causing provision of the one or more terms to the user device.

10. The system of claim 9, the operations further comprising: associating the one or more terms with a predetermined time period, wherein the predetermined time period indicates a maximum period of time the one or more terms will be included in the grammar absent receiving a subsequent utterance of the one or more terms by the user to perform the operation.

11. The system of claim 8, the operations further comprising: causing generation, by the server, of a data object that, when executed, causes the user device to initiate performance of the operation responsive to the utterance.

12. The system of claim 11, wherein causing an update to the grammar stored locally on the user device to be provided further comprises causing provision of the generated data object to the user device.

13. The system of claim 8, wherein the update to the grammar enables the user device to execute a data object that initiates performance of the operation in response to subsequently receiving the utterance, and without requesting information from the server.

14. The system of claim 8, wherein the operations are locally performed at the user device and/or the other user devices.

15. The system of claim 8, wherein the operations are remotely performed at the server.

16. The system of claim 8, wherein a first set of the operations are locally performed at the user device and/or the other user devices, and a second set of the operations are remotely performed at a server.

17. A non-transitory computer-readable medium storing software comprising instructions executable by one or more computers which, upon such execution, cause the one or more

computers to perform operations comprising: receiving a transcription of an utterance that was detected via a first user interface input, wherein the transcription of the utterance includes one or more terms corresponding to the utterance; identifying, based on the terms of the transcription of the utterance, an operation to be performed responsive to the utterance, wherein the operation to be performed responsive to the utterance comprises changing a state of a user device comprising the first user interface; determining, whether the operation can be performed at the user device responsive to the utterance and without requesting information from a server, wherein determining that the operation can be performed by the user device responsive to the utterance and without requesting information from the server comprises: determining that the user device does not require data from the server to change the state responsive to the utterance, and determining that the user device includes one or more rules associated with changing the state responsive to the utterance; and based on determining, that the operation can be performed by the user device responsive to the utterance and without requesting information from the server: causing, based on the one or more terms of the transcription corresponding to the utterance and one or more of the rules associated with the operation to be performed responsive to the utterance, an update to a grammar stored locally on the user device to be provided, wherein, when it is determined, that the operation cannot be performed by the user device responsive to the utterance and without requesting information from the server, no update is provided to the grammar stored locally on the user device; subsequent to, and in response to, causing the update to the grammar stored locally on the user device to be provided: identifying a user profile that is associated with the user device, wherein identifying the user profile that is associated with the user device comprises identifying the user profile based on the user profile being associated with the user device and the user device being the user device to which the update to the grammar is provided; identifying, based on the user profile, one or more other user devices that are registered to a user of the user device, wherein identifying the one or more other user devices based on the user profile is in response to the user profile being identified as associated with the user device; determining, whether the operation can be performed by one or more of the other user devices responsive to the utterance and without requesting information from the server; and causing, based on the identification of an other user device of the other user devices that are registered to the user of the user device, based on determining the operation can be performed by the other user device responsive to the utterance and without requesting information from the server, and based on the one or more terms, an update to an additional grammar stored locally on the other user device, wherein, when it is determined, that the operation cannot be performed by the other user device responsive to the utterance and without requesting information from the server, no update is provided to the additional grammar stored locally on the other user device.

18. The non-transitory computer-readable medium of claim 17, wherein causing the update to the grammar stored locally on the user device to be provided includes causing provision of the one or more terms to the user device.

19. The non-transitory computer-readable medium of claim 18, the operations further comprising: associating the one or more terms with a predetermined time period, wherein the predetermined time period indicates a maximum period of time the one or more terms will be included in the grammar absent receiving a subsequent utterance of the one or more terms by the user to perform the operation.

20. The non-transitory computer-readable medium of claim 17, the operations further comprising: causing generation, by the server, of a data object that, when executed, causes the user device to initiate performance of the operation responsive to the utterance.
