

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent

12388711

Kind Code

B2

Date of Patent

August 12, 2025

Inventor(s)

Bosanac; Jonathan Michael et al.

Policy management across multiple cloud computing environments within a network

Abstract

A system for providing policy-controlled communication between a plurality of different cloud computing environments includes a user interface that receives configuration settings to be applied to a plurality of first instances and a plurality of second instances. A plurality of collectors of the system that retrieve information from a first cloud computing environment and a second cloud computing environment, and a controller determines policies for the plurality of first instances and the plurality of second instances. A configurator of the system applies the policies to the plurality of first instances and the plurality of second instances, a first tester that inspects operations of the plurality of first instances and detects violations of the policies, and an enforcer responds to the detected violations. The controller instructs the configurator to apply the first policy to the first instance again, shut down the first instance or cut off communications with the first instance.

Inventors: Bosanac; Jonathan Michael (San Francisco, CA), Geering; Christopher Robert (Berkeley, CA), Eggleston; Jason (Newport Beach, CA), Jasinskyj; Lonhyn (Palo Alto, CA), Sengenberger; John (Meridian, ID)

Applicant: Netskope, Inc. (Santa Clara, CA)

Family ID: 1000008750750

Assignee: Netskope, Inc. (Santa Clara, CA)

Appl. No.: 17/701467

Filed: March 22, 2022

Prior Publication Data

Document Identifier

Publication Date

US 20220217050 A1

Jul. 07, 2022

Related U.S. Application Data

Publication Classification

Int. Cl.: **H04L41/0866** (20220101); **H04L41/0893** (20220101); **H04L41/0894** (20220101)

U.S. Cl.:

CPC **H04L41/0866** (20130101); **H04L41/0893** (20130101); H04L41/0894 (20220501)

Field of Classification Search

CPC: H04L (41/0893); H04L (41/0866); H04L (41/0894); H04L (63/20); H04L (41/0853); H04L (43/50); H04L (43/0817); H04L (41/28); H04L (41/0681); H04L (41/0873)

USPC: 709/220

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
5440723	12/1994	Arnold et al.	N/A	N/A
6353866	12/2001	Fensore	710/104	G06F 13/426
6513122	12/2002	Magdych et al.	N/A	N/A
6622248	12/2002	Hirai	N/A	N/A
7000006	12/2005	Chen	N/A	N/A
7080408	12/2005	Pak et al.	N/A	N/A
7298864	12/2006	Jones	N/A	N/A
7376719	12/2007	Shafer et al.	N/A	N/A
7735116	12/2009	Gauvin	N/A	N/A
7966654	12/2010	Crawford	N/A	N/A
8000329	12/2010	Fendick et al.	N/A	N/A
8112379	12/2011	Voskuil et al.	N/A	N/A
8296178	12/2011	Hudis et al.	N/A	N/A
8719804	12/2013	Jain	N/A	N/A
8793151	12/2013	DelZoppo et al.	N/A	N/A
8839417	12/2013	Jordan	N/A	N/A
9197601	12/2014	Pasdar	N/A	N/A
9225734	12/2014	Hastings	N/A	N/A
9231968	12/2015	Fang et al.	N/A	N/A
9280678	12/2015	Redberg	N/A	N/A
9282118	12/2015	Elzur	N/A	N/A
9729465	12/2016	Labocki et al.	N/A	N/A
9811662	12/2016	Sharpe et al.	N/A	N/A
10051042	12/2017	Diwakar et al.	N/A	N/A
10084825	12/2017	Xu	N/A	N/A
10237282	12/2018	Nelson et al.	N/A	N/A
10334442	12/2018	Vaughn et al.	N/A	N/A
10382468	12/2018	Dods	N/A	N/A

10437625	12/2018	Kaufman et al.	N/A	N/A
10484334	12/2018	Lee et al.	N/A	N/A
10498693	12/2018	Strauss	N/A	H04L 61/50
10826941	12/2019	Jain et al.	N/A	N/A
11032301	12/2020	Mandrychenko et al.	N/A	N/A
11036856	12/2020	Graun et al.	N/A	N/A
11281775	12/2021	Burdett et al.	N/A	N/A
2002/0099666	12/2001	Dryer et al.	N/A	N/A
2003/0055994	12/2002	Herrmann et al.	N/A	N/A
2003/0063321	12/2002	Inoue et al.	N/A	N/A
2003/0172292	12/2002	Judge	N/A	N/A
2003/0204632	12/2002	Willebeek-LeMair et al.	N/A	N/A
2004/0015719	12/2003	Lee et al.	N/A	N/A
2005/0010593	12/2004	Fellenstein et al.	N/A	N/A
2005/0044197	12/2004	Lai	709/223	H04L 67/02
2005/0271246	12/2004	Sharma et al.	N/A	N/A
2006/0156401	12/2005	Newstadt et al.	N/A	N/A
2007/0204018	12/2006	Chandra et al.	N/A	N/A
2007/0237147	12/2006	Quinn et al.	N/A	N/A
2008/0069480	12/2007	Aarabi et al.	N/A	N/A
2008/0134332	12/2007	Keohane et al.	N/A	N/A
2009/0144818	12/2008	Kumar et al.	N/A	N/A
2009/0249470	12/2008	Litvin et al.	N/A	N/A
2009/0300351	12/2008	Lei et al.	N/A	N/A
2010/0017436	12/2009	Wolge	N/A	N/A
2010/0318642	12/2009	Dozier	N/A	N/A
2011/0119481	12/2010	Auradkar et al.	N/A	N/A
2011/0145594	12/2010	Jho et al.	N/A	N/A
2012/0185913	12/2011	Martinez	726/1	H04L 63/20
2012/0278896	12/2011	Fang et al.	N/A	N/A
2013/0159694	12/2012	Chiueh et al.	N/A	N/A
2013/0298190	12/2012	Sikka et al.	N/A	N/A
2013/0347085	12/2012	Hawthorn et al.	N/A	N/A
2014/0013112	12/2013	Cidon et al.	N/A	N/A
2014/0068030	12/2013	Chambers et al.	N/A	N/A
2014/0068705	12/2013	Chambers et al.	N/A	N/A
2014/0259093	12/2013	Narayanaswamy et al.	N/A	N/A
2014/0282843	12/2013	Buruganahalli et al.	N/A	N/A
2014/0359282	12/2013	Shikfa et al.	N/A	N/A
2014/0366079	12/2013	Pasdar	N/A	N/A
2015/0100357	12/2014	Seese et al.	N/A	N/A
2016/0218926	12/2015	Johnson et al.	N/A	N/A
2016/0323318	12/2015	Terrill et al.	N/A	N/A
2016/0350145	12/2015	Botzer et al.	N/A	N/A
2017/0064005	12/2016	Lee	N/A	N/A
2017/0093917	12/2016	Chandra et al.	N/A	N/A
2017/0250877	12/2016	Seyvet et al.	N/A	N/A
2017/0250951	12/2016	Wang et al.	N/A	N/A

2018/0069948	12/2017	Blank et al.	N/A	N/A
2018/0115463	12/2017	Sinha et al.	N/A	N/A
2018/0173604	12/2017	Bhat et al.	N/A	N/A
2018/0234459	12/2017	Kung	N/A	H04L 63/0263
2018/0241751	12/2017	Kruse	N/A	H04L 63/102
2020/0004589	12/2019	Geiger et al.	N/A	N/A
2020/0028894	12/2019	Memon	N/A	G06F 3/0605
2020/0050686	12/2019	Kamalapuram et al.	N/A	N/A
2020/0295999	12/2019	Anandam et al.	N/A	N/A
2020/0364078	12/2019	Potter	N/A	N/A

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
1063833	12/1999	EP	N/A

OTHER PUBLICATIONS

Yifei Yuan, Rajeev Alur, and Boon Thau Loo. 2014. NetEgg: Programming Network Policies by Examples. In Proceedings of the 13th ACM Workshop on Hot Topics in Networks (HotNets-XIII). Association for Computing Machinery, New York, NY, USA, 1-7. (Year: 2014). cited by examiner

Tang, Y., Cheng, G., Xu, Z. et al. Automatic belief network modeling via policy inference for SDN fault localization. J Internet Serv Appl 7, 1 (2016) (Year: 2016). cited by examiner

Martin, Victoria “Cooperative Security Fabric,” The Fortinet Cookbook, Jun. 8, 2016, 6 pgs., archived Jul. 28, 2016 at <https://web.archive.org/web/20160728170025/http://cookbook.fortinet.com/cooperative-security-fabric-54>. cited by applicant

Huckaby, Jeff Ending Clear Text Protocols, Rackaid.com, Dec. 9, 2008, 3 pgs. cited by applicant

Nevvton, Harry “fabric,” Newton's Telecom Dictionary, 30th Updated, Expanded, Anniversary Edition, 2016, 3 pgs. cited by applicant

Fortinet, “Fortinet Security Fabric Earns 100% Detection Scores Across Several Attack Vectors in NSS Labs' Latest Breach Detection Group Test [press release]”, Aug. 2, 2016, 4 pgs, available at <https://www.fortinet.com/de/corporate/about-us/newsroom/press-releases/2016/security-fabric-earns-100-percent-breach-detection-scores-nss-labs>. cited by applicant

Fortinet, “Fortinet Security Fabric Named 2016 CRN Network Security Product of the Year [press release]”, Dec. 5, 2016, 4 pgs, available at <https://www.fortinet.com/corporate/about-us/newsroom/press-releases/2016/fortinet-security-fabric-named-2016-crn-network-security-product>. cited by applicant

McCullagh, Declan, “How safe is instant messaging? A security and privacy survey,” CNET, Jun. 9, 2008, 14 pgs. cited by applicant

Beck et al. “IBM and Cisco: Together for a World Class Data Center,” IBM Redbooks, Jul. 2013, 654 pgs. cited by applicant

Martin, Victoria “Installing internal FortiGates and enabling a security fabric,” The Fortinet Cookbook, Jun. 8, 2016, 11 pgs, archived Aug. 28, 2016 at <https://web.archive.org/web/20160828235831/http://cookbook.fortinet.com/installing-isfw-fortigate-enabling-csf-54>. cited by applicant

Zetter, Kim, “Revealed: The Internet's Biggest Security Hole,” Wired, Aug. 26, 2008, 13 pgs. cited by applicant

Adya et al., Farsite: Federated, available, and reliable storage for an incompletely trusted environment, SIGOPS Oper. Syst. Rev. 36, SI, Dec. 2002, pp. 1-14. cited by applicant

Agrawal et al., "Order preserving encryption for numeric data," In Proceedings of the 2004 ACM SIGMOD international conference on Management of data, Jun. 2004, pp. 563-574. cited by applicant

Balakrishnan et al., "A layered naming architecture for the Internet," ACM SIGCOMM Computer Communication Review, 34(4), 2004, pp. 343-352. cited by applicant

Downing et al. , Naming Dictionary of Computer and Internet Terms, (11th Ed.) Barron's, 2013, 6 pgs. cited by applicant

Downing et al., Dictionary of Computer and Internet Terms, (10th Ed.) Barron's, 2009, 4 pgs. cited by applicant

Zoho Mail, "Email Protocols: What they are & their different types," 2006, 7 pgs. available at <https://www.zoho.com/mail/glossary/email-protocols.html#:~:text=mode of communication.-,What are the different email protocols%3F,and also has defined functions.> cited by applicant

NIIT, Special Edition Using Storage Area Networks, Que, 2002, 6 pgs. cited by applicant

Chapple, Mike, "Firewall redundancy: Deployment scenarios and benefits," Tech Target, 2005, 5 pgs. available at https://www.techtarget.com/searchsecurity/tip/Firewall-redundancy-Deployment-scenarios-and-benefits?%20Offer=abt_pubpro_AI-Insider. cited by applicant

Fortinet, FortiGate—3600 User Manual (vol. 1 , Version 2.50 MR2) Sep. 5, 2003, 329 pgs. cited by applicant

Fortinet, FortiGate SOHO and SMB Configuration Example, (Version 3.0 MR5), Aug. 24, 2007, 54 pgs. cited by applicant

Fortinet, FortiSandbox—Administration Guide, (Version 2.3.2), Nov. 9, 2016, 191 pgs. cited by applicant

Fortinet, FortiSandbox Administration Guide, (Version 4.2.4) Jun. 12, 2023, 245 pgs. available at https://fortinetweb.s3.amazonaws.com/docs.fortinet.com/v2/attachments/fba32b46-b7c0-11ed-8e6d-fa163e15d75b/FortiSandbox-4.2.4-Administration_Guide.pdf. cited by applicant

Fortinet, FortiOS—Administration Guide, (Versions 6.4.0), Jun. 3, 2021, 1638 pgs. cited by applicant

Heady et al., "The Architecture of a Network Level Intrusion Detection System," University of New Mexico, Aug. 15, 1990, 21 pgs. cited by applicant

Kephart et al., "Fighting Computer Viruses," Scientific American (vol. 277, No. 5) Nov. 1997, pp. 88-93. cited by applicant

Wang, L., Chapter 5: Cooperative Security in D2D Communications, "Physical Layer Security in Wireless Cooperative Networks," 41 pgs. first online on Sep. 1, 2017 at https://link.springer.com/chapter/%2010.1007/978-3-319-61863-0_5. cited by applicant

Lee et al., "A Data Mining Framework for Building Intrusion Detection Models," Columbia University, n.d. 13 pgs, 1999. cited by applicant

Merriam-Webster Dictionary, 2004, 5 pgs. cited by applicant

Microsoft Computer Dictionary, (5th Ed.), Microsoft Press, 2002, 8 pgs. cited by applicant

Microsoft Computer Dictionary, (4th Ed.), Microsoft Press, 1999, 5 pgs. cited by applicant

Mika et al. "Metadata Statistics for a Large Web Corpus," LDOW2012, Apr. 16, 2012, 6 pgs. cited by applicant

Oxford Dictionary of Computing (6th Ed.), 2008, 5 pgs. cited by applicant

Paxson, Vern, "Bro: a System for Detecting Network Intruders in Real-Time," Proceedings of the 7th USENIX Security Symposium, Jan. 1998, 22 pgs. cited by applicant

Fortinet Inc., U.S. Appl. No. 62/503,252, "Building a Cooperative Security Fabric of Hierarchically Interconnected Network Security Devices." n.d., 87 pgs. cited by applicant

Song et al., "Practical techniques for searches on encrypted data," In Proceeding 2000 IEEE symposium on security and privacy. S&P 2000, May 2000, pp. 44-55. cited by applicant

Dean, Tamara, Guide to Telecommunications Technology, Course Technology, 2003, 5 pgs. cited by applicant

U.S. Appl. No. 60/520,577, “Device, System, and Method for Defending a Computer Network,” Nov. 17, 2003, 21 pgs. cited by applicant
U.S. Appl. No. 60/552,457, “Fortinet Security Update Technology,” Mar. 2004, 6 pgs. cited by applicant
Tittel, Ed, Unified Threat Management for Dummies, John Wiley & Sons, Inc., 2012, 76 pgs. cited by applicant
Fortinet, FortiOS Handbook: UTM Guide (Version 2), Oct. 15, 2010, 188 pgs. cited by applicant
Full Definition of Security, Wayback Machine Archive of Merriam-Webster on Nov. 17, 2016, 1 pg. cited by applicant
Definition of Cooperative, Wayback Machine Archive of Merriam-Webster on Nov. 26, 2016, 1 pg. cited by applicant
Pfaffenberger, Bryan, Webster's New World Computer Dictionary, (10th Ed.), 2003, 5 pgs. cited by applicant

Primary Examiner: Elfervig; Taylor A

Attorney, Agent or Firm: MUGHAL GAUDRY & FRANKLIN PC

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS (1) This application is a continuation of U.S. patent application Ser. No. 17/101,892, entitled “MULTI-ENVIRONMENT NETWORKING MANAGEMENT SYSTEM,” filed on Nov. 23, 2020, the contents of which are incorporated by reference in their entirety for all purposes.

BACKGROUND

(1) This disclosure relates in general to systems and methods for managing policies across multiple cloud computing environments within a network.

(2) Different cloud computing environments use different languages, data sources, commands, and protocols. For example, each cloud computing provider may use a different method to allocate subnets and IP addresses to instances within the respective cloud computing environment. This may cause instances within a network having multiple cloud computing environments to have overlapping subnets and IP addresses. Further, each cloud computing provider may require different instructions for applying policies to instances within the respective cloud computing environment.

SUMMARY

(3) Exemplary embodiments of the invention provide systems and methods for determining, enforcing, and managing policies across different cloud computing environments within a network. According to an aspect of the invention, a system includes a user interface that receives configuration settings to be applied to a plurality of first instances and a plurality of second instances. A plurality of collectors of the system that retrieve information from a first cloud computing environment and a second cloud computing environment, and a controller determines policies for the plurality of first instances and the plurality of second instances. A configurator of the system applies the policies to the plurality of first instances and the plurality of second instances, a first tester that inspects operations of the plurality of first instances and detects violations of the policies, and an enforcer responds to the detected violations by receiving a notification from the first tester that a first instance from the plurality of first instances violated a first policy. The controller instructs the configurator to apply the first policy to the first instance again, shut down the first instance or cut off communications with the first instance.

(4) According to another aspect of the invention, a method may include receiving configuration settings to be applied to a plurality of first instances within a first cloud computing environment and a plurality of second instances within a second cloud computing environment. In one step, information from the first cloud computing environment and the second cloud computing environment is retrieved. The information comprises a plurality of functionalities of the first cloud computing environment and the second cloud computing environment. The method further includes determining policies for the plurality of first instances within the first cloud computing environment and the plurality of second instances within the second cloud computing environment as functions of the configuration settings and the information. The policies are applied to the plurality of first instances within the first cloud computing environment and the plurality of second instances within the second cloud computing environment. Operations of the plurality of first instances within the first cloud computing environment and the plurality of second instances within the second cloud computing environment are inspected and violations of the policies by the plurality of first instances within the first cloud computing environment and the plurality of second instances within the second cloud computing environment are detected. In addition, the method also includes responding to the detected violations by receiving a notification that a first instance from the plurality of first instances violated a first policy. A controller instructs a configurator to apply the first policy to the first instance again, shut down the first instance or cut off communications with the first instance.

(5) According to another aspect of the invention, a system may include a user interface that receives configuration settings to be applied to a plurality of first instances within a first cloud computing environment and a plurality of second instances within a second cloud computing environment. The system may also include a plurality of collectors that retrieve information from the first cloud computing environment and the second cloud computing environment. The information comprises a plurality of functionalities of the first cloud computing environment and the second cloud computing environment. In addition, the system may include a controller that determines policies for the plurality of first instances within the first cloud computing environment and the plurality of second instances within the second cloud computing environment as functions of the configuration settings and the information. Further, the system may include a configurator that applies the policies to the plurality of first instances within the first cloud computing environment and the plurality of second instances within the second cloud computing environment, a first tester that inspects operations of the plurality of first instances within the first cloud computing environment and detects violations of the policies by the plurality of first instances within the first cloud computing environment, and an enforcer that responds to the detected violations by receiving a notification from the first tester that a first instance from the plurality of first instances violated a first policy. The controller instructs the configurator to apply the first policy to the first instance again, shut down the first instance or cut off communications with the first instance.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) The present disclosure is described in conjunction with the appended figures:
- (2) FIG. 1 depicts a block diagram of an embodiment of a network;
- (3) FIG. 2 depicts a block diagram of an embodiment of a cloud computing environment that includes a control system;
- (4) FIG. 3 depicts a block diagram of an embodiment of a portion of the network shown in FIG. 1;
- (5) FIG. 4 depicts a block diagram of a cloud Open Systems Interconnection (OSI) model for cloud computing environments;

- (6) FIG. 5 depicts a flowchart of an embodiment of a method; and
- (7) FIG. 6 depicts a flowchart of an embodiment of a portion of the method shown in FIG. 5.
- (8) In the appended figures, similar components and/or features may have the same reference label. Further, various components of the same type may be distinguished by following the reference label by a dash and a second label that distinguishes among the similar components. If only the first reference label is used in the specification, the description is applicable to any one of the similar components having the same first reference label irrespective of the second reference label.

DETAILED DESCRIPTION

- (9) The ensuing description provides preferred exemplary embodiment(s) only, and is not intended to limit the scope, applicability or configuration of the disclosure. Rather, the ensuing description of the preferred exemplary embodiment(s) will provide those skilled in the art with an enabling description for implementing a preferred exemplary embodiment. It is understood that various changes may be made in the function and arrangement of elements without departing from the spirit and scope as set forth in the appended claims.
- (10) Referring first to FIG. 1, a block diagram of an embodiment of a network **100** is shown. The network **100** may include a first cloud computing environment **140a**, a second cloud computing environment **140b**, and a third cloud computing environment **140c** that communicate via a public Internet **125**. The first cloud computing environment **140a**, the second cloud computing environment **140b**, and the third cloud computing environment **140c** may be public clouds. Some examples of the first cloud computing environment **140a**, the second cloud computing environment **140b**, and the third cloud computing environment **140c** include Amazon Web Services (AWS)®, Google Cloud Platform (GCP)®, and Microsoft Azure®. Some or all of the first cloud computing environment **140a**, the second cloud computing environment **140b**, and the third cloud computing environment **140c** may be different from each other. For example, the first cloud computing environment **140a** may run Amazon Web Services (AWS)®, the second cloud computing environment **140b** may run Google Cloud Platform (GCP)®, and the third cloud computing environment **140c** may run Microsoft Azure®. Although three cloud computing environments are shown, any suitable number of cloud computing environments may be provided.
- (11) Each of the cloud computing environments may communicate with the Internet via a secure connection. For example, the first cloud computing environment **140a** may communicate with the public Internet **125** via a virtual private network (VPN) **190a**, the second cloud computing environment **140b** may communicate with the public Internet **125** via a VPN **190b**, and the third cloud computing environment **140c** may communicate with the public Internet **125** via a VPN **190c**.
- (12) A plurality of enterprises **198** may also communicate with the public Internet **125** via a VPN **190d**. Some examples of the enterprises **198** may include corporations, educational facilities, governmental entities, and private consumers. In addition, the plurality of enterprises **198** may communicate with a plurality of first domain users **195a** via a VPN **190f**, a plurality of second domain users **195b** via a VPN **190g**, and a plurality of third domain users **195c** via a VPN **190h**. Some examples of the first domain users **195a**, the second domain users **195b**, and the third domain users **195c** may include individual users that are authorized to use computing resources of the enterprises **198**.
- (13) Further, a control system **185** may communicate with the public Internet **125** via a VPN **190e**. As discussed in further detail below, the control system **185** may configure, test, and enforce policies across the first cloud computing environment **140a**, the second cloud computing environment **140b**, and the third cloud computing environment **140c**. For example, the control system **185** may ensure that the policies are consistent across the first cloud computing environment **140a**, the second cloud computing environment **140b**, and the third cloud computing environment **140c**.
- (14) With reference to FIG. 2, a block diagram of an embodiment of a cloud computing

environment **200** that includes the control system **185** is shown. The cloud computing environment **200** may be a private cloud. A firewall **235** may be provided for the control system **185**. Some examples of the firewall **235** may include a proxy firewall, a stateful inspection firewall, a unified threat management (UTM) firewall, a next-generation firewall (NGFW), a threat-focused NGFW, and a virtual firewall.

(15) The control system **185** may include a plurality of user interfaces **205**. The user interfaces **205** may allow users to provide input to the control system **185**. Some examples of the user interfaces **205** may include a keyboard, a mouse, a touchpad or touch screen on a display, a scroll wheel, a keypad, and an audio input device. For example, the user interfaces **205** may receive configuration settings to be applied to instances within the first cloud computing environment **140a**, the second cloud computing environment **140b**, and/or the third cloud computing environment **140c**.

(16) The control system **185** may also include a controller **215**, a plurality of collectors **220**, a configurator **255**, a plurality of testers **265**, an enforcer **270**, and a reporter **280**. The controller **215**, the configurator **255**, the testers **265**, the enforcer **270**, and the reporter **280** may be modules within a computing system or may be separate computing systems that are communicatively coupled. The computing systems may have various components such as processors, storage subsystems, and communications subsystems. Some examples of the computing systems may include personal computers, workstations, mainframes, server racks, and handheld portable devices.

(17) The collectors **220** may retrieve information from the first cloud computing environment **140a**, the second cloud computing environment **140b**, and/or the third cloud computing environment **140c**. The information retrieved by the collectors **220** may include functionalities of the first cloud computing environment **140a**, the second cloud computing environment **140b**, and/or the third cloud computing environment **140c**, such as network configurations, a firewall rules, cloud application programming interfaces (APIs), resources, cloud service providers, and data sets. The information retrieved by the collectors **220** may also include data input types, data types, data sizes, or data ages of the first cloud computing environment **140a**, the second cloud computing environment **140b**, and/or the third cloud computing environment **140c**.

(18) The controller **215** may determine policies for the instances within the first cloud computing environment **140a**, the second cloud computing environment **140b**, and/or the third cloud computing environment **140c**. For example, the policies may include firewall rules, forwarding rules, network configurations, cross-cloud routing rules, IP addressing rules, cross-cloud peering rules, security group management rules, storage bucket access rules, resource management rules, or subnet configurations. The controller **215** may determine general policies for all of the instances as functions of the configuration settings that are received from the user interfaces **205**. Because the first cloud computing environment **140a**, the second cloud computing environment **140b**, and/or the third cloud computing environment **140c** may use different languages, data sources, commands, and protocols, the controller **215** may also translate the general policies into specific policies for instances within each different cloud computing environment. For example, the controller **215** may determine specific policies according to the information retrieved from the first cloud computing environment **140a**, the second cloud computing environment **140b**, and/or the third cloud computing environment **140c**.

(19) The configurator **255** may receive the policies from the controller **215** and apply the policies to the instances. For each specific policy, the configurator **255** may retrieve a script from an API and execute the script in order to apply the specific policy to the instances within one of the cloud computing environments. The script may include instructions for the instances to implement the specific policy. Any number of specific policies may be applied to the instances within the cloud computing environment. In this example, the configurator **255** pushes the policies directly to the instances in order to update the configurations of the instances.

(20) In another example, a metadata endpoint **250** may be provided for the configurator **255**. The metadata endpoint **250** may receive the policies from the configurator **255** and host changes to the

instances according to the policies. In this example, the instances retrieve the changes from the metadata service endpoint and apply the changes. The instances may subscribe to the metadata endpoint **250** and periodically check the metadata endpoint **250** for any updates. In this example, the instances pull the policies from the metadata endpoint **250** in order to update the configurations of the instances.

(21) The testers **265** may inspect operations of the instances and detect violations of the policies by the instances. In some examples, a different tester **265** may be provided for each cloud computing environment. In other examples, a single tester **265** may be provided for a plurality of cloud computing environments. As described in further detail below, the testers **265** may detect violations of one or more specific policies by any of the instances.

(22) The enforcer **270** may respond to the violations in a variety of ways. For example, the enforcer **270** may send a notification of the detected violation to the controller **215**, which may direct the configurator **255** to apply the policy that was violated to the non-complying instance. In one example, if the policy requires the instances to have non-overlapping IP addresses and the testers **265** identify an overlap between the IP addresses, the enforcer **270** may direct the configurator **255** to request new IP addresses from at least one of the cloud computing environments. This procedure may be repeated until the number of overlapping IP addresses has been reduced or eliminated.

(23) Alternatively or in addition, the reporter **280** may send a notification of the detected violation to at least one of the user interfaces **205**. The notification may identify the instance that violated the policy, the cloud computing environment in which the instance is located, and the policy that was violated. The reporter **280** may send the notification via short message service (SMS), email, API call, or another notification method.

(24) Referring next to FIG. **3**, a block diagram of an embodiment of a portion **300** of the network **100** is shown. The first cloud computing environment **140a** may include a plurality of first instances **345a**, the second cloud computing environment **140b** may include a plurality of second instances **345b**, and the third cloud computing environment **140c** may include a plurality of third instances **345c**. Some examples of the first instances **345a**, the second instances **345b**, and the third instances **345c** may include virtual machines that emulate computer systems. The virtual machines may run various software packages. The first instances **345a**, the second instances **345b**, and the third instances **345c** may be examples of the instances discussed above with respect to FIG. **2**.

(25) A metadata endpoint **350** may be provided for one, some, or all of the instances, such as a fourth instance **345d** within the third cloud computing environment **140c**. As discussed above, the metadata endpoint **350** may receive instructions from the configurator **255** and apply the policies to the fourth instance **345d**.

(26) With reference to FIG. **4**, a block diagram of a cloud Open Systems Interconnection (OSI) model **400** for cloud computing environments is shown. The cloud OSI model **400** for cloud computing environments partitions the flow of data in a communication system into six layers of abstraction. The cloud OSI model **400** for cloud computing environments may include, in order, an application layer **410**, a service layer **415**, an image layer **420**, a software-defined data center layer **425**, a hypervisor layer **430**, and an infrastructure layer **435**. Each layer serves a class of functionality to the layer above it and is served by the layer below it. Classes of functionality may be realized in software by various communication protocols.

(27) The infrastructure layer **435** may include hardware, such as physical devices in a data center, that provides the foundation for the rest of the layers. The infrastructure layer **435** may transmit and receive unstructured raw data between a device and a physical transmission medium. For example, the infrastructure layer **435** may convert the digital bits into electrical, radio, or optical signals.

(28) The hypervisor layer **430** may perform virtualization, which may allow the physical devices to be divided into virtual machines that can be bin packed onto physical machines for greater efficiency. The hypervisor layer **430** may provide virtualized compute, storage, and networking. For example, OpenStack® software that is installed on bare metal servers in a data center may

provide virtualization cloud capabilities. The OpenStack® software may provide various infrastructure management capabilities to cloud operators and administrators, and may utilize the Infrastructure-as-Code concept for deployment and lifecycle management of a cloud data center. In the Infrastructure-as-Code concept, the infrastructure elements are described in definition files. Changes in the files are reflected in the configuration of data center hosts and cloud services.

(29) The software-defined data center layer **425** may provide resource pooling, usage tracking, and governance on top of the hypervisor layer **430**. The software-defined data center layer **425** may enable the creation virtualization for the Infrastructure-as-Code concept by using representational state transfer (REST) APIs. The management of block storage devices may be virtualized, and end users may be provided with a self-service API to request and consume those resources without requiring any knowledge of where the storage is actually deployed or on what type of device. Various compute nodes may be balanced for storage.

(30) The image layer **420** may use various operating systems and other pre-installed software components. Patch management may be used to identify, acquire, install, and verify patches for products and systems. Patches may be used to correct security and functionality problems in software. Patches may also be used to add new features to operating systems, including security capabilities. The image layer **420** may focus on the compute instead of storage and networking. The instances within the cloud computing environments may be provided at the image layer **420**.

(31) The service layer **415** may provide middleware, such as functional components that applications use in tiers. In some examples, the middleware components may include databases, load balancers, web servers, message queues, email services, or other notification methods. The middleware components may be defined at the service layer **415** on top of particular images from the image layer **420**. Different cloud computing environment providers may have different middleware components.

(32) The application layer **420** may interact with software applications that implement a communicating component. The application layer **420** is the layer that is closest to the end user. Functions of the application layer **420** may include identifying communication partners, determining resource availability, and synchronizing communication. Applications within the application layer **420** may include custom code that makes use of middleware defined in the service layer **415**.

(33) Various features discussed above may be performed at one or more layers of the cloud OSI model **400** for cloud computing environments. For example, translating the general policies into specific policies for different cloud computing environments may be performed at the service layer **415** and the software-defined data center layer **425**. Various scripts may be updated across the service layer **415**, the image layer **420**, and the software-defined data center layer **425**. Further, APIs and policies may operate at the software-defined data center layer **425** and the hypervisor layer **430**.

(34) Each of the different cloud computing environments may have different service layers **415**, image layers **420**, software-defined data center layers **425**, hypervisor layers **430**, and infrastructure layers **435**. Further, each of the different cloud computing environments may have an application layer **410** that can make calls to the specific policies in the service layer **415** and the software-defined data center layer **425**. The application layer **410** may have substantially the same format and operation for each of the different cloud computing environments. Accordingly, developers for the application layer **410** may not need to understand the peculiarities of how each of the cloud computing environments operates in the other layers.

(35) Referring next to FIG. 5, a flowchart of an embodiment of a method **500** is shown. The method **500** begins at block **510** where configuration settings to be applied to instances within a plurality of different cloud computing environments are received. For example, the configuration settings may be received via the user interfaces **205** and provided to the controller **215**. The configuration settings may be intended for the first instances **345a** within the first cloud computing

environment **140a** and the second instances **345b** within the second cloud computing environment **140b**. The first cloud computing environment **140a** and the second cloud computing environment **140b** may be run by different providers and may use different languages, data sources, commands, and protocols.

(36) The method **500** continues at block **515** where information from the different cloud computing environments is retrieved. For example, information from the first cloud computing environment **140a** and the second cloud computing environment **140b** may be retrieved by the collectors **220**. One of the collectors **220** may retrieve the information from the first cloud computing environment **140a**, and another one of the collectors may retrieve the information from the second cloud computing environment **140b**. Both the configuration settings and the information may be provided to the controller **215**.

(37) The method **500** continues at block **520** where general policies are determined as functions of the configuration settings. For example, the controller **215** may determine general policies that apply to the first instances **345a** within the first cloud computing environment **140a** and the second instances **345b** within the second cloud computing environment **140b**. The controller **215** may determine general policies that apply to all of the instances within some or all of the cloud computing environments within a network. This may ensure that the policies are consistent across the different cloud computing environments. For example, the policies may include firewall rules, forwarding rules, network configurations, cross-cloud routing rules, IP addressing rules, cross-cloud peering rules, security group management rules, storage bucket access rules, resource management rules, or subnet configurations.

(38) The method **500** continues at block **525** where the general policies are translated to specific policies for instances within the different cloud computing environments. For example, the controller **215** may use the information about the first cloud computing environment **140a** and the second cloud computing environment **140b** to determine specific policies for instances within the first cloud computing environment **140a** and the second cloud computing environment **140b**, respectively. This translation may be performed at the service layer **415** and the software-defined data center layer **425** of the cloud OSI model **400** for cloud computing environments.

(39) The method **500** continues at block **530** where scripts corresponding to the specific policies for the instances within the first cloud computing environment **140a** and the second cloud computing environment **140b** are retrieved. For example, the configurator **225** may receive the specific policies from the controller **215** and retrieve scripts corresponding to the specific policies from an API. The scripts may be written in JavaScript Object Notation (JSON)®. For each of the first cloud computing environment **140a** and the second cloud computing environment **140b**, a separate script may be retrieved for each of the specific policies. The following is an example of a script that may be used to establish a firewall rule for the first instances **345a** within the first cloud computing environment **140a**, where the first cloud computing environment **140a** is run by Amazon Web Services (AWS)®:

(40) Further, the following is an example of a script that may be used to establish the same firewall rule for the second instances **345b** within the second cloud computing environment **140b**, where the second cloud computing environment **140b** is run by Google Cloud Platform (GCP)®:

(41) TABLE-US-00001 C02XW1WCJGH6:~ username\$ gcloud compute instances list --quiet --project example-project --account username@netskope.com --format="json" [{
 "canIpForward": false, "cpuPlatform": "Intel Broadwell", "creationTimestamp":
 "2020-05-07T21:42:51.931-07:00", "deletionProtection": false, "description": "",
 "disks": [{ "autoDelete": true, "boot": true,
 "deviceName": "example-sample", "diskSizeGb": "500",
 "guestOsFeatures": [{ "type":
"VIRTIO_SCSI_MULTIQUEUE" }], "index": 0,
 "interface": "SCSI", "kind": "compute#attachedDisk", "licenses": [

```
    "https://www.googleapis.com/compute/v1/projects/ubuntu-os-
cloud/global/licenses/ubuntu-1604-xenial"    ],    "mode": "READ_WRITE",
    "source": "https://www.googleapis.com/compute/v1/projects/example-project/zones/us-
west1-b/disks/example-sample",    "type": "PERSISTENT"    }    ],
    "displayDevice": {    "enableDisplay": false    },    "fingerprint":
"4GTjmojZZPQ=",    "id": "123456789098765432",    "kind": "compute#instance",
    "labelFingerprint": "L_2Mh2A8RIA=",    "labels": {    "env": "test",    "team":
"data"    },    "lastStartTimestamp": "2020-05-07T21:42:56.861-07:00",
    "machineType": "https://www.googleapis.com/compute/v1/projects/example-
project/zones/us- west1-b/machineTypes/custom-8-32768",    "metadata": {
    "fingerprint": "MMMMMMMMMMMMMM=",    "kind": "compute#metadata"    },
    "name": "example-sample",    "networkInterfaces": [    {    "accessConfigs":
[    {    "kind": "compute#accessConfig",    "name": "External
NAT",
    "natIP": "252.1.2.3",    "networkTier": "PREMIUM",
    "type": "ONE_TO_ONE_NAT"    }    ],    "fingerprint":
"00_ffffff=",    "kind": "compute#networkInterface",    "name": "nic0",
    "network": "https://www.googleapis.com/compute/v1/projects/example-
project/global/networks/sample-example",    "networkIP": "10.240.128.18",
    "subnetwork": "https://www.googleapis.com/compute/v1/projects/example-
project/regions/us-west1/subnetworks/sample-example-priv-vpn-subnet"    }    ],
    "reservationAffinity": {    "consumeReservationType": "ANY_RESERVATION"    },
    "scheduling": {    "automaticRestart": true,    "onHostMaintenance": "MIGRATE",
    "preemptible": false    },    "selfLink":
"https://www.googleapis.com/compute/v1/projects/example-project/zones/us- west1-
b/instances/example-sample",    "serviceAccounts": [    {    "email":
"1000000000000-compute@developer.gserviceaccount.com",    "scopes": [
    "https://www.googleapis.com/auth/devstorage.read_only",
    "https://www.googleapis.com/auth/logging.write",
    "https://www.googleapis.com/auth/monitoring.write",
    "https://www.googleapis.com/auth/servicecontrol",
    "https://www.googleapis.com/auth/service.management.readonly",
    "https://www.googleapis.com/auth/trace.append"    ]    }    ],
    "startRestricted": false,    "status": "RUNNING",    "tags": {    "fingerprint": "1-
888_XXXXX=",    "items": [    "https-server",    "nspublic"    ]    },
    "zone": "https://www.googleapis.com/compute/v1/projects/example-project/zones/us-west1-
b"    },    {    "canIpForward": false,    "cpuPlatform": "Intel Broadwell",
    "creationTimestamp": "2020-04-17T14:54:58.261-07:00",    "deletionProtection": false,
    "description": "",    "disks": [    {    "autoDelete": true,    "boot": true,
    "deviceName": "test-vm",    "diskSizeGb": "500",    "guestOsFeatures":
[    {    "type": "VIRTIO_SCSI_MULTIQUEUE"    }    ],
    "index": 0,    "interface": "SCSI",    "kind": "compute#attachedDisk",
    "licenses": [    "https://www.googleapis.com/compute/v1/projects/ubuntu-os-
cloud/global/licenses/ubuntu-1604-xenial"    ],    "mode": "READ_WRITE",
    "source": "https://www.googleapis.com/compute/v1/projects/example-project/zones/us-
west1-b/disks/test-vm",    "type": "PERSISTENT"    }    ],    "displayDevice":
{    "enableDisplay": false    },    "fingerprint": "ZZZZZZZZZZZZ=",    "id":
"55555555555555555555",    "kind": "compute#instance",    "labelFingerprint":
"XBXXBXXBXXBXX=",    "labels": {    "env": "test",    "team": "webui"    },
    "lastStartTimestamp": "2020-04-29T10:59:48.190-07:00",    "lastStopTimestamp": "2020-
04-29T10:59:30.707-07:00",    "machineType":
```

```

“https://www.googleapis.com/compute/v1/projects/example-project/zones/us-west1-
b/machineTypes/custom-8-44544”,      “metadata”: {      “fingerprint”: “LMFFFFFF-FFF=”,
      “items”: [      {      “key”: “serial-port-enable”,      “value”: “true”
      }      ],      “kind”: “compute#metadata”      },      “name”: “test-vm”,
“networkInterfaces”: [      {      “accessConfigs”: [      {
      “kind”: “compute#accessConfig”,      “name”: “External NAT”,
      “natIP”: “35.1.1.15”,      “networkTier”: “PREMIUM”,
      “type”: “ONE_TO_ONE_NAT”      }      ]      },      “fingerprint”:
“BBBBBrIIII=”,      “kind”: “compute#networkInterface”,      “name”: “nic0”,
      “network”: “https://www.googleapis.com/compute/v1/projects/example-
project/global/networks/sample-example”,      “networkIP”: “ 10.240.131.209”,
      “subnetwork”: “https://www.googleapis.com/compute/v1/projects/example-
project/regions/us-west1/subnetworks/sample-example-priv-vpn-subnet”      }      ],
“reservationAffinity”: {      “consumeReservationType”: “ANY_RESERVATION”      },
“scheduling”: {      “automaticRestart”: true,      “onHostMaintenance”: “MIGRATE”,
      “preemptible”: false      },      “selfLink”:
“https://www.googleapis.com/compute/v1/projects/example-project/zones/us-
west1-b/instances/test-vm”,      “serviceAccounts”: [      {      “email”: “1000000000000-
compute@developer.gserviceaccount.com”,      “scopes”: [
      “https://www.googleapis.com/auth/devstorage.read_only”,
      “https://www.googleapis.com/auth/logging.write”,
      “https://www.googleapis.com/auth/monitoring.write”,
      “https://www.googleapis.com/auth/servicecontrol”,
      “https://www.googleapis.com/auth/service.management.readonly”,
      “https://www.googleapis.com/auth/trace.append”      ]      }      ],
      “startRestricted”: false,      “status”: “RUNNING”,      “tags”: {      “fingerprint”: “1-
888_XXXXX=”,      “items”: [      “https-server”,      “nspublic”      ]      },
      “zone”: “https://www.googleapis.com/compute/v1/projects/example-project/zones/us-west1-
b”      },      {      “canIpForward”: false,      “cpuPlatform”: “Intel Broadwell”,
      “creationTimestamp”: “2020-04-27T16:32:35.259-07:00”,      “deletionProtection”: false,
      “description”: “”,      “disks”: [      {      “autoDelete”: true,      “boot”: true,
      “deviceName”: “example-vm-test”,      “diskSizeGb”: “500”,
      “guestOsFeatures”: [      {      “type”:
“VIRTIO_SCSI_MULTIQUEUE”      }      ],      “index”: 0,
      “interface”: “SCSI”,      “kind”: “compute#attachedDisk”,      “licenses”: [
      “https://www.googleapis.com/compute/v1/projects/ubuntu-os-
cloud/global/licenses/ubuntu-1604-xenial”      ],      “mode”: “READ_WRITE”,
      “source”: “https://www.googleapis.com/compute/v1/projects/example-project/zones/us-
west1-b/disks/example-vm-test”,      “type”: “PERSISTENT”      }      ],
      “displayDevice”: {      “enableDisplay”: false      },      “fingerprint”: “1111111111=”,
      “id”: “8978978978978978978”,      “kind”: “compute#instance”,      “labelFingerprint”:
“4242424242=”,      “lastStartTimestamp”: “2020-04-27T16:32:41.002-07:00”,
      “machineType”: “https://www.googleapis.com/compute/v1/projects/example-
project/zones/us-west1-b/machineTypes/n1-standard-8”,      “metadata”: {      “fingerprint”:
“MMMMMMMMMMMMMM=”,      “kind”: “compute#metadata”      },      “name”: “example-
vm-test”,      “networkInterfaces”: [      {      “accessConfigs”: [

```

(42) The method **500** continues at block **535** where the scripts corresponding to the specific policies for the instances within the first cloud computing environment **140a** and the second cloud computing environment **140b** are confirmed and executed. For example, the configurator **225** may confirm that the first script quoted above will instruct the first instances **345a** within the first cloud

computing environment **140a** to establish the desired firewall rule. The configurator **225** may then execute the first script in order to apply the firewall rule to the first instances **345a** within the first cloud computing environment **140a**. The firewall rule may be applied to the first instances **345a** within the first cloud computing environment **140a** simultaneously or in sequence. The configurator **225** may then confirm that the second script quoted above will instruct the second instances **345b** within the second cloud computing environment **140b** to establish the desired firewall rule. The configurator **225** may then execute the second script in order to apply the firewall rule to the second instances **345b** within the second cloud computing environment **140b**. Again, the firewall rule may be applied to the second instances **345b** within the second cloud computing environment **140b** simultaneously or in sequence. This procedure may be repeated until the firewall rule has been applied to all of the instances in the network.

(43) The method **500** continues at block **540** where operations of the instances within the cloud computing environments are inspected. For example, the testers **265** may inspect operations of the first instances **345a** within the first cloud computing environment **140a** and the second instances **345b** within the second cloud computing environment **140b**. One or more of the testers **265** may be provided for each of the cloud computing environments, or a single tester **265** may be provided for all of the cloud computing environments.

(44) As one example of the testing procedure that may be performed at block **540**, the first tester **265** may determine whether the first instance **345a** was correctly set up as a load balancer. For example, the first tester **265** may inspect the first instance **345a** to determine whether it complies with various policies, such as whether the load balancer was set up, whether the load balancer is running correctly, whether the load balancer has enough processing power, and whether the load balancer has the correct IP address. More specific details of embodiments of the testing procedure are provided below.

(45) Referring next to FIG. **6**, a flowchart of an embodiment of block **540** of the method **500** described in FIG. **5** is shown. The example shown in FIG. **6** describes a method for testing a first instance for compliance with a first policy. However, this example may be expanded to test the first instance for compliance with a plurality of policies. This example may also be expanded to test additional instances for compliance with the first policy and/or the plurality of policies. Any such testing may be performed in parallel and/or in sequence.

(46) The method **540** begins at block **610** where a first policy is identified for testing compliance of the first instance **345a** within the first cloud computing environment **140a**. For example, a first tester **265** may identify the first policy by referencing a table that stores a list of each specific policy that has been applied to each instance within the network **100**. The table may be stored within a data storage component within the control system **185**. The first policy may be identified by a variety of methods, such as a random selection from the specific policies that have been applied to the first instance **345a**, or a selection of the specific policy that was applied to the first instance **345a** at the earliest time.

(47) The method **540** continues at block **615** where a testing schedule is retrieved for the first instance **345a**. For example, the testing schedule may indicate a frequency of testing for the first instance **345a** for compliance with each of the specific policies. The frequency of testing may be constant or may change as a function of time. The testing schedule may also indicate a predetermined time at which to start testing the first instance **345a** after each of the specific policies was applied to the first instance **345a**. The testing schedule may be the same for one or more of the specific policies, or may be different for one or more of the specific policies. The first tester **265** may retrieve the testing schedule from the data storage component within the control system **185**.

(48) The method **540** continues at block **620** where a test module to test the compliance of the first instance **345a** within the first cloud computing environment **140a** with the first policy is retrieved. A separate test module may be stored within the data storage component for each of the specific policies. Each test module may be written for a specific cloud computing environment, and may be

written as a JavaScript Object Notation (JSON)® script. For example, the first tester **265** may retrieve a first test module to test the compliance of the first instance **345a** within the first cloud computing environment **140a** for compliance with the first policy.

(49) The method **540** continues at block **625** where it is determined whether a predetermined time has elapsed since the first policy was applied to the first instance **345a** or the first instance **345a** was last tested for compliance with the first policy. For example, the first tester **265** may refer to the testing schedule for the first instance **345a**, along with a table that stores a list of the times at which the first instance **345a** was tested for compliance with the specific policies. If the first instance **345a** has already been tested for compliance with the first policy, the first tester **265** may determine whether a predetermined time has elapsed since the first instance **345a** was last tested for compliance with the first policy. If the predetermined time has not elapsed, the method **540** may proceed to block **630**. If the predetermined time has elapsed, the method may proceed to block **640**. Similarly, if the first instance **345a** has not yet been tested for compliance with the first policy, the first tester **265** may determine whether another predetermined time has elapsed since the first policy was applied to the first instance **345a**. If the other predetermined time has not elapsed, the method **540** may proceed to block **630**. If the other predetermined time has elapsed, the method may proceed to block **640**.

(50) At block **630** it is determined whether an error message has been received from the first instance **345a**. For example, after the first policy is applied to the first instance **345a**, the first instance **345a** may send an error message to the control system **185** indicating that the first instance **345a** was unable to implement the first policy or maintain compliance with the first policy.

(51) At block **635** it is determined whether a notification of a change to the first cloud computing environment **140a** has been received. For example, after the first policy is applied to the first instance **345a**, the first cloud computing environment **140a** may send a message to the control system **185** indicating a change in its operation or configuration that may affect the compliance of the first instance **345a** with the first policy. If no notifications have been received from the first cloud computing environment **140a**, the method **540** may return to block **625**. If a notification has been received from the first cloud computing environment **140a**, the method **540** may proceed to block **640**.

(52) At block **640** the first instance **345a** is tested for compliance with the first policy by calling the test module corresponding to the first policy. For example, the first tester **265** may call the test module that was retrieved at block **610** to test the first instance **345a** for compliance with the first policy. The testing may determine that the first instance **345a** is complying with the first policy or violating the first policy.

(53) Returning to FIG. 5, the method **500** continues at block **545** where the control system **185** responds to any detected violations. For example, the enforcer **270** may receive a notification from the first tester **265** that the first instance **345a** violated the first policy. The enforcer **270** may then send a notification of the violation to the controller **215**, which may receive the notification and require the first instance **345a** to comply with the first policy. For example, the controller **215** may instruct the configurator **255** to apply the first policy to the first instance **345a** again. Alternatively, the controller **215** may shut down the first instance **345a** or cut off communications with the first instance **345a**. Alternatively or in addition, the reporter **280** may send a notification of the violation to at least one of the user interfaces **205**. For example, the reporter **280** may send the notification of the violation to the user interface **205** that provided the configuration settings for the first instance **345a**.

(54) As another example of the testing and response procedures discussed with respect to FIG. 5, the first tester **265** may inspect the IP addresses that have been assigned to the plurality of first instances **345a** within the first cloud computing environment **140a** at block **540**. The first tester **265** may then determine whether there is a violation of a policy to have non-overlapping IP addresses for the plurality of first instances **345a** at block **540**. If such a violation exists, the first tester **265**

may send a notification to the enforcer **270**. The enforcer **270** may then send a notification of the violation to the controller **215**, which may receive the notification and request the allocation of new IP addresses from the first cloud computing environment **140a** at block **545**. This testing and response procedure may be repeated until the first instances **345a** have non-overlapping IP addresses, or until the number of overlapping IP addresses has been reduced below a threshold. For example, the threshold may be a percentage of the number of the first instances **345a** within the first cloud computing environment **140a**. Alternatively, the controller **215** may shut down any instances having overlapping IP addresses or cut off communications with any instances having overlapping IP addresses. Alternatively or in addition, the reporter **280** may send a notification of the overlapping IP addresses, including a list of the affected instances, to at least one of the user interfaces **205**.

(55) Further, this testing and response procedure may be conducted to inspect the IP addresses that have been assigned to all of the instances within the network **100**. For example, one or more of the testers **265** may inspect the IP addresses that have been assigned to the plurality of first instances **345a** within the first cloud computing environment **140a**, the plurality of second instances **345b** within the first cloud computing environment **140b**, and the plurality of third instances **345c** within the third cloud computing environment **140c**. The remainder of the testing and response procedure may be the same as discussed above.

(56) Specific details are given in the above description to provide a thorough understanding of the embodiments. However, it is understood that the embodiments may be practiced without these specific details. For example, circuits may be shown in block diagrams in order not to obscure the embodiments in unnecessary detail. In other instances, well-known circuits, processes, algorithms, structures, and techniques may be shown without unnecessary detail in order to avoid obscuring the embodiments.

(57) Implementation of the techniques, blocks, steps and means described above may be done in various ways. For example, these techniques, blocks, steps and means may be implemented in hardware, software, or a combination thereof. For a hardware implementation, the processing units may be implemented within one or more application specific integrated circuits (ASICs), digital signal processors (DSPs), digital signal processing devices (DSPDs), programmable logic devices (PLDs), field programmable gate arrays (FPGAs), processors, controllers, micro-controllers, microprocessors, other electronic units designed to perform the functions described above, and/or a combination thereof.

(58) Also, it is noted that the embodiments may be described as a process which is depicted as a flowchart, a flow diagram, a swim diagram, a data flow diagram, a structure diagram, or a block diagram. Although a depiction may describe the operations as a sequential process, many of the operations can be performed in parallel or concurrently. In addition, the order of the operations may be re-arranged. A process is terminated when its operations are completed, but could have additional steps not included in the figure. A process may correspond to a method, a function, a procedure, a subroutine, a subprogram, etc. When a process corresponds to a function, its termination corresponds to a return of the function to the calling function or the main function.

(59) Furthermore, embodiments may be implemented by hardware, software, scripting languages, firmware, middleware, microcode, hardware description languages, and/or any combination thereof. When implemented in software, firmware, middleware, scripting language, and/or microcode, the program code or code segments to perform the necessary tasks may be stored in a machine readable medium such as a storage medium. A code segment or machine-executable instruction may represent a procedure, a function, a subprogram, a program, a routine, a subroutine, a module, a software package, a script, a class, or any combination of instructions, data structures, and/or program statements. A code segment may be coupled to another code segment or a hardware circuit by passing and/or receiving information, data, arguments, parameters, and/or memory contents. Information, arguments, parameters, data, etc. may be passed, forwarded, or transmitted

via any suitable means including memory sharing, message passing, token passing, network transmission, etc.

(60) For a firmware and/or software implementation, the methodologies may be implemented with modules (e.g., procedures, functions, and so on) that perform the functions described herein. Any machine-readable medium tangibly embodying instructions may be used in implementing the methodologies described herein. For example, software codes may be stored in a memory. Memory may be implemented within the processor or external to the processor. As used herein the term “memory” refers to any type of long term, short term, volatile, nonvolatile, or other storage medium and is not to be limited to any particular type of memory or number of memories, or type of media upon which memory is stored.

(61) Moreover, as disclosed herein, the term “storage medium” may represent one or more memories for storing data, including read only memory (ROM), random access memory (RAM), magnetic RAM, core memory, magnetic disk storage mediums, optical storage mediums, flash memory devices and/or other machine readable mediums for storing information. The term “machine-readable medium” includes, but is not limited to portable or fixed storage devices, optical storage devices, and/or various other storage mediums capable of storing that contain or carry instruction(s) and/or data.

(62) While the principles of the disclosure have been described above in connection with specific apparatuses and methods, it is to be clearly understood that this description is made only by way of example and not as limitation on the scope of the disclosure.

Claims

1. A system for providing policy-controlled communication over the Internet between a plurality of different cloud computing environments, detecting violations of policies and responding to the violations, the system comprising: a user interface that receives configuration settings to be applied to a plurality of first instances within a first cloud computing environment and a plurality of second instances within a second cloud computing environment, wherein: the first cloud computing environment comprises one or more first processors and one or more first memories, and the second cloud computing environment comprises one or more second processors and one or more memories; a plurality of collectors that retrieve information from the first cloud computing environment and the second cloud computing environment, wherein the information comprises a plurality of functionalities of the first cloud computing environment and the second cloud computing environment; a controller that determines policies for the plurality of first instances within the first cloud computing environment and the plurality of second instances within the second cloud computing environment as functions of the configuration settings and the information; a configurator that applies the policies to the plurality of first instances within the first cloud computing environment and the plurality of second instances within the second cloud computing environment; a first tester that inspects operations of the plurality of first instances within the first cloud computing environment and detects violations of the policies by the plurality of first instances within the first cloud computing environment, wherein the first tester inspects the operations based on a testing schedule which indicates a frequency of testing for the first instance for compliance after each specific policies was applied to the first instance and a table that stores a list of a number of times at which the first instance was tested for the compliance with the specific policies; and an enforcer that responds to the detected violations by receiving a notification from the first tester that a first instance from the plurality of first instances violated a first policy, wherein the controller instructs the configurator to apply the first policy to the first instance again, shut down the first instance or cut off communications with the first instance.

2. The system of claim 1, wherein the plurality of functionalities of the first cloud computing environment and the second cloud computing environment comprises at least one of network

configurations, firewall rules, cloud application programming interfaces (APIs), resources, cloud service providers, or data sets.

3. The system of claim 1, wherein the information further comprises at least one of a data input type, a data type, a data size, or a data age, and the plurality of collectors are further configured to normalize the information to have a common format.
4. The system of claim 1, wherein the policies comprise at least one of firewall rules, forwarding rules, network configurations, cross-cloud routing rules, IP addressing rules, cross-cloud peering rules, security group management rules, storage bucket access rules, resource management rules, or subnet configurations.
5. A method for providing policy-controlled communication over the Internet between a plurality of different cloud computing environments, detecting violations of policies and responding to the violations, the method comprising: receiving configuration settings to be applied to a plurality of first instances within a first cloud computing environment and a plurality of second instances within a second cloud computing environment; retrieving information from the first cloud computing environment and the second cloud computing environment, wherein the information comprises a plurality of functionalities of the first cloud computing environment and the second cloud computing environment; determining policies for the plurality of first instances within the first cloud computing environment and the plurality of second instances within the second cloud computing environment as functions of the configuration settings and the information; applying the policies to the plurality of first instances within the first cloud computing environment and the plurality of second instances within the second cloud computing environment; inspecting operations of the plurality of first instances within the first cloud computing environment and the plurality of second instances within the second cloud computing environment and detecting violations of the policies by the plurality of first instances within the first cloud computing environment and the plurality of second instances within the second cloud computing environment; and responding to the detected violations by receiving a notification that a first instance from the plurality of first instances violated a first policy, wherein a controller instructs a configurator to apply the first policy to the first instance again, shut down the first instance or cut off communications with the first instance.
6. The method of claim 5, wherein determining the policies comprises determining general policies as functions of the configuration settings and translating the general policies to specific policies for the first cloud computing environment and the second cloud computing environment by using the information.
7. The method of claim 6, wherein applying the policies comprises retrieving scripts corresponding to the specific policies and applying the specific policies to the plurality of first instances within the first cloud computing environment and the plurality of second instances within the second cloud computing environment by executing the scripts.
8. The method of claim 5, further comprising sending the notification of a detected violation of a policy by an instance of the plurality of first instances within the first cloud computing environment to a user interface.
9. The method of claim 5, wherein a first tester may inspect the first instance to determine whether it complies with various policies, such as whether a load balancer was set up, whether the load balancer is running correctly, whether the load balancer has enough processing power, and whether the load balancer has a correct IP address.
10. A system for providing policy-controlled communication over the Internet between a plurality of different cloud computing environments, detecting violations of policies and responding to the violations, the system comprising one or more processors and one or more memories with code for: a user interface that is configured to receive configuration settings to be applied to a plurality of first instances within a first cloud computing environment and a plurality of second instances within a second cloud computing environment, wherein: the first cloud computing environment

comprises one or more first processors and one or more first memories, and the second cloud computing environment comprises one or more second processors and one or more memories; a plurality of collectors that are configured to retrieve information from the first cloud computing environment and the second cloud computing environment, wherein the information comprises a plurality of functionalities of the first cloud computing environment and the second cloud computing environment; a controller that is configured to determine policies for the plurality of first instances within the first cloud computing environment and the plurality of second instances within the second cloud computing environment as functions of the configuration settings and the information; a configurator that is configured to apply the policies to the plurality of first instances within the first cloud computing environment and the plurality of second instances within the second cloud computing environment; a first tester that is configured to inspect operations of the plurality of first instances within the first cloud computing environment, wherein the first tester inspects the operations based on a testing schedule which indicates a frequency of testing for the first instance for compliance after each specific policies was applied to the first instance and a table that stores a list of a number of times at which the first instance was tested for the compliance with the specific policies, and a second tester that inspect operations of the plurality of second instances within the second cloud computing environment to detect violations of the policies by the plurality of first instances within the first cloud computing environment and violations of the policies by the plurality of second instances within the second cloud computing environment, respectively; and an enforcer that is configured to respond to the detected violations by receiving a notification from the first tester that a first instance from the plurality of first instances violated a first policy, wherein the controller instructs the configurator to apply the first policy to the first instance again, shut down the first instance or cut off communications with the first instance.

11. The system of claim 10, wherein the controller is configured to determine the policies by determining general policies as functions of the configuration settings and translating the general policies to specific policies for the first cloud computing environment and the second cloud computing environment by using the information.

12. The system of claim 11, wherein the controller is configured to translate the general policies to the specific policies at a service layer and a software-defined data center layer.

13. The system of claim 11, wherein the configurator is configured to retrieve scripts corresponding to the specific policies and to apply the specific policies to the plurality of first instances within the first cloud computing environment and the plurality of second instances within the second cloud computing environment by executing the scripts.

14. The system of claim 10, further comprising a reporter is configured to send the notification via a short message service (SMS), an email, an Application Programming Interface (API) call, or another notification method.

15. The system of claim 10, further comprising a metadata service endpoint that is configured to receive the policies from the configurator and to host changes to the plurality of first instances and the plurality of second instances, wherein the plurality of first instances within the first cloud computing environment and the plurality of second instances within the second cloud computing environment are configured to retrieve the changes from the metadata service endpoint and to apply the changes.

16. The system of claim 10, wherein the enforcer is further configured to send the notification of the detected violation to the controller, which directs the configurator to apply the first policy that was violated to a non-complying instance.

17. The system of claim 10, wherein the notification identifies the first instance that violated the first policy, the first cloud computing environment in which the first instance is located, and the first policy that was violated.

18. The system of claim 17, wherein the controller is further configured to receive the notification of the detected violation and to require the first instance of the plurality of first instances within the

first cloud computing environment to comply with the first policy that was violated.

19. The system of claim 10, wherein the first tester is further configured to inspect existing IP addresses of the plurality of first instances and to send a notification to the enforcer to request new IP addresses from the first cloud computing environment upon identifying an overlap between the existing IP addresses.

20. The system of claim 10, wherein the first tester is further configured to inspect the operations of the plurality of first instances based on a testing schedule which indicates a frequency of testing for the plurality of first instances for compliance with the policies.
