

# US Patent & Trademark Office

## Patent Public Search | Text View

---

United States Patent	12388863
Kind Code	B2
Date of Patent	August 12, 2025
Inventor(s)	Murphy; Brian P. et al.

---

### Threat mitigation system and method

---

#### Abstract

A computer-implemented method, computer program product and computing system for establishing connectivity with a plurality of security-relevant subsystems within a computing platform; receiving an initial notification of a security event from one of the security-relevant subsystems, wherein the initial notification includes a computer-readable language portion that defines one or more specifics of the security event; processing the initial notification using a generative AI model and a formatting script to produce a summarized human-readable report for the initial notification, wherein the summarized human-readable report defines one or more recommended next steps; and automatically executing some or all of the recommended next steps to define one or more recommended actions.

---

**Inventors:** Murphy; Brian P. (Tampa, FL), Partlow; Joe (Tampa, FL), O'Connor; Colin (Tampa, FL), Pfeiffer; Jason (Tampa, FL), Murphy; Brian Philip (St. Petersburg, FL), Echavarria; Jonathan R. (Tampa, FL)

**Applicant:** ReliaQuest Holdings, LLC (Tampa, FL)

**Family ID:** 1000008748015

**Assignee:** RELIAQUEST HOLDINGS, LLC (Tampa, FL)

**Appl. No.:** 18/585120

**Filed:** February 23, 2024

#### Prior Publication Data

Document Identifier	Publication Date
US 20240291850 A1	Aug. 29, 2024

#### Related U.S. Application Data

us-provisional-application US 63486617 20230223

---

#### Publication Classification

**Int. Cl.:** **H04L9/40** (20220101); **G06F16/34** (20250101); **G06F21/55** (20130101); **G06F21/56** (20130101); **G06F40/103** (20200101); **G06F40/154** (20200101); **G06F40/56** (20200101); **G06N3/0475** (20230101); **H04L41/16** (20220101)

**U.S. Cl.:**

**CPC** **H04L63/1441** (20130101); **G06F16/345** (20190101); **G06F21/552** (20130101); **G06F21/554** (20130101); **G06F21/566** (20130101); **G06F40/103** (20200101); **G06F40/154** (20200101); **G06F40/56** (20200101); **G06N3/0475** (20230101); **H04L41/16** (20130101); **H04L63/1416** (20130101); **H04L63/1425** (20130101); **G06F2221/034** (20130101)

## Field of Classification Search

**USPC:** None

---

## References Cited

### U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
7159237	12/2006	Schneier et al.	N/A	N/A
8745703	12/2013	Lambert et al.	N/A	N/A
9292496	12/2015	Carbonell et al.	N/A	N/A
9436820	12/2015	Gleichauf et al.	N/A	N/A
9641544	12/2016	Treat et al.	N/A	N/A
10397255	12/2018	Bhalotra et al.	N/A	N/A
10681060	12/2019	Scheidler et al.	N/A	N/A
10778714	12/2019	Shi	N/A	N/A
11023576	12/2020	Griffin et al.	N/A	N/A
11119630	12/2020	Marchetti et al.	N/A	N/A
11336689	12/2021	Miramirkhani et al.	N/A	N/A
11363036	12/2021	Levin et al.	N/A	N/A
11388040	12/2021	Mdini et al.	N/A	N/A
11416504	12/2021	Pelloin	N/A	N/A
11558408	12/2022	Eisenkot et al.	N/A	N/A
11743287	12/2022	Nagarajegowda	726/23	H04L 63/1458
11750634	12/2022	Chowdhary	726/25	H04L 63/1416
11886585	12/2023	Davis	N/A	N/A
11916767	12/2023	Wu et al.	N/A	N/A
12081568	12/2023	Koide et al.	N/A	N/A
12184697	12/2023	Crabtree et al.	N/A	N/A
2003/0051026	12/2002	Carter	709/224	H04L 63/0263
2005/0254654	12/2004	Rockwell et al.	N/A	N/A
2006/0236395	12/2005	Barker et al.	N/A	N/A
2009/0172773	12/2008	Moore	N/A	N/A
2012/0158615	12/2011	Spring	N/A	N/A
2014/0249760	12/2013	Proud et al.	N/A	N/A
2015/0067835	12/2014	Chari et al.	N/A	N/A
2017/0006058	12/2016	Murphy	N/A	H04L 63/1408
2017/0006061	12/2016	Murphy	N/A	H04L 63/1441
2017/0054745	12/2016	Zhang et al.	N/A	N/A
2017/0286455	12/2016	Li	N/A	G06F 16/212
2018/0020021	12/2017	Gilmore et al.	N/A	N/A
2018/0285479	12/2017	Mackay	N/A	G06F 16/90335
2018/0288070	12/2017	Price et al.	N/A	N/A

2018/0307833	12/2017	Noeth et al.	N/A	N/A
2019/0124109	12/2018	Foster et al.	N/A	N/A
2019/0164224	12/2018	Hunter et al.	N/A	N/A
2019/0205395	12/2018	Bonin et al.	N/A	N/A
2019/0208363	12/2018	Shapiro et al.	N/A	N/A
2019/0228297	12/2018	Shen et al.	N/A	N/A
2019/0258716	12/2018	Becker et al.	N/A	N/A
2019/0260764	12/2018	Humphrey et al.	N/A	N/A
2019/0260769	12/2018	Sharon et al.	N/A	N/A
2019/0260779	12/2018	Bazalgette	N/A	G06N 20/10
2019/0312889	12/2018	Philip et al.	N/A	N/A
2019/0377876	12/2018	Murphy	N/A	H04L 63/1416
2019/0379705	12/2018	Murphy et al.	N/A	N/A
2019/0394226	12/2018	Murphy	N/A	H04L 67/563
2020/0023846	12/2019	Husain et al.	N/A	N/A
2020/0067985	12/2019	Bhargava	N/A	G06F 3/0481
2020/0125725	12/2019	Petersen et al.	N/A	N/A
2020/0134103	12/2019	Mankovskii	N/A	N/A
2020/0285977	12/2019	Brebner	N/A	N/A
2020/0329054	12/2019	Bjarnason	N/A	H04L 63/205
2020/0344249	12/2019	Mohan et al.	N/A	N/A
2020/0358792	12/2019	Bazalgette	N/A	G06N 20/00
2020/0396231	12/2019	Krebs et al.	N/A	N/A
2020/0410320	12/2019	Almazn et al.	N/A	N/A
2021/0014256	12/2020	Malhotra	N/A	N/A
2021/0026954	12/2020	Murphy et al.	N/A	N/A
2021/0029159	12/2020	Murphy et al.	N/A	N/A
2021/0035116	12/2020	Berrington	N/A	G06N 5/04
2021/0037043	12/2020	Lee et al.	N/A	N/A
2021/0042662	12/2020	Pu et al.	N/A	N/A
2021/0073389	12/2020	Murphy	N/A	H04L 67/60
2021/0073390	12/2020	Murphy	N/A	G06F 11/3409
2021/0075818	12/2020	Murphy	N/A	H04L 67/34
2021/0075819	12/2020	Murphy	N/A	G06N 5/04
2021/0075820	12/2020	Murphy	N/A	H04L 67/30
2021/0133670	12/2020	Cella	N/A	G06N 3/044
2021/0150411	12/2020	Coenders et al.	N/A	N/A
2021/0160258	12/2020	Murphy	N/A	G06F 16/2452
2021/0160274	12/2020	Murphy	N/A	H04L 63/20
2021/0168161	12/2020	Dunn et al.	N/A	N/A
2021/0168175	12/2020	Crabtree et al.	N/A	N/A
2021/0194924	12/2020	Heinemeyer et al.	N/A	N/A
2021/0273954	12/2020	Linton	N/A	H04L 63/1425
2021/0273960	12/2020	Humphrey	N/A	G06N 7/01
2021/0287800	12/2020	Ghosh et al.	N/A	N/A
2021/0319090	12/2020	Demir et al.	N/A	N/A
2021/0359980	12/2020	Crabtree	N/A	H04L 63/1433
2021/0366586	12/2020	Ryan et al.	N/A	N/A
2022/0036153	12/2021	O'Malia et al.	N/A	N/A
2022/0058273	12/2021	Rathore et al.	N/A	N/A
2022/0148397	12/2021	Schoeman	N/A	H04L 12/2823
2022/0156380	12/2021	Pradzynski et al.	N/A	N/A
2022/0187847	12/2021	Cella et al.	N/A	N/A
2022/0237368	12/2021	Tran	N/A	N/A
2022/0245641	12/2021	Wintle et al.	N/A	N/A

2022/0261478	12/2021	Yankovskiy et al.	N/A	N/A
2022/0329630	12/2021	Li	N/A	G06F 21/577
2022/0382611	12/2021	Kapish et al.	N/A	N/A
2022/0400131	12/2021	Shao et al.	N/A	N/A
2023/0009704	12/2022	Boyle et al.	N/A	N/A
2023/0095415	12/2022	Boyer et al.	N/A	N/A
2023/0124288	12/2022	Meier-Hellstern et al.	N/A	N/A
2023/0132501	12/2022	Phillippe et al.	N/A	N/A
2023/0164158	12/2022	Fellows	726/23	H04L 63/1441
2023/0164567	12/2022	Fellows	455/410	H04L 63/1416
2023/0171266	12/2022	Brunner	726/23	H04L 41/16
2023/0179628	12/2022	Porras	726/22	G06F 40/30
2023/0229937	12/2022	Takata et al.	N/A	N/A
2023/0245234	12/2022	Sumant	N/A	N/A
2023/0252147	12/2022	Gervais et al.	N/A	N/A
2023/0259860	12/2022	Sarkar	N/A	N/A
2023/0283629	12/2022	Boyer	726/22	H04L 63/1425
2023/0308474	12/2022	Thompson	N/A	N/A
2023/0315722	12/2022	Saxe et al.	N/A	N/A
2023/0319071	12/2022	Durbin	726/23	H04L 63/1466
2023/0368321	12/2022	Bringsjord	N/A	G06Q 50/26
2023/0396641	12/2022	Hebbagodi et al.	N/A	N/A
2023/0403296	12/2022	Pickman	N/A	H04L 63/1458
2023/0403587	12/2022	Lee et al.	N/A	N/A
2024/0045990	12/2023	Boyer	N/A	G06F 40/279
2024/0046318	12/2023	Muriqi	N/A	G06Q 20/389
2024/0061937	12/2023	Camp	N/A	G06F 40/237
2024/0070270	12/2023	Mace	N/A	G06F 21/554
2024/0112791	12/2023	Simmons et al.	N/A	N/A
2024/0152624	12/2023	Helyar	N/A	G06F 21/563
2024/0168864	12/2023	Hamlin et al.	N/A	N/A
2024/0171451	12/2023	Soni et al.	N/A	N/A
2024/0256582	12/2023	Jain et al.	N/A	N/A
2024/0273227	12/2023	Thompson	N/A	N/A

## OTHER PUBLICATIONS

International Search Report and Written Opinion issued in related Application Serial No. PCT/US2024/017020 on Jun. 21, 2024. cited by applicant

International Search Report and Written Opinion issued in related Application Serial No. PCT/US2024/017024 on Jun. 21, 2024. cited by applicant

International Search Report and Written Opinion issued in related Application Serial No. PCT/US2024/017029 on Jun. 21, 2024. cited by applicant

International Search Report and Written Opinion issued in related Application Serial No. PCT/US2024/017032 on Jun. 21, 2024. cited by applicant

International Search Report and Written Opinion issued in related Application Serial No. PCT/US2024/017037 on Jun. 21, 2024. cited by applicant

International Search Report and Written Opinion issued in related Application Serial No. PCT/US2024/017041 on Jun. 26, 2024. cited by applicant

International Search Report and Written Opinion issued in related Application Serial No. PCT/US2024/017044 on Jun. 26, 2024. cited by applicant

International Search Report and Written Opinion issued in related Application Serial No. PCT/US2024/017056 on Jun. 26, 2024. cited by applicant

International Search Report and Written Opinion issued in related Application Serial No. PCT/US2024/017063 on Jul. 24, 2024. cited by applicant

International Search Report and Written Opinion issued in related Application Serial No.

PCT/US2024/017066 on Jun. 4, 2024. cited by applicant  
International Search Report and Written Opinion issued in related Application Serial No. PCT/US2024/017070 on Jun. 26, 2024. cited by applicant  
Non Final Office Action issued in related U.S. Appl. No. 18/584,994 on May 31, 2024. cited by applicant  
Non Final Office Action issued in related U.S. Appl. No. 18/584,989 on Jun. 7, 2024. cited by applicant  
Non Final Office Action issued in related U.S. Appl. No. 18/585,000 on May 22, 2024. cited by applicant  
Non Final Office Action issued in related U.S. Appl. No. 18/585,128 on Jun. 5, 2024. cited by applicant  
Non Final Office Action issued in related U.S. Appl. No. 18/585,186 on Jun. 6, 2024. cited by applicant  
Non Final Office Action issued in related U.S. Appl. No. 18/585,428 on Jun. 17, 2024. cited by applicant  
Non Final Office Action issued in related U.S. Appl. No. 18/585,510 on May 30, 2024. cited by applicant  
Non-Final Office Action issued in related U.S. Appl. No. 18/584,996 on Jun. 13, 2024. cited by applicant  
Non Final Office Action issued in related U.S. Appl. No. 18/585,125 on May 22, 2024. cited by applicant  
Final Office Action issued in related U.S. Appl. No. 18/584,994 on Sep. 11, 2024. cited by applicant  
Final Office Action issued in related U.S. Appl. No. 18/585,125 on Sep. 17, 2024. cited by applicant  
Final Office Action issued in related U.S. Appl. No. 18/585,128 on Sep. 24, 2024. cited by applicant  
Final Office Action issued in related U.S. Appl. No. 18/585,186 on Sep. 28, 2024. cited by applicant  
Final Office Action issued in related U.S. Appl. No. 18/585,510 on Oct. 22, 2024. cited by applicant  
Final Office Action issued in related U.S. Appl. No. 18/585,687 on Oct. 31, 2024. cited by applicant  
Chu et al., "Analysis of network security and privacy security based on AI in IOT environment", Sep. 2021, IEEE 4th International Conference on Information Systems and Computer Aided Education, pp. 390-393 (Year: 2021). cited by applicant  
Final Office Action issued in related U.S. Appl. No. 18/584,989 on Nov. 27, 2024. cited by applicant  
Final Office Action issued in related U.S. Appl. No. 18/584,996 on Jan. 29, 2025. cited by applicant  
Final Office Action issued in related U.S. Appl. No. 18/585,000 on Dec. 30, 2024. cited by applicant  
Final Office Action issued in related U.S. Appl. No. 18/585,428 on Dec. 13, 2024. cited by applicant  
Non-Final Office Action issued in related U.S. Appl. No. 18/585,128 on Jan. 15, 2025. cited by applicant  
Notice of Allowance issued in related U.S. Appl. No. 18/585,125 on Jan. 13, 2025. cited by applicant  
Notice of Allowance issued in related U.S. Appl. No. 18/585,186 on Feb. 12, 2025. cited by applicant  
Notice of Allowance issued in related U.S. Appl. No. 18/585,186 on Jan. 16, 2025. cited by applicant  
Non-Final Office Action issued in related U.S. Appl. No. 18/584,994 on Mar. 3, 2025. cited by applicant  
Non-Final Office Action issued in related U.S. Appl. No. 18/585,510 on Feb. 21, 2025. cited by applicant  
Notice of Allowance issued in related U.S. Appl. No. 18/584,989 on Feb. 24, 2025. cited by applicant  
Notice of Allowance issued in related U.S. Appl. No. 18/584,989 on Mar. 19, 2025. cited by applicant  
Notice of Allowance issued in related U.S. Appl. No. 18/585,000 on Apr. 21, 2025. cited by applicant  
Notice of Allowance issued in related U.S. Appl. No. 18/585,125 on Apr. 24, 2025. cited by applicant  
Notice of Allowance issued in related U.S. Appl. No. 18/585,128 on May 8, 2025. cited by applicant  
Notice of Allowance issued in related U.S. Appl. No. 18/585,186 on Apr. 14, 2025 cited by applicant  
Notice of Allowance issued in related U.S. Appl. No.18/585,186 on Apr. 24, 2025. cited by applicant  
Notice of Allowance issued in related U.S. Appl. No. 18/585,428 on Apr. 16, 2025. cited by applicant  
Notice of Allowance issued in related U.S. Appl. No. 18/585,687 on Mar. 20, 2025. cited by applicant  
Non Final Office Action issued in related U.S. Appl. No. 18/584,996 on May 21, 2025. cited by applicant  
Final Office Action issued in related U.S. Appl. No. 18/584,994 on Jun. 30, 2025. cited by applicant  
Notice of Allowance issued in related U.S. Appl. No. 18/584,989 on Jun. 3, 2025. cited by applicant  
Notice of Allowance issued in related U.S. Appl. No. 18/584,989 on May 23, 2025. cited by applicant  
Notice of Allowance issued in related U.S. Appl. No. 18/585,125 on May 30, 2025. cited by applicant  
Notice of Allowance issued in related U.S. Appl. No. 18/585,186 on Jun. 11, 2025. cited by applicant

---

*Primary Examiner:* Zand; Kambiz

*Assistant Examiner:* Ahmed; Mahabub S

*Attorney, Agent or Firm:* Holland & Knight LLP

---

## Background/Summary

RELATED APPLICATION(S) (1) This application claims the benefit of U.S. Provisional Patent Application No. 63/486,617, filed on 23 Feb. 2023, the entire contents of which are herein incorporated by reference.

### TECHNICAL FIELD

(1) This disclosure relates to threat mitigation systems and, more particularly, to threat mitigation systems that utilize a universal query language.

### BACKGROUND

(2) In the computer world, there is a constant battle occurring between bad actors that want to attack computing platforms and good actors who try to prevent the same. Unfortunately, the complexity of such computer attacks is constantly increasing, so technology needs to be employed that understands the complexity of these attacks and is capable of addressing the same.

(3) Threat mitigation systems may utilize and/or communicate with a plurality of security-relevant subsystems, wherein these security-relevant subsystems may gather information concerning such computer attacks. Unfortunately and in order to obtain such gathered information from these security-relevant subsystems, the user of the threat mitigation system would often be required to formulate a unique query for each security-relevant subsystem.

### SUMMARY OF DISCLOSURE

(4) Auto-Execution of Recommended Next Steps

(5) In one implementation, a computer-implemented method is executed on a computing device and includes: establishing connectivity with a plurality of security-relevant subsystems within a computing platform; receiving an initial notification of a security event from one of the security-relevant subsystems, wherein the initial notification includes a computer-readable language portion that defines one or more specifics of the security event; processing the initial notification using a generative AI model and a formatting script to produce a summarized human-readable report for the initial notification, wherein the summarized human-readable report defines one or more recommended next steps; and automatically executing some or all of the recommended next steps to define one or more recommended actions.

(6) One or more of the following features may be included. Receiving an initial notification of a security event from one of the security-relevant subsystems may include: receiving the initial notification of the security event from an agent executed on one of the security-relevant subsystems. Processing the initial notification using a generative AI model and a formatting script to produce a summarized human-readable report for the initial notification may include: processing the initial notification using the generative AI model, the formatting script and/or one or more tools to produce the summarized human-readable report for the initial notification. The one or more tools may include one or more of: a decoding tool to decode an encoded initial notification; a decompression tool to decompress a compressed initial notification; and an identification tool to identify an owner of a domain associated with the initial notification. Processing the initial notification using a generative AI model and a formatting script to produce a summarized human-readable report for the initial notification may include: processing the initial notification using a large language model. Processing the initial notification using a generative AI model and a formatting script to produce a summarized human-readable report for the initial notification may include: utilizing prompt engineering to produce the summarized human-readable report for the initial notification. Processing the initial notification using a generative AI model and a formatting script to produce a summarized human-readable report for the initial notification may include: utilizing several loops and/or nested loops to produce the summarized human-readable report for the initial notification. Automatically executing some or all of the recommended next steps to define one or more recommended actions may include: automatically performing one or more investigative operations concerning the security event.

(7) In another implementation, a computer program product resides on a computer readable medium and has a plurality of instructions stored on it. When executed by a processor, the instructions cause the processor to perform operations including establishing connectivity with a plurality of security-relevant subsystems within a computing platform; receiving an initial notification of a security event from one of the security-relevant subsystems, wherein the initial notification includes a computer-readable language portion that defines one or more specifics of the security event; processing the initial notification using a generative AI model and a formatting script to produce a summarized human-readable report for the initial notification,

wherein the summarized human-readable report defines one or more recommended next steps; and automatically executing some or all of the recommended next steps to define one or more recommended actions.

(8) One or more of the following features may be included. Receiving an initial notification of a security event from one of the security-relevant subsystems may include: receiving the initial notification of the security event from an agent executed on one of the security-relevant subsystems. Processing the initial notification using a generative AI model and a formatting script to produce a summarized human-readable report for the initial notification may include: processing the initial notification using the generative AI model, the formatting script and/or one or more tools to produce the summarized human-readable report for the initial notification. The one or more tools may include one or more of: a decoding tool to decode an encoded initial notification; a decompression tool to decompress a compressed initial notification; and an identification tool to identify an owner of a domain associated with the initial notification. Processing the initial notification using a generative AI model and a formatting script to produce a summarized human-readable report for the initial notification may include: processing the initial notification using a large language model. Processing the initial notification using a generative AI model and a formatting script to produce a summarized human-readable report for the initial notification may include: utilizing prompt engineering to produce the summarized human-readable report for the initial notification. Processing the initial notification using a generative AI model and a formatting script to produce a summarized human-readable report for the initial notification may include: utilizing several loops and/or nested loops to produce the summarized human-readable report for the initial notification. Automatically executing some or all of the recommended next steps to define one or more recommended actions may include: automatically performing one or more investigative operations concerning the security event.

(9) In another implementation, a computing system includes a processor and a memory system configured to perform operations including establishing connectivity with a plurality of security-relevant subsystems within a computing platform; receiving an initial notification of a security event from one of the security-relevant subsystems, wherein the initial notification includes a computer-readable language portion that defines one or more specifics of the security event; processing the initial notification using a generative AI model and a formatting script to produce a summarized human-readable report for the initial notification, wherein the summarized human-readable report defines one or more recommended next steps; and automatically executing some or all of the recommended next steps to define one or more recommended actions.

(10) One or more of the following features may be included. Receiving an initial notification of a security event from one of the security-relevant subsystems may include: receiving the initial notification of the security event from an agent executed on one of the security-relevant subsystems. Processing the initial notification using a generative AI model and a formatting script to produce a summarized human-readable report for the initial notification may include: processing the initial notification using the generative AI model, the formatting script and/or one or more tools to produce the summarized human-readable report for the initial notification. The one or more tools may include one or more of: a decoding tool to decode an encoded initial notification; a decompression tool to decompress a compressed initial notification; and an identification tool to identify an owner of a domain associated with the initial notification. Processing the initial notification using a generative AI model and a formatting script to produce a summarized human-readable report for the initial notification may include: processing the initial notification using a large language model. Processing the initial notification using a generative AI model and a formatting script to produce a summarized human-readable report for the initial notification may include: utilizing prompt engineering to produce the summarized human-readable report for the initial notification. Processing the initial notification using a generative AI model and a formatting script to produce a summarized human-readable report for the initial notification may include: utilizing several loops and/or nested loops to produce the summarized human-readable report for the initial notification. Automatically executing some or all of the recommended next steps to define one or more recommended actions may include: automatically performing one or more investigative operations concerning the security event.

(11) The details of one or more implementations are set forth in the accompanying drawings and the description below. Other features and advantages will become apparent from the description, the drawings, and the claims,

---

## Description

### BRIEF DESCRIPTION OF THE DRAWINGS

- (1) FIG. 1 is a diagrammatic view of a distributed computing network including a computing device that executes a threat mitigation process according to an embodiment of the present disclosure;
- (2) FIG. 2 is a diagrammatic view of an exemplary probabilistic model rendered by a probabilistic process of the threat mitigation process of FIG. 1 according to an embodiment of the present disclosure;
- (3) FIG. 3 is a diagrammatic view of the computing platform of FIG. 1 according to an embodiment of the present disclosure;
- (4) FIG. 4 is a flowchart of an implementation of the threat mitigation process of FIG. 1 according to an embodiment of the present disclosure;
- (5) FIGS. 5-6 are diagrammatic views of screens rendered by the threat mitigation process of FIG. 1 according to an embodiment of the present disclosure;
- (6) FIGS. 7-9 are flowcharts of other implementations of the threat mitigation process of FIG. 1 according to an embodiment of the present disclosure;
- (7) FIG. 10 is a diagrammatic view of a screen rendered by the threat mitigation process of FIG. 1 according to an embodiment of the present disclosure;
- (8) FIG. 11 is a flowchart of another implementation of the threat mitigation process of FIG. 1 according to an embodiment of the present disclosure;
- (9) FIG. 12 is a diagrammatic view of a screen rendered by the threat mitigation process of FIG. 1 according to an embodiment of the present disclosure;
- (10) FIG. 13 is a flowchart of another implementation of the threat mitigation process of FIG. 1 according to an embodiment of the present disclosure;
- (11) FIG. 14 is a diagrammatic view of a screen rendered by the threat mitigation process of FIG. 1 according to an embodiment of the present disclosure;
- (12) FIG. 15 is a flowchart of another implementation of the threat mitigation process of FIG. 1 according to an embodiment of the present disclosure;
- (13) FIG. 16 is a diagrammatic view of screens rendered by the threat mitigation process of FIG. 1 according to an embodiment of the present disclosure;
- (14) FIGS. 17-23 are flowcharts of other implementations of the threat mitigation process of FIG. 1 according to an embodiment of the present disclosure;
- (15) FIG. 24 is a diagrammatic view of a screen rendered by the threat mitigation process of FIG. 1 according to an embodiment of the present disclosure;
- (16) FIGS. 25-31 are flowcharts of other implementations of the threat mitigation process of FIG. 1 according to an embodiment of the present disclosure;
- (17) FIG. 32 is a diagrammatic view of data field mapping according to an embodiment of the present disclosure;
- (18) FIG. 33 is a flowchart of another implementation of the threat mitigation process of FIG. 1 according to an embodiment of the present disclosure;
- (19) FIG. 34 is a flowchart of another implementation of the threat mitigation process of FIG. 1 according to an embodiment of the present disclosure;
- (20) FIG. 35 is a flowchart of another implementation of the threat mitigation process of FIG. 1 according to an embodiment of the present disclosure;
- (21) FIG. 36 is a flowchart of another implementation of the threat mitigation process of FIG. 1 according to an embodiment of the present disclosure;
- (22) FIG. 37 is a flowchart of another implementation of the threat mitigation process of FIG. 1 according to an embodiment of the present disclosure;
- (23) FIG. 38 is a flowchart of another implementation of the threat mitigation process of FIG. 1 according to an embodiment of the present disclosure;
- (24) FIG. 39 is a flowchart of another implementation of the threat mitigation process of FIG. 1 according to an embodiment of the present disclosure;
- (25) FIG. 40 is a flowchart of another implementation of the threat mitigation process of FIG. 1 according to an embodiment of the present disclosure;
- (26) FIG. 41 is a flowchart of another implementation of the threat mitigation process of FIG. 1 according to



an embodiment of the present disclosure;

(27) FIG. 42 is a flowchart of another implementation of the threat mitigation process of FIG. 1 according to an embodiment of the present disclosure;

(28) FIG. 43 is a flowchart of another implementation of the threat mitigation process of FIG. 1 according to an embodiment of the present disclosure; and

(29) FIG. 44 is a diagrammatic view of a threat mitigation platform for effectuating the threat mitigation process of FIG. 1 according to an embodiment of the present disclosure.

(30) Like reference symbols in the various drawings indicate like elements.

## DETAILED DESCRIPTION OF THE PREFERRED EMBODIMENTS

### (31) System Overview

(32) Referring to FIG. 1, there is shown threat mitigation process 10. Threat mitigation process 10 may be implemented as a server-side process, a client-side process, or a hybrid server-side/client-side process. For example, threat mitigation process 10 may be implemented as a purely server-side process via threat mitigation process 10s. Alternatively, threat mitigation process 10 may be implemented as a purely client-side process via one or more of threat mitigation process 10c1, threat mitigation process 10c2, threat mitigation process 10c3, and threat mitigation process 10c4. Alternatively still, threat mitigation process 10 may be implemented as a hybrid server-side/client-side process via threat mitigation process 10s in combination with one or more of threat mitigation process 10c1, threat mitigation process 10c2, threat mitigation process 10c3, and threat mitigation process 10c4. Accordingly, threat mitigation process 10 as used in this disclosure may include any combination of threat mitigation process 10s, threat mitigation process 10c1, threat mitigation process 10c2, threat mitigation process, and threat mitigation process 10c4.

(33) Threat mitigation process 10s may be a server application and may reside on and may be executed by computing device 12, which may be connected to network 14 (e.g., the Internet or a local area network). Examples of computing device 12 may include, but are not limited to: a personal computer, a laptop computer, a personal digital assistant, a data-enabled cellular telephone, a notebook computer, a television with one or more processors embedded therein or coupled thereto, a cable/satellite receiver with one or more processors embedded therein or coupled thereto, a server computer, a series of server computers, a mini computer, a mainframe computer, or a cloud-based computing network.

(34) The instruction sets and subroutines of threat mitigation process 10s, which may be stored on storage device 16 coupled to computing device 12, may be executed by one or more processors (not shown) and one or more memory architectures (not shown) included within computing device 12. Examples of storage device 16 may include but are not limited to: a hard disk drive; a RAID device; a random-access memory (RAM); a read-only memory (ROM); and all forms of flash memory storage devices.

(35) Network 14 may be connected to one or more secondary networks (e.g., network 18), examples of which may include but are not limited to: a local area network; a wide area network; or an intranet, for example.

(36) Examples of threat mitigation processes 10c1, 10c2, 10c3, 10c4 may include but are not limited to a client application, a web browser, a game console user interface, or a specialized application (e.g., an application running on e.g., the Android™ platform or the iOS™ platform). The instruction sets and subroutines of threat mitigation processes 10c1, 10c2, 10c3, 10c4, which may be stored on storage devices 20, 22, 24, 26 (respectively) coupled to client electronic devices 28, 30, 32, 34 (respectively), may be executed by one or more processors (not shown) and one or more memory architectures (not shown) incorporated into client electronic devices 28, 30, 32, 34 (respectively). Examples of storage device 16 may include but are not limited to: a hard disk drive; a RAID device; a random-access memory (RAM); a read-only memory (ROM); and all forms of flash memory storage devices.

(37) Examples of client electronic devices 28, 30, 32, 34 may include, but are not limited to, data-enabled, cellular telephone 28, laptop computer 30, personal digital assistant 32, personal computer 34, a notebook computer (not shown), a server computer (not shown), a gaming console (not shown), a smart television (not shown), and a dedicated network device (not shown). Client electronic devices 28, 30, 32, 34 may each execute an operating system, examples of which may include but are not limited to Microsoft Windows™, Android™, WebOS™, iOS™, Redhat Linux™, or a custom operating system.

(38) Users 36, 38, 40, 42 may access threat mitigation process 10 directly through network 14 or through secondary network 18. Further, threat mitigation process 10 may be connected to network 14 through secondary network 18, as illustrated with link line 44.

(39) The various client electronic devices (e.g., client electronic devices 28, 30, 32, 34) may be directly or

indirectly coupled to network **14** (or network **18**). For example, data-enabled, cellular telephone **28** and laptop computer **30** are shown wirelessly coupled to network **14** via wireless communication channels **46**, **48** (respectively) established between data-enabled, cellular telephone **28**, laptop computer **30** (respectively) and cellular network/bridge **50**, which is shown directly coupled to network **14**. Further, personal digital assistant **32** is shown wirelessly coupled to network **14** via wireless communication channel **52** established between personal digital assistant **32** and wireless access point (i.e., WAP) **54**, which is shown directly coupled to network **14**. Additionally, personal computer **34** is shown directly coupled to network **18** via a hardwired network connection.

(40) WAP **54** may be, for example, an IEEE 802.11a, 802.11b, 802.11g, 802.11n, Wi-Fi, and/or Bluetooth device that is capable of establishing wireless communication channel **52** between personal digital assistant **32** and WAP **54**. As is known in the art, IEEE 802.11x specifications may use Ethernet protocol and carrier sense multiple access with collision avoidance (i.e., CSMA/CA) for path sharing. The various 802.11x specifications may use phase-shift keying (i.e., PSK) modulation or complementary code keying (i.e., CCK) modulation, for example. As is known in the art, Bluetooth is a telecommunications industry specification that allows e.g., mobile phones, computers, and personal digital assistants to be interconnected using a short-range wireless connection.

(41) Artificial Intelligence/Machines Learning Overview:

(42) Assume for illustrative purposes that threat mitigation process **10** includes AI/ML process **56** (e.g., an artificial intelligence/machine learning process) that is configured to process information (e.g., information **58**). As will be discussed below in greater detail, examples of information **58** may include but are not limited to platform information (e.g., structured or unstructured content) being scanned to detect security events (e.g., access auditing; anomalies; authentication; denial of services; exploitation; malware; phishing; spamming; reconnaissance; and/or web attack) within a monitored computing platform (e.g., computing platform **60**).

(43) As is known in the art, structured content may be content that is separated into independent portions (e.g., fields, columns, features) and, therefore, may have a pre-defined data model and/or is organized in a pre-defined manner. For example, if the structured content concerns an employee list: a first field, column or feature may define the first name of the employee; a second field, column or feature may define the last name of the employee; a third field, column or feature may define the home address of the employee; and a fourth field, column or feature may define the hire date of the employee.

(44) Further and as is known in the art, unstructured content may be content that is not separated into independent portions (e.g., fields, columns, features) and, therefore, may not have a pre-defined data model and/or is not organized in a pre-defined manner. For example, if the unstructured content concerns the same employee list: the first name of the employee, the last name of the employee, the home address of the employee, and the hire date of the employee may all be combined into one field, column or feature.

(45) For the following illustrative example, assume that information **58** is unstructured content, an example of which may include but is not limited to unstructured user feedback received by a company (e.g., text-based feedback such as text-messages, social media posts, and email messages; and transcribed voice-based feedback such as transcribed voice mail, and transcribed voice messages).

(46) When processing information **58**, AI/ML process **56** may use probabilistic modeling to accomplish such processing, wherein examples of such probabilistic modeling may include but are not limited to discriminative modeling, generative modeling, or combinations thereof.

(47) As is known in the art, probabilistic modeling may be used within modern artificial intelligence systems (e.g., AI/ML process **56**), in that these probabilistic models may provide artificial intelligence systems with the tools required to autonomously analyze vast quantities of data (e.g., information **58**).

(48) Examples of the tasks for which probabilistic modeling may be utilized may include but are not limited to: predicting media (music, movies, books) that a user may like or enjoy based upon media that the user has liked or enjoyed in the past; transcribing words spoken by a user into editable text; grouping genes into gene clusters; identifying recurring patterns within vast data sets; filtering email that is believed to be spam from a user's inbox; generating clean (i.e., non-noisy) data from a noisy data set; analyzing (voice-based or text-based) customer feedback; and diagnosing various medical conditions and diseases.

(49) For each of the above-described applications of probabilistic modeling, an initial probabilistic model may be defined, wherein this initial probabilistic model may be subsequently (e.g., iteratively or continuously) modified and revised, thus allowing the probabilistic models and the artificial intelligence systems (e.g., AI/ML process **56**) to “learn” so that future probabilistic models may be more precise and may

explain more complex data sets.

(50) Accordingly, AI/ML process **56** may define an initial probabilistic model for accomplishing a defined task (e.g., the analyzing of information **58**). For the illustrative example, assume that this defined task is analyzing customer feedback (e.g., information **58**) that is received from customers of e.g., store **62** via an automated feedback phone line. For this example, assume that information **58** is initially voice-based content that is processed via e.g., a speech-to-text process that results in unstructured text-based customer feedback (e.g., information **58**).

(51) With respect to AI/ML process **56**, a probabilistic model may be utilized to go from initial observations about information **58** (e.g., as represented by the initial branches of a probabilistic model) to conclusions about information **58** (e.g., as represented by the leaves of a probabilistic model).

(52) As used in this disclosure, the term “branch” may refer to the existence (or non-existence) of a component (e.g., a sub-model) of (or included within) a model. Examples of such a branch may include but are not limited to: an execution branch of a probabilistic program or other generative model, a part (or parts) of a probabilistic graphical model, and/or a component neural network that may (or may not) have been previously trained.

(53) While the following discussion provides a detailed example of a probabilistic model, this is for illustrative purposes only and is not intended to be a limitation of this disclosure, as other configurations are possible and are considered to be within the scope of this disclosure. For example, the following discussion may concern any type of model (e.g., be it probabilistic or other) and, therefore, the below-described probabilistic model is merely intended to be one illustrative example of a type of model and is not intended to limit this disclosure to probabilistic models.

(54) Additionally, while the following discussion concerns word-based routing of messages through a probabilistic model, this is for illustrative purposes only and is not intended to be a limitation of this disclosure, as other configurations are possible and are considered to be within the scope of this disclosure. Examples of other types of information that may be used to route messages through a probabilistic model may include: the order of the words within a message; and the punctuation interspersed throughout the message.

(55) For example and referring also to FIG. 2, there is shown one simplified example of a probabilistic model (e.g., probabilistic model **100**) that may be utilized to analyze information **58** (e.g., unstructured text-based customer feedback) concerning store **62**. The manner in which probabilistic model **100** may be automatically-generated by AI/ML process **56** will be discussed below in detail. In this particular example, probabilistic model **100** may receive information **58** (e.g., unstructured text-based customer feedback) at branching node **102** for processing. Assume that probabilistic model **100** includes four branches off of branching node **102**, namely: service branch **104**; selection branch **106**; location branch **108**; and value branch **110** that respectively lead to service node **112**, selection node **114**, location node **116**, and value node **118**.

(56) As stated above, service branch **104** may lead to service node **112**, which may be configured to process the portion of information **58** (e.g., unstructured text-based customer feedback) that concerns (in whole or in part) feedback concerning the customer service of store **62**. For example, service node **112** may define service word list **120** that may include e.g., the word service, as well as synonyms of (and words related to) the word service (e.g., cashier, employee, greeter and manager). Accordingly and in the event that a portion of information **58** (e.g., a text-based customer feedback message) includes the word cashier, employee, greeter and/or manager, that portion of information **58** may be considered to be text-based customer feedback concerning the service received at store **62** and (therefore) may be routed to service node **112** of probabilistic model **100** for further processing. Assume for this illustrative example that probabilistic model **100** includes two branches off of service node **112**, namely: good service branch **122** and bad service branch **124**.

(57) Good service branch **122** may lead to good service node **126**, which may be configured to process the portion of information **58** (e.g., unstructured text-based customer feedback) that concerns (in whole or in part) good feedback concerning the customer service of store **62**. For example, good service node **126** may define good service word list **128** that may include e.g., the word good, as well as synonyms of (and words related to) the word good (e.g., courteous, friendly, lovely, happy, and smiling). Accordingly and in the event that a portion of information **58** (e.g., a text-based customer feedback message) that was routed to service node **112** includes the word good, courteous, friendly, lovely, happy, and/or smiling, that portion of information **58** may be considered to be text-based customer feedback indicative of good service received at store **62** (and, therefore, may be routed to good service node **126**).

(58) Bad service branch **124** may lead to bad service node **130**, which may be configured to process the portion of information **58** (e.g., unstructured text-based customer feedback) that concerns (in whole or in part) bad feedback concerning the customer service of store **62**. For example, bad service node **130** may define bad service word list **132** that may include e.g., the word bad, as well as synonyms of (and words related to) the word bad (e.g., rude, mean, jerk, miserable, and scowling). Accordingly and in the event that a portion of information **58** (e.g., a text-based customer feedback message) that was routed to service node **112** includes the word bad, rude, mean, jerk, miserable, and/or scowling, that portion of information **58** may be considered to be text-based customer feedback indicative of bad service received at store **62** (and, therefore, may be routed to bad service node **130**).

(59) As stated above, selection branch **106** may lead to selection node **114**, which may be configured to process the portion of information **58** (e.g., unstructured text-based customer feedback) that concerns (in whole or in part) feedback concerning the selection available at store **62**. For example, selection node **114** may define selection word list **134** that may include e.g., words indicative of the selection available at store **62**. Accordingly and in the event that a portion of information **58** (e.g., a text-based customer feedback message) includes any of the words defined within selection word list **134**, that portion of information **58** may be considered to be text-based customer feedback concerning the selection available at store **62** and (therefore) may be routed to selection node **114** of probabilistic model **100** for further processing. Assume for this illustrative example that probabilistic model **100** includes two branches off of selection node **114**, namely: good selection branch **136** and bad selection branch **138**.

(60) Good selection branch **136** may lead to good selection node **140**, which may be configured to process the portion of information **58** (e.g., unstructured text-based customer feedback) that concerns (in whole or in part) good feedback concerning the selection available at store **62**. For example, good selection node **140** may define good selection word list **142** that may include words indicative of a good selection at store **62**. Accordingly and in the event that a portion of information **58** (e.g., a text-based customer feedback message) that was routed to selection node **114** includes any of the words defined within good selection word list **142**, that portion of information **58** may be considered to be text-based customer feedback indicative of a good selection available at store **62** (and, therefore, may be routed to good selection node **140**).

(61) Bad selection branch **138** may lead to bad selection node **144**, which may be configured to process the portion of information **58** (e.g., unstructured text-based customer feedback) that concerns (in whole or in part) bad feedback concerning the selection available at store **62**. For example, bad selection node **144** may define bad selection word list **146** that may include words indicative of a bad selection at store **62**. Accordingly and in the event that a portion of information **58** (e.g., a text-based customer feedback message) that was routed to selection node **114** includes any of the words defined within bad selection word list **146**, that portion of information **58** may be considered to be text-based customer feedback indicative of a bad selection being available at store **62** (and, therefore, may be routed to bad selection node **144**).

(62) As stated above, location branch **108** may lead to location node **116**, which may be configured to process the portion of information **58** (e.g., unstructured text-based customer feedback) that concerns (in whole or in part) feedback concerning the location of store **62**. For example, location node **116** may define location word list **148** that may include e.g., words indicative of the location of store **62**. Accordingly and in the event that a portion of information **58** (e.g., a text-based customer feedback message) includes any of the words defined within location word list **148**, that portion of information **58** may be considered to be text-based customer feedback concerning the location of store **62** and (therefore) may be routed to location node **116** of probabilistic model **100** for further processing. Assume for this illustrative example that probabilistic model **100** includes two branches off of location node **116**, namely: good location branch **150** and bad location branch **152**.

(63) Good location branch **150** may lead to good location node **154**, which may be configured to process the portion of information **58** (e.g., unstructured text-based customer feedback) that concerns (in whole or in part) good feedback concerning the location of store **62**. For example, good location node **154** may define good location word list **156** that may include words indicative of store **62** being in a good location. Accordingly and in the event that a portion of information **58** (e.g., a text-based customer feedback message) that was routed to location node **116** includes any of the words defined within good location word list **156**, that portion of information **58** may be considered to be text-based customer feedback indicative of store **62** being in a good location (and, therefore, may be routed to good location node **154**).

(64) Bad location branch **152** may lead to bad location node **158**, which may be configured to process the portion of information **58** (e.g., unstructured text-based customer feedback) that concerns (in whole or in

part) bad feedback concerning the location of store **62**. For example, bad location node **158** may define bad location word list **160** that may include words indicative of store **62** being in a bad location. Accordingly and in the event that a portion of information **58** (e.g., a text-based customer feedback message) that was routed to location node **116** includes any of the words defined within bad location word list **160**, that portion of information **58** may be considered to be text-based customer feedback indicative of store **62** being in a bad location (and, therefore, may be routed to bad location node **158**).

(65) As stated above, value branch **110** may lead to value node **118**, which may be configured to process the portion of information **58** (e.g., unstructured text-based customer feedback) that concerns (in whole or in part) feedback concerning the value received at store **62**. For example, value node **118** may define value word list **162** that may include e.g., words indicative of the value received at store **62**. Accordingly and in the event that a portion of information **58** (e.g., a text-based customer feedback message) includes any of the words defined within value word list **162**, that portion of information **58** may be considered to be text-based customer feedback concerning the value received at store **62** and (therefore) may be routed to value node **118** of probabilistic model **100** for further processing. Assume for this illustrative example that probabilistic model **100** includes two branches off of value node **118**, namely: good value branch **164** and bad value branch **166**.

(66) Good value branch **164** may lead to good value node **168**, which may be configured to process the portion of information **58** (e.g., unstructured text-based customer feedback) that concerns (in whole or in part) good value being received at store **62**. For example, good value node **168** may define good value word list **170** that may include words indicative of receiving good value at store **62**. Accordingly and in the event that a portion of information **58** (e.g., a text-based customer feedback message) that was routed to value node **118** includes any of the words defined within good value word list **170**, that portion of information **58** may be considered to be text-based customer feedback indicative of good value being received at store **62** (and, therefore, may be routed to good value node **168**).

(67) Bad value branch **166** may lead to bad value node **172**, which may be configured to process the portion of information **58** (e.g., unstructured text-based customer feedback) that concerns (in whole or in part) bad value being received at store **62**. For example, bad value node **172** may define bad value word list **174** that may include words indicative of receiving bad value at store **62**. Accordingly and in the event that a portion of information **58** (e.g., a text-based customer feedback message) that was routed to value node **118** includes any of the words defined within bad value word list **174**, that portion of information **58** may be considered to be text-based customer feedback indicative of bad value being received at store **62** (and, therefore, may be routed to bad value node **172**).

(68) Once it is established that good or bad customer feedback was received concerning store **62** (i.e., with respect to the service, the selection, the location or the value), representatives and/or agents of store **62** may address the provider of such good or bad feedback via e.g., social media postings, text-messages and/or personal contact.

(69) Assume for illustrative purposes that user **36** uses data-enabled, cellular telephone **28** to provide feedback **64** (e.g., a portion of information **58**) to an automated feedback phone line concerning store **62**. Upon receiving feedback **64** for analysis, AI/ML process **56** may identify any pertinent content that is included within feedback **64**.

(70) For illustrative purposes, assume that user **36** was not happy with their experience at store **62** and that feedback **64** provided by user **36** was “my cashier was rude and the weather was rainy”. Accordingly and for this example, AI/ML process **56** may identify the pertinent content (included within feedback **64**) as the phrase “my cashier was rude” and may ignore/remove the irrelevant content “the weather was rainy”. As (in this example) feedback **64** includes the word “cashier”, AI/ML process **56** may route feedback **64** to service node **112** via service branch **104**. Further, as feedback **64** also includes the word “rude”, AI/ML process **56** may route feedback **64** to bad service node **130** via bad service branch **124** and may consider feedback **64** to be text-based customer feedback indicative of bad service being received at store **62**.

(71) For further illustrative purposes, assume that user **36** was happy with their experience at store **62** and that feedback **64** provided by user **36** was “the clothing I purchased was classy but my cab got stuck in traffic”. Accordingly and for this example, AI/ML process **56** may identify the pertinent content (included within feedback **64**) as the phrase “the clothing I purchased was classy” and may ignore/remove the irrelevant content “my cab got stuck in traffic”. As (in this example) feedback **64** includes the word “clothing”, AI/ML process **56** may route feedback **64** to selection node **114** via selection branch **106**. Further, as feedback **64** also includes the word “classy”, AI/ML process **56** may route feedback **64** to good selection

node **140** via good selection branch **136** and may consider feedback **64** to be text-based customer feedback indicative of a good selection being available at store **62**.

(72) Model Generation Overview:

(73) While the following discussion concerns the automated generation of a probabilistic model, this is for illustrative purposes only and is not intended to be a limitation of this disclosure, as other configurations are possible and are considered to be within the scope of this disclosure. For example, the following discussion of automated generation may be utilized on any type of model. For example, the following discussion may be applicable to any other form of probabilistic model or any form of generic model (such as Dempster Shaffer theory or fuzzy logic).

(74) As discussed above, probabilistic model **100** may be utilized to categorize information **58**, thus allowing the various messages included within information **58** to be routed to (in this simplified example) one of eight nodes (e.g., good service node **126**, bad service node **130**, good selection node **140**, bad selection node **144**, good location node **154**, bad location node **158**, good value node **168**, and bad value node **172**). For the following example, assume that store **62** is a long-standing and well-established shopping establishment. Further, assume that information **58** is a very large quantity of voice mail messages (>10,000 messages) that were left by customers of store **62** on a voice-based customer feedback line. Additionally, assume that this very large quantity of voice mail messages (>10,000) have been transcribed into a very large quantity of text-based messages (>10,000).

(75) AI/ML process **56** may be configured to automatically define probabilistic model **100** based upon information **58**. Accordingly, AI/ML process **56** may receive content (e.g., a very large quantity of text-based messages) and may be configured to define one or more probabilistic model variables for probabilistic model **100**. For example, AI/ML process **56** may be configured to allow a user to specify such probabilistic model variables. Another example of such variables may include but is not limited to values and/or ranges of values for a data flow variable. For the following discussion and for this disclosure, examples of a “variable” may include but are not limited to variables, parameters, ranges, branches and nodes.

(76) Specifically and for this example, assume that AI/ML process **56** defines the initial number of branches (i.e., the number of branches off of branching node **102**) within probabilistic model **100** as four (i.e., service branch **104**, selection branch **106**, location branch **108** and value branch **110**). The defining of the initial number of branches (i.e., the number of branches off of branching node **102**) within probabilistic model **100** as four may be effectuated in various ways (e.g., manually or algorithmically). Further and when defining probabilistic model **100** based, at least in part, upon information **58** and the one or more model variables (i.e., defining the number of branches off of branching node **102** as four), AI/ML process **56** may process information **58** to identify the pertinent content included within information **58**. As discussed above, AI/ML process **56** may identify the pertinent content (included within information **58**) and may ignore/remove the irrelevant content.

(77) This type of processing of information **58** may continue for all of the very large quantity of text-based messages (>10,000) included within information **58**. And using the probabilistic modeling technique described above, AI/ML process **56** may define a first version of the probabilistic model (e.g., probabilistic model **100**) based, at least in part, upon pertinent content found within information **58**. Accordingly, a first text-based message included within information **58** may be processed to extract pertinent information from that first message, wherein this pertinent information may be grouped in a manner to correspond (at least temporarily) with the requirement that four branches originate from branching node **102** (as defined above).

(78) As AI/ML process **56** continues to process information **58** to identify pertinent content included within information **58**, AI/ML process **56** may identify patterns within these text-based message included within information **58**. For example, the messages may all concern one or more of the service, the selection, the location and/or the value of store **62**. Further and e.g., using the probabilistic modeling technique described above, AI/ML process **56** may process information **58** to e.g.: a) sort text-based messages concerning the service into positive or negative service messages; b) sort text-based messages concerning the selection into positive or negative selection messages; c) sort text-based messages concerning the location into positive or negative location messages; and/or d) sort text-based messages concerning the value into positive or negative service messages. For example, AI/ML process **56** may define various lists (e.g., lists **128**, **132**, **142**, **146**, **156**, **160**, **170**, **174**) by starting with a root word (e.g., good or bad) and may then determine synonyms for these words and use those words and synonyms to populate lists **128**, **132**, **142**, **146**, **156**, **160**, **170**, **174**.

(79) Continuing with the above-stated example, once information **58** (or a portion thereof) is processed by AI/ML process **56**, AI/ML process **56** may define a first version of the probabilistic model (e.g., probabilistic

model **100**) based, at least in part, upon pertinent content found within information **58**. AI/ML process **56** may compare the first version of the probabilistic model (e.g., probabilistic model **100**) to information **58** to determine if the first version of the probabilistic model (e.g., probabilistic model **100**) is a good explanation of the content.

(80) When determining if the first version of the probabilistic model (e.g., probabilistic model **100**) is a good explanation of the content, AI/ML process **56** may use an ML algorithm to fit the first version of the probabilistic model (e.g., probabilistic model **100**) to the content, wherein examples of such an ML algorithm may include but are not limited to one or more of: an inferencing algorithm, a learning algorithm, an optimization algorithm, and a statistical algorithm.

(81) For example and as is known in the art, probabilistic model **100** may be used to generate messages (in addition to analyzing them). For example and when defining a first version of the probabilistic model (e.g., probabilistic model **100**) based, at least in part, upon pertinent content found within information **58**, AI/ML process **56** may define a weight for each branch within probabilistic model **100** based upon information **58**. For example, threat mitigation process **10** may equally weight each of branches **104**, **106**, **108**, **110** at 25%. Alternatively, if e.g., a larger percentage of information **58** concerned the service received at store **62**, threat mitigation process **10** may equally weight each of branches **106**, **108**, **110** at 20%, while more heavily weighting branch **104** at 40%.

(82) Accordingly and when AI/ML process **56** compares the first version of the probabilistic model (e.g., probabilistic model **100**) to information **58** to determine if the first version of the probabilistic model (e.g., probabilistic model **100**) is a good explanation of the content, AI/ML process **56** may generate a very large quantity of messages e.g., by auto-generating messages using the above-described probabilities, the above-described nodes & node types, and the words defined in the above-described lists (e.g., lists **128**, **132**, **142**, **146**, **156**, **160**, **170**, **174**), thus resulting in generated information **58'**. Generated information **58'** may then be compared to information **58** to determine if the first version of the probabilistic model (e.g., probabilistic model **100**) is a good explanation of the content. For example, if generated information **58'** exceeds a threshold level of similarity to information **58**, the first version of the probabilistic model (e.g., probabilistic model **100**) may be deemed a good explanation of the content. Conversely, if generated information **58'** does not exceed a threshold level of similarity to information **58**, the first version of the probabilistic model (e.g., probabilistic model **100**) may be deemed not a good explanation of the content.

(83) If the first version of the probabilistic model (e.g., probabilistic model **100**) is not a good explanation of the content, AI/ML process **56** may define a revised version of the probabilistic model (e.g., revised probabilistic model **100'**). When defining revised probabilistic model **100'**, AI/ML process **56** may e.g., adjust weighting, adjust probabilities, adjust node counts, adjust node types, and/or adjust branch counts to define the revised version of the probabilistic model (e.g., revised probabilistic model **100'**). Once defined, the above-described process of auto-generating messages (this time using revised probabilistic model **100'**) may be repeated and this newly-generated content (e.g., generated information **58''**) may be compared to information **58** to determine if e.g., revised probabilistic model **100'** is a good explanation of the content. If revised probabilistic model **100'** is not a good explanation of the content, the above-described process may be repeated until a proper probabilistic model is defined.

(84) The Threat Mitigation Process

(85) As discussed above, threat mitigation process **10** may include AI/ML process **56** (e.g., an artificial intelligence/machine learning process) that may be configured to process information (e.g., information **58**), wherein examples of information **58** may include but are not limited to platform information (e.g., structured or unstructured content) that may be scanned to detect security events (e.g., access auditing; anomalies; authentication; denial of services; exploitation; malware; phishing; spamming; reconnaissance; and/or web attack) within a monitored computing platform (e.g., computing platform **60**).

(86) Referring also to FIG. **3**, the monitored computing platform (e.g., computing platform **60**) utilized by business today may be a highly complex, multi-location computing system/network that may span multiple buildings/locations/countries. For this illustrative example, the monitored computing platform (e.g., computing platform **60**) is shown to include many discrete computing devices, examples of which may include but are not limited to: server computers (e.g., server computers **200**, **202**), desktop computers (e.g., desktop computer **204**), and laptop computers (e.g., laptop computer **206**), all of which may be coupled together via a network (e.g., network **208**), such as an Ethernet network. Computing platform **60** may be coupled to an external network (e.g., Internet **210**) through WAF (i.e., Web Application Firewall) **212**. A wireless access point (e.g., WAP **214**) may be configured to allow wireless devices (e.g., smartphone **216**) to

access computing platform **60**. Computing platform **60** may include various connectivity devices that enable the coupling of devices within computing platform **60**, examples of which may include but are not limited to: switch **216**, router **218** and gateway **220**. Computing platform **60** may also include various storage devices (e.g., NAS **222**), as well as functionality (e.g., API Gateway **224**) that allows software applications to gain access to one or more resources within computing platform **60**.

(87) In addition to the devices and functionality discussed above, other technology (e.g., security-relevant subsystems **226**) may be deployed within computing platform **60** to monitor the operation of (and the activity within) computing platform **60**. Examples of security-relevant subsystems **226** may include but are not limited to: CDN (i.e., Content Delivery Network) systems; DAM (i.e., Database Activity Monitoring) systems; UBA (i.e., User Behavior Analytics) systems; MDM (i.e., Mobile Device Management) systems; IAM (i.e., Identity and Access Management) systems; DNS (i.e., Domain Name Server) systems, antivirus systems, operating systems, data lakes; data logs; security-relevant software applications; security-relevant hardware systems; and resources external to the computing platform.

(88) Each of security-relevant subsystems **226** may monitor and log their activity with respect to computing platform **60**, resulting in the generation of platform information **228**. For example, platform information **228** associated with a client-defined MDM (i.e., Mobile Device Management) system may monitor and log the mobile devices that were allowed access to computing platform **60**.

(89) Further, SIEM (i.e., Security Information and Event Management) system **230** may be deployed within computing platform **60**. As is known in the art, SIEM system **230** is an approach to security management that combines SIM (security information management) functionality and SEM (security event management) functionality into one security management system. The underlying principles of a SIEM system is to aggregate relevant data from multiple sources, identify deviations from the norm and take appropriate action. For example, when a security event is detected, SIEM system **230** might log additional information, generate an alert and instruct other security controls to mitigate the security event. Accordingly, SIEM system **230** may be configured to monitor and log the activity of security-relevant subsystems **226** (e.g., CDN (i.e., Content Delivery Network) systems; DAM (i.e., Database Activity Monitoring) systems; UBA (i.e., User Behavior Analytics) systems; MDM (i.e., Mobile Device Management) systems; IAM (i.e., Identity and Access Management) systems; DNS (i.e., Domain Name Server) systems, antivirus systems, operating systems, data lakes; data logs; security-relevant software applications; security-relevant hardware systems; and resources external to the computing platform).

(90) Computing Platform Analysis & Reporting

(91) As will be discussed below in greater detail, threat mitigation process **10** may be configured to e.g., analyze computing platform **60** and provide reports to third-parties concerning the same. Further and since security-relevant subsystems **226** may monitor and log activity with respect to computing platform **60** and computing platform **60** may include a wide range of computing devices (e.g., server computers **200**, **202**, desktop computer **204**, laptop computer **206**, network **208**, web application firewall **212**, wireless access point **214**, switch **216**, router **218**, gateway **220**, NAS **222**, and API Gateway **224**), threat mitigation process **10** may provide holistic monitoring of the entirety of computing platform **60** (e.g., both central devices and end point devices), generally referred to as XDR (extended detection and response) functionality. As defined by analyst firm Gartner, Extended Detection and Response (XDR) is “a SaaS-based, vendor-specific, security threat detection and incident response tool that natively integrates multiple security products into a cohesive security operations system that unifies all licensed components.”

(92) Referring also to FIGS. **4-6**, threat mitigation process **10** may be configured to obtain and combine information from multiple security-relevant subsystem to generate a security profile for computing platform **60**. For example, threat mitigation process **10** may obtain **330** first system-defined platform information (e.g., system-defined platform information **232**) concerning a first security-relevant subsystem (e.g., the number of operating systems deployed) within computing platform **60** and may obtain **332** at least a second system-defined platform information (e.g., system-defined platform information **234**) concerning at least a second security-relevant subsystem (e.g., the number of antivirus systems deployed) within computing platform **60**.

(93) The first system-defined platform information (e.g., system-defined platform information **232**) and the at least a second system-defined platform information (e.g., system-defined platform information **234**) may be obtained from one or more log files defined for computing platform **60**.

(94) Specifically, system-defined platform information **232** and/or system-defined platform information **234** may be obtained from SIEM system **230**, wherein (and as discussed above) SIEM system **230** may be configured to monitor and log the activity of security-relevant subsystems **226** (e.g., CDN (i.e., Content



Delivery Network) systems; DAM (i.e., Database Activity Monitoring) systems; UBA (i.e., User Behavior Analytics) systems; MDM (i.e., Mobile Device Management) systems; IAM (i.e., Identity and Access Management) systems; DNS (i.e., Domain Name Server) systems, antivirus systems, operating systems, data lakes; data logs; security-relevant software applications; security-relevant hardware systems; and resources external to the computing platform).

(95) Alternatively, the first system-defined platform information (e.g., system-defined platform information **232**) and the at least a second system-defined platform information (e.g., system-defined platform information **234**) may be obtained from the first security-relevant subsystem (e.g., the operating systems themselves) and the at least a second security-relevant subsystem (e.g., the antivirus systems themselves). Specifically, system-defined platform information **232** and/or system-defined platform information **234** may be obtained directly from the security-relevant subsystems (e.g., the operating systems and/or the antivirus systems), which (as discussed above) may be configured to self-document their activity.

(96) Threat mitigation process **10** may combine **334** the first system-defined platform information (e.g., system-defined platform information **232**) and the at least a second system-defined platform information (e.g., system-defined platform information **234**) to form system-defined consolidated platform information **236**. Accordingly and in this example, system-defined consolidated platform information **236** may independently define the security-relevant subsystems (e.g., security-relevant subsystems **226**) present on computing platform **60**.

(97) Threat mitigation process **10** may generate **336** a security profile (e.g., security profile **350**) based, at least in part, upon system-defined consolidated platform information **236**. Through the use of security profile (e.g., security profile **350**), the user/owner/operator of computing platform **60** may be able to see that e.g., they have a security score of 605 out of a possible score of 1,000, wherein the average customer has a security score of 237. While security profile **350** is shown in the example to include several indicators that may enable a user to compare (in this example) computing platform **60** to other computing platforms, this is for illustrative purposes only and is not intended to be a limitation of this disclosure, as it is understood that other configurations are possible and are considered to be within the scope of this disclosure.

(98) Naturally, the format, appearance and content of security profile **350** may be varied greatly depending upon the design criteria and anticipated performance/use of threat mitigation process **10**. Accordingly, the appearance, format, completeness and content of security profile **350** is for illustrative purposes only and is not intended to be a limitation of this disclosure, as other configurations are possible and are considered to be within the scope of this disclosure. For example, content may be added to security profile **350**, removed from security profile **350**, and/or reformatted within security profile **350**.

(99) Additionally, threat mitigation process **10** may obtain **338** client-defined consolidated platform information **238** for computing platform **60** from a client information source, examples of which may include but are not limited to one or more client-completed questionnaires (e.g., questionnaires **240**) and/or one or more client-deployed platform monitors (e.g., client-deployed platform monitor **242**, which may be configured to effectuate SIEM functionality). Accordingly and in this example, client-defined consolidated platform information **238** may define the security-relevant subsystems (e.g., security-relevant subsystems **226**) that the client believes are present on computing platform **60**.

(100) When generating **336** a security profile (e.g., security profile **350**) based, at least in part, upon system-defined consolidated platform information **236**, threat mitigation process **10** may compare **340** the system-defined consolidated platform information (e.g., system-defined consolidated platform information **236**) to the client-defined consolidated platform information (e.g., client-defined consolidated platform information **238**) to define differential consolidated platform information **352** for computing platform **60**.

(101) Differential consolidated platform information **352** may include comparison table **354** that e.g., compares computing platform **60** to other computing platforms. For example and in this particular implementation of differential consolidated platform information **352**, comparison table **354** is shown to include three columns, namely: security-relevant subsystem column **356** (that identifies the security-relevant subsystems in question); system-defined consolidated platform information column **358** (that is based upon system-defined consolidated platform information **236** and independently defines what security-relevant subsystems are present on computing platform **60**); and client-defined consolidated platform column **360** (that is based upon client-defined platform information **238** and defines what security-relevant subsystems the client believes are present on computing platform **60**). As shown within comparison table **354**, there are considerable differences between that is actually present on computing platform **60** and what is believed to be present on computing platform **60** (e.g., 1 IAM system vs. 10 IAM systems; 4,000 operating systems vs.

10,000 operating systems, 6 DNS systems vs. 10 DNS systems; 0 antivirus systems vs. 1 antivirus system, and 90 firewalls vs. 150 firewalls).

(102) Naturally, the format, appearance and content of differential consolidated platform information **352** may be varied greatly depending upon the design criteria and anticipated performance/use of threat mitigation process **10**. Accordingly, the appearance, format, completeness and content of differential consolidated platform information **352** is for illustrative purposes only and is not intended to be a limitation of this disclosure, as other configurations are possible and are considered to be within the scope of this disclosure. For example, content may be added to differential consolidated platform information **352**, removed from differential consolidated platform information **352**, and/or reformatted within differential consolidated platform information **352**.

(103) Referring also to FIG. 7, threat mitigation process **10** may be configured to compare what security relevant subsystems are actually included within computing platform **60** versus what security relevant subsystems were believed to be included within computing platform **60**. As discussed above, threat mitigation process **10** may combine **334** the first system-defined platform information (e.g., system-defined platform information **232**) and the at least a second system-defined platform information (e.g., system-defined platform information **234**) to form system-defined consolidated platform information **236**.

(104) Threat mitigation process **10** may obtain **400** system-defined consolidated platform information **236** for computing platform **60** from an independent information source, examples of which may include but are not limited to: one or more log files defined for computing platform **60** (e.g., such as those maintained by SIEM system **230**); and two or more security-relevant subsystems (e.g., directly from the operating system security-relevant subsystem and the antivirus security-relevant subsystem) deployed within computing platform **60**.

(105) Further and as discussed above, threat mitigation process **10** may obtain **338** client-defined consolidated platform information **238** for computing platform **60** from a client information source, examples of which may include but are not limited to one or more client-completed questionnaires (e.g., questionnaires **240**) and/or one or more client-deployed platform monitors (e.g., client-deployed platform monitor **242**, which may be configured to effectuate SIEM functionality).

(106) Additionally and as discussed above, threat mitigation process **10** may compare **402** system-defined consolidated platform information **236** to client-defined consolidated platform information **238** to define differential consolidated platform information **352** for computing platform **60**, wherein differential consolidated platform information **352** may include comparison table **354** that e.g., compares computing platform **60** to other computing platforms.

(107) Threat mitigation process **10** may process **404** system-defined consolidated platform information **236** prior to comparing **402** system-defined consolidated platform information **236** to client-defined consolidated platform information **238** to define differential consolidated platform information **352** for computing platform **60**. Specifically, threat mitigation process **10** may process **404** system-defined consolidated platform information **236** so that it is comparable to client-defined consolidated platform information **238**.

(108) For example and when processing **404** system-defined consolidated platform information **236**, threat mitigation process **10** may homogenize **406** system-defined consolidated platform information **236** prior to comparing **402** system-defined consolidated platform information **236** to client-defined consolidated platform information **238** to define differential consolidated platform information **352** for computing platform **60**. Such homogenization **406** may result in system-defined consolidated platform information **236** and client-defined consolidated platform information **238** being comparable to each other (e.g., to accommodate for differing data nomenclatures/headers).

(109) Further and when processing **404** system-defined consolidated platform information **236**, threat mitigation process **10** may normalize **408** system-defined consolidated platform information **236** prior to comparing **402** system-defined consolidated platform information **236** to client-defined consolidated platform information **238** to define differential consolidated platform information **352** for computing platform **60** (e.g., to accommodate for data differing scales/ranges).

(110) Referring also to FIG. 8, threat mitigation process **10** may be configured to compare what security relevant subsystems are actually included within computing platform **60** versus what security relevant subsystems were believed to be included within computing platform **60**.

(111) As discussed above, threat mitigation process **10** may obtain **400** system-defined consolidated platform information **236** for computing platform **60** from an independent information source, examples of which may include but are not limited to: one or more log files defined for computing platform **60** (e.g., such as those

maintained by SIEM system **230**); and two or more security-relevant subsystems (e.g., directly from the operating system security-relevant subsystem and the antivirus security-relevant subsystem) deployed within computing platform **60**

(112) Further and as discussed above, threat mitigation process **10** may obtain **338** client-defined consolidated platform information **238** for computing platform **60** from a client information source, examples of which may include but are not limited to one or more client-completed questionnaires (e.g., questionnaires **240**) and/or one or more client-deployed platform monitors (e.g., client-deployed platform monitor **242**, which may be configured to effectuate SIEM functionality).

(113) Threat mitigation process **10** may present **450** differential consolidated platform information **352** for computing platform **60** to a third-party, examples of which may include but are not limited to the user/owner/operator of computing platform **60**.

(114) Additionally and as discussed above, threat mitigation process **10** may compare **402** system-defined consolidated platform information **236** to client-defined consolidated platform information **238** to define differential consolidated platform information **352** for computing platform **60**, wherein differential consolidated platform information **352** may include comparison table **354** that e.g., compares computing platform **60** to other computing platforms, wherein (and as discussed above) threat mitigation process **10** may process **404** (e.g., via homogenizing **406** and/or normalizing **408**) system-defined consolidated platform information **236** prior to comparing **402** system-defined consolidated platform information **236** to client-defined consolidated platform information **238** to define differential consolidated platform information **352** for computing platform **60**.

(115) Computing Platform Analysis & Recommendation

(116) As will be discussed below in greater detail, threat mitigation process **10** may be configured to e.g., analyze & display the vulnerabilities of computing platform **60**.

(117) Referring also to FIG. **9**, threat mitigation process **10** may be configured to make recommendations concerning security relevant subsystems that are missing from computing platform **60**. As discussed above, threat mitigation process **10** may obtain **500** consolidated platform information for computing platform **60** to identify one or more deployed security-relevant subsystems **226** (e.g., CDN (i.e., Content Delivery Network) systems; DAM (i.e., Database Activity Monitoring) systems; UBA (i.e., User Behavior Analytics) systems; MDM (i.e., Mobile Device Management) systems; IAM (i.e., Identity and Access Management) systems; DNS (i.e., Domain Name Server) systems, antivirus systems, operating systems, data lakes; data logs; security-relevant software applications; security-relevant hardware systems; and resources external to the computing platform). This consolidated platform information may be obtained from an independent information source (e.g., such as SIEM system **230** that may provide system-defined consolidated platform information **236**) and/or may be obtained from a client information source (e.g., such as questionnaires **240** that may provide client-defined consolidated platform information **238**).

(118) Referring also to FIG. **10**, threat mitigation process **10** may process **506** the consolidated platform information (e.g., system-defined consolidated platform information **236** and/or client-defined consolidated platform information **238**) to identify one or more non-deployed security-relevant subsystems (within computing platform **60**) and may then generate **508** a list of ranked & recommended security-relevant subsystems (e.g., non-deployed security-relevant subsystem list **550**) that ranks the one or more non-deployed security-relevant subsystems.

(119) For this particular illustrative example, non-deployed security-relevant subsystem list **550** is shown to include column **552** that identifies six non-deployed security-relevant subsystems, namely: a CDN subsystem, a WAF subsystem, a DAM subsystem; a UBA subsystem; an API subsystem, and an MDM subsystem.

(120) When generating **508** a list of ranked & recommended security-relevant subsystems (e.g., non-deployed security-relevant subsystem list **550**) that ranks the one or more non-deployed security-relevant subsystems, threat mitigation process **10** may rank **510** the one or more non-deployed security-relevant subsystems (e.g., a CDN subsystem, a WAF subsystem, a DAM subsystem; a UBA subsystem; a API subsystem, and an MDM subsystem) based upon the anticipated use of the one or more non-deployed security-relevant subsystems within computing platform **60**. This ranking **510** of the non-deployed security-relevant subsystems (e.g., a CDN subsystem, a WAF subsystem, a DAM subsystem; a UBA subsystem; a API subsystem, and an MDM subsystem) may be agnostic in nature and may be based on the functionality/effectiveness of the non-deployed security-relevant subsystems and the anticipated manner in which their implementation may impact the functionality/security of computing platform **60**.

(121) Threat mitigation process **10** may provide **512** the list of ranked & recommended security-relevant subsystems (e.g., non-deployed security-relevant subsystem list **550**) to a third-party, examples of which may include but are not limited to a user/owner/operator of computing platform **60**.

(122) Additionally, threat mitigation process **10** may identify **514** a comparative for at least one of the non-deployed security-relevant subsystems (e.g., a CDN subsystem, a WAF subsystem, a DAM subsystem; a UBA subsystem; an API subsystem, and an MDM subsystem) defined within the list of ranked & recommended security-relevant subsystems (e.g., non-deployed security-relevant subsystem list **550**). This comparative may include vendor customers in a specific industry comparative and/or vendor customers in any industry comparative.

(123) For example and in addition to column **552**, non-deployed security-relevant subsystem list **550** may include columns **554**, **556** for defining the comparatives for the six non-deployed security-relevant subsystems, namely: a CDN subsystem, a WAF subsystem, a DAM subsystem; a UBA subsystem; an API subsystem, and an MDM subsystem. Specifically, column **554** is shown to define comparatives concerning vendor customers that own the non-deployed security-relevant subsystems in a specific industry (i.e., the same industry as the user/owner/operator of computing platform **60**). Additionally, column **556** is shown to define comparatives concerning vendor customers that own the non-deployed security-relevant subsystems in any industry (i.e., not necessarily the same industry as the user/owner/operator of computing platform **60**). For example and concerning the comparatives of the WAF subsystem: 33% of the vendor customers in the same industry as the user/owner/operator of computing platform **60** deploy a WAF subsystem; while 71% of the vendor customers in any industry deploy a WAF subsystem.

(124) Naturally, the format, appearance and content of non-deployed security-relevant subsystem list **550** may be varied greatly depending upon the design criteria and anticipated performance/use of threat mitigation process **10**. Accordingly, the appearance, format, completeness and content of non-deployed security-relevant subsystem list **550** is for illustrative purposes only and is not intended to be a limitation of this disclosure, as other configurations are possible and are considered to be within the scope of this disclosure. For example, content may be added to non-deployed security-relevant subsystem list **550**, removed from non-deployed security-relevant subsystem list **550**, and/or reformatted within non-deployed security-relevant subsystem list **550**.

(125) Referring also to FIG. **11**, threat mitigation process **10** may be configured to compare the current capabilities to the possible capabilities of computing platform **60**. As discussed above, threat mitigation process **10** may obtain **600** consolidated platform information to identify current security-relevant capabilities for computing platform **60**. This consolidated platform information may be obtained from an independent information source (e.g., such as SIEM system **230** that may provide system-defined consolidated platform information **236**) and/or may be obtained from a client information source (e.g., such as questionnaires **240** that may provide client-defined consolidated platform information **238**). Threat mitigation process **10** may then determine **606** possible security-relevant capabilities for computing platform **60** (i.e., the difference between the current security-relevant capabilities of computing platform **60** and the possible security-relevant capabilities of computing platform **60**). For example, the possible security-relevant capabilities may concern the possible security-relevant capabilities of computing platform **60** using the currently-deployed security-relevant subsystems. Additionally/alternatively, the possible security-relevant capabilities may concern the possible security-relevant capabilities of computing platform **60** using one or more supplemental security-relevant subsystems.

(126) Referring also to FIG. **12** and as will be explained below, threat mitigation process **10** may generate **608** comparison information **650** that compares the current security-relevant capabilities of computing platform **60** to the possible security-relevant capabilities of computing platform **60** to identify security-relevant deficiencies. Comparison information **650** may include graphical comparison information, such as multi-axial graphical comparison information that simultaneously illustrates a plurality of security-relevant deficiencies.

(127) For example, comparison information **650** may define (in this particular illustrative example) graphical comparison information that include five axes (e.g. axes **652**, **654**, **656**, **658**, **660**) that correspond to five particular types of computer threats. Comparison information **650** includes origin **662**, the point at which computing platform **60** has no protection with respect to any of the five types of computer threats that correspond to axes **652**, **654**, **656**, **658**, **660**. Accordingly, as the capabilities of computing platform **60** are increased to counter a particular type of computer threat, the data point along the corresponding axis is proportionately displaced from origin **652**.

(128) As discussed above, threat mitigation process **10** may obtain **600** consolidated platform information to identify current security-relevant capabilities for computing platform **60**. Concerning such current security-relevant capabilities for computing platform **60**, these current security-relevant capabilities are defined by data points **664, 666, 668, 670, 672**, the combination of which define bounded area **674**. Bounded area **674** (in this example) defines the current security-relevant capabilities of computing platform **60**.

(129) Further and as discussed above, threat mitigation process **10** may determine **606** possible security-relevant capabilities for computing platform **60** (i.e., the difference between the current security-relevant capabilities of computing platform **60** and the possible security-relevant capabilities of computing platform **60**).

(130) As discussed above, the possible security-relevant capabilities may concern the possible security-relevant capabilities of computing platform **60** using the currently-deployed security-relevant subsystems. For example, assume that the currently-deployed security relevant subsystems are not currently being utilized to their full potential. Accordingly, certain currently-deployed security relevant subsystems may have certain features that are available but are not utilized and/or disabled. Further, certain currently-deployed security relevant subsystems may have expanded features available if additional licensing fees are paid. Therefore and concerning such possible security-relevant capabilities of computing platform **60** using the currently-deployed security-relevant subsystems, data points **676, 678, 680, 682, 684** may define bounded area **686** (which represents the full capabilities of the currently-deployed security-relevant subsystems within computing platform **60**).

(131) Further and as discussed above, the possible security-relevant capabilities may concern the possible security-relevant capabilities of computing platform **60** using one or more supplemental security-relevant subsystems. For example, assume that supplemental security-relevant subsystems are available for the deployment within computing platform **60**. Therefore and concerning such possible security-relevant capabilities of computing platform **60** using such supplemental security-relevant subsystems, data points **688, 690, 692, 694, 696** may define bounded area **698** (which represents the total capabilities of computing platform **60** when utilizing the full capabilities of the currently-deployed security-relevant subsystems and any supplemental security-relevant subsystems).

(132) Naturally, the format, appearance and content of comparison information **650** may be varied greatly depending upon the design criteria and anticipated performance/use of threat mitigation process **10**. Accordingly, the appearance, format, completeness and content of comparison information **650** is for illustrative purposes only and is not intended to be a limitation of this disclosure, as other configurations are possible and are considered to be within the scope of this disclosure. For example, content may be added to comparison information **650**, removed from comparison information **650**, and/or reformatted within comparison information **650**.

(133) Referring also to FIG. **13**, threat mitigation process **10** may be configured to generate a threat context score for computing platform **60**. As discussed above, threat mitigation process **10** may obtain **600** consolidated platform information to identify current security-relevant capabilities for computing platform **60**. This consolidated platform information may be obtained from an independent information source (e.g., such as SIEM system **230** that may provide system-defined consolidated platform information **236**) and/or may be obtained from a client information source (e.g., such as questionnaires **240** that may provide client-defined consolidated platform information **238**). As will be discussed below in greater detail, threat mitigation process **10** may determine **700** comparative platform information that identifies security-relevant capabilities for a comparative platform, wherein this comparative platform information may concern vendor customers in a specific industry (i.e., the same industry as the user/owner/operator of computing platform **60**) and/or vendor customers in any industry (i.e., not necessarily the same industry as the user/owner/operator of computing platform **60**).

(134) Referring also to FIG. **14** and as will be discussed below, threat mitigation process **10** may generate **702** comparison information **750** that compares the current security-relevant capabilities of computing platform **60** to the comparative platform information determined **700** for the comparative platform to identify a threat context indicator for computing platform **60**, wherein comparison information **750** may include graphical comparison information **752**.

(135) Graphical comparison information **752** (which in this particular example is a bar chart) may identify one or more of: a current threat context score **754** for a client (e.g., the user/owner/operator of computing platform **60**); a maximum possible threat context score **756** for the client (e.g., the user/owner/operator of computing platform **60**); a threat context score **758** for one or more vendor customers in a specific industry

(i.e., the same industry as the user/owner/operator of computing platform **60**); and a threat context score **760** for one or more vendor customers in any industry (i.e., not necessarily the same industry as the user/owner/operator of computing platform **60**).

(136) Naturally, the format, appearance and content of comparison information **750** may be varied greatly depending upon the design criteria and anticipated performance/use of threat mitigation process **10**.

Accordingly, the appearance, format, completeness and content of comparison information **750** is for illustrative purposes only and is not intended to be a limitation of this disclosure, as other configurations are possible and are considered to be within the scope of this disclosure. For example, content may be added to comparison information **750**, removed from comparison information **750**, and/or reformatted within comparison information **750**.

(137) Computing Platform Monitoring & Mitigation

(138) As will be discussed below in greater detail, threat mitigation process **10** may be configured to e.g., monitor the operation and performance of computing platform **60**.

(139) Referring also to FIG. **15**, threat mitigation process **10** may be configured to monitor the health of computing platform **60** and provide feedback to a third-party concerning the same. Threat mitigation process **10** may obtain **800** hardware performance information **244** concerning hardware (e.g., server computers, desktop computers, laptop computers, switches, firewalls, routers, gateways, WAPs, and NASs), deployed within computing platform **60**. Hardware performance information **244** may concern the operation and/or functionality of one or more hardware systems (e.g., server computers, desktop computers, laptop computers, switches, firewalls, routers, gateways, WAPs, and NASs) deployed within computing platform **60**.

(140) Threat mitigation process **10** may obtain **802** platform performance information **246** concerning the operation of computing platform **60**. Platform performance information **246** may concern the operation and/or functionality of computing platform **60**.

(141) When obtaining **802** platform performance information concerning the operation of computing platform **60**, threat mitigation process **10** may (as discussed above): obtain **400** system-defined consolidated platform information **236** for computing platform **60** from an independent information source (e.g., SIEM system **230**); obtain **338** client-defined consolidated platform information **238** for computing platform **60** from a client information (e.g., questionnaires **240**); and present **450** differential consolidated platform information **352** for computing platform **60** to a third-party, examples of which may include but are not limited to the user/owner/operator of computing platform **60**.

(142) When obtaining **802** platform performance information concerning the operation of computing platform **60**, threat mitigation process **10** may (as discussed above): obtain **500** consolidated platform information for computing platform **60** to identify one or more deployed security-relevant subsystems **226** (e.g., CDN (i.e., Content Delivery Network) systems; DAM (i.e., Database Activity Monitoring) systems; UBA (i.e., User Behavior Analytics) systems; MDM (i.e., Mobile Device Management) systems; IAM (i.e., Identity and Access Management) systems; DNS (i.e., Domain Name Server) systems, antivirus systems, operating systems, data lakes; data logs; security-relevant software applications; security-relevant hardware systems; and resources external to the computing platform); process **506** the consolidated platform information (e.g., system-defined consolidated platform information **236** and/or client-defined consolidated platform information **238**) to identify one or more non-deployed security-relevant subsystems (within computing platform **60**); generate **508** a list of ranked & recommended security-relevant subsystems (e.g., non-deployed security-relevant subsystem list **550**) that ranks the one or more non-deployed security-relevant subsystems; and provide **514** the list of ranked & recommended security-relevant subsystems (e.g., non-deployed security-relevant subsystem list **550**) to a third-party, examples of which may include but are not limited to a user/owner/operator of computing platform **60**.

(143) When obtaining **802** platform performance information concerning the operation of computing platform **60**, threat mitigation process **10** may (as discussed above): obtain **600** consolidated platform information to identify current security-relevant capabilities for the computing platform; determine **606** possible security-relevant capabilities for computing platform **60**; and generate **608** comparison information **650** that compares the current security-relevant capabilities of computing platform **60** to the possible security-relevant capabilities of computing platform **60** to identify security-relevant deficiencies.

(144) When obtaining **802** platform performance information concerning the operation of computing platform **60**, threat mitigation process **10** may (as discussed above): obtain **600** consolidated platform information to identify current security-relevant capabilities for computing platform **60**; determine **700** comparative platform information that identifies security-relevant capabilities for a comparative platform;

and generate **702** comparison information **750** that compares the current security-relevant capabilities of computing platform **60** to the comparative platform information determined **700** for the comparative platform to identify a threat context indicator for computing platform **60**.

(145) Threat mitigation process **10** may obtain **804** application performance information **248** concerning one or more applications (e.g., operating systems, user applications, security application, and utility application) deployed within computing platform **60**. Application performance information **248** may concern the operation and/or functionality of one or more software applications (e.g., operating systems, user applications, security application, and utility application) deployed within computing platform **60**.

(146) Referring also to FIG. **16**, threat mitigation process **10** may generate **806** holistic platform report (e.g., holistic platform reports **850**, **852**) concerning computing platform **60** based, at least in part, upon hardware performance information **244**, platform performance information **246** and application performance information **248**. Threat mitigation process **10** may be configured to receive e.g., hardware performance information **244**, platform performance information **246** and application performance information **248** at regular intervals (e.g., continuously, every minute, every ten minutes, etc.).

(147) As illustrated, holistic platform reports **850**, **852** may include various pieces of content such as e.g., thought clouds that identify topics/issues with respect to computing platform **60**, system logs that memorialize identified issues within computing platform **60**, data sources providing information to computing system **60**, and so on. The holistic platform report (e.g., holistic platform reports **850**, **852**) may identify one or more known conditions concerning the computing platform; and threat mitigation process **10** may effectuate **808** one or more remedial operations concerning the one or more known conditions.

(148) For example, assume that the holistic platform report (e.g., holistic platform reports **850**, **852**) identifies that computing platform **60** is under a DoS (i.e., Denial of Services) attack. In computing, a denial-of-service attack (DOS attack) is a cyber-attack in which the perpetrator seeks to make a machine or network resource unavailable to its intended users by temporarily or indefinitely disrupting services of a host connected to the Internet. Denial of service is typically accomplished by flooding the targeted machine or resource with superfluous requests in an attempt to overload systems and prevent some or all legitimate requests from being fulfilled.

(149) In response to detecting such a DOS attack, threat mitigation process **10** may effectuate **808** one or more remedial operations. For example and with respect to such a DOS attack, threat mitigation process **10** may effectuate **808** e.g., a remedial operation that instructs WAF (i.e., Web Application Firewall) **212** to deny all incoming traffic from the identified attacker based upon e.g., protocols, ports or the originating IP addresses.

(150) Threat mitigation process **10** may also provide **810** the holistic report (e.g., holistic platform reports **850**, **852**) to a third-party, examples of which may include but are not limited to a user/owner/operator of computing platform **60**.

(151) Naturally, the format, appearance and content of the holistic platform report (e.g., holistic platform reports **850**, **852**) may be varied greatly depending upon the design criteria and anticipated performance/use of threat mitigation process **10**. Accordingly, the appearance, format, completeness and content of the holistic platform report (e.g., holistic platform reports **850**, **852**) is for illustrative purposes only and is not intended to be a limitation of this disclosure, as other configurations are possible and are considered to be within the scope of this disclosure. For example, content may be added to the holistic platform report (e.g., holistic platform reports **850**, **852**), removed from the holistic platform report (e.g., holistic platform reports **850**, **852**), and/or reformatted within the holistic platform report (e.g., holistic platform reports **850**, **852**).

(152) Referring also to FIG. **17**, threat mitigation process **10** may be configured to monitor computing platform **60** for the occurrence of a security event and (in the event of such an occurrence) gather artifacts concerning the same. For example, threat mitigation process **10** may detect **900** a security event within computing platform **60** based upon identified suspect activity. Examples of such security events may include but are not limited to: DDOS events, DoS events, phishing events, spamming events, malware events, web attacks, and exploitation events.

(153) When detecting **900** a security event (e.g., DDOS events, DoS events, phishing events, spamming events, malware events, web attacks, and exploitation events) within computing platform **60** based upon identified suspect activity, threat mitigation process **10** may monitor **902** a plurality of sources to identify suspect activity within computing platform **60**.

(154) For example, assume that threat mitigation process **10** detects **900** a security event within computing platform **60**. Specifically, assume that threat mitigation process **10** is monitoring **902** a plurality of sources

(e.g., the various log files maintained by SIEM system **230**). And by monitoring **902** such sources, assume that threat mitigation process **10** detects **900** the receipt of inbound content (via an API) from a device having an IP address located in Uzbekistan; the subsequent opening of a port within WAF (i.e., Web Application Firewall) **212**; and the streaming of content from a computing device within computing platform **60** through that recently-opened port in WAF (i.e., Web Application Firewall) **212** and to a device having an IP address located in Moldova.

(155) Upon detecting **900** such a security event within computing platform **60**, threat mitigation process **10** may gather **904** artifacts (e.g., artifacts **250**) concerning the above-described security event. When gathering **904** artifacts (e.g., artifacts **250**) concerning the above-described security event, threat mitigation process **10** may gather **906** artifacts concerning the security event from a plurality of sources associated with the computing platform, wherein examples of such plurality of sources may include but are not limited to the various log files maintained by SIEM system **230**, and the various log files directly maintained by the security-relevant subsystems.

(156) Once the appropriate artifacts (e.g., artifacts **250**) are gathered **904**, threat mitigation process **10** may assign **908** a threat level to the above-described security event based, at least in part, upon the artifacts (e.g., artifacts **250**) gathered **904**.

(157) When assigning **908** a threat level to the above-described security event, threat mitigation process **10** may assign **910** a threat level using artificial intelligence/machine learning. As discussed above and with respect to artificial intelligence/machine learning being utilized to process data sets, an initial probabilistic model may be defined, wherein this initial probabilistic model may be subsequently (e.g., iteratively or continuously) modified and revised, thus allowing the probabilistic models and the artificial intelligence systems (e.g., AI/ML process **56**) to “learn” so that future probabilistic models may be more precise and may explain more complex data sets. As further discussed above, AI/ML process **56** may define an initial probabilistic model for accomplishing a defined task (e.g., the analyzing of information **58**), wherein the probabilistic model may be utilized to go from initial observations about information **58** (e.g., as represented by the initial branches of a probabilistic model) to conclusions about information **58** (e.g., as represented by the leaves of a probabilistic model). Accordingly and through the use of AI/ML process **56**, massive data sets concerning security events may be processed so that a probabilistic model may be defined (and subsequently revised) to assign **910** a threat level to the above-described security event.

(158) Once assigned **910** a threat level, threat mitigation process **10** may execute **912** a remedial action plan (e.g., remedial action plan **252**) based, at least in part, upon the assigned threat level.

(159) For example and when executing **912** a remedial action plan, threat mitigation process **10** may allow **914** the above-described suspect activity to continue when e.g., threat mitigation process **10** assigns **908** a “low” threat level to the above-described security event (e.g., assuming that it is determined that the user of the local computing device is streaming video of his daughter's graduation to his parents in Moldova).

(160) Further and when executing **912** a remedial action plan, threat mitigation process **10** may generate **916** a security event report (e.g., security event report **254**) based, at least in part, upon the artifacts (e.g., artifacts **250**) gathered **904**; and provide **918** the security event report (e.g., security event report **254**) to an analyst (e.g., analyst **256**) for further review when e.g., threat mitigation process **10** assigns **908** a “moderate” threat level to the above-described security event (e.g., assuming that it is determined that while the streaming of the content is concerning, the content is low value and the recipient is not a known bad actor).

(161) Further and when executing **912** a remedial action plan, threat mitigation process **10** may autonomously execute **920** a threat mitigation plan (shutting down the stream and closing the port) when e.g., threat mitigation process **10** assigns **908** a “severe” threat level to the above-described security event (e.g., assuming that it is determined that the streaming of the content is very concerning, as the content is high value and the recipient is a known bad actor).

(162) Additionally, threat mitigation process **10** may allow **922** a third-party (e.g., the user/owner/operator of computing platform **60**) to manually search for artifacts within computing platform **60**. For example, the third-party (e.g., the user/owner/operator of computing platform **60**) may be able to search the various information resources include within computing platform **60**, examples of which may include but are not limited to the various log files maintained by SIEM system **230**, and the various log files directly maintained by the security-relevant subsystems within computing platform **60**.

(163) Computing Platform Aggregation & Searching

(164) As will be discussed below in greater detail, threat mitigation process **10** may be configured to e.g., aggregate data sets and allow for unified search of those data sets.



(165) Referring also to FIG. 18, threat mitigation process **10** may be configured to consolidate multiple separate and discrete data sets to form a single, aggregated data set. For example, threat mitigation process **10** may establish **950** connectivity with a plurality of security-relevant subsystems (e.g., security-relevant subsystems **226**) within computing platform **60**. As discussed above, examples of security-relevant subsystems **226** may include but are not limited to: CDN (i.e., Content Delivery Network) systems; DAM (i.e., Database Activity Monitoring) systems; UBA (i.e., User Behavior Analytics) systems; MDM (i.e., Mobile Device Management) systems; IAM (i.e., Identity and Access Management) systems; DNS (i.e., Domain Name Server) systems, Antivirus systems, operating systems, data lakes; data logs; security-relevant software applications; security-relevant hardware systems; and resources external to the computing platform. (166) When establishing **950** connectivity with a plurality of security-relevant subsystems, threat mitigation process **10** may utilize **952** at least one application program interface (e.g., API Gateway **224**) to access at least one of the plurality of security-relevant subsystems. For example, a 1st API gateway may be utilized to access CDN (i.e., Content Delivery Network) system; a 2nd API gateway may be utilized to access DAM (i.e., Database Activity Monitoring) system; a 3rd API gateway may be utilized to access UBA (i.e., User Behavior Analytics) system; a 4th API gateway may be utilized to access MDM (i.e., Mobile Device Management) system; a 5th API gateway may be utilized to access IAM (i.e., Identity and Access Management) system; and a 6th API gateway may be utilized to access DNS (i.e., Domain Name Server) system.

(167) Threat mitigation process **10** may obtain **954** at least one security-relevant information set (e.g., a log file) from each of the plurality of security-relevant subsystems (e.g., CDN system; DAM system; UBA system; MDM system; IAM system; and DNS system), thus defining plurality of security-relevant information sets **258**. As would be expected, plurality of security-relevant information sets **258** may utilize a plurality of different formats and/or a plurality of different nomenclatures. Accordingly, threat mitigation process **10** may combine **956** plurality of security-relevant information sets **258** to form an aggregated security-relevant information set **260** for computing platform **60**.

(168) When combining **956** plurality of security-relevant information sets **258** to form aggregated security-relevant information set **260**, threat mitigation process **10** may homogenize **958** plurality of security-relevant information sets **258** to form aggregated security-relevant information set **260**. For example, threat mitigation process **10** may process one or more of security-relevant information sets **258** so that they all have a common format, a common nomenclature, and/or a common structure.

(169) Once threat mitigation process **10** combines **956** plurality of security-relevant information sets **258** to form an aggregated security-relevant information set **260** for computing platform **60**, threat mitigation process **10** may enable **960** a third-party (e.g., the user/owner/operator of computing platform **60**) to access aggregated security-relevant information set **260** and/or enable **962** a third-party (e.g., the user/owner/operator of computing platform **60**) to search aggregated security-relevant information set **260**.

(170) Referring also to FIG. 19, threat mitigation process **10** may be configured to enable the searching of multiple separate and discrete data sets using a single search operation. For example and as discussed above, threat mitigation process **10** may establish **950** connectivity with a plurality of security-relevant subsystems (e.g., security-relevant subsystems **226**) within computing platform **60**. As discussed above, examples of security-relevant subsystems **226** may include but are not limited to: CDN (i.e., Content Delivery Network) systems; DAM (i.e., Database Activity Monitoring) systems; UBA (i.e., User Behavior Analytics) systems; MDM (i.e., Mobile Device Management) systems; IAM (i.e., Identity and Access Management) systems; DNS (i.e., Domain Name Server) systems, Antivirus systems, operating systems, data lakes; data logs; security-relevant software applications; security-relevant hardware systems; and resources external to the computing platform.

(171) When establishing **950** connectivity with a plurality of security-relevant subsystems, threat mitigation process **10** may utilize **952** at least one application program interface (e.g., API Gateway **224**) to access at least one of the plurality of security-relevant subsystems. For example, a 1st API gateway may be utilized to access CDN (i.e., Content Delivery Network) system; a 2nd API gateway may be utilized to access DAM (i.e., Database Activity Monitoring) system; a 3rd API gateway may be utilized to access UBA (i.e., User Behavior Analytics) system; a 4th API gateway may be utilized to access MDM (i.e., Mobile Device Management) system; a 5th API gateway may be utilized to access IAM (i.e., Identity and Access Management) system; and a 6th API gateway may be utilized to access DNS (i.e., Domain Name Server) system.

(172) Threat mitigation process **10** may receive **1000** unified query **262** from a third-party (e.g., the

user/owner/operator of computing platform **60**) concerning the plurality of security-relevant subsystems. As discussed above, examples of security-relevant subsystems **226** may include but are not limited to: CDN (i.e., Content Delivery Network) systems; DAM (i.e., Database Activity Monitoring) systems; UBA (i.e., User Behavior Analytics) systems; MDM (i.e., Mobile Device Management) systems; IAM (i.e., Identity and Access Management) systems; DNS (i.e., Domain Name Server) systems, Antivirus systems, operating systems, data lakes; data logs; security-relevant software applications; security-relevant hardware systems; and resources external to the computing platform.

(173) Threat mitigation process **10** may distribute **1002** at least a portion of unified query **262** to the plurality of security-relevant subsystems, resulting in the distribution of plurality of queries **264** to the plurality of security-relevant subsystems. For example, assume that a third-party (e.g., the user/owner/operator of computing platform **60**) wishes to execute a search concerning the activity of a specific employee.

Accordingly, the third-party (e.g., the user/owner/operator of computing platform **60**) may formulate the appropriate unified query (e.g., unified query **262**) that defines the employee name, the computing device(s) of the employee, and the date range of interest. Unified query **262** may then be parsed to form plurality of queries **264**, wherein a specific query (within plurality of queries **264**) may be defined for each of the plurality of security-relevant subsystems and provided to the appropriate security-relevant subsystems. For example, a 1st query may be included within plurality of queries **264** and provided to CDN (i.e., Content Delivery Network) system; a 2nd query may be included within plurality of queries **264** and provided to DAM (i.e., Database Activity Monitoring) system; a 3rd query may be included within plurality of queries **264** and provided to UBA (i.e., User Behavior Analytics) system; a 4th query may be included within plurality of queries **264** and provided to MDM (i.e., Mobile Device Management) system; a 5th query may be included within plurality of queries **264** and provided to IAM (i.e., Identity and Access Management) system; and a 6th query may be included within plurality of queries **264** and provided to DNS (i.e., Domain Name Server) system.

(174) Threat mitigation process **10** may effectuate **1004** at least a portion of unified query **262** on each of the plurality of security-relevant subsystems to generate plurality of result sets **266**. For example, the 1st query may be executed on CDN (i.e., Content Delivery Network) system to produce a 1st result set; the 2nd query may be executed on DAM (i.e., Database Activity Monitoring) system to produce a 2nd result set; the 3rd query may be executed on UBA (i.e., User Behavior Analytics) system to produce a 3rd result set; the 4th query may be executed on MDM (i.e., Mobile Device Management) system to produce a 4th result set; the 5th query may be executed on IAM (i.e., Identity and Access Management) system to produce a 5th result set; and the 6th query may be executed on DNS (i.e., Domain Name Server) system to produce a 6th result set.

(175) Threat mitigation process **10** may receive **1006** plurality of result sets **266** from the plurality of security-relevant subsystems. Threat mitigation process **10** may then combine **1008** plurality of result sets **266** to form unified query result **268**. When combining **1008** plurality of result sets **266** to form unified query result **268**, threat mitigation process **10** may homogenize **1010** plurality of result sets **266** to form unified query result **268**. For example, threat mitigation process **10** may process one or more discrete result sets included within plurality of result sets **266** so that the discrete result sets within plurality of result sets **266** all have a common format, a common nomenclature, and/or a common structure. Threat mitigation process **10** may then provide **1012** unified query result **268** to the third-party (e.g., the user/owner/operator of computing platform **60**).

(176) Referring also to FIG. **20**, threat mitigation process **10** may be configured to utilize artificial intelligence/machine learning to automatically consolidate multiple separate and discrete data sets to form a single, aggregated data set. For example and as discussed above, threat mitigation process **10** may establish **950** connectivity with a plurality of security-relevant subsystems (e.g., security-relevant subsystems **226**) within computing platform **60**. As discussed above, examples of security-relevant subsystems **226** may include but are not limited to: CDN (i.e., Content Delivery Network) systems; DAM (i.e., Database Activity Monitoring) systems; UBA (i.e., User Behavior Analytics) systems; MDM (i.e., Mobile Device Management) systems; IAM (i.e., Identity and Access Management) systems; DNS (i.e., Domain Name Server) systems, Antivirus systems, operating systems, data lakes; data logs; security-relevant software applications; security-relevant hardware systems; and resources external to the computing platform.

(177) As discussed above and when establishing **950** connectivity with a plurality of security-relevant subsystems, threat mitigation process **10** may utilize **952** at least one application program interface (e.g., API Gateway **224**) to access at least one of the plurality of security-relevant subsystems. For example, a 1st API gateway may be utilized to access CDN (i.e., Content Delivery Network) system; a 2nd API gateway may be

utilized to access DAM (i.e., Database Activity Monitoring) system; a 3rd API gateway may be utilized to access UBA (i.e., User Behavior Analytics) system; a 4th API gateway may be utilized to access MDM (i.e., Mobile Device Management) system; a 5th API gateway may be utilized to access IAM (i.e., Identity and Access Management) system; and a 6th API gateway may be utilized to access DNS (i.e., Domain Name Server) system.

(178) As discussed above, threat mitigation process **10** may obtain **954** at least one security-relevant information set (e.g., a log file) from each of the plurality of security-relevant subsystems (e.g., CDN system; DAM system; UBA system; MDM system; IAM system; and DNS system), thus defining plurality of security-relevant information sets **258**. As would be expected, plurality of security-relevant information sets **258** may utilize a plurality of different formats and/or a plurality of different nomenclatures.

(179) Threat mitigation process **10** may process **1050** plurality of security-relevant information sets **258** using artificial learning/machine learning to identify one or more commonalities amongst plurality of security-relevant information sets **258**. As discussed above and with respect to artificial intelligence/machine learning being utilized to process data sets, an initial probabilistic model may be defined, wherein this initial probabilistic model may be subsequently (e.g., iteratively or continuously) modified and revised, thus allowing the probabilistic models and the artificial intelligence systems (e.g., AI/ML process **56**) to “learn” so that future probabilistic models may be more precise and may explain more complex data sets. As further discussed above, AI/ML process **56** may define an initial probabilistic model for accomplishing a defined task (e.g., the analyzing of information **58**), wherein the probabilistic model may be utilized to go from initial observations about information **58** (e.g., as represented by the initial branches of a probabilistic model) to conclusions about information **58** (e.g., as represented by the leaves of a probabilistic model). Accordingly and through the use of AI/ML process **56**, plurality of security-relevant information sets **258** may be processed so that a probabilistic model may be defined (and subsequently revised) to identify one or more commonalities (e.g., common headers, common nomenclatures, common data ranges, common data types, common formats, etc.) amongst plurality of security-relevant information sets **258**. When processing **1050** plurality of security-relevant information sets **258** using artificial learning/machine learning to identify one or more commonalities amongst plurality of security-relevant information sets **258**, threat mitigation process **10** may utilize **1052** a decision tree (e.g., probabilistic model **100**) based, at least in part, upon one or more previously-acquired security-relevant information sets.

(180) Threat mitigation process **10** may combine **1054** plurality of security-relevant information sets **258** to form aggregated security-relevant information set **260** for computing platform **60** based, at least in part, upon the one or more commonalities identified.

(181) When combining **1054** plurality of security-relevant information sets **258** to form aggregated security-relevant information set **260** for computing platform **60** based, at least in part, upon the one or more commonalities identified, threat mitigation process **10** may homogenize **1056** plurality of security-relevant information sets **258** to form aggregated security-relevant information set **260**. For example, threat mitigation process **10** may process one or more of security-relevant information sets **258** so that they all have a common format, a common nomenclature, and/or a common structure.

(182) Once threat mitigation process **10** combines **1054** plurality of security-relevant information sets **258** to form an aggregated security-relevant information set **260** for computing platform **60**, threat mitigation process **10** may enable **1058** a third-party (e.g., the user/owner/operator of computing platform **60**) to access aggregated security-relevant information set **260** and/or enable **1060** a third-party (e.g., the user/owner/operator of computing platform **60**) to search aggregated security-relevant information set **260**.

(183) Threat Event Information Updating

(184) As will be discussed below in greater detail, threat mitigation process **10** may be configured to be updated concerning threat event information.

(185) Referring also to FIG. **21**, threat mitigation process **10** may be configured to receive updated threat event information for security-relevant subsystems **226**. For example, threat mitigation process **10** may receive **1100** updated threat event information **270** concerning computing platform **60**, wherein updated threat event information **270** may define one or more of: updated threat listings; updated threat definitions; updated threat methodologies; updated threat sources; and updated threat strategies. Threat mitigation process **10** may enable **1102** updated threat event information **270** for use with one or more security-relevant subsystems **226** within computing platform **60**. As discussed above, examples of security-relevant subsystems **226** may include but are not limited to: CDN (i.e., Content Delivery Network) systems; DAM (i.e., Database Activity Monitoring) systems; UBA (i.e., User Behavior Analytics) systems; MDM (i.e.,

Mobile Device Management) systems; IAM (i.e., Identity and Access Management) systems; DNS (i.e., Domain Name Server) systems, Antivirus systems, operating systems, data lakes; data logs; security-relevant software applications; security-relevant hardware systems; and resources external to the computing platform.

(186) When enabling **1102** updated threat event information **270** for use with one or more security-relevant subsystems **226** within computing platform **60**, threat mitigation process **10** may install **1104** updated threat event information **270** on one or more security-relevant subsystems **226** within computing platform **60**.

(187) Threat mitigation process **10** may retroactively apply **1106** updated threat event information **270** to previously-generated information associated with one or more security-relevant subsystems **226**.

(188) When retroactively apply **1106** updated threat event information **270** to previously-generated information associated with one or more security-relevant subsystems **226**, threat mitigation process **10** may: apply **1108** updated threat event information **270** to one or more previously-generated log files (not shown) associated with one or more security-relevant subsystems **226**; apply **1110** updated threat event information **270** to one or more previously-generated data files (not shown) associated with one or more security-relevant subsystems **226**; and apply **1112** updated threat event information **270** to one or more previously-generated application files (not shown) associated with one or more security-relevant subsystems **226**.

(189) Additionally, /alternatively, threat mitigation process **10** may proactively apply **1114** updated threat event information **270** to newly-generated information associated with one or more security-relevant subsystems **226**.

(190) When proactively applying **1114** updated threat event information **270** to newly-generated information associated with one or more security-relevant subsystems **226**, threat mitigation process **10** may: apply **1116** updated threat event information **270** to one or more newly-generated log files (not shown) associated with one or more security-relevant subsystems **226**; apply **1118** updated threat event information **270** to one or more newly-generated data files (not shown) associated with one or more security-relevant subsystems **226**; and apply **1120** updated threat event information **270** to one or more newly-generated application files (not shown) associated with one or more security-relevant subsystems **226**.

(191) Referring also to FIG. **22**, threat mitigation process **10** may be configured to receive updated threat event information **270** for security-relevant subsystems **226**. For example and as discussed above, threat mitigation process **10** may receive **1100** updated threat event information **270** concerning computing platform **60**, wherein updated threat event information **270** may define one or more of: updated threat listings; updated threat definitions; updated threat methodologies; updated threat sources; and updated threat strategies. Further and as discussed above, threat mitigation process **10** may enable **1102** updated threat event information **270** for use with one or more security-relevant subsystems **226** within computing platform **60**. As discussed above, examples of security-relevant subsystems **226** may include but are not limited to: CDN (i.e., Content Delivery Network) systems; DAM (i.e., Database Activity Monitoring) systems; UBA (i.e., User Behavior Analytics) systems; MDM (i.e., Mobile Device Management) systems; IAM (i.e., Identity and Access Management) systems; DNS (i.e., Domain Name Server) systems, Antivirus systems, operating systems, data lakes; data logs; security-relevant software applications; security-relevant hardware systems; and resources external to the computing platform.

(192) As discussed above and when enabling **1102** updated threat event information **270** for use with one or more security-relevant subsystems **226** within computing platform **60**, threat mitigation process **10** may install **1104** updated threat event information **270** on one or more security-relevant subsystems **226** within computing platform **60**.

(193) Sometimes, it may not be convenient and/or efficient to immediately apply updated threat event information **270** to security-relevant subsystems **226**. Accordingly, threat mitigation process **10** may schedule **1150** the application of updated threat event information **270** to previously-generated information associated with one or more security-relevant subsystems **226**.

(194) When scheduling **1150** the application of updated threat event information **270** to previously-generated information associated with one or more security-relevant subsystems **226**, threat mitigation process **10** may: schedule **1152** the application of updated threat event information **270** to one or more previously-generated log files (not shown) associated with one or more security-relevant subsystems **226**; schedule **1154** the application of updated threat event information **270** to one or more previously-generated data files (not shown) associated with one or more security-relevant subsystems **226**; and schedule **1156** the application of updated threat event information **270** to one or more previously-generated application files (not shown) associated with one or more security-relevant subsystems **226**.

(195) Additionally, /alternatively, threat mitigation process **10** may schedule **1158** the application of the

updated threat event information to newly-generated information associated with the one or more security-relevant subsystems.

(196) When scheduling **1158** the application of updated threat event information **270** to newly-generated information associated with one or more security-relevant subsystems **226**, threat mitigation process **10** may: schedule **1160** the application of updated threat event information **270** to one or more newly-generated log files (not shown) associated with one or more security-relevant subsystems **226**; schedule **1162** the application of updated threat event information **270** to one or more newly-generated data files (not shown) associated with one or more security-relevant subsystems **226**; and schedule **1164** the application of updated threat event information **270** to one or more newly-generated application files (not shown) associated with one or more security-relevant subsystems **226**.

(197) Referring also to FIGS. **23-24**, threat mitigation process **10** may be configured to initially display analytical data, which may then be manipulated/updated to include automation data. For example, threat mitigation process **10** may display **1200** initial security-relevant information **1250** that includes analytical information (e.g., thought cloud **1252**). Examples of such analytical information may include but is not limited to one or more of: investigative information; and hunting information.

(198) Investigative Information (a portion of analytical information): Unified searching and/or automated searching, such as e.g., a security event occurring and searches being performed to gather artifacts concerning that security event.

(199) Hunt Information (a portion of analytical information): Targeted searching/investigations, such as the monitoring and cataloging of the videos that an employee has watched or downloaded over the past 30 days.

(200) Threat mitigation process **10** may allow **1202** a third-party (e.g., the user/owner/operator of computing platform **60**) to manipulate initial security-relevant information **1250** with automation information.

(201) Automate Information (a portion of automation): The execution of a single (and possibly simple) action one time, such as the blocking an IP address from accessing computing platform **60** whenever such an attempt is made.

(202) Orchestrate Information (a portion of automation): The execution of a more complex batch (or series) of tasks, such as sensing an unauthorized download via an API and a) shutting down the API, adding the requesting IP address to a blacklist, and closing any ports opened for the requestor.

(203) When allowing **1202** a third-party (e.g., the user/owner/operator of computing platform **60**) to manipulate initial security-relevant information **1250** with automation information, threat mitigation process **10** may allow **1204** a third-party (e.g., the user/owner/operator of computing platform **60**) to select the automation information to add to initial security-relevant information **1250** to generate revised security-relevant information **1250'**. For example and when allowing **1204** a third-party (e.g., the user/owner/operator of computing platform **60**) to select the automation information to add to initial security-relevant information **1250** to generate revised security-relevant information **1250'**, threat mitigation process **10** may allow **1206** the third-party (e.g., the user/owner/operator of computing platform **60**) to choose a specific type of automation information from a plurality of automation information types.

(204) For example, the third-party (e.g., the user/owner/operator of computing platform **60**) may choose to add/initiate the automation information to generate revised security-relevant information **1250'**. Accordingly, threat mitigation process **10** may render selectable options (e.g., selectable buttons **1254**, **1256**) that the third-party (e.g., the user/owner/operator of computing platform **60**) may select to manipulate initial security-relevant information **1250** with automation information to generate revised security-relevant information **1250'**. For this particular example, the third-party (e.g., the user/owner/operator of computing platform **60**) may choose two different options to manipulate initial security-relevant information **1250**, namely: "block ip" or "search", both of which will result in threat mitigation process **10** generating **1208** revised security-relevant information **1250'** (that includes the above-described automation information).

(205) When generating **1208** revised security-relevant information **1250'** (that includes the above-described automation information), threat mitigation process **10** may combine **1210** the automation information (that results from selecting "block IP" or "search") and initial security-relevant information **1250** to generate and render **1212** revised security-relevant information **1250'**.

(206) When rendering **1212** revised security-relevant information **1250'**, threat mitigation process **10** may render **1214** revised security-relevant information **1250'** within interactive report **1258**.

(207) Training Routine Generation and Execution

(208) As will be discussed below in greater detail, threat mitigation process **10** may be configured to allow for the manual or automatic generation of training routines, as well as the execution of the same.

(209) Referring also to FIG. 25, threat mitigation process **10** may be configured to allow for the manual generation of testing routine **272**. For example, threat mitigation process **10** may define **1300** training routine **272** for a specific attack (e.g., a Denial of Services attack) of computing platform **60**. Specifically, threat mitigation process **10** may generate **1302** a simulation of the specific attack (e.g., a Denial of Services attack) by executing training routine **272** within a controlled test environment, an example of which may include but is not limited to virtual machine **274** executed on a computing device (e.g., computing device **12**).

(210) When generating **1302** a simulation of the specific attack (e.g., a Denial of Services attack) by executing training routine **272** within the controlled test environment (e.g., virtual machine **274**), threat mitigation process **10** may render **1304** the simulation of the specific attack (e.g., a Denial of Services attack) on the controlled test environment (e.g., virtual machine **274**).

(211) Threat mitigation process **10** may allow **1306** a trainee (e.g., trainee **276**) to view the simulation of the specific attack (e.g., a Denial of Services attack) and may allow **1308** the trainee (e.g., trainee **276**) to provide a trainee response (e.g., trainee response **278**) to the simulation of the specific attack (e.g., a Denial of Services attack). For example, threat mitigation process **10** may execute training routine **272**, which trainee **276** may “watch” and provide trainee response **278**.

(212) Threat mitigation process **10** may then determine **1310** the effectiveness of trainee response **278**, wherein determining **1310** the effectiveness of the trainee response may include threat mitigation process **10** assigning **1312** a grade (e.g., a letter grade or a number grade) to trainee response **278**.

(213) Referring also to FIG. 26, threat mitigation process **10** may be configured to allow for the automatic generation of testing routine **272**. For example, threat mitigation process **10** may utilize **1350** artificial intelligence/machine learning to define training routine **272** for a specific attack (e.g., a Denial of Services attack) of computing platform **60**.

(214) As discussed above and with respect to artificial intelligence/machine learning being utilized to process data sets, an initial probabilistic model may be defined, wherein this initial probabilistic model may be subsequently (e.g., iteratively or continuously) modified and revised, thus allowing the probabilistic models and the artificial intelligence systems (e.g., AI/ML process **56**) to “learn” so that future probabilistic models may be more precise and may explain more complex data sets. As further discussed above, AI/ML process **56** may define an initial probabilistic model for accomplishing a defined task (e.g., the analyzing of information **58**), wherein the probabilistic model may be utilized to go from initial observations about information **58** (e.g., as represented by the initial branches of a probabilistic model) to conclusions about information **58** (e.g., as represented by the leaves of a probabilistic model). Accordingly and through the use of AI/ML process **56**, information may be processed so that a probabilistic model may be defined (and subsequently revised) to define training routine **272** for a specific attack (e.g., a Denial of Services attack) of computing platform **60**.

(215) When using **1350** artificial intelligence/machine learning to define training routine **272** for a specific attack (e.g., a Denial of Services attack) of computing platform **60**, threat mitigation process **10** may process **1352** security-relevant information to define training routine **272** for specific attack (e.g., a Denial of Services attack) of computing platform **60**. Further and when using **1350** artificial intelligence/machine learning to define training routine **272** for a specific attack (e.g., a Denial of Services attack) of computing platform **60**, threat mitigation process **10** may utilize **1354** security-relevant rules to define training routine **272** for a specific attack (e.g., a Denial of Services attack) of computing platform **60**. Accordingly, security-relevant information that e.g., defines the symptoms of e.g., a Denial of Services attack and security-relevant rules that define the behavior of e.g., a Denial of Services attack may be utilized by threat mitigation process **10** when defining training routine **272**.

(216) As discussed above, threat mitigation process **10** may generate **1302** a simulation of the specific attack (e.g., a Denial of Services attack) by executing training routine **272** within a controlled test environment, an example of which may include but is not limited to virtual machine **274** executed on a computing device (e.g., computing device **12**).

(217) Further and as discussed above, when generating **1302** a simulation of the specific attack (e.g., a Denial of Services attack) by executing training routine **272** within the controlled test environment (e.g., virtual machine **274**), threat mitigation process **10** may render **1304** the simulation of the specific attack (e.g., a Denial of Services attack) on the controlled test environment (e.g., virtual machine **274**).

(218) Threat mitigation process **10** may allow **1306** a trainee (e.g., trainee **276**) to view the simulation of the specific attack (e.g., a Denial of Services attack) and may allow **1308** the trainee (e.g., trainee **276**) to provide a trainee response (e.g., trainee response **278**) to the simulation of the specific attack (e.g., a Denial

of Services attack). For example, threat mitigation process **10** may execute training routine **272**, which trainee **276** may “watch” and provide trainee response **278**.

(219) Threat mitigation process **10** may utilize **1356** artificial intelligence/machine learning to revise training routine **272** for the specific attack (e.g., a Denial of Services attack) of computing platform **60** based, at least in part, upon trainee response **278**.

(220) As discussed above, threat mitigation process **10** may then determine **1310** the effectiveness of trainee response **278**, wherein determining **1310** the effectiveness of the trainee response may include threat mitigation process **10** assigning **1312** a grade (e.g., a letter grade or a number grade) to trainee response **278**.

(221) Referring also to FIG. **27**, threat mitigation process **10** may be configured to allow a trainee to choose their training routine. For example mitigation process **10** may allow **1400** a third-party (e.g., the user/owner/operator of computing platform **60**) to select a training routine for a specific attack (e.g., a Denial of Services attack) of computing platform **60**, thus defining a selected training routine. When allowing **1400** a third-party (e.g., the user/owner/operator of computing platform **60**) to select a training routine for a specific attack (e.g., a Denial of Services attack) of computing platform **60**, threat mitigation process **10** may allow **1402** the third-party (e.g., the user/owner/operator of computing platform **60**) to choose a specific training routine from a plurality of available training routines. For example, the third-party (e.g., the user/owner/operator of computing platform **60**) may be able to select a specific type of attack (e.g., DDOS events, DoS events, phishing events, spamming events, malware events, web attacks, and exploitation events) and/or select a specific training routine (that may or may not disclose the specific type of attack).

(222) Once selected, threat mitigation process **10** may analyze **1404** the requirements of the selected training routine (e.g., training routine **272**) to determine a quantity of entities required to effectuate the selected training routine (e.g., training routine **272**), thus defining one or more required entities. For example, assume that training routine **272** has three required entities (e.g., an attacked device and two attacking devices). According, threat mitigation process **10** may generate **1406** one or more virtual machines (e.g., such as virtual machine **274**) to emulate the one or more required entities. In this particular example, threat mitigation process **10** may generate **1406** three virtual machines, a first VM for the attacked device, a second VM for the first attacking device and a third VM for the second attacking device. As is known in the art, a virtual machine (VM) is a virtual emulation of a physical computing system. Virtual machines may be based on computer architectures and may provide the functionality of a physical computer, wherein their implementations may involve specialized hardware, software, or a combination thereof.

(223) Threat mitigation process **10** may generate **1408** a simulation of the specific attack (e.g., a Denial of Services attack) by executing the selected training routine (e.g., training routine **272**). When generating **1408** the simulation of the specific attack (e.g., a Denial of Services attack) by executing the selected training routine (e.g., training routine **272**), threat mitigation process **10** may render **1410** the simulation of the specific attack (e.g., a Denial of Services attack) by executing the selected training routine (e.g., training routine **272**) within a controlled test environment (e.g., such as virtual machine **274**).

(224) As discussed above, threat mitigation process **10** may allow **1306** a trainee (e.g., trainee **276**) to view the simulation of the specific attack (e.g., a Denial of Services attack) and may allow **1308** the trainee (e.g., trainee **276**) to provide a trainee response (e.g., trainee response **278**) to the simulation of the specific attack (e.g., a Denial of Services attack). For example, threat mitigation process **10** may execute training routine **272**, which trainee **276** may “watch” and provide trainee response **278**.

(225) Further and as discussed above, threat mitigation process **10** may then determine **1310** the effectiveness of trainee response **278**, wherein determining **1310** the effectiveness of the trainee response may include threat mitigation process **10** assigning **1312** a grade (e.g., a letter grade or a number grade) to trainee response **278**.

(226) When training is complete, threat mitigation process **10** may cease **1412** the simulation of the specific attack (e.g., a Denial of Services attack), wherein ceasing **1412** the simulation of the specific attack (e.g., a Denial of Services attack) may include threat mitigation process **10** shutting down **1414** the one or more virtual machines (e.g., the first VM for the attacked device, the second VM for the first attacking device and the third VM for the second attacking device).

(227) Information Routing

(228) As will be discussed below in greater detail, threat mitigation process **10** may be configured to route information based upon whether the information is more threat-pertinent or less threat-pertinent.

(229) Referring also to FIG. **28**, threat mitigation process **10** may be configured to route more threat-pertinent content in a specific manner. For example, threat mitigation process **10** may receive **1450** platform

information (e.g., log files) from a plurality of security-relevant subsystems (e.g., security-relevant subsystems **226**). As discussed above, examples of security-relevant subsystems **226** may include but are not limited to: CDN (i.e., Content Delivery Network) systems; DAM (i.e., Database Activity Monitoring) systems; UBA (i.e., User Behavior Analytics) systems; MDM (i.e., Mobile Device Management) systems; IAM (i.e., Identity and Access Management) systems; DNS (i.e., Domain Name Server) systems, Antivirus systems, operating systems, data lakes; data logs; security-relevant software applications; security-relevant hardware systems; and resources external to the computing platform.

(230) Threat mitigation process **10** may process **1452** this platform information (e.g., log files) to generate processed platform information. And when processing **1452** this platform information (e.g., log files) to generate processed platform information, threat mitigation process **10** may: parse **1454** the platform information (e.g., log files) into a plurality of subcomponents (e.g., columns, rows, etc.) to allow for compensation of varying formats and/or nomenclature; enrich **1456** the platform information (e.g., log files) by including supplemental information from external information resources; and/or utilize **1458** artificial intelligence/machine learning (in the manner described above) to identify one or more patterns/trends within the platform information (e.g., log files).

(231) Threat mitigation process **10** may identify **1460** more threat-pertinent content **280** included within the processed content, wherein identifying **1460** more threat-pertinent content **280** included within the processed content may include processing **1462** the processed content to identify actionable processed content that may be used by a threat analysis engine (e.g., SIEM system **230**) for correlation purposes. Threat mitigation process **10** may route **1464** more threat-pertinent content **280** to this threat analysis engine (e.g., SIEM system **230**).

(232) Referring also to FIG. **29**, threat mitigation process **10** may be configured to route less threat-pertinent content in a specific manner. For example and as discussed above, threat mitigation process **10** may receive **1450** platform information (e.g., log files) from a plurality of security-relevant subsystems (e.g., security-relevant subsystems **226**). As discussed above, examples of security-relevant subsystems **226** may include but are not limited to: CDN (i.e., Content Delivery Network) systems; DAM (i.e., Database Activity Monitoring) systems; UBA (i.e., User Behavior Analytics) systems; MDM (i.e., Mobile Device Management) systems; IAM (i.e., Identity and Access Management) systems; DNS (i.e., Domain Name Server) systems, Antivirus systems, operating systems, data lakes; data logs; security-relevant software applications; security-relevant hardware systems; and resources external to the computing platform

(233) Further and as discussed above, threat mitigation process **10** may process **1452** this platform information (e.g., log files) to generate processed platform information. And when processing **1452** this platform information (e.g., log files) to generate processed platform information, threat mitigation process **10** may: parse **1454** the platform information (e.g., log files) into a plurality of subcomponents (e.g., columns, rows, etc.) to allow for compensation of varying formats and/or nomenclature; enrich **1456** the platform information (e.g., log files) by including supplemental information from external information resources; and/or utilize **1458** artificial intelligence/machine learning (in the manner described above) to identify one or more patterns/trends within the platform information (e.g., log files).

(234) Threat mitigation process **10** may identify **1500** less threat-pertinent content **282** included within the processed content, wherein identifying **1500** less threat-pertinent content **282** included within the processed content may include processing **1502** the processed content to identify non-actionable processed content that is not usable by a threat analysis engine (e.g., SIEM system **230**) for correlation purposes. Threat mitigation process **10** may route **1504** less threat-pertinent content **282** to a long-term storage system (e.g., long term storage system **284**). Further, threat mitigation process **10** may be configured to allow **1506** a third-party (e.g., the user/owner/operator of computing platform **60**) to access and search long term storage system **284**.

(235) Automated Analysis

(236) As will be discussed below in greater detail, threat mitigation process **10** may be configured to automatically analyze a detected security event.

(237) Referring also to FIG. **30**, threat mitigation process **10** may be configured to automatically classify and investigate a detected security event. As discussed above and in response to a security event being detected, threat mitigation process **10** may obtain **1550** one or more artifacts (e.g., artifacts **250**) concerning the detected security event. Examples of such a detected security event may include but are not limited to one or more of: access auditing; anomalies; authentication; denial of services; exploitation; malware; phishing; spamming; reconnaissance; and web attack. These artifacts (e.g., artifacts **250**) may be obtained **1550** from a plurality of sources associated with the computing platform, wherein examples of such plurality of sources



may include but are not limited to the various log files maintained by SIEM system **230**, and the various log files directly maintained by the security-relevant subsystems

(238) Threat mitigation process **10** may obtain **1552** artifact information (e.g., artifact information **286**) concerning the one or more artifacts (e.g., artifacts **250**), wherein artifact information **286** may be obtained from information resources include within (or external to) computing platform **60**.

(239) For example and when obtaining **1552** artifact information **286** concerning the one or more artifacts (e.g., artifacts **250**), threat mitigation process **10** may obtain **1554** artifact information **286** concerning the one or more artifacts (e.g., artifacts **250**) from one or more investigation resources (such as third-party resources that may e.g., provide information on known bad actors).

(240) Once the investigation is complete, threat mitigation process **10** may generate **1556** a conclusion (e.g., conclusion **288**) concerning the detected security event (e.g., a Denial of Services attack) based, at least in part, upon the detected security event (e.g., a Denial of Services attack), the one or more artifacts (e.g., artifacts **250**), and artifact information **286**. Threat mitigation process **10** may document **1558** the conclusion (e.g., conclusion **288**), report **1560** the conclusion (e.g., conclusion **288**) to a third-party (e.g., the user/owner/operator of computing platform **60**). Further, threat mitigation process **10** may obtain **1562** supplemental artifacts and artifact information (if needed to further the investigation).

(241) While the system is described above as being computer-implemented, this is for illustrative purposes only and is not intended to be a limitation of this disclosure, as other configurations are possible and are considered to be within the scope of this disclosure. For example, some or all of the above-described system may be implemented by a human being.

(242) Unified Searching

(243) As discussed above, threat mitigation process **10** may be configured to e.g., analyze a monitored computing platform (e.g., computing platform **60**) and provide information to third-parties concerning the same. Further and as discussed above, such a monitored computing platform (e.g., computing platform **60**) may be a highly complex, multi-location computing system/network that may span multiple buildings/locations/countries.

(244) For this illustrative example, the monitored computing platform (e.g., computing platform **60**) is shown to include many discrete computing devices, examples of which may include but are not limited to: server computers (e.g., server computers **200**, **202**), desktop computers (e.g., desktop computer **204**), and laptop computers (e.g., laptop computer **206**), all of which may be coupled together via a network (e.g., network **208**), such as an Ethernet network. Computing platform **60** may be coupled to an external network (e.g., Internet **210**) through WAF (i.e., Web Application Firewall) **212**. A wireless access point (e.g., WAP **214**) may be configured to allow wireless devices (e.g., smartphone **216**) to access computing platform **60**. Computing platform **60** may include various connectivity devices that enable the coupling of devices within computing platform **60**, examples of which may include but are not limited to: switch **216**, router **218** and gateway **220**. Computing platform **60** may also include various storage devices (e.g., NAS **222**), as well as functionality (e.g., API Gateway **224**) that allows software applications to gain access to one or more resources within computing platform **60**.

(245) In addition to the devices and functionality discussed above, other technology (e.g., security-relevant subsystems **226**) may be deployed within computing platform **60** to monitor the operation of (and the activity within) computing platform **60**. Examples of security-relevant subsystems **226** may include but are not limited to: CDN (i.e., Content Delivery Network) systems; DAM (i.e., Database Activity Monitoring) systems; UBA (i.e., User Behavior Analytics) systems; MDM (i.e., Mobile Device Management) systems; IAM (i.e., Identity and Access Management) systems; DNS (i.e., Domain Name Server) systems, antivirus systems, operating systems, data lakes; data logs; security-relevant software applications; security-relevant hardware systems; and resources external to the computing platform. Each of security-relevant subsystems **226** may monitor and log their activity with respect to computing platform **60**, resulting in the generation of platform information **228**. For example, platform information **228** associated with a client-defined MDM (i.e., Mobile Device Management) system may monitor and log the mobile devices that were allowed access to computing platform **60**.

(246) Further, SIEM (i.e., Security Information and Event Management) system **230** may be deployed within computing platform **60**. As is known in the art, SIEM system **230** is an approach to security management that combines SIM (security information management) functionality and SEM (security event management) functionality into one security management system. The underlying principles of a SIEM system is to aggregate relevant data from multiple sources, identify deviations from the norm and take appropriate action.

For example, when a security event is detected, SIEM system **230** might log additional information, generate an alert and instruct other security controls to mitigate the security event. Accordingly, SIEM system **230** may be configured to monitor and log the activity of security-relevant subsystems **226** (e.g., CDN (i.e., Content Delivery Network) systems; DAM (i.e., Database Activity Monitoring) systems; UBA (i.e., User Behavior Analytics) systems; MDM (i.e., Mobile Device Management) systems; IAM (i.e., Identity and Access Management) systems; DNS (i.e., Domain Name Server) systems, antivirus systems, operating systems, data lakes; data logs; security-relevant software applications; security-relevant hardware systems; and resources external to the computing platform).

(247) Referring also to FIGS. **31-32**, threat mitigation process **10** may be configured to enable the querying of multiple separate and discrete subsystems (e.g., security-relevant subsystems **226**) using a single query operation. For example, threat mitigation process **10** may establish **1600** connectivity with a plurality of security-relevant subsystems (e.g., security-relevant subsystems **226**) within computing platform **60**.

(248) As discussed above, examples of security-relevant subsystems **226** may include but are not limited to: CDN (i.e., Content Delivery Network) systems; DAM (i.e., Database Activity Monitoring) systems; UBA (i.e., User Behavior Analytics) systems; MDM (i.e., Mobile Device Management) systems; IAM (i.e., Identity and Access Management) systems; DNS (i.e., Domain Name Server) systems, Antivirus systems, operating systems, data lakes; data logs; security-relevant software applications; security-relevant hardware systems; and resources external to the computing platform.

(249) When establishing **1600** connectivity with a plurality of security-relevant subsystems (e.g., security-relevant subsystems **226**), threat mitigation process **10** may utilize at least one application program interface (e.g., API Gateway **224**) to access at least one of the plurality of security-relevant subsystems. For example, a 1st API gateway may be utilized to access CDN (i.e., Content Delivery Network) system; a 2nd API gateway may be utilized to access DAM (i.e., Database Activity Monitoring) system; a 3rd API gateway may be utilized to access UBA (i.e., User Behavior Analytics) system; a 4th API gateway may be utilized to access MDM (i.e., Mobile Device Management) system; a 5th API gateway may be utilized to access IAM (i.e., Identity and Access Management) system; and a 6th API gateway may be utilized to access DNS (i.e., Domain Name Server) system.

(250) In order to enable the querying of multiple separate and discrete subsystems (e.g., security-relevant subsystems **226**) using a single query operation, threat mitigation process **10** may map **1602** one or more data fields of unified platform **290** (e.g., a platform effectuated by threat mitigation process **10**) to one or more data fields of each of the plurality of security-relevant subsystems (e.g., security-relevant subsystems **226**).

(251) For example, unified platform **290** may be a platform that enables a third-party (e.g., the user/owner/operator of computing platform **60**) to query multiple security-relevant subsystems (within security-relevant subsystems **226**), such as security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**. As discussed above, examples of such security-relevant subsystem (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**) may include but are not limited to: CDN (i.e., Content Delivery Network) systems; DAM (i.e., Database Activity Monitoring) systems; UBA (i.e., User Behavior Analytics) systems; MDM (i.e., Mobile Device Management) systems; IAM (i.e., Identity and Access Management) systems; DNS (i.e., Domain Name Server) systems, antivirus systems, operating systems, data lakes; data logs; security-relevant software applications; security-relevant hardware systems; and resources external to the computing platform.

(252) Each of these security-relevant subsystem (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**) may include a plurality of data fields that enable the third-party (e.g., the user/owner/operator of computing platform **60**) to search for and obtain information from these security-relevant subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**). For example: security-relevant subsystem **1650** is shown to include data fields **1656**, **1658**, **1660**, **1662**; security-relevant subsystem **1652** is shown to include data fields **1664**, **1666**, **1668**, **1670**; and security-relevant subsystem **1654** is shown to include data fields **1672**, **1674**, **1676**, **1678**.

(253) These data fields (e.g., data fields **1656**, **1658**, **1660**, **1662**, **1664**, **1666**, **1668**, **1670**, **1672**, **1674**, **1676**, **1678**) may be populatable by the third-party (e.g., the user/owner/operator of computing platform **60**) to enable such searching. For example, the third-party (e.g., the user/owner/operator of computing platform **60**) may populate these data fields by typing information into some of these data fields (e.g., data fields **1656**, **1658**, **1660**, **1666**, **1668**, **1670**, **1672**, **1674**, **1676**). Additionally/alternatively, the third-party (e.g., the user/owner/operator of computing platform **60**) may populate these data fields via a drop-down menu

available within these data fields (e.g., data fields **1662**, **1664**, **1678**). For example, data field **1662** is shown to be populatable via drop down menu **1680**, data field **1664** is shown to be populatable via drop down menu **1682**, and data field **1678** is shown to be populatable via drop down menu **1684**.

(254) Through the use of such data fields, the third-party (e.g., the user/owner/operator of computing platform **60**) may populate one or more of these data fields to define a query that may be effectuated on the information contained/available within these security-relevant subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**) so that the pertinent information may be obtained.

(255) Naturally, the subject matter of these individual data fields may vary depending upon the type of information available via these security-relevant subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**). As (in this example) these are security-relevant subsystems, the information available from these security-relevant subsystems concerns the security of computing platform **60** and/or any security events (e.g., access auditing; anomalies; authentication; denial of services; exploitation; malware; phishing; spamming; reconnaissance; and/or web attack) occurring therein. For example, some of these data fields may concern e.g., user names, user IDs, device locations, device types, device IP addresses, source IP addresses, destination IP addresses, port addresses, deployed operating systems, utilized bandwidth, etc.

(256) As discussed above, in order to enable the querying of multiple separate and discrete subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**) using a single query operation, threat mitigation process **10** may map **1602** one or more data fields of unified platform **290** (e.g., a platform effectuated by threat mitigation process **10**) to one or more data fields of each of the plurality of security-relevant subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**).

(257) In this particular example, unified platform **290** (e.g., a platform effectuated by threat mitigation process **10**) is shown to include four data fields (e.g., data fields **1686**, **1688**, **1690**, **1692**), wherein: data field **1686** within unified platform **290** concerns a user ID (and is entitled USER\_ID); data field **1688** within unified platform **290** concerns a device IP address (and is entitled DEVICE\_IP); data field **1690** within unified platform **290** concerns a destination IP address (and is entitled DESTINATION\_IP); and data field **1692** within unified platform **290** concerns a query result set (and is entitled QUERY\_RESULT). When mapping **1602** data fields within unified platform **290** (e.g., a platform effectuated by threat mitigation process **10**) to data fields within each of the plurality of security-relevant subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**), threat mitigation process **10** may only map **1602** data fields that are related with respect to subject matter.

(258) As discussed above, data field **1686** within unified platform **290** (e.g., a platform effectuated by threat mitigation process **10**) concerns a user ID (and is entitled USER\_ID). For this example, assume that: data field **1656** within security-relevant subsystem **1650** also concerns a user ID and is entitled USER; data field **1666** within security-relevant subsystem **1652** also concerns a user ID and is entitled ID; and data field **1676** within security-relevant subsystem **1654** also concerns a user ID and is entitled USER\_ID.

(259) Accordingly, threat mitigation process **10** may map **1602** data field **1686** of unified platform **290** (e.g., a platform effectuated by threat mitigation process **10**) to: data field **1656** of security-relevant subsystem **1650**; data field **1666** of security-relevant subsystem **1652**; and data field **1676** of security-relevant subsystem **1654**.

(260) As discussed above, data field **1688** within unified platform **290** (e.g., a platform effectuated by threat mitigation process **10**) concerns a device IP address (and is entitled DEVICE\_IP). For this example, assume that: data field **1660** within security-relevant subsystem **1650** also concerns a device IP address and is entitled DEV\_IP; data field **1670** within security-relevant subsystem **1652** also concerns a device IP address and is entitled IP\_DEVICE; and data field **1674** within security-relevant subsystem **1654** also concerns a device IP address and is entitled IP\_DEV.

(261) Accordingly, threat mitigation process **10** may map **1602** data field **1688** of unified platform **290** (e.g., a platform effectuated by threat mitigation process **10**) to: data field **1660** of security-relevant subsystem **1650**; data field **1670** of security-relevant subsystem **1652**; and data field **1674** of security-relevant subsystem **1654**.

(262) As discussed above, data field **1690** within unified platform **290** (e.g., a platform effectuated by threat mitigation process **10**) concerns a destination IP address (and is entitled DESTINATION\_IP). For this example, assume that: data field **1658** within security-relevant subsystem **1650** also concerns a destination IP

address and is entitled DEST\_IP; data field **1668** within security-relevant subsystem **1652** also concerns a destination IP address and is entitled IP\_DEST; and data field **1672** within security-relevant subsystem **1654** also concerns a destination IP address and is entitled IP\_DEST.

(263) Accordingly, threat mitigation process **10** may map **1602** data field **1690** of unified platform **290** (e.g., a platform effectuated by threat mitigation process **10**) to: data field **1658** of security-relevant subsystem **1650**; data field **1668** of security-relevant subsystem **1652**; and data field **1672** of security-relevant subsystem **1654**.

(264) As discussed above, data field **1692** within unified platform **290** (e.g., a platform effectuated by threat mitigation process **10**) concerns a query result (and is entitled QUERY\_RESULT). For this example, assume that: data field **1662** within security-relevant subsystem **1650** also concerns a query result and is entitled RESULT; data field **1664** within security-relevant subsystem **1652** also concerns a query result and is entitled Q\_RESULT; and data field **1678** within security-relevant subsystem **1654** also concerns a query result and is entitled RESULT\_Q.

(265) Accordingly, threat mitigation process **10** may map **1602** data field **1692** of unified platform **290** (e.g., a platform effectuated by threat mitigation process **10**) to: data field **1662** of security-relevant subsystem **1650**; data field **1664** of security-relevant subsystem **1652**; and data field **1678** of security-relevant subsystem **1654**.

(266) Through the use of threat mitigation process **10**, a query (e.g., query **1694**) may be defined within one or more of data fields **1686**, **1688**, **1690** of unified platform **290** (e.g., a platform effectuated by threat mitigation process **10**), wherein this query (e.g., query **1694**) may be provided (via the above-described mappings) to the appropriate data fields within the security-relevant subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**).

(267) Accordingly and when mapping **1602** one or more data fields of the unified platform (e.g., unified platform **290**) to one or more data fields of each of the plurality of security-relevant subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**), threat mitigation process **10** may map **1604** one or more data fields within a query structure of the unified platform (e.g., unified platform **290**) to one or more data fields within a query structure of each of the plurality of security-relevant subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**).

(268) Therefore, if a query (e.g., query **1694**) was defined on unified platform **290** (e.g., a platform effectuated by threat mitigation process **10**) that specified a user ID within data field **1686**, a device IP address within data field **1688**, and a destination IP address within data field **1690**; by mapping **1604** one or more data fields of the unified platform (e.g., unified platform **290**) to one or more data fields of each of the plurality of security-relevant subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**), this structured query (e.g., query **1694**) may be provided to the plurality of security-relevant subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**) in a fashion that enables the plurality of security-relevant subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**) to effectuate the structured query (e.g., query **1694**).

(269) Upon effectuating such a structured query (e.g., query **1694**), the plurality of security-relevant subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**) may each generate a subsystem-specific result set. For example, security-relevant subsystem **1650** may generate subsystem-specific result set **1696**, security-relevant subsystem **1652** may generate subsystem-specific result set **1698**, and security-relevant subsystem **1654** may generate subsystem-specific result set **1700**.

(270) Through the use of threat mitigation process **10**, subsystem-specific result sets (e.g., subsystem-specific result sets **1696**, **1698**, **1700**) may be defined within one or more of data fields (e.g., data fields **1662**, **1664**, **1678**) of the plurality of security-relevant subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**), wherein these subsystem-specific result sets (e.g., subsystem-specific result sets **1696**, **1698**, **1700**) may be provided (via the above-described mappings) to the appropriate data fields within the unified platform (e.g., unified platform **290**).

(271) Accordingly and when mapping **1602** one or more data fields of the unified platform (e.g., unified platform **290**) to one or more data fields of each of the plurality of security-relevant subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**), threat mitigation process **10** may map **1606** one or more data fields within a result set structure of each of the

plurality of security-relevant subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**) to one or more data fields within a result set structure of the unified platform (e.g., unified platform **290**).

(272) Therefore, by mapping **1606** one or more data fields within a result set structure of each of the plurality of security-relevant subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**) to one or more data fields within a result set structure of the unified platform (e.g., unified platform **290**), these subsystem-specific result sets (e.g., subsystem-specific result sets **1696**, **1698**, **1700**) may be provided to the unified platform (e.g., unified platform **290**) in a fashion that enables the unified platform (e.g., unified platform **290**) to properly process these subsystem-specific result sets (e.g., subsystem-specific result sets **1696**, **1698**, **1700**).

(273) It is foreseeable that over time, the data fields within the plurality of security-relevant subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**) may change. For example, additional data fields may be added to and/or certain data fields may be deleted from the plurality of security-relevant subsystems. Accordingly and in order to ensure that the above-described mapping remain current and accurate, such mappings may be periodically refreshed.

(274) Accordingly and when mapping **1602** one or more data fields of the unified platform (e.g., unified platform **290**) to one or more data fields of each of the plurality of security-relevant subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**), threat mitigation process **10** may map **1608** one or more data fields of the unified platform (e.g., unified platform **290**) to one or more data fields of each of the plurality of security-relevant subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**) at a defined periodicity.

(275) Therefore, at a certain frequency (e.g., every few minutes, every few hours, every few days, every few weeks or every few months), the above-describe mapping process may be reperformed to ensure that the above-described mappings are up to date.

(276) Further and when mapping **1602** one or more data fields of the unified platform (e.g., unified platform **290**) to one or more data fields of each of the plurality of security-relevant subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**), threat mitigation process **10** may proactively map **1610** one or more data fields of the unified platform (e.g., unified platform **290**) to one or more data fields of each of the plurality of security-relevant subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**).

(277) For example, the above-described mapping process may be proactively done, wherein threat mitigation process **10** actively monitors the security-relevant subsystems within computing platform **60** so that the data fields within these security-relevant subsystems may be proactively mapped **1610** prior to a third-party (e.g., the user/owner/operator of computing platform **60**) defining a query within unified platform **290**.

(278) Additionally and when mapping **1602** one or more data fields of the unified platform (e.g., unified platform **290**) to one or more data fields of each of the plurality of security-relevant subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**), threat mitigation process **10** may reactively map **1612** one or more data fields of the unified platform (e.g., unified platform **290**) to one or more data fields of each of the plurality of security-relevant subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**).

(279) For example, the above-described mapping process may be reactively performed, wherein threat mitigation process **10** may not actively monitor the security-relevant subsystems within computing platform **60** and the data fields within these security-relevant subsystems may be reactively mapped **1612** after a third-party (e.g., the user/owner/operator of computing platform **60**) defines a query within unified platform **290**.

(280) As discussed above, threat mitigation process **10** may allow a third-party (e.g., the user/owner/operator of computing platform **60**) to define **1614** a unified query (e.g., query **1694**) on a unified platform (e.g., unified platform **290**) concerning security-relevant subsystems **226** (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**).

(281) As discussed above, examples of security-relevant subsystems **226** may include but are not limited to: CDN (i.e., Content Delivery Network) systems; DAM (i.e., Database Activity Monitoring) systems; UBA (i.e., User Behavior Analytics) systems; MDM (i.e., Mobile Device Management) systems; IAM (i.e., Identity and Access Management) systems; DNS (i.e., Domain Name Server) systems, Antivirus systems, operating systems, data lakes; data logs; security-relevant software applications; security-relevant hardware systems; and resources external to the computing platform.

(282) Threat mitigation process **10** may denormalize **1616** the unified query (e.g., query **1694**) to define a subsystem-specific query for each of the plurality of security-relevant subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**), thus defining a plurality of subsystem-specific queries (e.g., subsystem-specific queries **1702**, **1704**, **1706**).

(283) As discussed above, unified platform **290** (e.g., a platform effectuated by threat mitigation process **10**) is shown to include four data fields (e.g., data fields **1686**, **1688**, **1690**, **1692**), wherein a third-party (e.g., the user/owner/operator of computing platform **60**) may utilize these data fields to define the unified query (e.g., query **1694**). As this unified query (e.g., query **1694**) may be used as the basis to search for pertinent information on (in this example) three entirely separate and discrete subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**), it is foreseeable that these subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**) may require queries to be structured differently.

(284) Accordingly and when denormalizing **1616** the unified query (e.g., query **1694**) to define a subsystem-specific query for each of the plurality of security-relevant subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**), thus defining a plurality of subsystem-specific queries (e.g., subsystem-specific queries **1702**, **1704**, **1706**), threat mitigation process **10** may translate **1618** a syntax of the unified query (e.g., query **1694**) to a syntax of each of the plurality of subsystem-specific queries (e.g., subsystem-specific queries **1702**, **1704**, **1706**). For example: security-relevant subsystem **1650** may only be capable of processing queries having a first structure and/or utilizing a first nomenclature; security-relevant subsystem **1652** may only be capable of processing queries having a second structure and/or utilizing a second nomenclature; and security-relevant subsystem **1654** may only be capable of processing queries having a third structure and/or utilizing a third nomenclature.

(285) Accordingly and when denormalizing **1616** the unified query (e.g., query **1694**) to define a plurality of subsystem-specific queries (e.g., subsystem-specific queries **1702**, **1704**, **1706**), threat mitigation process **10** may translate **1618** the syntax of the unified query (e.g., query **1694**) so that: subsystem-specific query **1702** has a first structure and/or utilizes a first nomenclature; subsystem-specific query **1704** has a second structure and/or utilizes a second nomenclature; subsystem-specific query **1706** has a third structure and/or utilizes a third nomenclature.

(286) Threat mitigation process **10** may provide **1620** the plurality of subsystem-specific queries (e.g., subsystem-specific queries **1702**, **1704**, **1706**) to the plurality of security-relevant subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**).

(287) The plurality of subsystem-specific queries (e.g., subsystem-specific queries **1702**, **1704**, **1706**) may be effectuated on the appropriate security-relevant subsystem. For example, subsystem-specific query **1702** may be effectuated on security-relevant subsystem **1650**, subsystem-specific query **1704** may be effectuated on security-relevant subsystem **1652**, and subsystem-specific query **1706** may be effectuated on security-relevant subsystem **1654**; resulting in the generation of subsystem-specific result sets. For example, security-relevant subsystem **1650** may generate subsystem-specific result set **1696**, security-relevant subsystem **1652** may generate subsystem-specific result set **1698**, and security-relevant subsystem **1654** may generate subsystem-specific result set **1700**.

(288) Threat mitigation process **10** may receive **1622** a plurality of subsystem-specific results sets (e.g., subsystem-specific result sets **1696**, **1698**, **1700**) from the plurality of security-relevant subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**, respectively) that were generated in response to the plurality of subsystem-specific queries (e.g., subsystem-specific queries **1702**, **1704**, **1706**).

(289) Threat mitigation process **10** may normalize **1624** the plurality of subsystem-specific results sets (e.g., subsystem-specific result sets **1696**, **1698**, **1700**) received from the plurality of security-relevant subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**, respectively) to define a unified result set (e.g., unified result set **1708**). For example, threat mitigation process **10** may process the plurality of subsystem-specific results sets (e.g., subsystem-specific result sets **1696**, **1698**, **1700**) so that the subsystem-specific results sets all have a common format, a common nomenclature, and/or a common structure.

(290) Accordingly and when normalizing **1624** the plurality of subsystem-specific results sets (e.g., subsystem-specific result sets **1696**, **1698**, **1700**) received from the plurality of security-relevant subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**, respectively) to define a unified result set (e.g., unified result set **1708**), threat mitigation process **10**

may translate **1626** a syntax of the plurality of subsystem-specific results sets (e.g., subsystem-specific result sets **1696, 1698, 1700**) to a syntax of the unified result set (e.g., unified result set **1708**). (291) As discussed above: security-relevant subsystem **1650** may only be capable of processing queries having a first structure and/or utilizing a first nomenclature; security-relevant subsystem **1652** may only be capable of processing queries having a second structure and/or utilizing a second nomenclature; and security-relevant subsystem **1654** may only be capable of processing queries having a third structure and/or utilizing a third nomenclature.

(292) Accordingly and when producing a result set: security-relevant subsystem **1650** may only be capable producing a result set (e.g., subsystem-specific result set **1696**) having a first structure and/or utilizing a first nomenclature; security-relevant subsystem **1652** may only be capable producing a result set (e.g., subsystem-specific result set **1698**) having a second structure and/or utilizing a second nomenclature; and security-relevant subsystem **1654** may only be capable producing a result set (e.g., subsystem-specific result set **1700**) having a third structure and/or utilizing a third nomenclature.

(293) Accordingly and when normalizing **1624** the plurality of subsystem-specific results sets (e.g., subsystem-specific result sets **1696, 1698, 1700**) received from the plurality of security-relevant subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**, respectively) to define a unified result set (e.g., unified result set **1708**), threat mitigation process **10** may translate **1626** the syntax of: subsystem-specific result set **1696** from a first structure/first nomenclature to a unified syntax of the unified result set (e.g., unified result set **1708**); subsystem-specific result set **1698** from a second structure/second nomenclature to the unified syntax of the unified result set (e.g., unified result set **1708**); subsystem-specific result set **1700** from a third structure/third nomenclature to a unified syntax of the unified result set (e.g., unified result set **1708**).

(294) Once normalized **1624, 1626**, threat mitigation process **10** may combine the subsystem-specific results sets (e.g., subsystem-specific result sets **1696, 1698, 1700**) to form the unified result set (e.g., unified result set **1708**), wherein threat mitigation process **10** may then provide **1628** the unified result set (e.g., unified result set **1708**) to a third-party (e.g., the user/owner/operator of computing platform **60**).

(295) Threat Hunting

(296) Referring also to FIG. **33**, threat mitigation process **10** may establish **1800** connectivity with a plurality of security-relevant subsystems (e.g., security-relevant subsystems **226**) within computing platform **60**, wherein examples of security-relevant subsystems **226** may include but are not limited to: CDN (i.e., Content Delivery Network) systems; DAM (i.e., Database Activity Monitoring) systems; UBA (i.e., User Behavior Analytics) systems; MDM (i.e., Mobile Device Management) systems; IAM (i.e., Identity and Access Management) systems; DNS (i.e., Domain Name Server) systems, Antivirus systems, operating systems, data lakes; data logs; security-relevant software applications; security-relevant hardware systems; and resources external to the computing platform.

(297) When establishing **1800** connectivity with a plurality of security-relevant subsystems (e.g., security-relevant subsystems **226**), threat mitigation process **10** may utilize at least one application program interface (e.g., API Gateway **224**) to access at least one of the plurality of security-relevant subsystems. For example, a 1st API gateway may be utilized to access CDN (i.e., Content Delivery Network) system; a 2nd API gateway may be utilized to access DAM (i.e., Database Activity Monitoring) system; a 3rd API gateway may be utilized to access UBA (i.e., User Behavior Analytics) system; a 4th API gateway may be utilized to access MDM (i.e., Mobile Device Management) system; a 5th API gateway may be utilized to access IAM (i.e., Identity and Access Management) system; and a 6th API gateway may be utilized to access DNS (i.e., Domain Name Server) system.

(298) As discussed above, threat mitigation process **10** may allow a third-party (e.g., the user/owner/operator of computing platform **60**) to define **1802** a unified query (e.g., query **1694**) on a unified platform (e.g., unified platform **290**) concerning security-relevant subsystems **226** (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**). In order to enable the querying of these separate and discrete subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654** within security-relevant subsystems **226**) using a single query operation, threat mitigation process **10** may map (in the manner discussed above) one or more data fields of unified platform **290** (e.g., a platform effectuated by threat mitigation process **10**) to one or more data fields of each of the plurality of security-relevant subsystems (e.g., e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654** within security-relevant subsystems **226**).

(299) Threat mitigation process **10** may denormalize **1804** the unified query (e.g., query **1694**) to define a subsystem-specific query for each of the plurality of security-relevant subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**), thus defining a plurality of subsystem-specific queries (e.g., subsystem-specific queries **1702**, **1704**, **1706**).

(300) One or more of the plurality of subsystem-specific queries (e.g., subsystem-specific queries **1702**, **1704**, **1706**) may have a defined execution schedule (e.g., defined execution schedule **1702S** for subsystem-specific query **1702**, defined execution schedule **1704S** for subsystem-specific query **1704**, and defined execution schedule **1706S** for subsystem-specific query **1706**). The defined execution schedule (e.g., defined execution schedule **1702S**, **1704S**, **1706S**) may include one or more of: a defined execution time; a defined execution date; a defined execution frequency; and a defined execution scope. Defined Execution Time: The defined execution schedule (e.g., defined execution schedule **1702S**, **1704S**, **1706S**) may define a particular time that a task is performed. For example, the defined execution schedule (e.g., defined execution schedule **1702S**, **1704S**, **1706S**) may define that an MDM (i.e., Mobile Device Management) system provide a device access report at midnight (local time) every day. Defined Execution Date: The defined execution schedule (e.g., defined execution schedule **1702S**, **1704S**, **1706S**) may define a particular date that a task is performed. For example, the defined execution schedule (e.g., defined execution schedule **1702S**, **1704S**, **1706S**) may define that a router provide a port opening report at COB every Friday (local time). Defined Execution Frequency: The defined execution schedule (e.g., defined execution schedule **1702S**, **1704S**, **1706S**) may define a particular frequency that a task is performed. For example, the defined execution schedule (e.g., defined execution schedule **1702S**, **1704S**, **1706S**) may define that a CDN (i.e., Content Delivery Network) system provide a quantity delivered report every hour. Defined Execution Scope: The defined execution schedule (e.g., defined execution schedule **1702S**, **1704S**, **1706S**) may define a particular scope for a task being performed. For example, the defined execution schedule (e.g., defined execution schedule **1702S**, **1704S**, **1706S**) may define that a switch provide an activity report for a specific port within the switch.

(301) These defined execution schedules (e.g., defined execution schedule **1702S**, **1704S**, **1706S**) may be a default execution schedule that is configured to be revisable by a third-party (e.g., the user/owner/operator of computing platform **60**). For example and with respect to these defined execution schedules (e.g., defined execution schedule **1702S**, **1704S**, **1706S**): the default time may be midnight, which may be revisable by the third-party (e.g., the user/owner/operator of computing platform **60**); the default date may be the 1.sup.st of the month, which may be revisable by the third-party (e.g., the user/owner/operator of computing platform **60**); the default frequency may be once, which may be revisable by the third-party (e.g., the user/owner/operator of computing platform **60**); and the default scope may be a narrower scope, which may be revisable by the third-party (e.g., the user/owner/operator of computing platform **60**).

(302) As discussed above, unified platform **290** (e.g., a platform effectuated by threat mitigation process **10**) is shown to include four data fields (e.g., data fields **1686**, **1688**, **1690**, **1692**), wherein a third-party (e.g., the user/owner/operator of computing platform **60**) may utilize these data fields to define the unified query (e.g., query **1694**). As this unified query (e.g., query **1694**) may be used as the basis to search for pertinent information on (in this example) three entirely separate and discrete subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**), it is foreseeable that these subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**) may require queries to be structured differently.

(303) Accordingly and when denormalizing **1804** the unified query (e.g., query **1694**) to define a subsystem-specific query for each of the plurality of security-relevant subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**), thus defining a plurality of subsystem-specific queries (e.g., subsystem-specific queries **1702**, **1704**, **1706**), threat mitigation process **10** may translate **1806** a syntax of the unified query (e.g., query **1694**) to a syntax of each of the plurality of subsystem-specific queries (e.g., subsystem-specific queries **1702**, **1704**, **1706**). For example: security-relevant subsystem **1650** may only be capable of processing queries having a first structure and/or utilizing a first nomenclature; security-relevant subsystem **1652** may only be capable of processing queries having a second structure and/or utilizing a second nomenclature; and security-relevant subsystem **1654** may only be capable of processing queries having a third structure and/or utilizing a third nomenclature.

(304) Accordingly and when denormalizing **1804** the unified query (e.g., query **1694**) to define a plurality of subsystem-specific queries (e.g., subsystem-specific queries **1702**, **1704**, **1706**), threat mitigation process **10** may translate **1806** the syntax of the unified query (e.g., query **1694**) so that: subsystem-specific query **1702** has a first structure and/or utilizes a first nomenclature; subsystem-specific query **1704** has a second structure



and/or utilizes a second nomenclature; subsystem-specific query **1706** has a third structure and/or utilizes a third nomenclature.

(305) Threat mitigation process **10** may provide **1808** the plurality of subsystem-specific queries (e.g., subsystem-specific queries **1702**, **1704**, **1706**) to the plurality of security-relevant subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**).

(306) The plurality of subsystem-specific queries (e.g., subsystem-specific queries **1702**, **1704**, **1706**) may be effectuated on the appropriate security-relevant subsystem. For example, subsystem-specific query **1702** may be effectuated on security-relevant subsystem **1650**, subsystem-specific query **1704** may be effectuated on security-relevant subsystem **1652**, and subsystem-specific query **1706** may be effectuated on security-relevant subsystem **1654**; resulting in the generation of subsystem-specific result sets. For example, security-relevant subsystem **1650** may generate subsystem-specific result set **1696**, security-relevant subsystem **1652** may generate subsystem-specific result set **1698**, and security-relevant subsystem **1654** may generate subsystem-specific result set **1700**.

(307) Threat mitigation process **10** may receive **1810** a plurality of subsystem-specific results sets (e.g., subsystem-specific result sets **1696**, **1698**, **1700**) from the plurality of security-relevant subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**, respectively) that were generated in response to the plurality of subsystem-specific queries (e.g., subsystem-specific queries **1702**, **1704**, **1706**).

(308) And by mapping (in the manner discussed above) one or more data fields within a result set structure of each of the plurality of security-relevant subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**) to one or more data fields within a result set structure of the unified platform (e.g., unified platform **290**), these subsystem-specific result sets (e.g., subsystem-specific result sets **1696**, **1698**, **1700**) may be provided to the unified platform (e.g., unified platform **290**) in a fashion that enables the unified platform (e.g., unified platform **290**) to properly process these subsystem-specific result sets (e.g., subsystem-specific result sets **1696**, **1698**, **1700**).

(309) Threat mitigation process **10** may normalize **1812** the plurality of subsystem-specific results sets (e.g., subsystem-specific result sets **1696**, **1698**, **1700**) received from the plurality of security-relevant subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**, respectively) to define a unified result set (e.g., unified result set **1708**). For example, threat mitigation process **10** may process the plurality of subsystem-specific results sets (e.g., subsystem-specific result sets **1696**, **1698**, **1700**) so that the subsystem-specific results sets all have a common format, a common nomenclature, and/or a common structure.

(310) Accordingly and when normalizing **1812** the plurality of subsystem-specific results sets (e.g., subsystem-specific result sets **1696**, **1698**, **1700**) received from the plurality of security-relevant subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**, respectively) to define a unified result set (e.g., unified result set **1708**), threat mitigation process **10** may translate **1814** a syntax of each of the plurality of subsystem-specific results sets (e.g., subsystem-specific result sets **1696**, **1698**, **1700**) to a syntax of the unified result set (e.g., unified result set **1708**).

(311) As discussed above: security-relevant subsystem **1650** may only be capable of processing queries having a first structure and/or utilizing a first nomenclature; security-relevant subsystem **1652** may only be capable of processing queries having a second structure and/or utilizing a second nomenclature; and security-relevant subsystem **1654** may only be capable of processing queries having a third structure and/or utilizing a third nomenclature.

(312) Accordingly and when producing a result set: security-relevant subsystem **1650** may only be capable producing a result set (e.g., subsystem-specific result set **1696**) having a first structure and/or utilizing a first nomenclature; security-relevant subsystem **1652** may only be capable producing a result set (e.g., subsystem-specific result set **1698**) having a second structure and/or utilizing a second nomenclature; and security-relevant subsystem **1654** may only be capable producing a result set (e.g., subsystem-specific result set **1700**) having a third structure and/or utilizing a third nomenclature.

(313) Accordingly and when normalizing **1812** the plurality of subsystem-specific results sets (e.g., subsystem-specific result sets **1696**, **1698**, **1700**) received from the plurality of security-relevant subsystems (e.g., security-relevant subsystem **1650**, security-relevant subsystem **1652** and security-relevant subsystem **1654**, respectively) to define a unified result set (e.g., unified result set **1708**), threat mitigation process **10** may translate **1814** the syntax of: subsystem-specific result set **1696** from a first structure/first nomenclature to a unified syntax of the unified result set (e.g., unified result set **1708**); subsystem-specific result set **1698**

from a second structure/second nomenclature to the unified syntax of the unified result set (e.g., unified result set **1708**); subsystem-specific result set **1700** from a third structure/third nomenclature to a unified syntax of the unified result set (e.g., unified result set **1708**).

(314) As could be imagined, it is foreseeable that e.g., one or more of security-relevant subsystems **226** may be offline when asked to perform a task (or go offline while performing a task). Therefore, one or more of subsystem-specific result sets **1696**, **1698**, **1700** may be missing/incomplete/defective. Accordingly, threat mitigation process **10** may be configured to determine **1816** whether one or more of the plurality of subsystem-specific queries (e.g., subsystem-specific queries **1702**, **1704**, **1706**) failed to execute properly, thus defining one or more failed subsystem-specific queries. And if one or more of the plurality of subsystem-specific queries (e.g., subsystem-specific queries **1702**, **1704**, **1706**) failed to execute properly, threat mitigation process **10** may reexecute **1818** the one or more failed subsystem-specific queries.

(315) As discussed above and in this example, threat mitigation process **10** provides **1808** subsystem-specific query **1702** to security-relevant subsystem **1650**; subsystem-specific query **1704** to security-relevant subsystem **1652**; and subsystem-specific query **1706** to security-relevant subsystem **1654**.

(316) Assume for this example that security-relevant subsystem **1650** went offline while executing subsystem-specific query **1702** and has since come back online. However, upon threat mitigation process **10** examining subsystem-specific result set **1696**, it is determined that subsystem-specific result set **1696** only contains 53,246 pieces of data (but is supposed to contain 100,000 pieces of data). Accordingly, threat mitigation process **10** may determine **1816** that subsystem-specific query **1702** failed to execute properly, thus defining subsystem-specific query **1702** as a failed subsystem-specific query. Accordingly, threat mitigation process **10** may reexecute **1818** the failed subsystem-specific query (e.g., subsystem-specific query **1702**) so the requested 100,000 pieces of data may be obtained from security-relevant subsystem **1650** (and the previously-obtained 53,246 pieces of data may be deleted).

(317) Once the plurality of subsystem-specific results sets (e.g., subsystem-specific result sets **1696**, **1698**, **1700**) are normalized **1812**, threat mitigation process **10** may combine the subsystem-specific results sets (e.g., subsystem-specific result sets **1696**, **1698**, **1700**) to form the unified result set (e.g., unified result set **1708**), wherein threat mitigation process **10** may then provide **1820** the unified result set (e.g., unified result set **1708**) to a third-party (e.g., the user/owner/operator of computing platform **60**).

(318) Generative AI/Large Language Model Utilization

(319) Threat mitigation process **10** may be configured to harness the power of Generative AI and Large Language Models (LLM). Generative AI models (e.g., AI/ML process **56**), as part of the broader artificial intelligence and machine learning landscape, are beginning to play a crucial role in enhancing network threat detection systems. Unlike traditional, discriminative models that classify input data into predefined categories (e.g., malicious or benign), generative models can learn to generate new data samples that are similar to the training data.

(320) Here's how these capabilities are being harnessed for network threat detection: Anomaly Detection: Generative models, such as Generative Adversarial Networks (GANs), can be trained on normal network traffic data to understand what typical network behavior looks like. Once trained, these models can generate new network traffic data that is expected to be similar to the "normal" traffic. By comparing real network traffic to these generated patterns, anomalies that could indicate potential threats, such as DDOS attacks or unauthorized access, can be detected more efficiently. Anomalies stand out because they deviate significantly from the generated "normal" patterns. Synthetic Data Generation: One of the challenges in training effective network threat detection systems is the scarcity of labeled data, especially for new and emerging threats. Generative AI models can help by creating large volumes of synthetic network traffic data, including both normal operations and various types of attack scenarios. This synthetic data can help in training more robust discriminative models (such as deep learning-based classifiers) by providing a richer, more varied dataset that covers a wider range of possible threats. Improving Data Privacy: In some contexts, using real network traffic data to train threat detection models can raise privacy concerns, especially if the data contains sensitive information. Generative models can be used to create synthetic data that mimics real network traffic without containing any actual user or proprietary information. This approach allows for the development and testing of threat detection systems in a manner that is respectful of privacy concerns. Evolving Threat Simulation: Cyber threats are constantly evolving, and keeping threat detection systems up to date can be challenging. Generative models can be used to simulate how threats might evolve over time, generating new, unseen threat patterns for testing the resilience of network systems. This proactive approach helps in identifying potential vulnerabilities before they are exploited in the wild. Training and Testing Environments:

Generative models can create realistic network environments for training cybersecurity professionals. By simulating various attack scenarios, these models provide a dynamic and challenging environment for cybersecurity training, allowing professionals to experience and respond to a range of threats in a controlled, risk-free setting. Limitations and Challenges: While generative AI models offer promising capabilities for network threat detection, there are also limitations and challenges. These include the complexity of training these models, the risk of generating misleading data, and the computational resources required. Additionally, as attackers also leverage AI, there's a continuous arms race between threat actors and defenders.

(321) Generally speaking, generative AI models are increasingly being explored for their potential to revolutionize network threat detection systems. By enhancing anomaly detection, enabling the generation of synthetic data, and simulating evolving threats, these models can significantly improve the ability of organizations to detect and respond to cyber threats more effectively and efficiently.

(322) As is known in the art, a large language model is an artificial intelligence system that is trained on massive amounts of text data to generate human-like responses to natural language inputs. These models use complex algorithms and neural networks to learn patterns and relationships in language data, enabling them to understand and generate responses to human language.

(323) The primary use of large language models is to improve natural language processing in a wide range of applications (e.g., virtual assistants, chatbots, search engines, and language translation tools). These models have made significant advances in recent years, and are now able to generate highly convincing and accurate responses to complex human language inputs.

(324) Large language models can be used to generate text in a variety of formats, including spoken language, written language, and code. They can also be used to summarize text, generate creative writing, and even create music or art. As the technology continues to improve, large language models are expected to play an increasingly important role in a wide range of industries, including healthcare, finance, and entertainment.

(325) Automated Generation of a Human-Readable Report

(326) Referring also to FIG. 34, threat mitigation process 10 may establish 1900 connectivity with a plurality of security-relevant subsystems (e.g., security-relevant subsystems 226) within a computing platform (e.g., computing platform 60).

(327) As discussed above, examples of security-relevant subsystems 226 may include but are not limited to: CDN (i.e., Content Delivery Network) systems; DAM (i.e., Database Activity Monitoring) systems; UBA (i.e., User Behavior Analytics) systems; MDM (i.e., Mobile Device Management) systems; IAM (i.e., Identity and Access Management) systems; DNS (i.e., Domain Name Server) systems, antivirus systems, operating systems, data lakes; data logs; security-relevant software applications; security-relevant hardware systems; and resources external to the computing platform.

(328) In a computing platform (e.g., computing platform 60), establishing connectivity between security-relevant subsystems (e.g., security-relevant subsystems 226)—such as firewalls, intrusion detection systems, intrusion prevention systems, and security information and event management systems—may require a multifaceted approach that encompasses network configuration, secure communication protocols, authentication, authorization mechanisms, and centralized management. Initially, each subsystem may be assigned a unique IP address, either statically or via DHCP, for identification and is often segmented into subnets to enhance both performance and security, with dedicated security subnets for these critical components.

(329) Secure communication among these subsystems may be paramount, utilizing protocols such as TLS/SSL for encryption, VPNs for creating secure connections over potentially insecure networks, and SSH for secure administrative actions and file transfers. The integrity and confidentiality of communications may be further ensured through the use of digital certificates within a Public Key Infrastructure, Access Control Lists, and Role-Based Access Control, which collectively authenticate devices and authorize only permitted interactions.

(330) The backbone of inter-subsystem connectivity may lie in network protocols like IPSec for securing IP communications and SNMPv3 for secure network management. These subsystems are typically managed through centralized consoles, allowing for uniform policy distribution and configuration across the network. Monitoring and logging may play crucial roles, with tools like Syslog and SIEM systems aggregating and analyzing log data for real-time security alerting.

(331) Moreover, network segmentation and the implementation of demilitarized zones (DMZs) may be strategies employed to further delineate and secure the network infrastructure. Firewalls may be meticulously configured to control traffic between these segments, enforcing security policies that dictate allowed and

blocked communications based on established rules.

(332) Through this comprehensive approach-integrating secure communication channels, robust authentication and authorization, and vigilant monitoring-security-relevant subsystems within a computer network can establish secure and efficient connectivity. This interconnectedness may be vital for the detection, prevention, and response to security threats, ensuring the overarching protection of information systems and data within an organization.

(333) Threat mitigation process **10** may receive **1902** an initial notification (e.g., initial notification **298**) of a security event from one of the security-relevant subsystems (e.g., security-relevant subsystems **226**). As discussed above, examples of such security events may include but are not limited to access auditing; anomalies; authentication; denial of services; exploitation; malware; phishing; spamming; reconnaissance; and/or web attack within a monitored computing platform (e.g., computing platform **60**).

(334) The initial notification (e.g., initial notification **298**) may include a computer-readable language portion that defines one or more specifics of the security event. An example of the computer-readable language portion (e.g., within initial notification **298** of the security event) may include but is not limited to a JSON portion.

(335) When receiving **1902** an initial notification (e.g., initial notification **298**) of a security event from one of the security-relevant subsystems (e.g., security-relevant subsystems **226**), threat mitigation process **10** may receive **1904** the initial notification (e.g., initial notification **298**) of the security event from an agent (e.g., agent **300**) executed on one of the security-relevant subsystems (e.g., security-relevant subsystems **226**).

(336) In the context of a threat mitigation process **10**, an agent (e.g., agent **300**) may refer to a software component that plays a crucial role in monitoring, detecting, and reporting potential security threats or malicious activities within a computing platform (e.g., computing platform **60**). These agents (e.g., agent **300**) may be deployed across various parts of a computing platform (e.g., computing platform **60**) to ensure comprehensive surveillance and protection.

(337) Functions of Agents (e.g., Agent **300**):

(338) Monitoring Network Traffic: Agents may continuously monitor network traffic for signs of unusual or suspicious behavior. This includes analyzing packets, inspecting protocols, and scrutinizing port activity, among other things. Detection of Anomalies: Agents may use predefined rules or sophisticated algorithms (including machine learning models) to identify deviations from normal network behavior, which could indicate an intrusion or an attempt at one. Log Activity: Agents may log network activity, providing a detailed record of traffic patterns, access attempts, and potentially malicious activities. This information is crucial for forensic analysis and understanding the nature of any attack. Alert Generation: Upon detecting suspicious activities, agents may generate alerts. These alerts can be configured according to severity levels and are sent to administrators or a central monitoring system for further action.

Types of Agents (e.g., Agent **300**): Passive Agents: These agents monitor and analyze network traffic in real-time without interfering with the network's operation. They passively watch for signs of intrusion and report findings to a central system or administrator. Active Agents: In addition to monitoring, active agents can take predefined actions when a threat is detected, such as blocking traffic, isolating affected network segments, or directly interacting with the threat to mitigate its impact.

Deployment Strategies for Agent (e.g., Agent **300**): Host-based Agents: These are installed on individual hosts or devices within the network. They monitor incoming and outgoing traffic from the device, along with system logs and operations, to detect potential intrusions. Network-based Agents: Deployed at strategic points within the network, such as at gateways or along backbone connections, these agents may monitor the flow of data across the network to identify suspicious patterns or anomalies.

Significance of Agents (e.g., Agent **300**): Scalability: Agents may allow a NIDS to scale effectively. By distributing the monitoring load across multiple points in the network, the system can handle large volumes of traffic without significant bottlenecks. Real-time Detection: The real-time monitoring capability of agents enables immediate detection of potential threats, allowing for quicker responses to mitigate damage.

Comprehensive Coverage: Deploying agents across different parts of a network ensures that both internal and external traffic is monitored, providing a more comprehensive defense mechanism against intrusions.

Flexibility: Agents may be tailored to specific network environments and requirements. This includes customizing the detection algorithms, adjusting sensitivity levels, and defining appropriate responses to detected threats.

(339) Generally speaking, agents (e.g., agent **300**) may function as the eyes and ears of threat mitigation

process **10**, providing the essential capabilities needed for the early detection of and response to cybersecurity threats. Their deployment and management may help maintain the integrity and security of networked systems.

(340) Threat mitigation process **10** may iteratively process **1906** the initial notification (e.g., initial notification **298**) using a generative AI model (e.g., generative AI model **302**) and a formatting script (e.g., formatting script **304**) to produce a summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**).

(341) The summarized human-readable report (e.g., summarized human-readable report **306**) may define recommended next steps, recommended actions and/or disclaimers. For example and in response to a security event that is based upon suspicious activity occurring on computing platform **60**: Recommended Next Steps may provide examples of additional investigations that may be implemented (e.g., port analysis/domain owner identification/perpetrator analysis) to further analyze the security event to gauge the risk/severity of the same. Recommended Actions may provide examples of responsive actions that may be implemented (e.g., port blocking/stream shutdown/perpetrator account disablement) to mitigate the negative impact of the security event. Disclaimers may provide explanations for why the suspicious activity of the security event may be benign and occurring for a legitimate (i.e., non-threatening) reason (e.g., such port traffic may occur during weekly backups, the person performing this operation is the president).

(342) As discussed above, a generative AI model (e.g., generative AI model **302**) is a type of artificial intelligence system designed to generate new, synthetic data that resembles its training data. It learns the patterns, features, and distributions of the input data and can produce novel outputs, such as images, text, or sound, that mimic the original dataset. These models are widely used for applications including content creation, data augmentation, and simulation. Examples include Generative Adversarial Networks (GANs) and Variational Autoencoders (VAEs), which have become foundational in fields requiring realistic and diverse data generation.

(343) A formatting script (e.g., formatting script **304**) may include a set of instructions or codes configured to structure, preprocess, or format data (input or output) in a way that's optimal for interaction with or processing by a large language model. This can include tasks like cleaning data, structuring prompts, or formatting the model's outputs for specific applications. The exact nature of formatting script **304** can vary widely depending on the requirements of the task at hand and the specifics of the model's interface.

(344) For example, in a web application that uses a large language model to generate content based on user inputs, a formatting script might: Preprocess User Inputs: Clean and structure user queries into a format that the model can more effectively understand and process. This could involve correcting typos, removing unnecessary punctuation, or structuring the input into a more coherent prompt. Format Model Prompts: Tailor prompts to fit specific use cases or to elicit more accurate responses from the model. This might include adding specific instructions or context to the prompt that guides the model in generating the desired output. Post-Process Model Outputs: Clean or format the text generated by the model to meet user expectations or application requirements. This could involve correcting grammar, structuring the output into a specific format (e.g., HTML, JSON), or truncating responses to fit length constraints. Handle Special Formatting: For certain applications, such as code generation or creating structured data from unstructured text, the script might include rules or templates to format the output in a specific syntax or schema.

(345) These formatting scripts (e.g., formatting script **304**) may help integrate large language models into broader applications or workflows, ensuring that the interaction between human users and the AI is as seamless and effective as possible. Formatting scripts (e.g., formatting script **304**) may be implemented in various programming languages, depending on the environment in which the large language model is being deployed (e.g., Python scripts for a server-side application or JavaScript for client-side processing in a web application).

(346) Accordingly and when iteratively processing **1906** the initial notification (e.g., initial notification **298**) using a generative AI model (e.g., generative AI model **302**) and a formatting script (e.g., formatting script **304**) to produce a summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**), threat mitigation process **10** may iteratively process **1908** the initial notification (e.g., initial notification **298**) using a large language model (e.g., large language model **308**).

(347) As discussed above, a large language model (e.g., large language model **308**) is an advanced artificial intelligence system designed to understand and generate human-like text, which is trained on vast amounts of text data, learning patterns and structures of language. These LLMs can perform various natural language

processing tasks, such as answering questions, generating text, translating languages, and more. LLMs work by processing input text, analyzing it, and generating appropriate responses based on learned patterns and context.

(348) Large language model (e.g., large language model **308**) are a specific subset of generative AI models that focus on understanding, generating, and manipulating natural language text. The relationship between large language models (e.g., large language model **308**) and generative AI models (e.g., generative AI model **302**) can be seen in terms of their foundational technologies, objectives, and the principles they employ to generate new data.

(349) Large language models (e.g., large language model **308**) relate to the broader category of generative AI models (e.g., generative AI model **302**) as follows:

(350) Shared Foundation in Generative Techniques

(351) Generative Principle: At their core, both LLMs and generative AI models are designed to generate new data samples that mimic the distribution of their training data. For generative AI models, this might mean creating new images, music, or text that resemble the original dataset. LLMs specifically focus on generating text that is coherent, contextually relevant, and stylistically similar to the text they were trained on. Modeling Data Distributions: Both LLMs and other generative AI models aim to model the underlying probability distribution of their training data. For LLMs, this involves predicting the likelihood of a sequence of words or tokens based on the vast corpus of text they were trained on. Other generative models, like Generative Adversarial Networks (GANs) or Variational Autoencoders (VAEs), learn to generate data in their respective domains (e.g., images) by modeling the distribution of the training data in those domains.

Use of Deep Learning Architectures Neural Network Architectures: Both LLMs and generative AI models leverage advanced neural network architectures to learn from their training data. Transformers, a type of neural network architecture, have proven particularly effective for LLMs due to their ability to handle long-range dependencies in text. Similarly, GANs utilize a duo of neural networks (generator and discriminator) to generate new data, while VAEs use encoder-decoder architectures for generating data. Advancements in AI: The development and refinement of these neural network architectures have propelled advancements in both fields. Innovations in training techniques, model architecture, and computational efficiency benefit both LLMs and generative AI models across different domains.

Specificity vs. Generality Domain-Specific vs. Domain-Generality: LLMs are domain-specific in that they are tailored for natural language processing tasks. In contrast, the term “generative AI models” encompasses a broader range of models designed for various types of data, including but not limited to text. This generality vs. specificity distinction highlights how LLMs fit within the larger ecosystem of generative AI by applying its principles to the specific domain of language.

Application and Impact Versatile Applications: Both LLMs and generative AI models have wide-ranging applications across industries. LLMs are particularly influential in areas requiring natural language understanding and generation, such as chatbots, content creation, and automated customer service. Other generative AI models find their applications in creating synthetic datasets, enhancing creative design processes, and even drug discovery. Enhancing Human Creativity and Efficiency: Both sets of technologies augment human capabilities by automating creative processes, generating new content, and providing tools for decision-making and analysis.

(352) In conclusion, LLMs are a specialized form of generative AI models with a focus on natural language. They share the foundational approach of learning to generate new data that resembles their training input but apply these principles specifically to the domain of text. This relationship underscores the versatility and breadth of generative AI technologies and their profound impact on both specific industries and broader societal contexts.

(353) Accordingly and when iteratively processing **1906** the initial notification (e.g., initial notification **298**) using a generative AI model (e.g., generative AI model **302**) and a formatting script (e.g., formatting script **304**) to produce a summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**), threat mitigation process **10** may utilize **1910** prompt engineering to produce the summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**).

(354) As is also known in the art, prompt engineering is an essential aspect of working with large language models (e.g., large language model **308**), as it provides a way to guide the AI model's responses and ensure that they are accurate, relevant, and appropriate for the intended application.

(355) In general, prompt engineering involves designing and fine-tuning prompts (e.g., formatting script **304**)

that may be used to train or fine-tune a large language model, such as OpenAI's GPT-3. The prompts (e.g., formatting script **304**) can take a variety of forms, including natural language queries, prompts with specific keywords or phrases, or a combination of both.

(356) The goal of prompt engineering is to create a set of prompts (e.g., formatting script **304**) that are tailored to the specific use case or application, such as generating conversational responses, answering specific questions, or generating creative writing. By designing prompts (e.g., formatting script **304**) that are closely aligned with the intended use case, developers can improve the accuracy and relevance of the model's responses, resulting in more effective and engaging interactions.

(357) Once the prompts (e.g., formatting script **304**) have been designed and fine-tuned, they are used to train or fine-tune the large language model. During the training process, the model is exposed to the prompts (e.g., formatting script **304**) and learns to generate responses that are consistent with the patterns and relationships in the training data. As the model is fine-tuned with additional prompts, its performance improves, allowing it to generate more natural and effective responses over time.

(358) Overall, prompt engineering is an essential aspect of working with large language models (e.g., large language model **308**), as it enables developers to create more accurate and effective natural language processing applications.

(359) When iteratively processing **1906** the initial notification (e.g., initial notification **298**) using a generative AI model (e.g., generative AI model **302**) and a formatting script (e.g., formatting script **304**) to produce a summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**), threat mitigation process **10** may iteratively process **1912** the initial notification (e.g., initial notification **298**) using the generative AI model (e.g., generative AI model **302**), the formatting script (e.g., formatting script **304**) and/or one or more tools (e.g., tools **310**) to produce the summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**).

(360) The one or more tools (e.g., tools **310**) may include one or more of: a decoding tool to decode an encoded initial notification (e.g., initial notification **298**); a decompression tool to decompress a compressed initial notification (e.g., initial notification **298**); and an identification tool to identify an owner of a domain associated with the initial notification (e.g., initial notification **298**).

(361) In the context of managing and responding to security events within a computing platform (e.g., computing platform **60**), decoding tools, decompression tools, and identification tools serve distinct yet complementary purposes. These tools are part of the arsenal used by cybersecurity professionals to analyze, understand, and mitigate security incidents.

(362) Below is an explanation of each tool's purpose: Decoding Tool: Decoding tools are designed to convert data from a coded form into its original form. In the context of a security event, initial notification (e.g., initial notification **298**) may be encoded in a format (e.g., Base64) that is unreadable by threat mitigation process **10** in its native form. Accordingly, threat mitigation process **10** may utilize such a decoding tool to decode such an encoded initial notification. Decompression Tool: Decompression tools are used to expand compressed files back into their original form. In the context of a security event, initial notification (e.g., initial notification **298**) may be compressed in a format (e.g., ZIP, RAR, or custom compression algorithms) that is unreadable by threat mitigation process **10** in its native form. Accordingly, threat mitigation process **10** may utilize such a decompression tool to decompress such an encoded initial notification. Identification Tool: Identification tools concerning domain ownership are utilized to determine the registrants or owners of domains involved in a security event. This can include tools like WHOIS lookups, DNS query tools, or specialized software designed to trace domain affiliations and histories. When a security event involves network communication with suspicious or malicious domains (e.g., for data exfiltration, C2 communication, or phishing), understanding who owns these domains can provide crucial clues about the attackers. This information can help in assessing the credibility and intent behind the domains, tracking the source of the attack, and potentially identifying the attackers or their affiliations. Moreover, it aids in blacklisting domains, strengthening domain reputation checks, and enhancing overall network security posture.

(363) In summary, decoding and decompression tools help cybersecurity teams understand and analyze the content and nature of the threat by revealing the true form of data and files involved in a security event. Identification tools concerning domain ownership extend this analysis by providing insights into the actors behind the threats, enabling more targeted and effective responses. Together, these tools are essential for diagnosing, understanding, and mitigating security incidents in a computer platform (e.g., computing platform **60**).

(364) When iteratively processing **1906** the initial notification (e.g., initial notification **298**) using a generative AI model (e.g., generative AI model **302**) and a formatting script (e.g., formatting script **304**) to produce a summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**), threat mitigation process **10** may utilize **1914** several loops (not shown) and/or nested loops (not shown) to produce the summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**).

(365) In the intricate process of investigating security events on a computing platform (e.g., computing platform **60**), the strategic application of loops and nested loops within an iterative AI process (e.g., generative AI model **302**) proves to be immensely beneficial. These programming constructs allow for the automation of repetitive tasks, crucial in the analysis of vast volumes of network traffic data for potential security threats. A loop facilitates the sequential examination of collected data, enabling the AI system to methodically identify unusual patterns or signatures indicative of malicious activities. The complexity of network security investigations is further addressed through the implementation of nested loops, where a loop is embedded within another, thereby allowing for multi-layered analysis.

(366) For instance, in the detection of security incidents such as distributed denial-of-service (DDOS) attacks, an outer loop could iterate over specific time intervals, scrutinizing traffic data to spot abnormalities in volume that unfold over time. Within each identified time frame, an inner loop could delve deeper, examining individual data packets or sessions for more direct signs of compromise, such as suspicious request frequencies or known malware signatures. This dual-level approach, with an outer loop assessing broader temporal patterns and an inner loop focusing on granular data points, exemplifies the nuanced analysis possible with nested loops.

(367) Such a methodology not only enhances the thoroughness of the security assessment but also significantly accelerates the detection process. By automating the scrutiny of terabytes of network data, AI systems equipped with loop-based algorithms can identify threats with a precision and speed unattainable through manual analysis. The adaptability of loops and nested loops to various levels of data granularity ensures that complex, layered security events are effectively uncovered and addressed. Consequently, the use of iterative loops in AI-driven security event investigation stands as a cornerstone technique in bolstering the defense mechanisms of computer networks against an ever-evolving landscape of cyber threats.

(368) Intelligent Agent

(369) Referring also to FIG. **35**, threat mitigation process **10** may deploy **2000** an agent (e.g., agent **300**) to proactively monitor activity within a computing platform (e.g., computing platform **60**) and generate an initial notification (e.g., initial notification **298**) if a security event is detected.

(370) As discussed above, an agent (e.g., agent **300**) may refer to a software component that plays a crucial role in monitoring, detecting, and reporting potential security threats or malicious activities within a computing platform (e.g., computing platform **60**). These agents (e.g., agent **300**) may be deployed across various parts of a computing platform (e.g., computing platform **60**) to ensure comprehensive surveillance and protection.

(371) As discussed above, examples of such security events may include but are not limited to access auditing; anomalies; authentication; denial of services; exploitation; malware; phishing; spamming; reconnaissance; and/or web attack within a monitored computing platform (e.g., computing platform **60**). An example of the computer-readable language portion (e.g., within the notification of the security event) may include but is not limited to a JSON portion.

(372) As discussed above, the computing platform (e.g., computing platform **60**) may include a plurality of security-relevant subsystems (e.g., security-relevant subsystems **226**). Accordingly and when deploying **2000** an agent (e.g., agent **300**) to proactively monitor activity within a computing platform (e.g., computing platform **60**) and generate an initial notification (e.g., initial notification **298**) if a security event is detected, threat mitigation process **10** may deploy **2002** the agent (e.g., agent **300**) to proactively monitor activity within one or more of the security-relevant subsystems (e.g., security-relevant subsystems **226**) of the computing platform (e.g., computing platform **60**) and generate the initial notification (e.g., initial notification **298**) if the security event is detected.

(373) As discussed above, examples of security-relevant subsystems **226** may include but are not limited to: CDN (i.e., Content Delivery Network) systems; DAM (i.e., Database Activity Monitoring) systems; UBA (i.e., User Behavior Analytics) systems; MDM (i.e., Mobile Device Management) systems; IAM (i.e., Identity and Access Management) systems; DNS (i.e., Domain Name Server) systems, antivirus systems, operating systems, data lakes; data logs; security-relevant software applications; security-relevant hardware



systems; and resources external to the computing platform.

(374) Threat mitigation process **10** may receive **2004** the initial notification (e.g., initial notification **298**) of the security event from the agent (e.g., agent **300**), wherein the initial notification (e.g., initial notification **298**) includes a computer-readable language portion that defines one or more specifics of the security event, (375) As discussed above, examples of such security events may include but are not limited to access auditing; anomalies; authentication; denial of services; exploitation; malware; phishing; spamming; reconnaissance; and/or web attack within a monitored computing platform (e.g., computing platform **60**). An example of the computer-readable language portion (e.g., within the notification of the security event) may include but is not limited to a JSON portion.

(376) In the manner discussed above, threat mitigation process **10** may iteratively process **2006** the initial notification (e.g., initial notification **298**) using a generative AI model (e.g., generative AI model **302**) and a formatting script (e.g., formatting script **304**) to produce a summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**).

(377) As discussed above, the summarized human-readable report (e.g., summarized human-readable report **306**) may define recommended next steps, recommended actions and/or disclaimers. For example and in response to a security event that is based upon suspicious activity occurring on computing platform **60**: Recommended Next Steps may provide examples of additional investigations that may be implemented (e.g., port analysis/domain owner identification/perpetrator analysis) to further analyze the security event to gauge the risk/severity of the same. Recommended Actions may provide examples of responsive actions that may be implemented (e.g., port blocking/stream shutdown/perpetrator account disablement) to mitigate the negative impact of the security event. Disclaimers may provide explanations for why the suspicious activity of the security event may be benign and occurring for a legitimate (i.e., non-threatening) reason (e.g., such port traffic may occur during weekly backups, the person performing this operation is the president.

(378) When iteratively processing **2006** the initial notification (e.g., initial notification **298**) using a generative AI model (e.g., generative AI model **302**) and a formatting script (e.g., formatting script **304**) to produce a summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**), threat mitigation process **10** may iteratively process **2008** the initial notification (e.g., initial notification **298**) using the generative AI model (e.g., generative AI model **302**), the formatting script (e.g., formatting script **304**) and/or one or more tools (e.g., tools **310**) to produce the summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**).

(379) As discussed above, the one or more tools (e.g., tools **310**) may include one or more of: a decoding tool to decode an encoded initial notification (e.g., initial notification **298**); a decompression tool to decompress a compressed initial notification (e.g., initial notification **298**); and an identification tool to identify an owner of a domain associated with the initial notification (e.g., initial notification **298**).

(380) When iteratively processing **2006** the initial notification (e.g., initial notification **298**) using a generative AI model (e.g., generative AI model **302**) and a formatting script (e.g., formatting script **304**) to produce a summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**), threat mitigation process **10** may iteratively process **2010** the initial notification (e.g., initial notification **298**) using a large language model (e.g., large language model **308**).

(381) As discussed above, a large language model (e.g., large language model **308**) is an advanced artificial intelligence system designed to understand and generate human-like text, which is trained on vast amounts of text data, learning patterns and structures of language. These LLMs can perform various natural language processing tasks, such as answering questions, generating text, translating languages, and more. LLMs work by processing input text, analyzing it, and generating appropriate responses based on learned patterns and context.

(382) When iteratively processing **2006** the initial notification (e.g., initial notification **298**) using a generative AI model (e.g., generative AI model **302**) and a formatting script (e.g., formatting script **304**) to produce a summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**), threat mitigation process **10** may utilize **2012** prompt engineering to produce the summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**).

(383) As discussed above, prompt engineering is an essential aspect of working with large language models (e.g., large language model **308**), as it provides a way to guide the AI model's responses and ensure that they

are accurate, relevant, and appropriate for the intended application.

(384) As discussed above and when iteratively processing **2006** the initial notification (e.g., initial notification **298**) using a generative AI model (e.g., generative AI model **302**) and a formatting script (e.g., formatting script **304**) to produce a summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**), threat mitigation process **10** may utilize **2014** several loops and/or nested loops to produce the summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**).

(385) As discussed above, in the intricate process of investigating security events on a computing platform (e.g., computing platform **60**), the strategic application of loops and nested loops within an iterative AI process (e.g., generative AI model **302**) proves to be immensely beneficial. These programming constructs allow for the automation of repetitive tasks, crucial in the analysis of vast volumes of network traffic data for potential security threats. A loop facilitates the sequential examination of collected data, enabling the AI system to methodically identify unusual patterns or signatures indicative of malicious activities. The complexity of network security investigations is further addressed through the implementation of nested loops, where a loop is embedded within another, thereby allowing for multi-layered analysis.

(386) Threat mitigation process **10** may train **2016** the agent (e.g., agent **300**) to proactively monitor activity within a computing platform (e.g., computing platform **60**) and generate an initial notification (e.g., initial notification **298**) if a security event is detected based, at least in part, upon best practices defined via artificial intelligence (e.g., AI/ML process **56**). For example and during the operation of threat mitigation process **10**, data may be archived concerning activities that occurred within the computing platform (e.g., computing platform **60**). So over time, threat mitigation process **10** may build a data repository (e.g., data repository **312**) that identifies various examples of “concerning” activities within the computing platform (e.g., computing platform **60**) and whether those activities resulted in an actual security event or were simply false alarms. Accordingly, threat mitigation process **10** may train **2016** the agent (e.g., agent **300**) to proactively monitor activity within a computing platform (e.g., computing platform **60**) and generate an initial notification (e.g., initial notification **298**) if a security event is detected based, at least in part, upon the information contained within the data repository (e.g., data repository **312**). Additionally/alternatively, threat mitigation process **10** may train **2016** the agent (e.g., agent **300**) to proactively monitor activity within a computing platform (e.g., computing platform **60**) and generate an initial notification (e.g., initial notification **298**) if a security event is detected based, at least in part, upon supplemental information (e.g., supplemental information **314**) obtained from e.g., technical bulletins released by software houses, antivirus providers, hardware manufactures, etc.).

(387) Prompt Engineering

(388) Referring also to FIG. **36**, threat mitigation process **10** may define **2100** a formatting script (e.g., formatting script **304**) for use with a Generative AI model (e.g., generative AI model **302**). An example of such a formatting script (e.g., formatting script **304**) may include but is not limited to a group of one or more prompts that are tailored to the specific use case or application for which the Generative AI model (e.g., generative AI model **302**) is deployed. Specifically, the formatting script (e.g., formatting script **304**) may include one or more discrete instructions for the Generative AI model (e.g., generative AI model **302**) and/or the large language model (e.g., large language model **308**). Such instructions for the Generative AI model (e.g., generative AI model **302**) and/or the large language model (e.g., large language model **308**) may include: formatting instructions and/or content instructions.

(389) As discussed above, these formatting scripts (e.g., formatting script **304**) may help integrate large language models into broader applications or workflows, ensuring that the interaction between human users and the AI is as seamless and effective as possible. Formatting scripts (e.g., formatting script **304**) may be implemented in various programming languages, depending on the environment in which the large language model is being deployed (e.g., Python scripts for a server-side application or JavaScript for client-side processing in a web application).

(390) Threat mitigation process **10** may receive **2102** a notification of a security event, wherein the notification includes a computer-readable language portion that defines one or more specifics of the security event. As discussed above, examples of such security events may include but are not limited to access auditing; anomalies; authentication; denial of services; exploitation; malware; phishing; spamming; reconnaissance; and/or web attack within a monitored computing platform (e.g., computing platform **60**). An example of the computer-readable language portion (e.g., within the notification of the security event) may include but is not limited to a JSON portion.

(391) Below is an example of such a JSON portion:

```
(392) TABLE-US-00001 {  "timestamp": 1676573073400,  "formatVersion": 1,  "webaclId":  
"arn:aws:wafv2::051480436342:icnifuhtyzwa-SharedServices-Policy1606244930846/5a071c76-4c57-4971-  
9326-5c0c8a649b1c",  "terminatingRuleId": "IP-Whitelist-606244930846",  "terminatingRuleType":  
"GROUP",  "action": "ALLOW",  "terminatingRuleMatchDetails": [ ],  "httpSourceName": "ALB",  
  "httpSourceId": "223275863938-app/k8s-toolskon-f53b6065de/888885884d9c7626",  "ruleGroupList":  
[  {  "ruleGroupId": "arn:aws:wafv2::132154534106:nywk0s0jgn37-IP-Whitelist/6f83906e-  
e4c9-4b9e-b4ce-a83633520409",  "terminatingRule": {  "ruleId": "Public-IP-Whitelist",  
    "action": "ALLOW",  "ruleMatchDetails": null  },  
    "nonTerminatingMatchingRules": [ ],  "excludedRules": null,  "customerConfig": null  
  } ],  "rateBasedRuleList": [ ],  "nonTerminatingMatchingRules": [ ],  "requestHeadersInserted":  
null,  "responseCodeSent": null,  "httpRequest": {  "clientIp": "10.142.82.58",  "country":  
"US",  "headers": [  {  "name": "host",  "value": "site.example.com"  },  
    {  "name": "content-encoding",  "value": "snappy"  },  {  
      "name": "content-type",  "value": "application/x-protobuf"  },  {  
        "name": "user-agent",  "value": "GrafanaAgent/v0.26.1"  },  {  
          "name": "x-scope-orgid",  "value": "prod"  },  {  "name":  
"content-length",  "value": "40792"  } ],  "uri": "/api/v1/push",  "args": "",  
    "httpVersion": "HTTP/2.0",  "httpMethod": "POST",  "requestId": "1-63ee7991-  
4fb3b76547a55ccd5badf00d"  },  "oversizeFields": [  "REQUEST_BODY"  ] }
```

(393) Threat mitigation process **10** may process **2104** at least a portion of the computer-readable language portion of the notification using the large language model (e.g., large language model **308**) and the formatting script (e.g., formatting script **304**) to summarize the computer-readable language portion and generate a summarized human-readable report (e.g., summarized human-readable report **306**).

(394) Below is an example of such a summarized human-readable report (e.g., summarized human-readable report **306**):

(395) TABLE-US-00002 Summary & Analysis: At timestamp 1676573073400, the web ACL (arn:aws:wafv2::051480436342:icnifuhtyzwa-SharedServices-Policy1606244930846/5a071c76-4c57-4971-9326-5c0c8a649b1c) allowed an HTTP POST request from external IP 10.142.82.58 (hostname site.example.com, US) to URI '/api/v1/push'. This event could indicate malicious activity as the request includes an API key and the request body is over the size limit. Suggested Legitimate Activity: - Multiple requests sent in a burst - Sending information that is larger than average - Use of an API key Next Steps: - Analyze the source IP address using public resources to identify the owner and location. - Analyze the request body to identify any suspicious or malicious activity, such as attempts to gain access to sensitive information. - Check the headers to verify that the user-agent is legitimate and that the content-type is appropriate for the request.

(396) Threat mitigation process **10** may present **2106** the (above-illustrated) summarized human-readable report (e.g., summarized human-readable report **306**) to a user (e.g., analyst **256**).

(397) Through the use of the above-described formatting script (e.g., formatting script **304**), the above-illustrated summarized human-readable report (e.g., summarized human-readable report **306**) may be concise and easily digestible by the user (e.g., analyst **256**). For example and if the above-illustrated JSON portion was provided to the above-described Generative AI model (e.g., generative AI model **302**) without the above-described formatting script (e.g., formatting script **304**), the result produced would be much less concise and generally less readable.

(398) Below is an example of such a less-concise & less-readable summarized human-readable report (e.g., summarized human-readable report **306**):

```
(399) TABLE-US-00003 ##### Human Readable Output ### WebACL | timestamp | webaclId |  
terminatingRuleId | terminatingRuleType | action | httpSourceName | httpSourceId | | --- | --- | --- | --- | --- | ---  
| --- | | 1676573073400 | arn:aws:wafv2::051480436342:icnifuhtyzwa-SharedServices-  
Policy1606244930846/5a071c76-4c57-4971-9326-5c0c8a649b1c |  
arn:aws:wafv2::132154534106:nywk0s0jgn37-IP-Whitelist/6f83906e-e4c9-4b9e-b4ce-a83633520409 |  
GROUP | ALLOW | ALB | 223275863938-app/k8s-kong-toolskon-f53b6065de/888885884d9c7626 | ###  
Rule Group | ruleGroupId | | --- | | arn:aws:wafv2::132154534106:nywk0s0jgn37-IP-Whitelist/6f83906e-  
e4c9-4b9e-b4ce-a83633520409 | ### Terminating Rule | ruleId | action | | --- | --- | | SNOW-Public-IP-  
Whitelist | ALLOW | ### HTTP Request | clientIp | country | uri | args | httpVersion | httpMethod | requestId |
```

--- | --- | --- | --- | --- | 10.142.82.58 | US | /api/v1/push | | HTTP/2.0 | POST | 1-63ee7991-4fb3b76547a55ccd5badf00d | ### Headers | name | value | | --- | --- | | host | site.example.com | | content-encoding | snappy | content-type | application/x-protobuf | | user-agent | GrafanaAgent/v0.26.1 | | x-prometheus-remote-write-version | 0.1.0 | | x-scope-orgid | prod | | content-length | 40792 |

(400) Threat mitigation process **10** may prompt **2108** a user (e.g., analyst **256**) to provide feedback concerning the (above-illustrated) summarized human-readable report (e.g., summarized human-readable report **306**). And (if provided), threat mitigation process **10** may receive **2110** feedback concerning the summarized human-readable report (e.g., summarized human-readable report **306**) from a user (e.g., analyst **256**). For example, the user (e.g., analyst **256**) may be asked to give “thumbs-up/thumbs-down” feedback concerning the quality of the (above-illustrated) summarized human-readable report (e.g., summarized human-readable report **306**). In the event that the feedback provided is e.g., marginal or poor, threat mitigation process **10** may ask the user (e.g., analyst **256**) to provide additional commentary, examples of which may include but are not limited to: “the summary is too long”, “the summary is too short”, “I would appreciate a more detailed roadmap for remediation”, “more concise language would be helpful”, etc. And (if feedback is provided), threat mitigation process **10** may utilize **2112** the feedback to revise the above-described formatting script (e.g., formatting script **304**) so that the (above-illustrated) summarized human-readable report (e.g., summarized human-readable report **306**) may be tailored based upon such feedback.

(401) Chunking and Recombining to Overcome Token Limits

(402) Referring also to FIG. **37** and as is known in the art, the inputs to (and outputs from) a Generative AI model (e.g., generative AI model **302**) may be limited in scope. Accordingly and if multiple notifications (concerning security events) are received, it is often not practical to have those events simultaneously summarized by such a Generative AI model (e.g., generative AI model **302**). Specifically, large language models (e.g., large language model **308**) often specify such limits based upon a maximum number of tokens.

(403) As is known in the art, the token limits of a large language model (e.g., large language model **308**) refer to the maximum number of words or tokens that the model can process in a single input sequence. The specific token limit of a large language model depends on the architecture and specifications of the model. Depending on the model used, requests can use up to 4097 tokens shared between prompt and completion. If your prompt is 4000 tokens, your completion can be 97 tokens at most. The limit is currently a technical limitation, but there are often creative ways to solve problems within the limit, e.g., condensing your prompt, breaking the text into smaller pieces, etc.

(404) When an input sequence exceeds the token limit of a language model, it needs to be broken up into smaller segments or “chunks” that can be processed separately. This process is known as “chunking” or “windowing”. The chunks are then fed into the model sequentially, and the output from each chunk is combined to produce the final result. Chunking can introduce some challenges, as it requires careful management of the context and flow of the input sequence. In some cases, the output of a previous chunk may need to be taken into account when processing the next chunk, in order to maintain continuity and coherence.

(405) Overall, the token limits of large language models (e.g., large language model **308**) are an important consideration for developers and researchers working with natural language processing applications. By carefully managing the input sequence and chunking appropriately, it is possible to create highly effective and accurate language models that can process very large amounts of text data.

(406) As discussed above, threat mitigation process **10** may define **2200** a formatting script (e.g., formatting script **304**) for use with a Generative AI model (e.g., generative AI model **302**).

(407) As discussed above, these formatting scripts (e.g., formatting script **304**) may help integrate large language models into broader applications or workflows, ensuring that the interaction between human users and the AI is as seamless and effective as possible. Formatting scripts (e.g., formatting script **304**) may be implemented in various programming languages, depending on the environment in which the large language model is being deployed (e.g., Python scripts for a server-side application or JavaScript for client-side processing in a web application).

(408) Threat mitigation process **10** may receive **2202** a plurality of notifications (e.g., initial notification **298** and additional notification **316**) of a security event, wherein each of the plurality of notifications (e.g., initial notification **298** and additional notification **316**) includes a computer-readable language portion that defines one or more specifics of the security event, thus defining a plurality of computer-readable language portions.

(409) As discussed above, examples of such the security events may include but are not limited to access auditing; anomalies; authentication; denial of services; exploitation; malware; phishing; spamming;

reconnection; and/or web attack within computing platform (e.g., computing platform 60). As further discussed above, an example of the computer-readable language portions (e.g., within the notification of the security event) may include but is not limited to a JSON portion.

(410) Assume for the following example that threat mitigation process 10 receives two notifications of a security event.

(411) Below is an example of such a JSON portion for EVENT #1:

(412) TABLE-US-00004 { "eventVersion": "1.08", "userIdentity": { "type": "AssumedRole",  
"principalId": "AYEDVBV3CPSALNLBYZTE6:q5btsdo6lhqv@uyf0bn1fk303.com", "arn":  
"arn:aws:sts::996966753428:assumed-role/DevOps/q5btsdo6lhqv@uyf0bn1fk303.com", "accountId":  
"896966753408", "accessKeyId": "ASIA5B444444FTJIUG", "sessionContext": {  
"sessionIssuer": { "type": "Role", "principalId":  
"ARO5BV3CPAAAAABYZTE6", "arn": "arn:aws:iam::896966753408:role/DevOps",  
"accountId": "123966753123", "userName": "DevOps" },  
"webIdFederationData": { }, "attributes": { "creationDate": "2023-01-  
24T15:47:29Z", "mfaAuthenticated": "false" } }, "invokedBy":  
"amplifybackend.amazonaws.com" }, "eventTime": "2023-01-24T16:53:14Z", "eventSource":  
"iam.amazonaws.com", "eventName": "CreateRole", "awsRegion": "us-east-1", "sourceIPAddress":  
"amplifybackend.amazonaws.com", "userAgent": "amplifybackend.amazonaws.com",  
"requestParameters": { "roleName": "us-east-1\_F4tKzsOrI",  
"assumeRolePolicyDocument": "{ \"Version\": \"2012-10-17\", \"Statement\":  
[ { \"Sid\": \"CognitoAssumeRolePolicy\", \"Effect\": \"Allow\", \"Principal\": { \"Federated\": \"cognito-  
identity.amazonaws.com\" }, \"Action\": \"sts:AssumeRoleWithWebIdentity\", \"Condition\":  
{ \"StringEquals\": { \"cognito-identity.amazonaws.com:aud\": \"us-east-1:62444912-9f39-4eca-f00d-  
5ab4de99b55b\" }, \"ForAnyValue:StringLike\": { \"cognito-  
identity.amazonaws.com:amr\": \"authenticated\" } } ] } \" } \"responseElements\": { \"role\": {  
\"path\": \"/\", \"roleName\": \"us-east-1\_G8tKzsOrI\_Manage-only\", \"roleId\":  
\"ARO5BV3CPSAMOFIYG2AT\", \"arn\": \"arn:aws:iam::896966753408:role/us-east-  
1\_G8tKzsOrI\_Manage-only\", \"createDate\": \"Jan 24, 2023 4:53:14 PM\",  
\"assumeRolePolicyDocument\": \"%7B%22Version%22%3A%222012-10-  
17%22%2C%22Statement%22%3A%5B%7B%22Sid%22%3A%22CognitoAssumeRolePolicy%22%2C%22Effect%  
%22%3A%22Allow%22%2C%22Principal%22%3A%7B%22Federated%22%3A%22cognito-  
identity.amazonaws.com%22%7D%2C%22Action%22%3A%22sts%3AAssumeRoleWithWebIdentity%22%2C%  
%22Condition%22%3A%7B%22StringEquals%22%3A%7B%22cognito-  
identity.amazonaws.com%3Aud%22%3A%22us-east-1%3A62444912-9f39-4eca-f00d-  
5ab4de99b55b%22%7D%2C%22ForAnyValue%3AStringLike%22%3A%7B%22cognito-  
identity.amazonaws.com%3Aamr%22%3A%22authenticated%22%7D%7D%7D%5D%7D\" } }\",  
\"requestID\": \"01dce44c-e2cb-447f-b4df-00d4a3547842\", \"eventID\": \"aaafe757-bb5e-45cf-9f1c-  
6a64f4ee35d2\", \"readOnly\": \"false\", \"eventType\": \"AwsApiCall\", \"managementEvent\": \"true\",  
\"recipientAccountId\": \"896966755608\", \"eventCategory\": \"Management\",  
\"sessionCredentialFromConsole\": \"true\" } }

(413) Below is an example of such a JSON portion for EVENT #2:

(414) TABLE-US-00005 { "eventVersion": "1.08", "userIdentity": { "type": "AssumedRole",  
"principalId": "AYEDVBV3CPSALNLBYZTE6:q5btsdo6lhqv@uyf0bn1fk303.com", "arn":  
"arn:aws:sts::996966753428:assumed-role/DevOps/q5btsdo6lhqv@uyf0bn1fk303.com", "accountId":  
"896966753408", "accessKeyId": "ASIA5B444444FTJIUG", "sessionContext": {  
"sessionIssuer": { "type": "Role", "principalId":  
"ARO5BV3CPAAAAABYZTE6", "arn": "arn:aws:iam::896966753408:role/DevOps",  
"accountId": "123966753123", "userName": "DevOps" },  
"webIdFederationData": { }, "attributes": { "creationDate": "2023-01-  
24T15:47:29Z", "mfaAuthenticated": "false" } }, "invokedBy":  
"amplifybackend.amazonaws.com" "eventTime": "2023-01-24T16:53:14Z", "eventSource":  
"iam.amazonaws.com", "eventName": "CreateRole", "awsRegion": "us-east-1", "sourceIPAddress":  
"amplifybackend.amazonaws.com", "userAgent": "amplifybackend.amazonaws.com",  
"requestParameters": { "roleName": "us-east-1\_F4tKzsOrI", "assumeRolePolicyDocument":  
"{ \"Version\": \"2012-10-17\", \"Statement\":



security event are summarized (as shown above), threat mitigation process **10** may process **2206** at least a portion of each of the plurality of event summaries (illustrated above) using the Generative AI model (e.g., generative AI model **302**) and the above-described formatting script (e.g., formatting script **304**) to summarize the plurality of event summaries and generate a summarized human-readable report (e.g., summarized human-readable report **306**).

(421) Below is an example of such a summarized human-readable report (e.g., summarized human-readable report **306**) for EVENTS 1-2:

(422) TABLE-US-00008 Summary: 1. On Jan 24th, 2023 at 16:52:58Z, a role was created with the name us-east-1\_F4tKzs0rI by user q5btsdo6lhoy@uyf0n1fk303.com with the IP address amplifybackend.amazonaws.com. 2. At 16:53:14 on 2023-01-24, the user with accessKeyId 'ASIA5BV3CPSAFTJIUG' and assumed role AROA5BV3CPSALNLBYZTE6 attempted to PutRolePolicy for a role called 'us-east-1\_G8tKzs0rI\_Manage-only' from the source IP Address 'amplifybackend.amazonaws.com'. Impact on the Organization: The events indicate that a user is attempting to modify a role in the AWS IAM service, which could potentially grant additional privileges to the user and associated IP address. This could lead to unauthorized access to sensitive resources or data, or privilege escalation, resulting in financial loss or other damage to the organization. Relevant Artifacts: • User: q5btsdo6lhqv@uyf0bn1fk303.com • Access Key ID: ASIA5B444444FTJIUG • Assumed Role: AYEDVBV3CPSALNLBYZTE6 • IP Address: amplifybackend.amazonaws.com • Role Name: us-east-1\_G8tKzs0rI-authRole and us-east-1\_G8tKzs0rI-only • Account ID: 896966753408 • Location of External IP: Unknown Potential Security Threats Indicated by Events: The events indicate that a user is attempting to modify a role in the AWS IAM service, which could potentially grant additional privileges to the user and associated IP address. This could lead to unauthorized access to sensitive resources or data, or privilege escalation, resulting in financial loss or other damage to the organization. Indicators of Compromise (IOCs): • User identity associated with accessKeyId 'ASIA5B444444FTJIUG' • Policy document attempted to be applied to role • IP address amplifybackend.amazonaws.com • Unusually high API usage or unsuccessful authentication attempts from user or IP address • Attempts to access sensitive data or modifications to existing policies from user or IP address Legitimate Activity Contributing to False Positives: • Creation of a new role for a legitimate user • Creation of a new role for an application • Creation of a new role for an automated process • Updating the policy on an existing role to allow access to certain resources • Modifying an existing user's permissions • Creating new users or groups • Modifying existing groups or users Next Steps for Further Investigation: • Review the user identity associated with the event and look for suspicious activity that may be associated with the user. • Check for any changes in the IAM role that was created to ensure that it does not provide more access than intended. • Verify that the IP address associated with the event is a trusted source and that no suspicious activity has been observed from that IP in the past. • Look for any other events associated with the user or IP address that may indicate malicious or suspicious activity. • Confirm the identity of the user associated with the accessKeyId 'ASIA5B444444FTJIUG' by checking the IAM user records. • Analyze the policy document to ensure that the new policy does not grant more access than is necessary for the role. • Investigate any suspicious activity that could be associated with the user, such as unusually high API usage or unsuccessful authentication attempts. • Investigate any malicious activity that could be associated with the user, such as attempts to access sensitive data or modifications to existing policies. Recommend Actions: • Selective shutdown / suspension of user account(s). • Selective shutdown of impacted ports. • Selective shutdown of suspicious streams. • Quarantining of inbound file(s).

(423) As discussed above, threat mitigation process **10** may present **2208** the (above-illustrated) summarized human-readable report (e.g., summarized human-readable report **306**) to a user (e.g., analyst **256**) and may prompt **2210** the user (e.g., analyst **256**) to provide feedback concerning the (above-illustrated) summarized human-readable report (e.g., summarized human-readable report **306**).

(424) Threat mitigation process **10** may receive **2212** feedback concerning the (above-illustrated) summarized human-readable report (e.g., summarized human-readable report **306**) from a user (e.g., analyst **256**) and may utilize **2214** the feedback to revise the above-described formatting script (e.g., formatting script **304**) so that the (above-illustrated) summarized human-readable report (e.g., summarized human-readable report **306**) may be tailored based upon such feedback.

(425) Auto-Execution of Recommended Next Steps

(426) Referring also to FIG. **38** and as discussed above, threat mitigation process **10** may establish **2300** connectivity with a plurality of security-relevant subsystems (e.g., security-relevant subsystems **226**) within a computing platform (e.g., computing platform **60**).

(427) As discussed above, establishing connectivity between security-relevant subsystems (e.g., security-relevant subsystems **226**) may require a multifaceted approach that encompasses network configuration, secure communication protocols, authentication, authorization mechanisms, and centralized management. (428) As discussed above, examples of security-relevant subsystems **226** may include but are not limited to: CDN (i.e., Content Delivery Network) systems; DAM (i.e., Database Activity Monitoring) systems; UBA (i.e., User Behavior Analytics) systems; MDM (i.e., Mobile Device Management) systems; IAM (i.e., Identity and Access Management) systems; DNS (i.e., Domain Name Server) systems, antivirus systems, operating systems, data lakes; data logs; security-relevant software applications; security-relevant hardware systems; and resources external to the computing platform.

(429) Threat mitigation process **10** may receive **2302** an initial notification (e.g., initial notification **298**) of a security event from one of the security-relevant subsystems (e.g., security-relevant subsystems **226**), wherein the initial notification (e.g., initial notification **298**) includes a computer-readable language portion that defines one or more specifics of the security event. As discussed above, examples of such security events may include but are not limited to access auditing; anomalies; authentication; denial of services; exploitation; malware; phishing; spamming; reconnaissance; and/or web attack within a monitored computing platform (e.g., computing platform **60**). An example of the computer-readable language portion (e.g., within the notification of the security event) may include but is not limited to a JSON portion.

(430) When receiving **2302** an initial notification (e.g., initial notification **298**) of a security event from one of the security-relevant subsystems (e.g., security-relevant subsystems **226**), threat mitigation process **10** may receive **2304** the initial notification (e.g., initial notification **298**) of the security event from an agent (e.g., agent **300**) executed on one of the security-relevant subsystems (e.g., security-relevant subsystems **226**).

(431) As discussed above, an agent (e.g., agent **300**) may refer to a software component that plays a crucial role in monitoring, detecting, and reporting potential security threats or malicious activities within a computing platform (e.g., computing platform **60**). These agents (e.g., agent **300**) may be deployed across various parts of a computing platform (e.g., computing platform **60**) to ensure comprehensive surveillance and protection.

(432) Threat mitigation process **10** may process **2306** the initial notification (e.g., initial notification **298**) using a generative AI model (e.g., generative AI model **302**) and a formatting script (e.g., formatting script **304**) to produce a summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**), wherein the summarized human-readable report (e.g., summarized human-readable report **306**) defines one or more recommended next steps.

(433) With respect to the (above-illustrated) summarized human-readable report (e.g., summarized human-readable report **306**), examples of one or more recommended next steps defined therein are as follows:

(434) TABLE-US-00009 Next Steps for Further Investigation: • Review the user identity associated with the event and look for suspicious activity that may be associated with the user. • Check for any changes in the IAM role that was created to ensure that it does not provide more access than intended. • Verify that the IP address associated with the event is a trusted source and that no suspicious activity has been observed from that IP in the past. • Look for any other events associated with the user or IP address that may indicate malicious or suspicious activity. • Confirm the identity of the user associated with the accessKeyId 'ASIA5B444444FTJIUG' by checking the IAM user records. • Analyze the policy document to ensure that the new policy does not grant more access than is necessary for the role. • Investigate any suspicious activity that could be associated with the user, such as unusually high API usage or unsuccessful authentication attempts. • Investigate any malicious activity that could be associated with the user, such as attempts to access sensitive data or modifications to existing policies.

(435) When processing **2306** the initial notification (e.g., initial notification **298**) using a generative AI model (e.g., generative AI model **302**) and a formatting script (e.g., formatting script **304**) to produce a summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**), threat mitigation process **10** may process **2308** the initial notification (e.g., initial notification **298**) using the generative AI model (e.g., generative AI model **302**), the formatting script (e.g., formatting script **304**) and/or one or more tools (e.g., tools **310**) to produce the summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**).

(436) As discussed above, the one or more tools (e.g., tools **310**) includes one or more of: a decoding tool to decode an encoded initial notification (e.g., initial notification **298**); a decompression tool to decompress a



compressed initial notification (e.g., initial notification **298**); and an identification tool to identify an owner of a domain associated with the initial notification (e.g., initial notification **298**).

(437) When processing **2306** the initial notification (e.g., initial notification **298**) using a generative AI model (e.g., generative AI model **302**) and a formatting script (e.g., formatting script **304**) to produce a summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**), threat mitigation process **10** may process **2310** the initial notification (e.g., initial notification **298**) using a large language model (e.g., large language model **308**).

(438) As discussed above, a large language model (e.g., large language model **308**) is an advanced artificial intelligence system designed to understand and generate human-like text, which is trained on vast amounts of text data, learning patterns and structures of language. These LLMs can perform various natural language processing tasks, such as answering questions, generating text, translating languages, and more. LLMs work by processing input text, analyzing it, and generating appropriate responses based on learned patterns and context.

(439) When processing **2306** the initial notification (e.g., initial notification **298**) using a generative AI model (e.g., generative AI model **302**) and a formatting script (e.g., formatting script **304**) to produce a summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**), threat mitigation process **10** may utilize **2312** prompt engineering to produce the summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**).

(440) As discussed above, prompt engineering is an essential aspect of working with large language models (e.g., large language model **308**), as it provides a way to guide the AI model's responses and ensure that they are accurate, relevant, and appropriate for the intended application.

(441) When processing **2306** the initial notification (e.g., initial notification **298**) using a generative AI model (e.g., generative AI model **302**) and a formatting script (e.g., formatting script **304**) to produce a summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**), threat mitigation process **10** may utilize **2314** several loops and/or nested loops to produce the summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**).

(442) As discussed above, in the intricate process of investigating security events on a computing platform (e.g., computing platform **60**), the strategic application of loops and nested loops within an iterative AI process (e.g., generative AI model **302**) proves to be immensely beneficial. These programming constructs allow for the automation of repetitive tasks, crucial in the analysis of vast volumes of network traffic data for potential security threats. A loop facilitates the sequential examination of collected data, enabling the AI system to methodically identify unusual patterns or signatures indicative of malicious activities. The complexity of network security investigations is further addressed through the implementation of nested loops, where a loop is embedded within another, thereby allowing for multi-layered analysis.

(443) Threat mitigation process **10** may automatically execute **2316** some or all of the recommended next steps to define one or more recommended actions. Further and when automatically executing **2316** some or all of the recommended next steps to define one or more recommended actions, threat mitigation process **10** may automatically perform **2318** one or more investigative operations concerning the security event.

(444) As discussed above and with respect to the (above-illustrated) summarized human-readable report (e.g., summarized human-readable report **306**), examples of one or more recommended next steps defined therein are as follows:

(445) TABLE-US-00010 Next Steps for Further Investigation: • Review the user identity associated with the event and look for suspicious activity that may be associated with the user. • Check for any changes in the IAM role that was created to ensure that it does not provide more access than intended. • Verify that the IP address associated with the event is a trusted source and that no suspicious activity has been observed from that IP in the past. • Look for any other events associated with the user or IP address that may indicate malicious or suspicious activity. • Confirm the identity of the user associated with the accessKeyId 'ASIA5B444444FTJIUG' by checking the IAM user records. • Analyze the policy document to ensure that the new policy does not grant more access than is necessary for the role. • Investigate any suspicious activity that could be associated with the user, such as unusually high API usage or unsuccessful authentication attempts. • Investigate any malicious activity that could be associated with the user, such as attempts to access sensitive data or modifications to existing policies.

(446) Accordingly, threat mitigation process **10** may automatically execute **2316** some or all of these

recommended next steps to define one or more recommended actions. For example, threat mitigation process **10** may automatically execute **2316** this recommended next step:

(447) TABLE-US-00011 • Review the user identity associated with the event and look for suspicious activity that may be associated with the user

(448) Upon executing **2316** this recommended next step, threat mitigation process **10** may determine that User X is acting in a very suspicious manner. Accordingly, threat mitigation process **10** may automatically perform **2318** one or more investigative operations concerning User X with respect to the security event. For example, threat mitigation process **10** may automatically perform **2318** one or more investigative operations concerning the network usage of User X, the background of User X, the web browsing history of User X, etc. All of this research and investigation may result in threat mitigation process **10** defining the recommended action of disabling all accounts of User X.

(449) Auto-Execution of Recommended Actions

(450) Referring also to FIG. **39**, threat mitigation process **10** may establish **2400** connectivity with a plurality of security-relevant subsystems (e.g., security-relevant subsystems **226**) within a computing platform (e.g., computing platform **60**).

(451) As discussed above, establishing connectivity between security-relevant subsystems (e.g., security-relevant subsystems **226**) may require a multifaceted approach that encompasses network configuration, secure communication protocols, authentication, authorization mechanisms, and centralized management.

(452) As discussed above, examples of security-relevant subsystems **226** may include but are not limited to: CDN (i.e., Content Delivery Network) systems; DAM (i.e., Database Activity Monitoring) systems; UBA (i.e., User Behavior Analytics) systems; MDM (i.e., Mobile Device Management) systems; IAM (i.e., Identity and Access Management) systems; DNS (i.e., Domain Name Server) systems, antivirus systems, operating systems, data lakes; data logs; security-relevant software applications; security-relevant hardware systems; and resources external to the computing platform.

(453) Threat mitigation process **10** may receive **2402** an initial notification (e.g., initial notification **298**) of a security event from one of the security-relevant subsystems (e.g., security-relevant subsystems **226**), wherein the initial notification (e.g., initial notification **298**) includes a computer-readable language portion that defines one or more specifics of the security event. As discussed above, examples of such security events may include but are not limited to access auditing; anomalies; authentication; denial of services; exploitation; malware; phishing; spamming; reconnaissance; and/or web attack within a monitored computing platform (e.g., computing platform **60**). An example of the computer-readable language portion (e.g., within the notification of the security event) may include but is not limited to a JSON portion.

(454) When receiving **2402** an initial notification (e.g., initial notification **298**) of a security event from one of the security-relevant subsystems (e.g., security-relevant subsystems **226**), threat mitigation process **10** may receive **2404** the initial notification (e.g., initial notification **298**) of the security event from an agent (e.g., agent **300**) executed on one of the security-relevant subsystems (e.g., security-relevant subsystems **226**).

(455) As discussed above, an agent (e.g., agent **300**) may refer to a software component that plays a crucial role in monitoring, detecting, and reporting potential security threats or malicious activities within a computing platform (e.g., computing platform **60**). These agents (e.g., agent **300**) may be deployed across various parts of a computing platform (e.g., computing platform **60**) to ensure comprehensive surveillance and protection.

(456) Threat mitigation process **10** may process **2406** the initial notification (e.g., initial notification **298**) using a generative AI model (e.g., generative AI model **302**) and a formatting script (e.g., formatting script **304**) to produce a summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**), wherein the summarized human-readable report (e.g., summarized human-readable report **306**) defines one or more recommended actions.

(457) With respect to the (above-illustrated) summarized human-readable report (e.g., summarized human-readable report **306**), examples of one or more recommended actions defined therein are as follows:

(458) TABLE-US-00012 Recommend Actions: • Selective shutdown / suspension of user account(s). • Selective shutdown of impacted port(s). • Selective shutdown of suspicious stream(s). • Quarantining of inbound file(s).

(459) When processing **2406** the initial notification (e.g., initial notification **298**) using a generative AI model (e.g., generative AI model **302**) and a formatting script (e.g., formatting script **304**) to produce a summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification

(e.g., initial notification **298**), threat mitigation process **10** may process **2408** the initial notification (e.g., initial notification **298**) using the generative AI model (e.g., generative AI model **302**), the formatting script (e.g., formatting script **304**) and/or one or more tools (e.g., tools **310**) to produce the summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**).

(460) As discussed above, the one or more tools (e.g., tools **310**) includes one or more of: a decoding tool to decode an encoded initial notification (e.g., initial notification **298**); a decompression tool to decompress a compressed initial notification (e.g., initial notification **298**); and an identification tool to identify an owner of a domain associated with the initial notification (e.g., initial notification **298**).

(461) When processing **2406** the initial notification (e.g., initial notification **298**) using a generative AI model (e.g., generative AI model **302**) and a formatting script (e.g., formatting script **304**) to produce a summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**), threat mitigation process **10** may process **2410** the initial notification (e.g., initial notification **298**) using a large language model (e.g., large language model **308**).

(462) As discussed above, a large language model (e.g., large language model **308**) is an advanced artificial intelligence system designed to understand and generate human-like text, which is trained on vast amounts of text data, learning patterns and structures of language. These LLMs can perform various natural language processing tasks, such as answering questions, generating text, translating languages, and more. LLMs work by processing input text, analyzing it, and generating appropriate responses based on learned patterns and context.

(463) When processing **2406** the initial notification (e.g., initial notification **298**) using a generative AI model (e.g., generative AI model **302**) and a formatting script (e.g., formatting script **304**) to produce a summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**), threat mitigation process **10** may utilize **2412** prompt engineering to produce the summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**).

(464) As discussed above, prompt engineering is an essential aspect of working with large language models (e.g., large language model **308**), as it provides a way to guide the AI model's responses and ensure that they are accurate, relevant, and appropriate for the intended application.

(465) When processing **2406** the initial notification (e.g., initial notification **298**) using a generative AI model (e.g., generative AI model **302**) and a formatting script (e.g., formatting script **304**) to produce a summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**), threat mitigation process **10** may utilize **2414** several loops and/or nested loops to produce the summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**).

(466) As discussed above, in the intricate process of investigating security events on a computing platform (e.g., computing platform **60**), the strategic application of loops and nested loops within an iterative AI process (e.g., generative AI model **302**) proves to be immensely beneficial. These programming constructs allow for the automation of repetitive tasks, crucial in the analysis of vast volumes of network traffic data for potential security threats. A loop facilitates the sequential examination of collected data, enabling the AI system to methodically identify unusual patterns or signatures indicative of malicious activities. The complexity of network security investigations is further addressed through the implementation of nested loops, where a loop is embedded within another, thereby allowing for multi-layered analysis.

(467) Threat mitigation process **10** may automatically execute **2416** some or all of the recommended actions to address the security event. Further and when automatically executing **2416** some or all of the recommended actions, threat mitigation process **10** may automatically perform **2418** one or more remedial operations concerning the security event.

(468) As discussed above and with respect to the (above-illustrated) summarized human-readable report (e.g., summarized human-readable report **306**), examples of one or more recommended actions defined therein are as follows:

(469) TABLE-US-00013 Recommend Actions: • Selective shutdown / suspension of user account(s). • Selective shutdown of impacted port(s). • Selective shutdown of suspicious stream(s). • Quarantining of inbound file(s).

(470) Accordingly, threat mitigation process **10** may automatically execute **2416** some or all of these recommended actions to address the security event. For example, threat mitigation process **10** may

automatically execute **2416** this recommended action: Selective shutdown of impacted port (471) Upon executing **2416** this recommended action, threat mitigation process **10** may shut down Port A which is receiving data from BlackHat.RU and may shut down Port B which is providing data to BadActor.RU. Further, threat mitigation process **10** may automatically perform **2418** one or more remedial operations concerning the security event. For example, threat mitigation process **10** may automatically delete/quarantine any data that was received on Port A from BlackHat.RU.

(472) Model Registry

(473) Referring also to FIG. **40**, threat mitigation process **10** may maintain **2500** a model repository (e.g., model repository **318**) that defines a plurality of AI models (e.g., plurality of AI models **320**).

(474) Maintaining **2500** a model repository (e.g., model repository **318**) for use by threat mitigation process **10** may involve several activities centered around the creation, storage, management, and updating of AI models that are designed to identify and respond to suspicious or malicious activities within a computing platform (e.g., computing platform **60**). Generally speaking, Network Intrusion Detection Systems equipped with AI capabilities can significantly improve the detection of complex and evolving cyber threats. Here's what maintaining such a repository generally entails:

(475) Maintaining **2500** such a model repository (e.g., model repository **318**) may include various different functionalities, examples of which may include but are not limited to: Model Development and Training: Initially, AI models are developed and trained using historical data, which includes both normal network behavior and various types of intrusions or attacks. This phase involves feature selection, choosing appropriate machine learning algorithms, and training models to recognize patterns indicative of potential security breaches. Model Validation and Testing: Before deployment, models are validated and tested to ensure they accurately detect intrusions while minimizing false positives and false negatives. This step might involve using separate datasets not seen by the model during the training phase to evaluate performance. Repository Storage: The repository (e.g., model repository **318**) acts as a centralized library where these AI models are stored. It includes not only the models themselves but also metadata about the models, such as their type (e.g., decision trees, neural networks), performance metrics, intended use cases (e.g., detecting DDOS attacks, malware), and information on training datasets. Version Control: Similar to software development practices, maintaining a version control system for the AI models is crucial. This ensures that updates, improvements, and changes to the models are systematically managed, allowing for the rollback to previous versions if needed. Model Deployment: Models may be deployed into the operational environment of the NIDS so they can start analyzing network traffic and identifying potential threats. This might involve integrating models into existing NIDS frameworks or updating NIDS components to accommodate new AI capabilities. Monitoring and Updating: Cyber threats are constantly evolving; therefore, AI models require continuous monitoring and retraining to stay effective. This includes updating models with new data reflecting the latest threat patterns and re-deploying them. The repository (e.g., model repository **318**) must support these iterative cycles of retraining and updating. Access Control and Security: Given the sensitivity of the models and the data they process, maintaining proper access control and security measures for the repository (e.g., model repository **318**) is paramount. This ensures that only authorized personnel can access, modify, or deploy models. Compliance and Documentation: Ensuring that the repository (e.g., model repository **318**) and its models comply with relevant regulations and standards, and maintaining thorough documentation for each model may be of paramount importance. This documentation should cover the model's purpose, performance characteristics, training data sources, and any limitations or biases.

(476) By maintaining **2500** an AI model repository (e.g., model repository **318**) for a Network Intrusion Detection System, organizations can systematically manage the lifecycle of AI models (e.g., plurality of AI models **320**), from development to deployment, ensuring that their NIDS remains effective against the continuously changing landscape of network threats.

(477) Threat mitigation process **10** may establish **2502** connectivity with a plurality of security-relevant subsystems (e.g., security-relevant subsystems **226**) within a computing platform (e.g., computing platform **60**).

(478) As discussed above, establishing connectivity between security-relevant subsystems (e.g., security-relevant subsystems **226**) may require a multifaceted approach that encompasses network configuration, secure communication protocols, authentication, authorization mechanisms, and centralized management.

(479) Threat mitigation process **10** may receive **2504** an initial notification (e.g., initial notification **298**) of a security event from one of the security-relevant subsystems (e.g., security-relevant subsystems **226**), wherein the initial notification (e.g., initial notification **298**) includes a computer-readable language portion that

defines one or more specifics of the security event. As discussed above, examples of such security events may include but are not limited to access auditing; anomalies; authentication; denial of services; exploitation; malware; phishing; spamming; reconnaissance; and/or web attack within a monitored computing platform (e.g., computing platform **60**). An example of the computer-readable language portion (e.g., within the notification of the security event) may include but is not limited to a JSON portion.

(480) When receiving **2504** an initial notification (e.g., initial notification **298**) of a security event from one of the security-relevant subsystems (e.g., security-relevant subsystems **226**), threat mitigation process **10** may receive **2506** the initial notification (e.g., initial notification **298**) of the security event from an agent (e.g., agent **300**) executed on one of the security-relevant subsystems (e.g., security-relevant subsystems **226**).

(481) As discussed above, an agent (e.g., agent **300**) may refer to a software component that plays a crucial role in monitoring, detecting, and reporting potential security threats or malicious activities within a computing platform (e.g., computing platform **60**). These agents (e.g., agent **300**) may be deployed across various parts of a computing platform (e.g., computing platform **60**) to ensure comprehensive surveillance and protection.

(482) Threat mitigation process **10** may select **2508** a generative AI model (e.g., generative AI model **302**) for processing the initial notification (e.g., initial notification **298**) of the security event from the plurality of AI models (e.g., plurality of AI models **320**) defined within the model repository (e.g., model repository **318**), thus defining a selected generative AI model (e.g., generative AI model **302**).

(483) Examples of the plurality of AI models (e.g., plurality of AI models **320**) defined within the model repository (e.g., model repository **318**) may include but are not limited to: BERT (Bidirectional Encoder Representations from Transformers): Developed by Google, BERT is a powerful natural language processing model that has been influential in various NLP tasks, including question answering and sentiment analysis. OpenAI's GPT (Generative Pre-trained Transformer) Series: This includes GPT-2, GPT-3, GPT-4 and potentially future iterations. These models are developed by OpenAI and are known for their ability to generate human-like text across a wide range of topics. XLNet: Developed by Google, XLNet is a generalized autoregressive pretraining method that outperforms BERT on several NLP benchmarks. T5 (Text-to-Text Transfer Transformer): Also developed by Google, T5 is a versatile model capable of performing various NLP tasks by converting all tasks into a text-to-text format. BERT-based models from Hugging Face: Hugging Face provides pre-trained BERT-based models like ROBERTa, DistilBERT, and BERTweet, which are widely used in the NLP community. Microsoft's Turing Natural Language Generation (T-NLG): T-NLG is a large-scale AI language model developed by Microsoft Research, which competes in the domain of natural language generation and understanding. Facebook's ROBERTa (Robustly optimized BERT approach): ROBERTa is an optimized BERT model developed by Facebook AI Research, which achieves better performance on various NLP benchmarks. Tencent's ERNIE (Enhanced Representation through kNowledge Integration): ERNIE is a knowledge-enhanced language representation model developed by Tencent AI Lab, which integrates external knowledge for better understanding. Fast.ai's ULMFIT (Universal Language Model Fine-Tuning): ULMFIT is a transfer learning method developed by Fast.ai, which enables easy fine-tuning of pre-trained language models for specific tasks with limited data. Salesforce's CTRL (Conditional Transformer Language Model): CTRL is a large-scale autoregressive language model developed by Salesforce Research, which allows users to control the topic of the generated text.

(484) The plurality of AI models (e.g., plurality of AI models **320**) defined within the model repository (e.g., model repository **318**) may include multiple versions of the same model (e.g., ChatGPT 3.0 versus ChatGPT 3.5 versus ChatGPT 4.0) . . . wherein such different versions provide different levels of performance/operating cost.

(485) Accordingly, the plurality of AI models (e.g., plurality of AI models **320**) defined within the model repository (e.g., model repository **318**) may offer e.g., different features, operate on different cost structures or perform certain operations more efficiently. Therefore, threat mitigation process **10** may select **2508** a generative AI model (e.g., generative AI model **302**) from the plurality of AI models (e.g., plurality of AI models **320**) defined within the model repository (e.g., model repository **318**) based upon operation requirements. For example, Model A may be very fast and quite expensive to operate. However, it may be very skilled at generating synthetic speech. Accordingly, threat mitigation process **10** may select **2508** Model A when realistic synthetic speech is needed. Conversely, Model B may be slower and less expensive to operate. But it may be really good at translating text between languages. Accordingly, threat mitigation

process **10** may select **2508** Model B when translations are needed at a more leisurely pace.

(486) Threat mitigation process **10** may process **2510** the initial notification (e.g., initial notification **298**) using the selected generative AI model (e.g., generative AI model **302**) and a formatting script (e.g., formatting script **304**) to produce a summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**).

(487) When processing **2510** the initial notification (e.g., initial notification **298**) using the selected generative AI model (e.g., generative AI model **302**) and a formatting script (e.g., formatting script **304**) to produce a summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**), threat mitigation process **10** may process **2512** the initial notification (e.g., initial notification **298**) using the selected generative AI model (e.g., generative AI model **302**), the formatting script (e.g., formatting script **304**) and/or one or more tools (e.g., tools **310**) to produce the summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**).

(488) As discussed above, the one or more tools (e.g., tools **310**) includes one or more of: a decoding tool to decode an encoded initial notification (e.g., initial notification **298**); a decompression tool to decompress a compressed initial notification (e.g., initial notification **298**); and an identification tool to identify an owner of a domain associated with the initial notification (e.g., initial notification **298**).

(489) When processing **2510** the initial notification (e.g., initial notification **298**) using the selected generative AI model (e.g., generative AI model **302**) and a formatting script (e.g., formatting script **304**) to produce a summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**), threat mitigation process **10** may process **2514** the initial notification (e.g., initial notification **298**) using a large language model (e.g., large language model **308**).

(490) As discussed above, a large language model (e.g., large language model **308**) is an advanced artificial intelligence system designed to understand and generate human-like text, which is trained on vast amounts of text data, learning patterns and structures of language. These LLMs can perform various natural language processing tasks, such as answering questions, generating text, translating languages, and more. LLMs work by processing input text, analyzing it, and generating appropriate responses based on learned patterns and context.

(491) When processing **2510** the initial notification (e.g., initial notification **298**) using the selected generative AI model (e.g., generative AI model **302**) and a formatting script (e.g., formatting script **304**) to produce a summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**), threat mitigation process **10** may utilize **2516** prompt engineering to produce the summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**).

(492) As discussed above, prompt engineering is an essential aspect of working with large language models (e.g., large language model **308**), as it provides a way to guide the AI model's responses and ensure that they are accurate, relevant, and appropriate for the intended application.

(493) When processing **2510** the initial notification (e.g., initial notification **298**) using the selected generative AI model (e.g., generative AI model **302**) and a formatting script (e.g., formatting script **304**) to produce a summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**), threat mitigation process **10** may utilize **2518** several loops and/or nested loops to produce the summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**).

(494) As discussed above, in the intricate process of investigating security events on a computing platform (e.g., computing platform **60**), the strategic application of loops and nested loops within an iterative AI process (e.g., generative AI model **302**) proves to be immensely beneficial. These programming constructs allow for the automation of repetitive tasks, crucial in the analysis of vast volumes of network traffic data for potential security threats. A loop facilitates the sequential examination of collected data, enabling the AI system to methodically identify unusual patterns or signatures indicative of malicious activities. The complexity of network security investigations is further addressed through the implementation of nested loops, where a loop is embedded within another, thereby allowing for multi-layered analysis.

(495) Dynamic Decision Making

(496) Referring also to FIG. **41**, threat mitigation process **10** may establish **2600** connectivity with a plurality of security-relevant subsystems (e.g., security-relevant subsystems **226**) within a computing platform (e.g., computing platform **60**).

(497) As discussed above, establishing connectivity between security-relevant subsystems (e.g., security-relevant subsystems **226**) may require a multifaceted approach that encompasses network configuration, secure communication protocols, authentication, authorization mechanisms, and centralized management. (498) As discussed above, examples of security-relevant subsystems **226** may include but are not limited to: CDN (i.e., Content Delivery Network) systems; DAM (i.e., Database Activity Monitoring) systems; UBA (i.e., User Behavior Analytics) systems; MDM (i.e., Mobile Device Management) systems; IAM (i.e., Identity and Access Management) systems; DNS (i.e., Domain Name Server) systems, antivirus systems, operating systems, data lakes; data logs; security-relevant software applications; security-relevant hardware systems; and resources external to the computing platform.

(499) Threat mitigation process **10** may receive **2602** an initial notification (e.g., initial notification **298**) of a security event from one of the security-relevant subsystems (e.g., security-relevant subsystems **226**), wherein the initial notification (e.g., initial notification **298**) includes a computer-readable language portion that defines one or more specifics of the security event. As discussed above, examples of such security events may include but are not limited to access auditing; anomalies; authentication; denial of services; exploitation; malware; phishing; spamming; reconnaissance; and/or web attack within a monitored computing platform (e.g., computing platform **60**). An example of the computer-readable language portion (e.g., within the notification of the security event) may include but is not limited to a JSON portion.

(500) When receiving **2602** an initial notification (e.g., initial notification **298**) of a security event from one of the security-relevant subsystems (e.g., security-relevant subsystems **226**), threat mitigation process **10** may receive **2604** the initial notification (e.g., initial notification **298**) of the security event from an agent (e.g., agent **300**) executed on one of the security-relevant subsystems (e.g., security-relevant subsystems **226**).

(501) As discussed above, an agent (e.g., agent **300**) may refer to a software component that plays a crucial role in monitoring, detecting, and reporting potential security threats or malicious activities within a computing platform (e.g., computing platform **60**). These agents (e.g., agent **300**) may be deployed across various parts of a computing platform (e.g., computing platform **60**) to ensure comprehensive surveillance and protection.

(502) Threat mitigation process **10** may process **2606** the initial notification (e.g., initial notification **298**) using a generative AI model (e.g., generative AI model **302**) and a formatting script (e.g., formatting script **304**) to define one or more recommended actions.

(503) As discussed above and with respect to the (above-illustrated) summarized human-readable report (e.g., summarized human-readable report **306**), examples of one or more recommended actions defined therein are as follows:

(504) TABLE-US-00014 Recommend Actions: • Selective shutdown / suspension of user account(s). • Selective shutdown of impacted port(s). • Selective shutdown of suspicious stream(s). • Quarantining of inbound file(s).

(505) When processing **2606** the initial notification (e.g., initial notification **298**) using a generative AI model (e.g., generative AI model **302**) and a formatting script (e.g., formatting script **304**) to define one or more recommended actions, threat mitigation process **10** may process **2608** the initial notification (e.g., initial notification **298**) using the generative AI model (e.g., generative AI model **302**), the formatting script (e.g., formatting script **304**) and/or one or more tools (e.g., tools **310**) to define one or more recommended actions for the initial notification (e.g., initial notification **298**).

(506) As discussed above, the one or more tools (e.g., tools **310**) includes one or more of: a decoding tool to decode an encoded initial notification (e.g., initial notification **298**); a decompression tool to decompress a compressed initial notification (e.g., initial notification **298**); and an identification tool to identify an owner of a domain associated with the initial notification (e.g., initial notification **298**).

(507) When processing **2606** the initial notification (e.g., initial notification **298**) using a generative AI model (e.g., generative AI model **302**) and a formatting script (e.g., formatting script **304**) to define one or more recommended actions, threat mitigation process **10** may process **2610** the initial notification (e.g., initial notification **298**) using a large language model (e.g., large language model **308**).

(508) As discussed above, a large language model (e.g., large language model **308**) is an advanced artificial intelligence system designed to understand and generate human-like text, which is trained on vast amounts of text data, learning patterns and structures of language. These LLMs can perform various natural language processing tasks, such as answering questions, generating text, translating languages, and more. LLMs work by processing input text, analyzing it, and generating appropriate responses based on learned patterns and

context.

(509) When processing **2606** the initial notification (e.g., initial notification **298**) using a generative AI model (e.g., generative AI model **302**) and a formatting script (e.g., formatting script **304**) to define one or more recommended actions, threat mitigation process **10** may utilize **2612** prompt engineering to define one or more recommended actions for the initial notification (e.g., initial notification **298**).

(510) As discussed above, prompt engineering is an essential aspect of working with large language models (e.g., large language model **308**), as it provides a way to guide the AI model's responses and ensure that they are accurate, relevant, and appropriate for the intended application.

(511) When processing **2606** the initial notification (e.g., initial notification **298**) using a generative AI model (e.g., generative AI model **302**) and a formatting script (e.g., formatting script **304**) to define one or more recommended actions, threat mitigation process **10** may utilize **2614** several loops and/or nested loops to define one or more recommended actions for the initial notification (e.g., initial notification **298**).

(512) As discussed above, in the intricate process of investigating security events on a computing platform (e.g., computing platform **60**), the strategic application of loops and nested loops within an iterative AI process (e.g., generative AI model **302**) proves to be immensely beneficial. These programming constructs allow for the automation of repetitive tasks, crucial in the analysis of vast volumes of network traffic data for potential security threats. A loop facilitates the sequential examination of collected data, enabling the AI system to methodically identify unusual patterns or signatures indicative of malicious activities. The complexity of network security investigations is further addressed through the implementation of nested loops, where a loop is embedded within another, thereby allowing for multi-layered analysis.

(513) Threat mitigation process **10** may automatically generate **2616** a playbook (e.g., playbook **322**) to effectuate at least one of the above-discussed recommended actions. The playbook (e.g., playbook **322**) may define a set of procedures and/or guidelines configured to at least partially address the security event.

(514) In the context of a Network Intrusion Detection System (NIDS) and broader cybersecurity operations, a playbook (e.g., playbook **322**) refers to a predefined set of procedures or steps that are to be followed in response to specific types of alerts or indicators of compromise. These playbooks (e.g., playbook **322**) may be essential for ensuring that an organization's response to potential threats is swift, effective, and consistent.

(515) Examples of the roles and benefits of playbooks (e.g., playbook **322**) in a NIDS context are as follows: Standardizing Response Procedures: Playbooks provide a standardized method for responding to different types of security incidents. This standardization helps in minimizing errors and ensures that all necessary steps are taken to mitigate and analyze the threat. Automating Response Actions: Many modern NIDS and Security Orchestration, Automation, and Response (SOAR) platforms allow for the automation of certain playbook actions. For example, a playbook might automatically isolate a compromised system from the network, update firewall rules to block malicious traffic, or gather additional context about an alert without human intervention. Facilitating Quick Decision-Making: By having a set of predetermined actions, playbooks enable security analysts to make quick decisions in response to detected threats. This is crucial in minimizing the time an attacker has inside the network and reducing the potential damage they can cause. Enhancing Incident Management: Playbooks help in organizing the workflow of incident response, from initial detection to post-incident analysis. This includes specifying roles and responsibilities, documenting actions taken, and ensuring compliance with regulatory requirements. Improving Training and Readiness: Playbooks are also valuable training tools for security teams. They help in familiarizing new analysts with the typical response processes and can be used in tabletop exercises to simulate responses to hypothetical security incidents. Evolving with Threat Landscape: As new types of attacks emerge and organizations' network environments change, playbooks must be regularly updated. This ensures that the response strategies remain effective against the latest threats and are aligned with the current network architecture and business processes.

(516) In summary, playbooks (e.g., playbook **322**) in a Network Intrusion Detection System context may be critical for managing and responding to security incidents efficiently. They help in minimizing the impact of attacks, ensuring compliance with regulatory standards, and maintaining the overall security posture of an organization.

(517) When automatically generating **2616** a playbook (e.g., playbook **322**) to effectuate at least one of the recommended actions, threat mitigation process **10** may automatically generate **2618** a playbook (e.g., playbook **322**) based, at least in part, upon best practices defined via artificial intelligence (e.g., AI/ML process **56**).

(518) For example and during the operation of threat mitigation process **10**, data may be archived concerning



activities that occurred within the computing platform (e.g., computing platform **60**). So over time, threat mitigation process **10** may build a data repository (e.g., data repository **312**) that identifies various examples of “concerning” activities within the computing platform (e.g., computing platform **60**), the procedures employed to address these “concerning” activities, and whether such procedures were successful. Accordingly, threat mitigation process **10** may automatically generate **2618** a playbook (e.g., playbook **322**) based, at least in part, upon best practices extracted from data repository **312** via artificial intelligence (e.g., AI/ML process **56**). Accordingly and through the use of threat mitigation process **10**, playbooks need not be static and may be dynamic . . . wherein threat mitigation process **10** may automatically generate **2618** playbook **322** based, at least in part, upon best practices defined via artificial intelligence (e.g., AI/ML process **56**).

(519) Threat mitigation process **10** may process **2620** the playbook (e.g., playbook **322**) to address at least a portion of the security event, wherein processing **2620** the playbook (e.g., playbook **322**) to address at least a portion of the security event may include performing **2622** the set of procedures and/or guidelines defined within the playbook (e.g., playbook **322**). Examples of such procedures and/or guidelines defined within the playbook (e.g., playbook **322**) may include but are not limited to:

(520) TABLE-US-00015 • Selective shutdown / suspension of user account(s). • Selective shutdown of impacted port(s). • Selective shutdown of suspicious stream(s). • Quarantining of inbound file(s).

Adaptive Defense

(521) Referring also to FIG. **42**, threat mitigation process **10** may generate **2700** one or more detection rules (e.g., detection rules **324**) that are indicative of a security event, wherein the one or more detection rules are based upon historical suspect activity and/or historical security events. As discussed above, examples of such security events may include but are not limited to access auditing; anomalies; authentication; denial of services; exploitation; malware; phishing; spamming; reconnaissance; and/or web attack within a monitored computing platform (e.g., computing platform **60**).

(522) As discussed above and during the operation of threat mitigation process **10**, data may be archived concerning activities that occurred within the computing platform (e.g., computing platform **60**). So over time, threat mitigation process **10** may build a data repository (e.g., data repository **312**) that identifies various examples of “concerning” activities within the computing platform (e.g., computing platform **60**), the procedures employed to address these “concerning” activities, and whether such procedures were successful. Accordingly, threat mitigation process **10** may generate **2700** such detection rules (e.g., detection rules **324**) that are indicative of a security event based upon historical suspect activity and/or historical security events defined within data repository **312**.

(523) Threat mitigation process **10** may monitor **2702** activity within a computing platform (e.g., computing platform **60**), thus defining monitored activity (e.g., monitored activity **326**).

(524) The computing platform (e.g., computing platform **60**) may include a plurality of security-relevant subsystems (e.g., security-relevant subsystems **226**).

(525) As discussed above, examples of security-relevant subsystems **226** may include but are not limited to: CDN (i.e., Content Delivery Network) systems; DAM (i.e., Database Activity Monitoring) systems; UBA (i.e., User Behavior Analytics) systems; MDM (i.e., Mobile Device Management) systems; IAM (i.e., Identity and Access Management) systems; DNS (i.e., Domain Name Server) systems, antivirus systems, operating systems, data lakes; data logs; security-relevant software applications; security-relevant hardware systems; and resources external to the computing platform.

(526) Accordingly and when monitoring **2702** activity within a computing platform (e.g., computing platform **60**), threat mitigation process **10** may monitor **2704** activity within one or more of the plurality of security-relevant subsystems (e.g., security-relevant subsystems **226**) of the computing platform (e.g., computing platform **60**).

(527) Threat mitigation process **10** may compare **2706** such monitored activity (e.g., monitored activity **326**) to the one or more detection rules (e.g., detection rules **324**) to determine if such monitored activity (e.g., monitored activity **326**) includes suspect activity indicative of a security event.

(528) As discussed above, examples of such security events may include but are not limited to access auditing; anomalies; authentication; denial of services; exploitation; malware; phishing; spamming; reconnaissance; and/or web attack within a monitored computing platform (e.g., computing platform **60**).

(529) Threat mitigation process **10** may generate **2708** an initial notification (e.g., initial notification **298**) of the security event, wherein the initial notification (e.g., initial notification **298**) includes a computer-readable language portion that defines one or more specifics of the security event. An example of the computer-

readable language portion (e.g., within the notification of the security event) may include but is not limited to a JSON portion.

(530) Threat mitigation process **10** may iteratively process **2710** the initial notification (e.g., initial notification **298**) using a generative AI model (e.g., generative AI model **302**) and a formatting script (e.g., formatting script **304**) to produce a summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**).

(531) As discussed above, the summarized human-readable report (e.g., summarized human-readable report **306**) may define recommended next steps, recommended actions and/or disclaimers. For example and in response to a security event that is based upon suspicious activity occurring on computing platform **60**: Recommended Next Steps may provide examples of additional investigations that may be implemented (e.g., port analysis/domain owner identification/perpetrator analysis) to further analyze the security event to gauge the risk/severity of the same. Recommended Actions may provide examples of responsive actions that may be implemented (e.g., port blocking/stream shutdown/perpetrator account disablement) to mitigate the negative impact of the security event. Disclaimers may provide explanations for why the suspicious activity of the security event may be benign and occurring for a legitimate (i.e., non-threatening) reason (e.g., such port traffic may occur during weekly backups, the person performing this operation is the president).

(532) When iteratively processing **2710** the initial notification (e.g., initial notification **298**) using a generative AI model (e.g., generative AI model **302**) and a formatting script (e.g., formatting script **304**) to produce a summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**), threat mitigation process **10** may iteratively process **2712** the initial notification (e.g., initial notification **298**) using the generative AI model (e.g., generative AI model **302**), the formatting script (e.g., formatting script **304**) and/or one or more tools (e.g., tools **310**) to produce the summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**).

(533) As discussed above, the one or more tools (e.g., tools **310**) includes one or more of: a decoding tool to decode an encoded initial notification (e.g., initial notification **298**); a decompression tool to decompress a compressed initial notification (e.g., initial notification **298**); and an identification tool to identify an owner of a domain associated with the initial notification (e.g., initial notification **298**).

(534) When iteratively processing **2710** the initial notification (e.g., initial notification **298**) using a generative AI model (e.g., generative AI model **302**) and a formatting script (e.g., formatting script **304**) to produce a summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**), threat mitigation process **10** may iteratively process **2714** the initial notification (e.g., initial notification **298**) using a large language model (e.g., large language model **308**).

(535) As discussed above, a large language model (e.g., large language model **308**) is an advanced artificial intelligence system designed to understand and generate human-like text, which is trained on vast amounts of text data, learning patterns and structures of language. These LLMs can perform various natural language processing tasks, such as answering questions, generating text, translating languages, and more. LLMs work by processing input text, analyzing it, and generating appropriate responses based on learned patterns and context.

(536) When iteratively processing **2710** the initial notification (e.g., initial notification **298**) using a generative AI model (e.g., generative AI model **302**) and a formatting script (e.g., formatting script **304**) to produce a summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**), threat mitigation process **10** may utilize **2716** prompt engineering to produce the summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**).

(537) As discussed above, prompt engineering is an essential aspect of working with large language models (e.g., large language model **308**), as it provides a way to guide the AI model's responses and ensure that they are accurate, relevant, and appropriate for the intended application.

(538) When iteratively processing **2710** the initial notification (e.g., initial notification **298**) using a generative AI model (e.g., generative AI model **302**) and a formatting script (e.g., formatting script **304**) to produce a summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**), threat mitigation process **10** may utilize **2718** several loops and/or nested loops to produce the summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**).

(539) As discussed above, in the intricate process of investigating security events on a computing platform (e.g., computing platform **60**), the strategic application of loops and nested loops within an iterative AI process (e.g., generative AI model **302**) proves to be immensely beneficial. These programming constructs allow for the automation of repetitive tasks, crucial in the analysis of vast volumes of network traffic data for potential security threats. A loop facilitates the sequential examination of collected data, enabling the AI system to methodically identify unusual patterns or signatures indicative of malicious activities. The complexity of network security investigations is further addressed through the implementation of nested loops, where a loop is embedded within another, thereby allowing for multi-layered analysis.

(540) Threat mitigation process **10** may update **2720** the one or more detection rules (e.g., detection rules **324**) based upon current suspect activity, current security events, future suspect activity and/or future security events.

(541) As discussed above and as threat mitigation process **10** continues to operate, data may continue to be archived concerning activities that occurred within the computing platform (e.g., computing platform **60**). And as time continues to pass, threat mitigation process **10** may continue to build a data repository (e.g., data repository **312**) that identifies various examples of “concerning” activities within the computing platform (e.g., computing platform **60**), the procedures employed to address these “concerning” activities, and whether such procedures were successful. Accordingly, threat mitigation process **10** may update **2720** the one or more detection rules (e.g., detection rules **324**) based upon current suspect activity, current security events, future suspect activity and/or future security events.

(542) Next Generation Risk Modeling

(543) Referring also to FIG. **43**, threat mitigation process **10** may monitor **2800** activity within a computing platform (e.g., computing platform **60**), thus defining monitored activity (e.g., monitored activity **326**).

(544) As discussed above, examples of security-relevant subsystems **226** may include but are not limited to: CDN (i.e., Content Delivery Network) systems; DAM (i.e., Database Activity Monitoring) systems; UBA (i.e., User Behavior Analytics) systems; MDM (i.e., Mobile Device Management) systems; IAM (i.e., Identity and Access Management) systems; DNS (i.e., Domain Name Server) systems, antivirus systems, operating systems, data lakes; data logs; security-relevant software applications; security-relevant hardware systems; and resources external to the computing platform.

(545) The computing platform (e.g., computing platform **60**) may include a plurality of security-relevant subsystems (e.g., security-relevant subsystems **226**).

(546) As discussed above, examples of security-relevant subsystems **226** may include but are not limited to: CDN (i.e., Content Delivery Network) systems; DAM (i.e., Database Activity Monitoring) systems; UBA (i.e., User Behavior Analytics) systems; MDM (i.e., Mobile Device Management) systems; IAM (i.e., Identity and Access Management) systems; DNS (i.e., Domain Name Server) systems, antivirus systems, operating systems, data lakes; data logs; security-relevant software applications; security-relevant hardware systems; and resources external to the computing platform.

(547) Accordingly and when monitoring **2800** activity within a computing platform (e.g., computing platform **60**), threat mitigation process **10** may monitor **2802** activity within one or more of the plurality of security-relevant subsystems (e.g., security-relevant subsystems **226**) of the computing platform (e.g., computing platform **60**).

(548) Threat mitigation process **10** may associate **2804** the monitored activity (e.g., monitored activity **326**) with a user of the computing platform (e.g., computing platform **60**), thus defining an associated user (e.g., associated user **328**).

(549) Threat mitigation process **10** may assign **2806** a risk level to the monitored activity (e.g., monitored activity **326**) to determine if such monitored activity (e.g., monitored activity **326**) is indicative of a security event, wherein the assigned risk level is based, at least in part, upon the associated user (e.g., associated user **328**). Accordingly, if the associated user (e.g., associated user **328**) is the owner of the company, the assigned risk level may be reduced due to the position of associated user **328**. Conversely, if the associated user (e.g., associated user **328**) is a new hire of the company (or someone who has shown questionable judgement in the past), the assigned risk level may be increased.

(550) As discussed above, examples of such security events may include but are not limited to access auditing; anomalies; authentication; denial of services; exploitation; malware; phishing; spamming; reconnaissance; and/or web attack within a monitored computing platform (e.g., computing platform **60**).

(551) If such monitored activity (e.g., monitored activity **326**) is indicative of a security event, threat mitigation process **10** may generate **2808** an initial notification (e.g., initial notification **298**) of the security

event, wherein the initial notification (e.g., initial notification **298**) includes a computer-readable language portion that defines one or more specifics of the security event. An example of the computer-readable language portion (e.g., within the notification of the security event) may include but is not limited to a JSON portion.

(552) Threat mitigation process **10** may iteratively process **2810** the initial notification (e.g., initial notification **298**) using a generative AI model (e.g., generative AI model **302**) and a formatting script (e.g., formatting script **304**) to produce a summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**).

(553) When iteratively processing **2810** the initial notification (e.g., initial notification **298**) using a generative AI model (e.g., generative AI model **302**) and a formatting script (e.g., formatting script **304**) to produce a summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**), threat mitigation process **10** may iteratively process **2812** the initial notification (e.g., initial notification **298**) using the generative AI model (e.g., generative AI model **302**), the formatting script (e.g., formatting script **304**) and/or one or more tools (e.g., tools **310**) to produce the summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**).

(554) As discussed above, the one or more tools (e.g., tools **310**) includes one or more of: a decoding tool to decode an encoded initial notification (e.g., initial notification **298**); a decompression tool to decompress a compressed initial notification (e.g., initial notification **298**); and an identification tool to identify an owner of a domain associated with the initial notification (e.g., initial notification **298**).

(555) When iteratively processing **2810** the initial notification (e.g., initial notification **298**) using a generative AI model (e.g., generative AI model **302**) and a formatting script (e.g., formatting script **304**) to produce a summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**), threat mitigation process **10** may iteratively process **2814** the initial notification (e.g., initial notification **298**) using a large language model (e.g., large language model **308**).

(556) As discussed above, a large language model (e.g., large language model **308**) is an advanced artificial intelligence system designed to understand and generate human-like text, which is trained on vast amounts of text data, learning patterns and structures of language. These LLMs can perform various natural language processing tasks, such as answering questions, generating text, translating languages, and more. LLMs work by processing input text, analyzing it, and generating appropriate responses based on learned patterns and context.

(557) When iteratively processing **2810** the initial notification (e.g., initial notification **298**) using a generative AI model (e.g., generative AI model **302**) and a formatting script (e.g., formatting script **304**) to produce a summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**), threat mitigation process **10** may utilize **2816** prompt engineering to produce the summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**).

(558) As discussed above, prompt engineering is an essential aspect of working with large language models (e.g., large language model **308**), as it provides a way to guide the AI model's responses and ensure that they are accurate, relevant, and appropriate for the intended application.

(559) When iteratively processing **2810** the initial notification (e.g., initial notification **298**) using a generative AI model (e.g., generative AI model **302**) and a formatting script (e.g., formatting script **304**) to produce a summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**), threat mitigation process **10** may utilize **2818** several loops and/or nested loops to produce the summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**).

(560) As discussed above, in the intricate process of investigating security events on a computing platform (e.g., computing platform **60**), the strategic application of loops and nested loops within an iterative AI process (e.g., generative AI model **302**) proves to be immensely beneficial. These programming constructs allow for the automation of repetitive tasks, crucial in the analysis of vast volumes of network traffic data for potential security threats. A loop facilitates the sequential examination of collected data, enabling the AI system to methodically identify unusual patterns or signatures indicative of malicious activities. The complexity of network security investigations is further addressed through the implementation of nested loops, where a loop is embedded within another, thereby allowing for multi-layered analysis.

(561) AI-Based Threat Mitigation Platform

(562) Referring also to FIG. 44, there is shown threat mitigation platform **2900**. Threat mitigation platform **2900** may include an agent subsystem (e.g., an agent subsystem **2902**) configured to generate an initial notification (e.g., initial notification **298**) concerning a security event within a computing platform (e.g., computing platform **60**).

(563) As discussed above, examples of such security events may include but are not limited to access auditing; anomalies; authentication; denial of services; exploitation; malware; phishing; spamming; reconnaissance; and/or web attack within a monitored computing platform (e.g., computing platform **60**).

(564) The threat mitigation platform (e.g., threat mitigation platform **2900**) may include a generative AI-based planner subsystem (e.g., generative AI-based planner subsystem **2904**) configured to receive the initial notification (e.g., initial notification **298**) and generate a mitigation plan (e.g., mitigation plan **2906**) to address, in whole or in part, the security event within the computing platform (e.g., computing platform **60**).

(565) The generative AI-based planner subsystem (e.g., generative AI-based planner subsystem **2904**) may be configured to utilize one or more tools (e.g., tools **310**) available via tool kit **2908** to process the initial notification (e.g., initial notification **298**).

(566) As discussed above, the one or more tools (e.g., tools **310**) utilized by generative AI-based planner subsystem **2904** includes one or more of: a decoding tool to decode an encoded initial notification (e.g., initial notification **298**); a decompression tool to decompress a compressed initial notification (e.g., initial notification **298**); and an identification tool to identify an owner of a domain associated with the initial notification (e.g., initial notification **298**).

(567) The threat mitigation platform (e.g., threat mitigation platform **2900**) may include an executor subsystem (e.g., executor subsystem **2910**) configured to iteratively process the mitigation plan (e.g., mitigation plan **2906**) using a generative AI model (e.g., generative AI model **302**) to generate an output (e.g., output **2912**).

(568) The executor subsystem (e.g., executor subsystem **2910**) may be configured to utilize one or more tools (e.g., tools **310**) available via tool kit **2908** to process the mitigation plan (e.g., mitigation plan **2906**).

(569) As discussed above, the one or more tools (e.g., tools **310**) utilized by the executor subsystem **2910** includes one or more of: a decoding tool to decode an encoded initial notification (e.g., initial notification **298**); a decompression tool to decompress a compressed initial notification (e.g., initial notification **298**); and an identification tool to identify an owner of a domain associated with the initial notification (e.g., initial notification **298**).

(570) The executor subsystem (e.g., executor subsystem **2910**) may be configured to utilize several loops and/or nested loops to generate the output (e.g., output **2912**).

(571) As discussed above, in the intricate process of investigating security events on a computing platform (e.g., computing platform **60**), the strategic application of loops and nested loops within an iterative AI process (e.g., generative AI model **302**) proves to be immensely beneficial. These programming constructs allow for the automation of repetitive tasks, crucial in the analysis of vast volumes of network traffic data for potential security threats. A loop facilitates the sequential examination of collected data, enabling the AI system to methodically identify unusual patterns or signatures indicative of malicious activities. The complexity of network security investigations is further addressed through the implementation of nested loops, where a loop is embedded within another, thereby allowing for multi-layered analysis.

(572) The threat mitigation platform (e.g., threat mitigation platform **2900**) may include an output formatter subsystem (e.g., output formatter subsystem **2914**) configured to format the output (e.g., output **2912**) and generate a summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**).

(573) The output formatter subsystem (e.g., output formatter subsystem **2914**) may be configured to utilize a large language model (e.g., large language model **308**) to generate the summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**).

(574) As discussed above, a large language model (e.g., large language model **308**) is an advanced artificial intelligence system designed to understand and generate human-like text, which is trained on vast amounts of text data, learning patterns and structures of language. These LLMs can perform various natural language processing tasks, such as answering questions, generating text, translating languages, and more. LLMs work by processing input text, analyzing it, and generating appropriate responses based on learned patterns and context.

(575) The output formatter subsystem (e.g., output formatter subsystem **2914**) may be configured to utilize a

formatting script (e.g., formatting script **304**) to generate the summarized human-readable report (e.g., summarized human-readable report **306**) for the initial notification (e.g., initial notification **298**).

(576) As discussed above, the summarized human-readable report (e.g., summarized human-readable report **306**) may define recommended next steps, recommended actions and/or disclaimers. For example and in response to a security event that is based upon suspicious activity occurring on computing platform **60**: Recommended Next Steps may provide examples of additional investigations that may be implemented (e.g., port analysis/domain owner identification/perpetrator analysis) to further analyze the security event to gauge the risk/severity of the same. Recommended Actions may provide examples of responsive actions that may be implemented (e.g., port blocking/stream shutdown/perpetrator account disablement) to mitigate the negative impact of the security event. Disclaimers may provide explanations for why the suspicious activity of the security event may be benign and occurring for a legitimate (i.e., non-threatening) reason (e.g., such port traffic may occur during weekly backups, the person performing this operation is the president).

#### GENERAL

(577) As will be appreciated by one skilled in the art, the present disclosure may be embodied as a method, a system, or a computer program product. Accordingly, the present disclosure may take the form of an entirely hardware embodiment, an entirely software embodiment (including firmware, resident software, micro-code, etc.) or an embodiment combining software and hardware aspects that may all generally be referred to herein as a “circuit,” “module” or “system.” Furthermore, the present disclosure may take the form of a computer program product on a computer-usable storage medium having computer-usable program code embodied in the medium.

(578) Any suitable computer usable or computer readable medium may be utilized. The computer-usable or computer-readable medium may be, for example but not limited to, an electronic, magnetic, optical, electromagnetic, infrared, or semiconductor system, apparatus, device, or propagation medium. More specific examples (a non-exhaustive list) of the computer-readable medium may include the following: an electrical connection having one or more wires, a portable computer diskette, a hard disk, a random access memory (RAM), a read-only memory (ROM), an erasable programmable read-only memory (EPROM or Flash memory), an optical fiber, a portable compact disc read-only memory (CD-ROM), an optical storage device, a transmission media such as those supporting the Internet or an intranet, or a magnetic storage device. The computer-usable or computer-readable medium may also be paper or another suitable medium upon which the program is printed, as the program can be electronically captured, via, for instance, optical scanning of the paper or other medium, then compiled, interpreted, or otherwise processed in a suitable manner, if necessary, and then stored in a computer memory. In the context of this document, a computer-usable or computer-readable medium may be any medium that can contain, store, communicate, propagate, or transport the program for use by or in connection with the instruction execution system, apparatus, or device. The computer-usable medium may include a propagated data signal with the computer-usable program code embodied therewith, either in baseband or as part of a carrier wave. The computer usable program code may be transmitted using any appropriate medium, including but not limited to the Internet, wireline, optical fiber cable, RF, etc.

(579) Computer program code for carrying out operations of the present disclosure may be written in an object-oriented programming language such as Java, Smalltalk, C++ or the like. However, the computer program code for carrying out operations of the present disclosure may also be written in conventional procedural programming languages, such as the “C” programming language or similar programming languages. The program code may execute entirely on the user's computer, partly on the user's computer, as a stand-alone software package, partly on the user's computer and partly on a remote computer or entirely on the remote computer or server. In the latter scenario, the remote computer may be connected to the user's computer through a local area network/a wide area network/the Internet (e.g., network **14**).

(580) The present disclosure is described with reference to flowchart illustrations and/or block diagrams of methods, apparatus (systems) and computer program products according to embodiments of the disclosure. It will be understood that each block of the flowchart illustrations and/or block diagrams, and combinations of blocks in the flowchart illustrations and/or block diagrams, may be implemented by computer program instructions. These computer program instructions may be provided to a processor of a general-purpose computer/special purpose computer/other programmable data processing apparatus, such that the instructions, which execute via the processor of the computer or other programmable data processing apparatus, create means for implementing the functions/acts specified in the flowchart and/or block diagram block or blocks.

(581) These computer program instructions may also be stored in a computer-readable memory that may direct a computer or other programmable data processing apparatus to function in a particular manner, such that the instructions stored in the computer-readable memory produce an article of manufacture including instruction means which implement the function/act specified in the flowchart and/or block diagram block or blocks.

(582) The computer program instructions may also be loaded onto a computer or other programmable data processing apparatus to cause a series of operational steps to be performed on the computer or other programmable apparatus to produce a computer-implemented process such that the instructions which execute on the computer or other programmable apparatus provide steps for implementing the functions/acts specified in the flowchart and/or block diagram block or blocks.

(583) The flowcharts and block diagrams in the figures may illustrate the architecture, functionality, and operation of possible implementations of systems, methods and computer program products according to various embodiments of the present disclosure. In this regard, each block in the flowchart or block diagrams may represent a module, segment, or portion of code, which comprises one or more executable instructions for implementing the specified logical function(s). It should also be noted that, in some alternative implementations, the functions noted in the block may occur out of the order noted in the figures. For example, two blocks shown in succession may, in fact, be executed substantially concurrently, or the blocks may sometimes be executed in the reverse order, depending upon the functionality involved. It will also be noted that each block of the block diagrams and/or flowchart illustrations, and combinations of blocks in the block diagrams and/or flowchart illustrations, may be implemented by special purpose hardware-based systems that perform the specified functions or acts, or combinations of special purpose hardware and computer instructions.

(584) The terminology used herein is for the purpose of describing particular embodiments only and is not intended to be limiting of the disclosure. As used herein, the singular forms “a”, “an” and “the” are intended to include the plural forms as well, unless the context clearly indicates otherwise. It will be further understood that the terms “comprises” and/or “comprising,” when used in this specification, specify the presence of stated features, integers, steps, operations, elements, and/or components, but do not preclude the presence or addition of one or more other features, integers, steps, operations, elements, components, and/or groups thereof.

(585) The corresponding structures, materials, acts, and equivalents of all means or step plus function elements in the claims below are intended to include any structure, material, or act for performing the function in combination with other claimed elements as specifically claimed. The description of the present disclosure has been presented for purposes of illustration and description, but is not intended to be exhaustive or limited to the disclosure in the form disclosed. Many modifications and variations will be apparent to those of ordinary skill in the art without departing from the scope and spirit of the disclosure. The embodiment was chosen and described in order to best explain the principles of the disclosure and the practical application, and to enable others of ordinary skill in the art to understand the disclosure for various embodiments with various modifications as are suited to the particular use contemplated.

(586) A number of implementations have been described. Having thus described the disclosure of the present application in detail and by reference to embodiments thereof, it will be apparent that modifications and variations are possible without departing from the scope of the disclosure defined in the appended claims.

## Claims

1. A computer-implemented method executed on a computing device comprising: establishing connectivity with a plurality of security-relevant subsystems within a computing platform; receiving an initial notification of a security event from one of the security-relevant subsystems, wherein the initial notification includes a computer-readable language portion that defines one or more specifics of the security event; processing the initial notification using a generative AI model and a formatting script and one or more tools to produce a summarized human-readable report for the initial notification, wherein the summarized human-readable report defines one or more recommended next steps, wherein the one or more tools include one or more of: a decompression tool to decompress a compressed initial notification; and an identification tool to identify an owner of a domain associated with the initial notification; and automatically executing some or all of the recommended next steps to define one or more recommended actions.

2. The computer-implemented method of claim 1 wherein receiving an initial notification of a security event from one of the security-relevant subsystems includes: receiving the initial notification of the security event from an agent executed on one of the security-relevant subsystems.
3. The computer-implemented method of claim 1 wherein the one or more tools includes a decoding tool to decode an encoded initial notification.
4. The computer-implemented method of claim 1 wherein processing the initial notification using a generative AI model and a formatting script to produce a summarized human-readable report for the initial notification includes: processing the initial notification using a large language model.
5. The computer-implemented method of claim 1 wherein processing the initial notification using a generative AI model and a formatting script to produce a summarized human-readable report for the initial notification includes: utilizing prompt engineering to produce the summarized human-readable report for the initial notification.
6. The computer-implemented method of claim 1 wherein processing the initial notification using a generative AI model and a formatting script to produce a summarized human-readable report for the initial notification includes: utilizing one or more of several loops and nested loops to produce the summarized human-readable report for the initial notification.
7. The computer-implemented method of claim 1 wherein automatically executing some or all of the recommended next steps to define one or more recommended actions includes: automatically performing one or more investigative operations concerning the security event.
8. A computer program product residing on a non-transitory computer readable medium having a plurality of instructions stored thereon which, when executed by a processor, cause the processor to perform operations comprising: establishing connectivity with a plurality of security-relevant subsystems within a computing platform; receiving an initial notification of a security event from one of the security-relevant subsystems, wherein the initial notification includes a computer-readable language portion that defines one or more specifics of the security event; processing the initial notification using a generative AI model and a formatting script and one or more tools to produce a summarized human-readable report for the initial notification, wherein the summarized human-readable report defines one or more recommended next steps, wherein the one or more tools include one or more of: a decompression tool to decompress a compressed initial notification; and an identification tool to identify an owner of a domain associated with the initial notification; and automatically executing some or all of the recommended next steps to define one or more recommended actions.
9. The computer program product of claim 8 wherein receiving an initial notification of a security event from one of the security-relevant subsystems includes: receiving the initial notification of the security event from an agent executed on one of the security-relevant subsystems.
10. The computer program product of claim 8 wherein the one or more tools includes a decoding tool to decode an encoded initial notification.
11. The computer program product of claim 8 wherein processing the initial notification using a generative AI model and a formatting script to produce a summarized human-readable report for the initial notification includes: processing the initial notification using a large language model.
12. The computer program product of claim 8 wherein processing the initial notification using a generative AI model and a formatting script to produce a summarized human-readable report for the initial notification includes: utilizing prompt engineering to produce the summarized human-readable report for the initial notification.
13. The computer program product of claim 8 wherein processing the initial notification using a generative AI model and a formatting script to produce a summarized human-readable report for the initial notification includes: utilizing one or more of several loops and nested loops to produce the summarized human-readable report for the initial notification.
14. The computer program product of claim 8 wherein automatically executing some or all of the recommended next steps to define one or more recommended actions includes: automatically performing one or more investigative operations concerning the security event.
15. A computing system including a processor and memory configured to perform operations comprising: establishing connectivity with a plurality of security-relevant subsystems within a computing platform; receiving an initial notification of a security event from one of the security-relevant subsystems, wherein the initial notification includes a computer-readable language portion that defines one or more specifics of the security event; processing the initial notification using a generative AI model and a formatting script and one



or more tools to produce a summarized human-readable report for the initial notification, wherein the summarized human-readable report defines one or more recommended next steps, wherein the one or more tools include one or more of: a decompression tool to decompress a compressed initial notification; and an identification tool to identify an owner of a domain associated with the initial notification; and automatically executing some or all of the recommended next steps to define one or more recommended actions.

16. The computing system of claim 15 wherein receiving an initial notification of a security event from one of the security-relevant subsystems includes: receiving the initial notification of the security event from an agent executed on one of the security-relevant subsystems.

17. The computing system of claim 15 wherein the one or more tools includes a decoding tool to decode an encoded initial notification.

18. The computing system of claim 15 wherein processing the initial notification using a generative AI model and a formatting script to produce a summarized human-readable report for the initial notification includes: processing the initial notification using a large language model.

19. The computing system of claim 15 wherein processing the initial notification using a generative AI model and a formatting script to produce a summarized human-readable report for the initial notification includes: utilizing prompt engineering to produce the summarized human-readable report for the initial notification.

20. The computing system of claim 15 wherein processing the initial notification using a generative AI model and a formatting script to produce a summarized human-readable report for the initial notification includes: utilizing one or more of several loops and nested loops to produce the summarized human-readable report for the initial notification.

21. The computing system of claim 15 wherein automatically executing some or all of the recommended next steps to define one or more recommended actions includes: automatically performing one or more investigative operations concerning the security event.

---