

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12388908
Kind Code	B2
Date of Patent	August 12, 2025
Inventor(s)	Thompson; Cheryl

Cache retrieval based on tiered data

Abstract

A system, method, and machine-readable storage medium for retrieving data are provided. In some embodiments, a cache may receive a request for data from a client. The cache may determine that a first subset of the data is stored on a storage device and that a second subset of the data is stored at a cloud address located at a cloud storage endpoint. The cache may also receive from the storage device the first subset of data. The cache further receives from the cloud storage endpoint the second subset of data in response to transmitting a request for the second subset of data stored at the cloud address to the cloud storage endpoint. The cache then transmits to the client the first and second subsets of data from the various sources in response to the data request.

Inventors:	Thompson; Cheryl (Sunnyvale, CA)
Applicant:	NetApp, Inc. (San Jose, CA)
Family ID:	1000008749827
Assignee:	NETAPP, INC. (San Jose, CA)
Appl. No.:	18/617363
Filed:	March 26, 2024

Prior Publication Data

Document Identifier	Publication Date
US 20240305696 A1	Sep. 12, 2024

Related U.S. Application Data

continuation parent-doc US 16825946 20200320 US 11943314 child-doc US 18617363

Publication Classification

Int. Cl.: H04L67/568 (20220101); G06F3/06 (20060101); G06F12/0866 (20160101)

U.S. Cl.:

CPC H04L67/568 (20220501); G06F3/0605 (20130101); G06F3/0638 (20130101); G06F3/067 (20130101); G06F12/0866 (20130101);

Field of Classification Search

CPC: H04L (67/2842); G06F (3/0605); G06F (3/0638); G06F (3/067); G06F (12/0866)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
8832330	12/2013	Lancaster	N/A	N/A
10042768	12/2017	Karnowski et al.	N/A	N/A
10462012	12/2018	Rao	N/A	G06F 3/0604
11943314	12/2023	Thompson	N/A	N/A
2005/0210203	12/2004	Stenstrom	N/A	N/A
2005/0262103	12/2004	Stakutis et al.	N/A	N/A
2014/0032850	12/2013	Phelan	711/135	G06F 12/0873
2015/0254196	12/2014	Sukonik et al.	N/A	N/A
2017/0235485	12/2016	Davis et al.	N/A	N/A
2019/0129855	12/2018	Kelly	N/A	N/A
2021/0096997	12/2020	Krasner	N/A	G06F 3/0656

OTHER PUBLICATIONS

Notice of Allowance mailed on Nov. 17, 2023 for U.S. Appl. No. 16/825,946, filed Mar. 20, 2020, 10 pages. cited by applicant

Primary Examiner: Thammavong; Prasith

Attorney, Agent or Firm: HAYNES AND BOONE, LLP

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS (1) The present application is a continuation of U.S. patent application Ser. No. 16/825,946 filed Mar. 20, 2020, the disclosure of which is hereby incorporated by reference in its entirety.

TECHNICAL FIELD

(1) The present description relates to caching objects, and more specifically, to a system, method, and machine-readable storage medium for retrieving objects in a distributed storage system for cost savings, efficiency, and/or load balancing.

BACKGROUND

(2) Networks and distributed storage allow data and storage space to be shared between devices located anywhere a connection is available. These implementations may range from a single machine offering a shared drive over a home network to an enterprise-class cloud storage array

with multiple copies of data distributed throughout the world. Larger implementations may incorporate Network Attached Storage (NAS) devices, Storage Area Network (SAN) devices, and other configurations of storage elements and controllers to provide data and manage its flow. Storage nodes or servers in storage networks generally store data across multiple data storage devices that together include a data container, which may also be referred to herein as an aggregate. Storage nodes may employ various forms of local data storage devices, such as hard disk drives, solid state drives, flash drives, or tape devices, as well as remote data storage devices, such as cloud storage devices or repositories, for example.

(3) The data storage devices may host one or more data stores or volumes within the aggregates, which are associated with file systems that define an overall logical arrangement of storage space in a storage network. Volumes can be established to have tiering policies to store data in multiple performance tiers. For example, composite aggregates can include volumes that span relatively high-performance flash data storage devices and relatively low performance cloud storage devices for warmer and colder data, respectively.

(4) To improve performance of a storage network, a cache volume located on a storage node may act as a cache for an origin volume that is located on a different storage node. The cache volume may be used to, for example, speed access to remote data and/or to offload traffic from heavily accessed volumes. For example, a cache volume may directly serve read requests if it contains the data requested by a client. If the cache volume does not contain the requested data, the cache volume may request the data from the origin volume and store the data in the cache volume before serving the client request. Subsequent read requests for the data may then be served directly from the cache volume. While caching data can reduce latency for responding to client requests, current approaches have limitations in their ability for reducing latency for clients storing content in the distributed storage system.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) The present disclosure is best understood from the following detailed description when read with the accompanying figures.

(2) FIG. 1 is a schematic diagram of a computing architecture according to one or more aspects of the present disclosure.

(3) FIG. 2 is an illustrative example of a data storage system in accordance with one or more aspects of the present disclosure.

(4) FIG. 3 is a schematic diagram of a cache volume that retrieves data from an origin volume and data from a cloud storage endpoint to satisfy a single data request according to one or more aspects of the present disclosure.

(5) FIG. 4 is a signaling diagram illustrating a method of retrieving data from an origin volume and data from a cloud storage endpoint to satisfy a single data request according to one or more aspects of the present disclosure.

(6) FIG. 5 is a flow diagram of a method of retrieving data from an origin volume and a from cloud storage endpoint according to one or more aspects of the present disclosure.

DETAILED DESCRIPTION

(7) All examples and illustrative references are non-limiting and should not be used to limit the claims to specific implementations and embodiments described herein and their equivalents. For simplicity, reference numbers may be repeated between various examples. This repetition is for clarity only and does not dictate a relationship between the respective embodiments, unless noted otherwise. Finally, in view of this disclosure, features described in relation to one aspect or embodiment may be applied to other disclosed aspects or embodiments of the disclosure, even

though not specifically shown in the drawings or described in the text.

(8) Customers may want to store data from a small edge data center into the cloud and then access that data from one or more of their other data centers. As the pool of data in the cloud grows, interest in leveraging it for a variety of projects across a geographically distributed organization may grow. To allow customers to have easy and fast access to their data, caches may be provided for prioritizing retention of the working dataset (e.g., most recently used data). A cache may work in operation with a data fabric technology that enables automated tiering of data to low-cost object storage tiers either on or off premises.

(9) If the cache receives a request for data and the data is not stored in the cache (e.g., cache miss), then the cache may fetch the data from an origin source. The origin source may be part of the data fabric technology that enables data tiering out to the cloud. If a portion of the requested data is tiered out the cloud, the origin source may read the data from the cloud and then return the data to the cache. If a first subset of the requested data is stored at the origin source and a second subset of the requested data is stored at the cloud storage endpoint, then it may be desirable to allow the cache to retrieve the first subset from the origin source and the second subset from the cloud storage endpoint. Accordingly, the cache may access and retrieve the data from multiple locations (e.g., the cloud storage endpoint and the origin volume).

(10) The present application provides techniques for a cache to retrieve data from the origin volume/or and from the cloud storage endpoint to satisfy a single data request. An infrastructure that would enable such data retrieval may provide improved performance by providing load distribution, reduced latency by locating data closer to the point of client access, and/or enhanced availability by serving cached data in a network disconnection situation.

(11) FIG. 1 is a block diagram illustrating a clustered network environment **100** in accordance with one or more aspects of the present disclosure. The clustered network environment **100** includes data storage systems **102** and **104** that are coupled over a cluster fabric **106**, such as a computing network embodied as a private InfiniBand, Fiber Channel (FC), or Ethernet network facilitating communication between the data storage systems **102** and **104** (and one or more modules, components, etc. therein, such as, nodes **116** and **118**, for example). The data storage systems **102** and **104** may be computing devices that interact with other components via, for example, the cluster fabric **106**. It will be appreciated that while two data storage systems **102** and **104** and nodes **116** and **118** are illustrated in FIG. 1, any suitable number of such components is contemplated.

(12) In an example, nodes **116**, **118** include storage controllers (e.g., node **116** may include a primary or local storage controller and node **118** may include a secondary or remote storage controller) that provide client devices, such as host devices **108** and **110**, with access to data stored within data storage devices **128** and **130**. Similarly, unless specifically provided otherwise herein, the same is true for other modules, elements, features, items, etc. referenced herein and/or illustrated in the accompanying drawings. That is, a particular number of components, modules, elements, features, items, etc. disclosed herein is not meant to be interpreted in a limiting manner.

(13) It will be further appreciated that clustered networks are not limited to any particular geographic areas and can be clustered locally and/or remotely. Thus, in an embodiment a clustered network can be distributed over a plurality of storage systems and/or nodes located in a plurality of geographic locations; while in an embodiment a clustered network can include data storage systems (e.g., **102**, **104**) residing in a same geographic location (e.g., in a single onsite rack of data storage devices).

(14) In the example illustrated in FIG. 1, one or more host devices **108**, **110** which may include, for example, client devices, personal computers (PCs), computing devices used for storage (e.g., storage servers), and other computers or peripheral devices (e.g., printers), are coupled to the respective data storage systems **102**, **104** by storage network connections **112**, **114**. A network connection **112**, **114** may include a local area network (LAN) or wide area network (WAN), for example, that utilizes Network Attached Storage (NAS) protocols, such as a Common Internet File

System (CIFS) protocol or a Network File System (NFS) protocol to exchange data packets, a Storage Area Network (SAN) protocol, such as Small Computer System Interface (SCSI) or Fiber Channel Protocol (FCP), an object protocol, such as AMAZON S3®, etc.

(15) The host devices **108, 110** may be general-purpose computers running applications, and may interact with the data storage systems **102, 104** using a client/server model for exchange of information. For example, the host device **108** may request data from the data storage system **102, 104** (e.g., data on a storage device managed by a network storage control configured to process I/O commands issued by the host device for the storage device), and the data storage system **102, 104** may return results of the request to the host device via the storage network connection **112, 114**.

(16) The nodes **116, 118** on clustered data storage systems **102, 104** may include network or host nodes that are interconnected as a cluster to provide data storage and management services, such as to an enterprise having remote locations, cloud storage (e.g., a cloud storage endpoint **160** may be stored within a data cloud), etc., for example. Such a node in the clustered network environment **100** may be a device attached to the network as a connection point, redistribution point or communication endpoint, for example. A node may send, receive, and/or forward information over a network communications channel, and may include any device that meets any or all of these criteria. An example of a node may be a data storage and management server attached to a network, where the server may include a general purpose computer or a computing device particularly configured to operate as a server in a data storage and management system.

(17) In an example, a first cluster of nodes such as the nodes **116, 118** (e.g., a first set of storage controllers configured to provide access to a first storage aggregate including a first logical grouping of one or more storage devices) may be located on a first storage site. A second cluster of nodes, not illustrated, may be located at a second storage site (e.g., a second set of storage controllers configured to provide access to a second storage aggregate including a second logical grouping of one or more storage devices). The first cluster of nodes and the second cluster of nodes may be configured according to a disaster recovery configuration where a surviving cluster of nodes provides switchover access to storage devices of a disaster cluster of nodes in the event a disaster occurs at a disaster storage site comprising the disaster cluster of nodes (e.g., the first cluster of nodes provides client devices with switchover data access to storage devices of the second storage aggregate in the event a disaster occurs at the second storage site).

(18) As illustrated in the clustered network environment **100**, nodes **116, 118** may include various functional components that coordinate to provide a distributed storage architecture for the cluster. For example, the nodes may include network modules **120, 122** and disk modules **124, 126**. The network modules **120, 122** may be configured to allow the nodes **116, 118** (e.g., network storage controllers) to connect with host devices **108, 110** over the storage network connections **112, 114**, for example, allowing the host devices **108, 110** to access data stored in the distributed storage system. Further, the network modules **120, 122** may provide connections with one or more other components through the cluster fabric **106**. For example, in FIG. 1, the network module **120** of the node **116** may access a second data storage device by sending a request through the disk module **126** of the node **118**.

(19) Disk modules **124, 126** may be configured to connect one or more data storage devices **128, 130**, such as disks or arrays of disks, flash memory, or some other form of data storage, to the nodes **116, 118**. The nodes **116, 118** may be interconnected by the cluster fabric **106**, for example, allowing respective nodes in the cluster to access data on data storage devices **128, 130** connected to different nodes in the cluster. Disk modules **124, 126** may communicate with the data storage devices **128, 130** according to the SAN protocol, such as SCSI or FCP, for example. As seen from an operating system on nodes **116, 118**, the data storage devices **128, 130** may appear as locally attached to the operating system. Accordingly, different nodes **116, 118**, etc. may access data blocks through the operating system, rather than expressly requesting abstract files.

(20) It should be appreciated that, while the clustered network environment **100** illustrates an equal

number of network and disk modules, other embodiments may include a differing number of these modules. For example, there may be a plurality of network and disk modules interconnected in a cluster that does not have a one-to-one correspondence between the network and disk modules. That is, different nodes may have a different number of network and disk modules, and the same node may have a different number of network modules than disk modules.

(21) Further, host devices **108, 110** may be networked with the nodes **116, 118** in the cluster, over the storage networking connections **112, 114**. As an example, respective host devices **108, 110** that are networked to a cluster may request services (e.g., exchanging of information in the form of data packets) of nodes **116, 118** in the cluster, and the nodes **116, 118** may return results of the requested services to the host devices **108, 110**. In an embodiment, the host devices **108, 110** may exchange information with the network modules **120, 122** residing in the nodes **116, 118** (e.g., network hosts) in the data storage systems **102, 104**.

(22) In an embodiment, the data storage devices **128, 130** include volumes **132**, which may include an implementation of storage of information onto disk drives or disk arrays or other storage (e.g., flash) as a file system for data, for example. In an example, a disk array may include all traditional hard drives, all flash drives, or a combination of traditional hard drives and flash drives. Volumes may span a portion of a disk, a collection of disks, or portions of disks, for example, and typically define an overall logical arrangement of file storage on disk space in the storage system. In an embodiment a volume may include stored data as one or more files that reside in a hierarchical directory structure within the volume.

(23) Volumes are typically configured in formats that may be associated with particular storage systems, and respective volume formats typically include features that provide functionality to the volumes, such as providing an ability for volumes to form clusters. For example, a first storage system may utilize a first format for their volumes, and a second storage system may utilize a second format for their volumes, where the first and second formats are different from each other.

(24) In the clustered network environment **100**, the host devices **108, 110** may utilize the data storage systems **102, 104** to store and retrieve data from the volumes **132**. For example, the host device **108** may send data packets to the network module **120** in the node **116** within data storage system **102**. The node **116** may forward the data to the data storage device **128** using the disk module **124**, where the data storage device **128** includes a volume **132A**. In this example, the host device **108** may access the volume **132A**, to store and/or retrieve data, using the data storage system **102** connected by the storage network connection **112**. Further, the host device **110** may exchange data with the network module **122** in the node **118** within the data storage system **104** (e.g., which may be remote from the data storage system **102**). The node **118** may forward the data to the data storage device **130** using the disk module **126**, thereby accessing volume **132B** associated with the data storage device **130**. While host device **108** is illustrated as communicating with data storage system **102**, and similarly host device **110** with data storage system **104**, the host devices **108, 110** may communicate via the network (e.g., via fabric **106**) with other storage systems without requiring traversal through storage systems **102, 104** respectively (as just one example). Thus, if storage system **102** is down, then the host device **108** may still access data via storage system **104** or some other cluster at another site.

(25) The data storage system **102, 104** may further provide automated tiering of data to lower-cost object storage tiers either on or off premises to aid in lowering the cost of storage. For example, the data storage system **102, 104** may deliver the benefits of cloud economies by tiering to the cloud storage endpoint **160** (e.g., public clouds and/or private clouds). The data storage system **102, 104** may be associated with a cloud tier including a cloud storage endpoint **160**. The cloud storage endpoint **160** may be an external object store that is associated with a local tier (e.g., the data storage device **128** including the volume **132A**), creating a composite collection of discs. The external object store may store one or more objects. The term “object” may refer to a chunk of data (having one or more blocks of data and/or metadata) that is written together in an object storage

tier. Additionally or alternatively, the term “object” may refer to content or a data object. The cloud storage endpoint **160** may store cloud objects in any size. In some examples, the cloud storage endpoint **160** may store data as 4K blocks, and each object stored in the cloud storage endpoint **160** may be composed of 1,024 4 kilobyte (KB) blocks.

(26) To illustrate an example, the node **116** may tier data stored in the volume **132A** to the cloud storage endpoint **160**. For example, the node **116** may identify infrequently used data stored in the volume **132A** and move that data to a lower-cost object storage in the cloud storage endpoint **160**, leaving frequently used data on the higher-performing, data center storage system. For example, the frequently used data may remain in high-performance solid state drive (SSDs) or hard disk drives (HDDs) of the volume **132A**, allowing the system to reclaim space on the primary storage (e.g., volume **132A**). Although the examples provided may discuss a volume including one or more SSDs and may discuss SSD addresses, it should be understood that this discussion extends to a volume including one or more HDDs and HDD addresses.

(27) Volumes may take advantage of the tiering by keeping active (“hot”) data on the local tier and tiering inactive (“cold”) data to the cloud storage endpoint **160**. The volume **132A** may include one or more SSDs **168** (e.g., SSD **168A**, . . . , **168N**). When a block **166** is written to an SSD **168A** of the volume **132A**, the node **116** may assign the block a temperature value indicating that it is hot. Over time, the node **116** may scan the blocks stored in the SSD **168** (e.g., SSD **168A**, . . . , SSD **168N**) and based on the tiering policies, may keep the scanned block as hot (indicating that the block is frequently accessed), may change the block from hot to cool (indicating that the block is infrequently accessed), may mark cold blocks for tiering to the cloud storage endpoint **160**, and/or may tier marked blocks to the cloud storage endpoint **160**.

(28) The node **116** may concatenate marked blocks stored on the volume **132A** (e.g., SSD **168A**, . . . , SSD **168N**) into an object **170** and when the number of blocks in the object **170** reaches a threshold number (e.g., 1,024), the node **116** may write the object **170** to the cloud storage endpoint **160**. After the block **166** is moved to the cloud storage endpoint **160**, the block **166** may be removed from the SSD **168A**. As shown in FIG. 1, the object **170** is stored at a cloud address **172** and includes a plurality of blocks including the block **166**.

(29) The node **116** may continue to scan blocks stored in the volume **132A** to determine whether to tier data stored in the volume **132A** out to the cloud storage endpoint **160**. Although the following example describes the data storage system **102** tiering data from the volume **132A** to the cloud storage endpoint **160**, it should be understood that the data storage system **104** may perform similar actions as those discussed in the present disclosure in relation to the data storage **102** to tier data from the volume **132B** to the cloud storage endpoint **160** or to another cloud storage endpoint (or other cluster) not shown.

(30) FIG. 2 is an illustrative example of a data storage system **200** (e.g., data storage system **102**, **104** in FIG. 1), in accordance with one or more aspects of the present disclosure. The data storage system **200** includes a node **202** (e.g., nodes **116**, **118** in FIG. 1), and a data storage device **234** (e.g., data storage devices **128**, **130** in FIG. 1). The node **202** may be a general purpose computer, for example, or some other computing device particularly configured to operate as a storage server. A host device **205** (e.g., host device **108**, **110** in FIG. 1) may be connected to the node **202** over a network **216**, for example, to provide access to files and/or other data stored on the data storage device **234**. The node **202** may include a storage controller that provides client devices, such as the host device **205**, with access to data stored within data storage device **234**.

(31) The data storage device **234** can include mass storage devices, such as disks **224**, **226**, **228** of a disk array **218**, **220**, **222**. It will be appreciated that the techniques and systems, described herein, are not limited by the example illustrated in FIG. 2. For example, disks **224**, **226**, **228** may include any type of mass storage devices, including but not limited to magnetic disk drives, flash memory (e.g., SSDs), and any other similar media adapted to store information, including, for example, data (D) and/or parity (P) information.

(32) The node **202** includes one or more processors **204**, a memory **206**, a network adapter **210**, a cluster access adapter **212**, and a storage adapter **214** interconnected by a system bus **242**. The network adapter **210** may correspond to and/or be an example of the network module **120** in FIG. 1. The storage adapter **214** may correspond to and/or be an example of the disk module **124** in FIG. 1. The data storage system **200** also includes an operating system **208** installed in the memory **206** of the node **202** that can, for example, implement a Redundant Array of Independent (or Inexpensive) Disks (RAID) optimization technique, or error correction coding (to name just a few examples), to optimize a reconstruction process of data of a failed disk in an array. The operating system **208** may manage communications for the data storage system **200**, and communications between other data storage systems that may be in a clustered network, such as attached to a cluster fabric **215** (e.g., cluster fabric **106** in FIG. 1). Thus, the node **202**, such as a network storage controller, can respond to host device requests to manage data on the data storage device **234** (e.g., or additional clustered devices) in accordance with these host device requests.

(33) The operating system **208** may include several modules or “layers” executed by one or both of the network module **120** or the disk module **124**. These layers may include a file system **240** that keeps track of a hierarchical structure of the data stored in the storage devices and manages read/write operations (e.g., executes read/write operations on storage in response to client requests). The operating system **208** may establish one or more file systems on the data storage system **200**, where a file system can include software code and data structures that implement a persistent hierarchical namespace of files and directories, for example. The file system may logically organize stored information as a hierarchical structure for files/directories/objects at the storage devices. Each “on disk” file may be implemented as a set of blocks configured to store information, such as text, whereas a directory may be implemented as a specially formatted file in which other files and directories are stored. These data blocks may be organized within a volume block number (VBN) space that is maintained by a file system of the storage operating system **208**. The file system may also assign each data block in the file a corresponding “file offset” or a file block number (FBN). The file system may assign sequences of FBNs on a per-file basis, whereas VBNs may be assigned over a larger volume address space. The file system may organize the data blocks within the VBN space as a logical volume. The file system may be composed of a contiguous range of VBNs from zero to n , for a file system of size $n-1$ blocks, where n is a number greater than 1. In an example, when a new data storage device (not shown) is added to a clustered network system, the operating system **208** is informed where, in an existing directory tree, new files associated with the new data storage device are to be stored. This is often referred to as “mounting” a file system.

(34) In the example data storage system **200**, memory **206** may include storage locations that are addressable by the processors **204** and network adapter **210**, cluster access adapter **212**, and/or storage adapter **214** for storing related software application code and data structures. The processors **204**, the network adapter **210**, the cluster access adapter **212**, and/or the storage adapter **214** may, for example, include processing elements and/or logic circuitry configured to execute the software code and manipulate the data structures. The operating system **208**, portions of which are typically resident in the memory **206** and executed by the processing elements, functionally organizes the storage system by, among other things, invoking storage operations in support of a file service implemented by the storage system. It will be apparent that other processing and memory mechanisms, including various computer readable media, may be used for storing and/or executing application instructions pertaining to the techniques described herein. For example, the operating system may also utilize one or more control files (not shown) to aid in the provisioning of virtual machines.

(35) The network adapter **210** includes the mechanical, electrical and signaling circuitry for connecting the data storage system **200** to the host device **205** over the network **216**, which may include, among other things, a point-to-point connection or a shared medium, such as a LAN. The network adapter **210** may also connect the data storage system **200** to the cloud tier (e.g., cloud

storage endpoint **160** in FIG. 1). The host device **205** may be a general-purpose computer configured to execute applications. As described above, the host device **205** may interact with the data storage system **200** in accordance with a client/host model of information delivery.

(36) The storage adapter **214** cooperates with the operating system **208** executing on the node **202** to access information requested by the host device **205** (e.g., access data on a storage device managed by a network storage controller). The information may be stored on any type of attached array of writeable media such as magnetic disk drives, flash memory, and/or any other similar media adapted to store information. In the example data storage system **200**, the information may be stored in data blocks on the disks **224**, **226**, **228**. The storage adapter **214** can include input/output (I/O) interface circuitry that couples to the disks over an I/O interconnect arrangement, such as a storage area network (SAN) protocol (e.g., Small Computer System Interface (SCSI), iSCSI, hyperSCSI, Fiber Channel Protocol (FCP)). The information may be retrieved by the storage adapter **214** and, in some examples, processed by the one or more processors **204** (or the storage adapter **214** itself) prior to being forwarded over the system bus **242** to the network adapter **210** (and/or the cluster access adapter **212** if sending to another node in the cluster) where the information is formatted into a data packet and returned to the host device **205** over the network **216** (and/or returned to another node attached to the cluster over the cluster fabric **215**). In some examples, the network adapter **210** may format the information into a data packet and forward the data packet to the cloud tier (e.g., cloud storage endpoint **160** in FIG. 1).

(37) In an embodiment, storage of information on disk arrays **218**, **220**, **222** can be implemented as one or more storage volumes **230**, **232** that include a cluster of disks **224**, **226**, **228** defining an overall logical arrangement of disk space. The disks **224**, **226**, **228** that include one or more volumes may be organized as one or more groups of RAID(s) (while in other examples, error correction coding may be used). As an example, volume **230** includes an aggregate of disk arrays **218** and **220**, which include the cluster of disks **224** and **226**. In an example, to facilitate access to disks **224**, **226**, **228**, the operating system **208** may implement a file system (e.g., write anywhere file system) that logically organizes the information as a hierarchical structure of directories and files on the disks. Accordingly, respective files may be implemented as a set of disk blocks configured to store information, whereas directories may be implemented as specially formatted files in which information about other files and directories are stored.

(38) Whatever the underlying physical configuration within this data storage system **200**, data can be stored as files within physical and/or virtual volumes, which can be associated with respective volume identifiers, such as file system identifiers (FSIDs), which can be 32-bits in length in one example. A physical volume corresponds to at least a portion of physical storage devices whose address, addressable space, location, etc. does not change, such as at least some of one or more data storage devices **234** (e.g., a Redundant Array of Independent (or Inexpensive) Disks (RAID system)). In some examples, the location of the physical volume does not change in that the (range of) address(es) used to access it may generally remain constant. A virtual volume, in contrast, may be stored over an aggregate of disparate portions of different physical storage devices. The virtual volume may be a collection of different available portions of different physical storage device locations, such as some available space from each of the disks **224**, **226**, and/or **228**, and is not “tied” to any one particular storage device. Accordingly, a virtual volume may be said to include a layer of abstraction or virtualization, which allows it to be resized and/or flexible in some regards.

(39) Further, a virtual volume may include one or more logical unit numbers (LUNs) **238**, directories **236**, and/or Qtrees **235**. Among other things, these features may allow the disparate memory locations within which data is stored to be identified, for example, and grouped as a data storage unit. For example, the LUNs **238** may be characterized as constituting a virtual disk or drive upon which data within the virtual volume may be stored within the aggregate. LUNs may be referred to as virtual drives, such that they emulate a hard drive from a general purpose computer, while they actually include data blocks stored in various parts of a volume.

(40) One or more data storage devices **234** may have one or more physical ports, where each physical port may be assigned a target address (e.g., SCSI target address). To represent respective volumes stored on a data storage device, a target address on the data storage device **234** may be used to identify one or more LUNs **238**. For example, when the node **202** connects to a volume **230, 232** through the storage adapter **214**, a connection between the node **202** and the one or more LUNs **238** underlying the volume is created. Additionally or alternatively, respective target addresses may identify multiple LUNs, such that a target address may represent multiple volumes. The I/O interface, which may be implemented as circuitry and/or software in the storage adapter **214** or as executable code residing in memory **206** and executed by the processors **204**, for example, may connect to volume **230** by using one or more addresses that identify the one or more LUNs **238**.

(41) Data stored in a volume (e.g., volume **230, 232**) may also be stored in a cache, which may store frequently accessed portions of a source of data in a way that allows the data to be served faster and/or more efficiently than it would be by fetching the data from the source. Referring back to FIG. 1, data stored on the volume **132** may be cached at a cache volume. A cache volume may provide a remote caching capability for an origin volume (e.g., volume **132A**, volume **132B**, etc.), simplifying file distribution, reducing WAN latency, and/or lowering WAN bandwidth costs. In some examples, the cache volume may be beneficial in read-intensive environments where data is accessed more than once and is shared by multiple hosts.

(42) The cache volume may be populated as the host device reads data from the origin volume. For example, the host device may request data from the cache volume. On a first read of any data, the cache volume may fetch the requested data from the origin volume. The requested data may be returned to the cache volume, stored in the cache volume, and then passed back to the host device. As reads are passed through the cache volume, the cache volume may fill up by storing the requested data. In an example, the cache volume may write the data locally in the cache volume. If the host device requests data that is stored in the cache volume, the cache volume may serve the requested data back to the host device without spending time and resources accessing the original source of the data (e.g., the origin volume). Accordingly, the cache volume may serve frequently accessed data directly to the host device without fetching the data from the origin volume again.

(43) The cache volume may serve data faster, if for example, the data storage device on which the cache volume resides is faster than the data storage device on which the origin volume resides. In an example, the cache volume may have faster storage (e.g., FC versus SATA), increased processing power, and/or increased (or faster) memory compared to the origin volume. In another example, the storage space for the cache volume may be physically closer to the host device, such that it does not take as long to reach the data. The cache volume may provide improved performance by providing load distribution, may provide reduced latency by locating data closer to the point of client access, and/or may provide enhanced availability by serving cached data in a network disconnection situation.

(44) In some examples, the cache volume may be aware of the cloud storage endpoint **160** (in the example of FIG. 1) and retrieve at least some data from the cloud storage endpoint. For example, the cache volume may receive a request for data and satisfy the data request by retrieving a first subset of the data from the origin volume and a second subset of the data from the cloud storage endpoint **160**. In this example, the cache volume may bypass requesting the second subset of data from the origin volume and request the second subset of data directly from the cloud storage endpoint **160**, potentially reducing latency and the time it would take to satisfy the data request. For example, the origin volume may receive a large number of requests from host devices, and by requesting the second subset of data directly from the cloud storage endpoint **160** rather than through the origin volume, the cache volume may reduce the load on the origin volume and spread the load across portions of the network.

(45) FIG. 3 is a schematic diagram **300** of a cache volume that retrieves data from an origin volume

and data from a cloud storage endpoint to satisfy a single data request according to one or more aspects of the present disclosure. In the example illustrated in FIG. 3, a data storage system 302 is coupled to a data storage system 304 over a network 305. The network 305 may include, for example, a LAN or WAN. Additionally, the nodes 306 and 308 may communicate over the network 305 and/or over the cluster fabric 106 (shown in FIG. 1). The data storage system 302 and the data storage system 304 may be examples of data storage systems 102, 104 in FIG. 1 and/or data storage system 200 in FIG. 2. The data storage system 302 includes a node 306, and the data storage system 304 includes a node 308. The node 306 may be an example of node 116 in FIG. 1 and/or node 202 in FIG. 2. The node 308 that may be an example of node 118 in FIG. 1 and/or node 202 in FIG. 2. The nodes 306 and 308 may be in the same cluster or a different cluster from each other. Additionally, the node 306, 308 may allow other nodes to access data on data storage devices 310, 312. In the present disclosure, reference to a node 306 performing an action (e.g., receiving, transmitting, determining, storing, etc.) may refer to the data storage device 310 and/or the origin volume 320 performing such action. Likewise, reference to the data storage device 310 and/or the origin volume 320 performing an action may refer to the node 306 performing such action. Similarly, reference to a node 308 performing an action (e.g., receiving, transmitting, determining, storing, etc.) may refer to the data storage device 312 and/or the cache volume 322 performing such an action. Likewise, reference to the data storage device 312 and/or the cache volume 322 performing an action may refer to the node 308 performing such action.

(46) The node 306, 308 may be coupled to data storage device 310, 312, which may be examples of data storage device 128, 130 in FIG. 1 and/or data storage device 234 in FIG. 2. The data storage device 310 includes an origin volume 320 that is mapped to a cache volume 322, with the origin volume 320 storing the original source of data. The data storage device 312 includes the cache volume 322, which may be a temporary storage location that resides between a host device 314 and the origin volume 320. The host device 314 may be an example of host devices 108, 110 in FIG. 1 and/or host device 205 in FIG. 2.

(47) The origin volume 320 stores one or more blocks of data, and the cache volume 322 may be a destination volume that provides a remote cache of the origin volume 320. The origin volume 320 may include one or more SSDs 334 (e.g., 334A, . . . , 334N), which may refer to the local tier. The SSDs 334 may be examples of the SSDs 168A, . . . , 168N in FIG. 1 and/or the disk arrays 218, 220, 222 in FIG. 2. Each of the SSDs 334 may store data in, for example, 4K blocks. The SSD 334A stores a first data subset 330 at an SSD address 336 in the origin volume 320. The first data subset 330 may include one or more 4 KB blocks included in the data 316.

(48) The cache volume 322 may be a writable, persistent cache of the origin volume 320 in a location remote from the data storage device 310 on which the origin volume 320 resides. The cache volume 322 may be a sparse copy of the origin volume 320 and may store a cached subset of the data stored in the origin volume 320. Storage in the cache volume 322 may be used efficiently by prioritizing retention of the working dataset (e.g., most recently used data). The cache volume 322 may use a protocol to communicate with the origin volume 320, where the protocol links the cache volume 322 to the origin volume 320. In an example, the protocol may be remote access layer (RAL), which may include a feature that enables the cache volume 322 to have a revocable read/write or read-only cache granted on an inode by the origin volume 320 to the cache volume 322.

(49) As illustrated in FIG. 3, the cloud storage endpoint 160 may include one or more cloud storage devices 338. The cloud storage device 338 may store an object 340 at a cloud address 342 of the cloud storage device 338 in cloud storage endpoint 160. The object 340 may be an example of the object 170 in FIG. 1. The object 340 may include a plurality of blocks including a second data subset 344. The second data subset 344 may include one or more blocks (e.g., 4 KB blocks) included in the data 316 (e.g., “cold” data). As discussed above, the cache volume 322 may be populated as the host device 314 reads data from the origin volume 320.

(50) The host device **314** may desire to retrieve the data **316** and may transmit a data request **318** for the data **316** to the node **308**. The node **308** (as a cache target) may receive the data request **318** and search the cache volume **322** for the data **316**. It should be understood that the data **316** may include one or more data subsets stored at the origin volume **320** and/or the cloud storage endpoint **160**. The node **308** may determine, based on an identifier of the requested data, whether the requested data is stored in the cache volume **322**. For example, one or more blocks included in the data **316** may be stored at the cache volume **322**, one or more blocks included in the data **316** may be stored at the origin volume **320**, and/or one or more blocks included in the data **316** may be stored at the cloud storage endpoint **160**. In the example illustrated in FIG. 3, the cache volume **322** does not yet store any portions (or blocks) of the data **316**, and the requested data **316** may include the first data subset **330** and the second data subset **344**. In the illustrated example, the first data subset **330** is stored at the origin volume **320**, and the second data subset **344** is stored at the cloud storage endpoint **160**.

(51) FIGS. 3 and 4 may be discussed in relation to each other to better explain data retrieval by the cache volume **322** from the origin volume **320** and the cloud storage endpoint **160**. FIG. 4 is a signaling diagram illustrating a method **400** of retrieving data from an origin volume and data from a cloud storage endpoint to satisfy a single data request according to one or more aspects of the present disclosure. The method **400** may be implemented between the host device **314**, the nodes **306** and **308**, and the cloud storage endpoint **160** (e.g., located in the network **100**). The method **400** may employ similar data retrieval techniques as described with respect to aspects of FIGS. 1, 2, 3, and/or 5. As illustrated, the method **400** includes a number of enumerated actions, but embodiments of the method **400** may include additional actions before, after, and in between the enumerated actions. In some embodiments, one or more of the enumerated actions may be omitted or performed in a different order.

(52) At action **402**, the host device **314** may transmit a data request for data. The data request transmitted at action **402** may be an example of the data request **310** for the data **316** in FIG. 3. A file stored in the cache volume **322** and/or the origin volume **320** may be implemented as a set of blocks configured to store information (e.g., text). For example, the node **306**, **308** may assign each data block in the file a corresponding “file offset” or file block number (FBN). The data request **318** may include an identifier of the data **316**. For example, the data request **318** may include a file handle (e.g., file identifier), a file offset, and a length value of the data **316** that together may identify the data **316**. The data request **318** may be a request for X blocks (or more generally bytes) of data, where X is a positive number. The node **306**, **308** may understand the same file handle, offset, and length value of the data and understand the request as being for X blocks of data in the file identified by the file handle, starting at the offset of the file and spanning the indicated length value. Data may be addressed or organized in other ways while remaining within the scope of the present disclosure.

(53) As discussed above, the node **308** may determine whether the cache volume **322** stores a portion (e.g., one or more blocks) of the requested data **316**. In response to determining that the cache volume **322** stores at least some blocks included in the requested data **316**, the node **308** (e.g., a cache target) may return the applicable blocks to the host device **314**. In response to determining that the cache volume **322** does not store all the blocks included in the requested data **316**, the node **308** may attempt to retrieve such portions of the requested data **316** from the origin volume **320**. It may be desirable to retrieve only those portions of the requested data **316** that are stored at the origin volume **320** (e.g., SSDs **334A**, . . . , **334N**) and request the remaining portion of the requested data **316** directly from the cloud storage endpoint **160**.

(54) Referring back to FIG. 4, at action **404**, the node **308** may transmit a local data request for the data **316** to the node **306**. In response to determining that the cache volume **322** does not store all the blocks included in the requested data **316**, the node **308** may transmit the local data request for the data to the node **306**. The node **306** may receive the local data request from the node **308**.

(55) The local data request may specify to the receiving node **306** to retrieve those portions of the requested data **316** that are stored locally at the data storage device **310** (e.g., stored in the SSDs **334A**, **334N** of the origin volume **320**) and provide an object name of the object and one or more cloud addresses at which the remaining portion(s) of the requested data **316** are stored at the cloud storage endpoint **160**. Referring back to FIG. **3**, the node **306** may determine that the first data subset **330** is stored at the SSD address **336** and that the second data subset **344** (the remaining portion of the requested data **316** that is not stored in the origin volume **320**) is stored at the cloud address **342** of the cloud storage endpoint **160**. In an example, the SSD address **336** may be an SSD physical volume block number (PVBN) in the origin volume **320**, and the cloud address **342** may be a cloud PVBN. In some examples, the file system **240** allocates blocks, frees blocks, to and from a virtual volume of an aggregate. The aggregate, as discussed above, may be a physical volume including one or more groups of storage devices, such as RAID groups, underlying one or more virtual volumes of the storage system. The aggregate may have its own PVBN space and maintain metadata, such as block allocation bitmap structures, within the PVBN space. Each virtual volume may have its own virtual volume block number (VVBN) space and may maintain metadata, such as block allocation bitmap structures, within that PVBN space. PVBNs may be used as block pointers within buffer trees of files stored in a virtual volume. Systems and methods for creating or using PVBNs are described in further detail in U.S. patent application Ser. No. 14/994,924 filed Jan. 13, 2016, entitled “METHODS AND SYSTEMS FOR EFFICIENTLY STORING DATA,” which is incorporated herein by reference in its entirety.

(56) Referring to FIG. **4**, at action **406**, responsive to the local data request from the node **308**, the node **306** may transmit a response including the first data subset **330**, an SSD address of the first data subset **330**, a name of the object storing the second data subset **344**, and a cloud address of the second data subset **344**. Referring back to FIG. **3**, an SSD address of the first data subset **330** may identify the location at which the first data subset **330** is stored at the SSD **334A**. The node **306** may use the SSD address **336** of the SSD **334A** to find the first data subset **330**. The object name may provide an indication to the node **308** regarding how to read the object. A cloud address of the second data subset **344** may be, for example, a cloud PVBN that identifies the location at which the second data subset **344** is stored at the cloud storage endpoint **160**. In some examples, the origin volume **320** may use the cloud address to retrieve the second data subset **344** from the cloud storage endpoint **160**.

(57) The node **308** may receive the response from the node **306** in relation to the action **406**. Responsive to receiving the response from the node **306**, the node **308** may determine that the first data subset **330** is stored at the SSD address of the SSD **334A** and that the second data subset **344** is stored at the cloud address at the cloud storage device **338**, which is in the cloud storage endpoint **160**.

(58) It should be understood that the node **306** may include one or more SSD addresses of the first data subset **330**, one or more names of objects storing the second data subset **344**, and one or more cloud addresses of the second data subset **344**. Additionally, it should be understood that the node **306** may include in a response at least (but not all) the first data subset **330**, an SSD address of the first data subset **330**, a name of the object storing the second data subset **344**, and a cloud address of the second data subset **344**.

(59) In an example, the node **306** may transmit each of the data subsets of the requested data stored at the origin volume **320** and an SSD address of each of the respective data subsets. The data **316** may also, as already noted, include one or more data subsets of the data **316** stored at the cloud storage endpoint **160**. In an example, the node **306** may transmit each of the cloud addresses at which a respective data subset of the data **316** is stored at the cloud storage endpoint **160**, and for each of these cloud addresses, the node **308** may transmit a request for the data stored at the respective cloud address to the cloud storage endpoint **160** (e.g., as noted at action **410** below).

(60) Referring back to FIG. **4**, at action **408**, the node **308** may store the first data subset **330** into

the cache volume **322** and store the object name and cloud address of the second data subset **344** into the cache volume **322**. In some examples, the cache volume **322** may track the first data subset **330** by writing the data that the cache volume **322** receives at the correct offset for the inode it is attempting to fetch the data from. The cache volume **322** may determine the local address in the cache volume **322** and may disregard the origin address. The cache volume **322** may accordingly return the data to the host device **314** in response to a request from the host device **314** for the data.

(61) In some examples, a cloud PVBN may contain a bin number (e.g., three bits) indicating that the PVBN is for the cloud storage endpoint **160**, an object identifier providing a unique identifier (e.g., 34-bit value) of the object, and a slot number represented as a P-bit value (e.g., P=10). The slot number may indicate the location of a block within the object. In some examples, an object may contain 1,024 4K blocks such that 1,024 object PVBNs may have the same object identifier, but different slot numbers from each other.

(62) An object identifier used on the origin volume **320** may already be used by an aggregate on which the cache volume **322** resides. As a result, the cache volume **322** may assign the second data subset **344** a local object identifier that identifies the second data subset **344** in the aggregate on which the cache volume **322** resides. The cache volume **322** may create a mapping from the object identifier received from the origin volume **320** to the local object identifier and save the object name under the local object identifier. The mapping may allow the cache volume **322** to determine whether it has the object name stored for other cloud addresses (e.g., cloud PVBNs) using the same object identifier received from the origin volume **320**. The information about the object may then be stored under the local object identifier in the object information metafile of the cache volume **322**'s aggregate.

(63) Based on receiving the cloud address, the node **308** knows where the second data subset **344** is stored at the cloud storage endpoint **160** and may directly request the second data subset **344** from the cloud storage endpoint **160**. At action **410**, the node **308** may transmit a request for the second data subset **408** stored at the cloud address of the cloud storage device **338** located at the cloud storage endpoint **160**.

(64) If the cache volume **322** is scheduled to read blocks directly from the cloud storage endpoint **160** but has lost connectivity to the object store, then the cache volume **322** may instead transmit to the origin volume **320** a request to fetch the block from the cloud storage endpoint **160** rather than returning an error message that the data was unavailable. In an example, if the aggregate of the mirror volume is cloud mirrored, the cache volume **322** may send the request only if connectivity was lost to both the primary and the mirror.

(65) The cloud storage endpoint **160** may receive the request and retrieve the second data subset **344** stored at the cloud address. At action **412**, responsive to receiving the request for the second data subset **344** stored at the cloud addresses, the cloud storage endpoint **160** may transmit a response including the second data subset **344** to the node **308**. The node **308** may receive the response including the second data subset **344** from the node **306**.

(66) The node **308** may determine, based on a set of cloud storage policies, whether to store the second data subset **344** received from the cloud storage endpoint **160** at the cache volume **322**. A user may establish the set of cloud storage policies for one or more cache volumes. For example, a cloud storage policy may specify that data associated with random reads, but not sequential reads, are to be moved back and stored in the cache volume **322**. In this example, a cloud storage policy may specify that data associated with sequential reads are read, but not to be moved back and stored in the cache volume **322**. If data is moved back and stored in the cache volume **322**, the data may reside in both the cloud tier and in the cache volume. When the data is requested at a later point, the cache volume **322** may retrieve the data locally rather than request the data from the cloud tier.

(67) In another example, a cloud storage policy may specify that data associated with sequential reads, but not random reads, are to be moved back and stored in the cache volume **322**. In this

example, a cloud storage policy may specify that data associated with random reads are read, but not to be moved back and stored in the cache volume **322**. In another example, a cloud storage policy may specify that data associated with sequential reads and random reads are moved back and stored in the cache volume **322**. In another example, a cloud storage policy may specify that neither data associated with sequential reads nor random reads are moved back and stored in the cache volume **322**. If sequential reads of cloud data from the cache volume **322** only trigger reads from cloud and are not stored in the origin volume **320**, the I/O patterns may be the same and the cache volume **322** may consume less SSD space.

(68) If a cloud storage policy specifies the node **308** to store the second data subset **344** at the cache volume **322**, then the node **308** may store the second data subset **344** at the cache volume **322**. The node **308** may store the second data subset **344** to an appropriate inode offset in the cache volume **322**. The node **308** may thereafter serve a request for the second data subset **344** to the host device **314** without requesting the second data subset **344** from the origin volume **320** or the cloud storage endpoint **160**. Additionally, the node **308** may remove the cloud address of the second data subset **344** returned by the origin volume **320** (that had been received as part of action **406** of FIG. 4).

(69) At action **414**, responsive to the data request received in relation to action **402**, the node **308** may transmit a response including the first data subset **330** and the second data subset **344** to the host device **314**. In an example, the first data subset **330** and the second data subset **344** includes all of the blocks included in the requested data. In another example, the first data subset **330** and the second data subset **344** do not include all of the blocks included in the requested data. In this example, the cache volume **322** may retrieve additional data subsets of the requested data (e.g., one or more data subsets stored on the origin volume **320** and/or the cloud storage endpoint **160**) using aspects discussed in the present disclosure. The node **308** may wait until it has received each of the requested blocks before transmitting the response to the host device **314**.

(70) It should be understood that the requested data **316** may include one or more data subsets stored at the origin volume **320**, the cache volume **322** (e.g., fetched by the node **308** based on an earlier request), and/or the cloud storage endpoint **160**. In an example, the data **316** may include the first data subset that is referred to in action **406** and may also include another data subset that is stored in the origin volume **320**. In this example, the response at action **406** may also include the other data subset that is stored in the origin volume **320** along with an SSD address of the other data subset. In other words, for any given data subset that is requested by the client and stored in the origin volume **320**, the node **306** may transmit the given data subset and an SSD address of the given data subset in a response to the node **308**. At action **408**, the node **308** may accordingly store each of the given data subsets that were stored at the origin into the cache volume and track the SSDs address of the given data subsets.

(71) In another example, the data **316** may include the second data subset that is referred to in actions **406**, **408**, and **410** and that is stored at a cloud address. The data **316** may also include another cloud data subset that is stored in the cloud tier. In this example, the response at action **406** may include a name of an object storing the other cloud data subset that is stored in the cloud tier along with a cloud address of the other cloud data subset. In other words, for any given cloud data subset that is requested by the client and stored in the cloud tier, the node **306** may transmit the name of the object storing the given cloud data subset and a cloud address of the given cloud data subset in a response to the node **308**.

(72) At action **408**, the node **308** may accordingly store each of the object names and the cloud addresses associated with the cloud data sets that were provided in the response. At action **410**, the node **308** may accordingly transmit a request for each of the given cloud data subsets stored at their respective cloud addresses. An advantage of identifying the cloud addresses is that the node **308** may reuse these addresses to satisfy later requests from the host device **314** by retrieving data from the cloud storage endpoint **160** without requesting the data from the node **306**.

(73) As discussed above, the cache volume **322** may assign the second data subset **344** a local

object identifier that identifies the second data subset **344** in the aggregate on which the cache volume **322** resides. The cache volume **322** may create a mapping from the object identifier received from the origin volume **320** to the local object identifier and save the object name under the local object identifier. An object may be stored in a filesystem and store multiple blocks (e.g., 1,024 blocks). Accordingly, many (e.g., more than 1,024 references) may point to the object from the filesystem.

(74) The cache volume **322** may realize that an object is invalid through various means and accordingly redirect reads to the origin volume **320**. In some examples, as the data stored in the origin volume **320** is tiered to the cloud storage endpoint **160**, blocks of an object stored on the origin volume **320** may be removed, and the object may become fragmented. Over time, references to the object may decrease and reach a threshold. If the number of references reaches the threshold, the node **306** may determine to free the object. To free the object, the node **306** may send a command to the cache volume **322** to invalidate its information about that object. The cache volume **322** may receive the command and in response, mark the object as invalid in the cache volume **322**'s aggregate object information metafile and remove the mapping entry of that object identifier received from the origin volume **320** to the local object identifier. On a subsequent read request for the object, the cache volume **322** may determine that the object is invalid, and the read may be redirected to the origin volume **320**. The command to the cache volume **322** to invalidate its information about an object may be performed to speed up the redirection process and may be unnecessary for correctness. If the origin volume **320** determines that the cache volume **322** is unavailable, the origin volume **320** may proceed with freeing the object. If the cache volume **322** tries to read the object, then the cache volume **322** will receive an object not found error, mark the object as invalid in its metadata, remove the object identifier mapping, and direct the read to the origin volume **320**.

(75) In some examples, if an object is marked as invalid, the node **308** does not immediately free the object in the cache volume **322** or clean up the corresponding object information metafile because the container file may still contain object PVBNs referencing the object. If the node **308** were to free the object information, then the object identifier may be reused and those old PVBNs may direct the cache volume **322** to the wrong data. To avoid this, the cache volume **322** may wait until those stale PVBNs are freed before freeing the object information and the object identifier. To expedite this process, the node **308** may scan the cache volumes and free the PVBNs of invalid objects rather than performing tiering and defragmentation work.

(76) In some examples, if the cache volume **322** does not receive the command to invalidate an object, then the origin volume **320** may both free and reuse an object identifier before the cache volume **322** has realized that the original object is no longer valid. While the object is reused, the name of the new object may be different because an object name may contain a monotonically increasing sequence number. If the cache volume **322** transmits a request to the origin volume **320** to fetch a block and the origin volume **320** responds with a reused object identifier, then the cache volume **322** may determine that the object identifier mapping that it has is stale by comparing the sequence number received from the origin volume **320** with the sequence number that it stored in the object information metafile for the corresponding local object identifier. If they do not match, then the node **308** may determine that the object identifier mapping is stale. The node **306** may then mark the old object as invalid in its object information metafile, allocate a new local object identifier for the new object, and update the mapping.

(77) FIG. 5 is a flow diagram of a method **500** of retrieving data from an origin volume and from a cloud storage endpoint according to one or more aspects of the present disclosure. Blocks of the method **500** can be executed by a computing device (e.g., a processor, processing circuit, and/or other suitable component). For example, a data storage system such as the data storage system **304** may utilize one or more components, such as the node **308**, data storage device **312**, and/or the cache volume **322**, to execute the blocks of method **500** (as also discussed above with respect to

FIG. 4). As illustrated, the method 500 includes a number of enumerated blocks, but embodiments of the method 500 may include additional blocks before, after, and in between the enumerated blocks. In some embodiments, one or more of the enumerated blocks may be omitted or performed in a different order.

(78) At block 502, the method 500 includes receiving, by a cache from a client, a request for data. At block 504, the method 500 includes determining, by the cache, that a first subset of the data is stored on a storage device and that a second subset of the data is stored at a cloud address located at a cloud storage endpoint. At block 506, the method 500 includes receiving, by the cache from the storage device, the first subset of data in response to transmitting a local data request for the data stored on the storage device. At block 508, the method 500 includes receiving, by the cache from the cloud storage endpoint, the second subset of data in response to transmitting a request for the second subset of data stored at the cloud address to the cloud storage endpoint. At block 510, the method 500 includes transmitting, by the cache to the client, the first and second subsets of data in response to the request for data.

(79) The present embodiments can take the form of an entirely hardware embodiment, an entirely software embodiment, or an embodiment containing both hardware and software elements. Accordingly, it is understood that any operation of the computing systems of computing architecture 100 may be implemented by the respective computing system using corresponding instructions stored on or in a non-transitory computer readable medium accessible by the processing system. For the purposes of this description, a tangible computer-usable or computer-readable medium can be any apparatus that can store the program for use by or in connection with the instruction execution system, apparatus, or device. The medium may include non-volatile memory including magnetic storage, solid-state storage, optical storage, cache memory, and RAM.

(80) Thus, the present disclosure provides a system, method, and machine-readable storage medium for retrieving data in a clustered networking environment. In some embodiments, the method includes receiving, by a cache from a client, a request for data; determining, by the cache, that a first subset of the data is stored on a storage device and that a second subset of the data is stored at a cloud address located at a cloud storage endpoint; receiving, by the cache from the storage device, the first subset of data in response to transmitting a local data request for the data stored on the storage device; receiving, by the cache from the cloud storage endpoint, the second subset of data in response to transmitting a request for the second subset of data stored at the cloud address to the cloud storage endpoint; and transmitting, by the cache to the client, the first and second subsets of data in response to the request for data.

(81) In some examples, the method also includes storing, by the cache, the cloud address and an association between the cloud address and the second subset of data. The cloud address may include a cloud PVBN. In some examples, the method also includes storing, by the cache, the first subset of the data, the first subset being stored at an SSD address of an SSD of the storage device. In some examples, the method further includes transmitting, by the cache to the storage device, the local data request indicating to the storage device to return the first subset of data to the cache and to return a cloud address at which the second subset of data is stored at the cloud storage endpoint.

(82) In some examples, the method also includes transmitting, by the cache to the cloud storage endpoint, the request for the second subset of data stored at the cloud address. The second data subset may be stored on a cloud storage device located in the cloud storage endpoint. In some examples, receiving a request for data includes receiving, by the cache from the client, the request for one or more four kilobyte (KB) data blocks, where the first subset of data includes a first four KB data block that is stored in a first object on an SSD of the storage device, and the second subset of data includes a second four KB data block that is stored at the cloud address of the cloud storage endpoint. In some examples, the storage device includes a plurality of SSDs, and at least one SSD of the plurality stores the first subset of data.

(83) In yet further embodiments, the non-transitory machine-readable medium has instructions for

performing the method of retrieving data, including machine executable code, which when executed by at least one machine, causes the machine to: receive, by a cache from a host device, a request for data; transmit, by the cache to a storage device, a local data request specifying the storage device to return a first portion of the data stored at the storage device and to return a cloud address at which a second portion of the data is stored on a cloud storage endpoint; transmit, by the cache to the cloud storage endpoint, a request for the second portion stored at the cloud address; receive, by the cache, the first and second portions of the data; and transmit, by the cache to the host device, the first and second portions of the data in response to the request for data.

(84) In some examples, the non-transitory machine-readable medium has instructions for performing the method of retrieving data, including machine executable code, which when executed by at least one machine, causes the machine to transmit, based on a storage policy specifying storage of data retrieved from sequential reads from the cloud storage endpoint, the second subset of data to the storage device based on whether the second data subset is based on a sequential read. In some examples, the non-transitory machine-readable medium has instructions for performing the method of retrieving data, including machine executable code, which when executed by at least one machine, causes the machine to transmit, based on a storage policy specifying storage of data retrieved from random reads from the cloud storage endpoint, the second subset of data to the storage device based on whether the second data subset is based on a random read. The cloud address may include a cloud PVBN.

(85) In some examples, the non-transitory machine-readable medium has instructions for performing the method of retrieving data, including machine executable code, which when executed by at least one machine, causes the machine to store, by the cache, the first subset of the data, the first subset being stored at an SSD address of an SSD of the storage device. The second data subset may be stored on a cloud storage device located in the cloud storage endpoint. The cloud storage endpoint may store objects including one or more 4 KB blocks.

(86) In yet further embodiments, the computing device includes a memory containing a machine-readable medium comprising machine executable code having stored thereon instructions for performing a method of retrieving data and a processor coupled to the memory. The processor is configured to execute the machine executable code to: receive, by a cache from a client, a request for data; determine, by the cache, that a first subset of the data is stored on a storage device and that a second subset of the data is stored at a cloud address located at a cloud storage endpoint; receive, by the cache, the first subset from the storage device and the second subset from the cloud storage endpoint; and transmit, by the cache to the client, the first and second subsets of data in response to the request for data.

(87) In some examples, the processor may be configured to execute the machine executable code to receive, by the cache, the cloud address from the storage device. In some examples, the processor may be configured to execute the machine executable code to transmit, by the cache to the storage device, a local data request indicating to the storage device to return the first subset of data and to return a cloud address at which the second subset of data is stored at the cloud storage endpoint. In some examples, the processor may be configured to execute the machine executable code to transmit, by the cache to the cloud storage endpoint, a request for the second subset of the data stored at the cloud address after receiving the cloud address from the storage device.

(88) The foregoing outlines features of several embodiments so that those skilled in the art may better understand the aspects of the present disclosure. Those skilled in the art should appreciate that they may readily use the present disclosure as a basis for designing or modifying other processes and structures for carrying out the same purposes and/or achieving the same advantages of the embodiments introduced herein. Those skilled in the art should also realize that such equivalent constructions do not depart from the spirit and scope of the present disclosure, and that they may make various changes, substitutions, and alterations herein without departing from the spirit and scope of the present disclosure.

Claims

1. A method comprising: receiving, by a cache from a client, a request for data, the cache being remote from the client and located in a data storage device of a first storage system; transmitting, by the cache to an origin volume that is remote from the client and the cache and located in a data storage device of a second storage system remote from the first storage system, a local data request indicating to the second storage system to return a first subset of data to the cache and to return a cloud address at which a second subset of data is stored at a cloud storage endpoint that is remote from the client, the cache, and the origin volume; receiving, by the cache from the origin volume, the first subset of data in response to transmitting the local data request for the data stored on the origin volume; receiving, by the cache from the cloud storage endpoint, the second subset of data in response to transmitting a request for the second subset of data stored at the cloud address to the cloud storage endpoint; and transmitting, by the cache to the client, the first and second subsets of data in response to the request for data.
2. The method of claim 1, further comprising: storing, by the cache, the cloud address and an association between the cloud address and the second subset of data; and in response to a further request for data that includes the second subset of data, retrieving, by the cache, the second subset of data from the cloud storage endpoint using the cloud address and without requesting the cloud address from the origin volume.
3. The method of claim 1, wherein the cloud address includes a cloud physical volume block number (PVCBN).
4. The method of claim 1, further comprising: storing, by the cache, the first subset of the data, the first subset being stored at a solid-state drive (SSD) address of an SSD of the second storage system or a hard disk drive (HDD) address of an HDD of the second storage system.
5. The method of claim 1, further comprising: transmitting, by the cache to the cloud storage endpoint, the request for the second subset of data stored at the cloud address.
6. The method of claim 1, wherein the second data subset is stored on a cloud storage device located in the cloud storage endpoint.
7. The method of claim 1, wherein the second storage system includes a plurality of SSDs, and at least one SSD of the plurality stores the first subset of data.
8. The method of claim 1, further comprising: transmitting, based on a storage policy specifying storage of data retrieved from sequential reads from the cloud storage endpoint, the second subset of data to the origin volume based on whether the second data subset is based on a sequential read.
9. The method of claim 1, further comprising: determining, by the cache, that the first subset of data is stored on the origin volume and that the second subset of the data is stored at the cloud address located at the cloud storage endpoint.
10. The method of claim 1, further comprising: storing, based on a storage policy, the second data subset into the cache, wherein the second data subset is thereby stored in both the cloud storage endpoint and in the cache.
11. A non-transitory machine-readable medium having stored thereon instructions for performing a method of retrieving data, comprising machine executable code which when executed by at least one machine, causes the machine to: receive, by a cache from a host device, a request for data, wherein the cache is external to the host device and located in a data storage device of a first storage system; transmit, by the cache to an origin volume that is external to the cache and the host device and located in a data storage device of a second storage system external to the first storage system, a local data request specifying the second storage system to return a first portion of the data stored at the origin volume and to return a cloud address at which a second portion of the data is stored on a cloud storage endpoint; transmit, by the cache to the cloud storage endpoint that is remote from the cache, the origin volume, and the host device, a request for the second portion

stored at the cloud address; receive, by the cache, the first and second portions of the data; and transmit, by the cache to the host device, the first and second portions of the data in response to the request for data.

12. The non-transitory machine-readable medium of claim 11, further comprising code, which causes the machine to: transmit, based on a storage policy specifying storage of data retrieved from sequential reads from the cloud storage endpoint, the second portion of data to the origin volume based on whether the second data portion is based on a sequential read.

13. The non-transitory machine-readable medium of claim 11, wherein the cloud address includes a cloud physical volume block number (PVCBN).

14. The non-transitory machine-readable medium of claim 11, further comprising code, which causes the machine to: store, by the cache, the first portion of the data, the first portion being stored at a solid-state drive (SSD) address of an SSD of the second storage system.

15. The non-transitory machine-readable medium of claim 11, wherein the second data portion is stored on a cloud storage device located in the cloud storage endpoint.

16. A computing device comprising: a memory containing a machine-readable medium comprising machine executable code having stored thereon instructions for performing a method of retrieving data; and a processor coupled to the memory, the processor configured to execute the machine executable code to: receive, by a cache from a client, a request for data, the cache being remote from the client and located in a data storage device of a first storage system; transmit, by the cache to an origin volume that is remote from the cache and the client and located in a data storage device of a second storage system remote from the first storage system, a local data request indicating to the second storage system to return a first subset of data and to return a cloud address at which a second subset of data is stored at a cloud storage endpoint that is remote from the cache, the origin volume, and the client; receive, by the cache, the first subset from the origin volume and the second subset from the cloud storage endpoint; and transmit, by the cache to the client, the first and second subsets of data in response to the request for data.

17. The computing device of claim 16, wherein the processor is configured to execute the machine executable code to: receive, by the cache, the cloud address from the origin volume.

18. The computing device of claim 16, wherein the processor is configured to execute the machine executable code to: transmit, by the cache to the cloud storage endpoint, a request for the second subset of the data stored at the cloud address after receiving the cloud address from the origin volume.

19. The computing device of claim 16, wherein the processor is further configured to: determine that the first subset of data is stored on the origin volume and that the second subset of data is stored at the cloud address located at the cloud storage endpoint.

20. The computing device of claim 16, wherein: the local data request further indicates to the origin volume to return an object name of an object storing the second subset of data at the cloud storage endpoint.
