

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication	20250251922
Kind Code	A1
Publication Date	August 07, 2025
Inventor(s)	SEQUEIRA DE JUSTO TEIXEIRA; Pedro Miguel et al.

ARCHITECTURE FUNCTION POINTERS POINTERS HAVING CONSISTENT REFERENCE ADDRESSES

Abstract

A function is compiled against a first application binary interface (ABI) and a second ABI of a native first instruction set architecture (ISA). The second ABI defines context data not exceeding a size expected by a third ABI of a foreign second ISA, and uses a subset of registers of the first ISA that are mapped to registers of the second ISA. Use of the subset of registers by the second ABI results in some functions being foldable when compiled using both the first and second ABIs. First and second compiled versions of the function are identified as foldable, or not, based on whether the compiled versions match. Both the first and second compiled versions are emitted into a binary file when they are not foldable, and only one of the first or second compiled versions is emitted into the binary file when they are foldable.

Inventors: SEQUEIRA DE JUSTO TEIXEIRA; Pedro Miguel (Kirkland, WA), MIHOCKA; Darek Josip (Mercer Island, WA), BERRY; Jon Robert (Woodinville, WA), HADLEY; Russell Charles (Bellevue, WA), CLEARY; James David (Seattle, WA), DANG; Clarence Siu Yeen (Newcastle, WA)

Applicant: Microsoft Technology Licensing, LLC (Redmond, WA)

Family ID: 76060003

Appl. No.: 19/047011

Filed: February 06, 2025

Related U.S. Application Data

parent US continuation 18333709 20230613 parent-grant-document US 12229541 child US 19047011

parent US continuation 17550452 20211214 parent-grant-document US 11720335 child US 18333709

Publication Classification

Int. Cl.: G06F8/41 (20180101)

U.S. Cl.:

CPC G06F8/47 (20130101); G06F8/447 (20130101);

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS [0001] This application is a continuation of U.S. patent application Ser. No. 18/333,709, filed Jun. 13, 2023, and entitled “DUAL ARCHITECTURE FUNCTION POINTERS HAVING CONSISTENT REFERENCE ADDRESSES,” which was issued as U.S. Pat. No. _____ on _____, which is a continuation of U.S. patent application Ser. No. 17/550,452, filed Dec. 14, 2021, and entitled “HYBRID BINARIES SUPPORTING CODE STREAM FOLDING,” which was issued as U.S. Pat. No. 11,720,335 on Aug. 8, 2023, which is a continuation of U.S. patent application Ser. No. 17/008,270, filed Aug. 31, 2020, and entitled “NATIVE EMULATION COMPATIBLE APPLICATION BINARY INTERFACE FOR SUPPORTING EMULATION OF FOREIGN CODE,” which was issued as U.S. Pat. No. 11,231,918 on Jan. 25, 2022. The entire contents of each of the foregoing applications are incorporated by reference herein in their entireties.

BACKGROUND

[0002] In computing, each processor, or central processing unit (CPU), implements an instruction set architecture (ISA) that describes syntax and rules to which machine code instructions must adhere in order for those instructions to be natively executable on that processor. Examples of contemporary ISAs are the IA-32 ISA designed by INTEL CORPORATION, the x86-64 ISA designed by ADVANCED MICRO DEVICES, INC., the AArch64 ISA designed by ARM LIMITED, the POWER ISA designed by INTERNATIONAL BUSINESS MACHINES CORPORATION, and the RISC-V ISA designed by UNIVERSITY OF CALIFORNIA, BERKELEY, though many additional ISAs exist. Additionally, in computing, an application binary interface (ABI) is an interface between two binary program components that comprise machine code instructions targeting the same ISA. ABIs are often defined by an operating system (OS) and enable third-party software components to target, and therefore interface with, other libraries/components targeting the same ABI. For example, MICROSOFT WINDOWS provides a “Windows-X64” ABI that enables software to target WINDOWS devices comprising processor(s) implementing the x86-64 ISA, and provides an “ARM64” ABI that enables software to target WINDOWS devices comprising processor(s) implementing the AArch64 ISA.

[0003] In general, a software component is “native” to a device that provides (e.g., via an OS) an ABI to which the component is compiled, and which also comprises a processor having an ISA matching the machine code instructions defining that component. Thus, any software component that does not match an ABI provided by a given computing device can be considered to be “foreign” to that device. For example, an application targeting a Linux ABI for x86-64 cannot natively execute on a computing device having a processor targeting the x86-64 ISA but providing only the Windows-X64 ABI. Similarly, an application targeting the Windows-X64 ABI cannot natively execute on a computing device having a processor targeting the x86-64 ISA but providing

only a Linux ABI for x86-64.

[0004] Additionally, any software component that does not match an ISA of the device's processor(s), can be considered to be “foreign” to that device. For example, an application targeting the x86_64 ISA (regardless of the ABI targeted) cannot natively execute on any computing device having an AArch64 processor. Similarly, an application targeting the AArch64 ISA (regardless of the ABI targeted) cannot natively execute on any computing device having an x86-64 processor.

[0005] Some environments, such as WINDOWS, do make it possible to execute foreign software via software-based ISA emulation, with often significant performance penalties. Thus, for example, it may be possible for an application targeting the ARM64 ABI to be emulated on a device having an x86-64 processor via an AArch64 emulator, and/or it may be possible for an application targeting the Windows-X64 ABI to be emulated on a device having an AArch64 processor via an x86-64 emulator.

[0006] Even when a given software application is available in multiple ABIs, distribution and use of that application can be a problem. For example, average computer users are not aware of the ISAs and/or ABIs supported by their computing devices, and are therefore not able to make an informed decision as to which version of a program to install on his or her computer, leading to confusion and frustration. Additionally, it is burdensome for a developer to distribute and support multiple versions of the same application to target different ABIs/ISAs. One naïve solution that attempts to address these problems is to support the generation and OS support of what are known as “fat” binaries that include both “native” and “foreign” code streams. Thus, for example, a software vendor may generate and distribute a single fat binary for an application, which includes an AArch64 code stream that is used when the binary is loaded on an AArch64-based device, and an x86-64 code stream that is used when the binary is loaded on an x86-64-based device. While fat binaries largely hide the complexities of different processor ISAs from end-users, this is accomplished by essentially doubling the amount of code that is included in a binary (i.e., a 100% code size increase), half of which is not actually used by a given device. In an era in which portable devices that frequently possess limited amounts of storage space are common, fat binaries can quickly consume precious storage resources and lead to software faults and user dissatisfaction.

[0007] Given the broad and diverse ecosystem of devices available today, it is frequently desirable for users to be able to execute foreign software on their computing devices. As such, computer manufacturers and/or OS vendors may be motivated to make execution of foreign software a relatively seamless experience. Some solutions have enabled foreign software to execute via emulation, while enabling that foreign software to interface with native OS libraries. However, these solutions have been limited in that these applications still execute under emulation, rather than natively.

BRIEF SUMMARY

[0008] At least some embodiments described herein facilitate generating and/or consuming native binaries (e.g., applications, modules, libraries, etc.) comprising “emulation compatible” (EC) code that is configured for close interoperation with emulated foreign code, and that provides a high level of interoperability and compatibility to the foreign code. For example, the embodiments herein enable an application developer to target an application to an operating system (OS) provided native EC application binary interface (ABI) to produce and distribute applications that execute natively at a computing device.

[0009] In embodiments, the EC ABI exhibits many behaviors of a foreign ABI, thereby enabling native code targeting the EC ABI to call (via an emulator) and interoperate with foreign code. Thus, by virtue of the EC ABI, these natively-executed applications are enabled to efficiently and seamlessly (e.g., from the perspective of an end-user) interface with emulated foreign code, such as legacy libraries and plug-ins that are loaded within that application's execution context.

[0010] In embodiments, the EC ABI provides context and calling conventions (CCs) for native code that support exception unwinding and/or thread suspensions within emulated foreign code.

For example, in embodiments, the EC ABI uses context data and CCs that are compatible with the foreign ABI, while preserving compatibility with context data and CCs expected by an incumbent native ABI.

[0011] In embodiments, even though the EC ABI is compatible with the foreign ABI, it preserves enough compatibility with the incumbent native ABI that much of code compiled against the EC ABI is identical to code compiled against the incumbent native ABI. This facilitates “folding” of that code within a hybrid binary that supports both the incumbent native ABI (which enables the binary to be consumed by legacy systems that are aware of the incumbent native ABI but not the EC ABI) and the EC ABI (which enables enhanced interoperability with emulated foreign code on enlightened systems that are aware of the EC ABI).

[0012] Notably, in the embodiments herein, there are potentially three entry points of a function compiled to the EC ABI. First, the function could be entered by its ‘real’ address, if it is being called by another EC ABI function. Second, the function could be entered via an “entry thunk,” if it is being called by the foreign ABI. Third, the function could be entered via a “fast-forward sequence,” if the function has been modified at runtime by code executing under the foreign ABI. Regardless of the actual entry point taken to a function, when code compiled against the EC ABI and code compiled against the foreign ABI interact, embodiments permit functions to have consistent memory address references when crossing ABI boundaries, improving compatibility.

[0013] One or more embodiments are directed to methods, systems, and computer program products for compiling source code to a binary file targeting a native first instruction set architecture (ISA) while being operable with code compiled to a foreign second ISA. In these embodiments, a computer system identifies (i) a first ABI corresponding to the first ISA, (ii) a second ABI corresponding to the first ISA, and (iii) a third ABI corresponding to the second ISA. The second ABI (i) defines a first context data format that has a first size that does not exceed second size of a second context data format used by the third ABI, and (ii) defines use of a subset of registers from among a plurality of registers of the first ISA and that are used by the first ABI, the subset of registers being mapped to a set of registers of the second ISA that are used by the third ABI, use of the subset of registers by the second ABI resulting in at least a portion of functions being foldable when compiled using each of the first ABI and the second ABI. For a function defined in the source code, the computer system generates a first compiled version of the function that targets the first ABI, including generating a first set of instructions in the first ISA, and generates a second compiled version of the function that targets the second ABI, including generating a second set of instructions in the first ISA. The computer system determines whether the first compiled version of the function and the second compiled version of the function are foldable within the binary file, based at least on determining whether the first set of instructions and the second set of instructions match. Based at least on determining whether the first compiled version of the function and the second compiled version of the function are foldable within the binary file, the computer system either (i) emits both the first compiled version of the function and the second compiled version of the function into the binary file (i.e., when the first compiled version of the function and the second compiled version of the function are determined to not be foldable within the binary file), or (ii) emits only one of the first compiled version of the function or the second compiled version of the function into the binary file (i.e., when the first compiled version of the function and the second compiled version of the function are determined to be foldable within the binary file).

[0014] One or more embodiments are also directed to methods, systems, and computer program products for generating a hybrid binary image, the hybrid binary image being executable under both a native ABI and a compatibility ABI. In these embodiments, based at least on identifying a first machine type corresponding to the native ABI, a computer system emits the first machine type into a machine type field of the hybrid binary image. The machine type field is structured to be utilized when the hybrid binary image is loaded by a native process executing under the native

ABI. Based at least on identifying a non-foldable first function, the computer system also emits, into the hybrid binary image, both of (i) a first compiled version of the first function that is executable under the native ABI, and (ii) a second compiled version of the first function that is executable under the compatibility ABI. Based at least on identifying a foldable second function, the computer system also emits into the hybrid binary image a compiled version of the second function that is executable under both of the native ABI and the compatibility ABI. The compiled version of the second function is structured to call the first compiled version of the first function when the hybrid binary image is loaded by the native process. The computer system also emits, into the hybrid binary image, a fixup table that is structured to be utilized when the hybrid binary image is loaded by a compatibility process executing under the compatibility ABI. The fixup table defines a plurality of transformations to memory loaded from the hybrid binary image, including a first transformation that adjusts the machine type field to comprise a second machine type corresponding to the compatibility ABI, and a second transformation that configures the compiled version of the second function to call the second compiled version of the first function instead of the first compiled version of the first function.

[0015] One or more embodiments are also directed to methods, systems, and computer program products for consuming a hybrid binary image by a process executing under a compatibility ABI, the hybrid binary image being executable under both a native ABI and the compatibility ABI. In these embodiments, during loading of the hybrid binary image, a computer system determines that a first machine type stored in a machine type field of the hybrid binary image mismatches a second machine type corresponding to the compatibility ABI under which the process is executing. Based on determining that the first machine type mismatches the second machine type, the computer system locates, within the binary image, a fixup table defining a plurality of transformations to memory loaded from the hybrid binary image. The computer system then applies at least a portion of the plurality of transformations to the memory loaded from the hybrid binary image, including applying at least a first transformation that adjusts the machine type field to comprise a second machine type corresponding to the compatibility ABI, and a second transformation that modifies a call site calling a first compiled version of a non-folded function that is executable under the native ABI to instead call a second compiled version of the non-folded function that is executable under the compatibility ABI.

[0016] One or more embodiments are also directed to methods, systems, and computer program products for using a common reference memory address when processing calls within a process that supports execution of both (i) native code targeting a native ABI that corresponds to a native ISA and that has a first CC, and (ii) foreign code targeting a foreign ABI that corresponds to a foreign ISA and that has a second CC. In these embodiments, based at least on identifying a call that targets a reference memory address for a callee function, a computer system determines whether the callee function corresponds to the native ABI or to the foreign ABI. The callee function is determined to correspond to the native ABI based on one of (i) a lookup structure indicating that the reference memory address is contained within a first memory range storing native code, or (ii) a fast-forward sequence being identified at the reference memory address. The callee function is determined to correspond to the foreign ABI based at least on the lookup structure indicating that the reference memory address is contained within a second memory range not storing native code. Based at least on the determining, the computer system initiates execution of the callee function based on performing one of (A) when a caller function corresponds to the foreign ABI, and when the callee function is determined to correspond to the foreign ABI, directly calling the callee function using the reference memory address within an emulator; (B) when the caller function corresponds to the foreign ABI, and when the callee function is determined to correspond to the native ABI, calling an entry thunk that (i) adapts a second CC to the first CC and then (ii) directly calls the callee function using the reference memory address; (C) when the caller function corresponds to the native ABI, and when the callee function is determined to correspond to the

foreign ABI, calling an exit thunk that (i) adapts a first CC of the native ABI to a second CC of the foreign ABI and then (ii) invokes the emulator to directly call the callee function using the reference memory address; or (D) when the caller function corresponds to the native ABI, and when the callee function is determined to correspond to the native ABI, directly calling the callee function using the reference memory address.

[0017] This summary is provided to introduce a selection of concepts in a simplified form that are further described below in the Detailed Description. This Summary is not intended to identify key features or essential features of the claimed subject matter, nor is it intended to be used as an aid in determining the scope of the claimed subject matter.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0018] In order to describe the manner in which the above-recited and other advantages and features of the invention can be obtained, a more particular description of the invention briefly described above will be rendered by reference to specific embodiments thereof which are illustrated in the appended drawings. Understanding that these drawings depict only typical embodiments of the invention and are not therefore to be considered to be limiting of its scope, the invention will be described and explained with additional specificity and detail through the use of the accompanying drawings in which:

[0019] FIG. 1 illustrates an example computer system that facilitates generating and/or consuming native binaries comprising emulation compatible (EC) code that is configured for close interoperation with emulated foreign code, and that provides a high level of interoperability and compatibility to the foreign code;

[0020] FIG. 2 illustrates an example of an EC native application binary interface (ABI) that defines context and calling conventions (CCs) that mirror a foreign ABI, while remaining compatible with an incumbent native ABI;

[0021] FIG. 3A illustrates an example of a dual-architecture hybrid binary;

[0022] FIG. 3B illustrates an example of a windowed view of an import and/or export table;

[0023] FIG. 4A illustrates calling behaviors for an “incumbent” native function compiled to target an incumbent ABI, as well as calling behaviors for an “EC” native function compiled to target an EC ABI;

[0024] FIG. 4B illustrates calling behaviors for an “EC” native function which has been folded because identical code resulted from compiling a source code function to both an incumbent ABI and an EC ABI;

[0025] FIG. 5 illustrates a flow chart of an example method for compiling source code to a binary file targeting a native first instruction set architecture (ISA) while being operable with code compiled to a foreign second ISA;

[0026] FIG. 6 illustrates a flow chart of an example method for generating a hybrid binary image that is executable under both a native ABI and a compatibility ABI;

[0027] FIG. 7 illustrates a flow chart of an example method for consuming a hybrid binary image by a process executing under a compatibility ABI, the hybrid binary image being executable under both a native ABI and the compatibility ABI; and

[0028] FIG. 8 illustrates a flow chart of an example method for using a common reference memory address when processing calls within a process that supports execution of both (i) native code targeting a native application binary interface (ABI) that corresponds to the native ISA and that has a first CC, and (ii) foreign code targeting a foreign ABI that corresponds to a foreign ISA and that has a second CC.

DETAILED DESCRIPTION

[0029] FIG. 1 illustrates an example computer system **101** that facilitates generating and/or consuming native binaries (e.g., applications, modules, libraries, etc.) comprising “emulation compatible” (EC) code that is configured for close interoperation with emulated foreign code, and that provides a high level of interoperability and compatibility to the foreign code. Computer system **101** additionally facilitates generating and/or consuming native binaries that, while comprising EC code, are still compatible with “legacy” computing devices that are not aware of the EC features contained within these binaries.

[0030] In general, computer system **101** operates based on an EC native application binary interface (ABI) that exhibits many behaviors of a foreign ABI, thereby enabling native code targeting the EC native ABI to call (via an emulator) and interoperate with foreign code, such as legacy plug-ins, libraries, etc. In embodiments, the EC native ABI provides context and calling conventions (CCs) for native code that support exception unwinding and/or thread suspensions within emulated foreign code. For example, in embodiments, the EC native ABI uses context data and CCs that are compatible with the foreign ABI, while preserving compatibility with context data and CCs expected by an incumbent native ABI. In doing so, much of code compiled against the EC native ABI is identical to code compiled against the incumbent native ABI, facilitating “folding” of that code within a hybrid binary that supports both the incumbent native ABI (which enables the binary to be consumed by legacy systems that are aware of the incumbent native ABI but not the EC native ABI) and the EC native ABI (which enables enhanced interoperability with emulated foreign code on enlightened systems that are aware of the EC native ABI). Additionally, when code compiled against the EC native ABI and code compiled against the foreign ABI interact, embodiments permit functions to have consistent memory address references when crossing ABI boundaries, improving compatibility.

[0031] In embodiments, computer system **101** comprises or utilizes special-purpose or general-purpose computer hardware, such as, for example, a processor **102** (or a plurality of processors), durable storage **103**, and system memory **104** which are communicatively coupled using a communication bus **105**. In computer system **101**, the processor **102** may implement any available processor instruction set architecture (ISA), such as x86-64, AArch64, POWER, RISC-V, etc., which is referred to herein as the “native” ISA of the processor **102**. Any ISA not matching the native ISA of the processor **102** is referred to herein as a “foreign” ISA. In general, the ISA of the processor **102** defines many hardware aspects of the processor **102**, such as the syntax of machine code instructions that are executable by the processor **102**, a set of registers that are exposed by the processor **102** for use by those machine code instructions, a memory model used by the processor **102**, and the like. Thus, for example, if processor **102** were to implement the AArch64 ISA, then it would execute a different set of machine code instructions (including, for example, one or more of instructions available, instruction format, etc.), and expose a different set of registers (including, for example, one or more of register name, register size, a number of registers, etc.), than if the processor **102** were to implement the x86-64 ISA.

[0032] The durable storage **103** stores computer-executable instructions and/or data structures representing executable software components; correspondingly, during execution of this software at the processor(s) **102**, one or more portions of these computer-executable instructions and/or data structures are loaded into system memory **104**. For example, the durable storage **103** is shown as potentially storing computer-executable instructions and/or data structures corresponding to an operating system **106**, a development environment **107**, a hybrid binary **108**, a foreign binary **109**, and source code **110**.

[0033] The system memory **104** is capable of storing a broad variety of data, which can be loaded from durable storage **103**, stored by the processor **102**, and/or sourced from some other location such as a network device (not shown). In general, computer system **101** operates by loading memory pages defined by one or more binary images stored on durable storage **103** (e.g., hybrid binary **108**, foreign binary **109**, etc.) into system memory **104**, and operating on those memory

pages (as loaded into system memory **104**) using the processor **102**. This includes including executing machine code instructions stored within those memory page(s) to operate on data stored within those memory page(s).

[0034] The operating system **106** (referred to hereinafter as OS **106**) includes libraries **106a**, such as libraries supporting execution of application binaries targeting one or more ABIs that use the native ISA of the processor **102**. As indicated by arrows originating from the development environment **107**, in some embodiments the libraries **106a** include support for an incumbent native ABI **107a** (referred to hereinafter as incumbent ABI **107a**) and an EC native ABI **107b** (referred to hereinafter as EC ABI **107b**), which are described in more detail later. The OS **106** also includes a loader **106c** for loading binary images into system memory **104**, and which is aware of both the incumbent ABI **107a** and the EC ABI **107b**. Thus, based on possessing libraries **106a** and loader **106c**, the OS **106** supports consumption of native binaries (e.g., applications, modules, libraries, etc.) that comprise code targeting the one, or both, of the incumbent ABI **107a** or the EC ABI **107b**.

[0035] In embodiments, the OS **106** also includes an emulator **106b** which, as indicated by arrows originating from the development environment **107**, supports execution of binaries targeting a foreign ABI **107c** (e.g., based on a foreign ISA), via emulation. In embodiments, the emulator **106b** is configured for interoperation with the libraries **106a**, including libraries providing the EC ABI **107b**. Thus, based on possessing emulator **106b**, the OS **106** supports execution of foreign binaries (e.g., applications, modules, libraries, etc.) that comprise code targeting the foreign ABI **107c**.

[0036] As used herein, the term “emulation” can encompass translation and/or interpretation. For example, with ahead-of-time (AOT) translation, foreign ISA instructions are translated into native ISA instructions and are persisted into storage (e.g., durable storage **103**); these translated instruction are then summoned when needed for runtime. In general, AOT translation happens before a program is requested to be executed, such as when the application is installed. In another example, with just-in-time (JIT) translation, foreign ISA instructions are translated into native ISA instruction as the execution of foreign ISA instructions is requested (e.g., when a user runs the program). With JIT translation, translated native code is immediately executed once it is translated from foreign code. In embodiments, JIT translation happens in pieces, as more foreign code is “discovered” for execution. In embodiments, JIT translation of the same block of foreign ISA code is only conducted once, so if the code is executed more than once the cost of translation is only incurred once. In yet another example, with interpretation foreign ISA instructions are read as execution is required, and the equivalent function is performed by an interpreter, but corresponding native ISA code is not generated. Since foreign code is not translated under interpretation, if the same foreign ISA function is executed twice, then the cost of interpretation is incurred twice.

[0037] The development environment **107** supports creation of binaries that target at least the EC ABI **107b**, but frequently also target the incumbent ABI **107a** (i.e., as a dual-architecture or “hybrid” binary image). In embodiments, the development environment **107** even supports creation of binaries that further target the foreign ABI **107c** (and are therefore directly executable on a computer system for which the foreign ABI **107c** is native). In embodiments, the incumbent ABI **107a** is a “legacy” ABI that targets the native ISA of the processor **102**, and the foreign ABI **107c** is an ABI that targets some foreign ISA. The EC ABI **107b**, on the other hand, is a native ABI for the processor **102**, but defines context data and CCs that mirror, or at least share some attributes with, the foreign ABI **107c**. Thus, in embodiments, the incumbent ABI **107a** can be viewed as defining fully “native” data structures and behaviors, and the EC ABI **107b**—while also being native—can be viewed as sharing at least some data structure attributes and behaviors with the foreign ABI **107c**. In one example, the incumbent ABI **107a** is the ARM64 ABI targeting the AArch64 ISA, the foreign ABI **107c** is the Windows-X64 ABI targeting the x86-64 ISA, and the EC ABI **107b** targets the AArch64 ISA but includes context data and CCs at least partially mirroring the Windows-X64 ABI. In embodiments, the EC ABI **107b** defines enough context data and CCs that mirror or map to the foreign ABI **107c** to enable emulated foreign code and EC native code to interact at a very low

level, but also has enough in common with the incumbent ABI **107a** (e.g., a set of available registers) to result in compilation of at least a portion of source functions to each of the incumbent ABI **107a** and the EC ABI **107b** to result in identical compiled function bodies. When compiled function bodies are identical, they can be considered “foldable” such that only one function body is actually included in a resulting binary, and that single function body is used by both a code stream targeting the incumbent ABI **107a** and a code stream targeting the EC ABI **107b** (e.g., via pointer aliasing).

[0038] In order to further demonstrate embodiments of the EC ABI **107b**, FIG. 2 illustrates an example **200** of an EC native ABI that defines context and CCs that mirror a foreign ABI, while remaining compatible with an incumbent native ABI. Initially, FIG. 2 shows a representation of each of the incumbent ABI **107a**, the EC ABI **107b**, and the foreign ABI **107c** from FIG. 1. As shown, the incumbent ABI **107a** defines a set of nine available registers **201a** comprising registers N:A to N:I, which are assumed in example **200** to be all available registers in the native ISA. The foreign ABI **107c**, on the other hand, defines a set of five available registers **201c** comprising registers F:V to F:Z, which are assumed in example **200** to be all available registers in the foreign ISA. As an analogy, the AArch64 ISA (e.g., the “native” ISA) defines 31 general-purpose 64-bit registers, while the x86-64 ISA (e.g., the “foreign” ISA) defines 16 general-purpose 64-bit registers. The EC ABI **107b** bridges the incumbent ABI **107a** and the foreign ABI **107c** by defining a set of five available registers **201b** comprising registers N:A, N:C, N:D, N:E, and N:G. Thus, even though the native ISA has nine registers available, in example **200** the EC ABI **107b** is restricted to only use five of these registers—which is a number corresponding to the number of registers available in the foreign ISA. As will be appreciated by one of ordinary skill in the art, even though native code targeting the EC ABI **107b** may not use all registers defined by the native ISA, that native code is still executable on the native ISA. Thus, the EC ABI **107b** can define use of available registers **201b** in a manner that is compatible with the incumbent ABI **107a**, such that native code targeting the EC ABI **107b** is also executable under a system supporting only the incumbent ABI **107a**.

[0039] In embodiments, the set of available registers **201b** is chosen to use registers most commonly used by code compiled against the incumbent ABI **107a**. Thus, even though, in example **200**, the set of available registers **201b** uses less than all of the set of available registers **201a**, code compiled while targeting the set of available registers **201a** may frequently only actually use registers that are selected from the set of available registers **201b**. In these situations, code compiled against each of the incumbent ABI **107a** and the EC ABI **107b** is identical, and can be folded withing a resulting binary. Notably while, in example **200**, the set of available registers **201b** of the EC ABI **107b** comprises less than all of the set of available registers **201a** of the incumbent ABI **107a**, alternate examples may define the set of available registers **201b** to use all of the set of available registers **201a** (for example, if the foreign ISA has a number of registers matching the native ISA).

[0040] In addition, arrows within FIG. 2 show that a mapping has been defined between the available registers **201b** of the EC ABI **107b** and the available registers **201c** of the foreign ABI **107c**. Thus, for example, the EC ABI **107b** uses register N:A in a manner that mirrors the foreign ABI **107c**'s use of register F:V, the EC ABI **107b** uses register N:C in a manner that mirrors the foreign ABI **107c**'s use of register F:W, and so on. In one example, the EC ABI **107b** defines use of available registers **201b** in a manner that mirrors CCs used by the foreign ABI **107c**. For example, the EC ABI **107b** may limit a number of registers that are available (i.e., as compared to the incumbent ABI **107a**) for passing parameters to a function in order to achieve behaviors of the foreign ABI **107c**. As an analogy, the X64 ABI only uses four registers to pass parameters, so the EC ABI **107b** may only permit four registers to be used for passing parameters, even though the incumbent ARM64 ABI is defined to use additional registers for parameter passing purposes. Additionally, the EC ABI **107b** may mirror stack use by the foreign ABI **107c** (e.g., for saving a

return address, for returning values from a function, etc.) rather than using registers for that purpose (as may be the case for the incumbent ABI **107a**). In embodiments, the mappings between registers includes a mapping of volatility between registers. Thus, for example, a register in the native ISA that is considered non-volatile under the incumbent ABI **107a** might be considered volatile by the EC ABI **107b** if it is mapped to a register considered volatile by the foreign ABI **107c**.

[0041] FIG. **2** also shows that the incumbent ABI **107a** defines context data **202a**, the EC ABI **107b** defines context data **202b**, and the foreign ABI **107c** defines context data **202c**. In embodiments, context data is a format defined by an ABI for storing a snapshot of processor context, such as registers and flags. This context data usable for a variety of purposes, such as for an application to observe its own state (e.g., as part of copy protection schemes, as part of virus or malware detection, etc.), for an exception unwinding routine to unwind a stack after an exception, or to facilitate thread suspension/resumption. As demonstrated by a visual size of each of context data **202a**, context data **202b**, and context data **202c**, in embodiments the EC ABI **107b** defines a format of context data **202b** (e.g., a selection and arrangement of available registers **201b**) that does not exceed a size of context data **202c** used by the foreign ABI **107c**—even though the incumbent ABI **107a** may define context data **202a** exceeding this size. By defining context data **202b** to not exceed a size of context data **202c**, interoperability between the EC ABI **107b** and the foreign ABI **107c** is enabled. For example, if an application executing under the foreign ABI **107c** allocates a memory buffer for storing a format of context data **202c**, and that memory buffer is filled by the EC ABI **107b** using a format of context data **202b** (e.g., as part of handing an exception), then the context data **202b** written by the EC ABI **107b** does not exceed the allocated memory buffer.

[0042] In some embodiments, context data **202b** is defined to have a format that is a blend of context data **202a** and context data **202c**. For example, a format of context data **202a** is visually represented with forward diagonal lines, while a format of context data **202c** is visually represented with backward diagonal lines. Then, a format of context data **202b** is visually represented with both forward and backward diagonal lines, representing at least a partial blending of formats. For example, even though context data **202b** defines an arrangement of registers of the native ISA, it may arrange those registers in a manner that is expected by the foreign ISA (e.g., based on the mappings shown between available registers **201b** and available registers **201c**).

[0043] Notably, when defining the EC ABI **107b** (including one or more of the available registers **201b**, the context data **202b**, or mappings between available registers **201b** and available registers **201c**), there can be tradeoffs between defining the EC ABI **107b** as more closely resembling the incumbent ABI **107a** versus defining the EC ABI **107b** as more closely resembling the foreign ABI **107c**. For example, the more closely the EC ABI **107b** resembles the incumbent ABI **107a**, the more likely it is that code compiled to the EC ABI **107b** can be “folded” with code compiled to the native ABI **107a**; however, this may also make it more likely that ABI translations (via entry and/or exit thunks, discussed later) will need to be employed when transitioning between the EC ABI **107b** and the foreign ABI **107c**. Conversely, the more closely the EC ABI **107b** resembles the foreign ABI **107c**, the more likely it is that native code in the EC ABI **107b** and foreign code in the foreign ABI can interact without use of ABI translations (thunks); however, this may also make it more likely that code compiled to the EC ABI **107b** cannot be “folded” with code compiled to the native ABI **107a**.

[0044] Returning to FIG. **1**, in some embodiments, the development environment **107** supports creation of “dual architecture” hybrid binaries that target two or more ABIs. For example, as indicated by various arrows, FIG. **1** shows that a compiler toolchain **107d** consumes source code **110** in order to generate a hybrid binary **108** that includes code targeting both the incumbent ABI **107a** (i.e., incumbent native code **108a**) and the EC ABI **107b** (i.e., EC native code **108b**). Thus, while containing only native code, the hybrid binary **108** can be viewed as targeting both incumbent native behavior (i.e., incumbent native code **108a**) and as targeting foreign behavior (i.e., EC native code **108b**). Although not shown, in some embodiments hybrid binaries also

include code targeting one or more additional ABIs, such as the foreign ABI **107c**.

[0045] In FIG. **1**, there is a broad arrow between incumbent native code **108a** and EC native code **108b**, indicating that the incumbent native code **108a** and the EC native code **108b** are at least partially “folded” together (i.e., such that a single set of instructions are used by both the incumbent native code **108a** and the EC native code **108b**). As discussed, even though the EC ABI **107b** targets at least some foreign behaviors, it has enough in common with the incumbent ABI **107a** to result in compilation of at least a portion of source functions to each of the incumbent ABI **107a** and the EC ABI **107b** to result in identical compiled function bodies. In embodiments, when compilation of a function to each of the incumbent ABI **107a** and the EC ABI **107b** results in the same compiled code, the compiler toolchain **107d** “folds” this function within the hybrid binary **108**—emitting only a single compiled version of this function for use by both the incumbent native code **108a** and the EC native code **108b**. For example, even though the compiler toolchain **107d** generates both an incumbent native “version” of compiled code and an EC native “version” of compiled code for the function, compiler toolchain **107d** emits only one of these versions into the hybrid binary **108** when those “versions” match (i.e., have identical native code). In testing, the inventors have observed that, due to code folding, supporting dual ABIs (i.e., incumbent ABI **107a** and EC ABI **107b**) in this manner has results in only a 10-30% increase in binary size versus binary targeting a single ABI. This is in stark contrast to traditional fat binaries, which—lacking the ability to fold code—would have close to a 100% increase in binary size.

[0046] In embodiments, when compiling source code **110** to target the EC ABI **107b**, the compiler toolchain **107d** follows source code definitions—such as preprocessor directives—as if the compiler toolchain **107d** were targeting the ISA of the foreign ABI **107c** (even though the compiler toolchain **107d** is generating native code). This is because the EC ABI **107b** exhibits behaviors of the foreign ABI **107c**. For example, the following function follows different logic paths depending on the target ABI:

```
TABLE-US-00001    int function1 (int x) {           #if defined (_ARM64_)           return x
+ 10;              #else                           return x - 2;           #endif }
```

In particular, preprocessor directives define that the value of 10 should be added to ‘x’ when targeting the ARM64 ISA, and that the value of 2 should be subtracted from ‘x’ when not targeting the ARM64 ISA. In embodiments, when compiling this function to target an incumbent ABI **107a** using the ARM64 ISA, the compiler toolchain **107d** generates instructions that add the value of 10 to ‘x’ (i.e., the ARM64 logic path); conversely, when compiling this function to target EC ABI **107b** the compiler toolchain **107d** generates instructions that subtract the value of 2 from ‘x’ (i.e., the non-ARM64 logic path).

[0047] In embodiments, the hybrid binary **108** is configured by the compiler toolchain **107d** to be natively parsed and utilized by a legacy loader (i.e., that is aware of the incumbent ABI **107a** but not the EC ABI **107b**), but to be usable by an enlightened loader (i.e., that is aware of the EC ABI **107b**). In embodiments, the hybrid binary **108** uses a layout/format that is expected by a legacy loader and thus “defaults” to executing under the incumbent ABI **107a**. However, the hybrid binary **108** also includes additional information, such as a fixup table **108c** that enables the hybrid binary **108** to also be consumed by the EC ABI **107b**.

[0048] FIG. **3A** illustrates an example **300a** of a “dual architecture” hybrid binary. More particularly, FIG. **3A** illustrates a more detailed representation of hybrid binary **108** from FIG. **1**, showing two interrelated code streams—code stream **301a** and code stream **301b**. In embodiments, code stream **301a** corresponds to incumbent native code **108a**, and code stream **301b** corresponds to EC native code **108b**. Thus, hybrid binary **108** targets two different native ABIs (i.e., incumbent ABI **107a** and EC ABI **107b**), together with their differing behaviors (i.e., full native behaviors and foreign behaviors). As indicated by arrows between code stream **301a** and code stream **301b**, there is a folding of functions between some functions. In particular, each of code stream **301a** and code stream **301b** includes “non-folded” functions delineated as boxes having solid lines, indicating that

a different compiled version of these functions exists in each of code stream **301a** and code stream **301b**. However, each of code stream **301a** and code stream **301b** also include “folded” functions delineated as boxes having broken lines and connected by arrows, indicating that only one version of each function exists for use by both code stream **301a** and code stream **301b** (e.g., via aliasing). FIG. 3A shows that the hybrid binary **108** might include an additional code stream **301c**, such as a foreign code stream. More typically, however, in embodiments the hybrid binary **108** interacts with a foreign binary **109** (e.g., corresponding to a legacy plug-in/library) that includes a foreign code stream **302**.

[0049] As mentioned, in embodiments, the hybrid binary **108** is configured for native execution by the incumbent ABI **107a**. For example, a machine identifier **303** field in the hybrid binary **108** identifies a machine type expected by the incumbent ABI **107a**, an entry point **304** specifies an entry point to code stream **301a**, and any import and/or export tables (input/export table **305**) provide a “native” view exposing functions (and their locations) that are relevant to the incumbent ABI **107a**. In addition, any “folded” functions that call a non-folded function are configured to call the “incumbent” version of the function in code stream **301a**, rather than the “EC” version in code stream **301b**. As such, a legacy loader need only load the hybrid binary **108** as it would any other compatible binary, in order to execute code stream **301a**.

[0050] In order to facilitate loading by an enlightened loader (e.g., loader **106c**), FIGS. 1 and 3A illustrate that the hybrid binary **108** includes a fixup table **108c**. In embodiments, the fixup table **108c** specifies one or more memory transformations to be applied by the loader **106c** to memory loaded from the hybrid binary **108** when the hybrid binary **108** is loaded into an emulated process **111** (e.g., a compatibility process which can execute EC native code **108b** natively at the processor **102**, as well as execute foreign code—such as foreign binary **109**—via the emulator **106b**).

[0051] In embodiments, the loader **106c** locates the fixup table **108c** upon recognizing a that the machine identifier **303** in the hybrid binary **108** is improper for the EC ABI **107b**. Then, the loader **106c** applies one or more transformations specified in the fixup table **108c** to portion(s) of system memory **104** containing memory page(s) loaded from the hybrid binary **108**, in order to execute code stream **301b** rather than code stream **301a**. In embodiments, each transformation specified in the fixup table **108c** identifies a memory location (e.g., by relative address, by absolute address, etc.), together with a transformation to be applied at that memory location (e.g., to replace one or more bytes at the memory location, to apply an arithmetic operation at the memory location, etc.). In embodiments, the fixup table **108c** comprises one or more fixups to adjust the machine identifier **303** to match an ISA of a process (e.g., emulated process **111**) into which the hybrid binary **108** is loaded. In embodiments, this is a foreign ISA corresponding to the foreign ABI **107c**. In embodiments, the fixup table **108c** comprises one or more fixups to adjust the entry point **304** to specify an entry point to code stream **301b**. In embodiments, the fixup table **108c** comprises one or more fixups to cause folded functions to call the EC version of a non-folded function in code stream **301b** rather than the incumbent version of the function in code stream **301a**. In embodiments, the fixup table **108c** comprises one or more fixups that cause an import/export table **305** to provide a “compatibility” view exposing functions (and their locations) that are relevant to the EC ABI **107b**.

[0052] In some embodiments, causing folded functions to call the EC version of a non-folded function in code stream **301b** rather than the incumbent version of the function in code stream **301a** comprises patching the call in the folded function, itself. In other embodiments, however, folded functions are configured to call non-folded functions indirectly via a dispatch table **306**. In these embodiments, the fixup table **108c** comprises one or more fixups to this dispatch table **306**, which replace an address or offset to the incumbent version of the function in code stream **301a** with an address or offset to the EC version of the function in code stream **301b**. In embodiments, use of a dispatch table **306** can provide efficiency by limiting the number of memory pages to which memory fixups need to be applied in order to execute code stream **301b** rather than code stream

301a.

[0053] In some embodiments, causing the import/export table **305** to provide a “compatibility” view exposing functions (and their locations) that are relevant to the EC ABI **107b** comprises patching a reference to the import/export table **305** to expose a different portion (window) of the import/export table **305**. To illustrate this concept, FIG. 3B illustrates an example **300b** of a windowed view of an import/export table. In particular, FIG. 3B illustrates an import/export table **305** that maps a plurality of functions with their corresponding memory location (e.g., address or offset). In example **300b**, import/export table **305** includes a first shaded portion/zone comprising functions (i.e., A, B, and C) that are applicable to only code stream **301a** (e.g., as non-folded functions), a non-shaded portion/zone comprising functions (i.e., D and E) that are applicable to both code stream **301a** and code stream **301b** (e.g., as folded functions), and a second shaded portion/zone comprising functions (i.e., A', B', and F) that are applicable to only code stream **301b** (e.g., as non-folded functions). Brackets delineate a native view **307a** of import/export table **305** and a compatibility view **307b** of import/export table **305**. In embodiments the native view is specified as a base reference to the first entry of the import/export table **305** and a size/count (e.g., of 5), which includes functions A, B, and C (applicable to only code stream **301a**) as well as functions D and E (applicable to both code streams). In embodiments the compatibility view is specified as a base reference to the fourth entry of the import/export table **305** and a size/count (e.g., of 5), which includes functions D and E (applicable to both code streams) as well as functions A', B', and F (applicable to only code stream **301b**). Thus, in embodiments, patching the base reference to the import/export table **305** to expose a different window of the import/export table **305** comprises patching one, or both, of the base reference or the size/count.

[0054] As mentioned, the EC ABI **107b** facilitates interoperability with the foreign ABI **107c**, such as to execute a legacy plugin/library (via emulation) within the context of an application executing natively under the EC ABI **107b**. This includes enabling functions in EC native code **108b** to call functions in foreign code (e.g., code stream **302** in foreign binary **109**), and enabling functions in foreign code to call functions in EC native code **108b**. Since foreign code executes under an emulator **106b**, there is a transition into, or out of, the emulator **106b** for each of these cross-ABI function calls. In embodiments, a hybrid binary **108** may include entry thunks for transitioning from the foreign ABI **107c** to the EC ABI **107b** (i.e., when code in the foreign binary **109** calls a function in EC native code **108b**), and exit thunks for transitioning from the EC ABI **107b** to the foreign ABI **107c** (i.e., when code in the EC native code **108b** calls a function in the foreign binary **109**). In some embodiments, each function in the EC native code **108b** that can be called by foreign code comprises one entry thunk (which could be zero in size), and an exit thunk for each call to a different foreign function. In embodiments, thunks adapt to differences between the EC ABI **107b** and the foreign ABI **107c**, such as by adapting a CC of the EC ABI **107b** to a CC of the foreign ABI **107c** (or vice versa).

[0055] In embodiments, each entry thunk ensures that parameters being passed to an EC native function by a foreign function are in appropriate location(s) for consumption by the EC native function. This may include, for example, moving a value from one register to another, moving a value from a stack location to a register, moving a value from a register to a stack location, etc. In some embodiments, the EC ABI **107b** may be defined in such a way that parameters passed to an EC native function by a foreign function are already in appropriate location(s) for consumption by the EC native function. In these situations, an entry thunk may be zero in size (and thus do nothing or be omitted). In some implementations, this may particularly be the case if the number of parameters being passed to the EC native function are below a threshold. For example, the inventors have observed that with an AArch64 native ISA and an x86_64 foreign ISA, it is possible to define the EC ABI **107b** (including mappings between AArch64 and x86_64 registers) in a manner that makes it possible to have zero-sized entry thunks if fewer or equal than four parameters are passed from a foreign function to an EC native function, and all of these parameters

are fundamental integers.

[0056] In embodiments, an exit thunk is utilized for each call by an EC native function to a foreign function, and thus a single EC native function may have zero or more exit thunks (depending on how many different foreign functions the EC native function calls). In embodiments, each exit thunk performs one or more of (i) saving a function return address to an appropriate location (e.g., stack or register) for returning to the EC native function, (ii) ensuring that parameters being passed from the EC native function to a called foreign function are in appropriate location(s) for consumption by the foreign function, or (iii) initiating execution of the foreign function within the emulator **106b**.

[0057] FIGS. **4A** and **4B** illustrate examples **400a** and **400b** of calls among native functions, and between native functions and foreign functions. In particular, FIG. **4A** illustrates calling behaviors for an incumbent native function **403** compiled to target the incumbent ABI **107a**, as well as calling behaviors for an EC native function **404** compiled to target the EC ABI **107b**. In embodiments, the incumbent native function **403** and the EC native function **404** are different compiled versions of the same source code function, but have not been folded together due to mismatch between the resulting compiled code. In embodiments, this mismatch could arise because the incumbent native function **403** uses one or more registers available under the incumbent ABI **107a** (but not under the EC ABI **107b**), while the EC native function **403** uses only a reduced set of registers that map to the foreign ABI **107c** (as discussed in connection with FIG. **3A**). Alternatively, this mismatch could arise because a conditional code compilation forces a code logic divergence when targeting the incumbent ABI **107a** and EC ABI **107b**. For instance, as discussed, in embodiments the compiler toolchain **107d** follows different preprocessor directives (and thus different logic paths) when compiling the same function to target each of the incumbent ABI **107a** and EC ABI **107b**. In particular, the compiler toolchain **107d** takes a “native ISA” logic path when targeting the incumbent ABI **107a**, but takes a “foreign ISA” logic path when targeting the EC ABI **107b**.

[0058] In FIG. **4A**, path #1 shows that the incumbent native function **403** is called using a native CC (“Native CC” in the Figures). Additionally, path #1 shows that the incumbent native function **403** can make calls using the Native CC. Path #2, on the other hand, shows that the EC native function **404** is indirectly called using a foreign CC (Foreign CC in the Figures) via an entry thunk **401**, which adapts any differences between the Foreign CC and the Native CC, and then invokes the EC native function **404**. Additionally, path #2 shows that the EC native function **404** may indirectly call a foreign function via an exit thunk **402** which adapts any differences between the Native CC and the Foreign CC, and then invokes the foreign function. Path #3 shows that foreign functions **405** (e.g., within code stream **301c** and/or code stream **302**) are called (e.g., by an exit thunk **402** or another foreign function) using the Foreign CC.

[0059] FIG. **4B** illustrates calling behaviors for an EC native function **406** which has been folded because identical code resulted from compiling a source code function to both the incumbent ABI **107a** and the EC ABI **107b**. In FIG. **4B**, path #1 shows that the EC native function **406** can be called using the Native CC, while path #2 shows that the EC native function **406** can also be indirectly called using the Foreign CC via an entry thunk **401**. Additionally, FIG. **4B** shows that the “EC native function **404** can make calls using the Native CC, either to other native functions or to foreign functions via an exit thunk **402**. Path #3 shows that foreign functions **405** are called (e.g., by an exit thunk **402** or another foreign function) using the Foreign CC.

[0060] In embodiments, computer system **101** enable function call sites to consistently call/reference the true memory address of callee functions, even when bridging native and foreign code (i.e., between the ABI **107b** and the foreign ABI **107c**). Notably, using the true memory addresses at call sites is a challenge when bridging native and foreign code, due to the use of entry and exit thunks to bridge ABI transitions. This is in contrast to prior solutions that bridge native and foreign code, in which call sites within two native functions may actually use different addresses to call the same foreign function. This is because, in these prior solutions, call sites are actually

calling the addresses of thunks, rather than the true address of the foreign functions. This behavior can introduce compatibility concerns if program logic in the native code relies on addresses comparisons (e.g., comparing a first pointer to a foreign function that was obtained by a first native function with a second pointer to the foreign function that was obtained by a second native function), if a pointer is passed from native code to foreign code (where it is invalid/unusable for the foreign code), or if a pointer is passed from foreign code to native code (where it is invalid/unusable for the native code). Using the solutions described herein, when obtaining (or “taking”) a memory address of a first foreign function (e.g., within foreign binary **109**), embodiments ensure that a first native function in the EC native code **108b** and a second native function in the EC native code **108b** both obtain the same memory address for the first foreign function—which is the true memory address at which the first foreign function begins. Additionally, embodiments also ensure that a second foreign function (e.g., within foreign binary **109**) also obtains that same true memory address for the first foreign function. Embodiments also ensure that the address of a native function is the same regardless of whether that address is taken by a foreign function or by another native function.

[0061] In embodiments, consistent memory address references are enabled by an EC lookup structure **307** within the hybrid binary **108** (e.g., which is emitted into the hybrid binary **108** by the compiler toolchain **107d**), together with a dispatcher **112** (i.e., dispatcher **112a** in libraries **106a** and dispatcher **112b** in emulator **106b**). In embodiments, the EC lookup structure **307** is any type of structure that is usable to determine which range(s) of memory addresses of a memory image defined by the hybrid binary **108** contain EC native code **108b**. In embodiments, the EC lookup structure **307** is a bitmap, which uses one bit value to indicate whether or not a corresponding range of memory (e.g., a memory page) comprises EC native code **108b**. However, the EC lookup structure **307** could comprise an alternate data structure type, such as a hash table or a binary tree.

[0062] In embodiments, when a call is being made from a caller function to a callee function, the dispatcher **112** uses the EC lookup structure **307** to determine whether a destination memory address for callee function is within EC native code. Then, with inherent knowledge of whether the caller function is native or foreign code, the dispatcher **112** dispatches the call as appropriate. In embodiments, the dispatcher **112** operates within at least four scenarios: a native caller and a native callee, a native caller and a foreign callee, a foreign caller and a foreign callee, and a foreign caller and a native callee.

[0063] In the first scenario, the caller function is a native function executing under the EC ABI **107b**, and the call is thus handled by dispatcher **112a** in libraries **106a**. The dispatcher **112a** uses the EC lookup structure **307** to determine that the callee's reference memory address is within a memory region corresponding to EC native code **108b**, and that the callee function is therefore also a native function executing under the EC ABI **107b**. In this situation, the caller is calling the true memory address of the callee function, and no thunk is needed, so the dispatcher **112a** directly invokes the callee function using the reference memory address.

[0064] In the second scenario, the caller function is a native function executing under the EC ABI **107b**, and the call is thus handled by dispatcher **112a** in libraries **106a**. The dispatcher **112a** uses the EC lookup structure **307** to determine that the callee's reference memory address is not within a memory region corresponding to EC native code **108b**, and that the callee function is therefore a foreign function executing under the foreign ABI **107c**. Referring to FIG. 4A, this is the situation of path #2, in which the EC native function **404** calls a foreign function via exit thunk **402**. As such, the dispatcher **112a** cannot call the reference memory address (i.e., of the callee function) directly, because that would bypass the exit thunk **402**. Instead, the dispatcher **112a** locates a new reference address to the exit thunk **402**, and invokes the exit thunk **402** using the new reference memory address. In embodiments, the exit thunk **402**, in turn, adapts a CC of the EC ABI **107b** to a CC of the foreign ABI **107c**, and then invokes the emulator **106b**. The emulator **106b**, in turn, directly calls the callee function using the original reference memory address.

[0065] In embodiments, the new reference address to the exit thunk is contained within the callee function, itself. In these embodiments, a call site is associated with two reference memory addresses: the original reference memory address of the callee function, and the new reference address to the exit thunk. When the hybrid binary **108** is loaded under the incumbent ABI **107a**, a dispatcher used by the incumbent ABI **107a** uses the original reference memory address of the callee function directly, ignoring the new reference address to the exit thunk. Notably, the incumbent ABI **107a** can ignore the new reference address to the exit thunk because it does not interact with foreign code. When the hybrid binary **108** is loaded under the EC ABI **107b**, on the other hand, the dispatcher **112a** also utilizes the new reference address to the exit thunk for interacting with foreign code. Notably, the compiler toolchain **107d** facilitates foldability by including both the original reference memory address and the new reference address within compiled code, regardless of whether a function is being targeted to the incumbent ABI **107a** or the EC ABI **107b**.

[0066] In the third scenario, the caller function is a foreign function executing under the foreign ABI **107c** within the emulator **106b**, and the call is thus handled by dispatcher **112b** in the emulator **106b**. The dispatcher **112b** uses the EC lookup structure **307** to determine that the callee's reference memory address is not within a memory region corresponding to EC native code **108b**, and that the callee function is therefore also a foreign function executing under the foreign ABI **107c**. In this situation, the caller is calling the true memory address of the callee function, and no thunk is needed, so the dispatcher **112b** directly invokes the callee function using the reference memory address within the emulator **106b**.

[0067] In the fourth scenario, the caller function is a foreign function executing under the foreign ABI **107c** within the emulator **106b**, and the call is thus handled by dispatcher **112b** within the emulator **106b**. The dispatcher **112b** uses the EC lookup structure **307** to determine that the callee's reference memory address is within a memory region corresponding to EC native code **108b**, and that the callee function is therefore a native function executing under the EC ABI **107b**. Referring to FIG. 4A, this is the situation of path #2, in which a foreign function is calling the EC native function **404** via entry thunk **401**. As such, the dispatcher **112b** cannot call the reference memory address (i.e., of the callee function) directly, because that would bypass the entry thunk **401**. Instead, the dispatcher **112b** locates a new reference address to the entry thunk **401**, and invokes the entry thunk **401** using the new reference memory address. In embodiments, the entry thunk **401**, in turn, adapts a CC of the foreign ABI **107c** to a CC of the EC ABI **107b**, and then invokes the callee function using the original reference memory address.

[0068] Notably, in the fourth scenario, the caller function may be legacy foreign code that has not been designed with awareness of, or compatibility with, the EC ABI **107b**. Thus—unlike the second scenario—the caller function cannot be modified (e.g., by a compiler toolchain) to contain the address of the entry thunk. In embodiments, the dispatcher **112b** obtains the new reference address for the entry thunk from a block of memory immediately preceding the original reference memory address of the callee function. In embodiments, data for obtaining new reference address for the entry thunk was inserted into this block of memory by the compiler toolchain **107d** during generation of the hybrid binary **108**. The particular data contained in this block of memory can vary, but in some embodiments, it is a memory offset (e.g., from a beginning of a memory image defined by the hybrid binary **108**). In other embodiments, it could be a direct address reference to the entry thunk. In other embodiments, the dispatcher **112b** obtains the new reference address in some alternate way, such as from a data-tree or an ordered array of addresses.

[0069] In embodiments, the EC ABI **107b**—together with the compiler toolchain **107d** and the dispatcher **112**—supports a “compatibility mode” that enables foreign code that calls native functions to successfully identify, and potentially patch, foreign code inserted at the beginning of the native functions. This enables foreign code to disassemble the beginning of a called function (e.g., as part of antivirus protection, copy protection, etc.) and/or to patch the beginning of a called

function (e.g., as part of a profiling redirection) in connection with calling that function. Since foreign code—which may not have awareness that it is being emulated—may expect to find recognizable foreign code at the beginning of the called function, in some embodiments, the compiler toolchain **107d** emits, into the hybrid binary **108**, a “fast-forward sequence” (a form of a thunk) for a native function, on which the foreign code can operate for disassembly and/or patching. In embodiments, a “fast-forward sequence” comprises placeholder foreign code that concludes with a reference or jump to the true address of the native function. In embodiments, this fast-forward sequence is stored in a memory page marked in the EC lookup structure **307** as not comprising EC native code. Thus, in embodiments, when a foreign caller calls a native function having a fast-forward sequence, the call is initially treated as a foreign-to-foreign call (i.e., as in the third scenario above) such that the emulator **106b** executes the fast-forward sequence and makes a foreign-to-native call at the conclusion of the fast-forward sequence (i.e., as in the fourth scenario above).

[0070] Notably, this arrangement could cause a double-thunking performance issue in which the emulator **106b** is invoked just to emulate a few instructions in the fast-forward thunk (e.g., as in the third scenario), only to then initiate a call to another thunk—an entry thunk to the native function (e.g., as in the fourth scenario). Some embodiments avoid this double-thunking performance issue by configuring the dispatcher **112b** (in the emulator **106b**) to perform an additional check after determining that a callee's reference memory address is not within a memory region corresponding to EC native code **108b** (i.e., using the EC lookup structure **307**). This additional check is to “peek” at the beginning of the callee function to determine if a signature of the fast-forward sequence is present. If so, the dispatcher **112b** obtains the true address of the native function from the fast-forward sequence and directly initiates a call to the native function (i.e., as in the fourth scenario above). By doing this additional check, processing of the fast-forward thunk by the emulator **106b** has been avoided.

[0071] In view of the foregoing introduction, the following discussion now refers to a number of methods and method acts. Although the method acts may be discussed in a certain order or may be illustrated in a flow chart as occurring in a particular order, no particular ordering is required unless specifically stated, or required because an act is dependent on another act being completed prior to the act being performed.

Native Emulation Compatible Application Binary Interface for Supporting Emulation of Foreign Code

[0072] FIG. 5 illustrates a flow chart of an example method **500** for compiling source code to a binary file targeting a native first ISA while being operable with code compiled to a foreign second ISA. Initially, method **500** comprises an act **501** of identifying a first and second ABI for a first ISA, and a third ABI for a second ISA, the second ABI defining context data that does not exceed a context data size of the third ABI, and defining use of a subset of registers that map to the third ABI. In embodiments, act **501** comprises identifying (i) a first ABI corresponding to the first ISA, (ii) a second ABI corresponding to the first ISA, and (iii) a third ABI corresponding to the second ISA. In an example, the compiler toolchain **107d** identifies each of the incumbent ABI **107a** (i.e., first ABI), the EC ABI **107b** (i.e., second ABI), and the foreign ABI (i.e., third ABI).

[0073] In these embodiments, the second ABI defines a first context data format that has a first size that does not exceed second size of a second context data format used by the third ABI. In reference to FIG. 3A, for example, a format of context data **202b** is defined by the EC ABI **107b** to have a size that does not exceed a size of context data **202c** (which is defined by the foreign ABI **107c**).

[0074] In these embodiments, the second ABI also defines use of a subset of registers from among a plurality of registers of the first ISA, and that are used by the first ABI. In some embodiments, the subset of registers comprises less than all of the plurality of registers of the first ISA (e.g., as in example **200**), though in other embodiments subset of registers comprises all of the plurality of

registers of the first ISA. In embodiments, this subset of registers is mapped to a set of registers of the second ISA that are used by the third ABI. In embodiments, the second ABI only permits use of the subset of registers, to the exclusion of one or more other registers permitted by the first ABI. Referring again to FIG. 3A, the EC ABI **107b** defines a set of available registers **201b**, which are a subset of the available registers **201a** used by the incumbent ABI **107a** (i.e., and which omit one or more of available registers **201a** used by the incumbent ABI **107a**). In addition, the available registers **201b** are mapped by the EC ABI **107b** to the set of available registers **201c** used by the foreign ABI **107c**.

[0075] In embodiments, use of the subset of registers by the second ABI results in at least a portion of functions being foldable when compiled using each of the first ABI and the second ABI. For example, in embodiments the selection of available registers **201b** mirrors available registers **201a** closely enough that, in at least some cases, functions compiled against each of the incumbent ABI **107a** and the EC ABI **107b** result in identical compiled code. In embodiments, use of the subset of registers by the second ABI also enables at least a portion of functions compiled to target the second ABI to be called from the third ABI without an entry thunk. For example, if the first ABI is the AArch64 ABI, and the third ABI is the Windows-X64 ABI, it may be possible to define the second ABI in a manner that can omit entry thunks for native functions that receive four or fewer input parameters.

[0076] Method **500** also comprises an act **502** of emitting a function into a binary. As shown, act **502** includes an act **503** of compiling the function while targeting the first ABI to create a first compiled version, and an act **504** of compiling the function while targeting the second ABI to create a second compiled version. As shown, no particular ordering is shown between acts **503** and **504**, and it will be appreciated that these acts could be performed in parallel or serially (in either order). In embodiments, for a function defined in the source code, act **503** comprises generating a second compiled version of the function that targets the second ABI, including generating a second set of instructions in the first ISA, while act **504** comprises generating a first compiled version of the function that targets the first ABI, including generating a first set of instructions in the first ISA. In an example, the compiler toolchain **107d** compiles a source code function twice—once while targeting the first ABI, and once while targeting the second ABI—resulting in two compiled versions of the function. As discussed, in embodiments the compiler toolchain **107d** takes a “native ISA” logic path when targeting the incumbent ABI **107a**, and takes a “foreign ISA” logic path when targeting the EC ABI **107b**. Thus, in some embodiments, generating the first compiled version of the function that targets the first ABI in act **502** comprises using source code definitions of the first ISA, and generating the second compiled version of the function that targets the second ABI in act **503** comprises using source code definitions of the second ISA.

[0077] Act **502** also comprises an act **505** of determining if the first and second compiled versions match and are therefore foldable. In embodiments, act **505** comprises determining whether the first compiled version of the function and the second compiled version of the function are foldable within the binary file, based at least on determining whether the first set of instructions and the second set of instructions match. In an example, the compiler toolchain **107d** performs a comparison between the two compiled versions of the function produced in act **503** and act **504**. If they match, then the two compiled versions of the function are foldable; otherwise, they are not foldable.

[0078] Depending on the outcome of act **505**, act **502** comprises either an act **506** (when the compiled functions are not foldable) of emitting both compiled version into the binary, or an act **507** (when the compiled functions are foldable) of emitting only one compiled version into the binary. In embodiments, act **506** comprises, based at least on determining whether the first compiled version of the function and the second compiled version of the function are foldable within the binary file, and when the first compiled version of the function and the second compiled version of the function are determined to not be foldable within the binary file, emitting both the

first compiled version of the function and the second compiled version of the function into the binary file. In an example, the compiler toolchain **107d** emits both an incumbent version of the function (e.g., such as incumbent native function **403**) and an EC version of the function (e.g., such as EC native function **404**) into hybrid binary **108** when the compiled functions are not foldable. In embodiments, act **507** comprises, based at least on determining whether the first compiled version of the function and the second compiled version of the function are foldable within the binary file, and when the first compiled version of the function and the second compiled version of the function are determined to be foldable within the binary file, emitting only one of the first compiled version of the function or the second compiled version of the function into the binary file. In an example, the compiler toolchain **107d** emits a single EC version of the function (e.g., such as EC native function **406**) into hybrid binary **108** when the compiled functions are foldable.

[0079] In embodiments, the second ABI enables the binary file to be natively executed on both of (i) a first computer system implementing the first ABI but not the second ABI, and (ii) a second computer system implementing the second ABI. For example, a hybrid binary **108** generated using method **500** is executable on a “legacy” computer system implementing only the incumbent ABI **107a**, or on a modern computer system also implementing the EC ABI **107b**.

[0080] In some embodiments, method **500** also comprises emitting an entry thunk into the binary file, the entry thunk comprising code in the first ISA that adapts the third ABI to the second ABI. In an example, when act **506** is performed, the compiler toolchain **107d** emits an entry thunk for the emitted EC version of the function (e.g., EC native function **404**), enabling foreign code to call the function. In another example, when act **507** is performed, the compiler toolchain **107d** emits an entry thunk for the single emitted function, which can be considered EC function (e.g., EC native function **406**), enabling foreign code to call the function.

[0081] In some embodiments, method **500** also comprises emitting one or more exit thunks into the binary file, each exit thunk comprising code in the first ISA that adapts the second ABI to the third ABI. In an example, when act **506** is performed, the compiler toolchain **107d** emits one or more exit thunks for the emitted EC version of the function (e.g., EC native function **404**), enabling the function to call foreign code. In another example, when act **507** is performed, the compiler toolchain **107d** emits one or more exit thunks for the single emitted function, enabling the function to call foreign code. In some embodiments, method **500** generates a different exit thunk for each different call by the function, but in other embodiments method **500** could generate a single exit thunk for each callee.

[0082] As mentioned, a hybrid binary may comprise additional foreign code streams, such as code stream **301c** comprising foreign code. Thus, method **500** could further comprise generating a third compiled version of the function that targets the third ABI, including generating a third set of instructions in the second ISA, and emitting the third compiled version of the function into the binary file.

[0083] Some variations of method **500** may omit the “incumbent” code stream, such that the emitted binary targets the EC ABI **107b**, but not the incumbent ABI **107a**. As will be appreciated, the resulting binary provides compatibility with code targeting the foreign ABI **107c**, but would lack backwards compatibility with programs implementing the incumbent ABI **107a** but not the EC ABI **107b**. In embodiments, omitting the “incumbent” code stream may be useful for binaries that are only intended to be used in an application that does not use the incumbent ABI **107a**. As one example, a binary implementing functionality (e.g., codec support for an obsolete video format) that has been deprecated for use with native applications, but which is retained for compatibility with emulated applications.

Hybrid Binaries Supporting Code Stream Folding

[0084] While FIG. 5 focused on a method for generating code using an EC ABI **107b** that enables code folding, a hybrid binary **108** may comprise additional features that enable backward compatibility with legacy systems (e.g., having only the incumbent ABI **107a**), while enabling

modern system (e.g., having the EC ABI **107b**) to take advantage of additional features for interacting with foreign code. To further describe creation of these dual-architecture hybrid binaries **108**, FIG. **6** illustrates a flow chart of an example method **600** for generating a hybrid binary image, the hybrid binary image being executable under both a native ABI and a compatibility ABI. Then, to further describe consumption of these dual-architecture hybrid binaries **108**, FIG. **7** illustrates a flow chart of an example method **700** for consuming a hybrid binary image by a process executing under a compatibility ABI, the hybrid binary image being executable under both a native ABI and the compatibility ABI.

[0085] Referring first to binary image creation, and to FIG. **6**, Method **600** comprises a plurality of acts (i.e., acts **601-604**) that, as indicated in the flow chart, can be performed in any order with respect to each other. As shown, method **600** comprises an act **601** of identifying and emitting a machine type of a native ABI into a machine type field that is read by the native ABI. In some embodiments, act **601** comprises, based at least on identifying a first machine type corresponding to the native ABI, emitting the first machine type into a machine type field of the hybrid binary image, the machine type field being structured to be utilized when the hybrid binary image is loaded by a native process executing under the native ABI. In an example, the compiler toolchain **107d** emits, into hybrid binary **108**, a machine identifier **303** that matches an identifier expected by the incumbent ABI **107a** when loading binaries targeting the incumbent ABI **107a**. By emitting an identifier expected by the incumbent ABI **107a**, the incumbent ABI **107a** sees what is expected when it loads the hybrid binary **108**, even if the hybrid binary **108** is also compatible with the EC ABI **107b**.

[0086] Method **600** also comprises an act **602** of identifying and emitting a non-foldable function. As shown, act **602** comprises both of (i) an act **602a** of emitting a first compiled version of the non-foldable function that is executed under the native ABI, and (ii) an act **602b** of emitting a second compiled version of the non-foldable function that is executed under a compatibility ABI. As shown, no particular ordering is shown between acts **602a** and **602b**, and it will be appreciated that these acts could be performed in parallel or serially (in either order). In some embodiments, act **602a** comprises, based at least on identifying a non-foldable first function, emitting, into the hybrid binary image, a first compiled version of the first function that is executable under the native ABI, while act **602b** comprises, based at least on identifying a non-foldable first function, emitting, into the hybrid binary image, a second compiled version of the first function that is executable under the compatibility ABI. In an example, based at least on compiling a source code function to different compiled versions (e.g., in act **503** and **504** of method **500**), and also based on determining that those compiled versions are not foldable (e.g., in act **505** of method **500**), the compiler toolchain **107d** emits both of those compiled versions into hybrid binary **108** in act **602a** and act **602b** (e.g., as part of act **506** of method **500**).

[0087] Method **600** also comprises an act **603** of identifying and emitting a foldable function that is executed under both the native ABI and the compatibility ABI, including emitting a compiled version of the foldable function with a call to the first compiled version of the non-foldable function. In some embodiments, act **603** comprises, based at least on identifying a foldable second function, emitting into the hybrid binary image a compiled version of the second function that is executable under both of the native ABI and the compatibility ABI. In an example, based at least on compiling a source code function to different compiled versions (e.g., in act **503** and **504** of method **500**), and also based on determining that those compiled versions are foldable (e.g., in act **505** of method **500**), the compiler toolchain **107d** emits only one of those compiled versions into hybrid binary **108** in act **603** (e.g., as part of act **507** of method **500**).

[0088] In embodiments, the compiled version of the second function is structured to call the first compiled version of the first function when the hybrid binary image is loaded by the native process. In an example, the compiler toolchain **107d** configures the hybrid binary **108** such that the code emitted in act **603** calls the code emitted in act **602a** “by default,” so that the code emitted in act

602a (rather than the code emitted in act **602b**) is executed when the hybrid binary **108** is loaded under the incumbent ABI **107a**.

[0089] Method **600** also comprises an act **604** of emitting a fixup table utilized by the compatibility ABI. In some embodiments, act **604** comprises emitting, into the hybrid binary image, a fixup table that is structured to be utilized when the hybrid binary image is loaded by a compatibility process (e.g., emulated process **111**) executing under the compatibility ABI, the fixup table defining a plurality of transformations to memory loaded from the hybrid binary image. In an example, the compiler toolchain **107d** emits, into hybrid binary **108**, a fixup table **108c** that includes transformations, applied when the hybrid binary **108** is loaded under the EC ABI **107b**, that cause EC features of the hybrid binary **108** to be utilized.

[0090] In embodiments, the plurality of transformations in the fixup table **108c** include a transformation that adjusts the machine type field to comprise a second machine type corresponding to the compatibility ABI. In an example, the compiler toolchain **107d** emits, into the fixup table **108c**, a memory transformation that replaces (in system memory **104**) a machine identifier loaded from the machine identifier **303** into system memory **104** with a machine identifier matching a foreign ISA corresponding to the foreign ABI **107c**.

[0091] In embodiments, the plurality of transformations in the fixup table **108c** include a transformation that configures the compiled version of the (foldable) second function to call the second compiled version of the (non-foldable) first function instead of the first compiled version of the (non-foldable) first function. In an example, the compiler toolchain **107d** emits, into the fixup table **108c**, a memory transformation that replaces a first memory address referencing the first compiled version of the (non-foldable) first function with a second memory address referencing the second compiled version of the (non-foldable) first function. While the first memory address could be replaced by the second memory address directly within the compiled version of the (foldable) second function directly, in embodiments, the compiled version of the (foldable) second function is structured to call the first compiled version of the (non-foldable) first function via a dispatch table **306** that references the first compiled version of the (non-foldable) first function. In these embodiments—rather than modifying the compiled version of the (foldable) second function, itself—this transformation modifies the dispatch table **306** to reference the second compiled version of the (non-foldable) first function instead of the first compiled version of the (non-foldable) first function. Further, in these embodiments, method **600** also comprises emitting the dispatch table into the hybrid binary image.

[0092] In embodiments, method **600** also comprises emitting, into the hybrid binary image, an entry point **304** referencing the first compiled version of the (foldable) first function. In embodiments, this entry point **304** is structured to be utilized when the hybrid binary image is loaded by the native process (e.g., using the incumbent ABI **107a**). In embodiments, the plurality of transformations in the fixup table **108c** include a transformation that adjusts the entry point to reference the second compiled version of the (foldable) first function instead. This adjusted entry point is thus utilized when the hybrid binary image is loaded by the compatibility process (e.g., using the EC ABI **107b**).

[0093] As discussed in connection with FIGS. 3A and 3B, a hybrid binary **108** may comprise one or more of an import table or an export table (import/export table **305**), with the hybrid binary **108** configured to show a “native” view exposing functions (and their locations) that are relevant to the incumbent ABI **107a**. For example, as demonstrated in FIG. 3B, the import/export table **305** could be referred to using a base reference size/count that exposes native view **307a**. Thus, in embodiments, method **600** comprises emitting, into the hybrid binary image, one or more tables that comprise (i) a first zone referencing at least the (non-folded) first function using the first compiled version of the (non-folded) first function, (ii) a second zone referencing at least the (folded) second function using the compiled version of the second function, and (iii) a third zone referencing at least the (non-folded) first function using the second compiled version of the (non-

folded) first function.

[0094] In embodiments, in order to provide a “native” view of these table(s), method **600** comprises emitting, into the hybrid binary image, a reference to the one or more tables that provides a native view of the one or more tables that includes the first zone and the second zone, while excluding the third zone. This reference is structured to be utilized when the hybrid binary image is loaded by the native process (e.g., corresponding to incumbent ABI **107a**). In these embodiments, the plurality of transformations in the fixup table **108c** then include a transformation that adjusts the reference to the one or more tables to provide a compatibility view of the one or more tables (for use by the EC ABI **107b**) that includes the second zone and the third zone, while excluding the first zone. Thus, in embodiments, the reference to the one or more tables provides the native view of the one or more tables by specifying an offset and a size, and a transformation in the fixup table **108c** adjusts the reference to provide the compatibility view of these table(s) by modifying one or more of the offset or the size.

[0095] In embodiments, the emitted table(s) comprise an import table; thus, in embodiments of method **600**, the one or more tables comprise one or more import tables, and the first, second, and third entries comprise first, second, and third function imports. In additional or alternative embodiments, the emitted table(s) comprise an export table; thus, in embodiments of method **600**, or the one or more tables comprise one or more export tables, and the first, second, and third entries comprise first, second, and third function exports.

[0096] Referring now to binary image consumption, and to FIG. 7, method **700** comprises an act **701** of initiating loading of a hybrid binary within a process that uses a compatibility ABI. In an example, the loader **106c** initiates loading of hybrid binary **108**, leveraging one of libraries **106a** that implement the EC ABI **107b**.

[0097] Method **700** also comprises an act **702** of determining that a machine type stored in the hybrid binary mismatches a machine type for the compatibility ABI. In some embodiments, act **702** comprises, during loading of the hybrid binary image, determining that a first machine type stored in a machine type field of the hybrid binary image mismatches a second machine type corresponding to the compatibility ABI under which the process is executing. In an example, as part of loading hybrid binary **108**, the loader **106c** copies a memory page containing the machine identifier **303** into system memory **104**. Then, the loader **106c** determines that a value of the machine identifier **303** mismatches a value expected for the EC ABI **107b** (e.g., a foreign ISA corresponding to the foreign ABI **107c**).

[0098] Method **700** also comprises an act **703** of, based on the mismatch, locating a fixup table. In some embodiments, act **703** comprises, based on determining that the first machine type mismatches the second machine type, locating, within the binary image, a fixup table defining a plurality of transformations to memory loaded from the hybrid binary image. In an example, the loader **106c** identifies fixup table **108c** within the hybrid binary **108**, such as by referring to a predefined address or offset within the hybrid binary **108**, or to a predefined address or offset within a portion of system memory **104** that is populated by memory page(s) loaded from the hybrid binary **108**.

[0099] Method **700** also comprises an act **704** of applying one or more transformation(s) within the fixup table to memory loaded from the hybrid binary. In some embodiments, act **704** comprises applying at least a portion of the plurality of transformations to the memory loaded from the hybrid binary image. In an example, the loader **106c** applies one or more transformations obtained from the fixup table **108c** to one or more portion(s) of system memory **104** that are populated by memory page(s) loaded from the hybrid binary **108**.

[0100] While the loader **106c** could apply all transformations in the fixup table **108c** at once, in embodiments the loader **106c** applies them on a page-by-page basis as those pages are loaded from the hybrid binary **108** (e.g., as part of a page fault handling routine). As such, in FIG. 7, act **704** is shown with an arrow leading back to act **704**, indicating that this act may be applied repeatedly as

memory pages are progressively loaded from the hybrid binary **108**. Thus, in embodiments, transformation(s) to the memory loaded from the hybrid binary image are applied to a memory page loaded from the hybrid binary image in connection with processing a memory page fault. [0101] In embodiments, the plurality of transformations obtained from the fixup table **108c** include a transformation that adjusts the machine type field to comprise a second machine type corresponding to the compatibility ABI. In an example, the loader **106c** applies a transformation to a memory location corresponding to the machine identifier **303** within system memory **104**, which adjusts the memory location to store a value matching a foreign ISA corresponding to the foreign ABI **107c**.

[0102] In embodiments, the plurality of transformations obtained from the fixup table **108c** include a transformation that modifies a call site calling a first compiled version of a non-folded function that is executable under the native ABI to instead call a second compiled version of the non-folded function that is executable under the compatibility ABI. In an example, the loader **106c** applies a memory location corresponding to the call site to replace a first memory address referencing the first compiled version of the non-folded function with a second memory address referencing a second compiled version of the non-folded function. While the call site, itself, could be transformed, in embodiments, the call site is structured to call the first compiled version of the non-folded function via a dispatch table **306** that references the first compiled version of the non-folded function. In these embodiments—rather than modifying the call site, itself—this transformation modifies the dispatch table **306** to reference the second compiled version of the non-folded function instead of the first compiled version of the non-folded function.

[0103] In embodiments, the hybrid binary image includes an entry point referencing the first compiled version of the non-folded function. In the embodiments, the plurality of transformations obtained from the fixup table **108c** (and applied in act **704**) include a transformation that adjusts the entry point to reference the second compiled version of the non-folded function.

[0104] As discussed in connection with FIG. **6**, in embodiments, method **600** emits one or more tables, such as corresponding to import and/or export tables, along with reference(s) to these table(s) that provide a native view. Thus, in the context of method **700**, in some embodiments the hybrid binary image includes one or more tables that comprise (i) a first zone referencing at least the non-folded function using the first compiled version of the non-folded function, (ii) a second zone referencing at least a folded function, and (iii) a third zone referencing at least the non-folded function using the second compiled version of the non-folded function. In these embodiments the hybrid binary image also includes a reference to the one or more tables that provides a native view of the one or more tables that includes the first zone and the second zone, while excluding the third zone. In order to provide a compatibility view when the binary image is loaded under the EC ABI **107b**, method **700** can comprise applying a transformation that adjusts the reference to the one or more tables to provide a compatibility view of the one or more tables that includes the second zone and the third zone, while excluding the first zone.

Dual Architecture Function Pointers Having Consistent Reference Addresses

[0105] FIG. **8** illustrates a flow chart of an example method **800** for using a common reference memory address when processing calls within a process that supports execution of both (i) native code targeting a native ABI that corresponds to the native ISA and that has a first CC, and (ii) foreign code targeting a foreign ABI that corresponds to a foreign ISA and that has a second CC.

[0106] As shown, method **800** comprises an act **801** of identifying a call to a callee function using a reference address. In one example, the dispatcher **112a** identifies a call from a native function in EC native code **108b** to a reference memory address. In another example, dispatcher **112b** identifies a call from a foreign function in foreign binary **109** (and which is being emulated by emulator **106b**) to a reference memory address.

[0107] Method **800** also comprises an act **802** of, using a lookup structure and the reference address, determining whether the callee function corresponds to a native ABI or to a foreign ABI.

In embodiments, act **802** comprises, based at least on identifying a call that targets a reference memory address for a callee function, determining whether the callee function corresponds to the native ABI or to the foreign ABI. In embodiments, the callee function is determined to correspond to the native ABI based a lookup structure indicating that the reference memory address is contained within a first memory range storing native code, and the callee function is determined to correspond to the foreign ABI based at least on the lookup structure indicating that the reference memory address is contained within a second memory range not storing native code. In an example, the dispatcher **112** (which could be either dispatcher **112a** when the caller is native, or dispatcher **112b** when the caller is foreign) consults EC lookup structure **307** to determine whether the reference memory address for the callee function is within a memory region corresponding to EC native code **108b** (in which case the callee is determined to correspond to the EC ABI **107b**), or whether the reference memory address for the callee function is not within a memory region corresponding to EC native code **108b** (in which case the callee is determined to correspond to foreign ABI **107c**).

[0108] As noted, in embodiments the EC lookup structure **307** is a bitmap, which uses one bit value to indicate whether or not a corresponding range of memory (e.g., a memory page) comprises EC native code **108b**. However, the EC lookup structure **307** could comprise an alternate data structure type, such as a hash table or a binary tree. Thus, in method **800**, the lookup structure comprises at least one of a bitmap, a hash table, or a binary tree.

[0109] As discussed, there could be an alternate scenario in which the callee function is determined to correspond to the native ABI—that is, when the dispatcher **112b** determines that the callee function contains a signature of a fast-forward sequence. Thus, in some embodiments of act **802**, the callee function is determined to correspond to the native ABI based on one of (i) a lookup structure indicating that the reference memory address is contained within a first memory range storing native code, or (ii) a fast-forward sequence being identified at the reference memory address. Additionally, in some embodiments of act **802**, when the lookup structure indicates that the reference memory address is contained within the second memory range not storing native code, act **802** comprises determining whether the fast-forward sequence is identifiable at the reference memory address. In these embodiments, the dispatcher **112b** obtains a new reference memory address from the fast-forward sequence (i.e., the true address of the native function) and uses that new reference memory address for a native call. Thus, in embodiments, when the fast-forward sequence is identified at the reference memory address, method **800** comprises updating the reference memory address with a new reference memory address obtained from the fast-forward sequence.

[0110] Method **800** also comprises an act **803** of initiating execution of the callee function. In embodiments act **803** comprises, based at least on the determining, initiating execution of the callee function. As shown, act **803** comprise performing one of an act **803a** of, when the caller is foreign and the callee is foreign, directly calling the reference address in an emulator; an act **803b** of, when the caller is foreign and the callee is native, locating and calling an entry thunk; an act **803c** of, when the caller is native and the callee is foreign, calling an exit thunk; or an act **803d** of, when the caller is native and the callee is native, directly calling the reference address.

[0111] In embodiments, act **803a** comprises, when a caller function corresponds to the foreign ABI, and when the callee function is determined to correspond to the foreign ABI, directly calling the callee function using the reference memory address within an emulator. In an example, based on a call from a foreign function, the dispatcher **112b** determines that the callee is also a foreign function. Thus, the caller is calling the true memory address of the callee function, and no thunk is needed, so the dispatcher **112b** directly invokes the callee function using the reference memory address within the emulator **106b**.

[0112] In embodiments, act **803b** comprises, when the caller function corresponds to the foreign ABI, and when the callee function is determined to correspond to the native ABI, calling an entry

thunk that (i) adapts a second CC to the first CC and then (ii) directly calls the callee function using the reference memory address. In an example, based on a call from a foreign function, the dispatcher **112b** determines that the callee is a native function. Thus, the dispatcher **112b** cannot call the reference memory address (i.e., of the callee function) directly, because that would bypass an entry thunk. Instead, the dispatcher **112b** locates a new reference address to the entry thunk and invokes the entry thunk using the new reference memory address. The entry thunk, in turn, adapts a CC of the foreign ABI **107c** to a CC of the EC ABI **107b**, and invokes the callee function using the original reference memory address.

[0113] As discussed, in embodiments, the dispatcher **112b** locates the new reference address to the entry thunk based on data (such as an offset or a direct address reference) contained in a block of memory immediately preceding the original reference memory address of the callee function. Thus, in embodiments, method **800** also comprises identifying a location of the entry thunk based at least on (i) reading a block of memory immediately preceding the reference memory address, and (ii) determining from the block of memory an offset or a pointer to the location of entry thunk.

[0114] In embodiments, act **803c** comprises, when the caller function corresponds to the native ABI, and when the callee function is determined to correspond to the foreign ABI, calling an exit thunk that (i) adapts a first CC of the native ABI to a second CC of the foreign ABI and then (ii) invokes the emulator to directly call the callee function using the reference memory address. In an example, based on a call from a native function, the dispatcher **112a** determines that the callee is a foreign function. Thus, the dispatcher **112a** cannot call the reference memory address (i.e., of the callee function) directly, because that would bypass an exit thunk. Instead, the dispatcher **112a** locates a new reference address to the exit thunk and invokes the exit thunk using the new reference memory address. The exit thunk, in turn, adapts a CC of the EC ABI **107b** to a CC of the foreign ABI **107c**, and invokes the emulator **106b**. The emulator **106b**, in turn, directly calls the callee function using the original reference memory address. As discussed, in embodiments the new reference address to the exit thunk is contained within the callee function, itself. Thus, in some embodiments of method **800**, a location of the exit thunk is contained within the caller function.

[0115] In embodiments, act **803d** comprises, when the caller function corresponds to the native ABI, and when the callee function is determined to correspond to the native ABI, directly calling the callee function using the reference memory address. In an example, based on a call from a native function, the dispatcher **112a** determines that the callee is also a native function. Thus, the caller is calling the true memory address of the callee function, and no thunk is needed, so the dispatcher **112a** directly invokes the callee function using the reference memory address.

[0116] Although the subject matter has been described in language specific to structural features and/or methodological acts, it is to be understood that the subject matter defined in the appended claims is not necessarily limited to the described features or acts described above, or the order of the acts described above. Rather, the described features and acts are disclosed as example forms of implementing the claims.

[0117] Embodiments of the present invention may comprise or utilize a special-purpose or general-purpose computer system (e.g., computer system **101**) that includes computer hardware, such as, for example, one or more processors (e.g., processor **102**) and system memory (e.g., system memory **104**), as discussed in greater detail below. Embodiments within the scope of the present invention also include physical and other computer-readable media for carrying or storing computer-executable instructions and/or data structures. Such computer-readable media can be any available media that can be accessed by a general-purpose or special-purpose computer system. Computer-readable media that store computer-executable instructions and/or data structures are computer storage media (e.g., durable storage **103**, system memory **104**). Computer-readable media that carry computer-executable instructions and/or data structures are transmission media. Thus, by way of example, and not limitation, embodiments of the invention can comprise at least two distinctly different kinds of computer-readable media: computer storage media and transmission

media.

[0118] Computer storage media are physical storage media that store computer-executable instructions and/or data structures. Physical storage media include computer hardware, such as RAM, ROM, EEPROM, solid state drives (“SSDs”), flash memory, phase-change memory (“PCM”), optical disk storage, magnetic disk storage or other magnetic storage devices, or any other hardware storage device(s) which can be used to store program code in the form of computer-executable instructions or data structures, which can be accessed and executed by a general-purpose or special-purpose computer system to implement the disclosed functionality of the invention.

[0119] Transmission media can include a network and/or data links which can be used to carry program code in the form of computer-executable instructions or data structures, and which can be accessed by a general-purpose or special-purpose computer system. A “network” is defined as one or more data links that enable the transport of electronic data between computer systems and/or modules and/or other electronic devices. When information is transferred or provided over a network or another communications connection (either hardwired, wireless, or a combination of hardwired or wireless) to a computer system, the computer system may view the connection as transmission media. Combinations of the above should also be included within the scope of computer-readable media.

[0120] Further, upon reaching various computer system components, program code in the form of computer-executable instructions or data structures can be transferred automatically from transmission media to computer storage media (or vice versa). For example, computer-executable instructions or data structures received over a network or data link can be buffered in RAM within a network interface module (e.g., a “NIC”), and then eventually transferred to computer system RAM and/or to less volatile computer storage media at a computer system. Thus, it should be understood that computer storage media can be included in computer system components that also (or even primarily) utilize transmission media.

[0121] Computer-executable instructions comprise, for example, instructions and data which, when executed at one or more processors, cause a general-purpose computer system, special-purpose computer system, or special-purpose processing device to perform a certain function or group of functions. Computer-executable instructions may be, for example, binaries, intermediate format instructions such as assembly language, or even source code.

[0122] Those skilled in the art will appreciate that the invention may be practiced in network computing environments with many types of computer system configurations, including, personal computers, desktop computers, laptop computers, message processors, hand-held devices, multi-processor systems, microprocessor-based or programmable consumer electronics, network PCs, minicomputers, mainframe computers, mobile telephones, PDAs, tablets, pagers, routers, switches, and the like. The invention may also be practiced in distributed system environments where local and remote computer systems, which are linked (either by hardwired data links, wireless data links, or by a combination of hardwired and wireless data links) through a network, both perform tasks. As such, in a distributed system environment, a computer system may include a plurality of constituent computer systems. In a distributed system environment, program modules may be located in both local and remote memory storage devices.

[0123] Those skilled in the art will also appreciate that the invention may be practiced in a cloud computing environment. Cloud computing environments may be distributed, although this is not required. When distributed, cloud computing environments may be distributed internationally within an organization and/or have components possessed across multiple organizations. In this description and the following claims, “cloud computing” is defined as a model for enabling on-demand network access to a shared pool of configurable computing resources (e.g., networks, servers, storage, applications, and services). The definition of “cloud computing” is not limited to any of the other numerous advantages that can be obtained from such a model when properly

deployed.

[0124] A cloud computing model can be composed of various characteristics, such as on-demand self-service, broad network access, resource pooling, rapid elasticity, measured service, and so forth. A cloud computing model may also come in the form of various service models such as, for example, Software as a Service (“SaaS”), Platform as a Service (“PaaS”), and Infrastructure as a Service (“IaaS”). The cloud computing model may also be deployed using different deployment models such as private cloud, community cloud, public cloud, hybrid cloud, and so forth.

[0125] Some embodiments, such as a cloud computing environment, may comprise a system that includes one or more hosts that are each capable of running one or more virtual machines. During operation, virtual machines emulate an operational computing system, supporting an OS and perhaps one or more other applications as well. In some embodiments, each host includes a hypervisor that emulates virtual resources for the virtual machines using physical resources that are abstracted from view of the virtual machines. The hypervisor also provides proper isolation between the virtual machines. Thus, from the perspective of any given virtual machine, the hypervisor provides the illusion that the virtual machine is interfacing with a physical resource, even though the virtual machine only interfaces with the appearance (e.g., a virtual resource) of a physical resource. Examples of physical resources including processing capacity, memory, disk space, network bandwidth, media drives, and so forth.

[0126] The present invention may be embodied in other specific forms without departing from its spirit or essential characteristics. The described embodiments are to be considered in all respects only as illustrative and not restrictive. The scope of the invention is, therefore, indicated by the appended claims rather than by the foregoing description. All changes which come within the meaning and range of equivalency of the claims are to be embraced within their scope. When introducing elements in the appended claims, the articles “a,” “an,” “the,” and “said” are intended to mean there are one or more of the elements. The terms “comprising,” “including,” and “having” are intended to be inclusive and mean that there may be additional elements other than the listed elements. The terms “set” and “subset” are intended to exclude an empty set, and thus “set” is defined as a non-empty set, and “subset” is defined as a non-empty subset.

Claims

1. (canceled)

2. A method, implemented in a computer system that includes a processor, the method comprising: determining, for a call that targets a reference memory address for a callee function, whether the callee function corresponds to a native application binary interface (ABI) that corresponds to a native instruction set architecture (ISA) of the processor and that has a first calling convention, or to a foreign ABI that corresponds to a foreign ISA and that has a second calling convention; and initiating execution of the callee function based on performing one of: directly calling the callee function using the reference memory address within an emulator, when a caller function corresponds to the foreign ABI and when the callee function is determined to correspond to the foreign ABI; calling an entry thunk that (i) adapts a second calling convention to the first calling convention and then (ii) directly calls the callee function using the reference memory address, when the caller function corresponds to the foreign ABI and when the callee function is determined to correspond to the native ABI; calling an exit thunk that (i) adapts a first calling convention of the native ABI to a second calling convention of the foreign ABI and then (ii) invokes the emulator to directly call the callee function using the reference memory address, when the caller function corresponds to the native ABI, and when the callee function is determined to correspond to the foreign ABI; or directly calling the callee function using the reference memory address, when the caller function corresponds to the native ABI, and when the callee function is determined to correspond to the native ABI.

3. The method of claim 2, wherein: the caller function corresponds to the foreign ABI and the callee function is determined to correspond to the foreign ABI; and initiating execution of the callee function comprises directly calling the callee function using the reference memory address within the emulator.
4. The method of claim 2, wherein: the caller function corresponds to the foreign ABI and the callee function is determined to correspond to the native ABI; and initiating execution of the callee function comprises calling the entry thunk.
5. The method of claim 4, wherein the method further comprises: identifying a location of the entry thunk based at least on (i) reading a block of memory immediately preceding the reference memory address, and (ii) determining from the block of memory an offset or a pointer to the location of entry thunk.
6. The method of claim 2, wherein: the caller function corresponds to the native ABI and the callee function is determined to correspond to the foreign ABI; and initiating execution of the callee function comprises calling the exit thunk.
7. The method of claim 6, wherein a location of the exit thunk is contained within the caller function.
8. The method of claim 2, wherein: the caller function corresponds to the native ABI and the callee function is determined to correspond to the native ABI; and initiating execution of the callee function comprises directly calling the callee function using the reference memory address.
9. The method of claim 2, wherein: the callee function is determined to correspond to the native ABI based on one of (i) a lookup structure indicating that the reference memory address is contained within a first memory range storing native code, or (ii) a fast-forward sequence being identified at the reference memory address; and the callee function is determined to correspond to the foreign ABI based at least on the lookup structure indicating that the reference memory address is contained within a second memory range not storing native code.
10. The method of claim 9, wherein the method further comprises: updating the reference memory address with a new reference memory address obtained from the fast-forward sequence, when the fast-forward sequence is identified at the reference memory address.
11. The method of claim 9, wherein the method further comprises: determining whether the fast-forward sequence is identifiable at the reference memory address, when the lookup structure indicates that the reference memory address is contained within the second memory range not storing native code.
12. A computer system comprising: a processor implementing a native instruction set architecture (ISA); and a hardware storage device that stores computer-executable instructions that are executable by the processor to cause the computer system to at least: determine, for a call that targets a reference memory address for a callee function, whether the callee function corresponds to a native application binary interface (ABI) that corresponds to the native ISA and that has a first calling convention, or to a foreign ABI that corresponds to a foreign ISA and that has a second calling convention; and initiate execution of the callee function based on performing one of: a direct call to the callee function using the reference memory address within an emulator, when a caller function corresponds to the foreign ABI and when the callee function is determined to correspond to the foreign ABI; a call to an entry thunk that (i) adapts a second calling convention to the first calling convention and then (ii) directly calls the callee function using the reference memory address, when the caller function corresponds to the foreign ABI and when the callee function is determined to correspond to the native ABI; a call to an exit thunk that (i) adapts a first calling convention of the native ABI to a second calling convention of the foreign ABI and then (ii) invokes the emulator to directly call the callee function using the reference memory address, when the caller function corresponds to the native ABI, and when the callee function is determined to correspond to the foreign ABI; or a direct call the callee function using the reference memory address, when the caller function corresponds to the native ABI, and when the callee function is

determined to correspond to the native ABI.

13. The computer system of claim 12, wherein: the caller function corresponds to the foreign ABI and the callee function is determined to correspond to the foreign ABI; and initiating execution of the callee function comprises directly calling the callee function using the reference memory address within the emulator.

14. The computer system of claim 12, wherein: the caller function corresponds to the foreign ABI and the callee function is determined to correspond to the native ABI; and initiating execution of the callee function comprises calling the entry thunk.

15. The computer system of claim 14, wherein the computer-executable instructions are also executable by the processor to cause the computer system to: identify a location of the entry thunk based at least on (i) reading a block of memory immediately preceding the reference memory address, and (ii) determining from the block of memory an offset or a pointer to the location of entry thunk.

16. The computer system of claim 12, wherein: the caller function corresponds to the native ABI and the callee function is determined to correspond to the foreign ABI; and initiating execution of the callee function comprises calling the exit thunk.

17. The computer system of claim 12, wherein: the caller function corresponds to the native ABI and the callee function is determined to correspond to the native ABI; and initiating execution of the callee function comprises directly calling the callee function using the reference memory address.

18. The computer system of claim 12, wherein: the callee function is determined to correspond to the native ABI based on one of (i) a lookup structure indicating that the reference memory address is contained within a first memory range storing native code, or (ii) a fast-forward sequence being identified at the reference memory address; and the callee function is determined to correspond to the foreign ABI based at least on the lookup structure indicating that the reference memory address is contained within a second memory range not storing native code.

19. The computer system of claim 18, wherein the computer-executable instructions are also executable by the processor to cause the computer system to: update the reference memory address with a new reference memory address obtained from the fast-forward sequence when the fast-forward sequence is identified at the reference memory address.

20. The computer system of claim 18, wherein the computer-executable instructions are also executable by the processor to cause the computer system to: determine whether the fast-forward sequence is identifiable at the reference memory address when the lookup structure indicates that the reference memory address is contained within the second memory range not storing native code.

21. A hardware storage device that stores computer-executable instructions that are executable by a processor to cause a computer system to at least: determine, for a call that targets a reference memory address for a callee function, whether the callee function corresponds to a native application binary interface (ABI) that corresponds to a native instruction set architecture (ISA) of the processor and that has a first calling convention, or to a foreign ABI that corresponds to a foreign ISA and that has a second calling convention; and initiate execution of the callee function based on performing one of: a direct call to the callee function using the reference memory address within an emulator, when a caller function corresponds to the foreign ABI and when the callee function is determined to correspond to the foreign ABI; a call to an entry thunk that (i) adapts a second calling convention to the first calling convention and then (ii) directly calls the callee function using the reference memory address, when the caller function corresponds to the foreign ABI and when the callee function is determined to correspond to the native ABI; a call to an exit thunk that (i) adapts a first calling convention of the native ABI to a second calling convention of the foreign ABI and then (ii) invokes the emulator to directly call the callee function using the reference memory address, when the caller function corresponds to the native ABI, and when the

callee function is determined to correspond to the foreign ABI; or a direct call the callee function using the reference memory address, when the caller function corresponds to the native ABI, and when the callee function is determined to correspond to the native ABI.
