

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250252650

Kind Code

A1

Publication Date

August 07, 2025

Inventor(s)

LABBE; HUGUES et al.

GRAPHICS ARCHITECTURE INCLUDING A NEURAL NETWORK PIPELINE

Abstract

One embodiment provides a graphics processor comprising a block of graphics cores and circuitry including a programmable neural network unit, the programmable neural network unit including one or more neural network hardware blocks, wherein a neural network hardware block includes circuitry to perform neural network operations and activation operations for a layer of a neural network, the programmable neural network unit addressable by cores within the block of graphics cores, wherein the programmable neural network unit is to configure one or more neural network hardware blocks with a meta-shader neural network, the meta-shader neural network to generate a texture for one of multiple types of terrain.

Inventors: LABBE; HUGUES (Granite Bay, CA), PALKE; DARREL (Portland, OR), ABDELHAK; SHERINE (BEAVERTON, OR), BOYCE; JILL (PORTLAND, OR), GEORGE; VARGHESE (FOLSOM, CA), JANUS; SCOTT (LOOMIS, CA), LAKE; ADAM (Portland, OR), LEI; ZHIJUN (HILLSBORO, OR), LI; ZHENGMIN (HILLSBORO, OR), MACPHERSON; MIKE (PORTLAND, OR), MARSHALL; CARL (PORTLAND, OR), PANNEER; SELVAKUMAR (PORTLAND, OR), SURTI; PRASOONKUMAR (FOLSOM, CA), VEERAMANI; KARTHIK (HILLSBORO, OR), VEMBAR; DEEPAK (PORTLAND, OR), SOMAYAZULU; VALLABHAJOSYULA SRINIVASA (PORTLAND, OR)

Applicant: Intel Corporation (Santa Clara, CA)

Family ID: 69406118

Assignee: Intel Corporation (Santa Clara, CA)

Appl. No.: 19/015406

Filed: January 09, 2025

Related U.S. Application Data

parent US continuation 18310015 20230501 parent-grant-document US 12229867 child US 19015406
parent US continuation 17500631 20211013 parent-grant-document US 11676322 child US 18310015
parent US continuation 16537140 20190809 parent-grant-document US 11151769 child US 17500631
us-provisional-application US 62717685 20180810
us-provisional-application US 62717593 20180810
us-provisional-application US 62717603 20180810

Publication Classification

Int. Cl.: **G06T15/00** (20110101); **G06N3/08** (20230101); **G06T1/20** (20060101); **G06T1/60** (20060101); **G06T15/40** (20110101); **G06T17/20** (20060101)

U.S. Cl.:

CPC **G06T15/005** (20130101); **G06N3/08** (20130101); **G06T1/20** (20130101); **G06T1/60** (20130101); **G06T15/40** (20130101); **G06T17/20** (20130101);

Background/Summary

CROSS-REFERENCE [0001] This application is a continuation of U.S. application Ser. No. 18/310,015, filed May 1, 2023, which is a continuation of U.S. application Ser. No. 17/500,631, filed Oct. 13, 2021, issued as U.S. Pat. No. 11,676,322, which is a continuation of U.S. application Ser. No. 16/537,140 filed on Aug. 9, 2019, issued as U.S. Pat. No. 11,151,769, which claims the benefit of priority to U.S. Provisional Application No. 62/717,685, U.S. Provisional Application No. 62/717,593, and U.S. Provisional Application No. 62/717,603, each of which were filed on Aug. 10, 2018 and are hereby incorporated herein by reference in their entirety.

FIELD

[0002] Embodiments relate generally to data processing and more particularly to scheduling and dispatching in a graphics architecture via neural networks

BACKGROUND OF THE DESCRIPTION

[0003] Conventional scheduling and dispatching in graphics hardware rely on prefetching subsequent consecutive cache lines anticipating the next instruction. These techniques can be improved upon via the use of neural networks.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0004] So that the manner in which the above recited features of the present embodiments can be understood in detail, a more particular description of the embodiments, briefly summarized above, may be had by reference to embodiments, some of which are illustrated in the appended drawings. It is to be noted, however, that the appended drawings illustrate only typical embodiments and are therefore not to be considered limiting of its scope.

[0005] FIG. 1 is a block diagram of a processing system, according to an embodiment;
[0006] FIG. 2 is a block diagram of a processor according to an embodiment;
[0007] FIG. 3 is a block diagram of a graphics processor, according to an embodiment;
[0008] FIG. 4 is a block diagram of a graphics processing engine of a graphics processor in accordance with some embodiments;
[0009] FIG. 5 is a block diagram of hardware logic of a graphics processor core, according to some embodiments described herein;
[0010] FIG. 6A-6B illustrate thread execution logic including an array of processing elements employed in a graphics processor core according to embodiments described herein;
[0011] FIG. 7 is a block diagram illustrating a graphics processor instruction formats according to some embodiments;
[0012] FIG. 8 is a block diagram of a graphics processor according to another embodiment;
[0013] FIG. 9A-9B illustrate a graphics processor command format and command sequence, according to some embodiments;
[0014] FIG. 10 illustrates exemplary graphics software architecture for a data processing system according to some embodiments;
[0015] FIG. 11A is a block diagram illustrating an IP core development system, according to an embodiment;
[0016] FIG. 11B illustrates a cross-section side view of an integrated circuit package assembly, according to some embodiments described herein;
[0017] FIG. 12 is a block diagram illustrating an exemplary system on a chip integrated circuit, according to an embodiment;
[0018] FIG. 13A-13B are block diagrams illustrating exemplary graphics processors for use within an SoC, according to embodiments described herein;
[0019] FIG. 14A-14B illustrate additional exemplary graphics processor logic according to embodiments described herein;
[0020] FIG. 15 illustrates a machine learning software stack, according to an embodiment;
[0021] FIG. 16A-16B illustrate layers of exemplary deep neural networks;
[0022] FIG. 17 illustrates an exemplary recurrent neural network;
[0023] FIG. 18 illustrates training and deployment of a deep neural network;
[0024] FIG. 19 is a block diagram illustrating distributed learning;
[0025] FIG. 20 is a block diagram of a computing device including a graphics processor, according to an embodiment;
[0026] FIG. 21 illustrates a tessellation module, according to an embodiment;
[0027] FIG. 22A-22B illustrate training and inference associated with an AI tessellation module;
[0028] FIG. 23A-23B illustrate training and inference for simulated tessellation using an AI tessellation module;
[0029] FIG. 24 illustrates a set of hardware blocks associated with a graphics processor described herein;
[0030] FIG. 25 illustrates a single layer hardware neural network block, according to an embodiment;
[0031] FIG. 26 illustrates a multiple layer hardware neural network block, according to an embodiment;
[0032] FIG. 27 illustrates a system in which geometry culling visibility is determined using machine learning, according to an embodiment;
[0033] FIG. 28 illustrates a meta-shader system, according to an embodiment.
[0034] FIG. 29 illustrates a graphics processing system including AI-based dynamic scheduling, according to an embodiment;
[0035] FIG. 30 illustrates a system for intelligent memory controller scheduling, according to an embodiment;

[0036] FIG. **31** illustrates a system having support for dynamic pipeline switching, according to embodiments described herein;

[0037] FIG. **32** illustrates a system having support for AI-based dynamic pipeline switching, according to embodiments described herein;

[0038] FIG. **33** illustrates a system to enable AI driven thread dispatch, according to an embodiment; and

[0039] FIG. **34A-34B** illustrate a system to enable AI-driven hardware memory prefetching, according to embodiments described herein.

DETAILED DESCRIPTION

[0040] Embodiments described herein provide techniques to improve the efficiency of GPU deep pipelines. A first embodiment provides for AI-based tessellation at the vertex and pixel level. A second embodiment provides for a processor including a neural network (NN) block that is addressable by graphics execution units (EU) within the processor. A third embodiment provides for geometry culling visibility using machine learning to avoid expensive pre-passes in fixed function hardware blocks. A fourth embodiment provides for a generative texture shader model in which a meta-shader can generate many different types of textures. A fifth embodiment provides for an AI-based dynamic scheduling on complex GPU architecture. A sixth embodiment provides for intelligent memory controller scheduling to support various types of memory requests. A seventh embodiment provides for an implementation of a neural network for a graphics pipeline. An eighth embodiment provides for a neural network switch to determine when to switch between a GPU pipeline and a neural network pipeline. A ninth embodiment provides for AI Driven Thread Dispatch. A tenth embodiment provides for AI-driven hardware memory prefetching.

[0041] For the purposes of explanation, numerous specific details are set forth to provide a thorough understanding of the various embodiments described below. However, it will be apparent to a skilled practitioner in the art that the embodiments may be practiced without some of these specific details. In other instances, well-known structures and devices are shown in block diagram form to avoid obscuring the underlying principles, and to provide a more thorough understanding of embodiments. Although some of the following embodiments are described with reference to a graphics processor, the techniques and teachings described herein may be applied to various types of circuits or semiconductor devices, including general purpose processing devices or graphic processing devices. Reference herein to “one embodiment” or “an embodiment” indicate that a particular feature, structure, or characteristic described in connection or association with the embodiment can be included in at least one of such embodiments. However, the appearances of the phrase “in one embodiment” in various places in the specification do not necessarily all refer to the same embodiment.

[0042] In the following description and claims, the terms “coupled” and “connected,” along with their derivatives, may be used. It should be understood that these terms are not intended as synonyms for each other. “Coupled” is used to indicate that two or more elements, which may or may not be in direct physical or electrical contact with each other, co-operate or interact with each other. “Connected” is used to indicate the establishment of communication between two or more elements that are coupled with each other.

[0043] In the description that follows, an overview of exemplary data processing system and processor logic is provided, along with details for the various embodiments presented herein. The following embodiments are described with reference to a graphics processor. However, similar techniques and teachings may be applied to other types of circuits or semiconductor devices, including but not limited to a many integrated core (MIC) processor, a CPUs, one or more instances of a field programmable gate array (FPGA), or other processing logic that is tailored to performing

System Overview

[0044] FIG. **1** is a block diagram of a processing system **100**, according to an embodiment. In

various embodiments the system **100** includes one or more processors **102** and one or more graphics processors **108**, and may be a single processor desktop system, a multiprocessor workstation system, or a server system having a large number of processors **102** or processor cores **107**. In one embodiment, the system **100** is a processing platform incorporated within a system-on-a-chip (SoC) integrated circuit for use in mobile, handheld, or embedded devices.

[0045] In one embodiment the system **100** can include or be incorporated within a server-based gaming platform, a game console, including a game and media console, a mobile gaming console, a handheld game console, or an online game console. In some embodiments the system **100** is a mobile phone, smart phone, tablet computing device or mobile Internet device. The processing system **100** can also include, couple with, or be integrated within a wearable device, such as a smart watch wearable device, smart eyewear device, augmented reality device, or virtual reality device. In some embodiments, the processing system **100** is a television or set top box device having one or more processors **102** and a graphical interface generated by one or more graphics processors **108**.

[0046] In some embodiments, the one or more processors **102** each include one or more processor cores **107** to process instructions which, when executed, perform operations for system and user software. In some embodiments, each of the one or more processor cores **107** is configured to process a specific instruction set **109**. In some embodiments, instruction set **109** may facilitate Complex Instruction Set Computing (CISC), Reduced Instruction Set Computing (RISC), or computing via a Very Long Instruction Word (VLIW). Multiple processor cores **107** may each process a different instruction set **109**, which may include instructions to facilitate the emulation of other instruction sets. Processor core **107** may also include other processing devices, such as a Digital Signal Processor (DSP).

[0047] In some embodiments, the processor **102** includes cache memory **104**. Depending on the architecture, the processor **102** can have a single internal cache or multiple levels of internal cache. In some embodiments, the cache memory is shared among various components of the processor **102**. In some embodiments, the processor **102** also uses an external cache (e.g., a Level-3 (L3) cache or Last Level Cache (LLC)) (not shown), which may be shared among processor cores **107** using known cache coherency techniques. A register file **106** is additionally included in processor **102** which may include different types of registers for storing different types of data (e.g., integer registers, floating point registers, status registers, and an instruction pointer register). Some registers may be general-purpose registers, while other registers may be specific to the design of the processor **102**.

[0048] In some embodiments, one or more processor(s) **102** are coupled with one or more interface bus(es) **110** to transmit communication signals such as address, data, or control signals between processor **102** and other components in the system **100**. The interface bus **110**, in one embodiment, can be a processor bus, such as a version of the Direct Media Interface (DMI) bus. However, processor busses are not limited to the DMI bus, and may include one or more Peripheral Component Interconnect buses (e.g., PCI, PCI Express), memory busses, or other types of interface busses. In one embodiment the processor(s) **102** include an integrated memory controller **116** and a platform controller hub **130**. The memory controller **116** facilitates communication between a memory device and other components of the system **100**, while the platform controller hub (PCH) **130** provides connections to I/O devices via a local I/O bus.

[0049] The memory device **120** can be a dynamic random access memory (DRAM) device, a static random access memory (SRAM) device, flash memory device, phase-change memory device, or some other memory device having suitable performance to serve as process memory. In one embodiment the memory device **120** can operate as system memory for the system **100**, to store data **122** and instructions **121** for use when the one or more processors **102** executes an application or process. Memory controller **116** also couples with an optional external graphics processor **112**, which may communicate with the one or more graphics processors **108** in processors **102** to

perform graphics and media operations. In some embodiments a display device **111** can connect to the processor(s) **102**. The display device **111** can be one or more of an internal display device, as in a mobile electronic device or a laptop device or an external display device attached via a display interface (e.g., DisplayPort, etc.). In one embodiment the display device **111** can be a head mounted display (HMD) such as a stereoscopic display device for use in virtual reality (VR) applications or augmented reality (AR) applications.

[0050] In some embodiments the platform controller hub **130** enables peripherals to connect to memory device **120** and processor **102** via a high-speed I/O bus. The I/O peripherals include, but are not limited to, an audio controller **146**, a network controller **134**, a firmware interface **128**, a wireless transceiver **126**, touch sensors **125**, a data storage device **124** (e.g., hard disk drive, flash memory, etc.). The data storage device **124** can connect via a storage interface (e.g., SATA) or via a peripheral bus, such as a Peripheral Component Interconnect bus (e.g., PCI, PCI Express). The touch sensors **125** can include touch screen sensors, pressure sensors, or fingerprint sensors. The wireless transceiver **126** can be a Wi-Fi transceiver, a Bluetooth transceiver, or a mobile network transceiver such as a 3G, 4G, or Long Term Evolution (LTE) transceiver. The firmware interface **128** enables communication with system firmware, and can be, for example, a unified extensible firmware interface (UEFI). The network controller **134** can enable a network connection to a wired network. In some embodiments, a high-performance network controller (not shown) couples with the interface bus **110**. The audio controller **146**, in one embodiment, is a multi-channel high definition audio controller. In one embodiment the system **100** includes an optional legacy I/O controller **140** for coupling legacy (e.g., Personal System 2 (PS/2)) devices to the system. The platform controller hub **130** can also connect to one or more Universal Serial Bus (USB) controllers **142** connect input devices, such as keyboard and mouse **143** combinations, a camera **144**, or other USB input devices.

[0051] It will be appreciated that the system **100** shown is exemplary and not limiting, as other types of data processing systems that are differently configured may also be used. For example, an instance of the memory controller **116** and platform controller hub **130** may be integrated into a discreet external graphics processor, such as the external graphics processor **112**. In one embodiment the platform controller hub **130** and/or memory controller **116** may be external to the one or more processor(s) **102**. For example, the system **100** can include an external memory controller **116** and platform controller hub **130**, which may be configured as a memory controller hub and peripheral controller hub within a system chipset that is in communication with the processor(s) **102**.

[0052] FIG. 2 is a block diagram of an embodiment of a processor **200** having one or more processor cores **202A-202N**, an integrated memory controller **214**, and an integrated graphics processor **208**. Those elements of FIG. 2 having the same reference numbers (or names) as the elements of any other figure herein can operate or function in any manner similar to that described elsewhere herein, but are not limited to such. Processor **200** can include additional cores up to and including additional core **202N** represented by the dashed lined boxes. Each of processor cores **202A-202N** includes one or more internal cache units **204A-204N**. In some embodiments each processor core also has access to one or more shared cached units **206**.

[0053] The internal cache units **204A-204N** and shared cache units **206** represent a cache memory hierarchy within the processor **200**. The cache memory hierarchy may include at least one level of instruction and data cache within each processor core and one or more levels of shared mid-level cache, such as a Level 2 (L2), Level 3 (L3), Level 4 (L4), or other levels of cache, where the highest level of cache before external memory is classified as the LLC. In some embodiments, cache coherency logic maintains coherency between the various cache units **206** and **204A-204N**.

[0054] In some embodiments, processor **200** may also include a set of one or more bus controller units **216** and a system agent core **210**. The one or more bus controller units **216** manage a set of peripheral buses, such as one or more PCI or PCI express busses. System agent core **210** provides

management functionality for the various processor components. In some embodiments, system agent core **210** includes one or more integrated memory controllers **214** to manage access to various external memory devices (not shown).

[0055] In some embodiments, one or more of the processor cores **202A-202N** include support for simultaneous multi-threading. In such embodiment, the system agent core **210** includes components for coordinating and operating cores **202A-202N** during multi-threaded processing. System agent core **210** may additionally include a power control unit (PCU), which includes logic and components to regulate the power state of processor cores **202A-202N** and graphics processor **208**.

[0056] In some embodiments, processor **200** additionally includes graphics processor **208** to execute graphics processing operations. In some embodiments, the graphics processor **208** couples with the set of shared cache units **206**, and the system agent core **210**, including the one or more integrated memory controllers **214**. In some embodiments, the system agent core **210** also includes a display controller **211** to drive graphics processor output to one or more coupled displays. In some embodiments, display controller **211** may also be a separate module coupled with the graphics processor via at least one interconnect, or may be integrated within the graphics processor **208**.

[0057] In some embodiments, a ring-based interconnect **212** is used to couple the internal components of the processor **200**. However, an alternative interconnect unit may be used, such as a point-to-point interconnect, a switched interconnect, or other techniques, including techniques well known in the art. In some embodiments, graphics processor **208** couples with the ring-based interconnect **212** via an I/O link **213**.

[0058] The exemplary I/O link **213** represents at least one of multiple varieties of I/O interconnects, including an on package I/O interconnect which facilitates communication between various processor components and a high-performance embedded memory module **218**, such as an eDRAM module. In some embodiments, each of the processor cores **202A-202N** and graphics processor **208** use embedded memory modules **218** as a shared Last Level Cache.

[0059] In some embodiments, processor cores **202A-202N** are homogenous cores executing the same instruction set architecture. In another embodiment, processor cores **202A-202N** are heterogeneous in terms of instruction set architecture (ISA), where one or more of processor cores **202A-202N** execute a first instruction set, while at least one of the other cores executes a subset of the first instruction set or a different instruction set. In one embodiment processor cores **202A-202N** are heterogeneous in terms of microarchitecture, where one or more cores having a relatively higher power consumption couple with one or more power cores having a lower power consumption. Additionally, processor **200** can be implemented on one or more chips or as an SoC integrated circuit having the illustrated components, in addition to other components.

[0060] FIG. **3** is a block diagram of a graphics processor **300**, which may be a discrete graphics processing unit, or may be a graphics processor integrated with a plurality of processing cores. In some embodiments, the graphics processor communicates via a memory mapped I/O interface to registers on the graphics processor and with commands placed into the processor memory. In some embodiments, graphics processor **300** includes a memory interface **314** to access memory. Memory interface **314** can be an interface to local memory, one or more internal caches, one or more shared external caches, and/or to system memory.

[0061] In some embodiments, graphics processor **300** also includes a display controller **302** to drive display output data to a display device **320**. Display controller **302** includes hardware for one or more overlay planes for the display and composition of multiple layers of video or user interface elements. The display device **320** can be an internal or external display device. In one embodiment the display device **320** is a head mounted display device, such as a virtual reality (VR) display device or an augmented reality (AR) display device. In some embodiments, graphics processor **300** includes a video codec engine **306** to encode, decode, or transcode media to, from, or between one

or more media encoding formats, including, but not limited to Moving Picture Experts Group (MPEG) formats such as MPEG-2, Advanced Video Coding (AVC) formats such as H.264/MPEG-4 AVC, as well as the Society of Motion Picture & Television Engineers (SMPTE) 421M/VC-1, and Joint Photographic Experts Group (JPEG) formats such as JPEG, and Motion JPEG (MJPEG) formats.

[0062] In some embodiments, graphics processor **300** includes a block image transfer (BLIT) engine **304** to perform two-dimensional (2D) rasterizer operations including, for example, bit-boundary block transfers. However, in one embodiment, 2D graphics operations are performed using one or more components of graphics processing engine (GPE **310**). In some embodiments, GPE **310** is a compute engine for performing graphics operations, including three-dimensional (3D) graphics operations and media operations.

[0063] In some embodiments, GPE **310** includes a 3D pipeline **312** for performing 3D operations, such as rendering three-dimensional images and scenes using processing functions that act upon 3D primitive shapes (e.g., rectangle, triangle, etc.). The 3D pipeline **312** includes programmable and fixed function elements that perform various tasks within the element and/or spawn execution threads to a 3D/Media sub-system **315**. While 3D pipeline **312** can be used to perform media operations, an embodiment of GPE **310** also includes a media pipeline **316** that is specifically used to perform media operations, such as video post-processing and image enhancement.

[0064] In some embodiments, media pipeline **316** includes fixed function or programmable logic units to perform one or more specialized media operations, such as video decode acceleration, video de-interlacing, and video encode acceleration in place of, or on behalf of video codec engine **306**. In some embodiments, media pipeline **316** additionally includes a thread spawning unit to spawn threads for execution on 3D/Media sub-system **315**. The spawned threads perform computations for the media operations on one or more graphics execution units included in 3D/Media sub-system **315**.

[0065] In some embodiments, 3D/Media sub-system **315** includes logic for executing threads spawned by 3D pipeline **312** and media pipeline **316**. In one embodiment, the pipelines send thread execution requests to 3D/Media sub-system **315**, which includes thread dispatch logic for arbitrating and dispatching the various requests to available thread execution resources. The execution resources include an array of graphics execution units to process the 3D and media threads. In some embodiments, 3D/Media sub-system **315** includes one or more internal caches for thread instructions and data. In some embodiments, the subsystem also includes shared memory, including registers and addressable memory, to share data between threads and to store output data.

Graphics Processing Engine

[0066] FIG. **4** is a block diagram of a graphics processing engine of a graphics processor in accordance with some embodiments. In one embodiment, the graphics processing engine (GPE **410**) is a version of the GPE **310** shown in FIG. **3**. Elements of FIG. **4** having the same reference numbers (or names) as the elements of any other figure herein can operate or function in any manner similar to that described elsewhere herein, but are not limited to such. For example, the 3D pipeline **312** and media pipeline **316** of FIG. **3** are illustrated. The media pipeline **316** is optional in some embodiments of the GPE **410** and may not be explicitly included within the GPE **410**. For example, in at least one embodiment a separate media and/or image processor is coupled to the GPE **410**.

[0067] In some embodiments, GPE **410** couples with or includes a command streamer **403**, which provides a command stream to the 3D pipeline **312** and/or media pipelines **316**. In some embodiments, command streamer **403** is coupled with memory, which can be system memory, or one or more of internal cache memory and shared cache memory. In some embodiments, command streamer **403** receives commands from the memory and sends the commands to 3D pipeline **312** and/or media pipeline **316**. The commands are directives fetched from a ring buffer, which stores commands for the 3D pipeline **312** and media pipeline **316**. In one embodiment, the ring buffer can

additionally include batch command buffers storing batches of multiple commands. The commands for the 3D pipeline **312** can also include references to data stored in memory, such as but not limited to vertex and geometry data for the 3D pipeline **312** and/or image data and memory objects for the media pipeline **316**. The 3D pipeline **312** and media pipeline **316** process the commands and data by performing operations via logic within the respective pipelines or by dispatching one or more execution threads to a graphics core array **414**. In one embodiment the graphics core array **414** include one or more blocks of graphics cores (e.g., graphics core(s) **415A**, graphics core(s) **415B**), each block including one or more graphics cores. Each graphics core includes a set of graphics execution resources that includes general-purpose and graphics specific execution logic to perform graphics and compute operations, as well as fixed function texture processing and/or machine learning and artificial intelligence acceleration logic.

[0068] In various embodiments the 3D pipeline **312** includes fixed function and programmable logic to process one or more shader programs, such as vertex shaders, geometry shaders, pixel shaders, fragment shaders, compute shaders, or other shader programs, by processing the instructions and dispatching execution threads to the graphics core array **414**. The graphics core array **414** provides a unified block of execution resources for use in processing these shader programs. Multi-purpose execution logic (e.g., execution units) within the graphics core(s) **415A-414B** of the graphics core array **414** includes support for various 3D API shader languages and can execute multiple simultaneous execution threads associated with multiple shaders.

[0069] In some embodiments the graphics core array **414** also includes execution logic to perform media functions, such as video and/or image processing. In one embodiment, the execution units additionally include general-purpose logic that is programmable to perform parallel general-purpose computational operations, in addition to graphics processing operations. The general-purpose logic can perform processing operations in parallel or in conjunction with general-purpose logic within the processor core(s) **107** of FIG. **1** or core **202A-202N** as in FIG. **2**.

[0070] Output data generated by threads executing on the graphics core array **414** can output data to memory in a unified return buffer (URB **418**). The URB **418** can store data for multiple threads. In some embodiments the URB **418** may be used to send data between different threads executing on the graphics core array **414**. In some embodiments the URB **418** may additionally be used for synchronization between threads on the graphics core array and fixed function logic within the shared function logic **420**.

[0071] In some embodiments, graphics core array **414** is scalable, such that the array includes a variable number of graphics cores, each having a variable number of execution units based on the target power and performance level of GPE **410**. In one embodiment the execution resources are dynamically scalable, such that execution resources may be enabled or disabled as needed.

[0072] The graphics core array **414** couples with shared function logic **420** that includes multiple resources that are shared between the graphics cores in the graphics core array. The shared functions within the shared function logic **420** are hardware logic units that provide specialized supplemental functionality to the graphics core array **414**. In various embodiments, shared function logic **420** includes but is not limited to sampler **421**, math **422**, and inter-thread communication (ITC) **423** logic. Additionally, some embodiments implement one or more cache(s) **425** within the shared function logic **420**.

[0073] A shared function is implemented where the demand for a given specialized function is insufficient for inclusion within the graphics core array **414**. Instead a single instantiation of that specialized function is implemented as a stand-alone entity in the shared function logic **420** and shared among the execution resources within the graphics core array **414**. The precise set of functions that are shared between the graphics core array **414** and included within the graphics core array **414** varies across embodiments. In some embodiments, specific shared functions within the shared function logic **420** that are used extensively by the graphics core array **414** may be included within shared function logic **416** within the graphics core array **414**. In various embodiments, the

shared function logic **416** within the graphics core array **414** can include some or all logic within the shared function logic **420**. In one embodiment, all logic elements within the shared function logic **420** may be duplicated within the shared function logic **416** of the graphics core array **414**. In one embodiment the shared function logic **420** is excluded in favor of the shared function logic **416** within the graphics core array **414**.

[0074] FIG. **5** is a block diagram of hardware logic of a graphics processor core **500**, according to some embodiments described herein. Elements of FIG. **5** having the same reference numbers (or names) as the elements of any other figure herein can operate or function in any manner similar to that described elsewhere herein, but are not limited to such. The illustrated graphics processor core **500**, in some embodiments, is included within the graphics core array **414** of FIG. **4**. The graphics processor core **500**, sometimes referred to as a core slice, can be one or multiple graphics cores within a modular graphics processor. The graphics processor core **500** is exemplary of one graphics core slice, and a graphics processor as described herein may include multiple graphics core slices based on target power and performance envelopes. Each graphics core **500** can include a fixed function block **530** coupled with multiple sub-cores **501A-501F**, also referred to as sub-slices, that include modular blocks of general-purpose and fixed function logic.

[0075] In some embodiments the fixed function block **530** includes a geometry/fixed function pipeline **536** that can be shared by all sub-cores in the graphics processor **500**, for example, in lower performance and/or lower power graphics processor implementations. In various embodiments, the geometry/fixed function pipeline **536** includes a 3D fixed function pipeline (e.g., 3D pipeline **312** as in FIG. **3** and FIG. **4**) a video front-end unit, a thread spawner and thread dispatcher, and a unified return buffer manager, which manages unified return buffers, such as the URB **418** of FIG. **4**.

[0076] In one embodiment the fixed function block **530** also includes a graphics SoC interface **537**, a graphics microcontroller **538**, and a media pipeline **539**. The graphics SoC interface **537** provides an interface between the graphics core **500** and other processor cores within a system on a chip integrated circuit. The graphics microcontroller **538** is a programmable sub-processor that is configurable to manage various functions of the graphics processor **500**, including thread dispatch, scheduling, and pre-emption. The media pipeline **539** (e.g., media pipeline **316** of FIG. **3** and FIG. **4**) includes logic to facilitate the decoding, encoding, pre-processing, and/or post-processing of multimedia data, including image and video data. The media pipeline **539** implement media operations via requests to compute or sampling logic within the sub-cores **501-501F**.

[0077] In one embodiment the graphics SoC interface **537** enables the graphics core **500** to communicate with general-purpose application processor cores (e.g., CPUs) and/or other components within an SoC, including memory hierarchy elements such as a shared last level cache memory, the system RAM, and/or embedded on-chip or on-package DRAM. The graphics SoC interface **537** can also enable communication with fixed function devices within the SoC, such as camera imaging pipelines, and enables the use of and/or implements global memory atomics that may be shared between the graphics core **500** and CPUs within the SoC. The graphics SoC interface **537** can also implement power management controls for the graphics core **500** and enable an interface between a clock domain of the graphic core **500** and other clock domains within the SoC. In one embodiment the graphics SoC interface **537** enables receipt of command buffers from a command streamer and global thread dispatcher that are configured to provide commands and instructions to each of one or more graphics cores within a graphics processor. The commands and instructions can be dispatched to the media pipeline **539**, when media operations are to be performed, or a geometry and fixed function pipeline (e.g., geometry and fixed function pipeline **536**, geometry and fixed function pipeline **514**) when graphics processing operations are to be performed.

[0078] The graphics microcontroller **538** can be configured to perform various scheduling and management tasks for the graphics core **500**. In one embodiment the graphics microcontroller **538**

can perform graphics and/or compute workload scheduling on the various graphics parallel engines within execution unit (EU) arrays **502A-502F**, **504A-504F** within the sub-cores **501A-501F**. In this scheduling model, host software executing on a CPU core of an SoC including the graphics core **500** can submit workloads one of multiple graphic processor doorbells, which invokes a scheduling operation on the appropriate graphics engine. Scheduling operations include determining which workload to run next, submitting a workload to a command streamer, pre-empting existing workloads running on an engine, monitoring progress of a workload, and notifying host software when a workload is complete. In one embodiment the graphics microcontroller **538** can also facilitate low-power or idle states for the graphics core **500**, providing the graphics core **500** with the ability to save and restore registers within the graphics core **500** across low-power state transitions independently from the operating system and/or graphics driver software on the system. [0079] The graphics core **500** may have greater than or fewer than the illustrated sub-cores **501A-501F**, up to N modular sub-cores. For each set of N sub-cores, the graphics core **500** can also include shared function logic **510**, shared and/or cache memory **512**, a geometry/fixed function pipeline **514**, as well as additional fixed function logic **516** to accelerate various graphics and compute processing operations. The shared function logic **510** can include logic units associated with the shared function logic **420** of FIG. 4 (e.g., sampler, math, and/or inter-thread communication logic) that can be shared by each N sub-cores within the graphics core **500**. The shared and/or cache memory **512** can be a last-level cache for the set of N sub-cores **501A-501F** within the graphics core **500**, and can also serve as shared memory that is accessible by multiple sub-cores. The geometry/fixed function pipeline **514** can be included instead of the geometry/fixed function pipeline **536** within the fixed function block **530** and can include the same or similar logic units.

[0080] In one embodiment the graphics core **500** includes additional fixed function logic **516** that can include various fixed function acceleration logic for use by the graphics core **500**. In one embodiment the additional fixed function logic **516** includes an additional geometry pipeline for use in position only shading. In position-only shading, two geometry pipelines exist, the full geometry pipeline within the geometry/fixed function pipeline **536**, and a cull pipeline, which is an additional geometry pipeline which may be included within the additional fixed function logic **516**. In one embodiment the cull pipeline is a trimmed down version of the full geometry pipeline. The full pipeline and the cull pipeline can execute different instances of the same application, each instance having a separate context. Position only shading can hide long cull runs of discarded triangles, enabling shading to be completed earlier in some instances. For example, in one embodiment the cull pipeline logic within the additional fixed function logic **516** can execute position shaders in parallel with the main application and generally generates critical results faster than the full pipeline, as the cull pipeline fetches and shades only the position attribute of the vertices, without performing rasterization and rendering of the pixels to the frame buffer. The cull pipeline can use the generated critical results to compute visibility information for all the triangles without regard to whether those triangles are culled. The full pipeline (which in this instance may be referred to as a replay pipeline) can consume the visibility information to skip the culled triangles to shade only the visible triangles that are finally passed to the rasterization phase.

[0081] In one embodiment the additional fixed function logic **516** can also include machine-learning acceleration logic, such as fixed function matrix multiplication logic, for implementations including optimizations for machine learning training or inferencing.

[0082] Within each graphics sub-core **501A-501F** includes a set of execution resources that may be used to perform graphics, media, and compute operations in response to requests by graphics pipeline, media pipeline, or shader programs. The graphics sub-cores **501A-501F** include multiple EU arrays **502A-502F**, **504A-504F**, thread dispatch and inter-thread communication (TD/IC) logic **503A-503F**, a 3D (e.g., texture) sampler **505A-505F**, a media sampler **506A-506F**, a shader processor **507A-507F**, and shared local memory (SLM) **508A-508F**. The EU arrays **502A-502F**,

504A-504F each include multiple execution units, which are general-purpose graphics processing units capable of performing floating-point and integer/fixed-point logic operations in service of a graphics, media, or compute operation, including graphics, media, or compute shader programs. The TD/IC logic **503A-503F** performs local thread dispatch and thread control operations for the execution units within a sub-core and facilitate communication between threads executing on the execution units of the sub-core. The 3D sampler **505A-505F** can read texture or other 3D graphics related data into memory. The 3D sampler can read texture data differently based on a configured sample state and the texture format associated with a given texture. The media sampler **506A-506F** can perform similar read operations based on the type and format associated with media data. In one embodiment, each graphics sub-core **501A-501F** can alternately include a unified 3D and media sampler. Threads executing on the execution units within each of the sub-cores **501A-501F** can make use of shared local memory **508A-508F** within each sub-core, to enable threads executing within a thread group to execute using a common pool of on-chip memory.

Execution Units

[0083] FIG. **6A-6B** illustrate thread execution logic **600** including an array of processing elements employed in a graphics processor core according to embodiments described herein. Elements of FIG. **6A-6B** having the same reference numbers (or names) as the elements of any other figure herein can operate or function in any manner similar to that described elsewhere herein, but are not limited to such. FIG. **6A** illustrates an overview of thread execution logic **600**, which can include a variant of the hardware logic illustrated with each sub-core **501A-501F** of FIG. **5**. FIG. **6B** illustrates exemplary internal details of an execution unit.

[0084] As illustrated in FIG. **6A**, in some embodiments thread execution logic **600** includes a shader processor **602**, a thread dispatcher **604**, instruction cache **606**, a scalable execution unit array including a plurality of execution units **608A-608N**, a sampler **610**, a data cache **612**, and a data port **614**. In one embodiment the scalable execution unit array can dynamically scale by enabling or disabling one or more execution units (e.g., any of execution unit **608A**, **608B**, **608C**, **608D**, through **608N-1** and **608N**) based on the computational requirements of a workload. In one embodiment the included components are interconnected via an interconnect fabric that links to each of the components. In some embodiments, thread execution logic **600** includes one or more connections to memory, such as system memory or cache memory, through one or more of instruction cache **606**, data port **614**, sampler **610**, and execution units **608A-608N**. In some embodiments, each execution unit (e.g. **608A**) is a stand-alone programmable general-purpose computational unit that is capable of executing multiple simultaneous hardware threads while processing multiple data elements in parallel for each thread. In various embodiments, the array of execution units **608A-608N** is scalable to include any number individual execution units.

[0085] In some embodiments, the execution units **608A-608N** are primarily used to execute shader programs. A shader processor **602** can process the various shader programs and dispatch execution threads associated with the shader programs via a thread dispatcher **604**. In one embodiment the thread dispatcher includes logic to arbitrate thread initiation requests from the graphics and media pipelines and instantiate the requested threads on one or more execution unit in the execution units **608A-608N**. For example, a geometry pipeline can dispatch vertex, tessellation, or geometry shaders to the thread execution logic for processing. In some embodiments, thread dispatcher **604** can also process runtime thread spawning requests from the executing shader programs.

[0086] In some embodiments, the execution units **608A-608N** support an instruction set that includes native support for many standard 3D graphics shader instructions, such that shader programs from graphics libraries (e.g., Direct 3D and OpenGL) are executed with a minimal translation. The execution units support vertex and geometry processing (e.g., vertex programs, geometry programs, vertex shaders), pixel processing (e.g., pixel shaders, fragment shaders) and general-purpose processing (e.g., compute and media shaders). Each of the execution units **608A-608N** is capable of multi-issue single instruction multiple data (SIMD) execution and multi-

threaded operation enables an efficient execution environment in the face of higher latency memory accesses. Each hardware thread within each execution unit has a dedicated high-bandwidth register file and associated independent thread-state. Execution is multi-issue per clock to pipelines capable of integer, single and double precision floating point operations, SIMD branch capability, logical operations, transcendental operations, and other miscellaneous operations. While waiting for data from memory or one of the shared functions, dependency logic within the execution units **608A-608N** causes a waiting thread to sleep until the requested data has been returned. While the waiting thread is sleeping, hardware resources may be devoted to processing other threads. For example, during a delay associated with a vertex shader operation, an execution unit can perform operations for a pixel shader, fragment shader, or another type of shader program, including a different vertex shader.

[0087] Each execution unit in execution units **608A-608N** operates on arrays of data elements. The number of data elements is the “execution size,” or the number of channels for the instruction. An execution channel is a logical unit of execution for data element access, masking, and flow control within instructions. The number of channels may be independent of the number of physical Arithmetic Logic Units (ALUs) or Floating Point Units (FPUs) for a particular graphics processor. In some embodiments, execution units **608A-608N** support integer and floating-point data types.

[0088] The execution unit instruction set includes SIMD instructions. The various data elements can be stored as a packed data type in a register and the execution unit will process the various elements based on the data size of the elements. For example, when operating on a 256-bit wide vector, the 256 bits of the vector are stored in a register and the execution unit operates on the vector as four separate 64-bit packed data elements (Quad-Word (QW) size data elements), eight separate 32-bit packed data elements (Double Word (DW) size data elements), sixteen separate 16-bit packed data elements (Word (W) size data elements), or thirty-two separate 8-bit data elements (byte (B) size data elements). However, different vector widths and register sizes are possible.

[0089] In one embodiment one or more execution units can be combined into a fused graphics execution unit **609A-609N** having thread control logic (**607A-607N**) that is common to the fused EUs. Multiple EUs can be fused into an EU group. Each EU in the fused EU group can be configured to execute a separate SIMD hardware thread. The number of EUs in a fused EU group can vary according to embodiments. Additionally, various SIMD widths can be performed per-EU, including but not limited to SIMD8, SIMD16, and SIMD32. Each fused graphics execution unit **609A-609N** includes at least two execution units. For example, fused graphics execution unit **609A** includes a first EU **608A**, second EU **608B**, and thread control logic **607A** that is common to the first EU **608A** and the second EU **608B**. The thread control logic **607A** controls threads executed on the fused graphics execution unit **609A**, allowing each EU within the fused graphics execution units **609A-609N** to execute using a common instruction pointer register.

[0090] One or more internal instruction caches (e.g., **606**) are included in the thread execution logic **600** to cache thread instructions for the execution units. In some embodiments, one or more data caches (e.g., **612**) are included to cache thread data during thread execution. In some embodiments, a sampler **610** is included to provide texture sampling for 3D operations and media sampling for media operations. In some embodiments, sampler **610** includes specialized texture or media sampling functionality to process texture or media data during the sampling process before providing the sampled data to an execution unit.

[0091] During execution, the graphics and media pipelines send thread initiation requests to thread execution logic **600** via thread spawning and dispatch logic. Once a group of geometric objects has been processed and rasterized into pixel data, pixel processor logic (e.g., pixel shader logic, fragment shader logic, etc.) within the shader processor **602** is invoked to further compute output information and cause results to be written to output surfaces (e.g., color buffers, depth buffers, stencil buffers, etc.). In some embodiments, a pixel shader or fragment shader calculates the values of the various vertex attributes that are to be interpolated across the rasterized object. In some

embodiments, pixel processor logic within the shader processor **602** then executes an application programming interface (API)-supplied pixel or fragment shader program. To execute the shader program, the shader processor **602** dispatches threads to an execution unit (e.g., **608A**) via thread dispatcher **604**. In some embodiments, shader processor **602** uses texture sampling logic in the sampler **610** to access texture data in texture maps stored in memory. Arithmetic operations on the texture data and the input geometry data compute pixel color data for each geometric fragment or discards one or more pixels from further processing.

[0092] In some embodiments, the data port **614** provides a memory access mechanism for the thread execution logic **600** to output processed data to memory for further processing on a graphics processor output pipeline. In some embodiments, the data port **614** includes or couples to one or more cache memories (e.g., data cache **612**) to cache data for memory access via the data port.

[0093] As illustrated in FIG. **6B**, a graphics execution unit **608** can include an instruction fetch unit **637**, a general register file array (GRF) **624**, an architectural register file array (ARF) **626**, a thread arbiter **622**, a send unit **630**, a branch unit **632**, a set of SIMD floating point units (FPUs) **634**, and in one embodiment a set of dedicated integer SIMD ALUs **635**. The GRF **624** and ARF **626** includes the set of general register files and architecture register files associated with each simultaneous hardware thread that may be active in the graphics execution unit **608**. In one embodiment, per thread architectural state is maintained in the ARF **626**, while data used during thread execution is stored in the GRF **624**. The execution state of each thread, including the instruction pointers for each thread, can be held in thread-specific registers in the ARF **626**.

[0094] In one embodiment the graphics execution unit **608** has an architecture that is a combination of Simultaneous Multi-Threading (SMT) and fine-grained Interleaved Multi-Threading (IMT). The architecture has a modular configuration that can be fine-tuned at design time based on a target number of simultaneous threads and number of registers per execution unit, where execution unit resources are divided across logic used to execute multiple simultaneous threads.

[0095] In one embodiment, the graphics execution unit **608** can co-issue multiple instructions, which may each be different instructions. The thread arbiter **622** of the graphics execution unit **608** can dispatch the instructions to one of the send unit **630**, branch unit **632**, or SIMD FPU(s) **634** for execution. Each execution thread can access 128 general-purpose registers within the GRF **624**, where each register can store 32 bytes, accessible as a SIMD 8-clement vector of 32-bit data elements. In one embodiment, each execution unit thread has access to 4 Kbytes within the GRF **624**, although embodiments are not so limited, and greater or fewer register resources may be provided in other embodiments. In one embodiment up to seven threads can execute simultaneously, although the number of threads per execution unit can also vary according to embodiments. In an embodiment in which seven threads may access 4 Kbytes, the GRF **624** can store a total of 28 Kbytes. Flexible addressing modes can permit registers to be addressed together to build effectively wider registers or to represent strided rectangular block data structures.

[0096] In one embodiment, memory operations, sampler operations, and other longer-latency system communications are dispatched via “send” instructions that are executed by the message passing send unit **630**. In one embodiment, branch instructions are dispatched to a dedicated branch unit **632** to facilitate SIMD divergence and eventual convergence.

[0097] In one embodiment the graphics execution unit **608** includes one or more SIMD floating point units (FPU(s)) **634** to perform floating-point operations. In one embodiment, the FPU(s) **634** also support integer computation. In one embodiment the FPU(s) **634** can SIMD execute up to M number of 32-bit floating-point (or integer) operations, or SIMD execute up to 2M 16-bit integer or 16-bit floating-point operations. In one embodiment, at least one of the FPU(s) provides extended math capability to support high-throughput transcendental math functions and double precision 64-bit floating-point. In some embodiments, a set of 8-bit integer SIMD ALUs **635** are also present, and may be specifically optimized to perform operations associated with machine learning computations.

[0098] In one embodiment, arrays of multiple instances of the graphics execution unit **608** can be instantiated in a graphics sub-core grouping (e.g., a sub-slice). For scalability, product architects can choose the exact number of execution units per sub-core grouping. In one embodiment the graphics execution unit **608** can execute instructions across a plurality of execution channels. In a further embodiment, each thread executed on the graphics execution unit **608** is executed on a different channel.

[0099] FIG. 7 is a block diagram illustrating graphics processor instruction formats **700** according to some embodiments. In one or more embodiment, the graphics processor execution units support an instruction set having instructions in multiple formats. The solid lined boxes illustrate the components that are generally included in an execution unit instruction, while the dashed lines include components that are optional or that are only included in a sub-set of the instructions. In some embodiments, the graphics processor instruction formats **700** described and illustrated are macro-instructions, in that they are instructions supplied to the execution unit, as opposed to micro-operations resulting from instruction decode once the instruction is processed.

[0100] In some embodiments, the graphics processor execution units natively support instructions in a 128-bit instruction format **710**. A 64-bit compacted instruction format **730** is available for some instructions based on the selected instruction, instruction options, and number of operands. The native 128-bit instruction format **710** provides access to all instruction options, while some options and operations are restricted in the 64-bit format **730**. The native instructions available in the 64-bit format **730** vary by embodiment. In some embodiments, the instruction is compacted in part using a set of index values in an index field **713**. The execution unit hardware references a set of compaction tables based on the index values and uses the compaction table outputs to reconstruct a native instruction in the 128-bit instruction format **710**.

[0101] For each format, instruction opcode **712** defines the operation that the execution unit is to perform. The execution units execute each instruction in parallel across the multiple data elements of each operand. For example, in response to an add instruction the execution unit performs a simultaneous add operation across each color channel representing a texture element or picture element. By default, the execution unit performs each instruction across all data channels of the operands. In some embodiments, instruction control field **714** enables control over certain execution options, such as channels selection (e.g., predication) and data channel order (e.g., swizzle). For instructions in the 128-bit instruction format **710** an exec-size field **716** limits the number of data channels that will be executed in parallel. In some embodiments, exec-size field **716** is not available for use in the 64-bit compact instruction format **730**.

[0102] Some execution unit instructions have up to three operands including two source operands, src0 **720**, src1 **722**, and one destination **718**. In some embodiments, the execution units support dual destination instructions, where one of the destinations is implied. Data manipulation instructions can have a third source operand (e.g., SRC2 **724**), where the instruction opcode **712** determines the number of source operands. An instruction's last source operand can be an immediate (e.g., hard-coded) value passed with the instruction.

[0103] In some embodiments, the 128-bit instruction format **710** includes an access/address mode field **726** specifying, for example, whether direct register addressing mode or indirect register addressing mode is used. When direct register addressing mode is used, the register address of one or more operands is directly provided by bits in the instruction.

[0104] In some embodiments, the 128-bit instruction format **710** includes an access/address mode field **726**, which specifies an address mode and/or an access mode for the instruction. In one embodiment the access mode is used to define a data access alignment for the instruction. Some embodiments support access modes including a 16-byte aligned access mode and a 1-byte aligned access mode, where the byte alignment of the access mode determines the access alignment of the instruction operands. For example, when in a first mode, the instruction may use byte-aligned addressing for source and destination operands and when in a second mode, the instruction may use

16-byte-aligned addressing for all source and destination operands.

[0105] In one embodiment, the address mode portion of the access/address mode field **726** determines whether the instruction is to use direct or indirect addressing. When direct register addressing mode is used bits in the instruction directly provide the register address of one or more operands. When indirect register addressing mode is used, the register address of one or more operands may be computed based on an address register value and an address immediate field in the instruction.

[0106] In some embodiments instructions are grouped based on opcode **712** bit-fields to simplify Opcode decode **740**. For an 8-bit opcode, bits **4**, **5**, and **6** allow the execution unit to determine the type of opcode. The precise opcode grouping shown is merely an example. In some embodiments, a move and logic opcode group **742** includes data movement and logic instructions (e.g., move (mov), compare (cmp)). In some embodiments, move and logic group **742** shares the five most significant bits (MSB), where move (mov) instructions are in the form of 0000xxxxb and logic instructions are in the form of 0001xxxxb. A flow control instruction group **744** (e.g., call, jump (jmp)) includes instructions in the form of 0010xxxxb (e.g., 0x20). A miscellaneous instruction group **746** includes a mix of instructions, including synchronization instructions (e.g., wait, send) in the form of 0011xxxxb (e.g., 0x30). A parallel math instruction group **748** includes component-wise arithmetic instructions (e.g., add, multiply (mul)) in the form of 0100xxxxb (e.g., 0x40). The parallel math group **748** performs the arithmetic operations in parallel across data channels. The vector math group **750** includes arithmetic instructions (e.g., dp4) in the form of 0101xxxxb (e.g., 0x50). The vector math group performs arithmetic such as dot product calculations on vector operands.

Graphics Pipeline

[0107] FIG. **8** is a block diagram of another embodiment of a graphics processor **800**. Elements of FIG. **8** having the same reference numbers (or names) as the elements of any other figure herein can operate or function in any manner similar to that described elsewhere herein, but are not limited to such.

[0108] In some embodiments, graphics processor **800** includes a geometry pipeline **820**, a media pipeline **830**, a display engine **840**, thread execution logic **850**, and a render output pipeline **870**. In some embodiments, graphics processor **800** is a graphics processor within a multi-core processing system that includes one or more general-purpose processing cores. The graphics processor is controlled by register writes to one or more control registers (not shown) or via commands issued to graphics processor **800** via a ring interconnect **802**. In some embodiments, ring interconnect **802** couples graphics processor **800** to other processing components, such as other graphics processors or general-purpose processors. Commands from ring interconnect **802** are interpreted by a command streamer **803**, which supplies instructions to individual components of the geometry pipeline **820** or the media pipeline **830**.

[0109] In some embodiments, command streamer **803** directs the operation of a vertex fetcher **805** that reads vertex data from memory and executes vertex-processing commands provided by command streamer **803**. In some embodiments, vertex fetcher **805** provides vertex data to a vertex shader **807**, which performs coordinate space transformation and lighting operations to each vertex. In some embodiments, vertex fetcher **805** and vertex shader **807** execute vertex-processing instructions by dispatching execution threads to execution units **852A-852B** via a thread dispatcher **831**.

[0110] In some embodiments, execution units **852A-852B** are an array of vector processors having an instruction set for performing graphics and media operations. In some embodiments, execution units **852A-852B** have an attached L1 cache **851** that is specific for each array or shared between the arrays. The cache can be configured as a data cache, an instruction cache, or a single cache that is partitioned to contain data and instructions in different partitions.

[0111] In some embodiments, geometry pipeline **820** includes tessellation components to perform

hardware-accelerated tessellation of 3D objects. In some embodiments, a programmable hull shader **811** configures the tessellation operations. A programmable domain shader **817** provides back-end evaluation of tessellation output. A tessellator **813** operates at the direction of hull shader **811** and contains special purpose logic to generate a set of detailed geometric objects based on a coarse geometric model that is provided as input to geometry pipeline **820**. In some embodiments, if tessellation is not used, tessellation components (e.g., hull shader **811**, tessellator **813**, and domain shader **817**) can be bypassed.

[0112] In some embodiments, complete geometric objects can be processed by a geometry shader **819** via one or more threads dispatched to execution units **852A-852B**, or can proceed directly to the clipper **829**. In some embodiments, the geometry shader operates on entire geometric objects, rather than vertices or patches of vertices as in previous stages of the graphics pipeline. If the tessellation is disabled, the geometry shader **819** receives input from the vertex shader **807**. In some embodiments, geometry shader **819** is programmable by a geometry shader program to perform geometry tessellation if the tessellation units are disabled.

[0113] Before rasterization, a clipper **829** processes vertex data. The clipper **829** may be a fixed function clipper or a programmable clipper having clipping and geometry shader functions. In some embodiments, a rasterizer and depth test component **873** in the render output pipeline **870** dispatches pixel shaders to convert the geometric objects into per pixel representations. In some embodiments, pixel shader logic is included in thread execution logic **850**. In some embodiments, an application can bypass the rasterizer and depth test component **873** and access un-rasterized vertex data via a stream out unit **823**.

[0114] The graphics processor **800** has an interconnect bus, interconnect fabric, or some other interconnect mechanism that allows data and message passing amongst the major components of the processor. In some embodiments, execution units **852A-852B** and associated logic units (e.g., L1 cache **851**, sampler **854**, texture cache **858**, etc.) interconnect via a data port **856** to perform memory access and communicate with render output pipeline components of the processor. In some embodiments, sampler **854**, caches **851**, **858** and execution units **852A-852B** each have separate memory access paths. In one embodiment the texture cache **858** can also be configured as a sampler cache.

[0115] In some embodiments, render output pipeline **870** contains a rasterizer and depth test component **873** that converts vertex-based objects into an associated pixel-based representation. In some embodiments, the rasterizer logic includes a windower/masker unit to perform fixed function triangle and line rasterization. An associated render cache **878** and depth cache **879** are also available in some embodiments. A pixel operations component **877** performs pixel-based operations on the data, though in some instances, pixel operations associated with 2D operations (e.g. bit block image transfers with blending) are performed by the 2D engine **841**, or substituted at display time by the display controller **843** using overlay display planes. In some embodiments, a shared L3 cache **875** is available to all graphics components, allowing the sharing of data without the use of main system memory.

[0116] In some embodiments, graphics processor media pipeline **830** includes a media engine **837** and a video front-end **834**. In some embodiments, video front-end **834** receives pipeline commands from the command streamer **803**. In some embodiments, media pipeline **830** includes a separate command streamer. In some embodiments, video front-end **834** processes media commands before sending the command to the media engine **837**. In some embodiments, media engine **837** includes thread spawning functionality to spawn threads for dispatch to thread execution logic **850** via thread dispatcher **831**.

[0117] In some embodiments, graphics processor **800** includes a display engine **840**. In some embodiments, display engine **840** is external to processor **800** and couples with the graphics processor via the ring interconnect **802**, or some other interconnect bus or fabric. In some embodiments, display engine **840** includes a 2D engine **841** and a display controller **843**. In some

embodiments, display engine **840** contains special purpose logic capable of operating independently of the 3D pipeline. In some embodiments, display controller **843** couples with a display device (not shown), which may be a system integrated display device, as in a laptop computer, or an external display device attached via a display device connector.

[0118] In some embodiments, the geometry pipeline **820** and media pipeline **830** are configurable to perform operations based on multiple graphics and media programming interfaces and are not specific to any one application programming interface (API). In some embodiments, driver software for the graphics processor translates API calls that are specific to a particular graphics or media library into commands that can be processed by the graphics processor. In some embodiments, support is provided for the Open Graphics Library (OpenGL), Open Computing Language (OpenCL), and/or Vulkan graphics and compute API, all from the Khronos Group. In some embodiments, support may also be provided for the Direct3D library from the Microsoft Corporation. In some embodiments, a combination of these libraries may be supported. Support may also be provided for the Open Source Computer Vision Library (OpenCV). A future API with a compatible 3D pipeline would also be supported if a mapping can be made from the pipeline of the future API to the pipeline of the graphics processor.

Graphics Pipeline Programming

[0119] FIG. **9A** is a block diagram illustrating a graphics processor command format **900** according to some embodiments. FIG. **9B** is a block diagram illustrating a graphics processor command sequence **910** according to an embodiment. The solid lined boxes in FIG. **9A** illustrate the components that are generally included in a graphics command while the dashed lines include components that are optional or that are only included in a sub-set of the graphics commands. The exemplary graphics processor command format **900** of FIG. **9A** includes data fields to identify a client **902**, a command operation code (opcode) **904**, and data **906** for the command. A sub-opcode **905** and a command size **908** are also included in some commands.

[0120] In some embodiments, client **902** specifies the client unit of the graphics device that processes the command data. In some embodiments, a graphics processor command parser examines the client field of each command to condition the further processing of the command and route the command data to the appropriate client unit. In some embodiments, the graphics processor client units include a memory interface unit, a render unit, a 2D unit, a 3D unit, and a media unit. Each client unit has a corresponding processing pipeline that processes the commands. Once the command is received by the client unit, the client unit reads the opcode **904** and, if present, sub-opcode **905** to determine the operation to perform. The client unit performs the command using information in data field **906**. For some commands an explicit command size **908** is expected to specify the size of the command. In some embodiments, the command parser automatically determines the size of at least some of the commands based on the command opcode. In some embodiments commands are aligned via multiples of a double word.

[0121] The flow diagram in FIG. **9B** illustrates an exemplary graphics processor command sequence **910**. In some embodiments, software or firmware of a data processing system that features an embodiment of a graphics processor uses a version of the command sequence shown to set up, execute, and terminate a set of graphics operations. A sample command sequence is shown and described for purposes of example only as embodiments are not limited to these specific commands or to this command sequence. Moreover, the commands may be issued as batch of commands in a command sequence, such that the graphics processor will process the sequence of commands in at least partially concurrence.

[0122] In some embodiments, the graphics processor command sequence **910** may begin with a pipeline flush command **912** to cause any active graphics pipeline to complete the currently pending commands for the pipeline. In some embodiments, the 3D pipeline **922** and the media pipeline **924** do not operate concurrently. The pipeline flush is performed to cause the active graphics pipeline to complete any pending commands. In response to a pipeline flush, the

command parser for the graphics processor will pause command processing until the active drawing engines complete pending operations and the relevant read caches are invalidated. Optionally, any data in the render cache that is marked ‘dirty’ can be flushed to memory. In some embodiments, pipeline flush command **912** can be used for pipeline synchronization or before placing the graphics processor into a low power state.

[0123] In some embodiments, a pipeline select command **913** is used when a command sequence requires the graphics processor to explicitly switch between pipelines. In some embodiments, a pipeline select command **913** is required only once within an execution context before issuing pipeline commands unless the context is to issue commands for both pipelines. In some embodiments, a pipeline flush command **912** is required immediately before a pipeline switch via the pipeline select command **913**.

[0124] In some embodiments, a pipeline control command **914** configures a graphics pipeline for operation and is used to program the 3D pipeline **922** and the media pipeline **924**. In some embodiments, pipeline control command **914** configures the pipeline state for the active pipeline. In one embodiment, the pipeline control command **914** is used for pipeline synchronization and to clear data from one or more cache memories within the active pipeline before processing a batch of commands.

[0125] In some embodiments, commands for return buffer state **916** are used to configure a set of return buffers for the respective pipelines to write data. Some pipeline operations require the allocation, selection, or configuration of one or more return buffers into which the operations write intermediate data during processing. In some embodiments, the graphics processor also uses one or more return buffers to store output data and to perform cross thread communication. In some embodiments, the return buffer state **916** includes selecting the size and number of return buffers to use for a set of pipeline operations.

[0126] The remaining commands in the command sequence differ based on the active pipeline for operations. Based on a pipeline determination **920**, the command sequence is tailored to the 3D pipeline **922** beginning with the 3D pipeline state **930** or the media pipeline **924** beginning at the media pipeline state **940**.

[0127] The commands to configure the 3D pipeline state **930** include 3D state setting commands for vertex buffer state, vertex element state, constant color state, depth buffer state, and other state variables that are to be configured before 3D primitive commands are processed. The values of these commands are determined at least in part based on the particular 3D API in use. In some embodiments, 3D pipeline state **930** commands are also able to selectively disable or bypass certain pipeline elements if those elements will not be used.

[0128] In some embodiments, 3D primitive **932** command is used to submit 3D primitives to be processed by the 3D pipeline. Commands and associated parameters that are passed to the graphics processor via the 3D primitive **932** command are forwarded to the vertex fetch function in the graphics pipeline. The vertex fetch function uses the 3D primitive **932** command data to generate vertex data structures. The vertex data structures are stored in one or more return buffers. In some embodiments, 3D primitive **932** command is used to perform vertex operations on 3D primitives via vertex shaders. To process vertex shaders, 3D pipeline **922** dispatches shader execution threads to graphics processor execution units.

[0129] In some embodiments, 3D pipeline **922** is triggered via an execute **934** command or event. In some embodiments, a register write triggers command execution. In some embodiments execution is triggered via a ‘go’ or ‘kick’ command in the command sequence. In one embodiment, command execution is triggered using a pipeline synchronization command to flush the command sequence through the graphics pipeline. The 3D pipeline will perform geometry processing for the 3D primitives. Once operations are complete, the resulting geometric objects are rasterized and the pixel engine colors the resulting pixels. Additional commands to control pixel shading and pixel back end operations may also be included for those operations.

[0130] In some embodiments, the graphics processor command sequence **910** follows the media pipeline **924** path when performing media operations. In general, the specific use and manner of programming for the media pipeline **924** depends on the media or compute operations to be performed. Specific media decode operations may be offloaded to the media pipeline during media decode. In some embodiments, the media pipeline can also be bypassed, and media decode can be performed in whole or in part using resources provided by one or more general-purpose processing cores. In one embodiment, the media pipeline also includes elements for general-purpose graphics processor unit (GPGPU) operations, where the graphics processor is used to perform SIMD vector operations using computational shader programs that are not explicitly related to the rendering of graphics primitives.

[0131] In some embodiments, media pipeline **924** is configured in a similar manner as the 3D pipeline **922**. A set of commands to configure the media pipeline state **940** are dispatched or placed into a command queue before the media object commands **942**. In some embodiments, commands for the media pipeline state **940** include data to configure the media pipeline elements that will be used to process the media objects. This includes data to configure the video decode and video encode logic within the media pipeline, such as encode or decode format. In some embodiments, commands for the media pipeline state **940** also support the use of one or more pointers to “indirect” state elements that contain a batch of state settings.

[0132] In some embodiments, media object commands **942** supply pointers to media objects for processing by the media pipeline. The media objects include memory buffers containing video data to be processed. In some embodiments, all media pipeline states must be valid before issuing a media object command **942**. Once the pipeline state is configured and media object commands **942** are queued, the media pipeline **924** is triggered via an execute command **944** or an equivalent execute event (e.g., register write). Output from media pipeline **924** may then be post processed by operations provided by the 3D pipeline **922** or the media pipeline **924**. In some embodiments, GPGPU operations are configured and executed in a similar manner as media operations.

Graphics Software Architecture

[0133] FIG. **10** illustrates exemplary graphics software architecture for a data processing system **1000** according to some embodiments. In some embodiments, software architecture includes a 3D graphics application **1010**, an operating system **1020**, and at least one processor **1030**. In some embodiments, processor **1030** includes a graphics processor **1032** and one or more general-purpose processor core(s) **1034**. The graphics application **1010** and operating system **1020** each execute in the system memory **1050** of the data processing system.

[0134] In some embodiments, 3D graphics application **1010** contains one or more shader programs including shader instructions **1012**. The shader language instructions may be in a high-level shader language, such as the High Level Shader Language (HLSL) or the OpenGL Shader Language (GLSL). The application also includes executable instructions **1014** in a machine language suitable for execution by the general-purpose processor core **1034**. The application also includes graphics objects **1016** defined by vertex data.

[0135] In some embodiments, operating system **1020** is a Microsoft® Windows® operating system from the Microsoft Corporation, a proprietary UNIX-like operating system, or an open source UNIX-like operating system using a variant of the Linux kernel. The operating system **1020** can support a graphics API **1022** such as the Direct3D API, the OpenGL API, or the Vulkan API. When the Direct3D API is in use, the operating system **1020** uses a front-end shader compiler **1024** to compile any shader instructions **1012** in HLSL into a lower-level shader language. The compilation may be a just-in-time (JIT) compilation or the application can perform shader pre-compilation. In some embodiments, high-level shaders are compiled into low-level shaders during the compilation of the 3D graphics application **1010**. In some embodiments, the shader instructions **1012** are provided in an intermediate form, such as a version of the Standard Portable Intermediate Representation (SPIR) used by the Vulkan API.

[0136] In some embodiments, user mode graphics driver **1026** contains a back-end shader compiler **1027** to convert the shader instructions **1012** into a hardware specific representation. When the OpenGL API is in use, shader instructions **1012** in the GLSL high-level language are passed to a user mode graphics driver **1026** for compilation. In some embodiments, user mode graphics driver **1026** uses operating system kernel mode functions **1028** to communicate with a kernel mode graphics driver **1029**. In some embodiments, kernel mode graphics driver **1029** communicates with graphics processor **1032** to dispatch commands and instructions.

IP Core Implementations

[0137] One or more aspects of at least one embodiment may be implemented by representative code stored on a machine-readable medium which represents and/or defines logic within an integrated circuit such as a processor. For example, the machine-readable medium may include instructions which represent various logic within the processor. When read by a machine, the instructions may cause the machine to fabricate the logic to perform the techniques described herein. Such representations, known as “IP cores,” are reusable units of logic for an integrated circuit that may be stored on a tangible, machine-readable medium as a hardware model that describes the structure of the integrated circuit. The hardware model may be supplied to various customers or manufacturing facilities, which load the hardware model on fabrication machines that manufacture the integrated circuit. The integrated circuit may be fabricated such that the circuit performs operations described in association with any of the embodiments described herein.

[0138] FIG. **11A** is a block diagram illustrating an IP core development system **1100** that may be used to manufacture an integrated circuit to perform operations according to an embodiment. The IP core development system **1100** may be used to generate modular, re-usable designs that can be incorporated into a larger design or used to construct an entire integrated circuit (e.g., an SOC integrated circuit). A design facility **1130** can generate a software simulation **1110** of an IP core design in a high-level programming language (e.g., C/C++). The software simulation **1110** can be used to design, test, and verify the behavior of the IP core using a simulation model **1112**. The simulation model **1112** may include functional, behavioral, and/or timing simulations. A register transfer level (RTL) design **1115** can then be created or synthesized from the simulation model **1112**. The RTL design **1115** is an abstraction of the behavior of the integrated circuit that models the flow of digital signals between hardware registers, including the associated logic performed using the modeled digital signals. In addition to an RTL design **1115**, lower-level designs at the logic level or transistor level may also be created, designed, or synthesized. Thus, the particular details of the initial design and simulation may vary.

[0139] The RTL design **1115** or equivalent may be further synthesized by the design facility into a hardware model **1120**, which may be in a hardware description language (HDL), or some other representation of physical design data. The HDL may be further simulated or tested to verify the IP core design. The IP core design can be stored for delivery to a 3rd party fabrication facility **1165** using non-volatile memory **1140** (e.g., hard disk, flash memory, or any non-volatile storage medium). Alternatively, the IP core design may be transmitted (e.g., via the Internet) over a wired connection **1150** or wireless connection **1160**. The fabrication facility **1165** may then fabricate an integrated circuit that is based at least in part on the IP core design. The fabricated integrated circuit can be configured to perform operations in accordance with at least one embodiment described herein.

[0140] FIG. **11B** illustrates a cross-section side view of an integrated circuit package assembly **1170**, according to some embodiments described herein. The integrated circuit package assembly **1170** illustrates an implementation of one or more processor or accelerator devices as described herein. The package assembly **1170** includes multiple units of hardware logic **1172**, **1174** connected to a substrate **1180**. The logic **1172**, **1174** may be implemented at least partly in configurable logic or fixed-functionality logic hardware and can include one or more portions of any of the processor core(s), graphics processor(s), or other accelerator devices described herein. Each unit of logic

1172, 1174 can be implemented within a semiconductor die and coupled with the substrate **1180** via an interconnect structure **1173**. The interconnect structure **1173** may be configured to route electrical signals between the logic **1172, 1174** and the substrate **1180**, and can include interconnects such as, but not limited to bumps or pillars. In some embodiments, the interconnect structure **1173** may be configured to route electrical signals such as, for example, input/output (I/O) signals and/or power or ground signals associated with the operation of the logic **1172, 1174**. In some embodiments, the substrate **1180** is an epoxy-based laminate substrate. The package assembly **1170** may include other suitable types of substrates in other embodiments. The package assembly **1170** can be connected to other electrical devices via a package interconnect **1183**. The package interconnect **1183** may be coupled to a surface of the substrate **1180** to route electrical signals to other electrical devices, such as a motherboard, other chipset, or multi-chip module.

[0141] In some embodiments, the units of logic **1172, 1174** are electrically coupled with a bridge **1182** that is configured to route electrical signals between the logic **1172, 1174**. The bridge **1182** may be a dense interconnect structure that provides a route for electrical signals. The bridge **1182** may include a bridge substrate composed of glass or a suitable semiconductor material. Electrical routing features can be formed on the bridge substrate to provide a chip-to-chip connection between the logic **1172, 1174**.

[0142] Although two units of logic **1172, 1174** and a bridge **1182** are illustrated, embodiments described herein may include more or fewer logic units on one or more dies. The one or more dies may be connected by zero or more bridges, as the bridge **1182** may be excluded when the logic is included on a single die. Alternatively, multiple dies or units of logic can be connected by one or more bridges. Additionally, multiple logic units, dies, and bridges can be connected together in other possible configurations, including three-dimensional configurations.

Exemplary System on a Chip Integrated Circuit

[0143] FIG. **12**, FIG. **13A-13B**, and FIG. **14A-14B** illustrated exemplary integrated circuits and associated graphics processors that may be fabricated using one or more IP cores, according to various embodiments described herein. In addition to what is illustrated, other logic and circuits may be included, including additional graphics processors/cores, peripheral interface controllers, or general-purpose processor cores.

[0144] FIG. **12** is a block diagram illustrating an exemplary system on a chip integrated circuit **1200** that may be fabricated using one or more IP cores, according to an embodiment. Exemplary integrated circuit **1200** includes one or more application processor(s) **1205** (e.g., CPUs), at least one graphics processor **1210**, and may additionally include an image processor **1215** and/or a video processor **1220**, any of which may be a modular IP core from the same or multiple different design facilities. Integrated circuit **1200** includes peripheral or bus logic including a USB controller **1225**, UART controller **1230**, an SPI/SDIO controller **1235**, and an I2S/I2C controller **1240**.

Additionally, the integrated circuit can include a display device **1245** coupled to one or more of a high-definition multimedia interface (HDMI) controller **1250** and a mobile industry processor interface (MIPI) display interface **1255**. Storage may be provided by a flash memory subsystem **1260** including flash memory and a flash memory controller. Memory interface may be provided via a memory controller **1265** for access to SDRAM or SRAM memory devices. Some integrated circuits additionally include an embedded security engine **1270**.

[0145] FIG. **13A-13B** are block diagrams illustrating exemplary graphics processors for use within an SoC, according to embodiments described herein. FIG. **13A** illustrates an exemplary graphics processor **1310** of a system on a chip integrated circuit that may be fabricated using one or more IP cores, according to an embodiment. FIG. **13B** illustrates an additional exemplary graphics processor **1340** of a system on a chip integrated circuit that may be fabricated using one or more IP cores, according to an embodiment. Graphics processor **1310** of FIG. **13A** is an example of a low power graphics processor core. Graphics processor **1340** of FIG. **13B** is an example of a higher performance graphics processor core. Each of the graphics processors **1310, 1340** can be variants

of the graphics processor **1210** of FIG. **12**.

[0146] As shown in FIG. **13A**, graphics processor **1310** includes a vertex processor **1305** and one or more fragment processor(s) **1315A-1315N** (e.g., **1315A**, **1315B**, **1315C**, **1315D**, through **1315N-1**, and **1315N**). Graphics processor **1310** can execute different shader programs via separate logic, such that the vertex processor **1305** is optimized to execute operations for vertex shader programs, while the one or more fragment processor(s) **1315A-1315N** execute fragment (e.g., pixel) shading operations for fragment or pixel shader programs. The vertex processor **1305** performs the vertex processing stage of the 3D graphics pipeline and generates primitives and vertex data. The fragment processor(s) **1315A-1315N** use the primitive and vertex data generated by the vertex processor **1305** to produce a framebuffer that is displayed on a display device. In one embodiment, the fragment processor(s) **1315A-1315N** are optimized to execute fragment shader programs as provided for in the OpenGL API, which may be used to perform similar operations as a pixel shader program as provided for in the Direct 3D API.

[0147] Graphics processor **1310** additionally includes one or more memory management units (MMUs) **1320A-1320B**, cache(s) **1325A-1325B**, and circuit interconnect(s) **1330A-1330B**. The one or more MMU(s) **1320A-1320B** provide for virtual to physical address mapping for the graphics processor **1310**, including for the vertex processor **1305** and/or fragment processor(s) **1315A-1315N**, which may reference vertex or image/texture data stored in memory, in addition to vertex or image/texture data stored in the one or more cache(s) **1325A-1325B**. In one embodiment the one or more MMU(s) **1320A-1320B** may be synchronized with other MMUs within the system, including one or more MMUs associated with the one or more application processor(s) **1205**, image processor **1215**, and/or video processor **1220** of FIG. **12**, such that each processor **1205-1220** can participate in a shared or unified virtual memory system. The one or more circuit interconnect(s) **1330A-1330B** enable graphics processor **1310** to interface with other IP cores within the SoC, either via an internal bus of the SoC or via a direct connection, according to embodiments.

[0148] As shown FIG. **13B**, graphics processor **1340** includes the one or more MMU(s) **1320A-1320B**, cache(s) **1325A-1325B**, and circuit interconnect(s) **1330A-1330B** of the graphics processor **1310** of FIG. **13A**. Graphics processor **1340** includes one or more shader core(s) **1355A-1355N** (e.g., **1355A**, **1355B**, **1355C**, **1355D**, **1355E**, **1355F**, through **1355N-1**, and **1355N**), which provides for a unified shader core architecture in which a single core or type or core can execute all types of programmable shader code, including shader program code to implement vertex shaders, fragment shaders, and/or compute shaders. The exact number of shader cores present can vary among embodiments and implementations. Additionally, graphics processor **1340** includes an inter-core task manager **1345**, which acts as a thread dispatcher to dispatch execution threads to one or more shader cores **1355A-1355N** and a tiling unit **1358** to accelerate tiling operations for tile-based rendering, in which rendering operations for a scene are subdivided in image space, for example to exploit local spatial coherence within a scene or to optimize use of internal caches.

[0149] FIG. **14A-14B** illustrate additional exemplary graphics processor logic according to embodiments described herein. FIG. **14A** illustrates a graphics core **1400** that may be included within the graphics processor **1210** of FIG. **12** and may be a unified shader core **1355A-1355N** as in FIG. **13B**. FIG. **14B** illustrates a GPGPU **1430** suitable for deployment on a multi-chip module.

[0150] As shown in FIG. **14A**, the graphics core **1400** includes a shared instruction cache **1402**, a texture unit **1418**, and a cache/shared memory **1420** that are common to the execution resources within the graphics core **1400**. The graphics core **1400** can include multiple slices **1401A-1401N** or partition for each core, and a graphics processor can include multiple instances of the graphics core **1400**. The slices **1401A-1401N** can include support logic including a local instruction cache **1404A-1404N**, a thread scheduler **1406A-1406N**, a thread dispatcher **1408A-1408N**, and a set of registers **1410A**. To perform logic operations, the slices **1401A-1401N** can include a set of additional function units (AFUs **1412A-1412N**), floating-point units (FPU **1414A-1414N**), integer

arithmetic logic units (ALUs **1416-1416N**), address computational units (ACU **1413A-1413N**), double-precision floating-point units (DPFPU **1415A-1415N**), and matrix processing units (MPU **1417A-1417N**).

[0151] Some of the computational units operate at a specific precision. For example, the FPU **1414A-1414N** can perform single-precision (32-bit) and half-precision (16-bit) floating point operations, while the DPFPU **1415A-1415N** perform double precision (64-bit) floating point operations. The ALUs **1416A-1416N** can perform variable precision integer operations at 8-bit, 16-bit, and 32-bit precision, and can be configured for mixed precision operations. The MPUs **1417A-1417N** can also be configured for mixed precision matrix operations, including half-precision floating point and 8-bit integer operations. The MPUs **1417-1417N** can perform a variety of matrix operations to accelerate machine learning application frameworks, including enabling support for accelerated general matrix to matrix multiplication (GEMM). The AFUs **1412A-1412N** can perform additional logic operations not supported by the floating-point or integer units, including trigonometric operations (e.g., Sine, Cosine, etc.).

[0152] As shown in FIG. **14B**, a general-purpose processing unit (GPGPU **1430**) can be configured to enable highly parallel compute operations to be performed by an array of graphics processing units. Additionally, the GPGPU **1430** can be linked directly to other instances of the GPGPU to create a multi-GPU cluster to improve training speed for particularly deep neural networks. The GPGPU **1430** includes a host interface **1432** to enable a connection with a host processor. In one embodiment the host interface **1432** is a PCI Express interface. However, the host interface can also be a vendor specific communications interface or communications fabric. The GPGPU **1430** receives commands from the host processor and uses a global scheduler **1434** to distribute execution threads associated with those commands to a set of compute clusters **1436A-1436H**. The compute clusters **1436A-1436H** share a cache memory **1438**. The cache memory **1438** can serve as a higher-level cache for cache memories within the compute clusters **1436A-1436H**.

[0153] The GPGPU **1430** includes memory **1434A-1434B** coupled with the compute clusters **1436A-1436H** via a set of memory controllers **1442A-1442B**. In various embodiments, the memory **1434A-1434B** can include various types of memory devices including dynamic random access memory (DRAM) or graphics random access memory, such as synchronous graphics random access memory (SGRAM), including graphics double data rate (GDDR) memory.

[0154] In one embodiment the compute clusters **1436A-1436H** each include a set of graphics cores, such as the graphics core **1400** of FIG. **14A**, which can include multiple types of integer and floating point logic units that can perform computational operations at a range of precisions including suited for machine learning computations. For example, in one embodiment at least a subset of the floating point units in each of the compute clusters **1436A-1436H** can be configured to perform 16-bit or 32-bit floating point operations, while a different subset of the floating point units can be configured to perform 64-bit floating point operations.

[0155] Multiple instances of the GPGPU **1430** can be configured to operate as a compute cluster. The communication mechanism used by the compute cluster for synchronization and data exchange varies across embodiments. In one embodiment the multiple instances of the GPGPU **1430** communicate over the host interface **1432**. In one embodiment the GPGPU **1430** includes an I/O hub **1439** that couples the GPGPU **1430** with a GPU link **1440** that enables a direct connection to other instances of the GPGPU. In one embodiment the GPU link **1440** is coupled to a dedicated GPU-to-GPU bridge that enables communication and synchronization between multiple instances of the GPGPU **1430**. In one embodiment the GPU link **1440** couples with a high speed interconnect to transmit and receive data to other GPGPUs or parallel processors. In one embodiment the multiple instances of the GPGPU **1430** are located in separate data processing systems and communicate via a network device that is accessible via the host interface **1432**. In one embodiment the GPU link **1440** can be configured to enable a connection to a host processor in addition to or as an alternative to the host interface **1432**.

[0156] While the illustrated configuration of the GPGPU **1430** can be configured to train neural networks, one embodiment provides alternate configuration of the GPGPU **1430** that can be configured for deployment within a high performance or low power inferencing platform. In an inferencing configuration the GPGPU **1430** includes fewer of the compute clusters **1436A-1436H** relative to the training configuration. Additionally, the memory technology associated with the memory **1434A-1434B** may differ between inferencing and training configurations, with higher bandwidth memory technologies devoted to training configurations. In one embodiment the inferencing configuration of the GPGPU **1430** can support inferencing specific instructions. For example, an inferencing configuration can provide support for one or more 8-bit integer dot product instructions, which are commonly used during inferencing operations for deployed neural networks.

Machine Learning Overview

[0157] A machine learning algorithm is an algorithm that can learn based on a set of data. Embodiments of machine learning algorithms can be designed to model high-level abstractions within a data set. For example, image recognition algorithms can be used to determine which of several categories to which a given input belong; regression algorithms can output a numerical value given an input; and pattern recognition algorithms can be used to generate translated text or perform text to speech and/or speech recognition.

[0158] An exemplary type of machine learning algorithm is a neural network. There are many types of neural networks; a simple type of neural network is a feedforward network. A feedforward network may be implemented as an acyclic graph in which the nodes are arranged in layers. Typically, a feedforward network topology includes an input layer and an output layer that are separated by at least one hidden layer. The hidden layer transforms input received by the input layer into a representation that is useful for generating output in the output layer. The network nodes are fully connected via edges to the nodes in adjacent layers, but there are no edges between nodes within each layer. Data received at the nodes of an input layer of a feedforward network are propagated (i.e., “fed forward”) to the nodes of the output layer via an activation function that calculates the states of the nodes of each successive layer in the network based on coefficients (“weights”) respectively associated with each of the edges connecting the layers. Depending on the specific model being represented by the algorithm being executed, the output from the neural network algorithm can take various forms.

[0159] Before a machine learning algorithm can be used to model a particular problem, the algorithm is trained using a training data set. Training a neural network involves selecting a network topology, using a set of training data representing a problem being modeled by the network, and adjusting the weights until the network model performs with a minimal error for all instances of the training data set. For example, during a supervised learning training process for a neural network, the output produced by the network in response to the input representing an instance in a training data set is compared to the “correct” labeled output for that instance, an error signal representing the difference between the output and the labeled output is calculated, and the weights associated with the connections are adjusted to minimize that error as the error signal is backward propagated through the layers of the network. The network is considered “trained” when the errors for each of the outputs generated from the instances of the training data set are minimized.

[0160] The accuracy of a machine learning algorithm can be affected significantly by the quality of the data set used to train the algorithm. The training process can be computationally intensive and may require a significant amount of time on a conventional general-purpose processor. Accordingly, parallel processing hardware is used to train many types of machine learning algorithms. This is particularly useful for optimizing the training of neural networks, as the computations performed in adjusting the coefficients in neural networks lend themselves naturally to parallel implementations. Specifically, many machine learning algorithms and software

applications have been adapted to make use of the parallel processing hardware within general-purpose graphics processing devices.

[0161] FIG. 15 is a generalized diagram of a machine learning software stack 1500. A machine learning application 1502 can be configured to train a neural network using a training dataset or to use a trained deep neural network to implement machine intelligence. The machine learning application 1502 can include training and inference functionality for a neural network and/or specialized software that can be used to train a neural network before deployment. The machine learning application 1502 can implement any type of machine intelligence including but not limited to image recognition, mapping and localization, autonomous navigation, speech synthesis, medical imaging, or language translation.

[0162] Hardware acceleration for the machine learning application 1502 can be enabled via a machine learning framework 1504. The machine learning framework 1504 can provide a library of machine learning primitives. Machine learning primitives are basic operations that are commonly performed by machine learning algorithms. Without the machine learning framework 1504, developers of machine learning algorithms would be required to create and optimize the main computational logic associated with the machine learning algorithm, then re-optimize the computational logic as new parallel processors are developed. Instead, the machine learning application can be configured to perform the necessary computations using the primitives provided by the machine learning framework 1504. Exemplary primitives include tensor convolutions, activation functions, and pooling, which are computational operations that are performed while training a convolutional neural network (CNN). The machine learning framework 1504 can also provide primitives to implement basic linear algebra subprograms performed by many machine-learning algorithms, such as matrix and vector operations.

[0163] The machine learning framework 1504 can process input data received from the machine learning application 1502 and generate the appropriate input to a compute framework 1506. The compute framework 1506 can abstract the underlying instructions provided to the GPGPU driver 1508 to enable the machine learning framework 1504 to take advantage of hardware acceleration via the GPGPU hardware 1510 without requiring the machine learning framework 1504 to have intimate knowledge of the architecture of the GPGPU hardware 1510. Additionally, the compute framework 1506 can enable hardware acceleration for the machine learning framework 1504 across a variety of types and generations of the GPGPU hardware 1510.

Machine Learning Neural Network Implementations

[0164] The computing architecture provided by embodiments described herein can be configured to perform the types of parallel processing that is particularly suited for training and deploying neural networks for machine learning. A neural network can be generalized as a network of functions having a graph relationship. As is known in the art, there are a variety of types of neural network implementations used in machine learning. One exemplary type of neural network is the feedforward network, as previously described.

[0165] A second exemplary type of neural network is the Convolutional Neural Network (CNN). A CNN is a specialized feedforward neural network for processing data having a known, grid-like topology, such as image data. Accordingly, CNNs are commonly used for compute vision and image recognition applications, but they also may be used for other types of pattern recognition such as speech and language processing. The nodes in the CNN input layer are organized into a set of “filters” (feature detectors inspired by the receptive fields found in the retina), and the output of each set of filters is propagated to nodes in successive layers of the network. The computations for a CNN include applying the convolution mathematical operation to each filter to produce the output of that filter. Convolution is a specialized kind of mathematical operation performed by two functions to produce a third function that is a modified version of one of the two original functions. In convolutional network terminology, the first function to the convolution can be referred to as the input, while the second function can be referred to as the convolution kernel. The output may be

referred to as the feature map. For example, the input to a convolution layer can be a multidimensional array of data that defines the various color components of an input image. The convolution kernel can be a multidimensional array of parameters, where the parameters are adapted by the training process for the neural network.

[0166] Recurrent neural networks (RNNs) are a family of feedforward neural networks that include feedback connections between layers. RNNs enable modeling of sequential data by sharing parameter data across different parts of the neural network. The architecture for an RNN includes cycles. The cycles represent the influence of a present value of a variable on its own value at a future time, as at least a portion of the output data from the RNN is used as feedback for processing subsequent input in a sequence. This feature makes RNNs particularly useful for language processing due to the variable nature in which language data can be composed.

[0167] The figures described below present exemplary feedforward, CNN, and RNN networks, as well as describe a general process for respectively training and deploying each of those types of networks. It will be understood that these descriptions are exemplary and non-limiting as to any specific embodiment described herein and the concepts illustrated can be applied generally to deep neural networks and machine learning techniques in general.

[0168] The exemplary neural networks described above can be used to perform deep learning. Deep learning is machine learning using deep neural networks. The deep neural networks used in deep learning are artificial neural networks composed of multiple hidden layers, as opposed to shallow neural networks that include only a single hidden layer. Deeper neural networks are generally more computationally intensive to train. However, the additional hidden layers of the network enable multistep pattern recognition that results in reduced output error relative to shallow machine learning techniques.

[0169] Deep neural networks used in deep learning typically include a front-end network to perform feature recognition coupled to a back-end network which represents a mathematical model that can perform operations (e.g., object classification, speech recognition, etc.) based on the feature representation provided to the model. Deep learning enables machine learning to be performed without requiring hand crafted feature engineering to be performed for the model. Instead, deep neural networks can learn features based on statistical structure or correlation within the input data. The learned features can be provided to a mathematical model that can map detected features to an output. The mathematical model used by the network is generally specialized for the specific task to be performed, and different models will be used to perform different task.

[0170] Once the neural network is structured, a learning model can be applied to the network to train the network to perform specific tasks. The learning model describes how to adjust the weights within the model to reduce the output error of the network. Backpropagation of errors is a common method used to train neural networks. An input vector is presented to the network for processing. The output of the network is compared to the desired output using a loss function and an error value is calculated for each of the neurons in the output layer. The error values are then propagated backwards until each neuron has an associated error value which roughly represents its contribution to the original output. The network can then learn from those errors using an algorithm, such as the stochastic gradient descent algorithm, to update the weights of the of the neural network.

[0171] FIG. 16A-16B illustrate an exemplary convolutional neural network. FIG. 16A illustrates various layers within a CNN. As shown in FIG. 16A, an exemplary CNN used to model image processing can receive input **1602** describing the red, green, and blue (RGB) components of an input image. The input **1602** can be processed by multiple convolutional layers (e.g., first convolutional layer **1604**, second convolutional layer **1606**). The output from the multiple convolutional layers may optionally be processed by a set of fully connected layers **1608**. Neurons in a fully connected layer have full connections to all activations in the previous layer, as previously described for a feedforward network. The output from the fully connected layers **1608** can be used to generate an output result from the network. The activations within the fully

connected layers **1608** can be computed using matrix multiplication instead of convolution. Not all CNN implementations are configured to make use of fully connected layers **1608**. For example, in some implementations the second convolutional layer **1606** can generate output for the CNN. [0172] The convolutional layers are sparsely connected, which differs from traditional neural network configuration found in the fully connected layers **1608**. Traditional neural network layers are fully connected, such that every output unit interacts with every input unit. However, the convolutional layers are sparsely connected because the output of the convolution of a field is input (instead of the respective state value of each of the nodes in the field) to the nodes of the subsequent layer, as illustrated. The kernels associated with the convolutional layers perform convolution operations, the output of which is sent to the next layer. The dimensionality reduction performed within the convolutional layers is one aspect that enables the CNN to scale to process large images.

[0173] FIG. **16B** illustrates exemplary computation stages within a convolutional layer of a CNN. Input to a convolutional layer **1612** of a CNN can be processed in three stages of a convolutional layer **1614**. The three stages can include a convolution stage **1616**, a detector stage **1618**, and a pooling stage **1620**. The convolutional layer **1614** can then output data to a successive convolutional layer. The final convolutional layer of the network can generate output feature map data or provide input to a fully connected layer, for example, to generate a classification value for the input to the CNN.

[0174] In the convolution stage **1616** performs several convolutions in parallel to produce a set of linear activations. The convolution stage **1616** can include an affine transformation, which is any transformation that can be specified as a linear transformation plus a translation. Affine transformations include rotations, translations, scaling, and combinations of these transformations. The convolution stage computes the output of functions (e.g., neurons) that are connected to specific regions in the input, which can be determined as the local region associated with the neuron. The neurons compute a dot product between the weights of the neurons and the region in the local input to which the neurons are connected. The output from the convolution stage **1616** defines a set of linear activations that are processed by successive stages of the convolutional layer **1614**.

[0175] The linear activations can be processed by a detector stage **1618**. In the detector stage **1618**, each linear activation is processed by a non-linear activation function. The non-linear activation function increases the nonlinear properties of the overall network without affecting the receptive fields of the convolution layer. Several types of non-linear activation functions may be used. One particular type is the rectified linear unit (ReLU), which uses an activation function defined as $f(x)=\max(0, x)$, such that the activation is thresholded at zero.

[0176] The pooling stage **1620** uses a pooling function that replaces the output of the second convolutional layer **1606** with a summary statistic of the nearby outputs. The pooling function can be used to introduce translation invariance into the neural network, such that small translations to the input do not change the pooled outputs. Invariance to local translation can be useful in scenarios where the presence of a feature in the input data is more important than the precise location of the feature. Various types of pooling functions can be used during the pooling stage **1620**, including max pooling, average pooling, and l2-norm pooling. Additionally, some CNN implementations do not include a pooling stage. Instead, such implementations substitute an additional convolution stage having an increased stride relative to previous convolution stages.

[0177] The output from the convolutional layer **1614** can then be processed by the next layer **1622**. The next layer **1622** can be an additional convolutional layer or one of the fully connected layers **1608**. For example, the first convolutional layer **1604** of FIG. **16A** can output to the second convolutional layer **1606**, while the second convolutional layer can output to a first layer of the fully connected layers **1608**.

[0178] FIG. **17** illustrates an exemplary recurrent neural network. In a recurrent neural network

(RNN), the previous state of the network influences the output of the current state of the network. RNNs can be built in a variety of ways using a variety of functions. The use of RNNs generally revolves around using mathematical models to predict the future based on a prior sequence of inputs. For example, an RNN may be used to perform statistical language modeling to predict an upcoming word given a previous sequence of words. The illustrated RNN **1700** can be described as having an input layer **1702** that receives an input vector, hidden layers **1704** to implement a recurrent function, a feedback mechanism **1705** to enable a ‘memory’ of previous states, and an output layer **1706** to output a result. The RNN **1700** operates based on time-steps. The state of the RNN at a given time step is influenced based on the previous time step via the feedback mechanism **1705**. For a given time step, the state of the hidden layers **1704** is defined by the previous state and the input at the current time step. An initial input (x.sub.1) at a first time step can be processed by the hidden layer **1704**. A second input (x.sub.2) can be processed by the hidden layer **1704** using state information that is determined during the processing of the initial input (x.sub.1). A given state can be computed as $s_{sub.t} = f(Ux_{sub.t} + Ws_{sub.t-1})$, where U and W are parameter matrices. The function f is generally a nonlinearity, such as the hyperbolic tangent function (Tanh) or a variant of the rectifier function $f(x) = \max(0, x)$. However, the specific mathematical function used in the hidden layers **1704** can vary depending on the specific implementation details of the RNN **1700**.

[0179] In addition to the basic CNN and RNN networks described, variations on those networks may be enabled. One example RNN variant is the long short-term memory (LSTM) RNN. LSTM RNNs are capable of learning long-term dependencies that may be necessary for processing longer sequences of language. A variant on the CNN is a convolutional deep belief network, which has a structure similar to a CNN and is trained in a manner similar to a deep belief network. A deep belief network (DBN) is a generative neural network that is composed of multiple layers of stochastic (random) variables. DBNs can be trained layer-by-layer using greedy unsupervised learning. The learned weights of the DBN can then be used to provide pre-train neural networks by determining an optimal initial set of weights for the neural network.

[0180] FIG. **18** illustrates training and deployment of a deep neural network. Once a given network has been structured for a task the neural network is trained using a training dataset **1802**. Various training frameworks have been developed to enable hardware acceleration of the training process. For example, the machine learning framework **1504** of FIG. **15** may be configured as a training framework **1804**. The training framework **1804** can hook into an untrained neural network **1806** and enable the untrained neural net to be trained using the parallel processing resources described herein to generate a trained neural network **1808**. To start the training process the initial weights may be chosen randomly or by pre-training using a deep belief network. The training cycle then be performed in either a supervised or unsupervised manner.

[0181] Supervised learning is a learning method in which training is performed as a mediated operation, such as when the training dataset **1802** includes input paired with the desired output for the input, or where the training dataset includes input having known output and the output of the neural network is manually graded. The network processes the inputs and compares the resulting outputs against a set of expected or desired outputs. Errors are then propagated back through the system. The training framework **1804** can adjust to adjust the weights that control the untrained neural network **1806**. The training framework **1804** can provide tools to monitor how well the untrained neural network **1806** is converging towards a model suitable to generating correct answers based on known input data. The training process occurs repeatedly as the weights of the network are adjusted to refine the output generated by the neural network. The training process can continue until the neural network reaches a statistically desired accuracy associated with a trained neural network **1808**. The trained neural network **1808** can then be deployed to implement any number of machine learning operations to generate an inference result **1814** based on input of new data **1812**.

[0182] Unsupervised learning is a learning method in which the network attempts to train itself using unlabeled data. Thus, for unsupervised learning the training dataset **1802** will include input data without any associated output data. The untrained neural network **1806** can learn groupings within the unlabeled input and can determine how individual inputs are related to the overall dataset. Unsupervised training can be used to generate a self-organizing map, which is a type of trained neural network **1808** capable of performing operations useful in reducing the dimensionality of data. Unsupervised training can also be used to perform anomaly detection, which allows the identification of data points in an input dataset that deviate from the normal patterns of the data.

[0183] Variations on supervised and unsupervised training may also be employed. Semi-supervised learning is a technique in which in the training dataset **1802** includes a mix of labeled and unlabeled data of the same distribution. Incremental learning is a variant of supervised learning in which input data is continuously used to further train the model. Incremental learning enables the trained neural network **1808** to adapt to the new data **1812** without forgetting the knowledge instilled within the network during initial training.

[0184] Whether supervised or unsupervised, the training process for particularly deep neural networks may be too computationally intensive for a single compute node. Instead of using a single compute node, a distributed network of computational nodes can be used to accelerate the training process.

[0185] FIG. **19** is a block diagram illustrating distributed learning. Distributed learning is a training model that uses multiple distributed computing nodes to perform supervised or unsupervised training of a neural network. The distributed computational nodes can each include one or more host processors and one or more of the general-purpose processing nodes. As illustrated, distributed learning can be performed model parallelism **1902**, data parallelism **1904**, or a combination of model and data parallelism **1904**.

[0186] In model parallelism **1902**, different computational nodes in a distributed system can perform training computations for different parts of a single network. For example, each layer of a neural network can be trained by a different processing node of the distributed system. The benefits of model parallelism include the ability to scale to particularly large models. Splitting the computations associated with different layers of the neural network enables the training of very large neural networks in which the weights of all layers would not fit into the memory of a single computational node. In some instances, model parallelism can be particularly useful in performing unsupervised training of large neural networks.

[0187] In data parallelism **1904**, the different nodes of the distributed network have a complete instance of the model and each node receives a different portion of the data. The results from the different nodes are then combined. While different approaches to data parallelism are possible, data parallel training approaches all require a technique of combining results and synchronizing the model parameters between each node. Exemplary approaches to combining data include parameter averaging and update based data parallelism. Parameter averaging trains each node on a subset of the training data and sets the global parameters (e.g., weights, biases) to the average of the parameters from each node. Parameter averaging uses a central parameter server that maintains the parameter data. Update based data parallelism is similar to parameter averaging except that instead of transferring parameters from the nodes to the parameter server, the updates to the model are transferred. Additionally, update based data parallelism can be performed in a decentralized manner, where the updates are compressed and transferred between nodes.

[0188] Combined model and data parallelism **1906** can be implemented, for example, in a distributed system in which each computational node includes multiple GPUs. Each node can have a complete instance of the model with separate GPUs within each node are used to train different portions of the model.

[0189] Distributed training has increased overhead relative to training on a single machine.

However, the parallel processors and GPGPUs described herein can each implement various techniques to reduce the overhead of distributed training, including techniques to enable high bandwidth GPU-to-GPU data transfer and accelerated remote data synchronization.

Exemplary Machine Learning Applications

[0190] Machine learning can be applied to solve a variety of technological problems, including but not limited to computer vision, autonomous driving and navigation, speech recognition, and language processing. Computer vision has traditionally been one of the most active research areas for machine learning applications. Applications of computer vision range from reproducing human visual abilities, such as recognizing faces, to creating new categories of visual abilities. For example, computer vision applications can be configured to recognize sound waves from the vibrations induced in objects visible in a video. Parallel processor accelerated machine learning enables computer vision applications to be trained using significantly larger training dataset than previously feasible and enables inferencing systems to be deployed using low power parallel processors.

[0191] Parallel processor accelerated machine learning has autonomous driving applications including lane and road sign recognition, obstacle avoidance, navigation, and driving control. Accelerated machine learning techniques can be used to train driving models based on datasets that define the appropriate responses to specific training input. The parallel processors described herein can enable rapid training of the increasingly complex neural networks used for autonomous driving solutions and enables the deployment of low power inferencing processors in a mobile platform suitable for integration into autonomous vehicles.

[0192] Parallel processor accelerated deep neural networks have enabled machine learning approaches to automatic speech recognition (ASR). ASR includes the creation of a function that computes the most probable linguistic sequence given an input acoustic sequence. Accelerated machine learning using deep neural networks have enabled the replacement of the hidden Markov models (HMMs) and Gaussian mixture models (GMMs) previously used for ASR.

[0193] Parallel processor accelerated machine learning can also be used to accelerate natural language processing. Automatic learning procedures can make use of statistical inference algorithms to produce models that are robust to erroneous or unfamiliar input. Exemplary natural language processor applications include automatic machine translation between human languages.

[0194] The parallel processing platforms used for machine learning can be divided into training platforms and deployment platforms. Training platforms are generally highly parallel and include optimizations to accelerate multi-GPU single node training and multi-node, multi-GPU training, while deployed machine learning (e.g., inferencing) platforms generally include lower power parallel processors suitable for use in products such as cameras, autonomous robots, and autonomous vehicles.

[0195] Conventional scheduling and dispatching in graphics hardware rely on prefetching subsequent consecutive cache lines anticipating the next instruction. These techniques can be improved upon via the use of neural networks.

[0196] Embodiments described herein provide techniques to improve the efficiency of GPU deep pipelines. A first embodiment provides for an AI-based dynamic scheduling on complex GPU architecture. A second embodiment provides for intelligent memory controller scheduling to support various types of memory requests. A third embodiment provides for an implementation of a neural network for a graphics pipeline. A fourth embodiment provides for a neural network switch to determine when to switch between a GPU pipeline and a neural network pipeline. A fifth embodiment provides for AI Driven Thread Dispatch. A sixth embodiment provides for AI-driven hardware memory prefetching.

Graphics Processor Having a Neural Network Pipeline

[0197] FIG. 20 is a block diagram of a computing device 2000 including a graphics processor 2004, according to an embodiment. The computing device 2000 can be a computing device such as

the data processing system **100** as in of FIG. **1** and can include components shown in any of FIG. **2** through FIG. **10** and/or FIG. **14A-14B**. The computing device can be manufactured according to techniques illustrated in FIG. **11A-11B**. In one embodiment, computing device **2000** can also include graphics processors of FIG. **12** and/or FIG. **13A-13B**. The computing device **2000** may also be or be included within a communication device such as a set-top box (e.g., Internet-based cable television set-top boxes, etc.), global positioning system (GPS)-based devices, etc. The computing device **2000** may also be or be included within mobile computing devices such as cellular phones, smartphones, personal digital assistants (PDAs), tablet computers, laptop computers, e-readers, smart televisions, television platforms, wearable devices (e.g., glasses, watches, bracelets, smartcards, jewelry, clothing items, etc.), media players, etc. For example, in one embodiment, the computing device **2000** includes a mobile computing device employing an integrated circuit (“IC”), such as system on a chip (“SoC” or “SOC”), integrating various hardware and/or software components of computing device **2000** on a single chip. Computing device **2000** can also be included in smart wearable devices, virtual reality (VR) devices, head-mounted display (HMDs), mobile computers, Internet of Things (IoT) devices, laptop computers, desktop computers, server computers, or other types of computing devices.

[0198] The computing device **2000** includes a graphics processor **2004**. The graphics processor **2004** represents any graphics processor described herein. The graphics processor includes one or more graphics engine(s), graphics processor cores, and other graphics execution resources as described herein. Such graphics execution resources can be presented in the forms including but not limited to execution units, shader engines, fragment processors, vertex processors, streaming multiprocessors, graphics processor clusters, or any collection of computing resources suitable for the processing of graphics resources or image resources, or performing general purpose computational operations in a heterogeneous processor.

[0199] In one embodiment, the graphics processor **2004** includes a cache **2014**, which can be a single cache or divided into multiple segments of cache memory, including but not limited to any number of L1, L2, L3, or L4 caches, render caches, depth caches, sampler caches, and/or shader unit caches. In some embodiments, the graphics processor **2004** includes a GPGPU engine **2044** that includes a neural network pipeline **2042** and a compute block **2055**. In one embodiment the compute block **2055** includes an array of execution units as described herein. In other embodiments the compute block **2055** can include other designs of compute units or streaming multiprocessors that enable SIMD and/or SIMT (single instruction, multiple thread) processing of compute and/or graphics threads. The neural network pipeline **2042** includes fixed function and programmable logic elements that can accelerate operations associated with neural networks, including operations performed during training or inference. neural network pipeline **2042** can work in concert with the compute block **2055** or independently of the compute block **2055**.

[0200] In one embodiment the neural network pipeline **2042** can be used to perform operations artificial intelligence operations (AI) on behalf of AI tessellation logic **2024** and AI texture generation logic **2025**. The neural network pipeline **2042** can also perform operations for or operate in concert with an AI scheduler **2026**, an AI memory optimizer **2027**, and an AI visibility processor **2028**. The logic that is implemented by or operates in concert with the neural network pipeline **2042** can be implemented as software or firmware logic and may be partially and simultaneously hosted by multiple components of computing device **2000**, such as one or more of graphics driver logic **2022**, graphics processor **2004**, application processor **2006**, OS **2002**, and/or firmware within the application processor **2006** and/or graphics processor **2004**. Specific details of the neural network pipeline **2042** and the AI tessellation logic **2024**, AI texture generation logic **2025**, as well as the AI scheduler **2026**, AI memory optimizer **2027**, and AI visibility processor **2028** are illustrated in the figures described below.

[0201] In one embodiment, and in addition to the graphics processor **2004**, the computing device **2000** may further include any number and type of hardware components and/or software

components, including, but not limited to an application processor **2006**, memory **2008**, and input/output (I/O) sources **2010**. The application processor **2006** can interact with a hardware graphics pipeline, as illustrated with reference to FIG. 3, to share graphics pipeline functionality. Processed data is stored in a buffer in the hardware graphics pipeline and state information is stored in memory **2008**. The resulting data can be transferred to a display controller for output via a display device, such as the display device **320** of FIG. 3. The display device may be of various types, such as Cathode Ray Tube (CRT), Thin Film Transistor (TFT), Liquid Crystal Display (LCD), Organic Light Emitting Diode (OLED) array, etc., and may be configured to display information to a user via a graphical user interface.

[0202] The application processor **2006** can include one or processors, such as processor(s) **102** of FIG. 1 and may be the central processing unit (CPU) that is used at least in part to execute an operating system (OS) **2002** for the computing device **2000**. The OS **2002** can serve as an interface between hardware and/or physical resources of the computing device **2000** and one or more users. The OS **2002** can include driver logic for various hardware devices in the computing device **2000**, including graphics driver logic **2022**, such as the user mode graphics driver **1026** and/or kernel mode graphics driver **1029** of FIG. 10.

[0203] It is contemplated that in some embodiments the graphics processor **2004** may exist as part of the application processor **2006** (such as part of a physical CPU package) in which case, at least a portion of the memory **2008** may be shared by the application processor **2006** and graphics processor **2004**, although at least a portion of the memory **2008** may be exclusive to the graphics processor **2004**, or the graphics processor **2004** may have a separate store of memory. The memory **2008** may comprise a pre-allocated region of a buffer (e.g., framebuffer); however, it should be understood by one of ordinary skill in the art that the embodiments are not so limited, and that any memory accessible to the lower graphics pipeline may be used. The memory **2008** may include various forms of random access memory (RAM) (e.g., SDRAM, SRAM, etc.) comprising an application that makes use of the graphics processor **2004** to render a desktop or 3D graphics scene. A memory controller hub, such as memory controller **116** of FIG. 1, may access data in the memory **2008** and forward it to graphics processor **2004** for graphics pipeline processing. The memory **2008** may be made available to other components within the computing device **2000**. For example, any data (e.g., input graphics data) received from various I/O sources **2010** of the computing device **2000** can be temporarily queued into memory **2008** prior to their being operated upon by one or more processor(s) (e.g., application processor **2006**) in the implementation of a software program or application. Similarly, data that a software program determines should be sent from the computing device **2000** to an outside entity through one of the computing system interfaces, or stored into an internal storage element, is often temporarily queued in memory **2008** prior to its being transmitted or stored.

[0204] The I/O sources can include devices such as touchscreens, touch panels, touch pads, virtual or regular keyboards, virtual or regular mice, ports, connectors, network devices, or the like, and can attach via a platform controller hub **130** as referenced in FIG. 1. Additionally, the I/O sources **2010** may include one or more I/O devices that are implemented for transferring data to and/or from the computing device **2000** (e.g., a networking adapter); or, for a large-scale non-volatile storage within the computing device **2000** (e.g., hard disk drive). User input devices, including alphanumeric and other keys, may be used to communicate information and command selections to graphics processor **2004**. Another type of user input device is cursor control, such as a mouse, a trackball, a touchscreen, a touchpad, or cursor direction keys to communicate direction information and command selections to GPU and to control cursor movement on the display device. Camera and microphone arrays of the computing device **2000** may be employed to observe gestures, record audio and video and to receive and transmit visual and audio commands.

[0205] I/O sources **2010** configured as network interfaces can provide access to a network, such as a LAN, a wide area network (WAN), a metropolitan area network (MAN), a personal area network

(PAN), Bluetooth, a cloud network, a cellular or mobile network (e.g., 3.sup.rd Generation (3G), 4.sup.th Generation (4G), 5.sup.th Generation (5G) etc.), an intranet, the Internet, etc. Network interface(s) may include, for example, a wireless network interface having one or more antenna (c). Network interface(s) may also include, for example, a wired network interface to communicate with remote devices via network cable, which may be, for example, an Ethernet cable, a coaxial cable, a fiber optic cable, a serial cable, or a parallel cable.

[0206] Network interface(s) may provide access to a LAN, for example, by conforming to IEEE 802.11 standards, and/or the wireless network interface may provide access to a personal area network, for example, by conforming to Bluetooth standards. Other wireless network interfaces and/or protocols, including previous and subsequent versions of the standards, may also be supported. In addition to, or instead of, communication via the wireless LAN standards, network interface(s) may provide wireless communication using, for example, Time Division, Multiple Access (TDMA) protocols, Global Systems for Mobile Communications (GSM) protocols, Code Division, Multiple Access (CDMA) protocols, and/or any other type of wireless communications protocols.

[0207] It is to be appreciated that a lesser or more equipped system than the example described above may be preferred for certain implementations. Therefore, the configuration of the computing device **2000** may vary from implementation to implementation depending upon numerous factors, such as price constraints, performance requirements, technological improvements, or other circumstances. Examples include (without limitation) a mobile device, a personal digital assistant, a mobile computing device, a smartphone, a cellular telephone, a handset, a one-way pager, a two-way pager, a messaging device, a computer, a personal computer (PC), a desktop computer, a laptop computer, a notebook computer, a handheld computer, a tablet computer, a server, a server array or server farm, a web server, a network server, an Internet server, a work station, a mini-computer, a main frame computer, a supercomputer, a network appliance, a web appliance, a distributed computing system, multiprocessor systems, processor-based systems, consumer electronics, programmable consumer electronics, television, digital television, set top box, wireless access point, base station, subscriber station, mobile subscriber center, radio network controller, router, hub, gateway, bridge, switch, machine, or combinations thereof.

[0208] Embodiments may be implemented as any one, or a combination of one or more microchips or integrated circuits interconnected using a parent-board, hardwired logic, software stored by a memory device and executed by a microprocessor, firmware, an application specific integrated circuit (ASIC), and/or a field programmable gate array (FPGA). The term “logic” may include, by way of example, software or hardware and/or combinations of software and hardware.

[0209] Embodiments may be provided, for example, as a computer program product which may include one or more machine-readable media having stored thereon machine-executable instructions that, when executed by one or more machines such as a computer, network of computers, or other electronic devices, may result in the one or more machines carrying out operations in accordance with embodiments described herein. A machine-readable medium may include, but is not limited to, floppy diskettes, optical disks, CD-ROMs (Compact Disc-Read Only Memories), and magneto-optical disks, ROMs, RAMs, EPROMs (Erasable Programmable Read Only Memories), EEPROMs (Electrically Erasable Programmable Read Only Memories), magnetic or optical cards, flash memory, or other type of non-transitory machine-readable media suitable for storing machine-executable instructions.

[0210] Moreover, embodiments may be downloaded as a computer program product, wherein the program may be transferred from a remote computer (e.g., a server) to a requesting computer (e.g., a client) by way of one or more data signals embodied in and/or modulated by a carrier wave or other propagation medium via a communication link (e.g., a modem and/or network connection).

AI-Based Tessellation

[0211] According to one embodiment, artificial intelligence (AI) techniques are applied to a

graphics processing unit to enable an AI-based tessellation mechanism. AI-based tessellation mechanism can be performed using AI tessellation logic **2024** of the GPGPU engine **2044** of FIG. **20**. In one embodiment, the AI tessellation logic **2024** operates in concert with the neural network pipeline **2042** and compute block **2055** of FIG. **20**. The AI-based tessellation mechanism performs AI driven higher-order geometry and tessellation inference using a neural network that is trained based using a dataset that includes pre and post tessellated vertices. The tessellation module **2100** can train a network to recognize patterns from coarse geometric representations and infer higher-order and/or higher tessellation for said geometric shapes. In one embodiment, a higher-order geometric representations can be output in the form of non-uniform rational B-spline (NURB) curves, which are defined by an order, a set of weighted control points, and a knot vector, where the order of the curve defines the number of nearby control points that influence any given point on the curve.

[0212] In one embodiment, the AI tessellation logic **2024** may be used to replace tessellation logic within a graphics pipeline, with the output of the AI tessellation logic **2024** and/or neural network pipeline **2042** flowing through the remaining programmable & fixed-function portion of the graphics pipeline, (e.g., geometry shader, rasterization, output-merger, etc.) to enable the application of custom special effects to the tessellated output. In some implementations, the AI tessellation logic **2024** may eliminate the need for a fixed function tessellator. In other implementations, the AI tessellation logic **2024** can be enabled automatically when tessellation is not supported by an application.

[0213] Tessellation operates on a group of vertices known as patches, which provide a framework for use in interpolating vertices during tessellation. A shader program then transforms the generated vertices into a polygonal form. A vertex shader in the geometry pipeline provides an array of vertices to a tessellation module, along with attributes corresponding to various output variables. A tessellation control shader will then execute for each vertex and generate data sets that include control variables that the tessellation evaluator will use to interpret the additional vertices generated by the tessellation generator and tessellation values that the tessellation generator will use to generate new vertices.

[0214] FIG. **21** illustrates a tessellation module **2100**, according to an embodiment. The tessellation module **2100** includes a tessellation controller **2108**, a tessellation generator **2109** and a tessellation evaluator **2110**. In one embodiment, the tessellation controller **2108** and tessellation evaluator **2110** include programmable logic that can be programmed via shader logic. For example, the tessellation controller **2108** can be controlled via a hull shader program **2118**, while the tessellation evaluator **2110** can be programmed via a domain shader program. In such embodiment, the tessellation controller **2108** can function as the hull shader **811** of the geometry pipeline **820** of FIG. **8**. The tessellation evaluator **2110** can function as the domain shader **817** of the geometry pipeline **820** of FIG. **8**. The tessellation generator **2109** can be fixed-function logic that is initialized by binding a hull shader to the graphics pipeline. The tessellation generator **2109** can be configured to subdivide a domain (e.g., quad, triangle, line) into many smaller objects (triangles, points or lines).

[0215] During operation, a vertex shader **2107** can perform transform and lighting operations on vertices at the direction of a vertex program and then output patches of vertices to the tessellation controller **2108**, which uses the patches as input control points **2102**. The vertices are processed at the tessellation controller **2108** using a hull shader program **2118**. The hull shader program **2118** can generate two phases of threads consisting of control point threads **2119** and patch constant threads **2120**. The threads of each phase are generally executed in parallel. The control point threads **2119** read the input control points **2102** to generate output control points **2106**. The patch-constant phase operates once per patch and generates the patch constant data **2104** (e.g. outer and inner tessellation factors). The patch constant data **2104** is used as input to the tessellation generator **2109** and tessellation evaluator **2110**. The tessellation generator **2109** generates new vertices by creating new primitives inside of the patch of vertices. In one embodiment, triangles,

quads (e.g. rectangles and squares), or lines can be drawn within the vertex patch and then new vertices are generated by subdividing the polygons to make new, smaller polygons. New vertices are then interpolated based on the smaller polygons. The control point data can then be used by a tessellation evaluator **2110** to place the new vertex data into the 3D scene. The output control points **2106** and patch constant data **2104** can be used to define how the tessellation evaluator **2110** transforms the newly generated vertices **2113** for use within the 3D scene.

[0216] FIG. **22A** illustrates use of the tessellation module **2100** to train a neural network used by AI tessellation logic **2024** to perform AI-based tessellation. Using a training and deployment system similar to that illustrated in FIG. **18**, the tessellation module **2100** can be used to generate a training dataset **2202** that can be used to train an untrained neural network **2206** into a trained neural network **2208**. The training dataset **2202** can include sets of coarse vertex data that was input to the tessellation module **2100** and the resulting fine vertex data that is output by the tessellation module **2100**. After sufficient training, the trained neural network **2208** can be configured to receive coarse vertex data **2212** and generate appropriate fine vertex data **2214**.

[0217] FIG. **22B** illustrates inference using AI tessellation logic **2024** that is configured to use a trained neural network to generate fine vertex data **2214** upon receipt of coarse vertex data **2212**. The AI tessellation logic **2024** can initially operate on specific patterns of coarse vertex data **2212**. Over time, more complex models or a greater number of models can be used to widen the types of vertex patterns that can be tessellated via the AI tessellation logic **2024**.

[0218] In a further embodiment, a tessellation mechanism described herein can make use of a neural network to simulate tessellation of a 3D scene at the pixel level to increase the quality of rendered output. The tessellation mechanism can use a pre-trained neural network to simulate the look of a tessellated image without performing tessellation at the vertex or geometry level.

[0219] FIG. **23A** illustrates use of post-shader pixel data for tessellated geometry to train a neural network that may be used to simulate tessellation at the vertex level. Using a training and deployment system similar to that illustrated in FIG. **18**, post-shader pixel data for tessellated geometry **2301** can be used to generate a training dataset **2302** that can be used to train an untrained neural network **2306** into a trained neural network **2308**. The training dataset **2302** can include sets of pixel shaded coarse vertex data and pixel shaded tessellated vertex data. The shaded coarse vertex data for the training dataset **2302** can be generated by applying a pixel shader to coarse vertex data. The shaded fine vertex data for the training dataset **2302** can be generated by applying the pixel shader to tessellated geometry that is generated based on the coarse vertex data. After sufficient training, the trained neural network **2308** can be configured to receive pixel shaded coarse vertex data **2312** and output the appropriate shaded fine vertex data **2314**.

[0220] FIG. **23B** illustrates inference using the AI tessellation logic **2024** when the AI tessellation module is configured to simulate tessellation at the pixel level. In one embodiment the AI tessellation logic **2024** can be configured to output or generate shaded fine vertex data **2314** based on shaded coarse vertex data **2312** that is input to the AI tessellation logic **2024**. The AI tessellation logic **2024** can initially operate on specific patterns of shaded coarse vertex data **2312**. Over time, more complex models or a greater number of models can be used to widen the types of shaded vertex patterns that can be tessellated via the AI tessellation logic **2024**. In one embodiment, the AI tessellation logic **2024** can optionally operate in concert with the AI texture generation logic **2025** to assist in the pixel generation process for the shaded fine vertex data.

EU Addressable NN Block

[0221] Some embodiments described herein provide for a neural network block that is addressable from a shader EU in the same way a sampler is addressable. In one embodiment, a block of execution units can calculate particle movement using artificial intelligence logic implemented via the neural network block. In one embodiment, GPU-based actor AI decision making is enabled. In one embodiment, general-purpose neural network inferencing can be applied on a per-EU thread/lane basis.

[0222] FIG. 24 illustrates a set of hardware blocks **2400** associated with a graphics processor described herein. The hardware blocks **2400** include an EU block **2402**, a sampler **2404**, a neural network block (NN block **2405**), and a memory fetch unit. The EU block **2402** can be a variant of the compute block **2055** of FIG. 20. The neural network block **2405** can be a part of the neural network pipeline **2042** of FIG. 20. The EU block **2402** includes one or more graphics execution units as described herein, each execution unit including general-purpose logic that is programmable to perform parallel general-purpose computational operations in addition to graphics processing operations. The sampler **2404** can read texture or other 3D graphics related data into memory. The 3D sampler can read texture data differently based on a configured sample state and the texture format associated with a given texture. The NN block **2405** enables execution of a small neural network with a fixed maximum layer size and/or number of layers. The memory fetch unit **2406** enables the EU block **2402**, sampler, **2404**, and NN block **2405** to fetch portions of memory from a cache hierarchy, local graphics memory, or system memory.

[0223] The sampler **2404** takes input in the form of a texture resource and a sampler state. The texture resource is a pointer to texture data in memory. The sampler state includes details such as the sampler filter mode, addressing mode, and max anisotropy. The sampler generates output in the form of some color data that is augmented based on the selected sampler state. A similar concept can be applied to neural networks, specifically neural network used for inference. In one embodiment a pointer to input data is provided to the NN block **2405** which, based on a state block representing the number of layers, the size of each layer, and associated activation function for each layer, can perform neural network calculations and output an inference result to the EU block **2402**.

[0224] In various embodiments, multiple designs are possible for the NN block **2405**. In one embodiment, the NN block **2405** has a single layer design. In one embodiment, the NN block **2405** has a multiple layer design. For each layer, a user can program the NN block **2405** with a state block representing the number of layers, the size of each layer, and associated activation function for each layer. The user can also pass in a uniform buffer of weights for the neural network for each layer and any other associated state needed to operate the neural network.

[0225] Mapping processing to an execution unit within the EU block **2402** can be supported by either serializing the entire thread's lanes or by adding a SIMD path. The sampler **2404** can take somewhere between 8 and 64 pixels worth of texture coordinate data, determine the minimum set of affected cache lines, and parallelize the sampling operation. The sampler operations can be performed serially, once the data is present, or in parallel. The NN block **2405** can perform in the same manner as the sampler **2404** and can potentially reduce the required work if the inputs are identical.

[0226] In one embodiment the NN block **2405** is programmatically configurable to accelerate neural network primitive operations, such as but not limited to matrix multiply and accumulate operations and/or operations to accelerate steps of a convolution operation. The NN block **2405** can also be configured to accelerate activation functions, as described below.

[0227] FIG. 25 illustrates a single layer hardware neural network block **2500**, according to an embodiment. The single layer hardware neural network block **2500** is a variant of the NN block **2405** of FIG. 24. The single layer block includes a single neural network block **2520** including a source data buffer **2502**, a neural network & activation operations unit **2504**, and an output data buffer **2506**. The single layer hardware neural network block **2500** additionally includes a block programming unit **2508** and a weights cache **2510**.

[0228] The neural network block **2520** accepts and initial set of input data **2501**, which is stored in a source data buffer **2502**. The neural network & activation operations unit **2504** includes hardware logic to perform calculations associated with neural network operations and activations operations for a layer of a neural network. The neural network & activation operations unit **2504** in various embodiments, can be a single unit or separate units. The neural network & activation operations unit **2504** can include logic to perform a variety of operations to accelerate training or inference

operations for a neural network, including matrix multiply and accumulate operations and/or operations to accelerate steps of a convolution operation. The neural network & activation operations unit **2504** can also support a variety of activation functions including, but not limited to the rectified linear unit (ReLU) function of equation (1), the sigmoid function of equation (2), or the hard-sigmoid function of equation (3).

[00001] $f(x) = \max(0, x)$ (1) $\sigma(x) = \frac{1}{(1+e^{-x})}$ (2) $\sigma(x) = \max(0, \min(1, \frac{x+1}{2}))$ (3)

[0229] Output data from the neural network & activation operations unit **2504** is stored in the output data buffer **2506**.

[0230] The neural network block **2520** is programmed by the block programming unit **2508**. Programming data **2511** is provided to the block programming unit **2508**, where the programming data **2511** can be specified by program logic that is performing operations on a graphics processor, or other processing unit that contains the single layer hardware neural network block **2500**. The block programming unit **2508**, based on the programming data **2511**, configures the number of layers and layer state information. The layer state information is state information for one or more layers of a neural network to be processed by the neural network unit. Exemplary layer state information includes the number of input and output values, the number of weights, and activation operations to use for the configured layers of the neural network. Weights for the various layers are stored in the weights cache **2510**.

[0231] In one embodiment, the single layer hardware neural network block **2500** is programmed to wrap back upon itself as the block processes each layer. Each neuron is programmed with a specific set of weights and feeds inputs either from the initial input or the prior layer. As each layer is processed, the neural network block consumes consume a portion of a weight buffer, which can be stored in the weights cache **2510**, and applies the neuron calculation and activation function via the neural network & activation operations unit **2504**. When a layer is completed, the output buffer becomes the input buffer for the next layer. Input and output can then alternate between buffers in a ping-pong manner. Upon completion of the last layer, the results are returned to the shader.

[0232] FIG. **26** illustrates a multiple layer hardware neural network block **2600**, according to an embodiment. In one embodiment, the multiple layer hardware neural network block **2600** includes multiple neural network blocks **2520** illustrated in FIG. **25**. For example, the illustrated multiple layer hardware neural network block **2600** includes three neural network blocks **2520A-2520C**, although a different number of blocks can be included.

[0233] Although the multiple layer hardware neural network block **2600** consumes a greater die area than the single layer hardware neural network block **2500**, the multiple layer block can pipeline layer processing, enabling improved performance.

[0234] In one embodiment, the multiple layer hardware neural network block **2600** supports a fixed maximum layer size, represented by the number of neural network blocks **2520** included within the hardware. In one embodiment, the multiple layer hardware neural network block **2600** functions like a multi-block single layer design, with the ability to “wrap around” back to the start allowing the multiple layer hardware neural network block **2600** to scale to any neural network depth.

Geometry Culling Visibility Using Machine Learning

[0235] One embodiment provides for a method of geometry culling visibility using machine learning. Using machine learning can avoid expensive pre-passes in fixed function logic within the graphics processor when performing geometry culling. A coarse description of world or scene, including lighting data, opacity information, etc. can be processed before the tessellation stage of the graphics pipeline. A machine learning model can be trained on data with labels of 0 (culled) and 1 (non-culled). The trained machine learning model can then generate visibility information on a per-object basis. When generating visibility information on a per-object basis, the machine learning model can determine whether an object is unobstructed, partially obstructed, or completely obstructed. A list of completely obstructed objects can be sent to the graphics pipeline for culling.

[0236] FIG. 27 illustrates a system 2700 in which geometry culling visibility is determined using machine learning, according to an embodiment. In the illustrated system 2700, a viewpoint 2702 relative to a scene 2704 that includes geometry data, along with parameters including but not limited to the lighting and opacity data of the objects described by the geometry data, is provided to a visibility processing neural network 2705. The visibility processing neural network 2705, in one embodiment, is trained using a training data set that includes portion of geometry data in which a given object is either culled or not culled. The visibility processing neural network 2705 can be included within the AI visibility processor 2028 of FIG. 20. After training, the visibility processing neural network 2705 can be used to determine, for a given object in the scene 2704, whether the object will be visible given the provided lighting and opacity data. The visibility processing neural network 2705 can be pre-trained before deployment and the weights and model associated with the neural network can be provided to a neural network block described herein (e.g., NN block 2405) for execution. The visibility processing neural network 2705 can output culled geometry 2706 that can be submitted to a rendering pipeline 2708.

Generative Texture Shader Model

[0237] Embodiment described herein provide for a generative texture shader model in which a meta-shader can generate many different types of textures. The generative texture shader, in one embodiment, is implemented via fixed function logic within a sampler unit. In one embodiment, the generative texture shader is implemented using one or more variants of an EU addressable neural network block as described herein (e.g., NN block 2405).

[0238] In one embodiment the generative texture shader is a generative adversarial network (GAN). A GAN is a machine learning network that includes a generator and a discriminator, where the generator maps a latent encoding to a data space, while the discriminator distinguishes between samples generated by the generator and real data. The generator is trained to deceive the discriminator, while the discriminator is trained to avoid being deceived by the generator. The specific implementation of the generator can vary based on the GAN implementation, and different implementations can use different numbers and combinations of neural network layers.

[0239] FIG. 28 illustrates a meta-shader system 2800, according to an embodiment. The meta-shader system 2800 can be configured to generate many different types of textures and can be used as a replacement for a procedural texture shader. A procedural texture shader is a shader program that can execute on a shader EU (e.g., within EU block 2402). The procedural texture shader uses program logic to generate texture data that can be applied to geometry generated by a graphics pipeline. The meta-shader system 2800, instead of using a shader program to generate a specific type of texture, can be configured to generate many different types of textures.

[0240] In one embodiment the meta-shader system 2800 includes a GAN 2815 including a generator network 2803 and a discriminator network 2806. The GAN 2815 can be included in the AI texture generation logic 2025 of FIG. 20. The generator network 2803 is configured to transform random input data (e.g., noise vector 2801) into generated data 2805. The discriminator network 2806 is trained to distinguish between generated data 2805 output by the generator network 2803 and actual data within a dataset. The generator network 2803 and the discriminator network 2806 are trained together via a training module 2807. During training, the generator network 2803 learns to generate higher quality generated data 2805 based on feedback from the discriminator network 2806. During training, the discriminator network 406 learns to distinguish between authentic data and generated data 2805. Based on input parameters 2802, such as terrain data, the GAN 2815 can be configured to generate a variety of different types of generated texture data 2808 (e.g., Texture 1, Texture 2, . . . , Texture N). For example, textures for terrains including deserts, high mountains, green trees, or any other type of terrain can be generated by the meta-shader, instead of having different procedural texture shaders for different types of terrain.

AI-Based Dynamic Scheduling on Complex GPU Architectures

[0241] GPU deep pipelines may be difficult to program and/or to maintain performance during

dynamic workloads such as 3D gaming. One embodiment provides for the use of a pre-trained network from a GPU ecosystem, which can be locally pre-trained or continuously trained to improve performance. For a GPU hardware block that is statically programmed for performance, it may be possible that some tradeoff has been made in the graphics driver that may not be ideal for all cases in which the GPU hardware block is programmed. The graphics driver may not always have the full picture, as understanding the full picture of the GPU workload is too time consuming to perform in real time. Accordingly, heuristics are employed as to which tradeoff to make for a given hardware block.

[0242] Embodiments described herein introduce a neural network hardware block that is responsible for programming the behavior for other hardware blocks to obtain the higher overall performance. Specific blocks and/or logic units that can benefit from dynamic adjustments include but are not limited to scheduling blocks, memory controllers, caching hierarchy eviction algorithms, shared local memory banking, internal performance optimization hardware blocks, and power profiles. For example, a hardware scheduling block at the beginning of or within the GPU pipeline can route work to various GPU blocks. This scheduling block is statically partitioned in existing implementations. Embodiments described herein enable dynamic partitioning. Additionally, the memory controller can be configured to determine how data is pre-fetched based on access patterns. Caching hierarchy can be partitioned between other hardware units. The amount of cache that is given to each hardware function (e.g., instruction, data, etc.) or how the cache is mapped to a specific API concept (e.g., shader stage) within the GPU pipeline can be dynamically adjusted. For shared local memory banking, memory access patterns may dictate different scheduling algorithms to ensure fewer banking collisions. Furthermore, internal performance optimization of hardware blocks can be performed, for example, the serialization of pixel processing after rasterization to prevent unnecessary hardware or software operations. Additionally, power profiles can be dynamically optimized over time. Depending on performance needs, the amount of power routed to the GPU can be modified based on a determination made by a neural network.

[0243] FIG. 29 illustrates a graphics processing system 2900 including AI-based dynamic scheduling, according to an embodiment. The system includes a scheduler 2902, a GPU pipeline 2904, a memory controller 2906, a cache hierarchy 2908, a power unit 2909, and a neural network block (NN block 2910). The NN block 2910, based on external input 2901 and performance statistics 2905, can be configured to execute a neural network that can dynamically manage or configure any one or more of the scheduler 2902, GPU pipeline 2904, memory controller 2906, cache hierarchy 2908, and/or power unit 2909 using techniques described herein. The NN block 2910 can provide programming information 2911 to hardware blocks within the graphics processing system 2900 to configure those blocks based on the external input 2901, performance statistics 2905, and models executed by the NN block 2910.

[0244] In various embodiments, the NN block 2910 of the graphics processing system 2900 can be configured to implement embodiments of the various concepts described below.

Intelligent Memory Controller Scheduling to Support Various Types of Memory Requests

[0245] FIG. 30 illustrates a system 3000 for intelligent memory controller scheduling, according to an embodiment. In one embodiment the NN block 2910 can enable intelligent memory controller scheduling to support various types of memory requests. The memory controller 2906 receives request 3001 to access memory for different workloads including but not limited to 3D, general purpose compute, or AI workloads. For some of those requests 3001, the memory controller 2906 may require support for requests to access array of structures (AoS 3002) or structures of arrays (SoA 3003) from a single memory block. AoS 3002 and SoA 3003 are contrasting ways to arrange a sequence of records in memory with regard to interleaving. The memory controller 2906 is configured to work with an AI unit 3010 that can execute on a neural network block, such as NN block 2910. The AI unit 3010 can support requests better utilize the cache and execution units

based on the requests **3001**, to reduce or high latency associated with those requests. The AI unit **3010** can be trained based on cache utilization, page faults, latency, power, execution unit wait time, or other factors. In one embodiment the AI unit **3010** is included within the AI memory optimizer **2027**.

Implementation of a Neural Network for a Graphics Pipeline

[0246] One embodiment provides for hardware logic to switch between a traditional GPU pipeline and neural network-based image generation and processing pipeline. The switch can be performed at run time depending on desired tradeoffs between quality and performance.

[0247] FIG. **31** illustrates a system **3100** having support for dynamic pipeline switching, according to embodiments described herein. In one embodiment the system includes a GPU pipeline **2904** as described herein and a neural network pipeline **3110**, which includes one or more hardware blocks configured to execute neural network operations. The neural network pipeline **3110** can be programmed to use pre-trained or continuously trained neural networks to perform a variety of operations including classification or image (e.g., texture) generation. The GPU pipeline **2904** and the neural network pipeline **3110** can each output to and read from a framebuffer **3104**. The neural network pipeline **3110** can also be configured to use the framebuffer **3104** as input data for neural network processing.

[0248] Depending on the workload the GPU pipeline **2904** may generate higher quality output relative to the neural network pipeline **3110**, while the neural network pipeline **3110** may operate at a higher performance or lower latency for at least a subset of operations. Additionally, the neural network pipeline **3110** can perform inference operations based on data output to the framebuffer **3104** by the graphics pipeline **2904**. In some embodiments, the GPU pipeline **2904** and the neural network pipeline **3110** may share certain execution resources within the system **3100**. Accordingly, the GPU pipeline **2904** and the neural network pipeline **3110** may not be able to operate simultaneously (e.g., within the same clock cycle). In such embodiment, a switch **3101** can be enabled that allows runtime switching between the GPU pipeline **2904** and the neural network pipeline **3110**. For example, the GPU pipeline **2904** can be configured to generate image data, which can be written to the framebuffer **3104**. The switch **3101** can be activated at runtime to enable the neural network pipeline **3110**, which can use the data within the framebuffer **3104** as input to a neural network to perform an inference operation. In one embodiment the GPU pipeline **2904** and the neural network pipeline **3110** can be switched dynamically on a per-clock basis.

Neural Network-Based Switching Between a GPU Pipeline and a Neural Network Pipeline

[0249] FIG. **32** illustrates a system **3200** having support for AI-based dynamic pipeline switching, according to embodiments described herein. In one embodiment, the system **3200** of FIG. **32** can be similar to the system **3100** of FIG. **31**, with the addition of the neural network block (NN block **2910**) at the head of the GPU pipeline **2904** and neural network pipeline **3110**. The NN block **2910** can be configured to determine a pipeline to which workloads are to be dispatched for a given cycle. The NN block **2910** can dynamically configure the switch **3101** to select between the GPU pipeline **2904** and the neural network pipeline **3110** at run time depending on the desired quality or performance of a pending workload.

AI Driven Thread Dispatch

[0250] FIG. **33** illustrates a system **3300** to enable AI driven thread dispatch, according to an embodiment. In one embodiment the system **3300** includes a processor **3302** coupled to the NN block **2910**. The processor **3302** can be a CPU or GPU and can include CPU and/or GPU cores. The processor **3302** includes a thread scheduler **3304** to schedule threads for workloads to be executed on the processor. For example, a set of kernel programs **3310** to be dispatched for execution on the processor **3302** can be provided to a thread scheduler **3304**. The NN block **2910** can be used to determine an optimal number of threads **3306** to be dispatched to the processor **3302** or to be active on the processor **3302** during a given set of execution cycles. The NN block **2910** can execute a neural network that is trained from data sets of workloads executed on the processor

3302. The neural network can be pre-trained and/or can be continuously trained to determine the optimal number of threads **3306**. In one embodiment, AI driven thread dispatch is performed by the AI scheduler **2026** of FIG. **20**.

AI-Driven Hardware Memory Prefetching

[0251] Existing HW prefetchers may not reach optimal efficiency for all workloads. Some hardware prefetchers are configured to simply prefetch subsequent consecutive cache lines under the assumption that the next instruction will read or write data in a contiguous fashion.

Embodiments described herein can be configured to train a neural network to learn memory access patterns for a variety of workloads and use the learned patterns to infer the memory access pattern for a given workload. The trained neural network can then recognize a specific workload and adjust the prefetch pattern to be specific for the recognized workload, resulting in improved prefetch efficiency for data from memory.

[0252] FIG. **34A-34B** illustrate a system **3400** to enable AI-Driven Hardware Memory Prefetching, according to embodiments described herein. In one embodiment, AI-Driven Hardware Memory Prefetching is implemented via logic within the AI memory optimizer **2027** of FIG. **20**. FIG. **34A** illustrates hardware components of the system **3400**, where the hardware components include a graphics processor **3404**, the neural network block **2910**, a cache prefetch unit **3424**, and the cache hierarchy **2908** into which data will be prefetched by the cache prefetch unit **3424**. The cache hierarchy **2908** can be a multi-level cache including level 1, level 2, level 3, and/or level 4 caches. Other cache levels can also be included. In one embodiment the graphics processor **3404** can also share an additional last level cache with a CPU. Data can be prefetched into one or more levels of the cache hierarchy **2908** from memory **3430**, which can be memory that is local the graphics processor **3404** or can be system memory that is shared with a CPU. Where the memory **3430** is system memory, the graphics processor **3404** can have a unified memory architecture with the CPU or one or more other graphics processors or accelerators.

[0253] FIG. **34B** illustrates a cache lines **3450** associated with a block of memory, according to embodiments described herein. The cache lines **3450** are within the cache hierarchy **2908** and store data that is prefetched from the memory **3430** within the system. In one embodiment, NN block **2910** can enable AI logic to control the cache prefetch unit **3424** to prefetch memory into the cache lines **3450**. The illustrated cache lines **3450** compare memory **3456** prefetched by an AI hardware prefetcher to memory **3454** prefetched by a conventional hardware prefetcher, relative to the memory **3452** read by a workload. As illustrated, the conventional hardware prefetcher assumes a different read pattern than the pattern used by the workload. However, the AI prefetcher, having been trained on the memory pattern used by the workload, can recognize the workload and apply the proper prefetch pattern.

[0254] Embodiments described herein provide techniques to improve the efficiency of GPU deep pipelines. A first embodiment provides for AI-based tessellation at the vertex and pixel level. A second embodiment provides for a processor including a neural network (NN) block that is addressable by graphics compute units (e.g., execution units, compute units, multiprocessors, etc.) within the processor. A third embodiment provides for geometry culling visibility using machine learning to avoid expensive pre-passes in fixed function hardware blocks. A fourth embodiment provides for a generative texture shader model in which a meta-shader can generate many different types of textures. A fifth embodiment provides for an AI-based dynamic scheduling on complex GPU architecture. A sixth embodiment provides for intelligent memory controller scheduling to support various types of memory requests. A seventh embodiment provides for an implementation of a neural network for a graphics pipeline. An eight embodiment provides for a neural network switch to determine when to switch between a GPU pipeline and a neural network pipeline. A ninth embodiment provides for AI Driven Thread Dispatch. A tenth embodiment provides for AI-driven hardware memory prefetching.

[0255] The following clauses and/or examples pertain to specific embodiments or examples

thereof. Specifics in the examples may be used anywhere in one or more embodiments. The various features of the different embodiments or examples may be variously combined with some features included and others excluded to suit a variety of different applications. Examples may include subject matter such as a method, means for performing acts of the method, at least one machine-readable medium including instructions that, when performed by a machine cause the machine to perform acts of the method, or of an apparatus or system according to embodiments and examples described herein. Various components can be a means for performing the operations or functions described.

[0256] One embodiment provides for a graphics processor comprising a block of graphics compute units, a graphics processor pipeline coupled to the block of graphics compute units, and a programmable neural network unit including one or more neural network hardware blocks, where the programmable neural network unit is coupled with the block of graphics compute units and the graphics processor pipeline. The one or more neural network hardware blocks include hardware to perform neural network operations and activation operations for a layer of a neural network. The programmable neural network unit can configure settings of one or more hardware blocks within the graphics processor pipeline based on a machine learning model trained to optimize performance of a set of workloads.

[0257] In one embodiment the graphics processor additionally comprises a memory controller coupled with the block of graphics compute units. The programmable neural network unit is to configure the memory controller to support one of multiple types of memory requests from a shader program executed via the graphics processor pipeline.

[0258] The graphics processor can additionally include a neural network pipeline coupled with the block of graphics compute units, where the neural network pipeline to perform an inferencing operation based on data generated by the graphics processor pipeline. Additionally, the graphics processor can be configured dynamically switch between the neural network pipeline and the graphics processor pipeline. In one embodiment, the one or more neural network hardware blocks are to execute at least one layer of a neural network to determine when to dynamically switch between the neural network pipeline and the graphics processor pipeline.

[0259] In one embodiment the graphics processor additionally includes a thread scheduler to schedule a thread for execution on the graphics processor pipeline. The one or more neural network hardware blocks can execute at least one layer of a neural network to determine a number of threads to dispatch to the block of graphics compute units within a dispatch cycle.

[0260] In one embodiment the graphics processor additionally includes a cache memory and a cache memory prefetcher. The one or more neural network hardware blocks can execute at least one layer of a neural network to determine a pre-fetch pattern for the cache memory prefetcher.

[0261] One embodiment provides for a graphics processor comprising a block of graphics compute units, a memory fetch unit coupled with the block of graphics compute units, and a programmable neural network unit including one or more neural network hardware blocks, wherein a neural network hardware block includes hardware logic to perform neural network operations and activation operations for a layer of a neural network, the programmable neural network unit addressable by compute units within the block of compute units. In one embodiment the programmable neural network unit includes multiple neural network hardware blocks.

[0262] Various applications of the neural network block are provided by embodiments described herein. In one embodiment the programmable neural network unit is to determine visibility for a geometry culling operation via the neural network hardware block, the neural network hardware block to determine visibility on a per-object basis. In one embodiment the programmable neural network unit is to configure one or more neural network hardware blocks with a meta-shader neural network. The meta-shader neural network can generate a texture for one of multiple indicated types of terrain.

[0263] One embodiment provides for an electronic device comprising a block of graphics compute

units, a graphics processor pipeline coupled to the block of graphics compute units, and a programmable neural network unit including one or more neural network hardware blocks. The programmable neural network unit is coupled with the block of graphics compute units and the graphics processor pipeline. The electronic device also includes a tessellation module to configure the programmable neural network unit to generate tessellated output based on coarse input data. [0264] One embodiment provides a graphics processor comprising a block of execution resources; a cache memory; a cache memory prefetcher; and circuitry including a programmable neural network unit, the programmable neural network unit comprising a network hardware block including circuitry to perform neural network operations and activation operations for a layer of a neural network, the programmable neural network unit addressable by cores within the block of graphics cores and the neural network hardware block configured to perform operations associated with a neural network configured to determine a prefetch pattern for the cache memory prefetcher. [0265] Those skilled in the art will appreciate from the foregoing description that the broad techniques of the embodiments can be implemented in a variety of forms. Therefore, while the embodiments have been described in connection with particular examples thereof, the true scope of the embodiments should not be so limited since other modifications will become apparent to the skilled practitioner upon a study of the drawings, specification, and following claims.

Claims

1. A graphics processor comprising: a block of graphics cores; and circuitry including a programmable neural network unit, the programmable neural network unit including one or more neural network hardware blocks, wherein a neural network hardware block includes circuitry to perform neural network operations and activation operations for a layer of a neural network, the programmable neural network unit addressable by cores within the block of graphics cores, wherein the programmable neural network unit is to configure one or more neural network hardware blocks with a meta-shader neural network, the meta-shader neural network to generate a texture for one of multiple types of terrain.
2. The graphics processor of claim 1, wherein the meta-shader neural network is to generate a texture for one of multiple types of terrain based on one or more input parameters to the meta-shader neural network.
3. The graphics processor of claim 1, wherein the one or more neural network hardware blocks include a source data buffer, a neural network operations and activation operations block, and an output data buffer.
4. The graphics processor of claim 3, wherein the neural network operations and activation operations block is programmably configurable.
5. The graphics processor of claim 4, wherein the programmable neural network unit includes a block programming unit to configure layer state information for the one or more neural network hardware blocks, the layer state information associated with one or more layers of a neural network to be processed by the programmable neural network unit.
6. The graphics processor of claim 4, wherein the programmable neural network unit includes weights cache to cache weights associated with one or more layers of a neural network to be processed by the programmable neural network unit.
7. The graphics processor of claim 6, wherein the programmable neural network unit includes multiple neural network hardware blocks.
8. The graphics processor of claim 7, wherein multiple neural network hardware blocks are respectively associated with one or more layers of the neural network to be processed by the programmable neural network unit.
9. The graphics processor of claim 1, wherein the programmable neural network unit is to determine visibility for a geometry culling operation via the neural network hardware block and the

neural network hardware block is to determine visibility on a per-object basis.

10. The graphics processor of claim 1, further comprising a tessellation module to configure the programmable neural network unit to generate tessellated output based on coarse input data.

11. A method comprising: addressing a programmable neural network unit within circuitry of a graphics processor, the graphics processor including a block of graphics cores, the programmable neural network unit comprising one or more neural network hardware blocks, wherein a neural network hardware block includes circuitry to perform neural network operations and activation operations for a layer of a neural network; configuring the one or more neural network hardware blocks within the programmable neural network unit with a meta-shader neural network; and generating a texture for one of multiple types of terrain via the meta-shader neural network based on one or more input parameters to the meta-shader neural network.

12. The method of claim 11, wherein the one or more neural network hardware blocks include a source data buffer, a neural network operations and activation operations block, and an output data buffer, the method additionally comprising programmably configuring a neural network operations and activation operations block.

13. The method of claim 12, comprising configuring layer state information for the one or more neural network hardware blocks via a block programming unit of the programmable neural network unit, the layer state information associated with one or more layers of a neural network to be processed by the programmable neural network unit.

14. The method of claim 13, wherein multiple neural network hardware blocks are each associated with one or more layers of the neural network to be processed by the programmable neural network unit.

15. A data processing system comprising: a memory device; and an accelerator device coupled with the memory device, the accelerator device includes a processing cluster and circuitry including a programmable neural network unit, the programmable neural network unit including: one or more neural network hardware blocks, wherein a neural network hardware block includes circuitry to perform neural network operations and activation operations for a layer of a neural network, the programmable neural network unit addressable by cores within the processing cluster, wherein the programmable neural network unit is to configure one or more neural network hardware blocks with a meta-shader neural network, the meta-shader neural network to generate a texture for one of multiple types of terrain based on one or more input parameters to the meta-shader neural network.

16. The data processing system of claim 15, wherein the one or more neural network hardware blocks include a source data buffer, a neural network operations and activation operations block, and an output data buffer.

17. The data processing system of claim 16, wherein the neural network operations and activation operations block is programmably configurable.

18. The data processing system of claim 17, wherein the programmable neural network unit includes a block programming unit to configure layer state information for the one or more neural network hardware blocks, the layer state information associated with one or more layers of a neural network to be processed by the programmable neural network unit.

19. The data processing system of claim 17, wherein the programmable neural network unit includes weights cache to cache weights associated with one or more layers of a neural network to be processed by the programmable neural network unit.

20. The data processing system of claim 19, wherein the programmable neural network unit includes multiple neural network hardware blocks, the multiple neural network hardware blocks respectively associated with one or more layers of the neural network to be processed by the programmable neural network unit.
