

# US Patent & Trademark Office

## Patent Public Search | Text View

---

United States Patent Application Publication

20250252664

Kind Code

A1

Publication Date

August 07, 2025

Inventor(s)

Nonn; Thomas Ivan et al.

---

### **METHODS AND SYSTEMS FOR AUGMENTING DEPTH DATA FROM A DEPTH SENSOR, SUCH AS WITH DATA FROM A MULTIVIEW CAMERA SYSTEM**

---

#### **Abstract**

Methods of determining the depth of a scene and associated systems are disclosed herein. In some embodiments, a method can include augmenting depth data of a scene captured with a depth sensor with depth data from one or more images of the scene. For example, the method can include capturing image data of the scene with a plurality of cameras. The method can further include generating a point cloud representative of the scene based on the depth data from the depth sensor and identifying a missing region of the point cloud, such as a region occluded from the view of the depth sensor. The method can then include generating depth data for the missing region based on the image data. Finally, the depth data for the missing region can be merged with the depth data from the depth sensor to generate a merged point cloud representative of the scene.

---

**Inventors:** Nonn; Thomas Ivan (Kenmore, WA), Colmenares; David Julio (Seattle, WA), Youngquist; James Andrew (Seattle, WA), Jones; Adam Gabriel (Seattle, WA)

**Applicant:** Proprio, Inc. (Seattle, WA)

**Family ID:** 76857295

**Appl. No.:** 19/035856

**Filed:** January 24, 2025

#### **Related U.S. Application Data**

parent US continuation 18312430 20230504 parent-grant-document US 12243162 child US 19035856

parent US continuation 17154670 20210121 parent-grant-document US 11682165 child US 18312430

us-provisional-application US 62963717 20200121

---

## Publication Classification

**Int. Cl.:** **G06T17/00** (20060101); **G06T7/557** (20170101); **G06T17/20** (20060101); **G06V10/80** (20220101); **G06V20/20** (20220101); **G06V20/64** (20220101); **H04N13/00** (20180101); **H04N13/282** (20180101)

**U.S. Cl.:**

**CPC** **G06T17/00** (20130101); **G06T7/557** (20170101); **G06T17/20** (20130101); **G06V10/806** (20220101); **G06V20/20** (20220101); **G06V20/647** (20220101); **H04N13/282** (20180501); **G06T2207/10052** (20130101); **G06T2207/20104** (20130101); **G06T2210/41** (20130101); **H04N2013/0081** (20130101)

---

## Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS [0001] This application is a continuation application of U.S. patent application Ser. No. 18/312,430, filed on May 4, 2023, and titled “METHODS AND SYSTEMS FOR AUGMENTING DEPTH DATA FROM A DEPTH SENSOR, SUCH AS WITH DATA FROM A MULTIVIEW CAMERA SYSTEM,” which is a continuation application of U.S. patent application Ser. No. 17/154,670, now U.S. Pat. No. 11,682,165, filed on Jan. 21, 2021, and titled “METHODS AND SYSTEMS FOR AUGMENTING DEPTH DATA FROM A DEPTH SENSOR, SUCH AS WITH DATA FROM A MULTIVIEW CAMERA SYSTEM,” which claims the benefit of U.S. Provisional Patent Application No. 62/963,717, filed Jan. 21, 2020, and titled “METHODS AND SYSTEMS FOR AUGMENTING DEPTH DATA FROM A DEPTH SENSOR, SUCH AS WITH DATA FROM A MULTIVIEW CAMERA SYSTEM,” each of which is incorporated herein by reference in its entirety.

### TECHNICAL FIELD

[0002] The present technology generally relates to imaging systems, and more specifically, to imaging systems for generating a virtual perspective of a scene for a mediated-reality viewer.

### BACKGROUND

[0003] In a mediated reality system, an image processing system adds, subtracts, and/or modifies visual information representing an environment. For surgical applications, a mediated reality system can enable a surgeon to view a surgical site from a desired perspective together with contextual information that assists the surgeon in more efficiently and precisely performing surgical tasks. Such mediated reality systems rely on image data from multiple camera angles and depth information about the environment to reconstruct an image of the environment.

[0004] Depth information about the environment is typically obtained via a dedicated depth sensor, such as a structured light depth sensor. However, the capture of complex three-dimensional geometries using a structured light depth sensor requires that all surfaces of the environment be exposed to the depth sensor to acquire error-free and complete data. In practice, this is not feasible due to geometries and/or occluded surfaces hidden from the depth sensor.

---

## Description

### BRIEF DESCRIPTION OF THE DRAWINGS

[0005] Many aspects of the present disclosure can be better understood with reference to the following drawings. The components in the drawings are not necessarily to scale. Instead, emphasis

is placed on clearly illustrating the principles of the present disclosure.

[0006] FIG. 1 is a schematic view of an imaging system configured in accordance with embodiments of the present technology.

[0007] FIG. 2 is a perspective view of a surgical environment employing the imaging system of FIG. 1 for a surgical application in accordance with embodiments of the present technology.

[0008] FIG. 3 is an isometric view of a portion of the imaging system including a depth sensor and multiple cameras configured in accordance with embodiments of the present technology.

[0009] FIG. 4 is a schematic view of a point cloud generated by the imaging system in accordance with embodiments of the present technology.

[0010] FIG. 5 is a flow diagram of a process or method for augmenting depth data captured with the depth sensor using image data captured by the cameras in accordance with embodiments of the present technology.

[0011] FIG. 6 is a schematic view of a merged point cloud in which depth data has been filled in for missing regions of the point cloud of FIG. 4 in accordance with embodiments of the present technology.

[0012] FIG. 7 is an isometric view of a portion of the imaging system configured in accordance with embodiments of the present technology.

[0013] FIG. 8 is a flow diagram of a process or method for augmenting depth data captured with the depth sensor with depth data from a medical scan of a patient in accordance with embodiments of the present technology.

[0014] FIG. 9A is a schematic view of a point cloud generated by the imaging system in accordance with embodiments of the present technology.

[0015] FIG. 9B is a schematic view of CT scan data corresponding to the point cloud of FIG. 9A in accordance with embodiments of the present technology.

[0016] FIG. 9C is a schematic view of a merged point cloud in which the CT scan data shown in FIG. 9B has been merged with the point cloud shown in FIG. 9A in accordance with embodiments of the present technology.

#### DETAILED DESCRIPTION

[0017] Aspects of the present disclosure are directed generally to methods of determining the depth of a scene, such as surgical scene, and using the depth information to reconstruct a virtual camera perspective of the scene. In several of the embodiments described below, for example, a method includes augmenting depth data of the scene captured with a depth sensor with depth data from one or more images of the scene. For example, the method can include capturing (i) the depth data of the scene with the depth sensor and (ii) images of the scene with a plurality of cameras. The method can further include generating a point cloud representative of the scene based on the depth data from the depth sensor and identifying a missing region of the point cloud, such as a region occluded from the view of the depth sensor. The method can then include generating depth data for the missing region based on the images from the cameras. The images can be light field images containing information about the intensity of light rays emanating from the scene and also information about a direction the light rays are traveling through space. The method can further include merging (i) the depth data for the missing region derived from the images with (ii) the depth data from the depth sensor to generate a merged point cloud representative of the scene.

[0018] In one aspect of the present technology, the merged point cloud can have a greater accuracy and/or resolution than the point cloud generated from the depth data from the depth sensor alone. In another aspect of the present technology, depth information is determined quickly for as much of the scene as possible using a depth sensor, and light field processing is used only for the relatively small regions of the scene where depth information cannot be or cannot accurately be determined using the depth sensor (e.g., the missing regions). Accordingly, the present technology can provide real time or near real time depth and image processing while also providing improved accuracy. That is, the combined depth determination approach of the present technology can provide (i)

improved latency compared to light field processing alone and (ii) improved accuracy compared to depth sensor processing alone.

[0019] Specific details of several embodiments of the present technology are described herein with reference to FIGS. **1-9C**. The present technology, however, can be practiced without some of these specific details. In some instances, well-known structures and techniques often associated with camera arrays, light field cameras, image reconstruction, depth sensors, and the like have not been shown in detail so as not to obscure the present technology. The terminology used in the description presented below is intended to be interpreted in its broadest reasonable manner, even though it is being used in conjunction with a detailed description of certain specific embodiments of the disclosure. Certain terms can even be emphasized below; however, any terminology intended to be interpreted in any restricted manner will be overtly and specifically defined as such in this Detailed Description section.

[0020] The accompanying Figures depict embodiments of the present technology and are not intended to be limiting of its scope. The sizes of various depicted elements are not necessarily drawn to scale, and these various elements can be arbitrarily enlarged to improve legibility. Component details can be abstracted in the Figures to exclude details such as position of components and certain precise connections between such components when such details are unnecessary for a complete understanding of how to make and use the present technology. Many of the details, dimensions, angles, and other features shown in the Figures are merely illustrative of particular embodiments of the disclosure. Accordingly, other embodiments can have other details, dimensions, angles, and features without departing from the spirit or scope of the present technology.

[0021] The headings provided herein are for convenience only and should not be construed as limiting the subject matter disclosed.

## I. Selected Embodiments of Imaging Systems

[0022] FIG. **1** is a schematic view of an imaging system **100** (“system **100**”) configured in accordance with embodiments of the present technology. In the illustrated embodiment, the system **100** includes an image processing device **102** that is operably/communicatively coupled to one or more display devices **104**, one or more input controllers **106**, and a camera array **110**. In other embodiments, the system **100** can comprise additional, fewer, or different components. In some embodiments, the system **100** can include features that are generally similar or identical to those of the mediated-reality imaging systems disclosed in U.S. patent application Ser. No. 16/586,375, titled “CAMERA ARRAY FOR A MEDIATED-REALITY SYSTEM,” which is incorporated herein by reference in its entirety.

[0023] In the illustrated embodiment, the camera array **110** includes a plurality of cameras **112** (identified individually as cameras **112a-112n**) that are each configured to capture images of a scene **108** from a different perspective. In some embodiments, the cameras **112** are positioned at fixed locations and orientations relative to one another. For example, the cameras **112** can be structurally secured by/to a mounting structure (e.g., a frame) at predefined fixed locations and orientations. In some embodiments, the cameras **112** can be positioned such that neighboring cameras share overlapping views of the scene **108**. In some embodiments, the cameras **112** in the camera array **110** are synchronized to capture images of the scene **108** substantially simultaneously (e.g., within a threshold temporal error). In some embodiments, all or a subset of the cameras **112** can be light-field/plenoptic/RGB cameras that are configured to capture information about the light field emanating from the scene **108** (e.g., information about the intensity of light rays in the scene **108** and also information about a direction the light rays are traveling through space).

[0024] In the illustrated embodiment, the camera array **110** further comprises (i) one or more projectors **114** configured to project a structured light pattern onto/into the scene **108**, and (ii) one or more depth sensors **116** configured to estimate a depth of a surface in the scene **108**. In some embodiments, the depth sensor **116** can estimate depth based on the structured light pattern emitted

from the projector **114**.

[0025] The image processing device **102** is configured to (i) receive images (e.g., light-field images, light field image data) captured by the camera array **110** and depth information from the depth sensor **116**, and (ii) process the images and depth information to synthesize an output image corresponding to a virtual camera perspective. In the illustrated embodiment, the output image corresponds to an approximation of an image of the scene **108** that would be captured by a camera placed at an arbitrary position and orientation corresponding to the virtual camera perspective. More specifically, the depth information can be combined with the images from the cameras **112** to synthesize the output image as a three-dimensional rendering of the scene **108** as viewed from the virtual camera perspective. In some embodiments, the image processing device **102** can synthesize the output image using any of the methods disclosed in U.S. patent application Ser. No. 16/457,780, titled "SYNTHESIZING AN IMAGE FROM A VIRTUAL PERSPECTIVE USING PIXELS FROM A PHYSICAL IMAGER ARRAY WEIGHTED BASED ON DEPTH ERROR SENSITIVITY," which is incorporated herein by reference in its entirety.

[0026] The image processing device **102** can synthesize the output image from a subset (e.g., two or more) of the cameras **112** in the camera array **110**, but does not necessarily utilize images from all of the cameras **112**. For example, for a given virtual camera perspective, the image processing device **102** can select a stereoscopic pair of images from two of the cameras **112** that are positioned and oriented to most closely match the virtual camera perspective. In some embodiments, the image processing device **102** (and/or depth sensor **116**) is configured to estimate a depth for each surface point of the scene **108** and to generate a point cloud and/or three-dimensional (3D) mesh that represents the surface of the scene **108**. For example, in some embodiments the depth sensor **116** can detect the structured light projected onto the scene **108** by the projector **114** to estimate depth information of the scene **108**. Alternatively or additionally, the image processing device **102** can perform the depth estimation based on depth information received from the depth sensor **116**. As described in detail below, in some embodiments the image processing device **102** can estimate depth from multiview image data from the cameras **112** with or without utilizing information collected by the projector **114** or the depth sensor **116**.

[0027] In some embodiments, functions attributed to the image processing device **102** can be practically implemented by two or more physical devices. For example, in some embodiments a synchronization controller (not shown) controls images displayed by the projector **114** and sends synchronization signals to the cameras **112** to ensure synchronization between the cameras **112** and the projector **114** to enable fast, multi-frame, multi-camera structured light scans. Additionally, such a synchronization controller can operate as a parameter server that stores hardware specific configurations such as parameters of the structured light scan, camera settings, and camera calibration data specific to the camera configuration of the camera array **110**. The synchronization controller can be implemented in a separate physical device from a display controller that controls the display device **104**, or the devices can be integrated together.

[0028] The image processing device **102** can comprise a processor and a non-transitory computer-readable storage medium that stores instructions that when executed by the processor, carry out the functions attributed to the image processing device **102** as described herein. Although not required, aspects and embodiments of the present technology can be described in the general context of computer-executable instructions, such as routines executed by a general-purpose computer, e.g., a server or personal computer. Those skilled in the relevant art will appreciate that the present technology can be practiced with other computer system configurations, including Internet appliances, hand-held devices, wearable computers, cellular or mobile phones, multi-processor systems, microprocessor-based or programmable consumer electronics, set-top boxes, network PCs, mini-computers, mainframe computers and the like. The present technology can be embodied in a special purpose computer or data processor that is specifically programmed, configured or constructed to perform one or more of the computer-executable instructions explained in detail

below. Indeed, the term “computer” (and like terms), as used generally herein, refers to any of the above devices, as well as any data processor or any device capable of communicating with a network, including consumer electronic goods such as game devices, cameras, or other electronic devices having a processor and other components, e.g., network communication circuitry.

[0029] The invention can also be practiced in distributed computing environments, where tasks or modules are performed by remote processing devices, which are linked through a communications network, such as a Local Area Network (“LAN”), Wide Area Network (“WAN”), or the Internet. In a distributed computing environment, program modules or sub-routines can be located in both local and remote memory storage devices. Aspects of the invention described below can be stored or distributed on computer-readable media, including magnetic and optically readable and removable computer discs, stored as in chips (e.g., EEPROM or flash memory chips). Alternatively, aspects of the invention can be distributed electronically over the Internet or over other networks (including wireless networks). Those skilled in the relevant art will recognize that portions of the present technology can reside on a server computer, while corresponding portions reside on a client computer. Data structures and transmission of data particular to aspects of the present technology are also encompassed within the scope of the invention.

[0030] The virtual camera perspective can be controlled by an input controller **106** that provides a control input corresponding to the location and orientation of the virtual camera perspective. The output images corresponding to the virtual camera perspective are outputted to the display device **104**. The display device **104** is configured to receive output images (e.g., the synthesized three-dimensional rendering of the scene **108**) and to display the output images for viewing by one or more viewers. The image processing device **102** can beneficially process received inputs from the input controller **106** and process the captured images from the camera array **110** to generate output images corresponding to the virtual perspective in substantially real-time as perceived by a viewer of the display device **104** (e.g., at least as fast as the frame rate of the camera array **110**).

[0031] The display device **104** can comprise, for example, a head-mounted display device, a monitor, a computer display, and/or another display device. In some embodiments, the input controller **106** and the display device **104** are integrated into a head-mounted display device and the input controller **106** comprises a motion sensor that detects position and orientation of the head-mounted display device. The virtual camera perspective can then be derived to correspond to the position and orientation of the head-mounted display device **104** such that the virtual perspective corresponds to a perspective that would be seen by a viewer wearing the head-mounted display device **104**. Thus, in such embodiments the head-mounted display device **104** can provide a real-time rendering of the scene **108** as it would be seen by an observer without the head-mounted display device **104**. Alternatively, the input controller **106** can comprise a user-controlled control device (e.g., a mouse, pointing device, handheld controller, gesture recognition controller) that enables a viewer to manually control the virtual perspective displayed by the display device **104**.

[0032] FIG. 2 is a perspective view of a surgical environment employing the system **100** for a surgical application in accordance with embodiments of the present technology. In the illustrated embodiment, the camera array **110** is positioned over the scene **108** (e.g., a surgical site) and supported/positioned via a swing arm **222** that is operably coupled to a workstation **224**. In some embodiments, the swing arm **222** can be manually moved to position the camera array **110** while, in other embodiments, the swing arm **222** can be robotically controlled in response to the input controller **106** (FIG. 1) and/or another controller. In the illustrated embodiment, the display device **104** is embodied as a head-mounted display device (e.g., a virtual reality headset, augmented reality headset). The workstation **224** can include a computer to control various functions of the image processing device **102**, the display device **104**, the input controller **106**, the camera array **110**, and/or other components of the system **100** shown in FIG. 1. Accordingly, in some embodiments the image processing device **102** and the input controller **106** are each integrated in the workstation **224**. In some embodiments, the workstation **224** includes a secondary display **226** that can display

a user interface for performing various configuration functions, a mirrored image of the display on the display device **104**, and/or other useful visual images/indications.

## II. Selected Embodiments of Augmenting Depth Data from a Depth Sensor

[0033] FIG. **3** is an isometric view of a portion of the system **100** illustrating four of the cameras **112** and the depth sensor **116** in accordance with embodiments of the present technology. Other components of the system **100** (e.g., other portions of the camera array **110**, the image processing device **102**, etc.) are not shown in FIG. **3** for the sake of clarity. In the illustrated embodiment, each of the cameras **112** has a field of view **330** and is oriented such that the field of view **330** is aligned with a portion of the scene **108**. Likewise, the depth sensor **116** can have a field of view **332** aligned with a portion of the scene **108**. In some embodiments, a portion of some or all of the field of views **330**, **332** can overlap.

[0034] In the illustrated embodiment, a portion of a spine **309** of a patient (e.g., a human patient) is located in/at the scene **108**. Often, the spine **309** (or other surfaces located in the scene **108**) will have a complex 3D geometry such that is difficult to accurately determine the depth of its surfaces, and therefore difficult to accurately model with a point cloud, 3D mesh, and/or other mathematical representation. For example, if a portion of the surface of the scene **108** (e.g., a portion of the spine **309**) is occluded from the field of view **332** of the depth sensor **116**, the depth sensor **116** will be unable to determine the depth of the occluded region. Likewise, it can be difficult to accurately determine the depth along steep surfaces of the scene **108**. More specifically, for a structured light system to recover depth at a given location the structured light projector (e.g., the projector **114**) must illuminate that location with a pixel or block of pixels of structured illumination. Also, the imagers (e.g., the depth sensor **116**) measuring the projection must have a pixel/block that sees that illumination. Both conditions must be met to make a measurement of the location. In practice, it is typically not possible to achieve a fill rate of 100%—where every pixel of depth has a valid value. This is because real scenes have complex geometries that cause occlusion of the projector, the imagers/sensors, or both. Accordingly, if the system **100** uses only the depth sensor **116** to determine depth, a depth model generated by the system **100** can have missing regions (e.g., holes) corresponding to the portions (e.g., surfaces) of the scene **108** where depth information is unavailable.

[0035] In some embodiments, the system **100** is unable to adequately generate—or unable to accurately generate—an output image of the scene **108** for such portions of the scene **108** that have inadequate and/or inaccurate depth information. FIG. **4**, for example, is a schematic view of a point cloud **440** generated by the depth sensor **116** of the system **100** for the surfaces of the spine **309** shown in FIG. **3** in accordance with embodiments of the present technology. Referring to FIGS. **1-4** together, the point cloud **440** generally comprises a plurality (e.g., hundreds, thousands, millions, or more) of data points corresponding to a distance of the surfaces of the spine **309** and/or other features in the scene **108** relative to the sensor **116** (e.g., the depth of the spinal surfaces). The point cloud **440** therefore maps/represents the 3D surface of the spine **309** and can be used by the image processing device **102** to synthesize the images from the cameras **112** into the output image of the scene **108** rendered from any desired virtual perspective, as described in detail above. In the illustrated embodiment, the point cloud **440** includes one or more missing regions **442** corresponding to portions (e.g., surfaces) of the scene **108** where depth information is inadequate and/or not reliable (e.g., regions where the depth sensor **116** is occluded). Accordingly, the system **100** may not be able to render an accurate output image for those portions of the scene **108**.

[0036] In some embodiments, the image processing device **102** can process image data from one or more of the cameras **112** to determine the depth of the spinal surfaces at one or more of the locations where depth information from the depth sensor **116** is inadequate, unreliable, and/or inaccurate. That is, the image data from the cameras **112** can be used to “fill in” the missing regions **442** of the point cloud **440**. More specifically, FIG. **5** is a flow diagram of a process or method **550** for augmenting depth data captured with the depth sensor **116** using image data captured by the

cameras **112** in accordance with embodiments of the present technology. Although some features of the method **550** are described in the context of the embodiments shown in FIGS. **1-4** for the sake of illustration, one skilled in the art will readily understand that the method **550** can be carried out using other suitable systems and/or devices described herein.

[0037] At block **551**, the method **550** includes (a) capturing depth data of the scene **108** (e.g., of the spine **309**) from the depth sensor **116** and (b) image data of the scene from one or more of the cameras **112**. The depth data can include, for example, data about a structured light pattern projected onto/into the scene **108** (e.g., from the projector **114**). The image data can be light field data including data about the intensity of light rays emanating from the scene **108** and also information about a direction the light rays are traveling. In some embodiments, the depth sensor **116** and the cameras **112** can capture the depth data and the image data simultaneously or substantially simultaneously and/or in real time or near real time. In other embodiments, the depth data can be captured before the image data.

[0038] At block **552**, the method **550** includes generating a point cloud of the scene **108**, such as the point cloud **440**, based on the depth data from the depth sensor **116**. In some embodiments, the image processing device **102** can receive the depth data from the depth sensor **116** and generate the point cloud based on the depth data. In some embodiments, the method **550** can further include generating a 3D mesh instead of or in addition to a point cloud. In other embodiments, at block **552** the method **550** can include generating other mathematical representations of the physical geometry of the scene **108**.

[0039] At block **553**, the method **550** includes back projecting the point cloud and/or depth data associated with the point cloud to individual ones of the cameras **112** and/or to the image processing device **102**. Back projecting the point cloud to the cameras **112** allows an image of the scene **108** to be reconstructed. More specifically, back projection correlates a 2D pixel location in the images from the cameras **112** with a 3D position from the point cloud. By back projecting the point cloud to each of the images from the cameras **112**, each pixel in the images can be associated with a 3D point—or not if the 3D position cannot be determined for the reasons discussed in detail above. An even simpler classifier is to label each pixel in the 2D image as having a valid 3D correspondence or not. In some embodiments, this classification can be used to create a binary mask for each of the cameras **112** that indicates which pixels have a valid 3D point.

[0040] At block **554**, the method **550** includes identifying regions of missing data in the point cloud. For example, the method **550** can include identifying the missing regions **442** of the point cloud **440** where depth data is missing or incomplete. In some embodiments, identifying the missing data can include filtering the point cloud data and searching for holes that are greater than a predetermined threshold (e.g., a user-specified threshold) using, for example, an inverse Eulerian approach. In some embodiments, identifying missing data can include scanning the point cloud to determine regions with sparse or non-existent points. In some embodiments, a mesh can be generated for the point cloud (e.g., at block **552**), and holes can be identified in the mesh using, for example, a method that identifies triangles in the mesh having at least one edge that is not shared by another triangle. In yet other embodiments, the missing regions **442** can be identified by searching for regions of the images from the cameras **112** where no valid 3D correspondence exists (e.g., by examining the binary mask for each image). In some embodiments, blocks **553** and **554** can be executed using the same algorithm and/or as part of the same computational process.

[0041] In some embodiments, at block **554**, the method **550** can additionally or alternatively include identifying regions of invalid depth data, low confidence depth data, and/or other potentially problematic regions of the point cloud. For example, the depth sensor **116** can be configured to tag the depth data it captures with validation or confidence levels, and the method **550** can include identifying regions of the point cloud and/or mesh with validation or confidence levels that are below a predetermined threshold (e.g., a user-specified threshold). Such invalid or low confidence regions can be regions of the point cloud or mesh having discontinuities, sparse



depth data, badly behaved normal values, and the like. In some embodiments, the method 550 may not identify single missing pixels as missing or invalid regions, and/or conversely may identify as missing/invalid regions of missing pixels with some “valid” pixels interspersed.

[0042] In some embodiments, at block 554, the method 550 can further include determining depth data for areas surrounding the missing regions 442 of the point cloud. This surrounding depth data can help inform/predict the depth of the missing regions 442 if it is assumed that there are not large discontinuities between the missing regions 442 and the surrounding areas, such that the missing depth values can be expected to be close to the surrounding depths.

[0043] At block 555, the method 550 includes extracting/identifying image data corresponding to the missing or invalid regions of the point cloud or mesh. For example, the image processing device 102 can determine which of the cameras 112 have their field of view 330 aligned with the region of the scene 108 that corresponds to the missing region. In some embodiments, the image processing device 102 can make this determination based on a priori information about (i) the positions and orientations of the cameras (and thus the extent of their fields of view 330), (ii) the back projection of the depth data to the cameras 112 (block 553), (iii) processing of the point cloud or mesh, and/or (iv) other data. Moreover, in some embodiments the system 100 can identify and extract image data from only those of the cameras 112 that are determined to have adequate optical coverage of the missing regions. In some embodiments, at least some of the cameras 112 can have at least partially overlapping fields of view 330 such that it is very likely that at least one of the cameras 112 has a field of view 330 aligned with the region of the scene 108 that corresponds to the missing region-even when other ones of the cameras 112 are occluded. Accordingly, in one aspect of the present technology the system 100 is configured to robustly capture image data about the missing regions even where substantial occlusions exist in the scene 108. In some embodiments, blocks 553-555 can be executed using the same algorithm and/or as part of the same computational process.

[0044] At block 556, the method 550 includes processing the extracted image data to generate depth data for the missing or invalid regions. For example, the image processing device 102 can generate depth data for the missing regions using the disparity from the cameras 112 that have the missing regions within their field of view 330 (e.g., that are facing the missing regions). In other embodiments, other suitable image processing techniques (e.g., computational algorithms) for determining depth from light field data can be used. In some embodiments, determining depth by processing the image data from the cameras 112 can be more computationally expensive (e.g., slower) than determining depth using the depth sensor 116 because of the complex nature of computational algorithms for processing depth information from light field data. As a result, image data from less than all of the cameras 112 may be used to generate depth data for the missing or invalid regions. In some embodiments, depth information about the areas surrounding the missing or invalid regions (e.g., captured at block 554) can be used to accelerate processing of the extracted image data. Specifically, many depth processing algorithms iterate through depths to search for the true values. Accordingly, by limiting the depth range based on the depth of the surrounding areas, a smaller range of depths, disparities, planes, and so on have to be searched through. Thus, the search can avoid local minima that may exist outside this expected region/range-accelerating processing.

[0045] At block 557, the method 550 includes merging/fusing the depth data for the missing or invalid regions with the original depth data (e.g., captured at block 551) to generate a merged point cloud. FIG. 6, for example, is a schematic view of a merged point cloud 640 in which image-based depth data 644 has been filled into the missing regions 442 of the point cloud 440 shown in FIG. 4 in accordance with embodiments of the present technology. Accordingly, the merged point cloud 640 can provide a more accurate and robust depth map of the scene 108 that facilitates better reconstruction and synthesis of an output image of the scene 108 rendered from any desired virtual perspective, as described in detail above.

[0046] At block 558, the method 550 can optionally include generating a three-dimensional mesh

based on the merged point cloud. The 3D mesh can be used to reconstruct/synthesize the output image of the scene **108**. In some embodiments, the method **550** can return to block **551** to update the depth information of the scene **108**. In some embodiments, the method **550** can proceed to back project the merged point cloud to the cameras **112** (block **553**).

[0047] As noted above, determining depth by processing light field image data can be more computationally expensive than determining depth using a depth sensor. Indeed, if depth information for an entire scene were determined entirely through light field image processing, it would be difficult/impracticable to render output images in real time or near real-time because even very fast systems cannot measure and process the significant volume of data fast enough. However, in one aspect of the present technology depth information is determined quickly for as much of the scene as possible using a depth sensor, and light field processing is used only for the relatively small regions of the scene where there is inadequate and/or unreliable depth information from the depth sensor. Accordingly, the present technology can provide real time or near real time depth and image processing while also providing improved accuracy. That is, the combined depth determination approach of the present technology can provide (i) improved latency compared to light field processing alone and (ii) improved accuracy and resolution compared to depth sensor processing alone.

[0048] In some embodiments, the latency of the system **100** can be further improved by updating the depth information only for missing or invalid regions of the point cloud for which increased accuracy is desired. For example, FIG. 7 is an isometric view of a portion of the system **100** illustrating two of the cameras **112**, the display device **104**, and the generated point cloud **440** (FIG. 4) in accordance with embodiments of the present technology. Other components of the system **100** (e.g., other portions of the camera array **110**, the image processing device **102**, etc.) are not shown in FIG. 7 for the sake of clarity.

[0049] In the illustrated embodiment, the display device **104** is a head-mounted display device **104** (e.g., a headset) configured to be worn by a user (e.g., a surgeon) and having a field of view **736** that is aligned with only a portion of the scene **108** (e.g., a portion of the spine **309** shown in FIG. 3). The head-mounted display device **104** can include a display **705** configured to display the rendered output image of the scene **108** to the user. The display **705** can be opaque or partially transparent. In some embodiments, the field of view **736** of the head-mounted display device **104** corresponds to a foveated region that represents the relatively narrow field of view that the eyes of the user can perceive.

[0050] The system **100** (e.g., the image processing device **102**) can track the position and orientation of the field of view **736** relative to the scene **108** and can employ the method **550** (FIG. 5) to update only the missing regions **442** of the point cloud **440** that are within the field of view **736**—without updating regions outside of the field of view **736**. In some embodiments, the system **100** can identify the cameras **112** that have the best optical coverage of the portion of the scene **108** within the field of view **736**. When the user changes the position and/or orientation of the head-mounted display device **104**—and thus the field of view **736**—the system **100** can seamlessly update (e.g., fill-in) the missing regions **442** that are within the field of view **736** in real time or near real time. In one aspect of the present technology, the latency of the image presented to the user via the head-mounted display device **104** is decreased because the missing regions **442** that are outside the foveated region of the user are not updated.

[0051] Referring again to FIG. 1, in some embodiments the cameras **112** can have a higher resolution than the depth sensor **116** such that more depth detail of the scene **108** can be extracted from the cameras **112** than from the depth sensor **116**. Accordingly, even where depth information from the depth sensor **116** exists and is at least adequate to determine the general depth of the scene **108**, it can be advantageous to also include image data from the cameras **112** to increase the depth resolution and, correspondingly, the resolution of the image output to the user via the display device **104**. Therefore, in some embodiments the system **100** can process image data for particular

local regions of the scene **108** to supplement or replace the depth data captured by the depth sensor **116** for those local regions. In some embodiments, a background process running on the image processing device **102** can update the local regions of the scene **108** automatically if, for example, the depth data from the depth sensor **116** is of poor quality in those regions. In other embodiments, the user can select certain areas in which to improve the resolution. In yet other embodiments, the system **100** can improve the resolution by processing light field data corresponding to all or a portion of the foveated region **736** of the user as shown in FIG. 7.

[0052] In other embodiments, depth data captured by the depth sensor **116** can be supplemented or replaced with depth information obtained from means other than processing image data from the cameras **112**. For example, FIG. 8 is a flow diagram of a process or method **860** for augmenting depth data captured with the depth sensor **116** with depth data from one or more medical scans of a patient in accordance with embodiments of the present technology. Although some features of the method **860** are described in the context of the system **100** shown in FIG. 1 for the sake of illustration, one skilled in the art will readily understand that the method **860** can be carried out using other suitable systems and/or devices described herein. Moreover, while the method **860** is described in the context of augmenting depth data of the anatomy of a patient with medical scans of the patient, the method **860** can be practiced to update/augment depth data for other scenes and/or based on other data from other imaging/scanning techniques.

[0053] At block **861**, the method **860** includes capturing depth data of the scene **108** (e.g., live data) from the depth sensor **116**, such as data about a structured light pattern projected onto/into the scene **108**. The scene **108** can include, for example, a portion of a patient undergoing surgery. As one example, the portion of the patient can be a portion of the patient's spine exposed during spinal surgery. Block **862** of the method **860** can proceed generally similarly or identically to block **552** of the method **550** of FIG. 5 to, for example, generate a point cloud representation of the depth of the scene **108**.

[0054] At block **863**, the method **860** includes registering the point cloud with medical scan data (e.g., patient data). In some embodiments, the medical scan can be a computerized tomography (CT) scan of the patient's spine that provides a complete 3D data set for at least a portion of the scene **108**. The registration process matches points in the point cloud to corresponding 3D points in the medical scan. The system **100** can register the point cloud to the medical scan data by detecting positions of fiducial markers and/or feature points visible in both data sets. For example, where the volumetric data comprises CT data, rigid bodies of bone surface calculated from the CT data can be registered to the corresponding points/surfaces of the point cloud. In other embodiments, the system **100** can employ other registration processes based on other methods of shape correspondence, and/or registration processes that do not rely on fiducial markers (e.g., markerless registration processes). In some embodiments, the registration/alignment process can include features that are generally similar or identical to the registration/alignment processes disclosed in U.S. Provisional Patent Application No. 62/796,065, titled "ALIGNING PRE-OPERATIVE SCAN IMAGES TO REAL-TIME OPERATIVE IMAGES FOR A MEDIATED-REALITY VIEW OF A SURGICAL SITE," filed Jan. 23, 2019, which is incorporated herein by reference in its entirety, and which is attached hereto as Appendix A.

[0055] At block **864**, the method **860** includes identifying missing/invalid regions of the point cloud. In some embodiments, block **864** can proceed generally similarly or identically to block **554** of the method **550** of FIG. 5. In one aspect of the present technology, the medical scan data includes 3D depth data corresponding to the missing or invalid regions of the point cloud. For example, FIG. 9A is a schematic view of a point cloud **970** corresponding to a portion of a patient's spine and including missing regions **972**. FIG. 9B is a schematic view of corresponding CT scan data **974** of the patient's spine in accordance with embodiments of the present technology. Referring to FIGS. 9A and 9B together, the CT scan data **974** can include 3D volumetric depth data **976** corresponding to at least a portion of the missing regions **972** in the point cloud **970**.

[0056] At block **865**, the method **860** includes merging/fusing the 3D data of the medical scan with the original depth data (e.g., captured at block **861**) to generate a merged point cloud that includes data points for the missing or invalid regions. In general, the medical scan data can replace and/or supplement the data in the point cloud. For example, the medical scan data can replace data in regions of the point cloud where captured data is poor, and supplement (e.g., fill in) the missing regions of the point cloud. Accordingly, the merged point cloud can provide a more accurate and robust depth map of the scene **108** that facilitates better reconstruction and synthesis of an output image of the scene **108** rendered from any desired virtual perspective, as described in detail above. [0057] More specifically, in some embodiments data from the medical scan is filled-in only for the missing regions of the point cloud. FIG. **9C**, for example, is a schematic view of a merged point cloud **980** in which the CT scan data **974** shown in FIG. **9B** has been filled in for the missing regions **972** of the point cloud **970** shown in FIG. **9A**. In some embodiments, the appropriate regions of the CT scan data **974** corresponding to the missing regions **972** of the point cloud **970** can be found by comparing nearest neighbors between the registered CT scan data **974** and the point cloud **970**. That is, for example, points in the medical scan that have no neighbor (e.g., are below a threshold) in the registered point cloud can be identified for merger into/with the point cloud data. In other embodiments, as much of the original depth data (e.g., the point cloud **970**) as possible can be replaced with the registered medical scan data (e.g., the CT scan data **974**). In some embodiments, a nearest neighbors algorithm can be used to determine which regions of the original depth data to remove and replace. In yet other embodiments, the medical scan data and the point cloud can be directly merged with a volumetric (e.g., voxel) representation, such as a truncated signed distance function (TSDF).

[0058] At block **866**, the method **860** can optionally include generating a three-dimensional mesh based on the merged point cloud. The 3D mesh can be used to reconstruct/synthesize the output image of the scene **108**. In some embodiments, the method **860** can return to block **851** to update the depth information of the scene **108**. In some embodiments, when the medical scan data and original depth data are directly merged using a TSDF, the 3D mesh can be generated using a marching cubes or other suitable algorithm.

[0059] In some embodiments, the medical scan data is known a priori and thus does not require significant processing. Accordingly, in one aspect of the present technology the method **860** can quickly update (e.g., supplement and/or replace) the original depth based on the medical scan—allowing real time or near real time processing and generation of an output image of the scene **108**.

[0060] In other embodiments, the medical scan data can act as an initial state for a depth optimization process, where further refinement is possible. For example, the medical scan data can be registered to the live data to fill-in holes as described in detail with reference to FIGS. **8-9C**. However, in some embodiments the cameras **112** can have a higher resolution/accuracy than the medical scan data. Accordingly, the merged depth information from the depth sensor **116** and the medical scan data can be used to initialize a 3D reconstruction process using images from the cameras **112**. The depth information from the images can then be merged with or replace the depth information from the medical scan. In some embodiments, the depth processing of the image data is accelerated because a depth range for the missing/invalid regions is known based on the medical scan data—thus minimizing the range of iterations needed to determine true depth values from the image data.

### III. Further Examples

[0061] The following examples are illustrative of several embodiments of the present technology:

[0062] 1. A method of determining the depth of a scene, the method comprising: [0063] capturing depth data of the scene with a depth sensor; [0064] capturing image data of the scene with a plurality of cameras; [0065] generating a point cloud representative of the scene based on the depth data; [0066] identifying a region of the point cloud; [0067] generating depth data for the region based on the image data; and [0068] merging the depth data for the region with the depth data from

the depth sensor to generate a merged point cloud representative of the scene. [0069] 2. The method of example 1 wherein the region of the point cloud is a missing region of the point cloud in which the point cloud includes no data or sparse data. [0070] 3. The method of example 2 wherein identifying the missing region of the point cloud includes determining that the missing region of the point cloud has fewer than a predetermined threshold number of data points. [0071] 4. The method of example 2 or example 3 wherein identifying the missing region of the point cloud includes identifying a hole in the point cloud that is larger than a user-defined threshold. [0072] 5. The method of any one of examples 2-4 wherein generating the depth data for the missing region is further based on a portion of the depth data captured by the depth sensor that surrounds the missing region. [0073] 6. The method of any one of examples 1-5 wherein the depth data for the region has a greater resolution than the depth data captured with the depth sensor. [0074] 7. The method of any one of examples 1-6 wherein the method further comprises generating a three-dimensional mesh representative of the scene based on the merged point cloud. [0075] 8. The method of any one of examples 1-7 wherein the scene is a surgical scene. [0076] 9. The method of any one of examples 1-8 wherein the plurality of cameras each have a different position and orientation relative to the scene, and wherein the image data is light field image data. [0077] 10. The method of any one of examples 1-9 wherein the method further comprises: [0078] processing the image data and the merged point cloud to synthesize an output image of the scene corresponding to a virtual camera perspective; and [0079] transmitting the output image to the display for display to a user. [0080] 11. The method of example 10 wherein the display is a head-mounted display worn by the user, and wherein identifying the region of the scene is based on at least one of a position and an orientation of the head-mounted display. [0081] 12. A system for imaging a scene, comprising: [0082] multiple cameras arranged at different positions and orientations relative to the scene and configured to capture image data of the scene; [0083] a depth sensor configured to capture depth data about a depth of the scene; and [0084] a computing device communicatively coupled to the cameras and the depth sensor, wherein the computing device has a memory containing computer-executable instructions and a processor for executing the computer-executable instructions contained in the memory, and wherein the computer-executable instructions include instructions for — [0085] receiving the image data from the cameras; [0086] receiving the depth data from the depth sensor; [0087] generating a point cloud representative of the scene based on the depth data; [0088] identifying a region of the point cloud; [0089] generating depth data for the region based on the image data; and [0090] merging the depth data for the region with the depth data from the depth sensor to generate a merged point cloud representative of the scene. [0091] 13. The system of example 12 wherein the region of the point cloud is a missing region of the point cloud in which the point cloud includes no data or sparse data. [0092] 14. The system of example 12 or example 13 wherein the region of the point cloud is user-selected. [0093] 15. The system of any one of examples 12-14, further comprising a display, wherein the computing device is communicatively coupled to the display, and wherein the computer-executable instructions further include instructions for— [0094] processing the image data and the merged point cloud to synthesize an output image of the scene corresponding to a virtual camera perspective; and [0095] transmitting the output image to the display for display to a user. [0096] 16. The system of example 15 wherein identifying the region of the scene is based on at least one of a position and an orientation of the display. [0097] 17. A method of determining the depth of a scene, the method comprising: [0098] capturing depth data of the scene with a depth sensor; [0099] generating a point cloud representative of the scene based on the depth data; [0100] identifying a region of the point cloud; [0101] registering the point cloud with three-dimensional (3D) medical scan data; and [0102] merging at least a portion of the 3D medical scan data with the depth data from the depth sensor to generate a merged point cloud representative of the scene. [0103] 18. The method of example 17 wherein the region of the point cloud is a missing region of the point cloud in which the point cloud includes no data or sparse data. [0104] 19. The method of example 18 wherein the scene is a

medical scene including a portion of a patient, wherein the missing region of the point cloud corresponds to the portion of the patient, and wherein the portion of the 3D medical scan data portion corresponds to the same portion of the patient. [0105] 20. The method of any one of examples 17-19 wherein the 3D medical scan data is a computed tomography (CT) data.

#### IV. CONCLUSION

[0106] The above detailed description of embodiments of the technology are not intended to be exhaustive or to limit the technology to the precise form disclosed above. Although specific embodiments of, and examples for, the technology are described above for illustrative purposes, various equivalent modifications are possible within the scope of the technology as those skilled in the relevant art will recognize. For example, although steps are presented in a given order, alternative embodiments can perform steps in a different order. The various embodiments described herein can also be combined to provide further embodiments.

[0107] From the foregoing, it will be appreciated that specific embodiments of the technology have been described herein for purposes of illustration, but well-known structures and functions have not been shown or described in detail to avoid unnecessarily obscuring the description of the embodiments of the technology. Where the context permits, singular or plural terms can also include the plural or singular term, respectively.

[0108] Moreover, unless the word “or” is expressly limited to mean only a single item exclusive from the other items in reference to a list of two or more items, then the use of “or” in such a list is to be interpreted as including (a) any single item in the list, (b) all of the items in the list, or (c) any combination of the items in the list. Additionally, the term “comprising” is used throughout to mean including at least the recited feature(s) such that any greater number of the same feature and/or additional types of other features are not precluded. It will also be appreciated that specific embodiments have been described herein for purposes of illustration, but that various modifications can be made without deviating from the technology. Further, while advantages associated with some embodiments of the technology have been described in the context of those embodiments, other embodiments can also exhibit such advantages, and not all embodiments need necessarily exhibit such advantages to fall within the scope of the technology. Accordingly, the disclosure and associated technology can encompass other embodiments not expressly shown or described herein.

## Claims

**1-20.** (canceled)

**21.** A method of determining depth within a surgical scene, the method comprising: capturing depth data of anatomy of a patient within the surgical scene undergoing a surgical procedure with a depth sensor; capturing image data of the surgical scene with a plurality of cameras different than the depth sensor; receiving an input selecting a region of the anatomy for which higher resolution depth information is desired; processing the image data to generate additional depth data for the region of the anatomy for which higher resolution depth information is desired; and merging the depth data of the anatomy from the depth sensor with the additional depth data to generate a merged depth map of the anatomy.

**22.** The method of claim 21 wherein the image data is light field image data.

**23.** The method of claim 21 wherein the merged depth map is a point cloud.

**24.** The method of claim 21 wherein the method further comprises: generating a point cloud representative of the anatomy based on the depth data from the depth sensor; determining at least one depth value from the point cloud for an area adjacent to the region of the anatomy for which higher resolution depth information is desired; and processing the image data with a depth processing algorithm to generate the additional depth data, wherein processing the image data with the depth processing algorithm includes limiting the depth processing algorithm to determine the additional depth data in a depth range surrounding the at least one depth value.

25. The method of claim 21 wherein the input comprises a user selection of the region of the anatomy for which higher resolution depth information is desired.
26. The method of claim 21 wherein the input comprises tracking data of one or more instruments within the surgical scene.
27. The method of claim 21 wherein the method further comprises registering the merged depth map with medical scan data of the anatomy.
28. The method of claim 27 wherein the medical scan data comprises computed tomography (CT) scan data.
29. The method of claim 21 wherein the method further comprises: generating a point cloud representative of the anatomy based on the depth data from the depth sensor; and registering the point cloud with medical scan data of the anatomy.
30. The method of claim 29 wherein processing the image data comprises processing the image data with a depth processing algorithm to generate the additional depth data, and wherein processing the image data with the depth processing algorithm includes limiting the depth processing algorithm to determine the additional depth data in a depth range surrounding the registered medical scan data in the region for which higher resolution depth information is desired.
31. The method of claim 30 wherein the medical scan data comprises computed tomography (CT) scan data.
32. The method of claim 21 wherein the additional depth data has a higher resolution than the depth data from the depth sensor.
33. The method of claim 21 wherein the cameras and the depth sensor are rigidly mounted to a common frame and fixed in position relative to one another.
34. The method of claim 33 wherein the cameras are RGB cameras.
35. A system for imaging a surgical scene, comprising: multiple cameras arranged at different positions and orientations relative to the surgical scene and configured to capture image data of anatomy of a patient within the surgical scene undergoing a surgical procedure; a depth sensor configured to capture depth data of the anatomy; and a computing device communicatively coupled to the cameras and the depth sensor, wherein the computing device has a memory containing computer-executable instructions and a processor for executing the computer-executable instructions contained in the memory, and wherein the computer-executable instructions, when executed by the processor, cause the processor to: receive the depth data of the anatomy from the depth sensor; receive the image data of the anatomy from the cameras; receive an input selecting a region of the anatomy for which higher resolution depth information is desired; process the image data to generate additional depth data for the region of the anatomy for which higher resolution depth information is desired; and merge the depth data of the anatomy from the depth sensor with the additional depth data to generate a merged depth map of the anatomy.
36. The system of claim 35 wherein the cameras and the depth sensor are rigidly mounted to a common frame and fixed in position relative to one another.
37. The system of claim 36 wherein the cameras are RGB cameras.
38. The system of claim 35 wherein the input comprises a user selection of the region of the anatomy for which higher resolution depth information is desired.
39. The system of claim 35 wherein the input comprises tracking data of one or more instruments within the surgical scene.
40. The system of claim 35 wherein the computer-executable instructions, when executed by the processor, further cause the processor to: generate a point cloud representative of the anatomy based on the depth data from the depth sensor; determine at least one depth value from the point cloud for an area adjacent to the region of the anatomy for which higher resolution depth information is desired; process the image data with a depth processing algorithm to generate the additional depth data; and limit the depth processing algorithm to determine the additional depth data in a depth range surrounding the at least one depth value.

