

# US Patent & Trademark Office

## Patent Public Search | Text View

United States Patent Application Publication

20250254213

Kind Code

A1

Publication Date

August 07, 2025

Inventor(s)

Pinczel; Balázs et al.

### CHECKING STATE CONSISTENCY IN A LOG REPLICATION CONFIGURATION

#### Abstract

According to an aspect there is provided a computer-implemented method of operating a leader server in a leader server-follower server log replication configuration. The method comprises (i) storing **(801)** a current master state of a system and maintaining a master log relating to the current master state, wherein the master log comprises a plurality of log entries each having a respective log index, wherein each log entry indicates a respective change to the master state of the system; (ii) performing **(803)** a log compaction procedure to remove from the master log any log entry for which the respective change to the master state of the system has been overridden by a change to the master state in a subsequent log entry; (iii) storing **(805)** a compaction index that is the highest log index of the log entries removed in the log compaction procedure; (iv) storing **(807)** an override index that is the highest log index of the log entries for which the respective change to the master state overrode a change to the master state by an earlier log entry; and (v) sending **(809)** the compaction index and the override index to one or more follower servers that are to store a replica of the master log relating to the current master state.

**Inventors:** Pinczel; Balázs (Budapest, HU), Géhberger; Dániel (Montreal, CA), Mátray; Péter (Budapest, HU), Németh; Gábor (Budapest, HU), Nagy; Mate Imre (Budapest, HU)

**Applicant:** Telefonaktiebolaget LM Ericsson (PUBL) (Stockholm, SE)

**Family ID:** 73040091

**Appl. No.:** 19/188785

**Filed:** April 24, 2025

#### Related U.S. Application Data

parent US division 18033998 20230426 parent-grant-document US 12289368 WO division PCT/EP2020/080431 20201029 child US 19188785

## Publication Classification

Int. Cl.: H04L67/1095 (20220101); G06F16/23 (20190101); G06F16/27 (20190101)

U.S. Cl.:

CPC H04L67/1095 (20130101); G06F16/2358 (20190101); G06F16/27 (20190101);

---

## Background/Summary

### TECHNICAL FIELD

[0001] This disclosure relates to techniques for use with a log replication configuration in which a leader server maintains a log relating to a current state and sends updates to one or more follower servers that maintain a replica log.

### BACKGROUND

[0002] In recent years, the telecommunications industry has gradually shifted towards applying cloud-based solutions. With the advent of 5.sup.th Generation (5G) networks, and the early visions of 6.sup.th Generation (6G) networks, this trend is set to continue, and the areas in which such solutions can efficiently be utilised is increasing. As the functions of the network—which were previously implemented as individual physical nodes—are converted to services running in the cloud, new challenges arise in the reliable storage and retrieval of their state, which is data that is required for their (and other nodes) correct operation. To fully benefit from the use of cloud technologies, these services should be small, mostly independent from each other, and most importantly, starting, stopping, or relocating service instances should not disrupt their operation.

[0003] A common solution for managing state data for these services while meeting these requirements is to utilise state externalisation. Here, a separate utility is responsible for holding (storing) and managing the state, so that the services themselves can be implemented as stateless ‘workers’ which turn to (i.e. consult) this utility when they have to read or update their state. Such a utility—while logically presenting itself as centralised—is usually also built from multiple instances, which each hold a replica of the data. Since read and update requests come in to each of these replicas independently, some form of consensus protocol is required to keep the replicas consistent with each other. Consistency is especially important for telecommunications systems (and control plane functions in particular), since they rely heavily on distributed transactions.

[0004] To maintain consensus on an evolving state, a family of solutions uses a replicated log (also known as a ‘distributed log’). A log is a series of log entries, each describing a specific change to the state. For example if the state is a complete database, a log entry may describe a single update to one data item in the database.

[0005] FIG. 1 shows an exemplary log 12, having three log entries 14, 16, 18 for a key-value database. Each log entry 14, 16, 18 is labelled with a respective log index i. The state of the database after each log entry 14 has been applied to the database is shown above the log entries 14. Thus, the first log entry 14 (with log index ‘1’) indicates that a value “x” was added to field A in the database (note that a field in a key-value database is also known as a “key”). After the first log entry 14, the database has a state 20 in which there is a value “x” in field A. The second log entry 16 (with log index ‘2’) indicates that a value “y” was added to data field B in the database. After the second log entry 16, the database has a state 22 in which there is a value “x” in field A and a value “y” in field B. The third log entry 18 (with log index ‘3’) adds a value “z” to data field A in the database. As data field A already has a value “x”, the value “z” overwrites the value “x” in data field A. After the third log entry 18, the database has a state 24 in which there is a value “z” in field

A and value “y” in field B.

[0006] There can be multiple servers in the system, and each server maintains a replica of this log. Each server also maintains a replica of the state, which is initially identical on all servers (usually set to empty) and is evolved by performing the updates described by the log entries.

[0007] Consensus protocols used with replicated logs aim to keep the logs identical, even as new entries are added to them. This ensures that the state replica goes through the same history on every server. A common way to synchronise the evolving logs between the servers is to use a leader/follower configuration, where one server has the ‘leader’ role, which means that it is the only server allowed to add new entries to the log by itself. The rest of the servers are ‘followers’, and they can only add entries to their log when the leader informs them of a new log entry. Requests to update the state can originate from any server, but these requests have to be forwarded to the leader. When the leader receives a request, it creates a log entry describing the change, and adds it to its replica of the log. Periodically, the leader updates the followers by sending them any newly added log entries. Clients of the system can issue read requests to any of the servers to read the state, and their queries can be answered by any server (follower or leader) based on the local state replica maintained at each server. A popular consensus protocol of this type is known as Raft, as described in “In search of an understandable consensus algorithm” by D. Ongaro & J. Ousterhout, 2014 USENIX ATC 14, pp. 305-319.

[0008] Most consensus protocols also have features which enable the transfer of the leader role on-the-fly, without losing consistency of the replicated state and with no (or only minimal) loss of availability, which means that update requests can continue to arrive, even during this transition. This is of course important for practical applications, and most of the challenge in designing consensus protocols comes from this feature.

[0009] As noted above, the combination of the distributed log and the leader/follower pattern can achieve consensus on a constantly evolving state, resulting in the same history on all servers. This means that each replica of the state will go through the same values, in the same order. This, however, doesn't necessarily mean that they progress through this state history at the same pace. For example, if the communication between the leader and a follower is temporarily degraded (e.g. due to network congestion) or is unavailable, the follower's state may lag behind that of the others. Once the communication problem is resolved, the leader has to replay (resend) the most recent part of the log (i.e. the most recent log entries) to this follower. In cases the entire log has to be replayed (resent). For example, when a new server is added to the group of follower servers, it has to start from the initial (empty) state and go through the entire history of state changes to become up to date.

[0010] For applications where the state goes through frequent updates, the size of the log can easily become a problem. Not only does it occupy space (memory and/or persistent storage) on the servers, but it also increases the time it takes to replay some (or all) of the log, making it harder for lagging or newly joined followers to catch up.

[0011] As a solution, various ‘log compaction’ schemes are available. The common idea in log compaction schemes is to discard or alter some of the old history in the log, yielding a log which is smaller, but still results in the same replicated state if replayed all the way up to the recent entries.

[0012] Most log compaction schemes are variants of either a so-called ‘snapshotting’, or ‘log cleaning’ approach.

[0013] In the snapshotting approach, the servers create a read-only copy of the replicated state, discard the log entries up to the point where this snapshot was made, and replace these log entries with a single “big” log entry which describes or corresponds to the snapshot. This snapshot can be created, for example, periodically, or whenever the log grows over some threshold size (in terms of memory or storage requirements). After the first time a snapshot is made, the log always has the following structure: it begins with a snapshot entry describing a relatively recent state, and then it continues with normal log entries describing updates that happened after (or during) the creation of

the snapshot. Newly-joined servers first have to receive and install the snapshot, which could take a while since it usually contains much more data than normal log entries, but in doing so the server has immediately caught up to a recent state.

[0014] While the snapshotting approach removes the entire early portion of the log and replaces it with equivalent but more compact data in the form of the snapshot, the log cleaning approach aims to be more granular. In the log cleaning approach, only individual log entries can be removed, and only if they no longer contribute to the current state. In that case there is no need to replace these log entries with anything. The exact interpretation of “an entry no longer contributing to the current state” depends on the application, but it usually means that the log already contains a newer (later) entry which overrides the effect or data of the older log entry.

[0015] FIG. 2 shows the effect of a log cleaning approach on the example of FIG. 1. Here, as the effect of the first log entry **14** has been overridden by a later log entry (i.e. the log entry with log index ‘3’ replaces the value of data field A that was set by the first log entry **14**), it is fine if a newly-joined server only receives the second and third log entries **16**, **18**, as the server would still arrive at the correct state (A=“z” and B=“y”). Thus, the first log entry **14** can be removed from the log.

[0016] Other log compaction schemes are available that build on the same underlying principles as snapshotting and log cleaning, with these schemes differing in the specifics of, e.g., the data structure used for a snapshot, how/when snapshotting or log cleaning is scheduled, and how they solve the challenge of serving incoming updates during these operations.

[0017] However, the snapshotting approach to log compaction is not lightweight enough to be feasible for the use case of low-latency state externalisation. Creating a snapshot of the entire state is a resource-intensive task and can take a long time if the state is large, during which the normal operation of the consensus protocol (and the application it is embedded in) should also be maintained without suffering any adverse effects. This requires either permanently over-dimensioning the resources for the worst-case scenario of these heavyweight snapshotting events or using a cloud environment with auto-scaling of the virtualised resources to match the periodically increased load. Experience shows that auto-scaling solutions are not necessarily good at adapting to a workload like this, or not at the required time-scale, but even if they improve in the future, it is good practice to avoid relying on them if there are other solutions.

[0018] In contrast, if the principles of log cleaning are used, the load of compaction can be spread over time to be more uniform, since the obsolete log entries can be removed as part of normal operation (e.g. immediately after an overriding log entry becomes part of the log), eliminating the need for a periodic heavyweight cleaning task. However, existing log cleaning solutions fail to provide full consistency when replaying from a compacted log. A lagging or newly-joined follower sees only the older parts of the log, where some obsolete log entries may have already been removed, but the corresponding overriding log entries from the recent parts of the log may also be missing until the replay progresses far enough.

[0019] For example, consider the key-value database example in FIG. 2. A log cleaning solution may remove the first log entry **14** (the log entry with log index ‘1’) as soon as the third log entry **18** (the log entry with log index ‘3’) becomes part of the log. If a new follower joined at this point, it would only ever receive the second and third log entries **16**, **18**, as shown in FIG. 3. This means that a newly-joined follower, after updating the state according to the second log entry **16**, would have a state **26** in which there is a value “y” in field B, and no value for field A. However, this state never existed for the leader, since after updating the state according to the second log entry **16** the state had a value “x” in field A and a value “y” in field B. Only once the new follower has received the third log entry **18** and updated the state accordingly does the new follower have a state that is consistent with that of the leader.

[0020] While the leader can send log entries in batches rather than one at a time, an earlier batch of log entries may not include all the log entries that overrode other log entries (e.g. the first batch of

log entries may only include the second log entry **16** in the example of FIG. 3), and the state replica maintained at follower will be inconsistent until the required log entries have been received and acted upon.

[0021] Therefore, with a log cleaning approach, the requirement for consistency in the replicated logs and states is violated temporarily.

[0022] Thus, neither snapshotting nor log cleaning variants satisfy both the requirements of lightweight and consistent operation, which is needed by the low-latency state externalisation use-case.

## SUMMARY

[0023] Therefore there is a need for improved techniques for a configuration where a leader server performs log compaction and maintains a compacted log relating to a current state and sends updates to one or more follower servers that maintain a replica of the log and current state. In particular there is a need for improvements in maintaining consistency (or identifying inconsistencies) in replicated log-based consensus protocols that use log compaction.

[0024] According to a first aspect, there is provided a computer-implemented method of operating a leader server in a leader server-follower server log replication configuration. The method comprises (i) storing a current master state of a system and maintaining a master log relating to the current master state, wherein the master log comprises a plurality of log entries each having a respective log index, wherein each log entry indicates a respective change to the master state of the system; (ii) performing a log compaction procedure to remove from the master log any log entry for which the respective change to the master state of the system has been overridden by a change to the master state in a subsequent log entry; (iii) storing a compaction index that is the highest log index of the log entries removed in the log compaction procedure; (iv) storing an override index that is the highest log index of the log entries for which the respective change to the master state overrode a change to the master state by an earlier log entry; and (v) sending the compaction index and the override index to one or more follower servers that are to store a replica of the master log relating to the current master state.

[0025] According to a second aspect, there is provided a computer-implemented method of operating a follower server in a leader server-follower server log replication configuration. The method comprises (i) storing a replica state of a system and maintaining a replica log relating to the replica state, wherein the replica log comprises a plurality of log entries each having a respective log index, wherein each log entry indicates a respective change to the replica state of the system; (ii) receiving a compaction index and an override index from a leader server that stores a master state of the system and maintains a master log relating to the current master state, wherein the compaction index and the override index relate to a compacted master log resulting from a log compaction procedure at the leader server, wherein the compaction index is the highest log index of any log entries removed in the log compaction procedure, and wherein the override index is the highest log index of any log entries for which the respective change to the master state overrode a change to the master state of an earlier log entry; and (iii) determining if the stored replica state is consistent with a current or previous master state maintained by the leader server based on a highest log index in the stored replica log and one or both of the received compaction index and the received override index.

[0026] According to a third aspect, there is provided a computer program product comprising a computer readable medium having computer readable code embodied therein, the computer readable code being configured such that, on execution by a suitable computer or processor, the computer or processor is caused to perform the method according to the first aspect, the second aspect, or any embodiments thereof.

[0027] According to a fourth aspect, there is provided an apparatus for use as a leader server in a leader server-follower server log replication configuration. The apparatus is configured to (i) store a current master state of a system and maintaining a master log relating to the current master state,

wherein the master log comprises a plurality of log entries each having a respective log index, wherein each log entry indicates a respective change to the master state of the system; (ii) perform a log compaction procedure to remove from the master log any log entry for which the respective change to the master state of the system has been overridden by a change to the master state in a subsequent log entry; (iii) store a compaction index that is the highest log index of the log entries removed in the log compaction procedure; (iv) store an override index that is the highest log index of the log entries for which the respective change to the master state overrode a change to the master state by an earlier log entry; and (v) send the compaction index and the override index to one or more follower servers that are to store a replica of the master log relating to the current master state.

[0028] According to a fifth aspect, there is provided an apparatus for use as a follower server in a leader server-follower server log replication configuration. The apparatus is configured to: (i) store a replica state of a system and maintaining a replica log relating to the replica state, wherein the replica log comprises a plurality of log entries each having a respective log index, wherein each log entry indicates a respective change to the replica state of the system; (ii) receive a compaction index and an override index from a leader server that stores a master state of the system and maintains a master log relating to the current master state, wherein the compaction index and the override index relate to a compacted master log resulting from a log compaction procedure at the leader server, wherein the compaction index is the highest log index of any log entries removed in the log compaction procedure, and wherein the override index is the highest log index of any log entries for which the respective change to the master state overrode a change to the master state of an earlier log entry; and (iii) determine if the stored replica state is consistent with a current or previous master state maintained by the leader server based on a highest log index in the stored replica log and one or both of the received compaction index and the received override index.

[0029] According to a sixth aspect, there is provided an apparatus for use as a leader server in a leader server-follower server log replication configuration. The apparatus comprises a processor and a memory, said memory containing instructions executable by said processor whereby said apparatus is operative to (i) store a current master state of a system and maintaining a master log relating to the current master state, wherein the master log comprises a plurality of log entries each having a respective log index, wherein each log entry indicates a respective change to the master state of the system; (ii) perform a log compaction procedure to remove from the master log any log entry for which the respective change to the master state of the system has been overridden by a change to the master state in a subsequent log entry; (iii) store a compaction index that is the highest log index of the log entries removed in the log compaction procedure; (iv) store an override index that is the highest log index of the log entries for which the respective change to the master state overrode a change to the master state by an earlier log entry; and (v) send the compaction index and the override index to one or more follower servers that are to store a replica of the master log relating to the current master state.

[0030] According to a seventh aspect, there is provided an apparatus for use as a follower server in a leader server-follower server log replication configuration. The apparatus comprises a processor and a memory, said memory containing instructions executable by said processor whereby said apparatus is operative to (i) store a replica state of a system and maintaining a replica log relating to the replica state, wherein the replica log comprises a plurality of log entries each having a respective log index, wherein each log entry indicates a respective change to the replica state of the system; (ii) receive a compaction index and an override index from a leader server that stores a master state of the system and maintains a master log relating to the current master state, wherein the compaction index and the override index relate to a compacted master log resulting from a log compaction procedure at the leader server, wherein the compaction index is the highest log index of any log entries removed in the log compaction procedure, and wherein the override index is the highest log index of any log entries for which the respective change to the master state overrode a

change to the master state of an earlier log entry; and (iii) determine if the stored replica state is consistent with a current or previous master state maintained by the leader server based on a highest log index in the stored replica log and one or both of the received compaction index and the received override index.

[0031] The above techniques enable follower servers to be aware of the consistency of their replica of the state while a leader server is replaying a compacted log to the follower server. Based on this awareness, it is possible to run servers in a fully consistent mode even when using log compaction, at least from the point of view of clients that can request information about the state. The proposed techniques satisfy both the consistency and lightweight operation requirements of a low-latency state externalisation use-case.

---

## Description

### BRIEF DESCRIPTION OF THE DRAWINGS

[0032] Some of the embodiments contemplated herein will now be described more fully with reference to the accompanying drawings, in which:

[0033] FIG. 1 shows an exemplary log having a number of log entries and the corresponding state of the database after the log entries have been applied;

[0034] FIG. 2 is an illustration of the application of a log cleaning approach to log compaction for the exemplary log of FIG. 1;

[0035] FIG. 3 is an illustration of problems with replaying a compacted log to a follower;

[0036] FIG. 4 is an illustration of a set of servers in which the techniques described herein can be implemented and one or more client devices;

[0037] FIG. 5 is an illustration of a consistent status and inconsistent status for a follower server;

[0038] FIG. 6 is a diagram illustrating exemplary signalling between a leader server, a follower server and a client device;

[0039] FIG. 7 is a diagram illustrating further exemplary signalling between a leader server, a follower server and a client device;

[0040] FIG. 8 is a flow chart illustrating a method of operating a leader server according to various embodiments;

[0041] FIG. 9 is a flow chart illustrating a method of operating a follower server according to various embodiments; and

[0042] FIG. 10 is a simplified block diagram of an apparatus that can be used to implement a leader server and/or a follower server according to various embodiments.

### DETAILED DESCRIPTION

[0043] Some of the embodiments contemplated herein will now be described more fully with reference to the accompanying drawings. Other embodiments, however, are contained within the scope of the subject matter disclosed herein, the disclosed subject matter should not be construed as limited to only the embodiments set forth herein; rather, these embodiments are provided by way of example to convey the scope of the subject matter to those skilled in the art.

[0044] As noted above, in consensus protocols which replicate a log of state-changing events, a lightweight way of keeping the log compact is to remove obsolete log entries after a new overriding log entry becomes part of the log. That is, when a new log entry overrides a change to the state caused by an earlier log entry, that earlier log entry is now obsolete (as the state change the earlier log entry related to is no longer part of the current state). Replayng a compacted log eventually results in the same state being stored in the follower server as if the log was not compacted.

However, when a replay hasn't progressed far enough, it may temporarily result in inconsistent states in the follower server which never existed and are just an artefact of the compaction.

[0045] In the proposed solution, the leader server for the consensus protocol can share additional

information with the follower servers, which allows them to recognise when they may have an inconsistent state while the log is being replayed. Followers that determine that their state is inconsistent can in turn decide to redirect client requests to another server until it reaches a consistent state again.

[0046] FIG. 4 is an illustration of a server-client arrangement in which the techniques described herein can be implemented. FIG. 4 shows a plurality of servers 50 that provide services for or towards one or more client devices 52. Each of the servers 50 maintains and stores a state for a system, and also stores a log that includes log entries relating to the evolution of the state. Thus, each log entry indicates a respective change to the state of the system. Each log entry has a respective log index  $i$  identifying the log entry. The log index also relates to the order in which the log entries were added to the log. For example, log index 1 is the first log entry, log index 2 is the second (next) log entry, log index 3 is the third (next) log entry, and so on. The system that the state relates to may be or include one or more of the client devices 52, or the state may relate to a separate system not shown in FIG. 4. In either case, the servers 50 are providing state externalisation for the system to which the state relates.

[0047] The set of servers 50 are configured to operate according to a leader/follower configuration in which one server has the leader role and the rest of the servers are followers. Thus, one of the servers 50 is designated as the leader server 54, and the remaining servers 56 are follower servers 58. The leader server 54 stores the state 60 for the system, and also stores a log 62 that comprises one or more log entries relating to the state 60. The state 60 stored by the leader server 54 is also referred to herein as the 'master state', and the log 62 stored by the leader server 54 is also referred to herein as the 'master log'. New log entries received by the leader server 54 are used to update the master state 60 and are stored in the master log 62. As the leader server 54, the leader server 54 is the only server allowed to add new log entries to the master log 62 by itself.

[0048] Each follower server 58 stores a replica state 64 for the system, and also stores a replica log 66 that comprises one or more log entries relating to the replica state 64. Each of the follower servers 58 are only able to add entries to their replica log 66 when the leader server 54 informs them of a new log entry, as shown by arrow 68. Each follower server 58 uses the log entries in its replica log 66 to update the stored replica state 64.

[0049] It should be noted that the replica state 64 and the replica log 66 stored by a particular follower server 58 are internally consistent, i.e. the replica state 64 at the follower server 58 implements all of the changes indicated by the replica log 66 at that follower server 58, but it is possible that the replica state 64 and the replica log 66 stored by a particular follower server 58 does not match the master state 60 and the master log 62 stored by the leader server 54. In particular, in an ideal situation, each of the follower servers 58 maintains an exact replica of the master state 64 and master log 66 (i.e. the replica state 64 matches the master state 60 at the leader server 54 and the replica log 66 matches the master log 62 at the leader server 54), which means that any of the follower servers 58 can accurately respond to requests from any of the client devices 52 about the current state. However, as noted above, there can be communication issues between the leader server 54 and one or more of the follower servers 58 which means that messages from the leader server 54 to the follower servers 58 about updates the state 60 and log 62 may not be received by that/those follower server(s) 58. As a result, the replica state 64 and the replica log 66 of a follower server 58 can fall behind the master state 60 and master log 62 of the leader server 54. In that case, the leader server 54 has to 'replay' (resend) part of the log 62 to the follower server 58 so the follower server 58 can 'catch up' and allow the replica state 64 and the replica log 66 to match the state 60 and log 62 stored in the leader server 54. Likewise, when a new follower server 58 is created or set up, that new follower server 58 needs to have the full log 62 resent or replayed by the leader server 54 so that it can build the correct replica state 64 and the correct replica log 66.

[0050] The leader server 54 implements a log compaction scheme in order to reduce the size of the log 62, and in the embodiments described herein, the leader server 54 uses a log cleaning approach



to remove obsolete log entries after a new overriding log entry becomes part of the log **62**. That is, when a new log entry overrides a change to the state **60** caused by an earlier log entry, that earlier log entry is now obsolete (as the state change the earlier log entry related to is no longer part of the current state **60**).

[0051] As noted above with reference to FIG. 3, in the case of a log cleaning approach, when the master log **62** is being replayed to a follower server **58**, the replica state **64** in the follower server **58** may temporarily be inconsistent with the previous states **60** in the leader server **54** until all of the overriding log entries have been received. It should be noted that the terms ‘consistent’ and ‘inconsistent’ as used herein with reference to the master state **60** and the replica state **64** refer to whether the stored replica state **64** matches the master state **60** at the time of the relevant log entry. Thus, the replica state **64** is considered to be inconsistent with the master state **60** when the replica state **64** at a log index a does not match the master state **60** at log index a. However, the replica state **64** is considered to be consistent with the master state **60** if the replica state **64** at a log index b does match the master state **60** at log index b, even if log index b is not the current (latest) log index. That is, the replica state **64** does not have to match the current master state **60** in order for the replica state **64** and the master state **60** to be considered consistent.

[0052] The techniques described herein apply to the log cleaning approach for log compaction and enable the follower servers **58** to determine when their replica state **66** is not consistent with the master state **60**. As in the existing log cleaning solutions, log entries in the log **62** are removed if they no longer contribute to the current state **60**. Specifically, a log entry at log index i should be removed if there is a log entry in the log **62** at some later index k (where  $k > i$ ) which overrides the effector data described by log entry i. In the example shown in FIG. 2, the log entry with index  $k=3$  overrides the effect of the log entry with index  $i=1$ , and so the log entry with index  $i=1$  is removed. Techniques for keeping track of which log entry overrides which other log entry(ies) can vary from application to application and are not described further herein. For most applications, identifying the earlier log entries that have been overridden can be inferred from the contents of the log entries themselves, e.g. as in the example shown in FIGS. 1-3.

[0053] The techniques provide for the leader server **54** to share additional information with the follower servers **58**, which allows the follower servers **58** to recognise when they may have an inconsistent state while the master log **62** is being replayed. Followers that determine that their replica state **64** is inconsistent can in turn decide to redirect client requests to another server (follower or leader) until its replica state **64** is consistent with the master state **60**.

[0054] In particular to allow a follower server **58** to detect inconsistency of its replica state **64** during log replays by the leader server **54**, the leader server **54** stores two pieces of information relating to the log cleaning process. The information comprises two indices relating to how far in the master log **62** the compaction (removal of entries) has progressed

[0055] The first index is referred to herein as a ‘compaction index’ **70**. The compaction index **70** indicates the highest log index where an overridden log entry was removed. That is, of the log entries removed from the master log **62** by the log cleaning process, the compaction index **70** is set to the highest log index of those removed log entries. In the example shown in FIG. 2, the compaction index **70** would be ‘1’, corresponding to the log entry **14** with log index 1, as this was the last (=highest log index) log entry that has been removed due to log cleaning.

[0056] The second index is referred to herein as an ‘override index’ **72**. The override index **72** indicates the highest log index of the log entries for which the respective change to the master state **60** overrode a change to the master state **60** by an earlier log entry. The override index **72** therefore keeps track of the log entry with the highest log index that overrode one or more earlier log entries and caused the removal of those one or more earlier log entries from the master log **62**. In the example shown in FIG. 2, the override index **72** would be ‘3’, corresponding to the log entry **18** with log index 3, as this was the last (=highest log index) log entry that caused the removal of an earlier log entry (the log entry with log index 1) due to log cleaning.

[0057] The compaction index **70** and the override index **72** are initially set to a default value, lower than the log index to be used for the first entry in the master log **62**. For example, the default value for the compaction index **70** and the override index **72** can be 0 in the case of a numbered log index system.

[0058] In some embodiments, the leader server **54** performs log cleaning each time that a new log entry is added to the master log **62**. In alternative embodiments, the leader server **54** can perform log cleaning periodically or intermittently (e.g. once a threshold number of log entries have been added to the master log **62** since the last log clean procedure). In the latter embodiments, the leader server **54** can perform the log cleaning sequentially, for example by evaluating log entries in the order in which they were added to the master log **62**. Alternatively, the leader server **54** can perform the log cleaning using a parallel approach (e.g. where the leader server **54** considers multiple log entries for log cleaning at the same time), or non-sequentially (e.g. where the leader server **54** does not evaluate log entries in the order in which they were added to the master log **62**).

[0059] Whenever a log entry with log index  $k$  is identified in the master log **62** that overrides one or more earlier log entries  $i_{\text{sub}.1}, \dots, i_{\text{sub}.n}$ , the compaction index **70** and override index **72** should be set as follows:

[00001]  $\text{Compactionindex} = \max(\text{currentcompactionindex}, i_1, \dots, i_n)$  (1)

$\text{Overrideindex} = \max(\text{currentoverrideindex}, k)$  (2)

[0060] Thus, according to (1), the new compaction index **70** is set to the higher of the existing value of the compaction index **70** and  $i_{\text{sub}.1}, \dots, i_{\text{sub}.n}$ . According to (2), the override index **72** is set to the higher of the existing value of the override index **72** and  $k$ . Where the leader server **54** performs log cleaning sequentially, a log entry  $k$  that overrides one or more earlier log entries will exceed the current override index **72** and thus the override index **72** is set to  $k$ . If the log entry with log index  $k$  does not override any changes caused by previous log entries, the compaction index **70** and the override index **72** remain unchanged from their previously-set values.

[0061] It will be appreciated that the compaction index **70** and the override index **72** can be set according to (1) and (2) regardless of whether the leader server **54** performs log cleaning for each received log entry or periodically or intermittently.

[0062] As in the existing leader server-follower server log replication configurations, the leader server **54** periodically or intermittently sends the new log entries to the follower servers **58**. This message is denoted 'AddLogEntries' herein. In the proposed solution, the leader server **54** also sends the compaction index **70** and the override index **72** to the follower server **58**. The compaction index **70** and the override index **72** can be sent in the AddLogEntries message, or in a separate message sent at the same or a similar time to the AddLogEntries message.

[0063] A follower server **58** receiving the compaction index **70** and the override index **72** from the leader server **54** can then decide whether its replica state **64** and the replica log **66** are consistent based on the received indices.

[0064] If the received compaction index **70** is greater than the highest index existing in the follower server's replica log **66** (before acting on any state changes in any received AddLogEntries message), then this indicates that a compacted part of the master log **62** is being replayed by the leader server **54**. Therefore, the replica state **64** built from this follower server's replica log **66** may experience the inconsistency issues described above with reference to FIG. 3. In applications that require strict consistency in the replica state **64**, the follower server **58** should not trust its replica state **64**, and the follower server **58** should therefore redirect or refuse to serve requests that require a consistent log and state. For example, for an application that uses key-value databases, a follower server **58** with an inconsistent replica state **64** and replica log **66** should not answer read requests.

[0065] After processing any state changes indicated in a received AddLogEntries message, if the received override index **72** is less than or equal to the highest index existing in the replica log **66** (including the newly added log entries), then the replica state **64** is consistent with the master state

**60**, and the follower server **58** can be sure of the consistency of its replica state **64**. In that case the follower server **58** can serve requests that require a consistent log and state. If the follower server **58** was previously redirecting or refusing read requests, it can now switch back to normal operation and respond to those requests.

[0066] Thus, based on the above, the follower server **58** can determine that its replica state **64** is inconsistent with the master state **60** if:

[00002] ‘inconsistentstatus’ if  $a < \text{compactionindex}$  (3)

where  $a$  is the highest log index of the log entries stored in the replica log **66** before acting on any new log entries received from the leader server **54** (and thus before adding those log entries to the replica log **66** and before updating the replica state **64** based on those log entries).

[0067] Also based on the above, the follower server **58** can determine that its replica state **64** is consistent with the master state **60** if:

[00003] ‘consistentstatus’ if  $b \geq \text{overrideindex}$  (4)

where  $b$  is the highest log index of the log entries stored in the replica log **66** after acting on any new log entries received from the leader server **54** (and thus after adding those new log entries to the replica log **66** and updating the replica state **64** accordingly).

[0068] Each follower server **58** maintains a consistency status with respect to its replica state **64**. The consistency status can be ‘consistent’ or ‘inconsistent’. The follower server **58** evaluates the received compaction index **70** and override index **72** to determine if the current consistency status should be maintained or changed according to criterion (3) or criterion (4) above.

[0069] The cycling between a consistent status and an inconsistent status for a follower server **58** is illustrated in FIG. 5. FIG. 5 shows a consistent status **501** for the follower server **58** where the replica state **64** is consistent with the master state **60**, and an inconsistent status **503** for the follower server **58** where the replica state **64** is inconsistent or possibly inconsistent with the master state **60**. The default state for a follower server **58** is the consistent status **501**.

[0070] A follower server **58** with a previously-determined consistent status **501** for the replica state **64** will maintain that consistent status **501** until criterion (3) is satisfied, i.e. until the highest log entry in the replica log **66** is determined to be below a received compaction index **70**. In that case the follower server **58** updates its status to the inconsistent status **503**. When a follower server **58** has a consistent status **501** for the replica state **64**, the override index **72** (and thus criterion (4)) does not need to be evaluated by the follower server **58**.

[0071] A follower server **58** with an inconsistent status **503** for the replica state **64** maintains that inconsistent status **503** until criterion (4) is satisfied, i.e. until the highest log entry in the replica log **66** is determined to be equal to or above a received override index **72**. In that case the follower server **58** updates its status to the consistent status **501**. When a follower server **58** has an inconsistent status **503** for the replica state **64**, the compaction index **70** (and thus criterion (3)) does not need to be evaluated by the follower server **58**.

[0072] The compaction index **70** and the override index **72** are thus used as log entry index thresholds for changing to an inconsistent status and consistent status respectively.

[0073] The techniques described herein can be applied to any type of state information about a system, where that state changes over time via updates to the state. In one particular example which relates to a key-value database, the state represents the total information known about the currently existing keys. For example, for each key the database stores metadata about the key and/or its current value. Each log entry can then represent a single update to a single key, e.g. a log entry can signal that a key is initialised (or updated) to a particular value. When applying the log compaction procedure described herein, an overriding log entry is a newer log entry relating to the setting of a key value that overrides older log entries about setting a key value for the same key.

[0074] FIG. 6 illustrates exemplary signalling between a leader server **54**, a follower server **58** and a client device **52**. The signalling in FIG. 6 relates to the example provided in FIGS. 1 and 2.

Initially, the master log **62** in the leader server **54** has two log entries (with log indices **1** and **2**), and the follower server **58** has an empty replica state **64** and empty replica log **66** (and thus the highest log index stored in the replica log **66** is 0). The status of the follower server **58** is initially set to the 'consistent status'.

[0075] The leader server **54** sends an AddLogEntries message **601** to the follower server **58** that includes the first log entry **14** and the second log entry **16** from FIG. **1**. The leader server **54** also sends the compaction index **70** and override index **72** to the follower server **58**, for example in the AddLogEntries message **601**, or in a separate message that is sent at or around the same time as the AddLogEntries message **601**. As no log cleaning has yet taken place, the compaction index **70** has a value of 0 and the override index **72** has a value of 0.

[0076] As the follower server **58** has a consistent status, it evaluates criterion (3) and checks the received compaction index **70** against its highest log index in the replica log **66**. The highest log index 0 matches the received compaction index 0 and so the follower server **58** maintains the consistent status for its replica state **64**. The follower server **58** then adds the received log entries to the replica log **66** and updates the replica state **64** based on the received log entries.

[0077] The follower server **58** sends an acknowledgement **603** to the leader server **54** to acknowledge receipt of the AddLogEntries message **601**.

[0078] Subsequently, the client device **52** sends a query **605** for information about the state. This query **605** is received by the follower server **58**. As the follower server **58** has a consistent status, the follower server **58** responds to the query **605** with a response message **607** containing the requested information.

[0079] At step **609**, the master state **60** is updated and a corresponding third log entry **18** is added to the master log **62**. The leader server **54** is able to apply the log cleaning approach to the updated master log **62** with the result that the first log entry **14** (with log index 1) is removed from the master log **62**. This log cleaning sets the value of the compaction index **70** to 1 and sets the value of the override index **72** to 3 (since the log entry with log index 3 overwrote the log entry with log index 1).

[0080] The leader server **54** sends another AddLogEntries message **611** to the follower server **58** that includes the third log entry **18** from FIG. **1**. The leader server **54** also sends the updated compaction index **70** and override index **72** to the follower server **58**, for example in the AddLogEntries message **611**, or in a separate message that is sent at or around the same time as the AddLogEntries message **611**.

[0081] The follower server **58** still has a consistent status, so before updating the replica state **64** and the replica log **66**, the follower server **58** evaluates criterion (3) using the received compaction index **70** and the current highest log index in the replica log **66**. The highest log index 2 is higher than the received compaction index 1 so the follower server **58** maintains the consistent status for its replica state **64**. The follower server **58** then adds the received third log entry **18** to the replica log **66** and updates the replica state **64** based on the received log entries. As the follower server **58** has a consistent status, the follower server **58** does not need to evaluate criterion (4).

[0082] The follower server **58** sends an acknowledgement **613** to the leader server **54** to acknowledge receipt of the AddLogEntries message **611**.

[0083] As noted, FIG. **6** relates to a situation where the follower server **58** is not lagging behind the leader server **54**. FIG. **7** illustrates a situation where network degradation causes the follower server **58** to lag far enough behind the leader server **54** that it receives an already compacted part of the log after the communication is restored.

[0084] Thus, FIG. **7** illustrates exemplary signalling between a leader server **54**, a follower server **58** and a client device **52**. The signalling in FIG. **7** relates to the example provided in FIGS. **1** and **2**. Initially, the master log **62** in the leader server **54** is empty and contains no log entries. The replica log **66** in the follower server **58** is also empty and contains no log entries. The status of the follower server **58** is initially set to the 'consistent status'.

[0085] At step **701** there is network degradation between the leader server **54** and the follower server **58**. This network degradation can take any form, but is sufficient to prevent messages sent by the leader server **54** from being received by the follower server **58**.

[0086] The master log **62** in the leader server **54** is updated to include two log entries (with log indices **1** and **2**), and the leader server **54** attempts to send an AddLogEntries message **703** to the follower server **58** that includes the first log entry **14** and the second log entry **16** from FIG. **1**. The leader server **54** also attempts to send the compaction index **70** and override index **72** to the follower server **58**, for example in the AddLogEntries message **703**, or in a separate message that is sent at or around the same time as the AddLogEntries message **703**. As no log cleaning has yet taken place, the compaction index **70** has a value of 0 and the override index **72** has a value of 0.

[0087] Due to the ongoing network degradation, the follower server **58** does not receive the AddLogEntries message **703** (or any separate message containing the compaction index **70** and override index **72**). The leader server **54** therefore does not receive an acknowledgement from the follower server **58** relating to message **703**. Optionally, the leader server **54** can resend the message **703** (and any separate message containing the compaction index **70** and override index **72**) to the follower server **58**. This is indicated by message **705**. Again, the follower server **58** does not receive the message(s) **705** due to the network degradation. At this stage, the replica log **66** in the follower server **58** is still empty and contains no log entries. The status of the follower server **58** remains a 'consistent status'.

[0088] At step **707**, the master state **60** is updated and a corresponding third log entry **18** is added to the master log **62**. The leader server **54** is able to apply the log cleaning approach to the updated master log **62** with the result that the first log entry **14** (with log index 1) is removed from the master log **62**. Following the log cleaning the value of the compaction index **70** is set to 1 and the value of the override index **72** is set to 3 (since the log entry with log index 3 overwrote the log entry with log index 1). At this time the leader server **54** may not send an AddLogEntries message to the follower server **58** as it is aware of the communication issues therebetween, for example due to the lack of acknowledgements to messages **703** and **705**.

[0089] At step **709**, the connection to the follower server **58** is restored, which enables messages to be sent between the leader server **54** and the follower server **58**. The leader server **54** therefore starts to replay (resend) the parts of the master log **62** that the follower server **58** was previously unable to receive. However, the log entry with log index 1 was removed from the master log **62** in step **707**, so the leader server **54** is only able to resend the log entries remaining in the master log **62**. Thus, the leader server **54** send the second log entry **16** (the log entry with log index 2) to the follower server **58** in an AddLogEntries message **711**. The leader server **54** also sends the compaction index **70** and override index **72** to the follower server **58**, for example in the AddLogEntries message **711**, or in a separate message that is sent at or around the same time as the AddLogEntries message **711**. The value of the compaction index **70** is 1 and the value of the override index **72** is 3.

[0090] As the follower server **58** currently has a consistent status, it evaluates criterion (3) and checks the received compaction index **70** against its highest log index in the replica log **66**. The follower server **58** still has an empty replica log **66**, and so the highest log index 0 is less than the received compaction index 1. Therefore the replica state **64** stored in the follower server **58** is potentially inconsistent with the master state **60** and so the follower server **58** updates its status to an inconsistent status for its replica state **64**, as indicated by step **713**.

[0091] The follower server **58** then adds the received second log entry to the replica log **66** and updates the replica state **64** based on the received log entry. As the follower server **58** now has an 'inconsistent' status, after updating the replica state **64** and the replica log **66** based on the received second log entry **16**, the follower server **58** evaluates criterion (4) using the received override index **72** and the current highest log index in the replica log **66**. After updating the replica log **66**, the highest log index is 2, and this is less than the received override index **72** (which is also 3), and so

criterion (4) is not satisfied and the status of the follower server **58** remains 'inconsistent'.

[0092] The follower server **58** then sends an acknowledgement **715** to the leader server **54** to acknowledge receipt of the AddLogEntries message **711**.

[0093] Subsequently, the client device **52** sends a query **717** for information about the state. This query **717** is received by the follower server **58**. As the follower server **58** has an inconsistent status, the follower server **58** responds to the query **717** with a rejection message or redirection message **719**. This message **719** can indicate that the follower server **58** is not able to respond to the query **717**, and, in the case of a redirection message, the message **719** can indicate another server (e.g. the leader server **54** or another follower server **58**) that the client device **52** can request the information from.

[0094] As shown by query **721** and response **723**, the client device **52** can request and receive the required state information from another server.

[0095] Some time after acknowledgement message **715**, the leader server **54** sends an AddLogEntries message **725** to the follower server **58** that includes the third log entry **18** that was added in step **707**. The leader server **54** also sends the current compaction index **70** and override index **72** to the follower server **58**, for example in the AddLogEntries message **725**, or in a separate message that is sent at or around the same time as the AddLogEntries message **725**. It will be noted that the follower server **58** also received these same compaction index **70** and override index **72** values in or alongside message **711**.

[0096] The follower server **58** still has an 'inconsistent' status, so after updating the replica state **64** and the replica log **66** based on the received third log entry **18**, the follower server **58** evaluates criterion (4) using the received override index **72** and the current highest log index in the replica log **66**. After updating the replica log **66**, the highest log index is 3, and this is the same as the received override index **72** (which is also 3), and so criterion (4) is satisfied and the status of the follower server **58** is updated to 'consistent'.

[0097] The follower server **58** sends an acknowledgement **727** to the leader server **54** to acknowledge receipt of the AddLogEntries message **725**.

[0098] Subsequently, the follower server **58** receives a query **729** from a client device **52** for information about the state. As the follower server **58** now has a consistent status, the follower server **58** responds to the query **729** with a response message **731** containing the requested information.

[0099] Thus, by using the compaction index **70**, the follower server **58** is able to identify the possible inconsistency in the replica state **64**, and it can redirect read requests to some other follower server **58** or the leader server **54**, until a comparison with the override index **72** in a subsequent message indicates that the follower server **58** can resume normal operation.

[0100] It will be appreciated that the examples in FIGS. **6** and **7** are simplified cases, with only a few log entries and the inconsistent status resolved as soon as possible. Real cases might be more complex, for example the inconsistent state may last a long time, and new log entries may get added to the leader server's master log **62** in the meantime, causing the override index **72** to be increased, making the affected follower server **58** require even more catching up in order to achieve a consistent status again.

[0101] The flow chart in FIG. **8** illustrates an exemplary method of operating a leader server **54** according to various embodiments.

[0102] In step **801**, the leader server **54** stores a current master state **60** of a system. In this step the leader server **54** also maintains (e.g. stores) a master log **62** relating to the current master state **62**. As noted above, the master log **62** comprises a plurality of log entries each having a respective log index, with each log entry indicating a respective change to the master state **62** of the system.

[0103] In step **803** the leader server **54** performs a log compaction procedure to reduce the size of the master log **62**. The log compaction procedure is a log cleaning procedure in which any log entry in the master log **62** for which the respective change to the master state **60** of the system has been

overridden by a change to the master state **60** in a subsequent log entry is removed from the master log **62**.

[0104] In step **805** the leader server **54** stores a compaction index **70**. The compaction index **70** is the highest log index of the log entries removed in the log compaction procedure. In embodiments where the log compaction procedure is performed multiple times (e.g. after new log entries are added to the master log **62**), the compaction index **70** is the highest log index of any of the log entries removed in any iteration of the log compaction procedure.

[0105] In step **807**, the leader server **54** stores an override index **72**. The override index **72** is the highest log index of the log entries for which the respective change to the master state **60** overrode a change to the master state **60** by an earlier log entry. In embodiments where the log compaction procedure is performed multiple times (e.g. after new log entries are added to the master log **62**), the override index **72** is the highest log index of any of the log entries for which the respective change to the master state **60** overrode a change to the master state **60** by an earlier log entry.

[0106] Steps **803**, **805** and **807** can typically be performed at the same time, and it should be appreciated that steps **805** and **807** do not have to occur in the order shown in FIG. **8**.

[0107] Finally, in step **809**, the leader server **54** sends the compaction index **70** and the override index **72** to one or more follower servers **58** that are to store a replica of the master log **62** relating to the current master state **60**.

[0108] In some embodiments, the compaction index **70** and the override index **72** are sent to the one or more follower servers **58** when there are changes to the master state **60**. In particular, the leader server **54** can receive change information indicating one or more new changes to the master state **60** of the system. The leader server **54** can update the stored master state **60** of the system based on the received change information and update the master log **62** to include one or more respective new log entries for the one or more new changes to the master state **60**. The leader server **54** then repeats step **803** on the updated master log **62** to determine if any log entries can be removed from the master log **62**, and updates the compaction index **70** and the override index **72** if appropriate for the updated master log **62** (i.e. if any further log entries have been removed by repeating step **803**). The leader server **54** then sends the new log entries, the current compaction index **70** and the current override index **72** to the one or more follower servers **58**. The respective new log entries, the current compaction index **70** and the current override index **72** can be sent in the same message. This message can be an AddLogEntries message. Alternatively, the current compaction index **70** and the current override index **72** can be sent in a different message to the message containing the new log entries.

[0109] In some embodiments, the leader server **54** can replay or resend at least a part of the (compacted) master log **62** to one or more of the follower servers **58**. Replaying or resending the compacted master log **62** can comprise sending any log entries added to the master log **62** since the follower server **58** last successfully received a log entry message from the leader server **54**. In the case of a new follower server **58**, replaying or resending the compacted master log **62** can comprise sending all of the log entries to the new follower server **58**. The log entries can be sent in one batch or in multiple batches, starting from the oldest/earliest log entries remaining in the compacted master log **62**.

[0110] The flow chart in FIG. **9** illustrates an exemplary method of operating a follower server **58** according to various embodiments. In a configuration with multiple follower servers **58**, each of the follower servers **58** can operate according to the method in FIG. **9**.

[0111] In step **901**, the follower server **58** stores a replica state **64** of a system and maintains a replica log **66** relating to the replica state **64**. The replica log **66** comprises a plurality of log entries each having a respective log index, with each log entry indicating a respective change to the replica state **64** of the system.

[0112] In step **903**, the follower server **58** receives a compaction index **70** and an override index **72** from a leader server **54** that stores a master state **60** of the system and maintains a master log **62**

relating to the current master state **62**. The compaction index **70** and the override index **72** relate to a compacted master log **62** resulting from a log compaction procedure at the leader server **54**. The compaction index **70** is the highest log index of any log entries removed in the log compaction procedure, and the override index **72** is the highest log index of any log entries for which the respective change to the master state **60** overrode a change to the master state **60** of an earlier log entry.

[0113] In step **905**, the follower server **58** determines if the stored replica state **64** is consistent with a current or previous master state **60** maintained by the leader server **54** based on a highest log index in the stored replica log **66** and one or both of the received compaction index **70** and the received override index **72**.

[0114] In some embodiments, the follower server **58** can maintain a consistency status with respect to the current master state **60** or a previous master state **60**. This consistency status can be either 'consistent' or 'inconsistent'. When the follower server's consistency status is 'consistent', the stored replica state **64** is deemed consistent with a current or previous master state **60** maintained by the leader server **54**. When the follower server's consistency status is 'inconsistent', the stored replica state **64** is deemed inconsistent, or possibly inconsistent, with a current or previous master state **60** maintained by the leader server **54**.

[0115] In some embodiments, step **905** comprises determining that the replica state **64** is inconsistent with the master state **60** if the highest log index in the stored replica log **66** is less than the received compaction index **70**. In embodiments where the follower server **58** maintains a consistency status and the status is 'consistent', the follower server's consistency status can remain 'consistent' until the highest log index in the stored replica log **66** is less than the received compaction index **70**. Where the follower server's consistency status is already 'inconsistent', the follower server **58** does not need to evaluate the relationship between the highest log index in the stored replica log **66** and the received compaction index **70**.

[0116] In some embodiments, step **905** comprises determining that the replica state **64** is consistent with the master state **60** if the highest log index in the stored replica log **66** is equal to or greater than the received override index **72**. In embodiments where the follower server **58** maintains a consistency status and the status is 'inconsistent', the follower server's consistency status can remain 'inconsistent' until the highest log index in the stored replica log **66** is equal to or greater than the received override index **72**. Where the follower server's consistency status is already 'consistent', the follower server **58** does not need to evaluate the relationship between the highest log index in the stored replica log **66** and the received override index **72**.

[0117] In some embodiments, the method can further comprise the follower server **58** receiving one or more new log entries from the leader server **54**. These new log entry(ies) can be received in the same message as the compaction index **70** and override index **72** in step **903**, or in a different message. In these embodiments, the compaction index **70** and override index **72** received in step **903** relate to the compacted master log **62** following the addition of those new log entry(ies) to the master log **62**.

[0118] Where one or more new log entries are received, the follower server **58** can update the replica state **64** and the replica log **66** based on the received one or more new log entries. In these embodiments, step **905** can comprise determining that the updated replica state is consistent with the master state if the highest log index in the updated replica log is equal to or greater than the received override index **72**.

[0119] In some embodiments, the follower server **58** can receive a request from a client device **52** for information relating to the state of the system. The follower server **58** can determine whether to provide the requested information to the client device **52** based on whether the stored replica state **64** is determined to be consistent with the master state **60**. In these embodiments, the follower server **58** can provide the requested information to the client device **52** if the stored replica state **64** is determined to be consistent with the master state **60**. The follower server **58** can reject or redirect



the request if the stored replica state **64** is not determined to be consistent with the master state **60**. [0120] FIG. **10** illustrates an apparatus **1000** according to various embodiments that can be used to implement a leader server **54** or a follower server **58** as described herein. Multiple apparatuses **1000** can be provided that each implement a respective leader server **54** or follower server **58**. [0121] The apparatus **1000** comprises processing circuitry (or logic) **1001**. It will be appreciated that the apparatus **1000** may comprise one or more virtual machines running different software and/or processes. The apparatus **1000** may therefore comprise one or more servers, switches and/or storage devices and/or may comprise cloud computing infrastructure that runs the software and/or processes. [0122] The processing circuitry **1001** controls the operation of the apparatus **1000** and can implement the methods described herein in relation to a leader server **54** and/or a follower server **58**. The processing circuitry **1001** can comprise one or more processors, processing units, multi-core processors or modules that are configured or programmed to control the apparatus **1000** in the manner described herein. In particular implementations, the processing circuitry **1001** can comprise a plurality of software and/or hardware modules that are each configured to perform, or are for performing, individual or multiple steps of the method described herein in relation to the apparatus **1000**. [0123] In some embodiments, the apparatus **1000** may optionally comprise a communications interface **1002**. The communications interface **1002** can be for use in communicating with other apparatuses, such as other servers, and/or one or more client devices. For example, the communications interface **1002** can be configured to transmit to and/or receive from other servers or client devices requests, acknowledgements, log entries, the compaction index **70**, the override index **72**, other information, data, signals, or similar. The processing circuitry **1001** may be configured to control the communications interface **1002** to transmit to and/or receive from other servers or client devices requests, acknowledgements, log entries, the compaction index **70**, the override index **72**, other information, data, signals, or similar. [0124] Optionally, the apparatus **1000** may comprise a memory **1003**. In some embodiments, the memory **1003** can be configured to store program code that can be executed by the processing circuitry **1001** to perform the method described herein in relation to the leader server **54** or the follower server **58**. Alternatively or in addition, the memory **1003** can be configured to store any requests, acknowledgements, resources, log entries, the master state **60**, the master log **62**, the replica state **64**, the replica log **66**, the compaction index **70**, the override index **72**, or other information, data, signals, or similar that are described herein. The processing circuitry **1001** may be configured to control the memory **1003** to store such information therein. [0125] The foregoing merely illustrates the principles of the disclosure. Various modifications and alterations to the described embodiments will be apparent to those skilled in the art in view of the teachings herein. It will thus be appreciated that those skilled in the art will be able to devise numerous systems, arrangements, and procedures that, although not explicitly shown or described herein, embody the principles of the disclosure and can be thus within the scope of the disclosure. Various exemplary embodiments can be used together with one another, as well as interchangeably therewith, as should be understood by those having ordinary skill in the art.

## Claims

**1.-13.** (canceled)

**14.** An apparatus for use as a leader server in a leader server-follower server log replication configuration, the apparatus configured to: (i) store a current master state of a system and maintaining a master log relating to the current master state, wherein the master log comprises a plurality of log entries each having a respective log index, wherein each log entry indicates a respective change to the master state of the system; (ii) perform a log compaction procedure to

remove from the master log any log entry for which the respective change to the master state of the system has been overridden by a change to the master state in a subsequent log entry; (iii) store a compaction index that is the highest log index of the log entries removed in the log compaction procedure; (iv) store an override index that is the highest log index of the log entries for which the respective change to the master state overrode a change to the master state by an earlier log entry; and (v) send the compaction index and the override index to one or more follower servers that are to store a replica of the master log relating to the current master state.

**15.** An apparatus as claimed in claim 14, wherein the apparatus is further configured to: (vi) receive change information indicating one or more new changes to the master state of the system; (vii) update the current master state of the system based on the received change information and updating the master log to include respective new log entries for the one or more new changes to the master state; (viii) repeat operation (ii) on the updated master log, and updating the override index and compaction index, if appropriate, for the updated master log; and (ix) send the respective new log entries, the current compaction index and the current override index to the one or more follower servers.

**16.** An apparatus as claimed in claim 15, wherein the respective new log entries, the current compaction index and the current override index are sent in the same message.

**17.** An apparatus as claimed in claim 14, wherein the apparatus is further configured to: replay or resend at least a part of the master log to one or more of the follower servers.

**18.** An apparatus for use as a follower server in a leader server-follower server log replication configuration, the apparatus configured to: (i) store a replica state of a system and maintaining a replica log relating to the replica state, wherein the replica log comprises a plurality of log entries each having a respective log index, wherein each log entry indicates a respective change to the replica state of the system; (ii) receive a compaction index and an override index from a leader server that stores a master state of the system and maintains a master log relating to the current master state, wherein the compaction index and the override index relate to a compacted master log resulting from a log compaction procedure at the leader server, wherein the compaction index is the highest log index of any log entries removed in the log compaction procedure, and wherein the override index is the highest log index of any log entries for which the respective change to the master state overrode a change to the master state of an earlier log entry; and (iii) determine if the stored replica state is consistent with a current or previous master state maintained by the leader server based on a highest log index in the stored replica log and one or both of the received compaction index and the received override index.

**19.** An apparatus as claimed in claim 18, wherein the apparatus is further configured to: maintain a consistency status with respect to the current or a previous master state, wherein the consistency status is consistent when the stored replica state is considered consistent with a current or a previous master state, and the consistency status is inconsistent when the stored replica state is considered inconsistent, or possibly inconsistent, with a current or a previous master state.

**20.** An apparatus as claimed in claim 18, wherein the apparatus is configured to determine that the replica state is inconsistent with the master state if the highest log index in the stored replica log is less than the received compaction index.

**21.** An apparatus as claimed in claim 18, wherein the apparatus is configured to determine that the replica state is consistent with the master state if the highest log index in the stored replica log is equal to or greater than the received override index.

**22.** An apparatus as claimed in claim 18, wherein operation (ii) further comprises receiving one or more new log entries from the leader server; and wherein the apparatus is further configured to update the replica state of the system and the replica log based on the received one or more new log entries.

**23.** An apparatus as claimed in claim 22, wherein: the apparatus is configured to determine that the replica state is consistent with the master state if the highest log index in the stored replica log is

equal to or greater than the received override index; and operation (iii) comprises determining that the updated replica state is consistent with the master state if the highest log index in the updated replica log is equal to or greater than the received override index.

**24.** An apparatus as claimed in claim 18, wherein the apparatus is further configured to: receive a request from a client device for information relating to the state of the system; and determine whether to provide the requested information to the client device based on whether the stored replica state is determined to be consistent with the master state.

**25.** An apparatus as claimed in claim 24, wherein the apparatus is configured to provide the requested information to the client device if the stored replica state is determined to be consistent with the master state, and reject or redirect the request if the stored replica state is not determined to be consistent with the master state.

**26-37.** (canceled)

---