

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication	20250258600
Kind Code	A1
Publication Date	August 14, 2025
Inventor(s)	CONFALONIERI; Emanuele

LOSSY COMPRESSION IN MEMORY

Abstract

In some implementations, a controller may receive and from a host device, a command to write data to a memory location. The controller may compress the data using a lossy compression operation to obtain compressed data. The controller may cause the compressed data to be written to the memory location.

Inventors:	CONFALONIERI; Emanuele (Segrate, IT)
Applicant:	Micron Technology, Inc. (Boise, ID)
Family ID:	1000008391881
Appl. No.:	19/012074
Filed:	January 07, 2025

Related U.S. Application Data

us-provisional-application US 63551778 20240209

Publication Classification

Int. Cl.: G06F3/06 (20060101)

U.S. Cl.:

CPC G06F3/0608 (20130101); G06F3/0644 (20130101); G06F3/0673 (20130101);

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATION [0001] This Patent Application claims priority to U.S. Provisional Patent Application No. 63/551,778, filed on Feb. 9, 2024, and entitled

“LOSSY COMPRESSION IN MEMORY.” The disclosure of the prior Application is considered part of and is incorporated by reference into this Patent Application.

TECHNICAL FIELD

[0002] The present disclosure generally relates to memory devices, memory device operations, and, for example, to lossy compression in memory.

BACKGROUND

[0003] Memory devices are widely used to store information in various electronic devices. A memory device includes memory cells. A memory cell is an electronic circuit capable of being programmed to a data state of two or more data states. For example, a memory cell may be programmed to a data state that represents a single binary value, often denoted by a binary “1” or a binary “0.” As another example, a memory cell may be programmed to a data state that represents a fractional value (e.g., 0.5, 1.5, or the like). To store information, an electronic device may write to, or program, a set of memory cells. To access the stored information, the electronic device may read, or sense, the stored state from the set of memory cells.

[0004] Various types of memory devices exist, including random access memory (RAM), read only memory (ROM), dynamic RAM (DRAM), static RAM (SRAM), synchronous dynamic RAM (SDRAM), ferroelectric RAM (FeRAM), magnetic RAM (MRAM), resistive RAM (RRAM), holographic RAM (HRAM), flash memory (e.g., NAND memory and NOR memory), and others. A memory device may be volatile or non-volatile. Non-volatile memory (e.g., flash memory) can store data for extended periods of time even in the absence of an external power source. Volatile memory (e.g., DRAM) may lose stored data over time unless the volatile memory is refreshed by a power source. In some examples, a memory device may be associated with a compute express link (CXL). For example, the memory device may be a CXL compliant memory device and/or may include a CXL interface.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0005] FIG. 1 is a diagram illustrating an example system capable of lossy compression in memory.

[0006] FIG. 2 is a diagram illustrating an example system capable of lossy compression in memory.

[0007] FIG. 3 is a diagram illustrating an example CXL fabric configuration capable of lossy compression in memory.

[0008] FIGS. 4A-4C are diagrams of an example of lossy compression in memory.

[0009] FIG. 5 is a diagram of an example of a controller.

[0010] FIG. 6 is a diagram of an example system with memory pooling.

[0011] FIG. 7 is a flowchart of an example method associated with lossy compression in memory.

DETAILED DESCRIPTION

[0012] Emerging computing applications, such as artificial intelligence, cloud computing, or virtual reality, may involve the use of very large datasets. The efficient processing of large datasets requires a significant amount of memory (e.g., random-access memory (RAM)). In some cases, virtual memory, which uses storage (e.g., disk storage) to provide a virtual address space that can exceed a capacity of main memory, is one technique to accommodate large datasets. However, virtual memory may be associated with diminished system performance. Furthermore, expanding a system's main memory may be cost prohibitive.

[0013] Compression enables data size to be reduced, thereby increasing memory capacity, lowering energy costs, and reducing bandwidth consumption. For example, compression may alleviate memory bandwidth pressure by reducing a volume of data transferred between a main memory and a processor. In general, main memory compression is performed at a system level (e.g., by a host system). Moreover, main memory compression generally employs lossless compression methods,

which are not associated with high compression ratios.

[0014] Some implementations described herein exploit memory disaggregation to enable data compression to be performed by a memory system (e.g., a CXL memory module), rather than compression being performed at a host processor. This enables a relatively smaller amount of memory to provide increased memory capacity. In some implementations, the compression performed by the memory system may be a lossy compression (e.g., used in approximated computing), which can be utilized with large classes of applications that are tolerant to approximation. For example, the memory system may include a lossy compression and decompression engine. “Lossy compression” refers to a data compression technique that reduces the data size of content by selectively and permanently discarding portions of the original data of the content.

[0015] At a system level, one or more memory systems enabled with lossy compression may be part of a memory pool (e.g., along with one or more memory systems that are not enabled with lossy compression). In some implementations, a host system may flag pages allocated for an application as being approximable or not approximable in accordance with a configuration for the application that indicates whether particular data regions are approximable. “Approximable” information can be adequately represented or approximated with an acceptable degree of precision, enabling the removal of less-significant details to achieve compression. A memory system enabled with lossy compression may be used to compress and store approximable pages, thereby efficiently utilizing resources of the memory pool. In this way, large amounts of data can be managed in memory without the use of storage (e.g., disk storage), thereby improving system performance. Furthermore, the high compression ratios associated with lossy compression enable increased memory capacity in a relatively smaller amount of memory. Accordingly, due to the flexibility of disaggregated systems, these high-capacity memory resources can be made available to a rack or a fabric and may be used to address the needs of applications that are tolerant to approximation.

[0016] FIG. 1 is a diagram illustrating an example system **100** capable of lossy compression in memory. The system **100** may include one or more devices, apparatuses, and/or components for performing operations described herein. For example, the system **100** may include a host system **105** and a memory system **110**. The memory system **110** may include a memory system controller **115** and one or more memory devices **120**, shown as memory devices **120-1** through **120-N** (where $N \geq 1$). A memory device may include a local controller **125** and one or more memory arrays **130**. The host system **105** may communicate with the memory system **110** (e.g., the memory system controller **115** of the memory system **110**) via a host interface **140**. The memory system controller **115** and the memory devices **120** may communicate via respective memory interfaces **145**, shown as memory interfaces **145-1** through **145-N** (where $N \geq 1$).

[0017] The system **100** may be any electronic device configured to store data in memory. For example, the system **100** may be a computer, a mobile phone, a wired or wireless communication device, a network device, a server, a device in a data center, a device in a cloud computing environment, a vehicle (e.g., an automobile or an airplane), and/or an Internet of Things (IoT) device. The host system **105** may include a host processor **150**. The host processor **150** may include one or more processors configured to execute instructions and store data in the memory system **110**. For example, the host processor **150** may include a central processing unit (CPU), a graphics processing unit (GPU), a field-programmable gate array (FPGA), an application-specific integrated circuit (ASIC), and/or another type of processing component.

[0018] The memory system **110** may be any electronic device or apparatus configured to store data in memory. For example, the memory system **110** may be a hard drive, a solid-state drive (SSD), a flash memory system (e.g., a NAND flash memory system or a NOR flash memory system), a universal serial bus (USB) drive, a memory card (e.g., a secure digital (SD) card), a secondary storage device, a non-volatile memory express (NVMe) device, an embedded multimedia card (eMMC) device, a dual in-line memory module (DIMM), and/or a RAM device, such as a dynamic

RAM (DRAM) device or a static RAM (SRAM) device.

[0019] The memory system controller **115** may be any device configured to control operations of the memory system **110** and/or operations of the memory devices **120**. For example, the memory system controller **115** may include control logic, a memory controller, a system controller, an ASIC, an FPGA, a processor, a microcontroller, and/or one or more processing components. In some implementations, the memory system controller **115** may communicate with the host system **105** and may instruct one or more memory devices **120** regarding memory operations to be performed by those one or more memory devices **120** based on one or more instructions from the host system **105**. For example, the memory system controller **115** may provide instructions to a local controller **125** regarding memory operations to be performed by the local controller **125** in connection with a corresponding memory device **120**.

[0020] A memory device **120** may include a local controller **125** and one or more memory arrays **130**. In some implementations, a memory device **120** includes a single memory array **130**. In some implementations, each memory device **120** of the memory system **110** may be implemented in a separate semiconductor package or on a separate die that includes a respective local controller **125** and a respective memory array **130** of that memory device **120**. The memory system **110** may include multiple memory devices **120**.

[0021] A local controller **125** may be any device configured to control memory operations of a memory device **120** within which the local controller **125** is included (e.g., and not to control memory operations of other memory devices **120**). For example, the local controller **125** may include control logic, a memory controller, a system controller, an ASIC, an FPGA, a processor, a microcontroller, and/or one or more processing components. In some implementations, the local controller **125** may communicate with the memory system controller **115** and may control operations performed on a memory array **130** coupled with the local controller **125** based on one or more instructions from the memory system controller **115**. As an example, the memory system controller **115** may be an SSD controller, and the local controller **125** may be a NAND controller.

[0022] A memory array **130** may include an array of memory cells configured to store data. For example, a memory array **130** may include a non-volatile memory array (e.g., a NAND memory array or a NOR memory array) or a volatile memory array (e.g., an SRAM array or a DRAM array). In some implementations, the memory system **110** may include one or more volatile memory arrays **135**. A volatile memory array **135** may include an SRAM array and/or a DRAM array, among other examples. The one or more volatile memory arrays **135** may be included in the memory system controller **115**, in one or more memory devices **120**, and/or in both the memory system controller **115** and one or more memory devices **120**. In some implementations, the memory system **110** may include both non-volatile memory capable of maintaining stored data after the memory system **110** is powered off and volatile memory (e.g., a volatile memory array **135**) that requires power to maintain stored data and that loses stored data after the memory system **110** is powered off. For example, a volatile memory array **135** may cache data read from or to be written to non-volatile memory, and/or may cache instructions to be executed by a controller of the memory system **110**.

[0023] The host interface **140** enables communication between the host system **105** (e.g., the host processor **150**) and the memory system **110** (e.g., the memory system controller **115**). The host interface **140** may include, for example, a Small Computer System Interface (SCSI), a Serial-Attached SCSI (SAS), a Serial Advanced Technology Attachment (SATA) interface, a Peripheral Component Interconnect Express (PCIe) interface, an NVMe interface, a USB interface, a Universal Flash Storage (UFS) interface, an eMMC interface, a double data rate (DDR) interface, and/or a DIMM interface.

[0024] In some examples, the memory device **120** may be a compute express link (CXL) compliant memory device **120**. For example, the memory device **120** may include a PCIe/CXL interface (e.g., the host interface **140** may be associated with a PCIe/CXL interface). CXL is a high-speed CPU-to-

device and CPU-to-memory interconnect designed to accelerate next-generation performance. CXL technology maintains memory coherency between the CPU memory space and memory on attached devices, which allows resource sharing for higher performance, reduced software stack complexity, and lower overall system cost. CXL is designed to be an industry open standard interface for high-speed communications. CXL technology is built on the PCIe infrastructure, leveraging PCIe physical and electrical interfaces to provide advanced protocol in areas such as input/output (I/O) protocol, memory protocol, and coherency interface.

[0025] The memory interface **145** enables communication between the memory system **110** and the memory device **120**. The memory interface **145** may include a non-volatile memory interface (e.g., for communicating with non-volatile memory), such as a NAND interface or a NOR interface. Additionally, or alternatively, the memory interface **145** may include a volatile memory interface (e.g., for communicating with volatile memory), such as a DDR interface.

[0026] Although the example memory system **110** described above includes a memory system controller **115**, in some implementations, the memory system **110** does not include a memory system controller **115**. For example, an external controller (e.g., included in the host system **105**) and/or one or more local controllers **125** included in one or more corresponding memory devices **120** may perform the operations described herein as being performed by the memory system controller **115**. Furthermore, as used herein, a “controller” may refer to the memory system controller **115**, a local controller **125**, or an external controller. In some implementations, a set of operations described herein as being performed by a controller may be performed by a single controller. For example, the entire set of operations may be performed by a single memory system controller **115**, a single local controller **125**, or a single external controller. Alternatively, a set of operations described herein as being performed by a controller may be performed by more than one controller. For example, a first subset of the operations may be performed by the memory system controller **115** and a second subset of the operations may be performed by a local controller **125**. Furthermore, the term “memory apparatus” may refer to the memory system **110** or a memory device **120**, depending on the context.

[0027] A controller (e.g., the memory system controller **115**, a local controller **125**, or an external controller) may control operations performed on memory (e.g., a memory array **130**), such as by executing one or more instructions. For example, the memory system **110** and/or a memory device **120** may store one or more instructions in memory as firmware, and the controller may execute those one or more instructions. Additionally, or alternatively, the controller may receive one or more instructions from the host system **105** and/or from the memory system controller **115**, and may execute those one or more instructions. In some implementations, a non-transitory computer-readable medium (e.g., volatile memory and/or non-volatile memory) may store a set of instructions (e.g., one or more instructions or code) for execution by the controller. The controller may execute the set of instructions to perform one or more operations or methods described herein. In some implementations, execution of the set of instructions, by the controller, causes the controller, the memory system **110**, and/or a memory device **120** to perform one or more operations or methods described herein. In some implementations, hardwired circuitry is used instead of or in combination with the one or more instructions to perform one or more operations or methods described herein. Additionally, or alternatively, the controller may be configured to perform one or more operations or methods described herein. An instruction is sometimes called a “command.”

[0028] For example, the controller (e.g., the memory system controller **115**, a local controller **125**, or an external controller) may transmit signals to and/or receive signals from memory (e.g., one or more memory arrays **130**) based on the one or more instructions, such as to transfer data to (e.g., write or program), to transfer data from (e.g., read), to erase, and/or to refresh all or a portion of the memory (e.g., one or more memory cells, pages, sub-blocks, blocks, or planes of the memory). Additionally, or alternatively, the controller may be configured to control access to the memory and/or to provide a translation layer between the host system **105** and the memory (e.g., for

mapping logical addresses to physical addresses of a memory array **130**). In some implementations, the controller may translate a host interface command (e.g., a command received from the host system **105**) into a memory interface command (e.g., a command for performing an operation on a memory array **130**).

[0029] The number and arrangement of components shown in FIG. **1** are provided as an example. In practice, there may be additional components, fewer components, different components, or differently arranged components than those shown in FIG. **1**. Furthermore, two or more components shown in FIG. **1** may be implemented within a single component, or a single component shown in FIG. **1** may be implemented as multiple, distributed components. Additionally, or alternatively, a set of components (e.g., one or more components) shown in FIG. **1** may perform one or more operations described as being performed by another set of components shown in FIG. **1**.

[0030] FIG. **2** is a diagram illustrating an example system **200** capable of lossy compression in memory. The system **200** may include one or more devices, apparatuses, and/or components for performing operations described herein. In some implementations, the system **200** is a CXL system. For example, the system **200** may be configured to utilize memory pooling with multiple logical devices.

[0031] As shown, the system **200** may include a plurality of host devices **205** and a plurality of memory devices **210**. The host devices **205** may include CPUs. In some implementations, each host device **205** corresponds to the host system **105** and/or the host processor **150**. The memory devices **210** may include CXL devices (e.g., CXL compliant devices). In some implementations, the memory devices **210** may include Type 3 CXL devices. For example, the memory devices **210** may include memory modules and/or memory expansion devices (e.g., volatile memory devices). As an example, a memory device **210** may include a memory controller (e.g., a CXL ASIC controller) and one or more memory components (e.g., memory packages) coupled to the memory controller. In some implementations, each memory device **210** may correspond to a memory system **110** or a memory device **120**.

[0032] The host devices **205** and the memory devices **210** may be communicatively coupled via a switch **215** (e.g., one or more switches **215**). In some implementations, the switch **215** is a CXL switch. The switch **215** may include a hardware bridge, such as a PCIe bridge. In some implementations, the system **200** (e.g., the host devices **205**, the memory devices **210**, and/or the switch **215**) may be in a fabric configuration. For example, the system **200** (e.g., the host devices **205**, the memory devices **210**, and/or the switch **215**) may be in a configuration of a CXL fabric.

[0033] The system **200** may include a manager component **220**. The manager component **220** may be a device or implemented in a device. For example, the manager component **220** is shown as a component of the switch **215**, but may reside outside of the switch **215** in some examples. The manager component **220** may be configured to manage the fabric of the system **200** (e.g., the manager component **220** may be a fabric manager). The manager component **220** may be implemented in hardware and/or software.

[0034] The memory devices **210** may form a memory pool for the host devices **205**. In some implementations, the memory (e.g., the memory components) of a memory device **210** may be mapped to one or more (e.g., multiple) logical memory devices **225**. At a system level, a logical memory device **225** may correspond to a range of memory addresses. The logical memory devices **225** of a memory device **210** may be available to one or more of the host devices **205**. For example, all of the logical memory devices **225** of a single memory device **210** may be available to a particular host device **205**, or a subset of the logical memory devices **225** of a single memory device **210** may be available to a particular host device **205**. In some implementations, the logical memory devices **225** available to a particular host device **205** may be distributed across multiple memory devices **210**.

[0035] The number and arrangement of components shown in FIG. **2** are provided as an example.

In practice, there may be additional components, fewer components, different components, or differently arranged components than those shown in FIG. 2. Furthermore, two or more components shown in FIG. 2 may be implemented within a single component, or a single component shown in FIG. 2 may be implemented as multiple, distributed components. Additionally, or alternatively, a set of components (e.g., one or more components) shown in FIG. 2 may perform one or more operations described as being performed by another set of components shown in FIG. 2.

[0036] FIG. 3 is a diagram illustrating an example CXL fabric configuration **300** capable of lossy compression in memory. As shown, a plurality of devices **305** (e.g., CXL compliant devices) may be interconnected via the CXL fabric **310**. The plurality of devices **305** may include one or more host devices (e.g., host devices **205**), one or more Type 1 CXL devices (e.g., specialized accelerators, such as network interface controllers), one or more Type 2 CXL devices (e.g., general-purpose accelerators, such as GPUs, ASICs, or FPGAs), and/or one or more Type 3 CXL devices (e.g., global fabric attached memory (GFAM) devices, such as memory devices **210**). The CXL fabric **310** may be configured to interconnect the devices **305** to enable communication (e.g., peer-to-peer communication) between the devices **305**. The CXL fabric may be implemented in hardware and/or software. As further shown, the CXL fabric configuration **300** may include a manager component **315**, similar to the manager component **220** described herein.

[0037] In some implementations, one or more systems, devices, apparatuses, components, and/or controllers of FIGS. 1-3 may be configured to receive, from a host device, a command to write data to a memory location; compress the data to obtain compressed data, responsive to the memory location corresponding to a logical memory device that is configured with lossy compression enabled; and cause the compressed data to be written to the memory location.

[0038] In some implementations, one or more systems, devices, apparatuses, components, and/or controllers of FIGS. 1-3 may be configured to receive, from a host device, a command to write data to a memory location; compress the data using a lossy compression operation to obtain compressed data, responsive to the memory location corresponding to a logical memory device that is configured with lossy compression enabled; and write the compressed data to the memory location.

[0039] In some implementations, one or more systems, devices, apparatuses, components, and/or controllers of FIGS. 1-3 may be configured to receive, from a host device, a command to write data to a memory location; compress the data using a lossy compression operation to obtain compressed data; and cause the compressed data to be written to the memory location.

[0040] The number and arrangement of components shown in FIG. 3 are provided as an example. In practice, there may be additional components, fewer components, different components, or differently arranged components than those shown in FIG. 3. Furthermore, two or more components shown in FIG. 3 may be implemented within a single component, or a single component shown in FIG. 3 may be implemented as multiple, distributed components. Additionally, or alternatively, a set of components (e.g., one or more components) shown in FIG. 3 may perform one or more operations described as being performed by another set of components shown in FIG. 3.

[0041] FIGS. 4A-4C are diagrams of an example **400** of lossy compression in memory. As shown, example **400** includes a host device **405** and a memory system **410**. The host device **405** may correspond to the host system **105**, the host processor **150**, a host device **205**, and/or a device **305**. The memory system **410** may correspond to the memory system **110**, a memory device **120**, a memory device **210**, and/or a device **305**. For example, the host device **405** and the memory system **410** may be in the system **200** and/or in the CXL fabric configuration **300**.

[0042] In some implementations, the memory system **410** is a Type 3 CXL device. For example, the memory system **410** may include a memory module and/or a memory expansion device. As shown, the memory system **410** includes a controller **412** and one or more memory components **414** (e.g., memory packages). The controller **412** may include an ASIC, an FPGA, or the like. In

some implementations, the controller **412** may include a lossy compression/decompression engine **416**. The lossy compression/decompression engine **416** enables the memory system **410** to compress data using a lossy compression operation. The lossy compression/decompression engine **416** may be implemented in hardware and/or software.

[0043] As shown, the memory system **410** may have a logical characterization **418**. In the logical characterization **418** of the memory system **410**, the memory components **414** may be configured as multiple logical memory devices **420**, as described herein. In some implementations, one or more of the logical memory devices **420** may be configured with lossy compression enabled (e.g., logical memory devices **420-2** and **420-3** in example **400**). Moreover, one or more of the logical memory devices **420** may be configured with lossy compression disabled (e.g., logical memory devices **420-1**, **420-4**, and **420-n** in example **400**). In other words, the memory system **410** is configurable such that lossy compression (e.g., the lossy compression/decompression engine **416**) may be selectively enabled or disabled in each logical memory device **420** (e.g., enabled or disabled on a per-logical-memory-device basis). In some implementations, the memory components **414** may be configured as a single logical memory device **420** configured with lossy compression enabled or disabled. In some implementations, the logical memory devices **420** may be selectively configured with lossy compression enabled or disabled at an initialization (e.g., a booting) of the memory system **410**.

[0044] As further shown, the logical characterization **418** may include one or more logical ports **422** (also referred to as “CXL ports” or “heads”) configured to handle requests to the logical memory devices **420**. For example, the logical memory devices **420** may be mapped to the one or more logical ports **422**. In some implementations, the logical characterization **418** may include multiple logical ports **422**.

[0045] The host device **405** may implement one or more applications (e.g., artificial intelligence applications, graphics applications, or the like). Particular data regions of an application may be configured (e.g., annotated) as being approximable (e.g., by a programmer of the application, using a specialized system call). For example, in a graphics application, data relating to a background of an image may be configured as being approximable. In some implementations, approximable data may be configured with one or more error thresholds (e.g., by a programmer of the application). For example, a first threshold may limit an allowable error that is introduced by a single compression event, and a second threshold may limit a total accumulated error across the full application lifetime.

[0046] The host device **405** may allocate one or more pages for an application executing on the host device **405**. The host device **405** may mark an allocated page as being approximable if the page is to contain application data configured as being approximable. For example, the host device **405** may flag pages as being approximable or not approximable. As an example, each entry for a page in a page table and/or a translation lookaside buffer (TLB) may include an approximable indication (e.g., an additional bit) to indicate whether the page is approximable.

[0047] If data to be written to memory is associated with a flag indicating that the data is approximable (e.g., in the page table and/or the TLB), then the host device **405** may map the data to a memory location (e.g., a memory address) associated with a logical memory device **420** configured with lossy compression enabled. For example, the host device **405** may store information indicating one or more address ranges associated with logical memory devices **420** configured with lossy compression enabled and/or indicating one or more address ranges associated with logical memory devices **420** configured with lossy compression disabled (or that are incapable of lossy compression). In this way, one or more logical memory devices **420** configured with lossy compression enabled can be used to store approximable pages.

[0048] As shown in FIG. 4A, and by reference number **450**, the host device **405** may transmit, and the controller **412** may receive, a command to write the data to the memory location. The controller **412** may identify that the memory location corresponds to a logical memory device **420** that is

configured with lossy compression enabled (e.g., logical memory device **420-2** or **420-3**), which indicates that the data is to be compressed before being written.

[0049] As shown by reference number **455**, responsive to the memory location corresponding to a logical memory device **420** that is configured with lossy compression enabled, the controller **412** may compress the data using a lossy compression operation to obtain compressed data. For example, the controller **412** may be configured to compress the data using the lossy compression/decompression engine **416**. The lossy compression operation may reduce the size of the data. The lossy compression operation may use a lossy compression model or algorithm. For example, the lossy compression operation may use a fast, error-bounded, lossy high-performance computing (HPC) data compression, such as a Squeeze (SZ) algorithm, an SZx algorithm, or another algorithm.

[0050] As shown by reference number **460**, the controller **412** may cause the compressed data to be written to the memory location (e.g., the controller **412** may write the compressed data to the memory location). For example, the controller **412** may cause the compressed data to be written to one or more memory components **414** mapped to the logical memory device **420**. Writing the compressed data to memory, rather than the original uncompressed data, increases the virtual memory capacity of the memory system **410**.

[0051] In some implementations, prior to writing the compressed data to the memory location, the controller **412** may determine an error value for the compressed data. To determine the error value, the controller **412** may decompress the compressed data to obtain decompressed data (e.g., using the lossy compression/decompression engine **416**) and compare the decompressed data to the original data. For example, the error value may be in accordance with a difference between the decompressed data and the original data. Accordingly, the controller **412** may cause the compressed data to be written to the memory location responsive to the error value satisfying a threshold (e.g., the error value does not exceed the threshold). Otherwise, if the error value does not satisfy the threshold (e.g., the error value exceeds the threshold), the controller **412** may discard the compressed data and may cause the original uncompressed data to be written to the memory location (e.g., compression is skipped for the data).

[0052] In some implementations, the controller **412** may cause the compressed data to be written to the memory location if a relative error (e.g., computed as described above) of each individual value of a block does not exceed a first percentage threshold. Additionally, or alternatively, the controller **412** may cause the compressed data to be written to the memory location if the average relative error across all values of the block does not exceed a second percentage threshold.

[0053] In some implementations, the controller **412** may determine an updated virtual memory capacity of the logical memory devices **420** resulting from writing the compressed data to the memory location. For example, compressing the data may result in an increased virtual memory capacity. The controller **412** may output an indication of the updated virtual memory capacity to the host device **405**. In some implementations, the memory system **410** may utilize a CXL dynamic capacity feature to manage capacity changes due to compression.

[0054] As shown in FIG. **4B**, and by reference number **465**, the host device **405** may transmit, and the controller **412** may receive an additional command to read from the memory location. The controller **412** may identify that the memory location corresponds to a logical memory device **420** that is configured with lossy compression enabled (e.g., logical memory device **420-2** or **420-3**), which indicates that the data being read is to be decompressed.

[0055] As shown by reference number **470**, responsive to the memory location corresponding to a logical memory device **420** that is configured with lossy compression enabled, the controller **412** may read the compressed data from the memory location and decompress the compressed data using a decompression operation to obtain decompressed data. For example, the controller **412** may be configured to decompress the data using the lossy compression/decompression engine **416**. The decompression operation may correlate to the lossy compression operation (e.g., the decompression

operation may provide decompression of data compressed by the lossy compression operation). As shown by reference number **475**, the memory system **410** (e.g., the controller **412**) may then output the decompressed data to the host device **405**.

[0056] As shown in FIG. **4C**, and by reference number **480**, the host device **405** may transmit, and the controller **412** may receive an additional command to write additional data to an additional memory location. The controller **412** may identify that the additional memory location corresponds to a different logical memory device **420** that is configured with lossy compression disabled (e.g., logical memory device **420-1**, **420-4**, or **420-n**), which indicates that the additional data is to be written uncompressed. Accordingly, as shown by reference number **485**, responsive to the memory location corresponding to the different logical memory device **420** that is configured with lossy compression disabled, the controller **412** may cause the additional data to be written to the additional memory location without using the lossy compression operation (e.g., the controller **412** may write the additional data to the additional memory location without lossy compression). For example, the controller **412** may cause the additional data to be written, without lossy compression, to one or more memory components **414** mapped to the different logical memory device **420**.

[0057] As indicated above, FIG. **4** is provided as an example. Other examples may differ from what is described with regard to FIG. **4**.

[0058] FIG. **5** is a diagram of an example of the controller **412**. As shown, a PCIe/CXL interconnect **502** (e.g., based on the PCIe physical link layer and the CXL protocol) may connect to a CXL physical layer (PHY) component **504** of the controller **412**. The controller **412** may include a CXL control component **506**, the lossy compression/decompression engine **416**, a channel interleaver component **508** (e.g., configured to remap an address to a physical memory address), one or more error manager components **510**, one or more DDR control components **512**, and/or one or more DDR PHY components **514**. The DDR PHY components **514** may be configured to drive a memory bus interface via one or more DDR interconnects **516**. The controller **412** may also include one or more peripherals, such as an inter-integrated circuit (I2C) and/or improved inter-integrated circuit (I3C) component **518**, a Joint Test Action Group (JTAG) component **520**, a quad serial peripheral interface (QSPI) component **522**, and/or a CPU **524** (e.g., for peripheral computing operations, such as security operations), among other examples.

[0059] As indicated above, FIG. **5** is provided as an example. Other examples may differ from what is described with regard to FIG. **5**.

[0060] FIG. **6** is a diagram of an example system **600** with memory pooling. The system **600** is illustrated as a rack system, but other configurations for the system **600** may be employed. As shown, the system **600** may include one or more computing servers **605** (e.g., each corresponding to a host device **405** or collectively corresponding to a host device **405**), a memory pool **610**, and a switch **615** (e.g., a top-of-rack (TOR) switch, which may correspond to switch **215**).

[0061] The memory pool **610** may include a plurality of memory systems **620**. The memory systems **620** may be Type 3 CXL devices. For example, the memory systems **620** may include memory modules and/or memory expansion devices. The memory systems **620** may include one or more non-compression memory systems **620a** and/or one or more compression memory systems **620b**. A non-compression memory system **620a** may lack a capability to perform lossy compression. For example, a non-compression memory system **620a** may lack a lossy compression/decompression engine, as described herein. A compression memory system **620b** may have a capability to perform lossy compression. For example, a compression memory system **620b** may correspond to the memory system **410**. As an example, a compression memory system **620b** may be characterized as a plurality of logical memory devices, and lossy compression may be enabled or disabled on a per-logical-memory-device basis, as described herein.

[0062] As indicated above, FIG. **6** is provided as an example. Other examples may differ from what is described with regard to FIG. **6**.

[0063] FIG. **7** is a flowchart of an example method **700** associated with lossy compression in

memory. In some implementations, a controller (e.g., the controller **412**) may perform or may be configured to perform the method **700**. In some implementations, another device or a group of devices separate from or including the controller (e.g., the memory system **410**) may perform or may be configured to perform the method **700**. Additionally, or alternatively, one or more components of the controller (e.g., the lossy compression/decompression engine **416**) may perform or may be configured to perform the method **700**. Thus, means for performing the method **700** may include the controller and/or one or more components of the controller. Additionally, or alternatively, a non-transitory computer-readable medium may store one or more instructions that, when executed by the controller, cause the controller to perform the method **700**.

[0064] As shown in FIG. 7, the method **700** may include receiving, from a host device, a command to write data to a memory location (block **710**). As further shown in FIG. 7, the method **700** may include compressing the data using a lossy compression operation to obtain compressed data (block **720**). As further shown in FIG. 7, the method **700** may include causing the compressed data to be written to the memory location (block **730**).

[0065] The method **700** may include additional aspects, such as any single aspect or any combination of aspects described below and/or described in connection with one or more other methods or operations described elsewhere herein.

[0066] In a first aspect, a memory system includes one or more memory components configured as one or more logical memory devices.

[0067] In a second aspect, alone or in combination with the first aspect, the data is compressed and responsive to the memory location corresponding to a logical memory device, of the one or more logical memory devices, that is configured with lossy compression enabled.

[0068] In a third aspect, alone or in combination with one or more of the first and second aspects, the method **700** includes receiving an additional command to write additional data to an additional memory location, and causing the additional data to be written to the additional memory location without using the lossy compression operation, responsive to the additional memory location corresponding to an additional logical memory device, of the one or more logical memory devices, that is configured with lossy compression disabled.

[0069] In a fourth aspect, alone or in combination with one or more of the first through third aspects, the one or more logical memory devices include multiple logical memory devices mapped to multiple logical ports.

[0070] In a fifth aspect, alone or in combination with one or more of the first through fourth aspects, the method **700** includes receiving an additional command to read from the memory location, decompressing the compressed data using a decompression operation to obtain decompressed data, and outputting the decompressed data.

[0071] Although FIG. 7 shows example blocks of a method **700**, in some implementations, the method **700** may include additional blocks, fewer blocks, different blocks, or differently arranged blocks than those depicted in FIG. 7. Additionally, or alternatively, two or more of the blocks of the method **700** may be performed in parallel. The method **700** is an example of one method that may be performed by one or more devices described herein. These one or more devices may perform or may be configured to perform one or more other methods based on operations described herein.

[0072] In some implementations, a memory system includes one or more memory components configurable as one or more logical memory devices, where the one or more logical memory devices are to be configured with lossy compression enabled or lossy compression disabled; and a controller configured to: receive, from a host device, a command to write data to a memory location; compress the data to obtain compressed data, responsive to the memory location corresponding to a logical memory device, of the one or more logical memory devices, that is configured with lossy compression enabled; and cause the compressed data to be written to the memory location.

[0073] In some implementations, a system includes one or more host devices; and a memory pool,

including: a non-compression memory system; and a compression memory system, including: one or more memory components configurable as one or more logical memory devices, where the one or more logical memory devices are to be configured with lossy compression enabled or lossy compression disabled; and a controller configured to: receive, from a host device, a command to write data to a memory location; compress the data using a lossy compression operation to obtain compressed data, responsive to the memory location corresponding to a logical memory device, of the one or more logical memory devices, that is configured with lossy compression enabled; and write the compressed data to the memory location.

[0074] In some implementations, a method includes receiving, by a controller of a memory system, and from a host device, a command to write data to a memory location; compressing, by the controller, the data using a lossy compression operation to obtain compressed data; and causing, by the controller, the compressed data to be written to the memory location.

[0075] The foregoing disclosure provides illustration and description but is not intended to be exhaustive or to limit the implementations to the precise forms disclosed. Modifications and variations may be made in light of the above disclosure or may be acquired from practice of the implementations described herein.

[0076] As used herein, “satisfying a threshold” may, depending on the context, refer to a value being greater than the threshold, greater than or equal to the threshold, less than the threshold, less than or equal to the threshold, equal to the threshold, not equal to the threshold, or the like.

[0077] Even though particular combinations of features are recited in the claims and/or disclosed in the specification, these combinations are not intended to limit the disclosure of implementations described herein. Many of these features may be combined in ways not specifically recited in the claims and/or disclosed in the specification. For example, the disclosure includes each dependent claim in a claim set in combination with every other individual claim in that claim set and every combination of multiple claims in that claim set. As used herein, a phrase referring to “at least one of” a list of items refers to any combination of those items, including single members. As an example, “at least one of: a, b, or c” is intended to cover a, b, c, a+b, a+c, b+c, and a+b+c, as well as any combination with multiples of the same element (e.g., a+a, a+a+a, a+a+b, a+a+c, a+b+b, a+c+c, b+b, b+b+b, b+b+c, c+c, and c+c+c, or any other ordering of a, b, and c).

[0078] When “a component” or “one or more components” (or another element, such as “a controller” or “one or more controllers”) is described or claimed (within a single claim or across multiple claims) as performing multiple operations or being configured to perform multiple operations, this language is intended to broadly cover a variety of architectures and environments. For example, unless explicitly claimed otherwise (e.g., via the use of “first component” and “second component” or other language that differentiates components in the claims), this language is intended to cover a single component performing or being configured to perform all of the operations, a group of components collectively performing or being configured to perform all of the operations, a first component performing or being configured to perform a first operation and a second component performing or being configured to perform a second operation, or any combination of components performing or being configured to perform the operations. For example, when a claim has the form “one or more components configured to: perform X; perform Y; and perform Z,” that claim should be interpreted to mean “one or more components configured to perform X; one or more (possibly different) components configured to perform Y; and one or more (also possibly different) components configured to perform Z.”

[0079] No element, act, or instruction used herein should be construed as critical or essential unless explicitly described as such. Also, as used herein, the articles “a” and “an” are intended to include one or more items and may be used interchangeably with “one or more.” Further, as used herein, the article “the” is intended to include one or more items referenced in connection with the article “the” and may be used interchangeably with “the one or more.” Where only one item is intended, the phrase “only one,” “single,” or similar language is used. Also, as used herein, the terms “has,”

“have,” “having,” or the like are intended to be open-ended terms that do not limit an element that they modify (e.g., an element “having” A may also have B). Further, the phrase “based on” is intended to mean “based, at least in part, on” unless explicitly stated otherwise. As used herein, the term “multiple” can be replaced with “a plurality of” and vice versa. Also, as used herein, the term “or” is intended to be inclusive when used in a series and may be used interchangeably with “and/or,” unless explicitly stated otherwise (e.g., if used in combination with “either” or “only one of”).

Claims

1. A memory system, comprising: one or more memory components configurable as one or more logical memory devices, wherein the one or more logical memory devices are to be configured with lossy compression enabled or lossy compression disabled; and a controller configured to: receive, from a host device, a command to write data to a memory location; compress the data to obtain compressed data, responsive to the memory location corresponding to a logical memory device, of the one or more logical memory devices, that is configured with lossy compression enabled; and cause the compressed data to be written to the memory location.
2. The memory system of claim 1, wherein the controller is configured to compress the data using the lossy compression operation using a lossy compression and decompression engine of the controller.
3. The memory system of claim 1, wherein the controller is further configured to: receive an additional command to write additional data to an additional memory location; and cause the additional data to be written to the additional memory location without using the lossy compression operation, responsive to the additional memory location corresponding to an additional logical memory device, of the one or more logical memory devices, that is configured with lossy compression disabled.
4. The memory system of claim 1, wherein the controller is further configured to: receive an additional command to read from the memory location; decompress the compressed data using a decompression operation to obtain decompressed data; and output the decompressed data.
5. The memory system of claim 1, wherein the controller is further configured to: determine an updated virtual memory capacity of the one or more logical memory devices resulting from the compressed data being written to the memory location; and output an indication of the updated virtual memory capacity.
6. The memory system of claim 1, wherein the data is associated with a flag indicating that the data is approximable.
7. The memory system of claim 1, wherein the one or more memory components comprise volatile memory.
8. A system, comprising: one or more host devices; and a memory pool, comprising: a non-compression memory system; and a compression memory system, comprising: one or more memory components configurable as one or more logical memory devices, wherein the one or more logical memory devices are to be configured with lossy compression enabled or lossy compression disabled; and a controller configured to: receive, from a host device, a command to write data to a memory location; compress the data using a lossy compression operation to obtain compressed data, responsive to the memory location corresponding to a logical memory device, of the one or more logical memory devices, that is configured with lossy compression enabled; and write the compressed data to the memory location.
9. The system of claim 8, wherein the one or more logical memory devices comprise multiple logical memory devices.
10. The system of claim 8, wherein the controller is further configured to: receive an additional command to write additional data to an additional memory location; and cause the additional data

to be written to the additional memory location without using the lossy compression operation, responsive to the additional memory location corresponding to an additional logical memory device, of the one or more logical memory devices, that is configured with lossy compression disabled.

11. The system of claim 8, wherein the controller is further configured to: receive an additional command to read from the memory location; decompress the compressed data using a decompression operation to obtain decompressed data; and output the decompressed data.

12. The system of claim 8, wherein the controller is further configured to: determine an error value for the compressed data, wherein the controller is configured to cause the compressed data to be written to the memory location responsive to the error value satisfying a threshold.

13. The system of claim 12, wherein the controller, to determine the error value, is configured to: decompress the compressed data to obtain decompressed data; and compare the decompressed data to the data, wherein the error value is in accordance with a difference between the decompressed data and the data.

14. The system of claim 8, wherein the compression memory system is a Type 3 Compute Express Link (CXL) device.

15. A method, comprising: receiving, by a controller of a memory system, and from a host device, a command to write data to a memory location; compressing, by the controller, the data using a lossy compression operation to obtain compressed data; and causing, by the controller, the compressed data to be written to the memory location.

16. The method of claim 15, wherein the memory system comprises one or more memory components configured as one or more logical memory devices.

17. The method of claim 16, wherein the data is compressed and responsive to the memory location corresponding to a logical memory device, of the one or more logical memory devices, that is configured with lossy compression enabled.

18. The method of claim 16, further comprising: receiving an additional command to write additional data to an additional memory location; and causing the additional data to be written to the additional memory location without using the lossy compression operation, responsive to the additional memory location corresponding to an additional logical memory device, of the one or more logical memory devices, that is configured with lossy compression disabled.

19. The method of claim 16, wherein the one or more logical memory devices comprise multiple logical memory devices mapped to multiple logical ports.

20. The method of claim 15, further comprising: receiving an additional command to read from the memory location; decompressing the compressed data using a decompression operation to obtain decompressed data; and outputting the decompressed data.
