

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250258710

Kind Code

A1

Publication Date

August 14, 2025

Inventor(s)

SUN; Xiaoyu et al.

ARTIFICIAL INTELLIGENCE ACCELERATOR DEVICE

Abstract

An artificial intelligence (AI) accelerator device may include a plurality of on-chip mini buffers that are associated with a processing element (PE) array. Each mini buffer is associated with a subset of rows or a subset of columns of the PE array. Partitioning an on-chip buffer of the AI accelerator device into the mini buffers described herein may reduce the size and complexity of the on-chip buffer. The reduced size of the on-chip buffer may reduce the wire routing complexity of the on-chip buffer, which may reduce latency and may reduce access energy for the AI accelerator device. This may increase the operating efficiency and/or may increase the performance of the AI accelerator device. Moreover, the mini buffers may increase the overall bandwidth that is available for the mini buffers to transfer data to and from the PE array.

Inventors: SUN; Xiaoyu (San Jose, CA), PENG; Xiaochen (San Jose, CA), AKARVARDAR; Murat Kerem (Hsinchu, TW)

Applicant: Taiwan Semiconductor Manufacturing Co., Ltd. (Hsinchu, TW)

Family ID: 1000008571886

Appl. No.: 19/170096

Filed: April 04, 2025

Related U.S. Application Data

parent US continuation 17900471 20220831 parent-grant-document US 12293229 child US 19170096

Publication Classification

Int. Cl.: G06F9/50 (20060101); G06F7/50 (20060101); G06F15/80 (20060101)

U.S. Cl.:

Background/Summary

RELATED APPLICATION [0001] This application is a continuation of U.S. patent application Ser. No. 17/900,471, filed Aug. 31, 2022, which is incorporated herein by reference in its entirety.

BACKGROUND

[0002] An artificial intelligence (AI) accelerator device is a class of specialized hardware that may be used to accelerate machine learning workloads, which are typically modeled as a neural network that involves large-scale memory accesses and highly-parallelized computations. An AI accelerator device may be implemented in one or more types of hardware, such as a central processing unit (CPU), a graphics processing unit (GPU), a field programmable gate array (FPGA), and/or an application-specific integrated circuit (ASIC), among other examples.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0003] Aspects of the present disclosure are best understood from the following detailed description when read with the accompanying figures. It is noted that, in accordance with the standard practice in the industry, various features are not drawn to scale. In fact, the dimensions of the various features may be arbitrarily increased or reduced for clarity of discussion.

[0004] FIG. 1 is a diagram of an example environment in which systems and/or methods described herein may be implemented.

[0005] FIGS. 2-9 are diagrams of example implementations of an artificial intelligence (AI) accelerator device described herein.

[0006] FIGS. 10A-10C are diagrams of an example implementation described herein.

[0007] FIG. 11 is a diagram of an example implementation of a processing element circuit described herein.

[0008] FIG. 12 is a diagram of an example neural network described herein.

[0009] FIGS. 13A-13C are diagrams of example implementations of neural network dataflows described herein.

[0010] FIG. 14 is a diagram of an example implementation of energy consumption in an AI accelerator device described herein.

[0011] FIGS. 15A and 15B are diagrams of an example implementations of design parameters for an AI accelerator device described herein.

[0012] FIG. 16 is a diagram of example components of a device described herein.

[0013] FIG. 17 is a flowchart of an example process associated with operation of an AI accelerator device described herein.

DETAILED DESCRIPTION

[0014] The following disclosure provides many different embodiments, or examples, for implementing different features of the provided subject matter. Specific examples of components and arrangements are described below to simplify the present disclosure. These are, of course, merely examples and are not intended to be limiting. For example, the formation of a first feature over or on a second feature in the description that follows may include embodiments in which the first and second features are formed in direct contact, and may also include embodiments in which additional features may be formed between the first and second features, such that the first and second features may not be in direct contact. In addition, the present disclosure may repeat

reference numerals and/or letters in the various examples. This repetition is for the purpose of simplicity and clarity and does not in itself dictate a relationship between the various embodiments and/or configurations discussed.

[0015] An application specific integrated circuit (ASIC)-based AI accelerator device is a type of hardware device that may include a processing element (PE) array that includes a plurality of PEs spatially and/or temporally arranged to perform a particular type of operation, such as a multiply-and-accumulate (MAC) operation in a neural network. A MAC operation is a type of weighted sum operation in which a multiply operation is performed on two or more inputs, and the output from the multiply operation is accumulated with a partial sum in an accumulate operation in order to generate an updated partial sum.

[0016] The PE array of an ASIC-based AI accelerator device may be connected with relatively large monolithic on-chip buffers that include memory devices such as static random access memory (SRAM) memory devices. An on-chip buffer may enable data (e.g., weights, activations, partial sums) to be stored locally near the PE array to reduce latency in operation of the ASIC-based AI accelerator device and to reduce the need for frequently accessing slower dynamic random access memory (DRAM) of the ASIC-based AI accelerator device. However, as the size of neural networks implemented by ASIC-based AI accelerator devices increases, so too do the size and complexity of the associated on-chip buffers.

[0017] Larger monolithic on-chip buffers may suffer from increased latency and/or greater access energy due to the increased wire routing complexity in the on-chip buffers, which may reduce the operating efficiency and may reduce the performance of an ASIC-based AI accelerator that includes the monolithic on-chip buffers. Moreover, larger and more complex monolithic on-chip buffers may need greater memory bandwidth to transfer data to and from the PE array of an ASIC-based AI accelerator. If the memory bandwidth between the monolithic on-chip buffers and the PE array is limited, the monolithic on-chip buffers may not be capable of transferring a sufficient amount of data to and from the PE array, which may result in underutilization of the PEs in the PE array and reduced performance for the ASIC-based AI accelerator.

[0018] Some implementations described herein include an ASIC-based AI accelerator device in which on-chip buffers associated with a PE array are partitioned into a plurality of mini buffers (e.g., mini input buffers, mini weight buffers, and/or mini accumulator buffers). Each mini buffer is associated with a subset of rows and a subset of columns of the PE array. The ASIC-based AI accelerator device may include a distributor circuit to direct data to the appropriate mini buffers of the ASIC-based AI accelerator device. Partitioning the on-chip buffers of the ASIC-based AI accelerator device into the mini buffers described herein may reduce the size of the on-chip buffers of the ASIC-based AI accelerator device. The reduced size of the on-chip buffers may reduce the wire routing complexity of the on-chip buffers, which may reduce latency and reduce access energy for the ASIC-based AI accelerator device. This may increase the operating efficiency and/or may increase the performance of the ASIC-based AI accelerator device. Moreover, connecting the mini buffers of the ASIC-based AI accelerator device to subsets of rows and columns of the PE array of the ASIC-based AI accelerator device may increase the overall bandwidth that is available for the mini buffers to transfer data to and from the PE array, which may increase the utilization of the PEs in the PE array and may increase performance for the ASIC-based AI accelerator device.

[0019] In this way, an ASIC-based AI accelerator device may include on-chip buffers associated with a PE array are partitioned into a plurality of mini buffers (e.g., mini input buffers, mini weight buffers, and/or mini accumulator buffers). Each mini buffer is associated with a subset of rows and a subset of columns of the PE array. The ASIC-based AI accelerator device may include a distributor circuit to direct data to the appropriate mini buffers of the ASIC-based AI accelerator device. Partitioning the on-chip buffers of the ASIC-based AI accelerator device into the mini buffers described herein may reduce the size of the on-chip buffers of the ASIC-based AI accelerator device. The reduced size of the on-chip buffers may reduce the wire routing complexity

of the on-chip buffers, which may reduce latency and reduce access energy for the ASIC-based AI accelerator device. This may increase the operating efficiency and/or may increase the performance of the ASIC-based AI accelerator device. Moreover, connecting the mini buffers of the ASIC-based AI accelerator device to subsets of rows and columns of the PE array of the ASIC-based AI accelerator device may increase the overall bandwidth that is available for the mini buffers to transfer data to and from the PE array, which may increase the utilization of the PEs in the PE array and may increase performance for the ASIC-based AI accelerator device.

[0020] FIG. 1 is a diagram of an example environment **100** in which systems and/or methods described herein may be implemented. As shown in FIG. 1, environment **100** may include an edge computing system **101**, which may include one or more elements of and/or may execute within a cloud computing system **102**. The cloud computing system **102** may include one or more elements **103-113**, as described in more detail below. As further shown in FIG. 1, environment **100** may include a network **120** and one or more client devices **130-170**. Devices and/or elements of environment **100** may interconnect via wired connections and/or wireless connections.

[0021] The cloud computing system **102** includes computing hardware **103**, a resource management component **104**, a host operating system (OS) **105**, and/or one or more virtual computing systems **106**. The resource management component **104** may perform virtualization (e.g., abstraction) of computing hardware **103** to create the one or more virtual computing systems **106**. Using virtualization, the resource management component **104** enables a single computing device (e.g., a computer or a server) to operate like multiple computing devices, such as by creating multiple isolated virtual computing systems **106** from computing hardware **103** of the single computing device. In this way, computing hardware **103** can operate more efficiently, with lower power consumption, higher reliability, higher availability, higher utilization, greater flexibility, and lower cost than using separate computing devices.

[0022] Computing hardware **103** includes hardware and corresponding resources from one or more computing devices. For example, computing hardware **103** may include hardware from a single computing device (e.g., a single server) or from multiple computing devices (e.g., multiple servers), such as multiple computing devices in one or more data centers. As shown, computing hardware **103** may include one or more processors **107**, one or more memories **108**, one or more networking components **109**, and one or more AI accelerator devices **110**. Examples of a processor, a memory, and a networking component (e.g., a communication component) are described in connection with FIG. 16.

[0023] An AI accelerator device **110** may include a hardware device that includes application-specific circuitry configured to perform operations associated with a neural network (e.g., an artificial neural network or a deep neural network). The AI accelerator device **110** may include an ASIC device in which hardware (e.g., drivers, transistors, repeaters, complementary metal oxide semiconductor (CMOS) logic) is specifically configured to implement nodes (or artificial neurons) of a neural network to perform AI operations, and therefore is referred to as a hardware accelerator.

[0024] The edge computing system **101** may perform one or more operations and/or processes associated with AI and/or neural network processing. In particular, the edge computing system **101** may use the AI accelerator device(s) **110** to perform AI-specific computing tasks and workloads. An example AI-specific computing task may include receiving an input (e.g., an image) from a client device of the client devices **130-170** through the network **120** and generating, using the AI accelerator device(s) **110**, an output (e.g., an inference or a classification) based on the input. Here, the AI-specific computing task may involve identifying objects in the image, tagging the image with one or more tags, classifying the content of the image, and/or performing an image search based on the image, among other examples.

[0025] The generation of the output may involve the use of a multiply-and-accumulate (MAC) operation, in which inputs are multiplied with weights and accumulated with a partial sum to generate an updated partial sum. In some implementations, a node of an AI accelerator device **110**

may have a threshold such that the updated partial sum is provided only if the updated partial sum satisfies a threshold. This is referred to as “firing” of the node.

[0026] The weights applied by the nodes of an AI accelerator device **110** may be adjusted or “trained” using machine learning. This involves supervised and/or unsupervised training, in which large data sets are provided to the nodes of an AI accelerator device **110** to “learn” or form connections between inputs and outputs of the large data sets.

[0027] The resource management component **104** includes a virtualization application (e.g., executing on hardware, such as computing hardware **103**) capable of virtualizing computing hardware **103** to start, stop, and/or manage one or more virtual computing systems **106**. For example, the resource management component **104** may include a hypervisor (e.g., a bare-metal or Type **1** hypervisor, a hosted or Type **2** hypervisor, or another type of hypervisor) or a virtual machine monitor, such as when the virtual computing systems **106** are virtual machines. Additionally, or alternatively, the resource management component **104** may include a container manager, such as when the virtual computing systems **106** are containers **112**. In some implementations, the resource management component **104** executes within and/or in coordination with a host operating system **105**.

[0028] A virtual computing system **106** includes a virtual environment that enables cloud-based execution of operations and/or processes described herein using computing hardware **103**. As shown, a virtual computing system **106** may include a virtual machine **111**, a container **112**, or a hybrid environment **113** that includes a virtual machine and a container, among other examples. A virtual computing system **106** may execute one or more applications using a file system that includes binary files, software libraries, and/or other resources required to execute applications on a guest operating system (e.g., within the virtual computing system **106**) or the host operating system **105**.

[0029] Although the edge computing system **101** may include one or more elements **103-113** of the cloud computing system **102**, may execute within the cloud computing system **102**, and/or may be hosted within the cloud computing system **102**, in some implementations, the edge computing system **101** may not be cloud-based (e.g., may be implemented outside of a cloud computing system) or may be partially cloud-based. For example, the edge computing system **101** may include one or more devices that are not part of the cloud computing system **102**, such as device **1600** of FIG. **16**, which may include a standalone server or another type of computing device.

[0030] Network **120** includes one or more wired and/or wireless networks. For example, network **120** may include a cellular network, a public land mobile network (PLMN), a local area network (LAN), a wide area network (WAN), a private network, the Internet, and/or a combination of these or other types of networks. The network **120** enables communication among the devices of environment **100**.

[0031] The client devices **130-170** may include one or more devices capable of receiving, generating, storing, processing, and/or providing information. The client device **130-170** may include a communication device and/or a computing device. For example, the client device **130-170** may include a wireless communication device, a mobile phone, a user equipment, a laptop computer, a tablet computer, a desktop computer, a wearable communication device (e.g., a smart wristwatch, a pair of smart eyeglasses, a head mounted display, or a virtual reality headset), or a similar type of device.

[0032] The number and arrangement of devices and networks shown in FIG. **1** are provided as an example. In practice, there may be additional devices and/or networks, fewer devices and/or networks, different devices and/or networks, or differently arranged devices and/or networks than those shown in FIG. **1**. Furthermore, two or more devices shown in FIG. **1** may be implemented within a single device, or a single device shown in FIG. **1** may be implemented as multiple, distributed devices. Additionally, or alternatively, a set of devices (e.g., one or more devices) of environment **100** may perform one or more functions described as being performed by another set

of devices of environment **100**.

[0033] FIG. 2 is a diagram of an example implementation **200** of an AI accelerator device **110** described herein. The AI accelerator device **110** may include an ASIC-based AI accelerator device (e.g., an AI accelerator implemented in an application specific integrated circuit) or another type of AI hardware accelerator. The AI accelerator device **110** may be configured to perform artificial neural network and/or deep neural network operations, machine learning operations (e.g., training on large data sets to adjust the weights of the artificial neural network and/or deep neural network), and/or to perform other AI operations. The AI accelerator device **110** may include components **202-222**, which may be included on a same semiconductor die or on a same packaged semiconductor device.

[0034] As shown in FIG. 2, in the example implementation **200**, the AI accelerator device **110** may include a processing element array **202** that includes a plurality of processing element circuits **204a-204n** arranged in a grid or in a matrix (e.g., a systolic array, a vector engine arrangement). In particular, the processing element circuits **204a-204n** are arranged into a plurality of columns **206a-206l** and a plurality of rows **208a-208m**. For example, processing element circuits **204a**, **204d**, and **204g** may be included in column **206a**, processing element circuits **204b**, **204e**, and **204h** may be included in column **206b**, processing element circuits **204c**, **204f**, and **204n** may be included in column **206l**, and so on. As another example, processing element circuits **204a-204c** may be included in row **208a**, processing element circuits **204d-204f** may be included in row **208b**, processing element circuits **204g-204n** may be included in row **208m**, and so on.

[0035] In some implementations, the processing element array **202** includes a systolic array, in which computation is fully pipelined, and data is provided to the edges of the processing element array **202** and permitted to propagate through the processing element array **202** in a wave-like fashion through processing element interconnect. This type of processing element array configuration may increase and/or maximize data reuse within the processing element array **202**. In some implementations, the processing element array **202** includes a vector engine, in which one operand (e.g., a multiplication operand, an accumulation operand) is stationary in the processing element circuits **204a-204n**, and the other operand is multicasted to the same rows and/or columns. The partial sums are collected in a combinational manner through an adder tree.

[0036] A processing element circuit may include circuitry that is configured to perform a neural network operation. For example, each of the processing element circuits **204a-204n** may be configured to perform a multiply-and-accumulate (MAC) operation. Accordingly, each of the processing element circuits **204a-204n** may include multiplication circuitry, accumulator circuitry, input circuitry, output circuitry, and routing between circuit components (e.g., metallization lines), among other examples. In some implementations, the processing element circuits **204a-204n** are connected such that the output of one processing element is provided as input to one or more other processing element circuits in the processing element array **202**. The processing element circuits **204a-204n** are arranged in a grid or in a matrix to enable the processing element array **202** to process highly parallelized computational tasks in an efficient manner.

[0037] As further shown in FIG. 2, in the example implementation **200**, the AI accelerator device **110** may include a plurality of weight buffers (e.g., a plurality of mini weight buffers) **210a-210k**. The weight buffers **210a-210k** may be coupled with a weight buffer multiplexer circuit **212** and may be configured to buffer weight data (e.g., weights for the processing element circuits **204a-204n**) that is to be provided from a memory device **214** (e.g., a DRAM memory device or another type of memory device) to the processing element circuits **204a-204n** of the processing element array **202**.

[0038] Each of the weight buffers **210a-210k** may include one or more SRAM devices (or another type of memory devices having faster access times relative to the memory device **214**, which may be a magnetic random access memory (MRAM), resistive random access memory (RRAM), phase change material random access memory (PCM-RAM), and/or ferroelectric random access memory

(FeRAM)), drivers, inverters, repeaters, and associated routing (e.g., metallization lines) that connects the weight buffers **210a-210k** to the processing element array **202**. The use of the weight buffers **210a-210k**, as opposed to direct memory access to the memory device **214**, reduces access times and latency for providing weights to the processing element circuits **204a-204n**. The weight buffer multiplexer (MUX) circuit **212** may include logic circuitry (e.g., AND gates, OR gates, drivers, inverters, repeaters) that is configured to enable selection of specific ones of the weight buffers **210a-210k** and to enable weights to provide increased flexibility in writing data to different weight buffers **210a-210k**.

[0039] Each of the weight buffers **210a-210k** may be associated with, and may provide weights to, a subset of the columns **206a-206l** of processing element circuits **204a-204n** included in the processing element array **202**. In some implementations, each of the weight buffers **210a-210k** may be associated with, and may provide weights to, a single respective column of the columns **206a-206l**. For example, the weight buffer **210a** may be associated with, and may provide weights to, the column **206a**; the weight buffer **210b** may be associated with, and may receive provide weights to, the column **206b**; and so on. In some implementations, each of the weight buffers **210a-210k** may be associated with, and may provide weights to, a plurality of columns of the columns **206a-206l**.

[0040] Including a plurality of weight buffers (e.g., a plurality of mini weight buffers) in the AI accelerator device **110** may reduce the latency and energy consumption of the weight buffers of the AI accelerator device **110** relative to the use of a single monolithic weight buffer. The reduced latency and energy consumption is achieved through the reduced complexity of each of the weight buffers **210a-210k**. In particular, including plurality of weight buffers (e.g., a plurality of mini weight buffers) reduces the size and complexity of routing between SRAM devices of the weight buffers **210a-210k** and reduces the size and complexity of routing between the SRAM devices and the processing element array **202**. The reduced size and complexity of routing in the weight buffers **210a-210k** and between the weight buffers **210a-210k** and the processing element array **202** reduces the propagation distance of weight data in the weight buffers **210a-210k** and between the weight buffers **210a-210k** and the processing element array **202**. The reduced propagation distance results in faster propagation times (thus, lower latency), increased available bandwidth, and/or reduced access energy.

[0041] The quantity of weight buffers **210a-210k** included in the AI accelerator device **110**, and the quantity of the columns **206a-206l** associated with each of the weight buffers **210a-210k**, may be selected to satisfy one or more parameters, such as energy consumption of the weight buffers **210a-210k**, latency of the weight buffers **210a-210k**, and/or the physical footprint of the weight buffers **210a-210k**. For example, increasing the quantity of weight buffers **210a-210k** included in the AI accelerator device **110** (which decreases the quantity of the columns **206a-206l** associated with each of the weight buffers **210a-210k**) may decrease latency and energy consumption in the AI accelerator device **110** at the expense of increased size of the AI accelerator device **110** (e.g., because of the additional isolation needed between adjacent mini weight buffers).

[0042] As another example, decreasing the quantity of weight buffers **210a-210k** included in the AI accelerator device **110** (which increases the quantity of the columns **206a-206l** associated with each of the weight buffers **210a-210k**) may decrease the size of the AI accelerator device **110** at the expense of increased latency and energy consumption in the AI accelerator device **110**.

[0043] In some cases, a design constraint may be used to determine the quantity of the columns **206a-206l** that is to be associated with each of the weight buffers **210a-210k**. For example, if the quantity of weight buffers **210a-210k** is constrained to N, the quantity of the columns **206a-206l** that is to be associated with each of the weight buffers **210a-210k** may be determined as:

$$[00001] S = \frac{L}{N}$$

where S corresponds to the quantity of columns **206a-206l** per weight buffer, and L corresponds to the quantity of columns **206a-206l** in the processing element array **202**.

[0044] As further shown in FIG. 2, in the example implementation **200**, the AI accelerator device

110 may include a distributor circuit **216** that is coupled with the weight buffer multiplexer circuit **212**. The distributor circuit **216** may include circuitry (e.g., state-machine based controller circuitry) that is configured to provide one or more select signals (e.g., weight select signals) to the weight buffer multiplexer circuit **212**. The distributor circuit **216** may provide select signal(s) to the weight buffer multiplexer circuit **212** to select specific ones of the weight buffers **210a-210k**. In this way, the weight buffer multiplexer circuit **212** enables the distributor circuit **216** to independently control the weight buffers **210a-210k** to provide weights to specific columns of the processing element array **202**.

[0045] A single monolithic activation buffer **218** is included and is associated with the rows **208a-208m** of the processing element array **202**. The single monolithic activation buffer **218** is configured to provide input data to processing element circuits **204a-204n** in the rows **208a-208m**. The input data may include partial sums, values that are to be multiplied, and/or other input data. The single monolithic activation buffer **218** may include one or more SRAM devices, drivers, inverters, repeaters, and associated routing (e.g., metallization lines) that connects the single monolithic activation buffer **218** to the processing element array **202**.

[0046] A single monolithic accumulator buffer **220** is included and is associated with the columns **206a-206l** of the processing element array **202**. The single monolithic accumulator buffer **220** is configured to receive output data (e.g., updated partial sums) from the processing element circuits **204a-204n** in the columns **206a-206l**. The single monolithic accumulator buffer **220** may include one or more SRAM devices, drivers, inverters, repeaters, and associated routing (e.g., metallization lines) that connects the single monolithic accumulator buffer **220** to the processing element array **202**.

[0047] The single monolithic accumulator buffer **220** may be further configured to provide the updated partial sums to peripheral circuitry **222**, which may include circuitry (e.g., DRAM, logic circuitry, routing) that is configured to store and/or provide the updated partial sums to the single monolithic activation buffer **218** for use in subsequent neural network operations.

[0048] As indicated above, FIG. 2 is provided as an example. Other examples may differ from what is described with regard to FIG. 2.

[0049] FIG. 3 is a diagram of an example implementation **300** of an AI accelerator device **110** described herein. The example implementation **300** of the AI accelerator device **110** is similar to the example implementation **200** of the AI accelerator device **110**, and includes components that are similar to the components of the example implementation **200** of the AI accelerator device **110**. However, in the example implementation **300** of the AI accelerator device **110**, a single monolithic weight buffer is included as opposed to a plurality of weight buffers (e.g., a plurality of mini weight buffers), and therefore a weight multiplexer circuit may be omitted from the AI accelerator device **110** in the example implementation **300**. Moreover, a plurality of activation buffers (e.g., a plurality of mini activation buffers) are included as opposed to a single monolithic activation buffer. The plurality of activation buffers (e.g., a plurality of mini activation buffers) may be included to reduce the latency and energy consumption of the activation buffer of the AI accelerator device **110** relative to the latency and energy consumption of a single monolithic activation buffer.

[0050] As shown in FIG. 3, in the example implementation **300**, the AI accelerator device **110** may include a processing element array **302** that includes a plurality of processing element circuits **304a-304n** arranged in a grid or in a matrix. In particular, the processing element circuits **304a-304n** are arranged into a plurality of columns **306a-306l** and a plurality of rows **308a-308m**. The single monolithic weight buffer **310** is associated with the plurality of columns **306a-306l**, and is associated with a memory device **314** that is configured to store the weights for the processing element circuits **304a-304n**.

[0051] As shown in FIG. 3, in the example implementation **300**, the AI accelerator device **110** may include a distributor circuit **316**, a plurality of activation buffers **318a-318x**, a single monolithic accumulator circuit **320**, and peripheral circuitry **322**. Each activation buffer **318a-318x** may

include an independent circuit (e.g., independent from the other activation buffer **318a-318x**), which may reduce the complexity of the metallization routing and layout for the activation buffers **318a-318x**. Each of the activation buffers **318a-318x** may be associated with, and may provide input data to, a subset of the rows **308a-308m** of processing element circuits **304a-304n** included in the processing element array **302**. In some implementations, each of the activation buffers **318a-318x** may be associated with, and may provide input data to, a single respective row of the rows **308a-308m**. For example, the activation buffer **318a** may be associated with, and may provide input data to, the row **308a**; the activation buffer **318b** may be associated with, and may provide input data to, the row **308b**; and so on. In some implementations, each of the activation buffers **318a-318x** may be associated with, and may provide input data to, a plurality of rows of the rows **308a-308m**.

[0052] Including a plurality of activation buffers (e.g., a plurality of mini activation buffers) in the AI accelerator device **110** may reduce the latency and energy consumption of the activation buffer of the AI accelerator device **110** relative to the use of a single monolithic activation buffer. The reduced latency and energy consumption is achieved through the reduced complexity of each of the activation buffers **318a-318x**. In particular, including a plurality of activation buffers (e.g., a plurality of mini activation buffers) reduces the size and complexity of routing between SRAM devices of the activation buffers **318a-318x** and reduces the size and complexity of routing between the SRAM devices and the processing element array **302**. The reduced size and complexity of routing in the activation buffers **318a-318x** and between the activation buffers **318a-318x** and the processing element array **302** reduces the propagation distance of input data in the activation buffers **318a-318x** and between the activation buffers **318a-318x** and the processing element array **302**. The reduced propagation distance results in faster propagation times (thus, lower latency), increased available bandwidth, and/or reduced access energy.

[0053] The quantity of activation buffers **318a-318x** included in the AI accelerator device **110**, and the quantity of the rows **308a-308m** associated with each activation buffer, may be selected to satisfy one or more parameters, such as energy consumption of the activation buffers **318a-318x**, latency of the activation buffers **318a-318x**, and/or the physical footprint of the activation buffers **318a-318x**. For example, increasing the quantity of activation buffers **318a-318x** included in the AI accelerator device **110** (which decreases the quantity of the rows **308a-308m** associated with each of the activation buffers **318a-318x**) may decrease latency and energy consumption in the AI accelerator device **110** at the expense of increased size of the AI accelerator device **110**.

[0054] As another example, decreasing the quantity of activation buffers **318a-318x** included in the AI accelerator device **110** (which increases the quantity of the rows **308a-308m** associated with each of the activation buffers **318a-318x**) may decrease the size of the AI accelerator device **110** at the expense of increased latency and energy consumption in the AI accelerator device **110**.

[0055] In some cases, a design constraint may be used to determine the quantity of the rows **308a-308m** that is to be associated with each of the activation buffers **318a-318x**. For example, if the quantity of activation buffers **318a-318x** is constrained to N, the quantity of the rows **308a-308m** that is to be associated with each of the activation buffers **318a-318x** may be determined as:

$$[00002] S = \frac{L}{N}$$

where S corresponds to the quantity of rows **308a-308m** per activation buffer, and L corresponds to the quantity of rows **308a-308m** in the processing element array **302**.

[0056] As shown in FIG. 3, in the example implementation **300**, the AI accelerator device **110** may include an activation buffer multiplexer circuit **324** coupled with the distributor circuit **316** and the activation buffers **318a-318x**. The distributor circuit **316** may be configured to provide one or more select signals (e.g., activation select signals) to the activation buffer multiplexer circuit **324**. The distributor circuit **316** may provide select signal(s) to the activation buffer multiplexer circuit **324** to select specific ones of the activation buffers **318a-318x**. In this way, the activation buffer multiplexer circuit **324** enables distributor circuit **316** to independently control the activation

buffers **318a-318x** to provide input data to specific rows of the processing element array **302**.
[0057] The single monolithic accumulator circuit **320** may be further configured to provide the updated partial sums to the peripheral circuitry **322**, which may include circuitry (e.g., DRAM, logic circuitry, routing) that is configured to store and/or provide the updated partial sums to the activation buffer multiplexer circuit **324** for use in subsequent neural network operations.
[0058] As indicated above, FIG. 3 is provided as an example. Other examples may differ from what is described with regard to FIG. 3.

[0059] FIG. 4 is a diagram of an example implementation **400** of an AI accelerator device **110** described herein. The example implementation **400** of the AI accelerator device **110** is similar to the example implementation **200** of the AI accelerator device **110**, and includes components that are similar to the components of the example implementation **200** of the AI accelerator device **110**. However, in the example implementation **400** of the AI accelerator device **110**, a single monolithic weight buffer is included as opposed to a plurality of weight buffers **210a-210k** (e.g., a plurality of mini weight buffers). Accordingly, a weight buffer multiplexer circuit and a distributor circuit may be omitted from the AI accelerator device **110** in the example implementation **400**. Moreover, a plurality of accumulator buffers (e.g., a plurality of mini accumulator buffers) are included as opposed to a single monolithic accumulator buffer. The plurality of accumulator buffers (e.g., a plurality of mini accumulator buffers) may be included to reduce the latency and energy consumption of the accumulator buffer of the AI accelerator device **110** relative to the latency and energy consumption of a single monolithic accumulator buffer.

[0060] As shown in FIG. 4, in the example implementation **400**, the AI accelerator device **110** may include a processing element array **402** that includes a plurality of processing element circuits **404a-404n** arranged in a grid or in a matrix. In particular, the processing element circuits **404a-404n** are arranged into a plurality of columns **406a-406l** and a plurality of rows **408a-408m**. A single monolithic weight buffer **410** is associated with the plurality of columns **406a-406l**, and is associated with a memory device **414** that is configured to store the weights for the processing element circuits **404a-404n**. A single monolithic activation buffer **418** is included and is associated with the rows **408a-408m** of the processing element array **402**. The single monolithic activation buffer **418** is configured to provide input data to processing element circuits **404a-404n** in the rows **408a-408m**.

[0061] As shown in FIG. 4, in the example implementation **400**, the AI accelerator device **110** may include a plurality of accumulator buffers **420a-420y** (e.g., a plurality of mini accumulator buffers) coupled with the processing element array **402** and peripheral circuitry **422**. Each of the accumulator buffers **420a-420y** may be associated with, and may receive output data (e.g., a partial sum) from, a subset of the columns **406a-406l** of processing element circuits **404a-404n** included in the processing element array **402**. In some implementations, each of the accumulator buffers **420a-420y** may be associated with, and may receive an output from, a single respective column of the columns **406a-406l**. For example, the accumulator buffer **420a** may be associated with, and may receive an output from, the column **406a**; the accumulator buffer **420b** may be associated with, and may receive an output from, the column **406b**; and so on. In some implementations, each of the accumulator buffers **420a-420y** may be associated with, and may receive an output from, a plurality of columns of the columns **406a-406l**. The accumulator buffers **420a-420y** may be configured to provide outputs to the peripheral circuitry **422**, which provides the outputs to the single monolithic activation buffer **418**.

[0062] Including a plurality of accumulator buffers (e.g., a plurality of mini accumulator buffers) in the AI accelerator device **110** may reduce the latency and energy consumption of the accumulator buffer of the AI accelerator device **110** relative to the use of a single monolithic accumulator buffer. The reduced latency and energy consumption is achieved through the reduced complexity of each of the accumulator buffers **420a-420y**. In particular, including plurality of accumulator buffers (e.g., a plurality of mini accumulator buffers) reduces the size and complexity of routing between

SRAM devices of the accumulator buffers **420a-420y** and reduces the size and complexity of routing between the SRAM devices and the processing element array **402**. The reduced size and complexity of routing in the accumulator buffers **420a-420y** and between the accumulator buffers **420a-420y** and the processing element array **402** reduces the propagation distance of output data in the accumulator buffers **420a-420y** and between the accumulator buffers **420a-420y** and the processing element array **402**. The reduced propagation distance results in faster propagation times (thus, lower latency), increased available bandwidth, and/or reduced access energy.

[0063] The quantity of accumulator buffers **420a-420y** included in the AI accelerator device **110**, and the quantity of the columns **406a-406l** associated with each activation buffer, may be selected to satisfy one or more parameters, such as energy consumption of the accumulator buffers **420a-420y**, latency of the accumulator buffers **420a-420y**, and/or the physical footprint of the accumulator buffers **420a-420y**. For example, increasing the quantity of accumulator buffers **420a-420y** included in the AI accelerator device **110** (which decreases the quantity of the columns **406a-406l** associated with each of the accumulator buffers **420a-420y**) may decrease latency and energy consumption in the AI accelerator device **110** at the expense of increased size of the AI accelerator device **110**.

[0064] As another example, decreasing the quantity of accumulator buffers **420a-420y** included in the AI accelerator device **110** (which increases the quantity of the columns **406a-406l** associated with each of the accumulator buffers **420a-420y**) may decrease the size of the AI accelerator device **110** at the expense of increased latency and energy consumption in the AI accelerator device **110**.

[0065] In some cases, a design constraint may be used to determine the quantity of the columns **406a-406l** that is to be associated with each of the accumulator buffers **420a-420y**. For example, if the quantity of accumulator buffers **420a-420y** is constrained to N, the quantity of the columns **406a-406l** that is to be associated with each of the accumulator buffers **420a-420y** may be determined as:

$$[00003] S = \frac{L}{N}$$

where S corresponds to the quantity of columns **406a-406l** per accumulator buffer, and L corresponds to the quantity of columns **406a-406l** in the processing element array **402**.

[0066] As indicated above, FIG. 4 is provided as an example. Other examples may differ from what is described with regard to FIG. 4.

[0067] FIG. 5 is a diagram of an example implementation **500** of an AI accelerator device **110** described herein. The example implementation **500** of the AI accelerator device **110** is similar to the example implementation **200** of the AI accelerator device **110**, and includes components **502-522** that are similar to the components **202-222** of the example implementation **200** of the AI accelerator device **110**. However, in the example implementation **500** of the AI accelerator device **110**, a plurality of weight buffers **510a-510k** (e.g., a plurality of mini weight buffers) and a plurality of activation buffers **518a-518x** (e.g., a plurality of mini activation buffers) are both included. Accordingly, the AI accelerator device **110** in the example implementation **500** may further include an activation buffer multiplexer circuit **524** coupled with the distributor circuit **516** and the activation buffers **518a-518x**. The plurality of weight buffers **510a-510k** (e.g., a plurality of mini weight buffers) and the plurality of activation buffers **518a-518x** (e.g., a plurality of mini activation buffers) may be included to reduce the latency and energy consumption of both the weight buffer and the activation buffer of the AI accelerator device **110** relative to the latency and energy consumption of a single monolithic weight buffer and a single monolithic activation buffer.

[0068] Each of the weight buffers **510a-510k** may be associated with, and may provide weights to, a subset of the columns **506a-506l** of processing element circuits **504a-504n** included in the processing element array **502**. In some implementations, each of the weight buffers **510a-510k** may be associated with, and may provide weights to, a single respective column of the columns **506a-506l**. For example, the weight buffer **510a** may be associated with, and may provide weights to, the column **506a**; the weight buffer **510b** may be associated with, and may receive provide weights to,

the column **506b**; and so on. In some implementations, each of the weight buffers **510a-510k** may be associated with, and may provide weights to, a plurality of columns of the columns **506a-506l**. [0069] The quantity of weight buffers **510a-510k** included in the AI accelerator device **110**, and the quantity of the columns **506a-506l** associated with each weight buffer, may be selected to satisfy one or more parameters, such as energy consumption of the weight buffers **510a-510k**, latency of the weight buffers **510a-510k**, and/or the physical footprint of the weight buffers **510a-510k**. For example, increasing the quantity of weight buffers **510a-510k** included in the AI accelerator device **110** (which decreases the quantity of the columns **506a-506l** associated with each of the weight buffers **510a-510k**) may decrease latency and energy consumption in the AI accelerator device **110** at the expense of increased size of the AI accelerator device **110**. As another example, decreasing the quantity of weight buffers **510a-510k** included in the AI accelerator device **110** (which increases the quantity of the columns **506a-506l** associated with each of the weight buffers **510a-510k**) may decrease the size of the AI accelerator device **110** at the expense of increased latency and energy consumption in the AI accelerator device **110**.

[0070] Each of the activation buffers **518a-518x** may be associated with, and may provide input data to, a subset of the rows **508a-508m** of processing element circuits **504a-504n** included in the processing element array **502**. In some implementations, each of the activation buffers **518a-518x** may be associated with, and may provide input data to, a single respective row of the rows **508a-508m**. For example, the activation buffer **518a** may be associated with, and may provide input data to, the row **508a**; the activation buffer **518b** may be associated with, and may provide input data to, the row **508b**; and so on. In some implementations, each of the activation buffers **518a-518x** may be associated with, and may provide input data to, a plurality of rows of the rows **508a-508m**.

[0071] The quantity of activation buffers **518a-518x** included in the AI accelerator device **110**, and the quantity of the rows **508a-508m** associated with each activation buffer, may be selected to satisfy one or more parameters, such as energy consumption of the activation buffers **518a-518x**, latency of the activation buffers **518a-518x**, and/or the physical footprint of the activation buffers **518a-518x**. For example, increasing the quantity of activation buffers **518a-518x** included in the AI accelerator device **110** (which decreases the quantity of the rows **508a-508m** associated with each of the activation buffers **518a-518x**) may decrease latency and energy consumption in the AI accelerator device **110** at the expense of increased size of the AI accelerator device **110**. As another example, decreasing the quantity of activation buffers **518a-518x** included in the AI accelerator device **110** (which increases the quantity of the rows **508a-508m** associated with each of the activation buffers **518a-518x**) may decrease the size of the AI accelerator device **110** at the expense of increased latency and energy consumption in the AI accelerator device **110**.

[0072] As indicated above, FIG. 5 is provided as an example. Other examples may differ from what is described with regard to FIG. 5.

[0073] FIG. 6 is a diagram of an example implementation **600** of an AI accelerator device **110** described herein. The example implementation **600** of the AI accelerator device **110** is similar to the example implementation **200** of the AI accelerator device **110**, and includes components **602-622** that are similar to the components **202-222** of the example implementation **200** of the AI accelerator device **110**. However, in the example implementation **600** of the AI accelerator device **110**, a plurality of weight buffers **610a-610k** (e.g., a plurality of mini weight buffers) and a plurality of accumulator buffers **620a-620y** (e.g., a plurality of mini accumulator buffers) are both included. The plurality of weight buffers **610a-610k** (e.g., a plurality of mini weight buffers) and the plurality of accumulator buffers **620a-620y** (e.g., a plurality of mini accumulator buffers) may be included to reduce the latency and energy consumption of both the weight buffer and the accumulator buffer of the AI accelerator device **110** relative to the latency and energy consumption of a single monolithic weight buffer and a single monolithic accumulator buffer.

[0074] Each of the weight buffers **610a-610k** may be associated with, and may receive an output (e.g., a partial sum) from, a subset of the columns **606a-606l** of processing element circuits **604a-**

604n included in the processing element array **602**. In some implementations, each of the weight buffers **610a-610k** may be associated with, and may provide weights to, a single respective column of the columns **606a-606l**. For example, the weight buffer **610a** may be associated with, and may provide weights to, the column **606a**; the weight buffer **610b** may be associated with, and may receive provide weights to, the column **606b**; and so on. In some implementations, each of the weight buffers **610a-610k** may be associated with, and may provide weights to, a plurality of columns of the columns **606a-606l**.

[0075] The quantity of weight buffers **610a-610k** included in the AI accelerator device **110**, and the quantity of the columns **606a-606l** associated with each weight buffer, may be selected to satisfy one or more parameters, such as energy consumption of the weight buffers **610a-610k**, latency of the weight buffers **610a-610k**, and/or the physical footprint of the weight buffers **610a-610k**. For example, increasing the quantity of weight buffers **610a-610k** included in the AI accelerator device **110** (which decreases the quantity of the columns **606a-606l** associated with each of the weight buffers **610a-610k**) may decrease latency and energy consumption in the AI accelerator device **110** at the expense of increased size of the AI accelerator device **110**. As another example, decreasing the quantity of weight buffers **610a-610k** included in the AI accelerator device **110** (which increases the quantity of the columns **606a-606l** associated with each of the weight buffers **610a-610k**) may decrease the size of the AI accelerator device **110** at the expense of increased latency and energy consumption in the AI accelerator device **110**.

[0076] Each of the accumulator buffers **620a-620y** may be associated with, and may receive an output (e.g., a partial sum) from, a subset of the columns **606a-606l** of processing element circuits **604a-604n** included in the processing element array **602**. In some implementations, each of the accumulator buffers **620a-620y** may be associated with, and may receive an output from, a single respective column of the columns **606a-606l**. For example, the accumulator buffer **620a** may be associated with, and may receive an output from, the column **606a**; the accumulator buffer **620b** may be associated with, and may receive an output from, the column **606b**; and so on. In some implementations, each of the accumulator buffers **620a-620y** may be associated with, and may receive an output from, a plurality of columns of the columns **606a-606l**.

[0077] The quantity of accumulator buffers **620a-620y** included in the AI accelerator device **110**, and the quantity of the columns **606a-606l** associated with each accumulator buffer, may be selected to satisfy one or more parameters, such as energy consumption of the accumulator buffers **620a-620y**, latency of the accumulator buffers **620a-620y**, and/or the physical footprint of the accumulator buffers **620a-620y**. For example, increasing the quantity of accumulator buffers **620a-620y** included in the AI accelerator device **110** (which decreases the quantity of the columns **606a-606l** associated with each of the accumulator buffers **620a-620y**) may decrease latency and energy consumption in the AI accelerator device **110** at the expense of increased size of the AI accelerator device **110**. As another example, decreasing the quantity of accumulator buffers **620a-620y** included in the AI accelerator device **110** (which increases the quantity of the columns **606a-606l** associated with each of the accumulator buffers **620a-620y**) may decrease the size of the AI accelerator device **110** at the expense of increased latency and energy consumption in the AI accelerator device **110**.

[0078] As indicated above, FIG. 6 is provided as an example. Other examples may differ from what is described with regard to FIG. 6.

[0079] FIG. 7 is a diagram of an example implementation **700** of an AI accelerator device **110** described herein. The example implementation **700** of the AI accelerator device **110** is similar to the example implementation **200** of the AI accelerator device **110**, and includes components **702-710** and **714-724** that are similar to the components **202-210** and **214-224** of the example implementation **200** of the AI accelerator device **110**. However, in the example implementation **700** of the AI accelerator device **110**, a single monolithic weight buffer **710** is included as opposed to a plurality of weight buffers **210a-210k** (e.g., a plurality of mini weight buffers), and a plurality of

activation buffers **718a-718x** (e.g., a plurality of mini activation buffers) and a plurality of accumulator buffers **720a-720y** (e.g., a plurality of mini accumulator buffers) are both included. The plurality of activation buffers **718a-718x** (e.g., a plurality of mini activation buffers) and the plurality of accumulator buffers **720a-720y** (e.g., a plurality of mini accumulator buffers) may be included to reduce the latency and energy consumption of both the activation buffer and the accumulator buffer of the AI accelerator device **110** relative to the latency and energy consumption of a single monolithic activation buffer and a single monolithic accumulator buffer.

[0080] Each of the activation buffers **718a-718x** may be associated with, and may provide input data to, a subset of the rows **708a-708m** of processing element circuits **704a-704n** included in the processing element array **702**. In some implementations, each of the activation buffers **718a-718x** may be associated with, and may provide input data to, a single respective row of the rows **708a-708m**. For example, the activation buffer **718a** may be associated with, and may provide input data to, the row **708a**; the activation buffer **718b** may be associated with, and may provide input data to, the row **708b**; and so on. In some implementations, each of the activation buffers **718a-718x** may be associated with, and may provide input data to, a plurality of rows of the rows **708a-708m**.

[0081] The quantity of activation buffers **718a-718x** included in the AI accelerator device **110**, and the quantity of the rows **708a-708m** associated with each activation buffer, may be selected to satisfy one or more parameters, such as energy consumption of the activation buffers **718a-718x**, latency of the activation buffers **718a-718x**, and/or the physical footprint of the activation buffers **718a-718x**. For example, increasing the quantity of activation buffers **718a-718x** included in the AI accelerator device **110** (which decreases the quantity of the rows **708a-708m** associated with each of the activation buffers **718a-718x**) may decrease latency and energy consumption in the AI accelerator device **110** at the expense of increased size of the AI accelerator device **110**. As another example, decreasing the quantity of activation buffers **718a-718x** included in the AI accelerator device **110** (which increases the quantity of the rows **708a-708m** associated with each of the activation buffers **718a-718x**) may decrease the size of the AI accelerator device **110** at the expense of increased latency and energy consumption in the AI accelerator device **110**.

[0082] Each of the accumulator buffers **720a-720y** may be associated with, and may receive an output (e.g., a partial sum) from, a subset of the columns **706a-706l** of processing element circuits **704a-704n** included in the processing element array **702**. In some implementations, each of the accumulator buffers **720a-720y** may be associated with, and may receive an output from, a single respective column of the columns **706a-706l**. For example, the accumulator buffer **720a** may be associated with, and may receive an output from, the column **706a**; the accumulator buffer **720b** may be associated with, and may receive an output from, the column **706b**; and so on. In some implementations, each of the accumulator buffers **720a-720y** may be associated with, and may receive an output from, a plurality of columns of the columns **706a-706l**.

[0083] The quantity of accumulator buffers **720a-720y** included in the AI accelerator device **110**, and the quantity of the columns **706a-706l** associated with each activation buffer, may be selected to satisfy one or more parameters, such as energy consumption of the accumulator buffers **720a-720y**, latency of the accumulator buffers **720a-720y**, and/or the physical footprint of the accumulator buffers **720a-720y**. For example, increasing the quantity of accumulator buffers **720a-720y** included in the AI accelerator device **110** (which decreases the quantity of the columns **706a-706l** associated with each of the accumulator buffers **720a-720y**) may decrease latency and energy consumption in the AI accelerator device **110** at the expense of increased size of the AI accelerator device **110**. As another example, decreasing the quantity of accumulator buffers **720a-720y** included in the AI accelerator device **110** (which increases the quantity of the columns **706a-706l** associated with each of the accumulator buffers **720a-720y**) may decrease the size of the AI accelerator device **110** at the expense of increased latency and energy consumption in the AI accelerator device **110**.

[0084] As indicated above, FIG. 7 is provided as an example. Other examples may differ from what

is described with regard to FIG. 7.

[0085] FIG. 8 is a diagram of an example implementation **800** of an AI accelerator device **110** described herein. The example implementation **800** of the AI accelerator device **110** is similar to the example implementation **200** of the AI accelerator device **110**, and includes components **802-824** that are similar to the components **202-224** of the example implementation **200** of the AI accelerator device **110**. However, in the example implementation **800** of the AI accelerator device **110**, a plurality of weight buffers **810a-810k** (e.g., a plurality of mini weight buffers), a plurality of activation buffers **818a-818x** (e.g., a plurality of mini activation buffers), a plurality of accumulator buffers **820a-820y** (e.g., a plurality of mini accumulator buffers) are both included. The plurality of weight buffers **810a-810k** (e.g., a plurality of mini weight buffers), the plurality of activation buffers **818a-818x** (e.g., a plurality of mini activation buffers), and the plurality of accumulator buffers **820a-820y** (e.g., a plurality of mini accumulator buffers) may be included to reduce the latency and energy consumption of the weight buffer, the activation buffer, and the accumulator buffer of the AI accelerator device **110** relative to the latency and energy consumption of a single monolithic weight buffer, a single monolithic activation buffer, and a single monolithic accumulator buffer.

[0086] Each of the weight buffers **810a-810k** may be associated with, and may receive an output (e.g., a partial sum) from, a subset of the columns **806a-806l** of processing element circuits **804a-804n** included in the processing element array **802**. In some implementations, each of the weight buffers **810a-810k** may be associated with, and may provide weights to, a single respective column of the columns **806a-806l**. For example, the weight buffer **810a** may be associated with, and may provide weights to, the column **806a**; the weight buffer **810b** may be associated with, and may receive provide weights to, the column **806b**; and so on. In some implementations, each of the weight buffers **810a-810k** may be associated with, and may provide weights to, a plurality of columns of the columns **806a-806l**.

[0087] The quantity of weight buffers **810a-810k** included in the AI accelerator device **110**, and the quantity of the columns **806a-806l** associated with each weight buffer, may be selected to satisfy one or more parameters, such as energy consumption of the weight buffers **810a-810k**, latency of the weight buffers **810a-810k**, and/or the physical footprint of the weight buffers **810a-810k**. For example, increasing the quantity of weight buffers **810a-810k** included in the AI accelerator device **110** (which decreases the quantity of the columns **806a-806l** associated with each of the weight buffers **810a-810k**) may decrease latency and energy consumption in the AI accelerator device **110** at the expense of increased size of the AI accelerator device **110**. As another example, decreasing the quantity of weight buffers **810a-810k** included in the AI accelerator device **110** (which increases the quantity of the columns **806a-806l** associated with each of the weight buffers **810a-810k**) may decrease the size of the AI accelerator device **110** at the expense of increased latency and energy consumption in the AI accelerator device **110**.

[0088] Each of the activation buffers **818a-818x** may be associated with, and may provide input data to, a subset of the rows **808a-808m** of processing element circuits **804a-804n** included in the processing element array **802**. In some implementations, each of the activation buffers **818a-818x** may be associated with, and may provide input data to, a single respective row of the rows **808a-808m**. For example, the activation buffer **818a** may be associated with, and may provide input data to, the row **808a**; the activation buffer **818b** may be associated with, and may provide input data to, the row **808b**; and so on. In some implementations, each of the activation buffers **818a-818x** may be associated with, and may provide input data to, a plurality of rows of the rows **808a-808m**.

[0089] The quantity of activation buffers **818a-818x** included in the AI accelerator device **110**, and the quantity of the rows **808a-808m** associated with each activation buffer, may be selected to satisfy one or more parameters, such as energy consumption of the activation buffers **818a-818x**, latency of the activation buffers **818a-818x**, and/or the physical footprint of the activation buffers **818a-818x**. For example, increasing the quantity of activation buffers **818a-818x** included in the AI

accelerator device **110** (which decreases the quantity of the rows **808a-808m** associated with each of the activation buffers **818a-818x**) may decrease latency and energy consumption in the AI accelerator device **110** at the expense of increased size of the AI accelerator device **110**. As another example, decreasing the quantity of activation buffers **818a-818x** included in the AI accelerator device **110** (which increases the quantity of the rows **808a-808m** associated with each of the activation buffers **818a-818x**) may decrease the size of the AI accelerator device **110** at the expense of increased latency and energy consumption in the AI accelerator device **110**.

[0090] Each of the accumulator buffers **820a-820y** may be associated with, and may receive an output (e.g., a partial sum) from, a subset of the columns **806a-806l** of processing element circuits **804a-804n** included in the processing element array **802**. In some implementations, each of the accumulator buffers **820a-820y** may be associated with, and may receive an output from, a single respective column of the columns **806a-806l**. For example, the accumulator buffer **820a** may be associated with, and may receive an output from, the column **806a**; the accumulator buffer **820b** may be associated with, and may receive an output from, the column **806b**; and so on. In some implementations, each of the accumulator buffers **820a-820y** may be associated with, and may receive an output from, a plurality of columns of the columns **806a-806l**.

[0091] The quantity of accumulator buffers **820a-820y** included in the AI accelerator device **110**, and the quantity of the columns **806a-806l** associated with each accumulator buffer, may be selected to satisfy one or more parameters, such as energy consumption of the accumulator buffers **820a-820y**, latency of the accumulator buffers **820a-820y**, and/or the physical footprint of the accumulator buffers **820a-820y**. For example, increasing the quantity of accumulator buffers **820a-820y** included in the AI accelerator device **110** (which decreases the quantity of the columns **806a-806l** associated with each of the accumulator buffers **820a-820y**) may decrease latency and energy consumption in the AI accelerator device **110** at the expense of increased size of the AI accelerator device **110**. As another example, decreasing the quantity of accumulator buffers **820a-820y** included in the AI accelerator device **110** (which increases the quantity of the columns **806a-806l** associated with each of the accumulator buffers **820a-820y**) may decrease the size of the AI accelerator device **110** at the expense of increased latency and energy consumption in the AI accelerator device **110**.

[0092] In some implementations, the mapping between weight buffers **810a-810k** and the columns **806a-806l**, the mapping between activation buffers **818a-818x** and the rows **808a-808m**, and the mapping between accumulator buffers **820a-820y** and the columns **806a-806l** may be the same mapping. For example, each of the weight buffers **810a-810k** may be associated with a single column of the columns **806a-806l**, each of the activation buffers **818a-818x** may be associated with a single row of the rows **808a-808m**, and each of the accumulator buffers **820a-820y** may be associated with a single column of the columns **806a-806l**. As another example, each of the weight buffers **810a-810k** may be associated with a first quantity of the columns **806a-806l**, each of the activation buffers **818a-818x** may be associated with a second quantity of the rows **808a-808m**, and each of the accumulator buffers **820a-820y** may be associated with a third quantity of the columns **806a-806l**, where the first quantity, the second quantity, and the third quantity are the same quantity.

[0093] In some implementations, two or more of the mapping between weight buffers **810a-810k** and the columns **806a-806l**, the mapping between activation buffers **818a-818x** and the rows **808a-808m**, and/or the mapping between accumulator buffers **820a-820y** and the columns **806a-806l** may be different mappings. To illustrate, for example, each of the weight buffers **810a-810k** may be associated with a single column of the columns **806a-806l**, each of the activation buffers **818a-818x** may be associated with three (3) rows of the rows **808a-808m**, and each of the accumulator buffers **820a-820y** may be associated with five (5) columns of the columns **806a-806l**.

[0094] In some implementations, the mapping between two or more of the weight buffers **810a-810k** and the columns **806a-806l**, the mapping between two or more of the activation buffers **818a-**

818x and the rows **808a-808m**, and/or the mapping between two or more of the accumulator buffers **820a-820y** and the columns **806a-806l** may be different mappings. For example, the weight buffer **810a** may be associated with a single column of the columns **806a-806l**, the weight buffer **810b** may be associated with six (6) columns of the columns **806a-806l**, and the weight buffer **810k** may be associated with eight (8) columns of the columns **806a-806l**.

[0095] In some implementations, the quantity of columns mapped to each of the weight buffers **810a-810k** is greater relative to the quantity of rows mapped to each of the activation buffers **818a-818x** and/or is greater relative to the quantity of columns mapped to each of the accumulator buffers **820a-820y**. In some implementations, the quantity of rows mapped to each of the activation buffers **818a-818x** is greater relative to the quantity of columns mapped to each of the weight buffers **810a-810k** and/or is greater relative to the quantity of columns mapped to each of the accumulator buffers **820a-820y**. In some implementations, the quantity of columns mapped to each of the accumulator buffers **820a-820y** is greater relative to the quantity of rows mapped to each of the activation buffers **818a-818x** and/or is greater relative to the quantity of columns mapped to each of the weight buffers **810a-810k**.

[0096] In some implementations, the quantity of weight buffers **810a-810k** is greater relative to the quantity of activation buffers **818a-818x** and/or is greater relative to the quantity of accumulator buffers **820a-820y**. In some implementations, the quantity of activation buffers **818a-818x** is greater relative to the quantity of weight buffers **810a-810k** and/or is greater relative to the quantity of accumulator buffers **820a-820y**. In some implementations, the quantity of accumulator buffers **820a-820y** is greater relative to the quantity of activation buffers **818a-818x** and/or is greater relative to the quantity of weight buffers **810a-810k**.

[0097] As indicated above, FIG. 8 is provided as an example. Other examples may differ from what is described with regard to FIG. 8.

[0098] FIG. 9 is a diagram of an example implementation **900** of an AI accelerator device **110** described herein. The example implementation **900** of the AI accelerator device **110** is similar to the example implementation **800** of the AI accelerator device **110**, and includes components **902-924** that are similar to the components **802-824** of the example implementation **800** of the AI accelerator device **110**.

[0099] In the example implementation **900** of the AI accelerator device **110**, the weight buffer **910a** is associated with a subset **906a** of the columns of the processing element array **902**, the weight buffer **910k** is associated with a subset **906l** of the columns of the processing element array **902**, and so on, where each subset **906a-906k** includes a plurality of columns. Each column in the subset **906a** is exclusive to the subset **906a** and is not included in another subset of columns, each column in the subset **906l** is exclusive to the subset **906l** and is not included in another subset of columns, and so on.

[0100] In the example implementation **900** of the AI accelerator device **110**, the activation buffer **918a** is associated with a subset **908a** of the rows of the processing element array **902**, the activation buffer **918x** is associated with a subset **908m** of the rows of the processing element array **902**, and so on, where each subset **908a-908m** includes a plurality of rows. Each row in the subset **908a** is exclusive to the subset **908a** and is not included in another subset of rows, each row in the subset **908m** is exclusive to the subset **908m** and is not included in another subset of rows, and so on.

[0101] In the example implementation **900** of the AI accelerator device **110**, the accumulator buffer **920a** is associated with a subset **906a** of the columns of the processing element array **902**, the accumulator buffer **910x** is associated with a subset **906l** of the columns of the processing element array **902**, and so on, where each subset **906a-906l** includes a plurality of columns. Each column in the subset **906a** is exclusive to the subset **906a** and is not included in another subset of columns, each column in the subset **906l** is exclusive to the subset **906l** and is not included in another subset of columns, and so on.

[0102] As indicated above, FIG. 9 is provided as an example. Other examples may differ from what is described with regard to FIG. 9.

[0103] FIGS. 10A-10C are diagrams of an example implementation 1000 described herein. The example implementation 1000 includes an example operational flow of an AI accelerator device 110 described herein. The example implementation 1000 is illustrated in connection with the example implementation 800 of the AI accelerator device 110 of FIG. 8. However, the example operational flow (or portions thereof) of an AI accelerator device 110 may be performed by another example implementation of the AI accelerator device 110 described herein.

[0104] In some implementations, the operational flow is performed to perform neural network processing operations in the AI accelerator device 110. For example, the operational flow may be performed as a part of a machine learning operation to train the processing element circuits 804a-804n and to update or modify the weights associated with the processing element circuits 804a-804n. As another example, the operation flow may be performed to perform an inference in which the AI accelerator device 110 classifies objects in a large data set, in which the AI accelerator device 110 performs computer vision operations, and/or in which the AI accelerator device 110 performs other complex and/or large-scale operations using a neural network. In some implementations, the operation flow is performed based on receiving a request or a query from a client device (e.g., from one or more of the client devices 130-170) over a network 120. The request may be provided to the edge computing system 101, and the edge computing system 101 may use the AI accelerator device 110 to satisfy the request.

[0105] As shown in FIG. 10A, the AI accelerator device 110 may use the plurality of weight buffers 810a-810k to provide weights to the columns 808a-808m of processing element circuits 804a-804n of the processing element array 802. For example, the weight buffer 810a may provide weights to the processing element circuits 804a, 804d, and 804g in the column 806a; the weight buffer 810b may provide weights to the processing element circuits 804b, 804e, and 804h in the column 806b; the weight buffer 810k may provide weights to the processing element circuits 804c, 804f, and 804n in the column 806l; and so on. The weights may be stored in the memory device 814. The plurality of weight buffers 810a-810k may obtain the weights from the memory device 814 and may buffer (e.g., temporarily store) the weights in the plurality of weight buffers 810a-810k to reduce latency in the operational flow.

[0106] As further shown in FIG. 10A, at 1002, the weights may be provided from the memory device 814 to the plurality of weight buffers 810a-810k through the weight buffer multiplexer circuit 812. The weight buffer multiplexer circuit 812 enables the weights to be directed to specific ones of the plurality of weight buffers 810a-810k and, therefore, specific ones of the columns 806a-806l. The distributor circuit 816 may receive a dataflow select signal (DF_SEL) (e.g., from the one or more processors 107) as input, and may generate one or more types of select signals based on the dataflow select signal. For example, the distributor circuit 816 may generate one or more weight select signals and may provide the one or more weight select signals to the weight buffer multiplexer circuit 812. The one or more weight select signals may cause the weight buffer multiplexer circuit 812 to provide weights to specific ones of the plurality of weight buffers 810a-810k. This enables the weights to be loaded from external memory (e.g., the memory device 814), to the plurality of weight buffers 810a-810k, and then from the plurality of weight buffers 810a-810k, to the specific ones of the columns 806a-806l.

[0107] As shown in FIG. 10B, at 1004, the AI accelerator device 110 may use the plurality of activation buffers 818a-818x to provide input data (e.g., partial sums, values to be multiplied) to the rows 808a-808m of processing element circuits 804a-804n of the processing element array 802. For example, the activation buffer 818a may provide input data to the processing element circuits 804a-804c in the row 808a; the activation buffer 818b may provide input data to the processing element circuits 804d-804f in the row 808b; the activation buffer 818x may provide input data to the processing element circuits 804g-804n in the row 808m; and so on. The activation buffers

818a-818x may obtain the input data from the peripheral circuitry **822** and may buffer (e.g., temporarily store) the input data in the activation buffers **818a-818x** to reduce latency in the operational flow.

[0108] As further shown in FIG. **10B**, at **1004**, the input may be provided to the activation buffers **818a-818x** through the activation buffer multiplexer circuit **824**. The activation buffer multiplexer circuit **824** enables the input data to be directed to specific ones of the plurality of activation buffers **818a-818x** and, therefore, specific ones of the rows **808a-808m**. The distributor circuit **816** may generate one or more types of select signals based on the dataflow select signal. For example, the distributor circuit **816** may generate one or more activation select signals and may provide the one or more activation select signals to the activation buffer multiplexer circuit **824**. The one or more activation select signals may cause the activation buffer multiplexer circuit **824** to provide the input data to specific ones of the plurality of activation buffers **818a-818x**.

[0109] As shown in FIG. **10C**, at **1006**, each of the processing element circuits **804a-804n** in the processing element array **802** may perform one or more neural network operations (e.g., a MAC operation and/or another type of neural network operation) to generate an updated partial sum. The element circuits **804a-804n** may perform the one or more neural network operations in a parallel or semi-parallel manner. The updated partial sums are buffered (e.g., temporarily stored) in the plurality of accumulator buffers **820a-820y**. For example, the updated partial sums from the processing element circuits **804a**, **804d**, and **804g** in the column **806a** may be provided to (and buffered in) the accumulator buffer **820a**; the updated partial sums from processing element circuits **804b**, **804e**, and **804h** in the column **806b** may be provided to (and buffered in) the accumulator buffer **820b**; the updated partial sums from processing element circuits **804c**, **804f**, and **804n** in the column **806l** may be provided to (and buffered in) the accumulator buffer **820y**; and so on. The plurality of accumulator buffers **820a-820y** may provide the updated partial sums to the peripheral circuitry **822**.

[0110] As indicated above, FIGS. **10A-10C** are provided as an example. Other examples may differ from what is described with regard to FIGS. **10A-10C**.

[0111] FIG. **11** is a diagram of an example implementation **1100** of a processing element circuit described herein. For example, the example implementation **1100** may include an example MAC operation performed by a processing element circuit of the processing element circuits **202a-202n**, the processing element circuits **302a-302n**, the processing element circuits **402a-402n**, the processing element circuits **502a-502n**, the processing element circuits **602a-602n**, the processing element circuits **702a-702n**, the processing element circuits **802a-802n**, and/or the processing element circuits **904**, among other examples. The MAC operation (also referred to as a weight sum operation) may be performed, for example, as part of the operational flow of the example implementation **1000** (e.g., at **1006**).

[0112] As shown in FIG. **11**, a processing element circuit **1102** may include a multiplier circuit **1104** and an accumulator circuit **1106**. The multiplier circuit **1104** may receive a plurality of inputs **1108** and **1110**, and may perform the multiplication aspect of the MAC operation based on the inputs **1108** and **1110**. In particular, the multiplier circuit **1104** may multiply the inputs **1108** and **1110** together to generate an intermediate output **1112**. The inputs **1108** and **1110** may include values to be multiplied, a weight, and/or another type of input.

[0113] The intermediate output **1112** may be provided to the accumulator circuit **1106**. The accumulator circuit **1106** may perform an accumulate operation, in which the intermediate output **1112** and an input **1114** (e.g., a partial sum) are added together to generate an updated partial sum **1116**. The input **1114** (e.g., the partial sum) may be provided spatially from a neighboring processing element circuit **1102** or temporally from a previous MAC operation performed by the processing element circuit **1102**.

[0114] As indicated above, FIG. **11** is provided as an example. Other examples may differ from what is described with regard to FIG. **11**.

[0115] FIG. 12 is a diagram of an example neural network **1200** described herein. FIG. 12 illustrates a conceptual diagram of the neural network **1200** (e.g., an artificial neural network, a deep neural network), which may be implemented by an AI accelerator device **110** described herein.

[0116] As shown in FIG. 12, the neural network **1200** includes an input layer **1202**, a hidden layer **1204**, and an output layer **1206**. The input layer **1202** may include nodes **1208**, which provide inputs to one or more activation layers **1210**. The activation layers **1210** include nodes **1212** that are implemented by processing element circuits described herein, such as the processing element circuits **202a-202n**, the processing element circuits **302a-302n**, the processing element circuits **402a-402n**, the processing element circuits **502a-502n**, the processing element circuits **602a-602n**, the processing element circuits **702a-702n**, the processing element circuits **802a-802n**, the processing element circuits **904**, and/or the processing element circuit **1102**, among other examples. Each of the nodes **1212** may receive an input and may perform a MAC operation or another type of neural network operation based on the input. Each of the nodes **1212** may provide an output (e.g., an updated partial sum) to the next layer in the neural network, which may include another activation layer **1210** or the output layer **1206** that includes nodes **1214**.

[0117] As indicated above, FIG. 12 is provided as an example. Other examples may differ from what is described with regard to FIG. 12.

[0118] FIGS. 13A-13C are diagrams of example implementations of neural network dataflows described herein. In some implementations, an AI accelerator device **110** described herein may be configured to perform one or more of the neural network dataflows described in connection with FIGS. 13A-13C.

[0119] FIG. 13A illustrates an example implementation **1300** of a weight stationary dataflow. In the weight stationary dataflow, input feature maps (IFMAPs) **1302** may be arranged in rows. Each row of IFMAPs **1302** may be stored in a respective activation buffer (e.g., in a respective activation mini buffer described herein) and provided to a corresponding row of processing element circuits in a processing element array **1304**. Weights may be provided (e.g., by a weight buffer or a plurality of weight mini buffers described herein) to the processing element array **1304** and pre-filled in each processing element circuit in the processing element array **1304** prior to the start of a computation using the weights and the IFMAPs **1302**. The weights may be pre-filled such that all of the elements of a given filter in the processing element array **1304** are allocated along a column, meaning that each column of the processing element array **1304** receives weights from a dedicated weight buffer.

[0120] Partial sums are generated by the processing element circuits based on the IFMAPs **1302** and the pre-filled weights. In some implementations, the processing element circuits may generate one partial sum per cycle. The generated partial sums may then be reduced across the rows of the processing element array **1304**, along with each column in parallel to generate one output feature map (OFMAP) **1306** pixel per column. Each OFMAP column may be assigned to a respective accumulator buffer (e.g., a respective accumulator mini buffer described herein).

[0121] FIG. 13B illustrates an example implementation **1308** of an input stationary dataflow. The input stationary dataflow is similar to the weight stationary dataflow, except that the weights **1310** are streamed into the processing element array **1312** in which the IFMAPs are pre-filled in the processing element circuits of the processing element array **1312**. The weights **1310** are streamed from the edge of the processing element array **1312**, and each processing element circuit of the processing element array **1312** generates a partial sum each cycle based on the weights **1310** and the unrolled IFMAPs.

[0122] The generated partial sums may then be reduced across the rows of the processing element array **1312**, along with each column in parallel to generate one OFMAP **1314** pixel per column. Each OFMAP column may be assigned to a respective accumulator buffer (e.g., a respective accumulator mini buffer described herein).

[0123] FIG. 13C illustrates an example implementation 1316 of an output stationary dataflow. An output stationary dataflow refers to a mapping in which each processing element circuit performs all of the computations for one OFMAP while IFMAPs 1318 and weights 1320 are fed from the edges of the processing element array 1322 and distributed to the processing element circuits using processing element interconnects. Partial sums are generated and reduced within each processing element circuit in the processing element array 1322. Once all of the processing element circuits in the processing element array 1322 complete the generation of OFMAPs, the results are transferred out of the processing element array 1322 through the processing element interconnects.

[0124] As indicated above, FIGS. 13A-13C is provided as an example. Other examples may differ from what is described with regard to FIGS. 13A-13C.

[0125] FIG. 14 is a diagram of example implementation 1400 of energy consumption in an AI accelerator device 110 described herein. FIG. 14 illustrates an example 1402 of an AI accelerator device 110 that is configured to perform a weight stationary dataflow (e.g., according to the example implementation 1300), an example 1404 of an AI accelerator device 110 that is configured to perform an input stationary dataflow (e.g., according to the example implementation 1308), and an example 1406 of an AI accelerator device 110 that is configured to perform an output stationary dataflow (e.g., according to the example implementation 1316).

[0126] In the example 1402, energy consumption 1408 is illustrated for an AI accelerator device 1410 that includes a single monolithic weight buffer, a single monolithic activation buffer, and a single monolithic accumulator buffer. Energy consumption 1408 is also illustrated for an AI accelerator device 110 described herein. The overall energy consumption 1408 may be reduced for an AI accelerator device 110 described herein by including mini buffers for one or more of the weight buffer, the activation buffer, and/or the accumulator buffer of the AI accelerator device 110. In particular, while a compute energy consumption 1412 (e.g., the energy consumption of processing element circuits) remains relatively comparable for the AI accelerator device 1410 and the AI accelerator device 110, a weight buffer energy consumption 1414 may be reduced by including a plurality of weight buffers (such as in the example implementations 200, 500, 600, 800, and/or 900) in the AI accelerator device 110. Additionally and/or alternatively, an activation buffer energy consumption 1416 may be reduced by including a plurality of activation buffers (such as in the example implementations 300, 500, 700, 800, and/or 900) in the AI accelerator device 110. Additionally and/or alternatively, an accumulator buffer energy consumption 1418 may be reduced by including a plurality of accumulator buffers (such as in the example implementations 400, 600, 700, 800, and/or 900) in the AI accelerator device 110.

[0127] For the weight stationary dataflow in the example 1402, the energy consumption reduction may be greatest for the activation buffer energy consumption 1416 and for the accumulator buffer energy consumption 1418. Accordingly, the AI accelerator device 110 may be configured to include a plurality of activation buffers and a plurality of accumulator buffers to reduce energy consumption 1408 of the AI accelerator device 110, while including a single monolithic weight buffer to minimize the size or area of the weight buffer. This may correspond to the example implementation 700. However, other implementations described herein may be used instead of the example implementation 700.

[0128] In the example 1404, the overall energy consumption 1408 may be reduced for the AI accelerator device 110 described herein for the input buffer dataflow by including mini buffers for one or more of the weight buffer, the activation buffer, and/or the accumulator buffer of the AI accelerator device 110. In particular, while compute energy consumption 1412 (e.g., the energy consumption of processing element circuits) remains relatively comparable for the AI accelerator device 1410 and the AI accelerator device 110, the weight buffer energy consumption 1414 may be reduced by including a plurality of weight buffers (such as in the example implementations 200, 500, 600, 800, and/or 900) in the AI accelerator device 110. Additionally and/or alternatively, the activation buffer energy consumption 1416 may be reduced by including a plurality of activation

buffers (such as in the example implementations **300**, **500**, **700**, **800**, and/or **900**) in the AI accelerator device **110**. Additionally and/or alternatively, the accumulator buffer energy consumption **1418** may be reduced by including a plurality of accumulator buffers (such as in the example implementations **400**, **600**, **700**, **800**, and/or **900**) in the AI accelerator device **110**.

[0129] For the input stationary dataflow in the example **1404**, the energy consumption reduction may be greatest for the accumulator buffer energy consumption **1418**. Accordingly, the AI accelerator device **110** may be configured to include a plurality of accumulator buffers to reduce energy consumption **1408** of the AI accelerator device **110**, while including a single monolithic weight buffer and a single monolithic activation buffer to minimize the size or area of the weight buffer and the activation buffer. This may correspond to the example implementation **400**. However, other implementations described herein may be used instead of the example implementation **400**.

[0130] In the example **1406**, the overall energy consumption **1408** may be reduced for the AI accelerator device **110** described herein for the output buffer dataflow by including mini buffers for the weight buffer and/or the activation buffer of the AI accelerator device **110**. In particular, while compute energy consumption **1412** (e.g., the energy consumption of processing element circuits) remains relatively comparable for the AI accelerator device **1410** and the AI accelerator device **110**, the weight buffer energy consumption **1414** may be reduced by including a plurality of weight buffers (such as in the example implementations **200**, **500**, **600**, **800**, and/or **900**) in the AI accelerator device **110**. Additionally and/or alternatively, the activation buffer energy consumption **1416** may be reduced by including a plurality of activation buffers (such as in the example implementations **300**, **500**, **700**, **800**, and/or **900**) in the AI accelerator device **110**.

[0131] For the output stationary dataflow in the example **1406**, the energy consumption reduction may be greatest for the activation buffer energy consumption **1416**. Accordingly, the AI accelerator device **110** may be configured to include a plurality of activation buffers to reduce energy consumption **1408** of the AI accelerator device **110**, while including a single monolithic weight buffer and a single monolithic accumulator buffer to minimize the size or area of the weight buffer and the accumulator buffer. This may correspond to the example implementation **300**. However, other implementations described herein may be used instead of the example implementation **300**.

[0132] As indicated above, FIG. **14** is provided as an example. Other examples may differ from what is described with regard to FIG. **14**.

[0133] FIGS. **15A** and **15B** are diagrams of example implementations of design parameters for an AI accelerator device **110** described herein.

[0134] FIG. **15A** illustrates an example implementation **1500** of design parameters including normalized area **1502** and normalized energy **1504** as a function of processing element density **1506** in a buffer (e.g., quantity of processing element circuits assigned to a single buffer, such as a weight buffer, an activation buffer, an accumulator buffer).

[0135] As shown in FIG. **15A**, read energy **1508** and write energy **1510** decrease as processing element density **1506** in a buffer decreases, meaning that read energy **1508** and write energy **1510** is generally lower for a buffer that is assigned to a lesser quantity of processing element circuits. Accordingly, increasing the quantity of buffers (e.g., the quantity of weight buffers, the quantity of activation buffers, and/or the quantity of accumulator buffers) in an AI accelerator device **110** described herein such that each buffer is assigned to fewer processing element circuits may reduce read energy **1508** and write energy **1510** in the AI accelerator device **110**. However, buffer area **1512** increases as the processing element density **1506** in a buffer also increases. Accordingly, the buffers in an AI accelerator device **110** described herein may be configured such that the processing element density **1506**, the read energy **1508**, the write energy **1510**, and the buffer area **1512** of the buffers falls within an optimization window **1514**.

[0136] FIG. **15B** illustrates an example implementation **1516** of design parameters including normalized area **1518** and normalized latency **1520** as a function of processing element density

1522 in a buffer (e.g., quantity of processing element circuits assigned to a single buffer, such as a weight buffer, an activation buffer, an accumulator buffer).

[0137] As shown in FIG. **15B**, buffer latency **1524** decreases as processing element density **1522** in a buffer decreases, meaning that buffer latency **1524** is generally lower for a buffer that is assigned to a lesser quantity of processing element circuits. Accordingly, increasing the quantity of buffers (e.g., the quantity of weight buffers, the quantity of activation buffers, and/or the quantity of accumulator buffers) in an AI accelerator device **110** described herein such that each buffer is assigned to fewer processing element circuits may reduce buffer latency **1524** in the AI accelerator device **110**. However, buffer area **1526** increases as the processing element density **1522** in a buffer also increases. Accordingly, the buffers in an AI accelerator device **110** described herein may be configured such that the processing element density **1522**, the buffer latency **1524**, and the buffer area **1526** of the buffers falls within an optimization window **1528**.

[0138] As indicated above, FIGS. **15A** and **15B** is provided as an example. Other examples may differ from what is described with regard to FIGS. **15A** and **15B**.

[0139] FIG. **16** is a diagram of example components of a device **1600** described herein. Device **1600** may correspond to one or more components included in the edge computing system **101** and/or one or more of the client devices **130-170**. In some implementations, one or more components included in the edge computing system **101** and/or one or more of the client devices **130-170** may include one or more devices **1600** and/or one or more components of device **1600**. As shown in FIG. **16**, device **1600** may include a bus **1610**, a processor **1620**, a memory **1630**, an input component **1640**, an output component **1650**, and a communication component **1660**.

[0140] Bus **1610** may include one or more components that enable wired and/or wireless communication among the components of device **1600**. Bus **1610** may couple together two or more components of FIG. **16**, such as via operative coupling, communicative coupling, electronic coupling, and/or electric coupling. Processor **1620** may include a central processing unit, a graphics processing unit, a microprocessor, a controller, a microcontroller, a digital signal processor, a field-programmable gate array, an application-specific integrated circuit, and/or another type of processing component. Processor **1620** is implemented in hardware, firmware, or a combination of hardware and software. In some implementations, processor **1620** may include one or more processors capable of being programmed to perform one or more operations or processes described elsewhere herein.

[0141] Memory **1630** may include volatile and/or nonvolatile memory. For example, memory **1630** may include random access memory (RAM), read only memory (ROM), a hard disk drive, and/or another type of memory (e.g., a flash memory, a magnetic memory, and/or an optical memory). Memory **1630** may include internal memory (e.g., RAM, ROM, or a hard disk drive) and/or removable memory (e.g., removable via a universal serial bus connection). Memory **1630** may be a non-transitory computer-readable medium. Memory **1630** stores information, instructions, and/or software (e.g., one or more software applications) related to the operation of device **1600**. In some implementations, memory **1630** may include one or more memories that are coupled to one or more processors (e.g., processor **1620**), such as via bus **1610**.

[0142] Input component **1640** enables device **1600** to receive input, such as user input and/or sensed input. For example, input component **1640** may include a touch screen, a keyboard, a keypad, a mouse, a button, a microphone, a switch, a sensor, a global positioning system sensor, an accelerometer, a gyroscope, and/or an actuator. Output component **1650** enables device **1600** to provide output, such as via a display, a speaker, and/or a light-emitting diode. Communication component **1660** enables device **1600** to communicate with other devices via a wired connection and/or a wireless connection. For example, communication component **1660** may include a receiver, a transmitter, a transceiver, a modem, a network interface card, and/or an antenna.

[0143] Device **1600** may perform one or more operations or processes described herein. For example, a non-transitory computer-readable medium (e.g., memory **1630**) may store a set of

instructions (e.g., one or more instructions or code) for execution by processor **1620**. Processor **1620** may execute the set of instructions to perform one or more operations or processes described herein. In some implementations, execution of the set of instructions, by one or more processors **1620**, causes the one or more processors **1620** and/or the device **1600** to perform one or more operations or processes described herein. In some implementations, hardwired circuitry is used instead of or in combination with the instructions to perform one or more operations or processes described herein. Additionally, or alternatively, processor **1620** may be configured to perform one or more operations or processes described herein. Thus, implementations described herein are not limited to any specific combination of hardware circuitry and software.

[0144] The number and arrangement of components shown in FIG. **16** are provided as an example. Device **1600** may include additional components, fewer components, different components, or differently arranged components than those shown in FIG. **16**. Additionally, or alternatively, a set of components (e.g., one or more components) of device **1600** may perform one or more functions described as being performed by another set of components of device **1600**.

[0145] FIG. **17** is a flowchart of an example process **1700** associated with operation of an AI accelerator device described herein. In some implementations, one or more process blocks of FIG. **17** are performed by an AI accelerator device (e.g., the AI accelerator device **110**).

[0146] As shown in FIG. **17**, process **1700** may include providing a plurality of weights to a plurality of weight buffers of the AI accelerator device (block **1710**). For example, the AI accelerator device **110** may provide a plurality of weights to a plurality of weight buffers (e.g., weight buffers **810a-810k**, weight buffers **910a-910k**) of the AI accelerator device **110**, as described herein.

[0147] As further shown in FIG. **17**, process **1700** may include providing, using the plurality of weight buffers, a plurality of subsets of the plurality of weights to respective columns of processing element circuits of a processing element array of the AI accelerator device (block **1720**). For example, the AI accelerator device **110** may provide, using the plurality of weight buffers, a plurality of subsets of the plurality of weights to respective columns (e.g., columns **806a-806l**, columns **906a-906l**) of processing element circuits (e.g., processing element circuits **804a-804n**, processing element circuits **904**) of a processing element array (e.g., processing element array **802**, processing element array **902**) of the AI accelerator device **110**, as described herein.

[0148] As further shown in FIG. **17**, process **1700** may include providing activation data to a plurality of activation buffers of the AI accelerator device (block **1730**). For example, the AI accelerator device **110** may provide activation data to a plurality of activation buffers (e.g., activation buffers **818a-818x**, activation buffers **918a-918x**) of the AI accelerator device, as described herein.

[0149] As further shown in FIG. **17**, process **1700** may include providing, using the plurality of activation buffers, a plurality of subsets of the activation data to respective rows of processing element circuits of the processing element array (block **1740**). For example, the AI accelerator device **110** may provide, using the plurality of activation buffers, a plurality of subsets of the activation data to respective rows (e.g., rows **808a-808m**, rows **908a-908m**) of processing element circuits of the processing element array, as described herein.

[0150] As further shown in FIG. **17**, process **1700** may include providing a plurality of partial sums from the processing element array to respective accumulator buffers of a plurality of accumulator buffers of the AI accelerator device (block **1750**). For example, the AI accelerator device **110** may provide a plurality of partial sums from the processing element array to respective accumulator buffers (e.g., accumulator buffers **820a-820y**, accumulator buffers **920a-920y**) of a plurality of accumulator buffers of the AI accelerator device, as described herein. In some implementations, the plurality of partial sums are based on a MAC operation performed by the processing element array on the plurality of weights and the activation data.

[0151] As further shown in FIG. **17**, process **1700** may include providing the plurality of partial

sums from the respective accumulator buffers to peripheral circuitry of the AI accelerator device (block **1760**). For example, the AI accelerator device **110** may provide the plurality of partial sums from the respective accumulator buffers to peripheral circuitry (e.g., peripheral circuitry **822**, peripheral circuitry **922**) of the AI accelerator device **110**, as described herein.

[0152] Process **1700** may include additional implementations, such as any single implementation or any combination of implementations described below and/or in connection with one or more other processes described elsewhere herein.

[0153] In a first implementation, providing the plurality of weights to the plurality of weight buffers includes providing the plurality of weights from a memory device (e.g., memory device **814**, memory device **914**) of the AI accelerator device **110** to the plurality of weight buffers through a weight buffer multiplexer circuit (e.g., weight buffer multiplexer circuit **812**, weight buffer multiplexer circuit **912**), and the process **1700** includes generating a select signal, and providing the select signal to the weight buffer multiplexer circuit to cause the weight buffer multiplexer circuit to provide the plurality of weights to the plurality of weight buffers.

[0154] In a second implementation, along or in combination with the first implementation, providing the activation data to the plurality of activation buffers includes providing the activation data from peripheral circuitry (e.g., peripheral circuitry **822**, peripheral circuitry **922**) of the AI accelerator device **110** to the plurality of activation buffers through an activation buffer multiplexer circuit (e.g., activation buffer multiplexer circuit **824**, activation buffer multiplexer circuit **924**), and process **1700** includes generating a select signal, and providing the select signal to the activation buffer multiplexer to cause the activation buffer multiplexer to provide the activation data to the plurality of activation buffers.

[0155] In a third implementation, along or in combination with the first or second implementation, the plurality of partial sums include a plurality of updated partial sums, and wherein the activation data comprises a plurality of input partial sums, and inputting data to be multiplied with the plurality of weights and accumulated with the plurality of input partial sums to form the updated partial sums. In a fourth implementation, alone or in combination with one or more of the first through third implementations, process **1700** includes providing the plurality of input partial sums from the peripheral circuitry to the activation buffer multiplexer circuit. In a fifth implementation, alone or in combination with one or more of the first through fourth implementations, providing the plurality of partial sums from the processing element array to the respective accumulator buffers includes providing the plurality of partial sums from the respective columns of processing element circuits of the processing element array to the respective accumulator buffers.

[0156] Although FIG. **17** shows example blocks of process **1700**, in some implementations, process **1700** includes additional blocks, fewer blocks, different blocks, or differently arranged blocks than those depicted in FIG. **17**. Additionally, or alternatively, two or more of the blocks of process **1700** may be performed in parallel.

[0157] In this way, an AI accelerator device may include a plurality of on-chip mini buffer that are associated with a PE array. Each mini buffer is associated with a subset of rows or a subset of columns of the PE array. Partitioning an on-chip buffer of the AI accelerator device into the mini buffers described herein may reduce the size and complexity of the on-chip buffer. The reduced size of the on-chip buffer may reduce the wire routing complexity of the on-chip buffer, which may reduce latency and may reduce access energy for the AI accelerator device. This may increase the operating efficiency and/or may increase the performance of the AI accelerator device. Moreover, the mini buffers may increase the overall bandwidth that is available for the mini buffers to transfer data to and from the PE array. As described in greater detail above, some implementations described herein provide an AI accelerator device. The AI accelerator device includes a processing element array, including a plurality of columns of processing element circuits a plurality of rows of processing element circuits. The AI accelerator device includes a plurality of weight buffers associated with the processing element array, where the plurality of weight buffers are associated

with respective subsets of columns of the plurality of columns of processing element circuits of the processing element array.

[0158] As described in greater detail above, some implementations described herein provide an AI accelerator device. The AI accelerator device includes a processing element array, including a plurality of columns of processing element circuits a plurality of rows of processing element circuits. The AI accelerator device includes a plurality of activation buffers associated with the processing element array, where the plurality of activation buffers are associated with a respective subset of rows of the plurality of rows of processing element circuits of the processing element array.

[0159] As described in greater detail above, some implementations described herein provide a method. The method includes providing, by an AI accelerator device, a plurality of weights to a plurality of weight buffers of the AI accelerator device. The method includes providing, by the AI accelerator device and using the plurality of weight buffers, a plurality of subsets of the plurality of weights to respective columns of processing element circuits of a processing element array of the AI accelerator device. The method includes providing, by the AI accelerator device, activation data to a plurality of activation buffers of the AI accelerator device. The method includes providing, by the AI accelerator device and using the plurality of activation buffers, a plurality of subsets of the activation data to respective rows of processing element circuits of the processing element array. The method includes providing, by the AI accelerator device, a plurality of partial sums from the processing element array to respective accumulator buffers of a plurality of accumulator buffers of the AI accelerator device, where the plurality of partial sums are based on a MAC operation performed by the processing element array on the plurality of weights and the activation data. The method includes providing, by the AI accelerator device, the plurality of partial sums from the respective accumulator buffers to peripheral circuitry of the AI accelerator device.

[0160] The foregoing outlines features of several embodiments so that those skilled in the art may better understand the aspects of the present disclosure. Those skilled in the art should appreciate that they may readily use the present disclosure as a basis for designing or modifying other processes and structures for carrying out the same purposes and/or achieving the same advantages of the embodiments introduced herein. Those skilled in the art should also realize that such equivalent constructions do not depart from the spirit and scope of the present disclosure, and that they may make various changes, substitutions, and alterations herein without departing from the spirit and scope of the present disclosure.

Claims

1. An artificial intelligence (AI) accelerator device, comprising: a processing element array, comprising: a plurality of columns of processing element circuits; and a plurality of rows of processing element circuits; and a plurality of accumulator buffers associated with the processing element array, wherein the plurality of accumulator buffers are associated with respective subsets of columns of the plurality of columns of processing element circuits.
2. The AI accelerator device of claim 1, further comprising: a single monolithic weight buffer associated with the plurality of columns of processing element circuits.
3. The AI accelerator device of claim 1, further comprising: a single monolithic activation buffer associated with the plurality of rows of processing element circuits.
4. The AI accelerator device of claim 3, wherein the single monolithic activation buffer is configured to provide input data to the plurality of rows of processing element circuits.
5. The AI accelerator device of claim 3, further comprising: a periphery circuit configured to provide output to the single monolithic activation buffer, wherein the plurality of accumulator buffers are configured to provide output to the periphery circuit.
6. The AI accelerator device of claim 1, wherein each of the plurality of accumulator buffers is

configured to receive output data from a subset of the plurality of columns of processing element circuits.

7. The AI accelerator device of claim 1, wherein a first accumulator buffer, of the plurality of accumulator buffers, is configured to receive output from a first column of processing element circuits of the plurality of columns of processing element circuits; and wherein a second accumulator buffer, of the plurality of accumulator buffers, is configured to receive output from a second column of processing element circuits of the plurality of columns of processing element circuits.

8. An artificial intelligence (AI) accelerator device, comprising: a processing element array, comprising: a plurality of columns of processing element circuits; and a plurality of rows of processing element circuits; a plurality of accumulator buffers associated with the processing element array; and a plurality of weight buffers associated with the processing element array, wherein the plurality of weight buffers are associated with respective subsets of columns of the plurality of columns of processing element circuits.

9. The AI accelerator device of claim 8, wherein a first weight buffer, of the plurality of weight buffers, is configured to provide weights to a first column of processing element circuits of the plurality of columns of processing element circuits; and wherein a second weight buffer, of the plurality of weight buffers, is configured to provide weights to a second column of processing element circuits of the plurality of columns of processing element circuits.

10. The AI accelerator device of claim 8, wherein each of the plurality of weight buffers is associated with an output from a subset of the plurality of columns of processing element circuits.

11. The AI accelerator device of claim 8, wherein each of the plurality of weight buffers is configured to provide weights to a single column of processing element circuits of the plurality of columns of processing element circuits.

12. The AI accelerator device of claim 8, further comprising: a weight buffer multiplexer circuit coupled with the plurality of weight buffers.

13. The AI accelerator device of claim 8, further comprising: a single monolithic activation buffer associated with the plurality of rows of processing element circuits.

14. The AI accelerator device of claim 13, further comprising: a periphery circuit configured to provide output to the single monolithic activation buffer, wherein the plurality of accumulator buffers are configured to provide output to the periphery circuit.

15. A method, comprising: providing, by an artificial intelligence (AI) accelerator device, a plurality of weights to a plurality of weight buffers of the AI accelerator device; providing, by the AI accelerator device and using the plurality of weight buffers, a plurality of subsets of the plurality of weights to respective columns of processing element circuits of a processing element array of the AI accelerator device; providing, by the AI accelerator device, activation data to an activation buffer of the AI accelerator device; providing, by the AI accelerator device and using the activation buffer, the activation data to a row of processing element circuits of the processing element array; providing, by the AI accelerator device, a plurality of partial sums from the processing element array to respective accumulator buffers of a plurality of accumulator buffers of the AI accelerator device, wherein the plurality of partial sums are based on a multiply and accumulate (MAC) operation performed by the processing element array on the plurality of weights and the activation data; and providing, by the AI accelerator device, the plurality of partial sums from the respective accumulator buffers to peripheral circuitry of the AI accelerator device.

16. The method of claim 15, wherein the plurality of weights are provided to the plurality of weight buffers via a weight buffer multiplexer circuit.

17. The method of claim 15, wherein providing the plurality of subsets of the plurality of weights to the respective columns of processing element circuits comprises: providing, by a first weight buffer of the plurality of weight buffers, a first subset of the plurality of weights to a first column of processing element circuits; and providing, by a second weight buffer of the plurality of weight

buffers, a second subset of the plurality of weights to a second column of processing element circuits.

18. The method of claim 15, further comprising: performing, using the processing element array, neural network operations in parallel.

19. The method of claim 18, wherein the neural network operations include the MAC operation.

20. The method of claim 15, wherein providing the plurality of partial sums from the processing element array to respective accumulator buffers of the plurality of accumulator buffers comprises: providing a first subset of the plurality of partials sums from a first column of processing element circuits to a first accumulator buffer of the plurality of accumulator buffers; and providing a second subset of the plurality of partials sums from a second column of processing element circuits to a second accumulator buffer of the plurality of accumulator buffers.
