

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250258711

Kind Code

A1

Publication Date

August 14, 2025

Inventor(s)

TANDON; Prateek et al.

HARDWARE ACCELERATION FOR FUNCTION PROCESSING

Abstract

A function processing service may receive a request to execute source code. The source code may include instructions to perform a function. The function processing service may determine whether at least one hardware acceleration condition has been satisfied for the function. If at least one hardware acceleration condition has been satisfied, the instructions in the source code may be translated into hardware-specific code corresponding to a hardware circuit. The hardware circuit may be configured based on the hardware-specific code, and the hardware circuit may perform the function. The function processing service may then provide the result obtained from the hardware circuit to the requesting entity.

Inventors: TANDON; Prateek (Issaquah, WA), CORELL; Brian Jacob (Sammamish, WA)

Applicant: Microsoft Technology Licensing, LLC (Redmond, WA)

Family ID: 1000008572048

Appl. No.: 19/170880

Filed: April 04, 2025

Related U.S. Application Data

parent US continuation 17586434 20220127 parent-grant-document US 11900165 child US 18404705

parent US division 18404705 20240104 parent-grant-document US 12293230 child US 19170880

parent US division 16555927 20190829 parent-grant-document US 11237873 child US 17586434

Publication Classification

Int. Cl.: G06F9/50 (20060101); G06F9/38 (20180101); G06F9/48 (20060101); G06F13/362 (20060101)

U.S. Cl.:

CPC **G06F9/5027** (20130101); **G06F9/3836** (20130101); **G06F9/4843** (20130101);
G06F13/3625 (20130101);

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS [0001] This application is a Divisional of U.S. application Ser. No. 18/404,705, filed on Jan. 4, 2024, which is a Continuation of U.S. application Ser. No. 17/586,434, filed on Jan. 27, 2022, now U.S. Pat. No. 11,900,165, which is a Divisional of U.S. application Ser. No. 16/555,927, filed on Aug. 29, 2019, now U.S. Pat. No. 11,237,873, the entireties of which are incorporated herein by reference.

BACKGROUND

[0002] Cloud computing is the delivery of computing services (e.g., servers, storage, databases, networking, software, analytics) over the Internet. Broadly speaking, a cloud computing system includes two sections, a front end and a back end, that are in electronic communication with one another via the Internet. The front end includes the interface that users encounter through a client device. The back end includes the resources that deliver cloud-computing services, including processors, memory, storage, and networking hardware. The back end of a cloud computing system typically includes one or more data centers, which may be located in different geographical areas. Each data center typically includes a large number (e.g., hundreds or thousands) of computing devices, which may be referred to as host machines.

[0003] Many different types of services may be provided by a cloud computing system. A function processing service is a cloud computing service that allows users to execute source code without the complexity of building and maintaining the infrastructure that is typically associated with developing and launching an application. A function processing service automatically manages the computing resources that are required to execute the source code. Therefore, a function processing service makes it possible for software developers to execute source code without having to worry about provisioning or managing servers.

[0004] As the name “function processing” implies, a function processing service may be utilized to execute source code that includes instructions to perform one or more functions. Function processing is sometimes referred to as serverless computing or function as a service (FaaS).

[0005] As function processing services grow in popularity, the source code that is being executed is growing more complex and data intensive. For instance, machine learning or artificial intelligence applications can utilize data intensive scripts that process massive amounts of data. These data intensive scripts can greatly increase processing time and consume significant amounts of computing resources. Users, however, expect function processing services to run efficiently and quickly. Accordingly, benefits may be realized by techniques for enhancing the efficiency of a function processing service in a cloud computing system.

SUMMARY

[0006] In accordance with one aspect of the present disclosure, a computer-implemented method is disclosed that includes receiving, at a function processing service, a request to execute source code. The request is received from a requesting entity. The source code includes instructions to perform a function. The method further includes determining that a hardware acceleration condition has been satisfied for the function, translating the instructions in the source code into hardware-specific code corresponding to a hardware circuit, configuring the hardware circuit based on the hardware-specific code, causing the hardware circuit to perform the function, and providing a result obtained from the hardware circuit to the requesting entity.

[0007] As an example, the hardware circuit may include a field programmable gate array (FPGA), and the hardware-specific code may include hardware description language (HDL) code corresponding to the FPGA. As another example, the hardware circuit may include a graphics processing unit (GPU), and the hardware-specific code may include GPU code. As another example, the hardware circuit may include a digital signal processor (DSP), and the hardware-specific code may include DSP code. As another example, the hardware circuit may include an application specific integrated circuit (ASIC), and the hardware-specific code may include hardware description language (HDL) code.

[0008] Determining that the hardware acceleration condition has been satisfied may include counting a number of requests to perform the function that the function processing service receives within a time interval and determining that the number of requests exceeds a threshold value.

[0009] In some embodiments, the function may belong to a family of functions. In such embodiments, determining that the hardware acceleration condition has been satisfied may include counting a number of requests to perform any function within the family of functions that the function processing service receives within a time interval and determining that the number of requests exceeds a threshold value.

[0010] In some embodiments, determining that the hardware acceleration condition has been satisfied may include determining an estimated cost savings associated with performing the function via the hardware circuit compared to performing the function via a general-purpose processor and determining that the estimated cost savings exceeds a threshold value.

[0011] In some embodiments, determining that the hardware acceleration condition has been satisfied may include determining that the function is included in a function library and may further include adding new functions to the function library.

[0012] In accordance with another aspect of the present disclosure, a computer-implemented method is disclosed that includes determining that a hardware acceleration condition has been satisfied for a first function, configuring a hardware circuit to perform the first function, and adding the first function to a function library that is implemented in hardware circuits. The method further includes receiving, at a function processing service, a request to execute source code. The source code includes instructions to perform the first function. The method further includes causing the first function to be performed via the hardware circuit.

[0013] In some embodiments, the method may further include configuring the hardware circuit to perform a second function in parallel with the first function. The second function may be different from the first function. In other embodiments, the method may further include configuring the hardware circuit to perform a plurality of instances of the hardware circuit in parallel.

[0014] In accordance with another aspect of the present disclosure, a computer-implemented method includes analyzing first historical information related to a first plurality of requests received by a function processing service during a first time interval, configuring a set of hardware circuits to perform a first set of functions during the first time interval based on the first historical information, and analyzing second historical information related to a second plurality of requests received by the function processing service during a second time interval. The second time interval is different from the first time interval. The method further includes configuring the set of hardware circuits to perform a second set of functions during the second time interval based on the second historical information. The second set of functions is different from the first set of functions.

[0015] The method may further include receiving a first request to execute first source code. The first request may be received during the first time interval. The first source code may further include first instructions to perform a first function. The first set of functions may include the first function. The method may further include causing the first function to be performed via a first hardware circuit in the set of hardware circuits and receiving a second request to execute second source code. The second request may be received during the second time interval. The second source code may include second instructions to perform a second function. The second set of

functions may include the second function performing the second function during the second time interval via a second hardware circuit in the set of hardware circuits.

[0016] This Summary is provided to introduce a selection of concepts in a simplified form that are further described below in the Detailed Description. This Summary is not intended to identify key features or essential features of the claimed subject matter, nor is it intended to be used as an aid in determining the scope of the claimed subject matter.

[0017] Additional features and advantages will be set forth in the description that follows. Features and advantages of the disclosure may be realized and obtained by means of the systems and methods that are particularly pointed out in the appended claims. Features of the present disclosure will become more fully apparent from the following description and appended claims, or may be learned by the practice of the disclosed subject matter as set forth hereinafter.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0018] In order to describe the manner in which the above-recited and other features of the disclosure can be obtained, a more particular description will be rendered by reference to specific embodiments thereof which are illustrated in the appended drawings. For better understanding, the like elements have been designated by like reference numbers throughout the various accompanying figures. Understanding that the drawings depict some example embodiments, the embodiments will be described and explained with additional specificity and detail through the use of the accompanying drawings in which:

[0019] FIG. 1 illustrates an example of a function processing service that is configured to selectively utilize hardware acceleration in accordance with the present disclosure.

[0020] FIG. 2 illustrates aspects of a function processing service that implements a hardware acceleration condition based on the number of requests to perform a function that have been received.

[0021] FIG. 3 illustrates aspects of a function processing service that implements a hardware acceleration condition based on the number of requests that have been received within a defined time interval, where counters are used to keep track of the number of requests that have been received in connection with various functions.

[0022] FIG. 4 illustrates aspects of a function processing service that implements a hardware acceleration condition based on the number of requests that have been received within a defined time interval, where historical information is stored about previous requests that have been received.

[0023] FIG. 5 illustrates aspects of a function processing service that implements a hardware acceleration condition based on the number of requests to perform a function that have been received in connection with any function in a family of related functions.

[0024] FIG. 6 illustrates aspects of a function processing service that implements a hardware acceleration condition based on an estimated cost savings associated with the use of hardware acceleration to perform a function.

[0025] FIG. 7 illustrates aspects of a function processing service that maintains a function library, and that implements a hardware acceleration condition based on whether a particular function is already included in the function library.

[0026] FIG. 8 illustrates an example of a function processing service for which a plurality of hardware acceleration conditions have been defined.

[0027] FIG. 9 illustrates an example of a method that may be performed by a condition evaluator in a function processing system.

[0028] FIG. 9A illustrates an example of an alternative method that may be implemented by a

condition evaluator in an embodiment where a user is given the opportunity to choose whether or not hardware acceleration can be performed.

[0029] FIG. **10** illustrates an example of a function processing service that maintains a function library.

[0030] FIG. **11** illustrates an example of a method that may be performed by the function processing service shown in FIG. **10**.

[0031] FIG. **12** illustrates an example of a hardware circuit that has been configured to perform two different functions.

[0032] FIG. **13** illustrates an example of a hardware circuit that has been configured to perform two different instances of the same function.

[0033] FIGS. **14A-B** illustrate an example of a function processing service that periodically reconfigures a set of hardware circuits based on expected demand.

[0034] FIG. **15** illustrates an example of a method that may be performed by the function processing service in FIGS. **14A-B**.

[0035] FIG. **16** illustrates certain components that may be included within a computing system.

DETAILED DESCRIPTION

[0036] The present disclosure is generally related to enhancing the efficiency of a function processing service in a cloud computing system. In accordance with one aspect of the present disclosure, a function processing service may be configured to selectively utilize hardware acceleration. More specifically, a function processing service may be configured to utilize hardware acceleration to perform a function when a determination is made that it would be beneficial to do so. Such a determination may be made based on various factors, such as how frequently requests to perform the function are received, whether the speed of processing can be increased by a meaningful amount, and/or whether the cost of processing can be reduced by a meaningful amount.

8

[0037] In this context, the term “hardware acceleration” may refer to the use of a hardware circuit to perform a function, where the hardware circuit has been specifically designed or programmed to perform the function so that the function can be performed faster than if it were specified in software and performed on a general-purpose processor. The term “hardware circuit” may in general refer to any configuration of electrically or electromagnetically connected components or devices other than a general-purpose processor. A hardware circuit may be an integrated circuit. Some examples of hardware circuits that may be used in accordance with the present disclosure include field programmable gate arrays (FPGAs), graphics processing units (GPUs), digital signal processors (DSPs), and application specific integrated circuits (ASICs).

[0038] There are many different types of functions for which hardware acceleration may be beneficial. Some examples include data processing, compression or encryption, location tracking, video processing, financial transactions, log analysis, text processing, fast filtering of streaming data, web hosting, machine learning training (e.g., linear regression), machine learning inference (e.g., image recognition), signal processing (e.g., using a DSP), and so forth.

[0039] In some embodiments, one or more conditions may be defined that indicate when hardware acceleration should be used to perform a function. Such conditions may be referred to herein as hardware acceleration conditions. When a function processing service receives a request to execute source code that includes instructions to perform a function, the function processing service may determine whether any of the defined hardware acceleration conditions have been satisfied. If at least one hardware acceleration condition has been satisfied, then hardware acceleration may be utilized to perform the function.

[0040] In some embodiments, hardware acceleration may be performed automatically, without requiring the user to specifically request hardware acceleration or to provide hardware-specific code, such as register-transfer logic (RTL) code for an FPGA. For example, a user's request to perform a particular function may take the form of a request to execute source code. The source

code is typically in a high-level programming language such as Java, C++, Python, or the like. A function processing system that receives such a request may automatically determine whether it would be beneficial to use hardware acceleration to perform the function. If it is determined that hardware acceleration would be beneficial, then the source code may be automatically translated into hardware-specific code, without requiring additional input from the user.

[0041] In some embodiments, before a function processing service uses hardware acceleration to perform a function, a user may be notified about the possibility of using hardware acceleration and given the opportunity to choose whether or not hardware acceleration can be performed.

[0042] FIG. 1 illustrates an example of a function processing service **102** that is configured to selectively utilize hardware acceleration in accordance with the present disclosure. The function processing service **102** is shown receiving a request **104** to perform a particular function. However, the request **104** is not a request to perform a function via a hardware circuit. Instead, the request **104** is to execute source code **106**. The request **104** may be received from a client device **108**. The source code **106** includes instructions **110** that, if executed by a general-purpose processor **112**, cause a particular function to be performed.

[0043] In the depicted example, a plurality of hardware acceleration conditions **114** have been defined for the function processing service **102**. In addition, the function processing service **102** includes a component that is configured to determine whether any of the hardware acceleration conditions **114** have been satisfied. This component may be referred to herein as a condition evaluator **116**.

[0044] In response to receiving the request **104**, the condition evaluator **116** determines whether any of the hardware acceleration conditions **114** have been satisfied with respect to the function that is specified by the instructions **110** in the source code **106**. If none of the hardware acceleration conditions **114** have been satisfied, then the function processing service **102** fulfills the request **104** by causing the instructions **110** in the source code **106** to be executed by a general-purpose processor **112**.

[0045] However, if at least one of the hardware acceleration conditions **114** has been satisfied, then hardware acceleration may be utilized to perform the function. The function processing service **102** includes a component that is configured to facilitate the use of hardware acceleration. This component may be referred to herein as a hardware acceleration component **124**. The hardware acceleration component **124** may be configured to translate the instructions **110** in the source code **106** into hardware-specific code **120** corresponding to a hardware circuit **122**. Once the hardware-specific code **120** has been generated, the hardware acceleration component **124** may then configure the hardware circuit **122** based on the hardware-specific code **120** and cause the hardware circuit **122** to perform the function. The hardware circuit **122** may then perform the function and generate a particular result **126**. The function processing service **102** may then provide this result **126** to the requesting entity, which is the client device **108** in the depicted example.

[0046] There are many different types of hardware acceleration conditions that may be defined in accordance with the present disclosure. For example, a hardware acceleration condition related to a particular function may be based on how many requests to perform the function have been received. In some embodiments, a hardware acceleration condition may be defined in terms of a threshold value that indicates the minimum number of requests to perform a particular function that should be received in order to justify the use of hardware acceleration for that function. The number of such requests that are received by the function processing service may then be counted and compared with the threshold value. If the number of such requests that are received exceeds the threshold value, then the hardware acceleration condition is satisfied, and hardware acceleration may be used to perform the function.

[0047] FIG. 2 illustrates aspects of a function processing service **202** that implements a hardware acceleration condition **214** based on the number of requests to perform a function that have been received.

[0048] The function processing service **202** may define a threshold value **228** that indicates the minimum number of requests to perform a particular function that should be received in order to justify the use of hardware acceleration for the function. The function processing service **202** may also keep track of the number of requests that have been received in connection with various functions. In the depicted embodiment, the function processing service **202** utilizes counters **232** corresponding to a plurality of functions that may be performed by the function processing service **202**. FIG. 2 shows the function processing service **202** with a counter **232a** indicating the number of requests that have been received to perform function A, a counter **232b** indicating the number of requests that have been received to perform function B, and a counter **232c** indicating the number of requests that have been received to perform function C.

[0049] The function processing service **202** may also include a component that is configured to update the counters **232** when requests to perform particular functions are received. This component may be referred to herein as a counter update component **234**. When the function processing service **202** receives a request to perform a particular function (e.g., function A), the counter update component **234** may increment the corresponding counter (e.g., the function A counter **232a**).

[0050] The function processing service **202** includes a condition evaluator **216** that is configured to determine whether the hardware acceleration condition **214** has been satisfied. The hardware acceleration condition **214** may be defined such that hardware acceleration is performed for a particular function if the value of the corresponding counter **232** exceeds the defined threshold value **228**. Thus, when the function processing service **202** receives a request to perform a particular function (e.g., function A), the condition evaluator **216** may compare the value of the counter **232** corresponding to that function (e.g., the function A counter **232a**) to the threshold value **228**. If the value of the corresponding counter **232** exceeds the threshold value **228**, the condition evaluator **216** may conclude that the hardware acceleration condition **214** has been satisfied and hardware acceleration may be used to perform the function. If, however, the value of the corresponding counter **232** is less than the threshold value **228**, the condition evaluator **216** may conclude that the hardware acceleration condition **214** has not been satisfied. If there is not another hardware acceleration condition that has been satisfied for the function, then the function may be carried out using a general-purpose processor.

[0051] In some embodiments, a hardware acceleration condition may be defined such that hardware acceleration is performed for a particular function only if a certain number of requests to perform the function have been received within a defined time interval. In such embodiments, a function processing service may define both a threshold value and a relevant time interval. The threshold value may indicate the minimum number of requests to perform a particular function that should be received within the time interval in order to justify the use of hardware acceleration for a function.

[0052] For example, suppose that the threshold value is N (where N represents a positive integer), and the time interval is one week. In this case, if more than N requests to perform a particular function are received within the same week, then the hardware acceleration condition may be deemed to be satisfied and hardware acceleration may be performed. However, if more than N requests to perform a particular function are received but not within the same week, then it may be concluded that the hardware acceleration condition has not been satisfied.

[0053] There are various ways that such a temporal constraint may be implemented. In some embodiments, a function processing service may utilize counters corresponding to functions that may be performed by the function processing service, and the counters may be periodically reset based on the time interval that has been defined for the hardware acceleration condition. For example, if the defined time interval is one week, then the counters may be reset once each week.

[0054] In some embodiments, a function processing service may store historical information that includes information about previous requests that have been received. The information that is

stored about a particular request may include information identifying the requested function, as well as a time stamp indicating when the request was received. In such embodiments, the stored historical information may be queried to determine how many requests to perform a particular function have been received during a certain period of time. If, for example, a request to perform a particular function is received at time *t*, the historical information may be queried to determine how many requests to perform that function have been received during a time period corresponding to the defined time interval (e.g., the defined time interval immediately preceding time *t*). The result of that query may then be compared with the threshold value to determine whether the hardware acceleration condition is satisfied.

[0055] FIG. 3 illustrates aspects of a function processing service **302** that implements a hardware acceleration condition **314** based on the number of requests that have been received within a defined time interval **330**, where counters **332** are used to keep track of the number of requests that have been received in connection with various functions. FIG. 3 shows the function processing service **302** with a counter **332a** indicating the number of requests that have been received to perform function A, a counter **332b** indicating the number of requests that have been received to perform function B, and a counter **332c** indicating the number of requests that have been received to perform function C. The function processing service **302** includes a counter update component **334** that is configured to update the counters **332** when requests to perform particular functions are received.

[0056] In the depicted embodiment, the threshold value **328** indicates the minimum number of requests to perform a particular function that should be received within the defined time interval **330** in order to justify the use of hardware acceleration for the function. The counters **332** may be periodically reset based on the time interval **330** that has been defined for the hardware acceleration condition **314**.

[0057] The function processing service **302** may utilize a timer **336** to keep track of when the counters **332** should be reset. In some embodiments, the timer **336** may be a countdown timer, and the value of the countdown timer may be set to the defined time interval **330**. The counters **332** may be reset whenever the countdown timer expires. Alternatively, the timer **336** may be a countup timer, and the counters **332** may be reset whenever the countup timer reaches a value that is equivalent to the defined time interval **330**.

[0058] When the function processing service **302** receives a request to perform a particular function (e.g., function A), a condition evaluator **316** may compare the value of the counter **332** corresponding to that function (e.g., the function A counter **332a**) to the threshold value **328**. If the value of the corresponding counter **332** exceeds the threshold value **328**, hardware acceleration may be used to perform the function. If, however, the value of the corresponding counter **332** is less than the threshold value **328**, the function may be carried out using a general-purpose processor.

[0059] FIG. 4 illustrates aspects of a function processing service **402** that implements a hardware acceleration condition **414** based on the number of requests that have been received within a defined time interval **430**, where historical information is stored about previous requests that have been received.

[0060] As before, the function processing service **402** may define a threshold value **428** that indicates the minimum number of requests to perform a particular function that should be received within a defined time interval **430** in order to justify the use of hardware acceleration for the function.

[0061] The historical information may be stored in a database **448**. The historical information may include a plurality of records **440**, and each record **440** may correspond to a particular request that the function processing service **402** has received. More specifically, each record **440** may include information identifying the requested function, as well as a time stamp indicating when the request was received. FIG. 4 shows the historical information with a record **440a** corresponding to a

request to process function A, a record **440b** corresponding to a request to process function B, a record **440c** corresponding to another request to process function A, and a record **440d** corresponding to a request to process function C. The records **440a-d** include time stamps. The time stamp in a particular record **440** indicates the time at which the corresponding request was received.

[0062] The function processing service **402** includes a condition evaluator **416** that is configured to determine whether the hardware acceleration condition **414** has been satisfied. The condition evaluator **416** may include a querying component **446** that is capable of querying the database **448** to determine how many requests to perform a particular function have been received during a certain period of time. If, for example, a request to perform a particular function is received at time *t*, the querying component **446** may query the database **448** to determine how many requests to perform that function have been received during a time period corresponding to the defined time interval **430** (e.g., the defined time interval **430** immediately preceding time *t*). The result of that query may then be compared with the threshold value **428** to determine whether the hardware acceleration condition **414** is satisfied.

[0063] Instead of defining a hardware acceleration condition in terms of a threshold value for a particular function, a hardware acceleration condition may instead be defined in terms of a threshold value for a family of related functions. In other words, a hardware acceleration condition may be defined such that whenever a request is received to perform any function within a family of related functions, that counts toward satisfying the threshold value. If the number of such requests that are received within the time interval exceeds the threshold value, then the hardware acceleration condition may be deemed to have been satisfied, and hardware acceleration may be used to perform any function within the family of related functions.

[0064] In embodiments where a hardware acceleration condition is defined in terms of a threshold value for a family of related functions, the hardware acceleration condition may be implemented with or without a temporal constraint. In other words, in some embodiments the threshold value may be defined in terms of the total number of requests that have been received in connection with any function in a family of related functions, regardless of when those requests were received. Alternatively, in other embodiments the threshold value may be defined in terms of the total number of requests that have been received in connection with any function in a family of related functions during a defined time interval. In embodiments where a hardware acceleration condition is defined in terms of a threshold value for a family of related functions and a temporal constraint is utilized, the temporal constraint may be implemented in a variety of ways, including the ways discussed above (e.g., counters and historical information).

[0065] FIG. 5 illustrates aspects of a function processing service **502** that implements a hardware acceleration condition **514** based on the number of requests to perform a function that have been received in connection with any function in a family of related functions. For the sake of simplicity, no temporal constraint is utilized in the depicted embodiment.

[0066] The function processing service **502** may define a threshold value **528** that indicates the minimum number of requests to perform any function in a family of related functions that should be received in order to justify the use of hardware acceleration for the functions in that family. The function processing service **502** may also keep track of the number of requests that have been received in connection with various families of functions. In the depicted embodiment, the function processing service **502** utilizes counters **532** corresponding to different families of functions that may be performed by the function processing service **502**. It will be assumed that at least two families of functions have been defined: (i) a first family of functions that includes function A, function B, and function C, and (ii) a second family of functions that includes function D, function E, and function F. FIG. 5 shows the function processing service **502** with a first counter **532a** and a second counter **532b**. The first counter **532a** indicates the number of requests that have been received to perform any of the functions in the first family of functions (function A, function B, or

function C). The second counter **532b** indicates the number of requests that have been received to perform any of the functions in the second family of functions (function D, function E, or function F).

[0067] The function processing service **502** may also include a counter update component **534** that is configured to update the counters **532** when requests to perform particular functions are received. When the function processing service **502** receives a request to perform a particular function, the counter update component **534** may increment the corresponding counter **532**. Because the counters **532** correspond to particular families of functions, requests to perform different functions may cause the same counter **532** to be updated. For example, if the function processing service **502** receives a request to perform function A, the first counter **532a** may be updated because the first counter **532a** corresponds to the same family of functions to which function A belongs. If the function processing service **502** subsequently receives a request to perform function B, the first counter **532a** may once again be updated because function B belongs to the same family as function A. If, however, the function processing service **502** subsequently receives a request to perform function D, a different counter (specifically, the second counter **532b**) may be updated because function D belongs to a different family than function A and function B.

[0068] The function processing service **502** includes a condition evaluator **516** that is configured to determine whether the hardware acceleration condition **514** has been satisfied. The hardware acceleration condition **514** may be defined such that hardware acceleration is performed for a particular function if the value of the relevant counter **532** (i.e., the counter **532** corresponding to the family of functions to which the specified function belongs) exceeds the defined threshold value **528**. Thus, when the function processing service **502** receives a request to perform a particular function (e.g., function A), the condition evaluator **516** may compare the value of the counter **532** corresponding to the family of functions to which the specified function belongs (e.g., the first counter **532a**) to the threshold value **528**. If the value of the corresponding counter **532** exceeds the threshold value **528**, the condition evaluator **516** may conclude that the hardware acceleration condition **514** has been satisfied and hardware acceleration may be used to perform the function. If, however, the value of the corresponding counter **532** is less than the threshold value **528**, the condition evaluator **516** may conclude that the hardware acceleration condition **514** has not been satisfied. If no other hardware acceleration condition has been satisfied for the function, the function may be carried out using a general-purpose processor.

[0069] In some embodiments, a hardware acceleration condition may be based on an estimated cost savings associated with the use of hardware acceleration to perform a function compared to performing the function via a general-purpose processor. Generally speaking, less energy is required to perform a function via hardware acceleration than to perform the function via a general-purpose processor. This energy savings can lead to cost savings as well. If the cost savings is sufficiently high, then this could justify the use of hardware acceleration to perform the function.

[0070] In some embodiments, a hardware acceleration condition may be defined in terms of a threshold value that indicates the minimum estimated cost savings that would justify the use of hardware acceleration. When hardware acceleration is being considered for a particular function, the cost savings for using hardware acceleration may be estimated. This estimated cost savings may then be compared with the threshold value. If the estimated cost savings exceeds the threshold value, then it may be concluded that the hardware acceleration condition is satisfied, and hardware acceleration may be used to perform the function.

[0071] Some examples of factors that can be considered in connection with determining the cost savings for using hardware acceleration include (i) the cost of the general-purpose processor per second of runtime, (ii) the cost of the energy that is consumed, and (iii) an energy-delay product. Each of these factors can be weighted independently in a cost model.

[0072] As an example, assume that the baseline function takes I instructions and T seconds to complete on a general-purpose processor, the cost of the processor is C, and the energy needed to

complete the operation is E joules. Using hardware acceleration, the operations might be completed in $T_a = T/k_1$ seconds and use E/k_2 joules of energy. Also, the cost of the accelerator time may be $C/(k_3 \cdot T_a)$. The parameter k_1 indicates how much faster the function can be performed using hardware acceleration as compared to using a general-purpose processor. The parameter k_2 indicates how much energy can be saved using hardware acceleration as compared to using a general-purpose processor. The parameter k_3 indicates how much less expensive it is to use hardware acceleration as compared to using a general-purpose processor. The parameters k_1 , k_2 , and k_3 can be measured for certain operations, and those measurements can be used to estimate these parameters for other operations.

[0073] FIG. 6 illustrates aspects of a function processing service **602** that implements a hardware acceleration condition **614** based on an estimated cost savings associated with the use of hardware acceleration to perform a function. As before, the function processing service **602** may define a threshold value **628** that indicates the minimum estimated cost savings that would justify the use of hardware acceleration.

[0074] The function processing service **602** may include a component that is configured to obtain estimates **654** of the cost savings for using hardware acceleration in connection with particular functions. This component may be referred to herein as a cost savings estimator **652**. FIG. 6 shows several examples of cost savings estimates **654**, including a cost savings estimate **654a** for function A, a cost savings estimate **654b** for function B, and a cost savings estimate **654c** for function C.

[0075] The function processing service **602** includes a condition evaluator **616** that is configured to determine whether the hardware acceleration condition **614** has been satisfied. The hardware acceleration condition **614** may be defined such that hardware acceleration is performed for a particular function if the value of the corresponding cost savings estimate **654** exceeds the defined threshold value **628**. Thus, when the function processing service **602** receives a request to perform a particular function (e.g., function A), the condition evaluator **616** may compare the value of the cost savings estimate **654** corresponding to that function (e.g., the cost savings estimate **654a** corresponding to function A) to the threshold value **628**. If the value of the corresponding cost savings estimate **654** exceeds the threshold value **628**, the condition evaluator **616** may conclude that the hardware acceleration condition **614** has been satisfied and hardware acceleration may be used to perform the function. If, however, the value of the corresponding cost savings estimate **654** is less than the threshold value **628**, the condition evaluator **616** may conclude that the hardware acceleration condition **614** has not been satisfied. If no other hardware acceleration condition is satisfied for the function, the function may be carried out using a general-purpose processor.

[0076] In some embodiments, a function library may be maintained. Functions can be added to the function library on an ongoing basis. In this context, the term “function library” may refer to a library of functions that are implemented in hardware circuits. In such embodiments, a hardware acceleration condition may be based on whether a particular function is already included in the function library. For example, when hardware acceleration is being considered for a particular function, a determination may be made about whether that function is already included in the function library. If it is, then it may be concluded that the hardware acceleration condition is satisfied, and hardware acceleration may be used to perform the function.

[0077] FIG. 7 illustrates aspects of a function processing service **702** that maintains a function library, and that implements a hardware acceleration condition **714** based on whether a particular function is already included in the function library.

[0078] In the depicted embodiment, the function processing service **702** maintains a function library index **756**, which includes information about the functions that are presently included in the function library. The function library index **756** may include a plurality of records **758**, and each record **758** may identify a particular function in the function library as well as a hardware circuit that has been configured to perform that function. FIG. 7 shows the function library index **756** with a first record **758a**, a second record **758b**, and a third record **758c**. The first record **758a** includes

identifying information **764a** about function A as well as an identifier **762a** corresponding to a first hardware circuit **722a** that has been configured to perform function A. Similarly, the second record **758b** includes identifying information **764b** about function B as well as an identifier **762b** corresponding to a second hardware circuit **722b** that has been configured to perform function B. The third record **758c** includes identifying information **764c** about function C as well as an identifier **762c** corresponding to a third hardware circuit **722c** that has been configured to perform function C.

[0079] The function processing service **702** includes a condition evaluator **716** that is configured to determine whether the hardware acceleration condition **714** has been satisfied. The hardware acceleration condition **714** may be defined such that hardware acceleration is performed for a particular function if that function is already included in the function library. To determine whether a particular function is already included in the function library, the condition evaluator **716** may be configured to determine whether identifying information about that function is included in the function library index **756**. Thus, when the function processing service **702** receives a request to perform a particular function (e.g., function A), the condition evaluator **716** may search for identifying information about the function (e.g., identifying information **764a** about function A) in the function library index **756**. If the condition evaluator **716** finds identifying information about the requested function in the function library index **756**, the condition evaluator **716** may conclude that the hardware acceleration condition **714** has been satisfied and hardware acceleration may be used to perform the function. If, however, the condition evaluator **716** does not find identifying information about the function in the function library index **756**, the condition evaluator **716** may conclude that the hardware acceleration condition **714** has not been satisfied. If no other hardware acceleration condition has been satisfied for the function, the function may be carried out using a general-purpose processor.

[0080] As indicated above, functions can be added to the function library on an ongoing basis. In some embodiments, whenever hardware acceleration is used to perform a function, the function may be added to the function library as long as there is space to implement the function on the available hardware circuits. When available space becomes limited, then higher-priority functions (e.g., based on amount of usage) may be retained in the function library.

[0081] As indicated above, a function processing service may be utilized to execute source code that includes instructions to perform one or more functions. When a function processing service receives a request to execute source code that includes instructions to perform a particular function, a determination may be made about whether a hardware acceleration condition is satisfied. In some embodiments, a plurality of hardware acceleration conditions may be defined. In embodiments where a plurality of hardware acceleration conditions are defined, hardware acceleration may be utilized in connection with a particular function if any of the hardware acceleration conditions are satisfied. If, however, none of the hardware acceleration conditions are satisfied, the function may be performed using a general-purpose processor.

[0082] FIG. **8** illustrates an example of a function processing service **802** for which a plurality of hardware acceleration conditions **814** have been defined. For example, a hardware acceleration condition **814a** has been defined based on whether a particular function is already included in a function library. In the discussion that follows, this type of hardware acceleration condition **814a** may be referred to as a function library-based hardware acceleration condition **814a**. An example of a function library-based hardware acceleration condition **814a** was discussed above in connection with FIG. **7**.

[0083] In addition, a hardware acceleration condition **814b** has been defined based on the number of requests to perform a function that have been received. In the discussion that follows, this type of hardware acceleration condition **814b** may be referred to as a demand-based hardware acceleration condition **814b**. Several examples of demand-based hardware acceleration conditions **814b** were discussed above in connection with FIGS. **2-5**. Although just a single demand-based

hardware acceleration condition **814b** is shown in FIG. 8, in some embodiments more than one demand-based hardware acceleration condition **814b** may be defined.

[0084] In addition, a hardware acceleration condition **814c** has been defined based on an estimated cost savings associated with the use of hardware acceleration to perform a function. In the discussion that follows, this type of hardware acceleration condition **814c** may be referred to as a cost-based hardware acceleration condition **814c**. An example of a cost-based hardware acceleration condition **814c** was discussed above in connection with FIG. 6.

[0085] The function processing service **802** includes a condition evaluator **816** that is configured to determine whether any of the hardware acceleration conditions **814a-c** have been satisfied. The function processing service **802** also includes a hardware acceleration component **824** that is configured to facilitate the use of hardware acceleration. FIG. 9 illustrates an example of a method **900** that may be performed by the condition evaluator **816** and the hardware acceleration component **824**.

[0086] The method **900** includes receiving **902** a request **804** to perform a function. More specifically, the condition evaluator **816** may receive a request **804** to execute source code **806**, and the source code **806** may include instructions **810** that, if executed by a general-purpose processor **812**, cause a particular function to be performed.

[0087] In response to receiving **902** the request, the condition evaluator **816** may determine whether any of the defined hardware acceleration conditions **814** have been satisfied. In the depicted example, the condition evaluator **816** initially determines **904** whether the function library-based hardware acceleration condition **814a** has been satisfied. If the condition evaluator **816** determines **904** that the function library-based hardware acceleration condition **814a** has been satisfied, this means that a hardware circuit **822** has already been configured to perform the function. Therefore, the method **900** may then include causing **906** the hardware circuit **822** to perform the function.

[0088] If, however, the condition evaluator **816** determines **904** that the function library-based hardware acceleration condition **814a** has not been satisfied, the condition evaluator **816** may then determine whether another one of the defined hardware acceleration conditions **814** has been satisfied. In the depicted example, the condition evaluator **816** next determines **908** whether the demand-based hardware acceleration condition **814b** has been satisfied. If the demand-based hardware acceleration condition **814b** has been satisfied, then the method **900** may include configuring **910** a hardware circuit **822** to perform the function (because it was previously determined, as part of evaluating the function library-based hardware acceleration condition **814a**, that a hardware circuit **822** has not already been configured to perform the function). The method **900** may then include causing **912** the hardware circuit **822** to perform the function.

[0089] If, however, the condition evaluator **816** determines **908** that the demand-based hardware acceleration condition **814b** has not been satisfied, the condition evaluator **816** may then determine whether another one of the defined hardware acceleration conditions **814** has been satisfied. In the depicted example, the condition evaluator **816** next determines **914** whether the cost-based hardware acceleration condition **814c** has been satisfied. If the demand-based hardware acceleration condition **814b** has been satisfied, then the method **900** may include configuring **910** a hardware circuit **822** to perform the function and causing **912** the hardware circuit **822** to perform the function.

[0090] If, however, the condition evaluator **816** determines **914** that the cost-based hardware acceleration condition **814c** has not been satisfied, this means that none of the defined hardware acceleration conditions **814a-c** has been satisfied. In this case, the method **900** may then involve causing **916** the function to be performed using a general-purpose processor **812**.

[0091] In some embodiments, a function processing service may be configured so that a user (i.e., a person or entity who submits request(s) to execute source code) is able to choose whether the option of hardware acceleration is implemented. When a user submits a request to perform a

function to a function processing service, the user typically assumes that the function will be carried out using a general-purpose processor. In some embodiments, before a function processing service uses hardware acceleration to perform a function, a user may be notified about the possibility of using hardware acceleration and given the opportunity to choose whether or not hardware acceleration can be performed.

[0092] FIG. 9A illustrates an example of an alternative method **900A** that may be implemented by a condition evaluator **816** in an embodiment where a user is given the opportunity to choose whether or not hardware acceleration can be performed. The method **900A** is similar to the method **900** shown in FIG. 9, except as indicated below.

[0093] As before, the method **900A** includes receiving **902** a request **804** to perform a function. However, the method **900A** also includes determining **903** whether the user has indicated a preference for using hardware acceleration when it would be beneficial to do so (e.g., when one or more hardware acceleration conditions **814a-c** are satisfied). In some embodiments, determining **903** whether the user has indicated a preference for using hardware acceleration may involve accessing stored information about user preferences. As another example, determining **903** whether the user has indicated a preference for using hardware acceleration may involve prompting the user for authorization to perform hardware acceleration.

[0094] If it is determined **903** that the user has not indicated a preference for using hardware acceleration, then the method **900A** may proceed to the action of causing **916** the function to be performed using a general-purpose processor **812**. If, however, it is determined **903** that the user has indicated a preference for using hardware acceleration, then the method **900A** involves determining whether the hardware acceleration conditions **814a-c** have been satisfied, in the manner described above.

[0095] The specific order in which the hardware acceleration conditions **814a-c** are evaluated in the methods **900**, **900A** shown in FIGS. 9 and 9A is provided for purposes of example only and should not be interpreted as limiting the scope of the present disclosure. In alternative embodiments, the hardware acceleration conditions **814a-c** may be evaluated in a different order. In some embodiments, once a determination has been made that a hardware acceleration condition is satisfied and hardware acceleration is going to be utilized in connection with the function, a determination may then be made about whether a hardware circuit has already been configured for the function.

[0096] To configure a hardware circuit to perform a particular function, the source code that includes the instructions for performing the function may be translated into hardware-specific code corresponding to the type of hardware circuit that is going to be used. For example, if an FPGA is going to be used to perform hardware acceleration, then the source code may be translated into hardware description language (HDL) code corresponding to the FPGA. If a GPU is going to be used to perform hardware acceleration, then the source code may be translated into GPU code. If a programmable ASIC is going to be used to perform hardware acceleration, then a custom compiler may take the source code and generate machine code that is specific to the programmable ASIC. If a DSP is going to be used to perform hardware acceleration, then the source code may be translated into DSP code.

[0097] Once the source code has been translated into hardware-specific code corresponding to the type of hardware circuit that is going to be used, the hardware-specific code may then be provided to the hardware circuit. The function may then be performed via the hardware circuit. The result that is obtained from performing the function may then be provided to the entity that requested the function to be performed.

[0098] As indicated above, a function processing service may maintain a function library. In some embodiments, a function processing service may add a function to the library whenever a determination is made that a hardware acceleration condition corresponding to the function has been satisfied. In other words, once a determination has been made that a hardware acceleration

condition has been satisfied for a particular function, a hardware circuit may be configured to perform that function and the function may be added to the library. Then, upon subsequently receiving a request to execute source code that includes instructions to perform the function, the function may be performed via the hardware circuit.

[0099] FIG. **10** illustrates an example of a function processing service **1002** that maintains a function library. FIG. **11** illustrates an example of a method **1100** that may be performed by the function processing service **1002** in accordance with the present disclosure. The various components of the function processing service **1002** will be described in connection with the method **1100**.

[0100] The method **1100** includes determining **1102** that a hardware acceleration condition **1014** has been satisfied for a particular function. This determination may be made by a condition evaluator **1016** in the function processing service **1002**. In some embodiments, this determination may be made in response to receiving a request to execute source code. Alternatively, in other embodiments, this determination may be made independently of receiving a specific request to execute source code. For example, the function processing service **1002** may determine that a cost-based hardware acceleration condition is satisfied for a particular function in advance of receiving a request to perform that function.

[0101] In response to determining **1102** that a hardware acceleration condition **1014** has been satisfied for a particular function, a hardware circuit **1022** may be configured **1104** to perform that function. This operation may be performed by a hardware acceleration component **1024** in the function processing service **1002**. In addition, the function may be added **1106** to the function library that is maintained by the function processing service **1002**. This may involve adding a record **1058** to the function library index **1056**. The record **1058** may include identifying information **1064** about the function as well as an identifier **1062** corresponding to the hardware circuit **1022** that has been configured to perform the function.

[0102] At some point after the function has been added to the function library, a request **1004** may be received **1108** to perform the function. More specifically, a request **1004** to execute source code **1006** may be received. The source code **1006** may include instructions **1010** that, if executed by a general-purpose processor **1012**, cause the function to be performed.

[0103] In response to receiving the request **1004**, the function processing service **1002** may cause **1110** the function to be performed via the hardware circuit **1022**. This may involve providing input data to the hardware circuit **1022**, as well as instructions to commence processing the input data. In some embodiments, the function processing service **1002** may cause **1110** the function to be performed via the hardware circuit **1022** in response to receiving the request **1004** and determining that the function is included in the function library that is maintained by the function processing service **1002**. Alternatively, in other embodiments, the function processing service **1002** may cause **1110** the function to be performed via the hardware circuit **1022** in response to receiving the request **1004**, determining that the function is included in the function library that is maintained by the function processing service **1002**, and determining that at least one additional hardware acceleration condition has been satisfied (e.g., that the user has expressed a preference for hardware acceleration).

[0104] In some embodiments, a hardware circuit may be configured to perform a plurality of functions in parallel. For example, a hardware circuit may be configured to perform a plurality of different functions in parallel with one another. FIG. **12** illustrates an example of a hardware circuit **1222** that has been configured to perform two different functions, function A and function B, in parallel with one another. In particular, a first section **1266a** of the hardware circuit **1222** has been configured to perform function A, and a second section **1266b** of the hardware circuit **1222** has been configured to perform function B.

[0105] There are various scenarios in which such a hardware circuit **1222** may be utilized. As an example, suppose that a function processing service **1202** receives multiple requests **1204a-b** to

execute source code **1206a-b**. The source code **1206a** in the first request **1204a** may include instructions **1210a** that, if executed by a general-purpose processor, cause function A to be performed. The source code **1206b** in the second request **1204b** may include instructions **1210b** that, if executed by the general-purpose processor, cause function B to be performed. The requests **1204a-b** may be received at substantially the same time, and/or the function processing service **1202** may attempt to fulfill the requests **1204a-b** at substantially the same time. In response to a condition evaluator **1216** determining that at least one hardware acceleration condition **1214** is satisfied for both function A and function B, a hardware acceleration component **1224** in the function processing service **1202** may cause the hardware circuit **1222** to perform both function A and function B in parallel. In particular, the function processing service **1202** may cause the first section **1266a** of the hardware circuit **1222** to perform function A, and the second section **1266b** of the hardware circuit **1222** to perform function B.

[0106] Although the example shown in FIG. **12** only shows two different functions being performed in parallel, this should not be interpreted as limiting the scope of the present disclosure. More than two functions may be performed in parallel in accordance with the present disclosure. In general, there may be N pipelines running in parallel (where the term “pipeline” refers to a section of a hardware circuit that is configured to execute a particular function, and where N may be any positive integer). These N pipelines may be configured to execute up to N different functions in parallel.

[0107] As another example, a hardware circuit may be configured to perform a plurality of different instances of the same function in parallel with one another. FIG. **13** illustrates an example of a hardware circuit **1322** that has been configured to perform two instances of the same function in parallel with one another. In particular, a first section **1366a** of the hardware circuit **1322** has been configured to perform a first instance of the function, and a second section **1366b** of the hardware circuit **1322** has been configured to perform a second instance of the function.

[0108] An example will be described that is similar to the example discussed above in connection with the hardware circuit **1222** shown in FIG. **12**. Suppose that a function processing service **1302** receives multiple requests **1304a**, **1304b** to execute source code **1306a**, **1306b**. Further suppose that the source code **1306a**, **1306b** in both requests **1304a**, **1304b** includes instructions **1310a**, **1310b** that, if executed by a general-purpose processor, cause a particular function to be performed. The requests **1304a-b** may be received at substantially the same time, and/or the function processing service **1302** may attempt to fulfill the requests **1304a**, **1304b** at substantially the same time. In response to a condition evaluator **1316** determining that at least one hardware acceleration condition **1314** is satisfied for the function, a hardware acceleration component **1324** within the function processing service **1302** may cause the hardware circuit **1322** to perform two different instances of the function in parallel. In particular, the function processing service **1302** may cause the first section **1366a** of the hardware circuit **1322** to perform a first instance of the function, and the second section **1366b** of the hardware circuit **1322** to perform a second instance of the function.

[0109] Although the example shown in FIG. **13** only shows two instances of a function being performed in parallel, this should not be interpreted as limiting the scope of the present disclosure. More than two instances of a particular function may be performed in parallel in accordance with the present disclosure. In general, there may be N pipelines running in parallel (where the term “pipeline” refers to a section of a hardware circuit that is configured to execute the function, and where N may be any positive integer). These N pipelines may be configured to execute up to N instances of the function in parallel.

[0110] In accordance with another aspect of the present disclosure, hardware acceleration may be utilized in connection with different functions during different time intervals. In other words, a particular set of hardware circuits may be periodically reconfigured to perform different functions. This may be advantageous in a scenario where a function processing service has a particular set of hardware circuits available for use, and the function processing service receives different types of

requests at different points in time (e.g., different times of the day). In this case, the set of hardware circuits may be configured to perform one set of functions during one particular time interval (e.g., one part of the day) and to perform another set of functions during another time interval (e.g., another part of the day).

[0111] To identify the functions that should utilize hardware acceleration during a particular time interval, historical information may be analyzed about requests that have previously been received by the function processing service during that time interval. If this analysis indicates that requests to perform a particular set of functions have been frequently received during that time interval, then the available set of hardware circuits may be configured to perform that set of functions.

[0112] FIGS. **14A-B** illustrate an example of a function processing service **1402** that periodically reconfigures a set of hardware circuits **1422** based on expected demand. FIG. **15** illustrates an example of a method **1500** that may be performed by the function processing service **1402** in accordance with the present disclosure. The various components of the function processing service **1402** will be described in connection with the method **1500**.

[0113] For the sake of simplicity, only two different time intervals are defined in the function processing service **1402** shown in FIGS. **14A-B**. In some embodiments, the time intervals may be non-overlapping time intervals, such that the first time interval does not overlap with the second time interval. In other embodiments, the time intervals may at least partially overlap with one another. The different time intervals may correspond to different parts of the same day, week, month, etc.

[0114] The method **1500** includes analyzing **1502** historical information **1468a** related to requests received during the first time interval. The function processing service **1402** is shown with a request analyzer **1474** that is configured to analyze the historical information **1468a**. Analyzing **1502** the historical information **1468a** may involve identifying the types of requests that have been received most frequently during the first time interval. Based on the analysis of this historical information **1468a**, a first set of functions may be identified as the functions that are performed most frequently during the first time interval and that could benefit from hardware acceleration. The method **1500** then includes a hardware acceleration component **1424** in the function processing service **1402** configuring **1504** the available hardware circuits **1422** to perform the first set of functions during the first time interval. FIG. **14A** shows the hardware circuits **1422** in a first configuration **1472a**, in which the hardware circuits **1422** have been configured to perform the first set of functions. FIG. **14B** shows the hardware circuits **1422** in a second configuration **1472b**, in which the hardware circuits **1422** have been configured to perform the second set of functions.

[0115] Suppose that the function processing service **1402** receives a request **1404a** during the first time interval to execute source code **1406a** that includes instructions **1410a** to perform a function within the first set of functions. As indicated above, the hardware circuits **1422** have been configured to perform the first set of functions during the first time interval. Therefore, the function processing service **1402** may cause the function to be performed by one of the available hardware circuits **1422** that has been configured to perform the relevant function.

[0116] The method **1500** also includes analyzing **1506** historical information **1468b** related to requests received during the second time interval. This may involve identifying the types of requests that have been received most frequently during the second time interval. Based on the analysis of this historical information **1468b**, a second set of functions may be identified as the functions that are performed most frequently during the second time interval and that could benefit from hardware acceleration. The method **1500** then includes the hardware acceleration component **1424** configuring **1508** the available hardware circuits **1422** to perform the second set of functions during the second time interval.

[0117] Suppose that the function processing service **1402** receives a request **1404b** during the second time interval to execute source code **1406b** that includes instructions **1410b** to perform a function within the second set of functions. As indicated above, the hardware circuits **1422** have

been configured to perform the second set of functions during the second time interval. Therefore, the function processing service **1402** may cause the function to be performed by one of the available hardware circuits **1422** that has been configured to perform the relevant function. [0118] One or more computing systems may be used to implement a function processing service as disclosed herein. FIG. **16** illustrates certain components that may be included within a computing system **1600**.

[0119] The computing system **1600** includes a processor **1601**. The processor **1601** may be a general purpose single- or multi-chip microprocessor (e.g., an Advanced RISC (Reduced Instruction Set Computer) Machine (ARM)), a special purpose microprocessor (e.g., a digital signal processor (DSP)), a microcontroller, a programmable gate array, etc. The processor **1601** may be referred to as a central processing unit (CPU). Although just a single processor **1601** is shown in the computing system **1600** of FIG. **16**, in an alternative configuration, a combination of processors (e.g., an ARM and DSP) could be used.

[0120] The computing system **1600** also includes memory **1603** in electronic communication with the processor **1601**. The memory **1603** may be any electronic component capable of storing electronic information. For example, the memory **1603** may be embodied as random access memory (RAM), read-only memory (ROM), magnetic disk storage media, optical storage media, flash memory devices in RAM, on-board memory included with the processor **1601**, erasable programmable read-only memory (EPROM), electrically erasable programmable read-only memory (EEPROM) memory, registers, and so forth, including combinations thereof.

[0121] Instructions **1605** and data **1607** may be stored in the memory **1603**. The instructions **1605** may be executable by the processor **1601** to implement some or all of the methods, steps, operations, actions, or other functionality that is disclosed herein. Executing the instructions **1605** may involve the use of the data **1607** that is stored in the memory **1603**. Unless otherwise specified, any of the various examples of modules and components described herein may be implemented, partially or wholly, as instructions **1605** stored in memory **1603** and executed by the processor **1601**. Any of the various examples of data described herein may be among the data **1607** that is stored in memory **1603** and used during execution of the instructions **1605** by the processor **1601**.

[0122] The computing system **1600** may also include one or more communication interfaces **1609** for communicating with other electronic devices. The communication interface(s) **1609** may be based on wired communication technology, wireless communication technology, or both. Some examples of communication interfaces **1609** include a Universal Serial Bus (USB), an Ethernet adapter, a wireless adapter that operates in accordance with an Institute of Electrical and Electronics Engineers (IEEE) 802.11 wireless communication protocol, a Bluetooth® wireless communication adapter, and an infrared (IR) communication port.

[0123] A computing system **1600** may also include one or more input devices **1611** and one or more output devices **1613**. Some examples of input devices **1611** include a keyboard, mouse, microphone, remote control device, button, joystick, trackball, touchpad, and lightpen. One specific type of output device **1613** that is typically included in a computing system **1600** is a display device **1615**. Display devices **1615** used with embodiments disclosed herein may utilize any suitable image projection technology, such as liquid crystal display (LCD), light-emitting diode (LED), gas plasma, electroluminescence, or the like. A display controller **1617** may also be provided, for converting data **1607** stored in the memory **1603** into text, graphics, and/or moving images (as appropriate) shown on the display device **1615**. The computing system **1600** may also include other types of output devices **1613**, such as a speaker, a printer, etc.

[0124] The various components of the computing system **1600** may be coupled together by one or more buses, which may include a power bus, a control signal bus, a status signal bus, a data bus, etc. For the sake of clarity, the various buses are illustrated in FIG. **16** as a bus system **1619**.

[0125] The techniques described herein may be implemented in hardware, software, firmware, or any combination thereof, unless specifically described as being implemented in a specific manner.

Any features described as modules, components, or the like may also be implemented together in an integrated logic device or separately as discrete but interoperable logic devices. If implemented in software, the techniques may be realized at least in part by a non-transitory computer-readable medium having computer-executable instructions stored thereon that, when executed by at least one processor, perform some or all of the steps, operations, actions, or other functionality disclosed herein. The instructions may be organized into routines, programs, objects, components, data structures, etc., which may perform particular tasks and/or implement particular data types, and which may be combined or distributed as desired in various embodiments.

[0126] The steps, operations, and/or actions of the methods described herein may be interchanged with one another without departing from the scope of the claims. In other words, unless a specific order of steps, operations, and/or actions is required for proper functioning of the method that is being described, the order and/or use of specific steps, operations, and/or actions may be modified without departing from the scope of the claims.

[0127] In an example, the term “determining” (and grammatical variants thereof) encompasses a wide variety of actions and, therefore, “determining” can include calculating, computing, processing, deriving, investigating, looking up (e.g., looking up in a table, a database or another data structure), ascertaining and the like. Also, “determining” can include receiving (e.g., receiving information), accessing (e.g., accessing data in a memory) and the like. Also, “determining” can include resolving, selecting, choosing, establishing and the like.

[0128] The terms “comprising,” “including,” and “having” are intended to be inclusive and mean that there may be additional elements other than the listed elements. Additionally, it should be understood that references to “one embodiment” or “an embodiment” of the present disclosure are not intended to be interpreted as excluding the existence of additional embodiments that also incorporate the recited features. For example, any element or feature described in relation to an embodiment herein may be combinable with any element or feature of any other embodiment described herein, where compatible.

[0129] The present disclosure may be embodied in other specific forms without departing from its spirit or characteristics. The described embodiments are to be considered as illustrative and not restrictive. The scope of the disclosure is, therefore, indicated by the appended claims rather than by the foregoing description. Changes that come within the meaning and range of equivalency of the claims are to be embraced within their scope.

Claims

1. A computer-implemented method, comprising: determining that a hardware acceleration condition has been satisfied for a first function by: determining an estimated cost savings associated with performing a number of instructions of the first function via a hardware circuit compared to performing the first function via a general-purpose processor; and determining that the estimated cost savings exceed a hardware acceleration condition threshold value; based on the hardware acceleration condition threshold value being exceeded, configuring the hardware circuit to perform the first function; and causing the first function to be performed via the hardware circuit.
2. The computer-implemented method of claim 1, further comprising adding the first function to a function library that is implemented in the hardware circuit.
3. The computer-implemented method of claim 1, further comprising receiving, at a function processing service, a request to execute source code that includes instructions to perform the first function.
4. The computer-implemented method of claim 1, wherein the hardware circuit comprises: a field programmable gate array (FPGA) that uses hardware description language (HDL) code to perform the first function; a graphics processing unit (GPU) that uses GPU code to perform the first

function; a programmable application specific integrated circuit (ASIC) that uses customed compiled machine code to perform the first function; or a digital signal processor (DSP) that uses DSP code to perform the first function.

5. The computer-implemented method of claim 1, further comprising configuring the hardware circuit to perform a second function in parallel with the first function, wherein the second function is different from the first function.

6. The computer-implemented method of claim 5, further comprising configuring the hardware circuit to operate a plurality of instances of the hardware circuit in parallel.

7. The computer-implemented method of claim 1, wherein the estimated cost savings associated with performing the number of instructions of the first function via the hardware circuit compared to performing the first function via the general-purpose processor includes an amount of saved energy by performing the first function via the hardware circuit compared to the general-purpose processor.

8. The computer-implemented method of claim 1, wherein the estimated cost savings for the hardware circuit include: a first cost of the general-purpose processor per second of runtime; a second cost of consumed energy; and an energy-delay product.

9. A system, comprising: a processor; and a computer memory comprising instructions that, when executed by the processor, cause the system to perform operations comprising: determining that a hardware acceleration condition has been satisfied for a first function by: determining an estimated cost savings associated with performing a number of instructions of the first function via a hardware circuit compared to performing the first function via a general-purpose processor; and determining that the estimated cost savings exceed a hardware acceleration condition threshold value; based on the hardware acceleration condition threshold value being exceeded, configuring the hardware circuit to perform the first function; and causing the first function to be performed via the hardware circuit.

10. The system of claim 9, further comprising additional instructions that, when executed by the processor, cause the system to perform the operations comprising adding the first function to a function library that is implemented in the hardware circuit.

11. The system of claim 9, further comprising additional instructions that, when executed by the processor, cause the system to perform the operations comprising receiving, at a function processing service, a request to execute source code that includes instructions to perform the first function.

12. The system of claim 9, wherein the hardware circuit comprises: a field programmable gate array (FPGA) that uses hardware description language (HDL) code to perform the first function; a graphics processing unit (GPU) that uses GPU code to perform the first function; a programmable application specific integrated circuit (ASIC) that uses customed compiled machine code to perform the first function; or a digital signal processor (DSP) that uses DSP code to perform the first function.

13. The system of claim 9, further comprising additional instructions that, when executed by the processor, cause the system to perform the operations comprising configuring the hardware circuit to perform a second function in parallel with the first function, wherein the second function is different from the first function.

14. The system of claim 13, further comprising additional instructions that, when executed by the processor, cause the system to perform the operations comprising configuring the hardware circuit to operate a plurality of instances of the hardware circuit in parallel.

15. The system of claim 9, wherein the estimated cost savings associated with performing the number of instructions of the first function via the hardware circuit compared to performing the first function via the general-purpose processor includes an amount of saved energy by performing the first function via the hardware circuit compared to the general-purpose processor.

16. A non-transitory computer-readable storage medium comprising instructions that, when

executed by a processor, cause a computing device to carry out operations comprising: determining that a hardware acceleration condition has been satisfied for a first function by: determining an estimated cost savings associated with performing a number of instructions of the first function via a hardware circuit compared to performing the first function via a general-purpose processor; and determining that the estimated cost savings exceed a hardware acceleration condition threshold value; based on the hardware acceleration condition threshold value being exceeded, configuring the hardware circuit to perform the first function; and causing the first function to be performed via the hardware circuit.

17. The non-transitory computer-readable storage medium of claim 16, further comprising adding the first function to a function library that is implemented in the hardware circuit.

18. The non-transitory computer-readable storage medium of claim 16, further comprising receiving, at a function processing service, a request to execute source code that includes instructions to perform the first function.

19. The non-transitory computer-readable storage medium of claim 16, further comprising configuring the hardware circuit to perform a second function in parallel with the first function, wherein the second function is different from the first function.

20. The non-transitory computer-readable storage medium of claim 19, further comprising configuring the hardware circuit to operate a plurality of instances of the hardware circuit in parallel.
