

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250258914

Kind Code

A1

Publication Date

August 14, 2025

Inventor(s)

LAGO, III; Serafin et al.

DETECTING ANOMALOUS LOADING OF DYNAMIC-LINK LIBRARIES

Abstract

Anomalous loading of a dynamic-link library (DLL) is detected on a desktop computer. When the file of the DLL is not a known normal file, the locality sensitive (LSH) values of files of known versions of a known normal DLL that corresponds to the DLL are obtained from an LSH repository and compared to the LSH value of the file of the DLL. The smallest distance between the LSH value of the file of the DLL and each of the LSH values from the repository is selected for comparison to a risk threshold. The Loading of the DLL is detected to be an anomaly when the smallest distance is not less than the risk threshold.

Inventors: LAGO, III; Serafin (Manila, PH), ROAN; John David (Manila, PH), GATBUNTON; Manuel (Bulacan, PH)

Applicant: Trend Micro Incorporated (Tokyo, JP)

Family ID: 1000008283911

Appl. No.: 18/441784

Filed: February 14, 2024

Publication Classification

Int. Cl.: G06F21/56 (20130101); G06F9/445 (20180101)

U.S. Cl.:

CPC G06F21/566 (20130101); G06F9/44521 (20130101); G06F2221/034 (20130101)

Background/Summary

TECHNICAL FIELD

[0001] The present disclosure is directed to cybersecurity.

BACKGROUND

[0002] A Windows-based system is a desktop computer that runs the MICROSOFT WINDOWS operating system. Although, there are different versions of the MICROSOFT WINDOWS operating system, these versions share common features to maintain compatibility. The MICROSOFT WINDOWS operating system is currently the most popular desktop operating system in the world and is thus the target of numerous cyberattacks.

[0003] A dynamic-link library (DLL) is a module that contains functions or data that can be called by another module, such as an application program or another DLL. A DLL is not directly executed, but is instead loaded by an application program. The MICROSOFT WINDOWS operating system searches for DLLs to be loaded in a particular order, which is exploited by some malware to load a malicious DLL by an attack technique referred to as “DLL side-loading”. Examples of malware that perform DLL side-loading include qakbot, icedid, bumbleloader, and blackcat.

[0004] Conventional solutions against DLL side-loading are reactive (as opposed to proactive) and/or may not be effective in detecting DLL side-loading. For example, behavior monitoring detects malware based on process sequence or application programming interface (API) invocations. Behavior monitoring is an ineffective or reactive solution against DLL side-loading because the application program that loads the malicious DLL is typically legitimate and executes normally.

[0005] Malicious DLLs may be detected by file signature matching. This is a reactive solution in that there are numerous DLL export-dependent application programs, and any one of those DLLs may be modified to perform DLL side-loading.

[0006] Machine learning techniques may be employed to detect malicious DLLs based on file features. This solution is applicable to malicious DLLs that are heavily modified versions of legitimate DLLs, but will be ineffective if the modification to the legitimate DLL is minimal.

[0007] Malicious DLLs may be detected by locality-sensitive hash comparison, e.g., using the TLSH algorithm. This is a reactive solution because samples of malicious DLLs must first be obtained.

BRIEF SUMMARY

[0008] In one embodiment, a method of detecting an anomalous loading of a dynamic-link library (DLL) includes detecting loading of the DLL by an application program on a desktop computer. It is determined if a file of the DLL is a known normal file. When the file of the DLL is not a known normal file, a locality sensitive hash (LSH) value of the file of the DLL is calculated and a plurality of centroids of a known normal DLL having a file with the same filename as the file of the DLL is retrieved, each centroid of the plurality of centroids is representative of a cluster of a plurality of clusters, each of the plurality of clusters comprising LSH values of files of a known version of the known normal DLL. A distance between the LSH value of the file of the DLL and each of the plurality of centroids is calculated. A smallest distance between the LSH value of the file of the DLL and each of the plurality of centroids is selected. The smallest distance is compared to a risk threshold. The loading of the DLL is detected as an anomaly in response to the smallest distance is not less than the risk threshold.

[0009] In another embodiment, a system for detecting anomalous loading of a DLL includes a desktop computer and a backend system. The desktop computer is configured to: detect loading of a DLL by an application program on the desktop computer; determine if a file of the DLL is a known normal file; when the file of the DLL is not a known normal file, calculate an LSH value of the file of the DLL and provide the LSH value and a filename of the file of the DLL to a backend system. The backend system is configured to: retrieve a plurality of centroids of a known normal

DLL that has a same filename as the file of the DLL, each centroid of the plurality of centroids is representative of a cluster of a plurality of clusters, each of the plurality of clusters comprising LSH values of files of a known version of the known normal DLL; calculate a distance between the LSH value of the file of the DLL and each of the plurality of centroids; select a smallest distance between the LSH value of the file of the DLL and each of the plurality of centroids; compare the smallest distance to a risk threshold; and detect that the loading of the DLL is an anomaly in response to the smallest distance is not less than the risk threshold.

[0010] In yet another embodiment, a method of detecting an anomalous loading of a DLL includes detecting loading of the DLL on a desktop computer. An LSH value of the file of the DLL is calculated. A plurality of centroids of a known normal DLL is retrieved from a repository, each centroid of the plurality of centroids is representative of a cluster of a plurality of clusters, each of the plurality of clusters comprising LSH values of files of a known version of the known normal DLL. A distance between the LSH value of the file of the DLL and each of the plurality of centroids is calculated. A smallest distance between the LSH value of the file of the DLL and each of the plurality of centroids is selected. The smallest distance is compared to a risk threshold. The loading of the DLL is detected as an anomaly in response to the smallest distance is not less than the risk threshold.

[0011] These and other features of the present disclosure will be readily apparent to persons of ordinary skill in the art upon reading the entirety of this disclosure, which includes the accompanying drawings and claims.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0012] A more complete understanding of the subject matter may be derived by referring to the detailed description and claims when considered in conjunction with the following figures, wherein like reference numbers refer to similar elements throughout the figures.

[0013] FIG. 1 shows a block diagram of a Windows-based desktop computer, in accordance with an embodiment of the present invention.

[0014] FIGS. 2-4 show a flow diagram of a method of detecting an anomalous loading of a DLL, in accordance with an embodiment of the present invention.

[0015] FIG. 5 shows a flow diagram of a method of detecting an anomalous loading of a DLL, in accordance with an embodiment of the present invention.

[0016] FIG. 6 shows a block diagram of a computer system that may be employed with embodiments of the present invention.

DETAILED DESCRIPTION

[0017] In the present disclosure, numerous specific details are provided, such as examples of systems, components, and methods, to provide a thorough understanding of embodiments of the invention. Persons of ordinary skill in the art will recognize, however, that the invention can be practiced without one or more of the specific details. In other instances, well-known details are not shown or described to avoid obscuring aspects of the invention.

[0018] FIG. 1 shows a block diagram of a Windows-based desktop computer **100**, in accordance with an embodiment of the present invention. The desktop computer **100** may be a laptop, personal computer, or other computer that runs a desktop operating system. The desktop computer **100** comprises at least one processor **101** (e.g., an INTEL processor; AMD processor), a data storage device **102** (e.g., spinning hard disk drive, solid state drive), a main memory **103** (e.g., Dynamic Random Access Memory (DRAM)), a computer network interface **104** (e.g., network interface card (NIC)), and other devices **105** (e.g., display screen, keyboard, graphics card, mouse). The at least one processor **101** may have a plurality of processing cores. Instructions of program code (i.e.,

software) are stored in the main memory **103** for execution by the at least one processor **101** to cause the desktop computer **100** to perform steps in accordance with the program code. In one embodiment, the desktop computer **100** belongs to an enterprise computer network. The enterprise computer network includes a plurality of desktop computers, but only one is shown for clarity of illustration.

[0019] The desktop computer **100** is Windows-based in that it runs a desktop operating system **115** that is the MICROSOFT WINDOWS operating system, which is currently a version 11. As can be appreciated, embodiments of the present invention are equally applicable to other versions of the MICROSOFT WINDOWS operating system, such as the version 10 and future compatible versions.

[0020] The data storage device **102** stores files of the desktop computer **100**, which in the example of FIG. **1** includes files of an application program **112**, a dynamic-link library (DLL) **113**, antivirus program **114**, and the desktop operating system **115**. In the present disclosure, references to a locality sensitive hash (LSH) value or filename of a DLL are with respect to the file of the DLL. That is, an LSH value of a DLL and a filename of the DLL refers to the LSH value and the filename, respectively, of the file of the DLL. There are several DLLs and application programs on the desktop computer **100**, but only one of each is shown for clarity of illustration.

[0021] The antivirus program **114** comprises one or more components for performing cybersecurity on the desktop computer **100**. The antivirus program **114** may be configured to perform spam email detection, file scanning, and other conventional cybersecurity operations. In the example of FIG. **1**, the antivirus program **114** further performs anomalous DLL loading detection as described herein. As its name indicates, anomalous DLL loading detection involves determining whether the loading of a DLL is an anomaly. The anomaly detection does not provide an indication of whether the loaded DLL is in fact malicious. However, detecting that the loading of the DLL is an anomaly advantageously allows a user of the desktop computer **100** or an administrator (e.g., security operations center (SOC) personnel, information technology (IT) personnel) of the enterprise computer network to be alerted. This allows the loading of the DLL to be further investigated for possible DLL side-loading.

[0022] An executable program can be started automatically (e.g., on startup of the desktop computer) or manually (e.g., by double clicking on an icon of the executable program). When the executable program is started, the desktop operating system creates a process of the executable program and loads the image of the process into the allocated memory space of the process. The executable program, more specifically its process, thereafter runs on the desktop computer.

[0023] In the example of FIG. **1**, the application program **112** and the antivirus program **114** are executable programs. A process **130** is a process of the application program **112**. An anomaly detector **132** is a process of the antivirus program **114** for detecting anomalous DLL loading.

[0024] In one embodiment, the antivirus program **114** works in conjunction with a goodwill collection system **140** and a backend system **150** to detect an anomalous DLL loading. The desktop computer **100** may communicate with the goodwill collection system **140** and the backend system **150** over the public Internet.

[0025] For purposes of the present disclosure, “known” refers to something that is known to cybersecurity experts or other trustworthy entity. For example, a known normal file is a file that has been verified or vetted by cybersecurity experts or other trustworthy entities to be safe/legitimate. Similarly, a known normal DLL is a DLL that has been verified or vetted by cybersecurity experts or other trustworthy entities to be safe/legitimate.

[0026] The goodwill collection system **140** comprises a cloud computing platform or a dedicated server computer system that collects, from various sources, samples of known normal files that are maintained in a collection **141** of known normal files. The collection **141** may be stored on local storage, network attached storage, or cloud storage of the goodwill collection system **140**. A target (i.e., being evaluated) file may be provided to the goodwill collection system **140**, which compares

the target file to known normal files in the collection **141** to determine if the target file is a known normal file. The goodwill collection system **140** may be a conventional goodwill collection system, such as the Normal File Collection (NFC)/Goodware Resource and Information Database (GRID) of Trend Micro Incorporated.

[0027] The backend system **150** comprises a computer system that hosts an LSH repository **151**. The LSH repository **151** may be implemented on local storage, network attached storage, or cloud storage of the backend system **150**. The backend system **150** may be implemented on a dedicated server computer system or a cloud computing platform (e.g., Amazon Web Services (AWS)TM platform), for example. The backend system **150** includes program code for performing the functions of the backend system **150** as described herein.

[0028] The LSH repository **151** comprises a collection of clustered LSH values of known versions of known normal DLLs. More particularly, each known normal DLL may have one or more known versions, and each version may have one or more associated files, each with its own LSH value in the repository **151**. Accordingly, each known normal DLL may have several clusters of LSH values, with each cluster comprising LSH values of a particular version of the known normal DLL. Each cluster may have a centroid, which is an LSH value that is representative of the LSH values in the cluster.

[0029] Locality sensitive hashes are well known. Unlike other types of hashes, such as cryptographic hashes, small changes to a file will result in different but very similar LSH values. That is, with a suitable locality sensitive hashing algorithm, such as the TLSH algorithm, a file and small changes to the file will likely yield different but very similar LSH values. The mathematical distance (“distance”) between LSH values of two files may be calculated to determine similarity of the two files. The smaller the distance, the more similar the LSH values and thus the files. LSH values may be deemed to be similar when they are within a threshold distance; the threshold distance may be adjusted based on false positive/false negative requirements.

[0030] In one embodiment, all LSH values described herein are calculated using the TLSH algorithm. Open source program code of the TLSH algorithm for calculating the TLSH value of a file is generally available on the Internet and other sources. Open source program code of the TLSH algorithm may also include a distance calculation function, which may be used to calculate similarity of two TLSH values.

[0031] In an example operation, the anomaly detector **132** is loaded to execute in the memory **103** when the antivirus program **114** is started (see arrow **121**). Similarly, the process **130** is loaded to execute in the memory **103** when the application program **112** is started (see arrow **122**). The process **130** thereafter loads the DLL **113** (see arrow **123**). The DLL **113** as loaded in the memory **103** is also referred to herein as a loaded DLL **131**. The application program **114** may be a trusted or signed program, and accordingly allowed by the desktop operating system **115** to execute. However, the DLL **113** is not necessarily normal. That is, the DLL **113** is potentially malicious and is side-loaded by the application program **114** to avoid detection.

[0032] The anomaly detector **132** detects the loading of the DLL **113** (see arrow **124**). The anomaly detector **132** may detect the loading of a DLL a variety of conventional ways without detracting from the merits of the present invention. For example, the loading of the DLL **113** may be detected by hooking a suitable API, such as the LdrLoadDll function in the MICROSOFT WINDOWS operating system.

[0033] The anomaly detector **132** identifies the DLL **113** as having been loaded. The file of the DLL **113**, including its file path and filename, may be identified from the hooked function, for example. The anomaly detector **132** consults the goodwill collection system **140** to determine whether or not the file of the DLL **113** is a known normal file (see arrow **125**). In one embodiment, the anomaly detector **132** provides the file or hash of the DLL **113** to the goodwill collection system **140** over the public Internet. The goodwill collection system **140** compares the file or hash of the DLL **113** to those of known DLLs to determine if the file of the DLL **113** is normal.

[0034] When the goodwill collection system **140** does not indicate that the file of the DLL **113** is normal, the anomaly detector **132** consults the backend system **150** to determine whether the loading of the DLL **113** is anomalous (see arrow **126**). In one embodiment, the anomaly detector **132** calculates the LSH value of the DLL **113**, and provides the LSH value and filename of the DLL **113** to the backend system **150** over the public Internet. The backend system **150** uses the LSH value of the DLL **113** to determine whether the loading of the DLL **113** is an anomaly. The anomaly detector **132**, other component of the antivirus program **114**, or the backend system **150** may perform a response action in response to detecting an anomalous DLL loading. The response action may include raising an alert, such as sending a notification to a user or administrator. The notification may be displayed as a message on a display screen, sent by email or text message, etc.

[0035] The following is an example algorithm that may be employed to implement the anomaly detector **132** working in conjunction with the backend system **150**. [0036] (1) Hook the DLL loading function within a process [0037] (2) Get DLL information from the parameters of the DLL loading function (filename, path) [0038] (3) Compute the hash value of the loaded DLL [0039] (4) Connect to Goodware Collection System [0040] (a) Send the computed hash value of the loaded DLL to check if known DLL to the Goodware Collection System [0041] (5) Calculate LSH value of the loaded DLL if not a known DLL file [0042] (6) Check for anomaly [0043] (a) Query filename if exist within LSH repository [0044] (a.1) Pass if filename does not exist [0045] (b) Calculate the distances of the loaded DLL from the centroids of the LSH Forests [0046] (c) Get the smallest value from the distances [0047] (d) Compare the smallest value with the normal distance threshold [0048] (d.1) Return normal if below the normal distance threshold [0049] (e) Compare the smallest value with the low risk threshold [0050] (e.1) Return low risk if below the low risk threshold [0051] (f) Compare the value with the medium risk threshold [0052] (f.1) Return medium risk if below the medium risk threshold [0053] (g) Return high risk

[0054] In one embodiment, for continuous detection improvement, the anomaly detector **132** sends the backend system **150** feedback data regarding anomalies or security events detected on the desktop computer **100**.

[0055] FIGS. 2-4 show a flow diagram of a method **200** of detecting anomalous loading of a DLL, in accordance with an embodiment of the present invention. The method **200** is explained with certain functionalities performed by the desktop computer **100** and other functionalities performed by the backend system **150** for illustration purposes only. As can be appreciated, other components may also be employed, and functionalities may be distributed among different components.

[0056] The method **200** includes an initial filtering stage (FIG. 2, **220**) that includes detecting the loading of a DLL and determining whether the file of the DLL is a known normal file. The method **200** further includes an anomaly detection stage (FIG. 3, **300**), which detects whether the loading of a DLL is an anomaly.

[0057] Referring to FIG. 2, the method **200** begins when the application program **112** is started on the desktop computer **100**, causing the desktop operating system **115** to create the process **130** of the application program **112** (step **201**). The desktop operating system **115** loads the image of the process **130** in the memory **103** to start execution of the process **130** (step **202**). The process **130** loads the DLL **113**, which is detected by the anomaly detector **132**. The anomaly detector **132** consults the goodwill collection system **140** to determine if the file of the DLL **113** is a known normal file (step **203**). When the file of the DLL **113** is a known normal file, the anomaly detector **132** allows the DLL **113** to pass (step **203** to step **206**), i.e., does not perform any other action on the DLL **113**.

[0058] Otherwise, when the file of the DLL **113** is not a known normal file, the anomaly detector **132** determines, prior to the anomaly detection stage, if the LSH repository **151** has a known normal DLL that corresponds to the DLL **113**. In one embodiment, the correspondence is determined by filename. More particularly, the anomaly detector **132** determines if the LSH repository **151** includes a known normal DLL with a same filename as the DLL **113** (step **204**). For

example, the anomaly detector **132** may provide the filename of DLL **113** to the backend system **150**, which determines whether the LSH repository **151** has an entry for the filename. Alternatively, a listing of filenames of known normal DLLs in the LSH repository **151** may be maintained locally on the desktop computer **100** to facilitate access by the anomaly detector **132**. This allows the anomaly detector **132** to determine if the LSH repository **151** includes the filename of the DLL **113** before expending network bandwidth to communicate with the backend system **150**.

[0059] When the LSH repository **151** has an entry for a known normal DLL with the same filename as the DLL **113**, the method **200** continues to the anomaly detection stage (see arrow **205**) to determine if the loading of the DLL **113** is an anomaly. Otherwise, if the LSH repository **151** does not have an entry for a known normal DLL with the same filename as the DLL **113**, the anomaly detector **132** simply allows the DLL **113** to pass (step **204** to step **206**).

[0060] Referring to FIG. 3, the method **200** continues to the anomaly detection stage when the LSH repository **151** has an entry for a known normal DLL with the same filename as the DLL **113** (see arrow **301**). The LSH value of the DLL **113** is calculated as part of the anomaly detection stage (step **302**). The calculation of the LSH value of the DLL **113** may be performed by the anomaly detector **132** on the desktop computer **100** and provided to the backend system **150**. Alternatively, the anomaly detector **132** may provide the file of the DLL **113** to the backend system **150**, which then calculates the LSH value of the file. In one embodiment, all LSH values described herein are calculated using the TLSH algorithm. That is, the LSH values of the embodiments are TLSH values.

[0061] In the anomaly detection stage, the backend system **150** receives the filename of the DLL **113** from the anomaly detector **132**. The backend system **150** queries the LSH repository **151** for centroids of a known normal DLL that has the same filename as the DLL **113** (step **303**). The LSH repository **151** is now further explained with reference to FIG. 4.

[0062] Referring to FIG. 4, the LSH repository **151** comprises a plurality of LSH forests **401** (i.e., **401-1**, **401-2**, . . . , **401-n**). Each LSH forest **401** comprises a plurality of clusters **402** (i.e., **402-1**, **402-2**, . . . , **402-n**), with each cluster **402** having a centroid **403** (i.e., **403-1**, **403-2**, . . . , **403-n**).

[0063] Each LSH forest **401** is for a particular known normal DLL, which in one embodiment is identified by its filename. Each cluster **402** is for a particular version of the known normal DLL. LSH values of a version of a known normal DLL are clustered together in a same cluster **402**. Each cluster **402** has a centroid **403**. A centroid **403**, which is also an LSH value, is representative of the LSH values of the corresponding cluster **402**. A centroid **403** may be described as an average, median, or some other relationship between the members (i.e., LSH values) of the cluster **402**, depending on the clustering algorithm employed.

[0064] As a particular example, the LSH forest **401-1** may be for a DLL having a filename “sample.dll”. The “sample.dll” has different known versions 1, 2, etc., with each version having been found to have different files. The clusters **402-1**, **402-2**, **402-n** are a result of applying a clustering algorithm on the LSH values of the different versions of “sample.dll”. The centroids **403-1**, **403-2**, **403-n** are the representative of their corresponding clusters resulting from the application of the clustering algorithm on the LSH values.

[0065] Given a target filename (see arrow **410**), the LSH repository **151** may be queried for all centroids **403** of an LSH forest **401** of a known normal DLL whose filename is the same as the target filename (see arrow **411**). For example, assuming the LSH forest **401-1** is for a known normal DLL with the same filename as the DLL **113**, the centroids **403-1**, **403-2**, . . . , **403-n** of the LSH forest **401-1** are returned as a response to a query that has the filename of the DLL **113** as input.

[0066] Referring back to FIG. 3, in response to the query with the filename of the DLL **113** as input (see arrow **304**), the backend system **150** receives all centroids **403** of an LSH forest **401** of a known normal DLL that has the same filename as the DLL **113** (see arrow **305**). The backend system **150** calculates the distance between the LSH value of the DLL **113** and each of the

centroids **403** retrieved from the LSH repository **151** (step **306**) and selects the smallest distance (step **307**). That is, the backend system **150** calculates the distance between the LSH value of the DLL **113** and a centroid **403-1**, the distance between the LSH value of the DLL **113** and a centroid **403-2**, etc. returned by the LSH repository **151**, and selects the smallest calculated distance.

[0067] In one embodiment, different risk thresholds are predetermined for different risk levels, which in one embodiment are normal, low risk, medium risk, and high risk. The medium risk threshold is greater than the low risk threshold, which is greater than the normal threshold. Each risk level indicates a degree of security risk posed by an anomalous DLL loading. In one embodiment, a loading that poses a low risk, medium risk, or high risk is deemed to be an anomaly.

[0068] The backend system **150** compares the smallest distance to the normal threshold (step **308**). When the smallest distance is less than the normal threshold, the backend system **150** deems the loading of the DLL **113** to be normal and so notifies the anomaly detector **132** (see arrow **309**), which allows the DLL **113** to pass (see FIG. 2, arrow **207**). Otherwise, when the smallest distance is not less than the normal threshold, the backend system **150** compares the smallest distance to the low risk threshold (step **308** to step **310**).

[0069] When the smallest distance is less than the low risk threshold, the backend system **150** flags the loading of the DLL **113** as a low risk anomaly (step **311**). Otherwise, when the smallest distance is not less than the low risk threshold, the backend system **150** compares the smallest distance to the medium risk threshold (step **310** to step **312**).

[0070] When the smallest distance is less than the medium risk threshold, the backend system **150** flags the loading of the DLL **113** as a medium risk anomaly (step **313**). Otherwise, when the smallest distance is not less than the medium risk threshold, the backend system flags the loading of the DLL **113** as a high risk anomaly (step **314**).

[0071] The loading of the DLL **113** is an anomaly when the smallest distance is not less than the normal threshold. The risk level indicates the degree of the anomaly to allow for prioritization or other purpose. The backend system **150** informs the anomaly detector **132** of the result of the evaluation of the loading of the DLL **113**, for example as a return value of an API of the backend system **150**.

[0072] FIG. 5 shows a flow diagram of a method of detecting an anomalous loading of a DLL, in accordance with an embodiment of the present invention. The method of FIG. 5 may be performed using previously-discussed components. As can be appreciated, other components may also be employed without detracting from the merits of the present invention.

[0073] In step **501**, the loading of a DLL by a process of an application program is detected.

[0074] In step **502**, it is determined if a file of the DLL is a known normal file.

[0075] In step **503**, when the file of the DLL is not a known normal file, the LSH value of the file of the DLL is calculated. In one embodiment, the LSH value of the file of the DLL is calculated using the TLSH algorithm. As can be appreciated, the LSH value of the file of the DLL may be calculated before or after determining whether the file of the DLL is a known normal file. Generally, the LSH value of the file of the DLL may be calculated before the distance calculation of step **505**.

[0076] In step **504**, a plurality of centroids of a known normal DLL that corresponds to the DLL are retrieved from an LSH repository. In one embodiment, the correspondence between the known normal DLL and the DLL is by filename, i.e., the known normal DLL has a file with the same filename as the file of the DLL. Each centroid of the plurality of centroids is representative of a cluster of a plurality of clusters, with each of the plurality clusters comprising LSH values of files of a known version of the known normal DLL. In one embodiment, LSH values in the LSH repository, including the plurality of centroids, are TLSH values.

[0077] In step **505**, a distance between the LSH value of the file of the DLL and each of the centroids is calculated. This results in a plurality of distances, with each distance being between the LSH value of the file of the DLL and a centroid of the plurality of centroids.

[0078] In step **506**, a smallest distance among the plurality of distances is selected.

[0079] In step **507**, the smallest distance is compared to one or more risk thresholds.

[0080] In step **508**, the loading of the DLL is detected to be an anomaly responsive to the smallest distance is not less than a risk threshold.

[0081] In step **509**, an alert is raised responsive to detecting that the loading of the DLL is an anomaly.

[0082] FIG. **6** shows a block diagram of a computer system **600** that may be employed with embodiments of the present invention. The computer system **600** may be employed as a desktop computer, a goodware collection system, backend system, or other computer for performing functionality described herein. The computer system **600** may have fewer or more components to meet the needs of a particular cybersecurity application. The computer system **600** may include at least one processor **601**. The computer system **600** may have one or more buses **603** coupling its various components. The computer system **600** may include one or more user input devices **602** (e.g., keyboard, mouse), one or more data storage devices **606**, a display screen **604**, a computer network interface **605**, and a main memory **608**. The computer network interface **605** may be coupled to a computer network **607**, which in this example includes the public Internet.

[0083] The computer system **600** is a particular machine as programmed with one or more software modules **609**, comprising instructions of program code stored non-transitory in the main memory **608** for execution by at least one processor **601** to cause the computer system **600** to perform corresponding programmed steps. An article of manufacture may be embodied as computer-readable storage medium including instructions that when executed by at least one processor **601** cause the computer system **600** to be operable to perform the functions of the one or more software modules **609**.

[0084] While specific embodiments of the present invention have been provided, it is to be understood that these embodiments are for illustration purposes and not limiting. Many additional embodiments will be apparent to persons of ordinary skill in the art reading this disclosure

Claims

1. A method of detecting an anomalous loading of a dynamic-link library (DLL), the method comprising: detecting loading of a DLL by an application program on a desktop computer; determining if a file of the DLL is a known normal file; in response to the file of the DLL is not a known normal file, calculating a locality sensitive hash (LSH) value of the file of the DLL; retrieving a plurality of centroids of a known normal DLL that corresponds to the DLL, each centroid of the plurality of centroids is representative of a cluster of a plurality of clusters, each of the plurality of clusters comprising LSH values of files of a known version of the known normal DLL; calculating a distance between the LSH value of the file of the DLL and each of the plurality of centroids to generate a plurality of distances; selecting a smallest distance among the plurality of distances; comparing the smallest distance to a first risk threshold; and detecting that the loading of the DLL is an anomaly in response to the smallest distance is not less than the first risk threshold.
2. The method of claim 1, further comprising: raising an alert in response to detecting that the loading of the DLL is an anomaly.
3. The method of claim 2, wherein raising the alert includes sending a notification to a user of the desktop computer or an administrator of an enterprise computer network that includes the desktop computer.
4. The method of claim 1, further comprising: comparing the smallest distance to a second risk threshold that is greater than the first risk threshold; and raising an alert that the loading of the DLL is a low risk anomaly in response to the smallest distance is not less than the first risk threshold but is less than the second risk threshold.
5. The method of claim 1, further comprising: comparing the smallest distance to a second risk

threshold that is greater than the first risk threshold; in response to the smallest distance is not less than the second risk threshold, comparing the smallest distance to a third risk threshold that is greater than the second risk threshold; and raising an alert that the loading of the DLL is a medium risk anomaly in response to the smallest distance is less than the third risk threshold.

6. The method of claim 1, further comprising: comparing the smallest distance to a second risk threshold that is greater than the first risk threshold; in response to the smallest distance is not less than the second risk threshold, comparing the smallest distance to a third risk threshold that is greater than the second risk threshold; and raising an alert that the loading of the DLL is a high risk anomaly in response to the smallest distance is not less than the third risk threshold.

7. The method of claim 1, wherein each of the plurality of centroids and the LSH value of the file of the DLL is a TLSH value.

8. A system for detecting anomalous loading of a dynamic-link library (DLL) on a desktop computer, the system comprising: a desktop computer comprising at least one processor and a memory, the memory of the desktop computer storing instructions that when executed by the at least one processor of the desktop computer cause the desktop computer to: detect loading of a DLL by an application program; determine if a file of the DLL is a known normal file; calculate a locality sensitive hash (LSH) value of the file of the DLL; and in response to the file of the DLL is not a known normal file, provide the LSH value and a filename of the file of the DLL to a backend system; and the backend system comprising at least one processor and a memory, the memory of the backend system storing instructions that when executed by the at least one processor of the backend system cause the backend system to: retrieve a plurality of centroids of a known normal DLL having a file that has a same filename as the file of the DLL, each centroid of the plurality of centroids is representative of a cluster of a plurality of clusters, each of the plurality of clusters comprising LSH values of files of a known version of the known normal DLL; calculate a distance between the LSH value of the file of the DLL and each of the plurality of centroids to generate a plurality of distances; select a smallest distance among the plurality of distances; compare the smallest distance to a risk threshold, and detect that the loading of the DLL is an anomaly in response to the smallest distance is not less than the risk threshold.

9. The system of claim 8, wherein the instructions stored in the memory of the backend system, when executed by the at least one processor of the backend system, cause the backend system to raise an alert in response to detecting that the loading of the DLL is an anomaly.

10. The system of claim 8, wherein the instructions stored in the memory of the backend system, when executed by the at least one processor of the backend system, cause the backend system to raise an alert by sending a notification to a user of the desktop computer or an administrator of an enterprise computer network that includes the desktop computer.

11. The system of claim 8, wherein each of the plurality of centroids and the LSH value of the file of the DLL is a TLSH value.

12. The system of claim 8, further comprising a goodwill collection system that determines whether the file of the DLL is a known normal file.

13. A method of detecting an anomalous loading of a dynamic-link library (DLL), the method comprising: detecting loading of a DLL by an application program on a desktop computer; calculating a locality sensitive hash (LSH) value of the file of the DLL; receiving a plurality of centroids of a known normal DLL, each centroid of the plurality of centroids is representative of a cluster of a plurality of clusters, each of the plurality of clusters comprising LSH values of files of a known version of the known normal DLL; generating a plurality of distances by calculating a distance between the LSH value of the file of the DLL and each of the plurality of centroids; selecting a smallest distance among the plurality of distances; comparing the smallest distance to a risk threshold; and detecting that the loading of the DLL is an anomaly in response to the smallest distance is not less than the risk threshold.

14. The method of claim 13, further comprising: raising an alert in response to detecting that the

loading of the DLL is an anomaly.

15. The method of claim 14, wherein raising the alert includes sending a notification to a user of the desktop computer or an administrator of an enterprise computer network that includes the desktop computer.

16. The method of claim 13, wherein each of the plurality of centroids and the LSH value of the file of the DLL is a TLSH value.
