

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250259054

Kind Code

A1

Publication Date

August 14, 2025

Inventor(s)

Zhang; Jianguo et al.

SYSTEMS AND METHODS FOR A UNIFIED TRAINING FRAMEWORK OF LARGE LANGUAGE MODELS

Abstract

Embodiments described herein provide a unified LLM training pipeline that handles the diversity of various data structures and formats involving LLMs agent trajectories. These pipelines are specifically designed to transform incoming data into a standardized representation, ensuring compatibility across varied formats. Furthermore, the data collection undergoes a filtering process to ensure high-quality trajectories, adding an additional layer of refinement to the dataset. In this way, the training pipeline not only unifies trajectories across environments but also enhances the overall quality and reliability of the collected data for LLM training.

Inventors: Zhang; Jianguo (San Jose, CA), Lan; Tian (San Francisco, CA), Murthy; Rithesh (San Francisco, CA), Liu; Zhiwei (Palo Alto, CA), Yao; Weiran (Palo Alto, CA), Tan; Juntao (San Francisco, CA), Heinecke; Shelby (San Francisco, CA), Feng; Yihao (Austin, TX), Wang; Huan (Palo Alto, CA), Niebles; Juan Carlos (Mountain View, CA), Savarese; Silvio (Palo Alto, CA), Xiong; Caiming (Menlo Park, CA)

Applicant: Salesforce, Inc. (San Francisco, CA)

Family ID: 1000007909571

Appl. No.: 18/658899

Filed: May 08, 2024

Related U.S. Application Data

us-provisional-application US 63552164 20240211

Publication Classification

Int. Cl.: G06N3/08 (20230101)

Background/Summary

CROSS REFERENCE(S) [0001] The instant application is a nonprovisional of and claims priority under 35 U.S.C. 119 to U.S. provisional application No. 63/552,164, filed Feb. 11, 2024, which is hereby expressly incorporated by reference herein in its entirety.

TECHNICAL FIELD

[0002] The embodiments relate generally to neural networks and machine learning systems, and more specifically to a unified training framework for large language models (LLMs).

BACKGROUND

[0003] LLMs have been widely used as conversational agents in a variety of applications, such as information technology support, customer service, education, and/or the like. However, training LLMs for various tasks can be challenging because of various non-standardized data formats sourced from diverse dataset collections, especially those featuring interactions of multi-turns, as commonly observed in agent-relevant data. The heterogeneity in data structures, syntaxes, labeling conventions, and processing methods across datasets poses significant challenges, complicating the training and fine-tuning processes of LLMs.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0004] FIG. 1 is a simplified diagram illustrating an example architecture of a neural network training pipeline, according to embodiments described herein.

[0005] FIGS. 2A-2B provide a simplified diagram illustrating an example process of generating a training dataset for a hardware training platform of neural networks from heterogeneous multi-turn data, according to embodiments described herein.

[0006] FIGS. 3A-3B are example heterogeneous data records that are transformed into a homogeneous format as an agent trajectory, according to embodiments described herein.

[0007] FIG. 4 is an example diagram illustrating an example prompt for an LLM to generate a rating given an input of agent trajectory, according to embodiments described herein.

[0008] FIG. 5 is a simplified diagram illustrating a computing device implementing the LLM training pipeline framework described in FIGS. 1-4, according to one embodiment described herein.

[0009] FIG. 6 is a simplified diagram illustrating the neural network structure implementing the LLM training pipeline module described in FIG. 5, according to some embodiments.

[0010] FIG. 7 is a simplified block diagram of a networked system suitable for implementing the LLM training pipeline framework described in FIGS. 1-4 and other embodiments described herein.

[0011] FIG. 8 is an example logic flow diagram illustrating an example method of training an intelligent agent for carrying out actions in response to user interactions, according to embodiments described herein.

[0012] FIGS. 9A-9B are example diagrams illustrating example pseudo code segments for implementing the method shown in FIG. 8, according to embodiments described herein.

[0013] Embodiments of the disclosure and their advantages are best understood by referring to the detailed description that follows. It should be appreciated that like reference numerals are used to identify like elements illustrated in one or more of the figures, wherein showings therein are for

purposes of illustrating embodiments of the disclosure and not for purposes of limiting the same.

DETAILED DESCRIPTION

[0014] As used herein, the term “network” may comprise any hardware or software-based framework that includes any artificial intelligence network or system, neural network or system and/or any training or learning models implemented thereon or therewith.

[0015] As used herein, the term “module” may comprise hardware or software-based framework that performs one or more functions. In some embodiments, the module may be implemented on one or more neural networks.

[0016] As used herein, the term “Large Language Model” (LLM) may refer to a neural network based deep learning system designed to understand and generate human languages. An LLM may adopt a Transformer architecture that often entails a significant amount of parameters (neural network weights) and computational complexity. For example, LLM such as Generative Pre-trained Transformer (GPT) 3 has 175 billion parameters, Text-to-Text Transfer Transformers (T5) has around 11 billion parameters.

[0017] In view of the need for an improved training framework for LLMs on different datasets of diverse data formats, embodiments described herein provide a unified LLM training pipeline that trains an LLM using diverse training data having different data structures and formats on different tasks. Specifically, diverse training data involving different types of multi-turn interactions and/or executions may be transformed to a homogeneous multi-turn data format representing a user-agent interaction record (referred to as “trajectory”). A neural network based rating model may then rate each agent trajectory. Based on the ratings, a customized training dataset may be generated by sampling agent trajectories. The customized training datasets are then loaded to train the LLM.

[0018] In one embodiment, the training pipeline may transform incoming data into a homogeneous representation, ensuring compatibility across varied formats. Furthermore, the data collection undergoes a filtering process based on the rating of each training trajectory to ensure high-quality trajectories, adding an additional layer of refinement to the dataset. In this way, agent trajectories across environments may be efficiently used in training a neural network model such as an LLM in one training process or one training epoch, instead of multiple training processes that cost significant computational resources. With the computationally efficient training pipeline, neural network models can be trained using a large amount of diverse datasets to improve performance. Neural network technology is thus improved.

[0019] FIG. 1 is a simplified diagram illustrating an example architecture of a neural network training pipeline **100**, according to embodiments described herein. In one embodiment, neural network training pipeline **100** may comprise an agent trajectory standardizer **110**, an agent rater **120**, and a training data generator **130**. In one embodiment, agent trajectory standardizer **110**, agent rater **120**, and training data generator **130**. In one embodiment, one or more of the agent trajectory standardizer **110**, the agent rater **120**, and the training data generator **130** may comprise one or more neural network models that are housed at the same or different servers. For example, the agent trajectory standardizer **110** may be housed at one of data vendor servers **745**, **770** or **780** as shown in FIG. 7, while agent rater **120**, and the training data generator **130** may be situated at server **730** in FIG. 7.

[0020] In one embodiment, raw data in multi-turn environments **102** from heterogeneous data sources, such as an electronic commerce website activity log **102a**, an application programming interface (API) data activity log **102b**, a user-agent question and answering conversation **102c**, and/or the like, may be received.

[0021] The formats of agent data **102a-c** may vary significantly across different environments, posing significant difficulties and challenges in unifying data, training, and analyzing neural network models. For example, as shown in FIGS. 3A-3B, trajectories **302** and **304** from two distinct environments show different data formats different environments.

[0022] With reference back to FIG. 1, such trajectories **102a-c** of different formats from

heterogeneous data sources may be converted into a homogeneous multi-turn data format. In one embodiment, the agent trajectory standardizer **110** may construct a homogeneous JSON dictionary format to encapsulate all relevant content of each trajectory **102a-c**. For example, such format incorporates all important elements such as user query to store the initial user query, model name to identify the corresponding model and score to log the available model performance score. These elements can be used to differentiate between models and are poised to facilitate the development of pairwise samples for training methodologies such as DPO (Rafailov et al., 2023), self-reward (Yuan et al., 2024) and AI feedback (Guo et al., 2024) LLMs.

[0023] In one embodiment, the homogenous format further comprise an “other information” segment. Auxiliary trajectory information or specific notes are saved into other information, providing a reference for further analysis or model improvement initiatives.

[0024] In one embodiment, the homogenous format may comprise a “step” segment that captures the details of each interaction turn. A step such as **110a-110c** comprises three main components: input, output, and next observation. The input component consolidates the current prompt and a historical record of past interactions, serving as a comprehensive context for the interaction. The output component captures the model's predictions, detailing its decision-making and planning. The next observation component records the environment's feedback, essential for the feedback loop and system adaption.

[0025] In one embodiment, interaction history may be aggregated within the input component, effectively concatenating inputs and outputs from previous steps to construct a comprehensive context. Specifically, at the i .sup.th step, the input is formatted as input of step 1, Action: output of step 1, Observation: next observation of step 1, . . . , input of step $i-1$, Action: output of step $i-1$, Observation: next observation of step $i-1$. In this way, a detailed chronological account of interactions is constructed, facilitating a nuanced understanding of the trajectory.

[0026] In one embodiment, other customization of data compilation methods may be applied to aggregates interaction steps **110a-c**.

[0027] The challenge with agent trajectories extends to the evaluation of performance and quality. While some environments offer rewards as feedback for an agent's trajectory, such rewards are often tied to the task's final outcome rather than reflecting the quality of the trajectory itself. Consequently, a high reward does not necessarily indicate a flawless trajectory. For example, an agent might generate invalid actions during intermediate steps of a task. Here is partial trajectory from the Webshop environment (Yao et al., Webshop: Towards scalable real-world web interaction with grounded language agents, Advances in Neural Information Processing Systems, 35:20744-20757, 2022): model output of step 4: “click [old world style]”, observation of step 4: “You have clicked old world style.”; model output of step 5: “click [rope sausage]”, observation of step 5: “Invalid action!”; model output of step 6: “”, observation of step 6: “Invalid action!”; model output of step 7: “click [Buy Now]”, observation of step 7: “Your score (min 0.0, max 1.0): 1.0”, where the agent randomly clicks other buttons in Webshop website or generates empty responses before buying an item.

[0028] In one embodiment, as agent trajectories comprise scenarios when an agent interacts with complex environments such as websites, APIs, and tools. Such complexity can be heightened by the agent's capacity to communicate with other agents, navigate through diverse interfaces, and undertake tasks that require a sequence of interactions rather than single or limited exchanges. Thus, the challenge with agent trajectories extends to the evaluation of performance and quality. While some environments offer rewards as feedback for an agent's trajectory, such rewards are often tied to the task's final outcome rather than reflecting the quality of the trajectory itself. Consequently, a high reward does not necessarily indicate a flawless trajectory. For example, an agent might generate invalid actions during intermediate steps of a task. For instance, given a partial trajectory from an environment of an electronic commerce website: model output of step 4: “click [old world style]”, observation of step 4: “You have clicked old world style.”; model output

of step 5: “click [rope sausage]”, observation of step 5: “Invalid action!”; model output of step 6: “”, observation of step 6: “Invalid action!”; model output of step 7: “click [Buy Now]”, observation of step 7: “Your score (min 0.0, max 1.0): 1.0”, where the agent randomly clicks other buttons in the e-commerce website or generates empty responses before buying an item.

[0029] In view of the quality issues with agent trajectories, the Agent Rater **120** may receive and rate the standardized agent trajectory. In one embodiment, agent rater **120** may comprise a neural network model that is deployed and trained on a server (e.g., **730** in FIG. 7) using a training dataset of agent trajectory samples and annotated scores for each agent trajectory. Example training process may be described in relation to FIG. 6.

[0030] In one embodiment, agent rater **120** may adopt an external neural network model, such as a public LLM **121** such as Mistral, or a model API **122** such as ChatGPT, to perform the rating task. Unlike existing approaches that rate each triplet of (instruction, input, response) pair on general instruction data as there are usually single-turn or short conversations, agent rater **120** may rate the entire trajectory. For example, as shown in FIG. 4, an example input prompt may be used to include the agent trajectory, and then send to agent rater **120** to generate a rate the trajectory with a score 0-5 and an explanation, and they can be used to further develop better other rating models.

[0031] In one embodiment, the standardized agent trajectories with rating scores may be fed to a generator **130** to generate a training dataset. For example, generator **130** may comprise a combiner **131** to combine trajectories, a shuffler **132** to shuffle trajectories for random sampling, a customized filter **133** to filter trajectories with a lower than threshold rating score, an interleaver **134** to interleave trajectories, and/or the like.

[0032] In one embodiment, generator **130** may generate individual datasets by loading individual trajectories from agent rater **120**. Then, for each dataset, the filter **133** may be applied to further customize the selection of data just before feeding it into the training framework. For instance, data trajectories with relatively low scores will be further evaluated and removed. This dataset may be randomly shuffled with controlled seeding to ensure randomness and reproducibility.

[0033] In one embodiment, multiple individual datasets may be combined, e.g., by employing an `init.device.seed` function to diversify the controlled seeds based on the process ID when data parallelism across multiple devices. For example, a balanced distribution of data may be obtained in the training dataset by partitioning, shuffling and interleaving data across devices while preserving randomness, thus enhancing the robustness and reproducibility of training.

[0034] FIGS. 2A-2B provide a simplified diagram illustrating an example process of generating a training dataset for a hardware training platform of neural networks from heterogeneous multi-turn data, according to embodiments described herein. As illustrated in FIG. 2, heterogeneous multi-turn trajectories **202** may be converted to homogeneous trajectories **205**. The conversion may be performed by parsing and extracting fields from heterogeneous multi-turn trajectories **202** and re-arranging the format according to format **205**. In one implementation, the conversion may be performed by a neural network language model which receives an input of heterogeneous multi-turn trajectories **202** and generates an output of homogeneous trajectories **205**.

[0035] In one embodiment, the homogeneous trajectories **205** may be given an agent rating score (e.g., by agent rater **120** in FIG. 1). The output of agent rater **120** may comprise a homogenous data format of “user query,” “rater score,” “rate model” and “model name” showing the model such as GPT-4.0 that generates the rating score for the homogeneous trajectories **205**.

[0036] In one embodiment, homogeneous trajectories with agent rating scores **206** may then be loaded onto distributed trainers **210** for training one or more neural network models. Data loading process **208** may comprise combination, shuffling, filtering, interleaving data trajectories, as described in FIG. 1.

[0037] In one embodiment, distributed trainer **210** may comprise one or more hardware and/or software platform for training neural network models to perform different types of multi-turn agent tasks in different environments. For example, an agent model, such as a pre-trained Mixtral-8x7B-

Instruct-v0.1 model (Jiang et al., 2024) may be trained or finetuned using the generated dataset from the data loading process **208**. The distributed trainer **210** may be built on 8 Nvidia H100 GPUs, utilizing the QLoRA framework (Dettmers et al., 2023).

[0038] In one embodiment, during the training process, the agent model may be trained traversing each individual training dataset approximately **3** times on average. This multi-epoch training regimen facilitated comprehensive exposure to each individual dataset, enabling the model to effectively learn intricate patterns present in the training data.

[0039] FIGS. **3A-3B** are example heterogeneous data records that are transformed into a homogeneous format as an agent trajectory, according to embodiments described herein. For instance, as shown in FIG. **3A**, the trajectory **302** is taken from HotPotQA, which is a question answering dataset collected on the English Wikipedia, containing about 113K crowd-sourced questions that are constructed to require the introduction paragraphs of two Wikipedia articles to answer. Trajectory **302** consolidates the whole target trajectory into a single string under the prompt key. To process trajectory **302**, data fields such as user query, Thought, Model Action: along with Env Observation: i for each step $i \in [1, N]$ may be parsed and retrieved from a single string.

[0040] In another example, as shown in FIG. **3B**, ToolAIPaca is a framework that automatically generate a diverse tool-use corpus and learn generalized tool-use abilities on compact language models with minimal human intervention. Trajectory **304** from ToolAIPaca corpus may require the identification and matching of prompt inputs, model outputs, and observations at each step, followed by the accurate aggregation of trajectory history prior to proceeding to the next step.

[0041] For example, as shown in FIGS. **3A-3B**, the transformation of trajectories from environments such as HotpotQA and ToolAIPaca may comprise the defined step structure **110a-c**, where output aligns with the format specifications in the initial prompt input.

[0042] FIG. **4** is an example diagram illustrating an example prompt for an LLM to generate a rating given an input of agent trajectory, according to embodiments described herein. An example input prompt may be used to include the agent trajectory, and then send to agent rater **120** to generate a rate the trajectory with a score 0-5 and an explanation, and they can be used to further develop better other rating models.

Computer and Network Environment

[0043] FIG. **5** is a simplified diagram illustrating a computing device implementing the unified LLM training pipeline described in FIGS. **1-4**, according to one embodiment described herein. As shown in FIG. **5**, computing device **500** includes a processor **510** coupled to memory **520**.

Operation of computing device **500** is controlled by processor **510**. And although computing device **500** is shown with only one processor **510**, it is understood that processor **510** may be representative of one or more central processing units, multi-core processors, microprocessors, microcontrollers, digital signal processors, field programmable gate arrays (FPGAs), application specific integrated circuits (ASICs), graphics processing units (GPUs) and/or the like in computing device **500**. Computing device **500** may be implemented as a stand-alone subsystem, as a board added to a computing device, and/or as a virtual machine.

[0044] Memory **520** may be used to store software executed by computing device **500** and/or one or more data structures used during operation of computing device **500**. Memory **520** may include one or more types of machine-readable media. Some common forms of machine-readable media may include floppy disk, flexible disk, hard disk, magnetic tape, any other magnetic medium, CD-ROM, any other optical medium, punch cards, paper tape, any other physical medium with patterns of holes, RAM, PROM, EPROM, FLASH-EPROM, any other memory chip or cartridge, and/or any other medium from which a processor or computer is adapted to read.

[0045] Processor **510** and/or memory **520** may be arranged in any suitable physical arrangement. In some embodiments, processor **510** and/or memory **520** may be implemented on a same board, in a same package (e.g., system-in-package), on a same chip (e.g., system-on-chip), and/or the like. In some embodiments, processor **510** and/or memory **520** may include distributed, virtualized, and/or

containerized computing resources. Consistent with such embodiments, processor **510** and/or memory **520** may be located in one or more data centers and/or cloud computing facilities. [0046] In some examples, memory **520** may include non-transitory, tangible, machine readable media that includes executable code that when run by one or more processors (e.g., processor **510**) may cause the one or more processors to perform the methods described in further detail herein. For example, as shown, memory **520** includes instructions for LLM training pipeline module **530** that may be used to implement and/or emulate the systems and models, and/or to implement any of the methods described further herein. LLM training pipeline module **530** may receive input **540** such as an input text via the data interface **515** and generate an output **550** which may be a natural language processing task output.

[0047] The data interface **515** may comprise a communication interface, a user interface (such as a voice input interface, a graphical user interface, and/or the like). For example, the computing device **500** may receive the input **540** (such as a training text input) from a networked database via a communication interface. Or the computing device **500** may receive the input **540**, such as a user utterance, from a user via the user interface.

[0048] In some embodiments, the LLM training pipeline module **530** is configured to train one or more LLMs using consolidated standardized data from multiple data sources. The LLM training pipeline module **530** may further include a LLM submodule **531**, a data consolidation submodule **532**, a large action submodule **533** and and/or the like.

[0049] Some examples of computing devices, such as computing device **500** may include non-transitory, tangible, machine readable media that include executable code that when run by one or more processors (e.g., processor **510**) may cause the one or more processors to perform the processes of method. Some common forms of machine-readable media that may include the processes of method are, for example, floppy disk, flexible disk, hard disk, magnetic tape, any other magnetic medium, CD-ROM, any other optical medium, punch cards, paper tape, any other physical medium with patterns of holes, RAM, PROM, EPROM, FLASH-EPROM, any other memory chip or cartridge, and/or any other medium from which a processor or computer is adapted to read.

[0050] FIG. **6** is a simplified diagram illustrating the neural network structure implementing the LLM training pipeline module **530** described in FIG. **5**, according to some embodiments. In some embodiments, the LLM training pipeline module **330** and/or one or more of its submodules **631-435** may be implemented at least partially via an artificial neural network structure shown in FIG. **7**. The neural network comprises a computing system that is built on a collection of connected units or nodes, referred to as neurons (e.g., **644**, **645**, **646**). Neurons are often connected by edges, and an adjustable weight (e.g., **651**, **652**) is often associated with the edge. The neurons are often aggregated into layers such that different layers may perform different transformations on the respective input and output transformed input data onto the next layer.

[0051] For example, the neural network architecture may comprise an input layer **641**, one or more hidden layers **642** and an output layer **643**. Each layer may comprise a plurality of neurons, and neurons between layers are interconnected according to a specific topology of the neural network topology. The input layer **641** receives the input data (e.g., **640** in FIG. **6A**), such as an input image and an input text. The number of nodes (neurons) in the input layer **641** may be determined by the dimensionality of the input data (e.g., the length of a vector of a latent feature of the input image). Each node in the input layer represents a feature or attribute of the input.

[0052] The hidden layers **642** are intermediate layers between the input and output layers of a neural network. It is noted that two hidden layers **642** are shown in FIG. **6B** for illustrative purpose only, and any number of hidden layers may be utilized in a neural network structure. Hidden layers **642** may extract and transform the input data through a series of weighted computations and activation functions.

[0053] For example, as discussed in FIG. **6**, the LLM training pipeline module **530** receives an

input **640** of an input image and transforms the input into an output **650** of an image representation. To perform the transformation, each neuron receives input signals, performs a weighted sum of the inputs according to weights assigned to each connection (e.g., **651**, **652**), and then applies an activation function (e.g., **661**, **662**, etc.) associated with the respective neuron to the result. The output of the activation function is passed to the next layer of neurons or serves as the final output of the network. The activation function may be the same or different across different layers. Example activation functions include but not limited to Sigmoid, hyperbolic tangent, Rectified Linear Unit (ReLU), Leaky ReLU, Softmax, and/or the like. In this way, after a number of hidden layers, input data received at the input layer **641** is transformed into rather different values indicative data characteristics corresponding to a task that the neural network structure has been designed to perform.

[0054] The output layer **643** is the final layer of the neural network structure. It produces the network's output or prediction based on the computations performed in the preceding layers (e.g., **641**, **642**). The number of nodes in the output layer depends on the nature of the task being addressed. For example, in a binary classification problem, the output layer may consist of a single node representing the probability of belonging to one class. In a multi-class classification problem, the output layer may have multiple nodes, each representing the probability of belonging to a specific class.

[0055] Therefore, the LLM training pipeline module **530** and/or one or more of its submodules **631-335** may comprise the transformative neural network structure of layers of neurons, and weights and activation functions describing the non-linear transformation at each neuron. Such a neural network structure is often implemented on one or more hardware processors **610**, such as a graphics processing unit (GPU). An example neural network may be a Transformer model, and/or the like.

[0056] In one embodiment, the LLM training pipeline module **330** and its submodules **631-335** may be implemented by hardware, software and/or a combination thereof. For example, the LLM training pipeline module **330** and its submodules **331-333** may comprise a specific neural network structure implemented and run on various hardware platforms **660**, such as but not limited to CPUs (central processing units), GPUs (graphics processing units), FPGAs (field-programmable gate arrays), Application-Specific Integrated Circuits (ASICs), dedicated AI accelerators like TPUs (tensor processing units), and specialized hardware accelerators designed specifically for the neural network computations described herein, and/or the like. Example specific hardware for neural network structures may include, but not limited to Google Edge TPU, Deep Learning Accelerator (DLA), NVIDIA AI-focused GPUs, and/or the like. The hardware **660** used to implement the neural network structure is specifically configured based on factors such as the complexity of the neural network, the scale of the tasks (e.g., training time, input data scale, size of training dataset, etc.), and the desired performance.

[0057] In one embodiment, the neural network based LLM training pipeline module **330** and one or more of its submodules **631-435** may be trained by iteratively updating the underlying parameters (e.g., weights **651**, **652**, etc., bias parameters and/or coefficients in the activation functions **661**, **662** associated with neurons) of the neural network based on the loss. For example, during forward propagation, the training data such as a training image or a training text are fed into the neural network. The data flows through the network's layers **641**, **642**, with each layer performing computations based on its weights, biases, and activation functions until the output layer **643** produces the network's output **650**. In some embodiments, output layer **643** produces an intermediate output on which the network's output **650** is based.

[0058] The output generated by the output layer **643** is compared to the expected output (e.g., a "ground-truth") from the training data, to compute a loss function that measures the discrepancy between the predicted output and the expected output. Given the loss, the negative gradient of the loss function is computed with respect to each weight of each layer individually. Such negative

gradient is computed one layer at a time, iteratively backward from the last layer **643** to the input layer **641** of the neural network. These gradients quantify the sensitivity of the network's output to changes in the parameters. The chain rule of calculus is applied to efficiently calculate these gradients by propagating the gradients backward from the output layer **643** to the input layer **641**. [0059] Parameters of the neural network are updated backwardly from the last layer to the input layer (backpropagating) based on the computed negative gradient using an optimization algorithm to minimize the loss. The backpropagation from the last layer **643** to the input layer **641** may be conducted for a number of training samples in a number of iterative training epochs. In this way, parameters of the neural network may be gradually updated in a direction to result in a lesser or minimized loss, indicating the neural network has been trained to generate a predicted output value closer to the target output value with improved prediction accuracy. Training may continue until a stopping criterion is met, such as reaching a maximum number of epochs or achieving satisfactory performance on the validation data. At this point, the trained network can be used to make predictions on new, unseen data, such as image animation.

[0060] Neural network parameters may be trained over multiple stages. For example, initial training (e.g., pre-training) may be performed on one set of training data, and then an additional training stage (e.g., fine-tuning) may be performed using a different set of training data. In some embodiments, all or a portion of parameters of one or more neural-network model being used together may be frozen, such that the “frozen” parameters are not updated during that training phase. This may allow, for example, a smaller subset of the parameters to be trained without the computing cost of updating all of the parameters.

[0061] Therefore, the training process transforms the neural network into an “updated” trained neural network with updated parameters such as weights, activation functions, and biases. The trained neural network thus improves neural network technology in applications of intelligent agents.

[0062] FIG. 7 is a simplified block diagram of a networked system **700** suitable for implementing the LLM training pipeline framework described in FIGS. 1-4 and other embodiments described herein. In one embodiment, system **700** includes the user device **710** which may be operated by user **740**, data vendor servers **745**, **770** and **780**, server **730**, and other forms of devices, servers, and/or software components that operate to perform various methodologies in accordance with the described embodiments. Exemplary devices and servers may include device, stand-alone, and enterprise-class servers which may be similar to the computing device **300** described in FIG. 3, operating an OS such as a MICROSOFT® OS, a UNIX® OS, a LINUX® OS, or other suitable device and/or server-based OS. It can be appreciated that the devices and/or servers illustrated in FIG. 7 may be deployed in other ways and that the operations performed, and/or the services provided by such devices and/or servers may be combined or separated for a given embodiment and may be performed by a greater number or fewer number of devices and/or servers. One or more devices and/or servers may be operated and/or maintained by the same or different entities. [0063] The user device **710**, data vendor servers **745**, **770** and **780**, and the server **730** may communicate with each other over a network **760**. User device **710** may be utilized by a user **740** (e.g., a driver, a system admin, etc.) to access the various features available for user device **710**, which may include processes and/or applications associated with the server **730** to receive generated LLM outputs.

[0064] User device **710**, data vendor server **745**, and the server **730** may each include one or more processors, memories, and other appropriate components for executing instructions such as program code and/or data stored on one or more computer readable mediums to implement the various applications, data, and steps described herein. For example, such instructions may be stored in one or more computer readable media such as memories or data storage devices internal and/or external to various components of system **700**, and/or accessible over network **760**.

[0065] User device **710** may be implemented as a communication device that may utilize

appropriate hardware and software configured for wired and/or wireless communication with data vendor server **745** and/or the server **730**. For example, in one embodiment, user device **710** may be implemented as an autonomous driving vehicle, a personal computer (PC), a smart phone, laptop/tablet computer, wristwatch with appropriate computer hardware resources, eyeglasses with appropriate computer hardware (e.g., GOOGLE GLASS®), other type of wearable computing device, implantable communication devices, and/or other types of computing devices capable of transmitting and/or receiving data, such as an IPAD® from APPLE®. Although only one communication device is shown, a plurality of communication devices may function similarly.

[0066] User device **710** of FIG. 7 contains a user interface (UI) application **712**, and/or other applications **716**, which may correspond to executable processes, procedures, and/or applications with associated hardware. For example, the user device **710** may receive a message indicating an LLM output from the server **730** and display the message via the UI application **712**. In other embodiments, user device **710** may include additional or different modules having specialized hardware and/or software as required.

[0067] In various embodiments, user device **710** includes other applications **716** as may be desired in particular embodiments to provide features to user device **710**. For example, other applications **716** may include security applications for implementing client-side security features, programmatic client applications for interfacing with appropriate application programming interfaces (APIs) over network **760**, or other types of applications. Other applications **716** may also include communication applications, such as email, texting, voice, social networking, and IM applications that allow a user to send and receive emails, calls, texts, and other notifications through network **760**. For example, the other application **716** may be an email or instant messaging application that receives a forecast result from the server **730**. Other applications **716** may include device interfaces and other display modules that may receive input and/or output information. For example, other applications **716** may contain software programs for asset management, executable by a processor, including a graphical user interface (GUI) configured to provide an interface to the user **740** to view the visualized output.

[0068] User device **710** may further include database **718** stored in a transitory and/or non-transitory memory of user device **710**, which may store various applications and data and be utilized during execution of various modules of user device **710**. Database **718** may store user profile relating to the user **740**, predictions previously viewed or saved by the user **740**, historical data received from the server **730**, and/or the like. In some embodiments, database **718** may be local to user device **710**. However, in other embodiments, database **718** may be external to user device **710** and accessible by user device **710**, including cloud storage systems and/or databases that are accessible over network **760**.

[0069] User device **710** includes at least one network interface component **717** adapted to communicate with data vendor server **745** and/or the server **730**. In various embodiments, network interface component **717** may include a DSL (e.g., Digital Subscriber Line) modem, a PSTN (Public Switched Telephone Network) modem, an Ethernet device, a broadband device, a satellite device and/or various other types of wired and/or wireless network communication devices including microwave, radio frequency, infrared, Bluetooth, and near field communication devices.

[0070] Data vendor server **745** may correspond to a server that hosts database **719** to provide training datasets including training images/texts to the server **730**. The database **719** may be implemented by one or more relational database, distributed databases, cloud databases, and/or the like.

[0071] The data vendor server **745** includes at least one network interface component **726** adapted to communicate with user device **710** and/or the server **730**. In various embodiments, network interface component **726** may include a DSL (e.g., Digital Subscriber Line) modem, a PSTN (Public Switched Telephone Network) modem, an Ethernet device, a broadband device, a satellite device and/or various other types of wired and/or wireless network communication devices

including microwave, radio frequency, infrared, Bluetooth, and near field communication devices. For example, in one implementation, the data vendor server **745** may send asset information from the database **719**, via the network interface **726**, to the server **730**.

[0072] The server **730** may be housed with the LLM training pipeline module **330** and its submodules described in FIG. **4**. In some implementations, LLM training pipeline module **330** may receive data from database **719** at the data vendor server **745** via the network **760** to generate an LLM output. The generated output may also be sent to the user device **710** for review by the user **740** via the network **760**.

[0073] The database **732** may be stored in a transitory and/or non-transitory memory of the server **730**. In one implementation, the database **732** may store data obtained from the data vendor server **745**. In one implementation, the database **732** may store parameters of the LLM training pipeline module **530**. In one implementation, the database **732** may store previously generated tensor vectors, and the corresponding input feature vectors.

[0074] In some embodiments, database **732** may be local to the server **730**. However, in other embodiments, database **732** may be external to the server **730** and accessible by the server **730**, including cloud storage systems and/or databases that are accessible over network **760**.

[0075] The server **730** includes at least one network interface component **733** adapted to communicate with user device **710** and/or data vendor servers **745**, **770** or **780** over network **760**. In various embodiments, network interface component **733** may comprise a DSL (e.g., Digital Subscriber Line) modem, a PSTN (Public Switched Telephone Network) modem, an Ethernet device, a broadband device, a satellite device and/or various other types of wired and/or wireless network communication devices including microwave, radio frequency (RF), and infrared (IR) communication devices.

[0076] Network **760** may be implemented as a single network or a combination of multiple networks. For example, in various embodiments, network **760** may include the Internet or one or more intranets, landline networks, wireless networks, and/or other appropriate types of networks. Thus, network **760** may correspond to small scale communication networks, such as a private or local area network, or a larger scale network, such as a wide area network or the Internet, accessible by the various components of system **700**.

Example Work Flow

[0077] FIG. **8** is an example logic flow diagram illustrating an example method of training an intelligent agent for carrying out actions in response to user interactions, according to embodiments described herein. One or more of the processes of method **800** may be implemented, at least in part, in the form of executable code stored on non-transitory, tangible, machine-readable media that when run by one or more processors may cause the one or more processors to perform one or more of the processes. In some embodiments, method **600** corresponds to the operation of the LLM training pipeline module **530** (e.g., FIGS. **3** and **5**).

[0078] As illustrated, the method **800** includes a number of enumerated steps, but aspects of the method **800** may include additional steps before, after, and in between the enumerated steps. In some aspects, one or more of the enumerated steps may be omitted or performed in a different order.

[0079] At step **801**, a plurality of multi-turn user-agent interaction records (e.g., **202** in FIG. **2**, **302**, **304** in FIGS. **3A-3B**) corresponding to a plurality of different formats may be received via a communication interface (e.g., **515** in FIG. **5**, **733** in FIG. **7**) at a server (e.g., **730** in FIG. **7**) and from a plurality of data source servers (e.g., **745**, **770**, **780** in FIG. **7**). For example, the received plurality of multi-turn user-agent interaction records comprise a first data format indicating a first user-agent interaction to complete a first type of task in a first environment, e.g., responding to an API call, etc., and a second data format indicating a second user-agent interaction to complete a second type of task in a second environment, e.g., responding to a customer service task, etc.

[0080] At step **803**, the plurality of multi-turn user-agent interaction records may be converted into

a plurality of agent trajectories having a homogeneous data format (e.g., **205** in FIG. **2**). For example, a multi-turn user-agent interaction record may be parsed into a plurality of turns of interactions. For each turn of interactions, an interaction input indicating a context for the respective turn of interactions (e.g., “input” field), an interaction output indicating an agent action executed at the respective turn of interactions (e.g., “output” field) and an observation indicating a response from an environment at which the agent action is executed (e.g., “observation” field) may be extracted from the multi-turn user-agent interaction record. Turns of interaction inputs, interaction outputs and observations may be aggregated in the homogeneous data format.

[0081] In another implementation, a neural network based language model such as an LLM may be used to generate the plurality of agent trajectories using a prompt describing the homogeneous data format.

[0082] At step **805**, a neural network rating model (e.g., agent rater **120** in FIG. **1**) may generate a plurality of rating scores corresponding to the plurality of agent trajectories. For example, the neural network rating model is trained with a dataset of agent trajectories and annotated ratings. For example, the neural network rating model is a LLM external to the server and accessible via an application programming interface (API). A prompt input may be generated for the LLM, comprising at least one agent trajectory to be rated, and an instruction on evaluating the at least one agent trajectory.

[0083] At step **807**, a training dataset may be aggregated based on the plurality of agent trajectories and the plurality of rating scores. The training dataset is aggregated by filtering agent trajectories having rating scores lower than a threshold and shuffling remaining agent trajectories randomly.

[0084] At step **809**, a neural network based language model such as an agent model may be trained to generate a predicted agent action in response to a user query using the training dataset.

[0085] At step **811**, the trained neural network based model may be deployed on a hardware platform as an agent application for handling different user queries in different user environments. For example, the trained neural network based model may generate respective agent actions in response to different user queries from the different user environments.

[0086] FIGS. **9A-9B** are example diagrams illustrating example pseudo code segments for implementing the method shown in FIG. **8**, according to embodiments described herein. For example, as shown in FIG. **9A**, training data may be generated by the “filter” and “shuffling” functions as described in relation to FIG. **2**. As shown in FIG. **9B**, an interleave function may be performed to sample multi-turn data records.

Example Data Performance

[0087] Experimental evaluations may be conducted across four benchmarks: Webshop (Yao et al., Webshop: Towards scalable real-world web interaction with grounded language agents, Advances in Neural Information Processing Systems, 35:20744-20757, 2022), HotpotQA (Yang et al., HotpotQA: A dataset for diverse, explainable multi-hop question answering, in Proceedings of Empirical Methods in Natural Language Processing (EMNLP), 2018), ToolEval and MINT-Bench (Wang et al., MINT: Evaluating LLMs in multi-turn interaction with tools and language feedback, arXiv preprint arXiv: 2309.10691, 2023). For example, Webshop creates an online shopping environment simulating product purchases, while HotpotQA involves multi-hop question-answering tasks requiring logical reasoning across Wikipedia passages via the Wikipedia API. The BOLAA's framework (Liu et al., 2023b), comprising five single-agent settings and a multi-agent scenario, may be adopted to evaluate model performance. For the Webshop benchmark, BOLAA comprises 900 user queries, of which we serve 200 as a test subset. For HotpotQA, 300 user questions are sampled into three difficulty levels-easy, medium, and hard-with each category containing 100 questions. These questions are exclusively reserved for model testing to ensure a rigorous evaluation process. Evaluation metrics may comprise an average reward for Webshop and F1 score for HotpotQA, to measure model performance. In Webshop, the reward metric assesses model accuracy based on the attributes overlapping between the purchased and the ground-truth

items, while in HotpotQA, it quantifies the accuracy of agent-predicted answers against ground-truth responses.

[0088] In one embodiment, ToolEval is designed for real-time assessment of functional calling capabilities via RapidAPI, and GPT-4-0125-preview has been used as the evaluator. The evaluation employs the default depth-first search-based decision tree methodology, augmented by the Pass Rate metric to assess an LLM's ability to execute instructions, a fundamental criterion for optimal tool usage.

[0089] In one embodiment, MINT-Bench benchmark evaluates LLMs' ability to solve tasks with multi-turn interactions by using tools and leveraging natural language feedback. The benchmark focuses on reasoning, coding, and decision-making through a diverse set of established evaluation datasets, and carefully curate them into a compact subset for efficient evaluation. The benchmark asks LLMs to solve tasks with different interaction limits from 1 to 5 step and quantify LLMs' tool-augmented task-solving capability by absolute performance success rate, which measures the percentage of successful task instances as a function of interaction steps.

[0090] As shown in Table 1, the performance of agent model (referred to as xLAM-v0.1) trained using data loader loading training data **208** in FIG. 2 within the Webshop environment. The agent model consistently outperforms both GPT-3.5-Turbo and GPT-3.5-Turbo-Instruct across all agent configurations. Moreover, it surpasses GPT-4-0613 in five out of six settings, with the latter demonstrating superior planning capabilities but lower performance in reasoning, self-reflection, and multi-agent interactions.

TABLE-US-00001 TABLE 1 Average reward on the WebShop environment. **Bold** and Underline results denote the best result and the second best result for each setting, respectively. LAA
Architecture LLM ZS ZST ReAct PlanAct PlanReAct BOLAA Llama-2-70b-chat 0.0089 0.0102
0.4273 0.2809 0.3966 0.4986 Vicuna-33b 0.1527 0.2122 0.1971 0.3766 0.4032 0.5618 Mixtral-
8x7B-Instruct-v0.1 0.4634 0.4592 0.5638 0.4738 0.3339 0.5342 GPT-3.5-Turbo 0.4851 0.5058
0.5047 0.4930 0.5436 0.6354 GPT-3.5-Turbo-Instruct 0.3785 0.4195 0.4377 0.3604 0.4851 0.5811
GPT-4-0613 0.5002 0.4783 0.4616 **0.7950** 0.4635 0.6129 xLAM-v0.1 **0.5201** **0.5268** **0.6486**
0.6573 **0.6611** **0.6556**

[0091] As shown in Table 2, in the HotpotQA environment, the agent model trained using distributed trainer **210** in FIG. 2 (referred as “xLAM”) exhibits superior performance relative to GPT-3.5-Turbo and Mixtral-8x7B-Instruct-v0.1 across all settings.

TABLE-US-00002 TABLE 2 Average reward on the HotpotQA environment. **Bold** and Underline results denote the best result and the second best result for each setting, respectively. LAA
Architecture LLM ZS ZST ReAct PlanAct PlanReAct Mixtral-8x7B- 0.3912 0.3971 0.3714 0.3195
0.3039 Instruct-v0.1 GPT-3.5-Turbo 0.4196 0.3937 0.3868 0.4182 0.3960 GPT-4-0613 **0.5801**
0.5709 **0.6129** **0.5778** **0.5716** xLAM-v0.1 0.5492 0.4776 0.5020 0.5583 0.5030

[0092] As shown in Table 3, the performance results on ToolEval show the agent model (referred to as “xLAM-v0.1”) surpasses both TooLlama V2 and GPT-3.5-Turbo-0125 across all evaluated scenarios, and outperforms GPT-4-0125-preview in two out of the three settings. This performance indicates xLAM-v0.1's superior capabilities in function calling and handling complex tool usage tasks.

TABLE-US-00003 TABLE 3 Pass Rate on ToolEval on three distinct scenarios. Bold and Underline results denote the best result and the second best result for each setting, respectively.
Unseen Insts Unseen Tools Unseen Tools & Same Set & Seen Cat & Unseen Cat TooLlama V2
0.4385 0.4300 0.4350 GPT-3.5-Turbo-0125 0.5000 0.5150 0.4900 GPT-4-0125-preview **0.5462**
0.5450 0.5050 xLAM-v0.1 0.5077 **0.5650** **0.5200**

[0093] Table 4 presents the testing results on the challenging and comprehensive MINT-Bench, with baseline comparisons drawn from the official leaderboard. The agent model xLAM-v0.1 model secures the third rank in this rigorous benchmark, outperforming other agent-based models such as Lemur-70b-Chat-v1 and AgentLM-70b, as well as general large language models (LLMs)

including Claude-2 and GPT-3.5-Turbo-0613. These results highlight the exceptional capability of our model to navigate the complexities of multi-turn interactions and task resolution.

TABLE-US-00004 TABLE 4 Testing results on MINT-Bench with different interaction limits from 1 to 5 step.

1-step	2-step	3-step	4-step	5-step	GPT-4-0613	nan	nan	nan	nan	69.45	Claude-Instant-1
12.12	32.25	39.25	44.37	45.90	xLAM-v0.1	4.10	28.50	36.01	42.66	43.96	Claude-2
36.01	39.76	39.93	Lemur-70b-Chat-v1	3.75	26.96	35.67	37.54	37.03	GPT-3.5-Turbo-0613	2.73	
16.89	24.06	31.74	36.18	AgentLM-70b	6.48	17.75	24.91	28.16	28.67	CodeLlama-34b	0.17
23.04	25.94	28.16	Llama-2-70b-chat	4.27	14.33	15.70	16.55	17.92			

[0094] This description and the accompanying drawings that illustrate inventive aspects, embodiments, implementations, or applications should not be taken as limiting. Various mechanical, compositional, structural, electrical, and operational changes may be made without departing from the spirit and scope of this description and the claims. In some instances, well-known circuits, structures, or techniques have not been shown or described in detail in order not to obscure the embodiments of this disclosure. Like numbers in two or more figures represent the same or similar elements.

[0095] In this description, specific details are set forth describing some embodiments consistent with the present disclosure. Numerous specific details are set forth in order to provide a thorough understanding of the embodiments. It will be apparent, however, to one skilled in the art that some embodiments may be practiced without some or all of these specific details. The specific embodiments disclosed herein are meant to be illustrative but not limiting. One skilled in the art may realize other elements that, although not specifically described here, are within the scope and the spirit of this disclosure. In addition, to avoid unnecessary repetition, one or more features shown and described in association with one embodiment may be incorporated into other embodiments unless specifically described otherwise or if the one or more features would make an embodiment non-functional.

[0096] Although illustrative embodiments have been shown and described, a wide range of modification, change and substitution is contemplated in the foregoing disclosure and in some instances, some features of the embodiments may be employed without a corresponding use of other features. One of ordinary skill in the art would recognize many variations, alternatives, and modifications. Thus, the scope of the invention should be limited only by the following claims, and it is appropriate that the claims be construed broadly and, in a manner, consistent with the scope of the embodiments disclosed herein.

Claims

1. A method of training an intelligent agent for carrying out actions in response to user interactions, the method comprising: receiving, via a communication interface at a server and from a plurality of data source servers, a plurality of multi-turn user-agent interaction records corresponding to a plurality of different formats; converting the plurality of multi-turn user-agent interaction records into a plurality of agent trajectories having a homogeneous data format; generating, by a neural network rating model, a plurality of rating scores corresponding to the plurality of agent trajectories; aggregating a training dataset based on the plurality of agent trajectories and the plurality of rating scores; training a neural network based language models to generate a predicted agent action in response to a user query using the training dataset; and deploying the trained neural network based model on a hardware platform as an agent application for handling different user queries in different user environments.

2. The method of claim 1, wherein the received plurality of multi-turn user-agent interaction records comprise a first data format indicating a first user-agent interaction to complete a first type of task in a first environment and a second data format indicating a second user-agent interaction to complete a second type of task in a second environment.

3. The method of claim 1, wherein the converting the plurality of multi-turn user-agent interaction records into the plurality of agent trajectories having the homogeneous data format comprises: parsing a multi-turn user-agent interaction record into a plurality of turns of interactions; extracting, for each turn of interactions, an interaction input indicating a context for the respective turn of interactions, an interaction output indicating an agent action executed at the respective turn of interactions and an observation indicating a response from an environment at which the agent action is executed, from the multi-turn user-agent interaction record; aggregating turns of interaction inputs, interaction outputs and observations in the homogeneous data format.
4. The method of claim 1, wherein the converting the plurality of multi-turn user-agent interaction records into the plurality of agent trajectories having the homogeneous data format comprises: generating, by a neural network based language model, the plurality of agent trajectories using a prompt describing the homogeneous data format.
5. The method of claim 1, wherein the neural network rating model is trained with a dataset of agent trajectories and annotated ratings.
6. The method of claim 1, wherein the neural network rating model is a large language model external to the server and accessible via an application programming interface (API), and wherein the generating, by the neural network rating model, the plurality of rating scores comprises: generating a prompt input for the large language model, comprising at least one agent trajectory to be rated, and an instruction on evaluating the at least one agent trajectory.
7. The method of claim 1, wherein the training dataset is aggregated by filtering agent trajectories having rating scores lower than a threshold and shuffling remaining agent trajectories randomly.
8. The method of claim 1, wherein the training dataset comprises agent trajectories originated in different environments from the different data sources.
9. The method of claim 1, further comprising: generating, by the trained neural network based model, respective agent actions in response to different user queries from the different user environments.
10. A system of training an intelligent agent for carrying out actions in response to user interactions, the system comprising: a communication interface receiving, at a server and from a plurality of data source servers, a plurality of multi-turn user-agent interaction records corresponding to a plurality of different formats; a memory storing a plurality of processor-executable instructions; and one or more processors executing the plurality of processor-executable instructions to perform operations comprising: converting the plurality of multi-turn user-agent interaction records into a plurality of agent trajectories having a homogeneous data format; generating, by a neural network rating model, a plurality of rating scores corresponding to the plurality of agent trajectories; aggregating a training dataset based on the plurality of agent trajectories and the plurality of rating scores; training a neural network based language models to generate a predicted agent action in response to a user query using the raining dataset; and deploying the trained neural network based model on a hardware platform as an agent application for handling different user queries in different user environments.
11. The system of claim 10, wherein the received plurality of multi-turn user-agent interaction records comprise a first data format indicating a first user-agent interaction to complete a first type of task in a first environment and a second data format indicating a second user-agent interaction to complete a second type of task in a second environment.
12. The system of claim 10, wherein the operation of converting the plurality of multi-turn user-agent interaction records into the plurality of agent trajectories having the homogeneous data format comprises: parsing a multi-turn user-agent interaction record into a plurality of turns of interactions; extracting, for each turn of interactions, an interaction input indicating a context for the respective turn of interactions, an interaction output indicating an agent action executed at the respective turn of interactions and an observation indicating a response from an environment at which the agent action is executed, from the multi-turn user-agent interaction record; aggregating

turns of interaction inputs, interaction outputs and observations in the homogeneous data format.

13. The system of claim 10, wherein the operation of converting the plurality of multi-turn user-agent interaction records into the plurality of agent trajectories having the homogeneous data format comprises: generating, by a neural network based language model, the plurality of agent trajectories using a prompt describing the homogeneous data format.

14. The system of claim 10, wherein the neural network rating model is trained with a dataset of agent trajectories and annotated ratings.

15. The system of claim 10, wherein the neural network rating model is a large language model external to the server and accessible via an application programming interface (API), and wherein the operation of generating, by the neural network rating model, the plurality of rating scores comprises: generating a prompt input for the large language model, comprising at least one agent trajectory to be rated, and an instruction on evaluating the at least one agent trajectory.

16. The system of claim 10, wherein the training dataset is aggregated by filtering agent trajectories having rating scores lower than a threshold and shuffling remaining agent trajectories randomly.

17. The system of claim 10, wherein the training dataset comprises agent trajectories originated in different environments from the different data sources.

18. The system of claim 10, wherein the operations further comprise: generating, by the trained neural network based model, respective agent actions in response to different user queries from the different user environments.

19. A non-transitory processor-readable storage medium storing a plurality of processor-executable instructions for training an intelligent agent for carrying out actions in response to user interactions, the instructions being executed by one or more processors to perform operations comprising: receiving, via a communication interface at a server and from a plurality of data source servers, a plurality of multi-turn user-agent interaction records corresponding to a plurality of different formats; converting the plurality of multi-turn user-agent interaction records into a plurality of agent trajectories having a homogeneous data format; generating, by a neural network rating model, a plurality of rating scores corresponding to the plurality of agent trajectories; aggregating a training dataset based on the plurality of agent trajectories and the plurality of rating scores; training a neural network based language models to generate a predicted agent action in response to a user query using the training dataset; and deploying the trained neural network based model on a hardware platform as an agent application for handling different user queries in different user environments.

20. The non-transitory processor-readable storage medium of claim 19, wherein the operation of converting the plurality of multi-turn user-agent interaction records into the plurality of agent trajectories having the homogeneous data format comprises: parsing a multi-turn user-agent interaction record into a plurality of turns of interactions; extracting, for each turn of interactions, an interaction input indicating a context for the respective turn of interactions, an interaction output indicating an agent action executed at the respective turn of interactions and an observation indicating a response from an environment at which the agent action is executed, from the multi-turn user-agent interaction record; aggregating turns of interaction inputs, interaction outputs and observations in the homogeneous data format.
