

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250259370

Kind Code

A1

Publication Date

August 14, 2025

Inventor(s)

Ummenhofer; Benjamin et al.

LEARNING NEURAL REFLECTANCE SHADERS FROM IMAGES

Abstract

Described herein are techniques for learning neural reflectance shaders from images. A set of one or more machine learning models can be trained to optimize an illumination latent code and a set of reflectance latent codes for an object within a set of input images, such as lighting-based material effects for the object. A shader can then be generated based on a machine learning model of the one or more machine learning models. The shader is configured to sample the illumination latent code and the set of reflectance latent codes for the object. A 3D representation of the object can be rendered using the generated shader.

Inventors: Ummenhofer; Benjamin (Unterhaching, DE), Wang; Shenlong (Santa Clara, CA), Agrawal; Sanskar (Santa Clara, CA), Lao; Yixing (Santa Clara, CA), Zhang; Kai (Santa Clara, CA), Richter; Stephan (Neubiberg, DE), Koltun; Vladlen (Santa Clara, CA)

Applicant: Intel Corporation (Santa Clara, CA)

Family ID: 1000008560555

Assignee: Intel Corporation (Santa Clara, CA)

Appl. No.: 19/059523

Filed: February 21, 2025

Related U.S. Application Data

parent US continuation 18426740 20240130 parent-grant-document US 12260483 child US 19059523

parent US continuation 17849055 20220624 parent-grant-document US 11972519 child US 18426740

us-provisional-application US 63334588 20220425

Publication Classification

Int. Cl.: **G06T15/00** (20110101); **G06N3/045** (20230101); **G06T15/04** (20110101); **G06T15/50** (20110101); **G06T17/20** (20060101)

U.S. Cl.:

CPC **G06T15/005** (20130101); **G06N3/045** (20230101); **G06T15/04** (20130101); **G06T15/506** (20130101); **G06T17/20** (20130101); G06T2200/08 (20130101); G06T2210/52 (20130101)

Background/Summary

CROSS-REFERENCE [0001] The present patent application is a continuation of U.S. application Ser. No. 18/426,740, filed Jan. 30, 2024, which is a continuation of U.S. application Ser. No. 17/849,055, filed Jun. 24, 2022, issued as U.S. Pat. No. 11,972,519, which claims priority from U.S. Provisional Application No. 63/334,588 filed Apr. 25, 2022, which is hereby incorporated herein by reference.

FIELD

[0002] This disclosure relates generally to neural network data processing and more particularly to neural network data processing via a general-purpose graphics processing unit.

BACKGROUND OF THE DISCLOSURE

[0003] Realistic rendering of reconstructed 3D objects requires accurate material estimation. Accurate material estimation requires estimating parameters for existing reflectance models, which is a non-trivial task. Neural Radiance Fields (NeRF) have been proposed for photorealistic view generation. However, methods like NeRF that directly reconstruct the appearance do not allow for relighting the reconstructed scene or object. Related work that disentangles lighting and reflectance in the NeRF framework allows relighting but has a runtime overhead due to the volume rendering. Furthermore, neural radiance fields cannot be imported easily into common real-time rendering pipelines. Works that rely on controlled lighting or HDR images introduce overhead during capturing of the scene and may not work in some scenarios, e.g., outdoor scenes.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0004] The present invention is illustrated by way of example and not limitation in the figures of the accompanying drawings in which like references indicate similar elements, and in which:

[0005] FIG. 1 is a block diagram of a processing system, according to an embodiment;

[0006] FIG. 2A is a block diagram of an embodiment of a processor having one or more processor cores, an integrated memory controller, and an integrated graphics processor;

[0007] FIG. 2B is a block diagram of hardware logic of a graphics processor core block, according to some embodiments described herein;

[0008] FIG. 2C illustrates a graphics processing unit (GPU) that includes dedicated sets of graphics processing resources arranged into multi-core groups;

[0009] FIG. 2D is a block diagram of general-purpose graphics processing unit (GPGPU) that can be configured as a graphics processor and/or compute accelerator, according to embodiments described herein;

[0010] FIG. 3A is a block diagram of a graphics processor, which may be a discrete graphics

processing unit, or may be a graphics processor integrated with a plurality of processing cores, or other semiconductor devices such as, but not limited to, memory devices or network interfaces;

[0011] FIG. 3B illustrates a graphics processor having a tiled architecture, according to embodiments described herein;

[0012] FIG. 3C illustrates a compute accelerator, according to embodiments described herein;

[0013] FIG. 4 is a block diagram of a graphics processing engine of a graphics processor in accordance with some embodiments;

[0014] FIG. 5A illustrates graphics core cluster, according to an embodiment;

[0015] FIG. 5B illustrates a vector engine of a graphics core, according to an embodiment;

[0016] FIG. 5C illustrates a matrix engine of a graphics core, according to an embodiment;

[0017] FIG. 6 illustrates a tile of a multi-tile processor, according to an embodiment;

[0018] FIG. 7 is a block diagram illustrating graphics processor instruction formats according to some embodiments;

[0019] FIG. 8 is a block diagram of another embodiment of a graphics processor;

[0020] FIG. 9A is a block diagram illustrating a graphics processor command format that may be used to program graphics processing pipelines according to some embodiments;

[0021] FIG. 9B is a block diagram illustrating a graphics processor command sequence according to an embodiment;

[0022] FIG. 10 illustrates an exemplary graphics software architecture for a data processing system according to some embodiments;

[0023] FIG. 11A is a block diagram illustrating an IP core development system that may be used to manufacture an integrated circuit to perform operations according to an embodiment;

[0024] FIG. 11B illustrates a cross-section side view of an integrated circuit package assembly, according to some embodiments described herein;

[0025] FIG. 11C illustrates a package assembly that includes multiple units of hardware logic chiplets connected to a substrate;

[0026] FIG. 11D illustrates a package assembly including interchangeable chiplets, according to an embodiment;

[0027] FIG. 12 is a block diagram illustrating an exemplary system on a chip integrated circuit that may be fabricated using one or more IP cores, according to an embodiment;

[0028] FIG. 13A illustrates an exemplary graphics processor of a system on a chip integrated circuit that may be fabricated using one or more IP cores, according to an embodiment;

[0029] FIG. 13B illustrates an additional exemplary graphics processor of a system on a chip integrated circuit that may be fabricated using one or more IP cores, according to an embodiment;

[0030] FIG. 14 illustrates a machine learning software stack;

[0031] FIG. 15 illustrates a general-purpose graphics processing unit;

[0032] FIG. 16 illustrates a multi-GPU computing system;

[0033] FIG. 17A-17B illustrate layers of exemplary deep neural networks;

[0034] FIG. 18 illustrates an exemplary recurrent neural network;

[0035] FIG. 19 illustrates training and deployment of a deep neural network;

[0036] FIG. 20A is a block diagram illustrating distributed learning;

[0037] FIG. 20B is a block diagram illustrating a programmable network interface and data processing unit;

[0038] FIG. 21 illustrates an exemplary inferencing system on a chip (SOC) suitable for performing inferencing using a trained model;

[0039] FIG. 22 illustrates input and output of the system to reconstruct the shape and reflectance of objects from images;

[0040] FIG. 23 illustrates an overview of the processing pipeline for the system;

[0041] FIG. 24 illustrates an overview of the Reflectance Net and the Illumination Net, according to an embodiment;

[0042] FIG. 25 illustrates an overview of the geometric information that is input to the network;
[0043] FIG. 26 shows example output for a given latent code of the network;
[0044] FIG. 27 illustrates translation of the Reflectance Net to an HLSL shader;
[0045] FIG. 28 illustrates exemplary reference images and reconstructed objects;
[0046] FIG. 29 shows the architecture of the Reflectance Net, according to an embodiment;
[0047] FIG. 30 shows the architecture of the Illumination Net, according to an embodiment;
[0048] FIG. 31 shows a method to learn neural reflectance shaders from images; and
[0049] FIG. 32 is a block diagram of a computing device including a graphics processor, according to an embodiment.

DETAILED DESCRIPTION

[0050] Described herein is a technique to estimate reflectance properties from natural images while supporting real-time rendering. The estimate can be performed using a pair of multilayer perceptron (MLP) networks that have been trained on a large dataset of images is used to represent illumination and surface reflectance. MLP networks are deep fully-connected neural networks without convolution layers. The MLP networks described herein are sufficiently lightweight to be implemented as a shading program that is suitable for real-time rendering applications. The techniques described herein enable tools to create photorealistic 3D assets for use in constructing virtual models of real-world objects.

[0051] In the following description, numerous specific details are set forth to provide a more thorough understanding. However, it will be apparent to one of skill in the art that the embodiments described herein may be practiced without one or more of these specific details. In other instances, well-known features have not been described to avoid obscuring the details of the present embodiments.

System Overview

[0052] FIG. 1 is a block diagram of a processing system **100**, according to an embodiment. Processing system **100** may be used in a single processor desktop system, a multiprocessor workstation system, or a server system having a large number of processors **102** or processor cores **107**. In one embodiment, the processing system **100** is a processing platform incorporated within a system-on-a-chip (SoC) integrated circuit for use in mobile, handheld, or embedded devices such as within Internet-of-things (IoT) devices with wired or wireless connectivity to a local or wide area network.

[0053] In one embodiment, processing system **100** can include, couple with, or be integrated within: a server-based gaming platform; a game console, including a game and media console; a mobile gaming console, a handheld game console, or an online game console. In some embodiments the processing system **100** is part of a mobile phone, smart phone, tablet computing device or mobile Internet-connected device such as a laptop with low internal storage capacity. Processing system **100** can also include, couple with, or be integrated within: a wearable device, such as a smart watch wearable device; smart eyewear or clothing enhanced with augmented reality (AR) or virtual reality (VR) features to provide visual, audio or tactile outputs to supplement real world visual, audio or tactile experiences or otherwise provide text, audio, graphics, video, holographic images or video, or tactile feedback; other augmented reality (AR) device; or other virtual reality (VR) device. In some embodiments, the processing system **100** includes or is part of a television or set top box device. In one embodiment, processing system **100** can include, couple with, or be integrated within a self-driving vehicle such as a bus, tractor trailer, car, motor or electric power cycle, plane, or glider (or any combination thereof). The self-driving vehicle may use processing system **100** to process the environment sensed around the vehicle.

[0054] In some embodiments, the one or more processors **102** each include one or more processor cores **107** to process instructions which, when executed, perform operations for system or user software. In some embodiments, at least one of the one or more processor cores **107** is configured to process a specific instruction set **109**. In some embodiments, instruction set **109** may facilitate

Complex Instruction Set Computing (CISC), Reduced Instruction Set Computing (RISC), or computing via a Very Long Instruction Word (VLIW). One or more processor cores **107** may process a different instruction set **109**, which may include instructions to facilitate the emulation of other instruction sets. Processor core **107** may also include other processing devices, such as a Digital Signal Processor (DSP).

[0055] In some embodiments, the processor **102** includes cache memory **104**. Depending on the architecture, the processor **102** can have a single internal cache or multiple levels of internal cache. In some embodiments, the cache memory is shared among various components of the processor **102**. In some embodiments, the processor **102** also uses an external cache (e.g., a Level-3 (L3) cache or Last Level Cache (LLC)) (not shown), which may be shared among processor cores **107** using known cache coherency techniques. A register file **106** can be additionally included in processor **102** and may include different types of registers for storing different types of data (e.g., integer registers, floating point registers, status registers, and an instruction pointer register). Some registers may be general-purpose registers, while other registers may be specific to the design of the processor **102**.

[0056] In some embodiments, one or more processor(s) **102** are coupled with one or more interface bus(es) **110** to transmit communication signals such as address, data, or control signals between processor **102** and other components in the processing system **100**. The interface bus **110**, in one embodiment, can be a processor bus, such as a version of the Direct Media Interface (DMI) bus. However, processor busses are not limited to the DMI bus, and may include one or more Peripheral Component Interconnect buses (e.g., PCI, PCI express), memory busses, or other types of interface busses. In one embodiment the processor(s) **102** include a memory controller **116** and a platform controller hub **130**. The memory controller **116** facilitates communication between a memory device and other components of the processing system **100**, while the platform controller hub (PCH) **130** provides connections to I/O devices via a local I/O bus.

[0057] The memory device **120** can be a dynamic random-access memory (DRAM) device, a static random-access memory (SRAM) device, flash memory device, phase-change memory device, or some other memory device having suitable performance to serve as process memory. In one embodiment the memory device **120** can operate as system memory for the processing system **100**, to store data **122** and instructions **121** for use when the one or more processors **102** executes an application or process. The memory controller **116** also couples with an optional external graphics processor **118**, which may communicate with the one or more graphics processors **108** in processors **102** to perform graphics and media operations. In some embodiments, graphics, media, and or compute operations may be assisted by an accelerator **112** which is a coprocessor that can be configured to perform a specialized set of graphics, media, or compute operations. For example, in one embodiment the accelerator **112** is a matrix multiplication accelerator used to optimize machine learning or compute operations. In one embodiment the accelerator **112** is a ray-tracing accelerator that can be used to perform ray-tracing operations in concert with the graphics processor **108**. In one embodiment, an external accelerator **119** may be used in place of or in concert with the accelerator **112**.

[0058] In some embodiments a display device **111** can connect to the processor(s) **102**. The display device **111** can be one or more of an internal display device, as in a mobile electronic device or a laptop device or an external display device attached via a display interface (e.g., DisplayPort, etc.). In one embodiment the display device **111** can be a head mounted display (HMD) such as a stereoscopic display device for use in virtual reality (VR) applications or augmented reality (AR) applications.

[0059] In some embodiments the platform controller hub **130** enables peripherals to connect to memory device **120** and processor **102** via a high-speed I/O bus. The I/O peripherals include, but are not limited to, an audio controller **146**, a network controller **134**, a firmware interface **128**, a wireless transceiver **126**, touch sensors **125**, a data storage device **124** (e.g., non-volatile memory,

volatile memory, hard disk drive, flash memory, NAND, 3D NAND, 3D XPoint, etc.). The data storage device **124** can connect via a storage interface (e.g., SATA) or via a peripheral bus, such as a Peripheral Component Interconnect bus (e.g., PCI, PCI express). The touch sensors **125** can include touch screen sensors, pressure sensors, or fingerprint sensors. The wireless transceiver **126** can be a Wi-Fi transceiver, a Bluetooth transceiver, or a mobile network transceiver such as a 3G, 4G, 5G, or Long-Term Evolution (LTE) transceiver. The firmware interface **128** enables communication with system firmware, and can be, for example, a unified extensible firmware interface (UEFI). The network controller **134** can enable a network connection to a wired network. In some embodiments, a high-performance network controller (not shown) couples with the interface bus **110**. The audio controller **146**, in one embodiment, is a multi-channel high-definition audio controller. In one embodiment the processing system **100** includes an optional legacy I/O controller **140** for coupling legacy (e.g., Personal System 2 (PS/2)) devices to the system. The platform controller hub **130** can also connect to one or more Universal Serial Bus (USB) controllers **142** connect input devices, such as keyboard and mouse **143** combinations, a camera **144**, or other USB input devices.

[0060] It will be appreciated that the processing system **100** shown is exemplary and not limiting, as other types of data processing systems that are differently configured may also be used. For example, an instance of the memory controller **116** and platform controller hub **130** may be integrated into a discreet external graphics processor, such as the external graphics processor **118**. In one embodiment the platform controller hub **130** and/or memory controller **116** may be external to the one or more processor(s) **102** and reside in a system chipset that is in communication with the processor(s) **102**.

[0061] For example, circuit boards (“sleds”) can be used on which components such as CPUs, memory, and other components are placed are designed for increased thermal performance. In some examples, processing components such as the processors are located on a top side of a sled while near memory, such as DIMMs, are located on a bottom side of the sled. As a result of the enhanced airflow provided by this design, the components may operate at higher frequencies and power levels than in typical systems, thereby increasing performance. Furthermore, the sleds are configured to blindly mate with power and data communication cables in a rack, thereby enhancing their ability to be quickly removed, upgraded, reinstalled, and/or replaced. Similarly, individual components located on the sleds, such as processors, accelerators, memory, and data storage drives, are configured to be easily upgraded due to their increased spacing from each other. In the illustrative embodiment, the components additionally include hardware attestation features to prove their authenticity.

[0062] A data center can utilize a single network architecture (“fabric”) that supports multiple other network architectures including Ethernet and Omni-Path. The sleds can be coupled to switches via optical fibers, which provide higher bandwidth and lower latency than typical twisted pair cabling (e.g., Category 5, Category 5e, Category 6, etc.). Due to the high bandwidth, low latency interconnections and network architecture, the data center may, in use, pool resources, such as memory, accelerators (e.g., GPUs, graphics accelerators, FPGAs, ASICs, neural network and/or artificial intelligence accelerators, etc.), and data storage drives that are physically disaggregated, and provide them to compute resources (e.g., processors) on an as needed basis, enabling the compute resources to access the pooled resources as if they were local.

[0063] A power supply or source can provide voltage and/or current to processing system **100** or any component or system described herein. In one example, the power supply includes an AC to DC (alternating current to direct current) adapter to plug into a wall outlet. Such AC power can be renewable energy (e.g., solar power) power source. In one example, power source includes a DC power source, such as an external AC to DC converter. In one example, power source or power supply includes wireless charging hardware to charge via proximity to a charging field. In one example, power source can include an internal battery, alternating current supply, motion-based

power supply, solar power supply, or fuel cell source.

[0064] FIGS. 2A-2D illustrate computing systems and graphics processors provided by embodiments described herein. The elements of FIGS. 2A-2D having the same reference numbers (or names) as the elements of any other figure herein can operate or function in any manner similar to that described elsewhere herein, but are not limited to such.

[0065] FIG. 2A is a block diagram of an embodiment of a processor **200** having one or more processor cores **202A-202N**, an integrated memory controller **214**, and an integrated graphics processor **208**. Processor **200** can include additional cores up to and including additional core **202N** represented by the dashed lined boxes. Each of processor cores **202A-202N** includes one or more internal cache units **204A-204N**. In some embodiments each processor core also has access to one or more shared cache units **206**. The internal cache units **204A-204N** and shared cache units **206** represent a cache memory hierarchy within the processor **200**. The cache memory hierarchy may include at least one level of instruction and data cache within each processor core and one or more levels of shared mid-level cache, such as a Level 2 (L2), Level 3 (L3), Level 4 (L4), or other levels of cache, where the highest level of cache before external memory is classified as the LLC. In some embodiments, cache coherency logic maintains coherency between the various cache units **206** and **204A-204N**.

[0066] In some embodiments, processor **200** may also include a set of one or more bus controller units **216** and a system agent core **210**. The one or more bus controller units **216** manage a set of peripheral buses, such as one or more PCI or PCI express busses. System agent core **210** provides management functionality for the various processor components. In some embodiments, system agent core **210** includes one or more integrated memory controllers **214** to manage access to various external memory devices (not shown).

[0067] In some embodiments, one or more of the processor cores **202A-202N** include support for simultaneous multi-threading. In such embodiment, the system agent core **210** includes components for coordinating and operating cores **202A-202N** during multi-threaded processing. System agent core **210** may additionally include a power control unit (PCU), which includes logic and components to regulate the power state of processor cores **202A-202N** and graphics processor **208**.

[0068] In some embodiments, processor **200** additionally includes graphics processor **208** to execute graphics processing operations. In some embodiments, the graphics processor **208** couples with the set of shared cache units **206**, and the system agent core **210**, including the one or more integrated memory controllers **214**. In some embodiments, the system agent core **210** also includes a display controller **211** to drive graphics processor output to one or more coupled displays. In some embodiments, display controller **211** may also be a separate module coupled with the graphics processor via at least one interconnect, or may be integrated within the graphics processor **208**.

[0069] In some embodiments, a ring-based interconnect **212** is used to couple the internal components of the processor **200**. However, an alternative interconnect unit may be used, such as a point-to-point interconnect, a switched interconnect, a mesh interconnect, or other techniques, including techniques well known in the art. In some embodiments, graphics processor **208** couples with the ring-based interconnect **212** via an I/O link **213**.

[0070] The exemplary I/O link **213** represents at least one of multiple varieties of I/O interconnects, including an on package I/O interconnect which facilitates communication between various processor components and a high-performance embedded memory module **218**, such as an eDRAM module or a high-bandwidth memory (HBM) module. In some embodiments, each of the processor cores **202A-202N** and graphics processor **208** can use the embedded memory module **218** as a shared Last Level Cache.

[0071] In some embodiments, processor cores **202A-202N** are homogenous cores executing the same instruction set architecture. In another embodiment, processor cores **202A-202N** are heterogeneous in terms of instruction set architecture (ISA), where one or more of processor cores **202A-202N** execute a first instruction set, while at least one of the other cores executes a subset of

the first instruction set or a different instruction set. In one embodiment, processor cores **202A-202N** are heterogeneous in terms of microarchitecture, where one or more cores having a relatively higher power consumption couple with one or more power cores having a lower power consumption. In one embodiment, processor cores **202A-202N** are heterogeneous in terms of computational capability. Additionally, processor **200** can be implemented on one or more chips or as an SoC integrated circuit having the illustrated components, in addition to other components. [0072] FIG. **2B** is a block diagram of hardware logic of a graphics processor core block **219**, according to some embodiments described herein. In some embodiments, elements of FIG. **2B** having the same reference numbers (or names) as the elements of any other figure herein may operate or function in a manner similar to that described elsewhere herein. The graphics processor core block **219** is exemplary of one partition of a graphics processor. The graphics processor core block **219** can be included within the integrated graphics processor **208** of FIG. **2A** or a discrete graphics processor, parallel processor, and/or compute accelerator. A graphics processor as described herein may include multiple graphics core blocks based on target power and performance envelopes. Each graphics processor core block **219** can include a function block **230** coupled with multiple graphics cores **221A-221F** that include modular blocks of fixed function logic and general-purpose programmable logic. The graphics processor core block **219** also includes shared/cache memory **236** that is accessible by all graphics cores **221A-221F**, rasterizer logic **237**, and additional fixed function logic **238**.

[0073] In some embodiments, the function block **230** includes a geometry/fixed function pipeline **231** that can be shared by all graphics cores in the graphics processor core block **219**. In various embodiments, the geometry/fixed function pipeline **231** includes a 3D geometry pipeline a video front-end unit, a thread spawner and global thread dispatcher, and a unified return buffer manager, which manages unified return buffers. In one embodiment the function block **230** also includes a graphics SoC interface **232**, a graphics microcontroller **233**, and a media pipeline **234**. The graphics SoC interface **232** provides an interface between the graphics processor core block **219** and other core blocks within a graphics processor or compute accelerator SoC. The graphics microcontroller **233** is a programmable sub-processor that is configurable to manage various functions of the graphics processor core block **219**, including thread dispatch, scheduling, and pre-emption. The media pipeline **234** includes logic to facilitate the decoding, encoding, pre-processing, and/or post-processing of multimedia data, including image and video data. The media pipeline **234** implement media operations via requests to compute or sampling logic within the graphics cores **221-221F**. One or more pixel backends **235** can also be included within the function block **230**. The pixel backends **235** include a cache memory to store pixel color values and can perform blend operations and lossless color compression of rendered pixel data.

[0074] In one embodiment the graphics SoC interface **232** enables the graphics processor core block **219** to communicate with general-purpose application processor cores (e.g., CPUs) and/or other components within an SoC or a system host CPU that is coupled with the SoC via a peripheral interface. The graphics SoC interface **232** also enables communication with off-chip memory hierarchy elements such as a shared last level cache memory, system RAM, and/or embedded on-chip or on-package DRAM. The SoC interface **232** can also enable communication with fixed function devices within the SoC, such as camera imaging pipelines, and enables the use of and/or implements global memory atomics that may be shared between the graphics processor core block **219** and CPUs within the SoC. The graphics SoC interface **232** can also implement power management controls for the graphics processor core block **219** and enable an interface between a clock domain of the graphics processor core block **219** and other clock domains within the SoC. In one embodiment the graphics SoC interface **232** enables receipt of command buffers from a command streamer and global thread dispatcher that are configured to provide commands and instructions to each of one or more graphics cores within a graphics processor. The commands and instructions can be dispatched to the media pipeline **234** when media operations are to be

performed, the geometry and fixed function pipeline **231** when graphics processing operations are to be performed. When compute operations are to be performed, compute dispatch logic can dispatch the commands to the graphics cores **221A-221F**, bypassing the geometry and media pipelines.

[0075] The graphics microcontroller **233** can be configured to perform various scheduling and management tasks for the graphics processor core block **219**. In one embodiment the graphics microcontroller **233** can perform graphics and/or compute workload scheduling on the various vector engines **222A-222F**, **224A-224F** and matrix engines **223A-223F**, **225A-225F** within the graphics cores **221A-221F**. In this scheduling model, host software executing on a CPU core of an SoC including the graphics processor core block **219** can submit workloads one of multiple graphics processor doorbells, which invokes a scheduling operation on the appropriate graphics engine. Scheduling operations include determining which workload to run next, submitting a workload to a command streamer, pre-empting existing workloads running on an engine, monitoring progress of a workload, and notifying host software when a workload is complete. In one embodiment the graphics microcontroller **233** can also facilitate low-power or idle states for the graphics processor core block **219**, providing the graphics processor core block **219** with the ability to save and restore registers within the graphics processor core block **219** across low-power state transitions independently from the operating system and/or graphics driver software on the system.

[0076] The graphics processor core block **219** may have greater than or fewer than the illustrated graphics cores **221A-221F**, up to N modular graphics cores. For each set of N graphics cores, the graphics processor core block **219** can also include shared/cache memory **236**, which can be configured as shared memory or cache memory, rasterizer logic **237**, and additional fixed function logic **238** to accelerate various graphics and compute processing operations.

[0077] Within each graphics cores **221A-221F** is set of execution resources that may be used to perform graphics, media, and compute operations in response to requests by graphics pipeline, media pipeline, or shader programs. The graphics cores **221A-221F** include multiple vector engines **222A-222F**, **224A-224F**, matrix acceleration units **223A-223F**, **225A-225D**, cache/shared local memory (SLM), a sampler **226A-226F**, and a ray tracing unit **227A-227F**.

[0078] The vector engines **222A-222F**, **224A-224F** are general-purpose graphics processing units capable of performing floating-point and integer/fixed-point logic operations in service of a graphics, media, or compute operation, including graphics, media, or compute/GPGPU programs. The vector engines **222A-222F**, **224A-224F** can operate at variable vector widths using SIMD, SIMT, or SIMT+SIMD execution modes. The matrix acceleration units **223A-223F**, **225A-225D** include matrix-matrix and matrix-vector acceleration logic that improves performance on matrix operations, particularly low and mixed precision (e.g., INT8, FP16, BF16) matrix operations used for machine learning. In one embodiment, each of the matrix acceleration units **223A-223F**, **225A-225D** includes one or more systolic arrays of processing elements that can perform concurrent matrix multiply or dot product operations on matrix elements.

[0079] The sampler **226A-226F** can read media or texture data into memory and can sample data differently based on a configured sampler state and the texture/media format that is being read. Threads executing on the vector engines **222A-222F**, **224A-224F** or matrix acceleration units **223A-223F**, **225A-225D** can make use of the cache/SLM **228A-228F** within each execution core. The cache/SLM **228A-228F** can be configured as cache memory or as a pool of shared memory that is local to each of the respective graphics cores **221A-221F**. The ray tracing units **227A-227F** within the graphics cores **221A-221F** include ray traversal/intersection circuitry for performing ray traversal using bounding volume hierarchies (BVHs) and identifying intersections between rays and primitives enclosed within the BVH volumes. In one embodiment the ray tracing units **227A-227F** include circuitry for performing depth testing and culling (e.g., using a depth buffer or similar arrangement). In one implementation, the ray tracing units **227A-227F** perform traversal and

intersection operations in concert with image denoising, at least a portion of which may be performed using an associated matrix acceleration unit **223A-223F**, **225A-225D**.

[0080] FIG. 2C illustrates a graphics processing unit (GPU) **239** that includes dedicated sets of graphics processing resources arranged into multi-core groups **240A-240N**. The details of multi-core group **240A** are illustrated. Multi-core groups **240B-240N** may be equipped with the same or similar sets of graphics processing resources.

[0081] As illustrated, a multi-core group **240A** may include a set of graphics cores **243**, a set of tensor cores **244**, and a set of ray tracing cores **245**. A scheduler/dispatcher **241** schedules and dispatches the graphics threads for execution on the various cores **243**, **244**, **245**. In one embodiment the tensor cores **244** are sparse tensor cores with hardware to enable multiplication operations having a zero-value input to be bypassed. The graphics cores **243** of the GPU **239** of FIG. 2C differ in hierarchical abstraction level relative to the graphics cores **221A-221F** of FIG. 2B, which are analogous to the multi-core groups **240A-240N** of FIG. 2C. The graphics cores **243**, tensor cores **244**, and ray tracing cores **245** of FIG. 2C are analogous to, respectively, the vector engines **222A-222F**, **224A-224F**, matrix engines **223A-223F**, **225A-225F**, and ray tracing units **227A-227F** of FIG. 2B.

[0082] A set of register files **242** can store operand values used by the cores **243**, **244**, **245** when executing the graphics threads. These may include, for example, integer registers for storing integer values, floating point registers for storing floating point values, vector registers for storing packed data elements (integer and/or floating-point data elements) and tile registers for storing tensor/matrix values. In one embodiment, the tile registers are implemented as combined sets of vector registers.

[0083] One or more combined level 1 (L1) caches and shared memory units **247** store graphics data such as texture data, vertex data, pixel data, ray data, bounding volume data, etc., locally within each multi-core group **240A**. One or more texture units **247** can also be used to perform texturing operations, such as texture mapping and sampling. A Level 2 (L2) cache **253** shared by all or a subset of the multi-core groups **240A-240N** stores graphics data and/or instructions for multiple concurrent graphics threads. As illustrated, the L2 cache **253** may be shared across a plurality of multi-core groups **240A-240N**. One or more memory controllers **248** couple the GPU **239** to a memory **249** which may be a system memory (e.g., DRAM) and/or a dedicated graphics memory (e.g., GDDR6 memory).

[0084] Input/output (I/O) circuitry **250** couples the GPU **239** to one or more I/O devices **252** such as digital signal processors (DSPs), network controllers, or user input devices. An on-chip interconnect may be used to couple the I/O devices **252** to the GPU **239** and memory **249**. One or more I/O memory management units (IOMMUs) **251** of the I/O circuitry **250** couple the I/O devices **252** directly to the memory **249**. In one embodiment, the IOMMU **251** manages multiple sets of page tables to map virtual addresses to physical addresses in memory **249**. In this embodiment, the I/O devices **252**, CPU(s) **246**, and GPU **239** may share the same virtual address space.

[0085] In one implementation, the IOMMU **251** supports virtualization. In this case, it may manage a first set of page tables to map guest/graphics virtual addresses to guest/graphics physical addresses and a second set of page tables to map the guest/graphics physical addresses to system/host physical addresses (e.g., within memory **249**). The base addresses of each of the first and second sets of page tables may be stored in control registers and swapped out on a context switch (e.g., so that the new context is provided with access to the relevant set of page tables). While not illustrated in FIG. 2C, each of the cores **243**, **244**, **245** and/or multi-core groups **240A-240N** may include translation lookaside buffers (TLBs) to cache guest virtual to guest physical translations, guest physical to host physical translations, and guest virtual to host physical translations.

[0086] In one embodiment, the CPUs **246**, GPU **239**, and I/O devices **252** are integrated on a single

semiconductor chip and/or chip package. The memory **249** may be integrated on the same chip or may be coupled to the memory controllers **248** via an off-chip interface. In one implementation, the memory **249** comprises GDDR6 memory which shares the same virtual address space as other physical system-level memories, although the underlying principles of the embodiments described herein are not limited to this specific implementation.

[0087] In one embodiment, the tensor cores **244** include a plurality of functional units specifically designed to perform matrix operations, which are the fundamental compute operation used to perform deep learning operations. For example, simultaneous matrix multiplication operations may be used for neural network training and inferencing. The tensor cores **244** may perform matrix processing using a variety of operand precisions including single precision floating-point (e.g., 32 bits), half-precision floating point (e.g., 16 bits), integer words (16 bits), bytes (8 bits), and half-bytes (4 bits). In one embodiment, a neural network implementation extracts features of each rendered scene, potentially combining details from multiple frames, to construct a high-quality final image.

[0088] In deep learning implementations, parallel matrix multiplication work may be scheduled for execution on the tensor cores **244**. The training of neural networks, in particular, requires a significant number of matrix dot product operations. In order to process an inner-product formulation of an $N \times N \times N$ matrix multiply, the tensor cores **244** may include at least N dot-product processing elements. Before the matrix multiply begins, one entire matrix is loaded into tile registers and at least one column of a second matrix is loaded each cycle for N cycles. Each cycle, there are N dot products that are processed.

[0089] Matrix elements may be stored at different precisions depending on the particular implementation, including 16-bit words, 8-bit bytes (e.g., INT8) and 4-bit half-bytes (e.g., INT4). Different precision modes may be specified for the tensor cores **244** to ensure that the most efficient precision is used for different workloads (e.g., such as inferencing workloads which can tolerate quantization to bytes and half-bytes).

[0090] In one embodiment, the ray tracing cores **245** accelerate ray tracing operations for both real-time ray tracing and non-real-time ray tracing implementations. In particular, the ray tracing cores **245** include ray traversal/intersection circuitry for performing ray traversal using bounding volume hierarchies (BVHs) and identifying intersections between rays and primitives enclosed within the BVH volumes. The ray tracing cores **245** may also include circuitry for performing depth testing and culling (e.g., using a Z buffer or similar arrangement). In one implementation, the ray tracing cores **245** perform traversal and intersection operations in concert with the image denoising techniques described herein, at least a portion of which may be executed on the tensor cores **244**. For example, in one embodiment, the tensor cores **244** implement a deep learning neural network to perform denoising of frames generated by the ray tracing cores **245**. However, the CPU(s) **246**, graphics cores **243**, and/or ray tracing cores **245** may also implement all or a portion of the denoising and/or deep learning algorithms.

[0091] In addition, as described above, a distributed approach to denoising may be employed in which the GPU **239** is in a computing device coupled to other computing devices over a network or high-speed interconnect. In this embodiment, the interconnected computing devices share neural network learning/training data to improve the speed with which the overall system learns to perform denoising for different types of image frames and/or different graphics applications.

[0092] In one embodiment, the ray tracing cores **245** process all BVH traversal and ray-primitive intersections, saving the graphics cores **243** from being overloaded with thousands of instructions per ray. In one embodiment, each ray tracing core **245** includes a first set of specialized circuitry for performing bounding box tests (e.g., for traversal operations) and a second set of specialized circuitry for performing the ray-triangle intersection tests (e.g., intersecting rays which have been traversed). Thus, in one embodiment, the multi-core group **240A** can simply launch a ray probe, and the ray tracing cores **245** independently perform ray traversal and intersection and return hit

data (e.g., a hit, no hit, multiple hits, etc.) to the thread context. The other cores **243**, **244** are freed to perform other graphics or compute work while the ray tracing cores **245** perform the traversal and intersection operations.

[0093] In one embodiment, each ray tracing core **245** includes a traversal unit to perform BVH testing operations and an intersection unit which performs ray-primitive intersection tests. The intersection unit generates a “hit”, “no hit”, or “multiple hit” response, which it provides to the appropriate thread. During the traversal and intersection operations, the execution resources of the other cores (e.g., graphics cores **243** and tensor cores **244**) are freed to perform other forms of graphics work.

[0094] In one particular embodiment described below, a hybrid rasterization/ray tracing approach is used in which work is distributed between the graphics cores **243** and ray tracing cores **245**.

[0095] In one embodiment, the ray tracing cores **245** (and/or other cores **243**, **244**) include hardware support for a ray tracing instruction set such as Microsoft's DirectX Ray Tracing (DXR) which includes a DispatchRays command, as well as ray-generation, closest-hit, any-hit, and miss shaders, which enable the assignment of unique sets of shaders and textures for each object. Another ray tracing platform which may be supported by the ray tracing cores **245**, graphics cores **243** and tensor cores **244** is Vulkan 1.1.85. Note, however, that the underlying principles of the embodiments described herein are not limited to any particular ray tracing ISA.

[0096] In general, the various cores **245**, **244**, **243** may support a ray tracing instruction set that includes instructions/functions for ray generation, closest hit, any hit, ray-primitive intersection, per-primitive and hierarchical bounding box construction, miss, visit, and exceptions. More specifically, one embodiment includes ray tracing instructions to perform the following functions:

[0097] Ray Generation—Ray generation instructions may be executed for each pixel, sample, or other user-defined work assignment. [0098] Closest Hit—A closest hit instruction may be executed to locate the closest intersection point of a ray with primitives within a scene. [0099] Any Hit—An any hit instruction identifies multiple intersections between a ray and primitives within a scene, potentially to identify a new closest intersection point. [0100] Intersection—An intersection instruction performs a ray-primitive intersection test and outputs a result. [0101] Per-primitive Bounding box Construction—This instruction builds a bounding box around a given primitive or group of primitives (e.g., when building a new BVH or other acceleration data structure). [0102] Miss—Indicates that a ray misses all geometry within a scene, or specified region of a scene. [0103] Visit—Indicates the child volumes a ray will traverse. [0104] Exceptions—Includes various types of exception handlers (e.g., invoked for various error conditions).

[0105] In one embodiment the ray tracing cores **245** may be adapted to accelerate general-purpose compute operations that can be accelerated using computational techniques that are analogous to ray intersection tests. A compute framework can be provided that enables shader programs to be compiled into low level instructions and/or primitives that perform general-purpose compute operations via the ray tracing cores. Exemplary computational problems that can benefit from compute operations performed on the ray tracing cores **245** include computations involving beam, wave, ray, or particle propagation within a coordinate space. Interactions associated with that propagation can be computed relative to a geometry or mesh within the coordinate space. For example, computations associated with electromagnetic signal propagation through an environment can be accelerated via the use of instructions or primitives that are executed via the ray tracing cores. Diffraction and reflection of the signals by objects in the environment can be computed as direct ray-tracing analogies.

[0106] Ray tracing cores **245** can also be used to perform computations that are not directly analogous to ray tracing. For example, mesh projection, mesh refinement, and volume sampling computations can be accelerated using the ray tracing cores **245**. Generic coordinate space calculations, such as nearest neighbor calculations can also be performed. For example, the set of points near a given point can be discovered by defining a bounding box in the coordinate space

around the point. BVH and ray probe logic within the ray tracing cores **245** can then be used to determine the set of point intersections within the bounding box. The intersections constitute the origin point and the nearest neighbors to that origin point. Computations that are performed using the ray tracing cores **245** can be performed in parallel with computations performed on the graphics cores **243** and tensor cores **244**. A shader compiler can be configured to compile a compute shader or other general-purpose graphics processing program into low level primitives that can be parallelized across the graphics cores **243**, tensor cores **244**, and ray tracing cores **245**.

[0107] FIG. 2D is a block diagram of general-purpose graphics processing unit (GPGPU) **270** that can be configured as a graphics processor and/or compute accelerator, according to embodiments described herein. The GPGPU **270** can interconnect with host processors (e.g., one or more CPU(s) **246**) and memory **271**, **272** via one or more system and/or memory busses. In one embodiment the memory **271** is system memory that may be shared with the one or more CPU(s) **246**, while memory **272** is device memory that is dedicated to the GPGPU **270**. In one embodiment, components within the GPGPU **270** and memory **272** may be mapped into memory addresses that are accessible to the one or more CPU(s) **246**. Access to memory **271** and **272** may be facilitated via a memory controller **268**. In one embodiment the memory controller **268** includes an internal direct memory access (DMA) controller **269** or can include logic to perform operations that would otherwise be performed by a DMA controller.

[0108] The GPGPU **270** includes multiple cache memories, including an L2 cache **253**, L1 cache **254**, an instruction cache **255**, and shared memory **256**, at least a portion of which may also be partitioned as a cache memory. The GPGPU **270** also includes multiple compute units **260A-260N**, which represent a hierarchical abstraction level analogous to the graphics cores **221A-221F** of FIG. 2B and the multi-core groups **240A-240N** of FIG. 2C. Each compute unit **260A-260N** includes a set of vector registers **261**, scalar registers **262**, vector logic units **263**, and scalar logic units **264**. The compute units **260A-260N** can also include local shared memory **265** and a program counter **266**. The compute units **260A-260N** can couple with a constant cache **267**, which can be used to store constant data, which is data that will not change during the run of kernel or shader program that executes on the GPGPU **270**. In one embodiment the constant cache **267** is a scalar data cache and cached data can be fetched directly into the scalar registers **262**.

[0109] During operation, the one or more CPU(s) **246** can write commands into registers or memory in the GPGPU **270** that has been mapped into an accessible address space. The command processors **257** can read the commands from registers or memory and determine how those commands will be processed within the GPGPU **270**. A thread dispatcher **258** can then be used to dispatch threads to the compute units **260A-260N** to perform those commands. Each compute unit **260A-260N** can execute threads independently of the other compute units. Additionally, each compute unit **260A-260N** can be independently configured for conditional computation and can conditionally output the results of computation to memory. The command processors **257** can interrupt the one or more CPU(s) **246** when the submitted commands are complete.

[0110] FIGS. 3A-3C illustrate block diagrams of additional graphics processor and compute accelerator architectures provided by embodiments described herein. The elements of FIGS. 3A-3C having the same reference numbers (or names) as the elements of any other figure herein can operate or function in any manner similar to that described elsewhere herein, but are not limited to such.

[0111] FIG. 3A is a block diagram of a graphics processor **300**, which may be a discrete graphics processing unit, or may be a graphics processor integrated with a plurality of processing cores, or other semiconductor devices such as, but not limited to, memory devices or network interfaces. In some embodiments, the graphics processor communicates via a memory mapped I/O interface to registers on the graphics processor and with commands placed into the processor memory. In some embodiments, graphics processor **300** includes a memory interface **314** to access memory. Memory interface **314** can be an interface to local memory, one or more internal caches, one or more shared

external caches, and/or to system memory.

[0112] In some embodiments, graphics processor **300** also includes a display controller **302** to drive display output data to a display device **318**. Display controller **302** includes hardware for one or more overlay planes for the display and composition of multiple layers of video or user interface elements. The display device **318** can be an internal or external display device. In one embodiment the display device **318** is a head mounted display device, such as a virtual reality (VR) display device or an augmented reality (AR) display device. In some embodiments, graphics processor **300** includes a video codec engine **306** to encode, decode, or transcode media to, from, or between one or more media encoding formats, including, but not limited to Moving Picture Experts Group (MPEG) formats such as MPEG-2, Advanced Video Coding (AVC) formats such as H.264/MPEG-4 AVC, H.265/HEVC, Alliance for Open Media (AOMedia) VP8, VP9, as well as the Society of Motion Picture & Television Engineers (SMPTE) 421M/VC-1, and Joint Photographic Experts Group (JPEG) formats such as JPEG, and Motion JPEG (MJPEG) formats.

[0113] In some embodiments, graphics processor **300** includes a block image transfer (BLIT) engine **304** to perform two-dimensional (2D) rasterizer operations including, for example, bit-boundary block transfers. However, in one embodiment, 2D graphics operations are performed using one or more components of graphics processing engine (GPE) **310**. In some embodiments, GPE **310** is a compute engine for performing graphics operations, including three-dimensional (3D) graphics operations and media operations.

[0114] In some embodiments, GPE **310** includes a 3D pipeline **312** for performing 3D operations, such as rendering three-dimensional images and scenes using processing functions that act upon 3D primitive shapes (e.g., rectangle, triangle, etc.). The 3D pipeline **312** includes programmable and fixed function elements that perform various tasks within the element and/or spawn execution threads to a 3D/Media subsystem **315**. While 3D pipeline **312** can be used to perform media operations, an embodiment of GPE **310** also includes a media pipeline **316** that is specifically used to perform media operations, such as video post-processing and image enhancement.

[0115] In some embodiments, media pipeline **316** includes fixed function or programmable logic units to perform one or more specialized media operations, such as video decode acceleration, video de-interlacing, and video encode acceleration in place of, or on behalf of video codec engine **306**. In some embodiments, media pipeline **316** additionally includes a thread spawning unit to spawn threads for execution on 3D/Media subsystem **315**. The spawned threads perform computations for the media operations on one or more graphics cores included in 3D/Media subsystem **315**.

[0116] In some embodiments, 3D/Media subsystem **315** includes logic for executing threads spawned by 3D pipeline **312** and media pipeline **316**. In one embodiment, the pipelines send thread execution requests to 3D/Media subsystem **315**, which includes thread dispatch logic for arbitrating and dispatching the various requests to available thread execution resources. The execution resources include an array of graphics cores to process the 3D and media threads. In some embodiments, 3D/Media subsystem **315** includes one or more internal caches for thread instructions and data. In some embodiments, the subsystem also includes shared memory, including registers and addressable memory, to share data between threads and to store output data.

[0117] FIG. 3B illustrates a graphics processor **320** having a tiled architecture, according to embodiments described herein. In one embodiment the graphics processor **320** includes a graphics processing engine cluster **322** having multiple instances of the graphics processing engine **310** of FIG. 3A within a graphics engine tile **310A-310D**. Each graphics engine tile **310A-310D** can be interconnected via a set of tile interconnects **323A-323F**. Each graphics engine tile **310A-310D** can also be connected to a memory module or memory device **326A-326D** via memory interconnects **325A-325D**. The memory devices **326A-326D** can use any graphics memory technology. For example, the memory devices **326A-326D** may be graphics double data rate (GDDR) memory. The memory devices **326A-326D**, in one embodiment, are HBM modules that can be on-die with their

respective graphics engine tile **310A-310D**. In one embodiment the memory devices **326A-326D** are stacked memory devices that can be stacked on top of their respective graphics engine tile **310A-310D**. In one embodiment, each graphics engine tile **310A-310D** and associated memory **326A-326D** reside on separate chiplets, which are bonded to a base die or base substrate, as described on further detail in FIGS. **11B-11D**.

[0118] The graphics processor **320** may be configured with a non-uniform memory access (NUMA) system in which memory devices **326A-326D** are coupled with associated graphics engine tiles **310A-310D**. A given memory device may be accessed by graphics engine tiles other than the tile to which it is directly connected. However, access latency to the memory devices **326A-326D** may be lowest when accessing a local tile. In one embodiment, a cache coherent NUMA (ccNUMA) system is enabled that uses the tile interconnects **323A-323F** to enable communication between cache controllers within the graphics engine tiles **310A-310D** to maintain a consistent memory image when more than one cache stores the same memory location.

[0119] The graphics processing engine cluster **322** can connect with an on-chip or on-package fabric interconnect **324**. In one embodiment the fabric interconnect **324** includes a network processor, network on a chip (NoC), or another switching processor to enable the fabric interconnect **324** to act as a packet switched fabric interconnect that switches data packets between components of the graphics processor **320**. The fabric interconnect **324** can enable communication between graphics engine tiles **310A-310D** and components such as the video codec engine **306** and one or more copy engines **304**. The copy engines **304** can be used to move data out of, into, and between the memory devices **326A-326D** and memory that is external to the graphics processor **320** (e.g., system memory). The fabric interconnect **324** can also couple with one or more of the tile interconnects **323A-323F** to facilitate or enhance the interconnection between the graphics engine tiles **310A-310D**. The fabric interconnect **324** is also configurable to interconnect multiple instances of the graphics processor **320** (e.g., via the host interface **328**), enabling tile-to-tile communication between graphics engine tiles **310A-310D** of multiple GPUs. In one embodiment, the graphics engine tiles **310A-310D** of multiple GPUs can be presented to a host system as a single logical device.

[0120] The graphics processor **320** may optionally include a display controller **302** to enable a connection with the display device **318**. The graphics processor may also be configured as a graphics or compute accelerator. In the accelerator configuration, the display controller **302** and display device **318** may be omitted.

[0121] The graphics processor **320** can connect to a host system via a host interface **328**. The host interface **328** can enable communication between the graphics processor **320**, system memory, and/or other system components. The host interface **328** can be, for example a PCI express bus or another type of host system interface. For example, the host interface **328** may be an NVLink or NVSwitch interface. The host interface **328** and fabric interconnect **324** can cooperate to enable multiple instances of the graphics processor **320** to act as single logical device. Cooperation between the host interface **328** and fabric interconnect **324** can also enable the individual graphics engine tiles **310A-310D** to be presented to the host system as distinct logical graphics devices.

[0122] FIG. **3C** illustrates a compute accelerator **330**, according to embodiments described herein. The compute accelerator **330** can include architectural similarities with the graphics processor **320** of FIG. **3B** and is optimized for compute acceleration. A compute engine cluster **332** can include a set of compute engine tiles **340A-340D** that include execution logic that is optimized for parallel or vector-based general-purpose compute operations. In some embodiments, the compute engine tiles **340A-340D** do not include fixed function graphics processing logic, although in one embodiment one or more of the compute engine tiles **340A-340D** can include logic to perform media acceleration. The compute engine tiles **340A-340D** can connect to memory **326A-326D** via memory interconnects **325A-325D**. The memory **326A-326D** and memory interconnects **325A-325D** may be similar technology as in graphics processor **320** or can be different. The graphics

compute engine tiles **340A-340D** can also be interconnected via a set of tile interconnects **323A-323F** and may be connected with and/or interconnected by a fabric interconnect **324**. Cross-tile communications can be facilitated via the fabric interconnect **324**. The fabric interconnect **324** (e.g., via the host interface **328**) can also facilitate communication between compute engine tiles **340A-340D** of multiple instances of the compute accelerator **330**. In one embodiment the compute accelerator **330** includes a large L3 cache **336** that can be configured as a device-wide cache. The compute accelerator **330** can also connect to a host processor and memory via a host interface **328** in a similar manner as the graphics processor **320** of FIG. 3B.

[0123] The compute accelerator **330** can also include an integrated network interface **342**. In one embodiment the network interface **342** includes a network processor and controller logic that enables the compute engine cluster **332** to communicate over a physical layer interconnect **344** without requiring data to traverse memory of a host system. In one embodiment, one of the compute engine tiles **340A-340D** is replaced by network processor logic and data to be transmitted or received via the physical layer interconnect **344** may be transmitted directly to or from memory **326A-326D**. Multiple instances of the compute accelerator **330** may be joined via the physical layer interconnect **344** into a single logical device. Alternatively, the various compute engine tiles **340A-340D** may be presented as distinct network accessible compute accelerator devices.

Graphics Processing Engine

[0124] FIG. 4 is a block diagram of a graphics processing engine **410** of a graphics processor in accordance with some embodiments. In one embodiment, the graphics processing engine (GPE) **410** is a version of the GPE **310** shown in FIG. 3A and may also represent a graphics engine tile **310A-310D** of FIG. 3B. Elements of FIG. 4 having the same reference numbers (or names) as the elements of any other figure herein can operate or function in any manner similar to that described elsewhere herein, but are not limited to such. For example, the 3D pipeline **312** and media pipeline **316** of FIG. 3A are illustrated. The media pipeline **316** is optional in some embodiments of the GPE **410** and may not be explicitly included within the GPE **410**. For example and in at least one embodiment, a separate media and/or image processor is coupled to the GPE **410**.

[0125] In some embodiments, GPE **410** couples with or includes a command streamer **403**, which provides a command stream to the 3D pipeline **312** and/or media pipelines **316**. Alternatively or additionally, the command streamer **403** may be directly coupled to a unified return buffer **418**. The unified return buffer **418** may be communicatively coupled to a graphics core cluster **414**. In some embodiments, command streamer **403** is coupled with memory, which can be system memory, or one or more of internal cache memory and shared cache memory. In some embodiments, command streamer **403** receives commands from the memory and sends the commands to 3D pipeline **312** and/or media pipeline **316**. The commands are directives fetched from a ring buffer, which stores commands for the 3D pipeline **312** and media pipeline **316**. In one embodiment, the ring buffer can additionally include batch command buffers storing batches of multiple commands. The commands for the 3D pipeline **312** can also include references to data stored in memory, such as but not limited to vertex and geometry data for the 3D pipeline **312** and/or image data and memory objects for the media pipeline **316**. The 3D pipeline **312** and media pipeline **316** process the commands and data by performing operations via logic within the respective pipelines or by dispatching one or more execution threads to a graphics core cluster **414**. In one embodiment the graphics core cluster **414** include one or more blocks of graphics cores (e.g., graphics core block **415A**, graphics core block **415B**), each block including one or more graphics cores. Each graphics core includes a set of graphics execution resources that includes general-purpose and graphics specific execution logic to perform graphics and compute operations, as well as fixed function texture processing and/or machine learning and artificial intelligence acceleration logic, such as matrix or AI acceleration logic.

[0126] In various embodiments the 3D pipeline **312** can include fixed function and programmable logic to process one or more shader programs, such as vertex shaders, geometry shaders, pixel

shaders, fragment shaders, compute shaders, or other shader and/or GPGPU programs, by processing the instructions and dispatching execution threads to the graphics core cluster **414**. The graphics core cluster **414** provides a unified block of execution resources for use in processing these shader programs. Multi-purpose execution logic within the graphics core blocks **415A-415B** of the graphics core cluster **414** includes support for various 3D API shader languages and can execute multiple simultaneous execution threads associated with multiple shaders.

[0127] In some embodiments, the graphics core cluster **414** includes execution logic to perform media functions, such as video and/or image processing. In one embodiment, the graphics cores include general-purpose logic that is programmable to perform parallel general-purpose computational operations, in addition to graphics processing operations. The general-purpose logic can perform processing operations in parallel or in conjunction with general-purpose logic within the processor core(s) **107** of FIG. **1** or core **202A-202N** as in FIG. **2A**.

[0128] Output data generated by threads executing on the graphics core cluster **414** can output data to memory in a unified return buffer (URB) **418**. The URB **418** can store data for multiple threads. In some embodiments the URB **418** may be used to send data between different threads executing on the graphics core cluster **414**. In some embodiments the URB **418** may additionally be used for synchronization between threads on the graphics core array and fixed function logic within the shared function logic **420**.

[0129] In some embodiments, graphics core cluster **414** is scalable, such that the cluster includes a variable number of graphics cores, each having a variable number of graphics cores based on the target power and performance level of GPE **410**. In one embodiment the execution resources are dynamically scalable, such that execution resources may be enabled or disabled as needed.

[0130] The graphics core cluster **414** couples with shared function logic **420** that includes multiple resources that are shared between the graphics cores in the graphics core array. The shared functions within the shared function logic **420** are hardware logic units that provide specialized supplemental functionality to the graphics core cluster **414**. In various embodiments, shared function logic **420** may include, but is not limited to sampler **421**, math **422**, and inter-thread communication (ITC) **423** logic. Additionally, some embodiments implement one or more cache(s) **425** within the shared function logic **420**. The shared function logic **420** can implement the same or similar functionality as the additional fixed function logic **238** of FIG. **2B**.

[0131] A shared function is implemented at least in a case where the demand for a given specialized function is insufficient for inclusion within the graphics core cluster **414**. Instead, a single instantiation of that specialized function is implemented as a stand-alone entity in the shared function logic **420** and shared among the execution resources within the graphics core cluster **414**. The precise set of functions that are shared between the graphics core cluster **414** and included within the graphics core cluster **414** varies across embodiments. In some embodiments, specific shared functions within the shared function logic **420** that are used extensively by the graphics core cluster **414** may be included within shared function logic **416** within the graphics core cluster **414**. In various embodiments, the shared function logic **416** within the graphics core cluster **414** can include some or all logic within the shared function logic **420**. In one embodiment, all logic elements within the shared function logic **420** may be duplicated within the shared function logic **416** of the graphics core cluster **414**. In one embodiment the shared function logic **420** is excluded in favor of the shared function logic **416** within the graphics core cluster **414**.

Graphics Processing Resources

[0132] FIG. **5A-5C** illustrate execution logic including an array of processing elements employed in a graphics processor, according to embodiments described herein. FIG. **5A** illustrates graphics core cluster, according to an embodiment. FIG. **5B** illustrates a vector engine of a graphics core, according to an embodiment. FIG. **5C** illustrates a matrix engine of a graphics core, according to an embodiment. Elements of FIG. **5A-5C** having the same reference numbers as the elements of any other figure herein may operate or function in any manner similar to that described elsewhere

herein, but are not limited as such. For example, the elements of FIG. 5A-5C can be considered in the context of the graphics processor core block **219** of FIG. 2B, and/or the graphics core blocks **415A-415B** of FIG. 4. In one embodiment, the elements of FIG. 5A-5C have similar functionality to equivalent components of the graphics processor **208** of FIG. 2A, the GPU **239** of FIG. 2C or the GPGPU **270** of FIG. 2D.

[0133] As shown in FIG. 5A, in one embodiment the graphics core cluster **414** includes a graphics core block **415**, which may be graphics core block **415A** or graphics core block **415B** of FIG. 4. The graphics core block **415** can include any number of graphics cores (e.g., graphics core **515A**, graphics core **515B**, through graphics core **515N**). Multiple instances of the graphics core block **415** may be included. In one embodiment the elements of the graphics cores **515A-515N** have similar or equivalent functionality as the elements of the graphics cores **221A-221F** of FIG. 2B. In such embodiment, the graphics cores **515A-515N** each include circuitry including but not limited to vector engines **502A-502N**, matrix engines **503A-503N**, memory load/store units **504A-504N**, instruction caches **505A-505N**, data caches/shared local memory **506A-506N**, ray tracing units **508A-508N**, samplers **510A-510N**. The circuitry of the graphics cores **515A-515N** can additionally include fixed function logic **512A-512N**. The number of vector engines **502A-502N** and matrix engines **503A-503N** within the graphics cores **515A-515N** of a design can vary based on the workload, performance, and power targets for the design.

[0134] With reference to graphics core **515A**, the vector engine **502A** and matrix engine **503A** are configurable to perform parallel compute operations on data in a variety of integer and floating-point data formats based on instructions associated with shader programs. Each vector engine **502A** and matrix engine **503A** can act as a programmable general-purpose computational unit that is capable of executing multiple simultaneous hardware threads while processing multiple data elements in parallel for each thread. The vector engine **502A** and matrix engine **503A** support the processing of variable width vectors at various SIMD widths, including but not limited to SIMD8, SIMD16, and SIMD32. Input data elements can be stored as a packed data type in a register and the vector engine **502A** and matrix engine **503A** can process the various elements based on the data size of the elements. For example, when operating on a 256-bit wide vector, the 256 bits of the vector are stored in a register and the vector is processed as four separate 64-bit packed data elements (Quad-Word (QW) size data elements), eight separate 32-bit packed data elements (Double Word (DW) size data elements), sixteen separate 16-bit packed data elements (Word (W) size data elements), or thirty-two separate 8-bit data elements (byte (B) size data elements). However, different vector widths and register sizes are possible. In one embodiment, the vector engine **502A** and matrix engine **503A** are also configurable for SIMT operation on warps or thread groups of various sizes (e.g., 8, 16, or 32 threads).

[0135] Continuing with graphics core **515A**, the memory load/store unit **504A** services memory access requests that are issued by the vector engine **502A**, matrix engine **503A**, and/or other components of the graphics core **515A** that have access to memory. The memory access request can be processed by the memory load/store unit **504A** to load or store the requested data to or from cache or memory into a register file associated with the vector engine **502A** and/or matrix engine **503A**. The memory load/store unit **504A** can also perform prefetching operations. In one embodiment, the memory load/store unit **504A** is configured to provide SIMT scatter/gather prefetching or block prefetching for data stored in memory **610**, from memory that is local to other tiles via the tile interconnect **608**, or from system memory. Prefetching can be performed to a specific L1 cache (e.g., data cache/shared local memory **506A**), the L2 cache **604** or the L3 cache **606**. In one embodiment, a prefetch to the L3 cache **606** automatically results in the data being stored in the L2 cache **604**.

[0136] The instruction cache **505A** stores instructions to be executed by the graphics core **515A**. In one embodiment, the graphics core **515A** also includes instruction fetch and prefetch circuitry that fetches or prefetches instructions into the instruction cache **505A**. The graphics core **515A** also

includes instruction decode logic to decode instructions within the instruction cache **505A**. The data cache/shared local memory **506A** can be configured as a data cache that is managed by a cache controller that implements a cache replacement policy and/or configured as explicitly managed shared memory. The ray tracing unit **508A** includes circuitry to accelerate ray tracing operations. The sampler **510A** provides texture sampling for 3D operations and media sampling for media operations. The fixed function logic **512A** includes fixed function circuitry that is shared between the various instances of the vector engine **502A** and matrix engine **503A**. Graphics cores **515B-515N** can operate in a similar manner as graphics core **515A**.

[0137] Functionality of the instruction caches **505A-505N**, data caches/shared local memory **506A-506N**, ray tracing units **508A-508N**, samplers **510A-2710N**, and fixed function logic **512A-512N** corresponds with equivalent functionality in the graphics processor architectures described herein. For example, the instruction caches **505A-505N** can operate in a similar manner as instruction cache **255** of FIG. 2D. The data caches/shared local memory **506A-506N**, ray tracing units **508A-508N**, and samplers **510A-2710N** can operate in a similar manner as the cache/SLM **228A-228F**, ray tracing units **227A-227F**, and samplers **226A-226F** of FIG. 2B. The fixed function logic **512A-512N** can include elements of the geometry/fixed function pipeline **231** and/or additional fixed function logic **238** of FIG. 2B. In one embodiment, the ray tracing units **508A-508N** include circuitry to perform ray tracing acceleration operations performed by the ray tracing cores **245** of FIG. 2C.

[0138] As shown in FIG. 5B, in one embodiment the vector engine **502** includes an instruction fetch unit **537**, a general register file array (GRF) **524**, an architectural register file array (ARF) **526**, a thread arbiter **522**, a send unit **530**, a branch unit **532**, a set of SIMD floating point units (FPUs) **534**, and in one embodiment a set of integer SIMD ALUs **535**. The GRF **524** and ARF **526** includes the set of general register files and architecture register files associated with each hardware thread that may be active in the vector engine **502**. In one embodiment, per thread architectural state is maintained in the ARF **526**, while data used during thread execution is stored in the GRF **524**. The execution state of each thread, including the instruction pointers for each thread, can be held in thread-specific registers in the ARF **526**.

[0139] In one embodiment the vector engine **502** has an architecture that is a combination of Simultaneous Multi-Threading (SMT) and fine-grained Interleaved Multi-Threading (IMT). The architecture has a modular configuration that can be fine-tuned at design time based on a target number of simultaneous threads and number of registers per graphics core, where graphics core resources are divided across logic used to execute multiple simultaneous threads. The number of logical threads that may be executed by the vector engine **502** is not limited to the number of hardware threads, and multiple logical threads can be assigned to each hardware thread.

[0140] In one embodiment, the vector engine **502** can co-issue multiple instructions, which may each be different instructions. The thread arbiter **522** can dispatch the instructions to one of the send unit **530**, branch unit **532**, or SIMD FPU(s) **534** for execution. Each execution thread can access **128** general-purpose registers within the GRF **524**, where each register can store 32 bytes, accessible as a variable width vector of 32-bit data elements. In one embodiment, each thread has access to 4 Kbytes within the GRF **524**, although embodiments are not so limited, and greater or fewer register resources may be provided in other embodiments. In one embodiment the vector engine **502** is partitioned into seven hardware threads that can independently perform computational operations, although the number of threads per vector engine **502** can also vary according to embodiments. For example, in one embodiment up to 16 hardware threads are supported. In an embodiment in which seven threads may access 4 Kbytes, the GRF **524** can store a total of 28 Kbytes. Where 16 threads may access 4 Kbytes, the GRF **524** can store a total of 64 Kbytes. Flexible addressing modes can permit registers to be addressed together to build effectively wider registers or to represent strided rectangular block data structures.

[0141] In one embodiment, memory operations, sampler operations, and other longer-latency

system communications are dispatched via “send” instructions that are executed by the message passing send unit **530**. In one embodiment, branch instructions are dispatched to a dedicated branch unit **532** to facilitate SIMD divergence and eventual convergence.

[0142] In one embodiment the vector engine **502** includes one or more SIMD floating point units (FPU(s)) **534** to perform floating-point operations. In one embodiment, the FPU(s) **534** also support integer computation. In one embodiment the FPU(s) **534** can execute up to M number of 32-bit floating-point (or integer) operations, or execute up to 2M 16-bit integer or 16-bit floating-point operations. In one embodiment, at least one of the FPU(s) provides extended math capability to support high-throughput transcendental math functions and double precision 64-bit floating-point. In some embodiments, a set of 8-bit integer SIMD ALUs **535** are also present and may be specifically optimized to perform operations associated with machine learning computations. In one embodiment, the SIMD ALUs are replaced by an additional set of SIMD FPUs **534** that are configurable to perform integer and floating-point operations. In one embodiment, the SIMD FPUs **534** and SIMD ALUs **535** are configurable to execute SIMT programs. In one embodiment, combined SIMD+SIMT operation is supported.

[0143] In one embodiment, arrays of multiple instances of the vector engine **502** can be instantiated in a graphics core. For scalability, product architects can choose the exact number of vector engines per graphics core grouping. In one embodiment the vector engine **502** can execute instructions across a plurality of execution channels. In a further embodiment, each thread executed on the vector engine **502** is executed on a different channel.

[0144] As shown in FIG. 5C, in one embodiment the matrix engine **503** includes an array of processing elements that are configured to perform tensor operations including vector/matrix and matrix/matrix operations, such as but not limited to matrix multiply and/or dot product operations. The matrix engine **503** is configured with M rows and N columns of processing elements (PE **552AA**-PE **552MN**) that include multiplier and adder circuits organized in a pipelined fashion. In one embodiment, the processing elements **552AA**-PE **552MN** make up the physical pipeline stages of an N wide and M deep systolic array that can be used to perform vector/matrix or matrix/matrix operations in a data-parallel manner, including matrix multiply, fused multiply-add, dot product or other general matrix-matrix multiplication (GEMM) operations. In one embodiment the matrix engine **503** supports 16-bit floating point operations, as well as 8-bit, 4-bit, 2-bit, and binary integer operations. The matrix engine **503** can also be configured to accelerate specific machine learning operations. In such embodiments, the matrix engine **503** can be configured with support for the bfloat (brain floating point) 16-bit floating point format or a tensor float 32-bit floating point format (TF32) that have different numbers of mantissa and exponent bits relative to Institute of Electrical and Electronics Engineers (IEEE) 754 formats.

[0145] In one embodiment, during each cycle, each stage can add the result of operations performed at that stage to the output of the previous stage. In other embodiments, the pattern of data movement between the processing elements **552AA**-**552MN** after a set of computational cycles can vary based on the instruction or macro-operation being performed. For example, in one embodiment partial sum loopback is enabled and the processing elements may instead add the output of a current cycle with output generated in the previous cycle. In one embodiment, the final stage of the systolic array can be configured with a loopback to the initial stage of the systolic array. In such embodiment, the number of physical pipeline stages may be decoupled from the number of logical pipeline stages that are supported by the matrix engine **503**. For example, where the processing elements **552AA**-**552MN** are configured as a systolic array of M physical stages, a loopback from stage M to the initial pipeline stage can enable the processing elements **552AA**-PE**552MN** to operate as a systolic array of, for example, 2M, 3M, 4M, etc., logical pipeline stages.

[0146] In one embodiment, the matrix engine **503** includes memory **541A**-**541N**, **542A**-**542M** to store input data in the form of row and column data for input matrices. Memory **542A**-**542M** is configurable to store row elements (**A0**-**Am**) of a first input matrix and memory **541A**-**541N** is

configurable to store column elements (B0-Bn) of a second input matrix. The row and column elements are provided as input to the processing elements 552AA-552MN for processing. In one embodiment, row and column elements of the input matrices can be stored in a systolic register file 540 within the matrix engine 503 before those elements are provided to the memory 541A-541N, 542A-542M. In one embodiment, the systolic register file 540 is excluded and the memory 541A-541N, 542A-542M is loaded from registers in an associated vector engine (e.g., GRF 524 of vector engine 502 of FIG. 5B) or other memory of the graphics core that includes the matrix engine 503 (e.g., data cache/shared local memory 506A for matrix engine 503A of FIG. 5A). Results generated by the processing elements 552AA-552MN are then output to an output buffer and/or written to a register file (e.g., systolic register file 540, GRF 524, data cache/shared local memory 506A-506N) for further processing by other functional units of the graphics processor or for output to memory. [0147] In some embodiments, the matrix engine 503 is configured with support for input sparsity, where multiplication operations for sparse regions of input data can be bypassed by skipping multiply operations that have a zero-value operand. In one embodiment, the processing elements 552AA-552MN are configured to skip the performance of certain operations that have zero value input. In one embodiment, sparsity within input matrices can be detected and operations having known zero output values can be bypassed before being submitted to the processing elements 552AA-552MN. The loading of zero value operands into the processing elements can be bypassed and the processing elements 552AA-552MN can be configured to perform multiplications on the non-zero value input elements. The matrix engine 503 can also be configured with support for output sparsity, such that operations with results that are pre-determined to be zero are bypassed. For input sparsity and/or output sparsity, in one embodiment, metadata is provided to the processing elements 552AA-552MN to indicate, for a processing cycle, which processing elements and/or data channels are to be active during that cycle.

[0148] In one embodiment, the matrix engine 503 includes hardware to enable operations on sparse data having a compressed representation of a sparse matrix that stores non-zero values and metadata that identifies the positions of the non-zero values within the matrix. Exemplary compressed representations include but are not limited to compressed tensor representations such as compressed sparse row (CSR), compressed sparse column (CSC), compressed sparse fiber (CSF) representations. Support for compressed representations enable operations to be performed on input in a compressed tensor format without requiring the compressed representation to be decompressed or decoded. In such embodiment, operations can be performed only on non-zero input values and the resulting non-zero output values can be mapped into an output matrix. In some embodiments, hardware support is also provided for machine-specific lossless data compression formats that are used when transmitting data within hardware or across system busses. Such data may be retained in a compressed format for sparse input data and the matrix engine 503 can use the compression metadata for the compressed data to enable operations to be performed on only non-zero values, or to enable blocks of zero data input to be bypassed for multiply operations.

[0149] In various embodiments, input data can be provided by a programmer in a compressed tensor representation, or a codec can compress input data into the compressed tensor representation or another sparse data encoding. In addition to support for compressed tensor representations, streaming compression of sparse input data can be performed before the data is provided to the processing elements 552AA-552MN. In one embodiment, compression is performed on data written to a cache memory associated with the graphics core cluster 414, with the compression being performed with an encoding that is supported by the matrix engine 503. In one embodiment, the matrix engine 503 includes support for input having structured sparsity in which a pre-determined level or pattern of sparsity is imposed on input data. This data may be compressed to a known compression ratio, with the compressed data being processed by the processing elements 552AA-552MN according to metadata associated with the compressed data.

[0150] FIG. 6 illustrates a tile 600 of a multi-tile processor, according to an embodiment. In one

embodiment, the tile **600** is representative of one of the graphics engine tiles **310A-310D** of FIG. 3B or compute engine tiles **340A-340D** of FIG. 3C. The tile **600** of the multi-tile graphics processor includes an array of graphics core clusters (e.g., graphics core cluster **414A**, graphics core cluster **414B**, through graphics core cluster **414N**), with each graphics core cluster having an array of graphics cores **515A-515N**. The tile **600** also includes a global dispatcher **602** to dispatch threads to processing resources of the tile **600**.

[0151] The tile **600** can include or couple with an L3 cache **606** and memory **610**. In various embodiments, the L3 cache **606** may be excluded or the tile **600** can include additional levels of cache, such as an L4 cache. In one embodiment, each instance of the tile **600** in the multi-tile graphics processor has an associated memory **610**, such as in FIG. 3B and FIG. 3C. In one embodiment, a multi-tile processor can be configured as a multi-chip module in which the L3 cache **606** and/or memory **610** reside on separate chiplets than the graphics core clusters **414A-414N**. In this context, a chiplet is an at least partially packaged integrated circuit that includes distinct units of logic that can be assembled with other chiplets into a larger package. For example, the L3 cache **606** can be included in a dedicated cache chiplet or can reside on the same chiplet as the graphics core clusters **414A-414N**. In one embodiment, the L3 cache **606** can be included in an active base die or active interposer, as illustrated in FIG. 11C.

[0152] A memory fabric **603** enables communication among the graphics core clusters **414A-414N**, L3 cache **606**, and memory **610**. An L2 cache **604** couples with the memory fabric **603** and is configurable to cache transactions performed via the memory fabric **603**. A tile interconnect **608** enables communication with other tiles on the graphics processors and may be one of tile interconnects **323A-323F** of FIGS. 3B and 3C. In embodiments in which the L3 cache **606** is excluded from the tile **600**, the L2 cache **604** may be configured as a combined L2/L3 cache. The memory fabric **603** is configurable to route data to the L3 cache **606** or memory controllers associated with the memory **610** based on the presence or absence of the L3 cache **606** in a specific implementation. The L3 cache **606** can be configured as a per-tile cache that is dedicated to processing resources of the tile **600** or may be a partition of a GPU-wide L3 cache.

[0153] FIG. 7 is a block diagram illustrating graphics processor instruction formats **700** according to some embodiments. In one or more embodiment, the graphics processor cores support an instruction set having instructions in multiple formats. The solid lined boxes illustrate the components that are generally included in a graphics core instruction, while the dashed lines include components that are optional or that are only included in a sub-set of the instructions. In some embodiments, the graphics processor instruction format **700** described and illustrated are macro-instructions, in that they are instructions supplied to the graphics core, as opposed to micro-operations resulting from instruction decode once the instruction is processed. Thus, a single instruction may cause hardware to perform multiple micro-operations.

[0154] In some embodiments, the graphics processor natively supports instructions in a 128-bit instruction format **710**. A 64-bit compacted instruction format **730** is available for some instructions based on the selected instruction, instruction options, and number of operands. The native 128-bit instruction format **710** provides access to all instruction options, while some options and operations are restricted in the 64-bit format **730**. The native instructions available in the 64-bit format **730** vary by embodiment. In some embodiments, the instruction is compacted in part using a set of index values in an index field **713**. The graphics core hardware references a set of compaction tables based on the index values and uses the compaction table outputs to reconstruct a native instruction in the 128-bit instruction format **710**. Other sizes and formats of instruction can be used.

[0155] For each format, instruction opcode **712** defines the operation that the graphics core is to perform. The graphics cores execute each instruction in parallel across the multiple data elements of each operand. For example, in response to an add instruction the graphics core performs a simultaneous add operation across each color channel representing a texture element or picture element. By default, the graphics core performs each instruction across all data channels of the

operands. In some embodiments, instruction control field **714** enables control over certain execution options, such as channels selection (e.g., predication) and data channel order (e.g., swizzle). For instructions in the 128-bit instruction format **710** an exec-size field **716** limits the number of data channels that will be executed in parallel. In some embodiments, exec-size field **716** is not available for use in the 64-bit compact instruction format **730**.

[0156] Some graphics core instructions have up to three operands including two source operands, src**0** **720**, src**1** **722**, and one destination **718**. In some embodiments, the graphics cores support dual destination instructions, where one of the destinations is implied. Data manipulation instructions can have a third source operand (e.g., SRC**2** **724**), where the instruction opcode **712** determines the number of source operands. An instruction's last source operand can be an immediate (e.g., hard-coded) value passed with the instruction.

[0157] In some embodiments, the 128-bit instruction format **710** includes an access/address mode field **726** specifying, for example, whether direct register addressing mode or indirect register addressing mode is used. When direct register addressing mode is used, the register address of one or more operands is directly provided by bits in the instruction.

[0158] In some embodiments, the 128-bit instruction format **710** includes an access/address mode field **726**, which specifies an address mode and/or an access mode for the instruction. In one embodiment the access mode is used to define a data access alignment for the instruction. Some embodiments support access modes including a 16-byte aligned access mode and a 1-byte aligned access mode, where the byte alignment of the access mode determines the access alignment of the instruction operands. For example, when in a first mode, the instruction may use byte-aligned addressing for source and destination operands and when in a second mode, the instruction may use 16-byte-aligned addressing for all source and destination operands.

[0159] In one embodiment, the address mode portion of the access/address mode field **726** determines whether the instruction is to use direct or indirect addressing. When direct register addressing mode is used bits in the instruction directly provide the register address of one or more operands. When indirect register addressing mode is used, the register address of one or more operands may be computed based on an address register value and an address immediate field in the instruction.

[0160] In some embodiments instructions are grouped based on opcode **712** bit-fields to simplify Opcode decode **740**. For an 8-bit opcode, bits **4**, **5**, and **6** allow the graphics core to determine the type of opcode. The precise opcode grouping shown is merely an example. In some embodiments, a move and logic opcode group **742** includes data movement and logic instructions (e.g., move (mov), compare (cmp)). In some embodiments, move and logic group **742** shares the five most significant bits (MSB), where move (mov) instructions are in the form of 0000xxxxb and logic instructions are in the form of 0001xxxxb. A flow control instruction group **744** (e.g., call, jump (jmp)) includes instructions in the form of 0010xxxxb (e.g., 0x20). A miscellaneous instruction group **746** includes a mix of instructions, including synchronization instructions (e.g., wait, send) in the form of 0011xxxxb (e.g., 0x30). A parallel math instruction group **748** includes component-wise arithmetic instructions (e.g., add, multiply (mul)) in the form of 0100xxxxb (e.g., 0x40). The parallel math instruction group **748** performs the arithmetic operations in parallel across data channels. The vector math group **750** includes arithmetic instructions (e.g., dp4) in the form of 0101xxxxb (e.g., 0x50). The vector math group performs arithmetic such as dot product calculations on vector operands. The illustrated opcode decode **740**, in one embodiment, can be used to determine which portion of a graphics core will be used to execute a decoded instruction. For example, some instructions may be designated as systolic instructions that will be performed by a systolic array. Other instructions, such as ray-tracing instructions (not shown) can be routed to a ray-tracing core or ray-tracing logic within a slice or partition of execution logic.

Graphics Pipeline

[0161] FIG. **8** is a block diagram of another embodiment of a graphics processor **800**. Elements of

FIG. 8 having the same reference numbers (or names) as the elements of any other figure herein can operate or function in any manner similar to that described elsewhere herein, but are not limited to such.

[0162] In some embodiments, graphics processor **800** includes a geometry pipeline **820**, a media pipeline **830**, a display engine **840**, thread execution logic **850**, and a render output pipeline **870**. In some embodiments, graphics processor **800** is a graphics processor within a multi-core processing system that includes one or more general-purpose processing cores. The graphics processor is controlled by register writes to one or more control registers (not shown) or via commands issued to graphics processor **800** via a ring interconnect **802**. In some embodiments, ring interconnect **802** couples graphics processor **800** to other processing components, such as other graphics processors or general-purpose processors. Commands from ring interconnect **802** are interpreted by a command streamer **803**, which supplies instructions to individual components of the geometry pipeline **820** or the media pipeline **830**.

[0163] In some embodiments, command streamer **803** directs the operation of a vertex fetcher **805** that reads vertex data from memory and executes vertex-processing commands provided by command streamer **803**. In some embodiments, vertex fetcher **805** provides vertex data to a vertex shader **807**, which performs coordinate space transformation and lighting operations to each vertex. In some embodiments, vertex fetcher **805** and vertex shader **807** execute vertex-processing instructions by dispatching execution threads to graphics cores **852A-852B** via a thread dispatcher **831**.

[0164] In some embodiments, graphics cores **852A-852B** are an array of vector processors having an instruction set for performing graphics and media operations. In some embodiments, graphics cores **852A-852B** have an attached L1 cache **851** that is specific for each array or shared between the arrays. The cache can be configured as a data cache, an instruction cache, or a single cache that is partitioned to contain data and instructions in different partitions.

[0165] In some embodiments, geometry pipeline **820** includes tessellation components to perform hardware-accelerated tessellation of 3D objects. In some embodiments, a programmable hull shader **811** configures the tessellation operations. A programmable domain shader **817** provides back-end evaluation of tessellation output. A tessellator **813** operates at the direction of hull shader **811** and contains special purpose logic to generate a set of detailed geometric objects based on a coarse geometric model that is provided as input to geometry pipeline **820**. In some embodiments, if tessellation is not used, tessellation components (e.g., hull shader **811**, tessellator **813**, and domain shader **817**) can be bypassed. The tessellation components can operate based on data received from the vertex shader **807**.

[0166] In some embodiments, complete geometric objects can be processed by a geometry shader **819** via one or more threads dispatched to graphics cores **852A-852B** or can proceed directly to the clipper **829**. In some embodiments, the geometry shader operates on entire geometric objects, rather than vertices or patches of vertices as in previous stages of the graphics pipeline. If the tessellation is disabled the geometry shader **819** receives input from the vertex shader **807**. In some embodiments, geometry shader **819** is programmable by a geometry shader program to perform geometry tessellation if the tessellation units are disabled.

[0167] Before rasterization, a clipper **829** processes vertex data. The clipper **829** may be a fixed function clipper or a programmable clipper having clipping and geometry shader functions. In some embodiments, a rasterizer and depth test component **873** in the render output pipeline **870** dispatches pixel shaders to convert the geometric objects into per pixel representations. In some embodiments, pixel shader logic is included in thread execution logic **850**. In some embodiments, an application can bypass the rasterizer and depth test component **873** and access un-rasterized vertex data via a stream out unit **823**.

[0168] The graphics processor **800** has an interconnect bus, interconnect fabric, or some other interconnect mechanism that allows data and message passing amongst the major components of

the processor. In some embodiments, graphics cores **852A-852B** and associated logic units (e.g., L1 cache **851**, sampler **854**, texture cache **858**, etc.) interconnect via a data port **856** to perform memory access and communicate with render output pipeline components of the processor. In some embodiments, sampler **854**, caches **851**, **858** and graphics cores **852A-852B** each have separate memory access paths. In one embodiment the texture cache **858** can also be configured as a sampler cache.

[0169] In some embodiments, render output pipeline **870** contains a rasterizer and depth test component **873** that converts vertex-based objects into an associated pixel-based representation. In some embodiments, the rasterizer logic includes a windower/masker unit to perform fixed function triangle and line rasterization. An associated render cache **878** and depth cache **879** are also available in some embodiments. A pixel operations component **877** performs pixel-based operations on the data, though in some instances, pixel operations associated with 2D operations (e.g., bit block image transfers with blending) are performed by the 2D engine **841**, or substituted at display time by the display controller **843** using overlay display planes. In some embodiments, a shared L3 cache **875** is available to all graphics components, allowing the sharing of data without the use of main system memory.

[0170] In some embodiments, media pipeline **830** includes a media engine **837** and a video front-end **834**. In some embodiments, video front-end **834** receives pipeline commands from the command streamer **803**. In some embodiments, media pipeline **830** includes a separate command streamer. In some embodiments, video front-end **834** processes media commands before sending the command to the media engine **837**. In some embodiments, media engine **837** includes thread spawning functionality to spawn threads for dispatch to thread execution logic **850** via thread dispatcher **831**.

[0171] In some embodiments, graphics processor **800** includes a display engine **840**. In some embodiments, display engine **840** is external to processor **800** and couples with the graphics processor via the ring interconnect **802**, or some other interconnect bus or fabric. In some embodiments, display engine **840** includes a 2D engine **841** and a display controller **843**. In some embodiments, display engine **840** contains special purpose logic capable of operating independently of the 3D pipeline. In some embodiments, display controller **843** couples with a display device (not shown), which may be a system integrated display device, as in a laptop computer, or an external display device attached via a display device connector.

[0172] In some embodiments, the geometry pipeline **820** and media pipeline **830** are configurable to perform operations based on multiple graphics and media programming interfaces and are not specific to any one application programming interface (API). In some embodiments, driver software for the graphics processor translates API calls that are specific to a particular graphics or media library into commands that can be processed by the graphics processor. In some embodiments, support is provided for the Open Graphics Library (OpenGL), Open Computing Language (OpenCL), and/or Vulkan graphics and compute API, all from the Khronos Group. In some embodiments, support may also be provided for the Direct3D library from the Microsoft Corporation. In some embodiments, a combination of these libraries may be supported. Support may also be provided for the Open Source Computer Vision Library (OpenCV). A future API with a compatible 3D pipeline would also be supported if a mapping can be made from the pipeline of the future API to the pipeline of the graphics processor.

Graphics Pipeline Programming

[0173] FIG. **9A** is a block diagram illustrating a graphics processor command format **900** that may be used to program graphics processing pipelines according to some embodiments. FIG. **9B** is a block diagram illustrating a graphics processor command sequence **910** according to an embodiment. The solid lined boxes in FIG. **9A** illustrate the components that are generally included in a graphics command while the dashed lines include components that are optional or that are only included in a sub-set of the graphics commands. The exemplary graphics processor command

format **900** of FIG. **9A** includes data fields to identify a client **902**, a command operation code (opcode) **904**, and a data field **906** for the command. A sub-opcode **905** and a command size **908** are also included in some commands.

[0174] In some embodiments, client **902** specifies the client unit of the graphics device that processes the command data. In some embodiments, a graphics processor command parser examines the client field of each command to condition the further processing of the command and route the command data to the appropriate client unit. In some embodiments, the graphics processor client units include a memory interface unit, a render unit, a 2D unit, a 3D unit, and a media unit. Each client unit has a corresponding processing pipeline that processes the commands. Once the command is received by the client unit, the client unit reads the opcode **904** and, if present, sub-opcode **905** to determine the operation to perform. The client unit performs the command using information in data field **906**. For some commands an explicit command size **908** is expected to specify the size of the command. In some embodiments, the command parser automatically determines the size of at least some of the commands based on the command opcode. In some embodiments commands are aligned via multiples of a double word. Other command formats can be used.

[0175] The flow diagram in FIG. **9B** illustrates an exemplary graphics processor command sequence **910**. In some embodiments, software or firmware of a data processing system that features an embodiment of a graphics processor uses a version of the command sequence shown to set up, execute, and terminate a set of graphics operations. A sample command sequence is shown and described for purposes of example only as embodiments are not limited to these specific commands or to this command sequence. Moreover, the commands may be issued as batch of commands in a command sequence, such that the graphics processor will process the sequence of commands in at least partially concurrence.

[0176] In some embodiments, the graphics processor command sequence **910** may begin with a pipeline flush command **912** to cause any active graphics pipeline to complete the currently pending commands for the pipeline. In some embodiments, the 3D pipeline **922** and the media pipeline **924** do not operate concurrently. The pipeline flush is performed to cause the active graphics pipeline to complete any pending commands. In response to a pipeline flush, the command parser for the graphics processor will pause command processing until the active drawing engines complete pending operations and the relevant read caches are invalidated. Optionally, any data in the render cache that is marked 'dirty' can be flushed to memory. In some embodiments, pipeline flush command **912** can be used for pipeline synchronization or before placing the graphics processor into a low power state.

[0177] In some embodiments, a pipeline select command **913** is used when a command sequence requires the graphics processor to explicitly switch between pipelines. In some embodiments, a pipeline select command **913** is required only once within an execution context before issuing pipeline commands unless the context is to issue commands for both pipelines. In some embodiments, a pipeline flush command **912** is required immediately before a pipeline switch via the pipeline select command **913**.

[0178] In some embodiments, a pipeline control command **914** configures a graphics pipeline for operation and is used to program the 3D pipeline **922** and the media pipeline **924**. In some embodiments, pipeline control command **914** configures the pipeline state for the active pipeline. In one embodiment, the pipeline control command **914** is used for pipeline synchronization and to clear data from one or more cache memories within the active pipeline before processing a batch of commands.

[0179] In some embodiments, commands related to the return buffer state **916** are used to configure a set of return buffers for the respective pipelines to write data. Some pipeline operations require the allocation, selection, or configuration of one or more return buffers into which the operations write intermediate data during processing. In some embodiments, the graphics processor also uses

one or more return buffers to store output data and to perform cross thread communication. In some embodiments, the return buffer state **916** includes selecting the size and number of return buffers to use for a set of pipeline operations.

[0180] The remaining commands in the command sequence differ based on the active pipeline for operations. Based on a pipeline determination **920**, the command sequence is tailored to the 3D pipeline **922** beginning with the 3D pipeline state **930** or the media pipeline **924** beginning at the media pipeline state **940**.

[0181] The commands to configure the 3D pipeline state **930** include 3D state setting commands for vertex buffer state, vertex element state, constant color state, depth buffer state, and other state variables that are to be configured before 3D primitive commands are processed. The values of these commands are determined at least in part based on the particular 3D API in use. In some embodiments, 3D pipeline state **930** commands are also able to selectively disable or bypass certain pipeline elements if those elements will not be used.

[0182] In some embodiments, 3D primitive **932** command is used to submit 3D primitives to be processed by the 3D pipeline. Commands and associated parameters that are passed to the graphics processor via the 3D primitive **932** command are forwarded to the vertex fetch function in the graphics pipeline. The vertex fetch function uses the 3D primitive **932** command data to generate vertex data structures. The vertex data structures are stored in one or more return buffers. In some embodiments, 3D primitive **932** command is used to perform vertex operations on 3D primitives via vertex shaders. To process vertex shaders, 3D pipeline **922** dispatches shader programs to the graphics cores.

[0183] In some embodiments, 3D pipeline **922** is triggered via an execute **934** command or event. In some embodiments, a register write triggers command execution. In some embodiments execution is triggered via a ‘go’ or ‘kick’ command in the command sequence. In one embodiment, command execution is triggered using a pipeline synchronization command to flush the command sequence through the graphics pipeline. The 3D pipeline will perform geometry processing for the 3D primitives. Once operations are complete, the resulting geometric objects are rasterized and the pixel engine colors the resulting pixels. Additional commands to control pixel shading and pixel back-end operations may also be included for those operations.

[0184] In some embodiments, the graphics processor command sequence **910** follows the media pipeline **924** path when performing media operations. In general, the specific use and manner of programming for the media pipeline **924** depends on the media or compute operations to be performed. Specific media decode operations may be offloaded to the media pipeline during media decode. In some embodiments, the media pipeline can also be bypassed and media decode can be performed in whole or in part using resources provided by one or more general-purpose processing cores. In one embodiment, the media pipeline also includes elements for general-purpose graphics processor unit (GPGPU) operations, where the graphics processor is used to perform SIMD vector operations using computational shader programs that are not explicitly related to the rendering of graphics primitives.

[0185] In some embodiments, media pipeline **924** is configured in a similar manner as the 3D pipeline **922**. A set of commands to configure the media pipeline state **940** are dispatched or placed into a command queue before the media object commands **942**. In some embodiments, commands for the media pipeline state **940** include data to configure the media pipeline elements that will be used to process the media objects. This includes data to configure the video decode and video encode logic within the media pipeline, such as encode or decode format. In some embodiments, commands for the media pipeline state **940** also support the use of one or more pointers to “indirect” state elements that contain a batch of state settings.

[0186] In some embodiments, media object commands **942** supply pointers to media objects for processing by the media pipeline. The media objects include memory buffers containing video data to be processed. In some embodiments, all media pipeline states must be valid before issuing a

media object command **942**. Once the pipeline state is configured and media object commands **942** are queued, the media pipeline **924** is triggered via an execute command **944** or an equivalent execute event (e.g., register write). Output from media pipeline **924** may then be post processed by operations provided by the 3D pipeline **922** or the media pipeline **924**. In some embodiments, GPGPU operations are configured and executed in a similar manner as media operations.

Graphics Software Architecture

[0187] FIG. **10** illustrates an exemplary graphics software architecture for a data processing system **1000** according to some embodiments. In some embodiments, software architecture includes a 3D graphics application **1010**, an operating system **1020**, and at least one processor **1030**. In some embodiments, processor **1030** includes a graphics processor **1032** and one or more general-purpose processor core(s) **1034**. The graphics application **1010** and operating system **1020** each execute in the system memory **1050** of the data processing system.

[0188] In some embodiments, 3D graphics application **1010** contains one or more shader programs including shader instructions **1012**. The shader language instructions may be in a high-level shader language, such as the High-Level Shader Language (HLSL) of Direct3D, the OpenGL Shader Language (GLSL), and so forth. The application also includes executable instructions **1014** in a machine language suitable for execution by the general-purpose processor core **1034**. The application also includes graphics objects **1016** defined by vertex data.

[0189] In some embodiments, operating system **1020** is a Microsoft® Windows® operating system from the Microsoft Corporation, a proprietary UNIX-like operating system, or an open source UNIX-like operating system using a variant of the Linux kernel. The operating system **1020** can support a graphics API **1022** such as the Direct3D API, the OpenGL API, or the Vulkan API. When the Direct3D API is in use, the operating system **1020** uses a front-end shader compiler **1024** to compile any shader instructions **1012** in HLSL into a lower-level shader language. The compilation may be a just-in-time (JIT) compilation or the application can perform shader pre-compilation. In some embodiments, high-level shaders are compiled into low-level shaders during the compilation of the 3D graphics application **1010**. In some embodiments, the shader instructions **1012** are provided in an intermediate form, such as a version of the Standard Portable Intermediate Representation (SPIR) used by the Vulkan API.

[0190] In some embodiments, user mode graphics driver **1026** contains a back-end shader compiler **1027** to convert the shader instructions **1012** into a hardware specific representation. When the OpenGL API is in use, shader instructions **1012** in the GLSL high-level language are passed to a user mode graphics driver **1026** for compilation. In some embodiments, user mode graphics driver **1026** uses operating system kernel mode functions **1028** to communicate with a kernel mode graphics driver **1029**. In some embodiments, kernel mode graphics driver **1029** communicates with graphics processor **1032** to dispatch commands and instructions.

IP Core Implementations

[0191] One or more aspects of at least one embodiment may be implemented by representative code stored on a machine-readable medium which represents and/or defines logic within an integrated circuit such as a processor. For example, the machine-readable medium may include instructions which represent various logic within the processor. When read by a machine, the instructions may cause the machine to fabricate the logic to perform the techniques described herein. Such representations, known as “IP cores,” are reusable units of logic for an integrated circuit that may be stored on a tangible, machine-readable medium as a hardware model that describes the structure of the integrated circuit. The hardware model may be supplied to various customers or manufacturing facilities, which load the hardware model on fabrication machines that manufacture the integrated circuit. The integrated circuit may be fabricated such that the circuit performs operations described in association with any of the embodiments described herein.

[0192] FIG. **11A** is a block diagram illustrating an IP core development system **1100** that may be used to manufacture an integrated circuit to perform operations according to an embodiment. The

IP core development system **1100** may be used to generate modular, re-usable designs that can be incorporated into a larger design or used to construct an entire integrated circuit (e.g., an SOC integrated circuit). A design facility **1130** can generate a software simulation **1110** of an IP core design in a high-level programming language (e.g., C/C++). The software simulation **1110** can be used to design, test, and verify the behavior of the IP core using a simulation model **1112**. The simulation model **1112** may include functional, behavioral, and/or timing simulations. A register transfer level (RTL) design **1115** can then be created or synthesized from the simulation model **1112**. The RTL design **1115** is an abstraction of the behavior of the integrated circuit that models the flow of digital signals between hardware registers, including the associated logic performed using the modeled digital signals. In addition to an RTL design **1115**, lower-level designs at the logic level or transistor level may also be created, designed, or synthesized. Thus, the particular details of the initial design and simulation may vary.

[0193] The RTL design **1115** or equivalent may be further synthesized by the design facility into a hardware model **1120**, which may be in a hardware description language (HDL), or some other representation of physical design data. The HDL may be further simulated or tested to verify the IP core design. The IP core design can be stored for delivery to a 3rd party fabrication facility **1165** using non-volatile memory **1140** (e.g., hard disk, flash memory, or any non-volatile storage medium). Alternatively, the IP core design may be transmitted (e.g., via the Internet) over a wired connection **1150** or wireless connection **1160**. The fabrication facility **1165** may then fabricate an integrated circuit that is based at least in part on the IP core design. The fabricated integrated circuit can be configured to perform operations in accordance with at least one embodiment described herein.

[0194] FIG. **11B** illustrates a cross-section side view of an integrated circuit package assembly **1170**, according to some embodiments described herein. The integrated circuit package assembly **1170** illustrates an implementation of one or more processor or accelerator devices as described herein. The package assembly **1170** includes multiple units of hardware logic **1172**, **1174** connected to a substrate **1180**. The logic **1172**, **1174** may be implemented at least partly in configurable logic or fixed-functionality logic hardware, and can include one or more portions of any of the processor core(s), graphics processor(s), or other accelerator devices described herein. Each unit of logic **1172**, **1174** can be implemented within a semiconductor die and coupled with the substrate **1180** via an interconnect structure **1173**. The interconnect structure **1173** may be configured to route electrical signals between the logic **1172**, **1174** and the substrate **1180**, and can include interconnects such as, but not limited to bumps or pillars. In some embodiments, the interconnect structure **1173** may be configured to route electrical signals such as, for example, input/output (I/O) signals and/or power or ground signals associated with the operation of the logic **1172**, **1174**. In some embodiments, the substrate **1180** is an epoxy-based laminate substrate. The substrate **1180** may include other suitable types of substrates in other embodiments. The package assembly **1170** can be connected to other electrical devices via a package interconnect **1183**. The package interconnect **1183** may be coupled to a surface of the substrate **1180** to route electrical signals to other electrical devices, such as a motherboard, other chipset, or multi-chip module.

[0195] In some embodiments, the units of logic **1172**, **1174** are electrically coupled with a bridge **1182** that is configured to route electrical signals between the logic **1172**, **1174**. The bridge **1182** may be a dense interconnect structure that provides a route for electrical signals. The bridge **1182** may include a bridge substrate composed of glass or a suitable semiconductor material. Electrical routing features can be formed on the bridge substrate to provide a chip-to-chip connection between the logic **1172**, **1174**.

[0196] Although two units of logic **1172**, **1174** and a bridge **1182** are illustrated, embodiments described herein may include more or fewer logic units on one or more dies. The one or more dies may be connected by zero or more bridges, as the bridge **1182** may be excluded when the logic is included on a single die. Alternatively, multiple dies or units of logic can be connected by one or

more bridges. Additionally, multiple logic units, dies, and bridges can be connected together in other possible configurations, including three-dimensional configurations.

[0197] FIG. **11C** illustrates a package assembly **1190** that includes multiple units of hardware logic chiplets connected to a substrate **1180**. A graphics processing unit, parallel processor, and/or compute accelerator as described herein can be composed from diverse silicon chiplets that are separately manufactured. A diverse set of chiplets with different IP core logic can be assembled into a single device. Additionally, the chiplets can be integrated into a base die or base chiplet using active interposer technology. The concepts described herein enable the interconnection and communication between the different forms of IP within the GPU. IP cores can be manufactured using different process technologies and composed during manufacturing, which avoids the complexity of converging multiple IPs, especially on a large SoC with several flavors IPs, to the same manufacturing process. Enabling the use of multiple process technologies improves the time to market and provides a cost-effective way to create multiple product SKUs. Additionally, the disaggregated IPs are more amenable to being power gated independently, components that are not in use on a given workload can be powered off, reducing overall power consumption.

[0198] In various embodiments a package assembly **1190** can include components and chiplets that are interconnected by a fabric **1185** and/or one or more bridges **1187**. The chiplets within the package assembly **1190** may have a 2.5D arrangement using Chip-on-Wafer-on-Substrate stacking in which multiple dies are stacked side-by-side on a silicon interposer **1189** that couples the chiplets with the substrate **1180**. The substrate **1180** includes electrical connections to the package interconnect **1183**. In one embodiment the silicon interposer **1189** is a passive interposer that includes through-silicon vias (TSVs) to electrically couple chiplets within the package assembly **1190** to the substrate **1180**. In one embodiment, silicon interposer **1189** is an active interposer that includes embedded logic in addition to TSVs. In such embodiment, the chiplets within the package assembly **1190** are arranged using 3D face to face die stacking on top of the active interposer **1189**. The active interposer **1189** can include hardware logic for I/O **1191**, cache memory **1192**, and other hardware logic **1193**, in addition to interconnect fabric **1185** and a silicon bridge **1187**. The fabric **1185** enables communication between the various logic chiplets **1172**, **1174** and the logic **1191**, **1193** within the active interposer **1189**. The fabric **1185** may be an NoC interconnect or another form of packet switched fabric that switches data packets between components of the package assembly. For complex assemblies, the fabric **1185** may be a dedicated chiplet enables communication between the various hardware logic of the package assembly **1190**.

[0199] Bridge structures **1187** within the active interposer **1189** may be used to facilitate a point-to-point interconnect between, for example, logic or I/O chiplets **1174** and memory chiplets **1175**. In some implementations, bridge structures **1187** may also be embedded within the substrate **1180**. The hardware logic chiplets can include special purpose hardware logic chiplets **1172**, logic or I/O chiplets **1174**, and/or memory chiplets **1175**. The hardware logic chiplets **1172** and logic or I/O chiplets **1174** may be implemented at least partly in configurable logic or fixed-functionality logic hardware and can include one or more portions of any of the processor core(s), graphics processor(s), parallel processors, or other accelerator devices described herein. The memory chiplets **1175** can be DRAM (e.g., GDDR, HBM) memory or cache (SRAM) memory. Cache memory **1192** within the active interposer **1189** (or substrate **1180**) can act as a global cache for the package assembly **1190**, part of a distributed global cache, or as a dedicated cache for the fabric **1185**.

[0200] Each chiplet can be fabricated as separate semiconductor die and coupled with a base die that is embedded within or coupled with the substrate **1180**. The coupling with the substrate **1180** can be performed via an interconnect structure **1173**. The interconnect structure **1173** may be configured to route electrical signals between the various chiplets and logic within the substrate **1180**. The interconnect structure **1173** can include interconnects such as, but not limited to bumps or pillars. In some embodiments, the interconnect structure **1173** may be configured to route

electrical signals such as, for example, input/output (I/O) signals and/or power or ground signals associated with the operation of the logic, I/O, and memory chiplets. In one embodiment, an additional interconnect structure couples the active interposer **1189** with the substrate **1180**.

[0201] In some embodiments, the substrate **1180** is an epoxy-based laminate substrate. The substrate **1180** may include other suitable types of substrates in other embodiments. The package assembly **1190** can be connected to other electrical devices via a package interconnect **1183**. The package interconnect **1183** may be coupled to a surface of the substrate **1180** to route electrical signals to other electrical devices, such as a motherboard, other chipset, or multi-chip module.

[0202] In some embodiments, a logic or I/O chiplet **1174** and a memory chiplet **1175** can be electrically coupled via a bridge **1187** that is configured to route electrical signals between the logic or I/O chiplet **1174** and a memory chiplet **1175**. The bridge **1187** may be a dense interconnect structure that provides a route for electrical signals. The bridge **1187** may include a bridge substrate composed of glass or a suitable semiconductor material. Electrical routing features can be formed on the bridge substrate to provide a chip-to-chip connection between the logic or I/O chiplet **1174** and a memory chiplet **1175**. The bridge **1187** may also be referred to as a silicon bridge or an interconnect bridge. For example, the bridge **1187**, in some embodiments, is an Embedded Multi-die Interconnect Bridge (EMIB). In some embodiments, the bridge **1187** may simply be a direct connection from one chiplet to another chiplet.

[0203] FIG. **11D** illustrates a package assembly **1194** including interchangeable chiplets **1195**, according to an embodiment. The interchangeable chiplets **1195** can be assembled into standardized slots on one or more base chiplets **1196**, **1198**. The base chiplets **1196**, **1198** can be coupled via a bridge interconnect **1197**, which can be similar to the other bridge interconnects described herein and may be, for example, an EMIB. Memory chiplets can also be connected to logic or I/O chiplets via a bridge interconnect. I/O and logic chiplets can communicate via an interconnect fabric. The base chiplets can each support one or more slots in a standardized format for one of logic or I/O or memory/cache.

[0204] In one embodiment, SRAM and power delivery circuits can be fabricated into one or more of the base chiplets **1196**, **1198**, which can be fabricated using a different process technology relative to the interchangeable chiplets **1195** that are stacked on top of the base chiplets. For example, the base chiplets **1196**, **1198** can be fabricated using a larger process technology, while the interchangeable chiplets can be manufactured using a smaller process technology. One or more of the interchangeable chiplets **1195** may be memory (e.g., DRAM) chiplets. Different memory densities can be selected for the package assembly **1194** based on the power, and/or performance targeted for the product that uses the package assembly **1194**. Additionally, logic chiplets with a different number of type of functional units can be selected at time of assembly based on the power, and/or performance targeted for the product. Additionally, chiplets containing IP logic cores of differing types can be inserted into the interchangeable chiplet slots, enabling hybrid processor designs that can mix and match different technology IP blocks.

Exemplary System on a Chip Integrated Circuit

[0205] FIGS. **12-13B** illustrate exemplary integrated circuits and associated graphics processors that may be fabricated using one or more IP cores, according to various embodiments described herein. In addition to what is illustrated, other logic and circuits may be included, including additional graphics processors/cores, peripheral interface controllers, or general-purpose processor cores.

[0206] FIG. **12** is a block diagram illustrating an exemplary system on a chip integrated circuit **1200** that may be fabricated using one or more IP cores, according to an embodiment. Exemplary integrated circuit **1200** includes one or more application processor(s) **1205** (e.g., CPUs), at least one graphics processor **1210**, and may additionally include an image processor **1215** and/or a video processor **1220**, any of which may be a modular IP core from the same or multiple different design facilities. Integrated circuit **1200** includes peripheral or bus logic including a USB controller **1225**,

UART controller **1230**, an SPI/SDIO controller **1235**, and an I.sup.2S/I.sup.2C controller **1240**. Additionally, the integrated circuit can include a display device **1245** coupled to one or more of a high-definition multimedia interface (HDMI) controller **1250** and a mobile industry processor interface (MIPI) display interface **1255**. Storage may be provided by a flash memory subsystem **1260** including flash memory and a flash memory controller. Memory interface may be provided via a memory controller **1265** for access to SDRAM or SRAM memory devices. Some integrated circuits additionally include an embedded security engine **1270**.

[0207] FIGS. **13A-13B** are block diagrams illustrating exemplary graphics processors for use within an SoC, according to embodiments described herein. FIG. **13A** illustrates an exemplary graphics processor **1310** of a system on a chip integrated circuit that may be fabricated using one or more IP cores, according to an embodiment. FIG. **13B** illustrates an additional exemplary graphics processor **1340** of a system on a chip integrated circuit that may be fabricated using one or more IP cores, according to an embodiment. Graphics processor **1310** of FIG. **13A** is an example of a low power graphics processor core. Graphics processor **1340** of FIG. **13B** is an example of a higher performance graphics processor core. Each of graphics processor **1310** and graphics processor **1340** can be variants of the graphics processor **1210** of FIG. **12**.

[0208] As shown in FIG. **13A**, graphics processor **1310** includes a vertex processor **1305** and one or more fragment processor(s) **1315A-1315N** (e.g., **1315A**, **1315B**, **1315C**, **1315D**, through **1315N-1**, and **1315N**). Graphics processor **1310** can execute different shader programs via separate logic, such that the vertex processor **1305** is optimized to execute operations for vertex shader programs, while the one or more fragment processor(s) **1315A-1315N** execute fragment (e.g., pixel) shading operations for fragment or pixel shader programs. The vertex processor **1305** performs the vertex processing stage of the 3D graphics pipeline and generates primitives and vertex data. The fragment processor(s) **1315A-1315N** use the primitive and vertex data generated by the vertex processor **1305** to produce a framebuffer that is displayed on a display device. In one embodiment, the fragment processor(s) **1315A-1315N** are optimized to execute fragment shader programs as provided for in the OpenGL API, which may be used to perform similar operations as a pixel shader program as provided for in the Direct 3D API.

[0209] Graphics processor **1310** additionally includes one or more memory management units (MMUs) **1320A-1320B**, cache(s) **1325A-1325B**, and circuit interconnect(s) **1330A-1330B**. The one or more MMU(s) **1320A-1320B** provide for virtual to physical address mapping for the graphics processor **1310**, including for the vertex processor **1305** and/or fragment processor(s) **1315A-1315N**, which may reference vertex or image/texture data stored in memory, in addition to vertex or image/texture data stored in the one or more cache(s) **1325A-1325B**. In one embodiment the one or more MMU(s) **1320A-1320B** may be synchronized with other MMUs within the system, including one or more MMUs associated with the one or more application processor(s) **1205**, image processor **1215**, and/or video processor **1220** of FIG. **12**, such that each processor **1205-1220** can participate in a shared or unified virtual memory system. The one or more circuit interconnect(s) **1330A-1330B** enable graphics processor **1310** to interface with other IP cores within the SoC, either via an internal bus of the SoC or via a direct connection, according to embodiments.

[0210] As shown FIG. **13B**, graphics processor **1340** includes the one or more MMU(s) **1320A-1320B**, cache(s) **1325A-1325B**, and circuit interconnect(s) **1330A-1330B** of the graphics processor **1310** of FIG. **13A**. Graphics processor **1340** includes one or more shader core(s) **1355A-1355N** (e.g., **1355A**, **1355B**, **1355C**, **1355D**, **1355E**, **1355F**, through **1355N-1**, and **1355N**), which provides for a unified shader core architecture in which a single core or type or core can execute all types of programmable shader code, including shader program code to implement vertex shaders, fragment shaders, and/or compute shaders. The unified shader core architecture is also configurable to execute direct compiled high-level GPGPU programs (e.g., CUDA). The exact number of shader cores present can vary among embodiments and implementations. Additionally, graphics processor **1340** includes an inter-core task manager **1345**, which acts as a thread dispatcher to dispatch

execution threads to one or more shader cores 1355A-1355N and a tiling unit 1358 to accelerate tiling operations for tile-based rendering, in which rendering operations for a scene are subdivided in image space, for example to exploit local spatial coherence within a scene or to optimize use of internal caches.

Machine Learning Overview

[0211] The architecture described above can be applied to perform training and inference operations using machine learning models. Machine learning has been successful at solving many kinds of tasks. The computations that arise when training and using machine learning algorithms (e.g., neural networks) lend themselves naturally to efficient parallel implementations. Accordingly, parallel processors such as general-purpose graphics processing units (GPGPUs) have played a significant role in the practical implementation of deep neural networks. Parallel graphics processors with single instruction, multiple thread (SIMT) architectures are designed to maximize the amount of parallel processing in the graphics pipeline. In an SIMT architecture, groups of parallel threads attempt to execute program instructions synchronously together as often as possible to increase processing efficiency. The efficiency provided by parallel machine learning algorithm implementations allows the use of high-capacity networks and enables those networks to be trained on larger datasets.

[0212] A machine learning algorithm is an algorithm that can learn based on a set of data. For example, machine learning algorithms can be designed to model high-level abstractions within a data set. For example, image recognition algorithms can be used to determine which of several categories to which a given input belong; regression algorithms can output a numerical value given an input; and pattern recognition algorithms can be used to generate translated text or perform text to speech and/or speech recognition.

[0213] An exemplary type of machine learning algorithm is a neural network. There are many types of neural networks; a simple type of neural network is a feedforward network. A feedforward network may be implemented as an acyclic graph in which the nodes are arranged in layers. Typically, a feedforward network topology includes an input layer and an output layer that are separated by at least one hidden layer. The hidden layer transforms input received by the input layer into a representation that is useful for generating output in the output layer. The network nodes are fully connected via edges to the nodes in adjacent layers, but there are no edges between nodes within each layer. Data received at the nodes of an input layer of a feedforward network are propagated (i.e., “fed forward”) to the nodes of the output layer via an activation function that calculates the states of the nodes of each successive layer in the network based on coefficients (“weights”) respectively associated with each of the edges connecting the layers. Depending on the specific model being represented by the algorithm being executed, the output from the neural network algorithm can take various forms.

[0214] Before a machine learning algorithm can be used to model a particular problem, the algorithm is trained using a training data set. Training a neural network involves selecting a network topology, using a set of training data representing a problem being modeled by the network, and adjusting the weights until the network model performs with a minimal error for all instances of the training data set. For example, during a supervised learning training process for a neural network, the output produced by the network in response to the input representing an instance in a training data set is compared to the “correct” labeled output for that instance, an error signal representing the difference between the output and the labeled output is calculated, and the weights associated with the connections are adjusted to minimize that error as the error signal is backward propagated through the layers of the network. The network is considered “trained” when the errors for each of the outputs generated from the instances of the training data set are minimized.

[0215] The accuracy of a machine learning algorithm can be affected significantly by the quality of the data set used to train the algorithm. The training process can be computationally intensive and

may require a significant amount of time on a conventional general-purpose processor. Accordingly, parallel processing hardware is used to train many types of machine learning algorithms. This is particularly useful for optimizing the training of neural networks, as the computations performed in adjusting the coefficients in neural networks lend themselves naturally to parallel implementations. Specifically, many machine learning algorithms and software applications have been adapted to make use of the parallel processing hardware within general-purpose graphics processing devices.

[0216] FIG. 14 is a generalized diagram of a machine learning software stack **1400**. A machine learning application **1402** is any logic that can be configured to train a neural network using a training dataset or to use a trained deep neural network to implement machine intelligence. The machine learning application **1402** can include training and inference functionality for a neural network and/or specialized software that can be used to train a neural network before deployment. The machine learning application **1402** can implement any type of machine intelligence including but not limited to image recognition, mapping and localization, autonomous navigation, speech synthesis, medical imaging, or language translation. Example machine learning applications **1402** include, but are not limited to, voice-based virtual assistants, image or facial recognition algorithms, autonomous navigation, and the software tools that are used to train the machine learning models used by the machine learning applications **1402**.

[0217] Hardware acceleration for the machine learning application **1402** can be enabled via a machine learning framework **1404**. The machine learning framework **1404** can provide a library of machine learning primitives. Machine learning primitives are basic operations that are commonly performed by machine learning algorithms. Without the machine learning framework **1404**, developers of machine learning algorithms would be required to create and optimize the main computational logic associated with the machine learning algorithm, then re-optimize the computational logic as new parallel processors are developed. Instead, the machine learning application can be configured to perform the necessary computations using the primitives provided by the machine learning framework **1404**. Exemplary primitives include tensor convolutions, activation functions, and pooling, which are computational operations that are performed while training a convolutional neural network (CNN). The machine learning framework **1404** can also provide primitives to implement basic linear algebra subprograms performed by many machine-learning algorithms, such as matrix and vector operations. Examples of a machine learning framework **1404** include, but are not limited to, TensorFlow, TensorRT, PyTorch, MXNet, Caffe, and other high-level machine learning frameworks.

[0218] The machine learning framework **1404** can process input data received from the machine learning application **1402** and generate the appropriate input to a compute framework **1406**. The compute framework **1406** can abstract the underlying instructions provided to the GPGPU driver **1408** to enable the machine learning framework **1404** to take advantage of hardware acceleration via the GPGPU hardware **1410** without requiring the machine learning framework **1404** to have intimate knowledge of the architecture of the GPGPU hardware **1410**. Additionally, the compute framework **1406** can enable hardware acceleration for the machine learning framework **1404** across a variety of types and generations of the GPGPU hardware **1410**. Exemplary compute frameworks **1406** include the CUDA compute framework and associated machine learning libraries, such as the CUDA Deep Neural Network (cuDNN) library. The machine learning software stack **1400** can also include communication libraries or frameworks to facilitate multi-GPU and multi-node compute. GPGPU Machine Learning Acceleration

[0219] FIG. 15 illustrates a general-purpose graphics processing unit **1500**, which may include the architecture and features of any graphics processor described herein. The general-purpose processing unit (GPGPU) **1500** may be configured to provide support for hardware acceleration of primitives provided by a machine learning framework to accelerate the processing the type of computational workloads associated with training deep neural networks. Additionally, the GPGPU

1500 can be linked directly to other instances of the GPGPU to create a multi-GPU cluster to improve training speed for particularly deep neural networks. Primitives are also supported to accelerate inference operations for deployed neural networks.

[0220] The GPGPU **1500** includes a host interface **1502** to enable a connection with a host processor. The host interface **1502** may be a PCI Express interface. However, the host interface can also be a vendor specific communications interface or communications fabric. The GPGPU **1500** receives commands from the host processor and uses a global scheduler **1504** to distribute execution threads associated with those commands to a set of processing clusters **1506A-1506H**. The processing clusters **1506A-1506H** share a cache memory **1508**. The cache memory **1508** can serve as a higher-level cache for cache memories within the processing clusters **1506A-1506H**. The illustrated processing clusters **1506A-1506H** may each correspond with a graphics core cluster **414** as described herein.

[0221] The GPGPU **1500** includes memory **1514A-1514B** coupled with the processing clusters **1506A-1506H** via a set of memory controllers **1512A-1512B**. The memory **1514A-1514B** can include various types of memory devices including dynamic random-access memory (DRAM) or graphics random access memory, such as synchronous graphics random access memory (SGRAM), including graphics double data rate (GDDR) memory. The memory **1514A-1514B** may also include 3D stacked memory, including but not limited to high bandwidth memory (HBM).

[0222] Each of the processing clusters **1506A-1506H** may include multiple graphics multiprocessors, which can may include components illustrated with respect to the multi-core group **240A-240N** as in FIG. 2C, compute units **260A-260N** of FIG. 2D, or the graphics cores **221A-221F** of FIG. 2B. The graphics multiprocessors of the compute cluster include multiple types of integer and floating-point logic units that can perform computational operations at a range of precisions including suited for machine learning computations. For example, at least a subset of the floating-point units in each of the processing clusters **1506A-1506H** can be configured to perform 16-bit, 16-bit, and/or 32-bit floating point operations, while a different subset of the floating-point units can be configured to perform 144-bit floating point operations.

[0223] Multiple instances of the GPGPU **1500** can be configured to operate as a compute cluster. The communication mechanism used by the compute cluster for synchronization and data exchange varies across embodiments. For example, the multiple instances of the GPGPU **1500** communicate over the host interface **1502**. In one embodiment the GPGPU **1500** includes an I/O hub **1509** that couples the GPGPU **1500** with a GPU link **1510** that enables a direct connection to other instances of the GPGPU. The GPU link **1510** may be coupled to a dedicated GPU-to-GPU bridge that enables communication and synchronization between multiple instances of the GPGPU **1500**. Optionally, the GPU link **1510** couples with a high-speed interconnect to transmit and receive data to other GPGPUs or parallel processors. The multiple instances of the GPGPU **1500** may be located in separate data processing systems and communicate via a network device that is accessible via the host interface **1502**. The GPU link **1510** may be configured to enable a connection to a host processor in addition to or as an alternative to the host interface **1502**.

[0224] While the illustrated configuration of the GPGPU **1500** can be configured to train neural networks, an alternate configuration of the GPGPU **1500** can be configured for deployment within a high performance or low power inferencing platform. In an inferencing configuration, the GPGPU **1500** includes fewer of the processing clusters **1506A-1506H** relative to the training configuration. Additionally, memory technology associated with the memory **1514A-1514B** may differ between inferencing and training configurations. In one embodiment, the inferencing configuration of the GPGPU **1500** can support inferencing specific instructions. For example, an inferencing configuration can provide support for one or more 16-bit integer dot product instructions, which are commonly used during inferencing operations for deployed neural networks.

[0225] FIG. 16 illustrates a multi-GPU computing system **1600**. The multi-GPU computing system **1600** can include a processor **1602** coupled to multiple GPGPUs **1606A-1606D** via a host interface

switch **1604**. The host interface switch **1604** may be a PCI express switch device that couples the processor **1602** to a PCI express bus over which the processor **1602** can communicate with the set of GPGPUs **1606A-1606D**. Each of the multiple GPGPUs **1606A-1606D** can be an instance of the GPGPU **1500** of FIG. **15**. The GPGPUs **1606A-1606D** can interconnect via a set of high-speed point to point GPU to GPU links **1616**. The high-speed GPU to GPU links can connect to each of the GPGPUs **1606A-1606D** via a dedicated GPU link, such as the GPU link **1510** as in FIG. **15**. The P2P GPU links **1616** enable direct communication between each of the GPGPUs **1606A-1606D** without requiring communication over the host interface bus to which the processor **1602** is connected. With GPU-to-GPU traffic directed to the P2P GPU links, the host interface bus remains available for system memory access or to communicate with other instances of the multi-GPU computing system **1600**, for example, via one or more network devices. While in FIG. **16** the GPGPUs **1606A-1606D** connect to the processor **1602** via the host interface switch **1604**, the processor **1602** may alternatively include direct support for the P2P GPU links **1616** and connect directly to the GPGPUs **1606A-1606D**. In one embodiment the P2P GPU link **1616** enable the multi-GPU computing system **1600** to operate as a single logical GPU.

Machine Learning Neural Network Implementations

[0226] The computing architecture described herein can be configured to perform the types of parallel processing that is particularly suited for training and deploying neural networks for machine learning. A neural network can be generalized as a network of functions having a graph relationship. As is well-known in the art, there are a variety of types of neural network implementations used in machine learning. One exemplary type of neural network is the feedforward network, as previously described.

[0227] A second exemplary type of neural network is the Convolutional Neural Network (CNN). A CNN is a specialized feedforward neural network for processing data having a known, grid-like topology, such as image data. Accordingly, CNNs are commonly used for compute vision and image recognition applications, but they also may be used for other types of pattern recognition such as speech and language processing. The nodes in the CNN input layer are organized into a set of “filters” (feature detectors inspired by the receptive fields found in the retina), and the output of each set of filters is propagated to nodes in successive layers of the network. The computations for a CNN include applying the convolution mathematical operation to each filter to produce the output of that filter. Convolution is a specialized kind of mathematical operation performed by two functions to produce a third function that is a modified version of one of the two original functions. In convolutional network terminology, the first function to the convolution can be referred to as the input, while the second function can be referred to as the convolution kernel. The output may be referred to as the feature map. For example, the input to a convolution layer can be a multidimensional array of data that defines the various color components of an input image. The convolution kernel can be a multidimensional array of parameters, where the parameters are adapted by the training process for the neural network.

[0228] Recurrent neural networks (RNNs) are a family of feedforward neural networks that include feedback connections between layers. RNNs enable modeling of sequential data by sharing parameter data across different parts of the neural network. The architecture for an RNN includes cycles. The cycles represent the influence of a present value of a variable on its own value at a future time, as at least a portion of the output data from the RNN is used as feedback for processing subsequent input in a sequence. This feature makes RNNs particularly useful for language processing due to the variable nature in which language data can be composed.

[0229] The figures described below present exemplary feedforward, CNN, and RNN networks, as well as describe a general process for respectively training and deploying each of those types of networks. It will be understood that these descriptions are exemplary and non-limiting as to any specific embodiment described herein and the concepts illustrated can be applied generally to deep neural networks and machine learning techniques in general.

[0230] The exemplary neural networks described above can be used to perform deep learning. Deep learning is machine learning using deep neural networks. The deep neural networks used in deep learning are artificial neural networks composed of multiple hidden layers, as opposed to shallow neural networks that include only a single hidden layer. Deeper neural networks are generally more computationally intensive to train. However, the additional hidden layers of the network enable multistep pattern recognition that results in reduced output error relative to shallow machine learning techniques.

[0231] Deep neural networks used in deep learning typically include a front-end network to perform feature recognition coupled to a back-end network which represents a mathematical model that can perform operations (e.g., object classification, speech recognition, etc.) based on the feature representation provided to the model. Deep learning enables machine learning to be performed without requiring hand crafted feature engineering to be performed for the model. Instead, deep neural networks can learn features based on statistical structure or correlation within the input data. The learned features can be provided to a mathematical model that can map detected features to an output. The mathematical model used by the network is generally specialized for the specific task to be performed, and different models will be used to perform different task.

[0232] Once the neural network is structured, a learning model can be applied to the network to train the network to perform specific tasks. The learning model describes how to adjust the weights within the model to reduce the output error of the network. Backpropagation of errors is a common method used to train neural networks. An input vector is presented to the network for processing. The output of the network is compared to the desired output using a loss function and an error value is calculated for each of the neurons in the output layer. The error values are then propagated backwards until each neuron has an associated error value which roughly represents its contribution to the original output. The network can then learn from those errors using an algorithm, such as the stochastic gradient descent algorithm, to update the weights of the of the neural network.

[0233] FIG. 17A-17B illustrate an exemplary convolutional neural network. FIG. 17A illustrates various layers within a CNN. As shown in FIG. 17A, an exemplary CNN used to model image processing can receive input **1702** describing the red, green, and blue (RGB) components of an input image. The input **1702** can be processed by multiple convolutional layers (e.g., convolutional layer **1704**, convolutional layer **1706**). The output from the multiple convolutional layers may optionally be processed by a set of fully connected layers **1708**. Neurons in a fully connected layer have full connections to all activations in the previous layer, as previously described for a feedforward network. The output from the fully connected layers **1708** can be used to generate an output result from the network. The activations within the fully connected layers **1708** can be computed using matrix multiplication instead of convolution. Not all CNN implementations make use of fully connected layers **1708**. For example, in some implementations the convolutional layer **1706** can generate output for the CNN.

[0234] The convolutional layers are sparsely connected, which differs from traditional neural network configuration found in the fully connected layers **1708**. Traditional neural network layers are fully connected, such that every output unit interacts with every input unit. However, the convolutional layers are sparsely connected because the output of the convolution of a field is input (instead of the respective state value of each of the nodes in the field) to the nodes of the subsequent layer, as illustrated. The kernels associated with the convolutional layers perform convolution operations, the output of which is sent to the next layer. The dimensionality reduction performed within the convolutional layers is one aspect that enables the CNN to scale to process large images.

[0235] FIG. 17B illustrates exemplary computation stages within a convolutional layer of a CNN. Input to a convolutional layer **1712** of a CNN can be processed in three stages of a convolutional layer **1714**. The three stages can include a convolution stage **1716**, a detector stage **1718**, and a pooling stage **1720**. The convolutional layer **1714** can then output data to a successive

convolutional layer. The final convolutional layer of the network can generate output feature map data or provide input to a fully connected layer, for example, to generate a classification value for the input to the CNN.

[0236] In the convolution stage **1716** performs several convolutions in parallel to produce a set of linear activations. The convolution stage **1716** can include an affine transformation, which is any transformation that can be specified as a linear transformation plus a translation. Affine transformations include rotations, translations, scaling, and combinations of these transformations. The convolution stage computes the output of functions (e.g., neurons) that are connected to specific regions in the input, which can be determined as the local region associated with the neuron. The neurons compute a dot product between the weights of the neurons and the region in the local input to which the neurons are connected. The output from the convolution stage **1716** defines a set of linear activations that are processed by successive stages of the convolutional layer **1714**.

[0237] The linear activations can be processed by a detector stage **1718**. In the detector stage **1718**, each linear activation is processed by a non-linear activation function. The non-linear activation function increases the nonlinear properties of the overall network without affecting the receptive fields of the convolution layer. Several types of non-linear activation functions may be used. One particular type is the rectified linear unit (ReLU), which uses an activation function defined as $f(x)=\max(0, x)$, such that the activation is thresholded at zero.

[0238] The pooling stage **1720** uses a pooling function that replaces the output of the convolutional layer **1706** with a summary statistic of the nearby outputs. The pooling function can be used to introduce translation invariance into the neural network, such that small translations to the input do not change the pooled outputs. Invariance to local translation can be useful in scenarios where the presence of a feature in the input data is more important than the precise location of the feature. Various types of pooling functions can be used during the pooling stage **1720**, including max pooling, average pooling, and 12-norm pooling. Additionally, some CNN implementations do not include a pooling stage. Instead, such implementations substitute an additional convolution stage having an increased stride relative to previous convolution stages.

[0239] The output from the convolutional layer **1714** can then be processed by the next layer **1722**. The next layer **1722** can be an additional convolutional layer or one of the fully connected layers **1708**. For example, the first convolutional layer **1704** of FIG. 17A can output to the second convolutional layer **1706**, while the second convolutional layer can output to a first layer of the fully connected layers **1708**.

[0240] FIG. 18 illustrates an exemplary recurrent neural network **1800**. In a recurrent neural network (RNN), the previous state of the network influences the output of the current state of the network. RNNs can be built in a variety of ways using a variety of functions. The use of RNNs generally revolves around using mathematical models to predict the future based on a prior sequence of inputs. For example, an RNN may be used to perform statistical language modeling to predict an upcoming word given a previous sequence of words. The illustrated RNN **1800** can be described as having an input layer **1802** that receives an input vector, hidden layers **1804** to implement a recurrent function, a feedback mechanism **1805** to enable a 'memory' of previous states, and an output layer **1806** to output a result. The RNN **1800** operates based on time-steps. The state of the RNN at a given time step is influenced based on the previous time step via the feedback mechanism **1805**. For a given time step, the state of the hidden layers **1804** is defined by the previous state and the input at the current time step. An initial input ($x_{sub.1}$) at a first-time step can be processed by the hidden layer **1804**. A second input ($x_{sub.2}$) can be processed by the hidden layer **1804** using state information that is determined during the processing of the initial input ($x_{sub.1}$). A given state can be computed as $s_{sub.t}=f(Ux_{sub.t}+Ws_{sub.t-1})$, where U and W are parameter matrices. The function f is generally a nonlinearity, such as the hyperbolic tangent function (\tanh) or a variant of the rectifier function $f(x)=\max(0,x)$. However, the specific

mathematical function used in the hidden layers **1804** can vary depending on the specific implementation details of the RNN **1800**.

[0241] In addition to the basic CNN and RNN networks described, acceleration for variations on those networks may be enabled. One example RNN variant is the long short term memory (LSTM) RNN. LSTM RNNs are capable of learning long-term dependencies that may be necessary for processing longer sequences of language. A variant on the CNN is a convolutional deep belief network, which has a structure similar to a CNN and is trained in a manner similar to a deep belief network. A deep belief network (DBN) is a generative neural network that is composed of multiple layers of stochastic (random) variables. DBNs can be trained layer-by-layer using greedy unsupervised learning. The learned weights of the DBN can then be used to provide pre-train neural networks by determining an optimal initial set of weights for the neural network. In further embodiments, acceleration for reinforcement learning is enabled. In reinforcement learning, an artificial agent learns by interacting with its environment. The agent is configured to optimize certain objectives to maximize cumulative rewards.

[0242] FIG. **19** illustrates training and deployment of a deep neural network. Once a given network has been structured for a task the neural network is trained using a training dataset **1902**. Various training frameworks **1904** have been developed to enable hardware acceleration of the training process. For example, the machine learning framework **1404** of FIG. **14** may be configured as a training framework **1904**. The training framework **1904** can hook into an untrained neural network **1906** and enable the untrained neural net to be trained using the parallel processing resources described herein to generate a trained neural network **1908**.

[0243] To start the training process the initial weights may be chosen randomly or by pre-training using a deep belief network. The training cycle then be performed in either a supervised or unsupervised manner.

[0244] Supervised learning is a learning method in which training is performed as a mediated operation, such as when the training dataset **1902** includes input paired with the desired output for the input, or where the training dataset includes input having known output and the output of the neural network is manually graded. The network processes the inputs and compares the resulting outputs against a set of expected or desired outputs. Errors are then propagated back through the system. The training framework **1904** can adjust to adjust the weights that control the untrained neural network **1906**. The training framework **1904** can provide tools to monitor how well the untrained neural network **1906** is converging towards a model suitable to generating correct answers based on known input data. The training process occurs repeatedly as the weights of the network are adjusted to refine the output generated by the neural network. The training process can continue until the neural network reaches a statistically desired accuracy associated with a trained neural network **1908**. The trained neural network **1908** can then be deployed to implement any number of machine learning operations to generate an inference result **1914** based on input of new data **1912**.

[0245] Unsupervised learning is a learning method in which the network attempts to train itself using unlabeled data. Thus, for unsupervised learning the training dataset **1902** will include input data without any associated output data. The untrained neural network **1906** can learn groupings within the unlabeled input and can determine how individual inputs are related to the overall dataset. Unsupervised training can be used to generate a self-organizing map, which is a type of trained neural network **1908** capable of performing operations useful in reducing the dimensionality of data. Unsupervised training can also be used to perform anomaly detection, which allows the identification of data points in an input dataset that deviate from the normal patterns of the data.

[0246] Variations on supervised and unsupervised training may also be employed. Semi-supervised learning is a technique in which in the training dataset **1902** includes a mix of labeled and unlabeled data of the same distribution. Incremental learning is a variant of supervised learning in which input data is continuously used to further train the model. Incremental learning enables the trained neural

network **1908** to adapt to the new data **1912** without forgetting the knowledge instilled within the network during initial training.

[0247] Whether supervised or unsupervised, the training process for particularly deep neural networks may be too computationally intensive for a single compute node. Instead of using a single compute node, a distributed network of computational nodes can be used to accelerate the training process.

[0248] FIG. **20A** is a block diagram illustrating distributed learning. Distributed learning is a training model that uses multiple distributed computing nodes to perform supervised or unsupervised training of a neural network. The distributed computational nodes can each include one or more host processors and one or more of the general-purpose processing nodes, such as the highly parallel general-purpose graphics processing unit **1500** as in FIG. **15**. As illustrated, distributed learning can be performed with model parallelism **2002**, data parallelism **2004**, or a combination of model and data parallelism **2006**.

[0249] In model parallelism **2002**, different computational nodes in a distributed system can perform training computations for different parts of a single network. For example, each layer of a neural network can be trained by a different processing node of the distributed system. The benefits of model parallelism include the ability to scale to particularly large models. Splitting the computations associated with different layers of the neural network enables the training of very large neural networks in which the weights of all layers would not fit into the memory of a single computational node. In some instances, model parallelism can be particularly useful in performing unsupervised training of large neural networks.

[0250] In data parallelism **2004**, the different nodes of the distributed network have a complete instance of the model and each node receives a different portion of the data. The results from the different nodes are then combined. While different approaches to data parallelism are possible, data parallel training approaches all require a technique of combining results and synchronizing the model parameters between each node. Exemplary approaches to combining data include parameter averaging and update-based data parallelism. Parameter averaging trains each node on a subset of the training data and sets the global parameters (e.g., weights, biases) to the average of the parameters from each node. Parameter averaging uses a central parameter server that maintains the parameter data. Update based data parallelism is similar to parameter averaging except that instead of transferring parameters from the nodes to the parameter server, the updates to the model are transferred. Additionally, update-based data parallelism can be performed in a decentralized manner, where the updates are compressed and transferred between nodes.

[0251] Combined model and data parallelism **2006** can be implemented, for example, in a distributed system in which each computational node includes multiple GPUs. Each node can have a complete instance of the model with separate GPUs within each node are used to train different portions of the model.

[0252] Distributed training has increased overhead relative to training on a single machine. However, the parallel processors and GPGPUs described herein can each implement various techniques to reduce the overhead of distributed training, including techniques to enable high bandwidth GPU-to-GPU data transfer and accelerated remote data synchronization.

[0253] FIG. **20B** is a block diagram illustrating a programmable network interface **2010** and data processing unit. The programmable network interface **2010** is a programmable network engine that can be used to accelerate network-based compute tasks within a distributed environment. The programmable network interface **2010** can couple with a host system via host interface **2070**. The programmable network interface **2010** can be used to accelerate network or storage operations for CPUs or GPUs of the host system. The host system can be, for example, a node of a distributed learning system used to perform distributed training, for example, as shown in FIG. **20A**. The host system can also be a data center node within a data center.

[0254] In one embodiment, access to remote storage containing model data can be accelerated by

the programmable network interface **2010**. For example, the programmable network interface **2010** can be configured to present remote storage devices as local storage devices to the host system. The programmable network interface **2010** can also accelerate remote direct memory access (RDMA) operations performed between GPUs of the host system with GPUs of remote systems. In one embodiment, the programmable network interface **2010** can enable storage functionality such as, but not limited to NVME-oF. The programmable network interface **2010** can also accelerate encryption, data integrity, compression, and other operations for remote storage on behalf of the host system, allowing remote storage to approach the latencies of storage devices that are directly attached to the host system.

[0255] The programmable network interface **2010** can also perform resource allocation and management on behalf of the host system. Storage security operations can be offloaded to the programmable network interface **2010** and performed in concert with the allocation and management of remote storage resources. Network-based operations to manage access to the remote storage that would otherwise be performed by a processor of the host system can instead be performed by the programmable network interface **2010**.

[0256] In one embodiment, network and/or data security operations can be offloaded from the host system to the programmable network interface **2010**. Data center security policies for a data center node can be handled by the programmable network interface **2010** instead of the processors of the host system. For example, the programmable network interface **2010** can detect and mitigate against an attempted network-based attack (e.g., DDoS) on the host system, preventing the attack from compromising the availability of the host system.

[0257] The programmable network interface **2010** can include a system on a chip (SoC **2020**) that executes an operating system via multiple processor cores **2022**. The processor cores **2022** can include general-purpose processor (e.g., CPU) cores. In one embodiment the processor cores **2022** can also include one or more GPU cores. The SoC **2020** can execute instructions stored in a memory device **2040**. A storage device **2050** can store local operating system data. The storage device **2050** and memory device **2040** can also be used to cache remote data for the host system. Network ports **2060A-2060B** enable a connection to a network or fabric and facilitate network access for the SoC **2020** and, via the host interface **2070**, for the host system. The programmable network interface **2010** can also include an I/O interface **2075**, such as a USB interface. The I/O interface **2075** can be used to couple external devices to the programmable network interface **2010** or as a debug interface. The programmable network interface **2010** also includes a management interface **2030** that enables software on the host device to manage and configure the programmable network interface **2010** and/or SoC **2020**. In one embodiment the programmable network interface **2010** may also include one or more accelerators or GPUs **2045** to accept offload of parallel compute tasks from the SoC **2020**, host system, or remote systems coupled via the network ports **2060A-2060B**.

Exemplary Machine Learning Applications

[0258] Machine learning can be applied to solve a variety of technological problems, including but not limited to computer vision, autonomous driving and navigation, speech recognition, and language processing. Computer vision has traditionally been one of the most active research areas for machine learning applications. Applications of computer vision range from reproducing human visual abilities, such as recognizing faces, to creating new categories of visual abilities. For example, computer vision applications can be configured to recognize sound waves from the vibrations induced in objects visible in a video. Parallel processor accelerated machine learning enables computer vision applications to be trained using significantly larger training dataset than previously feasible and enables inferencing systems to be deployed using low power parallel processors.

[0259] Parallel processor accelerated machine learning has autonomous driving applications including lane and road sign recognition, obstacle avoidance, navigation, and driving control.

Accelerated machine learning techniques can be used to train driving models based on datasets that define the appropriate responses to specific training input. The parallel processors described herein can enable rapid training of the increasingly complex neural networks used for autonomous driving solutions and enables the deployment of low power inferencing processors in a mobile platform suitable for integration into autonomous vehicles.

[0260] Parallel processor accelerated deep neural networks have enabled machine learning approaches to automatic speech recognition (ASR). ASR includes the creation of a function that computes the most probable linguistic sequence given an input acoustic sequence. Accelerated machine learning using deep neural networks have enabled the replacement of the hidden Markov models (HMMs) and Gaussian mixture models (GMMs) previously used for ASR.

[0261] Parallel processor accelerated machine learning can also be used to accelerate natural language processing. Automatic learning procedures can make use of statistical inference algorithms to produce models that are robust to erroneous or unfamiliar input. Exemplary natural language processor applications include automatic machine translation between human languages.

[0262] The parallel processing platforms used for machine learning can be divided into training platforms and deployment platforms. Training platforms are generally highly parallel and include optimizations to accelerate multi-GPU single node training and multi-node, multi-GPU training. Exemplary parallel processors suited for training include the general-purpose graphics processing unit **1500** of FIG. **15** and the multi-GPU computing system **1600** of FIG. **16**. On the contrary, deployed machine learning platforms generally include lower power parallel processors suitable for use in products such as cameras, autonomous robots, and autonomous vehicles.

[0263] Additionally, machine learning techniques can be applied to accelerate or enhance graphics processing activities. For example, a machine learning model can be trained to recognize output generated by a GPU accelerated application and generate an upscaled version of that output. Such techniques can be applied to accelerate the generation of high-resolution images for a gaming application. Various other graphics pipeline activities can benefit from the use of machine learning. For example, machine learning models can be trained to perform tessellation operations on geometry data to increase the complexity of geometric models, allowing fine-detailed geometry to be automatically generated from geometry of relatively lower detail.

[0264] FIG. **21** illustrates an exemplary inferencing system on a chip (SOC) **2100** suitable for performing inferencing using a trained model. The SOC **2100** can integrate processing components including a media processor **2102**, a vision processor **2104**, a GPGPU **2106** and a multi-core processor **2108**. The GPGPU **2106** may be a GPGPU as described herein, such as the GPGPU **1500**, and the multi-core processor **2108** may be a multi-core processor described herein. The SOC **2100** can additionally include on-chip memory **2105** that can enable a shared on-chip data pool that is accessible by each of the processing components. The processing components can be optimized for low power operation to enable deployment to a variety of machine learning platforms, including autonomous vehicles and autonomous robots. For example, one implementation of the SOC **2100** can be used as a portion of the main control system for an autonomous vehicle. Where the SOC **2100** is configured for use in autonomous vehicles the SOC is designed and configured for compliance with the relevant functional safety standards of the deployment jurisdiction.

[0265] During operation, the media processor **2102** and vision processor **2104** can work in concert to accelerate computer vision operations. The media processor **2102** can enable low latency decode of multiple high-resolution (e.g., **4K**, **16K**) video streams. The decoded video streams can be written to a buffer in the on-chip memory **2105**. The vision processor **2104** can then parse the decoded video and perform preliminary processing operations on the frames of the decoded video in preparation of processing the frames using a trained image recognition model. For example, the vision processor **2104** can accelerate convolution operations for a CNN that is used to perform image recognition on the high-resolution video data, while back-end model computations are performed by the GPGPU **2106**.

[0266] The multi-core processor **2108** can include control logic to assist with sequencing and synchronization of data transfers and shared memory operations performed by the media processor **2102** and the vision processor **2104**. The multi-core processor **2108** can also function as an application processor to execute software applications that can make use of the inferencing compute capability of the GPGPU **2106**. For example, at least a portion of the navigation and driving logic can be implemented in software executing on the multi-core processor **2108**. Such software can directly issue computational workloads to the GPGPU **2106**, or the computational workloads can be issued to the multi-core processor **2108**, which can offload at least a portion of those operations to the GPGPU **2106**.

[0267] The GPGPU **2106** can include compute clusters such as a low power configuration of the processing clusters **1506A-1506H** within general-purpose graphics processing unit **1500**. The compute clusters within the GPGPU **2106** can support instruction that are specifically optimized to perform inferencing computations on a trained neural network. For example, the GPGPU **2106** can support instructions to perform low precision computations such as 16-bit and 4-bit integer and 16-bit and 16-bit floating-point operations.

Learning Neural Reflectance Shaders from Images

[0268] Described herein is a system to reconstruct the shape and reflectance of objects from images. The shape and reflectance of objects can be reconstructed using real-time rendering techniques via a pair of multilayer perceptron (MLP) networks that have been trained on a large dataset of images is used to represent illumination and surface reflectance. The training allows the efficient estimation of latent codes for illumination and reflectance for material surfaces through the networks with their weights fixed. The MLP for the surface reflectance is lightweight and can be converted to a shading program suitable for real-time applications. The use of this learned reflectance model can improve photorealism for computer generated scenes. Tools for creating photorealistic 3D assets are crucial in the graphics domain and for creating applications in virtual worlds. Furthermore, data-driven shaders in real-time graphics have the potential to influence the graphics pipeline and its compute patterns.

[0269] FIG. **22** illustrates input and output of a system to reconstruct the shape and reflectance of objects from images. Using the techniques described herein, images of a real-world object **2202** can be used to generate data that can be used to render a virtual image **2204** of that real-world object. The images of the real-world object **2202** can be provided as input to the system. The system will then output mesh and reflectance latent codes that can be used with a corresponding neural shader for real-time rendering. The mesh and latent codes can be used as input to the neural shader to render the virtual image **2204**.

[0270] FIG. **23** illustrates an overview of the processing pipeline **2300** for the system described herein. In the first stage, input images **2302** are processed via structure from motion **2304** to generate 3D meshes and camera poses **2306**. Structure from motion is a photogrammetric technique for estimating three-dimensional structures from a sequence of two-dimensional images. The motion represented by the sequence of images allows the application of techniques that enables the estimation of the structure of the objects depicted in the images. For the purpose of the techniques described herein, reconstruction via structure from motion **2304** can be performed using one or more off-the-shelf methods that enable the generation of the 3D mesh and camera poses **2306**. For example, an additional machine learning model can be configured to generate the 3D mesh and camera poses **2306** based on the input images **2302**.

[0271] In the second stage, optimization network models **3408** is applied to generate optimized estimation of latent codes **2310** for reflectance and illumination. In the last stage, the estimated latent codes **2310** are passed to a renderer **2312**, which can generate novel images **2314** using a multilayer perceptron (MLP) as a shader. In one embodiment, the system consists of two networks, Reflectance Net and Illumination Net, which can be used to optimize the latent codes for reflectance and illumination as well as rendering an image in a differentiable manner.

[0272] FIG. 24 illustrates an overview of the Reflectance Net 2412 and the Illumination Net 2414, according to an embodiment. The Reflectance Net 2412 implements a version of a bidirectional reflectance distribution function (BRDF). The BRDF is a function that defines how light is reflected at an opaque surface, as shown in FIG. 25. The Illumination Net 2414 implements the incoming light similar to image-based lighting (IBL) with high dynamic range images (HDRIs). IBL is a 3D rendering technique which includes capturing an omnidirectional representation of real-world light information as an image. This image is then used to simulate the lighting for the objects in the scene. In one embodiment, input data 2402 that includes input geometry and camera pose data can be processed to generate reflectance codes $z_{sub.i}$ 2404 and an illumination code w 2410. The illumination code w 2410 is used per scene, while the reflectance codes $z_{sub.i}$ 2404 are computed for each mesh vertex or for each texel if texture coordinates are available for the mesh. In addition to the latent codes, geometric information is passed to the networks.

[0273] With reference to FIG. 25, the geometric information includes the view angle for the camera 2504 relative to the normal 2506 of a surface 2502, which is represented as the incident angle q 2406 of a reflection, and the reflection vector $\omega_{sub.r}$ 2408 of the resulting reflection. The geometric information is passed to the corresponding network. The incident angle ϕ 2406 may be determined based on the camera pose data provided with the input data 2402.

[0274] Returning to FIG. 24, dot product operations 2416 are used to combine the output of the Reflectance Net 2412 and the Illumination Net 2414 during a forward shading 2420 pass to generate parameters to provide to the MLP Shader. The MLP shader is used to generate a rendering 2422. The difference between the rendering 2422 and the observation 2424 is computed for use during a backward optimization pass 2421 for both the Reflectance Net 2412 and the Illumination Net 2414 to generate optimized reflectance codes 2404 and an optimized illumination code 2410.

[0275] The Reflectance Net 2412 outputs 12 channels organized as a 4-tuple of RGB values, which are combined with the outputs of the Illumination Net via dot product operations 2416 to generate a single output for each color channel. The output of the Illumination Net is a preintegrated light intensity value based on 4 predefined specular lobe parameters $r_{sub.0,1,2,3}$, which define the level of preintegration. This preintegration scheme follows the split-sum approximation used in modern real-time rendering pipelines like the Unreal Engine. The specular lobe parameter is related to the surface roughness parameter in physically-based rendering (PBR) materials. FIG. 26 shows example output 2600 for a given latent code w of the Illumination Net 2414. Output of the Illumination Net 2414 includes equirectangular images. Four exemplary images are shown using four different predefined specular lobe parameters $r_{sub.0}=0$, $r_{sub.1}=0.3$, $r_{sub.2}=1$, $r_{sub.3}=5$.

Latent Code Estimation

[0276] Reflectance and illumination latent codes for an object/material are estimated by backpropagation through the fixed and trained networks. Given a set of points $x_{sub.i}$ on the mesh surface and a set of images $I_{sub.k}$ with camera poses $\xi_{sub.k}$, the system can optimize the objective:

$$[00001] \min_{\{z_i\}, w} \mathcal{L}_{recon} + \mathcal{L}_{reg},$$

[0277] The reconstruction loss is defined as:

$$[00002] \mathcal{L}_{recon} = \sum_i \sum_k \delta_{i,k} \| \pi(x_i; \xi_k) - L_o(\omega_{o,i,k}, z_i, w, n_i) \|^2,$$

with $\delta_{sub.i,k}$ as a visibility test indicating if point $x_{sub.i}$ is visible in image $I_{sub.k}$; π as a projection from 3D to the image space; $L_{sub.o}$ as the combined output of the networks; $\omega_{sub.o,i,k}$ as the ray from point $x_{sub.i}$ to the camera of image $I_{sub.k}$; and $n_{sub.i}$ as the surface normal.

[0278] The regularization loss is defined as:

[00003]

$$\mathcal{L}_{reg} = \sum_i \| z_i \|^2 + \lambda \sum_l \| w \|^2 + \sum_l \sum_{i \in l} \| \hat{z}_l - z_i \|^2,$$

with $\{\text{circumflex over } (z)\}_{sub.l}$ as the average latent code of cluster l . A k-means clustering of

surface points is performed with respect to the darkest observed color in all images and use the average latent code for each cluster as regularization. This regularization stabilizes the optimization for points with only few observations. The optimal number of clusters depends on the scene, with one exemplary implementation using ten clusters.

[0279] In one embodiment, the system uses an Adam optimizer to compute the latent codes. The system can alternate between optimizing reflectance and illumination codes every 500 iterations for the first 1750 iterations and switch to joint optimization in the last 250 iterations, totaling 2000 iterations.

Rendering

[0280] In one embodiment, the system is configured to render images by evaluating both networks for each pixel of the output network. Further, the system can render the reconstructed objects with new illumination by replacing the illumination code or constructing preintegrated environment map images as shown in FIG. 26. For integration into real-time graphics pipelines, the system can translate the Reflectance Net 2412 and its parameters into an HLSL shader, as shown in FIG. 27.

[0281] FIG. 27 illustrates translation of Reflectance Net 2412 to an HLSL shader 2704 for integration into a 3D engine. In one embodiment, a Reflectance Net Implementation 2702 can be translated or converted into an HLSL shader 2704 by converting the layers of the Reflectance Net Implementation 2702 into shader operations in a high level shader language. In one embodiment, the Reflectance Net 2412 is implemented as a PyTorch MLP, although embodiments are not limited to any particular implementation. The system described herein can translate the linear layers of the MLP to a sequence of block wise matrix multiplications using data types such as, for example, a float 4×4 format, which can be used to store a 4×4 floating-point matrix of various precisions.

[0282] Estimated reflectance codes $z_{sub.i}$ 2404 consist of 15 channels, which are quantized and converted to four red, green, blue, alpha (RGBA) textures with eight bits per channel. The latent code textures can be read with the texture sampler of a GPU and are dequantized inside the generated HLSL shader 2704. In one embodiment, the default texture sampler of the GPU can be used.

[0283] FIG. 28 illustrates exemplary reference images 2810, 2820, 2830 and rendered images 2812, 2822, 2832 that include reconstructed objects within the reference images. The reconstructed objects are rendered using the HLSL shader 2704. In one embodiment, the system can reconstruct the reflectance by backpropagation and images rendered inside a 3D engine such as, for example, Unreal Engine, or another engine in which the Reflectance Net Implementation 2702 can be executed as an HLSL shader 2704 to perform dynamic illumination.

Reflectance Net Details

[0284] FIG. 29 shows the architecture 2900 of the Reflectance Net 2412, according to an embodiment. In one embodiment, the reflectance code 2404 has 15 elements, which are split into three parts having seven, four, and four elements respectively. The three parts are used as input to different layers of the MLP of the Reflectance Net 2412. The input to the first linear layer 2906 is a vector with eight elements of which seven are from the reflectance code 2902 and one element is the incident angle 2406. The first linear layer 2906 generates 16 elements, which are processed via a linear RELU activation function and provided to the second layer 2908. The second layer 2908 processes the 16 elements and four elements of the reflectance code 2902 and applies a linear RELU to generate 20 elements that are input to the third layer 2910. The third layer 2910 processes the 20 elements and an additional four elements of the reflectance code 2902. Output from the third layer 2910 is processed via a linear RELU to generate 24 elements that are processed by the fourth layer 2912. The fourth layer 2912 uses SoftPlus as activation instead of the ReLU to improve gradient flow. In one embodiment, the SoftPlus activation is implemented as $\log(1+e^{\sup{x}})$. In one embodiment, the SoftPlus activation is implemented as $\max(0,x)+\log(1+e^{\sup{-|x|}})$.

[0285] The output 2914 is a 12 element vector which is interpreted as a 4-tuple of RGB values. The Reflectance Net 2412 is designed to be compact to avoid overfitting and facilitate translation to

shader programming languages. The Reflectance Net **2412** is trained using an image dataset of scenes or objects. During inferencing for a scene or object, 100 images may be used, in one embodiment, to provide good coverage of the objects surface to generate sufficient data to provide an optimal data set to be used during rendering.

Illumination Net Details

[0286] FIG. **30** shows the architecture **3000** of the Illumination Net **2414**, according to an embodiment. The Illumination Net **2414** is used during optimization of the reflectance codes to disentangle lighting and reflectance. In one embodiment, the Illumination Net **2414** includes seven layers **3008A-3008G**, each with 256 elements, and generates a three element output **3010**. The illumination code **2410** is provided as input to the Illumination Net **2414** to define the frequency and phase shift used in the linear Sine activation function associated with the second layer **3008B**, fourth layer **3008D**, and sixth layer **3008F**. The first layer **3008A**, third layer **3008C**, and fifth layer **3000E** are associated with a linear RELU activation function. The seventh layer **3008G** applies a linear SoftPlus activation function. The reflection vector **2408** is provided as input to the Illumination Net **2414**, as well as a specular lobe parameter r **3006**. The integral position encoding steers the level of preintegration, as shown in FIG. **26**. The specular lobe parameter r **3006** acts as the standard deviation of normal distribution. The preintegration is defined as:

$$[00004] \quad (r, r) = \begin{bmatrix} \sin(P \cdot r) \circ \exp(-\frac{1}{2}r^2 \text{diag}(PP^T)) \\ \cos(P \cdot r) \circ \exp(-\frac{1}{2}r^2 \text{diag}(PP^T)) \end{bmatrix},$$

where \circ as the Hadamard product and matrix $P \in \text{custom-character.sup.128} \times 3$ with its entries drawn from a standard normal distribution.

[0287] A dataset of HDRIs is used for training and fixing of the network weights. In one embodiment, system can use approximately 280 HDRIs to train the network weights, although other embodiments may use a dataset including a different number of images. Optimizing the latent code can performed using only images for the reconstructed object or scene and other images are not required to be used.

[0288] The above details with respect to the Reflectance Net **2412** and the Illumination Net **2414** are exemplary of one embodiment and are not limiting as to all embodiments. Other embodiments are configurable for use with different implementations of the Reflectance Net **2412** and the Illumination Net **2414**;

[0289] FIG. **31** illustrates a method **3100** to learn neural reflectance shaders from images, according to an embodiment. The method **3100** can be implemented via compute resources of a graphics processor or parallel processor described herein using the Reflectance Net **2412** and the Illumination Net **2414** described above, which are implemented as MLPs. The compute resources can be configured to process, via one or more machine learning models, a set of input images (**3102**). In one embodiment, the one or more machine learning models can include a single model that combines the reflectance and illumination networks. In one embodiment, the one or more machine learning models include separate reflectance networks and illumination networks. The set of input images can include a plurality of input images.

[0290] The compute resources can then generate, via the one or more machine learning models, an illumination latent code and a set of reflectance latent codes for an object within the set of input images (**3104**). For each scene or object within the scene, the set of reflectance latent codes can include a latent code for each vertex of the mesh. If the mesh has associated texture coordinates, the set of reflectance latent codes can include a latent code for each texel associated with the mesh. A single illumination code can be used for each scene or object within the scene.

[0291] The compute resources can then generate a shader based on a machine learning model of the one or more machine learning models, where the shader is configured to sample the illumination latent code and the set of reflectance latent codes for the object (**3106**). In one embodiment, the shader is generated based on the reflectance network. The shader can be generated by converting a

representation of the MLP network into a shader or general-purpose compute program in a high-level shading language. The compute resources can then render a three-dimensional (3D) object by shading a mesh via the shader (**3108**). In one embodiment, the 3D object is rendered via a real-time renderer and the mesh is shaded in an interactive environment. The mesh can be a 3D mesh that is generated based on the set of input images. The 3D mesh can be generated along with data associated camera poses associated with the set of input images. In one embodiment, the compute resources execute instructions to generate the 3D mesh and the camera poses via structure from motion analysis of the set of input images. In one embodiment, an apparatus or data processing system is provided that includes a graphics processor that is configured to implement method **3100**. Additional Exemplary Computing Device

[0292] FIG. **32** is a block diagram of a computing device **3200** including a graphics processor **3204**, according to an embodiment. Versions of the computing device **3200** may be or be included within a communication device such as a set-top box (e.g., Internet-based cable television set-top boxes, etc.), global positioning system (GPS)-based devices, etc. The computing device **3200** may also be or be included within mobile computing devices such as cellular phones, smartphones, personal digital assistants (PDAs), tablet computers, laptop computers, e-readers, smart televisions, television platforms, wearable devices (e.g., glasses, watches, bracelets, smartcards, jewelry, clothing items, etc.), media players, etc. For example, in one embodiment, the computing device **3200** includes a mobile computing device employing an integrated circuit (“IC”), such as system on a chip (“SoC” or “SOC”), integrating various hardware and/or software components of computing device **3200** on a single chip. The computing device **3200** can be a computing device such as the processing system **100** as in FIG. **1** and can include components to implement functionality provided by the various embodiments described herein.

[0293] The computing device **3200** includes a graphics processor **3204**. The graphics processor **3204** represents any graphics processor described herein. In one embodiment, the graphics processor **3204** includes a cache **3214**, which can be a single cache or divided into multiple segments of cache memory, including but not limited to any number of L1, L2, L3, or L4 caches, render caches, depth caches, sampler caches, and/or shader unit caches. In one embodiment the cache **3214** may be a last level cache that is shared with the application processor **3206**.

[0294] In one embodiment the graphics processor **3204** includes a graphics microcontroller **3215** that implements control and scheduling logic for the graphics processor. The graphics microcontroller **3215** may be, for example, any of the graphics microcontrollers **3802A-3802B**, **4102A-4102B** described herein. The control and scheduling logic can be firmware executed by the graphics microcontroller **3215**. The firmware may be loaded at boot by the graphics driver logic **3222**. The firmware may also be programmed to an electronically erasable programmable read only memory or loaded from a flash memory device within the graphics microcontroller **3215**. The firmware may enable a GPU OS **3216** that includes device management logic **3217**, device driver logic **3218**, and a scheduler **3219**. The GPU OS **3216** may also include a graphics memory manager **3220** that can supplement or replace the graphics memory manager **3221** within the graphics driver logic **3222**, and generally enables the offload of various graphics driver functionality from the graphics driver logic **3222** to the GPU OS **3216**.

[0295] The virtual memory address management for compression data described herein can be implemented, in various embodiments, by the graphics memory manager **3220** of the GPU OS **3216**, the graphics memory manager **3221** within the graphics driver logic **3222**, or another component of the GPU OS **3216** and/or graphics driver logic **3222**.

[0296] The graphics processor **3204** also includes a GPGPU engine **3244** that includes one or more graphics engine(s), graphics processor cores, and other graphics execution resources as described herein. Such graphics execution resources can be presented in the forms including but not limited to execution units, shader engines, fragment processors, vertex processors, streaming multiprocessors, graphics processor clusters, or any collection of computing resources suitable for the processing of

graphics resources or image resources or performing general purpose computational operations in a heterogeneous processor. The processing resources of the GPGPU engine **3244** can be included within multiple tiles of hardware logic connected to a substrate, as illustrated in FIG. **11B-11D**. The GPGPU engine **3244** can include GPU tiles **3245** that include graphics processing and execution resources, caches, samplers, etc. The GPU tiles **3245** may also include local volatile memory or can be coupled with one or more memory tiles, for example, as shown in FIG. **3B-3C**.

[0297] The GPGPU engine **3244** can also include and one or more special tiles **3246** that include, for example, a non-volatile memory tile **3256**, a network processor tile **3257**, and/or a general-purpose compute tile **3258**. The GPGPU engine **3244** also includes a matrix multiply accelerator **3260**. The general-purpose compute tile **3258** may also include logic to accelerate matrix multiplication operations. The non-volatile memory tile **3256** can include non-volatile memory cells and controller logic. The controller logic of the non-volatile memory tile **3256** may be managed by the device management logic **3217** or the device driver logic **3218**. The network processor tile **3257** can include network processing resources that are coupled to a physical interface within the input/output (I/O) sources **3210** of the computing device **3200**. The network processor tile **3257** may be managed by one or more of device management logic **3217** or the device driver logic **3218**. Any of the GPU tiles **3245** or one or more special tiles **3246** may include an active base with multiple stacked chiplets, as described herein.

[0298] The matrix multiply accelerator **3260** is a modular scalable sparse matrix multiply accelerator. The matrix multiply accelerator **3260** can includes multiple processing paths, with each processing path including multiple pipeline stages. Each processing path can execute a separate instruction. In various embodiments, the matrix multiply accelerator **3260** can have architectural features of any one of more of the matrix multiply accelerators described herein. For example, in one embodiment, the matrix multiply accelerator **3260** is a four-deep systolic array with a feedback loop that is configurable to operate with a multiple of four number of logical stages (e.g., four, eight, twelve, sixteen, etc.). In one embodiment the matrix multiply accelerator **3260** includes one or more instances of a two-path matrix multiply accelerator with a four stage pipeline or a four-path matrix multiply accelerator with a two stage pipeline. The matrix multiply accelerator **3260** can be configured to operate only on non-zero values of at least one input matrix. Operations on entire columns or submatrices can be bypassed where block sparsity is present. The matrix multiply accelerator **3260** can also include any logic based on any combination of these embodiments, and particularly include logic to enable support for random sparsity, according to embodiments described herein.

[0299] As illustrated, in one embodiment, and in addition to the graphics processor **3204**, the computing device **3200** may further include any number and type of hardware components and/or software components, including, but not limited to an application processor **3206**, memory **3208**, and input/output (I/O) sources **3210**. The application processor **3206** can interact with a hardware graphics pipeline, as illustrated with reference to FIG. **3A**, to share graphics pipeline functionality. Processed data is stored in a buffer in the hardware graphics pipeline and state information is stored in memory **3208**. The resulting data can be transferred to a display controller for output via a display device, such as the display device **318** of FIG. **3A**. The display device may be of various types, such as Cathode Ray Tube (CRT), Thin Film Transistor (TFT), Liquid Crystal Display (LCD), Organic Light Emitting Diode (OLED) array, etc., and may be configured to display information to a user via a graphical user interface.

[0300] The application processor **3206** can include one or processors, such as processor core(s) **107** of FIG. **1** and may be the central processing unit (CPU) that is used at least in part to execute an operating system (OS) **3202** for the computing device **3200**. The OS **3202** can serve as an interface between hardware and/or physical resources of the computing device **3200** and one or more users. The OS **3202** can include driver logic for various hardware devices in the computing device **3200**. The driver logic can include graphics driver logic **3222**, which can include the user mode graphics

driver **1026** and/or kernel mode graphics driver **1029** of FIG. **10**. The graphics driver logic can include a graphics memory manager **3221** to manage a virtual memory address space for the graphics processor **3204**.

[0301] It is contemplated that in some embodiments the graphics processor **3204** may exist as part of the application processor **3206** (such as part of a physical CPU package) in which case, at least a portion of the memory **3208** may be shared by the application processor **3206** and graphics processor **3204**, although at least a portion of the memory **3208** may be exclusive to the graphics processor **3204**, or the graphics processor **3204** may have a separate store of memory. The memory **3208** may comprise a pre-allocated region of a buffer (e.g., framebuffer); however, it should be understood by one of ordinary skill in the art that the embodiments are not so limited, and that any memory accessible to the lower graphics pipeline may be used. The memory **3208** may include various forms of random-access memory (RAM) (e.g., SDRAM, SRAM, etc.) comprising an application that makes use of the graphics processor **3204** to render a desktop or 3D graphics scene. A memory controller, such as memory controller **116** of FIG. **1** or any other memory controller described herein, may access data in the memory **3208** and forward it to graphics processor **3204** for graphics pipeline processing. The memory **3208** may be made available to other components within the computing device **3200**. For example, any data (e.g., input graphics data) received from various I/O sources **3210** of the computing device **3200** can be temporarily queued into memory **3208** prior to their being operated upon by one or more processor(s) (e.g., application processor **3206**) in the implementation of a software program or application. Similarly, data that a software program determines should be sent from the computing device **3200** to an outside entity through one of the computing system interfaces, or stored into an internal storage element, is often temporarily queued in memory **3208** prior to its being transmitted or stored.

[0302] The I/O sources can include devices such as touchscreens, touch panels, touch pads, virtual or regular keyboards, virtual or regular mice, ports, connectors, network devices, or the like, and can attach via a platform controller hub **130** as referenced in FIG. **1**. Additionally, the I/O sources **3210** may include one or more I/O devices that are implemented for transferring data to and/or from the computing device **3200** (e.g., a networking adapter); or, for a large-scale non-volatile storage within the computing device **3200** (e.g., SSD/HDD). User input devices, including alphanumeric and other keys, may be used to communicate information and command selections to graphics processor **3204**. Another type of user input device is cursor control, such as a mouse, a trackball, a touchscreen, a touchpad, or cursor direction keys to communicate direction information and command selections to GPU and to control cursor movement on the display device. Camera and microphone arrays of the computing device **3200** may be employed to observe gestures, record audio and video and to receive and transmit visual and audio commands.

[0303] The I/O sources **3210** can include one or more network interfaces. The network interfaces may include associated network processing logic and/or be coupled with the network processor tile **3257**. The one or more network interface can provide access to a LAN, a wide area network (WAN), a metropolitan area network (MAN), a personal area network (PAN), Bluetooth, a cloud network, a cellular or mobile network (e.g., 3rd Generation (3G), 4th Generation (4G), 5th Generation (5G), etc.), an intranet, the Internet, etc. Network interface(s) may include, for example, a wireless network interface having one or more antenna (e). Network interface(s) may also include, for example, a wired network interface to communicate with remote devices via network cable, which may be, for example, an Ethernet cable, a coaxial cable, a fiber optic cable, a serial cable, or a parallel cable.

[0304] Network interface(s) may provide access to a LAN, for example, by conforming to IEEE 802.11 standards, and/or the wireless network interface may provide access to a personal area network, for example, by conforming to Bluetooth standards. Other wireless network interfaces and/or protocols, including previous and subsequent versions of the standards, may also be supported. In addition to, or instead of, communication via the wireless LAN standards, network

interface(s) may provide wireless communication using, for example, Time Division, Multiple Access (TDMA) protocols, Global Systems for Mobile Communications (GSM) protocols, Code Division, Multiple Access (CDMA) protocols, and/or any other type of wireless communications protocols.

[0305] It is to be appreciated that a lesser or more equipped system than the example described above may be preferred for certain implementations. Therefore, the configuration of the computing devices described herein may vary from implementation to implementation depending upon numerous factors, such as price constraints, performance requirements, technological improvements, or other circumstances. Examples include (without limitation) a mobile device, a personal digital assistant, a mobile computing device, a smartphone, a cellular telephone, a handset, a one-way pager, a two-way pager, a messaging device, a computer, a personal computer (PC), a desktop computer, a laptop computer, a notebook computer, a handheld computer, a tablet computer, a server, a server array or server farm, a web server, a network server, an Internet server, a work station, a mini-computer, a main frame computer, a supercomputer, a network appliance, a web appliance, a distributed computing system, multiprocessor systems, processor-based systems, consumer electronics, programmable consumer electronics, television, digital television, set top box, wireless access point, base station, subscriber station, mobile subscriber center, radio network controller, router, hub, gateway, bridge, switch, machine, or combinations thereof.

[0306] Embodiments may be provided, for example, as a computer program product which may include one or more machine-readable media having stored thereon machine-executable instructions that, when executed by one or more machines such as a computer, network of computers, or other electronic devices, may result in the one or more machines carrying out operations in accordance with embodiments described herein. A machine-readable medium may include, but is not limited to, floppy diskettes, optical disks, CD-ROMs (Compact Disc-Read Only Memories), and magneto-optical disks, ROMs, RAMs, EPROMs (Erasable Programmable Read Only Memories), EEPROMs (Electrically Erasable Programmable Read Only Memories), magnetic or optical cards, flash memory, or other type of media/machine-readable medium suitable for storing machine-executable instructions.

[0307] Moreover, embodiments may be downloaded as a computer program product, wherein the program may be transferred from a remote computer (e.g., a server) to a requesting computer (e.g., a client) by way of one or more data signals embodied in and/or modulated by a carrier wave or other propagation medium via a communication link (e.g., a modem and/or network connection).

[0308] Throughout the document, term “user” may be interchangeably referred to as “viewer”, “observer”, “person”, “individual”, “end-user”, and/or the like. It is to be noted that throughout this document, terms like “graphics domain” may be referenced interchangeably with “graphics processing unit”, “graphics processor”, or simply “GPU” and similarly, “CPU domain” or “host domain” may be referenced interchangeably with “computer processing unit”, “application processor”, or simply “CPU”.

[0309] It is to be noted that terms like “node”, “computing node”, “server”, “server device”, “cloud computer”, “cloud server”, “cloud server computer”, “machine”, “host machine”, “device”, “computing device”, “computer”, “computing system”, and the like, may be used interchangeably throughout this document. It is to be further noted that terms like “application”, “software application”, “program”, “software program”, “package”, “software package”, and the like, may be used interchangeably throughout this document. Also, terms like “job”, “input”, “request”, “message”, and the like, may be used interchangeably throughout this document.

[0310] It is contemplated that terms like “request”, “query”, “job”, “work”, “work item”, and “workload” may be referenced interchangeably throughout this document. Similarly, an “application” or “agent” may refer to or include a computer program, a software application, a game, a workstation application, etc., offered through an application programming interface (API), such as a free rendering API, such as Open Graphics Library (OpenGL®), Open Computing

Language (OpenCL®), CUDA®, DirectX® 11, DirectX® 12, etc., where “dispatch” may be interchangeably referred to as “work unit” or “draw” and similarly, “application” may be interchangeably referred to as “workflow” or simply “agent”. For example, a workload, such as that of a three-dimensional (3D) game, may include and issue any number and type of “frames” where each frame may represent an image (e.g., sailboat, human face). Further, each frame may include and offer any number and type of work units, where each work unit may represent a part (e.g., mast of sailboat, forehead of human face) of the image (e.g., sailboat, human face) represented by its corresponding frame. However, for the sake of consistency, each item may be referenced by a single term (e.g., “dispatch”, “agent”, etc.) throughout this document.

[0311] References herein to “one embodiment,” “an embodiment,” “an example embodiment,” etc., indicate that the embodiment described may include a particular feature, structure, or characteristic, but every embodiment may not necessarily include the particular feature, structure, or characteristic. Moreover, such phrases are not necessarily referring to the same embodiment. Further, when a particular feature, structure, or characteristic is described in connection with an embodiment, it is submitted that it is within the knowledge of one skilled in the art to affect such feature, structure, or characteristic in connection with other embodiments whether explicitly described.

[0312] In the various embodiments described above, unless specifically noted otherwise, disjunctive language such as the phrase “at least one of A, B, or C” is intended to be understood to mean either A, B, or C, or any combination thereof (e.g., A, B, and/or C). As such, disjunctive language is not intended to, nor should it be understood to, imply that a given embodiment requires at least one of A, at least one of B, or at least one of C to each be present. Similarly, items listed in the form of “at least one of A, B, or C” can mean (A); (B); (C); (A and B); (B and C); or (A, B, and C).

[0313] In some embodiments, terms like “display screen” and “display surface” may be used interchangeably referring to the visible portion of a display device while the rest of the display device may be embedded into a computing device, such as a smartphone, a wearable device, etc. It is contemplated and to be noted that embodiments are not limited to any particular computing device, software application, hardware component, display device, display screen or surface, protocol, standard, etc. For example, embodiments may be applied to and used with any number and type of real-time applications on any number and type of computers, such as desktops, laptops, tablet computers, smartphones, head-mounted displays and other wearable devices, and/or the like. Further, for example, rendering scenarios for efficient performance using this novel technique may range from simple scenarios, such as desktop compositing, to complex scenarios, such as 3D games, augmented reality applications, etc.

[0314] Described herein are techniques for learning neural reflectance shaders from images. A set of one or more machine learning models can be trained to optimize an illumination latent code and a set of reflectance latent codes for an object within a set of input images. A shader can then be generated based on a machine learning model of the one or more machine learning models. The shader is configured to sample the illumination latent code and the set of reflectance latent codes for the object. A 3D representation of the object can be rendered using the generated shader.

[0315] One embodiment provides a data processing system comprising a memory to store instructions and one or more processors including a graphics processor. The graphics processor includes circuitry to execute the instructions. The instructions can cause the circuitry to process, via one or more machine learning models, a set of input images, generate, via the one or more machine learning models, an illumination latent code and a set of reflectance latent codes for an object within the set of input images, generate a shader based on a machine learning model of the one or more machine learning models, the shader configured to sample the illumination latent code and the set of reflectance latent codes for the object, and render a three-dimensional (3D) object by shading a mesh via the shader. The circuitry can be configured to render the 3D object via a real-time

renderer.

[0316] In one embodiment, the machine learning models include a plurality of machine learning models. The plurality of machine learning models can include a first model to implement a reflectance network (e.g., Reflectance Net) and a second model to implement an illumination network (e.g., Illumination Net). Each of the first model and the second model may be implemented as a multilayer perceptron and the shader can be generated based on the multilayer perceptron associated with the reflectance network. In one embodiment, the illumination latent code is a per-scene latent code, while the set of reflectance latent codes include a latent code for each vertex of the mesh. In one embodiment, the instructions cause the circuitry to generate the mesh and a set of camera poses for the set of input images. The mesh can be a 3D mesh, with the mesh and the set of camera poses being generated via structure from motion analysis of the set of input images. The set of input images can include a plurality of input images.

[0317] One embodiment provides a method that is performed by the data processing system described above. In one embodiment, a non-transitory machine-readable medium can store instructions to implement such method via one or more processors that include a graphics processor.

[0318] The foregoing description and drawings are to be regarded in an illustrative rather than a restrictive sense. Persons skilled in the art will understand that various modifications and changes may be made to the embodiments described herein without departing from the broader spirit and scope of the features set forth in the appended claims.

Claims

1. An apparatus comprising: a host interface; one or more processors coupled with the host interface, the one or more processors including a graphics processor, wherein the one or more processors are configured to execute instructions stored in a memory, the instructions to cause the one or more processors to: generate, via at least one machine learning model, a set of latent codes to estimate a surface of a material for an object in a scene, wherein the set of latent codes includes an illumination latent code to represent lighting information for the surface of the material, and reflectance latent codes to represent surface reflectance for the material; and render a three-dimensional representation of the object via a real-time renderer by shading a mesh via sampled latent codes from the set of latent codes to represent an effect of the material.
2. The apparatus as in claim 1, wherein the illumination latent code is to represent lighting information for multiple objects within a scene.
3. The apparatus as in claim 2, wherein the reflectance latent codes include a latent code for respective vertices of the mesh.
4. The apparatus as in claim 2, wherein the at least one machine learning model includes a plurality of machine learning models, the plurality of machine learning models including a first model to implement a reflectance network and a second model to implement an illumination network.
5. The apparatus as in claim 4, wherein the graphics processor of the one or more processors includes a matrix accelerator configured to combine output of the first model and the second model in association with a forward shading operation.
6. The apparatus as in claim 5, wherein the first model and the second model are respectively implemented as a multilayer perceptron and the one or more processors to are configured to generate a shader based on the multilayer perceptron associated with the reflectance network, the shader configured to sample the set of latent codes for the object.
7. The apparatus as in claim 6, wherein the mesh has associated texture coordinates and the reflectance latent codes include a latent code for respective texels associated with the mesh.
8. The apparatus as in claim 1, wherein the instructions cause the one or more processors to generate the mesh and a set of camera poses for the scene.

- 9.** The apparatus as in claim 8, wherein the mesh is a three-dimensional mesh and the one or more processors are configured to generate the mesh and the set of camera poses via structure from motion analysis of the scene.
- 10.** The apparatus as in claim 9, wherein to generate the set of latent codes via the at least one machine learning model includes to execute an instruction to perform a set of matrix multiply operations via a tensor core of the graphics processor.
- 11.** A non-transitory machine-readable medium having instructions stored thereon, the instructions, when executed by one or more processors, cause the one or more processors to perform operations comprising: generating, via at least one machine learning model, a set of latent codes to estimate a surface of a material for an object in a scene, wherein the set of latent codes includes an illumination latent code to represent lighting information for the surface of the material, and reflectance latent codes to represent surface reflectance for the material, the set of latent codes generated via the at least one machine learning model via instructions executed by a graphics processor of the one or more processors; and rendering a three-dimensional representation of the object via a real-time renderer by shading a mesh via sampled latent codes from the set of latent codes to represent an effect of the material.
- 12.** The non-transitory machine-readable medium as in claim 11, wherein the illumination latent code is to represent lighting information for multiple objects within a scene.
- 13.** The non-transitory machine-readable medium as in claim 12, wherein the at least one machine learning model includes a plurality of machine learning models, the plurality of machine learning models including a first model to implement a reflectance network and a second model to implement an illumination network.
- 14.** The non-transitory machine-readable medium as in claim 13, wherein the first model and the second model are respectively implemented as a multilayer perceptron.
- 15.** The non-transitory machine-readable medium as in claim 14, the operations further comprising generating a shader based on the multilayer perceptron associated with the reflectance network, the shader configured to sample the set of latent codes for the object.
- 16.** The non-transitory machine-readable medium as in claim 15, wherein the reflectance latent codes include a latent code for respective vertices of the mesh.
- 17.** The non-transitory machine-readable medium as in claim 16, wherein the mesh has associated texture coordinates and the reflectance latent codes include a latent code for respective texels associated with the mesh.
- 18.** The non-transitory machine-readable medium as in claim 11, the operations further comprising generating the mesh and a set of camera poses for the scene.
- 19.** The non-transitory machine-readable medium as in claim 18, wherein the mesh is a three-dimensional mesh and the operations further comprise generating the mesh and the set of camera poses via structure from motion analysis of the scene.
- 20.** The non-transitory machine-readable medium as in claim 19, generating the set of latent codes via the at least one machine learning model includes executing an instruction to perform a set of matrix multiply operations via a tensor core of the graphics processor.
-