

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250259630

Kind Code

A1

Publication Date

August 14, 2025

Inventor(s)

KUDLIGI ANANTHA; Harsha et al.

METHODS AND SYSTEMS FOR PROVIDING INSIGHTS IN REAL-TIME DURING A CONVERSATION

Abstract

The disclosure describes systems, methods, and media for generating real-time insights in a voice over internet protocol (VOIP) conversation. According to the methods, an application server receives a transcript of one or more voice utterances of a participant in the VOIP conversation, and identifies a context of the VOIP conversation and a first state of the context based on the transcript. The application server further receives an intent of the participant from a conversation artificial intelligence (AI) engine based on the transcript provided to the conversation AI engine. The application server further formulates one or more queries based on the intent, the context, and the first state of the context to retrieve one or more insights from one or more backend databases, and transmits the one or more insights to a terminal of at least one participants of the VoIP conversation for display.

Inventors: KUDLIGI ANANTHA; Harsha (Pleasanton, CA), SAH; Subodh Kishorilal (Santa Clara, CA), SHEKAR; Rashmi (San Francisco, CA), PATIL; Shailesh (Lathrop, CA), SHANKAR; Shreyas (Campbell, CA), BUZA; Kyle (Minneapolis, MN), MOHANA KRISHNA; Jayanth (Sunnyvale, CA)

Applicant: Clari Inc. (Sunnyvale, CA)

Family ID: 1000008563570

Appl. No.: 19/194438

Filed: April 30, 2025

Related U.S. Application Data

parent US continuation 17823380 20220830 parent-grant-document US 12315505 child US 19194438

Publication Classification

Int. Cl.: G10L15/22 (20060101); G10L15/18 (20130101); G10L15/30 (20130101)

U.S. Cl.:

CPC G10L15/22 (20130101); G10L15/1815 (20130101); G10L15/30 (20130101);
G10L2015/228 (20130101)

Background/Summary

RELATED APPLICATIONS [0001] This application is a continuation of U.S. patent application Ser. No. 17/823,380, filed Aug. 30, 2022, which is incorporated by reference in its entirety.

TECHNICAL FIELD

[0002] Embodiments of the present invention relate generally to voice over internet protocol (VOIP). More particularly, embodiments of the invention relate to obtaining real-time insights based on conversation context changes during a VoIP conversation.

BACKGROUND

[0003] VOIP calls are becoming increasingly common these days during internal and external sales calls. During either type of sales calls, participants may need real-time insights (e.g., relevant information or talking points) to make the calls more productive.

[0004] For example, in an internal sales call with five participants and where multiple different accounts need to be discussed, some participants may need to live share their screens to show statistics for the different accounts. As another example, in a VoIP call between a sales representative and a customer, the sales representative may need to state the pros and cons of his products in real time for comparison whenever the name of a competitor is mentioned.

[0005] Existing VoIP applications do not have the features to provide real-time insights. Thus, it would be beneficial for sales teams to have a VoIP application with such features to better engage existing or potential customers and/or move the sales calls forward faster.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0006] Embodiments of the invention are illustrated by way of example and are not limited to the figures of the accompanying drawings, in which like references indicate similar elements.

[0007] FIG. 1 is a schematic diagram illustrating a system for generating real-time insights during a VoIP conversation according to an embodiment of the invention.

[0008] FIG. 2 illustrates a system for populating the conversation database according to an embodiment of the invention.

[0009] FIG. 3 illustrates an example of the conversation state cache according to an embodiment of the invention.

[0010] FIG. 4 is a flow diagram illustrating an example use case of generating real-time insights based on voice utterances of a participant in a VoIP conversation according to an embodiment of the invention.

[0011] FIG. 5 illustrates an interface of a VoIP application in which both a live transcript of a VoIP conversation and related insights are displayed according to an embodiment of the invention.

[0012] FIGS. 6A and 6B illustrate two different VOIP conversations with different participants

according to different embodiments of the invention.

[0013] FIG. 7 illustrates a process of generating real-time insights during a VoIP conversation according to an embodiment of the invention.

[0014] FIG. 8 is a block diagram illustrating an example of a data processing system which may be used with any embodiment of the invention.

DETAILED DESCRIPTION

[0015] Various embodiments and aspects of the inventions will be described with reference to the details discussed below, and the accompanying drawings will illustrate the various embodiments. The following description and drawings are illustrative of the invention and are not to be construed as limiting the invention. Numerous specific details are described to provide a thorough understanding of various embodiments of the present invention. However, in certain instances, well-known or conventional details are not described in order to provide a concise discussion of embodiments of the present inventions.

[0016] Reference in the specification to “one embodiment” or “an embodiment” means that a particular feature, structure, or characteristic described in conjunction with the embodiment can be included in at least one embodiment of the invention. The appearances of the phrase “in one embodiment” or “in an embodiment” in various places in the specification do not necessarily all refer to the same embodiment.

[0017] The disclosure describes systems, methods, and media for generating real-time insights in a voice over internet protocol (VOIP) conversation. The insights can function as cues to remind a participant of proper things to say in response to a question of another participant in the VoIP conversation, or provide additional information regarding a particular matter raised in the VOIP conversation. According to a method in an embodiment, an application server receives a transcript of one or more voice utterances of a participant in the VOIP conversation, and identifies a context of the VoIP conversation and a first state of the context based on the transcript. The application server further receives an intent of the participant from a conversation artificial intelligence (AI) engine, which identifies the intent based on the transcript provided to the conversation AI engine. The application server further formulates one or more queries based on the intent, the context, and the first state of the context to retrieve one or more insights from one or more backend databases, and transmits the one or more insights to a terminal of at least one participant in the VoIP conversation for display.

[0018] In an embodiment, each of the one or more participants is a participant with one or more predetermined attributes, for example, an internal employee of the company that owns the backend application server. The one or more insights are not displayed on a terminal device of at least one participant in the VoIP conversation, and at least one participant has one or more predetermined attributes that one or more participants do not have. For example, a customer, who is not an employee of the company that owns the backend application server typically will not see one or more insights due to data privacy and security.

[0019] In an embodiment, the context of the VoIP conversation is a data object that defines one or more of the following properties: participants of the VoIP conversation, a period that the VoIP conversation relates to, whether the participants include an external party, contact information of the external party, activities between the participants, whether a competitor is mentioned, or deal information among the participants. One or more additional properties can be added to the data object based on the one or more void utterances.

[0020] In an embodiment, the first state of the context is changed to a second state when each of the one or more additional properties is added to the data object or when a value of an existing property changes. The context of the VoIP conversation, a unique identifier of the VOIP conversation, and the first state and the second state of the context of the VOIP conversation are stored in a cache in a cloud environment, and the cache further includes a state machine that keeps tracks of each state of the context during the VoIP conversation. The intent generated by the

conversation AI engine is further refined based on the state machine.

[0021] In an embodiment, the backend application server formulates one or more queries based on the intent, the context, and the state of the context to retrieve the one or more insights from one or more backend databases in the cloud environment. Backend databases include a task database and a conversation database, wherein the task database includes information related to a plurality of tasks, and wherein the conversation database includes a plurality of entries, each entry mapping an insight to a combination of a context, a state of the context, and an intent.

[0022] The other embodiments, functions and advantages will be apparent from the accompanying drawings and from the detailed description that follows.

[0023] FIG. 1 is a schematic diagram illustrating a system **100** for generating real-time insights during a VoIP conversation according to an embodiment of the invention. As shown in FIG. 1, the system **100** includes one or more terminal devices, such as mobile phones, laptop computers, etc., and servers and databases residing in the cloud and/or on-premises environment. This embodiment uses one terminal device **115**, and servers and databases residing in a cloud environment **101** for the purpose of illustration.

[0024] The cloud environment **101** can include a frontend application server **103**, a backend application server **105**, a conversation artificial intelligence (AI) engine **111**, a conversation state cache **107**, a conversation database **109**, and a task database **113**.

[0025] The frontend application server **103** can include a number of services, such as a VOIP service **102**, a speech to text service **104**, and a live call listener service **106**. The VOIP service **102** converts voice utterances of a user of the terminal device **115** into a digital signal that travels across the internet, allowing the user to make a call directly from the terminal device **115**. The speech to text service **102** can be a speech recognition service that converts the voice utterances/or speech of the user of the terminal device **115** into transcripts. The live call listener service **106** allows a third party to listen to a live call managed by the frontend application server **104** without the knowledge of either party on the call.

[0026] The backend application **105** can receive a transcript **108** generated by the frontend application server **103**, and provide the transcript **108** as input to the conversation AI engine **111**, and receive intents **114** generated by the conversation AI engine **111**.

[0027] In an embodiment, the conversation AI engine **111** can be a natural language processing (NLP) engine that combines computational linguistics (i.e., rule-based modeling of human language) with statistical, machine learning, and deep learning models to understand the meaning of human language, including the intent of the speaker. The conversation AI engine **111** can include a number of components for extracting linguistic features of the transcript **108**, such as the basic grammar rules, word organizations, and relations between words. The number of components can include a machine learning model for inferring the meaning of each sentence in the transcript **108**, and an intent classifier for identifying one or more intents of each sentence in the transcripts **108**. The intent classifier can be a rule-based pattern matching algorithm, or a deep learning model (e.g., an artificial neural network).

[0028] In an embodiment, an intent can be a text expression indicating the intention of the speaker when uttering a sentence. For example, if the transcript **108** is “I would like to book a flight from San Francisco to New York City,” the conversation AI engine **114** would analyze the sentence to infer its meaning, and then classify the sequence of words under the label “book flight”. In this case, the text “book flight” is an intent of the speaker.

[0029] In an embodiment, the conversation AI engine **111** can generate multiple intents for a transcript. In such a case, the conversation AI engine **111** would rank the multiple intents based on a number of factors, such as how frequently each intent appears historically. A more frequently appearing intent is ranked higher than a less frequently appearing intent if all the other ranking factors are the same. The highest ranked intent will be selected as the intent of the sentence by the conversation AI engine **111**.

[0030] The conversation state cache **107** can store each VOIP conversation along with a context of the VoIP conversation and an identifier that uniquely identifies the VoIP conversation. The context of the VoIP conversation can be a data object, for example, a JavaScript Object Notation (JSON) object, and can include one or more of the following properties: participants of the VOIP conversation, a period of time that the VoIP conversation relates to, whether the participants include an external party, contact information of the external party, activities between the participants, whether a competitor is mentioned, or deal information among the participants. The properties in each context can be deleted or added, and the value of each property can be also updated. Each change in a value of a property or each property addition or deletion can change the state of the context. The conversation state cache **107** also keeps track of the state changes in the context of each VoIP conversation.

[0031] In an embodiment, a state change can be caused by a conversation participant raising a new topic, mentioning a new competitor in a new voice utterance during a VoIP conversation. Multiple voice utterances can include the same intent. Thus, one state change can correspond to multiple transcripts, each of which can correspond to one or more intents.

[0032] The conversation database **109** can be a NoSQL database that stores entries that each map a state of a context of a VoIP conversation and a corresponding intent to an insight. Thus, given a state of a context and an intent associated with the context state, an insight corresponding to the state and the intent can be retrieved from the conversation database **109**.

[0033] In an embodiment, the insights in the conversation database **109** can be derived from historical data, including activity data and account data. The derivation of insights is described in detail in FIG. 2.

[0034] The task database **113** can be a customer relationship management (CRM) system, and can store information about deals. Examples of the deal information include potential value of the deal, a stage of the deal (e.g., “commit” or “completed”), contact persons, and their contact information. The contact persons of a deal can include one or more source contact persons (e.g., sales representatives) and one or more target contact persons (e.g., customer contact persons).

[0035] When receiving a transcript from the frontend application server **103**, the backend application server **105** can determine a context based on the participants of the VoIP conversation, and information associated with the participants that is retrieved from the task database **113**.

[0036] By default, each context has a default state. The default state can be changed by an addition or a deletion of a property, or by a change in a value of an existing property. The backend application server **105** can formulate one or more queries based on the intent generated by the conversation engine **111**, a context of a VoIP conversation, and a state of the context, and can intelligently determine which database to query.

[0037] For example, in one embodiment, if the intent concerns a simple retrieval of information, such as contact information of target contacts, the backend application server **105** can query the task database **113**; otherwise, the backend application server **105** can query the conversation database **109**. In another embodiment, the backend application server **105** will query both the task database **113** and the conversation database **109**.

[0038] FIG. 2 illustrates a system **200** for populating the conversation database **109** according to an embodiment of the invention. As shown in FIG. 2, the system **200** can include a data analytics server **205** for retrieving data from a number of data sources, such as the task database **113**, an activity database **201**, and a domain name registry **203**. The system **200** can use a machine learning model **202** to identify insights from the retrieved data, and populate the conversation database **109** with the identified insights.

[0039] The activity database **201** can include activity data such as historical email communications, meetings and transcripts of phone calls (i.e., conversations), and persons who conducted the activity (i.e., participants of the activity). A database manager **204** running on the data analytics server **205** can provide one or more services to initiate or schedule processing threads to

synchronize data between the activity database **201** and the task database system **113**, such that each task/deal in the task database system **113** can have updated activity information for that task. The synchronization can be performed automatically and periodically using multiple processing threads. Through a first processing thread executed by a processor, the database manager **204** can access the activity database **201** via an application programming interface (API) over a network to retrieve a list of event objects that are associated with a number of activities in the activity database **201**.

[0040] Through a second processing thread, the database manager **204** can determine one or more participant identifiers (IDs) from each of the event objects. The participant IDs identify one or more participants in an event or activity associated with the event object. The database manager **204** can extract a domain ID from each participant ID. Alternatively, the database manager **204** can obtain the domain ID for each participant by searching a data structure or database, or through a third party such as a domain registry server **203**. The database manager **204** can identify one or more entity objects based on the domain ID, where the entity objects (e.g., an account object) are stored in the task database **113**.

[0041] In an embodiment, at least one attribute of at least one of the entity objects is modified based on the participant ID and the domain ID. The modification can generate at least one modified entity object. The database manager **204** can use a third processing thread to transmit one or more modified event objects to the task database **113** via a second API over the network. The first processing thread, the second processing thread, and the third processing thread can be executed independently.

[0042] In one embodiment, in modifying at least one attribute of at least one entity object, the database manager can determine whether there are multiple entity objects associated with the same domain ID. If there are, the database manager **204** can select a first entity object from the multiple entity objects based on a set of rules, and modify one or more attributes of the selected entity object based on the participant ID and the domain ID.

[0043] If the database manager **204** determines that a participant ID matches a creator ID or owner ID corresponding to a creator or owner of any of the entity objects, the database manager **204** can select the first entity object for modification if the first entity object is the only entity object whose creator ID matches the participant ID. Otherwise, if the first entity object is not the only matching entity object, the database manager **204** can designate the multiple entity objects as entity object candidates.

[0044] In addition, for each of the entity objects, the database manager **204** can determine whether one or more task objects are associated with the entity object. Each task object is associated with a task to be completed within a predetermined period of time. For each of the task objects, the database manager **204** can determine whether a participant ID matches a user ID of any user of a user group associated with the task object. If so, at least one attribute of the task object is modified based on the participant ID.

[0045] In an embodiment, a data store **206** can be provided in the data analytics server **205** for use to store or cache data for a variety of tasks or activities, which may be periodically updated from the corresponding data source(s) or data provider(s), in this example, database systems **113** and **201**.

[0046] The machine learning model **202** can use a set of predetermined rules to identify insights from the task data and the activity data. As discussed above, each task can be in one of a number of stages. The machine learning model **202** can examine activity data corresponding to each stage of the task, and determines which type of activity might have contributed to the progression of the task from one stage to the next stage, and identify a particular email or a particular phone call as insights.

[0047] For example, the machine learning model **202** can examine activity data of two similar tasks. If task A moves from one stage to a next stage while task B does not, and the only difference

between the two tasks is that the sales representative for task A addresses the customer with “You” frequently when requesting an in-person meeting while the sales representative for task B rarely address the customer directly, then the machine learning model can identify “directly addressing customers with ‘You’” as an insight.

[0048] For each identified insight, the machine learning model **202** can also identify a context, a state of the context, and an intent of the person associated with the activity. In the above example, the context would be a communication between a sales representative and a customer, the state of the context would be a default state, and the intent of the sales representative would be to request an in-person meeting.

[0049] The conversation database **109** can be populated with all insights identified from the task data and the activity data an insight, and intents, contexts, and context states that correspond to the insights.

[0050] As shown in FIG. 2, the conversation database **109** includes three insights **221**, **229**, and **237** for a context A **215** that respectively correspond to context states **217**, **225**, and **244**, and intents **219**, **227** and **235**. The three entries are provided for the purpose of illustration. In actual implementations, the conversation database **109** can include as many insights as identified by the machine learning model **202** from the data sources **113**, **201**, and **203**.

[0051] FIG. 3 illustrates an example of the conversation state cache **107** according to an embodiment of the invention. The conversation state cache **107** can store conversation contexts and all of their states. During a VoIP conversation, the backend application server **105** stores each context identified and each state of the context to the conversation state cache **107**, and also retrieves the context and its associated states during the conversation for use in identifying insights.

[0052] In an embodiment, the conversation state cache **107** can store a name-value pair for each conversation. The name can be a unique ID for a particular conversation, and the value can be the context of the conversation.

[0053] As shown in FIG. 3, a context **307** for conversation A **303** and another context for conversation B **305** are stored in the conversation state cache **107**, and each context is associated with a unique identifier **304** and **306**. Further, states **311**, **313**, and **315** for the context **307** and states **317**, **319**, and **321** for the context **308** are stored in the conversation state cache **107**.

[0054] The state information for each conversation constitutes a state machine that keeps track of all states of the context of a VoIP conversation, e.g., what is being asked previously in the conversation. The state machine enables the backend application server **105** to identify an intent of the person more accurately.

[0055] For example, if a current state is associated with an utterance of “contact information”, this utterance alone is insufficient for the backend application server **105** to determine the intent of the speaker. However, if one or more earlier states in the state machine is associated with “a customer”, then the combination of the current state and the one or more earlier states would be sufficient to identify the intent of the participant in the VoIP conversation as requesting customer contact information.

[0056] Thus, in an embodiment, an intent generated by the conversation AI engine **111** can be refined based on one or more states in the state machine stored in the conversation states cache **107**.

[0057] FIG. 4 is a flow diagram illustrating an example use case of generating real-time insights based on voice utterances of a participant in a VoIP conversation according to an embodiment of the invention. This diagram focuses on particular use case among multiple possible use cases. More specifically, this diagram focuses on a use case defined by the path that includes blocks **401**, **402**, **403**, and **404**, with the block **401** representing the beginning of the conversation, and each of the other blocks representing a step where one of the participants makes a meaningful voice utterance.

[0058] The path **401**, **402**, **403**, and **404** represents a VoIP conversation between a salesperson “John” and his manager “Matt”, and this VoIP conversation can have a context “overview” created by the backend application server **105** when the VOIP conversation is established. The context is in

a default state **405** at the block **401**, and can change to different states **406**, **407**, and **408** as either the value of its property changes or a new property is added. Each state change may have an insight generated by the system **100** and sent to John and Matt on their respective terminals for display. [0059] For example, after the conversation starts, Matt asks John “How are things going with ACME?” This voice utterance by Matt can be converted into a transcript, which can be provided to the conversation AI engine **111**. The conversation AI engine **111** can identify the intent of Matt as wanting to know if there are any issues and/or progress with ACME. This intent can be further refined based on historical states of the context reflecting what has been previously discussed in this conversation. The backend application server **105** can construct one or more queries from the refined intent and the context state **404**, and run the one or more queries against the conversation database **109** and/or the task database **113** to retrieve one or more insights.

[0060] For example, one of the queries may retrieve from the task database **113** all tasks/deals whose stages have changed, and send the stage change information, including the names of the tasks, to the terminal device of John and Matt. Another query may retrieve another insight from the conversation database **109** if there is an entry that matches the context state, the context, and the intent.

[0061] In an embodiment, it may happen that none of the queries may retrieve anything from either database. When that happens, no insight will be displayed on the terminal devices of John and Matt.

[0062] At the block **403**, the context state changes again when Matt, the manager, asked the question, “Have you been in touch with Rob?” In this case, the backend application **105** may go through the above process described with the block **403**, and formulate one or more queries to retrieve insights from the task database **113** and/or the conversation database **109**. Examples of the insights include information indicating that Rob is a customer contact; his email address; and that there are email exchanges between John and Rob.

[0063] At the block **404**, insights may similarly be retrieved from the task database **113** and/or the conversation database **109**. Examples of the insights include information indicating the total amount of annual recurring revenue (ARR) from AMCE.

[0064] In an embodiment, the VoIP conversation between John and Matt that starts at the block **401** may also take other paths involving other blocks **409-431**. The other paths (use cases) are not described herein detail.

[0065] FIG. 5 illustrates an interface of a VoIP application in which both a live transcript of a VoIP conversation and related insights are displayed according to an embodiment of the invention. In this figure, sentences **501** and **505** in the box for the live transcript **117** are translated from voice utterances of one party/participant (e.g., a sales manager) of the VoIP conversation, and sentences **503** and **507** are translated from voice utterances of another participant (e.g., a sales representative).

[0066] In the box for real-time insights **119**, messages **509** and **511** are related to the sentence **501**, and message **513** is related to the sentence **505**. The insights are displayed in real-time while the VoIP conversation is on-going. In an embodiment, the insights can provide cues for a call participant to answer a question asked by another participant, or can provide additional information so that all participants are on the same page regarding a particular matter.

[0067] FIGS. 6A and 6B illustrate two different VOIP conversations with different participants according to different embodiments of the invention. In FIG. 6A, one participant is a sales representative **609**, and the other party is a sales manager **601**. Since both parties are employees of the same company, real-time insights **607** are displayed to terminal devices **601** and **603** of both participants, along with a live transcript **605** of the VoIP conversation.

[0068] Further, if an employee of the same company listens in to the call, that employee's terminal device may also have the live transcript **604** and the real-time insights **607** displayed.

[0069] FIG. 6B relates to a different VoIP conversation between the sales representative **609** and a contact of customer **615** of the company, and insights may concern data security or trade secret of a

company. Thus, insights **613** may be displayed on a terminal device of the sales representative **609**, but not on a terminal device **604** of the contact of the customer **615**. Similarly, if another contact from the customer **615** listens to the call, that contact will also not see the insights **613** on his or her terminal device.

[0070] FIG. **7** illustrates a process **700** of generating real-time insights during a VoIP conversation according to an embodiment of the invention. The process **700** can be performed by processing logic that includes software, hardware, or a combination thereof. For example, the processing logic may be performed by the backend application server **105** and/or one or more other components in the cloud environment **101** as described in FIG. **1**.

[0071] The process **700** enables real-time insights to be generated and to be displayed in a terminal alongside a transcript of the VoIP conversation. The insights can function as cues to remind a participant of proper things to say in response to a question from another participant in the VoIP conversation, or provide additional information regarding a particular matter raised in the VOIP conversation.

[0072] Referring to FIG. **7**, in operation **701**, the processing logic receives a transcript of one or more voice utterances of a participant in the VoIP conversation. In operation **703**, the processing logic identifies a context of the VoIP conversation and a first state of the context based on the transcript. In operation **705**, the processing logic receives an intent of the participant from a conversation artificial intelligence (AI) engine, wherein the intent is identified by the conversation AI engine based on a sequence of words in the transcript. In operation **707**, the processing logic determines one or more insights based on the context of the VoIP conversation, the first state of the context, and the intent of the participant. In operation **709**, the processing logic transmits the one or more insights to a terminal device of each of one or more participants of the VoIP conversation, wherein the one or more insights are displayed in real time in the terminal device of each of the one or more participants.

[0073] In an embodiment, the insights are not displayed to an external participant in the VOIP conversation for security and privacy reasons.

[0074] FIG. **8** is a block diagram illustrating an example of a data processing system **800** which may be used with any embodiment of the invention. For example, system **800** may represent any of data processing systems described above, such as data analytics system **104**, task database system **105**, or activity database system **106**, etc., performing any of the processes or methods described above. System **800** can include many different components. These components can be implemented as integrated circuits (ICs), portions thereof, discrete electronic devices, or other modules adapted to a circuit board such as a motherboard or add-in card of the computer system, or as components otherwise incorporated within a chassis of the computer system.

[0075] System **800** may represent a desktop, a laptop, a tablet, a server, a mobile phone, a media player, a personal digital assistant (PDA), a Smartwatch, a personal communicator, a gaming device, a network router or hub, a wireless access point (AP) or repeater, a set-top box, or a combination thereof. Further, while only a single machine or system is illustrated, the term “machine” or “system” shall also be taken to include any collection of machines or systems that individually or jointly execute a set (or multiple sets) of instructions to perform any one or more of the methodologies discussed herein.

[0076] For one embodiment, system **800** includes processor **801**, memory **803**, and devices **805-808** via a bus or an interconnect **810**. Processor **801** may represent a single processor or multiple processors with a single processor core or multiple processor cores included therein. Processor **801** may represent one or more general-purpose processors such as a microprocessor, a central processing unit (CPU), or the like. More particularly, processor **801** may be a complex instruction set computing (CISC) microprocessor, reduced instruction set computing (RISC) microprocessor, very long instruction word (VLIW) microprocessor, or processor implementing other instruction sets, or processors implementing a combination of instruction sets. Processor **801** may also be one

or more special-purpose processors such as an application specific integrated circuit (ASIC), a cellular or baseband processor, a field programmable gate array (FPGA), a digital signal processor (DSP), a network processor, a graphics processor, a network processor, a communications processor, a cryptographic processor, a co-processor, an embedded processor, or any other type of logic capable of processing instructions.

[0077] Processor **801**, which may be a low power multi-core processor socket such as an ultra-low voltage processor, may act as a main processing unit and central hub for communication with the various components of the system. Such processor can be implemented as a system on chip (SoC). Processor **801** is configured to execute instructions for performing the operations and steps discussed herein. System **800** may further include a graphics interface that communicates with optional graphics subsystem **804**, which may include a display controller, a graphics processor, and/or a display device.

[0078] Processor **801** may communicate with memory **803**, which in one embodiment can be implemented via multiple memory devices to provide for a given amount of system memory. Memory **803** may include one or more volatile storage (or memory) devices such as random access memory (RAM), dynamic RAM (DRAM), synchronous DRAM (SDRAM), static RAM (SRAM), or other types of storage devices. Memory **803** may store information including sequences of instructions that are executed by processor **801**, or any other device. For example, executable code and/or data of a variety of operating systems, device drivers, firmware (e.g., input output basic system or BIOS), and/or applications can be loaded in memory **803** and executed by processor **801**. An operating system can be any kind of operating systems, such as, for example, Windows® operating system from Microsoft®, Mac OS®/iOS® from Apple, Android® from Google®, Linux®, Unix®, or other real-time or embedded operating systems such as VxWorks.

[0079] System **800** may further include IO devices such as devices **805-808**, including network interface device(s) **805**, optional input device(s) **807**, and other optional IO device(s) **807**. Network interface device **805** may include a wireless transceiver and/or a network interface card (NIC). The wireless transceiver may be a WiFi transceiver, an infrared transceiver, a Bluetooth transceiver, a WiMax transceiver, a wireless cellular telephony transceiver, a satellite transceiver (e.g., a global positioning system (GPS) transceiver), or other radio frequency (RF) transceivers, or a combination thereof. The NIC may be an Ethernet card.

[0080] Input device(s) **807** may include a mouse, a touch pad, a touch sensitive screen (which may be integrated with display device **804**), a pointer device such as a stylus, and/or a keyboard (e.g., physical keyboard or a virtual keyboard displayed as part of a touch sensitive screen). For example, input device **806** may include a touch screen controller coupled to a touch screen. The touch screen and touch screen controller can, for example, detect contact and movement or break thereof using any of a plurality of touch sensitivity technologies, including but not limited to capacitive, resistive, infrared, and surface acoustic wave technologies, as well as other proximity sensor arrays or other elements for determining one or more points of contact with the touch screen.

[0081] IO devices **807** may include an audio device. An audio device may include a speaker and/or a microphone to facilitate voice-enabled functions, such as voice recognition, voice replication, digital recording, and/or telephony functions. Other IO devices **807** may further include universal serial bus (USB) port(s), parallel port(s), serial port(s), a printer, a network interface, a bus bridge (e.g., a PCI-PCI bridge), sensor(s) (e.g., a motion sensor such as an accelerometer, gyroscope, a magnetometer, a light sensor, compass, a proximity sensor, etc.), or a combination thereof. Devices **807** may further include an imaging processing subsystem (e.g., a camera), which may include an optical sensor, such as a charged coupled device (CCD) or a complementary metal-oxide semiconductor (CMOS) optical sensor, utilized to facilitate camera functions, such as recording photographs and video clips. Certain sensors may be coupled to interconnect **810** via a sensor hub (not shown), while other devices such as a keyboard or thermal sensor may be controlled by an embedded controller (not shown), dependent upon the specific configuration or design of system

800.

[0082] To provide for persistent storage of information such as data, applications, one or more operating systems and so forth, a mass storage (not shown) may also couple to processor **801**. In various embodiments, to enable a thinner and lighter system design as well as to improve system responsiveness, this mass storage may be implemented via a solid state device (SSD). However, for other embodiments, the mass storage may primarily be implemented using a hard disk drive (HDD) with a smaller amount of SSD storage to act as an SSD cache to enable non-volatile storage of context state and other such information during power down events so that a fast power up can occur on re-initiation of system activities. A flash device may also be coupled to processor **801**, e.g., via a serial peripheral interface (SPI). This flash device may provide for non-volatile storage of system software, including a BIOS as well as other firmware of the system.

[0083] Storage device **808** may include computer-accessible storage medium **809** (also known as a machine-readable storage medium or a computer-readable medium) on which is stored one or more sets of instructions or software (e.g., module, unit, and/or logic **828**) embodying any one or more of the methodologies or functions described herein. Module/unit/logic **828** may represent any of the components described above. Module/unit/logic **828** may also reside, completely or at least partially, within memory **803** and/or within processor **801** during execution thereof by data processing system **800**, memory **803** and processor **801** also constituting machine-accessible storage media. Module/unit/logic **828** may further be transmitted or received over a network via network interface device **805**.

[0084] Computer-readable storage medium **809** may also be used to store some software functionalities described above persistently. While computer-readable storage medium **809** is shown in an exemplary embodiment to be a single medium, the term “computer-readable storage medium” should be taken to include a single medium or multiple media (e.g., a centralized or distributed database, and/or associated caches and servers) that store the one or more sets of instructions. The terms “computer-readable storage medium” shall also be taken to include any medium that is capable of storing or encoding a set of instructions for execution by the machine and that cause the machine to perform any one or more of the methodologies of the present invention. The term “computer-readable storage medium” shall accordingly be taken to include, but not be limited to, solid-state memories, and optical and magnetic media, or any other non-transitory machine-readable medium.

[0085] Module/unit/logic **828**, components and other features described herein can be implemented as discrete hardware components or integrated in the functionality of hardware components such as ASICs, FPGAs, DSPs or similar devices. In addition, module/unit/logic **828** can be implemented as firmware or functional circuitry within hardware devices. Further, module/unit/logic **828** can be implemented in any combination hardware devices and software components.

[0086] Some portions of the preceding detailed descriptions have been presented in terms of algorithms and symbolic representations of operations on data bits within a computer memory. These algorithmic descriptions and representations are the ways used by those skilled in the data processing arts to most effectively convey the substance of their work to others skilled in the art. An algorithm is here, and generally, conceived to be a self-consistent sequence of operations leading to a desired result. The operations are those requiring physical manipulations of physical quantities.

[0087] Embodiments of the invention also relate to an apparatus for performing the operations herein. Such a computer program is stored in a non-transitory computer readable medium. A machine-readable medium includes any mechanism for storing information in a form readable by a machine (e.g., a computer). For example, a machine-readable (e.g., computer-readable) medium includes a machine (e.g., a computer) readable storage medium (e.g., read only memory (“ROM”), random access memory (“RAM”), magnetic disk storage media, optical storage media, flash memory devices).

[0088] The processes or methods depicted in the preceding figures may be performed by processing logic that comprises hardware (e.g., circuitry, dedicated logic, etc.), software (e.g., embodied on a non-transitory computer readable medium), or a combination of both. Although the processes or methods are described above in terms of some sequential operations, it should be appreciated that some of the operations described may be performed in a different order. Moreover, some operations may be performed in parallel rather than sequentially.

[0089] Embodiments of the present invention are not described with reference to any programming language. It will be appreciated that a variety of programming languages may be used to implement the teachings of embodiments of the invention as described herein.

[0090] In the foregoing specification, embodiments of the invention have been described with reference to specific exemplary embodiments thereof. It will be evident that various modifications may be made thereto without departing from the broader spirit and scope of the invention as set forth in the following claims. The specification and drawings are, accordingly, to be regarded in an illustrative sense rather than a restrictive sense.

Claims

1. A computer-implemented method of generating real-time insights during a voice over internet protocol (VOIP) conversation, the method comprising: receiving, at a server, a transcript of one or more voice utterances of a participant in the VoIP conversation; identifying, by the server, a context of the VoIP conversation and a first state of the context based on the transcript; receiving, by the server, an intent of the participant from a conversation artificial intelligence (AI) engine, wherein the intent is identified by the conversation AI engine based on a sequence of words in the transcript; determining, by the server, one or more insights based on the context of the VoIP conversation, the first state of the context, and the intent of the participant; transmitting, by the server, the one or more insights to a terminal device of each of one or more participants of the VoIP conversation, wherein the one or more insights are displayed in real time on the terminal device of each of the one or more participants; and in response to determining that a property of the VoIP conversation has changed, updating the first state of the context to a second state of the context.
2. The method of claim 1, wherein each of the one or more participants is associated with one or more predetermined attributes.
3. The method of claim 2, wherein at least one participant is associated with one or more predetermined attributes that are different than one or more attributes of at least one of remaining ones of the one or more participants.
4. The method of claim 1, wherein the context of the VoIP conversation is represented by a data object, wherein the method further comprises adding one or more additional properties to the data object based on the one or more voice utterances.
5. The method of claim 4, wherein the data object includes one or more of following properties: participants of the VOIP conversation, a time period that the VoIP conversation relates to, whether the participants include an external party, contact information of the external party, activities between the participants, whether a competitor is mentioned, or deal information among the participants.
6. The method of claim 4, wherein the first state of the context is changed to the second state, in response to adding the one or more properties to the data object or when a value of an existing property changes.
7. The method of claim 4, wherein the context of the VoIP conversation, a unique identifier of the VoIP conversation, and the first state and the second state of the context of the VoIP conversation are stored in a cache in a cloud environment, wherein the cache further includes a state machine that keeps tracks of each state of the context during the VoIP conversation.
8. The method of claim 7, wherein the intent generated by the conversation AI engine is further

refined based on the state machine.

9. The method of claim 4, further comprising formulating one or more queries to retrieve the one or more insights from one or more databases in the cloud environment.

10. The method of claim 9, wherein the one or more databases comprise: a task database that includes information related to a plurality of tasks; and a conversation database that includes a plurality of entries, each entry mapping an insight to a combination of a context, a state of the context, and an intent.

11. A non-transitory machine-readable medium having instructions stored therein, which when executed by a processor, cause the processor to perform a method of generating real-time insights during a voice over internet protocol (VOIP) conversation, the method comprising: receiving a transcript of one or more voice utterances of a participant in the VoIP conversation; identifying a context of the VoIP conversation and a first state of the context based on the transcript; receiving an intent of the participant from a conversation artificial intelligence (AI) engine, wherein the intent is identified by the conversation AI engine based on a sequence of words in the transcript; determining one or more insights based on the context of the VOIP conversation, the first state of the context, and the intent of the participant; transmitting the one or more insights to a terminal device of each of one or more participants of the VoIP conversation, wherein the one or more insights are displayed in real time on the terminal device of each of the one or more participants; and in response to determining that a property of the VoIP conversation has changed, updating the first state of the context to a second state of the context.

12. The machine-readable medium of claim 11, wherein each of the one or more participants is associated with one or more predetermined attributes.

13. The machine-readable medium of claim 12, wherein at least one participant is associated with one or more predetermined attributes that are different than one or more attributes of at least one of remaining ones of the one or more participants.

14. The machine-readable medium of claim 11, wherein the context of the VOIP conversation is represented by a data object, wherein the method further comprises adding one or more additional properties to the data object based on the one or more voice utterances.

15. The machine-readable medium of claim 14, wherein the data object includes one or more of following properties: participants of the VoIP conversation, a time period that the VOIP conversation relates to, whether the participants include an external party, contact information of the external party, activities between the participants, whether a competitor is mentioned, or deal information among the participants.

16. The machine-readable medium of claim 14, wherein the first state of the context is changed to the second state, in response to adding the one or more properties to the data object or when a value of an existing property changes.

17. The machine-readable medium of claim 14, wherein the context of the VOIP conversation, a unique identifier of the VoIP conversation, and the first state and the second state of the context of the VoIP conversation are stored in a cache in a cloud environment, wherein the cache further includes a state machine that keeps tracks of each state of the context during the VoIP conversation.

18. The machine-readable medium of claim 17, wherein the intent generated by the conversation AI engine is further refined based on the state machine.

19. A data processing system, comprising: a processor; and a memory storing instructions, which when executed by the processor, cause the processor to perform a method of generating real-time insights during a voice over internet protocol (VOIP) conversation, the method comprising: receiving a transcript of one or more voice utterances of a participant in the VoIP conversation; identifying a context of the VoIP conversation and a first state of the context based on the transcript; receiving an intent of the participant from a conversation artificial intelligence (AI) engine, wherein the intent is identified by the conversation AI engine based on a sequence of words in the transcript; determining one or more insights based on the context of the VoIP conversation,

the first state of the context, and the intent of the participant; transmitting the one or more insights to a terminal device of each of one or more participants of the VoIP conversation, wherein the one or more insights are displayed in real time on the terminal device of each of the one or more participants; and in response to determining that a property of the VoIP conversation has changed, updating the first state of the context to a second state of the context.

20. The system of claim 19, wherein each of the one or more participants is associated with one or more predetermined attributes.
