

# US Patent & Trademark Office

## Patent Public Search | Text View

United States Patent Application Publication

20250260581

Kind Code

A1

Publication Date

August 14, 2025

Inventor(s)

Struttmann; Christopher Edward et al.

### USING A TREE STRUCTURE TO SEGMENT AND DISTRIBUTE RECORDS ACROSS ONE OR MORE DECENTRALIZED, ACYCLIC GRAPHS OF CRYPTOGRAPHIC HASH POINTERS

#### Abstract

Provided is a process including: receiving, with one or more processors, a first request to store a record from a computing entity; encoding, with one or more processors, the record in a first plurality of segments; arranging, with one or more processors, the first plurality of segments in respective content nodes of a first content graph, wherein at least some content nodes of the first content graph have two or more content edges of the first content graph pointing to two or more respective other content nodes of the first content graph; and storing, with one or more processors, the content nodes of the first content graph in a verification graph.

**Inventors:** Struttmann; Christopher Edward (Indialantic, FL), Beecham; James Douglas (Austin, TX)

**Applicant:** ALTR Solutions, Inc. (Austin, TX)

**Family ID:** 1000008575180

**Appl. No.:** 19/196463

**Filed:** May 01, 2025

#### Related U.S. Application Data

parent US continuation 18180078 20230307 parent-grant-document US 12301724 child US 19196463

parent US continuation 17584088 20220125 parent-grant-document US 11637706 child US 18180078

parent US continuation 17143009 20210106 parent-grant-document US 11265171 child US 17584088

parent US continuation 16669337 20191030 parent-grant-document US 10911241 child US 17143009

parent US continuation 16211057 20181205 parent-grant-document US 10505740 child US 16669337  
parent US continuation 15917616 20180310 parent-grant-document US 10193696 child US 16211057  
parent US continuation-in-part 15675519 20170811 parent-grant-document US 10121019 child US 15917616  
parent US continuation-in-part 15171347 20160602 parent-grant-document US 10581977 child US 15675519  
us-provisional-application US 62469742 20170310  
us-provisional-application US 62374307 20160812  
us-provisional-application US 62169823 20150602

---

## Publication Classification

**Int. Cl.:** **H04L9/32** (20060101); **G06F16/901** (20190101); **G06F21/60** (20130101); **G06F21/64** (20130101); **G06F21/78** (20130101); **H04L9/00** (20220101); **H04L9/06** (20060101)

**U.S. Cl.:**

**CPC** **H04L9/3242** (20130101); **G06F16/9024** (20190101); **G06F21/602** (20130101); **G06F21/64** (20130101); **G06F21/78** (20130101); **H04L9/0637** (20130101); **H04L9/3239** (20130101); H04L9/50 (20220501)

---

## Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS [0001] This patent application is a continuation of U.S. patent application Ser. No. 18/180,078, titled USING A TREE STRUCTURE TO SEGMENT AND DISTRIBUTE RECORDS ACROSS ONE OR MORE DECENTRALIZED, ACYCLIC GRAPHS OF CRYPTOGRAPHIC HASH POINTERS, filed Mar. 7, 2023, which is a continuation of U.S. patent application Ser. No. 17/548,088, titled USING A TREE STRUCTURE TO SEGMENT AND DISTRIBUTE RECORDS ACROSS ONE OR MORE DECENTRALIZED, ACYCLIC GRAPHS OF CRYPTOGRAPHIC HASH POINTERS, filed 25 Jan. 2022, which is a continuation of U.S. patent application Ser. No. 17/143,009, having the same title, filed 6 Jan. 2021, issued as U.S. Pat. No. 11,265,171, which is a continuation of U.S. patent application Ser. No. 16/669,337, having the same title, filed 30 Oct. 2019, issued as U.S. Pat. No. 10,911,241, which is a continuation of U.S. patent application Ser. No. 16/211,057, having the same title, filed 5 Dec. 2018, issued as U.S. Pat. No. 10,505,740, which is a continuation of U.S. patent application Ser. No. 15/917,616, titled USING A TREE STRUCTURE TO SEGMENT AND DISTRIBUTE RECORDS ACROSS ONE OR MORE DECENTRALIZED, ACRYLIC GRAPHS OF CRYPTOGRAPHIC HASH POINTERS, filed 10 Mar. 2018, issued as U.S. Pat. No. 10,193,696, which claims the benefit of U.S. Provisional Patent Application 62/469,742, titled UTILIZING A TREE-STRUCTURE TO SEGMENT AND DISTRIBUTE FILES ACROSS A SERIES OF BLOCKCHAINS, filed 10 Mar. 2017; and U.S. patent application Ser. No. 15/917,616 is a continuation-in-part of U.S. patent application Ser. No. 15/675,519, titled STORING DIFFERENTIALS OF FILES IN A DISTRIBUTED BLOCKCHAIN, filed 11 Aug. 2017, issued as U.S. Pat. No. 10,121,019, which claims the benefit of U.S. Provisional Patent Application 62/374,307, titled STORING DIFFERENTIALS OF FILES IN A DISTRIBUTED BLOCKCHAIN, filed 12 Aug. 2016; and U.S. patent application Ser. No. 15/675,519 is a continuation-in-part of U.S. patent application Ser. No. 15/171,347, titled COMPUTER

SECURITY AND USAGE-ANALYSIS SYSTEM, filed 2 Jun. 2016, issued as U.S. Pat. No. 10,581,977, which claims the benefit of U.S. Provisional Patent Application 62/169,823, titled PRIVATE AND SECURE AUTHENTICATION VIA A COMPUTER NETWORK, BOTH THE PUBLIC INTERNET OR PRIVATE INTRANET WITHOUT PASSING PERSONAL IDENTIFYING INFORMATION THAT DETECTS AND RESPONDS TO INTRUDERS AND/OR MALICIOUS BEHAVIOR, filed 2 Jun. 2015. The entire content of each aforementioned patent filing is hereby incorporated by reference.

## BACKGROUND

### 1. Field

[0002] The present disclosure relates generally to cybersecurity and, more, specifically to using a tree structure to segment and distribute records across one or more decentralized, acyclic graphs of cryptographic hash pointers.

### 2. Description of the Related Art

[0003] Datastores, such as document repositories, file systems, relational databases, non-relational database, memory images, key-value repositories, and the like, are used in a variety of different types of computing systems. Often, data to be stored is received by the datastore and then later retrieved during a read operation. In many cases, the datastore arranges the data in a manner that facilitates access based on an address of the data in the datastore (e.g., a file name) or content of the data (e.g., a select statement in a structured query language query).

[0004] In many cases, the security and integrity of the data in the datastore cannot be trusted. Often, an attacker who has penetrated a computer network will modify or exfiltrate records in a datastore that are intended to be confidential. Further, in many cases, the attacker may be a credentialed entity within a network, such as a rogue employee, making many traditional approaches to datastore security inadequate in some cases. Aggravating the risk, in many cases, such an attacker may attempt to mask their activity in a network by deleting access logs stored in datastores.

[0005] Often, attempts to mitigate the above-techniques result in data structures that, while potentially more secure, impose relatively high latency during read or write operations. For instance, retrieving individual portions of a record serially from a blockchain may impose relatively high latency as segment-by-segment serial retrieval may, in the aggregate, be relatively slow as time to retrieve each segment accumulates.

## SUMMARY

[0006] The following is a non-exhaustive listing of some aspects of the present techniques. These and other aspects are described in the following disclosure.

[0007] Some aspects include a process including: receiving, with one or more processors, a first request to store a record from a computing entity; encoding, with one or more processors, the record in a first plurality of segments; arranging, with one or more processors, the first plurality of segments in respective content nodes of a first content graph, wherein at least some content nodes of the first content graph have two or more content edges of the first content graph pointing to two or more respective other content nodes of the first content graph; and storing, with one or more processors, the content nodes of the first content graph in a verification graph.

[0008] Some aspects include a tangible, non-transitory, machine-readable medium storing instructions that when executed by a data processing apparatus cause the data processing apparatus to perform operations including the above-mentioned process.

[0009] Some aspects include a system, including: one or more processors; and memory storing instructions that when executed by the processors cause the processors to effectuate operations of the above-mentioned process.

---

## Description

### BRIEF DESCRIPTION OF THE DRAWINGS

[0010] The above-mentioned aspects and other aspects of the present techniques will be better understood when the present application is read in view of the following figures in which like numbers indicate similar or identical elements:

[0011] FIG. 1 is a logical and physical architecture block diagram that shows an example of a computing environment in which the present techniques may be implemented in accordance with some embodiments;

[0012] FIG. 2 is a flow chart that shows an example of a process to read and write a document to a tamper-evident, immutable data repository in accordance with some embodiments;

[0013] FIG. 3 is a data model block diagram that shows an example of a tamper-evident, immutable data repository that may be operated upon by the process of FIG. 2 in accordance with some embodiments;

[0014] FIG. 4 is a flow chart that shows an example of a process by which stored values may be fragmented and stored among a plurality of different tamper-evident, immutable directed acyclic graphs, such as those shown in FIG. 3, in accordance with some embodiments;

[0015] FIG. 5 is a data model block diagram that shows an example of a tamper-evident, immutable data repository storing fragments of data distributed among multiple directed acyclic graphs that may be operated upon by the processes of FIGS. 4 and 6 in accordance with some embodiments;

[0016] FIG. 6 is a flow chart that shows an example of a process by which data may be read from the data structure of FIG. 5 in accordance with some embodiments;

[0017] FIG. 7 is a flow chart that shows an example of a process by which cryptographic hash pointers in tamper-evident data repositories may be formed to render the data repositories modular in accordance with some embodiments;

[0018] FIG. 8 is a flow chart that shows an example of a process by which existing workload applications and user interfaces may be retrofit to securely store and access relatively high-security data in a transparent fashion in accordance with some embodiments;

[0019] FIG. 9 is a flow chart that shows an example of a process by which access requests to scattered data may be logged to a tamper-evident, immutable data repository in accordance with some embodiments;

[0020] FIG. 10 is a flow chart that shows an example of a process by which data in tamper-evident, immutable data repositories may be compressed by storing differences between versions of documents or other units of content in accordance with some embodiments;

[0021] FIG. 11 is a flow chart that shows an example of a process by which version graphs storing sequences of changes between versions of documents may be traversed to read current versions of documents in accordance with some embodiments;

[0022] FIG. 12 is a data model block diagram that shows an example of a data structure by which a version graph may be stored as a sequence of content graphs on a collection of verification graphs or other tamper-evident, immutable data repositories in accordance with some embodiments;

[0023] FIG. 13 is a physical architecture block diagram that shows an example of a computing device by which the present techniques may be implemented;

[0024] FIG. 14 is a flow chart that shows an example of a process by which data may be stored and read in a manner that facilitates concurrent retrieval of subsets of a record to expedite read operations;

[0025] FIG. 15 is a diagram of an example of a content graph that facilitates concurrent retrieval of subsets of a record to expedite read operations and reduces duplicative writes of portions of records that do not change between versions to expedite write operations;

[0026] FIG. 16 is a diagram of an example of the content graph of FIG. 15 after a new version of a record is stored; and

[0027] FIG. 17 is a diagram of an example of the content graph of FIG. 16 after a third version of a record is stored.

[0028] While the present techniques are susceptible to various modifications and alternative forms, specific embodiments thereof are shown by way of example in the drawings and will herein be described in detail. The drawings may not be to scale. It should be understood, however, that the drawings and detailed description thereto are not intended to limit the present techniques to the particular form disclosed, but to the contrary, the intention is to cover all modifications, equivalents, and alternatives falling within the spirit and scope of the present techniques as defined by the appended claims.

#### DETAILED DESCRIPTION OF CERTAIN EMBODIMENTS

[0029] To mitigate the problems described herein, the inventors had to both invent solutions and, in some cases just as importantly, recognize problems overlooked (or not yet foreseen) by others in the field of cybersecurity. Indeed, the inventors wish to emphasize the difficulty of recognizing those problems that are nascent and will become much more apparent in the future should trends in industry continue as the inventors expect. Further, because multiple problems are addressed, it should be understood that some embodiments are problem-specific, and not all embodiments address every problem with traditional systems described herein or provide every benefit described herein. That said, improvements that solve various permutations of these problems are described below.

[0030] A variety of problems relating to security of datastores and networks of computers used by organizations are addressed by various versions of techniques described below. These different techniques can be used together, synergistically in some cases, so their descriptions are grouped into a single description that will be filed in multiple patent applications with different claim sets targeting the different techniques and combinations thereof. In view of this approach, it should be emphasized that the techniques are also independently useful and may be deployed in isolation from one another or in any permutation combining the different subsets of techniques, none of which to suggest that any other description herein is limiting. Conceptually related groups of these techniques are preceded by headings below. These headings should not be read as suggesting that the subject matter underneath different headings may not be combined, that every embodiment described under the heading has all of the features of the heading, or that every feature under a given heading must be present in an embodiment consistent with the corresponding conceptually related group of techniques, again which is not to suggest that any other description is limiting.

[0031] These techniques are best understood in view of an example computing environment **10** shown in FIG. 1. The computing environment **10** is one example of many computing architectures in which the present techniques may be implemented. In some embodiments, the present techniques are implemented as a multi-tenant distributed application in which some computing hardware is shared by multiple tenants that access resources on the computing hardware in computing devices controlled by those tenants, for example, on various local area networks operated by the tenants. Or in some cases, a single tenant may execute each of the illustrated computational entities on privately-controlled hardware, with multiple instances of the computing environment **10** existing for different organizations. Or some embodiments may implement a hybrid approach in which multi-tenant computing resources (e.g., computers, virtual machines, containers, microkernels, or the like) are combined with on-premises computing resources or private cloud resources. In some embodiments, the computing environment **10** may include and extend upon the security features of a computing environment described in U.S. patent application Ser. No. 15/171,347, titled COMPUTER SECURITY AND USAGE-ANALYSIS SYSTEM (docket no. 043788-0447379), filed 2 Jun. 2016, the contents of which are hereby incorporated by reference.

[0032] In some embodiments, the computing environment **10** includes a plurality of client computing devices **12**, a lower-trust database **14**, secure distributed storage **16**, a domain name service **18**, and a translator server **20** (or elastically scalable collection of instances of translator servers disposed behind a load balancer). In some embodiments, each of these components may communicate with one another via the Internet **22** and various local area networks in some cases. In some embodiments, communication may be via virtual private networks overlaid on top of the public Internet. In some embodiments, the illustrated components may be geographically distributed, for example, more than 1 kilometer apart, more than 100 kilometers apart, more than a thousand kilometers apart, or further, for example distributed over the content event of North America, or the world. Or in some cases, the components may be co-located and hosted within a airgapped or non-airgapped private network. In some embodiments, each of the illustrated blocks that connects to the Internet **22** may be implemented with one or more of the computing devices described below with reference to FIG. **13**.

[0033] In some embodiments, each of the client computing devices **12** may be one of a plurality of computing devices operated by users or applications of an entity that wishes to securely store data. For example, a given business or governmental organization may have more than 10, more than 100, more than 1,000, or more than 10,000 users and applications, each having associated computing devices that access data stored in the lower-trust database **14** (or a collection of such databases or other types of datastores) and the secure distributed storage **16**. In some embodiments, multiple entities may access the system in the computing environment **10**, for example more than five, more than 50, more than 500, or more than 5000 different entities may access shared resources with respective client computing devices or may have their own instance of the computing environment **10**. In some embodiments, some of the client computing devices **12** are end-user devices, for example, executing a client-side component of a distributed application that stores data in the lower-trust database **14** and the secure distributed storage **16**, or reads is such data. Client computing devices may be laptops, desktops, tablets, smartphones, or rack-mounted computing devices, like servers. In some embodiments, the client-computing devices are Internet-of-things appliances, like smart televisions, set-top media payers, security cameras, smart locks, self-driving cars, autonomous drones, industrial sensors, industrial actuators (like electric motors), or in-store kiosks. In some embodiments, some of the client computing devices **12** may be headless computing entities, such as containers, microkernels, virtual machines, or rack-mounted servers that execute a monolithic application or one or more services in a service-oriented application, like a micro services architecture, that stores or otherwise access data in the lower-trust database **14** or the secure distributed storage **16**.

[0034] In some embodiments, the lower-trust database **14** and the secure distributed storage **16** may each store a portion of the data accessed with the client computing devices **12**, in some cases with pointers therebetween stored in one or both of these datastores. In some embodiments, as described below, this data may be stored in a manner that abstracts away the secure distributed storage **16** from a workload application through which the data is accessed (e.g., read or written). In some embodiments, data access operations may store or access data in the lower-trust database **14** and the secure distributed storage **16** with a workload application that is not specifically configured to access data in the secure distributed storage **16**, e.g., one that is configured to operate without regard to whether the secure distributed storage **16** is present, and for which the storage of data in the secure distributed storage **16** is transparent to the workload application storing content in the lower-trust database **14** and the secure distributed storage **16**. In some embodiments, such a workload application may be configured to, and otherwise designed to, interface only with the lower-trust database **14** when storing this data, and as described below, some embodiments may wrap interfaces for the lower-trust database **14** with additional logic that routes some of the data to the secure distributed storage **16** and retrieves that data from the secure distributed storage **16** in a manner that is transparent to the workload application accessing content (i.e., data written or read

by the workload application).

[0035] Content stored in the lower-trust database **14** and secure distributed storage **16** may be created or accessed with a variety of different types of applications, such as monolithic applications or multi-service distributed applications (e.g., implementing a microservices architecture in which each service is hosted by one of the client computing devices **12**). Examples include email, word processing systems, spreadsheet applications, version control systems, customer relationship management systems, human resources computer systems, accounting systems, enterprise resource management systems, inventory management systems, logistics systems, secure chat computer systems, industrial process controls and monitoring, trading platforms, banking systems, and the like. Such applications that generate or access content in the database **14** for purposes of serving the application's functionality are referred to herein as “workload applications,” to distinguish those applications from infrastructure code by which the present techniques are implemented, which is not to suggest that these bodies of code cannot be integrated in some embodiments into a single workload application having the infrastructure functionality. In some cases, several workload applications (e.g., more than 2, more than 10, or more than 50), such as selected among those in the preceding list, may share resources provided by the infrastructure code and functionality described herein.

[0036] In some embodiments, the lower-trust database **14** is one of the various types of datastores described above. In some cases, the lower-trust database **14** is a relational database, having a plurality of tables, each with a set of columns corresponding to different fields, or types of values, stored in rows, or records (i.e., a row in some implementations) in the table, in some cases, each record, corresponding to a row may be a tuple with a primary key that is unique within that respective table, one or more foreign keys that are primary keys in other tables, and one or more other values corresponding to different columns that specify different fields in the tuple. Or in some cases, the database may be a column-oriented database in which records are stored in columns, with different rows corresponding to different fields. In some embodiments, the lower-trust database **14** may be a relational database configured to be accessed with structured query language (SQL) commands, such as commands to select records satisfying criteria specified in the command, commands to join records from multiple tables, or commands to write values to records in these tables.

[0037] Or in some cases, the lower-trust database **14** may be another type of database, such as a noSQL database, like various types of non-relational databases. In some embodiments, the lower-trust database **14** is a document-oriented database, such as a database storing a plurality of serialized hierarchical data format documents, like JavaScript™ object notation (JSON) documents, or extensible markup language (XML) documents. Access requests in some case may take the form of xpath or JSON-path commands. In some embodiments, the lower-trust database **14** is a key-value data store having a collection of key-value pairs in which data is stored. Or in some cases, the lower-trust database **14** is any of a variety of other types of datastores, for instance, such as instances of documents in a version control system, memory images, a distributed or non-distributed file-system, or the like. A single lower-trust database **14** is shown, but embodiments are consistent with, and in commercial instances likely to include, substantially more, such as more than two, more than five, or more than 10 different databases, in some cases of different types among the examples described above. In some embodiments, some of the lower-trust databases may be database of a software-as-a-service application hosted by a third party and accessed via a third party application program interface via exchanges with, for instance, a user's web browser or another application. In some cases, the lower-trust database **14** is a mutable data store or an immutable data store.

[0038] In some cases, access to data in the lower-trust database **14**, and corresponding access to corresponding records in the secure distributed storage **16**, may be designated in part with roles and permissions stored in association with various user accounts of an application used to access that

data. In some embodiments, these permissions may be modified, for example, revoked, or otherwise adjusted, with the techniques described in U.S. patent application Ser. No. 15/171,347, titled COMPUTER SECURITY AND USAGE-ANALYSIS SYSTEM (docket no. 043788-0447379), filed 2 Jun. 2016, the contents of which are hereby incorporated by reference.

[0039] The database **14** is described as “lower-trust.” The term “lower-trust” does not require an absolute measure of trust or any particular state of mind with respect to any party, but rather serves to distinguish the database **14** from the secure distributed storage **16** which has certain security features in some implementations described below and in some cases may be referred to as a “higher-trust” database.

[0040] In some cases, some of the data that an application writes to, or has written to, the lower-trust database **14** may be intercepted or moved to the secure distributed storage **16** with techniques described below. Further, access requests from a workload application to the lower-trust database **14** may be intercepted, or responses from such access request may be intercepted, and data from the lower-trust database **14** may be merged with data from the secure distributed storage **16** that is responsive to the request before being presented to the application, as described in greater detail below. Further, read requests may be intercepted, modified, and iteratively executed in a manner that limits how much information in the secure distributed storage is revealed to a client computing device at any one time, as described below.

[0041] In some embodiments, the secure distributed storage **16** may include a collection of data centers **24**, which may be distributed geographically and be of heterogeneous architectures. In some embodiments, the data centers **24** may be various public or private clouds or on-premises data centers for one or more organization-users, such as tenants, of the computing environment **10**. In some embodiments, the data centers **24** may be geographically distributed over the United States, North America, or the world, in some cases with different data centers more than 100 or 1,000 kilometers apart, and in some cases with different data centers **24** in different jurisdictions. In some embodiments, each of the data centers **24** may include a distinct private subnet through which computing devices, such as rack-mounted computing devices in the subnet communicate, for example, via wrap top-of-rack switches within a data center, behind a firewall relative to the Internet **22**. In some embodiments, each of the data centers **24**, or different subsets of the data centers **24**, may be operated by a different entity, implementing a different security architecture and having a different application program interface to access computing resources, examples including Amazon Web Services™, Azure from Microsoft™, and Rack Space™. Three different data centers **24** are shown, but embodiments are consistent with, and in commercial implementations likely to include, more data centers, such as more than five, more than 15, or more than 50. In some cases, the datacenters may be from the same provider but in different regions.

[0042] In some embodiments, each of the data centers **24** includes a plurality of different hosts exposed by different computational entities, like microkernels, containers, virtual machines, or computing devices executing a non-virtualized operating system. Each host may have an Internet Protocol address on the subnet of the respective data center **24** and may listen to and transmit via a port assigned to an instance of an application described below by which data is stored in a distributed ledger. In some embodiments, each storage compute node **26** may correspond to a different network hosts, each network coast having a server that monitors a port, and configured to implement an instance of one of the below-described directed acyclic graphs with hash pointers implementing immutable, tamper-evident distributed ledgers, examples include block chains and related data structures. In some cases, these storage compute nodes **26** may be replicated, in some cases across data centers **24**, for example, with three or more instances serving as replicated instances, and some embodiments may implement techniques described below to determine consensus among these replicated instances as to state of stored data. Further, some embodiments may elastically scale the number of such instances based on amount of data stored, amounts of access requests, or the like.



[0043] Some embodiments may further include a domain name service (DNS) **18**, such as a private DNS that maps uniform resource identifiers (such as uniform resource locators) to Internet Protocol address/port number pairs, for example, of the storage compute nodes **26**, the translator **20**, and in some cases other client computing devices **12** or other resources in the computing environment **10**. In some embodiments, a client computing device **12**, a storage compute node **16**, the database **14**, or translator **20** may encounter a uniform resource identifier, such as a uniform resource locator, and that computing entity may be configured to access the DNS **18** at an IP address and port number pair of the DNS **18**. The entity may send a request to the DNS **18** with the uniform resource identifier, and the DNS **18** may respond with a network and process address, such as Internet Protocol address and port number pair corresponding to the uniform resource identifier. As a result, underlying computing devices may be replaced, replicated, moved, or otherwise adjusted, without impairing cross-references between information stored on different computing devices. Or some embodiments may achieve such flexibility without using a domain name service **18**, for example, by implementing a distributed hash table or load-balancing that consistently maps data based on data content, for example based on a prefix or suffix of a hash based on the data or identifiers of data to the appropriate computing device or host. For instance, some embodiments may implement a load balancer that routes requests to storage compute nodes **26** based on a prefix of a node identifier, such as a preceding or trailing threshold number of characters.

[0044] Some embodiments may further include a virtual machine or container manager configured to orchestrate or otherwise elastically scale instances of compute nodes and instances of the translator **20**, for instance, automatically applying corresponding images to provisioned resources within one or more data centers **24** responsive to need and spinning down instances as need diminishes.

[0045] In some embodiments, the translator **20** may be configured to execute a routine described in greater detail below that translates between an address space of the lower-trust database **14** and an address space of the secure distributed storage **16**. In some embodiments, the translator **20** may receive one or more records from the client computing device **12** that is going to be written to the lower-trust database **14**, or may receive such records from the lower-trust database **14**, and those records may be mapped to the below-describe segment identifiers (or other pointers, such as other node identifiers) in the secure distributed storage **16**. The translator **20** may then cause those records to be stored in the secure distributed storage **16** and the segment identifiers to be stored in place of those records in the lower-trust database **14**, such as in place of individual values in records. In some embodiments, translation may happen at the level of individual values corresponding to individual fields in individual records, like rows of a table in the database **14**, or some embodiments may translate larger collections of data, for example, accepting entire records, like entire rows, or plurality of columns, like a primary key and an individual value other than the primary key in a given row. Some embodiments may accept files or other binary larger objects (BLOBS). The translator **20** that may then replace those values in the lower-trust database **14** with a pointer, like a segment identifier in the secure distributed storage, in the manner described below, and then cause those that data to be stored in the secure distributed storage **16** in the manner described below. In some examples, documents may be stored, which may be relatively small stand-alone values to binary large objects encoding file-system objects like word-processing files, audio files, video files, chat logs, compressed directories, and the like. In some cases, a document may correspond to an individual value within a database, or document may correspond to a file or other binary large object. In some cases, documents may be larger than one byte, 100 bytes, 1 kB, 100 kB, 1 MB, or 1 GB. In some embodiments, documents may correspond to messages in a messaging system, or printable document format documents, Microsoft Word™ documents, audio files, video files or the like.

[0046] In some embodiments, the translator **20** may include code that receives requests from drivers and facilitates the translation of data. In some cases, the translator **20** may be one of an

elastically scaled set of translators **20** remotely hosted in a public or private cloud. The translator may, in some cases, implement the following functions:

#### 1. Validate Request

[0047] a. Using a database, some embodiments validate a combination of user supplied parameters such as predefined software IDs, client IDs, and machine specific identifiers registered at install time. This is compared against a known list and then further verified with IP address and/or other network specific parameters.

#### 2. Data Validate

[0048] a. Parsing the HTTP body and then decoding some embodiments determine the unique list of reference values to replace with plain text. Using a database, some embodiments first check if the requesting machine has the rights to access the data. Next using a database, some embodiments find the network name of the first hop of the piece of data and place into an array.

#### 3. Threshold Check

[0049] a. With the location of each unique requested segment (or node or document or content) identifier, some embodiments check against a series of threshold or rate objects. Some embodiments look for access rate, time window, or location based rules and apply the requested data against a mapping of rules. If any particular data is breaking a threshold then an anomaly in the system is generated resulting in notifications and logging in some embodiments.

#### 4. Jobs

[0050] a. The translator **20** may split up the data requests into jobs and places the job onto a work queue. The split may be done by a static per message job size and may use a deal-letter exchange to retry and finally fail messages

#### 5. Response Function

[0051] a. Data may be returned from the queue and plain text values may be matched and replaced with the corresponding pointers (such as segment, document, node, or unit-of-content identifiers, which is not to suggest that these or any other list of categories describe disjoint sets). Once all jobs have returned the response a response may be returned in some embodiments.

[0052] In some embodiments, the client computing devices **12** may each execute an operating system in which one or more applications **28** execute. These applications may include client-side portions of the above-described examples of workload applications, which may include business logic and other program code by which a service in a micro-services architecture is implemented. In some embodiments, the applications **28** may be different in different client computing devices, and an individual client computing device may execute a plurality of different applications. In some embodiments, the applications **28** may be configured to interface with the lower-trust database **14** via a database driver **32** executed within the operating system. The database driver **32** may be any of a variety of different types of drivers such as an ODBC driver, a JDBC driver, and the like. In some embodiments, the database driver **32** may be configured to access the lower-trust database **14** via a network interface **34** of the client computing device **12**, such as a network interface card connected to a physical media of a local area network by which the Internet **22** is accessed.

[0053] Some embodiments may further include a security driver **30** that interfaces between the application **28** and the database driver **32**. In some embodiments, the security driver **30** may be transparent to the application **28**, such that an application program interface of the database driver **32** is presented to the application **28** by the security driver **30**, and that application program interface may be unmodified from the perspective of the application **28** relative to that presented by the database driver **32** in some cases. In some embodiments, the security driver **30** may wrap an application program interface of the database driver **32**, such that the security driver **30** receives application program interface requests from the application **28** to the driver **32**, acts on those requests, and in some cases modifies those requests, and then provides the request in some cases with modifications to the database driver **32**. Similarly, responses back to the application **28** may be

provided by the security driver **30** and in a manner consistent with that provided by the driver **32**, as described in greater detail below.

[0054] In some embodiments, the security driver **30** is configured to engage the translator **20** after (or to perform) splitting data being written to (or attempting) the lower-trust database **14** by the application **28** into higher-security data and lower-security data. Again, the terms “lower-security” and “higher-security” serve to distinguish data classified differently for purposes of security and do not require measurement against an absolute security metric or a state of mind. The lower-security data may then be written by the database driver **32** to the lower-trust database **14** in the manner provided for by the application **28** without regard to whether the security driver **30** is present.

[0055] The higher-security data, on the other hand, may be stored in a manner described below by the translator **20** that renders that data relatively robust to attacks by malicious actors. When returning data to the application **28**, for example in response to receiving a read request, these operations may be reversed in some cases. Again, these operations are described in greater detail below. Generally, in some embodiments, the data from the lower-trust database **14** and the data from the secure distributed storage **16** may be merged by the security driver **30**, in some cases, before that data is presented to the application **28**. By acting on the higher-security data within the client computing device **12**, before that data leaves the client computing device **12**, some embodiments may reduce an attack service of the computing environment **10**. That said, not all embodiments provide this benefit, and some embodiments may implement the functionality of the security driver **30** outside of the client computing devices **12**, for example, in a database gateway, in a database management system implemented at the lower-trust database **14**, or on another standalone application executed in a computing device disposed between the lower-trust database **14** and the network and the client computing device **12** in a path to the lower-trust database **14**.

[0056] In some embodiments, the security driver **30** includes an outbound path and an inbound path. In some embodiments, the outbound path includes an out-parser **36**, a validator **38**, a data multiplexer **40**. The out-parser may classify values as higher-security or lower-security values applying one or more rules in a data policy described below. The validator may perform the statement validate function described below. The multiplexer may route data to the lower-trust database **14** or the translator **20** based on the security classification. In some embodiments, the inbound path includes an in parser **42**, and a data de-multiplexer **44**. The inbound path may include a parser **42** configured to detect pointers to data in query responses from the lower-trust database **14** that point to data in the secure distributed storage **16**. The parser **42** may call the translator **20** to request that pointers be replaced with more securely stored data. In some cases, the de-multiplexer **44** may merge data from the translator **20** with lower-security data in the same query response. In some cases, the security driver may implement a process described below with reference to FIG. **8** and perform the following functions:

#### 1. Statement Parse

[0057] a. For a SELECT statement, there could be a WHERE clause which is looking to match data in a protected column. During this phase, some embodiments parse the SELECT statement and check if there is a need to flip any plain text values in the WHERE clause into the reference space. The statement may be marked for processing and passed along.

[0058] b. For an INSERT or UPDATE statement, there could be data in either the statement body or the WHERE clause (INSERT). During this phase, some embodiments parse the statement and check if there is a need to flip any plain text values in the WHERE clause or body into the reference space. The statement may be marked for processing and passed along.

[0059] c. The security driver may use a locally kept copy of the current protection settings for a given client. In some embodiments, it is this locally kept and updated (e.g., periodically or constantly) table that the database, table, and column names in the statements are compared against. The time between getting a new state table is determined by various factors.

#### 2. Statement Validate

[0060] a. During the operation of a database command some embodiments check the statement for potential injection or other malicious SQL statements and block the query or log that the event happened. This is a locally supported operation that can be done by each driver in some cases.

### 3. Statement Process

[0061] a. Depending upon the results of Parse, the driver may make HTTP requests to a preset URL and asks for plain-text data to be switched into the reference space, e.g., by the translator **20**.

[0062] b. The statement may be updated with reference space data if needed and the statement may be delivered to the lower-trust database **14** server.

### 4. Result Set Process

[0063] a. For a SELECT statement the result set is processed and if columns in the returned data match any entries in the locally held table, the security driver **20** may perform HTTP requests to switch reference space data to plain text space.

[0064] b. The driver **30** may iterate over the data and selects distinct values to place into an HTTP body and requests made using a preset URL and system DNS **18**, e.g., by engaging the translator **20**.

[0065] c. Data may be returned and replaced for each occurrence in the result set and returned to the application **28** in some cases.

[0066] Various aspects of the system above, or other architecture may implement various techniques expanded upon below under distinct headings.

### Immutable Datastore for Low-Latency Reading and Writing of Large Data Sets

[0067] Generally, traditional databases do not adequately protect against threat actors or internal resources (employees, information-technology staff, etc.) tampering with the data. At best, such systems typically provide audit access and the ability to modify the stored data, but the audit logs typically are mutable and, thus, can be changed just as easily as the data.

[0068] Recent immutable examples of databases include blockchain-based databases, such as bitcoind and MultiChain. Blockchain systems are built upon ideas first described in a paper titled “Bitcoin: A Peer-to-Peer Electronic Cash System” under the pseudonym Satoshi Nakamoto in October 2008. These systems typically implement a peer-to-peer system based on some combination of encryption, consensus algorithms, and proof-of-X, where X is some aspect that is difficult to consolidate across the network, such as proof-of-work, proof-of-stake, proof-of-storage, etc. Typically, those actors on a network having proof-of-X arrive at a consensus regarding the validation of peer-to-peer transactions, often using various consensus algorithms like Paxos, Raft, or hashgraph. Or some private blockchains do not implement proof-of-X consensus, e.g., where the computing hardware implementing the blockchain is controlled by trusted parties. Chained cryptographic operations tie a sequence of such transactions into a chain that once validated, is typically prohibitively computationally expensive to falsify.

[0069] However, many extant blockchain-based databases are not well suited for certain use cases, particularly those involving latency-sensitive access (e.g., reading or writing) to large files (e.g., documents or other collections of binary data treated as a single entity, often called “blobs”), for instance in a blockchain-hosted filesystem. Indeed, many blockchain databases are not readily configured to store large objects and object files (e.g., on the order of 500 kilobytes or larger, depending on the use case and acceptable latency), as such systems are typically highly specialized for small-payload “transactional” applications. In such systems, when storing larger collections of binary data (e.g., files or blobs), the chain can dramatically slow as the chain gets bigger, particularly for write operations.

[0070] As noted above, blockchains generally allow for small bits of information to be stored in an immutable data structure, but the moment data in the chain is altered, the chain is broken and can no longer function in a manner that represents the data to be valid. The most common blockchain implementation is the publicly accessible Bitcoin ledger (for which Blockchains were designed).

[0071] As noted, such systems present certain problems for some use cases. These, and other

problems are mitigated by some embodiments of a system referred to below as “Docuchain.” Docuchain is a blockchain software suite specifically designed for low-latency put and get operations of Binary Large Objects (BLOBs). Docuchain, in some embodiments, uses an improved version of an underlying data structure called “Merkle Trees” to provide the immutable properties of the Blockchain. In some embodiments, Docuchain is operative to respond to commands within less than 50 milliseconds (ms), e.g., and many commercial implementations are expected to provide sub 10 ms operations as perceived by the application interacting with the chain, and sub 5 ms operations at the level of the chain. Further, such response times are expected to scale with the chain size. Some embodiments may scale response times at  $n \log(n)$ , wherein  $n$  is the number of entries in the chain.

[0072] Merkle Trees generally work by a tree of cryptographic hashes in which each element of the tree contains a hash of the information contained by its children, and leaf elements are hashes of the subject data stored in the Merkle Tree. In many traditional implementations of Merkle Trees for the purposes of storing data, like those in many earlier blockchain databases, the data is stored in a separate logical datastore, hashed, and just the hash is carried into the Merkle Trees. That is, the data being by the database stored is not part of the Merkle Tree, only a hash digest of the data is part of the Merkle Tree.

[0073] To mitigate some of the problems with traditional blockchain databases, some embodiments of Docuchain store the data directly in Merkle Trees, though embodiments are not limited to data storage in Merkle Trees, which is not to suggest that other descriptions are limiting. That is, when data is written to the database or read from the database, that data is written into specific fields of the elements (e.g., attributes of node content of nodes) of the Merkle Tree or read from specific fields of the elements of the Merkle Tree (rather than just a hash digest of the data residing in the Merkle Tree with the entire data residing in an external datastore). With this tight coupling, if the data is altered at all, the entire hash for the tree (as is distinct from a hash of a file stored in the database) is thrown off immediately and ultimately the chain will be broken and detected as such during validation operations. Additionally, some embodiments prevent the BLOB from ever being stored in an external (potentially mutable and less secure) datastore. Finally, the tight coupling of the data to the tree is expected to reduce the number of read/write operations necessary to transact with the chain, further reducing potential latency.

[0074] In some embodiments, Docuchain contains two components: LibDocuchain and DocuSocket.

[0075] In some embodiments, LibDocuchain is a software library that contains the following data structures: [0076] MerkleTree<H, T> [0077] DataBlock<H, T> [0078] BlockManager<H, T> [0079] where “H” represents a cryptographic hashing function, and “T” represents the data type being stored in the chain. Generally, hashing functions map data of arbitrary size to data of fixed size. Examples of cryptographic hash functions include SHA256, SHA-2, MD5, and the like (e.g., applied more than once Docuchain, in some embodiments, can use a variety of different types of hashing functions and this is expected to include later developed hashing algorithms (e.g., new post-quantum cryptographic hashes).

[0080] The BlockManager, in some embodiments, is responsible for taking data, placing into the appropriate block and inserting into the appropriate Merkle tree, writing to disk at the appropriate times, increasing (e.g., guaranteeing) chain durability (i.e., tamper-resistance). Additionally, in some embodiments, there is a JobManager that manages concurrent operations on the datastructures to deliver operation times.

[0081] In some embodiments, DocuSocket is a software library that maintains connections to clients, e.g., remote computing devices seeking to access data stored in the database. Connections from clients may be relatively stateless (and in some cases, after a session is confirmed there is no further state). A connection, in some embodiments, can accept an arbitrary amount of requests in an arbitrary order and will return responses to the requests in arbitrary order. Example, the connection

might receive write (A), write (B), read (C) operations, and in some embodiments, the Docusocket may respond to request C before A or B. However, once a transaction response comes for A or B, those operations in some embodiments are considered final and are committed to the chain and are instantly queryable. This “always-on” and highly-available connection is one differentiator that allows Docusocket to out-perform standard blockchains in comparable benchmarks. (That said, not all embodiments afford these benefits, as various independently useful inventions are described, and some of those inventions might be used to other ends.)

[0082] Write requests, in some embodiments, are accepted through Docusocket and passed off to BlockManager for writing into a MerkleTree object. In some cases, the result of this operation is a hash composed from the position in the MerkleTree, the hash from the previous block, and the contents of the data that is written; this is called the Transaction ID (TXID), which is an example of a node identifier or pointer (which is not to suggest that these or other categories are disjoint). In some embodiments, these TXIDs are stored in alternate data stores for later referencing and retrieval, e.g., in the lower-trust database **14** in place of the data to which they point, or in a lower-trust file system in place of documents referenced by the TXIDs. In some cases, the TXID's are segment pointers to a linked list of segment pointers and segments that encode the stored data.

[0083] As noted, earlier blockchain databases are capable of storing blobs, but the way in which the data is stored and accessed often imposes severe performance penalties when accessing larger collections of data (e.g., larger BLOBs). Such systems often rely on additional functionality originally intended to handle the transactive nature of bitcoins and other cryptocurrencies. This additional functionality is often provided by an “in-chain” scripting language which defines the transactions that enter the chain. For many use cases, this is a secure and efficient method of maintaining a ledger, however for larger chains, these features come at significant time complexity cost, particularly in write operations. As these chains grow, they become significantly slower. Blockchain document databases that store the documents in-chain are expected to grow very fast, thus putting significant limitations on what is practically possible with prior technology.

[0084] In contrast, because some embodiments of Docuchain store data in its exact representation (which should not be read to exclude various transformations and may include storing a compressed, encrypted copy of the data) directly in the Merkle Tree, those embodiments are expected to be able to circumvent the need to have the in-chain scripting language and are able to provide  $O(\log(n))$  get/put operations on a tree once it is loaded in memory. Further, storing the data being committed to the database, rather than just a hash digest, in the Merkle Tree is expected to impede (and in some cases defeat) hash collision attacks on the stored data. In such attacks, malicious content is selected and designed to produce the same hash as stored data, and that malicious data is substituted in the database for the authentic data. With traditional systems, the Merkle Tree will yield a hash value that validates the malicious content as authentic. In contrast, some embodiments circumvent this attack vector by storing the data committed to the database in the Merkle Tree. Further, some embodiments may detect changes without the need to continually verify each BLOB with its respective hash digest on a continual basis, unlike many systems that merely store a hash digest of the BLOB in the chain. That said, not all embodiments afford these benefits, e.g., some embodiments may avoid the use of traditional in-chain scripts to access hashed digests in the chain, without storing the entire document, thereby expediting operations.

[0085] These libraries may implement an example of a process shown in FIG. **2**, which shows an example of a process **50** that may be implemented to store and access documents in a, immutable, tamper-evident data structure, an example of which is shown in FIG. **3**. In some embodiments, the process **50** may be implemented in the computing environment **10** described above, but it should be emphasized that embodiments are not limited to that implementation, which is not to suggest that any other description is limiting. Indeed, some embodiments may implement versions of the process shown on a single computing device as part of a monolithic application, such as a file manager that stores documents locally in the described fashion.

[0086] In some embodiments, the operations of the process **50**, and the other processes described herein, may be implemented in a different order from that described, may include steps that are repeated, may include steps that are executed concurrently, may include steps that are omitted, or may be otherwise differently arranged from the exemplary arrangements described, none of which is to suggest that other descriptions are limiting. Further, the operations of the process **50** and the other processes and functionality described herein may be implemented by executing instructions stored on a tangible, non-transitory, machine-readable medium with one or more processors, such that when the instructions are executed, the described functionality is effectuated. In some cases, notwithstanding the use of the singular medium, the operations may be stored on multiple media in a distributed fashion, for example, with different subsets of the instructions stored in memory of different computing devices that execute those different respective instructions, an arrangement which is consistent with the singular “medium” as that term is used herein. Further, storage of data need not be persistent in the absence of power and may be effectuated with either persistent data storage or dynamic data storage, such as dynamic random access memory that may dissipate stored data when power is removed.

[0087] In some embodiments, the process **50** includes receiving a write command requesting that a document associated with the write command be stored in an immutable data structure, as indicated by block **52**, such as a tamper-evident immutable data structure. In some cases, the write command may be a SQL command to write to the above-described lower-trust data store **14** from the application **28** received by the security driver **30**, or the command may be a command to store a file or other document in a file system. For instance, in some cases, the write command may be a request received by a file manager of an operating system to store a file in a repository presented as a network drive within a user interface of the operating system, where some or all of the network drive is hosted in the immutable data structure like that described below with reference to FIGS. **3** and **5**. In some embodiments, the write command may be generated programmatically, or in some cases the write command may be generated responsive to a user input requesting that, for example, a document be stored. As noted above, the term “document” is used relatively expansively herein and may include storage of an individual bites of information, or larger binary large objects, such as files ranging up to multiple gigabits or larger. Files may include stored data associated with metadata, like author, file-size, creation date, modification date, a file-name, an extension, and the like. Such files may be consistent with the definitions of files in the Windows™, Android™, iOS™, Linux™, or Mac™ operating systems.

[0088] The term “immutable” in the phrase “immutable data structure” refers to an arrangement of data that the computing system and write the right command is configured to leave in place even after the information represented by the data changes. For example, the data might represent a user's telephone number, and embodiments using an immutable data structure may write a new record indicating a new telephone number, rather than overwrite an existing record. Thus, both the older version and the newer version are preserved (and may be labeled as the older or newer version) even after the value changes. Such embodiments may then reference a most recent representation of the value for that field to effectuate the change, rather than deleting the old instance. In some embodiments, the immutable data structure may be a tamper-evident data structure that is computationally infeasible to modify without the modification being detectable to a party auditing the data structure. Some embodiments may implement cryptographic hash pointers described below in a directed acyclic graph that make it computationally infeasible for a party to modify stored data without those modifications being evident in virtue of inconsistent hash values elsewhere in the data structure. Computational feasibility may depend on the use case, for example, whether economics and time available provide for certain amounts of competing resources to be brought to bear. In many cases, an operation requiring more than  $10^{\{circumflex over ( )\}}128$  hashes on average to manufacture a collision with altered data may be said to be computationally infeasible.

[0089] In some embodiments, the process **50** may include forming nodes of a directed acyclic graph having node content that includes the document, as indicated by block **54**. Forming nodes may include accessing existing nodes and inserting node content into those nodes or forming entire new data structures by which nodes are encoded. In some embodiments, each node may include node content with a collection of attributes, which in some cases may include a pointer to another node. In some embodiments, those node attributes may include a pointer that is an identifier of an adjacent node in the directed acyclic graph and a cryptographic hash value based upon one or more attributes of the adjacent node that is identified. In some embodiments, these last two pair of attributes (an identifier of another node and cryptographic hash value of at least some of that node's content) may correspond to a cryptographic hash pointer from one node to another. Cryptographic hash pointers may define edges of the graph in some embodiments. In some embodiments, an individual node may contain zero cryptographic hash pointers (such as a leaf node in a binary tree), a single cryptographic hash pointer (such as a link in a block chain or linked list), a pair of cryptographic hash pointers (such as in a non-leaf node of a binary tree directed acyclic graph), or various other amounts of cryptographic hash pointers, for example, in splay trees or skip lists.

[0090] In some embodiments, the node attributes include an identifier of the respective node. In some embodiments, the identifier may be a segment identifier of the type described below, or in some cases the identifier may be an identifier within a namespace of the directed acyclic graph, for example, a count that increments with the addition of each node. In some embodiments, the identifier may be arranged hierarchically, for example, indicating a block and a block chain and then a sequence of binary values that specify a path through a binary tree, as is used in prefix trees in some cases (e.g., the value 00101 may define a path from a root node to the left, left, right, left, and then right, with left expressed as 0 and right expressed as 1, in a binary tree). In some embodiments, the identifier of a node is content-based, for example, a cryptographic hash value or non-cryptographic hash value based on one or more or all of the attributes of the node, which may include hash pointers of that node.

[0091] In some embodiments, the document may be stored in a single node of the directed acyclic graph. Or in some cases, the document may be segmented and stored in multiple nodes, for example consistent with the techniques described below with reference to FIGS. **4** through **6**. In some embodiments, some nodes of the graph may store other documents or portions of other documents. In some embodiments, nodes may be added to the graph over time, in some cases with groups of nodes added sequentially. In some embodiments, attributes of nodes or blocks of nodes (which may themselves be characterized as nodes in some cases) may further include a timestamp at which the node was formed, an identifier of a tenant account or data repository where data in the note is stored in the lower-trust data store **14**, a date the node was created, and the like.

[0092] In some embodiments, groups of nodes may be added as a “block,” for instance with each block corresponding to a binary tree having documents stored in leaf nodes. Or blocks may be linked lists, skip lists, splay trees, or combinations thereof, for example. In some embodiments, an individual node may store multiple documents as attributes of that node. In some embodiments, blocks have an integer index, a block capacity, a cryptographic hash value based on all of the nodes in the block (like a Merkle root), the nodes within the block, and a cryptographic hash based on content of a previous block (e.g., based on all values in the block, based on a Merkle root of that block, or the like).

[0093] In some embodiments, forming the nodes of the directed acyclic graph includes forming (e.g., updating or creating) a sequence of nodes along one or more paths through the directed acyclic graph. In some embodiments, this may include calculating one or more cryptographic hash values to form one or more cryptographic hash pointers along this path, which in some cases may include or terminate with a node in which the document is stored or a portion of the document is stored, such that each cryptographic hash pointer along the path is based on the document.

[0094] Cryptographic hash pointers may be based upon a cryptographic hash function which may



take a plurality of inputs, such as one or more node attributes and produce an output of fixed size. These functions may have pre-image resistance, second pre-image resistance, and collision resistance. Examples include an SHA-256, BLAKE, BLAKE2, SHA-1, SHA-2, and SHA-3 hash function. In some embodiments, the cryptographic hash function may be a one way function in which a given string of input produces deterministically a string of output that is relatively difficult or impossible to reverse to determine the input from the output while being relatively easy to confirm that an input corresponds to the output. For example, it may be computationally infeasible to identify a hash collision in which different instances of the input produce a given output. In some embodiments, the cryptographic hash function may implement the Merkle-Damgård construction. [0095] In some embodiments, the cryptographic hash function may be based upon a compression function that accepts a fixed size input and produces a fixed sized output with a one-way compression function. In some embodiments, because the input to the cryptographic hash function may be a variety of different sizes, the transformation may be performed in a variety of iteration and a plurality of iterations. Some embodiments may determine a length of input, such as a number of bytes, accepted by the one-way compression function, a length of output of the one-way compression function and determine a difference between these two lengths. Some embodiments may then parse an input to the cryptographic hash function into sequences of a size of this difference and iteratively input the parsed sequence into the one-way compression function and then combine the output of that iteration with a next portion parsed portion from the input, for example, incrementing from a beginning to an end of an input and iteratively concatenating, for example, prepending or appending or otherwise intermingling the output of the previous iteration of the one-way compression function with a next parsed portion of the input to the cryptographic hash function. Some embodiments may repeat this until an end of the input to the cryptographic hash function is reached (e.g., reaching a point within some threshold number of bytes or the last byte), as indicated by block **56**. In some embodiments, for example, where a plurality of inputs are applied, some embodiments may combine these inputs in a variety of approaches, for example prepending or appending or otherwise intermingling these inputs to form a string upon which these operations may be performed to produce a fixed sized output that is based upon the entirety of the input. The resulting directed acyclic graph may be stored in memory.

[0096] Before completing discussion of FIG. **2**, it is helpful to describe an example of a directed acyclic graph like that used in the process of FIG. **2**. FIG. **3** shows an example of a directed acyclic graph **70** consistent with the present techniques. In some embodiments, the directed acyclic graph **70** may include a plurality of blocks in a sequence, such as blocks **72**, **74**, and **76**, which may be arranged in a linked list, with links formed by cryptographic hash pointers corresponding to the illustrated arrows between the illustrated blocks. In some embodiments, the block **72**, **74**, and **76** may be added sequentially, over time, for example as each block is completed. In the illustrated example, the block **72** may be a most recent block, while a block **76** may be an earliest block. In some embodiments, the cryptographic hash pointers between each of the blocks may be based upon node content in the preceding block, and that node content may include a cryptographic hash pointer based on node content in the preceding block to that block. Thus, a cryptographic hash value of the cryptographic block **72** may be based upon each of the preceding blocks' cryptographic hash values and all content rendered tamper-evident by the data structure.

[0097] In some embodiments, each block may include a binary tree with a root node **80**. In some embodiments, each of the arrows between blocks in the binary trees of each of the block **72**, **74**, and **76** may also be cryptographic hash pointers, for example, based on an identifier of the node to which the cryptographic hash pointer points and a cryptographic hash value based upon node content of that node, which may include an attribute of that node that is itself a cryptographic hash value of another hash pointer. Thus, in some cases, a cryptographic hash value of a root node **80** may be based upon node content of every node of the binary tree in each block **72**, **74**, or **76**. In some embodiments, the node **80** may include three such hash pointers, corresponding to six node

attributes, the six attributes including three pairs of node identifiers and cryptographic hash values based on node content of those nodes. In some embodiments, node content may further include a cryptographic hash value based upon each of these values, or such a cryptographic hash value may be stored in another node that points to that node. The illustrated graphs are acyclic. As that term is used herein, it may refer to an acyclic subset of a larger cyclic graph. Thus, claims to acyclic directed graphs may not be avoided simply by adding an un-used cycle.

[0098] In some embodiments, the binary tree structure may facilitate relatively fast access to records within a given binary tree once a given block **72**, **74**, or **76** is identified. In some embodiments, to expedite access to specific blocks, some embodiments may include a skip list data structure, for example, with another node or value within node **80** that includes a cryptographic hash pointer to some number of blocks earlier in the sequence, for example, to a block four positions earlier, another to a block eight positions earlier, another to a block **16** positions earlier, another to a block **32** positions earlier, and so on. Some embodiments may thus, skip over some portions of the sequence of blocks to access a specified block. Three blocks **72**, **74**, and **76**, are shown, but is expected that commercial embodiments will include substantially more, for example more than 50, more than 100, more than 500, more than a thousand, or more than 10,000 blocks in a given directed acyclic graph **70**. In some embodiments, an earliest block **76** may include a cryptographic hash pointer to a seed node **78**, which may include a random value (e.g., a pseudo random value) longer than 64 bytes as node content as an attribute to provide a minimum level of entropy for each succeeding cryptographic hash value that points directly or indirectly to the node **78**.

[0099] In the illustrated example, the blocks **72**, **74**, and **76** include binary trees with three levels of hierarchy, but embodiments are expected to include more, for example four levels, five levels, six levels, seven levels, eight levels, or more levels of hierarchy, each including twice as many blocks as a proceeding level in the hierarchy. In some embodiments, the binary trees may be balanced binary trees, or in some cases the binary trees may be unbalanced. In some embodiments, nodes in the binary trees below a lowest level, such as nodes **80**, **84**, and **86** that are non-leaf nodes, may include node content that is based on cryptographic hash pointers but not based on node attributes in those nodes that store content of documents, or in some cases documents may also be stored in these nodes. In some embodiments, content of documents, such as documents themselves or portions of documents may be stored in attributes of node content of leaf nodes **88**, **90**, **92**, and **94**, such as an individual one of these leaf nodes in some cases. Thus, for example a given document stored in leaf node **90** as an attribute in node content of leaf node **90** may cause a cryptographic hash value in a cryptographic hash pointer of nodes **84** and **80** to be based upon that document. As a result, a modification to the document stored in node **90** may cause the cryptographic hash pointers in nodes **80** and **84** to be inconsistent with that document, as cryptographically hashing the modified document is not expected to yield the same cryptographic hash value in the cryptographic hash pointers of these nodes **80** and **94**. In some embodiments, this verification operation may be repeated along a sequence in a path of such nodes connected by cryptographic hash pointers, in some cases from or to any leaf node up to a most recent root node added to the directed acyclic graph **70**, thereby relatively granularly identifying a node with data from a workload application that has been modified and relatively reliably making that data tamper evident due to the computational infeasibility of crafting hash collisions consistent with the modification to the data along the path. In some cases, leaf nodes may be arranged based on the content, e.g., with each leaf node from left to right storing a range of values, or each leaf node storing up to a threshold amount of data arranged alphabetically, temporally, or in some other arrangement.

[0100] In the illustrated directed acyclic graph **70**, a sequence of binary trees are shown, but embodiments are consistent with other types of directed acyclic graphs. In some embodiments, the directed acyclic graph is a graph with no unconnected sub graphs. In some embodiments, the directed acyclic graph is a single binary tree, or in some cases the directed acyclic graph is a splay

tree of binary trees or a binary tree of splay trees. In some embodiments, the directed acyclic graph includes a skip list of splay trees or binary trees. In some embodiments, the directed acyclic graph includes a binary tree of skip lists or splay trees or linked lists. (Some embodiments may facilitate re-arrangement of trees and other structures to facilitate faster access with abstract cryptographic hash functions described below, or hash pointers may be recalculated upon restructuring).

[0101] Thus, the directed acyclic graph **70** may include a plurality of nodes, some of which have node content that includes cryptographic hash pointers correspond to edges in the directed acyclic graph **70**, and some of which includes node content that includes as attributes workload application data that is secured by the directed acyclic graph **70**.

[0102] In some embodiments, the directed acyclic graph **70** may take the form of a distributed ledger, such as a block chain having a linked list of blocks, with each block including a Merkel tree having a Merkel root that serves as a node attribute of the respective block, and subsequent blocks having node attributes that include cryptographic hash values based on the Merkel root of each proceeding block.

[0103] In some embodiments, adding nodes or collections of nodes, such as blocks, to the directed acyclic graph **70** may be implemented in a un-trusted computing environment by un-trusted computing devices, in some cases based on a consensus protocol (like Paxos, Raft, or others) in which different computing devices perform or demonstrate some proof of something that is difficult to consolidate, such as proof of work or proof of storage, for instance, by calculating a hash collision to a threshold number of prefix or suffix digits of a cryptographic hash value. In some embodiments, a collection of untrusted computing devices may execute a consensus algorithm to determine whether blocks should be added or are valid, for example. Or to expedite operations, some embodiments may avoid the overhead associated with proof of work or storage techniques by executing within an environment in which more control is exercised over the hardware that performs the operations described herein.

[0104] It should be emphasized that a data structure need not be labeled as a graph or as having nodes or edges in program code to constitute a graph, as long as the data structure includes linking between data elements present in a graph. In some cases, graphs may be encoded as a list of links between data elements, for example, in a key-value pair, or in a hierarchical data structure, such as in a hierarchical data serialization format, like JSON or XML. Graphs may also be encoded as relationships in a relational database. Further, multiple data structures, such as different graphs, may be overlaid on one another, for example in the manner described below, while still preserving the existence of the distinct graphs.

[0105] In some embodiments, each of the above-described storage compute nodes **26** may store one or more instances of a directed acyclic graph **70** like that shown in FIG. 3. In some embodiments, a given directed acyclic graph may be replicated in multiple instances on multiple storage compute nodes **26**, for example, in multiple data centers **24**. In some cases, an odd number of replicated instances may be stored. When reading data back, some embodiments may implement a consensus protocol to determine an authoritative copy, for example identifying a version of a document returned by a plurality or a majority of the replicated instances following a read operation like that described below.

[0106] In some embodiments, a given directed acyclic graph **70**, which may be replicated in multiple instances, may store data from a single workload application from a single tenant, and in some cases a subset of data from a single workload application, or data may be mixed. In some embodiments, a given directed acyclic graph may store data from multiple applications of a given tenant. In some embodiments, a given directed acyclic graph may store data from multiple tenants, depending upon the implementation.

[0107] In some embodiments, as data is accumulated, when the directed acyclic graph **70** exceeds a threshold size, a new directed acyclic graph may be instantiated, for example, using a cryptographic hash value of a root node of a last block added to a previous directed acyclic graph

as a seed value 78 in a seed node.

[0108] Thus, FIG. 3 shows an example of a directed acyclic graph in which documents stored with the process of FIG. 2 may be stored in a tamper-evident fashion in an immutable data structure. As the document is a revised, new records may be added with new versions of the document to the data structure 70, and pointers to a most recent version of the document may be updated, for example, in the lower-trust database 14 to reference those newer versions. In some embodiments, these pointers may implement techniques described below by which pointers are substituted for records in a data store. In some embodiments, those pointers may point to an individual node in an individual directed acyclic graph. In some embodiments, the pointers are node identifiers or the below described segment identifiers.

[0109] In some embodiments, the translator 20 described above may maintain in memory an index that maps pointers to uniform resource identifiers of storage compute nodes 26 and directed acyclic graphs maintained by the storage compute nodes, and in some cases, the DNS 18 may map those uniform resource identifiers to Internet Protocol addresses and port numbers of the corresponding storage compute nodes 26. Thus, after a document is written to the directed acyclic graph 70, a pointer, such as a node identifier that distinguishes that node from among all other nodes accessible to a given application, tenant, or all tenants, may be stored in the lower-trust database 14 or other data store in place of the document.

[0110] In some embodiments, the directed acyclic graph 70 may be periodically verified for internal consistency to detect instances of tampering with documents. Some embodiments may recalculate cryptographic pointers along various paths, for example, each time a new document or other record is added to the directed acyclic graph 70, for example, each time a new block is added. Upon detecting a path in which the cryptographic hash pointers do not correspond to a document or other node content in that path, some embodiments may designate that document or other values as having been modified. Some embodiments may further emit an alarm, e.g., send an email, send a text message, and suspend various user accounts, such as users that have access the document responsive to detecting tampering. To detect tampering, some embodiments may recalculate the cryptographic hash values for each cryptographic hash pointer along each path to each document and determine whether the recalculate cryptographic hash values match those stored as node attributes of nodes storing the respective cryptographic hash pointers. Thus, if a given document, for example in node 90 of block 72 is tampered with after storage, the cryptographic hash pointer from block 84 to block 90 will not match the stored cryptographic hash value in the node attribute of node 84, nor will the cryptographic hash pointer from block 80 to block 84 match the stored value in the node 80. Further, if additional blocks were added subsequent to block 72, the cryptographic hash pointers between blocks subsequent to block 72 will also fail to match, thereby providing a traceable path from a most current node back to a node in which the node content was tampered with by an attacker. In some embodiments, these verification operations may be performed each time a node is written, each time a block is written, periodically, like hourly, daily, or weekly, or each time a read operation is performed.

[0111] FIG. 2 further includes operations by which a document may be read back from the secure distributed storage 16. In some embodiments, the operations of process 50 may include receiving a read command requesting that the document be retrieved from the immutable data structure, as indicated by block 60. In some cases, the read command may be received after the write command, e.g., substantially later, for instance more than an hour, day, or week later. In some cases, the read command may be received after multiple write commands for the same document in which different versions are written to different nodes in different blocks, and in some cases to different directed acyclic graphs like those described above with reference to FIG. 3. In some embodiments, the read command may reference an identifier of a document that indicates a most current version of the document is to be retrieved, or in some cases the read command may reference a particular version of the document. In some cases, receiving the read command may cause the security driver

**30** to access the lower-trust database **14** or other lower-trust data store and retrieve a pointer to a node or sequence of nodes in which the specified document is stored. As noted above, in some cases the pointer may be a value that maps to a uniform resource identifier that the DNS **18** maps to an Internet Protocol address and port number where the corresponding node is stored in a corresponding directed acyclic graph.

[0112] Next, some embodiments may access one or more nodes of the directed acyclic graph to retrieve the document, as indicated by block **62**. In some cases, this may include sending a pointer to a node at which at least part of the document is stored to the translator **20**, accessing a uniform resource identifier of a directed acyclic graph in which the node resides in an index stored by the translator **20**, converting the uniform resource identifier to an Internet Protocol address and port number pair with the DNS **18**, requesting the corresponding node from the storage compute node **26** hosting the specified directed acyclic graph at the Internet Protocol address and port number, and returning the requested document to the requesting application **28**.

[0113] Further in some cases, accessed documents may be verified. For example, some embodiments may recalculate cryptographic hash values based on the document along a path from a most current node of the directed acyclic graph back to the document (or a subset thereof) and determine whether any calculated cryptographic hash values do not match those stored in the respective nodes of the directed acyclic graph, as indicated by block **64**.

[0114] For clarity of terminology, it should be noted that the storage compute nodes **26** refer to computing entities (e.g., a service topology), while the nodes of the directed acyclic graph **70** refer to elements of a data structure.

[0115] Embodiments may determine whether these calculated values are internally consistent, as indicated by block **66**. This may include determining whether calculated values match those stored in the node pointers of the directed acyclic graph **70**. Upon determining that the values match, some embodiments may return the requested document, as indicated by block **69**. Upon determining that the values do not match, some embodiments may indicate that the document has been modified, as indicated by block **68**. Thus, an entity retrieving the document may have some confidence that the document has not been tampered with. Further, because some embodiments store the document in the nodes, the document may be retrieved relatively quickly (e.g., within 200 ms, 100 ms, or 50 ms), even when storing relatively large amounts of data, e.g., more than 100 GB or more than 10 TB. Though embodiments are not limited to systems affording these benefits, as independently useful techniques are described, which is not to suggest that any other description is limiting.

[0116] In some embodiments, documents or other values written to the secure distributed storage **16** in nodes of directed acyclic graphs **70** may be subject to various types of encoding prior to write operations. In some embodiments, data may be encoded with redundant information such that if some of the data is modified before writing or after writing, the encoding reveals errors and in some cases provides enough redundant information that data may be recovered. For example, some embodiments may encode the data with Hamming codes or Turbo codes such that if a subset of the data is lost, the document may be reconstructed based on redundant information in another part of the document. In another example, data may be compressed, for example, with entropy coding to facilitate reduced use of bandwidth and faster retrieval. In some embodiments, sequences of values in a document may be analyzed to identify frequencies with which sequences occur, and those sequences may be mapped to a dictionary in which relatively frequent characters are represented with relatively short values, while longer less frequent sequences are mapped to longer values. The mapping and dictionary may then be stored in the secure distributed storage in a compressed format. In some embodiments, after compression, some embodiments may encrypt the data as well, for example, with a symmetric encryption algorithm in which keys are maintained by the client computing device **12**, for instance, by XOR'ing the document with a random string that serves as the symmetric key. In some cases, the symmetric key may be stored in the distributed storage **16** in the manner described above.

## Fragmenting Data for the Purposes of Persistent Storage Across Multiple Immutable Data Structures

[0117] As noted above, traditional databases do not adequately protect against threat actors or internal resources (employees, information-technology staff, etc.) tampering with the data. Data residing in singular logical spaces presents an easy target for threat actors. Once the computing device upon which the data is stored is compromised, the entirety of the data is potentially at risk. Similar problems arise when a single computing device controls access to all other computing devices in distributed master-slave architectures, as once the master-device is compromised, data on other devices is potentially exposed.

[0118] Techniques exist for storing data in a distributed form in a peer-to-peer network, but such systems often present other problems. Examples include various systems built around the BitTorrent protocol. These systems often operate without a central authority to be compromised. But in these systems, when or if the data is altered (potentially for the purposes of hiding or obfuscating changes), there is no recourse or method to detect these changes. Further, many distributed peer-to-peer storage networks, like many implementing the BitTorrent protocol, maintain full copies of files (or other BLOBs) at each peer, potentially leaving the full file open to inspection by an untrusted party.

[0119] To mitigate these issues and others, in some embodiments, when writing data, a component referred to as an “arbiter” (which may be a piece of middleware and may be implemented in the translator **20**) may be capable of taking as an input a fully formed datum (e.g., a unit of content, like a document or value in a database), starting from the last byte of the data, fragmenting that data into N pieces (where N is an integer larger than 1, and in many commercially relevant cases larger than 5 or 10, and in some cases can be relative to the size of the initial data), and placing each piece on a physically (or virtually) separate blockchain-backed storage data structures, with each piece containing pointers to the next storage location of the fragmented data. In some embodiments, the arbiter (or other middleware) then returns the TXID (or other node identifier) of the last written fragment (the first byte) to the application which requested the fragmenting of the data.

[0120] When an application or resource requests the reassembly of fragmented data, an arbiter (or other piece of middleware) is supplied with the TXID (or other node identifier) of the first byte of the data, in some embodiments. After reading the first byte, in some embodiments, the arbiter or middleware then reads the subsequent pointers until a null character or end of sequence character is read. Once all of the pieces have been read into memory, the arbiter or middleware respond to the application with the resultant unfragmented datum.

[0121] In some embodiments, other data structures, such as files or databases can be used to store the fragments. Additionally, some embodiments pre-process the data and count the number of pieces that are required from the beginning before fragmenting the data.

[0122] An additional enhancement in some embodiments is the fragmenting of the data into tree-like structures instead of list-like structures, which, during read operations, are expected to accommodate concurrency and reduce operation time.

[0123] In some embodiments, when reading or writing scattered data, the execution operation along with other meta-meta data such as user credentials may be written to an ‘auxiliary’ chain for the purposes of auditing access to scattered data.

[0124] Advantageously, the blockchain data structures may impede or prevent modification of the stored data fragments. Further, some embodiments may store the data with some redundancy across fragments, e.g., with Hamming codes or Turbo codes, such that even if one blockchain is compromised, fragments from other changes would be inconsistent with modified data and signal the compromise. Fragmenting the data is expected to facilitate concurrent data validation on the different chains, thereby expediting read operations in some embodiments. Further, fragmenting the data is expected to make the data more robust to attacks on individual peer machines, as even if a peer machine is compromised, only a relatively small fragment is at risk. That said, not all

embodiments provide these benefits, as multiple independently useful inventions are described, and various engineering and cost tradeoffs may result in performance advantages in one aspect being deployed to improve another aspect, e.g., using faster algorithms to accommodate slower, cheaper hardware.

[0125] To implement these techniques and others, some embodiments may execute a process **100** shown in FIG. **4** to fragment documents and distribute the fragments across multiple storage compute nodes **26** having multiple, different directed acyclic graphs, in some cases with each of those directed acyclic graphs being replicated on multiple storage compute nodes. In some embodiments, the data being stored may be individual values and individual fields having individual portions of records of a database, such as a value at a given row and column position in a given table. In some cases, the data being stored as a unit may include metadata, like a primary key, table identifier, database identifier, and tenant identifier. In some embodiments, the data being stored may be a document like those described above that is larger than an individual value.

[0126] In some embodiments, the process **100** may include receiving a write request to write values to a database, as indicated by block **102**. In some embodiments, this operation may be performed by the security driver **30** described above. In some embodiments, the write request may be a request to write a value with a structured query language statement to the low-trust database **14**, or a request to write to a file in a file system. Or in some cases, the values subject to subsequent operations may be values already stored in a lower-trust database **14** or file system, for example, values processed during an installation process by which a lower-trust database **14** is secured by moving data into the secure distributed storage **16** described above.

[0127] Next, some embodiments may classify values as higher or lower security, as indicated by block **104**. For example, a write request may specify that a given record is to be added to a new row in a table, and that record may specify a tuple in which values are assigned to each of several fields, each corresponding to different columns in the table. In some cases, some of those fields may be lower-security, while other fields may be higher-security. In some embodiments, the security driver **30** may include a table that maps table/field combinations to security designations, for example, binary values indicating whether the values are lower or higher security values. Or some embodiments may use the other types of data security policies described below. For example, a Social Security number field may be a higher-security value, while a username in a public-facing account may be a lower-security value. In another example, credit card information like a credit card number may be a higher-security value, while a ZIP Code in a mailing address may be a lower-security value.

[0128] Next, some embodiments may cause the database to store the lower-security values, as indicated by block **106**. In some cases, this operation may include modifying a SQL statement or other database command received by the security driver **30** from the application **28** and providing the modified command to the database driver **32**. For example, some embodiments may modify portions of the received command that would otherwise specify that the higher-security values are to be written to the lower-trust database **14**. In some cases, the modification takes the form of modifying the values themselves rather than reserved terms in the command, such that modified values are written to the lower-trust database **14** in the positions that the corresponding higher-security values were going to otherwise be written by the application **28** (thereby maintaining relationships stored in the database). Or pointers may be written to directories in a file system in place of files or other documents. For example, some embodiments may cause node identifiers or other pointers to nodes to be written, for example, identifiers of segments of the data to be written that are formed in operations described below, like a first segment.

[0129] Next, some embodiments may determine whether there are more higher-security values that were identified to process, as indicated by block **108**. In some cases, some embodiments may iterate through each of the higher-security values, or process some or all of the higher-security values concurrently, and operation **108** may determine whether the full set has been processed.

Upon determining that all of the higher-security values that were classified have been processed, some embodiments may return to block **102** and wait for a next write request. Alternatively, some embodiments may select a next value that is higher security, as indicated by block **110**, for subsequent operations.

[0130] Next, some embodiments may segment the current value into a plurality of segments, as indicated by block **112**. As noted above, the current value in some cases may be an individual value for an individual field of an individual row of a database, or in some cases the value may include additional information, such as an entire row, or a value that associates a value in the database with a table and primary key for that row. Or in some cases, the value may be a document like that described above that may include additional data.

[0131] Next, some embodiments may segment the current value into a plurality of segments, as indicated by block **112**. In some cases, the value may be segmented by character position in the sequence or byte position in a sequence of bytes, for example, segmenting each sequential byte in a serialized representation of the value or each sequential four bytes. In some embodiments, segmenting may include segmenting with the entropy coding techniques described above. For example, some embodiments may segment according to a dictionary coding in an entropy coding of the value, for example, with each dictionary entry corresponding to a different segment. In another example, segmenting may include an operation that breaks up the value into multiple segments that each contain nonconsecutive portions of the value, for example, placing every fifth byte into a first segment starting with a first byte, placing every fifth byte in a second segment starting with a second byte, placing every fifth byte in a third segment starting with a third byte, and so on. In some embodiments, values may be segmented into a specified number of segments, for example, dividing the value into fourths, such that one fourth of the bytes by which the value is encoded go into a first segment, one fourth into a second segment, one fourth into a third segment, and one fourth into the fourth segment.

[0132] In some embodiments, values may be segmented with techniques from network coding (e.g., linear network coding) such that resulting segments are combined in a bit-wise operation to reconstruct the value. For instance, a given value 64 bytes long may be segmented into two 64 byte segments that when combined in a bit-wise XOR operation produce the original value. In some embodiments, some of the segments may contain redundant information relative to other segments. For example, some embodiments may encode data before segmenting or during segmenting with Hamming codes or Turbo codes, such that if two of three consecutive segments are available, the information and the other segment can be re-created, or such coding may be applied with less redundant information such that errors in one segment render that segment inconsistent with the redundant information in an adjacent consecutive segment. In some embodiments, the segments may be represented as a sequence or in some cases, the segments may be represented in other data structures. For example in some cases, the segments may be represented as a collection of paths through a prefix tree.

[0133] Next, some embodiments may cause the database, such as the lower-trust database **14**, to store a pointer to a first segment, as indicated by block **114**. In some cases, this operation may further include causing a record to be added to an index maintained by the translator **20** that associates that pointer, which may be a unique segment identifier, which in some cases may be an identifier of an individual node in a data structure like that directed acyclic graphs **70** described above, or an individual attribute of a node in that data structure, with an address of the directed acyclic graph. In some cases, the association may be with a plurality of addresses each corresponding to a different replicated instance of the directed acyclic graph, such as a plurality of different uniform resource identifiers that are mapped by the DNS **18** to Internet protocol addresses and port numbers of the corresponding storage compute nodes **26** storing those different instances of the directed acyclic graph and in some cases performing the above-described operations by which nodes are formed and data is verified. In some embodiments, causing the database to store a



pointer to the first segment may include performing the above-described operations by which a write command received from the application **28** is modified to replace the higher higher-security values with the pointers before providing that command to the database driver **32**.

[0134] Some embodiments may iterate through each of the segments, causing each of the segments to be stored and determining in each iteration whether there are more segments to process, as indicated by block **116**. Upon determining that there are more segments for a given value to be processed, some embodiments may select a next segment, as indicated by block **118**. Some embodiments may then select a computing device, such as a computing device executing one of the above-described storage compute nodes **26**, to store the current segment, as indicated by block **120**. In some cases, a list of the storage compute nodes **26** may be maintained in a circular linked list, and some embodiments may maintain a pointer that is advanced with each storage operation, causing the pointer to cycle through the list, and using the pointed-to storage node to store the current segment. Or some embodiments may select storage compute nodes **26**, for example, with a random number or with a pseudorandom value generator, such as a linear shift register. In some embodiments, storage compute nodes to store a given segment may be selected based on a unique identifier assigned to the segment. In some embodiments, the unique identifier may be a hash value based on content of the segment, and a prefix or suffix of the hash may uniquely identify one of the storage compute nodes **26**. In some embodiments, such a prefix or suffix may uniquely identify a directed acyclic graph **70**, which may be replicated on multiple stores compute nodes. In some embodiments, the content of the segment itself may be hashed, and a prefix or suffix may specify one of the above-describe storage compute nodes **26** or an instance of a directed acyclic graph that may be replicated on multiple stores compute nodes **26**. For example, 64 different directed acyclic graphs may be maintained, each replicated three times, and a trailing six digits of a binary representation of a hash of the content of the segment may specify a value that ranges between zero and 63, and that value may serve as an identifier of one of the directed acyclic graphs to which that segment is to be stored. In some embodiments, selecting a directed acyclic graph may cause one or more, such as three, computing devices to be selected, those being the computing devices that store and maintain the corresponding directed acyclic graph.

[0135] Next, some embodiments may instruct the selected computing device to store the segment in memory in association with a pointer to a previous segment, as indicated by block **122** (except for a first-segment, which may not have a pointer to a previous segment). In some embodiments, this may include the above-described translator **20** sending instruction to each of an odd numbered set of storage compute nodes **26** replicating one of the directed acyclic graphs **70** to form a node in which the segment content and the pointer serve as one or more attributes of that node, such as one of the above-described leaf nodes **88** through **94**. In some embodiments, forming this node may further include forming cryptographic hash pointers to that node through the above-described types of data structures by which different directed acyclic graphs may be formed, for example, through higher levels of a hierarchy in a binary tree, in some cases adding additional blocks to a linked list of blocks, each block having a different binary tree and link to previous blocks according to cryptographic hash pointers like described above.

[0136] Next, some embodiments may return to determine whether there are more segments in association with the current value. Upon determining that there are no more segments, some embodiments may return to block **108** and determine whether there are more higher-security values to process. In some embodiments, multiple values may be processed concurrently. In some embodiments, values may be segmented from front-to-back or vice versa.

[0137] FIG. **5** shows an example of a resulting data structure **130** in which a plurality of segments are stored in different directed acyclic graphs **132**, **134**, and **136**, which in some cases may each have the features of the directed acyclic graphs described with reference to FIG. **3**. In some embodiments, the segments of a given value may be stored in a content graph overlaid on each of these directed acyclic graphs **132**, **134**, and **136**, which may be characterized as a verification

graphs to distinguish an overlaid content graph (and as they serve as an immutable tamper-evident log of the values of the segments in some embodiments).

[0138] In this example, the segments in the content graph **138** form a linked list, with a first node of the segment content graph **138** being stored in a leaf node **139** of verification graph **132**, that first segment being designated as node **140** in the content graph. The node **140** in the content graph may be stored as an attribute in node content in the node **138** of the verification graph, and the content node **140** may include a pointer **142** to a next node in the content graph **138**, which may be stored on a different verification graph **134**. In some cases, the pointer **142** may be characterized as an edge in the content graph **138** and may be expressed as an identifier of a node **143** in the verification graph **134** or in some cases as an identifier of an attribute in that node where multiple segments are stored in different attributes of a given node of a verification graph. In some embodiments, in the content graph **138**, node **140** points to another segment **144**, which may then the point to another segment with pointer **146** in verification node **147** of verification graph **136**. Verification node **147** may include as node content one or more attributes that specify a final segment in content graph **138** designated with element number **148**. In some cases, node **148** may specify that there are no additional nodes in the value.

[0139] As a result, even if a given malicious actor somehow compromises one of the verification graphs **132**, **134**, or **136**, that attacker will only be able to access a set of segments of values and will not have access to other segments needed to complete the full value stored in the distributed storage **16**. Further, because the segments are stored in a tamper-evident directed acyclic graph with the above-described hash point cryptographic hash pointers, evidence of tampering will not be computationally feasible to conceal.

[0140] Thus, FIG. 5 shows a plurality of verification directed acyclic graphs **132**, **134**, and **136**, each of which may be replicated, and each of which has nodes that may store as node content data that encodes a content graph **138**, which in this case is a linked list of segments, where each segment in sequence points to the next segment and its corresponding address, and in some cases attribute identifier in the underlying verification graphs.

[0141] In this example, segments are arranged in a one-dimensional linked list, but embodiments are consistent with other arrangements of content graphs. For example, some segments may include pointers to multiple subsequent segments, for example, in a skip list to facilitate concurrent retrieval, and in some cases segments may be stored in association with a segment position identifier, for example, an order in which the segments are to be sequenced to reconstitute the segmented value by the translator **20** in a read operation. In another example, segments in a content graph encoding a plurality of segments of an individual value may be stored in a binary tree content graph, a skip list content graph, or a combination of binary trees, linked lists, skip lists, and the like.

[0142] Three segments for a given value are shown, but embodiments are expected to include substantially more in some cases. In some cases, binary data encoding a single text character may be segmented, for example with a given Unicode character being segmented into two or more segments, and a given value yielding 10 or more or 20 or more segments, which in some cases may each be stored in different distributed acyclic graphs, which in some cases may each be replicated multiple times, for example 3 or more times. Thus, a given value may be stored in part on 30 different storage compute nodes **26**. In some cases, different instances may be compared to determine an authoritative copy, e.g., selecting a stored and returned value according to a majority rule approach among the replicated instances. In some cases, e.g., where the replicated instances of the graphs are on permissioned computing devices, embodiments may vote for a given value without performing a proof of work or proof of storage operation, or where devices storing the graphs are untrusted, some embodiments may determine consensus with a proof of work, storage, or stake, e.g., according to a consensus protocol, like Paxos, Raft, or the like. In some embodiments, e.g., in untrusted systems, instances may be addressed according to Interplanetary

File System (IPFS) or with various distributed hash table approaches.

[0143] In the example of FIG. 5, each of the directed acyclic graphs **132**, **134**, and **136** is the same type of directed acyclic graph, in this case a linked list of binary trees, where edges are formed by cryptographic hash pointers. In other embodiments, a heterogeneous set of directed acyclic graphs may be combined, for example with different segments stored in different types of graphs. For example, an initial segment may be stored in a binary tree, while other segments may be stored in directed acyclic graphs like those shown in FIG. 5, for example, in linked lists of binary trees.

[0144] FIG. 6 shows an example of a process **150** by which data may be retrieved from the data structure **130** of FIG. 5 or other data structures written to with the process of FIG. 4. In some embodiments, the process **150** includes detecting an application querying a database, as indicated by block **152**. In some cases, this includes executing the security driver **30** shown in FIG. 1 and receiving a read request sent to the database driver **32** by the application **28** in FIG. 1. Some embodiments further include receiving a query response from the database, as indicated by block **154**. In some embodiments, receiving the query response may occur after the database driver **32** in FIG. 1 sends a query to the lower-trust database **14**, which may return a query response, for example, a response to a SQL statement selecting certain records that satisfy certain criteria. Or the response may be a document accessed in a file system. In some embodiments, the query response may be received by the database driver **32**, which may then send the query response to the application **28**. This response may be intercepted by the security driver **30** and modified by the security driver **30** before it is provided to the application **28**. In some embodiments, the security driver **30** may detect pointers to segments stored in the process of FIG. 4, for example, in a data structure of FIG. 5, and send those pointers to the translator **22** be translated back into the values that were sent to the data structure in the suit secure distributed storage **16**, for example in a plurality of segments. Thus, a given query response may include a lower-security portion of the query response corresponding to values in fields that, in the process of FIG. 4, are not classified as higher-security values, and higher-security values that are classified as such in the process of FIG. 4, or at least pointers to those values in the secure distributed storage **16**. In some embodiments, these pointers may be detected with a regular expression configured to detect a prefix or suffix labeling the pointers as such, for example, with a sequence of reserved characters.

[0145] Thus, some embodiments may include an operation of determining whether any pointers formed in the process of FIG. 4 are present, as indicated by block **156**, in the received query response. Upon determining that at least some pointers are present, some embodiments may determine whether there are more pointers to process in the query response, as indicated by block **158**. Some embodiments may iteratively or concurrently process each of the pointers present in the query response from the lower-trust database **14** to replace those pointers with corresponding value stored in the secure distributed storage **16**, for example, in the data structure of FIG. 5 with the process of FIG. 4.

[0146] Upon determining that there are more pointers to process, some embodiments may select a next pointer, as indicated by block **160**, and retrieve a segment and associated pointer if and associated pointer is stored in association with that segment, as indicated by block **162**. Segments may be retrieved in reverse order or vice versa relative to the order in the value that is segmented, depending on how the content graph is structured, e.g., based on the order in which the segments are written. In some embodiments, retrieving the segments and associated pointers may include causing the translator **20** to access an index that associates pointers, for example, segment identifiers, with one or more URLs of one or more computing devices storing a replicated instance of one of the directed acyclic graphs **132**, **134**, or **136** storing the segment to which the pointer points. In some embodiments, the index may further identify, in association with the pointer, a node identifier, for example, a block and path through a binary tree, to a specific node storing as node content in an attribute of that node the segment at issue and an associated pointer of present, for example if it is not the last segment to be retrieved. In some embodiments, the translator **20** may

access the DNS **18** to identify an Internet Protocol address and port number of the secure compute node **26** maintaining the identified directed acyclic graph, or each secure storage compute node **26** storing one of a replicated instance of that directed acyclic graph. Some embodiments of the translator **20** may send a request to those storage compute nodes to return the identified segment, and the receiving storage compute nodes **26** may traverse the edges specified by the cryptographic hash pointers in the corresponding directed acyclic graph along a path specified by the request to the specified node storing the segment. Or some embodiments may navigate directly to the segment, for example, based on a node identifier that identifies an address in memory of the segment without requiring traversal. In some embodiments, to expedite access, locations of nodes in memory may be stored in a binary tree corresponding to node identifier, or in a sorted list, to facilitate relatively fast access to content of nodes.

[0147] In some embodiments, the storage compute node **26** or set of storage compute nodes **26** replicating the directed acyclic graph, may return the specified segment to the translator **20**. In some embodiments, the translator **20** may receive multiple instances of the requested segment, and some embodiments may determine whether each of the segments match one another or determine a content of the segment based on a content returned by a plurality or majority of the responses, thereby preventing a modification of one of the instances from affecting the authoritative content of the segment determined by the system. That said, not all embodiments provide this benefit, which is not to suggest that any other feature described herein is limiting.

[0148] Some embodiments may maintain in memory a partially or fully formed encoding of the value to which the segments belong, and some embodiments may concatenate the retrieved segment to a sequence of segments accumulated in this value, as indicated by block **164**. For example, some embodiments may prepend or append each successive retrieved segment to this sequence of segments. Or in some big cases, the segments may come be combined with other techniques. For example, some embodiments may combine the segments by XOR'ing the segments to re-create the value. The term “concatenate” is used broadly herein to refer to each of these different operations by which information in different segments is combined to produce the value from which the segments were taken. In some embodiments, segments may be retrieved out of order, for example, with segments labeled with a value indicating a position in a sequence, and some embodiments may sort the segments according to this value to re-create the value from which the segments were taken.

[0149] Next, some embodiments may return to determine whether there are more pointers to process in block **158**. Upon determining that there are no more pointers to process, some embodiments may proceed to form a revised version of the query response, as indicated in block **166**, in which the pointers are replaced with the higher-security values segmented and stored in the secure distributed storage **16**, for example, in the data structure of FIG. **5** with the process of FIG. **4**. In some embodiments, the replacement operation may be performed by the above-described translator **20** or within the security driver **30**. In some embodiments, the revised version of the query response may include the replacement values in the same position, in the same formatting as the pointers within the query response, thereby maintaining associations between different values and metadata in the query response, for example, associations in which records are distinguished and associations in which values are indicated as applying to particular fields in those records. In some cases, where no pointers are present, the revised version of the query response may be an unaltered version of the query response, upon determining that no pointers are present in block **156**.

[0150] Next, some embodiments may provide the revised version of the query response to the requesting application, as indicated by block **168**. In some cases, the revised version of the query response may be advanced by the security driver **30** to the application **28** in FIG. **1** according to an application program interface of the database driver **32**, such that an application **28** configured to interface with the database driver **32**, but not designed with the security driver **30** in mind, will transparently receive the query data without needing to be reengineered to interface with the secure

distributed storage **16**.

[0151] In some embodiments, the write operations of FIG. 4 or the read operations of FIG. 6, or both, may be performed in different portions of the system of FIG. 1 (or other systems) than those described. For example, in some cases, these operations may be performed by a database management system or a database gateway by which client devices communicate with the database. In some embodiments, these operations may be performed by a dedicated appliance on a network through which client computing devices access a database. In some embodiments, these operations may be performed entirely within a single computing device, for example, as part of a monolithic application in which the features of the secure distributed storage are operated within a single computing device, or on a single device in which the different components execute in different containers or virtualized operating systems accessed via a loopback Internet protocol address.

#### Generation of Hash Values within a Blockchain

[0152] As noted above, blockchains generally allow for small bits of information to be stored in an immutable data structure; the moment data in the chain is altered, the chain is broken and generally can no longer function. The most common blockchain implementation is the publicly accessible Bitcoin ledger (for which blockchains were designed).

[0153] However, many extant blockchain systems are not well suited for certain use cases. Blockchains are typically immutable, and therefore once data is committed to the chain, under most circumstances it is considered permanent. This can lead to scalability challenges as the chain can grow to be very large with no ability to reduce the size of the chain (none of which is to suggest that systems with immutable blockchain are disclaimed for all embodiments or are inconsistent with some implementations).

[0154] These, and other problems, are mitigated by some embodiments of the system introduced above referred to below as “Docuchain.” Docuchain, in some embodiments, is a blockchain software suite for low-latency put and get operations of BLOBs. Docuchain, in some embodiments, uses an improved version of a hash function for a blockchain.

[0155] Docuchain, in some embodiments, makes use of abstract functions (a term which is used generically to also refer to methods and other routines) to define which properties of individual blocks are used in creating the block hash (sometimes referred to as Merkle root). By varying the properties that are used to construct this hash, some embodiments increase or decrease the ambiguity of the hash output. In some embodiments, a lambda function can be passed in to the constructor of BlockManager, which receives a block entry as a parameter and returns the values to be passed into the hash function, as defined by ‘H’ in the expression BlockManager<H,T>.

[0156] For example, if only the block’s data is used in the generation of its Merkle Root, it would be possible to remove the head of the chain without consequence of breaking the chain itself. Alternatively, if systems use the data, the position, and potentially even more specific properties such as read/write times, the chain becomes completely immutable and must (in some implementations) stay intact in its entirety in order to be functional.

[0157] Some embodiments add this degree of control by removing the in-chain scripting language component and adding the ability for the hashed value to be controlled as part of configuration rather than part of source code. In some embodiments, the hash function accepts data that is not contingent on a block’s position within a chain. In some embodiments, the hash function accepts as input only values that can be re-created in another chain.

[0158] In other embodiments, a similar output as a result of this functionality could be produced by conducting a sort of refactoring on a chain. This operation may include writing all of the data intended to be kept to a new chain, and then updating all referential TXIDs (or other pointers like those described above) in any external data stores that exist.

[0159] To implement these approaches and others, some embodiments may execute a process **170** by which a cryptographic hash function is configured to make it is feasible to recompose portions of directed acyclic graphs having edges specified by cryptographic hash pointers formed with the

cryptographic hash function, for example, to accommodate directed acyclic graphs that have grown larger than a threshold size and break those graphs into smaller portions or spawn new directed acyclic graphs that link back to those directed acyclic graphs.

[0160] In some embodiments, the process **170** may be performed by the storage compute nodes **26** of FIG. **1**, but is not limited to that implementation, which is not to suggest that any other description herein is limiting. Indeed, the operations of process **170** may be applied to other types of directed acyclic graphs having cryptographic hash pointers as edges, for example, in block chains used for cryptocurrencies and in distributed tamper-evident ledgers used for other purposes.

[0161] In some embodiments, the process **170** includes obtaining a tamper-evident log, as indicated by block **172**. In some embodiments, the tamper-evident log is one of the directed acyclic graphs discussed with reference to FIGS. **3** and **5** and may take a variety of different forms. In some embodiments, the tamper-evident log includes a sequence of nodes, or blocks that are collections of nodes, for example, each block being a binary tree of nodes. In some embodiments, the tamper-evident log may indicate through the topology of a directed acyclic graph a sequence with which entries in the log are added. In some embodiments, each entry may be characterized as a node or as a block in the tamper-evident log.

[0162] Next, some embodiments may obtain a plurality of records to be protected, as indicated by block **174**. In some cases, the records may be the segments or documents (which is not to suggest that a document cannot be a segment) discussed above. In some cases, the records to be protected may be obtained by the storage compute nodes **26**, for example, upon being received from a translator **20** like that described with reference to FIG. **1**.

[0163] Some embodiments may iteratively or concurrently process the records and determine whether there are more records to process, as indicated by block **176**. Upon determining that there are no more records to process, some embodiments may return to block **174** and wait for additional records to process. Upon determining that there are more records to process, some embodiments may proceed to select a next record, as indicated by block **178**, among the unprocessed records.

[0164] Some embodiments may then access content of a most-recent entry of the tamper-evident log, as indicated by block **180**. In some cases, the most-recent entry may be a last block written to the log or a last node within a component of a block, like a leaf node or higher-level node in a binary tree in a block. In some cases, the accessed content may be one or more attributes of the most recent entry, such as one or more attributes of node content. In some embodiments, those attributes may include those like which are described above with reference to nodes in the directed acyclic graphs described with reference to FIGS. **3** and **5**.

[0165] Next, some embodiments may access an abstract cryptographic hash function, as indicated by block **182**. The term “abstract” refers to a programming construct by which the specific implementation of a function or method is defined by other code than that having a placeholder designating the function as abstract (e.g., an object, class, subroutine, or function), and in some cases certain aspects of the abstract function are defined in other code, for example in a class definition in an object-oriented programming language, where a non-abstract implementation of the abstract function is defined upon implementing an object in the class. In another example, the abstract function may be specified as abstract by code configured to receive a lambda function from a calling body of code or from a body of code that calls that is called by the body of code that receives the lambda function. In some cases, the implementation may have access to the scope of the body of code having the reference to the abstract function, e.g., an instance of an object in an abstract class or a method that receives the lambda function as a parameter.

[0166] Some embodiments may determine that an abstract cryptographic hash function has been accessed, for example, upon detecting the presence of a reserved term indicating that this type of function or method has been accessed.

[0167] Some embodiments may then determine an implementation of the abstract cryptographic hash function, as indicated by block **184**. In some cases, the implementation may be determined as

part of instantiating an object within an abstract class, and in some cases, the implementation may be determined by receiving a lambda function defined by calling code or called code. In some embodiments, the implementation may specify a cryptographic hash algorithm, such as SHA-256 or the other examples described above, or in some cases the type of cryptographic hash function may be specified in the specification of the abstract cryptographic hash function, leaving other aspects of the algorithm to be configured in the implementation. In some embodiments, the implementation may specify which inputs are processed by the cryptographic hash function to determine an output, while the abstract representation of the cryptographic hash function may not specify which inputs are used in the implementation. Thus, a single abstract cryptographic hash function may be implemented a variety of different ways within a given body of code, for example, calling those different ways at different times in different scenarios based on different criteria.

[0168] In some embodiments, the implementation of the abstract cryptographic hash function that is determined in block **184** may specify that certain types of inputs are to be used in calculating the cryptographic hash value output. In some embodiments, those types of inputs may be selected among a subset of attributes in node content accessed in block **180**, such as attributes and node content of an adjacent node of a node that is to be added to the tamper-evident log. In some embodiments, the types of inputs may be position-agnostic inputs, such as node attributes that do not indicate a position of the node accessed in block **180** within the tamper-evident log obtained in block **172**. For example, the selected types of inputs may exclude timestamps, dates, counter values that indicate a position in a sequence that is specific to a graph, like an array index, or the like. In some embodiments, the position-agnostic inputs that are selected may include items like attributes that specify a segments content or document content stored in the most-recent entry. In another example, the attributes of the access node content used as the position-agnostic input to the implemented cryptographic hash function may include a cryptographic hash value of a cryptographic hash pointer to another node that is pointed to by the node accessed in block **180**, thereby preserving the benefit of chaining cryptographic hash pointers.

[0169] In some embodiments, the same implementation of the cryptographic hash function may be applied to calculate each cryptographic hash value in each cryptographic hash pointer of a directed acyclic graph encoding the tamper-evident log obtained in block **172**. Or in some cases, different implementations may be accessed at different times for different directed hash pointers. For example, each directed edge encoded by a cryptographic hash pointer may also include a value that identifies the implementation, for example, a value that indicates whether the edge is specified with a position-agnostic implementation of the abstract cryptographic hash function or a position-dependent implementation, the position-implement dependent implementation including as inputs the types of values described above as non-position-agnostic, for example, timestamps dates, position indices, and the like. In some embodiments, the position-agnostic implementation may be selected upon determining that the tamper-evident log in a given directed-acyclic graph, or portion thereof, has reached a threshold size and that a new directed acyclic graph storing subsequent entries to the tamper-evident log is to be created, while the other cryptographic hash function implementations that are not that are position dependent may be used at other times. Or to simplify the code, the same position-agnostic implementation may be used for each cryptographic hash pointer.

[0170] Next, some embodiments may calculate a cryptographic hash value with the implementation determined in block **184** based on the accessed content from block **180**, as indicated by block **186**. In some embodiments, the process **170** further includes forming content of a new entry to the camper-evident log that includes the hash value calculated in block **186** and the current record selected in block **178**, as indicated by block **188**. In some cases, forming new content may include modifying an existing entry or creating a new entry. In some embodiments, forming content may include adding values to attributes in the content of the new entry, such as attributes in the above-described nodes in the graphs of FIGS. **3** and **5**. In some embodiments, the hash value in the current

record may be separately labeled as distinct attributes in the formed content, or in some cases these values may be combined, for example, with a single cryptographic hash value based on both the accessed content from block **180** and the current record. In some embodiments, the current record may be stored remotely, while a hash digest, such as a cryptographic hash value based on that content may be stored in the tamper-evident log obtained in block **172**. Or some embodiments may store the current record in the log itself to expedite access in accordance with some of the techniques described above, for example, with reference to FIGS. **2** and **3**.

[0171] Next, some embodiments may prepend the new entries to the tamper-evident log, as indicated by block **190**. In some embodiments, this may include adding a new entry to a sequence of entries, such as to the sequence of blocks **172**, **174**, and **176** in a data structure **70** of FIG. **3**. The term “prepend” does not require any particular array index or position in some programming construct, as long as the new entry is designated implicitly or explicitly in the tamper-evident log as being newer than existing entries in the tamper-evident log or otherwise older-entries.

[0172] Some embodiments may then determine whether a size of the log is greater than a threshold, as indicated by block **182**. In some embodiments, this may include determining whether the size of the log is greater than or equal to the threshold, a configuration also consistent with the designation greater than, which is used generically herein. Upon determining that a size of the log is not greater than the threshold, some embodiments, may return to determine whether there are more records to process in block **176**. In some embodiments, the size may be measured by an amount of blocks, an amount of nodes, or an amount of data stored within the directed acyclic graph, for example, measured in megabytes.

[0173] Upon determining that a size of the log does exceed the threshold, some embodiments may proceed to split the tamper-evident log into an older tamper-evident log and a newer tamper-evident log, as indicated by block **194**. In some embodiments, the split tamper-evident logs may each be of the same types of directed acyclic graphs, in some cases with the older tamper-evident log being larger at the time of the split. In some embodiments, splitting may be performed by determining to not prepend the new entry or the next entry after the new entry to the tamper-evident log and instead instantiating a new directed acyclic graph to receive a subsequent entry. In some embodiments, the older and newer tamper-evident logs may each be of the same type (e.g., class of topologies) and have the attributes of the directed acyclic graphs described with reference to FIGS. **3** and **5**. In some embodiments, the system may cease adding new entries to the older tamper-evident log and continue adding new entries to the newer tamper-evident log until that log exceeds the threshold, in which case some embodiments may then split that newer tamper-evident log into a yet newer tamper-evident log, continuing indefinitely to grow the number of tamper-evident logs.

[0174] In some embodiments, an initial node or block in those tamper-evident logs may include a cryptographic hash pointer to the older tamper-evident log, for example, a cryptographic hash value based on a root node of a newest block of a consecutively older tamper-evident log. In some embodiments, this referenced content may be used to generate a seed value like in block **78** of FIG. **3** and may be associated with an identifier of the older tamper-evident log. In some embodiments, the newer tamper-evident log may be instantiated on a different storage compute node **26** in FIG. **1** from that of the older tamper-evident log, or it may be instantiated on the same storage compute node **26**.

[0175] Thus, some embodiments may reconfigure directed acyclic graphs while maintaining provable tamper resistance, thereby accommodating arbitrarily large data sets.

Transparent Client Application to Arbitrate Data Storage Between Mutable and Immutable Data Repositories

[0176] As noted, in many cases, traditional databases are not sufficiently secure. Many mainstream databases provide some level of permissioning, but oftentimes these permissions are broad and difficult to manage, which leads to the opportunity for attack and misuse. Further, most widely adopted databases implement transactional SQL (T-SQL) or a similar variant. In most



implementations, poor design practices and legacy code can present vulnerabilities such as SQL injection attacks, which trick the database into returning potentially sensitive and unintended data. Additionally, there is very little auditability and permissioning for individual cells within a table other types of individual table values. Moreover, most database management systems rely on a driver that lives on the database client's machine. Oftentimes, these clients are open source and easily available, making exploitation easier. (None of this is to suggest that some embodiments may not also be afflicted with subsets of these problems, as several inventions are described and those inventions can be used independently without addressing every problem described herein.) [0177] To address such concerns, various approaches have been implemented to secure stored data. Examples include air-gapping the database or carefully managing permission to access the database. Many approaches, however, are difficult to administer (e.g., using baroque permission schemas), are slow (e.g., air-gapping), or break backward compatibility with expensive enterprise applications that expect a particular database driver interface, none of which is to suggest that any of these approaches are disclaimed in all embodiments.

[0178] Some embodiments mitigate some of the problems described above with a specialized client application that is able to identify sensitive data (by column type, predefined column name, or other predefined methods). Some embodiments may capture the data as it is being transmitted to the database and route sensitive data to more secure storage mechanisms, like those mentioned above. Thus, more secure, though potentially slightly slower, storage techniques may be reserved for the most sensitive data, while less sensitive data may be processed with faster and cheaper less secure storage solutions, like traditional relational databases. Further, in some embodiments, the client application may abstract away from other applications this differential routing and variation in query formats between systems.

[0179] In some embodiments, a subset of more sensitive data may be stored by scattering the data among multiple blockchains, as described above. Similarly, when the client application detects that sensitive data is being queried (using the same method), the client application, in some embodiments, may take a TXID (or other pointer) as it is coming from the database, send it to an arbiter instance for reassembly, confirm that the request has valid permissioning and if so, place the reassembled data in place of the TXID.

[0180] Because some embodiments intercept the data in the data path, some embodiments are able to produce an additional audit log which shows all attempts to access the data, as described in greater detail below. In some cases, these access logs can be notated with request-specific information such as: username, geolocation, client machine IP address, etc.

[0181] Through this approach, it is expected that other applications that implement traditional database drivers will require little or no scatter or blockchain-specific configuration. In some cases, the process is completely transparent to other legacy applications. Further, permissioning complexity may be relaxed with secure data routed to distinct, immutable, secure data structures, as access to, and modification of, data may be readily detected.

[0182] Certain types of data are expected to be particularly amenable to use with the present techniques. Often system-access credentials, like user names and passwords, are particularly sensitive, as entire accounts may be compromised if such information is subject to unauthorized access. Storing passwords on a local machine or in a database where the entire password is accessible in one location provides an easy target for threat actors looking to manipulate, steal, or otherwise misuse authentication credentials. Other examples include credit card numbers, social security numbers, or health-related data.

[0183] Some embodiments interface with blockchains as a storage data structure with an arbiter or other piece of middleware that is capable of taking as an input the full text representation of a user credential, starting from the last byte of that credential, fragmenting that credential into N pieces, and placing each piece on a physically (or virtually) separate blockchain backed storage data structure, with each piece containing pointers to the next storage locations of the fragmented

credential. When an application or resource requests the reassembly of a fragmented credential, in some embodiments, an arbiter or piece of middleware is supplied with the location of the first byte of the credential. After reading the first byte, in some embodiments, the arbiter or middleware then reads the subsequent pointers until a null character or end of sequence character is read. Once all of the pieces have been read into memory, the arbiter or other middleware may respond to the application with the resultant unfragmented credential. Some embodiments may preprocess the credential and count the number of pieces that are required from the beginning before fragmenting the credential. Some embodiments may require that credentials yield a threshold number of fragments. Some embodiments may salt fragments or credentials before fragmentation to defeat or impair rainbow table attacks.

[0184] These and other techniques may be implemented with a process in FIG. 8, which shows an example of a process **200** that may be executed by a client computing device to transparently retrofit an existing workload application to interface with a heterogeneous mix of databases and, in particular, with a combination of databases that includes a higher-security database than that which the application is configured to interface with as originally written, such as databases like those described above. It should be emphasized, though, that the present techniques are not limited to embodiments drawing upon the above-types of more secure databases and, in some cases, may be used in conjunction with other types of databases, such as another relational database or other type of datastore, such as one that is deemed to be higher-security or lower latency than that which the application accessing data is configured to interface with. In some embodiments, the process **200** may be executed by the above-described security driver **30**, though it should be noted that in some cases, some or all of the functionality may be executed in the translator **20** in a database gateway, in a database management system, or in some other computing device.

[0185] Executing the process **200** in a client computing device, before data leaves the client computing device, or upon data arriving into the client computing device, is expected to yield certain security benefits in some use cases, where for example, the database that the workload application executing on the client computing device is configured to access has been compromised. In such scenarios, it is likely that an adversary may have compromised other computing devices on a network, and merging or splitting data at the client computing device, immediately before or after writing or reading respectively, is expected to reduce the attack surface of a network. That said, embodiments are not limited to systems providing these benefits, and in some cases, these operations may be performed in another computing device believed to be relatively secure on a network, which is not to suggest that any other feature described herein is limiting.

[0186] In some embodiments, the process **200** may be made transparent to a workload application executing on a client computing device, such as a service on one host of a plurality of hosts executing different services in a micro-services architecture, or an application executing as a monolithic application on a single computing device. In some embodiments, the process **200** may be made transparent to that application by registering the process in the operating system of the client computing device to appear to be the database driver that the workload application is configured to access and then wrapping an application program interface of the original database driver with the operations described below. Thus, some embodiments may be responsive to the same set of application program interface requests that a database driver is responsive to, while providing additional functionality. Further, some embodiments may then pass modified or unmodified application program interface exchanges between the workload application and the database driver. In many cases, source code of the workload application is unavailable or is expensive to modify. Thus, retrofitting existing workload applications in a manner that does not require changes to code of that application is expected to be particularly desirable. That said, the present techniques are also applicable in use cases in which the source code is available for the workload application and is modified to implement the present techniques, which again is not to

suggest that any other description is limiting.

[0187] In some embodiments, the process **200** includes registering a security driver that wraps a database driver, as indicated by block **202**. In some embodiments, the security driver may be registered in an operating system in which a workload application (e.g., application **28** above) making database access request described in subsequent operations of process **200** is executed, and this operating system may also be an environment in which a database driver (e.g., driver **34** above) described below operates. In some embodiments, as a result of the registration operation, when an application sends an application program interface request to the database driver, that request may be received by the security driver instead, and the security driver may be configured to then communicate with the database driver as an intermediary between the database driver and the application within a single operating system on a client computing device.

[0188] In some embodiments, other types of access may be handled in a similar fashion. For instance, some embodiments may wrap a filesystem driver to obtain exchanges between filesystem drivers and workload applications, e.g., operating on documents. In some cases, a file system filter driver may be instantiated that emits events indicative of application program interface exchanges with the filesystem driver and some embodiments may classify these events as pertaining to higher-security documents (or not) and, in some cases, modify the driver behavior in response, e.g., substituting a document stored in the secure distributed storage **16** for an access request that pertains to, e.g., a locally stored text file with a pointer to such a document.

[0189] Some embodiments may then include receiving a first write request, as indicated by block **204**. In some cases, this operation may be performed sometime after registering the security driver and may occur with relative frequency. In some embodiments, the write request is like those described above. In some embodiments, the write request is an application program interface request to the database driver from the workload application executing on the computing device. In some embodiments, the request may be to a first remote database, such as the lower-trust database **14** described above or some other data store. In some embodiments, the first write request may specify that a set of values are to be written to a set of fields in a set of records in a database, which may include adding new values to new fields to new records or modifying existing data. The first write request may be formatted in a schema specified by an application program interface of the database driver and may be operative to cause the database driver to respond regardless of whether the registration step of block **202** has occurred, though the application may then proceed to operate in a less secure fashion in some cases as a result of omitting operations from the security driver.

[0190] Next, some embodiments may obtain a data policy, as indicated by block **206**. In some cases, the security driver **30** may maintain in memory a set of one or more policies that each include a set of rules, such as a policy for each application or each lower-trust database **14**.

[0191] Some embodiments may include classifying values to be written as higher-security values or lower-security values, as indicated by block **208**. In some cases, this operation may include selecting one or more rules from one or more data policies, for example, based on an application writing data, a lower-trust database receiving the data or intended to receive the data, or other criteria. Some embodiments may include applying one or more of the above-described rules to each value in the write request to classify that value as higher-security or lower-security.

[0192] In some embodiments, the rules may each include one or more criteria by which data being written to a database is classified as lower-security data or higher-security data. In some cases, these rules may include rules that designate certain types of fields as lower-security or higher-security, such as text fields versus integer fields, or fields within a specified pattern. In some embodiments, the criteria may explicitly list higher-security fields and lower-security fields, and data may be classified as higher-security or lower-security in virtue of a write request attempting to place data into these fields. In some embodiments, the criteria may apply a regular expression to a field identifier, such as a field name to determine whether values within that field are higher-security or lower-security. In some embodiments, the rules may apply to the content of values being

written, such that some values in a given field may be higher-security, while other values within that given field may be lower-security. For example, some criteria may include a regular expression that pattern matches against the content of values to determine whether those values are higher-security or lower-security, for instance, designating values that have a format consistent with a phone number, credit card number, or Social Security number, regardless of the field in which they belong, as higher security.

[0193] Next, some embodiments may store the lower-security values in the first remote database, as indicated by block **210**. In some cases, the first remote database may be the database that the workload application is configured to interface with when initially installed or off the shelf. Thus, the workload application may be configured to interface with the first remote data face without retrofitting. In some cases, this may include writing values to the lower-trust database **14**, for instance, consistent with the operations described above with reference to block **106** of FIG. **4**.

[0194] Some embodiments may then store the higher-security values in a second remote database, as indicated by block **212**. In some cases, the second remote database may be one of the above-described higher-security databases, such as those hosted within the secure distributed storage **16** of FIG. **1**, implementing the data structures of FIGS. **3** and **5**. Or in some cases, the secure second remote database may be another relational database or other type of database, for instance, one implementing additional security features relative to the first remote database or simply being isolated from the first remote database. In some cases, storing the higher-security values may include the operations described above with reference to FIGS. **2** and **4**.

[0195] Next, some embodiments of the security driver may cause pointers to be stored, for example, storing pointers to the higher-security values in the first remote database, as indicated by block **214**. For example, this may include modifying the application program interface request from the workload application to the database driver to replace higher-security values with node identifiers, segment identifiers, document identifiers, or other types of identifiers like those described above, before advancing the modified application program interface request to the database driver. In some cases, the database driver may then translate the application program interface request into commands and a data stream appropriate to cause the first remote database to store the pointers that identify where the corresponding values are stored in the second remote database.

[0196] Next, embodiments may later read data back. In some cases, this may include receiving a query and then receiving a query response with the pointers included in the query response, as indicated by block **216**. In some embodiments, the query itself may be modified, for example, where a criterion in the query depends upon the content of higher-security values. For example, some embodiments may select all values with a modified query; then within the security driver or the translator described above, replace pointers with the corresponding values retrieved from the second remote database; and then apply criteria of the original query to those values to determine which values are responsive to the query issued by the workload application.

[0197] In some cases, a single query from a workload application may spawn a cascade of iterative, subsequent queries, for example, where join operations are performed, and in which data is merged from the first and second remote databases to determine intermediate query responses. For example, a workload application may request mailing addresses of all users with a credit card number that begins with the specified sequence, and the mailing addresses and credit card numbers may be maintained in different tables, with the credit card numbers designated as higher-security values, and the mailing addresses designated as lower-security values. In some cases, these two different tables may be linked by a primary key in one table that is referenced as a foreign key in another table, and a query from a workload application may specify a join. Some embodiments may retrieve, for example, every record in the first database having pointers to values in the second database that reflect the credit card numbers, merge those values, determine which foreign keys in the table having mailing addresses are responsive to the query criteria, and then issue a subsequent

query to the first remote database for those records. In some cases, to expedite these operations, an index may be maintained in which the pointers are associated with values that indicate whether the values are responsive to certain criteria (e.g., a threshold number of prefix characters or suffix characters), and embodiments may access this index to identify a subset of pointers for which values are retrieved from the secure datastore.

[0198] Some embodiments may interface with databases with the techniques described in U.S. Provisional Patent Application 62/527,330, titled REPLACING DISTINCT DATA IN A RELATIONAL DATABASE WITH A DISTINCT REFERENCE TO THAT DATA AND DISTINCT DE-REFERENCING OF DATABASE DATA (docket no. 043788-0453422), filed 30 Jun. 2017, the contents of which are hereby incorporated by reference.

[0199] In some embodiments, these operations may be expedited by assigning pointers or other types of unique identifiers that are based on the content of the values to which the pointers point, for example, based on cryptographic hash values based solely on the content of the values to which the pointers point. As a result, different instances, of the same value, for example, in different rows or other tuples of a database may correspond to the same pointer. These pointers may be said to be unique identifiers in the sense that they uniquely identify content, in some cases without revealing the semantic information in that content, for instance, with the cryptographic hash identifier, while still having the same unique identifier replicated for multiple instances of that value appearing at multiple rows and a database, for example.

[0200] In some cases, values may be cryptographically hashed in conjunction with a tenant identifier (e.g., by concatenating the values with a delimiter before inputting to a hash function), such as a random string of a certain amount of entropy, like longer than 64 bytes, so that the same value for a given tenant consistently hashes to the same pointer, while the same value for different tenants hash to different pointers. Or in some cases, the unique identifiers may be unique between instances as well. Thus, the same value appearing in two different rows may have a different unique identifier in each instance, though some embodiments may operate more slowly as a result. In some embodiments, the security driver may detect duplicate instances of a pointer when reading back data, and cause a single request from the second remote database for a single value to populate each of these positions held by the different instances of the same pointer. Thus, fields with a relatively low cardinality may still facilitate relatively fast joins even when those fields are designated as relatively high security, and a relatively small number of values may populate a relatively large number of rows.

[0201] Similar operations may be performed when writing, for example, by grouping data classified as high security according to the value, for example, by sorting the data and then detecting groups of instances in which the values are the same, or storing the data in a hash table and detecting duplicates with hash collisions where the same values are written to the same index of the hash table. In these examples, some embodiments may then assign the same unique identifier to each instance in the group where this value is the same, and cause that unique identifier, which may serve as a pointer, to be stored in place of those higher-security values in the first remote database.

[0202] Upon replacing pointers with the values from the second remote database, as indicated in block **218**, some embodiments may provide a modified query response to the querying application, as indicated by block **220**. In some embodiments, the querying application may be a different workload application from that which wrote the data, and in some cases, may be on a different computing device. In some cases, data be read multiple times for a given write, or data may be written multiple times before a most recent version of the value is read.

Immutable Logging of Access Requests to Distributed File Systems

[0203] It is becoming increasingly common to use third party off-site data storage platforms (e.g., Dropbox™ or Google Drive™) as well as on-site platforms (e.g., SharePoint™, Confluence™, OwnCloud™, etc.), but in these systems, auditing changes to these files becomes a difficult and

often impractical task. In some cases, files are stored in large contiguous chunks that pose easy access for threat-actors, and in many cases, changes to the data do not yield a reliable record by which such changes may be detected. For instance, upon penetrating such a system, and making a change, a threat actor may also doctor log records or other records by which a change would otherwise be detectable. (Again, none of which is to suggest that such approaches are disclaimed.) [0204] As noted above, traditional databases do not adequately protect against threat actors or internal resources (employees, information-technology staff, etc.) tampering the data. To the extent such systems have audit logs, those logs are often only as secure as the data for which access is logged, meaning that if the data is compromised, often so is the audit log. Some products offer “secure storage” through the use of password protected folders or data storage areas, however such products generally do not provide immutable and distributed properties and often-times their audit logs can be modified, particularly by employees with elevated-access credentials.

[0205] As noted above, by scattering (e.g., breaking up into segments and distributing) files into one or more blockchains, some embodiments are able to provide immutable and distributed properties to files. Further, since the files are distributed, in some embodiments, there is an application that reassembles the files. By making that reassembly process necessary, some embodiments form a control point by which the embodiment is able to produce a blockchain-backed audit trail of every access to the files (or database values or log entries) stored using this method.

[0206] As noted above, in some embodiments, data access is implemented in an application (e.g., a client-side application) that is capable of listening to events generated by an operating system's filesystem, reading in the changes that caused the notification and reporting them to an application referred to as “an arbiter” instance for scattering and storage in blockchains. In some embodiments, the scatter operation's resultant TXID is what is stored in place of the actual data on the client's filesystem, or similar approaches like those described above may be implemented with the security driver above.

[0207] When a read operation is requested by the filesystem, in some embodiments, the stored TXID (or other pointer) is sent to an arbiter instance for reassembly, loaded in place and then that file's default application handler is opened with that file (e.g., in some cases, a PDF file would generally be opened with Adobe Reader™). If the user does not desire to have a filesystem reader placed on an entire filesystem or subset of a filesystem (a directory), in some embodiments, a user could scatter a file by accessing a context menu and directing the file to be placed in ScatterFS.

[0208] Finally, in some embodiments, an application program interface (API) may be exposed so that any third party application can pass a file-handle, buffer or other data-stream for scattering and storage.

[0209] Thus, some embodiments may distribute data among multiple computing devices, in some cases in a peer-to-peer network hosting an immutable distributed data structure, like a blockchain, and those embodiments may log and store access records in another (or the same) blockchain, thereby monitoring a necessary control point to access the data and creating an immutable record of such access attempts. It should be appreciated that the present techniques are not limited to blockchain-based distributed databases, and similar techniques are contemplated for other distributed file systems, e.g., Network File System, Self-certifying File System, Server Message Block, MapR FS, Amazon S3, and Starfish distributed file systems.

[0210] In some embodiments, logging and log analysis may be implemented with a process **230** shown in FIG. **9**. In some embodiments, the process **230** may be executed at least in part by a system implementing one or more of the above-described techniques. In some embodiments, logging may be implemented on a computing device that serves as a chokepoint in the aggregation or disaggregation of information being stored into or read from a distributed data store like that described above or other types of distributed data stores. In some embodiments, the information may be distributed in such a way that each unit of content requires information from different

computing devices to be accessed (e.g., in virtue of segmenting or separation of encryption keys from cyphertexts), and that information from different computing devices may be aggregated by a computing device that causes logging in a tamper-evident log. Similarly, write operations may be logged by a computing device that causes the information to be broken up and distributed among the different computing devices (e.g. with segmenting or separation of cyphertexts from encryption keys, or combinations thereof). Thus, in some cases, logging occurs during each access operation in virtue of logging being effectuated by a computing device necessary to read or write information that is otherwise inaccessible unless the information passes through that computing device that causes logging. In some embodiments, the information may be logged in a tamper-evident, immutable log like that described above, in some cases in the same data structures that store the information being accessed, such as workload content like database entries and various types of documents in a file system. In some cases, logging may be effectuated, for example, caused by the above-described translator **20** of FIG. **1**, but embodiments are not limited to that implementation, which is not to suggest that any other description is limiting.

[0211] In some embodiments, the process **230** includes receiving a request to access a distributed data store, as indicated by block **230**. In some embodiments, the request may be a read request or write request. In some embodiments, the write request is accompanied with a unit of content, such as a value being displaced with a pointer in the lower-trust database **14** in accordance with the above-described techniques, or a document being replaced by a pointer in a lower-trust file system in accordance with the above techniques, which is not to suggest that documents cannot be stored in a database as values in that database or the any other description herein is limiting. Similarly, and some cases, a read request may be accompanied with a pointer stored in a lower-trust data store in place of the unit of content and read in accordance with the above techniques. In some cases, the unit of content to be accessed is referenced with a pointer to a segment that serves in a first segment in a content graft like that described above, such as a linked list of segments where different segments of the unit of content are distributed among different tamper-evident directed acyclic graphs.

[0212] In some embodiments, the process **230** may include aggregating or disaggregating units of content pertaining to the access request, as indicated by block **234**. In some embodiments, this may include the segmentation operations described above with reference to block **112** or the joining of segments described above with reference to block **164**. In FIGS. **4** and **4**, respectively.

Alternatively, or additionally, this may include encryption and separation for storage of cyphertexts from encryption keys or bringing separately stored encryption keys and cyphertexts together to effectual decryption.

[0213] Before or after aggregating or disaggregating (which is not to suggest that any other step herein is limited to the sequence described in the present example) some embodiments may cause logging of the access request in an entry in a tamper-evident log, and logging may be caused with a computing device participating in aggregating or disaggregating the units of content, as indicated by block **236**. In some embodiments, logging may be caused by a computing device necessary to make the units of content accessible through the aggregating or disaggregating. In some embodiments, logging may be caused by the above-described translator **20**, which is not to suggest that the translator **20** is required in all embodiments consistent with all of the present techniques, or that any other description herein is limiting.

[0214] In some cases, causing logging includes selecting one or more of the above-described directed acyclic graphs, like those discussed with reference to FIG. **5** and stored in the storage compute nodes **26** of FIG. **1**. In some cases, causing logging includes sending an instruction to one or more of those directed acyclic graphs (e.g., to a service monitoring a network socket and managing the graph responsive to such instructions) to store a record describing a logged event in the directed acyclic graphs, for example, storing the record as a document in accordance with the techniques described above with reference to FIGS. **2** and **3**, and in some cases fragmenting the

record into multiple segments in accordance with the techniques described above with reference to FIGS. 4 and 5. Or in some cases, the record may be stored outside of the tamper-evident log, and a cryptographic hash of the record and a timestamp of the record may be stored as node content of one of the above-described tamper-evident directed acyclic graphs having cryptographic hash pointers as edges. In some embodiments, in virtue of these graphs, modifications to records describing log entries may be computationally infeasible to conceal, as the chain sequence of cryptographic hash values in the directed acyclic graphs based upon those records may create an insurmountable computational challenge to calculate hash collisions along the entire sequence of a path through the directed acyclic graph that collides with the values produced by the original unaltered record. (It should be noted that attributes of nodes may be stored in edges or vice versa.)

[0215] In some embodiments, the process **230** includes logging entries that describe a variety of different aspects of the access request. In some embodiments, the entry is documented with a record that includes an identifier of a user account making the access request, an identifier of an application, such as a workload application through which the access request was submitted, a content of the access request, such as a command including content to be written or identifying units of content to be read. Such records may further include a timestamp, such as a date or date and time indicating when the access request was received. The records in some cases may include a geolocation of a computing device submitting the request. In some embodiments, such records documenting entries in the tamper-evident log may further include identifiers of computing devices through which the access requests were submitted to the above-describe system, such as MAC addresses, Internet Protocol addresses, browser fingerprints, hardware fingerprints, or the like.

[0216] Some embodiments may store the tamper-evident log in memory, as indicated by block **238**, which in some cases may include adding a new entry or modifying an entry in a tamper-evident log already stored in memory, which may be replicated in multiple instances in accordance with the techniques described above. This may include forming part of a block to be added at a later time to such a log, e.g., a blockchain.

[0217] Some embodiments of the process **230** may include validating the tamper-evident log, as indicated by **240**. Validation may include determining whether the tamper-evident log indicates that logged entries have been modified after being logged. In some embodiments, this operation may include accessing a current log version of an entry and then calculating a cryptographic hash value based on that current log entry. Some embodiments may then compare that current cryptographic hash value to one or more other cryptographic hash values in the tamper-evident log, for example, some embodiments may compare that cryptographic hash value to a cryptographic hash value stored in an adjacent node in one of the above-described directed acyclic graphs to determine whether the cryptographic hash values match or, if they do not match, indicate that the record was modified. In some embodiments, a sequence of cryptographic hash values based upon one another may be calculated to determine which match and identify a path through a directed acyclic graph to a node where the cause of a discrepancy resides. In some embodiments, validation may occur upon each write operation, periodically, upon each read operation, or according to some other schedule or event.

[0218] Detectability of tampering may deter threat actors from modifying log entries or maliciously accessing data in virtue of the difficulty of modifying log entries in a concealable manner. Further, some embodiments may fragment log entries in accordance with some of the above-describe techniques, further making modifications difficult, as a heterogeneous set of different computing devices on different networks may need to be compromised to modify each segment even if a threat actor somehow was able to compromise the mechanisms by which tampering is indicated.

[0219] In some embodiments, the process **230** includes determining whether the validation operation evinces tampering, as indicated by block **241**, for example indicating a mismatch between cryptographic hash values within a directed acyclic graph. Upon detecting tampering, some embodiments may emit an alarm, as indicated by **242**. In some embodiments, the alarm may



be emitted by sending an email, text message, or chat message. Some embodiments may further proceed to take certain operations to locked down portions of a system, for example, disabling credentials, reducing levels of access associated with credentials, adjusting some of the below-described thresholds to decrease an amount of data that may be accessed before subsequent access is slowed or blocked, or limiting access to lower-security data.

[0220] In some embodiments, upon determining that there is no evidence of tampering, or after emitting the alarm, some embodiments may proceed to determine a risk metric based on access requests documented in the tamper-evident log, as indicated by block **244**. Risk metrics may take a variety of different forms, and in some cases risk metrics may be calculated for a variety of different entities. For example, different risk metrics may be calculated for different users, workload applications, computing devices, or combinations thereof. In some embodiments, the risk metric may be an amount of access requests received within a trailing duration of time, a total amount of access requests, or combination thereof. In some embodiments, the risk metric is based on deviation from previous patterns of behavior. For example, some embodiments may train a machine learning algorithm, such as a hidden Markov model, recurrent neural network, or the like, based on historical log events, to predict the likelihood of various types of access requests (or sequences of such requests), such as to particular units of content, types of unit of content, amounts of access requests, frequencies of access requests, or the like, for a given user, workload application, computing device, portion of a network, or combination thereof. Some embodiments may then compare these predictions based on a current trailing sequence of logged events with later received access requests and determine a risk metric based on the likelihood, for example, a probability of the given access requests given previous behavior. Failures of such predictive models may be evidence of anomalous, malicious behavior. Some embodiments may use this probability is a risk metric or determine an aggregate of these probabilities over a plurality of risk metrics, such as a measure of central tendency, like a mean, median, or mode of these probabilities over a trailing duration or number of access requests.

[0221] In another example, the risk metric may be based on the content of units of content being written. For example, some embodiments may calculate an entropy of units of content being written and compare that entropy to measures of entropy associated with other units of content historically written to the distributed data store, for example, previous database entries or documents. In some embodiments, this difference may be indicative of a ransomware attack in which relatively high entropy encrypted versions of data are being written as a current version. (Though it should be noted that some implementations may use an immutable data store in which earlier values remain in place and systems may be rolled back to earlier values in the event of such an attack in some embodiments by replacing pointers in lower-trust storage to current versions to be pointers to last-known good versions.)

[0222] Next, some embodiments may determine whether the risk metric exceeds or otherwise satisfies a first threshold, as indicated by block **246**. Comparisons to thresholds described herein should not be limited to a comparison that depends upon the sign of the values applied, which is not to suggest that any other description is limiting. The term “satisfies” is used generally to refer to scenarios in which one threshold may be satisfied by exceeding (or being equal to in some cases) that threshold and an equivalent threshold in which values are was multiplied by  $-1$  may be satisfied by being less than (or being equal to in some cases) that threshold, or vice versa. Upon determining that the risk metric satisfies the first threshold, some embodiments may block subsequent responses to subsequent access requests, as indicated by block **248**. In some cases, blocking may be implemented by the computing device participating in aggregating or disaggregating the units of content discussed above with reference to block **236**.

[0223] If access is not blocked, some embodiments may compare the risk metric to a second threshold, as indicated by block **250**. In some embodiments, the second threshold may be less stringent than the first threshold, for example, corresponding to lower levels of risk. In some

embodiments, upon determining that the risk metric satisfies the second threshold, some embodiments may delay subsequent responses to subsequent access requests, as indicated by block **252**. In some embodiments, this may include starting a countdown timer and determining when a designated duration of time has elapsed before aggregating or returning units of content or disaggregating and writing units of content, in some cases for each unit of content in an access request pertaining to a plurality of units of content. Thus, some embodiments may implement a soft blocking mechanism by which functionality is provided at a certain level, while delaying in providing time for a human response in the event that a scripted attack is occurring. In some embodiments, upon satisfying the first or second thresholds in block **246** and **250**, some embodiments may emit an alarm using techniques like those described above to facilitate investigation and, if needed human intervention. Embodiments may then return to wait for the next request to access the distributed data store in block **232**.

#### Storing Differentials of Files in a Distributed Blockchain

[0224] Often, blockchain-based databases are not well suited for storage of large, frequently modified collections of data, like files or other binary blobs of data. Because of the immutable nature of blockchains, previous entries in a chain (which may represent files or values in a database) generally cannot be deleted or overwritten. Thus, each version of a file can add to the size of a blockchain ledger, and where the file is stored in the ledger and modified frequently, the size of a blockchain ledger can become too expansive to be efficiently accessed.

[0225] To combat the above-noted negative consequences of the immutable property of blockchains, some embodiments store only changes (differentials) of files over time, rather than entire copies of a file at each change.

[0226] Some embodiments may receive a write request for a modified file (or other blob of data) and determine whether the data has changed from a previous version. In some embodiments, upon reading the data initially, hash digest (e.g., a MD5 hash) of the data may be calculated and held in memory. Upon a write, a new hash may be calculated based on the data to be re-written to memory, and that hash may be compared to the earlier hash to determine whether the file has changed.

[0227] Upon detecting a change, some embodiments may determine a delta between an earlier version (e.g., a most recent version) and a current version. In some cases, the database may store a delta of an earlier change, and some embodiments may iterate through a sequence of deltas to re-create a previous version of a document or other file. Upon obtaining the most recent version, some embodiments may determine a delta with the new version, e.g., by determining a set of longest common subsequences between the versions and store the resultant diff, e.g., in unified format.

[0228] In most cases, it is expected that the resultant diff will be much smaller than the new version of the file (or other blob). As such, storing the diff in the blockchain is expected to be less computationally expensive than storing the entire new version. Further, because the file itself is ultimately stored in the blockchain (rather than just a hash digest), the system is expected to be more robust to various attacks, such as a hash collision attack. In such attacks, malicious content is selected and designed to produce the same hash as stored data, and that malicious data is substituted in the database for the authentic data. With traditional systems, the blockchain will yield a hash value that validates the malicious content as authentic. In contrast, some embodiments circumvent this attack vector by storing the data committed to the database in the blockchain. Further, in contrast to systems that merely store a hash digest of a document in the chain, some embodiments offer increased control of the file (or other BLOB). In these older systems that only store a hash digest, the system does not have no control of the file. Such systems could delete the file from the external data-store, and all that would be left with in a chain is the hash digest. That is, such systems can authenticate a file, but they cannot reproduce the file. That said, not all embodiments afford these benefits, as various engineering and cost tradeoffs are envisioned, and multiple independently useful inventions are described, which is not to suggest that any other description is limiting.

[0229] In some embodiments, the above techniques may be implemented with processes described below with reference to FIGS. **10** and **11** that operate upon a data structure described below with reference to FIG. **12**.

[0230] As shown in FIG. **10**, some embodiments may include a process **260** that writes a difference between a new version of a document and previously stored versions of a document to a tamper-evident, immutable data repository, such as a block chain or one of the above-described examples of directed acyclic graphs having cryptographic hash pointers. In some embodiments, the process **260** may be executed by the above-described security driver **30** in conjunction with the above-described translators **20** or by other components, for example, based on a gateway on a network residing between network storage or a database and a client computing device. In some embodiments, the process **260** may be executed by a client-side application that wraps or otherwise interfaces between a client-side workload application and a database driver or a filesystem driver, for example, a file system driver executing in the same operating system as the workload application and the process performing the operations of FIGS. **10** and **11**.

[0231] In some embodiments, a filesystem of a local computing device includes a hierarchy of directories having files arranged therein, for example, binary large objects with various examples of metadata, like file names, creation dates, modification dates, authors, permissions to access, and the like. In some embodiments, these directories may be local directories stored on the computing device executing a workload application or on a remote network attached storage device. In some embodiments, some of the files in the directories may be replaced with text files having computer readable pointers to, for example, documents stored with techniques like those described above, and some embodiments may intercept (e.g., receive directly, pass through, observe upon filtering, etc.) read or write access requests by a filesystem explorer application to a filesystem driver and detect these pointers disposed in the directories in place of the documents. Some embodiments may then, in response, effectuate the corresponding operations on documents stored in a tamper-evident, immutable data repository, like those described above or other types of remote storage. In some cases, this may include creating new versions, updating pointers stored in the directory, and various other operations. Further, these techniques may similarly be applied to database read and write operations, for example, storing differences between previously stored values and databases that are relatively large and new versions of those values in databases.

[0232] In some embodiments, the process **260** includes receiving a request to write a new version of a document to a tamper-evident, immutable data repository, as indicated by block **262**. In some embodiments, this operation may be performed by the above-described security driver or other client-site arbiters, for example, with the process of FIG. **8**. In some embodiments, the request may include a designation of the document as higher security and a pointer to a previous version of the document stored in the client-side computer accessible directory or database. For example, the user may have previously read the previous version into local memory of the client computing device workload application, transformed the previous version, and then requested to store the new version, causing the request to be received in the operation of block **162**.

[0233] Next, some embodiments may determine that the new version of the document is different from the previous version of the document, as indicated by block **264**. In some embodiments, this may be performed by retrieving the previous versions from the tamper-evident, immutable data repository, for example, with a read process like that described below with reference to FIG. **11**. Or some embodiments may expedite this determination by storing outside of the tamper-evident, immutable data repository, an index that associates pointers to documents with hash digests of those documents. In some embodiments, the hash digest is a cryptographic hash value based upon the content of the document, or a non-cryptographic hash value based upon the content of the document. It should be noted, that not all hash functions are cryptographic hash functions having the attributes described above is being exhibited by such functions. In some cases, non-cryptographic hash functions may be faster to compute than cryptographic hash functions, or

embodiments may use cryptographic hash functions to enhance security. Accordingly, some embodiments may make the determination of block **264** without retrieving the previous version of the document from the tamper-evident, immutable data repository, for example, by calculating a new hash digest based on the new version and comparing that new hash digest to a previously calculated hash digest of the previous version stored outside the tamper-evident, immutable data repository, for example, in association with a pointer to that previous version in an index. Upon determining that the hash values match, some embodiments, may terminate the process **260**, as the previous version may be determined to be the same as the new version.

[0234] Alternatively, upon determining that the new version of the document is different from the previous version of the document, some embodiments may proceed to obtain the previous version of the document from the tamper-evident, immutable data repository, as indicated by block **266**. In some embodiments, this may include retrieving the previous version with read operations, for example, like those described with reference to FIG. **11**, which in some cases may engage the processes described above with reference to FIGS. **2**, **5**, **6**, and **7**.

[0235] Next, some embodiments may determine a set of changes required to transform the previous version into the new version of the document, as indicated by block **268**. In some embodiments, this operation may be relatively computationally expensive and include determining a longest matching substring between the two documents in the course of determining a minimum set of changes required to transform the previous version into the new version. In some cases, these changes may be limited to deletions, appending text (e.g., prepending or postpending), and changing values of existing text. In some embodiments, the determination may be made with a diff operation in which the previous version and the new version are input into a diff function, which may return the set of changes. In some embodiments, the changes may be implemented with the Hunt-McIlroy algorithm, as described in a paper titled AN ALGORITHM FOR DIFFERENTIAL FILE COMPARISON by Hunt and McIlroy, Bell Laboratories Computing science technical report, **1976**, the contents of which are hereby incorporated by reference.

[0236] In some cases, to expedite comparisons, each line of the two documents may be transformed into a hash digest of that line that is relatively fast to operate upon, for example, converting each line with a hash function into a 64 bit or shorter, such as a 32 bit or shorter or 16 bit or shorter hash digest of the content of that respective line. Thus, each of the two documents may be transformed into an ordered list of hash digests, and some embodiments may then compare the two ordered list of hash digest to identify sequences of text that are likely to have not changed between the two versions. For example, some embodiments may identify a number of lines at a beginning and an end of the list where there are no changes. Upon detecting a line at which changes begin, some embodiments may then search forward in one list until a match is found, thereby potentially detecting insertions.

[0237] Some embodiments may then store the set of changes in the tamper-evident, immutable data repository, as indicated by block **270**. In some cases, this may include performing the operations of FIG. **2** on the data structure of FIG. **3** to store a document, or segmenting the document in accordance with the techniques described above with reference to FIGS. **5** and **6** and storing the document in that manner.

[0238] Some embodiments may further include storing a pointer to the previous version of the document in association with the set of changes, as indicated by block **272**. In some embodiments, the pointer to the previous version of the document may be stored in the tamper-evident immutable data repository as metadata to the set of changes, thereby forming a node and edge of a version graph, like that described below with reference to FIG. **12**. In some embodiments, documents may be encoded as a linked list of sets of changes with pointers between the sets of changes tracing back to an initial version of the document, for example, in the above-described data structures, and documents may be retrieved by retrieving each of the sets of changes and then iteratively applying them, as described in greater detail below with reference to FIG. **11**.

[0239] In some embodiments, an updated pointer to the newly stored set of changes from block **270** may be stored in the file system (or database cell) in place of the previous pointer submitted with the request to write the new version in block **262**. Thus, a subsequent read request for that file may identify that new pointer and retrieve the new set of changes and then trace back to the initial version through the previous version, as described in greater detail below with reference to FIG. **11**. Similar operations, like those described above, may be performed on a value stored in a database, replacing a pointer in a database cell to reference the new version.

[0240] In some embodiments, version graphs may branch. For example a user may request to write to a previous version of a document or to write to new versions to a previous version of the document, thereby forming parallel change of a version graph that share a subset of their version history with one another but then branch a part later in the version graph. This is expected to further compress data and increase the amount of documents that can be stored in a given amount of storage in a tamper-evident, immutable data repository.

[0241] FIG. **11** shows an example of a process **280** to read a document, or other unit of content, from a tamper-evident, immutable data repository storing version graphs of documents as sequences of differences between versions of documents tracing back to initial version. In some embodiments, the process **280** may be performed in the course of obtaining the previous version of the document in block **266** of FIG. **10**, or upon receiving a request by workload application to access a stored document, for example, upon intercepting a request to a file system driver to access a document replaced with a pointer to a document stored in one of the above-describe data structures.

[0242] In some embodiments, the process **280** includes receiving a read request identifying a pointer to a most recent version of a document in a tamper-evident, immutable data repository, as indicated by block **282**. In some cases, the pointer may be to a previous version, for example, upon a user requesting to roll back to a previous version. In some embodiments, the read request may be intercepted with the techniques described above by which write requests are intercepted.

[0243] Next, some embodiments may read the version located by the current pointer, as indicated by block **284**, for example by performing the read operations described above by which documents are retrieved in the data structures of FIG. **3** or FIG. **5**, or other database entries are received.

[0244] Some embodiments may then determine whether the red version is the initial version of the document or other unit of content, as indicated by block **286**. In some cases, an initial versions may be explicitly flagged as such when stored, or initial versions may be implicitly designated as such in virtue of lacking a pointer to an earlier version. As noted above, pointers may be stored in association with subsequent versions, as discussed in reference to block **272**.

[0245] Upon determining that the version retrieved was not the initial version, some embodiments may designate a pointer associated with the retrieved version, for example, stored as metadata of the retrieved version in the tamper-evident, immutable data repository, as the current pointer, as indicated by block **288**.

[0246] Next, some embodiments may add a set of changes associated with the retrieved version, for example, stored at a location or sequence of locations in the tamper-evident, immutable data repository identified by the current pointer, to a last in first out buffer, as indicated by block **290**. In some cases, the buffer may be characterized as a stack. Some embodiments may then return to block **284** to read the next earlier version located by the updated current pointer in block **284**.

[0247] Upon encountering an initial version in block **286**, some embodiments may proceed to block **292** and initialize a copy of the document to the initial version, for example, setting a working copy of the document to be equal to the initial version that was retrieved. Some embodiments may then determine whether there are more changes in the buffer to process, indicated as indicated by block **294**. In some cases, some embodiments may determine whether the buffer is empty or there are more values, that is sets of changes, stacked in the buffer.

[0248] Some embodiments may, upon determining that there are more changes, retrieve a next set

of changes from the buffer, as indicated by block **296**, in some cases this may include deleting that next set of changes from the buffer to update the buffer to reflect the next set of changes are to be applied. In some cases, this may be characterized as popping a set of changes from a stack. [0249] Next, some embodiments may apply the next set of changes to the current working copy of the document, as indicated by block **298**. In some cases, this may include accessing a set of deletion operations, change operations, and append operations. In some embodiments, these changes may each be associated with a line number to which the change is to be applied, a sequence with which the change is to be applied, and a content in some cases to be applied, for example indicating replacement text or appended text (or bits).

[0250] Some embodiments may then return to block **294** to determine whether there more changes in the buffer. Upon reaching the end of the buffer, for example, the bottom of the stack, some embodiments may proceed to block **300** and return the copy of the document. In some cases, the copy of the document may be returned to a workload application configured to access the document, for example, based on a file system extension mapped in the operating system to a particular workload application. In some embodiments, the read operation may be transparent to a user, and may appear to the user as if the user is operating on a locally stored or network stored copy of a document, with the user experience being identical or substantially identical to that experienced by a user who is not interfacing with the above-described secure distributed storage **16**, thereby providing higher security without imposing on users or in some cases requiring retrofits of workload applications that can be expensive. (That said, embodiments are not limited to systems affording these benefits, which is not to suggest that other descriptions are limiting.)

[0251] In some embodiments, the processes of FIGS. **10** and **11** may produce a data structure **310** like that shown in FIG. **12**. In some embodiments, the data structure **310** may include the verification graphs described above with reference the data structure **130** of FIG. **5**, along with the content graph **138** described above with reference to FIG. **5**. In this example, the content graph **138** may be a most recent stored version of a document, which may be a set of changes from an earlier version. In some embodiments, the earlier version may, for example, be an initial version of a document, stored in another content graph **314** shown in FIG. **12**. Content graph **314** may, like content graph **138** include a sequence of content nodes, such as segments, connected by pointers between the segments, by which an initial version of a document or a set of changes associated with a given version may be assembled during a read operation. In some embodiments, thus, the data structure **310** may store a third overlaid graph, which may be a version graph **312**. In this example, version graph **312** includes an edge **316**, which may be a pointer from the most current version stored in content graph **238** to an earlier, initial version stored in content graph **314**. In some embodiments, the pointer to may be to an initial segment of the content graph storing the earlier version.

[0252] FIG. **13** is a diagram that illustrates an exemplary computing system **1000** in accordance with embodiments of the present technique. Various portions of systems and methods described herein, may include or be executed on one or more computer systems similar to computing system **1000**. Further, processes and modules described herein may be executed by one or more processing systems similar to that of computing system **1000**.

[0253] Computing system **1000** may include one or more processors (e.g., processors **1010a-1010n**) coupled to system memory **1020**, an input/output I/O device interface **1030**, and a network interface **1040** via an input/output (I/O) interface **1050**. A processor may include a single processor or a plurality of processors (e.g., distributed processors). A processor may be any suitable processor capable of executing or otherwise performing instructions. A processor may include a central processing unit (CPU) that carries out program instructions to perform the arithmetical, logical, and input/output operations of computing system **1000**. A processor may execute code (e.g., processor firmware, a protocol stack, a database management system, an operating system, or a combination thereof) that creates an execution environment for program instructions. A processor may include a

programmable processor. A processor may include general or special purpose microprocessors. A processor may receive instructions and data from a memory (e.g., system memory **1020**).

Computing system **1000** may be a uni-processor system including one processor (e.g., processor **1010a**), or a multi-processor system including any number of suitable processors (e.g., **1010a-1010n**). Multiple processors may be employed to provide for parallel or sequential execution of one or more portions of the techniques described herein. Processes, such as logic flows, described herein may be performed by one or more programmable processors executing one or more computer programs to perform functions by operating on input data and generating corresponding output. Processes described herein may be performed by, and apparatus can also be implemented as, special purpose logic circuitry, e.g., an FPGA (field programmable gate array) or an ASIC (application specific integrated circuit). Computing system **1000** may include a plurality of computing devices (e.g., distributed computer systems) to implement various processing functions.

[0254] I/O device interface **1030** may provide an interface for connection of one or more I/O devices **1060** to computer system **1000**. I/O devices may include devices that receive input (e.g., from a user) or output information (e.g., to a user). I/O devices **1060** may include, for example, graphical user interface presented on displays (e.g., a cathode ray tube (CRT) or liquid crystal display (LCD) monitor), pointing devices (e.g., a computer mouse or trackball), keyboards, keypads, touchpads, scanning devices, voice recognition devices, gesture recognition devices, printers, audio speakers, microphones, cameras, or the like. I/O devices **1060** may be connected to computer system **1000** through a wired or wireless connection. I/O devices **1060** may be connected to computer system **1000** from a remote location. I/O devices **1060** located on remote computer system, for example, may be connected to computer system **1000** via a network and network interface **1040**.

[0255] Network interface **1040** may include a network adapter that provides for connection of computer system **1000** to a network. Network interface **1040** may facilitate data exchange between computer system **1000** and other devices connected to the network. Network interface **1040** may support wired or wireless communication. The network may include an electronic communication network, such as the Internet, a local area network (LAN), a wide area network (WAN), a cellular communications network, or the like.

[0256] System memory **1020** may be configured to store program instructions **1100** or data **1110**. Program instructions **1100** may be executable by a processor (e.g., one or more of processors **1010a-1010n**) to implement one or more embodiments of the present techniques. Instructions **1100** may include modules of computer program instructions for implementing one or more techniques described herein with regard to various processing modules. Program instructions may include a computer program (which in certain forms is known as a program, software, software application, script, or code). A computer program may be written in a programming language, including compiled or interpreted languages, or declarative or procedural languages. A computer program may include a unit suitable for use in a computing environment, including as a stand-alone program, a module, a component, or a subroutine. A computer program may or may not correspond to a file in a file system. A program may be stored in a portion of a file that holds other programs or data (e.g., one or more scripts stored in a markup language document), in a single file dedicated to the program in question, or in multiple coordinated files (e.g., files that store one or more modules, sub programs, or portions of code). A computer program may be deployed to be executed on one or more computer processors located locally at one site or distributed across multiple remote sites and interconnected by a communication network.

[0257] System memory **1020** may include a tangible program carrier having program instructions stored thereon. A tangible program carrier may include a non-transitory computer readable storage medium. A non-transitory computer readable storage medium may include a machine readable storage device, a machine readable storage substrate, a memory device, or any combination thereof. Non-transitory computer readable storage medium may include non-volatile memory (e.g., flash

memory, ROM, PROM, EPROM, EEPROM memory), volatile memory (e.g., random access memory (RAM), static random access memory (SRAM), synchronous dynamic RAM (SDRAM)), bulk storage memory (e.g., CD-ROM and/or DVD-ROM, hard-drives), or the like. System memory **1020** may include a non-transitory computer readable storage medium that may have program instructions stored thereon that are executable by a computer processor (e.g., one or more of processors **1010a-1010n**) to cause the subject matter and the functional operations described herein. A memory (e.g., system memory **1020**) may include a single memory device and/or a plurality of memory devices (e.g., distributed memory devices). Instructions or other program code to provide the functionality described herein may be stored on a tangible, non-transitory computer readable media. In some cases, the entire set of instructions may be stored concurrently on the media, or in some cases, different parts of the instructions may be stored on the same media at different times.

[0258] I/O interface **1050** may be configured to coordinate I/O traffic between processors **1010a-1010n**, system memory **1020**, network interface **1040**, I/O devices **1060**, and/or other peripheral devices. I/O interface **1050** may perform protocol, timing, or other data transformations to convert data signals from one component (e.g., system memory **1020**) into a format suitable for use by another component (e.g., processors **1010a-1010n**). I/O interface **1050** may include support for devices attached through various types of peripheral buses, such as a variant of the Peripheral Component Interconnect (PCI) bus standard or the Universal Serial Bus (USB) standard.

[0259] Embodiments of the techniques described herein may be implemented using a single instance of computer system **1000** or multiple computer systems **1000** configured to host different portions or instances of embodiments. Multiple computer systems **1000** may provide for parallel or sequential processing/execution of one or more portions of the techniques described herein.

[0260] Those skilled in the art will appreciate that computer system **1000** is merely illustrative and is not intended to limit the scope of the techniques described herein. Computer system **1000** may include any combination of devices or software that may perform or otherwise provide for the performance of the techniques described herein. For example, computer system **1000** may include or be a combination of a cloud-computing system, a data center, a server rack, a server, a virtual server, a desktop computer, a laptop computer, a tablet computer, a server device, a client device, a mobile telephone, a personal digital assistant (PDA), a mobile audio or video player, a game console, a vehicle-mounted computer, or a Global Positioning System (GPS), or the like. Computer system **1000** may also be connected to other devices that are not illustrated, or may operate as a stand-alone system. In addition, the functionality provided by the illustrated components may in some embodiments be combined in fewer components or distributed in additional components. Similarly, in some embodiments, the functionality of some of the illustrated components may not be provided or other additional functionality may be available.

Using a Tree Structure to Segment and Distribute Records Across One or More Decentralized, Acrylic Graphs of Cryptographic Hash Pointers

[0261] Many extant blockchain-based databases are not well suited for certain use cases, particularly those involving latency-sensitive access. Many file-related tasks within normal everyday computing are expected to occur within an accepted tolerance of latency. This is particularly true for certain operations as intensive as file-systems or file operations, where data stores preferably operate within accepted latency tolerances.

[0262] Storing data in a blockchain in a 'linear' fashion often provides immutable and distributed data storage, however 'read' operations often cannot be made concurrent. Without concurrency, reads must be sequential and due to the nature of typical blockchains, will be too slow for use in many applications (which is not to suggest that this arrangement or any other herein is disclaimed). Scattering data and storing it in a 'linked-list' format (also can be thought of as a 1-degree tree) means that when reading, the system only has knowledge of the next 1 'chunk' of data, causing read latency to scale at least linearly ( $O(n)$  where  $n$  is the number of segments in Big O notation)



with the number of segments of a record.

[0263] This problem is mitigated by some embodiments that arrange the data into a multiple-degree tree, for this example, a binary (2-degree) tree. An example is shown in FIG. 15, which is a data structure that may be implemented with the system of FIG. 1 and process of FIG. 14. Or some embodiments may store the data in higher-degree trees. Storing this way is expected to allow a scatter client (described above with reference to FIG. 1) or other computing entity to read the data in halves (or larger groups) concurrently, each “chunk” points to child nodes that contain the preceding two portions of data (as data may be written in reverse order). Because the operations to retrieve data may be systemically distinct and linearizable, these operations can be performed concurrently in some embodiments. In some embodiments, the theoretical optimal tree-degree is equal to the degrees of parallelism offered by the implementing system

[0264] In some embodiments, the techniques above may be expedited with data structures and processes described below with reference to FIGS. 14 through 17. Some embodiments may facilitate concurrent writing or reading of segments from the above-described verification graphs with a process 400 shown in FIG. 14 to produce and read from data structures like those described with below with reference to FIG. 15. In some embodiments, those data structures may be updated to form data structures like those shown in FIGS. 16 and 17 to facilitate relatively fast writes that avoid rewriting portions of a record that have not changed when a version is updated.

[0265] In some embodiments, these techniques may be implemented with data structures like those described above with reference to FIG. 5, like the above-described content graphs stored within the above-described verification graphs. In some embodiments, though, the content graphs may include content nodes with multiple directed content edges pointing two or more other content nodes to facilitate concurrent reads to or writes of those other content nodes. In some embodiments, the data structures may be implemented as binary trees, quad trees, k-ary trees (or other rooted trees), where k is an integer and indicates a maximum number of edges pointing from a given node to other nodes within a content graph. In some embodiments, more than 10%, more than 50%, more than 95%, or all non-leaf nodes in a content graph may include k pointers to other nodes. In some embodiments, the contents graphs may be hierarchical tree data structures, starting with a root node and extending through intermediate branching layers to leaf nodes, in some cases with segments of records being stored in each level of the tree, or in subsets of the levels, like leaf nodes.

[0266] It should be emphasized that the content graph in some embodiments is distinct from the above-described verification graphs, despite being stored within a verification graph in some embodiments. In some embodiments, there may be a one-to-one correspondence between content nodes and verification nodes, with one and only one content node stored in each verification node, but content pointers (i.e., content edges) between content nodes may be different from pointers between verification nodes, for example, pointing to different collections of data. In some embodiments, pointers between content nodes may refer to other content nodes by identifiers of verification nodes storing those other content nodes, or embodiments may implement a distinct namespace for identifiers of content nodes and include a data structure that maps between a namespace of identifiers of content nodes and a namespace of identifiers of verification nodes. Embodiments may reference this data structure to translate between the two names cases.

[0267] Or in some embodiments, a single verification node may store multiple content nodes, or a given content node may be split among multiple verification nodes. In some embodiments, a content graph may be distributed among multiple verification graphs, for example, with different content nodes in a given content graph stored on different verification graphs, and in some cases pointers therebetween may include identifiers of verification graphs along with identifiers of verification nodes within those verification graphs.

[0268] In some embodiments, the process 400 may be executed by one of the above-described security drivers, which in some cases may read receive read and write events, like access requests directed to a database or to a file system from an operating system of a client computing device,

like those described above with reference to FIG. 1. In some embodiments, the stored records may be database entries, documents, directories, binary large objects (blobs), or various other collections of data, such as files having metadata, which is not to suggest that any of these categories cannot overlap or the any other listing of examples herein refers to distinct categories. [0269] In some embodiments, the process **400** includes determining whether a write request has been received, as indicated by block **402**. Upon determining that a write request has been received, some embodiments may receive a request to store a record corresponding to that write request, as indicated by block **404**. Next, some embodiments may encode the record in segments, as indicated by block **406**. In some embodiments, this operation may include the above-described transformations by which records are segmented to form node content of content nodes of a content graph. Examples include segmenting a file by breaking up a bitstream into contiguous blocks of less than or equal to a threshold size, like two or more blocks, 10 or more blocks, or 50 or more blocks. Other examples include other ways of breaking up a record into parts that do not individually reveal the content of the record. In some cases, the encoding includes breaking up a record or a portion thereof, into two binary sequences that when combined in an XOR operation re-create a portion of the record. In other embodiments, breaking up the segments may include applying various linear coding techniques, like turbo-codes, Hamming codes, and run-length coding to encode portions of a record. Some embodiments may partition records to maintain less than a threshold amount of entropy in given segments by determining when a current segment of a bitstream exceeds the threshold and defining a segment end at a bit before where the threshold is crossed.

[0270] Next, some embodiments may determine whether the record to be stored is a new version of an existing record, as indicated by block **408**. In some cases, this may include referencing a stored index that maps a filename of the record to a root node to determine whether that file name has been stored, or referencing a format of the record to determine whether the received record is received in association with a pointer to a root node of an existing content graph storing a previous version of the record.

[0271] Upon determining that the request is not a request to store a new version of an existing record, some embodiments may assign the segments to levels of a hierarchy of a tree, as indicated by block **410**. In some embodiments, the tree may include two or more, three or more, five or more, or 10 or more levels, ranging from leaf nodes to a root node. Examples are described below with reference to FIGS. 15 through 17. Or in some embodiments, other content graphs that facilitate concurrent retrieval may be implemented, such as a content graph encoded as a complete or strongly connected graph, such that any content node may serve as an entry point for concurrent retrieval of every other node of the content graph.

[0272] In some embodiments, the segments may be assigned to levels according to the sequence with which the segments are assembled to form the record, such as a first collection of bits in a bitstream of the record, followed by a second collection of bits to be appended thereto, and so on through a last collection of bits to be appended to the others to re-create a bitstream of the record. In some embodiments, assigning the segments to a level may include assigning the segments to levels based on their position in the sequence, for example, by determining a number of segments, dividing that number by some value to produce a fraction, rounding that fraction down to the nearest integer to produce an amount for a current level, and then assigning that amount from the existing segments, for example starting from a prefix position of the segments to the current level. For example, where  $k$  is equal to two in a binary tree, some embodiments may assign half of the remaining segments that have not yet been assigned to a level to the current level. Some embodiments may repeat this process to determine how many are assigned to the current level, until a single segment or selected member of remainder remains to be assigned to a root node.

[0273] In some embodiments, assigning segments to a level of a hierarchy of a tree may include expressing the number of segments in a base- $k$  representation. In some embodiments, each digit of

the base-k representation may then be multiplied by k to the power of the digit's position in the representation minus 1. For example, 21 segments may be arranged in a quad tree by converting the value 21 into base four, which is 111. The first digit (1, in position 3) may then be multiplied by four to the power of 3-1 (i.e.,  $4^{\text{circumflex over } ( )}2$ ) to determine that the first 16 segments are to be assigned to a lowest level of the hierarchy, serving as leaf nodes of the tree. The second digit may then be multiplied by four to the power of 2-1 (i.e.,  $4^{\text{circumflex over } ( )}1$ ) to determine that four segments are assigned to an intermediate layer of the hierarchy. And the last digit may be multiplied by four to the power of 1-1 (i.e.,  $4^{\text{circumflex over } ( )}0$ ) to determine that a last single segment is assigned to a highest level. In instances in which more than a single segment is assigned to the highest level, some embodiments may shift one member of that otherwise highest level to an added even higher level to produce a single root node. Or some embodiments may execute an analogous operation in the initial assignment to a lowest level of nodes so that the next layer after that lowest level of nodes in each layer after that corresponds to the value of k taken to an integer power. For each node in a given level, k-members of the next lower level may be assigned and edges pointing to those nodes may be added, in some cases, with no overlap in sets assigned to each node of a given level.

[0274] In some cases, in operation **408**, it may be determined that the request is a request to store a new version of an existing record. Some embodiments may then determine which segments have change between the previous version and the new version, as indicated by block **412**. In some cases, this operation may include executing the below-described read operation to retrieve those segments from a verification graph and comparing the retrieved copies to the new segments to determine which changed. Alternatively, to expedite operations, some embodiments may store an index associated with each record with a list of cryptographic hash representations of each segment of a record. Some embodiments may then recalculate those cryptographic hash representations on the updated record, and compare for each segment, the cryptographic hash representations to determine which segments have changed and thereby produce a difference in the cryptographic hash values, without reading the records from the verification graph.

[0275] Some embodiments may then select content nodes on paths (through the content graph) to changed segments in block **414**, for example, as described below with reference to FIG. 16. In some cases, the operations of block **410** may then be executed on the segments selected in operations from blocks **412** and **414**.

[0276] Some embodiments may select an initial level and then, with each level-iteration, determine whether there are more levels of a hierarchy to process in block **416**. Upon determining that there are more levels, some embodiments may select a next level, as indicated by block **418**. Some embodiments may begin with a highest level of a hierarchy and increment downwards towards leaf nodes. Or to reduce the amount of memory consumed, some embodiments may start with a lowest level of a hierarchy and increment upwards from the leaf nodes towards the root node, such that when a given content node is added with a given segment, the content nodes to which that content node points are extant in program state.

[0277] Some embodiments may increment through the segments of a given level and perform certain operations. Thus, some embodiments may determine at each segment whether there are more segments remaining in the current level to process, as indicated by block **420**. Upon determining that there are more segments to process, some embodiments may select a next segment, as it can by block **422**, and form a content node with that segment, as indicated by block **424**. In some cases, this may include storing the segment as content of that content node that is formed. Some embodiments may then return to block **420**. These operations may continue until it is determined that there are no more segments in a current layer, in which case program flow may return to block **416** where it is determined whether there are more levels to process.

[0278] Upon determining that all the levels have been processed, some embodiments may proceed to block **426** to arrange the formed content nodes in a content graph. In some embodiments, this

operation may include forming data structures like those described below with reference to FIGS. 15 through 17, which may include defining edges of the content graph that point from one content node to another. Those edges may connect each node with the selected sets in the lower level as described above. In some embodiments, the edges and the content of the content nodes may be stored in a single data structure, like an object in an object-oriented programming language. [0279] The results of the arrangement may form a content graph, which may be stored in a verification graph, as indicated by block 430. In some cases, this may include the operations described above with reference to FIG. 5 and elsewhere by which content nodes are stored in immutable, tamper evident, directed acyclic graphs of cryptographic hash pointers, for example with the content nodes and the content graph edges (pointing from the respective content nodes) stored as content of verification nodes in the verification graph.

[0280] Some embodiments may then return a root content node identifier to a computing entity that requested the write operation, as indicated by block 432. In some cases, the content node identifier is an identifier of a verification graph and a verification node within that verification graph that includes the root content node of the content graph. In some embodiments, for example, where multiple content nodes are stored in an individual verification node, the identifier may further include an identifier of a content node within that verification node. In some embodiments, the computing entity may be an application executing on a computing device, another computing device that sent an application program interface request over a network, or the like.

[0281] Program flow may then return to block 402, where upon determining that a write request has not been received, embodiments may determine whether a read request is received, as indicated by block 431. Upon determining that no read request has been received, embodiments may return to block 402 and continue waiting for a read or write request.

[0282] Upon receiving a read request, some embodiments may proceed to block 434 and retrieve a record from the content graph and verification graphs. To this end, some embodiments may receive a request to read the record, as indicated by block 434. This request, in some cases, may include a root node identifier of the records to be retrieved, for example, of a version of the record to be retrieved, like a most current version. Some embodiments may proceed to retrieve the root content node from a verification graph in a verification node of the verification graph identified by the root node identifier, as indicated by block 436.

[0283] Some embodiments may then traverse the content graph from the root node, as indicated by block 438, retrieving from verification multiple nodes of the verification graph. Some embodiments may execute a depth first traversal or a breath first traversal of the content graph. Some embodiments may execute this traversal by recursively calling a function with arguments that include an identifier of a current node of the content graph, and that function may retrieve identifiers of adjacent content nodes to the current content node in the content graph pointed to by edges from the current content node. In some cases, the function may call itself with each of those identified pointed to content nodes as arguments in a function call to recursively traverse the content graph. The function may return each segment stored in each content node as node content to retrieve the segments of the version of the record.

[0284] Some embodiments may verify that the segments are consistent with the verification graph upon retrieval, as indicated by block 440. In some cases, this may include the above-described operations by which tampering is detected based upon internally inconsistent cryptographic hash pointers within the verification graph. Upon detecting an inconsistent cryptographic hash value indicative of tampering, some embodiments may emit an alarm, log an alarm, and otherwise signal that the record may have been compromised.

[0285] Some embodiments may then proceed to return the record to the calling computing entity, as indicated by block 442. In some cases, this may include decoding the segments into the record, reversing the operations of block 406, for example, by concatenating the segments, executing the above-described XOR operation, or otherwise reconstituting the record from the retrieved segments

that have been verified.

[0286] FIG. 15 shows an example of a content graph produced with the process 400. The illustrated example is a binary tree, in this case a balanced binary tree, but embodiments are consistent with other types of data structures, such as quad trees, like k-ary trees, skip lists, or combinations thereof. As noted, in some cases, each node of the content graph 500 may store a segment as content, or in some cases, a subset of the nodes may store content and other nodes may serve as indices identifying those nodes. Further, as noted, the nodes of the content graph 500 may be stored in one or more verification nodes like those described above with reference to FIG. 5.

[0287] In this example, the content graph 500 includes four different hierarchical levels, ranging from leaf nodes at level 502, to intermediate levels 504 and 506, to a single root node in level 508. In this example, each of the content nodes 510 includes an alphabetic letter indicating the sequence of the segment of a record stored in that content node. In this example, the first segment of a record is stored in the content node labeled with reference letter A, and the last segment is stored in the root content node with the reference letters O. Or in other embodiments, this arrangement may be reversed or otherwise differently arranged.

[0288] To read the record from the content graph 500, embodiments may receive an identifier of the root content node at level 508 and then execute a depth-first traversal or breath-first traversal of the tree to retrieve each segment. In some embodiments, each content graph stored in a verification graph may store one and only one record, or in some cases multiple records may be stored in a content graph, or a record may be divided among multiple content graphs, for example with multiple identifiers to different root nodes serving as identifiers of the record. A collection of verification graphs may store hundreds, thousands, millions, or more content graphs.

[0289] FIG. 16 shows an updated data structure with an updated content graph 512 after a new version of the record is stored. The illustrated content graph 512 includes the above-described content graph 500 as a sub-graph, and another sub-graph labeled 514. The sub-graph labeled 514 includes modified versions of segments that either changed or are stored in content nodes that point to segments that are changed. In this example, for example, the content nodes labeled E' and F' have changed, while the content nodes labeled K', N', and O' point to those nodes directly or indirectly and thus are included in sub-graph labeled 514. To form this sub-graph 514, some embodiments may determine which nodes of the graph 500 are on a path to content nodes storing segments that have changed.

[0290] In this example, the content node oh O' prime may serve as a root node by which the updated version of the record may be retrieved, and the updated version of the record may be written without rewriting all of the segments. For example, all but one of the segments may be maintained, more than 20% of the segments may be left unaltered, more than 80% of the segments may be left and altered, or more than 95% of the segments may be left unaltered in the updated version of the content graph 512. In this example, extant portion of the immutable, tamper evident data structure may be left unaltered, with added records serving as an encoding of the new version, and some of the old content nodes forming part of the encoding of the new version. Thus, the content graph 500 storing the initial version of the record may partially overlap the portions of the content graph 512 that are reachable from the root node 510.

[0291] Again, the updated version of the record may be retrieved by executing a depth-first or breadth-first traversal of the graph 512, beginning with the new root node O'. Or the old version of the record may be retrieved by executing the same traversal from the root node O. In both cases, the same record for, for example, content node M may serve to reconstruct both versions of the record.

[0292] FIG. 17 shows a third version of the record being written to form content graph 516, which has as sub-graphs the above-described content graphs 500 and 512. A new portion of the graph 518 includes the segments that have changed (or on paths to segments that have changed) since the previous version was stored in the sub-graph 512 described above. Again, the third version may be

retrieved by recursively traversing the graph **516** beginning with the root content node **O**", with some of the segments in the third record being retrieved from the second version and some of the segments being retrieved from the first version of the record that was stored. Thus, new versions may be written relatively quickly without rewriting portions of a record that have not change between versions. Or some embodiments may store a new content graph with each version and not reuse content nodes between versions, which is not to suggest that any other feature herein may not also be modified in some implementations.

[0293] In some cases, for example where pointers from content nodes point to multiple child nodes, retrieval of records may be expedited by concurrently retrieving those different records. In some embodiments, to facilitate concurrent operations, a content graph may be split among multiple verification drafts systematically. For example, at a threshold number of levels of a hierarchy, all child nodes below that level may be stored on a different verification graph. In some cases, read latency may scale with the number of segments at less than  $O(n)$ , where  $n$  is the number of segments in Big O notation.

[0294] Those skilled in the art will also appreciate that while various items are illustrated as being stored in memory or on storage while being used, these items or portions of them may be transferred between memory and other storage devices for purposes of memory management and data integrity. Alternatively, in other embodiments some or all of the software components may execute in memory on another device and communicate with the illustrated computer system via inter-computer communication. Some or all of the system components or data structures may also be stored (e.g., as instructions or structured data) on a computer-accessible medium or a portable article to be read by an appropriate drive, various examples of which are described above. In some embodiments, instructions stored on a computer-accessible medium separate from computer system **1000** may be transmitted to computer system **1000** via transmission media or signals such as electrical, electromagnetic, or digital signals, conveyed via a communication medium such as a network or a wireless link. Various embodiments may further include receiving, sending, or storing instructions or data implemented in accordance with the foregoing description upon a computer-accessible medium. Accordingly, the present techniques may be practiced with other computer system configurations.

[0295] In block diagrams, illustrated components are depicted as discrete functional blocks, but embodiments are not limited to systems in which the functionality described herein is organized as illustrated. The functionality provided by each of the components may be provided by software or hardware modules that are differently organized than is presently depicted, for example such software or hardware may be intermingled, conjoined, replicated, broken up, distributed (e.g. within a data center or geographically), or otherwise differently organized. The functionality described herein may be provided by one or more processors of one or more computers executing code stored on a tangible, non-transitory, machine readable medium. In some cases, notwithstanding use of the singular term "medium," the instructions may be distributed on different storage devices associated with different computing devices, for instance, with each computing device having a different subset of the instructions, an implementation consistent with usage of the singular term "medium" herein. In some cases, third party content delivery networks may host some or all of the information conveyed over networks, in which case, to the extent information (e.g., content) is said to be supplied or otherwise provided, the information may provided by sending instructions to retrieve that information from a content delivery network.

[0296] The reader should appreciate that the present application describes several independently useful techniques. Rather than separating those techniques into multiple isolated patent applications, applicants have grouped these techniques into a single document because their related subject matter lends itself to economies in the application process. But the distinct advantages and aspects of such techniques should not be conflated. In some cases, embodiments address all of the deficiencies noted herein, but it should be understood that the techniques are independently useful,

and some embodiments address only a subset of such problems or offer other, unmentioned benefits that will be apparent to those of skill in the art reviewing the present disclosure. Due to costs constraints, some techniques disclosed herein may not be presently claimed and may be claimed in later filings, such as continuation applications or by amending the present claims. Similarly, due to space constraints, neither the Abstract nor the Summary of the Invention sections of the present document should be taken as containing a comprehensive listing of all such techniques or all aspects of such techniques.

[0297] It should be understood that the description and the drawings are not intended to limit the present techniques to the particular form disclosed, but to the contrary, the intention is to cover all modifications, equivalents, and alternatives falling within the spirit and scope of the present techniques as defined by the appended claims. Further modifications and alternative embodiments of various aspects of the techniques will be apparent to those skilled in the art in view of this description. Accordingly, this description and the drawings are to be construed as illustrative only and are for the purpose of teaching those skilled in the art the general manner of carrying out the present techniques. It is to be understood that the forms of the present techniques shown and described herein are to be taken as examples of embodiments. Elements and materials may be substituted for those illustrated and described herein, parts and processes may be reversed or omitted, and certain features of the present techniques may be utilized independently, all as would be apparent to one skilled in the art after having the benefit of this description of the present techniques. Changes may be made in the elements described herein without departing from the spirit and scope of the present techniques as described in the following claims. Headings used herein are for organizational purposes only and are not meant to be used to limit the scope of the description.

[0298] As used throughout this application, the word “may” is used in a permissive sense (i.e., meaning having the potential to), rather than the mandatory sense (i.e., meaning must). The words “include”, “including”, and “includes” and the like mean including, but not limited to. As used throughout this application, the singular forms “a,” “an,” and “the” include plural referents unless the content explicitly indicates otherwise. Thus, for example, reference to “an element” or “a element” includes a combination of two or more elements, notwithstanding use of other terms and phrases for one or more elements, such as “one or more.” The term “or” is, unless indicated otherwise, non-exclusive, i.e., encompassing both “and” and “or.” Terms describing conditional relationships, e.g., “in response to X, Y,” “upon X, Y,” “if X, Y,” “when X, Y,” and the like, encompass causal relationships in which the antecedent is a necessary causal condition, the antecedent is a sufficient causal condition, or the antecedent is a contributory causal condition of the consequent, e.g., “state X occurs upon condition Y obtaining” is generic to “X occurs solely upon Y” and “X occurs upon Y and Z.” Such conditional relationships are not limited to consequences that instantly follow the antecedent obtaining, as some consequences may be delayed, and in conditional statements, antecedents are connected to their consequents, e.g., the antecedent is relevant to the likelihood of the consequent occurring. Statements in which a plurality of attributes or functions are mapped to a plurality of objects (e.g., one or more processors performing steps A, B, C, and D) encompasses both all such attributes or functions being mapped to all such objects and subsets of the attributes or functions being mapped to subsets of the attributes or functions (e.g., both all processors each performing steps A-D, and a case in which processor 1 performs step A, processor 2 performs step B and part of step C, and processor 3 performs part of step C and step D), unless otherwise indicated. Further, unless otherwise indicated, statements that one value or action is “based on” another condition or value encompass both instances in which the condition or value is the sole factor and instances in which the condition or value is one factor among a plurality of factors. Unless otherwise indicated, statements that “each” instance of some collection have some property should not be read to exclude cases where some otherwise identical or similar members of a larger collection do not have the property, i.e., each

does not necessarily mean each and every. Limitations as to sequence of recited steps should not be read into the claims unless explicitly specified, e.g., with explicit language like “after performing X, performing Y,” in contrast to statements that might be improperly argued to imply sequence limitations, like “performing X on items, performing Y on the X'ed items,” used for purposes of making claims more readable rather than specifying sequence. Statements referring to “at least Z of A, B, and C,” and the like (e.g., “at least Z of A, B, or C”), refer to at least Z of the listed categories (A, B, and C) and do not require at least Z units in each category. Unless specifically stated otherwise, as apparent from the discussion, it is appreciated that throughout this specification discussions utilizing terms such as “processing,” “computing,” “calculating,” “determining” or the like refer to actions or processes of a specific apparatus, such as a special purpose computer or a similar special purpose electronic processing/computing device.

[0299] In this patent, certain U.S. patents, U.S. patent applications, or other materials (e.g., articles) have been incorporated by reference. The text of such U.S. patents, U.S. patent applications, and other materials is, however, only incorporated by reference to the extent that no conflict exists between such material and the statements and drawings set forth herein. In the event of such conflict, the text of the present document governs.

## Claims

1. A tangible, non-transitory, machine-readable medium storing instructions that when executed by one or more processors effectuate operations configured to expedite retrieval of records written to persistent storage, the operations comprising: receiving, with one or more processors, a request to read data stored in a relational database, the request being received from a computing entity and comprising a user identifier; accessing, with one or more processors, one or more log entries associated with prior read requests made using the user identifier; determining, with one or more processors, whether the data to be returned in response to the request is to be obfuscated based on an access policy associated with the user identifier, the access policy comprising one or more access control rules that define conditions for access to values stored in respective fields of the relational database; and returning, with one or more processors, a response to the request comprising data from the relational database with one or more values in the response obfuscated in accordance with the access control rules.
2. The medium of claim 1, further comprising: monitoring, with one or more processors, a number of read requests associated with the user identifier, each read request having an associated time duration; and determining, with one or more processors, whether a number of read requests having time durations that exceed a defined duration threshold satisfies a predefined count threshold.
3. The medium of claim 2, further comprising providing, upon determining that the number of requests satisfies the duration threshold, an alarm.
4. The medium of claim 3, wherein providing the alarm comprises: creating a notification comprising an indication of the user identifier and the number of requests; and transmitting the notification to a designated recipient or system.
5. The medium of claim 1, further comprising: evaluating, with one or more processors, a set of access control rules associated with the data, wherein: the access control rules are defined in association with one or more user attributes stored in an access policy repository; and the user attributes include at least three of a department identifier, clearance level, jurisdictional scope, or group membership of the computing entity.
6. The medium of claim 5, wherein evaluating the set of access control rules comprises retrieving a policy from a repository stored separately from the relational database and applying the policy to determine access to one or more data entries in the relational database.
7. The medium of claim 1, wherein the obfuscation to be performed based on the access policy is applied to at least some of the returned data in response to the request.



8. The medium of claim 1, wherein the obfuscation comprises applying at least one of the following techniques to one or more values in the returned data: applying a cryptographic hash to a value, redacting the value, replacing the value with a constant, inserting a placeholder value in place of an original value, or omitting the value from the response.
9. The medium of claim 1, further comprising parsing, with one or more processors, a structured query language statement included in the request to identify one or more query fields associated with data stored in the relational database.
10. The medium of claim 1, wherein the data returned in response to the request is obfuscated in accordance with the access control rules, and wherein the returned data is obfuscated without modifying the data stored in the relational database.
11. The medium of claim 1, further comprising: parsing, with one or more processors, a structured query language statement included in the request to identify one or more fields associated with the data stored in the relational database; and determining, with one or more processors, whether any of the fields identified in the request are associated with a sensitivity classification defined in metadata associated with the relational database.
12. The medium of claim 11, further comprising enforcing, with one or more processors, a masking rule or access restriction based on the sensitivity classification prior to returning the response.
13. The medium of claim 11, further comprising selecting, with one or more processors, a masking rule or access restriction to be applied based on the sensitivity classification associated with the identified fields.
14. The medium of claim 1, wherein: the data stored in the relational database is stored in one or more content nodes of a content graph; the content graph comprises a plurality of content nodes connected by directed content edges, wherein each directed content edge includes a cryptographic hash pointer based on one or more attributes of a target content node, the cryptographic hash pointer comprising a hash of content of the target content node and an identifier of the target content node; modification of content stored in any content node results in one or more cryptographic hash pointers within the content graph being inconsistent with a previously calculated hash value; and the content graph is decentralized across a plurality of storage compute nodes, wherein different subsets of the content nodes are stored at different compute nodes that are configured to replicate and verify content using the cryptographic hash pointers.
15. The medium of claim 1, wherein the computing entity is an application executing on a user device.
16. The medium of claim 1, further comprising: monitoring, with one or more processors, a number of read requests associated with the user identifier that result in obfuscation of one or more values based on the access control rules; and determining, with one or more processors, whether a count of such read requests satisfies a threshold condition associated with the access policy.
17. The medium of claim 1, further comprising exporting, with one or more processors, one or more access log entries associated with the request to read the data to a storage object maintained in a cloud-based object storage service, wherein the access log entries comprise a user identifier, a timestamp, and a reference to a field of the relational database accessed in response to the request.
18. The medium of claim 1, wherein determining whether the data to be returned in response to the request is to be obfuscated based on an access policy comprises steps for classifying values as higher security values or lower security values.
19. The medium of claim 1, wherein determining whether the data to be returned in response to the request is to be obfuscated based on an access policy comprises steps for determining a risk metric.
20. A method to expedite retrieval of records written to persistent storage, comprising: receiving, with one or more processors, a request to read data stored in a relational database, the request being received from a computing entity and comprising a user identifier; accessing, with one or more processors, one or more log entries associated with prior read requests made using the user identifier; determining, with one or more processors, whether the data to be returned in response to

the request is to be obfuscated based on an access policy associated with the user identifier, the access policy comprising one or more access control rules that define conditions for access to values stored in respective fields of the relational database; and returning, with one or more processors, a response to the request comprising data from the relational database with one or more values in the response obfuscated in accordance with the access control rules.

---