

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250260878

Kind Code

A1

Publication Date

August 14, 2025

Inventor(s)

Phillips; Christopher

SYSTEMS AND METHODS FOR EFFICIENT MANAGEMENT OF RESOURCES FOR STREAMING INTERACTIVE MULTIMEDIA CONTENT

Abstract

Interactive video game content for a video game is streamed to a client device using a first streaming protocol configured for low-latency streaming. When an indication of non-interactive content to be inserted into the video game is detected, streaming of the interactive video game content to the client device is paused and the non-interactive content is streamed to the client device using a second streaming protocol configured for adaptive bitrate streaming. Upon receiving an indication that streaming of the non-interactive content to the client device has ended, streaming of the interactive content to the client device is resumed using the first streaming protocol.

Inventors: Phillips; Christopher (Hartwell, GA)

Applicant: Adeia Guides Inc. (San Jose, CA)

Family ID: 1000008575097

Appl. No.: 19/169858

Filed: April 03, 2025

Related U.S. Application Data

parent US continuation 18384074 20231026 parent-grant-document US 12301955 child US 19169858

parent US continuation 17721874 20220415 parent-grant-document US 11838453 child US 18384074

Publication Classification

Int. Cl.: H04N21/81 (20110101)

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS [0001] This application is a continuation of U.S. patent application Ser. No. 18/384,074, filed Oct. 26, 2023, which is a continuation of U.S. patent application Ser. No. 17/721,874, filed Apr. 15, 2022, now U.S. Pat. No. 11,838,453, the disclosures of which are hereby incorporated by reference herein in their entireties.

BACKGROUND

[0002] This disclosure is directed to reducing bandwidth requirements for interactive content. In particular, techniques are disclosed for reducing bandwidth requirements by delivering non-interactive content inserted into the interactive content using adaptive bitrate streaming.

SUMMARY

[0003] Online purchasing of games with the ability to download, install, and run the game locally has been around for a long time. Video game companies are now starting to offer services for users to subscribe and play games that run in the network. This is known as cloud gaming. Cloud gaming is gaining traction with game studios, internet companies and network operators. There are currently products on the market like Google's Stadia™, Nvidia's GeForce Now™ and Sony's PS Now™ which offer remote game streaming as a business. Amazon is also entering into the business of cloud gaming.

[0004] Cloud gaming uses a technique called remote rendering. Remote rendered games run somewhere in the network, typically at a network edge location, and encode the video, audio and haptics output of the rendered game. The encoded video, audio, and haptics are transported or streamed over the network using a low latency protocol such as RTP to the client device. The client device decodes the video, audio and haptics data streams and renders those to the display the game on a client device. The controller input is consumed on the local device and delivered over the network to the remotely running game engine. In remote rendered cloud gaming, the game engine is running at the edge and rendering the video output.

[0005] Cutscenes are included by game studios and provide a video cinematic experience to enhance the game experience and to mask long load times when transitioning from one main game play level/world to the next. Game level load times can be as much as 30-50 seconds when transitioning from one level to the next. Rather than show a busy cursor like an animated icon, game engines offer the ability to render a cutscene which ties to the game play and involves computer animated in-game characters acting out scenes which tell a story for the upcoming level. These cutscenes can be up to 5-10 minutes. The game engine will allow for bypassing the cutscene once the next game level has been loaded allowing the user to continue after the required loading time.

[0006] Many newer game cutscenes are now interactive and are customized to in-game purchases and upgrades. When games are running locally and the cutscenes are generated from the same game engine, this is basically an extension of the game, and it is straight forward to offer basic interactivity and generate a custom cutscene based on the customization of the player's character. In cloud gaming, it is helpful to offer the same experience without consuming the managed network resources and the bandwidth typically required to run a game in its main gaming mode.

[0007] Additionally, many extended reality (XR) (e.g., augmented reality (AR) or virtual reality (VR)) developers have begun taking advantage of remote rendering techniques to support richer and more enhanced AR and VR applications. Game engines such as Unreal Engine, Blender, and

Unity, are being leveraged by AR and VR developers to support AR and VR applications.

[0008] In managed or semi-managed networks (e.g., 5G mobile networks), a dedicated slice for cloud gaming can be set up at the edge of the network. The “edge of the network” or “network edge” often refers to a network node or nodes accessible by the client device providing the game to the end-user. These nodes may provide, to the client device, computational resources or infrastructure to execute the game or portions of the game. Load balancing may be an issue with cloud gaming and edge computing. Load bearers may be assigned with a priority or requested bandwidth for the video, audio, and haptics stream delivery. In some instances, cloud gaming systems that enable clients to take advantage of resources at the network edge enable a client device to provide a game to an end-user without having the game or other software installed on the client device. In some instances, a client device may have an app installed to access a cloud gaming library but may not need the games installed in order for the user to play the games on the client device. In any event, cloud gaming may demand a high level of edge computing and network communication. For example, delivering non-interactive cinematic cutscenes or advertisements using a low latency encoding scheme using an IP GOP structure and a minimal user equipment buffer used in active game play can result in a much higher bitrate encoding than necessary due to the ultra-low latency encoding requirement in remote rendered cloud gaming. This results in a much higher bandwidth usage over a managed network such as a 5G mobile network. Rendering and encoding cutscenes also requires more GPU resources at the network edge.

[0009] In light of problems stemming from the resource-intensive nature of cloud gaming and edge computing, particularly with respect to rendering and encoding cutscenes (e.g., in a gaming or XR context), what is needed is a technique for conserving bandwidth and computing resources that are allocated for cloud based real-time interactive content, such as gaming, AR applications, and VR applications. Such resources may be conserved by switching from network managed real-time encoding and delivery of media data to over-the-top (OTT) delivery of media data using adaptive bitrate (ABR) encoding for portions of the content that do not require ultra-low latency encoding. Examples of such content may include game cutscenes, supplemental content, and advertisements. This also significantly reduces the file size of the interactive content, as the data required to render non-interactive content is less than that of interactive content. Thus, the storage requirements at network edge locations are reduced. Additionally, this allows for faster distribution time for the interactive content to be distributed from the content provider to network edge locations.

[0010] Systems and methods are described herein for reducing bandwidth requirements for streaming content for an interactive multimedia presentation (e.g., a video game, an AR presentation, or a VR presentation). Interactive content is streamed to a client device using a first streaming protocol or technique. The first streaming protocol or technique may be an ultra low-latency protocol configured for low-latency real-time streaming of media data, such as the Real Time Transport Protocol (RTP). This type of protocol is ideal for real-time media delivery such as for interactive content, as any buffering or lagging of interactive content significantly impacts the user experience of such content. In some instances, the first streaming protocol or technique may be any suitable technique that relies on encoding content at a single bit rate. When an indication of non-interactive content, or supplemental content (e.g. an interactive cutscene for a video game), to be inserted into the interactive multimedia presentation is detected, streaming of the interactive content to the client device is paused. The non-interactive or supplemental content may be streamed to the client device using a second streaming protocol or technique, such as the Hypertext Transfer Protocol (HTTP) live streaming protocol (HLS). For example, the streaming of the interactive content may be monitored for an indication of non-interactive or supplemental content to be inserted, or for whether the streaming of the interactive content has ended or has temporarily paused. The second streaming protocol may be a protocol configured for adaptive bitrate streaming. Non-interactive or supplemental content does not require low latency, and buffering or lagging of non-interactive or supplemental content will not impact the media content (e.g., game

play). Therefore, more traditional adaptive bitrate streaming can be effectively utilized to deliver the highest quality non-interactive or supplemental content to the client device while consuming the least amount of bandwidth and other system resources. In some embodiments, the second streaming protocol or technique may be any suitable technique that relies on adaptive bit rate encoding to encode content at multiple bitrates. Upon receiving an indication that streaming of the non-interactive or supplemental content to the client device has ended, streaming of the interactive content to the client device is resumed using the first streaming protocol or technique.

[0011] In some embodiments, the indication that streaming of the non-interactive or supplemental content to the client device has ended is a user input to skip a portion of the non-interactive or supplemental content item. The user input is received during the streaming of the non-interactive or supplemental content. In response to receiving the user input, streaming of the non-interactive or supplemental content to the client device is stopped and streaming of the interactive video game content to the client device is then resumed.

[0012] In some cases, the non-interactive or supplemental content may be chosen or customized for each particular user. For example, the non-interactive or supplemental content may be an advertisement and so an advertisement for an item in which the user may be interested can be selected. As another example, the non-interactive or supplemental content may be a cutscene between two portions of interactive gameplay that may change based on the type of character the user selected. In a fantasy role-playing game (RPG), the user may be able to choose their character as one of several different races (e.g., elf, dwarf, mage, etc.), and cutscenes may include content that is specific to each race. Other examples include changes to the appearance of the user's character due to in-game acquisitions, upgrades, and purchases of clothing, weaponry, armor, or other items. Thus, once a game session is established, an identifier of the user may be identified. A profile is then retrieved based on the identifier. The profile may be a game-specific profile that includes information about the user's character in the video game or may be a broader profile that includes information used for targeted advertising. Based on the profile, the non-interactive content item is chosen for streaming to the client device. For customized and interactive supplemental content, such as cutscenes, two

[0013] interactive content engines may be used. One engine at the network edge runs the main game levels. Another engine (e.g., dedicated to playing cutscenes, advertisements, or other supplemental content, any of which may be interactive) runs at another location, either at an edge of a Hyper Cloud Provider (HCP) or at another location in the operator's network. The system automatically switches between engines based on needed cutscene playout. There are two means of delivering the customized or interactive cutscene. It can be streamed over an unmanaged network, internal to the operator or from an HCP, by using a real-time transport protocol such as RTP leveraging low latency, low loss, scalable throughput (L4S) or using a live adaptive bitrate (ABR) protocol (e.g., HLS, MPEG-DASH, Microsoft Smooth Streaming, Adobe Dynamic HTTP Streaming) leveraging Common Media Application Format (CMAF) for low latency ABR delivery.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0014] The above and other objects and advantages of the disclosure will be apparent upon consideration of the following detailed description, taken in conjunction with the accompanying drawings, in which:

[0015] FIG. 1 shows an example of an interactive multimedia presentation with supplemental content inserted therein, in accordance with some embodiments of the disclosure;

[0016] FIG. 2 is an example of a system for providing an interactive multimedia presentation in which interactive content and supplemental content are provided from different sources, in

accordance with some embodiments of the disclosure;

[0017] FIG. **3** is an example of a system for providing an interactive multimedia presentation in which interactive content and supplemental content are provided from the same source, in accordance with some embodiments of the disclosure;

[0018] FIG. **4** is an example of a content delivery network for providing an interactive multimedia presentation, in accordance with some embodiments of the disclosure;

[0019] FIG. **5** is a block diagram showing components and data flow therebetween of a system for providing an interactive multimedia presentation in which interactive content and supplemental content are provided from the same source, in accordance with some embodiments of the disclosure;

[0020] FIG. **6** is a block diagram showing components and data flow therebetween for a system for providing an interactive multimedia presentation in which interactive content and supplemental content are provided from different sources, in accordance with some embodiments of the disclosure;

[0021] FIG. **7** is a block diagram showing components and data flow therebetween of an interactive content server when supplemental content is to be inserted into an interactive multimedia presentation, in accordance with some embodiments of the disclosure;

[0022] FIG. **8** is a block diagram showing components and data flow therebetween of an edge content management server for providing supplemental content, in accordance with some embodiments of the disclosure;

[0023] FIG. **9** is a flowchart representing a process for providing an interactive multimedia presentation, in accordance with some embodiments of the disclosure;

[0024] FIG. **10** is a flowchart representing a process for determining whether supplemental content is to be inserted into an interactive multimedia presentation, in accordance with some embodiments of the disclosure;

[0025] FIG. **11** is a flowchart representing a process for skipping a portion of supplemental content, in accordance with some embodiments of the disclosure;

[0026] FIG. **12** is a flowchart representing a process for selecting supplemental content where the interactive multimedia presentation comprises a video game, in accordance with some embodiments of the disclosure;

[0027] FIG. **13** is a flowchart representing a process for delivering a video game cutscene prior to start of the video game, in accordance with some embodiments of the disclosure;

[0028] FIG. **14** is a flowchart representing a process for delivering a video game cutscene between levels of the video game, in accordance with some embodiments of the disclosure; and

[0029] FIG. **15** is a flowchart representing a process for managing bandwidth usage during delivery of a video game cutscene, in accordance with some embodiments of the disclosure.

DETAILED DESCRIPTION

[0030] FIG. **1** shows an example of an interactive multimedia presentation with supplemental content inserted therein, in accordance with some embodiments of the disclosure. Interactive multimedia presentations may include video games, AR content, and VR content. As AR and VR developers are increasingly leveraging existing game engines for their applications, the descriptions of various embodiments of this disclosure will be discussed in the context of video games. However, it will be apparent to anyone skilled in the art that the descriptions can be applicable to other types of interactive multimedia presentations.

[0031] Interactive multimedia content **100** may be a video game. In the example of FIG. **1**, interactive multimedia content **100** is a first-person shooter video game. Cloud-based interactive content server **102** (e.g., a game engine server) encodes and transmits **104** the video game content to a client device using a high-bitrate low-latency streaming protocol. Examples of such protocols include Web Real-Time Communication protocol (WebRTC), Real Time Messaging Protocol (RTMP), Real-Time Transport Protocol (RTP), Real-Time Streaming Protocol (RTSP), low-latency

HTTP Live Streaming protocol (HLS), low latency Common Media Application Format (CMAF) for MPEG-DASH, WebRTC, etc.

[0032] When supplemental content, such as an advertisement or a cutscene, is to be added into the interactive multimedia presentation, cloud-based interactive content server **102** pauses streaming of the video game content to the client device and encodes and transmits **106** the supplemental content **108** to the client device using an adaptive bitrate protocol. Examples of such protocols include MPEG-DASH, Apple HTTP Live Streaming (Apple HLS), Adobe HTTP Dynamic Streaming (HDS), or Microsoft Smooth Streaming (MSS), etc. The supplemental content may be customized for each user. For example, an advertisement inserted into the interactive multimedia presentation may be selected for each user based on each user's preferences, media consumption history, shopping history, etc. In some cases, video game cutscenes can be customized for each user as well. In many video games, in-game character selections, purchases, and upgrades can change the appearance of a user's avatar. Portions of cutscenes that include a user's avatar can thus be customized such that the appearance of the user's avatar accurately reflects these changes.

[0033] For supplemental content that is not customized, pre-rendered media may already be encoded in a format compatible with the second streaming protocol. FIG. 2 is an example of a system for providing an interactive multimedia presentation in which interactive content and supplemental content are provided from different sources, in accordance with some embodiments of the disclosure. Client device **200** receives streaming interactive content from interactive content server **202** and receives supplemental content from supplemental content server **204**. Client device **200** communicates with interactive content server **202** and supplemental content server **204** through communications network **206**. When a point in the interactive multimedia presentation at which supplemental content is to be inserted is reached, client device **200** may transmit **208** an indication of, or request for, supplemental content insertion to interactive content server **202**. Interactive content server **202** may in turn request supplemental content from supplemental content server **204** or may instruct supplemental content server **204** to provide supplemental content to client device **200**. Whether from interactive content server **202** or supplemental content server **204**, client device **200** receives **210** the supplemental content.

[0034] In some embodiments, interactive content server **202** may be operated by a first content provider while supplemental content server **204** may be operated by a second content provider. In some embodiments, interactive content server **202** and supplemental content server **204** may be operated by the same content provider but are physically separate servers and may be in geographically separate locations. In yet other embodiments, interactive content server **202** and supplemental content server **204** may be virtual server applications running on the same physical server or on two different physical servers.

[0035] Customized supplemental content, however, must be generated, rendered, and encoded before it can be streamed to the user. This may require significant system resources. The interactive content server **202** may temporarily reduce the allocation of system resources to other game sessions in order to generate, render, and encode the supplemental content. In some implementations, however, the task of generating, rendering, and encoding the customized supplemental content may be given to another interactive content engine within the content provider's network. FIG. 3 is an example of a system for providing an interactive multimedia presentation in which interactive content and supplemental content are provided from the same source, in accordance with some embodiments of the disclosure. Content provider network **300** includes interactive content server **302** and supplemental content server **304**. Supplemental content server **304** may have similar processing capabilities and system resources as interactive content server **302** to enable supplemental content server **304** to generate, render, and encode customized supplemental content. Interactive content server **302** and supplemental content server **304** communicate with client device **306** via communications network **308**. Client device **306** may, upon detecting that supplemental content is to be inserted into the interactive multimedia

presentation, transmit **310** a request for supplemental content to interactive content server **302** or to supplemental content server **304**.

[0036] Interactive content server **302** or supplemental content server **304** may access a user profile database **312** to retrieve customization parameters relevant to the supplemental content. For example, if the supplemental content is an advertisement, product purchase histories, Internet search histories, or other information useful in selecting an advertisement for a particular user may be retrieved. If the supplemental content is a game cutscene or other cinematic content, information related to the appearance of the user's in-game avatar such as clothing, skin color, size (e.g., height and weight), fictional race (e.g., elf, dwarf, extraterrestrial being, etc.) and the like may be retrieved. Supplemental content server **304** generates, renders, and encodes the supplemental content for the user using the retrieved information. The supplemental content is then streamed **314** to client device **306**.

[0037] FIG. **4** is an example of a content delivery network for providing an interactive multimedia presentation, in accordance with some embodiments of the disclosure. Content delivery network **400** may include a plurality of servers. Some servers are generally utilized for content origination and do not usually handle streaming of content to client devices. Other servers which are usually tasked with streaming content to client devices, known as edge servers, are located in geographically separate locations to better serve users in different areas. A given edge server handles streaming of content to client devices that are located closer to that edge server than to any other edge server.

[0038] Edge server **402** may be a game engine server and is responsible for streaming **404** real-time interactive content to one or more client devices. Edge server **402** may maintain several interactive content sessions simultaneously, each for a different user. Edge server **406** may be a second game engine server and may also be responsible for streaming **408** real-time interactive content to one or more other client devices. Edge server **406**, may, alternatively or additionally, be responsible for streaming supplemental content to one or more users. In some implementations, a server other than an edge server, such as server **410**, may be responsible for streaming **412** supplemental content to client devices. This allows edge servers **402** and **406** to commit their full resources to real-time processing and streaming of interactive content, thus providing users with the best quality of experience, while other servers are tasked with the potentially processor-intensive tasks of generating, rendering, and encoding supplemental content for each user.

[0039] FIG. **5** is a block diagram showing components and data flow therebetween of a system for providing an interactive multimedia presentation in which interactive content and supplemental content are provided from the same source, in accordance with some embodiments of the disclosure. More specifically, the components shown in FIG. **5** relate to implementations in which real-time interactive content and supplemental content are provided from the same content provider and in which a single device or application controls content delivery to client devices from an interactive content server or server application and a supplemental content server or server application. For simplicity, control circuitry **500** will be discussed as though it resides on a single server. However, those skilled in the art will appreciate that the functions described in connection with FIG. **5** as being performed by control circuitry **500** or any of its components may be distributed among one or more physical or virtual servers.

[0040] Control circuitry **500** may be based on any suitable processing circuitry and comprises control circuits and memory circuits, which may be disposed on a single integrated circuit or may be discrete components. As referred to herein, processing circuitry should be understood to mean circuitry based on one or more microprocessors, microcontrollers, digital signal processors, programmable logic devices, field-programmable gate arrays (FPGAs), application-specific integrated circuits (ASICs), etc., and may include a multi-core processor (e.g., dual-core, quad-core, hexa-core, or any suitable number of cores). In some embodiments, processing circuitry may be distributed across multiple separate processors or processing units, for example, multiple of the

same type of processing units (e.g., two Intel Core i7 processors) or multiple different processors (e.g., an Intel Core i5 processor and an Intel Core i7 processor).

[0041] Control circuitry **500** communicates with other devices (e.g. client devices and other servers) via network interface **502**. Network interface **502** may be any suitable communications interface for transmitting data from a server, or network of servers, to other devices across a communications network. In some embodiments, the communications network over which the interactive content and supplemental content are streamed to the client device is a mobile broadband network, such as a 5G network. In such embodiments, network interface **502** may be a mobile broadband transmitter or an interface with one or more distribution nodes of a mobile broadband network.

[0042] Control circuitry **500** includes streaming circuitry **504**. Streaming circuitry **504** includes any suitable processing circuitry for generating a media stream for interactive media content or supplemental content. For example, streaming circuitry **504** may be a graphical processing unit (GPU) such as an NVIDIA Geforce RTX, or other GPU solutions. Streaming circuitry **504** receives **506** interactive media content from interactive content processing circuitry **508**. Interactive content processing circuitry **508** may be, for example, a game engine configured to process game environment and level data along with user inputs to generate, render, and encode interactive content into a real-time media format. Streaming circuitry **504** then generates a real-time stream for the interactive content and transmits **510** the real-time stream to transceiver circuitry **512**.

Transceiver circuitry **512** comprises a network connection over which data can be transmitted to and received from remote devices, such as an ethernet connection, WiFi connection, USB connection, or connection employing any other suitable data transfer protocol or networking protocol. Transceiver circuitry **512** in turn transmits **514** the real-time stream to network interface **502**, where the stream is finally transmitted **516** over a communications network to a client device.

[0043] Control circuitry **500** monitors the streaming of the interactive content for an indication that supplemental content is to be inserted into the interactive multimedia presentation. For example, control circuitry **500** may monitor the interactive media content being received from interactive content processing circuitry **508**, or metadata thereof, for an indication that a supplemental content insertion point is approaching or has arrived. Alternatively or additionally, control circuitry **500** may monitor data received from a client device which may comprise a request for supplemental content. A client device may independently determine that a supplemental content insertion point is approaching and request a supplemental content item. In the case of a video game cutscene, the request may include an identifier of a specific cutscene. The request is received **518** from the client device at network interface **502**, which routes **520** the request to control circuitry **500**, where it is received using transceiver circuitry **512**. Transceiver circuitry **512** in turn transmits **522** the request to streaming circuitry **504**. Based on receiving the request, streaming circuitry **504** determines that supplemental content is to be inserted into the interactive multimedia presentation.

[0044] In response to determining that supplemental content is to be inserted into the interactive multimedia presentation, streaming circuitry **504** instructs **524** interactive content processing circuitry **508** to pause output of the interactive content. In some embodiments, interactive content processing circuitry **508** is responsible for processing of interactive content for multiple interactive multimedia sessions, each being associated with a different user or different client device. In such cases, streaming circuitry **504** may provide a session identifier in the instruction to pause output of the interactive content, thereby causing interactive content processing circuitry **508** to pause processing of only the session into which the supplemental content is to be inserted.

[0045] Streaming circuitry **504** then instructs **526** supplemental content processing circuitry **528** to begin processing and output of the supplemental content to be inserted into the interactive multimedia presentation. Supplemental content processing circuitry **528** may be a less powerful graphical processing unit than interactive content processing circuitry **508** and may be configured for insertion of non-interactive supplemental content such as advertisements. However,

supplemental content processing circuitry **528** may be similar to interactive content processing circuitry **508** for generating, rendering, and encoding customized supplemental content. In other embodiments, supplemental content processing circuitry **528** is merely responsible for identifying supplemental content items and providing an associated manifest file or URL from which the client device can stream the supplemental content. Supplemental content processing circuitry **528** may request, from a supplemental content database, a manifest file or URL associated with a specific supplemental content item to be inserted into the interactive multimedia presentation.

[0046] In embodiments in which supplemental content processing circuitry **528** is responsible for encoding and output of supplemental content directly, supplemental content processing circuitry **528** may encode the supplemental content in a format compatible with a streaming protocol configured for adaptive bitrate streaming. Supplemental content processing circuitry **528** may also receive, request, or otherwise access information regarding available bandwidth at the client device and encode the supplemental content in a format most suited to the available bandwidth. For example, while the communications network may be a 5G mobile broadband network, the client device may be in a location that does not have good 5G coverage and can only support 4G transmissions. Supplemental content processing circuitry **528** may therefore encode the supplemental content in a format such that the highest bitrate available to the client device does not exceed the available bandwidth. This may conserve system resources when encoding the supplemental content.

[0047] Whether it is a manifest file, a URL, or encoded media content, supplemental content processing circuitry **528** transmits **530** its output to streaming circuitry **504**. Streaming circuitry **504** then generates an adaptive bitrate (ABR) stream for the supplemental content and transmits **532** the ABR stream to transceiver circuitry **512**, which then transmits **534** the output to network interface **502** where it is transmitted **536** to the client device over the communications network.

[0048] During playback of the supplemental content, a user of the client device may select an option to skip a portion of the supplemental content. An indication of such a selection, or of a natural end of playback of the supplemental content, may be received **538** from the client device via network interface **502**. Network interface **502** may route **540** the indication to control circuitry **500**, where it is received using transceiver circuitry **512**. Transceiver circuitry **512** transmits **542** the indication to streaming circuitry **504**. In response to receiving the indication, streaming circuitry **504** instructs **544** supplemental content processing circuitry **528** to end processing and output of the supplemental content. Supplemental content processing circuitry **528** transmits **546a** an acknowledgment and/or an indication that output of the supplemental content has stopped to streaming circuitry **504**. Streaming circuitry **504** then instructs **548** interactive content processing circuitry **508** to resume generating, rendering, and encoding the interactive content for output to the client device. Alternatively or additionally, supplemental content processing circuitry **528** transmits **546b** such indications directly to interactive content processing circuitry **508**, in response to which interactive content processing circuitry **508** resumes generating, rendering, and encoding the interactive content for output to the client device. Interactive content processing circuitry **508** transmits **550** the interactive content to streaming circuitry **504**. Streaming circuitry **504** then generates a real-time stream for the interactive content and transmits **552** the real-time stream to transceiver circuitry **512**. Transceiver circuitry **512** in turn transmits **554** the real-time stream to network interface **502**, where it is then transmitted **556** to the client device over the communications network.

[0049] FIG. **6** is a block diagram showing components and data flow therebetween for a system for providing an interactive multimedia presentation in which interactive content and supplemental content are provided from different sources, in accordance with some embodiments of the disclosure. For example, the interactive content may be provided by an interactive content server while the supplemental content is provided by an edge content manager that resides on a different server. Interactive content server **600** is responsible for initializing and maintaining interactive

multimedia sessions (e.g. game sessions), and generating, rendering, and encoding interactive content. Interactive content server **600** includes control circuitry **602**. Like control circuitry **500**, control circuitry **602** may be based on any suitable processing circuitry and comprises control circuits and memory circuits, which may be disposed on a single integrated circuit or may be discrete components. In some embodiments, processing circuitry may be distributed across multiple separate processors or processing units, for example, multiple of the same type of processing units (e.g., two Intel Core i7 processors) or multiple different processors (e.g., an Intel Core i5 processor and an Intel Core i7 processor).

[0050] Control circuitry **602** includes streaming circuitry **604**. Like streaming circuitry **504**, streaming circuitry **604** includes any suitable processing circuitry for generating a media stream for interactive media content or supplemental content. For example, streaming circuitry **604** may be a graphical processing unit (GPU) such as an NVIDIA Geforce RTX, or other GPU solutions. Streaming circuitry **604** receives **606** interactive media content from interactive content processing circuitry **608**. Similar to interactive content processing circuitry **508**, interactive content processing circuitry **608** may be, for example, a game engine configured to process game environment and level data along with user inputs to generate, render, and encode interactive content into a real-time media format. Streaming circuitry **604** then generates a real-time stream for the interactive content and transmits **610** the real-time stream to transceiver circuitry **612**. Like transceiver circuitry **512**, transceiver circuitry **612** comprises a network connection over which data can be transmitted to and received from remote devices, such as an ethernet connection, WiFi connection, USB connection, or connection employing any other suitable data transfer protocol or networking protocol. Transceiver circuitry **612** in turn transmits **614** the real-time stream to network interface **616**, where the stream is transmitted **618** over communications network **620** and delivered **622** to a client device. Network interface **616**, like network interface **502**, may be any suitable communications interface for transmitting data from a server, or network of servers, to other devices across a communications network. In some embodiments, the communications network over which the interactive content and supplemental content are streamed to the client device is a mobile broadband network, such as a 5G network. In such embodiments, network interface **616** may be a mobile broadband transmitter or an interface with one or more distribution nodes of a mobile broadband network.

[0051] Control circuitry **602** monitors the streaming of the interactive content for an indication that supplemental content is to be inserted into the interactive multimedia presentation. For example, control circuitry **602** may monitor the interactive media content being received from interactive content processing circuitry **608**, or metadata thereof, for an indication that a supplemental content insertion point is approaching or has arrived. Alternatively or additionally, control circuitry **602** may monitor data received over communications network **620** from a client device which may comprise a request for supplemental content. A client device may independently determine that a supplemental content insertion point is approaching and request a supplemental content item. In the case of a video game cutscene, the request may include an identifier of a specific cutscene. The request is transmitted **624** from the client device, over communications network **620**, and received **626** from the client device at network interface **616**, which routes **628** the request to interactive content server **600**. Control circuitry **602** then receives the request using transceiver circuitry **612**. Transceiver circuitry **612** in turn transmits **630** the request to streaming circuitry **604**. Based on receiving the request, streaming circuitry **604** determines that supplemental content is to be inserted into the interactive multimedia presentation.

[0052] In response to determining that supplemental content is to be inserted into the interactive multimedia presentation, streaming circuitry **604** instructs **632** interactive content processing circuitry **608** to pause output of the interactive content. In some embodiments, interactive content processing circuitry **608** is responsible for processing of interactive content for multiple interactive multimedia sessions, each being associated with a different user or different client device. In such

cases, streaming circuitry **604** may provide a session identifier in the instruction to pause output of the interactive content, thereby causing interactive content processing circuitry **608** to pause processing of only the session into which the supplemental content is to be inserted.

[0053] Network interface **616** also routes **634** the request to edge content manager **636**. Edge content manager **636** is another server, separate from interactive content server **600**, that is responsible for providing supplemental content to client devices. Edge content manager **636** includes control circuitry **638**. Like control circuitry **500** and control circuitry **602**, control circuitry **638** may be based on any suitable processing circuitry and comprises control circuits and memory circuits, which may be disposed on a single integrated circuit or may be discrete components. In some embodiments, processing circuitry may be distributed across multiple separate processors or processing units, for example, multiple of the same type of processing units (e.g., two Intel Core i7 processors) or multiple different processors (e.g., an Intel Core i5 processor and an Intel Core i7 processor).

[0054] Control circuitry **638** receives the request using transceiver circuitry **640**. Similar to transceiver circuitry **512** and transceiver circuitry **612**, transceiver circuitry **640** comprises a network connection over which data can be transmitted to and received from remote devices, such as an ethernet connection, WiFi connection, USB connection, or connection employing any other suitable data transfer protocol or networking protocol. Transceiver circuitry **640** transmits **642** the request to streaming circuitry **644**. Like streaming circuitry **504** and streaming circuitry **604**, streaming circuitry **644** includes any suitable processing circuitry for generating a media stream for interactive media content or supplemental content. For example, streaming circuitry **644** may be a graphical processing unit (GPU) such as an NVIDIA Geforce RTX, or other GPU solutions. Streaming circuitry **644** transmits **646** the request to supplemental content processing circuitry **648**. Supplemental content processing circuitry **648**, like supplemental content processing circuitry **528**, may be a less powerful graphical processing unit than interactive content processing circuitry **608** and may be configured for insertion of non-interactive supplemental content such as advertisements. However, supplemental content processing circuitry **648** may be similar to interactive content processing circuitry **608** for generating, rendering, and encoding customized supplemental content. In other embodiments, supplemental content processing circuitry **648** is merely responsible for identifying supplemental content items and providing an associated manifest file or URL from which the client device can stream the supplemental content. Supplemental content processing circuitry **648** may request, from a supplemental content database, a manifest file or URL associated with a specific supplemental content item to be inserted into the interactive multimedia presentation.

[0055] In embodiments in which supplemental content processing circuitry **648** is responsible for encoding and output of supplemental content directly, supplemental content processing circuitry **648** may encode the supplemental content in a format compatible with a streaming protocol configured for adaptive bitrate streaming. Supplemental content processing circuitry **648** may also receive, request, or otherwise access information regarding available bandwidth at the client device and encode the supplemental content in a format most suited to the available bandwidth. For example, while the communications network may be a 5G mobile broadband network, the client device may be in a location that does not have good 5G coverage and can only support 4G transmissions. Supplemental content processing circuitry **648** may therefore encode the supplemental content in a format such that the highest bitrate available to the client device does not exceed the available bandwidth. This may conserve system resources when encoding the supplemental content.

[0056] Whether it is a manifest file, a URL, or encoded media content, supplemental content processing circuitry **648** transmits **650** its output to streaming circuitry **644**. Streaming circuitry **644** then generates an ABR stream for the supplemental content and transmits **652** the ABR stream to transceiver circuitry **640**, which then transmits **654** the output to network interface **616** where it

is transmitted **658** over communications network **620** and routed **658** to the client device.

[0057] During playback of the supplemental content, a user of the client device may select an option to skip a portion of the supplemental content. An indication of such a selection, or of a natural end of playback of the supplemental content, may be transmitted **660** by the client device over the communications network **620**. The indication of the selection is received **662** at network interface **616**. Network interface **616** may route **664** the indication to control circuitry **638** of edge content manager **636**, where it is received using transceiver circuitry **640**. Transceiver circuitry **640** transmits **666** the indication to streaming circuitry **644**. In response to receiving the indication, streaming circuitry **644** instructs **668** supplemental content processing circuitry **648** to end processing and output of the supplemental content. Network interface **616** also routes **670** the indication to control circuitry **602** of interactive content server **600**. Control circuitry **602** receives the indication using transceiver circuitry **612**. Transceiver circuitry **612** transmits **672** the indication to streaming circuitry **604**. Streaming circuitry **604** then instructs **674** interactive content processing circuitry **608** to resume generating, rendering, and encoding the interactive content for output to the client device. Interactive content processing circuitry **608** transmits **676** the interactive content to streaming circuitry **604**. Streaming circuitry **604** then generates a real-time stream for the interactive content and transmits **678** the real-time stream to transceiver circuitry **612**. Transceiver circuitry **612** in turn transmits **680** the real-time stream to network interface **616**, where it is then transmitted **682** over communications network **620** and routed **684** to the client device.

[0058] FIG. 7 is a block diagram showing components and data flow therebetween of an interactive content server when supplemental content is to be inserted into an interactive multimedia presentation, in accordance with some embodiments of the disclosure. Specifically, FIG. 7 depicts an interactive content server **700** that provides a manifest file for supplemental content to a client device. Interactive content server **700** includes control circuitry **702**, which may be similar to control circuitry **500** or control circuitry **638** and may be based on any suitable processing circuitry as described above in connection therewith. Interactive content server **700** receives **704**, using transceiver circuitry **706**, an indication that supplemental content is to be inserted into the interactive multimedia presentation. Transceiver circuitry **706** may be similar to transceiver circuitry **512** or transceiver circuitry **612**, described above in connection with FIGS. 5 and 6, respectively. Transceiver circuitry **706** transmits **708** the indication to streaming circuitry **710**, which may be similar to streaming circuitry **504** or streaming circuitry **604**, described above in connection with FIGS. 5 and 6, respectively. Streaming circuitry **710** instructs **712** interactive content processing circuitry **714**, similar to interactive content processing circuitry **508** or interactive content processing circuitry **608**, described above in connection with FIGS. 5 and 6, respectively, to pause generation, rendering, and encoding of the interactive content to allow for insertion of the supplemental content.

[0059] Control circuitry **702** may identify a specific supplemental content item, either based on information received from the client device, user profile data associated with the user of the client device, or from the interactive content. Transceiver circuitry **706** transmits **716** a request for the identified supplemental content item to supplemental content manifest database **718**. While shown in FIG. 7 as a separate component, supplemental content manifest database **718** may, in some embodiments, reside in a memory of interactive content server **700**. In other embodiments, supplemental content manifest database **718** may reside on a separate server. In response to the request, control circuitry **702** receives **720** from supplemental content manifest database **718** a manifest file for the supplemental content or, in some embodiments, a URL from which the client device can retrieve the manifest file directly. Transceiver circuitry **706** then transmits **722** the manifest file or URL to the client device.

[0060] FIG. 8 is a block diagram showing components and data flow therebetween of an edge content management server for providing supplemental content, in accordance with some embodiments of the disclosure. Specifically, FIG. 8 depicts an edge content manager **800** that

provides a manifest file for supplemental content to a client device. Edge content manager **800** includes control circuitry **802**, which may be similar to control circuitry **500** or control circuitry **602** and may be based on any suitable processing circuitry as described above in connection therewith. Edge content manager **800** receives **804**, using transceiver circuitry **806**, an indication that supplemental content is to be inserted into the interactive multimedia presentation. In some embodiments, edge content manager **800** receives an instruction from an interactive content server to prepare supplemental content for insertion into an interactive multimedia presentation. The instruction may include a session identifier associated with the interactive multimedia presentation, a user identifier of a user associated with the client device, or a supplemental content identifier. Transceiver circuitry **806** may be similar to transceiver circuitry **512** or transceiver circuitry **640**, described above in connection with FIGS. 5 and 6, respectively. Transceiver circuitry **806** transmits **808** the indication to streaming circuitry **810**, which may be similar to streaming circuitry **504** or streaming circuitry **644**, described above in connection with FIGS. 5 and 6, respectively. Streaming circuitry **810** instructs **812** supplemental content processing circuitry **814**, similar to supplemental content processing circuitry **528** or supplemental content processing circuitry **648**, described above in connection with FIGS. 5 and 6, respectively, to prepare supplemental content for insertion into the interactive multimedia presentation.

[0061] Supplemental content processing circuitry **814** may identify a specific supplemental content item, either based on information received from the client device, user profile data associated with the user of the client device, or from the interactive content. Supplemental content processing circuitry **814** transmits **816** an identifier of the specific supplemental content item to transceiver circuitry **806**, which in turn transmits **818** a request for the identified supplemental content item to supplemental content manifest database **820**. As in FIG. 7, above, supplemental content manifest database **820** may, in some embodiments, reside in a memory of edge content manager **800**. In other embodiments, supplemental content manifest database **820** may reside on a separate server. In response to the request, control circuitry **802** receives **822** from supplemental content manifest database **820** a manifest file for the supplemental content or, in some embodiments, a URL from which the client device can retrieve the manifest file directly. Transceiver circuitry **806** then transmits **824** the manifest file or URL to the client device.

[0062] FIG. 9 is a flowchart representing a process **900** for providing an interactive multimedia presentation, in accordance with some embodiments of the disclosure. Process **900** may be implemented on control circuitry **500**, **602**, **638**, **702**, **802**, or any combination thereof. In addition, one or more actions of process **900** may be incorporated into or combined with one or more actions of any other process or embodiment described herein.

[0063] At **902**, control circuitry (e.g., control circuitry **500**) streams, to a client device using a first streaming protocol, interactive content from an interactive content engine for an interactive multimedia presentation. The interactive content may be virtual reality content, augmented reality content, video game content, or any other real-time content. The interactive content engine may be a game engine, which may be used to provide video game content or may be leveraged to provide other real-time content such as virtual reality or augmented reality content. In some implementations, where the methods described herein are performed on a mobile broadband network, the interactive content engine may be a 5G Extended Reality (5G-XR) application as defined in the 3GPP technical report TR 26.928, entitled “Extended Reality (XR) in 5G.” The first streaming protocol is configured for high-bitrate low-latency streaming of real-time data. While such streams are not tolerant of dropped or corrupted packets, they allow for significantly reduced lag time which is of high importance in interactive content such as video games and virtual reality simulations.

[0064] At **904**, the control circuitry determines whether supplemental content is to be inserted in the interactive multimedia presentation. For example, a client device may send a request for supplemental content upon reaching a point in the interactive multimedia presentation at which

supplemental content is to be inserted. Where the interactive multimedia presentation is a video game, the client device may determine, based on the user's progress in the video game, that a cutscene is to be inserted. In some embodiments, the real-time content stream comprising the interactive content, e.g., the video game content, may contain metadata indicating that a supplemental content insertion point has been reached, or is approaching. If no indication of supplemental content insertion has been detected (“No” at **904**), then streaming of the interactive content continues.

[0065] If supplemental content is to be inserted into the interactive multimedia presentation (“Yes” at **904**), then, at **906**, the control circuitry streams the supplemental content to the client device using a second streaming protocol configured for adaptive bitrate streaming. For example, the control circuitry may determine an available amount of bandwidth at the client device and retrieve a precached version of the supplemental content having the highest quality, or bitrate, that can be supported by the available bandwidth. The control circuitry then generates an adaptive bitrate stream from the precached version of the content and transmits the stream to the client device. In some embodiments, the control circuitry generates or retrieves from a database a manifest file for the supplemental content and transmits the manifest file to the client device. The client device then requests segments of the supplemental content according to the manifest file and the control circuitry transmits each requested segment to the client device.

[0066] At **908**, the control circuitry determines whether an indication that the streaming of the supplemental content to the client device has ended has been received. For example, during playback of the supplemental content, the client device may receive a user input to skip a portion of the supplemental content. An indication of such an input is transmitted by the client device to the content server. Alternatively or additionally, the control circuitry may determine that a threshold amount of time has passed since the last segment request was received from the client device, or may determine that the last segment of the supplemental content has been transmitted to the client device. If no such indication has been received (“No” at **908**), then the control circuitry continues to stream the supplemental content in an ABR stream.

[0067] If an indication that streaming of the supplemental content to the client device has ended has been received (“Yes” at **908**), then, at **910**, the control circuitry resumes streaming of the interactive content to the client device using the first streaming protocol. In some cases, the supplemental content may be inserted within a scene or portion of the interactive content. In such cases, the control circuitry simply continues streaming of the content from the point at which the streaming of the interactive content was interrupted for insertion of the supplemental content. In other cases, the supplemental content may be inserted between scenes or portions of the interactive content. In such cases, the control circuitry may first identify a next portion of the interactive content to provide to the client device. For example, in a video game, a cutscene may be inserted between two levels of the video game. While the cutscene is being played back, the circuitry may use the time to load data for the next level and begin rendering portions of the next level for output. Once the cutscene is over, or when a user selection to skip the cutscene is received, the control circuitry begins streaming the new level to the client device.

[0068] The actions or descriptions of FIG. **9** may be used with any other embodiment of this disclosure. In addition, the actions and descriptions described in related to FIG. **9** may be done in suitable alternative orders or in parallel to further the purposes of this disclosure.

[0069] FIG. **10** is a flowchart representing a process **1000** for determining whether supplemental content is to be inserted into an interactive multimedia presentation, in accordance with some embodiments of the disclosure. Process **1000** may be implemented on control circuitry **500**, **602**, **638**, **702**, **802**, or any combination thereof. In addition, one or more actions of process **1000** may be incorporated into or combined with one or more actions of any other process or embodiment described herein.

[0070] At **1002**, the control circuitry monitors streaming of the interactive content. The control

circuitry may scan the interactive content as it is transmitted to the client device for indications that supplemental content is to be inserted. For example, one or more frames or packets of interactive content may include metadata indicating a supplemental content insertion point. Alternatively, the interactive content processing circuitry may automatically pause output of the interactive content when a supplemental content insertion point has been reached.

[0071] At **1004**, the control circuitry determines whether an indication of supplemental content to be inserted in the interactive multimedia presentation has been detected. If not (“No” at **1004**), then the control circuitry continues to monitor streaming of the interactive content. At **1006**, the control circuitry determines whether streaming of the interactive content has stopped. If not (“No” at **1006**), then control circuitry continues to monitor streaming of the interactive content.

[0072] If an indication of supplemental content to be inserted into the interactive multimedia presentation has been detected (“Yes” at **1004**) or if streaming of the interactive content has stopped (“Yes” at **1006**), then, at **1008**, the control circuitry determines that supplemental content is to be inserted in the interactive multimedia presentation.

[0073] The actions or descriptions of FIG. **10** may be used with any other embodiment of this disclosure. In addition, the actions and descriptions described in related to FIG. **10** may be done in suitable alternative orders or in parallel to further the purposes of this disclosure.

[0074] FIG. **11** is a flowchart representing a process **1100** for skipping a portion of supplemental content, in accordance with some embodiments of the disclosure. Process **1100** may be implemented on control circuitry **500**, **602**, **638**, **702**, **802**, or any combination thereof. In addition, one or more actions of process **1100** may be incorporated into or combined with one or more actions of any other process or embodiment described herein.

[0075] At **1102**, the control circuitry pauses streaming of the interactive content to the client device and, at **1104**, begins streaming the supplemental content to the client device using the second streaming protocol. At **1106**, the control circuitry monitors user inputs received from the client device. For example, the client device may transmit a signal every time a user input is received. Alternatively, the client device may transmit a periodic signal (e.g., every 10 milliseconds) representing the state of one or more control inputs such as game controller buttons, keyboard keys, mouse movements, joystick movements, accelerometer data, inertial movement data, biometric data, etc. At **1108**, the control circuitry determines whether a user input to skip a portion of the supplemental content has been received. For example, a signal representing user selection of a graphical control element corresponding to a skip command may be received. As another example, a specific control input or combination of inputs may be received that corresponds to a skip command. If no such input as been received (“No” at **1108**), then, at **1110**, the control circuitry determines whether the supplemental content has ended. For example, the control circuitry may monitor streaming of the supplemental content and determine when the supplemental content has stopped being output for streaming. If the supplemental content has not ended (“No” at **1110**), then processing returns to **1106** and the control circuitry continues to monitor user inputs.

[0076] If a user input to skip a portion of the supplemental content has been received (“Yes” at **1108**) or if the supplemental content has ended (“Yes” at **1110**), then, at **1112**, the control circuitry stops the streaming of the supplemental content to the client device and, at **1114**, resumes streaming of the interactive content to the client device.

[0077] The actions or descriptions of FIG. **11** may be used with any other embodiment of this disclosure. In addition, the actions and descriptions described in related to FIG. **11** may be done in suitable alternative orders or in parallel to further the purposes of this disclosure.

[0078] FIG. **12** is a flowchart representing a process **1200** for selecting supplemental content where the interactive multimedia presentation comprises a video game, in accordance with some embodiments of the disclosure. Process **1200** may be implemented on control circuitry **500**, **602**, **638**, **702**, **802**, or any combination thereof. In addition, one or more actions of process **1200** may be incorporated into or combined with one or more actions of any other process or embodiment

described herein.

[0079] At **1202**, the control circuitry established a game session. For example, the control circuitry may receive new game session request from a client device. The control circuitry authenticates the requesting device and initializes a new game session for the authenticated device. Once initialized, the control circuitry may share an identifier of the new game session with the authenticated client device for use in requesting content or other data for the duration of the game session.

[0080] At **1204**, the control circuitry identifies, based on the game session, a user identifier. For example, the client device may include a user identifier in the game session request. In some embodiments, the game session is initialized without any user data and a user of the client device must subsequently log in in order to access the video game content. Once authenticated, an identifier of the user (e.g., a user name, a user ID, or other data unique to the user) may be retrieved by the control circuitry from the game session.

[0081] At **1206**, the control circuitry determines whether a user profile associated with the user identifier exists. The control circuitry may access a user profile database and search for data matching the user identifier. For example, the control circuitry may generate an SQL “SELECT” command based on the user identifier. If no result is returned from the database, the control circuitry determines that no user profile exists for the identified user (“No” at **1206**) and, at **1208**, selects a supplemental content item for streaming to the client device. For example, if the supplemental content is an advertisement, the control circuitry may select an advertisement at random or according to an ad campaign. If the supplemental content is a cutscene, the control circuitry identifies and selects the particular cutscene for the point within the video game at which the cutscene is to be inserted. This may be done on the basis on a cutscene identifier requested by the client device or identified in metadata of the video game content.

[0082] If the control circuitry determines that a profile associated with the user identifier exists (“Yes” at **1206**), then, at **1210**, the control circuitry retrieves the profile of the user. At **1212**, the control circuitry determines what type of supplemental content is to be inserted. For example, a request for supplemental content received from the client device may indicate whether a cutscene or an advertisement is to be inserted. Alternatively, metadata of the video game content may include a flag or other indicator that identifies the type of supplemental content to be inserted.

[0083] If the supplemental content to be inserted is an advertisement, then, at **1214**, the control circuitry selects, based on the profile of the user, a supplemental content item for streaming to the client device. For example, the supplemental content may be an advertisement. The control circuitry may retrieve data from the user profile, such as online shopping history, social media data, browsing history, or content consumption history. Based on the retrieved data, the control circuitry selects an advertisement for streaming to the client device in which the user of the client device is likely to be interested.

[0084] If the supplemental content to be inserted is a cutscene, then, at **1216**, the control circuitry identifies visual parameters from the profile of the user. For example, the cutscene may include a visualization of the user's in-game avatar. In some video games, the user can select their avatar's initial appearance (e.g., height, weight, skin color, eye color, hair color, fictional race, etc.). The user can also acquire, through completing portions of the video game or through in-game purchases or other microtransactions, various equipment such as clothing, armor, weapons, etc. These parameters may be stored in the user's profile. At **1218**, the control circuitry renders the cutscene using the identified visual parameters. The control circuitry uses the appearance of the user's avatar, as stored in the user's profile, to render a third-person visualization of the user's avatar. This may be a 3D rendering or a 2D rendering and may be of the entirety of the avatar's body or some portion thereof. The rendering can then be inserted into the cutscene to create a customized cutscene.

[0085] The actions or descriptions of FIG. **12** may be used with any other embodiment of this disclosure. In addition, the actions and descriptions described in related to FIG. **12** may be done in suitable alternative orders or in parallel to further the purposes of this disclosure.

[0086] FIG. 13 is a flowchart representing a process 1300 for delivering a video game cutscene prior to start of the video game, in accordance with some embodiments of the disclosure. Process 1300 may be implemented on control circuitry 500, 602, 638, 702, 802, or any combination thereof. In addition, one or more actions of process 1300 may be incorporated into or combined with one or more actions of any other process or embodiment described herein.

[0087] At 1302, control circuitry of the edge gaming manager, e.g., control circuitry 638 of edge content manager 636, receives a client session startup request with a session ID. The session ID may be received from the client device, or may be assigned by the game engine (e.g., interactive content server 600). In response, at 1304, the edge gaming manager starts the game engine. For example, the edge gaming manager transmits an initialization instruction to the game engine. At 1306, immediately upon starting, the game engine sends the edge gaming manager a session start response with a manifest URL. At 1308, the edge gaming manager sends the client device a session startup response with the session ID. This allows the client device to join and maintain connection to the game session.

[0088] Also upon starting, the game engine, at 1310, begins startup and initial loading of game data. At 1312, the game engine determines whether loading of game data is complete. If not, (“No” at 1312), then the game engine continues to wait until loading is complete. When loading is complete (“Yes” at 1312), then, at 1314, the game engine sends a “level loaded” notification to the edge gaming manager. The game engine may transmit a ready message once the necessary game data has completed loading. This may be game data for an entire level, a first portion of a level, a menu screen, or any other portion of the video game.

[0089] At 1316, the edge gaming manager determines whether a cutscene response was received and whether the “level loaded” notification was received, indicating that the game engine is initialized and ready for gameplay start. The client device may transmit an acknowledgement of receipt of the cutscene manifest URL, or an indication that ABR playback of the cutscene has successfully begun. If either of the cutscene response or the “level loaded” notification has not been received (“No” at 1316), then the edge gaming manager continues to wait until they are both received.

[0090] Once both messages have been received (“Yes” at 1316), at 1318, the edge gaming manager sends the “level loaded” notification to the client device, where it is directed to a session handler of the client device. At 1320, the edge gaming manager determines whether an “ABR cutscene play complete” notification has been received or if an “ABR cutscene cancel request” notification was received. These notifications may be transmitted by the client device in response to playback of the cutscene having completed naturally or in response to user selection of a skip command, respectively. If no notifications have been received (“No” at 1320), then the edge gaming manager waits while playback of the cutscene continues. If either notification has been received (“Yes” at 1320), then, at 1322, edge gaming manager instructs the game engine to start (in the case of a new game) or resume (in the case of a cutscene between game levels) the game.

[0091] The actions or descriptions of FIG. 13 may be used with any other embodiment of this disclosure. In addition, the actions and descriptions described in related to FIG. 13 may be done in suitable alternative orders or in parallel to further the purposes of this disclosure.

[0092] FIG. 14 is a flowchart representing a process 1400 for delivering a video game cutscene between levels of the video game, in accordance with some embodiments of the disclosure. Process 1400 may be implemented on control circuitry 500, 602, 638, 702, 802, or any combination thereof. In addition, one or more actions of process 1400 may be incorporated into or combined with one or more actions of any other process or embodiment described herein.

[0093] At 1402, the edge gaming manager detects that a game level has ended. For example, the game engine may pause output of game content or may include metadata in the game content indicating that the end of a level has been reached. At 1404, the game engine sends an ABR cutscene play request with a cutscene ID and manifest URL for the cutscene to the edge gaming

manager. At **1406**, the edge gaming manager sends the ABR cutscene play request, with the cutscene ID and manifest URL, to an ABR session handler at the client device.

[0094] Upon receipt of the “ABR cutscene play request,” the ABR client session handler transmits a response, such as an acknowledgement (ACK) response, to the edge gaming manager. At **1408**, the edge gaming manager determines whether such a response has been received from the ABR client session handler. If not (“No” at **1408**), edge gaming manager continues to wait until the response is received. Once the response has been received (“Yes” at **1408**), at **1410**, the edge gaming manager sends a “pause delivery” request to a delivery manager. The delivery manager is a component of the game engine, such as streaming circuitry **604**, that is responsible for streaming game content to the client device.

[0095] In response to receiving the “pause delivery” request, at **1412**, the delivery manager sends a stop or pause encoding request to a low-latency encoder. The low-latency encoder may be part of streaming circuitry **604** or may be a separate component. At **1414**, the low-latency encoder pauses encoding of video, audio, and haptics data for the video game. In response to this pause, at **1416**, a bandwidth manager releases bandwidth reserved for game content in the current game session. The bandwidth manager may be a component of the game engine or the edge gaming manager that is responsible for allocating to and/or reserving portions of the available bandwidth for each active game session. In order to provide the best quality content, bandwidth is only allocated to a game session that is actively streaming a high-bitrate low-latency stream. This allows the maximum amount of bandwidth to be shared among all active game sessions when needed. Thus, after releasing the bandwidth reserved for/allocated to the game content for the current game session, at **1418**, the bandwidth manager recalculates bandwidth for all current in-play non-cutscene state game sessions. The amount of bandwidth reserved for/allocated to each current in-play non-cutscene state game session is then adjusted accordingly.

[0096] After sending the “pause delivery” request, at **1420**, the edge gaming manager sends the “ABR cutscene play” response received from the client device to the game engine. In response, at **1422**, the game engine starts loading data for the next level of gameplay. At **1424**, the game engine determines whether the game level is loaded. If not (“No” at **1424**), then the game engine waits until the game level is loaded. Once loaded (“Yes” at **1424**), processing proceeds as in step **1314** of FIG. **13**, above, wherein the game engine sends a “level loaded” notification to the edge gaming manager.

[0097] The actions or descriptions of FIG. **14** may be used with any other embodiment of this disclosure. In addition, the actions and descriptions described in related to FIG. **14** may be done in suitable alternative orders or in parallel to further the purposes of this disclosure.

[0098] FIG. **15** is a flowchart representing a process for managing bandwidth usage during delivery of a video game cutscene, in accordance with some embodiments of the disclosure. Process **1500** may be implemented on control circuitry **500**, **602**, **638**, **702**, **802**, or any combination thereof. In addition, one or more actions of process **1500** may be incorporated into or combined with one or more actions of any other process or embodiment described herein.

[0099] At **1502**, the edge gaming manager receives an “ABR cutscene play complete” notification or an “ABR cutscene cancel request” notification. As discussed above, the client device or ABR client session handler may transmit a notification to the edge gaming manager when playback of the cutscene is complete or when the user has selected a command to skip the cutscene. At **1504**, the edge gaming manager sends the delivery manager a “resume” request with the session ID for the current game session. In response, at **1506**, the deliver manager sends a bandwidth reservation request for the game associated with the session ID to the bandwidth manager. At **1508**, the bandwidth manager calculates new bitrates for each non-cutscene state game session, including the session corresponding to the session ID. At **1510**, the bandwidth manager sends a bandwidth reservation response to the delivery manager, which then, at **1512**, sends a “start” or “resume encoding” request to the low-latency encoder.

[0100] In response to the “start” or “resume encoding” request, at **1514**, the low-latency encoder begins waiting for an input stream. If an input stream has not been received (“No” at **1514**), then the low-latency encoder continues to wait for the game engine to provide an input stream. If an input stream has been received (“Yes” at **1514**), then, at **1516**, the low-latency encoder transmits encoded video, audio, and haptics to the client device.

[0101] At **1518**, once the delivery manager sends the “start” or “resume encoding” request to the low-latency encoder, the edge gaming manager sends to the game engine an “ABR cutscene play complete” notification. If it is the start or end of a game (“Yes” at **1520**), then, at **1522**, the game engine renders menu options for the game. If it is not the start or end of a game (“No” at **1520**), then, at **1524**, the game engine resumes gameplay of the loaded level.

[0102] The actions or descriptions of FIG. **15** may be used with any other embodiment of this disclosure. In addition, the actions and descriptions described in related to FIG. **15** may be done in suitable alternative orders or in parallel to further the purposes of this disclosure.

[0103] The processes described above are intended to be illustrative and not limiting. One skilled in the art would appreciate that the steps of the processes discussed herein may be omitted, modified, combined, and/or rearranged, and any additional steps may be performed without departing from the scope of the invention. More generally, the above disclosure is meant to be exemplary and not limiting. Only the claims that follow are meant to set bounds as to what the present invention includes. Furthermore, it should be noted that the features and limitations described in any one embodiment may be applied to any other embodiment herein, and flowcharts or examples relating to one embodiment may be combined with any other embodiment in a suitable manner, done in different orders, or done in parallel. In addition, the systems and methods described herein may be performed in real time. It should also be noted that the systems and/or methods described above may be applied to, or used in accordance with, other systems and/or methods.

Claims

1. A method for providing interactive cloud-based supplemental content for an interactive multimedia presentation, the method comprising: reserving bandwidth for use by a first interactive content engine for streaming an interactive multimedia presentation to a plurality of client devices; determining that supplemental content is to be inserted into the interactive multimedia presentation streamed to at least one client device of the plurality of client devices by the first interactive content engine; in response to determining that the supplemental content is to be inserted into the interactive multimedia presentation, releasing at least a portion of the bandwidth reserved for use by the first interactive content engine for streaming the interactive multimedia presentation; streaming, by a second interactive content engine using the released bandwidth, the supplemental content to the at least one client device of the plurality of client devices; upon receiving an indication that streaming of the supplemental content to the at least one client device of the plurality of client devices has ended, reserving the released bandwidth for use by the first interactive content engine; and resuming streaming, by the first interactive content engine using the reserved bandwidth, the interactive content.
2. The method of claim 1, further comprising: monitoring the streaming of the interactive multimedia presentation; wherein the determining that the supplemental content is to be inserted into the interactive multimedia presentation comprises detecting an indication of supplemental content to be inserted into the interactive multimedia presentation.
3. The method of claim 1, further comprising: monitoring the streaming of the interactive multimedia presentation; wherein the determining that the supplemental content is to be inserted into the interactive multimedia presentation comprises detecting, based at least in part on the monitoring, that the streaming of the interactive multimedia presentation has stopped.
4. The method of claim 1, wherein the interactive multimedia presentation comprises a video game.

5. The method of claim 4, wherein the supplemental content comprises a cutscene.
6. The method of claim 1, wherein the supplemental content comprises an advertisement.
7. The method of claim 1, wherein the supplemental content is customized for a user profile of the at least one client device, the method further comprising: accessing a user profile associated with the client device; and modifying a portion of the supplemental content based on the user profile.
8. The method of claim 7, wherein the interactive content is a video game, the supplemental content is a cutscene of the video game, and the portion comprises an avatar of a user of the client device, wherein the modifying the portion of the supplemental content further comprises: determining, based on the user profile associated with the client device, physical characteristics of the avatar; and generating, based on the physical characteristics, the avatar for inclusion in the supplemental content.
9. The method of claim 1, further comprising: receiving, during streaming of the supplemental content, a user input from the at least one client device of the plurality of client devices to skip a portion of the supplemental content; and based at least in part on the user input: stopping the streaming of the supplemental content to the at least one client device; and resuming the streaming of the interactive multimedia presentation to the at least one client device.
10. The method of claim 1, further comprising: determining an amount of the released bandwidth available to the at least one client device; and encoding the supplemental content based at least in part on the amount of the released bandwidth available to the client device.
11. A system for providing interactive cloud-based supplemental content for an interactive multimedia presentation, the system comprising: a network interface configuration to allow communications between a plurality of servers and a plurality of client devices; and control circuitry configured to: reserve bandwidth for use by a first interactive content engine for streaming an interactive multimedia presentation to the plurality of client devices; determine that supplemental content is to be inserted into the interactive multimedia presentation streamed to at least one client device of the plurality of client devices by the first interactive content engine; in response to determining that the supplemental content is to be inserted into the interactive multimedia presentation, release at least a portion of the bandwidth reserved for use by the first interactive content engine for streaming the interactive multimedia presentation; stream, by a second interactive content engine using the released bandwidth, the supplemental content to the at least one client device of the plurality of client devices; upon receiving an indication that streaming of the supplemental content to the at least one client device of the plurality of client devices has ended, reserve the released bandwidth for use by the first interactive content engine; and resume streaming, by the first interactive content engine using the reserved bandwidth, the interactive content.
12. The system of claim 11, wherein the control circuitry is further configured to: monitor the streaming of the interactive multimedia presentation; wherein the control circuitry is configured to determine that the supplemental content is to be inserted into the interactive multimedia presentation by detecting an indication of supplemental content to be inserted into the interactive multimedia presentation.
13. The system of claim 11, wherein the control circuitry is further configured to: monitor the streaming of the interactive multimedia presentation; wherein the control circuitry is configured to determine that the supplemental content is to be inserted into the interactive multimedia presentation by detecting, based at least in part on the monitoring, that the streaming of the interactive multimedia presentation has stopped.
14. The system of claim 11, wherein the interactive multimedia presentation comprises a video game.
15. The system of claim 14, wherein the supplemental content comprises a cutscene.
16. The system of claim 11, wherein the supplemental content comprises an advertisement.
17. The system of claim 11, wherein the supplemental content is customized for a user profile of

the at least one client device, wherein the control circuitry is further configured to: access a user profile associated with the client device; and modify a portion of the supplemental content based on the user profile.

18. The system of claim 17, wherein the interactive content is a video game, the supplemental content is a cutscene of the video game, and the portion comprises an avatar of a user of the client device, wherein the control circuitry is further configured to modify the portion of the supplemental content by: determining, based on the user profile associated with the client device, physical characteristics of the avatar; and generating, based on the physical characteristics, the avatar for inclusion in the supplemental content.

19. The system of claim 11, wherein the control circuitry is further configured to: receive, during streaming of the supplemental content, a user input from the at least one client device of the plurality of client devices to skip a portion of the supplemental content; and based at least in part on the user input: stop the streaming of the supplemental content to the at least one client device; and resume the streaming of the interactive multimedia presentation to the at least one client device.

20. The system of claim 11, wherein the control circuitry is further configured to: determine an amount of the released bandwidth available to the at least one client device; and encode the supplemental content based at least in part on the amount of the released bandwidth available to the client device.
