

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication
Kind Code
Publication Date
Inventor(s)

20250260949
A1
August 14, 2025
Sugar; Gary L. et al.

WIRELESS ROOM OCCUPANCY MONITOR

Abstract

A room occupancy monitor includes an antenna array configured to detect wireless transmissions from a tag device; a wireless transceiver configured to receive the wireless transmissions detected by the antenna array and produce receive signals; a processor configured to process the receive signals; and a motion sensor coupled to the processor and configured to wake up the processor in response to detecting when the tag device enters or exits a room. After the motion sensor wakes up the processor, the processor is configured to power on the wireless transceiver and run an algorithm on a sequence of estimates of a parameter, derived from the received signals from the tag, that is indicative of a distance between the tag and monitor, and array response vectors derived from the receive signals and determine when the tag device has entered or exited the room via the entryway.

Inventors: Sugar; Gary L. (Shaker Heights, OH), Tesfai; Yohannes (Silver Spring, MD), Vaidyanathan; Chandra (Rockville, MD), Diener; Neil R. (Hudson, OH)

Applicant: Emanate Wireless, Inc. (Ijamsville, MD)

Family ID: 1000008574548

Appl. No.: 19/180262

Filed: April 16, 2025

Related U.S. Application Data

parent US continuation-in-part 17730278 20220427 PENDING child US 19180262
us-provisional-application US 63701674 20241001
us-provisional-application US 63643131 20240506
us-provisional-application US 63182071 20210430
us-provisional-application US 63196276 20210603
us-provisional-application US 63236288 20210824
us-provisional-application US 63308160 20220209

Publication Classification

Int. Cl.: H04W4/029 (20180101); G01S13/75 (20060101); H04W4/33 (20180101); H04W4/80 (20180101)

U.S. Cl.:

CPC H04W4/029 (20180201); G01S13/75 (20130101); H04W4/33 (20180201); H04W4/80 (20180201);

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS [0001] This application claims priority to U.S. Provisional Application No. 63/701,674, filed Oct. 1, 2024, to U.S. Provisional Application No. 63/643,131, filed May 6, 2024, and is a continuation-in-part of U.S. application Ser. No. 17/730,278, filed Apr. 27, 2022, which in turn claims the benefit of U.S. Provisional Patent Application No. 63/182,071, filed Apr. 30, 2021; U.S. Provisional Patent Application No. 63/196,276, filed Jun. 3, 2021; U.S. Provisional Patent Application No. 63/236,288, filed Aug. 24, 2021; and U.S. Provisional Patent Application No. 63/308,160, filed Feb. 9, 2022, which are incorporated by reference herein in their entirety.

FIELD OF THE INVENTION

[0002] The present disclosure relates generally to real-time location systems (RTLs) and, more specifically, to monitoring systems that can locate active radio-frequency identification (RFID) tags in real-time.

BACKGROUND OF THE INVENTION

[0003] There is a strong market need for real-time location systems (RTLs) that can deliver room-level accuracy. Healthcare automation applications such as those used for hand hygiene enforcement or nurse call cancellation demonstrate this need. For example, when a patient presses a nurse call button in a hospital room, the nurse call corridor lights illuminate, and the nurse may receive a call on his or her wireless phone. When the nurse enters the patient's room, the RTL automatically detects and records their presence while canceling the call. Obviously, if the RTL did not very accurately detect that the nurse entered the correct room at the correct time, the RTL would not be very useful for this application.

[0004] There is also a desire for RTLs based on Bluetooth™ Low Energy (BLE) wireless technology. Some advantages of BLE-based RTLs are lower cost, longer battery life and device portability afforded by the pervasiveness of BLE. As of the time of this writing, however, the state-of-the-art RTLs that can provide room-level accuracy use infrared (IR) and ultrasound technology to locate the wireless tag devices that they locate. Both IR and ultrasound technologies are non-standardized, costly to deploy and less energy efficient than BLE-only RTLs. A BLE-based (or more generally, wireless radio frequency (RF)-based) RTL that can provide room-level accuracy would be quite valuable for current and future location system applications.

SUMMARY OF THE INVENTION

[0005] The present disclosure describes a wireless RF-based RTL that can deliver room-level accuracy. According to one aspect, a wireless room occupancy monitor is provided. The wireless room occupancy monitor includes: an antenna array configured to detect wireless transmissions from a tag device; a wireless transceiver configured to receive the wireless transmissions detected by the antenna array and produce receive signals; a processor configured to process the receive signals; and a motion sensor coupled to the processor and configured to wake up the processor in response to detecting when the tag device enters or exits a room. The antenna array and motion sensor are configured to be mounted on a ceiling of the room, just inside an entryway to the room. In operation, after the motion sensor wakes up the processor, the processor is configured to power on the wireless transceiver and run an algorithm on a sequence of estimates of a parameter, derived from the received signals from the tag, that is indicative of a distance between the tag and monitor, and array response vectors derived from the receive signals, and determine when the tag device has entered or exited the room via the entryway.

[0006] According to another aspect, a room occupancy detection system is provided. The room occupancy detection system includes: one or more room occupancy monitors configured to detect entries into a room and exits from the room of one or more tag devices, and to produce room occupancy detection events; and a server configured to receive the room occupancy detection events from the one or more room occupancy monitors. Each of the one or more room occupancy monitors comprises: an antenna array configured to detect wireless transmissions from the one or more tag devices; a wireless transceiver configured to receive the wireless transmissions detected by the antenna array and produce receive signals; a processor configured to process the receive signals; and a motion sensor coupled to the processor and configured to wake up the processor in response to detecting when one of the one or more tag devices enters or exits the room. Each of the one or more room occupancy monitors is configured to be mounted on a ceiling of the room, just inside an entryway to the room. In operation, after the motion sensor wakes up the processor on any one of the one or more room occupancy monitors, the processor is configured to power on the wireless transceiver and run an algorithm on a sequence of estimates of a parameter, derived from the received signals from the tag, that is indicative of a distance between the tag and monitor, and array response vectors derived from the receive signals, and determine when the tag device has entered or exited the room via the entryway.

[0007] In accordance with still another aspect, a method is provided for training a machine learning algorithm for room occupancy monitoring. The method includes: storing receive signals produced by one or more room occupancy monitors as one or more tag devices enter into and exit one or more rooms, wherein the one or more room occupancy monitors are installed on a ceiling inside an entry of each of the one or more rooms, and wherein each of the one or more room occupancy monitors produces the receive signals from wireless transmissions from the one or more tag devices detected by an antenna array of the one or more room occupancy monitors; generating ground truth information comprising a time when each of one or more persons or machines wearing, carrying or using one or more of the tag devices entered or exited a room of the one or more rooms, an identity of the one or more tag devices that entered or exited the room of the one or more rooms, and the identity of each room occupancy monitor that detected one or more tag devices entering or existing the room of the one or more rooms; and providing the ground truth information and data descriptive of the receive signals to a machine learning algorithm to train the machine learning algorithm to detect room entries or exits using the ground truth information and the receive signals.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0008] The foregoing and other features of the present disclosure will become apparent to those skilled in the art to which the present disclosure relates upon reading the following description with reference to the accompanying drawings.

[0009] FIG. 1 is an illustration showing a person wearing a wireless tag device approaching a room with a room occupancy monitor installed in the entryway, according to embodiments presented herein.

[0010] FIG. 2 is a top-level view of the scenario shown in FIG. 1.

[0011] FIG. 3 is an exemplary electrical block diagram of a room occupancy monitor, in accordance with one embodiment.

[0012] FIG. 4 is an exemplary block diagram showing the mechanical components of a room occupancy monitor, in accordance with one embodiment.

[0013] FIG. 5 is a diagram showing a view of an antenna array of a room occupancy monitor from a view looking at the monitor

from below, in accordance with one embodiment.

[0014] FIG. 6 is a diagram showing arrangement of an antenna array of a room occupancy monitor relative to a rectangular coordinate system.

[0015] FIG. 7 is a block diagram of a room occupancy monitoring system that includes a plurality of room occupancy monitors, in accordance with one embodiment.

[0016] FIG. 8 is a diagram illustrating the signal processing performed on a room occupancy monitor, in accordance with one embodiment.

[0017] FIG. 9 is an exemplary plot showing the output of a beamforming cost metric evaluated over a set of grid of positions centered around a room occupancy monitor.

[0018] FIG. 10 is an exemplary plot showing the output of a position estimation cost metric evaluated over a set of grid of positions centered around a room occupancy monitor.

[0019] FIG. 11 is a flow chart depicting a disambiguation procedure used when multiple monitors detect a room entry for the same tag device, according to an example embodiment.

[0020] FIG. 12 is a flow chart of a method for training a machine learning (ML) based room occupancy detection algorithm, according to an example embodiment.

[0021] FIG. 13 is a flow chart of a non-ML-based room occupancy detection algorithm, in accordance with one embodiment.

[0022] FIG. 14 is a block diagram of a room occupancy monitor with additional sensors to improve detection performance, according to an example embodiment.

[0023] FIG. 15A is an exemplary block diagram showing how a dual-zone infrared (IR) sensor could be used as an additional input signal source of a room occupancy detection monitor, in accordance with one embodiment.

[0024] FIG. 15B is an exemplary plot showing the voltage output over time of the dual-zone IR sensor of FIG. 15A as a person walks into a room.

[0025] FIG. 16 is a diagram showing an example of how one or more room occupancy monitors may be used to detect zone occupancy of one or more tag devices in a multi-zone room, in accordance with one embodiment.

[0026] FIG. 17 is a flow chart of a method for training a ML-based zone occupancy detection algorithm for use in multi-zone rooms, in accordance with one embodiment.

[0027] FIG. 18 is a flow chart of a non-ML-based zone occupancy detection algorithm for use in multi-zone rooms, in accordance with one embodiment.

[0028] FIG. 19 is a diagram showing an example of how a room occupancy monitor could be used to detect disinfectant dispenser usage in a hand hygiene monitoring application, in accordance with one embodiment.

[0029] FIG. 20 is a flow chart of a method for training a ML-based disinfectant dispenser usage detection algorithm, in accordance with one embodiment.

[0030] FIG. 21 is a flow chart of a non-ML-based disinfectant dispenser usage detection algorithm, in accordance with one embodiment.

[0031] FIG. 22 is flow chart of a computationally efficient method for winnowing tag devices before running a more computationally intensive room occupancy detection algorithm, according to an example embodiment.

[0032] FIG. 23 is a diagram showing differences in how the room occupancy monitor behaves in rooms that open to the left or right, according to an example embodiment.

[0033] FIG. 24 illustrates a hardware block diagram of a computing device that may perform functions associated with various operations described herein, according to an example embodiment.

[0034] FIG. 25 is an exemplary electrical block diagram of an ultra-wideband (UWB)-based room occupancy monitor, in accordance with one embodiment.

[0035] FIG. 26 is a flow chart of a room occupancy detection procedure for a UWB-based room occupancy monitor, in accordance with one embodiment.

[0036] FIG. 27 is a flow chart of a self-supervised learning procedure for training the parameters of an ML-based room occupancy detection algorithm after the monitor is installed in a room, in accordance with one embodiment.

[0037] FIG. 28 illustrates example output of an accelerometer motion sensor as a person wearing the motion sensor is walking, in accordance with one embodiment.

[0038] FIG. 29 illustrates how a millimeter-wave radar sensor may be used as an additional input signal source of a room occupancy monitor, in accordance with one embodiment.

[0039] FIG. 30 is an exemplary two-dimensional illustration of point-cloud data taken from a millimeter-wave radar sensor deployed in a room occupancy monitor, in accordance with one embodiment.

[0040] FIG. 31 is a functional block diagram showing primary and secondary room occupancy detection algorithms running on a room monitor's processor, in accordance with one embodiment.

[0041] FIG. 32 is a flowchart showing the procedural flow for the primary and

[0042] secondary room occupancy detection algorithms, in accordance with one embodiment,

[0043] FIG. 33 is a diagram showing a safe zone in a room for secondary room occupancy detection, in accordance with an example embodiment.

DETAILED DESCRIPTION OF THE INVENTION

[0044] From the above description, those skilled in the art will perceive improvements, changes and modifications. Such improvements, changes and modifications are within the skill of one in the art and are intended to be covered by the appended claims.

[0045] Presented herein is a Room Occupancy Monitor and a Room Occupancy Monitoring System. The Room Occupancy Monitor is, in one form, a battery-powered device that can be installed inside a room and configured to receive transmissions from one or more wireless tag devices to determine when any of the tag devices enter or exit the room. The monitor could be

installed in a hospital patient room, for example, and used in a nurse-call cancellation application/use case to determine when a nurse wearing a wireless (e.g., Bluetooth™ wireless 5.1) tag or badge has entered a room to visit a patient and automatically cancel the pending nurse call request at the nurse's station. It could alternatively be used in automated hand hygiene compliance applications to make sure doctors and nurses are disinfecting their hands each time they enter or exit a patient room. There are other uses of the Room Occupancy Monitor and Room Occupancy Monitoring System in addition to healthcare applications. The monitor could be installed on the ceiling in a hallway and used as a chokepoint monitor to detect a passersby in either direction, in hotel rooms for staff duress applications, in industrial or commercial environments to provide access control to various parts of a building using Bluetooth badges, in museums or retail stores to monitor customer behavior or to provide waypoint information. There are numerous other uses of the Room Occupancy Monitor and Room Occupancy Monitoring System, not specifically mentioned herein.

[0046] The tag devices could be carried or worn by one or more persons or could be carried, mounted, integrated in, or attached to one or more machines or equipment that may move in and out of a room or space of interest. As far as these techniques are applicable to tracking movement (in and out of a room or space of interest) of machines or equipment, the machines or equipment may be configured to move under its own power and control autonomously (e.g., robot) or by physical assistance from a human.

[0047] FIG. 1 illustrates how a room occupancy monitor could be installed and used to monitor a room, in accordance with one embodiment. The room occupancy monitor (“monitor”) 1 is installed just inside an entryway 7 to a room 4. For example, the monitor 1 is affixed to the ceiling 8 a predetermined distance d (e.g., approximately six inches) inside a threshold to the entryway 7. The monitor 1 is equipped with a motion sensor 6, such as a passive infrared (PIR) sensor with a narrow field of view (FOV) 5. The motion sensor 6 is used to wake up the monitor 1 from a low-current sleep state when a person 2 wearing or carrying a tag device 3, enters or exits the room 4. The tag device 3 may also be affixed to another object such as an infusion pump or ultrasound machine which is pushed in or out of the room 4 by person 2.

[0048] FIG. 2 shows a top-down view of the scenario illustrated in FIG. 1. This illustration shows how the motion detector's narrow FOV 5 spans the entire width of the room entryway 7 but extends only slightly into and/or outside of the room 4. The narrow FOV 5 allows motion detector to conserve the monitor's battery energy by triggering typically only when a person enters or leaves the room, and not triggering when there is motion inside the room or outside the room in a hallway. The monitor 1 may be mounted proximate to the room entryway 7.

[0049] Turning now to FIG. 3, a schematic block diagram of the monitor 1 is shown, in accordance with one embodiment. The monitor 1 contains an antenna array 10 to detect the wireless signals transmitted by a tag device and a radio frequency (RF)/Baseband transceiver 13 to downconvert, demodulate and decode the received signals according to a wireless standard, such as Bluetooth™ 5.1. The RF/Baseband transceiver passes the decoded output to processor 15. When the motion sensor 6 detects motion, it wakes up the processor 15 to begin processing the received signals from the transceiver 13 to determine whether (the person 2 wearing, carrying, or moving) the tag device 3 has entered or exited the room 4, as depicted in FIG. 1. The transceiver 13 may include an RF switch 12 to select which of the antenna array antenna elements of the antenna array 10 to use when receiving the wireless signal from a tag. The RF switch 12 could also be used to switch to using an omnidirectional antenna 11 when the RF/Baseband transceiver 13 needs to transmit a signal, to or receive a signal from, the tag device 3 or some other device without having to use the direction of arrival information that the antenna array 10 provides.

[0050] The monitor 1 may include one or more batteries 16 to provide power to the electronics without requiring a cable run during installation. Alternatively, the monitor 1 could be powered through a Power-over-Ethernet (POE) cable interface 17, or a standard DC power supply 18 which plugs into an external AC mains.

[0051] The motion sensor 6 could be a passive infrared (PIR) type of sensor, which may be configured to consume very little current when no motion is detected. A number of other variations for the motion sensor 6 include ultrasound, infrared or laser ranging sensors, a thermopile sensor, or an RF Doppler motion sensor.

[0052] Other sensors 19, such as a multi-zone IR thermopile sensor, a digital camera or an inertial motion unit (IMU), could be used to provide additional inputs to a room occupancy monitoring algorithm, described further below.

[0053] Turning now to FIG. 4, a physical illustration of the monitor 1 is shown, in accordance with one embodiment. In this embodiment, the monitor electronics, such as the RF/Baseband transceiver 13, processor 15 and RF switch 12, are mounted on a first printed circuit board (PCB) 20, which sits on top of a second PCB 22 containing the antenna array 10. Both PCBs 20 and 22 are mounted inside a plastic enclosure 21, along with three batteries 16. The motion sensor 6, in the arrangement shown in FIG. 4, may be a PIR sensor with a Fresnel lens mounted on the bottom of the enclosure 21. The enclosure 21 may have plastic mounting tabs 23 used to attach the monitor 1 to a ceiling tile ceiling with metal crossbars, and held in place using zip-ties, double-sided tape, Velcro® fastener, magnets, or some other mounting mechanism.

[0054] Turning now to FIG. 5, a view of the antenna array 10 inside the monitor 1 is shown from a view looking up at the monitor 1 from below, in accordance with one embodiment. In this embodiment, the antenna array 10 is implemented as a set of 16 square patch antennas 10A-1 to 10A-16, each of which is etched into the copper layers of PCB 22. Again, the monitor is mounted proximate the room entryway 7.

[0055] A monitor-centric rectangular coordinate system is shown in FIG. 5. The origin of the coordinate system is the center of the antenna array 10. The X and Y axes are co-planar with the array PCB 22, with the positive X axis 30 extending directly out the entryway 7, and the positive Y axis 31 pointing to the left of the entryway 7 when looking into the room from outside. The positive Z axis 32 is perpendicular to the PCB 22 and points down from the array center toward the floor.

[0056] There is a spacing 35 between the centers of two adjacent patch elements. This spacing 35 may be a third of a wavelength. In a multipath-free environment, this would ensure a measured phase shift of at most 120 degrees between any two adjacent antenna elements regardless of where the transmitter of the tag device 3 is positioned relative to the antenna array 10.

[0057] Reference is now made to FIG. 6, which illustrates a side perspective view of the room 4 with a three-dimensional (3D) perspective of the monitor-centric rectangular coordinate system. FIG. 6 also shows spherical coordinates as well, with the r axis

33 representing the distance from the origin, the θ axis 34 representing the azimuth angle, and the ϕ axis 35 representing the elevation angle to the tag device 3.

[0058] In other embodiments, the monitor 1 could be installed at an angle on the ceiling or wall of the room entryway, with the plane of its antenna array 10 aiming into the room 4 instead of straight down from the ceiling. The monitor 1 could be installed on the wall on either side of the entryway at the approximate height where the tag device 3 is worn by the user 2, with the plane of the antenna array 10 oriented parallel to the wall.

[0059] Referring now to FIG. 7, a room occupancy monitoring system 40 that operates with one or more monitors 1, is shown. Reference is also made to FIG. 3 in connection with the description of FIG. 7. In the system 40, a plurality of monitors 1 are installed in a number of different rooms 4. Each monitor 1 in the system could be used to monitor one or more tag devices 3 or other types of wireless emitters, called "other devices" 41. Other devices 41 may include badges, medical equipment with asset tags attached, wrist band devices, as well as consumer devices, such as Smartphones, tablets, Smart Watches, etc., that use a wireless standard, such as Bluetooth™ 5.1.

[0060] Each monitor 1 in system 40 performs a room monitoring algorithm to determine when tag devices 3 or other devices 41 enter or exit the monitored rooms. Whenever a monitor 1 detects a room entry or exit, it wirelessly transmits to one or more nearby gateway devices 42 a room occupancy detection event message including information describing or indicating the event type (i.e., entry or exit), the tag device ID (for example, Bluetooth™ or Media Access Control (MAC) address of the tag device 3 or other device 41), and the monitor ID (for example, the Bluetooth™ or MAC address of the monitor). The monitors 1 could transmit the room detection event message using the Bluetooth™ 5.1 RF/Baseband transceiver 13 that was used to receive transmissions from the tag device 3 via the antenna array 10. The monitors 1 could set the RF switch 12 to transmit the room detection event message using the omnidirectional antenna 11 to increase the chances that it is picked up by at least one gateway device 42. Each gateway device 42 that received the room detection event message relays the message to a server 43 via network 44. The server 43 removes any duplicate room detection event messages, and then formats and forwards a final room detection event to another server 45 (for example, a nurse call or hand hygiene monitoring server in a hospital) external to the room monitoring system 40.

[0061] Turning now to FIG. 8, a signal processing data flow 50 is described, in accordance with one embodiment. Reference is also made to FIG. 3 for purposes of this description of FIG. 8. Wireless transmissions sent from the tag device 3 are detected at the antenna array 10 of a monitor, and downconverted, demodulated and decoded in transceiver 13, producing receive signals which are then presented to processor 15. If the wireless signals sent by the tag device 3 are formatted according to the Bluetooth™ 5.1 wireless standard, they could include a Constant Tone Extension (CTE) segment at the end of the packet, allowing the RF/Baseband transceiver 13 to switch through all the elements of the antenna array 10 during the segment and to digitize the in-phase (I) and quadrature (Q) components of the CTE signal after each switching period. Each CTE received by the RF/Baseband transceiver 13 produces an array response vector 51 containing N complex I/Q measurements which can be written as follows:

$$[00001] \ z = [I_1 + jQ_1 \text{ .Math. } I_N + jQ_N], \quad (1)$$

where N is the number of antennas in the antenna array 10, and a received signal strength indicator (RSSI) measurement 52 containing the received signal strength of the received packet in dBm. The RSSI measurement 52 is typically taken on a single antenna, which is often the first antenna element of array 10.

[0062] Each I/Q measurement $I_{\text{sub}.n} + jQ_{\text{sub}.n}$ in the array contains the in-phase (I) and quadrature (Q) components of the CTE tone as received from the nth antenna. The amplitude and phase of the CTE tone can be derived from the I/Q samples using

$$[00002] \sqrt{I_n^2 + Q_n^2} \text{ and } \tan^{-1} Q_n / I_n,$$

[0063] respectively. Because the I/Q values depend on the gain of the receiver, which generally varies from packet to packet, the above definition of z in (1) only contains information about the relative amplitude of the CTE on each antenna element. The definition can be modified as follows to include absolute signal level information:

$$[00003] \ z = \frac{10^{R_1/20}}{\sqrt{I_1^2 + Q_1^2}} \text{ .Math. } [I_1 + jQ_1 \text{ .Math. } I_N + jQ_N]. \quad (2)$$

The $R_{\text{sub}.1}$ parameter in the above expression represents the RSSI estimate provided by

$$[00004] \frac{10^{R_1/20}}{\sqrt{I_1^2 + Q_1^2}}$$

in the transceiver 13 in dBm as seen through the first antenna. The scaling factor in the expression is used to scale all the elements of z by the gain of the receiver to make it so that the magnitude squared of any component of z has units of milliwatts (mW). Therefore, taking 20 times the base ten logarithm of the magnitude of any of the elements of z will yield an estimate of the power seen through that antenna path in decibels relative to one milliwatt, or dBm. In particular, for the first element $z_{\text{sub}.1}$ of z, one can verify that

$$[00005] 20 \text{ .Math. } \log_{10} \text{ .Math. } z_1 \text{ .Math. } = 20 \text{ .Math. } \log_{10} [10^{\frac{R_1}{20}}] = R_1$$

which yields $R_{\text{sub}.1}$, which is as expected. Another useful RSSI metric that can be computed is the average RSSI overall all N antennas, which can be written

[00006]

$$R_{\text{Av}} = 10 \text{ .Math. } \log_{10} \frac{1}{N} \text{ .Math. } \sum_{n=1}^N \text{ .Math. } z_n \text{ .Math. }^2 = R_1 + 10 \text{ .Math. } \log_{10} \frac{1}{N} (1 + \frac{\text{ .Math. } z_2 \text{ .Math. }^2}{\text{ .Math. } z_1 \text{ .Math. }^2} + \text{ .Math. } + \frac{\text{ .Math. } z_N \text{ .Math. }^2}{\text{ .Math. } z_1 \text{ .Math. }^2}), \quad (2)$$

which, like $R_{\text{sub}.1}$, also has units of dBm. In a highly reflective indoor wireless environment, the parameter $R_{\text{sub}.Av}$ defined in (2) has the advantage of having N times less variance than $R_{\text{sub}.1}$ or any other RSSI measurement from a single antenna. Thus, a new definition for the array response vector z based on the lower variance RSSI estimate in (2) is as follows:

$$[00007] \ z = \frac{10^{R_{Av}/20}}{\sqrt{I_1^2 + Q_1^2}} \cdot \text{Math.}[I_1 + jQ_1] \cdot \text{Math.}[I_N + jQ_N] \cdot (3)$$

[0064] Equation (3) shows a method for combining the I/Q samples **51** and the single-antenna RSSI **52** signals that are received from RF transceiver **13** for each received CTE from a tag into one combined array response vector.

[0065] The data processing flow **50** includes a spatial signal processing step **55**, a resampling step **56**, a data conditioning step **57**, and a room occupancy detection algorithm **58**, all of which may be performed on each monitor, to generate room occupancy detection events **59**. The room occupancy detection events **59** are provided to server **43**. The server **43** executes a disambiguation procedure **60** to generate final room occupancy events **61**.

[0066] The spatial signal processing step **55** is now described with reference to FIG. **8**, together with FIGS. **9** and **10**.

[0067] Spatial signal processing step **55** generates spatial positioning information for the tag device **3** from the received sequence of array response vectors to estimate the position of tag device **3** relative to the monitor **1**. The spatial positioning information is contained in likelihood vs. position vectors output from the spatial signal processing step **55**, with each likelihood vs. position vector containing information related to the likelihood (or probability) that the tag device **3** is positioned at a particular grid point position over a set of candidate grid point positions in the vicinity of monitor **1**.

[0068] Let p.sub.1, . . . , p.sub.M represent a set of M candidate points in 3D space using the monitor-centric coordinate system defined earlier. Each point p.sub.m can be written as a 3-vector

$$p_{\text{sub}.m} = [x, y, z],$$

where x, y, and z are the components of p.sub.m along the X, Y and Z axes **30**, **31** and **32** respectively. As an example, a grid of M=400 points could be defined over a 20 foot by 20-foot rectangular region centered at the monitor, where the grid points are uniformly spaced by one foot in either the X or Y directions, and each point is assigned a fixed Z axis value of 4.5 feet from the monitor. A mathematical description of this example grid is as follows:

$$[00008] \ p_m = [-20, -10, 4.5] + [2m - 20] \cdot \text{Math.}[1] \cdot \text{Math.}[0] \cdot \text{Math.}[0], m = 0, \dots, 399 \quad (4)$$

[0069] It should be noted that although the grid points in the above example are uniformly spaced, there is no requirement for them to be.

[0070] One well-known way of generating likelihood vs. position information for tag device **3** is to use beamforming. For each candidate grid position p.sub.m, the beamforming algorithm computes the following:

$$[00009] \ y_m = a_m^H R a_m \cdot (5)$$

[0071] The vector a.sub.m is an “expected array response vector” with the same dimension N as the array response vector z. It contains the (complex) array response that would be observed at the antenna array **10** for an open-space wireless transmission from a device positioned at grid point p.sub.m. The matrix R is a complex N×N Hermitian covariance matrix of the array response vector z. R can be estimated using single array response vector z as follows: R=zz.sup.H. Substituting this definition of R into equation (5) yields y.sub.m=|a.sub.m.sup.Hz|.sup.2. This form of beamforming can be viewed as a “spatial matched filtering” operation, since it is effectively cross-correlating a received array response vector z with a set of M hypothetical responses—one for each candidate spatial position. For a transmitting device positioned at, for example, the kth candidate grid position in open space, one would expect to see maximum power precisely at the kth grid point. This is illustrated in FIG. **9**, where the cost function (5) is evaluated at each point on the example grid defined in (4) and shown in a contour plot. It can be shown that for an open space wireless environment, equation (5) is the log-likelihood function for the position of the tag device **3** given a received array response vector from the tag. The maximum likelihood position can be found by finding the grid position p.sub.m0 that maximizes (5). For the example shown in FIG. **9**, that position p.sub.m0 is at [1.5, 1.5, 4.5] feet on the rectangular monitor-centric coordinate system defined earlier herein.

[0072] For a more robust beamforming implementation, instead of using only a single array response vector to estimate the covariance matrix R, one could improve the estimate by using a running average over the L most recently received array response vectors. Using this approach, the covariance matrix after the nth array response vector received would be computed as:

$$[00010] R_n = \frac{1}{L} \sum_{l=0}^{L-1} z_{n-l} z_{n-l}^H \cdot$$

[0073] Another popular array-based position estimation technique is Multiple Signal Classification (MUSIC). MUSIC is a so-called subspace technique that uses a singular value decomposition (SVD) of the covariance matrix R to find a basis for its null-space, and then uses a search to find spatial grid positions that have expected array responses that are orthogonal or nearly orthogonal to the null-space. The MUSIC cost function can be written

$$[00011] \ y_m = \frac{1}{a_m^H U U^H a_m} \quad (6)$$

where U is the N×(N-r) matrix containing the lowest order N-r left singular vectors for the covariance matrix R, and r is an estimate of the rank of R. Note that the columns of the U matrix form an orthonormal basis for the null-space of R. For grid positions that are close to the actual position of the target wireless transmitter, the denominator of (6) becomes small (and thus its reciprocal (6) becomes large) since at these positions the a_m vector will approximate the actual received array response vector, which, by definition, is orthogonal or nearly orthogonal to the columns of the U matrix.

[0074] Since the goal in MUSIC is to find grid positions that maximize the cost function (6), that cost function reflects in some way the likelihood or probability of the tag device **3** being positioned at grid point p.sub.m. The MUSIC approach is a popular alternative to beamforming-based position estimation algorithms because it is known to be able to find the target's position with higher spatial resolution. A contour plot of the MUSIC response (6) for the previous example is shown in FIG. **10**. Note from the figures that the MUSIC response contains a very sharp peak at the true target position—much sharper than the beamformer response (5) shown in FIG. **9**.

[0075] There are other well-known array-based position estimation techniques in addition to the classical beamforming and

MUSIC approaches discussed above, including Minimum Variance Distortionless Response (MVDR) beamforming and Estimation of Signal Parameters by Rotational Invariance Techniques (ESPRIT). Any of these techniques could be used to provide likelihood vs. position information for the tag device **3** as part of the spatial processing algorithm **55**.

[0076] Referring back to FIG. **8**, following the spatial signal processing step **55**, resampling step **56** may be used to interpolate and replenish any missing samples due to missed CTE packets. In one form, resampling could be done using a nearest-neighbor approach. Other more intricate but well-known forms of interpolation could be used to fill in any missing samples, including linear or quadratic interpolation, spline, sinc, and polynomial interpolation approaches. It should be noted that the resampling step **56** could be done before the spatial signal processing **55**, if desired.

[0077] Following the resampling step **56** is a data conditioning step **57** followed by execution of room occupancy detection algorithm **58**. The data conditioning step **57** is only required if the room occupancy detection algorithm **58** is a machine learning (ML) algorithm, such as a neural network. The data conditioning step **57** is used to adjust the mean and variance of the input parameters to ensure good performance. The output from room occupancy detection algorithm **58** consists of room occupancy detection events **59**, which consist of either entry or exit events that are sent when the tag device **3** is detected as having either entered or exited the monitored room **4**. The room occupancy detection events **59** detected at the monitor **1** could be sent wirelessly from the monitor to the server **43** via one or more wireless gateways **41** using the Bluetooth™ wireless data protocol. For these transmissions, the omni-directional antenna **11** may be used on the monitor **1** to give reliable performance.

[0078] Since wireless RF signals easily propagate through walls, when a person **2** carrying a tag device **3** enters a monitored room, it is possible for two or more monitors to detect the tag device **3** as having entered their rooms at or around the same time. This often happens when the person enters two neighboring rooms with entryways that are close to each other—sometimes as little as 4 feet between entryway centers, or when one room is located directly above another room on different floors of a multi-story building. It is for this reason that the server **43** from time to time performs the disambiguation procedure **60** to identify which of the two or more monitored rooms was entered.

[0079] Reference is now made to FIG. **11**, which shows a flow chart describing the disambiguation procedure **60**. In step **62**, the server **43** receives room entry detection events from the monitors **1** via the gateways **41**. The room entry detection events sent from each monitor in step **62** contain both the identity of the tag device and the room where the detection event was made, as well as a confidence metric quantifying the monitor's level of confidence that the tag device entered that room. If the room occupancy detection algorithm **58** is a machine learning algorithm such as a neural network, the confidence metric could be the soft output (in units of logits, for example) from the neural network. If the room occupancy detection algorithm **58** uses a non-machine learning process, the confidence metric could be based on one or more normalized and scaled algorithmic parameters, such as received signal strength or X/Y position. In step **63**, when a room entry detection event is received at the server indicating that a tag has entered a room, the server waits for a waiting period—typically 1-2 seconds—to determine if any other room entry detection events for that same tag device arrive from any other rooms. In step **64**, if two or more room entry detection events for the same tag device are received from different rooms within the 1-2 second waiting period, the server selects the room served by the monitor having the highest confidence metric as the room that was entered, thus disambiguating the multiple entry events. In step **65**, the server records that room in a database, reports to its outbound software APIs that the tag entered the highest-confidence room, and then returns to looking for more entry events in step **62**.

[0080] The room occupancy detection algorithm **58** could be implemented using a machine learning approach entirely, a wholly “traditional” (i.e., non-machine learning) algorithm, or a combination of both. If a machine learning algorithm is used in whole or in part, that algorithm is trained using training data that behaves as if it came from a real deployed system and operating environment.

[0081] Reference is now made to FIG. **12**, which shows a flow chart of a method **70** for training such a machine learning algorithm. Reference is also made to FIG. **7** for purposes of the description of FIG. **12**. In step **71**, one or a plurality of monitors **1** are installed on the ceiling inside one or multiple rooms **4** in an area that resembles where the system **40** is intended to be used once in production. The monitor **1** should be installed on the ceiling just as it would be installed in practice, i.e., about 6 inches from the entry threshold in the center of the entryway. The monitor **1** should also be configured to operate just as it would in practice, i.e., the monitor should only wake up from sleep and begin receiving and storing data after it has been woken up by the motion sensor **6**.

[0082] In step **72**, each of the monitors is configured to receive, downconvert and detect wireless transmissions from one or more tag devices **3** as people **2** wearing, carrying or moving the tag devices walk into, out of, around the outside and around the inside of the monitored rooms **4**. The monitors **1** are also configured to generate and record receive signals for each of the detected wireless transmissions. The receive signals could include the received array response vectors and RSSI estimates for each received transmission from the tags. The receive signals could also include the results of operations applied to the array response vectors or RSSI information, such as beamforming or MUSIC spatial processing. The operations could also include the results of other operations, such as resampling and data conditioning.

[0083] In step **73**, ground truth information containing the tag device ID, room or monitor ID, entry time and exit time every time a tag device enters or exits a room is recorded and stored. The ground truth could be obtained by recording a video of the people wearing the tags walking in and out of the monitored rooms in step **72** and later playing back the video to determine with precision when each tag entered or exited each monitored room. In step **74**, the recorded receive signals and ground truth information is stored in a training database.

[0084] In step **75**, steps **71-74** are repeated at numerous different locations having different ceiling heights, floor layouts and construction materials. The people used to wear, carry or move the tag devices could enter or exit the rooms at different walking speeds and entry/exit angles, and change the way they wear, carry or move the tag to emulate the way the tag device will be used in practice as much as possible.

[0085] In step **76**, the training procedure in steps **71-75** is terminated when it is determined that enough training data has been taken. This determination is usually made by partitioning the receive signals and ground truth information stored in the database

into a training data set and a test data set. In one example, there is enough training data when adequate performance is achieved on the data set, and the machine learning algorithm performs adequately on the test data set and on data taken from new rooms or buildings. If the algorithm does not perform as well on newly obtained data from a new room, person, tag or environment as it typically does in the test set, this could mean that the algorithm should be trained on that new room or environment.

[0086] One class of machine learning algorithm that performs well for this application are so-called stateful deep learning algorithms, i.e., neural networks (NNs) that update and propagate state information at each input/output time step. Any of the following types of neural networks could be used for this purpose: recurrent NN (RNN), long-short-term memory NN (LSTM), convolutional NN (CNN) or a gated recurrent unit NN (GRU). It is also possible to use stateless NNs for this application as well as non-deep learning machine learning algorithms, such as decision trees, random forests, support vector machines, K-nearest neighbors, and the like. The training procedure **70** can be used to train and test any of these machine learning algorithms-either stateful or stateless.

[0087] It is also possible to implement room occupancy detection algorithm **58** using a more traditional, non-machine learning approach. One such approach is illustrated in FIG. **13**, in accordance with one embodiment. FIG. **13** illustrates a flow chart for a (non-machine learning) room occupancy detection process **80**. Reference is also made to FIGS. **3**, **5** and **6** in connection with this description of FIG. **13**. In step **81**, the monitor **1** remains asleep in a low current state until it receives a wake-up signal from the motion sensor **6**. In step **82**, once the processor **15** wakes up, it powers on the RF/Baseband transceiver **13**, configuring it to pass the array response and RSSI information received from any tags to the processor **15**, and for each array response received, the processor **15** estimates the X axis position of the transmitting tag using the monitor-centric rectangular coordinate system defined earlier. The X axis position of the tag could be estimated by evaluating the beamforming or MUSIC cost functions (5) or (6) over a grid of candidate target (X,Y,Z) positions, selecting the most likely grid position by finding the minimizer of (5) or (6) over the grid, and noting the X position of the minimum-cost grid point.

[0088] In step **83**, the process **90** looks to see if the estimated X position of the tag crosses zero from positive to negative.

[0089] In step **84**, if a positive-to-negative X crossing was detected, other metrics are computed based on the sequence of received array response vectors, RSSI estimates, and any computed likelihood vs. position estimates up to that point. The other metrics could include the Y axis position at the time of the X axis zero crossing (which, if the tag went into the monitored room, should be between -1.5 and 1.5 feet for a 3 ft. wide doorway) which could be computed by noting the Y position of the minimum-cost grid point of equation (5) or (6), the elevation angle of the minimum-cost grid point at the time the X axis zero crossing (typically at least 80 degrees for an entry into the monitored room), the average RSSI across all antennas (see equation (2)). Another metric that could be used is a "beam width" metric, which can be defined as the spatial width (in feet) of the region in the beamforming heat map within 1-2 dB of the peak. The beam width can be shown to be typically much smaller when the tag is directly under the monitor vs. elsewhere. Any or all of the metrics computed above could be smoothed over time using a moving average process, a digital lowpass filter or a more sophisticated non-linear filter such as a Kalman or Particle filter.

[0090] In step **85**, the metrics computed in step **84** are compared against an appropriate set of thresholds to determine if a room entry was made, and if so, the monitor notifies the server of the entry in step **86**.

[0091] In step **87**, the process **90** looks to see if the estimated X position of the tag crosses zero from negative to positive, indicating a possible room exit has occurred. In step **88**, if a negative-to-positive X crossing was detected in step **87**, other metrics such as Y axis position at the time of X zero crossing, beam width, elevation angle, or RSSI at time of X zero crossing are computed. In step **89**, the other metrics computed in step **88** are compared against an appropriate set of thresholds to determine if a room exit was made, and if so, the monitor notifies the server of the exit in step **90**.

[0092] Reference is now made to FIG. **14**, which shows a functional design **100** of a monitor that is configured to be augmented to support additional sensors. Because of the additional information additional sensors provide about the position and direction of travel of the tag device beyond what is available using only the antenna array **10**, these sensors can be used to improve the performance of algorithm **58**. The top part of the FIG. **14** shows the RF/Baseband transceiver **13** communicating with RF switch **12** on the RF side, and outputting complex (I/Q) array response vectors **51** and RSSI **52** data to the spatial signal processing **55**, resampling **56** and data conditioning **57** blocks on the baseband side. The output from the data conditioning block **57** is presented to the room occupancy detection algorithm **58**. This is the data feed that is available to the room occupancy detection algorithm **58** when it uses only the antenna array signals.

[0093] The most basic additional sensor signal that could be provided as input to the algorithm is the output from a digital PIR sensor **101**. This could be the same physical sensor that is used for the motion sensor **6** described earlier that is used to wake up the processor **15** from sleep. The output of the digital PIR sensor **101** could be sampled by the processor **15** concurrently along with the array response **51** and RSSI **52** signals received from the transceiver **13** and provided as input to the room occupancy detection algorithm **58** along with those signals. The precise timing of the activity of the digital PIR sensor **101**, when combined with the I/Q samples **51** and RSSI samples **52** provides additional information that the room occupancy detection algorithm **58** could use to determine the room occupancy state of a tag device. For example, if the motion sensor **6** was triggered because of a person walking around inside the room but not entering or exiting, the timing characteristics of the motion sensor signal may align differently to the position vs. time information carried by the array signals than an actual room entry or exit.

[0094] Using an analog PIR **102** or IR thermopile **103** instead of the digital PIR sensor **101** would provide more information to the algorithm since the intensity of the IR signals is available. The intensity of the IR signal indicates how close the target is to the monitor or whether the target is inside the IR sensor's FOV.

[0095] Using two or more PIRs or IR thermopiles would further increase the information content delivered to the room occupancy detection algorithm **58**. This is depicted as a multi-zone IR sensor **104** in FIG. **14**. The relative timing of the peaks of these two signals can be used to infer a direction of travel for the person wearing or using the tag device. If only one peak is seen, this would suggest that the person that caused the IR activity remained either inside or outside the room and did not enter or exit. The two-zone IR sensor scenario is illustrated in FIGS. **15A** and **15B**. In this example, one of the two IR sensors in multi-zone IR sensor **104** points inside the room from the monitor while the other points outside. A person **2** wearing a tag device **3** is

about to enter the room **4**. The plot in FIG. **15B** shows a pulse of IR activity at 0.75 seconds as the person walks through outside zone **110**, and then another pulse at time 2.2 seconds when the person walks through the inside zone **111**.

[0096] Referring back to FIG. **14**, the highest performance realization of a multi-zone IR sensor is an IR array **105**, which could be implemented using either a digital camera with an IR lens and filter, or a thermopile array chip. The use of the array effectively allows the room occupancy detection algorithm **58** to process a sequence of digital thermal images, sampled concurrently with the received array response and RSSI information received from the transceiver **13**.

[0097] Another alternative sensor that could be deployed on the monitor **1** is a digital camera **106**. The digital camera **106** could be used to periodically digitize images in its field of view and provide the digitized images into room occupancy detection algorithm **58** along with the antenna array outputs (I/Q samples **51** and RSSI samples **52**). The images could be digitized using either a red-green-blue (RGB) color or black-and-white encoding scheme.

[0098] As an alternative to IR-based sensors, laser, radar or ultrasound-based range sensors **107** could be used. One or more of these range sensors **107** could be installed on the monitor **1** and configured to periodically report on the measured distance between the range sensor and the nearest intervening object or person. The distance measurements from either of these sensors could be provided as input into the room occupancy detection algorithm **58** along with the antenna array outputs.

[0099] Alternatively, a multi-zone range sensor **107a** could be used. The multi-zone range sensor **107a** could be configured to report on the measured distance between itself and the nearest intervening object or person within a plurality of angular zones around the sensor **107a**. For example, a 4×4 or 8×8 angular grid of zones could be used.

[0100] The additional sensors could also be deployed on the tag device **3**. A low-cost micro electromechanical system (MEMS) inertial motion unit (IMU) **108** containing a three-axis accelerometer, compass or gyroscope sensor could be installed on the tag device. The tag device could then be configured to periodically digitize the IMU sensor outputs and transmit them over-the-air to the monitor **1**. FIG. **14** shows the IMU **108** in dotted line to indicate that the IMU is actually not resident on the monitor **1**, but instead is part of the tag device. The monitor **1** could decode the IMU sensor readings **109** and feed them into the room occupancy detection algorithm **58** along with the antenna array outputs. If the Bluetooth™ 5.1 wireless protocol is used, the encoded IMU sensor readings could be included in the same packets containing the CTE that are already being transmitted by the tag.

[0101] If the IMU **108** on the tag has a three-axis compass, for that compass to be useful to the monitor **1**, the monitor would need to know its own orientation relative to the earth's magnetic north. This could be done as a calibration step after the monitor **1** is first installed, or by installing a 3-axis compass or IMU **108** on the monitor itself. In the latter case, the monitor could measure its own orientation relative to an earth-centered earth-fixed (ECEF) coordinate system, removing the need for a calibration step. Once the monitor knows its own orientation relative to ECEF, the compass readings from the tag could be transformed from ECEF bearings to monitor-based bearings—i.e., directions of travel on the monitor-centric coordinate system describe earlier-before being fed into room occupancy detection algorithm **58**.

[0102] One additional benefit of having an IMU **108** installed on the tag **3** is that the accelerometer on the IMU could be used as a motion sensor to conserve tag battery life. For example, the tag device could save battery by disabling its transmitter and entering a low power sleep state when no motion has been detected for the past one minute. The tag device could immediately begin transmitting again once motion is detected.

[0103] If the room occupancy detection algorithm **58** is a machine learning algorithm, adding one or more additional sensor signals as input to the algorithm is conceptually trivial; one would include the additional sensors signals with the antenna array outputs **51** and **52** and re-run the data gathering and training procedure **70** using all the input signals. If the room occupancy detection algorithm **58** is a non-machine learning algorithm, logic and thresholds would be added to the algorithm to determine the tag's occupancy status using the additional inputs.

[0104] The monitor **1** can be used for other applications in addition to room occupancy detection. For example, the monitor **1** could be used to detect when a wireless tag or badge device **3** enters or exits a multi-zone room containing multiple hospital beds. The monitor **1** could also be used in a single or multi-zone room to determine when a caregiver **2** wearing a wireless tag or badge is at a patient's bedside. Yet another application that will be discussed below is hand hygiene monitoring.

[0105] Reference is now made to FIG. **16**, which shows how the monitor **1** could be used in a multi-zone room application. A room **120**, such as a hospital intensive care unit (ICU) or emergency room (ER), containing multiple zones **121** to accommodate a plurality of patients is shown. In rooms such as these, multiple monitors **1** could be used to determine within which zone **121** a caregiver **2** or patient wearing a tag device (badge) **3** as an emitter is, at a given time. The monitors **1** may be installed on the ceiling with their antenna arrays pointing straight down, and oriented so that their X and Y axes are parallel to zone boundaries **122**. Each monitor may use a PIR with an omni-directional Fresnel lens as its motion sensor to detect motion in any of the monitored zones **121**, or near the boundaries **122** between monitored zones. Each time the PIR wakes up the monitor, the processor in the monitor **1** could wake up the transceiver of the monitor to listen for wireless transmissions to determine if any of the monitored emitters **3** has entered or exited any of the zones **121**.

[0106] To determine when an emitter **3** has entered or exited a zone **121**, an algorithm similar to the room occupancy detection algorithm **58**, described above, may be used. If a machine learning algorithm is desired, the procedure **130** shown in FIG. **17** may be used to train the algorithm. Reference is now made to FIG. **17**, with continued reference to FIG. **16**. In step **131**, monitors **1** are installed on the ceiling at the midpoint between two or more zones spanning one or more rooms in an area that resembles the environment where the system is intended to be used once in production. Each monitor should be configured to operate just as it would in practice; in particular, the monitor should only wake up from sleep and begin receiving and storing data after it has been woken up by its motion sensor.

[0107] In step **132**, each of the monitors **1** is configured to receive, downconvert and decode wireless transmissions from one or more tag devices **3** as people **2** wearing, carrying, or moving the tag devices walk into, out of, around the outside and around the inside of the zones **121** in all the monitored rooms **4**. The monitors are also configured to generate and record receive signals produced by the one or more monitors for each of the detected wireless transmissions. The receive signals could include the

received array response vectors and RSSI estimates for each received transmission from the tag devices **3**. The receive signals could also include the results of operations applied to the array response vectors or RSSI information, such as beamforming or MUSC spatial processing. The operations could also include the results of other operations, such as resampling and data conditioning.

[0108] In step **133**, ground truth information containing the tag device ID, zone ID, monitor ID, entry time and exit time every time a tag or badge device **3** enters or exits a zone is recorded and stored. The ground truth could be obtained by recording a video of the people wearing the tags walking in and out of the monitored rooms and zones in step **132** and later playing back the video to determine with precision when each tag entered or exited each monitored room or zone.

[0109] In step **134**, the recorded receive signals and ground truth information are stored in a training database.

[0110] In step **135**, steps **131-134** are repeated at numerous different locations having different ceiling heights, floor layouts and construction materials. The people used to wear, carry or move the tag or badge devices **3** could enter or exit the rooms or zones at different walking speeds and entry/exit angles, and change the way they wear, carry or move the tag or badge to emulate the way the tag device will be used in practice as much as possible.

[0111] In step **136**, the training procedure in steps **131-135** is terminated when it is determined that a sufficient amount of training data has been taken. This determination is usually made by partitioning the receive signals and ground truth information stored in the database into a training data set and a test data set. There is enough training data when adequate performance is attained on the data set, and the machine learning algorithm performs adequately on the test data set and on data taken from new rooms, zones or buildings. If the algorithm does not perform as well on newly obtained data from a new room, zone, person, tag or environment as it typically does in the test set, this could mean that the algorithm should be trained on that new room, zone or environment.

[0112] In step **137**, once it is determined that a sufficient amount of training data has been obtained and stored in the training database, that data is used to train the machine learning algorithm that runs on the monitor.

[0113] Reference is now made to FIG. **18**, which shows a flow chart **140** depicting a non-machine learning approach for zone-based occupancy detection. In step **141**, the monitor **1** remains asleep in a low current state until it receives a wake-up signal from its motion sensor. In step **142**, once the processor of the monitor wakes up, it wakes up its RF/Baseband transceiver, configuring it to pass the array response and RSSI information received from any tag transmissions up to the processor, and for each array response and RSSI received, estimates the X,Y position of the tag using the monitor-centric rectangular coordinate system defined earlier. The X, Y position of the tag could be estimated by evaluating the beamforming or MUSIC cost functions (5) or (6) over a grid of candidate target (X,Y,Z) positions selecting the most likely grid position by finding the minimizer of (5) or (6) over the grid, and noting the X, Y position of the minimum-cost grid point. To minimize the number of cost function evaluations, the Z component of the tag position could be held at a constant value—e.g., 5 feet, to model a typical height for most adults.

[0114] In step **143**, the algorithm tests whether the tag has already been detected as being currently in any of the monitored zones.

[0115] In step **144**, if the tag has not been detected as being in any zone, the algorithm looks to see if the X, Y position of the tag is sufficiently close to any of the zone center positions for some minimum period of time—T1 seconds. In step **146**, if the test condition of step **144** is true, the algorithm declares the tag to be inside of the zone at the closest distance to the estimated X, Y position of the tag.

[0116] In step **145**, if the tag has been detected as being in some zone, the algorithm looks to see if the estimated X, Y position of the tag is too far from the center position of the currently detected zone for some minimum period of time—T2 seconds.

[0117] In step **147**, if the test condition of step **145** is true, the algorithm declares the tag to be not inside any of the zones.

[0118] In step **148**, the algorithm tests to see if there has been any motion detected by the motion sensor **6** over some period of time—T3 seconds. If there has not been any motion over the past T3 seconds, the algorithm puts the monitor into a low current sleep state and returns to step **141**. Otherwise, if there has been motion detected in the past T3 seconds, the monitor returns to step **142**.

[0119] Reference is now made to FIG. **19**, which shows how a monitor **1** could be used for a hand hygiene compliance application. For this application, in addition to being used to detect room or zone entries in a hospital or clinic, the monitor **1** is also installed near (Bluetooth™-enabled) disinfectant dispensers **150** and used to detect when a healthcare provider **2** wearing tag device/badge **3** uses the dispenser **150** to administer disinfectant on their hands.

[0120] The monitor **1** could be installed either on the ceiling **151** directly above the place **152** where a person would stand when using the disinfectant dispenser **150**, on the wall just above or alongside **153** the dispenser **150** or integrated into the dispenser in some way (the latter is not shown in FIG. **19**). An algorithm running on processor of the monitor **1** could be used to determine when a person **2** wearing a Bluetooth™ 5.1 enabled tag device/badge **3** is within a certain distance of the dispenser **150**. The dispenser **150** itself could be equipped to send a message to the monitor **1** whenever a dispense has been given. The monitor **1** could note the badge ID of the closest badge to the monitor at the time of the dispense to determine the identity of the person **2** that disinfected their hands.

[0121] If the monitor is installed on the ceiling **151**, a motion sensor of the monitor with a narrow FOV aiming directly downward from the monitor could be used to wake the monitor from sleep. If the monitor is installed on the wall next to the dispenser, a PIR could still be used to wake up the monitor, a motion sensor with a very short range (e.g., 1-3 feet) may be used to prevent false wakeups caused by people moving more than 3 feet away from the dispenser who are not using it.

[0122] For this application, the goal of the software algorithm running on the monitor's processor is to determine with confidence the Bluetooth™ 5.1 badge ID worn by a user **2** of the dispenser **150**. If a machine learning algorithm is desired, a procedure **160** for training such an algorithm to determine the user's badge ID after a dispense, may be employed. The procedure **160** is now described with reference to FIG. **20**, and continued reference to FIG. **19**.

[0123] In step **161**, the monitor's installation position is selected among options (a), (b) or (c) below: (a) on the ceiling **151**,

directly above the place **153** a person would stand when using the dispenser **150**, with the antenna array **10** center positioned X1 inches in front of the front midpoint of the dispenser **150** for some known parameter X1 that is held constant throughout the entire training procedure **160**, (b) on the wall with the antenna array **10** center positioned X2 inches above the floor and X3 inches to the left of the top-center point of the dispenser for some known parameters X2 and X3 that are held constant throughout the entire training procedure **160**, wherein using a negative value of X3 means to the right of the dispenser, (c) integrated inside the dispenser.

[0124] In step **162**, the monitor **1** is installed at the position selected in step **161**.

[0125] In step **163**, the monitor **1** is configured to receive, downconvert and decode wireless transmissions from one or more wireless tag devices/badges **3** as people **2** wearing the badges walk around and occasionally use the dispenser **150**. The monitor **1** is also configured to generate and record receive signals associated with each of the detected wireless transmissions. The receive signals could include the received array response vectors and RSSI estimates for each received transmission from the badges. The receive signals could also include the results of operations applied to the array response vectors or RSSI information, such as beamforming or MUSC spatial processing. The operations could also include the results of other operations, such as resampling and data conditioning.

[0126] In step **164**, ground truth information containing the badge ID of the user **2** and the dispense time each time the dispenser was used is recorded and stored. The ground truth information could be obtained by recording a video of the people using the dispenser in step **163** and later playing back the video to determine with precision when each dispense occurred and what user's badge was associated with each dispense.

[0127] In step **165**, the recorded receive signals and ground truth information are stored in a training database.

[0128] In step **166**, steps **162-165** are repeated at numerous different dispenser locations having different ceiling heights, dispenser heights, surrounding wall configurations, and construction materials. The people wearing the badges should attempt to approach and use the dispenser with a broad range of walking speeds, distances to the dispenser, walking angles and people grouping arrangements.

[0129] In step **167**, the training procedure in steps **161-166** is terminated when it is determined that a sufficient amount of training data has been taken. This determination is usually made by partitioning the receive signals and ground truth information stored in the database into a training data set and a test data set. There is enough training data when adequate performance is attained on the data set, and the machine learning algorithm performs adequately on the test data set and on data taken from new dispenser locations. If the algorithm does not perform as well on newly obtained data from a new dispenser location or use case as it typically location does in the test set, this could mean that the algorithm should be trained on that new location or use case.

[0130] In step **168**, once it is determined that a sufficient amount of training data has been obtained and stored in the training database, that data is used to train the machine learning algorithm.

[0131] Reference is now made to FIG. **21**, which shows a flow chart depicting a non-machine learning procedure **170** for determining the badge ID of a disinfectant dispenser user. Reference is also made to FIG. **19** for purposes of the description of FIG. **21**. In step **171**, the monitor **1** remains asleep in a low current state until it receives a wake-up signal from its motion sensor. In step **172**, once the processor of the monitor wakes up, it wakes up the RF/Baseband transceiver in the monitor in, configuring it to pass the array response and RSSI information received from any badge transmissions up to the processor, and for each array response received, estimates the X, Y, Z position of the badge using the monitor-centric rectangular coordinate system defined earlier herein. The X, Y, Z position of the tag could be estimated by evaluating the beamforming or MUSIC cost functions (5) or (6) over a grid of candidate target (X,Y,Z) positions selecting the most likely grid position by finding the minimizer of (5) or (6) over the grid, and noting the X, Y, Z position of the minimum-cost grid point.

[0132] In step **173**, the algorithm tests to see whether a disinfectant dispense event has been triggered in the disinfectant dispenser **150**, indicating that disinfectant was administered at the dispenser **150**. The dispense signal could be received from the disinfectant dispenser **150** via either a hard-wired connection between the dispenser **150** and the monitor **1**, or a Bluetooth wireless message sent from the dispenser **150** to the monitor **1**. If such a dispense trigger signal was not received by the monitor **1**, control proceeds to step **174**. If such a dispense trigger signal was received by the monitor **1**, then control proceeds to step **175**.

[0133] In step **174**, the processor of the monitor **1** determines how much time has transpired since motion was last detected by the motion sensor of the monitor. If the time since motion was last detected exceeds a timeout period, the processor puts the monitor to sleep by returning to **171**. Otherwise, control returns to step **172**.

[0134] In step **175**, the processor of the monitor determines which badge is closest to the monitor **1** at the time the dispense signal was received by looking for the tag which has an (X,Y,Z) position that is closest in Euclidean distance to the (X, Y,Z) position of the monitor **1**, and sends a message (e.g., to the server **43** via gateways **41** as depicted in FIG. **7**) containing the badge ID of that badge. If two or more badges are determined to be close to the monitor **1** at the time of the dispense, the processor may send a special message to the server **43** via the gateways **41** indicating that this was the case. This would give the server the option of removing this dispense from consideration for hand hygiene monitoring, since there is ambiguity in which user used the dispenser.

[0135] It should be noted any of the additional sensors shown in FIG. **14** that could be used as additional inputs to improve the performance of the room occupancy detection algorithm **58**, could also be used in much the same way to improve the performance of the multi-zone room occupancy detection and disinfectant dispenser badge ID detection algorithms. For the multi-zone room application, a multi-zone digital or analog PIR signal could be used to detect motion in each room zone and provided as input into the room zone detection algorithm along with the antenna array outputs. The same is true of an IR array **105** or a digital camera **106**, multi-zone range sensor **107a** or an IMU **108** on the tag device **3**. For the disinfectant dispenser application, the digital or analog PIR sensor, IR array **105**, multi-zone range sensor **107a** or digital camera **106**, or IMU on the tag device **3** could be used to help detect the presence of the dispenser user as well. For the dispenser application, the dispenser output signal itself (signaling when the disinfectant was administered from the dispenser) could be provided as an additional

input to the badge ID detection algorithm.

[0136] Some variants of the Room Occupancy detection algorithm **58** can be quite computationally expensive. For example, some large neural networks are known to require on the order of 50 MFLOPS to process the received array response information from a single tag device. In an operational setting such as a hospital nursing unit, it is reasonable to have a single monitor to service up to 50 tag devices at the same time. If the same 50 MFLOPS per tag neural network is used, then $50 \times 50 = 2.5$ GFLOPS are required to support all 50 tags. This amount of processing throughput will have a significantly detrimental impact on the monitor's manufacturing cost and battery life.

[0137] Reference is now made to FIG. **22**, which shows a computationally efficient winnowing procedure **180** for identifying a subset of wireless tag devices which are reasonably likely to enter or exit a monitored room or zone. The winnowing procedure **180** can be shown to require much less than 50 MFLOPS per tag to identify likely entry or exit candidates. It is reasonable to expect winnowing procedure **180** to require 0.1 MFLOPS or less per tag. For example, if the winnowing procedure **180** runs on a monitor **1** that services 50 tag devices and can winnow from 50 to 5 the tags on which to run a 50 MFLOPS per tag neural network, this reduces the peak computational load from 2.5 GFLOPS to $5 \times 50 + 45 \times 0.1 = 0.26$ GFLOPS.

[0138] In step **181** of procedure **180**, the monitor **1** remains asleep in a low current state until it receives a wake-up signal from its motion sensor.

[0139] In step **182**, once the processor of the monitor wakes up, it wakes up the RF/Baseband transceiver, configuring it to pass the array response and RSSI information received from any tag transmissions up to the processor, and waits for an array response vector and RSSI measurement to be received.

[0140] In step **183**, the processor of the monitor estimates the X, Y position of the tag using the received array response vector and RSSI measurement. The X, Y position is calculated using the monitor-centric rectangular coordinate system defined earlier herein. The X, Y position of the tag could be estimated by evaluating the beamforming or MUSIC cost functions (5) or (6) over a grid of candidate target (X,Y,Z) positions selecting the most likely grid position by finding the minimizer of (5) or (6) over the grid, and noting the X, Y position of the minimum-cost grid point. To minimize the number of cost function evaluations over the (X,Y,Z) grid, the Z component of the tag position could be held at a constant value—e.g., 5 feet over this grid, to model a typical height for most adults.

[0141] In step **184**, the processor tests to see if the X position signal for the tag device has changed sign while the absolute value of the Y position estimate is smaller than a threshold of T1 feet. If the test condition is true, then at **185**, the processor puts the tag device on a “Run” list.

[0142] In step **186**, the processor tests to see if the tag is on the Run list. If it is, then at **187**, the room occupancy algorithm **58** is run on the current array response signal and all the array response information received from this tag device since the sleep wakeup of step **182**.

[0143] In step **188**, the processor tests to see if more than T2 seconds have transpired since motion was last detected via the motion sensor **6**, where T2 is an appropriate timeout duration—for example, 5 seconds. If so, the processor removes all tag devices from the Run list and goes to sleep. Otherwise, the processor proceeds back to step **182**.

[0144] The concepts described above herein describe a system in which the tag device **3** is configured as a wireless emitter, the monitor device is configured as a receiver wherein the monitor device **1** receives transmissions from the emitter on multiple receive antennas and uses observed phase differences on the receive antennas to determine the position of the tag device using angle-of-arrival techniques. It should be noted that because of the laws of physics regarding antenna reciprocity, any of these approaches could instead be implemented using the monitor **1** as the emitter, the tag device as the receiver, wherein the monitor sends the transmit signal on multiple antennas (in sequence), and wherein the tag receives the transmissions on a single antenna, computes phase differences from the receive signals, and uses the computed phase differences to determine the tag's position relative to the monitor. In the latter case, the room occupancy detection algorithm **58** would run on the tag instead of the monitor since the tag is where the phase differences are computed and readily available. Such a room occupancy detection system would include one or more room beacons/transmitters configured to send a transmit signal (on a repetitive continuous basis, for example) from multiple antennas (in a sequential manner via one of the antennas at a time), and one or more tag devices each configured to receive the transmissions (on a single antenna) and compute phase differences from receive signals derived from receiving the transmissions from the room beacon(s). The tag device uses the computed phase differences to determine its position relative to the room beacon(s)/transmitter(s). Again, the room occupancy detection algorithm would run on the tag devices instead of on the room monitor. A server may be provided that is configured to receive the room occupancy detection events from the one or more tag devices. In this system, each of the one or more room occupancy beacons/transmitters includes an antenna array or a plurality of antennas from which to transmit signals; a wireless transmitter configured to generate transmissions to be transmitted via the plurality of antennas (in a sequential manner via one of the antennas at a time). Each tag device may include a wireless receiver to produce receive signals from reception of the transmissions from a room beacon/transmitter, a processor configured to process the receive signals and to run an algorithm on a sequence of estimates of a parameter, derived from the received signals from the tag, that is indicative of a distance between the tag and monitor, and array response vectors derived from the receive signals, and determine when the tag device has entered or exited the room via the entryway.

Operational Efficiencies Based on Room Entryway Orientation

[0145] Reference is now made to FIG. **23**, which shows an example scenario in which monitors are installed in left and right opening rooms. It is well known that most rooms in an office or healthcare environment have an entryway near either side of the room, but rarely in the middle of the room. Nearly half of these rooms open to the right as shown in right-opening room **191** when facing into the room from outside, the other nearly half open to the left as shown in left-opening room **197**, and a much smaller portion of the rooms have the entryway at or near the midpoint between the side walls of the room. The set of valid entry paths **192** for people walking into the room **191** that opens to the right excludes paths that go to the left because the intervening wall **193** blocks travel in that direction. And the valid entry paths **196** into the room **197** that opens to the left exclude paths that

go to the right for the same reason. Statistically speaking, the entry paths **192**, **196** for rooms that open to the right and left tend to be mirror images of one another, with the plane of the wall **193** being their axis of symmetry.

[0146] The fact that the entryway paths vary depending on the room entryway orientation poses both an opportunity and a logistical challenge for the room occupancy detection algorithm **58**. When installing the room occupancy monitor, the installer could specify as a system configuration parameter whether the room opens to the left, right, or middle, and this selection can be presented as input to the room occupancy detection algorithm **58** to improve its performance by taking advantage of the known differences among their walking path statistics. The logistical challenge comes from the additional data gathering to develop (if a non-ML algorithm is used) or train (if a ML algorithm is used) and test the room occupancy detection algorithm **58** for the three different room types, since training and testing the algorithm on the three different room types requires approximately three times more data than for a single room type.

[0147] It is possible to avoid the need to gather extra data by exploiting the inherent symmetries among the three different room entryway orientations. First, assuming a noise-free and multipath-free RF environment, it can be shown that the array response vector z_A received by a monitor **1A** installed in a right-opening room **191** from an emitter positioned at a point (X, Y, Z) **194** relative to that monitor is identical to the response vector z_B received by a monitor **1B** installed in a left-opening room **197** from an emitter positioned at the mirror image point $(X, -Y, Z)$ **195** relative to monitor **1B**—provided that the antennas in monitor **1B** are re-numbered so as to make them function as if they were transposed about the monitor's X axis relative to the monitor **1A** numbering. The antennas on the right-opening monitor **1A** are labeled **10A-1**, **10A-2**, . . . , **10A-16** in FIG. **23**, those on the left-facing monitor **1B** are labeled—after transposing—**10B-1**, **10B-2**, . . . , **10B-16**.

[0148] The array transpose operation can be implemented in software by exchanging any readings taken from antennas on one side of the X axis with the readings from the antennas at the mirror image positions on the other side of the array's X axis. So, given an array response vector z with component readings $\{z_1, z_2, \dots, z_{16}\}$ taken from antennas **10A-1**, **10A-2**, . . . , **10A-16**, respectively, the array transpose operation is implemented as follows: [0149] components z_1, z_5, z_9 and z_{13} are exchanged with z_4, z_8, z_{12} , and z_{16} ; [0150] components z_2, z_6, z_{10} and z_{14} are exchanged with z_3, z_7, z_{11} , and z_{15} .

[0151] To summarize, the vector $\{z_1, z_2, \dots, z_{16}\}$ would map to $\{z_4, z_3, z_2, z_1, z_8, z_7, z_6, z_5, z_{12}, z_{11}, z_{10}, z_9, z_{16}, z_{15}, z_{14}, z_{13}\}$ after transposing.

[0152] Thus, it is possible to make array response data obtained from any left (right) opening room appear as if it was obtained from a right (left) opening room by re-arranging the array response vectors from the left (right) facing room using the array transpose operation described above. By the time the data gets to room occupancy detection algorithm **58**, it will appear as if all rooms opened in the same direction. This approach can be applied to not only rectangular antenna arrays, but any array that has symmetry about the X axis **30**.

[0153] Referring back to FIGS. **3** and **12**, the data gathering, training and test procedure **70** can be modified as follows so that only data from left or right opening rooms is used for training the room occupancy detection algorithm **58** to detect room occupancy for left, right or middle opening rooms, without any requirement being placed on the relative proportions of left, right or middle opening rooms.

[0154] In step **74**, before storing the receive signals in the database, the array transpose operation may be applied to the receive signals taken from left (or right) opening rooms to make it appear to the room occupancy detection algorithm **58** during training or testing as if all rooms opened to the right (or left).

[0155] An alternative to applying the array transpose operation before storing the receive signals as described above would be to include a left-right indication—an indication as to whether each room opens to the left or right—in the ground truth information stored in step **74**, and transpose the receive signals before presenting them to the room occupancy detection algorithm **58** during training step **77** or testing based on the left-right indication.

[0156] In step **75**, the installation, data gathering ground truth generation and storage steps **71-74** can be repeated for any number of rooms that either open to the left or to the right, without regard to the relative numbers of left or right opening rooms. Middle-opening rooms, however, should not be used as part of this procedure per se.

[0157] Since the algorithm running on the monitors would be trained to detect entries into left (or right) opening rooms but not both, the array transpose operation should be used by processor **15** to transpose the array response vectors before presenting them to the room occupancy detection algorithm **58** whenever the monitor is installed in a right (or left) opening room.

[0158] To train the room occupancy detection algorithm **58** to detect entries or exits into middle-opening rooms, step **77** may be modified to randomly transpose the receive signals to simulate a left-opening room half of the time, and a right-opening room the other half of the time. For middle opening rooms, the monitor should not be configured to transpose the array response vectors.

Improved Distance Estimates

[0159] In the description above herein, the array response vectors derived from signals received at the monitor from the tag device provide information about the tag's angular position relative to the monitor, and the received signal strength estimates provide an indication of the tag's distance to the monitor.

[0160] In an indoor environment such as a hospital, received signal strength readings can yield poor distance estimates. This is because when a wireless signal is transmitted indoors, the receiver receives not just one but many reflected copies of the transmitted signal arriving at its antenna, with each copy having a different amplitude, delay and phase-shift relative to the shortest-distance path. When these signals combine at the receiver's antenna, the resulting constructive or destructive interference causes large fluctuations of up to 40 dB in the received signal strength when the transmitter or receiver changes position by more than a few inches at 2.4 GHz. One well-known way to reduce the fluctuations is to average the received signal strength over multiple receive antennas, but even then, received signal strength is still known to be a poor proxy for distance indoors.

[0161] Some wireless standards such as UWB (IEEE 802.15.4z) and Bluetooth 6 support the use of round-trip-time (RTT) measurements to improve distance accuracy. If, for example, the tag and monitor communicated over-the-air using IEEE

802.15.4z, the improved distance measurements obtained from RTT could be fed into the room occupancy detection algorithm **58** along with the array response vectors, received signal strength measurements and other sensor outputs and used to improve the performance of the algorithm.

[0162] To measure RTT, a first device sends a wireless signal to a second device. After receiving the wireless signal, the second device waits a known fixed amount of time T , then sends a second signal back to the first device. After receiving the second signal, the first device calculates the RTT by subtracting the time it transmitted the first signal and the response time T from the time it received the second signal. The distance between the two devices can be estimated by dividing the measured RTT by twice the speed of light.

[0163] Devices using IEEE 802.15.4z perform RTT measurements by exchanging UWB messages and performing either a single-sided two-way ranging (SS-TWR) measurement or a double-sided two-way ranging measurement (DS-TWR). For SS-TWR, one device estimates the distance between the devices using RTT measurements; for DS-TWR, both devices perform the distance estimate.

[0164] Devices using Bluetooth 6 perform RTT measurements using a channel-sounding measurement, whereby an initiator device and a reflector device exchange information across 72 RF channels, each 1 MHz wide, with the initiator transmitting first, followed by the reflector's response on each RF channel. Bluetooth 6 allows two different techniques for the initiator and reflector to exchange information. For indoor environments, the more accurate technique is known as phase-based-ranging (PBR), in which the two devices measure the phase shift of the round-trip path between the two devices at each of the 72 different frequencies. Without multipath, the inverse Fourier transform of the 72 MHz wide PBR response would yield an impulse with a time delay equal to the round-trip propagation time between the two devices. In practice, even when there is multipath, the initiator device can estimate the impulse time delay by processing the PBR channel sounding data using any one of a set of well-researched methods, such as MUSIC, ESPRIT or Maximum-Likelihood.

[0165] If the two devices are synchronized, i.e., if they both have access to a common clock (e.g., via GPS), their distance can be estimated using a single one-way time-of-flight (ToF) measurement instead of RTT. For example, the first device could send a signal to the second device with the transmit time embedded in the signal itself. After receiving the signal and noting its time-of-arrival (ToA) with respect to the common clock, the second device could demodulate and decode the transmit time from the received signal and subtract it from the ToA to derive the one-way ToF. The distance between the two devices can then be estimated by dividing the ToA by the speed of light.

UWB Room Occupancy Monitor

[0166] Aside from supporting RTT distance measurements, an embodiment in which the tags and monitors communicate using UWB wireless signals would have several advantages over a non-UWB embodiment, such as with Bluetooth. Since UWB signals span at least 500 MHz of RF bandwidth, these signals are significantly less susceptible to multipath reflections off walls, floors, ceilings, and other objects indoors. The increased bandwidth allows the receiver to distinguish delayed reflections of the wireless transmissions from the shortest-distance-path. With non-UWB signals, the delayed reflections can be much more difficult to distinguish from the shortest distance path, giving UWB a significant advantage in terms of RTT accuracy.

[0167] The inherent multipath immunity of UWB over other wireless standards would allow a UWB-based monitor to use a simpler-likely non-ML-room occupancy detection algorithm, yielding improved detection accuracy, (possibly) improved battery life and (possibly) reduced manufacturing cost. The improved multipath immunity would also allow the monitor to use fewer receive antennas in the monitor (as few as 2, or ideally 3, as opposed to up to 16 for Bluetooth 5), thereby reducing the physical size of the monitor and further reducing its manufacturing cost. The improved multipath immunity could also allow the tag to transmit at a lower rate (e.g., 1-2 Hz) than for a non-UWB embodiment (up to 10 Hz), further improving the tag battery life. Some advantages of UWB embodiments are lower chipset cost, lower chipset power consumption, and improved reliability at range, since UWB involves very low radiated transmit power.

[0168] In FIG. 25, a schematic block diagram of a UWB-based room occupancy monitor **1'** is shown. The block diagram for UWB-based room occupancy monitor **1'** is nearly identical to that shown in FIG. 3 but is more tailored to a UWB implementation. In a UWB implementation, the RF switch **12** in FIG. 3 is no longer necessary since the receiver would be receiving and downconverting the signal simultaneously from multiple antenna paths. The switch **12** is more appropriate for a Bluetooth 5 implementation which uses antenna switching. After the signal from each antenna path in the antenna array **10** is down-converted, it is digitized and processed in the UWB RF transceiver/baseband (BB) processor **250**, which uses a correlator to identify the minimum-distance path and separate the minimum-distance path from other delayed, reflected paths, and stores the carrier phase at the minimum-distance path correlator output for each of its N antennas. The stored carrier phases from each antenna path are then used to determine the angular position of the tag. As discussed earlier, the distance between the tag and monitor is determined using RTT measurements. Also mentioned earlier, the number of antennas N in a UWB monitor could be significantly fewer than in a non-UWB implementation, due to the inherent multipath immunity of UWB.

[0169] A non-UWB Gateway Interface Transceiver **251** is used to communicate with the wireless gateway **42** (FIG. 7) through omni-directional antenna **11**. This transceiver will send Entry and Exit events to the wireless gateway **42** using non-UWB wireless signals such as Bluetooth to overcome the range limitations of UWB.

[0170] Reference is now made to FIG. 26, which shows a room occupancy detection procedure **2600** to be implemented by the processor **15** in the UWB monitor **1'** depicted in FIG. 25, in accordance with one embodiment. In this procedure, it is assumed the tag devices **3** periodically (e.g., twice a second) send UWB beacons over the air and listen for incoming RTT measurement requests from the monitor(s) **1'**.

[0171] In step **2605**, the monitor **1'** sleeps in a low power state, awaiting a wakeup event from the motion sensor **6**.

[0172] In step **2610**, after motion is detected, the processor wakes up the UWB transceiver/BB processor **250** and listens for incoming UWB beacons transmitted from one or more nearby tags.

[0173] In step **2615**, for each UWB beacon received from a tag, the processor sends a RTT measurement request to the tag and measures the round-trip-time of the subsequently received RTT response message, dividing the RTT by twice the speed of light

to estimate the distance from the tag to the monitor. The processor also computes the phase-difference-of-arrival (PDOA) of the incoming message on its two or more antennas to estimate the angular position of the tag relative to the monitor **1'**.

[0174] In step **2620**, the processor estimates the position of the tag(s) relative to position and orientation of the monitor **1'** using the distance and angular position estimates computed in the previous step. and using the position estimates, determines whether the tag(s) have entered or exited the monitored room.

[0175] In step **2625**, the processor sends room occupancy detection events to the server **43, 45** via the gateways **42** based on the detected entries or exits.

[0176] In step **2630**, the processor puts the monitor back to sleep if a sufficient time period has transpired since motion was last detected via motion sensor **6**, that is, the motion sensor has been idle for a time period. Otherwise, the procedure returns to step **2610** above.

Self-Supervised ML Training

[0177] The behavior of the received signal strength and array response signals received by the monitor as people wearing tag devices **3** walk in, out and around a monitored room can be shown to vary significantly from one room to the next. This room-to-room variation is due to the physical differences that invariably exist between rooms. The physical dimensions of the room and entry area, construction materials, and objects placed in and around the room affect the RF characteristics of the room. When UWB wireless signals are used, most of the reflections can be ignored because they can be differentiated from the shortest-distance path by cross-correlating the received signal with a known transmit waveform. For non-UWB signals such as Bluetooth, it is impossible to separate the reflections—they instead appear as a single reflected path that is “smeared/spread out over time”, thereby degrading RTT distance estimates and causing AoA spatial responses to appear as if the signal from the tag is coming from multiple different directions. The multi-directional nature of the AoA algorithms with narrowband signals makes using a non-ML algorithm much more difficult.

[0178] In such embodiments where non-UWB signals and a ML room occupancy detection algorithm are used, for the system to work reliably at scale, an identical copy of the firmware containing the ML weights and parameters should be deployed on each monitor. An alternative would be to train a different version of the ML algorithm on data taken from each room, but the data gathering and labeling effort required to do that would be highly impractical and too costly. If an identical copy of firmware is to be used, the ML algorithm is to be trained on data taken from many rooms of different shapes, sizes and physical characteristics. The downside of this “one size fits all” approach is that once the monitor is installed in a specific room, it will not perform as well as it would have if trained on data taken from that room only.

[0179] Reference is now made to FIG. **27**, which shows a fully automated, self-supervised learning procedure **2700** that can be used to overcome the “one size fits all” performance limitation. The basic idea is to ship all monitors from the factory with the “one size fits all” algorithm that performs reasonably well in any room, and then once a monitor has been installed, configure the monitor to incrementally update the parameters of the ML algorithm using self-supervised learning to optimize its performance for that specific room over time.

[0180] In step **2705** of the procedure, the monitor is sleeping, waiting for a person to trigger the motion sensor.

[0181] In step **2710**, once the motion sensor has triggered, the monitor begins receiving input data in the form of array response vectors, RSSI and/or distance measurements and optionally data from other sensors while one or more tags enters, exits or walks near the monitored room.

[0182] In step **2715**, the input data is processed through the ML room occupancy detection algorithm **58**, producing output predictions. Each output prediction is a 3-tuple containing the probabilities that a room entry, exit or neither has just occurred. Each output prediction 3-tuple is then mapped to a pseudo-label—an integer in the set {0, 1, 2}—by computing the index of the max probability prediction in the 3-tuple.

[0183] In step **2720**, the pseudo-labels are stored along with the input data as pseudo-labeled data segments. A counter is incremented to indicate whether enough pseudo-labeled data has been obtained for a training batch. If not, control returns to step **2705**.

[0184] In step **2725**, once a full training batch of pseudo-labeled data segments has been obtained, the cross-entropy loss between the output predictions and pseudo-labels is computed over the entire training batch, and the weights of the ML algorithm **58** are updated using a gradient descent calculation with backpropagation. In general, a smaller learning rate should be used for the gradient descent calculation than was used for the “one size fits all” algorithm.

[0185] In step **2730**, a count of the number of completed training batches is incremented. If the number exceeds an appropriate threshold such as 5,000 or 10,000, training stops.

[0186] Any one of the following variations could be added to the procedure **2700** to potentially improve its performance.

[0187] First, the procedure involves updating all ML parameters at the same time, which could consume a large block of random access memory (RAM) on the monitor. To save RAM, one could partition the ML parameters into subsets (e.g., based on Neural Net layer), and update the parameters one layer at a time in round-robin fashion, e.g., batch 1 updates subset 1 parameters, batch 2 updates subset 2 parameters, and so on.

[0188] Second, since none of these calculations is time-critical, they could be performed in a low priority operating system (OS) thread on the processor, assuming a real-time OS is used.

[0189] Third, one enhancement that could improve the performance and robustness of the above procedure is to maintain two copies of the ML parameters: param_set_1 which is updated conservatively and param_set_2 which is updated liberally. When the monitor is first installed, param_set_1 and param_set_2 are both initialized to the “one size fits all” parameters that ship with, or are configured in, the monitor. Once installed, only param_set_2 is updated. After several hundred batches, param_set_1 and param_set_2 are both run on some test data, their cross-entropy loss and hard-decision confusion matrices computed, and param_set_2 is copied into param_set_1 only if param_set_2 has a cross-entropy loss and confusion matrix performance that exceeds that of param_set_1 by some appropriate comparison criteria. For example, one could only update param_set_1 if it has a cross-entropy loss improvement of at least 20% and has zero false positive entries or exits and fewer false negative entries and

exits on the test data. The test data could be some combination of labeled data taken from many different rooms and pseudo-labeled data taken from the room in which the monitor is installed. The procedure could continue in this manner, only replacing param_set_1 with param_set_2 when it delivers a significant performance improvement after updating for several hundred batches. If, on the other hand param_set_2 is seen to significantly underperform param_set_1, it could be replaced by param_set_1.

Enhanced Winnowing Procedures

[0190] Winnowing procedure **180**, shown in FIG. **22**, is a technique for identifying a subset of tag devices that are likely to have entered or exited a monitored room, wherein to conserve CPU resources and/or to minimize false positive detections, only receive data obtained from tag devices in the subset of tag devices is processed using the room occupancy detection algorithm **58**. Winnowing procedure **180** estimates the X, Y position of each tag using AoA techniques, and winnows out tag devices that do not appear to cross the X=0 plane (i.e., the plane of the doorway) within a few seconds of when the motion sensor was triggered. Only data received from unwinnowed tag devices, i.e., tags that did appear to cross the X=0 plane, is processed through the ML algorithm.

[0191] Another winnowing technique operates on the received signal strength of the tag transmissions. More specifically, after the motion sensor sees motion under the monitor, only data from tags with a received signal strength exceeding some minimum threshold (e.g., -80 dBm) within a few seconds of the motion sensor activity is processed through the ML algorithm; all other tags are considered “nuisance tags” and their transmissions are winnowed out.

[0192] If the Bluetooth 5.2 or Bluetooth 6 wireless standard is being used, the RSSI could be averaged over multiple receive antennas to reduce the variance of the RSSI estimate, allowing more nuisance tags to be winnowed.

[0193] Another winnowing technique uses distance to winnow out nuisance tags instead of received signal strength, wherein the distance could be obtained by RTT measurements via Bluetooth 6 or IEEE 802.15.4z. In this case, any tag with a distance measurement that exceeds some appropriate threshold while the motion sensor was active could be considered a nuisance tag.

[0194] Yet another winnowing technique operates on IMU readings sent over-the-air by the tag devices, as described earlier herein. The idea is that if a tag device is truly entering or exiting a room, it is moving, and moving in a particular way, while the monitor's motion sensor is active indicating that someone has passed under the monitor.

[0195] Most low-cost tags are equipped with a 3D MEMS accelerometer. Reference is now made to FIG. **28**, which shows a plot of the accelerometer magnitude over time **2800** as a person wearing a tag with an integrated 3D accelerometer is walking into a room. This waveform was high pass filtered to remove the 9.8 m/s.sup.2 acceleration due to earth's gravity before plotting. An effective way to implement winnowing based on the tag's motion is for the processor on the room monitor to compute a motion activity metric from the accelerometer magnitude waveform and to winnow out transmissions from tags in which the motion activity metric does not indicate that the tag is moving while the motion sensor is active. The motion activity metric could be defined to indicate a level of motion that is consistent with human walking behavior. For example, the motion activity metric could be defined as a standard deviation of the accelerometer magnitude taken during a time period in the vicinity of when the monitor's motion sensor was active, or some similar statistic, such as a mean absolute deviation, median square deviation, root mean squared (RMS) deviation, median absolute deviation, or maximum of the magnitude of the tag acceleration over the same time interval.

[0196] Another more sophisticated motion activity metric could be defined to look for periodicity in the accelerometer waveform that is consistent with a human walking pattern. The winnowing procedure could be used to winnow out transmissions from tags that do not exhibit a sufficient level of periodicity in their accelerometer waveforms, or tags that do exhibit periodic accelerometer activity, but not at a period that is consistent with human walking behavior. The level of periodicity could be found by computing the autocorrelation function of the accelerometer magnitude waveform (after high pass filtering), and searching for a peak over some appropriate time range. The time offset of the autocorrelation peak is indicative of the period of the waveform, and the level of the peak relative to the energy in the accelerometer waveform indicates the level of periodicity at that period.

Millimeter-Wave Radar

[0197] The description above mentions that laser, radar or ultrasound sensors could be used to provide additional information to room occupancy detection algorithm **58** to help improve performance. Millimeter-wave (mmwave) (e.g., 60 GHz) radar sensors are an attractive option due to their wide availability, reduced size, low cost and low power consumption. These sensors are available in system-on-a-chip (SoC) configurations with multiple transmit and receive antennas, e.g., 2 Tx and 4 Rx antennas. Since the wavelengths are so small, the antennas can either be embedded on the SoC package itself (antenna-on-package, AoP) or implemented off-chip using PCB traces outside the SoC. The sensors are available as FCC-certified modules, with the sensor, antennas and additional support circuitry soldered onto a small PCB.

[0198] Reference is now made to FIG. **29**, which shows an illustration of a person **2** wearing a wireless tag device **3** approaching a room **4** with a room occupancy monitor **1** installed in the entryway, wherein the monitor **1** contains a radar module **2910** (radar sensor), in accordance with one embodiment. The shaded area **2920** below the monitor shows the FOV of the radar antennas. In this embodiment, the radar module PCB is angled differently than the main PCB **2930** containing the rest of the monitor electronics to give the radar access to a larger portion of the inside of the room **4**.

[0199] The radar sensors typically use FMCW chirp radar that can measure the distance, velocity and angular position of any target within their FOV. The sensors can be configured to output “point clouds” once per radar scan period. A point cloud is a set of data structures indicating where static and moving targets are physically located relative to the radar sensor. Each point in the point cloud typically contains an X, Y, Z spatial position relative to the sensor in cartesian coordinates, expressed in meters or feet; a radial velocity (i.e., the velocity vector projected on to the unit vector pointing from the X, Y, Z point location to the radar sensor) and a receive signal strength, expressed in dBm, in accordance with one embodiment.

[0200] Reference is now made to FIG. **30**, which shows a two-dimensional (2D) illustration of how point-cloud data obtained by a radar sensor/module might appear when a monitor **1** containing a radar module/radar sensor (as shown in FIG. **29**) is deployed

in a hospital room **4**, in accordance with one embodiment. In the figure, the points can be seen to be clustered on the walls **3010**, desk **3020**, bed **3030**, chair **3040**, hand hygiene dispenser sensor **3050**, as well as a person **3060** standing in the room.

[0201] In certain embodiments, either the raw point cloud points, or a two-dimensional (2D) or three-dimensional (3D) spatial histogram derived from those points, or the positions and dimensions of clusters derived from the those points using a clustering algorithm, are presented as input to the algorithm **58** along with the one or more of the RSSI measurements, RTT distance estimates, array response vectors, spatial responses derived from the array response vectors, and/or data from other sensors, as described earlier herein.

[0202] The processor could process the radar point cloud data to determine physical characteristics of the room that could be fed as input parameters into the algorithm. For example, walls **3010** can be identified in the radar scene as a set of points that are static (i.e., zero radial velocity, and persistently present in the same position and dimensions) that appear on the surface of a large plane that is parallel with either the monitor's XZ plane or YZ plane. As another example, the processor could identify the hand hygiene dispenser sensor **3050** as a static rectangular cuboid with some appropriate set of outer dimensions positioned along the wall near the door. Other useful physical characteristics that could be identified using similar techniques include: doorway width, doorway position on the entryway wall (left, right or center), room opening direction after entry (left, right or center), whether door is currently open or closed, bed position and dimensions, monitor height above floor, position and dimensions of chairs or furniture, position of hand hygiene sensor along wall inside or outside of room, and position/location and dimensions of other unidentified or unknown static objects.

[0203] When the radar sensor data is used in combination with the received signal strength measurements, RTT distance estimates and/or array response vectors derived from received transmissions from a tag, the radar sensor data will not only improve the monitor's room occupancy detection performance (i.e., in determining whether a person wearing a tag device is positioned (strictly) inside or outside of a monitored room), it will also give the monitor the ability to accurately locate a person wearing the tag device (i.e., determine its X, Y, Z position) within the room. The combined data will also give the processor the ability to make predictions about activity in the room. For example, a ML implementation of room occupancy detection algorithm **58** could be configured to also output, in addition to entry/exit detections, 2D or 3D position estimates for each tag device. The algorithm could use the combined sensor data to determine whether the person wearing the tag is lying or sitting in (proximate) the bed. If the processor has already identified the 3D position and dimensions of the bed, and if the algorithm **58** running on the combined sensor data sees the tag within the X/Y boundaries of the bed, the processor may conclude that the person wearing the tag is either sitting or lying on the bed. If the processor sees some minimum number of point cloud points at some minimum distance (e.g., 1.5 feet) above the bed, it could conclude the person is sitting on the bed; otherwise, the processor could conclude the person is lying on the bed. The processor could use similar techniques to identify other scenarios or conditions within the monitored room, such as: how much time a doctor or nurse has spent in the room, how much time a doctor or nurse has spent by the patient's bedside, whether the patient is lying on the floor in the room (i.e., if they have fallen) or standing, if a medical staff person is positioned in front of a hand hygiene dispenser when its disinfectant was delivered.

Secondary Room Occupancy Detection Algorithm

[0204] The room occupancy detection algorithm **58** described above herein, begins processing when a person wearing or carrying a tag emitter device walks into the relatively narrow FOV of the monitor's motion sensor, thereby waking up the monitor. The algorithm puts the monitor back to sleep when it has determined that no person has been in the motion sensor FOV for some time-typically 10 seconds. Because the monitor is awake only when a person walks directly under the monitor, it can be viewed as a "chokepoint transition detector" or a "chokepoint detection algorithm" that detects entries and exits by looking explicitly for transitions from outside to inside the room or vice-versa. The alternative to a chokepoint detector is an "occupancy detector" that periodically detects and reports on whether each tag is inside of the room or outside at any given time.

[0205] The key benefit of the chokepoint approach is power savings since the monitor is only awake when a person transitions into or out of the room. One disadvantage of the chokepoint approach is that in the event the algorithm fails to detect a true entry or exit, it could cause the system to mislocate a tag for a potentially long period of time. For example, if a monitor fails to detect when a tag enters the room and remains inside for 2 hours with no other room entries or exits thereafter, the system will incorrectly locate the tag as being outside the room for the entire 2 hours.

[0206] One way to reduce the probability of these missed detections is for the processor to periodically wake up the monitor from sleep to take brief snapshots of its sensor data to determine whether a tag has changed position. As long as the quantity and duty cycle of the snapshots can be kept sufficiently small, this approach can be used to reduce the number of missed detections without a significant reduction in the monitor's battery life.

[0207] The snapshots defined above are referred to herein as "secondary dwells." The algorithm used to process the secondary dwell data is referred to as a "secondary algorithm" or a "second chance algorithm", since it gives the monitor a second chance at correcting missed detections. The original chokepoint detection algorithm **58** described above herein that detects tag entries and exits is referred to as the "primary algorithm", and the time periods containing the data that it processes, starting with motion detection/monitor wakeup and ending 10 seconds after motion was last detected, are referred to as "primary dwells".

[0208] To minimize battery energy, and because the tags are often not moving once inside or outside of the room, it is not generally necessary or desirable to use a secondary dwell duration any longer than a second or two. Even a fraction of a second could suffice.

[0209] The spacing between consecutive secondary dwells may be at least 1-2 seconds. This increases the likelihood that the wireless channel impulse response between monitor and tag will have changed since the previous secondary dwell due to motion in the room, which will help to decorrelate the measurements.

[0210] Reference is now made to FIG. **31**, which shows a functional block diagram **3100** of how the primary and secondary algorithms can be used together to form a software system running on the monitor's processor **15**, in accordance with an example embodiment.

[0211] The raw sensor inputs to both algorithms include array response vectors **51** received from one or more tag devices, RSSI

measurements (ideally average RSSI across a plurality of antennas), **52** for each received array response vector, digitized PIR waveform samples (obtained from an analog PIR **102**), and IMU samples (from IMU **108**) received from the tag devices.

[0212] The array response vectors and RSSI measurements are processed through a Spatial Signal Processing block **3105** to produce a “heat map” indicating the likelihood of the tag being located at each location in a grid of candidate positions. The likelihood function could be computed using any one of the following well-known array signal processing algorithms: beamforming, MUSIC, ESPRIT, MVDR, etc.

[0213] The Spatial Signal Processing block **3105** may also be used to estimate and output additional parameters, including the X/Y position of the tag device (maximum likelihood X/Y grid position in the heat map), a near-far-ratio (ratio of power received under monitor to power received from other locations), or an inside/outside ratio (ratio of power received from grid positions likely to be inside the room to power received from other locations).

[0214] The PIR samples may be digitally lowpass filtered to remove noise, and high pass filtered to track any changes in DC offset from the PIR sensor output.

[0215] The Tag IMU samples may also be digitally lowpass filtered to remove noise or high pass filtered to remove gravitational effects.

[0216] The filtered PIR, Tag IMU and Spatial Signal Processing block outputs are fed through a resampler **3107**, which resamples the data using a common time base and fills in any missing samples.

[0217] The resampler output is fed into the primary algorithm **58** and secondary algorithm **3110**.

[0218] The filtered PIR signal is fed into a motion detector **3115** and then a Wakeup, Sleep and Scheduling block **3120**, the latter being used to wake up the processor from sleep when motion is detected, to put the processor to sleep upon lack-of-motion detection, and to schedule the execution of the primary and secondary algorithms.

[0219] It should be noted that each tag device within “earshot” of the monitor should be assigned a separate standalone instance of the primary and secondary algorithms. Although the PIR samples are common across all instances, the array response vectors, RSSI measurements and Tag IMU data are different for each tag and is managed independently. This can be done on the monitor's processor using object-oriented programming, i.e., common code but separate data structures for each tag instance.

[0220] Reference is now made to FIG. **32**, which shows a procedure **3200** for running an enhanced room occupancy detection algorithm using both the primary and secondary algorithms. The procedure **3200** starts at step **3205**.

[0221] In step **3210**, the monitor is in a low-power sleep state, awaiting a wakeup trigger from the motion sensor **6**.

[0222] In step **3215**, once a wakeup event has occurred, the processor begins receiving and storing data from the monitor's sensors, including one or more of array response vectors, received signal strength measurements, RTT distance measurements, digitized motion sensor samples, radar point-cloud data, and over-the-air IMU samples; processing the data in real-time using the primary algorithm **58**, and notifying the outside world of any tag entry or exit detections. The processor continues in this fashion until no motion sensor activity has been seen for an inactivity period such as 10 seconds.

[0223] In step **3220**, the processor puts the monitor in a semi-sleep state, wherein most of the time is spent sleeping, but the processor periodically wakes up to perform a block of N secondary dwells and runs the secondary algorithm on the sensor data obtained during each dwell. After all N secondary dwells have been processed, the secondary algorithm issues a decision—one of INSIDE if the tag is determined to be inside the room, OUTSIDE if the tag is determined to be outside, or NEITHER if it is unable to reliably determine whether the tag is inside or outside. If a motion sensor wakeup event occurs at any time during the N secondary dwells, those dwells are immediately terminated, and step **3215** is reinvoked.

[0224] In step **3225**, the monitor sleeps and waits for wakeup event from the motion sensor. If the wakeup event occurs, the procedure returns to step **3215**. If no wakeup event has been received after an appropriate timeout interval (e.g., 15 minutes), step **3220** of this procedure is re-invoked, yielding another block of N secondary dwell measurements.

[0225] The secondary algorithm could be a ML or non-ML algorithm. For example, a non-ML algorithm could operate on the spatial heat map generated via a 2D beamforming or MUSIC map, indicating the probability of tag's position in the vicinity of the monitor based on array response data captured during the secondary dwells to make occupancy detection decisions. and declare a tag as being INSIDE if the most likely position of the tag on the spatial heat map appears inside a “safe zone” within the room, where a “safe zone” is a region within the room where there is a high likelihood that the tag is truly inside the room. An example indoor “safe zone” is shown as the shaded “Zone A” in FIG. **33**. Note that the safe zone in this example does not include points that are close to the wall, as those points have a higher likelihood of being confused for outside points. A similar approach can be used for the OUTSIDE case. A NEITHER decision could be used if the tag position appears outside of a safe zone.

[0226] If an ML algorithm is desired, and the primary ML algorithm is a stateful neural network such as an RNN, LSTM or GRU, a good choice would be to use that same algorithm for the secondary ML algorithm. The same neural network could be used to process the data from primary dwells and secondary dwells. Alternatively, a multi-layer perceptron (MLP) or some other well-known stateless ML algorithm could be used to process each secondary dwell, with some logic to combine the results from multiple secondary dwells and make a final INSIDE, OUTSIDE or NEITHER decision. The advantage of using a stateful algorithm over stateless is that the former can use both the primary and secondary dwell data to make its room occupancy decisions, i.e., it has access to more information than the stateless algorithm has, since the latter operates only on the secondary dwell data.

[0227] The data used to train the secondary algorithm if an ML algorithm is used may consist of sensor data (e.g., array response vectors, RSSI measurements, PIR samples, RTT distance measurements, radar data, and any other sensor data mentioned above herein) for many secondary dwells obtained when a tag was truly inside or outside of a room. The secondary dwells should be labeled with either INSIDE, OUTSIDE or NEITHER to train the ML algorithm.

[0228] Referring back to FIG. **27**, it should be noted that the self-supervised learning procedure **2700** described earlier herein for tuning a monitor's primary ML algorithm for a specific room once it has been installed can also be used to tune the parameters of the secondary algorithm. The input data for the secondary algorithm is data obtained during secondary dwells while the room

monitor is normally used. The pseudo-labels used for each secondary dwell are set to INSIDE if the primary algorithm detected an entry prior to the secondary dwell, or OUTSIDE if the primary algorithm detected an exit prior to the secondary dwell. A NEITHER pseudo-label could be used when the primary algorithm detects the tag's X/Y position within the room to be within a foot or two feet of the door or walls within the room.

[0229] Referring back to FIG. 24, FIG. 24 illustrates a hardware block diagram of a computing device 200 that may perform functions associated with operations discussed herein in connection with the techniques depicted in FIGS. 1-23 and 25-33. In various embodiments, a computing device or apparatus, such as computing device 200 or any combination of computing devices 200, may be configured as any entity/entities (e.g., servers 43 and/or 45 described above) as discussed for the techniques depicted in connection with FIGS. 1-23 and 25-33 in order to perform operations of the various techniques discussed herein.

[0230] In at least one embodiment, the computing device 200 may be any apparatus that may include one or more processor(s) 202, one or more memory element(s) 204, storage 206, a bus 208, one or more network processor unit(s) 210 interconnected with one or more network input/output (I/O) interface(s) 212, one or more I/O interface(s) 214, and control logic 220. In various embodiments, instructions associated with logic for computing device 200 can overlap in any manner and are not limited to the specific allocation of instructions and/or operations described herein.

[0231] In at least one embodiment, processor(s) 202 is/are at least one hardware processor configured to execute various tasks, operations and/or functions for computing device 200 as described herein according to software and/or instructions configured for computing device 200. Processor(s) 202 (e.g., a hardware processor) can execute any type of instructions associated with data to achieve the operations detailed herein. In one example, processor(s) 202 can transform an element or an article (e.g., data, information) from one state or thing to another state or thing. Any of potential processing elements, microprocessors, digital signal processor, baseband signal processor, modem, PHY, controllers, systems, managers, logic, and/or machines described herein can be construed as being encompassed within the broad term 'processor'.

[0232] In at least one embodiment, memory element(s) 204 and/or storage 206 is/are configured to store data, information, software, and/or instructions associated with computing device 200, and/or logic configured for memory element(s) 204 and/or storage 206. For example, any logic described herein (e.g., control logic 220) can, in various embodiments, be stored for computing device 200 using any combination of memory element(s) 204 and/or storage 206. Note that in some embodiments, storage 206 can be consolidated with memory element(s) 204 (or vice versa), or can overlap/exist in any other suitable manner.

[0233] In at least one embodiment, bus 208 can be configured as an interface that enables one or more elements of computing device 200 to communicate in order to exchange information and/or data. Bus 208 can be implemented with any architecture designed for passing control, data and/or information between processors, memory elements/storage, peripheral devices, and/or any other hardware and/or software components that may be configured for computing device 200. In at least one embodiment, bus 208 may be implemented as a fast kernel-hosted interconnect, potentially using shared memory between processes (e.g., logic), which can enable efficient communication paths between the processes.

[0234] In various embodiments, network processor unit(s) 210 may enable communication between computing device 200 and other systems, entities, etc., via network I/O interface(s) 212 (wired and/or wireless) to facilitate operations discussed for various embodiments described herein. In various embodiments, network processor unit(s) 210 can be configured as a combination of hardware and/or software, such as one or more Ethernet driver(s) and/or controller(s) or interface cards, Fibre Channel (e.g., optical) driver(s) and/or controller(s), wireless receivers/transmitters/transceivers, baseband processor(s)/modem(s), and/or other similar network interface driver(s) and/or controller(s) now known or hereafter developed to enable communications between computing device 200 and other systems, entities, etc. to facilitate operations for various embodiments described herein. In various embodiments, network I/O interface(s) 212 can be configured as one or more Ethernet port(s), Fibre Channel ports, any other I/O port(s), and/or antenna(s)/antenna array(s) now known or hereafter developed. Thus, the network processor unit(s) 210 and/or network I/O interface(s) 212 may include suitable interfaces for receiving, transmitting, and/or otherwise communicating data and/or information in a network environment.

[0235] I/O interface(s) 214 allow for input and output of data and/or information with other entities that may be connected to computing device 200. For example, I/O interface(s) 214 may provide a connection to external devices such as a keyboard, keypad, a touch screen, and/or any other suitable input and/or output device now known or hereafter developed. In some instances, external devices can also include portable computer readable (non-transitory) storage media such as database systems, thumb drives, portable optical or magnetic disks, and memory cards. In still some instances, external devices can be a mechanism to display data to a user, such as, for example, a computer monitor, a display screen, or the like.

[0236] In various embodiments, control logic 220 can include instructions that, when executed, cause processor(s) 202 to perform operations, which can include, but not be limited to, providing overall control operations of computing device; interacting with other entities, systems, etc. described herein; maintaining and/or interacting with stored data, information, parameters, etc. (e.g., memory element(s), storage, data structures, databases, tables, etc.); combinations thereof; and/or the like to facilitate various operations for embodiments described herein.

[0237] The programs described herein (e.g., control logic 220) may be identified based upon application(s) for which they are implemented in a specific embodiment. However, it should be appreciated that any particular program nomenclature herein is used merely for convenience; thus, embodiments herein should not be limited to use(s) solely described in any specific application(s) identified and/or implied by such nomenclature.

[0238] In various embodiments, any entity or apparatus as described herein may store data/information in any suitable volatile and/or non-volatile memory item (e.g., magnetic hard disk drive, solid state hard drive, semiconductor storage device, random access memory (RAM), read only memory (ROM), erasable programmable read only memory (EPROM), application specific integrated circuit (ASIC), etc.), software, logic (fixed logic, hardware logic, programmable logic, analog logic, digital logic), hardware, and/or in any other suitable component, device, element, and/or object as may be appropriate. Any of the memory items discussed herein should be construed as being encompassed within the broad term 'memory element'. Data/information being tracked and/or sent to one or more entities as discussed herein could be provided in any database, table, register, list, cache,

storage, and/or storage structure: all of which can be referenced at any suitable timeframe. Any such storage options may also be included within the broad term ‘memory element’ as used herein.

[0239] Note that in certain example implementations, operations as set forth herein may be implemented by logic encoded in one or more tangible media that is capable of storing instructions and/or digital information and may be inclusive of non-transitory tangible media and/or non-transitory computer readable storage media (e.g., embedded logic provided in: an ASIC, digital signal processing (DSP) instructions, software [potentially inclusive of object code and source code], etc.) for execution by one or more processor(s), and/or other similar machine, etc. Generally, memory element(s) **204** and/or storage **206** can store data, software, code, instructions (e.g., processor instructions), logic, parameters, combinations thereof, and/or the like used for operations described herein. This includes memory element(s) **204** and/or storage **206** being able to store data, software, code, instructions (e.g., processor instructions), logic, parameters, combinations thereof, or the like that are executed to carry out operations in accordance with teachings of the present disclosure.

[0240] In some instances, software of the present embodiments may be available via a non-transitory computer useable medium (e.g., magnetic or optical mediums, magneto-optic mediums, CD-ROM, DVD, memory devices, etc.) of a stationary or portable program product apparatus, downloadable file(s), file wrapper(s), object(s), package(s), container(s), and/or the like. In some instances, non-transitory computer readable storage media may also be removable. For example, a removable hard drive may be used for memory/storage in some implementations. Other examples may include optical and magnetic disks, thumb drives, and smart cards that can be inserted and/or otherwise connected to a computing device for transfer onto another computer readable storage medium.

[0241] In some aspects, the techniques described herein relate to a wireless room occupancy monitor, including: an antenna array configured to detect wireless transmissions from a tag device; a wireless transceiver configured to receive the wireless transmissions detected by the antenna array and produce receive signals; a processor configured to process the receive signals; and a motion sensor coupled to the processor and configured to wake up the processor in response to detecting when the tag device enters or exits a room; wherein the antenna array and motion sensor are configured to be mounted proximate an entryway to the room; wherein after the motion sensor wakes up the processor, the processor is configured to: power on the wireless transceiver and run an algorithm on a sequence of estimates of a parameter, derived from the receive signals from the tag device, that is indicative of a distance between the tag device and the wireless room occupancy monitor, and array response vectors derived from the receive signals; and determine when the tag device has entered or exited the room via the entryway.

[0242] In some aspects, the techniques described herein relate to a wireless room occupancy monitor, wherein the parameter is a received signal strength of the receive signals, a time-of-flight of the receive signals, or a round-trip time between one or more signals transmitted from the monitor to the tag device and the receive signals.

[0243] In some aspects, the techniques described herein relate to a wireless room occupancy monitor, wherein the round-trip time is computed using either a round-trip time measurement procedure specified in the Ultra-Wideband (UWB) IEEE 802.15.4z wireless standard or a phase-based-ranging channel sounding procedure specified in the Bluetooth 6 wireless standard.

[0244] In some aspects, the techniques described herein relate to a wireless room occupancy monitor, wherein before the processor runs the algorithm, the processor executes a winnowing procedure to identify a subset of tag devices that are likely to have entered or exited the room, wherein the processor runs the algorithm using only receive signals obtained from tag devices in the subset of tag devices, and wherein the winnowing procedure uses one or more of the received signal strength, the time-of-flight, or the round-trip time to identify the subset of tag devices.

[0245] In some aspects, the techniques described herein relate to a wireless room occupancy monitor, wherein the algorithm is a machine learning algorithm and wherein after the wireless room occupancy monitor has been installed in a room, the processor uses self-supervised learning to incrementally train the machine learning algorithm by using outputs from the machine learning algorithm to generate pseudo-labels, computing an algorithmic loss between the outputs and the pseudo-labels, and updating weights of the machine learning algorithm using gradient descent to minimize the algorithmic loss over time.

[0246] In some aspects, the techniques described herein relate to a wireless room occupancy monitor, further including a radar sensor, wherein outputs from the radar sensor are provided as inputs to the algorithm along with the sequence of estimates of the parameter.

[0247] In some aspects, the techniques described herein relate to a wireless room occupancy monitor, wherein the processor uses outputs from the radar sensor to determine one or more of the following physical characteristics of the room and present them as input parameters to the algorithm: wall dimensions, ceiling height, doorway position within the room, doorway opening width, whether door is currently open or closed, room opening direction, bed size and position, desk and chair locations, medical equipment locations, hand hygiene sensor location, or unknown static object location.

[0248] In some aspects, the techniques described herein relate to a wireless room occupancy monitor, wherein the algorithm is further configured to estimate a two-dimensional (2D) or a three-dimensional (3D) position of the tag device relative to the wireless room occupancy monitor.

[0249] In some aspects, the techniques described herein relate to a room occupancy monitor, wherein the processor is further configured to use one or more of the outputs from the radar sensor and the sequence of estimates of the parameter to identify one or more of the following scenarios occurring within the room: whether a person wearing the tag device is proximate to a bed within the room; whether a person wearing the tag device is lying in a bed within the room; whether a person wearing the tag device is lying on a floor within the room or standing, or; whether a person wearing the tag device is standing in front of a hand hygiene dispenser inside or outside of the room.

[0250] In some aspects, the techniques described herein relate to a wireless room occupancy monitor, wherein the algorithm includes a primary algorithm and a secondary algorithm, wherein the primary algorithm runs after a motion sensor wakeup and detects when the tag device has transitioned from outside to inside or inside to outside of the room, and the secondary algorithm runs periodically when the motion sensor has been idle for a time period and detects whether the tag device is strictly inside or outside of the room.

[0251] In some aspects, the techniques described herein relate to a room occupancy detection system, including: one or more room occupancy monitors configured to detect entries into a room and exits from the room of one or more tag devices, and to produce room occupancy detection events; and a server configured to receive the room occupancy detection events from the one or more room occupancy monitors; wherein each of the one or more room occupancy monitors includes: an antenna array configured to detect wireless transmissions from the one or more tag devices; a wireless transceiver configured to receive the wireless transmissions detected by the antenna array and produce receive signals; a processor configured to process the receive signals; and a motion sensor coupled to the processor and configured to wake up the processor in response to detecting when one of the one or more tag devices enters or exits the room; wherein each of the one or more room occupancy monitors is configured to be mounted proximate an entryway to the room; and wherein after the motion sensor wakes up the processor on any one of the one or more room occupancy monitors, the processor is configured to: power on the wireless transceiver and to run an algorithm on a sequence of estimates of a parameter, derived from the receive signals from the tag device, that is indicative of a distance between the tag device and monitor, and array response vectors derived from the receive signals; and determine when the tag device has entered or exited the room via the entryway.

[0252] In some aspects, the techniques described herein relate to a room occupancy detection system, wherein the processor of one or more of the room occupancy monitors is configured to decode inertial motion sensor data contained in the wireless transmissions received from one or more tag devices to produce decoded inertial sensor data containing a tag acceleration component; wherein the processor of the one or more room occupancy monitors is configured to, before running the algorithm, execute a winnowing procedure to identify a subset of the tag devices that are likely to have entered or exited the room based on the decoded inertial sensor data, and wherein only receive signals obtained from tag devices in the subset of the tag devices is processed using the algorithm.

[0253] In some aspects, the techniques described herein relate to a room occupancy detection system, wherein the winnowing procedure includes calculating a magnitude of the tag acceleration component, calculating a motion activity metric from the magnitude of the tag acceleration component, and using the motion activity metric to identify tags in the subset.

[0254] In some aspects, the techniques described herein relate to a room occupancy detection system, wherein the motion activity metric includes a standard deviation, mean absolute deviation, median square deviation, root mean squared (RMS) deviation, median absolute deviation, or maximum of the magnitude of the tag acceleration component in a vicinity of a time while the motion sensor indicates a person is under the room occupancy monitor.

[0255] In some aspects, the techniques described herein relate to a room occupancy detection system, wherein the motion activity metric includes a level of periodicity calculation or an estimate of a period of the magnitude of the tag acceleration component.

[0256] In some aspects, the techniques described herein relate to a method performed by a wireless room occupancy monitor, including: detecting, at an antenna array, wireless transmissions from a tag device; producing, with a wireless transceiver, receive signals from wireless transmissions detected at the antenna array; detecting, with a motion sensor of the wireless room occupancy monitor, when the tag device enters or exits a room; waking up a processor of the wireless room occupancy monitor upon the motion sensor detecting when the tag devices enters or exists the room; upon waking up the processor: powering on the wireless transceiver and running an algorithm on a sequence of estimates of a parameter, derived from the receive signals from the tag device, that is indicative of a distance between the tag device and monitor, and array response vectors derived from the receive signals; and determining when the tag device has entered or exited the room via an entryway.

[0257] In some aspects, the techniques described herein relate to a method, wherein the parameter is a received signal strength of the receive signals, a time-of-flight of the receive signals, or a round-trip time between one or more signals transmitted from the monitor to the tag device and the receive signals.

[0258] In some aspects, the techniques described herein relate to a method, wherein the round-trip time is computed using either a round-trip time measurement procedure specified in the Ultra-Wideband (UWB) IEEE 802.15.4z wireless standard or a phase-based-ranging channel sounding procedure specified in the Bluetooth 6 wireless standard.

[0259] In some aspects, the techniques described herein relate to a method, wherein the powering on the wireless transceiver and running further includes executing a winnowing procedure to identify a subset of tag devices that are likely to have entered or exited the room, wherein the running is only performed on receive signals obtained from tag devices in the subset of tag devices, and wherein the winnowing procedure uses one or more of the received signal strength, the time-of-flight, or the round-trip time to identify the subset of tag devices.

[0260] In some aspects, the techniques described herein relate to a method, wherein the algorithm is a machine learning algorithm, and the method further includes performing self-supervised learning to incrementally train the machine learning algorithm by using outputs from the machine learning algorithm to generate pseudo-labels, computing an algorithmic loss between the outputs and the pseudo-labels, and updating weights of the machine learning algorithm using gradient descent to minimize the algorithmic loss over time.

Variations and Implementations

[0261] Embodiments described herein may include one or more networks, which can represent a series of points and/or network elements of interconnected communication paths for receiving and/or transmitting messages (e.g., packets of information) that propagate through the one or more networks. These network elements offer communicative interfaces that facilitate communications between the network elements. A network can include any number of hardware and/or software elements coupled to (and in communication with) each other through a communication medium. Such networks can include, but are not limited to, any local area network (LAN), virtual LAN (VLAN), wide area network (WAN) (e.g., the Internet), software defined WAN (SD-WAN), wireless local area (WLA) access network, wireless wide area (WWA) access network, metropolitan area network (MAN), Intranet, Extranet, virtual private network (VPN), Low Power Network (LPN), Low Power Wide Area Network (LPWAN), Machine to Machine (M2M) network, Internet of Things (IoT) network, Ethernet network/switching system, any other appropriate architecture and/or system that facilitates communications in a network environment, and/or any suitable

combination thereof.

[0262] Networks through which communications propagate can use any suitable technologies for communications including wireless communications (e.g., 4G/5G/nG, IEEE 802.11 (e.g., Wi-Fi®/Wi-Fi6®), IEEE 802.16 (e.g., Worldwide Interoperability for Microwave Access (WiMAX)), Radio-Frequency Identification (RFID), Near Field Communication (NFC), Bluetooth™, mm.wave, Ultra-Wideband (UWB), etc.), and/or wired communications (e.g., T1 lines, T3 lines, digital subscriber lines (DSL), Ethernet, Fibre Channel, etc.). Generally, any suitable means of communications may be used such as electric, sound, light, infrared, and/or radio to facilitate communications through one or more networks in accordance with embodiments herein. Communications, interactions, operations, etc. as discussed for various embodiments described herein may be performed among entities that may directly or indirectly connected utilizing any algorithms, communication protocols, interfaces, etc. (proprietary and/or non-proprietary) that allow for the exchange of data and/or information.

[0263] Communications in a network environment can be referred to herein as ‘messages’, ‘messaging’, ‘signaling’, ‘data’, ‘content’, ‘objects’, ‘requests’, ‘queries’, ‘responses’, ‘replies’, etc. which may be inclusive of packets. As referred to herein and in the claims, the term ‘packet’ may be used in a generic sense to include packets, frames, segments, datagrams, and/or any other generic units that may be used to transmit communications in a network environment. Generally, a packet is a formatted unit of data that can contain control or routing information (e.g., source and destination address, source and destination port, etc.) and data, which is also sometimes referred to as a ‘payload’, ‘data payload’, and variations thereof.

[0264] To the extent that embodiments presented herein relate to the storage of data, the embodiments may employ any number of any conventional or other databases, data stores or storage structures (e.g., files, databases, data structures, data or other repositories, etc.) to store information.

[0265] Note that in this Specification, references to various features (e.g., elements, structures, nodes, modules, components, engines, logic, steps, operations, functions, characteristics, etc.) included in ‘one embodiment’, ‘example embodiment’, ‘an embodiment’, ‘another embodiment’, ‘certain embodiments’, ‘some embodiments’, ‘various embodiments’, ‘other embodiments’, ‘alternative embodiment’, and the like are intended to mean that any such features are included in one or more embodiments of the present disclosure, but may or may not necessarily be combined in the same embodiments. Note also that a module, engine, client, controller, function, logic or the like as used herein in this Specification, can be inclusive of an executable file comprising instructions that can be understood and processed on a server, computer, processor, machine, compute node, combinations thereof, or the like and may further include library modules loaded during execution, object files, system files, hardware logic, software logic, or any other executable modules.

[0266] It is also noted that the operations and steps described with reference to the preceding figures illustrate only some of the possible scenarios that may be executed by one or more entities discussed herein. Some of these operations may be deleted or removed where appropriate, or these steps may be modified or changed considerably without departing from the scope of the presented concepts. In addition, the timing and sequence of these operations may be altered considerably and still achieve the results taught in this disclosure. The preceding operational flows have been offered for purposes of example and discussion. Substantial flexibility is provided by the embodiments in that any suitable arrangements, chronologies, configurations, and timing mechanisms may be provided without departing from the teachings of the discussed concepts.

[0267] Each example embodiment disclosed herein has been included to present one or more different features. However, all disclosed example embodiments are designed to work together as part of a single larger system or method. This disclosure explicitly envisions compound embodiments that combine multiple previously-discussed features in different example embodiments into a single system or method.

[0268] Additionally, unless expressly stated to the contrary, the terms ‘first’, ‘second’, ‘third’, etc., are intended to distinguish the particular nouns they modify (e.g., element, condition, node, module, activity, operation, etc.). Unless expressly stated to the contrary, the use of these terms is not intended to indicate any type of order, rank, importance, temporal sequence, or hierarchy of the modified noun.

[0269] One or more advantages described herein are not meant to suggest that any one of the embodiments described herein necessarily provides all of the described advantages or that all the embodiments of the present disclosure necessarily provide any one of the described advantages. Numerous other changes, substitutions, variations, alterations, and/or modifications may be ascertained to one skilled in the art and it is intended that the present disclosure encompass all such changes, substitutions, variations, alterations, and/or modifications as falling within the scope of the appended claims.

Claims

1. A wireless room occupancy monitor, comprising: an antenna array configured to detect wireless transmissions from a tag device; a wireless transceiver configured to receive the wireless transmissions detected by the antenna array and produce receive signals; a processor configured to process the receive signals; and a motion sensor coupled to the processor and configured to wake up the processor in response to detecting when the tag device enters or exits a room; wherein the antenna array and motion sensor are configured to be mounted proximate an entryway to the room; wherein after the motion sensor wakes up the processor, the processor is configured to: power on the wireless transceiver and run an algorithm on a sequence of estimates of a parameter, derived from the receive signals from the tag device, that is indicative of a distance between the tag device and the wireless room occupancy monitor, and array response vectors derived from the receive signals; and determine when the tag device has entered or exited the room via the entryway.
2. The wireless room occupancy monitor of claim 1, wherein the parameter is a received signal strength of the receive signals, a time-of-flight of the receive signals, or a round-trip time between one or more signals transmitted from the monitor to the tag device and the receive signals.
3. The wireless room occupancy monitor of claim 2, wherein the round-trip time is computed using either a round-trip time

measurement procedure specified in the Ultra-Wideband (UWB) IEEE 802.15.4z wireless standard or a phase-based-ranging channel sounding procedure specified in the Bluetooth 6 wireless standard.

4. The wireless room occupancy monitor of claim 2, wherein before the processor runs the algorithm, the processor executes a winnowing procedure to identify a subset of tag devices that are likely to have entered or exited the room, wherein the processor runs the algorithm using only receive signals obtained from tag devices in the subset of tag devices, and wherein the winnowing procedure uses one or more of the received signal strength, the time-of-flight, or the round-trip time to identify the subset of tag devices.

5. The wireless room occupancy monitor of claim 1, wherein the algorithm is a machine learning algorithm and wherein after the wireless room occupancy monitor has been installed in a room, the processor uses self-supervised learning to incrementally train the machine learning algorithm by using outputs from the machine learning algorithm to generate pseudo-labels, computing an algorithmic loss between the outputs and the pseudo-labels, and updating weights of the machine learning algorithm using gradient descent to minimize the algorithmic loss over time.

6. The wireless room occupancy monitor of claim 1, further comprising a radar sensor, wherein outputs from the radar sensor are provided as inputs to the algorithm along with the sequence of estimates of the parameter.

7. The wireless room occupancy monitor of claim 6, wherein the processor uses outputs from the radar sensor to determine one or more of the following physical characteristics of the room and present them as input parameters to the algorithm: wall dimensions, ceiling height, doorway position within the room, doorway opening width, whether door is currently open or closed, room opening direction, bed size and position, desk and chair locations, medical equipment locations, hand hygiene sensor location, or unknown static object location.

8. The wireless room occupancy monitor of claim 6, wherein the algorithm is further configured to estimate a two-dimensional (2D) or a three-dimensional (3D) position of the tag device relative to the wireless room occupancy monitor.

9. The room occupancy monitor of claim 8, wherein the processor is further configured to use one or more of the outputs from the radar sensor and the sequence of estimates of the parameter to identify one or more of the following scenarios occurring within the room: whether a person wearing the tag device is proximate to a bed within the room; whether a person wearing the tag device is lying in a bed within the room; whether a person wearing the tag device is lying on a floor within the room or standing, or; whether a person wearing the tag device is standing in front of a hand hygiene dispenser inside or outside of the room.

10. The wireless room occupancy monitor of claim 1, wherein the algorithm comprises a primary algorithm and a secondary algorithm, wherein the primary algorithm runs after a motion sensor wakeup and detects when the tag device has transitioned from outside to inside or inside to outside of the room, and the secondary algorithm runs periodically when the motion sensor has been idle for a time period and detects whether the tag device is strictly inside or outside of the room.

11. A room occupancy detection system, comprising: one or more room occupancy monitors configured to detect entries into a room and exits from the room of one or more tag devices, and to produce room occupancy detection events; and a server configured to receive the room occupancy detection events from the one or more room occupancy monitors; wherein each of the one or more room occupancy monitors comprises: an antenna array configured to detect wireless transmissions from the one or more tag devices; a wireless transceiver configured to receive the wireless transmissions detected by the antenna array and produce receive signals; a processor configured to process the receive signals; and a motion sensor coupled to the processor and configured to wake up the processor in response to detecting when one of the one or more tag devices enters or exits the room; wherein each of the one or more room occupancy monitors is configured to be mounted proximate an entryway to the room; and wherein after the motion sensor wakes up the processor on any one of the one or more room occupancy monitors, the processor is configured to: power on the wireless transceiver and to run an algorithm on a sequence of estimates of a parameter, derived from the receive signals from the tag device, that is indicative of a distance between the tag device and monitor, and array response vectors derived from the receive signals; and determine when the tag device has entered or exited the room via the entryway.

12. The room occupancy detection system of claim 11, wherein the processor of one or more of the room occupancy monitors is configured to decode inertial motion sensor data contained in the wireless transmissions received from one or more tag devices to produce decoded inertial sensor data containing a tag acceleration component; wherein the processor of the one or more room occupancy monitors is configured to, before running the algorithm, execute a winnowing procedure to identify a subset of the tag devices that are likely to have entered or exited the room based on the decoded inertial sensor data, and wherein only receive signals obtained from tag devices in the subset of the tag devices is processed using the algorithm.

13. The room occupancy detection system of claim 12, wherein the winnowing procedure includes calculating a magnitude of the tag acceleration component, calculating a motion activity metric from the magnitude of the tag acceleration component, and using the motion activity metric to identify tags in the subset.

14. The room occupancy detection system of claim 13, wherein the motion activity metric includes a standard deviation, mean absolute deviation, median square deviation, root mean squared (RMS) deviation, median absolute deviation, or maximum of the magnitude of the tag acceleration component in a vicinity of a time while the motion sensor indicates a person is under the room occupancy monitor.

15. The room occupancy detection system of claim 13, wherein the motion activity metric includes a level of periodicity calculation or an estimate of a period of the magnitude of the tag acceleration component.

16. A method performed by a wireless room occupancy monitor, comprising: detecting, at an antenna array, wireless transmissions from a tag device; producing, with a wireless transceiver, receive signals from wireless transmissions detected at the antenna array; detecting, with a motion sensor of the wireless room occupancy monitor, when the tag device enters or exits a room; waking up a processor of the wireless room occupancy monitor upon the motion sensor detecting when the tag devices enters or exists the room; upon waking up the processor: powering on the wireless transceiver and running an algorithm on a sequence of estimates of a parameter, derived from the receive signals from the tag device, that is indicative of a distance

between the tag device and monitor, and array response vectors derived from the receive signals; and determining when the tag device has entered or exited the room via an entryway.

17. The method of claim 16, wherein the parameter is a received signal strength of the receive signals, a time-of-flight of the receive signals, or a round-trip time between one or more signals transmitted from the monitor to the tag device and the receive signals.

18. The method of claim 17, wherein the round-trip time is computed using either a round-trip time measurement procedure specified in the Ultra-Wideband (UWB) IEEE 802.15.4z wireless standard or a phase-based-ranging channel sounding procedure specified in the Bluetooth 6 wireless standard.

19. The method of claim 18, wherein the powering on the wireless transceiver and running further comprises executing a winnowing procedure to identify a subset of tag devices that are likely to have entered or exited the room, wherein the running is only performed on receive signals obtained from tag devices in the subset of tag devices, and wherein the winnowing procedure uses one or more of the received signal strength, the time-of-flight, or the round-trip time to identify the subset of tag devices.

20. The method of claim 16, wherein the algorithm is a machine learning algorithm, and the method further comprises performing self-supervised learning to incrementally train the machine learning algorithm by using outputs from the machine learning algorithm to generate pseudo-labels, computing an algorithmic loss between the outputs and the pseudo-labels, and updating weights of the machine learning algorithm using gradient descent to minimize the algorithmic loss over time.
